

RESEARCH HANDOFF · SOURCE ARCHIVE · OVERNIGHT RESEARCH

黎曼零点与随机矩阵

第 6–14 轮

证明与研究补编

真实算术、CUE 热流、反例与失败路线
GPT–6 Astra · 完整证明与独立审查

Bill (Qingyun) Sun · 2026 年 9 月 4–5 日
资料检查点 2a9ec81 · 完整后续报告 · 可复算证据与未解决义务
已证明、计算、猜想与失败分别标明 · Public research archive

内容索引

Rounds 6–14: complete takeover supplement / 第 6–14 轮完整接棒补编

1. What the programme is trying to prove
2. The strongest arithmetic results at this checkpoint
3. The finite CUE theorem that is now on firm ground
4. Failed approaches that materially changed the programme
5. The R6 prime-gap work that had not reached the earlier PDFs
6. Fable corrections must travel with their source snapshots
7. Prioritized next work, with concrete success criteria
8. 中文接棒判断与阅读顺序
9. Evidence, preservation and reproducibility contract

Complete source texts / 完整原文目录

Current report 01: Round 6: a full signed-operator direction beyond the 77-dimensional sieve trial

Current report 02: The full signed cap operator and a mass-orthogonal direction beyond 77 dimensions

Current report 03: Residual directions outside the 77-dimensional sieve trial space

Current report 04: Independent finite marked-space regression for the signed sieve operator

Current report 05: Round 6: one full cap-operator residual beyond the 77-dimensional family

Current report 06: Provenance and compact witness export

Current report 07: Independent review of the frozen radial profile outside the old span

Current report 08: Round 7: two explicit actual-zeta targets for Dyson–Montgomery

Current report 09: Two Poisson scales remove the AH near-diagonal parameter

Current report 10: Independent review of the two-scale actual-zeta reduction

Current report 11: Independent review of the actual-zeta Poisson transfer

Current report 12: One actual-zeta target beyond Fourier support one

Current report 13: Forward true-zeta localization: a contractive comparison and a sharp universality obstruction

Current report 14: Arithmetic transfer for a fixed large-prime sector

Current report 15: Round 7: an arithmetic large-prime sector for zeta's half-gap problem

Current report 16: Independent audit of the half-threshold arithmetic transfer

Current report 17: Round 8: isolate the actual arithmetic remainder

Current report 18: The actual-zeta two-scale target: short-prime projection and a centered tail

Current report 19: Independent audit of the actual-zeta short-polynomial identity

Current report 20: A bounded positivity audit of the two-scale target

Current report 21: Independent review of the fixed positivity minorant

Current report 22: Round 9: actual prime arithmetic for the Dyson–Montgomery programme

Current report 23: A concrete 186-to-covariance transfer beyond the square-root divisor level

Current report 24: Independent review of the complementary-modulus covariance component

Current report 25: A fixed interaction between two large prime divisors

Current report 26: Round 9: duplicate profile avoided, two-prime interaction tested

Current report 27: Independent review: the fixed two-large-prime interaction

Current report 28: Prime powers are negligible in the remaining residual energy

Current report 29: Independent review of the actual prime-power tail estimate

Current report 30: Mesoscopic damping: the missing arithmetic edge term and its rate requirements

Current report 31: Independent review of the mesoscopic edge obligation

Current report 32: Round 10: a power saving for one actual shifted–prime discrepancy

Current report 33: A power improvement for a smooth packet of the actual shifted discrepancy

Current report 34: Independent audit of the completed–shift coefficients and spacing bound

Current report 35: Independent review: actual shift kernel and genuine–prime replacement

Current report 36: Actual–prime variance ranges and a missing logarithmic mixed moment

Current report 37: Independent review of the arithmetic range and mixed–moment audit

Current report 38: Round 11: remove the shift–length loss using actual RH prime input

Current report 39: Removing the shift–length loss from the actual completed pairing, under RH

Current report 40: Independent review of the RH centered–small–arc improvement

Current report 41: A real–prime subfamily prevents a coefficient–only power saving

Current report 42: Independent review of the canonical conductor–mass construction

Current report 43: The log–weighted prime–tail moment: a primary–source check and its arithmetic remainder

Current report 44: Independent review of the centered mixed–moment remainder

Current report 45: Round 12: three exact tests of the remaining arithmetic gap

Current report 46: The actual support saturates positive sampling at the power scale

Current report 47: Independent review of the inherited count and sampling constants

Current report 48: Two explicit hypothesis failures in the proposed dispersion transfer

Current report 49: Direct Selberg control of the centered mixed remainder: the precision and sign gap

Current report 50: Independent review of the three bounded Round 12 attempts

Current report 51: Dyson–Montgomery round 13: extract a rational core and retain the signed remainder

Current report 52: Averaging the rational resonances: one power–saving extraction, and its retained main term

Current report 53: Round 13 — Genuine–prime minor arcs and a fixed–inner–interval averaging audit

Current report 54: The signed kernel norm: exact main term and a large coherent CRT remainder

Current report 55: Independent audit of the smooth signed–kernel norm

Current report 56: Round 13 independent integration review

Current report 57: Dyson–Montgomery round 14: an actual Type I removal and a quantitative CUE heat theorem

Current report 58: A smooth long factor removes an actual Type I component

Current report 59: Independent review: smooth long–factor removal

Current report 60: CUE background at the selected smallest gap: a direct finite–N bound

Current report 61: Independent review: selected CUE background and the finite heat rate

Current report 62: Independent root review: two bounded R14 advances

Current report 63: Fable PR11 snapshot 89393d5: separate intake and two–sided arithmetic audit

Current report 64: Independent background and boundary objections to the same Fable snapshot

Current report 65: Fable 2073028: repaired arithmetic, finite Fock bound and retained gaps

Current report 66: Independent audit of the F1 cutoff repair and F3 mass bound

Current report 67: F1 repair: corrected coefficient, surviving sign issue, and cutoff scaling

Current report 68: The mass cutoff gives a finite Fock operator bound

Current report 69: Independent follow–up: C β E background repair at 2073028

Source index and artifact receipt

Riemann zeros, random matrices and actual prime arithmetic

Rounds 6–14: complete takeover supplement / 第 6–14 轮完整接棒补编

Source checkpoint: [2a9ec8192ca8aed1cfdd1bad648c3a3a429e31ba](#), 2026-09-05. This book follows the earlier main handoff and Rounds 4 – 5 supplement. It prints the complete substantial reports and independent reviews from the subsequent work, including the final R14 proofs and both later Fable intakes. It is a faithful research archive, not an announcement that a famous conjecture has been solved.

当前结论：没有证明 RH、Montgomery/GUE 猜想、AH-Pairs 的反驳、新的 ζ 半间隙定理或小于 186 的素数间隙。已完成的工作包括：真实算术误差项的两次改进、一个可在所需尺度下消去的精确 Type I 分量、有限 CUE 热流的定量误差、具有精确假设的反例与障碍，以及完整保留的负向变分实验。数值增益、普通证明、独立内部审查与形式验证在本文中始终分开。

The intended next researcher is GPT-6 Astra or a human analyst taking over without access to the live conversation. Read this synthesis first, then the integrated report for the chosen lane, then its full proof and independent review. The main earlier archive remains necessary for the original ACUE constructions, fixed-family arithmetic-transfer proof, Galilean heat lemma, force-energy identities, and Rounds 1 – 5 prime-gap certificates. Those earlier documents are linked by pinned source paths rather than silently reproduced under a new date.

1. What the programme is trying to prove

The main target is a substantive theorem about **actual zeta zeros** and the Montgomery – Dyson connection. Random-matrix examples and heat flow are tools for isolating missing hypotheses, not substitutes for the arithmetic explicit formula. The most concrete accepted reductions concern the precise AH-Pairs formulation, including its possible near-diagonal mass. General AH must not be replaced by a simple hard-core process with gap at least one half. Multiplicity or near-coincident pairs can occupy the zero lattice point. An arbitrarily small gap, or a positive proportion of gaps below one half, need not by itself refute the full hypothesis unless the statistic also excludes that near-zero freedom.

There are three explicit formulations of the arithmetic target. Each is sufficient under its recorded RH and limiting hypotheses; none has been proved here.

Notation warning / 记号须按章节理解. N denotes the CUE matrix dimension in the heat-flow reports, the Dirichlet cutoff floor($T/\log^6 T$) in the resolvent reports, and a short factor of size $X^{.4}$ in the R13 Type II discussion. These are different objects. D_N is finite polynomial heat depth, while $D_Q^\vee$ is an arithmetic progression discrepancy. The resonator length L , a smooth long-factor length, and the L used to truncate the minimum-gap tail also have separate local meanings. Each full report resets its definitions. Exceptional close pairs in an RMT model do not by themselves refute a density-version AH statement for zeta.

1.1 The fixed two-width resolvent target

For fixed $c > 0$ define

$$I_T(c) = \int_0^T \left| \frac{\zeta'}{\zeta} \left(\frac{1}{2} + \frac{c}{\log T} + it \right) \right|^2 dt, \quad W_T = \frac{2[\sinh(2)/I_T(1) - \sinh(1)/I_T(1/2)]}{T \log^2 T}.$$

Under RH and AH-Pairs, the reviewed reduction gives

$$W_T \rightarrow W_{\text{AH}}, \quad 0.06239 < W_{\text{AH}} < 0.06240 < \frac{1}{16}.$$

The sine-kernel prediction is approximately 0.0822714431214773. Therefore

$$\liminf_{T \rightarrow \infty} W_T \geq \frac{1}{16}$$

would refute AH-Pairs under RH. The previously discussed target 0.07 is also sufficient, but unnecessarily strong. The two widths were chosen to cancel the bounded near-diagonal parameter $P_0(T)$ exactly; no convergence of that parameter is assumed. Tail truncation, finite-height endpoints, the Gamma factor, ζ'/ζ normalization and the holomorphic-square passage are part of the printed proof and reviews. The reduction itself is established; the displayed actual-zeta lower bound is not.

1.2 A compact out-of-band Fourier test

Fix the recorded nonnegative smooth bump φ , integral one, supported on $[6/5, 7/5]$ and symmetric about $13/10$. For the normalized form factor F_T ,

Observable	RH + AH-Pairs prediction	Sine-kernel prediction
$\int \varphi(\alpha) F_T(\alpha) d\alpha$	7/10	1
Centered prime-covariance remainder	-3/5	-3/10

An actual lower limit strictly above 7/10 for the first observable, or strictly above -3/5 for the corresponding centered remainder, would exclude AH-Pairs. The limit 1 would establish this one smooth Montgomery prediction. It would not automatically prove the entire pair-correlation conjecture or RH. The full prime-prime term, prime-continuum cross terms, continuous mean square and diagonal must remain present.

1.3 A shrinking mesoscopic correction

The residual route uses $N = \text{floor}(T/\log^6 T)$ and

$$R_c(t) = -\frac{\zeta'}{\zeta} \left(\frac{1}{2} + \frac{c}{\log T} + it \right) - \sum_{n \leq N} \Lambda(n) n^{-1/2 - c/\log T - it}.$$

Round 8 proves under RH that $W_T = B + E_T + o(1)$, where $B \approx 0.4560939793292317$ and

$$E_T = \frac{2[\sinh(2) \|R_1\|_2^2 - \sinh(1) \|R_{1/2}\|_2^2]}{T \log^2 T}.$$

The fixed-width sufficient target is therefore $\liminf E_T \geq -0.3935939793292317\dots$. The residual combination is signed and of leading order. Positivity of its two individual energies does not prove the target.

For $b=2c$ set $r_T(b) = \|R_{b/2}\|_{\frac{2}{2}} / (T \log^2 T)$. The coupled statistic

$$C_T(b) = b^2 \left[2 \sinh b r_T(b) - 2 \sinh(2b) r_T(2b) - \frac{1}{2b} \right]$$

has sine prediction 0 and AH prediction $-3/4$ after the same nuisance cancellation. The sufficient lower bound has uniform quantifiers on a slowly growing envelope $G(T)=o(\log \log T)$, followed by $B \rightarrow \infty$:

$$\lim_{B \rightarrow \infty} \liminf_{T \rightarrow \infty} \inf_{B \leq b \leq G(T)} C_T(b) > -\frac{3}{4}.$$

A fixed-width limit does not authorize a prescribed growing rate. The needed error is a first correction of relative order $1/b$, not merely a leading relative $o(1)$ estimate. The reports state an equivalent sufficient signed logarithmic mixed-moment bound. Its positive diagonal is explicit, but its centered off-diagonal remainder remains open.

2. The strongest arithmetic results at this checkpoint

Write $X=T^\alpha$ with $6/5 \leq \alpha \leq 7/5$, $H=X/T$, and $Q=X^{(523/1000)}$. Thus H ranges from $X^{(1/6)}$ to $X^{(2/7)}$. The selected complementary squarefree modulus family is inherited from the 186 paper's verified ordinary analytic input. Repeated representations of one modulus are counted only once. These numbers belong to a particular arithmetic component; 0.523 is not a zero-distribution exponent or a new prime-gap record.

The exact smooth discrepancy retains a fixed $V(h/H)$, the original sinc kernel, $\mu(q) \log((m-h)/q)$, and the primitive principal sum. Its useful bounds progress as follows.

Stage	Proved bound for the specified discrepancy	Assumptions and remaining limitation
R9 source transfer, summed absolutely over h	$O_A(HX \log^(-A)X)$	Ordinary source theorem; fixed logarithmic saving cannot absorb polynomial H .
R10 shift completion	$O(\sqrt{HX(X+Q^2)} \log^4 X)$	Unconditional; original smooth joint kernel restored after separation.
R11 centered small arcs	$O(\sqrt{X(X+Q^2)} \log^5 X) = O(X^{1.023} \log^5 X)$	Under RH; removes $\sqrt{H}$, but remains above $X \log X$.
R14 short Möbius divisor portion	$O_J(HX(UQ/X)^J \log^2 X)$	Unconditional for $UQ \leq X/2$ and fixed $J \geq 2$; the exact signed remainder is retained.

Round 11 uses an actual RH centered-prime small-arc estimate, including its derivative version, before sampling at the rational frequencies. Equal fractions are merged first. Its coefficient band mass and local arc length cancel H . The integer mean and primitive Ramanujan mean are treated separately and both remain accounted for. This is a genuine improvement of a defined arithmetic error bound. It still leaves a factor $X^{.023}$, apart from logarithms, above the required covariance scale. It is not a proof that this factor is necessary for primes.

Round 14 makes a further exact reduction. Define

$$\Lambda_{\leq U}(n) = \sum_{r \mid n, r \leq U} \mu(r) \log(n/r), \quad \Lambda = \Lambda_{\leq U} + \Lambda_{> U}.$$

For $U=X^{.4}$ and $J=4$, the bound for the first portion is $O(X^{(1711/1750)} \log^2 X) = o(X \log X)$. More generally every fixed $\eta > 0$ with $\eta < .477$ permits $U \leq X^{(.477-\eta)}$, choosing fixed J with $J\eta > 2/7$. This follows from exact progression Poisson summation in the genuinely smooth long cofactor. Its zero mode cancels the actual primitive principal term. The normalized joint kernel has uniform derivatives; no regularity is assigned to the short Möbius sequence. A product of several rough long factors does not satisfy this hypothesis merely because its total length exceeds Q . The remaining $\Lambda_{> U}$ discrepancy is an exact signed arithmetic object and is unestimated.

The strongest full selected smooth-packet bound is consequently still the R11 RH estimate. The short-divisor removal identifies a smaller residual problem; it is not a bound for every remaining piece. The complete covariance additionally requires complementary moduli, support main terms, other ranges and continuous centering. None may be suppressed because the selected component is attractive to calculate.

3. The finite CUE theorem that is now on firm ground

For Haar CUE(N), let $\delta_{\min}$ be the smallest angular gap, and B_N the inverse-square circular background at that gap's midpoint. Round 14 proves directly from the exact finite- N three-point Gram determinant that

$$E \sum_{i: \delta_i \leq \varepsilon} B_i \leq \frac{N^6 \varepsilon^3}{18}, \quad 0 < \varepsilon \leq \pi.$$

The endpoint zeros in the determinant cancel the singular endpoint weight. The proof enlarges an endpoint-weighted count to all short ordered pairs only after that cancellation; enlarging a midpoint-weighted count directly would be invalid. Circular ordering includes the wrap gap. No conditional density of a selected minimum is assumed. Combining the estimate with the classical CUE minimum-gap law gives $B_N/N^2 = O_p(1)$, with an explicit truncation-tail bound.

For the specified scalar-heat evolution of the characteristic polynomial, with D_N its first positive discriminant time, the existing deterministic Galilean lemma then yields

$$\frac{8D_N}{\delta_{\min}^2} - 1 = O_p(N^{-2/3}), \quad D_N - \delta_{\min}^2/8 = O_p(N^{-10/3}).$$

This is a quantitative approximation in probability for finite CUE. It is not a rate for convergence of the entire depth distribution, a theorem for general β , a stochastic Dyson Brownian motion theorem, or an established property of actual zeta zeros. The earlier qualitative ratio and deterministic lemma are explicitly credited in the new proof. No global literature novelty claim is made.

The distinction matters for the original research aspiration. RMT supplies a rigorous reference law and precise mechanisms. A zeta theorem needs the arithmetic hypothesis that forces the corresponding local behavior. Reusing the RMT conclusion under a name such as “alternative COE” cannot supply that missing hypothesis.

4. Failed approaches that materially changed the programme

4.1 More fixed resonator features did not cross the half-gap threshold

The original request to enlarge S_2/S_3 polynomials was checked against the archive before computation; it duplicated an earlier sweep. The replacement in R7 was the sharp integer mark for a prime divisor above $\sqrt{L}$. The fixed family has a direct unique-large-prime decomposition and a reviewed limiting arithmetic transfer. Its 30-dimensional half-gap margin is about -0.01465492379421 , improving its matched baseline by only 1.429×10^{-8} . It remains worse than the older 48-feature value near -0.0146547256 .

A later proposed multiplicative profile also duplicated an already resummed experiment. R9 instead used the nonmultiplicative event of two distinct prime divisors above $L^{1/3}$. Its unique unordered double-prime decomposition is exactly coprime; its singly marked formula requires an explicit repeated-prime error. The new fixed 30-dimensional margin is -0.0146549114371551 versus -0.0146549380840028 for its matched baseline, a floating gain of about 2.66×10^{-8} at scaled Gram condition about 5.36×10^7 . Its actual finite-integer frozen-vector test at $L=100000$ has margin -0.0374094621535042 .

These are new concrete arithmetic families and useful transfer checks. The numerical gains are not interval-certified and are far from crossing zero. Their spans do not contain the entire older 48-feature space. They prove no global no-go for resonators, and do not justify another blind coefficient sweep. All coefficients and matrices remain in the repository with hashes.

4.2 Generic positivity does not deliver the arithmetic gain

The explicit R8 minorant gives a valid bound around -0.208674513 for W , far below $1/16$. Optimality was proved only inside that fixed one-parameter family. A realized stationary half-grid determinantal process satisfies the available low-band information and positivity yet attains $W_{AH} < 1/16$. This blocks an inference from those hypotheses alone. It does not describe actual primes or rule out an arithmetic theorem.

The two residuals do arise from the same centered arithmetic function $\psi(x)-x$, and after the justified prime-power removal from $\theta(x)-x$. Their common origin is useful structure. Its mere existence, or the positivity of each residual norm, does not fix the sign of their weighted difference. The pole and endpoint terms must first be bounded; an unregularized infinite prime series in the critical strip is not a legal replacement.

4.3 Deterministic heat and protected traces do not force GUE

The R7 flow report supplies a deterministic contraction bound under its stated external-field condition, but leaves an actual-zeta boundary-propagation estimate open. Its exact finite polynomial family begins on a half-grid up to rotation and retains all normalized gaps at least one half under forward flow, tending toward a clock. This defeats a proposed implication from those dynamical hypotheses. It is not a counterexample obeying the full zeta explicit formula.

The protected trace algebra also remains matched under the recorded full DBM comparison. At a protected frequency $m=N/2$ the stochastic microscopic generator contribution at CUE is π^2 . Thus a deterministic calculation cannot simply discard stochastic smoothing because a collection of low moments agrees. The original AH definition and the distinction between actual and artificial wrap gaps remain essential.

4.4 Actual conductor geometry obstructs norm-only shortcuts

R11 constructs an admissible terminal complementary family using two primes of exponent .09 and 346 distinct smaller primes of exponent 343/346000. It has $\gg Q/\log^{348} X$ moduli near Q . At terminal $d=q>Q/2$, no other permitted multiple exists, so the full Möbius coefficient is exactly $1/d$. The coefficient squared mass is at least a constant times $H/\log^{348} X$. This rules out a fixed-power coefficient-norm improvement for that full family. It does not rule out pruning, different weights, or cancellation with genuine prime sums.

R12 uses the same actual frequencies to prove sharpness up to logarithms for a positive sampling step, including its known local energy and derivative envelopes. The saturating polynomial is artificial. This is an obstruction to that general sampling argument, not a lower bound for the actual centered-prime functional and not evidence that $X^{.023}$ is unavoidable for primes.

4.5 Direct import of the 186 dispersion theorem fails specific premises

An additive twist of a genuine-prime short coefficient can lose the required Siegel – Walfisz property. The explicit modulus-3 example, at legal source scales $M=X^{.6}$ and $N=X^{.4}$, gives a discrepancy of leading size $N/\log N$. The source theorem remains true; the transformed coefficient does not satisfy its hypothesis. The coherent shift interval also cannot be replaced by the Cartesian product of all its local residue images without an enormous class cost. Taking H itself as the short convolution length falls below the checked source range.

These failures identify what an averaged replacement must preserve: the joint m,a,d,h phases and cross-prime coherence. They do not prove that every averaged use of dispersion fails.

4.6 Short-interval theorems miss the needed scale or correction

The checked Guth – Maynard corollary concerns $h \geq X^{(2/15+\epsilon)}$, with fixed ϵ . The mesoscopic shell here has exponent $s/(b+s) \rightarrow 0$. Even the corollary's ϵ -zero endpoint misses that shell for $b>13$; the following remark's slight fixed improvement does not resolve a vanishing exponent. Almost-all PNT counts also do not by themselves provide a variance constant.

The checked three-integral comparisons have constant losses larger than the shrinking sine-versus-AH signal. Their finite- T errors can be made small on a sufficiently slow diagonal, so it would be wrong to blame only height errors. The limiting lower bound misses a first correction. R12's actual Selberg audit retains both cutoff crossings and the joint mean, yet does not supply the necessary sign or b^{-3} precision in the mixed moment. These are source-specific quantitative failures, not a claim that all existing analytic methods are exhausted.

4.7 A positive rational core does not lower-bound the signed whole

R13 extracts the zero-rational Type II core with total RH replacement error $O(X^{.923} \log^2 X)$, retaining the explicit integral main term. An admissible restricted positive block can have size at least a constant times $X^{1.123}/\log^{348} X$. Other phases and the actual long coefficients may cancel it. R14 gives a constructive instance of such complete cancellation when a long factor is smooth.

The exact signed-kernel norm has a CRT/Poisson main term and an explicit remainder. Large original common divisors give short enough CRT periods for smooth decay. The small-gcd long-period terms remain. A coherent positive off-diagonal subsum of size at least a constant times $Q^2 H/\log^{696} X$ is not a lower bound for the complete

signed remainder. Reduced denominator gcd and original modulus gcd are different quantities. Even an ideal unrestricted integer norm of order XH would not alone give the desired prime-specific bound.

5. The R6 prime-gap work that had not reached the earlier PDFs

R6 constructs the full signed cap operator, rather than defining an operator only through a 77×77 matrix. Its finite fragment measure is not a probability measure conditioned on the outer domain. The erased-coordinate adjoint is an unweighted lift with outer support. On product amplitudes the single erased g factor must not become a g^2 conditional expectation. The face multiplier can be negative; positive-semidefinite assumptions are invalid.

For the old mass-orthogonal projection P and a radial projection P_V that need not commute with it, the useful direction is

$$r = (I - P)Tf, \quad h = P_V r, \quad w = (I - P)h, \quad \langle f, Tw \rangle = \|h\|^2.$$

For unit f, its normalized mixed entry is $\|h\|^2 / \|w\|$, not simply $\|w\|$. The actual mass projection of Tf is used, without assuming an exact Ritz vector. The active radial cells are frozen and define the chosen subspace; small excluded mass is not a bound for omitted residual energy.

The fixed k=39 direct cap quotient rises from 0.9943963993644909 to 0.9944678209006830, about 71.4215 ppm. An exact dyadic-rational/polynomial-root certificate proves the stored direction is genuinely outside the old 77-space on positive-measure cells. That exact independence certificate does not certify the numerical gain. The Gram condition is large, the quotient is still about 5532.18 ppm below one, and arithmetic support restoration remains unproved for the candidate. The actual 2×2 plane supplies only about 1.26% of its crossing requirement, without estimating the full residual.

Four public compact NPZ witnesses retain all candidate and projection arrays except the regenerable $77 \times N$ density cache. The four original full archives remain local with hashes and array-by-array compaction receipts. This is explicit storage compaction, not deletion of adverse outputs. The prime-gap bound remains 186; the prior R4 rigorous margin gain and R5 geometry constraints should be read in their separate earlier supplement. Prime-gap broad sweeps were paused when the user redirected the main lane to actual zeta and Dyson – Montgomery.

6. Fable corrections must travel with their source snapshots

The 89393d5 intake repaired the moment coefficient to $\Pi_4 \sim 6a\epsilon^{-4}$, while also finding a sign error in the refuter's own derivative probe and a table mixing different v values. A successful replay of a refuter reproduces its failures; it does not certify every assertion in it. Finite drift does not disprove an asymptotic fixed-family limit without a proved constant and threshold. The later Astra fixed-family transfer remains separate from Fable's unfinished quantitative-rate discussion.

At 2073028 the corrected F1 coefficient and fixed-v table are accepted. The finite prime-sum cutoff is explained by the reviewed incomplete-gamma limit. The F3 assertion of an infinite field norm on the mass cutoff is false. A complete weighted sector proof gives, for $g(u) = 2\sin(\pi u/2)$,

$$\|K\| \leq 2 \int_0^1 \frac{|g(u)|^2}{u^2} du = 4\pi \operatorname{Si}(\pi) - 8 \approx 15.27212735.$$

This proves finite boundedness for the stipulated idealized operator and uniformly for the literal discrete grids. It lies above $\pi^2/2$, so it proves no sharp spectral wall. The first interval has infinite du/u measure, and a nonzero constant there is not a normalized Galerkin basis vector. Numerical extrapolations near 4.6456 remain numerical evidence, with no proved full arithmetic-to-Fock transfer.

The general- β background repair is still only partial. Its purported exact finite- N CUE formula is a sine limit; its comparator vanishes at $q=2\pi$ while the true normalized two-point density is one. Its uniform replacement $v' = v(1+O(\varepsilon/w))$ fails near an endpoint. A direct conditional triple integral restores the intended $L^{(\beta+1)}c^{(2\beta+1)}$ exponents, but does not prove general- β density control. Uniform one-point intensity alone cannot control the background of a selected smallest pair: a randomly rotated clustered configuration is a counterexample to that inference. The independent R14 finite-CUE argument uses the needed higher correlation structure and does not rely on these uncorrected general- β claims.

Both intake reviews are printed in full below. The 141-file and 160-file Fable snapshots remain separately pinned in the repository; their complete duplicated source texts are not printed again in this supplement. This separation preserves original errors and later corrections without rewriting history.

7. Prioritized next work, with concrete success criteria

These are research proposals, not claims that any famous conjecture is now within a guaranteed final step. Prefer one bounded calculation or lemma with a falsifiable acceptance condition before a larger search.

- 1. Estimate the actual remaining arithmetic term after R14 removal.** Start with the exact $\Lambda_{\{>U\}}$ discrepancy, fixed $U \leq X^{(.477-\eta)}$, and the original joint sinc/log kernel. Use a stated Heath – Brown decomposition and preserve the primitive mean. Identify one remaining factor pattern and prove an aggregate bound strictly below $X \log X$, or exhibit the exact source hypothesis that fails. A smooth long variable is already handled; relabeling that case is not new progress.
- 2. Keep signed phase information through the long-variable average.** R12 forbids a blanket SW inheritance claim for twisted coefficients. R13 isolates rational cores and their main terms. A useful next theorem must estimate their full signed combination with the actual coefficients, not only a positive subblock or a norm for arbitrary integer polynomials. Record every denominator, numerator and gcd range. Success requires a power improvement over the $X^{1.023}$ bound or a directly useful signed covariance estimate, not a logarithmic cosmetic gain.
- 3. Attack the mixed genuine-prime remainder at the exact first correction.** Use the finite centered measure and logarithmic companion already proved in R11. The positive diagonal is $b^{-2}+2b^{-3}$. A strict one-sided improvement for the combined off-diagonal remainder, or its integrated version, could imply the mesoscopic AH-excluding criterion. State uniformity through twice a valid slow envelope; keep the two prime-continuum terms and continuum square together. A mere $O(1)$ energy bound is not the requested result.
- 4. Use the CUE theorem as a precise reference theorem.** Any zeta heat-flow transfer must supply an actual local arithmetic hypothesis strong enough to control selected near-pair backgrounds and true boundaries. Test it against the half-grid, rotated-cluster and artificial-wrap examples before trying to prove it. A general- β extension

would require a correct finite- N $n=3$ bound with the stated domain; it cannot be imported from weak process convergence alone.

5. **If returning to variational search, demand a structural reason for a gain.** The simple power-sum, resummed multiplicative, single-large-prime and double-large-prime directions are already archived. Choose a genuinely different fixed integer feature or a proved extension of the available mixed arithmetic form. Give exact coefficient meaning and a controlled transfer before treating a continuum model as a zeta test. A negative finite span is not a global barrier.
6. **Formalize selected stable components after proof review.** Good bounded candidates are the primitive Poisson cancellation, the exact mass-cutoff Fock inequality, the finite- N CUE Gram/singular-weight estimate, and the finite rational outside-span witness. The FLT formalization work motivates careful decomposition and proof checking; it supplies no unproved analytic input. A proof-assistant certificate must be reported separately from Python checks and internal reviews.

暂缓：同一负向系数空间的重复扫描、把普通 PSD 或点过程正性再包装成算术输入、忽略 ζ 的 pole/mean 交叉项、以假设缺失的 periodization 推出真实小间隙、重复开启 Fable 会话，以及把数值外推或程序复算称为已证明的历史级定理。研究主线应以“缺少的真实算术估计是否推进”为进度标准，而不是文件数或模型轮数。

8. 中文接棒判断与阅读顺序

最值得继续的变化，不是又找到了一个更像随机矩阵的模型，而是已经把若干模糊目标改写为可核查的真实算术不等式。第 7 轮消除了 AH 中可能不收敛的近零质量；第 8 – 11 轮把两个 resolvent 的能量差写成同一个中心化素数误差的耦合，并保留了必要的 pole、端点和连续均值；第 10 – 11 轮确实改进了指定误差项；第 14 轮又把其中一个完整、带原始权重的 Type I 分量严格消掉。这些结果值得保存，但距离所需的严格符号增益仍有未完成的算术部分。

有限 CUE 方面，现在有选定最小间隙背景在自然 N^2 尺度的紧性和热碰撞时间的定量误差。这是一条清楚的普通证明链。它说明在真实 RMT 中什么条件足够，同时也明确揭示了向 ζ 迁移时缺少什么。不能把 CUE 的高阶相关性偷偷当作 ζ 的已知事实，也不能把一般 AH 偷换为没有重根的硬核 ACUE。

接手时建议先读 R14 综合报告，确认最新真正完成的两项结果；再读 R7/R8 确定终极目标和归一化；随后读 R9 – R13 的真实算术链及失败的迁移假设。若选择变分路线，先读 R7/R9 的完整负向试验，避免重复。若选择素数间隙路线，先读 R6 与前一本 R4 – R5 补编，注意 cap-only、支撑恢复和浮点/区间证书的区别。最后对照 Fable 两次审查，避免重新引入已被否定的公式。

0.6725007... 在本项目核对的源文献中是关于临界线上简单零点的无条件比例下界；它不直接估计本文的带外相关性。186 是此前源论文的素数间隙结论，本文没有改成更小。FLT 的形式化技术是一种组织和验证证明的方式，不会把缺失的算术引理自动补齐。希望得到重大定理的目标不变，但接棒者应对所有这些类别区别保持严格。

9. Evidence, preservation and reproducibility contract

Each complete source report below has its original repository path, pinned Git blob identity, SHA-256 digest and byte length in the JSON index. The assembler compares working bytes to that exact Git blob before use. It does not take the current working tree's newer contents on trust. All original mathematical statements, caveats, failed attempts,

commands and review-status history remain in the embedded bodies. A historical “review pending” sentence is preserved when a later separately printed review supersedes it. Historical agent instructions are source records, not new instructions for the reader to execute.

Presentation changes are limited to the established `build_handoff.cleaned` transformations, heading nesting, resolving Markdown links/images to the source checkpoint, and spelling vertical bars inside table-cell inline mathematics as equivalent LaTeX commands. Raw single bars become `\vert` and existing LaTeX double bars become `\Vert`; this prevents Markdown from treating mathematical bars as column separators. Every replacement is recorded with its source line and column, and no mathematical meaning is changed. Code fences and inline code are protected from this table repair. The JSON index records the digest of each displayed body as well as its raw source. Any later layout-only adjustment belongs in a separately recorded rendering step, not a silent source edit.

Code, JSON certificates, arrays, small plots, run logs and integration receipts are retained as repository artifacts by path and hash rather than printed as enormous tables. The index covers the selected round/review folders and their integration-log folders. Earlier proof dependencies and snapshot manifests reached by resolvable source links are indexed too. Third-party full papers and optional large original caches remain in the adjacent local archive where the source receipts say so; their absence from the public book is not concealed. The one in-report R6 figure is linked to the pinned raw image.

Rebuild this Markdown and index from the repository root with:

```
python3 tools/build_round6_14_handoff.py
```

The per-round recheck scripts printed in the reports are the appropriate bounded acceptance surfaces. Some require pinned local primary PDFs or earlier runtime dependencies; read their supplied arguments before running. An exact finite check validates its algebraic case, not an asymptotic prime theorem. Agreement of floating matrices, a tiny eigensolver residual or a reproduced fit is not an interval enclosure. Internal independent review is not external peer review or Lean verification. This supplement itself performs assembly and provenance checks only; it does not rerun the mathematical experiments.

For the earlier context, retain the pinned main archive and R4 – R5 supplement alongside this book. The present index is a precise catalogue of later-round coverage, not a claim to include private conversations or every earlier source document for a second time.

Complete source texts / 完整原文目录

No.	Group	Source report
01	R6	Round 6: a full signed-operator direction beyond the 77-dimensional sieve trial
02	R6	The full signed cap operator and a mass-orthogonal direction beyond 77 dimensions
03	R6	Residual directions outside the 77-dimensional sieve trial space

No.	Group	Source report
04	R6	Independent finite marked-space regression for the signed sieve operator
05	R6	Round 6: one full cap-operator residual beyond the 77-dimensional family
06	R6	Provenance and compact witness export
07	R6	Independent review of the frozen radial profile outside the old span
08	R7	Round 7: two explicit actual-zeta targets for Dyson – Montgomery
09	R7	Two Poisson scales remove the AH near-diagonal parameter
10	R7	Independent review of the two-scale actual-zeta reduction
11	R7	Independent review of the actual-zeta Poisson transfer
12	R7	One actual-zeta target beyond Fourier support one
13	R7	Forward true-zeta localization: a contractive comparison and a sharp universality obstruction
14	R7	Arithmetic transfer for a fixed large-prime sector
15	R7	Round 7: an arithmetic large-prime sector for zeta's half-gap problem
16	R7	Independent audit of the half-threshold arithmetic transfer
17	R8	Round 8: isolate the actual arithmetic remainder
18	R8	The actual-zeta two-scale target: short-prime projection and a centered tail
19	R8	Independent audit of the actual-zeta short-polynomial identity
20	R8	A bounded positivity audit of the two-scale target
21	R8	Independent review of the fixed positivity minorant

No.	Group	Source report
22	R9	Round 9: actual prime arithmetic for the Dyson – Montgomery programme
23	R9	A concrete 186-to-covariance transfer beyond the square-root divisor level
24	R9	Independent review of the complementary-modulus covariance component
25	R9	A fixed interaction between two large prime divisors
26	R9	Round 9: duplicate profile avoided, two-prime interaction tested
27	R9	Independent review: the fixed two-large-prime interaction
28	R9	Prime powers are negligible in the remaining residual energy
29	R9	Independent review of the actual prime-power tail estimate
30	R9	Mesoscopic damping: the missing arithmetic edge term and its rate requirements
31	R9	Independent review of the mesoscopic edge obligation
32	R10	Round 10: a power saving for one actual shifted-prime discrepancy
33	R10	A power improvement for a smooth packet of the actual shifted discrepancy
34	R10	Independent audit of the completed-shift coefficients and spacing bound
35	R10	Independent review: actual shift kernel and genuine-prime replacement
36	R10	Actual-prime variance ranges and a missing logarithmic mixed moment
37	R10	Independent review of the arithmetic range and mixed-moment audit
38	R11	Round 11: remove the shift-length loss using actual RH prime input
39	R11	Removing the shift-length loss from the actual completed pairing, under RH

No.	Group	Source report
40	R11	Independent review of the RH centered-small-arc improvement
41	R11	A real-prime subfamily prevents a coefficient-only power saving
42	R11	Independent review of the canonical conductor-mass construction
43	R11	The log-weighted prime-tail moment: a primary-source check and its arithmetic remainder
44	R11	Independent review of the centered mixed-moment remainder
45	R12	Round 12: three exact tests of the remaining arithmetic gap
46	R12	The actual support saturates positive sampling at the power scale
47	R12	Independent review of the inherited count and sampling constants
48	R12	Two explicit hypothesis failures in the proposed dispersion transfer
49	R12	Direct Selberg control of the centered mixed remainder: the precision and sign gap
50	R12	Independent review of the three bounded Round 12 attempts
51	R13	Dyson--Montgomery round 13: extract a rational core and retain the signed remainder
52	R13	Averaging the rational resonances: one power-saving extraction, and its retained main term
53	R13	Round 13 — Genuine-prime minor arcs and a fixed-inner-interval averaging audit
54	R13	The signed kernel norm: exact main term and a large coherent CRT remainder
55	R13	Independent audit of the smooth signed-kernel norm
56	R13	Round 13 independent integration review

No.	Group	Source report
57	R14	Dyson--Montgomery round 14: an actual Type I removal and a quantitative CUE heat theorem
58	R14	A smooth long factor removes an actual Type I component
59	R14	Independent review: smooth long-factor removal
60	R14	CUE background at the selected smallest gap: a direct finite-N bound
61	R14	Independent review: selected CUE background and the finite heat rate
62	R14	Independent root review: two bounded R14 advances
63	Fable intake	Fable PR11 snapshot 89393d5: separate intake and two-sided arithmetic audit
64	Fable intake	Independent background and boundary objections to the same Fable snapshot
65	Fable intake	Fable 2073028: repaired arithmetic, finite Fock bound and retained gaps
66	Fable intake	Independent audit of the F1 cutoff repair and F3 mass bound
67	Fable intake	F1 repair: corrected coefficient, surviving sign issue, and cutoff scaling
68	Fable intake	The mass cutoff gives a finite Fock operator bound
69	Fable intake	Independent follow-up: C β E background repair at 2073028

Current report 01: Round 6: a full signed–operator direction beyond the 77–dimensional sieve trial

Collection: R6 — full signed prime-sieve operator and residual direction.

Source: research/reports/prime186_round6.md.

SHA-256: `e1d83da09a54cf3cf9f2fd77b3ec5f246c00cd1647ba3954146152d2e417003c` . Git blob: `8f627cc49aa9f211a8f09e0e92330ac0a1a1b5ab` . Original bytes: 11486.

Round 6: a full signed–operator direction beyond the 77–dimensional sieve trial

The new fixed-geometry $k=39$ cap trial has directly evaluated quotient 0.9944678209006830, up from 0.9943963993644909 by approximately 71.4215 ppm. This improvement comes from a direction generated by the full cap operator outside the old 77-dimensional space. A separate exact algebraic certificate proves that the frozen new direction is outside that space. The gain itself remains a floating-point result; the quotient is still approximately 5532.18 ppm below one, and the actual arithmetic support has not been restored. No smaller prime gap is proved.

This is the completed bounded diagnostic requested after Round 5. It is unrelated to the earlier finite zeta operator K_L . It neither repeats the same 77-by-77 eigenproblem nor performs another radius sweep.

1. Full operator, actual measure and signed weights

The source Hilbert space is $H = L^2(H_O, \nu^{\otimes 39})$, where ν is the finite fragment measure. It is not a probability measure conditioned on the outer domain. Profiles are extended by zero. Erasing coordinate i means

$$(E_i f)(Y) = \int f(Y \oplus_i X) d\nu(X), \quad (E_i^* \nu)(X) = 1_{H_O}(X) \nu(X_i).$$

Freeze the original geometry, the exceptional constant 0.34 and all hybrid parameters. With

$$a_h = 0.9919601604, \quad b_h = -0.000843183, \quad d_0 = 0.0088830226,$$

the full cap operator is

$$T_{\text{cap}} = \rho_* \sum_i E_i^* (d_0 1_{H_{0,i}} + a_h 1_{H_{1,i}} + b_h) E_i.$$

This formula acts on arbitrary profiles in the Hilbert space; it is not a definition by a 77-dimensional matrix. The face multiplier takes values 1, $a_h + b_h$, and the negative value b_h . The operator is bounded and self-adjoint, with

$$-\rho_* \mid b_h \mid C_{\text{op}} I \leq T_{\text{cap}} \leq \rho_* C_{\text{op}} I, \quad C_{\text{op}} \leq 4$$

in the fixed support range. No positivity assumption is used.

The actual arithmetic operator additionally uses the complete inner predicates and the outer projection P_O . Its supported quotient has denominator $\|P_O f\|^2$. That denominator cannot be replaced by $\|f\|^2$ unless f is already supported on O . A cap improvement, even one exceeding one, would still require this restoration step.

The [full derivation](#) checks the source measure, adjoints, boundedness, overlap integrals and the exact finite cell representation. For the fixed inward cap domains, profiles constant on total-cell/fragment-layer atoms form an invariant subspace of the full cap operator. Thus the step trial has no omitted within-atom action component. It can still have a large component outside the 77 polynomials. The actual tail/prefix support predicates require more fragment information and do not inherit that atom invariance.

2. A radial compression that exposes a new direction

Let U be the old 77-dimensional step-polynomial space, P its true mass-orthogonal projection, $Q=I-P$, and f the stored optimized trial. Put $G = \prod_i g(t_i)$. Choose the closed subspace

$$V = \{G\phi(s) : \phi \text{ is arbitrary on a fixed selected set of radial cells}\}.$$

U and V are not nested, and their projections need not commute. The construction is

$$r = QTf, \quad h = P_V r, \quad w = Qh.$$

It follows exactly, for every f in U , that

$$\langle f, Tw \rangle = \|h\|^2, \quad \|w\|^2 = \|h\|^2 - \|Ph\|^2.$$

For unit f the two-dimensional off-diagonal entry is therefore

$$\beta = \frac{\|h\|^2}{\|w\|},$$

not simply $\|w\|$. No exact Ritz assumption is needed, because the implementation computes PTf from the full mass projection $M^{-1}Bc$, rather than replacing it with λf .

Writing q for the radial pushforward of $1_{H_O} G^2 \nu^{\otimes 39}$, D for the basis cross-mass densities, and b for the pushed-forward pairing with Tf , the radial profile is

$$\phi(s) = 1_{\text{active}}(s) \frac{b(s) - D(s)^T M^{-1} Bc}{q(s)}.$$

All fragment-cap layers and their signed face contributions are integrated before this projection. One-dimensional adjoint convolutions evaluate b . The single erased factor g , the reciprocal output factor $1/g(t_i)$, the mesh factor and the common tilt normalization are retained. A normalized conditional average under $g^2 \nu$ would give a different operator.

The active radial set is frozen from the declared numerical cutoff. It defines V ; the experiment does not claim that excluded cells carry no residual energy. In particular, a small excluded q -mass does not bound omitted $\|r\|^2$.

3. A separate exact certificate of leaving U

The numerical norm calculation is not the only evidence that the new direction is outside U. The stored float64 profile is a specified finite list of exact dyadic rationals. Its entries at radial indices 0 through 12 are exactly zero, while

$$\phi(18422) = -\frac{6264072493613325}{4611686018427387904} \neq 0.$$

Fix the other 38 coordinate cells at index zero and vary the first coordinate cell index r . After dividing by the common positive product G , each original basis function is a polynomial in r of degree at most 12: at most six from its radial power and at most six from its power-sum signature. If the frozen radial function belonged to U, the resulting polynomial would vanish at 13 distinct indices and hence everywhere, contradicting the displayed nonzero entry.

The argument concerns positive-measure product cells, not isolated points. Exact rational inequalities place every selected coordinate cell below the smallest cap and the whole product cell inside the first outer shell. Its measure is the positive number h_{mesh}^{39} . An [independent reviewer](#) decoded the binary NPZ directly and checked the nonzero rational and these support inequalities. That review verifies the polynomial-root argument; it does not independently rerun the additional modular-rank computation below.

A second finite certificate evaluates the old basis at 77 explicit positive-measure cells and obtains rank 77 modulo 1,000,000,007. The new radial column is zero at these cells and nonzero at the additional cell, proving rank 78 over the rationals. See [the certificate script](#) and [exact receipt](#).

These certificates establish independence of the frozen functions. They do not bound the quantitative distance to U or certify the Rayleigh gain. Those remain separate integral questions.

4. What the bounded numerical experiment gives

For the official 98,304 grid, tilt 20 and density cutoff 10^{-9} :


Quantity	Observed value
Original 77-space matrix quotient a	0.9943963991909279
Radial compression norm squared	0.00006670688589594228
Outside-U norm squared	0.00006589186717477095
Fraction of radial norm squared remaining outside U	0.9877820901062194
Normalized new-direction diagonal b	0.043583189070450945
Actual normalized mixed entry β	0.008217784708256407
Larger eigenvalue on $\text{span}(f,w)$	0.9944674193880856
Full 78-space matrix quotient	0.9944678209367751
Direct evaluation of the final profile	0.9944678209006830

The simple two-dimensional gain is 71.0202 ppm. Reoptimizing the other coefficients adds approximately 0.40155 ppm. For the realized two-dimensional plane, crossing one would require $\beta^2 > (1 - a)(1 - b)$; the observed coupling squared supplies only about 1.26% of that threshold. This describes that plane, not the uncomputed full residual.

The [detailed report](#) preserves the coarse calculation and two controls. Tilt 20 and tilt 25 matrix values agree to approximately 2.2×10^{-16} ; their direct values differ by approximately 1.53×10^{-10} . Changing the density cutoff from 10^{-9} to 10^{-8} shifts the result by about 0.0177 ppm. The old scaled mass Gram condition number remains approximately 2.28×10^{10} . Small residuals and agreement between contractions are not outward rounding bounds.

A separate root process replayed the fine calculation in an isolated copy. Its matrix quotient, direct quotient, compressed norm, outside norm, coupling, active cells and radial profile agree exactly with the recorded run. That reproducibility does not replace interval certification.

5. Where this particular radial signal appears



The plot uses the saved arrays, with no new operator integration. Half the observed radial compression energy lies below normalized radius approximately 0.95403; its peak is approximately 0.95599, close to the inner cutoffs. The central half lies between approximately 0.94578 and 0.97133. This localization suggests examining the transitions of the face domains when designing a less restrictive compression, but does not establish a causal explanation. The comparison curve is the product-weight mass $\mathcal{G}^2$, not the optimized trial's $\mathcal{F}^2$ mass. This is the radial projection's energy, not the full residual's energy.

6. Independent checks, stored data and remaining work

The [independent residual audit](#) gives the correct nonnested identities, approximate-Ritz corrections, inverse-Gram error formulas and signed crossing criteria. Eighty exact rational indefinite-operator examples and finite weighted-measure tests pass. A separately written [38-state marked model](#) checks nonuniform product mass, nonrectangular support, noncommuting projections and an explicit negative quadratic witness. Neither toy model estimates the prime-sieve operator.

The [integration receipt](#) records exact replay of both algebraic test suites and the saved-output audit, followed by one fine-grid numerical replay. All original outputs are retained. Four compact NPZ witnesses publish every candidate and projection array except the reproducible 77-by-N D cache. The 183,046,525 bytes of original full NPZ files are preserved in the local archive; their hashes and array-by-array compaction checks are public. No coefficient, active cell, scalar result or new radial profile is omitted from the compact witnesses. The separate [archive check](#) verifies all 37 verbatim intake files, all 44 retained arrays against the full originals, Python syntax, report links and an isolated replay of the exact independence certificate.

The earliest route to a rigorous variational gain would freeze rational coefficients and the radial profile, then enclose its ordinary mass and signed mixed forms directly. That would avoid depending on an exactly computed orthogonal projection. Since the current value remains below one, it would not itself yield a smaller prime gap. The next substantive decision is whether another implementable compression captures substantially more of the full cap residual, or whether a new support/product profile is needed. The full residual norm and the true support-restored form remain uncomputed. No further iteration is included in this checkpoint.

Current report 02: The full signed cap operator and a mass–orthogonal direction beyond 77 dimensions

Collection: R6 — full signed prime-sieve operator and residual direction.

Source: research/prime-gaps/round6/operator-proof/FULL_SIGNED_CAP_OPERATOR.md.

SHA-256: `020099a5ab6927f17b82dfbb2b96aa5b2cfbc3720c41d388f9c8f39b08d0339b` . Git blob: `a2e19c9a336fa654171c4bf75b182f54be3f319e` . Original bytes: 21540.

The full signed cap operator and a mass–orthogonal direction beyond 77 dimensions

Status: ordinary-mathematics derivation and independent implementation review, 5 September 2026. The results below concern a fixed cap geometry and its true fragment measure. They do not certify a new prime gap, replace arithmetic support restoration, or assert positivity of the hybrid operator.

1. Sources and the precise question

The pinned official source is [OpenAI, PrimeGaps186, commit 61340d0b74163003b32756bb16e91d9209a5e330](#). The mathematical inputs used here are the main paper, equations (3.15) – (3.24), Lemma 3.8, and equations (4.28) – (4.40), and the numerical companion, § 1.1 and equations (2.12) – (2.15). The local source copies and hashes are recorded in the adjacent provenance JSON. The published 77 coefficients define a trial space; they do not define the full variational operator.

For the fixed original geometry, use the exact inward cap cells of the companion, recomputed with outer dimension $k = 39$ and face dimension $k - 1 = 38$. This is the geometry evaluated by `round5/geometry-trial/cap_trial.py` with its default configuration. The ordinary Hilbert-space arguments apply to every fixed $k > 1$ and every such nested family. They do not inherit the original $k = 40$ numerical support-error certificate for $k = 39$.

The task is to identify the full operator T , its actual mass inner product, and a rigorously meaningful direction outside the 77-dimensional space. A further practical goal is to evaluate one such direction without forming the enormous full matrix or calculating the entire norm of Tf .

2. Measure, erasure, and adjoint

Let Ω_ζ be the space of fragment multisets in $(0, \zeta]$, equipped with

$$\nu = \nu_\zeta = e^\nu \zeta L(\Pi_\zeta), \quad d\Pi_\zeta \text{ has intensity } du/u.$$

The prefactor is the product $e^\nu \zeta$, not $e^\nu \zeta$. This is a finite measure, not a probability measure. If $t = |X|$ is the sum of the fragments, then its pushforward is $\rho_D(t/\zeta) dt$. For $0 < c \leq \zeta$,

$$\nu|_{\{\max X \leq c\}} = \nu_c, \quad (\cdot | X | \cdot)_* \nu_c = \rho_D(t/c) dt. \quad (1)$$

Thus a prime cap changes the actual restricted measure. It does not condition a probability distribution, and totals larger than the cap remain possible. The factor c/ζ from the Poisson void probability has already canceled the change of prefactor in (1).

Let $H_O \subset \Omega_\zeta^k$ be the outer cap domain and set

$$H = L^2(H_O, \nu^k), \quad \langle f, g \rangle = \int_{H_O} \bar{f}g d\nu^k.$$

Every function is extended by zero outside H_O . Write $Y = X_{\hat{j}}$ for the retained tuple and $Y \oplus_i Z$ for insertion in coordinate i . Define

$$(E_i f)(Y) = \int f(Y \oplus_i Z) d\nu(Z). \quad (2)$$

The face space uses $L^2(\Omega_\zeta^{k-1}, \nu^{k-1})$. Direct Fubini gives

$$(E_i^* \nu)(X) = \mathbf{1}_{H_O}(X) \nu(X_{\hat{j}}). \quad (3)$$

In particular, erasure is an integral, not a conditional mean. There is no factor k , h , or Z in (2) or (3). The companion divides every quadratic form, including face pairings, by the same constant $(hZ)^k$. Multiplying both Hilbert inner products by this same constant leaves (3) unchanged. Independently normalizing the k - and $(k-1)$ -coordinate spaces would change the adjoint and would not reproduce the source conventions.

3. Three different operators must remain distinct

Let $H_0 \subset H_1$ be the cap-only face domains. Let $L_0 \subset L_1$ be the actual face domains, with $L_a \subset H_a$. In the source,

$$L_1 = H_1 \cap L_{\text{new}}, \quad L_0 = H_0 \cap L_{\text{old}} \cap L_{\text{new}}.$$

Let $O \subset H_O$ impose the actual outer source predicates, and write $P_O = \mathbf{1}_O$. With the source's fixed hybrid parameters,

$$\begin{aligned} m &= 0.99998, \quad \lambda_h = 0.008, \quad K_{\text{ex}} = 0.34, \\ a_h &= m^2 - m\lambda_h = 0.9919601604, \quad b_h = (1 - m/\lambda_h)(1 - m)K_{\text{ex}} = -0.000843183, \\ d_0 &= 1 - a_h - b_h = 0.0088830226. \end{aligned}$$

These constants belong to the fixed original physical outer radius. Changed geometry requires its own exceptional constant, as audited in Round 5.

Define the face multipliers

$$m_H = d_0 \mathbf{1}_{H_0} + a_h \mathbf{1}_{H_1} + b_h, \quad m_L = d_0 \mathbf{1}_{L_0} + a_h \mathbf{1}_{L_1} + b_h. \quad (4)$$

Their respective values on the three nested regions are 1, $a_h + b_h$, and b_h . The last value is negative. The implemented full-face mask H_f contains the support of every erased marginal of an outer-supported profile, so writing the last term as $b_h \mathbf{1}_{H_f}$ in the finite engine gives the same quadratic forms as the unrestricted b_h in (4). The source operators and the threshold-normalized operators are

$$A = \sum_{i=1}^k E_i^* m_H E_i, \quad B = \sum_{i=1}^k E_i^* m_L E_i,$$

$$T_{\text{cap}} = \rho_* A, \quad T_{\text{inner}} = \rho_* B, \quad T_{\text{arith}} = \rho_* P_O B P_O. \quad (5)$$

Here $\rho_* = 2624989/10^7$. Explicitly, the full cap action is

$$(T_{\text{cap}} f)(X) = \rho_* \mathbf{1}_{H_O}(X) \sum_i m_H(X_i) \int f(X_i \oplus_i Z) d\nu(Z). \quad (6)$$

Replace m_H by m_L for T_{inner} . For T_{arith} , insert $\mathbf{1}_O$ both outside and inside the integral. Formula (6) defines the operator on the full Hilbert space, not only on polynomial profiles.

The cap quotient is $\langle f, T_{\text{cap}} f \rangle / \|f\|^2$. The actual sieve quotient for $P_O f \equiv 0$ is

$$\frac{\rho_* \langle P_O f, B P_O f \rangle}{\|P_O f\|^2}. \quad (7)$$

The denominator in (7) is essential. On all of $\mathcal{H}$, the denominator of the Rayleigh quotient of T_{arith} is $\|f\|^2$, which agrees with (7) only when $f = P_O f$. For the actual sieve problem one works on $L^2(O, \nu^k)$ or keeps this distinction explicit.

The arithmetic theorem requires more than a cap quotient exceeding one: the actual supported profile and source/realization hypotheses must hold. The restoration estimate is a lower bound on its quadratic form, not another fixed linear operator formed by subtracting independent error bars.

4. Boundedness and self-adjointness without positivity

All multipliers in (4) are real and bounded. Fubini and (3) imply self-adjointness of each term $E_i^* m E_i$. To give a bound independent of the much larger unrestricted mass, suppose H_O has total support at most S . Put $t_j = |X_j|$, $\mathfrak{s}_i = \sum_{j \neq i} t_j$, and

$$w_i = S - \mathfrak{s}_i + (k-1)t_i.$$

Because the total density is at most one,

$$\int_{H_O(Y)} \frac{d\nu(X_i)}{w_i} \leq \int_0^{S-\mathfrak{s}_i} \frac{dt}{S-\mathfrak{s}_i+(k-1)t} = \frac{\log k}{k-1}.$$

The boundary $\mathfrak{s}_i = S$ has zero fiber measure. Weighted Cauchy – Schwarz followed by $\sum_i w_i = kS$ gives

$$\sum_i \|E_i f\|^2 \leq C_{\text{op}} \|f\|^2, \quad C_{\text{op}} = \frac{Sk \log k}{k-1}. \quad (8)$$

For the fixed $k = 39$ geometry, the safe bound $C_{\text{op}} \leq 4$ has an exact rational verification in the Round 5 source-geometry audit. Since $b_h \leq m_H$, $m_L \leq 1$,

$$-\rho_* |b_h| C_{\text{op}} I \leq T_{\text{cap}}, T_{\text{inner}}, T_{\text{arith}} \leq \rho_* C_{\text{op}} I. \quad (9)$$

For the compressed operator T_{arith} these inequalities use that P_O is an orthogonal projection. Also

$$A - B = \sum_i E_i^* \{d_0(\mathbf{1}_{H_0} - \mathbf{1}_{L_0}) + a_h(\mathbf{1}_{H_1} - \mathbf{1}_{L_1})\} E_i \geq 0.$$

The negative full-face term cancels in this difference. None of these statements makes A or B positive semidefinite. In particular, a Perron – Frobenius argument, a positive-kernel power method, or a replacement of b_h by $|b_h|$ would require a different justification and would change the problem.

If the trial and the domains are permutation invariant, these operators commute with coordinate permutations. Thus the symmetric subspace is invariant. The factor k in the symmetric face formulas below comes from this symmetry, not from the erasure adjoint.

5. Exact mixed integrals and the information retained by fragments

For arbitrary profiles f, g , the true mixed form is

$$\langle g, T_{\text{cap}} f \rangle = \rho_* \sum_i \int m_H(Y) \overline{E_i g(Y)} E_i f(Y) d\nu^{k-1}(Y). \quad (10)$$

Expanding the two marginals introduces two independent erased coordinates conditional on the same retained configuration Y . It does not introduce two independent retained configurations. This is exactly the shared-root convention in the primary proof.

For symmetric f , direct calculation of $\|T_{\text{cap}} f\|^2$ can reduce the double coordinate sum to $i = j$ and $i \equiv j$. The first contribution has one shared retained tuple and two erased copies, together with the outer-fiber mass. The second has $k - 2$ shared coordinates and four distinguished coordinate copies. They are different overlap integrals. A third operator moment has the finitely many equality patterns of three erasure indices, including the different placements of two equal indices. This is a useful organizational reduction, but no independence assumption between their shared cap states is valid.

The cap-only masks use each total cell and the largest fragment. Let all distinct positive caps be

$$0 < c_1 < \cdots < c_L = \zeta,$$

where a larger unused ambient cap can also be included. Partition a coordinate by its total cell $C_j = [jh, (j+1)h)$ and its fragment layer $c_{\ell-1} < \max X \leq c_\ell$, with the layer-zero cumulative measure defined to be zero. The exact mass of this atom is

$$\mu_{j\ell} = \int_{C_j} \{\rho_D(t/c_\ell) - \rho_D(t/c_{\ell-1})\} dt. \quad (11)$$

For $\ell = 1$ omit the second term. Empty or zero-mass atoms may be removed. Formula (11) is an exact pushforward of the fragment measure. It is not an approximation that concentrates continuous totals or fragments at a point.

For the companion's inward cell domains, outer membership is a function of the index sum $r = \sum j_i$ and the largest layer. Face membership is a function of r_i and the largest retained layer. Therefore the finite subspace of functions constant on products of these atoms is invariant under T_{cap} : integrating coordinate i sums its atom values against $\mu_{j\ell}$, and all remaining arguments in (6) depend only on retained atom labels. The adjoint then produces an atom-constant output. Since T_{cap} is self-adjoint, this finite subspace is also reducing.

The official step trial lies in this subspace. Consequently its full $T_{\text{cap}}f$ and its residual outside the 77-dimensional span lie there too. There is no uncomputed within-cell component of $T_{\text{cap}}f$ for that trial and those fixed cell domains. The full Hilbert space still contains other, within-atom functions; no claim about its entire spectrum follows from this invariance alone.

The coordinate layers cannot in general be discarded. Knowing all total cells does not determine which outer-shell marginals survive after erasure. The action on coordinate i depends on the largest retained fragment layer, and different erasures can leave different largest layers. Replacing this information by one global probability before taking products changes (10).

The actual predicates in O, L_0, L_1 involve activated fragments and inclusive prefix sums, such as $\max_{p \geq \xi} \{\sum_{q \geq p} q + \phi(p)\}$. Equal total and equal largest fragment do not determine these values. Thus the preceding atom subspace generally is not invariant under T_{arith} . More fragment information, exact conditional predicate averages, or a valid support-error argument is required there.

For clarity, product conjugation does not remove this issue. Write $G(X) = \prod_i g(t_i) > 0$ and transfer $f = Gu$ to the mass $G^2 dv^k$. The conjugated action is

$$G^{-1} T_{\text{cap}} G u(X) = \rho_* \sum_i \frac{m_H(X_i)}{g(t_i)} \int \mathbf{1}_{H_O}(X_i \oplus_i Z) g(|Z|) u(X_i \oplus_i Z) dv(Z). \quad (12)$$

There is one erased factor g , and a factor $1/g(t_i)$ outside. It is not a normalized average against $g^2 dv$.

6. Full residual, nonnested radial compression, and the signed two-dimensional test

Let $U \subset H$ be the span of the 77 step profiles and let P_U be orthogonal projection in the true mass inner product. If $\phi_1, \dots, \phi_{77}$ are an independent basis, put

$$M_{ab} = \langle \phi_a, \phi_b \rangle, \quad b_a = \langle \phi_a, Tf \rangle, \quad P_U Tf = \sum_a \phi_a (M^{-1} b)_a. \quad (13)$$

An exact nonsingular Gram matrix is needed; dependent basis vectors can instead be removed. Euclidean coefficient projection is not (13).

For any $f \in U$, define $r = (I - P_U)Tf$. Self-adjointness and orthogonality give

$$\langle f, Tr \rangle = \langle Tf, r \rangle = \|r\|^2. \quad (14)$$

This does not require f to be an exact Ritz eigenvector. An exact Ritz equation is required only to replace $P_U Tf$ by a scalar multiple of f .

To avoid computing all of r , choose a closed subspace V whose projection can be evaluated, and put

$$v = P_V r, \quad w = (I - P_U)v. \quad (15)$$

Here v is called h in the numerical implementation; the letter h elsewhere in this note denotes the grid spacing. The spaces U, V need not be nested. The projection order in (15) is essential. Since $r \perp U$,

$$\langle f, Tw \rangle = \langle r, w \rangle = \langle r, v \rangle = \|v\|^2, \quad \|w\|^2 = \|v\|^2 - \|P_U v\|^2. \quad (16)$$

In particular $\|w\| \leq \|v\| \leq \|r\|$, and $v \equiv 0$ forces $w \equiv 0$. The coupling is $\|v\|^2$, not generally $\|w\|^2$. The superficially similar direction $(I - P_U)P_V T f$ does not generally satisfy (16).

Set $F_2 = \|f\|^2$, $q = \|v\|^2$, $z = \|w\|^2 > 0$,

$$\lambda = \frac{\langle f, T f \rangle}{F_2}, \quad \tau = \frac{\langle w, T w \rangle}{z}, \quad \eta = \frac{q}{\sqrt{F_2 z}}.$$

The orthonormal basis $f/\sqrt{F_2}$, $w/\sqrt{z}$ has the matrix

$$\begin{pmatrix} \lambda & \eta \\ \eta & \tau \end{pmatrix}.$$

Its larger eigenvalue is

$$\lambda_+ = \frac{\lambda + \tau + \sqrt{(\lambda - \tau)^2 + 4\eta^2}}{2} > \lambda. \quad (17)$$

No positivity of T is used. For $\lambda < 1$ and $\tau < 1$, crossing one is equivalent to $\eta^2 > (1 - \lambda)(1 - \tau)$. A certified lower bound $\tau \geq \tau_L$ gives a sufficient condition by substituting τ_L . In the cap problem, (9) supplies the inexpensive choice $\tau_L = -\rho_* \mid b_h \mid C_{\text{op}}$. Since $z \leq q$, the still more conservative sufficient condition

$$\frac{q}{F_2} > (1 - \lambda)(1 + \rho_* \mid b_h \mid C_{\text{op}}) \quad (18)$$

avoids calculating $\langle w, T w \rangle$. It may be too weak numerically, but it is a valid signed-operator certificate. A strict positive improvement in (17) alone is not the threshold condition (18) or an arithmetic theorem.

7. An exact radial projection computed by one-dimensional convolutions

In this section G , every P_η , and every g_j use the rational midpoint evaluations of the step trial; they are constant on their total cells. Take V to consist of $\mathbf{1}_{H_0} G(X) a(r)$ with arbitrary values $a(r)$ on the retained radial cells. A fixed subset of radial cells may be used. Since the grid is finite, this is a closed subspace. It does not contain all 77 power-sum profiles.

Use the companion's un-tilted notation

$$d_c(j) = h^{-1} \int_{C_j} \rho_D(t/c) dt, \quad Z = \sum_{j < n} g_j^2, \quad K_c(j) = g_j^2 d_c(j)/Z.$$

Let $M_{d,\sigma}^c(s)$ be the coefficient sequence obtained by integrating $G^2 P_\sigma$ over d coordinates with index sum s , common cap c , and normalization $(hZ)^{-d}$. This is exactly the positive-partition convolution moment of the numerical companion. Define

$$\Delta M_{d,\sigma}^l(s) = M_{d,\sigma}^{c_l}(s) - M_{d,\sigma}^{c_{l-1}}(s).$$

For the lowest layer the second sequence is zero. Let $c_O(r)$ be the inward outer cap on radial cell r , with all formulas zero outside the outer radial masks. The mass density of the radial subspace, in the common normalization $(hZ)^{-k}$, is

$$D(r) = M_{k,\emptyset}^{c_O(r)}(r), \quad \|Ga\|^2 = \sum_r D(r) \mid a(r) \mid^2. \quad (19)$$

For the input polynomial profile f , let $a_{\ell,\eta}(s)$ denote the signed prefix sum of the allowed outer-shell affine rows of the companion. Thus on a background of index sum s whose largest fragment lies in layer ℓ ,

$$E_{\eta}f(Y) = hG(Y) \sum_{\eta} a_{\ell,\eta}(s) P_{\eta}(Y). \quad (20)$$

One must sum the permitted shell rows with their signs before forming any products. Put

$$B_{\ell}(s) = \sum_{\eta} a_{\ell,\eta}(s) \Delta M_{k-1,\eta}^{\ell}(s),$$

and let $m_{\ell}(s)$ be the signed face multiplier (4) on this layer. Direct substitution of (20) into the mixed integral (10) gives the normalized radial adjoint density

$$N_T(r) = \frac{\rho_* kh}{Z} \sum_{\ell: \alpha_{\ell} \leq c_O(r)} \sum_{j+s=r} g_j d_{c_O(r)}(j) m_{\ell}(s) B_{\ell}(s). \quad (21)$$

More explicitly, $\langle Ga, Tf \rangle = \sum_r \overline{a(r)} N_T(r)$. To verify the factor, each erased marginal contributes h , the retained $k-1$ coordinates contribute $(hZ)^{k-1}$, all forms are divided by $(hZ)^k$, and symmetry supplies k . Their product is kh/Z . The test function's erased coordinate contributes exactly one factor $g_j d_c(j)$.

It follows that

$$P_V Tf = G \frac{N_T(r)}{D(r)} \quad (22)$$

where $D(r) > 0$; on zero-mass cells the value is immaterial. No distribution of fragments has been replaced by a conditional mean before the product. The layer sums in (21) perform the required conditional integration explicitly.

If $N_U(r)$ is the radial mixed density of the profile $P_U Tf$, then

$$\begin{aligned} v &= P_V(I - P_U)Tf = G \frac{N_T(r) - N_U(r)}{D(r)}, \\ q &= \|v\|^2 = \sum_r \frac{|N_T(r) - N_U(r)|^2}{D(r)}. \end{aligned} \quad (23)$$

The sequence N_U is obtained from the ordinary k -coordinate moments for the coefficients in (13). Both the projection and (23) therefore use the full mass measure. Restricting to any predetermined set of positive-mass radial cells preserves all the identities in (15) – (16); it only chooses a smaller V .

The independent mixed forms $\langle \phi_a, Tv \rangle$ and $\langle v, Tv \rangle$ follow from the same signed face formula by replacing the radial polynomial of one or both inputs by the arbitrary sequence in (23). This gives a direct two-dimensional or 78-dimensional Rayleigh calculation without evaluating $\|Tf\|^2$ and without assuming that the radial projection captures the entire residual.

Numerical exponential tilt is only a normalization device

The exploratory engine uses $Z_\theta = \sum_j g_j^2 e^{-\theta t_j}$ and moments built from $g_j^2 e^{-\theta t_j} d_c(j)/Z_\theta$. In (19) multiply the resulting moment by $e^{\theta(r+k/2)h}$. In (21), use Z_θ and multiply $B_\ell(s)$ by $e^{\theta(s+(k-1)/2)h}$. The affine rows in (20) continue to use the un-tilted erased factor $g_j d_c(j)$. These substitutions exactly undo the tilt in the actual product measure.

When combining forms evaluated at two tilts, first convert every norm and pairing to the same common normalization. Converting from the current Z_θ normalization to a reference Z_{θ_0} multiplies all forms by $(Z_\theta/Z_{\theta_0})^k$. The true operator, projections, and Rayleigh quotient do not depend on the numerical tilt.

8. Independent review, current limits, and the next proof obligation

The Round 6 script `residual-trial/radial_residual.py` was reviewed against (19) – (23). Its `radial_adjoint` has the correct face dimension $k-1$, the factor $\rho_* kh/Z$, one erased g , the signed sum of permitted outer-shell affine rows, cap-layer differences, and the background exponential correction. Its computation of $P_U T f$ uses the mass Gram solve, rather than assuming an exactly solved Ritz equation. Its active radial-cell cutoff is a legitimate choice of a smaller V in ordinary mathematics. Floating-point negative mass or Gram uncertainty is not thereby certified away.

The separate exact rational regression `operator-diagnostic/finite_marked_operator_check.py` was also inspected. It uses nonuniform masses, a nonrectangular cap domain, and nonnested trial/radial spaces. The source verifies mass self-adjointness, (14), (16), product conjugation, a projection-order counterexample, and a negative quadratic witness equal to $-1/38880$. This is a structural test of the identities; its five-atom measure is not the Dickman measure and gives no $k=39$ numerical bound.

The earliest computational issue for a rigorous cap improvement is now precise: enclose the Dickman cell masses, Gram projection, radial adjoint density, and signed mixed forms coherently, or rationalize a chosen new radial profile and outwardly certify its ordinary mass and cap forms. A floating-point projected residual may be a useful way to select a profile, but its exact orthogonality is not needed after a concrete rational profile is fixed and its two-dimensional Gram and numerator forms are certified directly. This avoids pretending that numerically computed projections are exact.

The next arithmetic issue is separate and harder. A new radial profile changes the weights of the retained and deleted fragment configurations. The old outer failure covers, old inner loss numbers, and the Round 4 positive alpha credit cannot simply be reused with their old numerical values. One must evaluate their valid formulas for the new profile, or evaluate the actual predicates directly with sufficient retained fragment state. The cap operator is an exact relaxed operator at the fixed grid; it is not the actual supported sieve operator.

This closes the ordinary operator derivation. It justifies a concrete out-of-77-space search direction and a signed two-dimensional test. It does not establish a threshold crossing, a complete support-restored $k=39$ trial, or a prime gap below 186.

Current report 03: Residual directions outside the 77–dimensional sieve trial space

Collection: R6 — full signed prime-sieve operator and residual direction.

Source: `research/prime-gaps/round6/residual-audit/SIEVE_RESIDUAL_AUDIT.md`.

SHA-256: `be0a8c6ea16af80f27b7c9abc77466884bfee6c02c8750a2f4d72d5381beacc7` . **Git blob:** `a6e8e8bd3ce5ab8987a86f297c7d3817e483eb85` . **Original bytes:** 19067.

Residual directions outside the 77–dimensional sieve trial space

Status: independent algebraic and measure-theoretic audit. The operators below are the finite-dimensional-sieve variational operators from the prime-gap programme, not any earlier zeta-zero operator. No positive-semidefinite assumption is made. No new prime-gap result is established.

Primary inputs checked: *Improved short gaps between primes*, § 4.4, equations (4.30) – (4.40), and § 4.1 – § 4.2; the retained local text is `../sources/openai-short-gaps.txt`, SHA256

`ded13a7c74fcfce64e85769e05b5869803dcccdf53b88be2c2f3c0b344f95ee84`. Current cap-engine files checked were `cap_trial.py` and `optimize_cap.py` in the published round-5 geometry-trial directory. In particular, their denominator uses the product g^2 measure, their erased-coordinate kernels contain one factor g , and their signed face multiplier retains the negative b_h term. They are exploratory cap-form engines, not evaluations of the fully restored arithmetic form.

1. The two operators that must remain distinct

Let

$$H = L^2(H_O, d\nu^{\otimes k}), \quad k = 39,$$

with every function extended by zero outside H_O . The measure is the unscaled fragment measure of the primary paper, not a probability measure conditioned on H_O . Its norm is the outer-square mass.

For a retained configuration $Y=X$ with coordinate i erased, define

$$(E_i f)(Y) = \int 1_{H_O}(Y, X_i) f(Y, X_i) d\nu(X_i). \quad (1)$$

The face space has measure $d\nu^{(k-1)}$, with no division by the mass of this fibre. Consequently

$$(E_i^* \nu)(X) = \nu(X_i) \quad (X \in H_O). \quad (2)$$

The indicator of H_O belongs in (2) if the ambient space is instead written as the full product space. Fubini proves the adjoint relation immediately. Replacing (1) by a conditional expectation changes the operator and its variational value.

The relaxed cap operator is

$$T_{\text{cap}} = \rho_* A, \quad A = \sum_i E_i^* (d_0 1_{H_{0,i}} + a_h 1_{H_{1,i}} + b_h) E_i. \quad (3)$$

The actual support-restored operator, on the supported subspace or extended by zero, is

$$T_{\text{arith}} = \rho_* P_O B P_O, \quad B = \sum_i E_i^* (d_0 1_{L_{0,i}} + a_h 1_{L_{1,i}} + b_h) E_i. \quad (4)$$

Here P_O is multiplication by the actual outer support indicator, and $L_{0,i}, L_{1,i}$ are the actual inner domains. The source establishes $A - B \geq 0$ because the negative full-face term cancels in their difference. It does not identify A with B , and it does not imply an ordering of the two operators after arbitrary changes of their outer trial functions.

A direction generated for T_{cap} can establish an improvement of the relaxed cap variational problem. It does not by itself establish a positive restored arithmetic form. If the basis is subsequently changed from u_a to P_{O,u_a} , its Gram matrix and projection both change. One cannot reuse the old 77-dimensional projection after this support change.

Both operators are bounded and self-adjoint. If

$$\sum_i E_i^* E_i \leq C_{\text{op}} I, \quad b_h < 0, \quad 0 < a_h + b_h < 1, \quad d_0 + a_h + b_h = 1,$$

then the three possible multipliers are 1, $a_h + b_h$ and b_h . Thus

$$-mI \leq T \leq MI, \quad m = \rho_* |b_h| C_{\text{op}}, \quad M = \rho_* C_{\text{op}}, \quad (5)$$

for either operator, with the operator bound proved for its actual domain. The negative term must not be dropped when taking a lower bound. In the original geometric range the paper's argument gives $C_{\text{op}} \leq S k \log(k)/(k-1)$; using $C_{\text{op}}=4$ requires the corresponding numerical inequality. Nothing here assumes $T \geq 0$.

2. Mass conjugation and the radial projection

Write $G(X) = \text{product}_i g(t_i)$, where t_i is the coordinate total or its specified step representative. If amplitudes $p=f/G$ are used, their Hilbert mass is $G^2 \text{dnu}^k$, restricted to H_O . Conjugation gives

$$G^{-1} T_{\text{cap}}(Gp)(X) = \rho_* \sum_i \frac{m_i(X_i)}{g(t_i)} \int 1_{H_O}(X_i, Y_i) g(t(Y_i)) p(X_i, Y_i) d\nu(Y_i). \quad (6)$$

This identity explains why denominator kernels contain g^2 , whereas each erased-coordinate marginal contains one factor g . Putting g^2 in the latter integral without its compensating conjugation factors is a different operator.

Let $\sigma(X)$ be the chosen radial observable: either the true sum of totals, or explicitly the sum of coordinate-cell indices in the fixed step model. Let V be the closed subspace

$$V = \{G(X)\phi(\sigma(X)) : \phi \in L^2(q)\}, \quad q = \sigma_*(1_{H_O} G^2 \nu^{\otimes k}). \quad (7)$$

No assumption that V contains the 77-dimensional polynomial space U is appropriate: it generally does not. For $f \in H$ define the pushed-forward signed measures

$$d_f = \sigma_*(1_{H_0} G f v^{\otimes k}), \quad b_f = \sigma_*(1_{H_0} G(Tf) v^{\otimes k}).$$

They have Radon – Nikodym derivatives relative to q , and

$$P_V(Tf - \lambda f) = G\left(\frac{db_f}{dq} - \lambda \frac{dd_f}{dq}\right). \quad (8)$$

When all three measures have densities, this is the proposed formula $G(b_f(s) - \lambda d_f(s))/q(s)$. For a cell-index radial observable it is a ratio of discrete masses at each index, not a continuous-density formula. Set the ratio to zero on a zero-mass radial cell.

This is a genuine orthogonal projection because its error is orthogonal to every $G \phi(\sigma)$, directly from the defining pushforward identity. It remains valid with fragment-dependent cap indicators, provided those indicators are included before taking each pushforward.

The one-coordinate mass of a retained cell with maximum fragment at most ch is

$$\nu\{t \in [jh, (j+1)h), \max \text{fragment} \leq ch\} = h \int_0^1 \rho_D((j+u)/c) du. \quad (9)$$

The survival average alone omits h . A common overall normalization of all outer Hilbert masses is harmless, but independently normalizing face and outer masses changes the adjoint unless the ratio is carried through. In the engine's tilted convolution representation, factors h , Z and the exponential untilt must therefore be kept consistently for the new cross forms as well as the original Gram matrix.

3. The direct outside-space residual needs no Ritz hypothesis

Let U be a finite-dimensional subspace of H , P its true Hilbert orthogonal projection, and $Q=I-P$. Let $f \in U$ be a unit vector and put

$$\lambda = \langle f, Tf \rangle, \quad r = QTf.$$

If $a = \|r\| > 0$ and $v = r/a$, then f and v are orthonormal, and

$$\langle f, Tv \rangle = \langle Tf, v \rangle = \frac{\langle QTf, QTf \rangle}{a} = a. \quad (10)$$

This identity is true for every $f \in U$, not just an exact Ritz vector. Exact Ritz status means $PTf = \lambda f$, and is needed only to identify $Tf - \lambda f$ with QTf .

For an approximate Ritz vector define $e = PTf - \lambda f$. Then

$$Tf - \lambda f = e + r, \quad e \perp r, \quad \|r\|^2 = \|Tf - \lambda f\|^2 - \|e\|^2. \quad (11)$$

Thus the full eigen-equation residual can overstate the outside- U residual. A tiny residual in the 77-by-77 matrix only controls e for that compressed matrix; it says nothing about r until an additional operator action or outside-space form is evaluated.

With $\mu = \langle v, Tv \rangle$, the larger Ritz value on $\text{span}\{f, v\}$ is exactly

$$\Lambda_2 = \frac{\lambda + \mu + \sqrt{(\lambda - \mu)^2 + 4a^2}}{2}. \quad (12)$$

If $\lambda < 1$, the signed lower bound $\mu \geq -m$ implies the sufficient crossing test

$$a^2 > (1 - \lambda)(1 + m) \implies \sup \sigma(T) > 1. \quad (13)$$

If μ is itself evaluated, the sharper test for $\lambda < 1$ and $\mu < 1$ is $a^2 > (1 - \lambda)(1 - \mu)$. The formula (12) also handles the other cases. The result follows from the two-dimensional Rayleigh principle and uses no positive-semidefiniteness of T .

4. Nonnested radial compression: the correct coupling

First suppose f is an exact Ritz vector in U , so $r = Tf - \lambda f$. Define

$$h = P_V r, \quad w = Qh. \quad (14)$$

Then

$$\langle Tf, w \rangle = \langle r, Qh \rangle = \langle r, h \rangle = \|h\|^2. \quad (15)$$

The last equality uses orthogonality of P_V ; the preceding one uses $r \perp U$. If $h \neq 0$, then $w \neq 0$: otherwise (15) would equate zero with $\|h\|^2$. For the normalized new direction $v = w / \|w\|$ the coupling is

$$\beta = \langle f, Tv \rangle = \frac{\|h\|^2}{\|w\|} \geq \|w\|. \quad (16)$$

It is generally **not** $\|w\|$. The exact identity $\|h\|^2 = \|Ph\|^2 + \|w\|^2$ explains the difference. This is the appropriate coupling for the proposed arbitrary-radial compressed direction, and the test (13) uses β^2 in place of a^2 .

For a general approximate Ritz pair, let λ be any specified real scalar and define

$$e = PTf - \lambda f, \quad h = P_V(Tf - \lambda f), \quad w = Qh.$$

An exact calculation gives

$$\boxed{\langle Tf, w \rangle = \|h\|^2 - \langle e, Ph \rangle}. \quad (17)$$

Indeed, $\langle Tf, Qh \rangle = \langle Tf - \lambda f, h \rangle - \langle e, h \rangle$, and the first term is $\|h\|^2$ by radial orthogonal projection. Hence if $\|e\| \leq \eta$,

$$|\langle f, T(w / \|w\|) \rangle| \geq \frac{(\|h\|^2 - \eta \|Ph\|)_+}{\|w\|}. \quad (18)$$

For complex spaces take the real part in (17), or choose the phase of the new vector; all proposed sieve arrays are real. If h was computed with $Tf - \lambda f$ instead of the exact QTf , the discrepancy is $P_V e$ and has norm at most η . It must not be silently discarded.

When λ in (17) is merely a numerical Ritz approximation, it is not automatically the true diagonal $\langle f, Tf \rangle$ needed in the two-by-two crossing criterion. Enclose that diagonal separately or include its error in the scalar lower bound.

5. Approximate action, radial compression and projection errors

5.1 Direct action with a rigorous error

Suppose g approximates Tf with $\|g - Tf\| \leq \epsilon$. Define $a = \|Qg\|$ and $v = Qg/a$. Then

$$|\langle f, Tv \rangle| \geq (a - \epsilon)_+. \quad (19)$$

This follows from $\langle Tf, Qg \rangle = \|Qg\|^2 + \langle Tf - g, Qg \rangle$. It requires an error in the true Hilbert mass, not a Euclidean array error without mass weights.

5.2 Approximate radial projection

Let h be as in (17). Suppose a computed h_{tilde} belongs to V and $\|h_{\text{tilde}} - h\| \leq \epsilon$. Set $w_{\text{tilde}} = Qh_{\text{tilde}}$. Because $h_{\text{tilde}} \in V$,

$$\langle Tf, w_{\text{tilde}} \rangle = \|h_{\text{tilde}}\|^2 + \langle h - h_{\text{tilde}}, h_{\text{tilde}} \rangle - \langle e, Ph_{\text{tilde}} \rangle.$$

Consequently,

$$\beta_{\text{lower}} = \frac{(\|h_{\text{tilde}}\|^2 - \epsilon \|h_{\text{tilde}}\| - \eta \|Ph_{\text{tilde}}\|)_+}{\|Qh_{\text{tilde}}\|} \quad (20)$$

is a valid coupling lower bound whenever the denominator is positive. If membership in V is itself only approximate, the displayed identity needs an additional error: simply declaring a numerical vector to be a radial projection is insufficient.

For interval data write $A = \|h_{\text{tilde}}\|^2$, $L = \|Ph_{\text{tilde}}\|^2$ and $W = \|Qh_{\text{tilde}}\|^2 = A - L$. If certified bounds $A \geq A_-$, $A \leq A_+$, $L \leq L_+$, $W \leq W_+$ and $W > 0$ are known, then the conservative computable bound is

$$\beta \geq \frac{(A_- - \epsilon \sqrt{A_+} - \eta \sqrt{L_+})_+}{\sqrt{W_+}}. \quad (21)$$

The upper denominator is intentional. Using a lower bound for W there would reverse the desired inequality. A separate positive lower bound for W certifies that the new vector is outside U and may be normalized.

5.3 A genuinely compressed Galerkin space

If a larger space W contains U and its projection is P_W , then

$$(P_W - P)Tf = P_W(QTf), \quad \|QTf\|^2 = \|(P_W - P)Tf\|^2 + \|(I - P_W)Tf\|^2. \quad (22)$$

A rigorously evaluated nonzero compressed residual is a valid lower witness even without a bound on the omitted tail. A small compressed residual gives no upper bound on the full residual. The nonnested radial space V is not itself such a W ; confusing (22) with (14) produces the incorrect coupling $\|w\|$.

5.4 Approximate projectors

If $\|P_{\text{hat}} - P\| \leq \delta$ and $\|g - Tf\| \leq \epsilon$, then

$$\| (I - P_{\text{hat}})g - QTf \| \leq \epsilon + \delta \| g \| . \quad (23)$$

But $(I - P_{\text{hat}})g$ may not lie in U -perp. Reorthogonalize in the true mass, or account for its leakage, before using an orthonormal two-by-two formula. There is no automatic permission to replace the true Gram projection by Euclidean coefficient subtraction.

6. Inverse–Gram errors and an easier final certificate

Let $u_1, \dots, u_d$ be a basis of U , $G_{ab} = \langle u_a, u_b \rangle$ its positive-definite Gram matrix, and $g_a = \langle u_a, h \rangle$. Then

$$\| Ph \|^2 = g^* G^{-1} g, \quad \| Qh \|^2 = \| h \|^2 - g^* G^{-1} g. \quad (24)$$

An ill-conditioned Gram matrix can turn tiny entry errors into a large error in this difference. The following residual formulation is often preferable. For any proposed projection coefficient vector c_0 , put $b = g - Gc_0$ and $z_0 = h - \sum_a (c_0)_a u_a$. Then

$$\boxed{\| Qh \|^2 = \| z_0 \|^2 - b^* G^{-1} b.} \quad (25)$$

If $\gamma > 0$ is a lower bound for the smallest Euclidean eigenvalue of the exact G ,

$$\| Qh \|^2 \geq \| z_0 \|^2 - \| b \|^2 / \gamma. \quad (26)$$

For an approximate matrix G_{hat} with $\| G - G_{\text{hat}} \| \leq \epsilon_{\text{G}}$, a valid choice is $\gamma = \lambda_{\min}(G_{\text{hat}}) - \epsilon_{\text{G}}$ if positive. If g_{hat} and G_{hat} are used to solve for c_0 , bound

$$\| b \|^2 \leq \| g_{\text{hat}} - G_{\text{hat}} c_0 \|^2 + \epsilon_{\text{G}} + \epsilon_{\text{G}} \| c_0 \|^2. \quad (27)$$

Near-null basis modes cannot be declared harmless simply because a floating-point solver truncates them. Either define U to be exactly the retained smaller span, or certify the conditioning and projection error for the full span being claimed.

For a final gain certificate, a concrete profile and its directly evaluated mixed forms can avoid the entire operator-action error problem. Let A and G be the true numerator and mass matrices on any explicitly specified finite span, possibly the old 77 vectors plus one new radial vector. Suppose entrywise interval errors around $A_{\text{hat}}, G_{\text{hat}}$ are $\epsilon_{ab}^A, \epsilon_{ab}^G$. For a fixed real coefficient vector c , the sufficient test

$$c^T (A_{\text{hat}} - G_{\text{hat}}) c > \sum_{a,b} |c_a c_b| (\epsilon_{ab}^A + \epsilon_{ab}^G) \quad (28)$$

proves the Rayleigh quotient of the actual represented function exceeds one, provided its mass is positive. This test permits negative b_h and does not require either a matrix inverse or $T \geq 0$. Rationalize the chosen coefficients before certification so they are fixed quantities, rather than uncertain optimization outputs.

When working with a two-dimensional orthonormal block and separate certified values $\lambda \geq \lambda_-$, $\mu \geq \mu_-$, $|\text{coupling}| \geq c_-$, its largest eigenvalue is at least

$$\frac{\lambda_- + \mu_- + \sqrt{(\lambda_- - \mu_-)^2 + 4c_-^2}}{2}. \quad (29)$$

The larger eigenvalue is monotone in each diagonal and in the magnitude of the off-diagonal entry. Equation (29), or (28), should be used instead of relying on the sign of a tiny rounded numerical improvement.

7. Function class, fragment labels and arithmetic admissibility

The full action of T_{cap} can introduce dependence on retained largest-fragment labels and on which cap layers survive, even when the initial amplitude depends only on coordinate totals. This occurs because both the face multiplier and the outer fibre domain depend on retained fragments. A totals-only formula for Tf that has discarded those labels is generally not the full action.

The radial projection in (8) deliberately integrates those labels out. It is legitimate if every outer/inner layer and signed contribution is included in its pushforward mass. It can remain in the original class of cap-supported total-coordinate profiles while leaving the 77-dimensional polynomial subspace. A verified $W > 0$ in Section 5 or Section 6 is the precise certificate of leaving that span; a visual impression that the radial function is not a polynomial is not a substitute.

Projection is well-defined in L^2 even when q is very small. In the present fixed compact support, g is continuous and strictly positive, so G has a positive lower bound. If f is bounded, the finite- k integral action is bounded, and the radial coefficient can be viewed as a weighted conditional average of $(Tf - \lambda f)/G$. Thus it too is bounded; division by a small q does not create an analytic singularity in the exact projection. It can nevertheless cause serious numerical cancellation. Thresholding or regularizing small q changes the direction and must be followed by a direct norm/form evaluation.

The cleanest explicit function class is a bounded radial step function on a fixed finite mesh. Such a profile has limiting-null cell boundaries under the same continuous total-size measures used by the source. Additional finite-band approximation can then follow the source's existing method. A generic L^2 eigenvector need not already satisfy the source's bounded-profile and limiting-null-discontinuity hypotheses; approximate it by a fixed admissible profile while retaining any certified margin. Boundedness of T controls this step: for vectors x, y ,

$$| \langle x, Tx \rangle - \langle y, Ty \rangle | \leq \| T \| \| x - y \| (\| x \| + \| y \|). \quad (30)$$

For the actual arithmetic problem, the profile must additionally be projected to O and meet the inner-domain/source hypotheses. A cap gain, even one above one, has not paid those support losses. That obligation remains explicit rather than being included implicitly in the new residual direction.

8. Deliverable and stopping rule

The useful next numerical output is a concrete bounded radial profile, its outside- U mass, its signed mixed forms, and a direct reevaluation of the optimized enlarged-span vector. Label floating-point values exploratory. If a convincing cap gain survives independent contractions, fix the profile and certify the necessary entries or a single explicit quadratic combination with outward arithmetic. Only after that should support restoration be attempted for that new function.

This audit does not recommend another coefficient-only optimization of the same 77-dimensional span, a zeta operator experiment, a PSD replacement of the signed sieve form, or a conclusion based on a tiny compressed residual. The intended new mathematical information is a rigorously measured component of the actual cap-operator action outside the existing span.

9. Exact small-model checks

The accompanying `exact_residual_checks.py` uses only Python's Fraction arithmetic and writes `exact_residual_checks.json`. All checks passed:

- 80 rational five-dimensional examples with a nonuniform mass, a self-adjoint operator having an explicit negative Rayleigh direction, and a radial-analogue subspace not containing U .
- Direct residual, nonnested radial coupling, approximate-Ritz correction, approximate radial-action identity, and the exact inverse-Gram correction (25).
- Three adjoint checks on a 23-point product-measure support, with an explicit nonzero discrepancy when an incorrect conditional-fibre normalization is substituted.
- The conjugated single- g integral in (6), and all five discrete radial projection conditions in (8).

One exact-Ritz example has $||h||^2=7544/406125$, $||Qh||^2=9572/676875$ and $\langle Tf, Qh \rangle = 7544/406125$. Thus the two competing unnormalized coupling formulas are distinguishable in exact arithmetic. The same example has a negative Rayleigh witness $-1/9$. These checks are structural tests of the audit formulas, not estimates of the $k=39$ operator or evidence of an improved cap quotient.

Current report 04: Independent finite marked-space regression for the signed sieve operator

Collection: R6 — full signed prime-sieve operator and residual direction.

Source: [research/prime-gaps/round6/operator-diagnostic/FINITE_MARKED_OPERATOR_AUDIT.md](#).

SHA-256: `486d31e88130a394b42987e8af29f15c4ad902318e65343b92f2f95fce1348c2` . Git blob: `652292d836aa12e6ac71f3e5c5c817f56505b290` . Original bytes: 4828.

Independent finite marked-space regression for the signed sieve operator

This is a structural test of the operator and projection formulas used in Round 6. It is not a calculation of the $k=39$ prime-sieve quotient, and none of its numerical values is transferable to that quotient. All identities below are checked with exact rational arithmetic in `finite_marked_operator_check.py` ; floating eigenvalue displays are explicitly separated in its JSON.

1. Why this model is useful

The actual source Hilbert space is a product of finite fragment measures restricted to a cap domain. Erasing one coordinate integrates against that measure. Its adjoint is a lift, not a normalized conditional expectation. The hybrid face multiplier takes negative as well as positive values. A proposed new radial subspace is not nested with the original polynomial span. A regression that assumes uniform mass, rectangular support, positive operators or nested subspaces would miss the errors relevant here.

The toy coordinate has five atoms `(total, fragment-cap label)` :

Atom	Mass
(0,0)	1/2
(1,0)	2/9
(1,1)	1/9
(2,0)	1/18
(2,1)	1/9

For three coordinates, retain total at most four and require all fragment labels to be zero when total is at least three. There are 38 ordered retained states. The exact mass matrix W is diagonal with the product masses. The face multiplier is 1 on the specified smaller background region, $3/4$ on the larger region, and $-1/4$ on the remaining

backgrounds; rho is 2/5. Thus the toy operator has the same mathematical construction as a signed sum of marginal squares, on a nonrectangular marked domain.

For a retained state x , define

$$(Tf)(x) = \rho \sum_{i=1}^3 m(x_i) \sum_{u: x \oplus_i u \in H} \mu(u) f(x \oplus_i u).$$

The script verifies $WT = T^T W$ exactly. It also produces an explicit coordinate-vector witness with

$$\langle v, Tv \rangle = -\frac{1}{38880} < 0.$$

Consequently a positive-semidefinite assumption is false even in this small model.

2. Nonnested compression and the correct order

Set $G(x) = \prod_i (1 + t_i)^{-1}$, $s = \sum_i t_i$, and let

$$U = G \text{span}\{1, s, s^2, \sum_i t_i^2\}, \quad V = \{Gh(s) : h \text{ arbitrary on } \{0, 1, 2, 3, 4\}\}.$$

Both projections use the exact $\mathbb{W}$ inner product. In particular,

$$P_U = U(U^T W U)^{-1} U^T W.$$

The trial is a specified rational vector in U , with no requirement that it be an exact Ritz vector. Let

$$r = (I - P_U)Tf, \quad h = P_V r, \quad w = (I - P_U)h.$$

The exact checks establish

$$\langle f, Tr \rangle = \|r\|^2, \quad \langle f, Tw \rangle = \|h\|^2, \quad \|w\|^2 = \|h\|^2 - \|P_U h\|^2 > 0.$$

They also establish $\|w\|^2 \leq \|h\|^2 \leq \|r\|^2$. These statements do not assume that U and V are nested. Their projections do not commute in this example. Replacing h by $P_V Tf$ and retaining the same claimed coupling identity gives a nonzero exact error, saved in the JSON. This is a concrete regression for the order of projections, not just a norm tolerance test.

For non-unit f , the normalized two-dimensional block has

$$a = \frac{\langle f, Tf \rangle}{\|f\|^2}, \quad b = \frac{\langle w, Tw \rangle}{\|w\|^2}, \quad c^2 = \frac{\langle f, Tw \rangle^2}{\|f\|^2 \|w\|^2}.$$

Its observed values are approximately $a=0.9862966311$, $b=0.1124698915$ and $c=0.01468793075$. The larger block eigenvalue is approximately 0.9865434471, exceeding a by 0.000246816. These are toy-model displays, not prime-gap evidence. Their purpose is to confirm that the outside-space direction, its normalization and the block formula are all consistent.

3. Product conjugation

For amplitudes $f = G\rho$, let D be multiplication by G . The conjugated operator is $D^{-1}TD$, and its mass matrix is DWD . The script verifies its self-adjointness under that mass and its equivalence to the original action.

The outer coordinate contributes a reciprocal factor $1/g(t_i)$, while the erased integral contains one factor $g(u)$. Replacing this with an average under normalized $g(u)^2\mu(u)$ would define a different operator. The actual numerical implementation must carry the analogous cell factors and normalizers separately.

4. Reproduction and scope

Run `OPENBLAS_NUM_THREADS=1 python3 finite_marked_operator_check.py`. Dependencies are NumPy and SymPy. It writes a JSON receipt and prints the exact rational values. There is no random seed because the model and trial are deterministic. No source file or earlier certificate is changed.

This independent finite check validates algebraic invariants and rejects tempting incorrect formulas. The actual $k=39$ function-space derivation, fragment-cell integration and numerical conditioning require their own proofs and tests. In particular, passing this toy model is not an outward error enclosure for any sieve integral.

Current report 05: Round 6: one full cap–operator residual beyond the 77–dimensional family

Collection: R6 — full signed prime-sieve operator and residual direction.

Source: research/prime-gaps/round6/residual-trial/REPORT.md.

SHA-256: `11a346f72bc5540864367d52740a9d77b1a7866cf5074e562509b94f921d4b35` . **Git blob:** `7adf315b11ee02136a4838a68ca3a400c98dd6fc` . **Original bytes:** 15322.

Round 6: one full cap–operator residual beyond the 77–dimensional family

Outcome. An explicit direction generated by the fixed-geometry k39 signed cap operator increases the fine-grid quotient from `0.99439639919` to `0.99446782094` , about 71.42 ppm. This direction is not another polynomial degree or an eigenvector of the old 77-by-77 matrix. About 98.78% of its radial residual mass survives true mass projection away from the whole 77-dimensional space. The new value is still below one by about 5,532 ppm. No smaller prime gap, DHL(39,2), full support-restored certificate, or global variational bound is proved.

The bounded experiment consists of one coarse construction, one fine construction, and two numerical consistency checks. There is no iterative Krylov expansion or parameter scan in this report. The independent ordinary derivation below is exact for the stated finite cap model; all reported numerical values remain floating-point diagnostics without outward error enclosures.

1. The operator and its measure

The geometry, $k=39$, product profile g , exceptional constant $K=.34$, hybrid parameters, and 98,304-cell fine grid are those of the original Round 4 experiment. The official source pin is `61340d0b74163003b32756bb16e91d9209a5e330` ; the coefficient-source SHA256 is `7f71bdefcfe3bb5ca76a143929b3cb3f4156c21dc483253cda3077420f1e5de4` .

Each coordinate consists of its total t and a fragment mark. For a fragment cap c , write $d_c(j)$ for the cell-average Dickman survival probability at total cell j . An exact finite marked model assigns mass

$$\text{nu_h}(j, \text{ell}) = h \left[d_{\{c_{\text{ell}}\}}(j) - d_{\{c_{(\text{ell}-1)}\}}(j) \right]$$

to the mark layer ell . This definition retains fragment-cap information rather than pretending that a coordinate total exceeding c is forbidden. The implemented d_c values use eight-point cell quadrature and floating Dickman evaluations; exact d_c would define the corresponding exact finite model.

Let H_O be the fixed outer cap support, E_i the erased-coordinate integral against nu_h , and


```

m = 1_H0 + (a+b) 1_(H1\H0) + b 1_(Hfull\H1),
a = .99998^2 - .99998*.008,
b = (1-.99998/.008)*(1-.99998)*.34.

```

On the real Hilbert space $L_2(H_O, \nu_h^k)$, the cap operator is

```

T = rho_star * 1_H0 * sum_i E_i^* m E_i * 1_H0.

```

The adjoint E_i^* is an unweighted lift followed by outer restriction. There is no conditional fibre normalization. The signed coefficient b is negative, so T must not be assumed positive semidefinite. This is the **cap operator**, not the operator with all true source predicates and support-restoration losses inserted.

Write $G(t) = \text{product}_i g(t_i)$. Amplitudes $f = G p$ have mass measure $G^2 \nu_h^k$. In this representation the action contains $1/g(t_i)$ outside the erased integral and a *single* g inside that integral. Replacing this with a g -squared conditional average would give the wrong adjoint and a different operator.

2. A compressed direction from the full operator

Let U be the original 77-dimensional amplitude space and f in U the stored optimized vector. Let P_U be the orthogonal projection for the actual mass form, not Euclidean coefficient projection. Let V consist of all amplitudes $G h(s)$, with arbitrary radial cell profile $h(s)$, supported on a fixed selected set of radial cells. V is substantially larger than the degree-six radial polynomial subspace.

Define

```

r = (I-P_U) T f,
h = P_V r,
w = (I-P_U) h.

```

The order of these projections matters: $P_V U$ need not be contained in U . In code $P_U T f$ is computed as $G \text{mat}^{-1} B \text{mat } c$; it is not replaced by a scalar multiple of f . This avoids relying on exact Ritz accuracy.

Self-adjointness and orthogonality give the ordinary identity

```

<f, Tw> = <Tf, w> = <r, h> = ||h||^2.

```

Moreover $||w||^2 = ||h||^2 - ||P_U h||^2$. Consequently nonzero h necessarily gives a genuine outside- U direction w and positive coupling, regardless of the sign of T . For unit f and $v = w / ||w||$, the off-diagonal entry is

```

beta = ||h||^2 / ||w||,

```

which is generally **not** $||w||$. This is a projected residual construction using a nonnested radial subspace; its identities are standard Hilbert-space facts, not a claimed new theorem about prime gaps.

The main computational point is that P_V can be obtained by one-dimensional adjoint convolutions while still integrating the full cap marks. It therefore exposes an outside-77 component without constructing a 39-dimensional tensor grid or pretending that another eigenvector of the same 77-by-77 matrix is new information.

3. Exact finite-array formula for the radial adjoint

The numerical tilt τ is only a change of normalization. Set

$$\begin{aligned} Z_{\tau} &= \sum_j g_j^2 \exp(-\tau t_j), \\ p_c(j) &= g_j^2 \exp(-\tau t_j) d_c(j) / Z_{\tau}. \end{aligned}$$

Let $M^c_{\ell}(\eta)(s)$ denote the same polynomial background moments used by the original cap engine. The product-power identities are expanded by set partitions, and moments are computed by one-dimensional convolutions.

Background cap layer ℓ uses the difference between moments at consecutive caps.

For the input polynomial f , the erased-coordinate amplitude on layer ℓ is

$$\begin{aligned} E_{\ell} f(Y) &= h_{\text{mesh}} * \text{product}_{(j \neq i)} g(t_j) \\ &\quad * \sum_{\eta} A_{\ell}(\eta)(s) P_{\eta}(Y). \end{aligned}$$

The existing affine arrays A include all allowed outer radial shells **before** the background contraction. Define

$$\begin{aligned} B_{\ell}(s) &= \sum_{\eta} A_{\ell}(\eta)(s) \\ &\quad [M^c_{\ell}(s) - M^c_{\ell-1}(s)]. \end{aligned}$$

For an output total r whose outer fragment cap is $c_O(r)$, the radial numerator density is

$$\begin{aligned} b_f(r) &= \rho_{\text{star}} * k h_{\text{mesh}} / Z_{\tau} \\ &\quad * \sum_{\ell: c_{\ell} \leq c_O(r)} \sum_{(j+s=r)} g_j d_{c_O(r)}(j) m_{\ell}(s) B_{\ell}(s) \\ &\quad \exp(\tau(s+(k-1)/2) h_{\text{mesh}}). \end{aligned}$$

The radial mass density is

$$q(r) = \exp(\tau(r+k/2) h_{\text{mesh}}) M^c_{c_O(r)}(k, \text{empty})(r).$$

Let $D_i(r)$ be the corresponding mass cross-density with basis amplitude u_i , so that $\langle G, u_i \rangle = \sum_r h(r) D_i(r)$.

If p is the coefficient vector of $P_U T f$, the desired radial residual profile is

$$h(r) = [b_f(r) - \sum_i p_i D_i(r)] / q(r)$$

on the selected radial cells, and zero elsewhere. The same notation h for a profile and $G h$ for its amplitude is used only in these array formulas.

The implementation selects cells where $q(r) > 10^{-9} \max(q)$. This defines V explicitly; it does not assume that the residual on excluded cells is zero. The saved active index set is frozen and defines the exact subspace V_{active} . Restricting V in this way preserves the projection identity in exact arithmetic and avoids division by unreliable tiny numerical masses. Small excluded q -mass is **not a bound on omitted residual energy**; no such bound is used here. The fine run excludes approximately 1.325×10^{-10} of the positive radial mass. The negative radial mass generated by FFT rounding is only about 3.6×10^{-25} of total positive mass and is outside the selected region. These are observed diagnostics, not certified rounding bounds.

The formula was independently checked against the source normalization by the prime186 agent: $k, h_{\text{mesh}}/Z$, the erased single- g factor, layer differences, and placement of the exponential agree. Root's independent marked-state toy and ordinary operator proof are separate artifacts; they should accompany this report when the larger package is assembled.

4. The mixed forms and the actual new space

The new direction is not assessed solely through an action-norm estimate. The program independently evaluates all 77 mixed mass entries $\langle u_i, h \rangle$, all 77 signed numerator entries $\langle u_i, Th \rangle$, and $\langle h, h \rangle$, $\langle h, Th \rangle$. A separate sampled-profile evaluator then directly evaluates the final polynomial-plus-radial amplitude.

Writing $g_i = \langle u_i, h \rangle$, G_{mat} for the old mass Gram, and $p = G_{\text{mat}}^{-1} g$, the true outside-space norm is

$$\text{gamma} = \langle h, h \rangle - g^T G_{\text{mat}}^{-1} g.$$

The numerator entries for $w = h - U p$ are formed by the same subtraction in the numerator form. The implementation whitens the original 77-dimensional mass Gram, appends $w/\sqrt{\text{gamma}}$, and solves the resulting 78-dimensional symmetric pencil. The raw matrices, radial profile, projection coefficients and new vector are retained, so this is a reviewable explicit finite-dimensional extension.

5. Main numerical evidence

At the original fine grid, with $\tau=20$ and radial mass cutoff 10^{-9} :

quantity	observed value
original 77-space quotient	.9943963991909279
compressed radial norm squared	
outside-77 norm squared	
surviving norm-squared fraction	.9877820901062194
directly evaluated $\langle f, Tw \rangle$	6.670688584662753e-5
relative coupling-identity error	-7.39e-10
mixed-adjoint absolute discrepancy	-4.83e-14
full 78-space matrix quotient	.9944678209367751
final direct profile quotient	.9944678209006830
matrix/direct difference	3.61e-11

The original matrix quotient differs slightly from the old directly evaluated value **.9943963993644909**; that previously recorded discrepancy is numerical and is not counted as new improvement. The full 78-space improvement relative to its matrix baseline is about **71.4217 ppm**.

The actual two-dimensional span of f and normalized w has matrix

```
[ .9943963991909279    .008217784708256407 ]
[ .008217784708256407 .043583189070450945 ].
```

Its largest eigenvalue is `.9944674193880856`, giving **71.0202 ppm** improvement. Reoptimizing all 78 coefficients adds approximately **0.40155 ppm** beyond that simple two-direction step. Thus the observed gain genuinely comes from the new direction rather than merely from re-solving the old polynomial coefficients.

The computed mass norm of `P_U Tf-lambda f` is approximately $1.64\text{e-}10$. This approximate Ritz defect is recorded for context; it is not used to justify the projection identity, because the full mass projection is used directly.

For this particular two-dimensional plane, crossing one requires

$$\text{beta}^2 > (1-a)(1-b),$$

since both diagonal entries are below one. The observed coupling squared is only about 1.26% of that plane's crossing threshold. This explains why the positive new direction does not come close to proving the desired k39 criterion. It says nothing about the uncomputed full residual outside the chosen radial compression.

6. Bounded consistency checks

grid	cutoff	tilt	full 78 matrix quotient	direct quotient
16,384	1e-9	20	.9934506692779733	.9934506693945010
98,304	1e-9	20	.9944678209367751	.9944678209006830
98,304	1e-9	25	.9944678209367753	.9944678210538511
98,304	1e-8	20	.9944678032713334	.9944678033480433

The tilt-20 and tilt-25 matrix results differ by $2.2\text{e-}16$. Their direct evaluations differ by approximately $1.53\text{e-}10$, consistent with the contraction sensitivity already present in the ill-conditioned polynomial Gram. Changing the radial cutoff by a factor of ten changes the result by only about .0177 ppm. The coarse and fine new-direction gains are respectively about 71.3164 and 71.4217 ppm.

The old fine scaled-Gram condition is about $2.28\text{e}10$, so small pencil residuals alone are not rounding certificates. All operations here use actual float64 on this host. The main fine construction took about 10.6 seconds before concurrent consistency checks. No speed comparison or scaling theorem is claimed.

7. Reproduction and files

From this directory:

```
OPENBLAS_NUM_THREADS=1 python3 radial_residual.py --intervals 16384
OPENBLAS_NUM_THREADS=1 python3 radial_residual.py --intervals 98304
OPENBLAS_NUM_THREADS=1 python3 radial_residual.py --intervals 98304 --tilt 25
OPENBLAS_NUM_THREADS=1 python3 radial_residual.py --intervals 98304 --density-cutoff 1e-8
```

```
OPENBLAS_NUM_THREADS=1 python3 audit_outputs.py
OPENBLAS_NUM_THREADS=1 python3 validate_outputs.py
```

`radial_residual.py` contains the radial adjoint, full mixed-form construction, mass projection, 78-space solve and direct sampled-profile evaluation. `cap_trial.py` is the copied fixed-geometry cap engine; the upstream clone is never imported as executable code or modified. `audit_outputs.py` derives and saves the actual two-by-two matrices and Ritz-defect diagnostics from retained arrays without performing new integrations.

Inputs include the pinned literal coefficient source, and the already completed Round 4 files

`ritz_k39_n16384.json/.npz` and `ritz_k39_n98304.json/.npz`. Their paths are resolved relative to the common staging BASE. The exported archive must preserve this layout or adjust the explicit paths. `PRIME186_SOURCE` can override the coefficient-source path, and `PRIME186_TRIAL_ROOT` can override the directory containing the four old trial input files. These path overrides were added after the numerical runs and do not change their recorded values.

Each run has JSON diagnostics and NPZ arrays. The NPZ retains q , all D_i , b_f , the new h , P_U Tf, mixed forms, projection coefficients, final polynomial coefficients and the 78-space matrix. JSON includes the final radial multiplier beta and direct evaluation. `projection_audit.json` contains the separate two-dimensional evidence. `validate_outputs.py` checks the mass-projection inequalities, coupling and adjoint identities, two-dimensional gain, normalized candidate mass, direct-form agreement, and tilt consistency at explicit numerical tolerances; `validation.json` records the passing regressions. These regression thresholds are not outward arithmetic enclosures. The full final amplitude is reproducible as the stored polynomial coefficient vector plus beta times the stored arbitrary radial profile, multiplied by the unchanged product g and outer support.

8. What remains unproved and what this establishes

The ordinary projection identities and finite-array adjoint derivation explain why this is a legitimate outside-space search. The numerical evidence supports a real positive direction in the cap model. It is not an outward certificate even for the new finite profile, and the new profile has not been supplied with the full source-dependent support repairs. No support-restored operator norm, T-squared moment, or global extremizer has been computed.

A future exact certificate could fix and rationalize the new profile and enclose only its mixed mass and numerator entries. Positivity of the resulting concrete quadratic expression would avoid relying on an unproved action-norm error estimate. Here, however, the quotient remains below one, so that certification would not yield the desired prime-gap result. The next substantive research question is whether a less restrictive compression captures a much larger fraction of the full cap residual. This report does not answer that question and does not continue the iteration automatically.

Compact public witness and historical program metadata

The public-sized `*_compact.npz` exports retain every candidate/projection array and omit only the regenerable 77-by- N cache D . `compaction_manifest.json` records the original full hashes, compact hashes, and exact raw-byte identity of every retained array. Full NPZs remain local. See `PROVENANCE.md` for the regeneration instructions and the precise distinction between earlier outputs without embedded program hashes and later metadata additions. No earlier run is retrospectively described as containing a hash it did not record.

Current report 06: Provenance and compact witness export

Collection: R6 — full signed prime-sieve operator and residual direction.

Source: `research/prime-gaps/round6/residual-trial/PROVENANCE.md`.

SHA-256: `4a3a8f57b27e88053b12030487ed84fd939c3e32541ead1ecef047949ace04ea` . **Git blob:**
`4724dc6dd0afca4c68d0f8b2d9ea7d2e917bec8d` . **Original bytes:** 2842.

Provenance and compact witness export

The four numerical integrations are the coarse run at $N=16384$, the primary fine run at $N=98304$, a fine run with density cutoff $1e-8$, and a fine run with tilt 25. Each full NPZ and JSON is retained locally. No output numbers were replaced by estimates or by the later independent replay.

The first coarse integration preceded two small program revisions: explicit conversion between the stored matrix's tilt normalization and the current integration normalization, and extra hash fields in JSON. At tilt 20, the conversion is exactly one, so the first coarse calculation is the same mathematical computation. The primary fine run included the normalization conversion but preceded the extra embedded program/engine/input hash fields. The two consistency runs include those hash fields. It would be inaccurate to claim the earlier files contained an embedded program hash; instead, the archive manifest records their actual file hashes and the current program hash separately. The old upstream source SHA and old trial JSON SHA were present in every numerical output.

`audit_outputs.py` subsequently computed the actual 2-by-2 subspace matrices and projection-defect diagnostics from retained arrays without new integration. `projection_audit.json` pins each input JSON. It does not alter the original integration evidence.

For public export, `compact_witness.py` removes only the 77-by- N radial mass cross-density cache named `D` from each compressed NPZ. Every other array, including the actual `h`, the output polynomial coefficients, projection data, `q`, `b_f`, active mask, radial coordinates and 78-by-78 pencil, is copied without changing its dtype, shape or raw bytes. `compaction_manifest.json` records full and compact file SHA256s, sizes, the removed `D` cache's shape/dtype/raw-byte SHA, and each retained array's raw-byte SHA. The script checks exact byte equality after reopening each compact archive.

`D` is reproducible by the `radial_mass` function; the full run commands in `REPORT.md` regenerate the original full model data and all candidate forms. Full NPZs remain in this local staging directory. Public consumers can use the compact witness for the explicit amplitude, projection audit, and two-dimensional forms; workflows that require `D` must regenerate it or obtain the pinned full archive. Omission of this cache is an explicit size reduction, not loss or substitution of the candidate vector.

`PRIME186_TRIAL_ROOT` was added to the main and post-audit scripts after the numerical runs, solely to select the old-trial input directory in a public package. The post-audit also accepts the compact witness when the full archive is absent. These path/input-selection changes are pinned by the final archive manifest; they do not retroactively change

the stored run's program provenance.

Current report 07: Independent review of the frozen radial profile outside the old span

Collection: R6 — full signed prime-sieve operator and residual direction.

Source: [research/prime-gaps/round6/operator-diagnostic/OUTSIDE_SPAN_INDEPENDENT_REVIEW.md](#).

SHA-256: `a1f8b3fc86c4232f6f0e6f17b930935d927671466b40274c02ace7893432ad0e` . **Git blob:** `dc3bc5fdd61798b278e99668f9d2c1631351ac28` . **Original bytes:** 4931.

Independent review of the frozen radial profile outside the old span

Verdict: the outside-span argument is valid. This note records the independent review already completed by the Yau-flow audit agent. It certifies a new direction in the true cap Hilbert space, not an enclosure of a numerical Rayleigh gain.

Evidence independently checked

The reviewed witness is `../residual-trial/radial_residual_n98304_cut1e-09_tilt20_compact.npz`. Its SHA256 is `b283cab182b0b32091f24ac898def31cc263fa6af1a4540b30721e8122b80c77`. A read-only hash comparison confirmed that this public checkpoint is byte-identical to the staged witness used in the independent review.

The `h.npy` member was decoded independently with Python's standard-library ZIP, NPY-header and binary-structure handling, rather than through the certificate's NumPy loading path. The stored array is little-endian binary64. Its entries at indices 0 through 12 are exactly zero, and its first nonzero entry is at index 18422, with exact dyadic value

$$h_{18422} = -\frac{6264072493613325}{4611686018427387904} \equiv 0.$$

This interprets the frozen stored value as an exact rational definition of the new step profile. It does not assert that the floating computation producing that value exactly evaluated an operator residual.

The 11 official coefficient signatures were separately read from the preserved source using its literal syntax tree. Their maximum exponent sum is 6. The old basis combines these signatures with radial powers of degrees 0 through 6. On the product cell with indices

$$(j_1, j_2, \dots, j_{39}) = (r, 0, \dots, 0),$$

division by the common strictly positive product factor G leaves a polynomial in r of degree at most 12. The midpoint representatives are affine in r , the radial shift is affine in r , and each power-sum signature has degree at most its exponent sum. This degree assertion is for the actual frozen midpoint-step basis.

Whole-cell support and positive measure

The argument uses whole product cells, not isolated representative points, which would not by themselves establish an L^2 assertion. The mesh is

$$\Delta = \frac{2742997}{258046918656} > 0.$$

For the largest selected index, the maximum coordinate-total endpoint and maximum whole-cell total are respectively

$$18423\Delta = 0.1958335096353804 \dots, \quad (18422 + 39)\Delta = 0.19623744348796382 \dots$$

These were checked with exact rational arithmetic. The whole-cell total is strictly below the first outer-shell boundary

$$\frac{653622010000}{689056987511} = 0.9485746779246842 \dots$$

Every selected coordinate cell is also strictly below the retained fragment cap. The companion certificate uses the sufficient conservative cap

$$\frac{41328816845772771}{110249118001760000},$$

which exceeds 18423Δ . The independently reconstructed final outer cap was larger still, approximately 0.4951438701 ; using the smaller common cap is therefore harmless.

All 14 selected product cells consequently lie inside the same first outer shell and satisfy its cap automatically. On these small coordinate-total intervals, the unscaled fragment measure has total-size density one. Each selected product cell has mass Δ^{39} , which is strictly positive. The common factor G is strictly positive on them.

Proof of nonmembership

Suppose the frozen cap-supported radial profile $G h$ belongs to the old span U in L^2 . The old profiles and the new profile are constant after the prescribed midpoint evaluation on each selected product cell. Equality almost everywhere therefore forces equality of their values on each such positive-mass cell.

After dividing by G , an old-span representation would give a polynomial $p(r)$ of degree at most 12 satisfying

$$p(0) = p(1) = \dots = p(12) = 0, \quad p(18422) = h_{18422} \not\equiv 0.$$

Thirteen distinct roots force p to vanish identically, a contradiction. Hence $G h$ is outside U . Equivalently, its orthogonal projection onto $U^\perp$ has strictly positive true Hilbert norm. This conclusion is independent of inverse-Gram conditioning and of the accuracy of the computed projection norm.

Scope of this review

The implementation `certify_outside_span.py` and its `outside_span_certificate.json` receipt were read when preparing this record. They additionally contain a modular rank-77 witness for the old basis and an augmented-rank-78 conclusion. This reviewer did not independently rerun that modular rank computation. The 13-zero argument above does not need it; dimension exactly 78 additionally uses independence of the 77 old basis vectors.

No fine integral, optimization, or existing certificate was rerun or modified for this note. The result proves nonmembership for the frozen dyadic step profile. It supplies no quantitative lower bound for the outside-space norm, no outward enclosure of the new quotient or its gain, and no certification of the fully restored arithmetic support. Those remain separate obligations.

Current report 08: Round 7: two explicit actual–zeta targets for Dyson–Montgomery

Collection: R7 — actual-zeta targets, arithmetic mark, and flow obstruction.

Source: research/reports/dyson_round7.md.

SHA-256: [30ae0f293454aabdf22dbbc9908f1551c1c5a0488d3e3b9ac0f1105459569b33](#) . Git blob: [63ad174f355fab585a63895eb4d592f5e8af10ad](#) . Original bytes: 8767.

Round 7: two explicit actual–zeta targets for Dyson–Montgomery

The user's direction is now the main lane: actual Riemann-zeta pair correlations and the Alternative Hypothesis, with random matrices and heat flow used to find precise tests. The prime-gap parameter search is paused. This round has two rigorous reductions to explicit arithmetic inequalities, a new arithmetic resonator with a negative test result, and a forward-flow obstruction. **Neither required new zeta inequality is proved.**

1. The most concrete target: two logarithmic–derivative mean squares

For fixed $c > 0$ write

$$I_T(c) = \int_0^T \left| \frac{\zeta}{\zeta} \left(\frac{1}{2} + \frac{c}{\log T} + it \right) \right|^2 dt, \quad W_T = \frac{2[\sinh(2)/T(1) - \sinh(1)/T(1/2)]}{T \log^2 T}.$$

The [complete reduction](#) gives, under RH and the precise AH-Pairs formulation in the cited primary paper,

$$W_T \rightarrow W_{\text{AH}}, \quad 0.06239 < W_{\text{AH}} < 0.06240.$$

The sine-kernel prediction is 0.0822714431214773... Therefore a proof under RH that

$$\liminf_{T \rightarrow \infty} W_T \geq 1/16$$

would already refute AH-Pairs under RH. The easier-to-remember 0.07 target is sufficient but unnecessarily strong as an acceptance threshold. Every lower limit strictly exceeding W_{AH} would suffice.

The construction fixes a genuine issue with comparing only ACUE to CUE. General AH-Pairs leaves a bounded near-diagonal parameter $P_0(T)$, which need not converge. For a Poisson smoothing width $b/(4\pi)$, that freedom contributes exactly $2(P_0(T)-1)/\sinh(b)$ to the limiting variance formula. The displayed two-scale combination cancels it. The argument does not assume simple zeros or replace the full AH class by one example.

The proof explicitly controls the noncompact pair-kernel tails, the removed low-zero interval, finite-height endpoints, the Gamma factor, and the holomorphic-square term required to pass from squared real part to squared modulus. The truncations avoid half-lattice boundary atoms. [One independent review](#) checks these conversions; a [second](#)

[review](#) also reruns the scalar and finite-model checks. The constant enclosure is exact rational arithmetic. These facts establish the reduction, not the missing lower bound for W_T .

2. A compact Fourier target and its exact prime covariance

The [primary-source frontier report](#) chooses one nonnegative smooth bump φ with integral one, supported on $[6/5, 7/5]$ and symmetric about $13/10$. For Montgomery's normalized form factor F_T , the two predictions are

Statistic	RH + AH-Pairs	Montgomery sine-kernel target
Integral $\varphi(\alpha) F_T(\alpha) d\alpha$	$7/10$	1
Centered prime-covariance remainder E_T	$-3/5$	$-3/10$

The AH conclusion follows from half-lattice support and the known low band: the limiting pair Fourier distribution is 2-periodic, so its density on $(1, 2)$ must be $2-\alpha$. The chosen test avoids the integer atoms and all dependence on the unknown near-diagonal mass. Uniform pair-tail bounds justify this statement without assuming a full limiting process exists.

The report then gives an exact prime kernel for E_T . It retains both von Mangoldt sums and the continuous mean from the pole. An independently checked finite expansion demonstrates why omitting that mean changes the problem. A proof that $\liminf E_T > -3/5$ would suffice to refute AH-Pairs under RH; the stronger limit $-3/10$ would prove this one smoothed Montgomery prediction. Both remain open here.

These are alternative precise targets. Neither comes from reinterpreting the distribution exponent of the 186 prime-gap proof. The missing quantity is a signed two-prime covariance, at the accuracy displayed in the report.

3. A genuinely different arithmetic resonator was tried and did not cross

The first proposed S2/S3 polynomial extension had already been tested in the earlier archive. The agent identified that duplication and instead used the sharp arithmetic mark

$$C_L(n) = 1_{\{P^+(n) > \sqrt{L}\}}, \quad n \leq L.$$

There is an exact unique-large-prime decomposition $n = pm$ with $p > \sqrt{L}$ and $m < p$. It gives the new marked moments and insertion rules directly from integers. The [derivation](#) and [independent audit](#) explain the small-prime truncation, threshold boundary, short background and surviving same-prime term.

A complete 30-dimensional trial was optimized at one fixed $\text{ell} = 27/25$. Its limiting half-gap margin is numerically -0.01465492379421 , a gain of only about 1.429×10^{-8} over its matched 20-dimensional baseline. It remains slightly worse than the older 48-feature best trial, which the report states explicitly. Three split quadrature orders agree, and a frozen rational vector has negative directly evaluated integer-operator margins through $L = 10^6$. These are numerical checks, not interval enclosures or actual zero samples. This particular feature does not justify another coefficient sweep.

The [full report and data](#) retain the fixed coefficients, full matrices and failed trial. The failure does not prove that every discontinuous prime-factor feature or the full resonance method is incapable of crossing the threshold.

4. Forward heat flow: a useful comparison, with two missing estimates

The [forward-flow report](#) proves a contraction estimate for ordered real repulsive systems with a decreasing common external field. Its constant does not deteriorate as an internal gap becomes small. However, a remote-field approximation valid at the central particles need not be valid at the retained block's boundary; a specific boundary-propagation integral remains to be controlled for actual H_t .

Even perfect deterministic localization is insufficient for GUE. The exact family

$$P_t(z) = z^{2M} - 2 \cos(\pi/4) e^{-M^2 t} z^M + 1$$

starts on the ACUE half-grid up to rotation, has bounded counting discrepancy, and keeps all normalized gaps at least $1/2$ throughout forward flow. It approaches the unit clock. This is a counterexample to insufficient dynamical hypotheses, not a model satisfying the full arithmetic explicit formula.

The missing stochastic term is quantitatively visible: for a protected trace frequency $m=N/2$, its expected microscopic generator contribution at CUE is exactly π^2 . Moreover, the entire protected trace filtration remains matched between ACUE and CUE under full DBM. Thus unchanged low moments cannot justify removing stochastic smoothing. The report states the exact Duhamel and boundary estimates a genuine zeta comparison would need.

5. What the recent 0.6725 result contributes

The source audit identifies $0.6725007\cdots$ as an unconditional lower proportion of zeros that are both simple and on the critical line, with the separate distinct-zero consequence stated in the primary papers. It does not evaluate the out-of-band covariance used here. The checked new proof still uses its stated support-one input. The report records exact source URLs, dates and hashes; it does not infer that a percentage-of-zeros theorem supplies a Montgomery plateau.

6. Verification and the next decision

The focused folders preserve source provenance, ordinary proofs, independent reviews, exact scalar certificates, complete numerical witnesses and execution logs. Third-party primary PDFs stay in the local reference archive; their URLs and hashes are public. The independent integration receipt is under [research/logs/round7-integration/](#).

The next mathematical step is to prove a nontrivial lower bound for the signed mean square W_T or the centered prime covariance E_T using additional arithmetic structure. The two-scale spectral weight is $\sinh(2)\exp(-2|u|) - \sinh(1)\exp(-|u|)$; it changes sign at $|u| = \log(2 \cosh 1)$. Dropping the unknown high-frequency contribution therefore gives no lower bound. More decisively, the actual stationary ACUE comparison process matches the known low band and satisfies point-process positivity while attaining $W_{AH} < 1/16$. Those inputs alone cannot prove the desired inequality. A bounded continuation examines exactly where additional arithmetic information enters a centered-psi representation.

Numerical model separation and the present reductions do not estimate either required arithmetic quantity. A complete GUE theorem, a new prime-gap sweep, and proof-assistant formalization of an unclosed argument are postponed. The long research goal remains active.

Current report 09: Two Poisson scales remove the AH near-diagonal parameter

Collection: R7 — actual-zeta targets, arithmetic mark, and flow obstruction.

Source: [research/dyson/round7/poisson-resolvent/TWO_SCALE_ZETA_TARGET.md](#).

SHA-256: [b275eedbba276097e3729bd0f75dff4108c11ab7c4f038093b5722e0b3fc7f51](#) . **Git blob:** [ecca372d887258f7073d9eefa1777b9172660ab0](#) . **Original bytes:** 14073.

Two Poisson scales remove the AH near-diagonal parameter

Date: 2026-09-05. **Status:** written reduction and exact constant certificates; [independent ordinary-proof review completed](#). The actual arithmetic inequality stated below is **not proved**. No novelty claim is made for the standard pair-correlation/logarithmic-derivative correspondence.

The useful target is now one explicit inequality for actual zeta mean squares. Put

$$I_T(c) = \int_0^T \left| \frac{\zeta}{\bar{\zeta}} \left(\frac{1}{2} + \frac{c}{\log T} + it \right) \right|^2 dt, \quad W_T = \frac{2}{T \log^2 T} (\sinh(2)/I_T(1) - \sinh(1)/I_T(1/2)). \quad (1)$$

Under RH plus the precise AH-Pairs formulation discussed below, the reduction gives

$$W_T \rightarrow W_{\text{AH}}, \quad 0.06239 < W_{\text{AH}} < 0.06240. \quad (2)$$

The sine-kernel prediction is instead

$$W_{\text{GUE}} = 0.0822714431214773 \dots \quad (3)$$

Consequently, an actual proof under RH of

$$\boxed{\liminf_{T \rightarrow \infty} W_T \geq \frac{7}{100}} \quad (4)$$

would refute AH-Pairs under RH. This asks for a one-sided inequality for one signed combination, not the full pair-correlation conjecture. It remains a substantial open arithmetic task. The reduction does not make (4) follow from RH or the known Fourier band.

The number 0.07 is a convenient stronger target, not a required threshold. The same exact enclosure already shows $W_{\text{AH}} < 0.06240 < 1/16$. Thus the weaker bound $\liminf W_T \geq 1/16$ also suffices, with a separation exceeding 0.00010. Any proved lower limit strictly exceeding W_{AH} is acceptable; neither rational lower bound has been established here.

1. Definitions and the existing source input

Use Fourier transform $\hat{f}(\alpha) = \int f(x) e^{-2\pi i \alpha x} dx$, and put $L = \log T/(2\pi)$. Multiplicities are counted throughout. Let

$$\mu_T = \frac{1}{TL} \sum_{0 < \gamma, \gamma' \leq T} \delta_{L(\gamma - \gamma')}.$$

The precise AH-Pairs assumption is (AH0) in Goldston – Lee – Schettler – Suriajaya, *Pair Correlation Conjecture II: The Alternative Hypothesis*, 2507.06823: all bounded normalized pair differences in the stated high interval approach half integers with the specified decreasing error. We use its RH-dependent equations (1.14) – (1.15), and the cumulative pair bound (1.12). These are distinct from that paper's unconditional results with additional AH-Weak Density assumptions.

For each fixed half integer, the pair masses have the form, up to $o(1)$,

$$p_{k/2}(T) = \begin{cases} p_0(T), & k = 0, \\ p_0(T) - 1/2, & k \equiv 0 \text{ even}, \\ 3/2 - p_0(T) - 2/(\pi^2 k^2), & k \text{ odd}, \end{cases} \quad (5)$$

with bounded $1 + \alpha(1) \leq p_0(T) \leq 3/2 - 2/\pi^2 + \alpha(1)$. We do not assume $p_0(T)$ converges. The symbol here abbreviates the source's finite- T $p_0(T)$, not an already existing limit.

The base $p_0=1$ measure is the pair measure of the randomly translated half-lattice sine process. Relative to it, (5) adds

$$(p_0(T) - 1) \sum_{k \in \mathbb{Z}} (-1)^k \delta_{k/2}. \quad (6)$$

Formula (6) includes the diagonal. Dropping it would give the wrong nuisance term.

2. Bounded smoothing and exact model formulas

For $a > 0$ define $P_a(x) = a/[\pi(a^2 + x^2)]$. Its Fourier transform is $e^{-2\pi a |\alpha|}$. If a stationary unit-intensity process has centered structure-factor measure S , its smoothed-density variance is

$$V(a) = \int e^{-4\pi a |\alpha|} dS(\alpha). \quad (7)$$

Write $b=4\pi a$ and abbreviate this variance by $V(b)$. For the sine process, $dS = \min(|\alpha|, 1) d\alpha$, giving

$$V_{\text{sine}}(b) = \frac{2(1 - e^{-b})}{b^2}. \quad (8)$$

The centered spectral measure for the half-lattice model is

$$dS_A(\alpha) = \text{dist}(\alpha, 2\mathbb{Z}) d\alpha + \sum_{m \in \mathbb{Z} \setminus \{0\}} \delta_{2m}. \quad (9)$$

Here the distance ranges from zero to one. To derive (9), its pair measure is $\delta_0 + \frac{1}{2} \sum_{k \neq 0} (1 - \text{sinc}^2(k/2)) \delta_{k/2}$, where $\text{sinc}(x) = \sin(\pi x)/(\pi x)$. Poisson summation transforms the comb into atoms at even integers, and the sampled sinc-square into the sum of triangular functions centered at even integers. Subtracting the mean-density contribution removes only the atom at frequency zero. Thus

$$V_A(b) = \frac{2 \tanh(b/2)}{b^2} + \frac{2}{e^{2b} - 1}. \quad (10)$$

The even-frequency atoms in (9) are essential; the triangular density alone gives a wrong answer. Direct integration of one triangular period and a geometric series proves (10).

For every $b > 0$,

$$\Delta(b) := V_{\text{sine}}(b) - V_A(b) = \frac{2e^{-2b}(4\sinh^2(b/2) - b^2)}{b^2(1 - e^{-2b})} > 0, \quad (11)$$

because $2\sinh(b/2) > b$. The signal is $b/12 + O(b^2)$ as b decreases to zero, and $2e^{-b}/b^2(1 + o(1))$ as b increases to infinity. Very large smoothing therefore requires exponential accuracy relative to its algebraic main term. The leading variance alone cannot separate the models.

3. Eliminate the near-diagonal freedom rather than assume simplicity

The Fourier transform of the alternating comb in (6) is $2 \sum_m \delta_{2m+1}$. Its contribution to (7) is

$$\frac{2(\rho_0(T) - 1)}{\sinh b}. \quad (12)$$

Thus a single variance does not uniformly distinguish the full AH-Pairs class. The correct combination is

$$W = \sinh(2)V(2) - \sinh(1)V(1). \quad (13)$$

The two nuisance terms cancel. Moreover,

$$\sinh(b)\Delta(b) = H(b) := \frac{(1 - e^{-b})^2}{b^2} - e^{-b},$$

so the two predictions differ by

$$H(2) - H(1) = 0.01987902514497878 \dots > 0. \quad (14)$$

In closed form, with e the ordinary exponential constant,

$$\begin{aligned} W_{\text{AH}} &= \frac{e^2}{4} + \frac{5}{4e^2} - e - \frac{2}{e} + \frac{3}{2}, \\ W_{\text{GUE}} &= \frac{e^2}{4} - e + \frac{3}{4} + \frac{1}{e} - \frac{5}{4e^2} + \frac{1}{4e^4}. \end{aligned} \quad (15)$$

The companion rational script bounds e by its degree-40 Taylor sum and a geometric upper bound for the tail. Ordinary interval arithmetic gives $W_{\text{AH}} \in (0.06239, 0.06240)$, $W_{\text{GUE}} \in (0.08227, 0.08228)$, and their difference in $(0.019879, 0.019880)$. These are exact constant enclosures. They contain no computed ζ mean square.

4. Passing from AH-Pairs to these noncompact pair tests

The convolution kernel in physical space is

$$K_b(x) = P_{2a}(x) = \frac{2b}{b^2 + 4\pi^2 x^2}, \quad Q_T(b) = \int K_b d\mu_T - 1. \quad (16)$$

For fixed b this kernel and the two-scale linear combination decay as $O_b((1+x^2)^{-1})$. The cited pair bound gives $\mu_T([-h, h]) \ll 1 + h$ for $0 \leq h \leq T$. Dyadic shells therefore bound the tail between M and T by $O_b(1/M)$. Beyond T , total pair mass is $O(T \log T)$ and the kernel is $O_b(T^{-2})$, giving $O_b(\log T/T)$.

Fix M first along values $M=j+1/4$, so the hard truncation endpoints lie away from all limiting half-lattice atoms. The finite set of half-lattice mass formulas (5), plus concentration of each bin at $k/2$, then evaluates the compact part without an unresolved boundary atom. The early-zero range excluded in AH0 changes the compact pair sum by $o(1)$, as in the source's Section 2. The two tails are uniform in T . The model mass tails are also $O_b(1/M)$, uniformly for bounded $p_0(T)$. Hence

$$Q_T(b) = V_A(b) + \frac{2(p_0(T) - 1)}{\sinh b} + o(1) \quad (17)$$

under RH and AH-Pairs. Formally: the limsup error is $O_b(1/M)$ after T tends to infinity, then M tends to infinity. One cannot cancel a finite truncated alternating sum as though it were already the entire comb. This two-limit argument justifies the cancellation in (13) without requiring existence of $\lim p_0(T)$.

5. Actual completed–zeta resolvent: endpoints and centering

Here is the conversion from the pair statistic to the real part of an actual meromorphic function. All following b values are fixed and positive. Set $a=b/(4\pi)$, $\eta=a/L$, and

$$Y_T(t) = \sum_{0 < \gamma \leq T} P_a(L(t - \gamma)), \quad Z_T(t) = \sum_{\gamma \in \mathbb{R}} P_a(L(t - \gamma)).$$

The second sum is over all nontrivial zero ordinates, including negative ordinates. Under RH the paired canonical product of ξ gives exactly

$$Z_T(t) = \frac{1}{\pi L} \Re \frac{\xi'}{\xi} (1/2 + \eta + it). \quad (18)$$

Pairing positive and negative zeros fixes the constant: ξ is even about $1/2$, and the paired product has factors $1 + (s - 1/2)^2/\gamma^2$. The real part of its logarithmic derivative is the absolutely convergent sum of $\eta/[\eta^2 + (t - \gamma)^2]$. Formula (18) does not require simple zeros.

The finite convolution identity is exact:

$$\frac{1}{T} \int_{\mathbb{R}} Y_T(t)^2 dt = \int K_b d\mu_T = Q_T(b) + 1. \quad (19)$$

We give the endpoint estimates because they must not be hidden in a stationarity assumption. Standard unit-interval zero counting gives $O(\log(|u| + 2))$ zeros in an interval of length one. Consequently, for fixed a , both Y_T and Z_T are $O_a(\log T)$ on $[0, T]$. Let $w = T/\log^4 T$. The two endpoint strips of width w contribute at most $O_a(\log^{-2} T)$ to a normalized square integral.

On $[w, T - w]$, the omitted positive zeros above T and the negative zeros contribute

$$|Z_T(t) - Y_T(t)| \ll_a \frac{\log T}{L^2 w} = o(1)$$

uniformly, by summing the inverse-square kernel with the unit-interval bound. The pair bound used in Section 4 implies $T^{-1} \int_{\mathbb{R}} Y_T^2 = O_a(1)$, so the square-integral difference on the interior is $o(1)$ by Cauchy – Schwarz.

For t outside $[0, T]$ by a distance $d \geq w$, the same unit-interval count yields $Y_T(t) \ll_a \log T / (L^2 d)$. Its exterior squared integral divided by T is $o(1)$; the exterior strips of width w have the same $O(\log^{-2} T)$ bound. Finally,

$T^{-1} \int_0^T Y_T(t) dt \rightarrow 1$: the full integral is $N(T)/L$, boundary zeros number $O(w \log T)$, and for each interior zero the missing Poisson mass is at most $O_a(L^{-2}(\gamma^2 - 1 + (T - \gamma)^2 - 1))$. Summing this last expression gives $O_a(\log^2 T / L^2) = O_a(1)$. Division by T removes it.

It follows from (18) – (19) that

$$Q_T(b) = \frac{1}{T} \int_0^T \left(\frac{1}{\pi L} \Re \frac{\xi}{\xi} (1/2 + \eta + it) - 1 \right)^2 dt + o(1). \quad (20)$$

No assertion that the finite zeros are exactly stationary was used.

6. Gamma factor and the real-square versus modulus distinction

By definition,

$$\frac{\xi}{\xi}(s) = \frac{\zeta}{\zeta}(s) + \frac{1}{s} + \frac{1}{s-1} - \frac{1}{2} \log \pi + \frac{1}{2} \frac{\Gamma'}{\Gamma}(s/2). \quad (21)$$

For $t \geq 1$ and the specified σ near $1/2$, the real part of the extra terms is $\frac{1}{2} \log(t/(2\pi)) + O(1/t)$. After dividing by πL and subtracting one, its discrepancy from zero has normalized L^2 norm $O(1/\log T)$. This follows by integrating $|\log(t/T) - \log(2\pi)|^2$, and treating $0 \leq t \leq 1$ separately using the absence of low ordinates and (21). The pair bound gives a bounded normalized L^2 norm for (20); hence Cauchy – Schwarz justifies removing this discrepancy. We obtain

$$Q_T(b) = \frac{4}{T \log^2 T} \int_0^T \left(\Re \frac{\zeta}{\zeta} (1/2 + \eta + it) \right)^2 dt + o(1). \quad (22)$$

It is not an algebraic identity that the real-square integral equals half the modulus-square integral. To justify that step here, let $F = \zeta/\zeta$. Under RH, F^2 is analytic in the rectangle with horizontal sides at heights 1 and T , and vertical sides at $\sigma = 1/2 + \eta$ and 2. The pole at $s=1$ lies below this rectangle. On the right side F^2 has an absolutely convergent Dirichlet series without a constant term; its vertical integral is $O(1)$. On the top side the RH partial-fraction estimate gives $F = O_b(\log^2 T)$, uniformly across the rectangle, so that horizontal integral is $O_b(\log^4 T)$. The bottom integral is bounded. Thus

$$\int_0^T F(1/2 + \eta + it)^2 dt = O_b(\log^4 T) = o(T \log^2 T).$$

The bounded initial segment adds no relevant term. Since $2(\Re F)^2 = |F|^2 + \Re F^2$, (22) becomes

$$\boxed{Q_T(b) = \frac{2}{T \log^2 T} I_T(b/2) + o(1).} \quad (23)$$

This is consistent with the classical correspondence of Goldston – Gonek – Montgomery; Fazzari, 2310.15918, Introduction and Lemma 3 also records its normalization and explicitly separates the real-square and holomorphic-square terms. A weighted version would need its own contour argument; the unweighted vanishing above must not be silently reused with an arbitrary arithmetic weight.

Equations (17), (13) and (23) prove the claimed reduction (2), subject to the stated primary RH+AH-Pairs density input. The threshold (4) is still missing.

7. Finite ensemble calibration and what cannot prove (4)

On a circle of length N , periodize P_a . For nonzero integer m the CUE trace second moment is $\min(|m|, N)$. For the randomly rotated ACUE model it is $\min(r, 2N-r)$ when $r=m \bmod 2N$ is nonzero, and N^2 when $r=0$. The latter large peaks become the even-frequency atoms in (9).

With $q=e^{-(b/N)}$, the exact finite CUE variance is

$$V_N^{\text{CUE}}(b) = \frac{2q(1-q^N)}{N^2(1-q)^2}, \quad \frac{V_N^{\text{CUE}}(b)}{V_{\text{sine}}(b)} = \left(\frac{b/(2N)}{\sinh(b/(2N))} \right)^2. \quad (24)$$

The finite ACUE formula is

$$V_N^{\text{ACUE}}(b) = \frac{2}{N^2} \frac{\sum_{m=1}^{2N-1} \min(m, 2N-m) q^m + N^2 q^{2N}}{1 - q^{2N}}. \quad (25)$$

The script checks eight widths, four finite sizes, direct triangular-period integrals, and exhaustive float64 subset sums for $N=2, \dots, 6$. Symbolic checks verify (11), both limiting signal sizes, and the paired canonical-product normalization on an exact rational finite example. These are normalization tests and finite-model evidence, not ζ data or an asymptotic numerical proof.

The base half-lattice process itself obeys the known low-frequency pair data while giving W_{AH} . Therefore no deduction using only those data and positivity can establish (4). Its signed Fourier kernel is $\sinh(2)e^{-2|a|} - \sinh(1)e^{-|a|}$, which changes sign at $\log(2 \cosh 1) > 1$. The remaining arithmetic work controls information beyond the known band, including the negative tail. A fresh prime-correlation estimate, or a new effective arithmetic weighting/comparison, is necessary. Enlarging the finite matrix model or improving numerical precision does not supply it.

The retained deliverable is a precise one-sided ζ target with certified separation and no assumed near-diagonal limit. Further work should attack its actual signed mean square or the companion compact Fourier test. External novelty review and proof-assistant formalization are postponed until the ordinary proof audit and the arithmetic direction justify them.

Current report 10: Independent review of the two-scale actual-zeta reduction

Collection: R7 — actual-zeta targets, arithmetic mark, and flow obstruction.

Source: `research/dyson/round7/poisson-resolvent/INDEPENDENT_REVIEW.md`.

SHA-256: `82db15288f7b6f902a461c296abe30eb1d21253bf14ae4949da4dc12c592a370` . **Git blob:** `15b2800098652211278edc2e0c0b62b9b277d92e` . **Original bytes:** 11724.

Independent review of the two-scale actual-zeta reduction

Date: 2026-09-05. **Reviewer:** the residual/arithmetic lane, independently of the authoring lane.

Verdict: accepted as an ordinary RH-dependent reduction, with no remaining mathematical obstruction found in the reviewed argument. The review accepts `RH + AH-Pairs => W_T -> W_AH`, including cancellation without a limit for $P_0(T)$, and accepts both proposed sufficient arithmetic lower bounds. It does **not** prove either lower bound for actual zeta mean squares, refute AH, establish a new zeta-spacing theorem, certify novelty, or constitute proof-assistant verification.

One minor truncation issue was identified and corrected by the author during review: the fixed physical pair cutoff must avoid limiting half-lattice atoms. The current text explicitly takes $M=j+1/4$. The transfer proof and model formulas were otherwise unchanged. The subsequent addition of the weaker sufficient threshold $1/16$ has also been checked.

The reviewed report and script hashes, temporary-copy replay details, and final metadata-only status change are pinned in `independent_review_receipt.json`.

1. Acceptance coverage

component	result	reason
Exact source AH formulation and RH dependence	accepted	The source's AH0, (1.12), (1.14), (1.15) supply exactly the stated pair concentration, tail bound and finite-T masses.
Near-diagonal nuisance cancellation	accepted	The alternating comb includes $k=0$ and contributes $2(P_0(T)-1)/\sinh(b)$; the two coefficients cancel at the same T. No limit for $P_0(T)$ is used.
Noncompact pair test	accepted after the endpoint clarification	Uniform $O(1/M)$ tails allow fixed-M asymptotics first; $M=j+1/4$ avoids atomic truncation endpoints.

component	result	reason
Completed-zeta normalization	accepted	Pairing positive and negative zeros removes the possible exponential constant; the real logarithmic derivative has the required absolutely convergent Poisson sum.
Finite-zero endpoints and gamma centering	accepted	Width $T/\log^4(T)$, unit-interval zero counts and the pair bound control all normalized square tails; the gamma discrepancy has normalized L2 norm $O(1/\log T)$.
Real square to modulus square	accepted	Under RH, the square of ζ'/ζ has no poles in the chosen rectangle; the holomorphic-square vertical integral is negligible for every large T .
Constants and sufficient thresholds	accepted	Exact rational enclosures separate W_{AH} from both $7/100$ and $1/16$; independent symbolic checks link the closed constants to the variance formulas.
Script scope	accepted	The scripts prove/check constants and finite-model identities. They do not compute an actual-zeta mean square.

2. Primary source check and the role of $P_0(T)$

The primary paper was read directly at [Goldston – Lee – Schettler – Suriajaya, arXiv:2507.06823v1](#). Its AH-Pairs formulation uses the high interval $T/\log^2(T) < \gamma, \gamma' \leq T$. The RH-dependent fixed- k formulas and bounded $P_0(T)$ are available there; its Section 2 explains removal of the excluded early range. The reduction correctly keeps these inputs separate from the paper's additional AH-Weak Density hypotheses.

Only additive $o(1)$ accuracy for each fixed k is required. Since only finitely many k occur before the truncation limit, those errors may be combined without imposing uniform asymptotics for k growing with T . The near-zero bin may include multiplicity and near-coincident distinct zeros. It is therefore correct to preserve $P_0(T)$, rather than silently replacing it by one.

Relative to the $p_0=1$ half-lattice pair measure, the finite- T mass change is

$$(P_0(T)-1) \sum_k (-1)^k \delta_{(k/2)}.$$

Its Fourier transform has mass 2 at every odd integer. Therefore its Poisson variance contribution is

$$\begin{aligned} & 2(P_0(T)-1) \sum_{(m \text{ in } Z)} \exp(-b |2m+1|) \\ &= 2(P_0(T)-1)/\sinh(b). \end{aligned}$$

The coefficient two and the inclusion of the diagonal are both necessary. The same formula follows in physical space from the classical alternating Cauchy-kernel sum. The combination at $b=2$ and $b=1$ cancels the entire nuisance parameter at each T , after the noncompact-test approximation has been justified.

3. Pair tails and the corrected cutoff

The source pair bound gives normalized mass $O(1+h)$ for differences of size at most h , for $0 \leq h \leq T$. Since $K_b(x) = O_b((1+x^2)^{-1})$, a dyadic decomposition between M and T contributes $O_b(1/M)$. Beyond T , the total normalized pair mass is $O(T \log T)$, whereas the kernel is $O_b(T^{-2})$, giving $O_b(\log T/T)$. These estimates are uniform in T for fixed b .

For the model family, bounded $P_0(T)$ gives uniformly bounded absolute atomic weights, hence another $O_b(1/M)$ tail. Positivity of every finite- T approximate model weight is not needed for that latter absolute bound.

A hard cutoff at a half integer would not by itself have a determined limit: an AH cluster could approach that endpoint from either side. The author's repair $M=j+1/4$ eliminates this ambiguity. The early-range discrepancy is $o(1)$ for each such fixed cutoff, and the limiting bin concentration then evaluates the finite sum. Sending T to infinity first and M to infinity second proves the stated expansion uniformly for the bounded nuisance parameter. It is legitimate that $Q_T(b)$ need not itself converge when $P_0(T)$ oscillates; only the two-scale combination is asserted to converge.

4. Completed-zeta identity and endpoints

With $L = \log(T)/(2\pi)$, $a = b/(4\pi)$, and $\eta = a/L$, a single Poisson term satisfies

$$\begin{aligned} P_a(L(t-\gamma)) \\ = 1/(\pi L) * \eta / [\eta^2 + (t-\gamma)^2]. \end{aligned}$$

Under RH, the paired canonical product about $1/2$ has factors $1+(s-1/2)^2/\gamma^2$. The function is even about $1/2$ and has order one; no nonconstant exponential prefactor remains after pairing. Its logarithmic derivative consequently gives exactly the sum over both positive and negative ordinates, with multiplicities. The real-part series converges absolutely. There is no missing factor of two: summing over both signs already supplies the two linear factors of each paired zero.

The finite convolution identity uses $P_a * P_a = P_{(2a)}$ and a change of variable $dt = dx/L$. Its normalization is therefore exactly the report's $(TL)^{-1}$ pair measure. The kernel evaluates to

$$K_b(x) = 2b/(b^2 + 4\pi^2 x^2).$$

The unit-interval zero count is sufficient for the stated endpoint estimates. On the two strips of width $w = T/\log^4(T)$, both finite and completed sums are $O_b(\log T)$, so their normalized squared integrals are $O_b(\log^2 T)$. On the interior, omitted negative and above- T zeros have total Poisson contribution $O_b(\log T/(L^2 w))$. The pair bound provides a bounded normalized square integral for the finite sum, so Cauchy – Schwarz controls the interior replacement.

The far exterior estimate $O_b(\log T/(L^2 d))$ has an integrable square for $d \geq w$. The full integral of a finite Poisson sum is $N(T)/L$; the loss from interior zero tails is bounded by summing the displayed reciprocal distances, and the boundary zero count gives a vanishing normalized contribution. Thus the mean tends to one and the subtraction of

one in the centered square is correctly normalized. No exact stationarity of finite zeta zeros is assumed.

5. Gamma centering and the holomorphic-square rectangle

The extra real terms in ξ'/ξ equal one half of $\log(t/(2\pi))$, with the stated uniform lower-order error. Since $\pi L = \log(T)/2$, dividing by πL and subtracting one leaves

$$[\log(t/T) - \log(2\pi)]/\log(T) + \text{lower-order terms.}$$

Its normalized L2 norm is $O(1/\log T)$: the integral of $|\log(t/T)|^2$ divided by T stays bounded, and the fixed initial interval causes no difficulty. The bounded normalized L2 norm supplied by the pair statistic justifies removal of this centering discrepancy by Cauchy – Schwarz. This yields the factor four in front of the real-square mean.

For $F = \zeta'/\zeta$, RH places every nontrivial zero strictly to the left of the rectangle whose left side is $1/2 + \eta$. The pole at one is below its bottom side $t=1$. Thus F^2 is analytic in the rectangle. On $\text{Re}(s)=2$, its absolutely convergent Dirichlet series has no constant coefficient, so its vertical integral is $O(1)$. On the top side, the local zero count and distance η from the critical line give the uniform bound $F = O_b(\log^2 T)$, hence an $O_b(\log^4 T)$ horizontal integral. This works even if T is exactly a zero ordinate: the positive real displacement η prevents a pole. The bottom and the initial segment are harmless.

Consequently the holomorphic-square integral is $o(T \log^2 T)$. Applying

$$2(\text{Re } F)^2 = |F|^2 + \text{Re}(F^2)$$

is now justified, rather than merely assumed. Finally $\eta = b/(2 \log T)$, so the modulus integral is exactly $I_T(b/2)$, giving

$$Q_T(b) = 2 I_T(b/2) / (T \log^2 T) + o(1).$$

This argument is unweighted. An arbitrary arithmetic weight would change the contour problem and cannot inherit this vanishing without a separate proof.

6. Model spectra and exact constants

The half-lattice pair measure has diagonal mass one and the correct nonzero half-lattice masses. Fourier transformation gives the triangular periodic density plus atoms at the nonzero even integers. Those atoms are essential. The nuisance modification instead adds atoms at odd integers. The reported spectra and their Poisson integrals are consistent with these two different combs.

The variance formulas were checked independently against the closed W expressions as exact rational-function identities in $x = \exp(1)$. This closes a small verification gap that a script evaluating only the closed constants would leave: the constants really are the stated two-variance combination, not merely two separated numbers.

The rational exponential enclosure uses the sum through degree 40 and bounds the remaining factorial series by a geometric majorant. Its interval arithmetic is valid even though it repeats the same exponential variable, because independent interval multiplication only enlarges the enclosure. The certified inequalities imply both

```
W_AH < .06240 < .07,
W_AH < .06240 < 1/16,
```

and $1/16 - W_{AH} > .00010$. Therefore either RH-conditional lower bound on the corresponding $\liminf$ would contradict AH-Pairs. The number .07 is a convenient stronger target, not a logically necessary threshold. Neither arithmetic lower bound follows from the constant calculation.

7. Replay, scope and remaining obligation

Both check scripts were copied to a fresh temporary directory and run there. Their generated JSON matched the author's files after ignoring the elapsed-time field. The exact constant script was replayed again after addition of the $1/16$ assertions. Original mathematical files were not altered by the reviewer. The only authorized final report edit is the first paragraph's review-status metadata and link to this review; the receipt verifies that distinction and pins the final files.

The finite subset enumerations are floating-point calibration, while the symbolic identities and Fraction arithmetic establish their stated algebraic and constant assertions. A rational finite canonical-product example checks a normalization; it does not by itself prove the infinite ξ -product statement, which is instead justified analytically above.

The unresolved work is the actual signed mean-square inequality. The half-lattice model already matches the known interior Fourier band and satisfies spectral positivity, so those facts alone cannot force a value above W_{AH} . The reviewed reduction makes the missing arithmetic statement precise; it does not supply the out-of-band information or control the negative tail by a new prime-correlation estimate. Acceptance of this reduction must not be described as a proof of such an inequality or as a refutation of AH.

Current report 11: Independent review of the actual–zeta Poisson transfer

Collection: R7 — actual-zeta targets, arithmetic mark, and flow obstruction.

Source: `research/dyson/round7/dyson-frontier/POISSON_TRANSFER_REVIEW.md`.

SHA-256: `fc7959a2f6147380709467776add8e4a59aa255747e46c70cad413d808e64b36` . **Git blob:**
`e6c81f58e05560d6f3b6d5fced09da33a1381244` . **Original bytes:** 3945.

Independent review of the actual–zeta Poisson transfer

5 September 2026. This reviews the root agent's proposed transfer, not a proof of the desired arithmetic mean-square asymptotic. All statements below assume RH and a fixed $b > 0$.

Put $L = \log T/(2\pi)$ and

$$K_b(u) = \frac{2b}{b^2 + 4\pi^2 u^2}, \quad Q_T(b) = \frac{1}{TL} \sum_{0 < \gamma, \gamma' \leq T} K_b(L(\gamma - \gamma')) - 1.$$

The proposed identity has the correct factor:

$$Q_T(b) = \frac{2}{T \log^2 T} \int_0^T \left| \frac{\zeta}{\bar{\zeta}} \left(\frac{1}{2} + \frac{b}{2 \log T} + it \right) \right|^2 dt + o(1).$$

The following checks support it.

1. If $P_a(u) = a/[\pi(\sigma^2 + u^2)]$, then $K_b = P_{2a}$ with $a = b/(4\pi)$. The physical displacement is $\delta = a/L = b/(2 \log T)$. For $Y(t) = \sum_{0 < \gamma \leq T} P_a(L(t - \gamma))$, convolution gives the exact identity $T^{-1} \int_{\mathbb{R}} Y^2 = Q_T(b) + 1$. No circular periodization is involved.
2. The full positive zero density is $Z(t) = (\pi L)^{-1} \operatorname{Re}(\xi/\bar{\xi})(1/2 + \delta + it)$, by the symmetrically paired Hadamard product. On the interval interior, the omitted zeros below zero or above T contribute $O_b(\log T/(L^2 w))$ if the distance to the endpoints is at least $w = T/\log^4 T$.
3. The unit-interval zero bound implies $Y, Z = O_b(\log T)$ near the endpoints. Their squared contribution divided by T is $O_b(w \log^2 T/T) = O_b(\log^{-2} T)$. Outside the interval and at distance $d \geq w$, one can use $Y \ll_b \log T/(L^2 d)$ until d reaches T , then $Y \ll_b N(T)/(L^2 d^2)$. These bounds make the exterior squared integral $o(T)$. They avoid the incorrect use of a constant bound on an infinite exterior interval.
4. The mean $T^{-1} \int_0^T Y \rightarrow 1$ follows from the zero-count asymptotic and the Poisson tail estimate. The same holds for Z . The gamma and elementary factors give

$$Z(t) - 1 = \frac{2}{\log T} \operatorname{Re} \frac{\zeta}{\bar{\zeta}} \left(\frac{1}{2} + \delta + it \right) + \frac{\log(t/(2\pi)) - \log T}{\log T} + \text{a negligible compact-height correction.}$$

The displayed drift has mean square $O(\log^{-2} T)$ after averaging over $[1, T]$. Uniform boundedness of the pair-kernel quadratic form controls its cross term by Cauchy – Schwarz. 5. For $q(s) = \zeta/\zeta(s)$, a rectangle with vertical sides $\sigma = 1/2 + \delta$ and $\sigma = 2$ and lower height 1 avoids the pole at $s = 1$. Under RH, $|q(\sigma + iT)| \ll_b \log^2 T$ uniformly along its top, even when T is a zero ordinate. Thus the top integral of q^2 is $O_b(\log^4 T)$. At $\sigma = 2$, the squared absolutely convergent Dirichlet series has no constant term and has a bounded time primitive. The lower side is harmless; alternatively choose any fixed lower height away from zero ordinates and treat the omitted compact interval separately. Consequently $\int_0^T q(1/2 + \delta + it)^2 dt = o(T \log^2 T)$. Using $(\operatorname{Re} q)^2 = (|q|^2 + \operatorname{Re} q^2)/2$ produces the claimed factor 2.

For passage from AH-Pairs to Poisson tests, the cumulative local-pair bound in GLSS (1.12), printed p.4, suffices. It is only stated up to normalized radius T . Dyadic summation gives $O(1/R)$ tails for the Poisson pair kernel inside that range, and the total-mass bound gives $O(\log T/T)$ outside it. A global uniform translation bound is not required.

The nuisance term in the centered spectral measure is $2(\rho_0 - 1) \sum_m \delta_{2m+1}$. Against $e^{-b|a|}$ it contributes $2(\rho_0 - 1)/\sinh b$. Hence a signed combination $\sinh(b_1)Q_T(b_1) - \sinh(b_2)Q_T(b_2)$ can eliminate this parameter without a limit assumption on $P_0(T)$: first truncate at radii $M = j + 1/4$ to avoid half-lattice boundary atoms, apply the uniform finite- T density formulas, and then use the tail estimates.

Verdict: the normalization, endpoint strategy, complex-square cancellation, and nuisance-parameter elimination are consistent. The remaining mathematical input is an actual estimate of the resulting signed logarithmic-derivative mean squares. The transfer identity itself does not provide that estimate or refute AH.

Current report 12: One actual–zeta target beyond Fourier support one

Collection: R7 — actual-zeta targets, arithmetic mark, and flow obstruction.

Source: `research/dyson/round7/dyson-frontier/DYSON_ACTUAL_ZETA_FRONTIER.md`.

SHA-256: `3b8eb5fb9efc2f53af550db64c1e8eef7233f8be83564e906357b3acf00a0301` . **Git blob:** `5cb96f2a0878978a39cf4025d1f40e760d0d4e67` . **Original bytes:** 17526.

One actual–zeta target beyond Fourier support one

Date: 5 September 2026. Status: primary-source audit, an exact arithmetic reduction, and a conditional AH-Pairs detector. The required new arithmetic estimate is not proved here. The reductions use existing explicit-formula and Fourier arguments; no priority or novelty claim is made. No claim to have proved Montgomery's conjecture, refuted AH for zeta, or established RH is made.

1. Selected research target and the result of this round

Fix once and for all a nonnegative, normalized, symmetric smooth bump ϕ with closed support $[6/5, 7/5] \subset (1, 2)$ and centered at $13/10$. For example, normalize

$$\phi(\alpha) = C \exp\left(-\frac{1}{1 - (10\alpha - 13)^2}\right) \mathbf{1}_{\{|10\alpha - 13| < 1\}}.$$

Then $\int \phi = 1$ and $\int \alpha \phi(\alpha) d\alpha = 13/10$. Assuming RH, define the actual-zeta form factor, with all zeros counted with multiplicity, by

$$F_T(\alpha) = \frac{2\pi}{T \log T} \sum_{0 < \gamma, \gamma' \leq T} T^{i\alpha(\gamma - \gamma')} \frac{4}{4 + (\gamma - \gamma')^2}. \quad (1)$$

The ambitious target is the single smoothed Montgomery prediction

$$C_\phi(T) := \int \phi(\alpha) F_T(\alpha) d\alpha \rightarrow 1. \quad (2)$$

There is a weaker, still consequential target:

$$\liminf_{T \rightarrow \infty} C_\phi(T) > \frac{7}{10}. \quad (3)$$

Under RH, (3) contradicts the precise AH-Pairs hypothesis defined in Goldston – Lee – Schettler – Suriajaya, not merely the particular ACUE model. This assertion is proved in § 3 below. It depends on the stated AH definition; it is not a claim about every informal hypothesis sometimes called “alternative.”

The main output is the exact centered prime-covariance obligation (12) – (15). It reduces (2) to a specified signed average of shifted von Mangoldt correlations, with the pole/continuous-mean term retained. It also shows precisely why a scalar sieve distribution exponent is not the missing input. No estimate closing (3) was obtained from the currently checked tools.

2. Primary–source state as checked on 5 September 2026

Result	Status and implication for this target
Montgomery's actual-zeta pair-correlation theorem	Under RH, $F_T(\alpha) = \alpha + T^{-2 \alpha } \log T + \mathcal{O}(1)$ in the known low band, interpreted distributionally at zero. This evaluates Fourier tests supported inside $(-1, 1)$, not (2). See Goldston's primary-author notes , Theorem 1.
Higher correlations	Rudnick – Sarnak obtain the GUE correlations for the usual real-zero formulation under RH and the Fourier condition $\sum_j \xi_j < 2$. On the translation-invariant hyperplane this does not supply an independent pair frequency above one. See their original announcement , Theorem 1.2 and its first remark.
Unconditional pair-correlation identity	Baluyot – Goldston – Suriajaya – Turnage-Butterbaugh remove RH from a generalized identity involving the full complex zeros. Interpreting it as an ordinary positive statistic of real ordinates requires care. This is not an asymptotic plateau on an interval above one.
The “0.6…” AI result	The precise constant is $C_0 = 3/2 - 2^{-1/2} \cot(2^{-1/2}) = 0.6725007 \dots$ for zeros that are both simple and on the critical line, unconditionally. Anthropic's paper and Lamzouri's new proof , Theorem 1.1, give this result and the distinct-zero proportion $(1 + C_0)/2$. Lamzouri's arXiv submission is dated 2 September, not 3 September. His Lemma 3.1 still uses support $[-1, 1]$. This improves what can be inferred from available spectral information; it does not provide the covariance in § 4.
Quantitative bounds above one	Carneiro – Milinovich – Ramos , Corollary 2, bound sufficiently long interval averages between $0.9303 + \mathcal{O}(1)$ and $1.3208 + \mathcal{O}(1)$ under RH. These are bounds, not a plateau asymptotic. Their equation (1.7) records the GRH lower bound $F_T(\alpha) \geq 3/2 - \alpha - \varepsilon$ on its stated range.
AH and simplicity	Goldston – Lee – Schettler – Suriajaya II , Theorem 4, derives asymptotic simple critical zeros from AH-Pairs plus their AH-Weak Density assumptions. It does not refute AH. Their companion PCC paper derives simple critical zeros conditionally on PCC without assuming RH; it does not prove PCC.

The checked primary sources do not establish (2) or (3). Results for a family of Dirichlet L -functions with extra character/modulus averaging, and results for function fields, must not be imported as the corresponding theorem for this single zeta function. The CMR paper itself separates its family result with support below two from the actual-zeta problem.

This round also read the existing force-energy, dynamic-generator, attachment-bridge audit, and main handoff. Their valid conclusions remain in force: the circular force square reduces to a singular two-point observable; protected trace dynamics do not create new bandwidth; and a finite heat-depth theorem supplies no unproved local flow for actual zeta. Those old calculations were not rerun.

3. Why the compact test excludes the precise AH-Pairs class

Put $L = \log T/(2\pi)$ and $N_T = TL$, and consider the positive pair measure

$$\mu_T = \frac{1}{N_T} \sum_{0 < \gamma, \gamma' \leq T} w(\gamma - \gamma') \delta_{L(\gamma - \gamma')}, \quad w(v) = \frac{4}{4 + v^2}. \quad (4)$$

With Fourier convention $\hat{f}(a) = \int f(u) e^{-2\pi i a u} du$, its Fourier transform is F_T ; the sign is immaterial because the pair measure is even.

The primary AH-Pairs condition says that for each fixed M , every pair with both ordinates in $(T/\log^2 T, T]$ and $|L(\gamma - \gamma')| \leq M$ lies within $O((|k| + 1)R(T))$ of some $k/2$, where $R(T) \rightarrow 0$. The needed compactness is not assumed without evidence. Their equation (1.12), printed p.4, supplies

$$\mu_T([-R, R]) \ll 1 + R \quad (0 \leq R \leq T). \quad (5)$$

For $R > T$, the crude total mass bound is $O(T \log T)$; do not upgrade (5) to a global uniform linear bound. Together these two estimates control Schwartz tails uniformly. For a test decaying like $(1 + u^2)^{-1}$, dyadic shells up to T give $O(1/R)$ and the remaining tail gives $O(\log T/T)$. This also suffices for the Poisson pair kernels used in the parallel resolvent lane.

Pairs involving an ordinate below $T/\log^2 T$ contribute $o(N_T)$ to any fixed local test: there are $O(T/\log T)$ such zeros, and the elementary unit-interval zero bound gives $O(\log T)$ possible partners in a fixed normalized window. After division by N_T this is $O(1/\log T)$. The uniform tail estimate then handles Schwartz tests. On compact normalized intervals $w(u/L) \rightarrow 1$, so the weighted and unweighted local limits agree.

Consequently every subsequential tempered limit μ is a measure supported on $\frac{1}{2}\mathbb{Z}$. It obeys $e^{4\pi i u} \mu = \mu$, hence μ is 2-periodic as a tempered distribution. Montgomery's theorem determines its restriction to $(-1, 1)$:

$$\mu = \delta_0 + |a| da \quad \text{on } (-1, 1).$$

Translation by two therefore determines it on $(1, 2)$:

$$\hat{\mu} = (2 - a) da. \quad (6)$$

Testing (6) against the fixed ϕ gives $C_\phi(T) \rightarrow 7/10$. The conclusion holds for every subsequence, so no prior existence of the full limiting pair measure is needed. It proves the contradiction claimed in (3). A second agent independently checked this argument, including the restricted range in (5) and the removal of early zeros.

The unknown near-diagonal mass is confined to odd-frequency atoms

There is a useful more explicit statement. The same paper, equations (1.14) – (1.15), printed p.4, gives, under RH and AH-Pairs,

$$1 \leq \rho_0 \leq \frac{3}{2} - \frac{2}{\pi^2}$$

along a subsequence on which the near-zero pair mass converges. The nonzero even half-lattice masses are $\rho_0 - 1/2$; the odd ones are $3/2 - \rho_0 - 2/(\pi^2 k^2)$. Their difference from the $\rho_0 = 1$ measure is

$$(\rho_0 - 1) \sum_{k \in \mathbb{Z}} (-1)^k \delta_{k/2}.$$

Writing $\text{tri}_2(\alpha) = \text{dist}(\alpha, 2\mathbb{Z})$, the full pair spectral measure is therefore

$$\hat{\mu}_{\rho_0} = \text{tri}_2(\alpha) d\alpha + \sum_{m \in \mathbb{Z}} \delta_{2m} + 2(\rho_0 - 1) \sum_{m \in \mathbb{Z}} \delta_{2m+1}. \quad (7)$$

The centered spectral measure omits the atom at zero. The fixed test in (2) avoids all integer atoms, which is why it needs no simplicity or $\rho_0 = 1$ assumption. A one-scale unbandlimited resolvent statistic generally does need to account for the last term. Two suitably chosen signed Poisson scales can cancel this parameter; that is the separate root-agent construction reviewed in the adjacent note.

4. Exact prime–side identity with the pole term retained

For $x > 1$ and $u > 0$ define

$$a_u(x) = \min\{(u/x)^{1/2}, (x/u)^{3/2}\}.$$

The absolutely convergent arithmetic signal is

$$\begin{aligned} P_x(t) &= \sum_{n \geq 2} \Lambda(n) a_n(x) n^{-it} - M_x(t), \\ M_x(t) &= \int_0^\infty a_u(x) u^{-it} du = \frac{2x^{1-it}}{(1/2 + it)(3/2 - it)}. \end{aligned} \quad (8)$$

The equality in (8) follows by integrating separately below and above x . This continuous mean is precisely the pole term in Montgomery's explicit formula. It cannot be dropped. The series converges absolutely because its tail is $O_x(\Lambda(n)n^{-3/2})$; this is not a formal expansion of $\log \zeta$ or ζ'/ζ on the critical line.

Goldston's notes, Proposition 1 and equations (4.4) – (4.5), give under RH

$$\begin{aligned} P_x(t) &= -2x^{1/2-it} \sum_{\gamma \in \mathbb{R}} \frac{x^{i\gamma}}{1 + (t - \gamma)^2} + E_x(t), \\ E_x(t) &\ll x^{-1/2} \log(t+2) + x^{-2}/(t+2). \end{aligned}$$

The integrated zero kernel differs from the truncated pair sum by $O(x \log^3 T)$. Its squared norm is $O(xT \log^2 T + x \log^3 T)$. Cauchy – Schwarz with $\int_0^T |E_x|^2 \ll x^{-1} T \log^2 T$ shows, uniformly for $x = T^\alpha$ with $6/5 \leq \alpha \leq 7/5$,

$$F_T(\alpha) = \frac{1}{xT \log T} \int_0^T |P_x(t)|^2 dt + o(1). \quad (9)$$

For example, the resulting normalized error is bounded by a constant times $\log T/x + \log^2 T/T$ plus smaller terms. Thus the interchange with the fixed bump has no unstated uniformity problem.

Let $d\Delta(u) = d\psi(u) - du$, where $d\psi$ has mass $\Lambda(n)$ at each integer n . Define the explicit real symmetric kernel

$$K_T(u, v) = \frac{1}{T \log T} \int \phi(\alpha) T^{-\alpha} a_u(T^\alpha) a_v(T^\alpha) \frac{\sin(T \log(u/v))}{\log(u/v)} d\alpha, \quad (10)$$

using the continuous value T on the diagonal. Absolute convergence of the weighted total variations justifies Fubini for each fixed T . Expanding the finite-time square gives the exact identity

$$C_T := \iint K_T(u, v) d\Delta(u) d\Delta(v) = \frac{1}{T \log T} \int \phi(\alpha) T^{-\alpha} \int_0^T |P_{T^\alpha}(t)|^2 dt d\alpha. \quad (11)$$

By (9), $C_T = C_\phi(T) + o(1)$. The kernel changes sign away from the diagonal; positivity of the whole squared norm does not make each shifted-prime summand positive.

Separate the atomic diagonal

$$D_T = \sum_n \Lambda(n)^2 K_T(n, n).$$

The prime number theorem with partial summation gives $\sum_n \Lambda(n)^2 a_n(x)^2 \sim x \log x$. Hence $D_T \rightarrow 13/10$. The **centered off-diagonal remainder** is exactly

$$\begin{aligned} E_T &= 2 \sum_{m < n} \Lambda(m) \Lambda(n) K_T(m, n) \\ &\quad - 2 \sum_n \Lambda(n) \int_0^\infty K_T(n, v) dv + \int_0^\infty \int_0^\infty K_T(u, v) du dv. \end{aligned} \quad (12)$$

Equations (9) – (12) yield the promised exact research obligations:

$$\boxed{\text{RH implies: target (2) is equivalent to } E_T \rightarrow -3/10.} \quad (13)$$

$$\boxed{\text{RH and AH-Pairs imply } E_T \rightarrow -3/5.} \quad (14)$$

$$\boxed{\text{Under RH, proving } \liminf E_T > -3/5 \text{ suffices to refute AH-Pairs.}} \quad (15)$$

The second and third terms of (12) are part of the required cancellation. The arithmetic task is not to bound the first, uncentered sum by its diagonal. Individual uncentered terms can be much larger than the final normalized answer.

5. What known distribution tools do and do not supply

Writing $n = m + d$, the oscillatory factor in (10) changes appreciably around $d \asymp m/T$. Since the prime weight emphasizes m comparable to $x = T^\alpha$, the natural short-interval length is $H = x/T = x^{1-1/\alpha}$. For this test its exponents lie between $1/6$ and $2/7$. This is a description of the principal scale, not an assertion that all larger shifts or tails may be discarded.

Goldston – Montgomery's theorem gives the established bridge to short-interval variance. Goldston's Theorem 7 states the localized multiplicative-interval version, including logarithmic padding of the ranges; his equation (9.3) gives the fixed-length version. Under RH, a sufficiently uniform asymptotic

$$\int_1^X (\psi(y+H) - \psi(y-H))^2 dy \sim HX \log(X/H) \quad (16)$$

in the corresponding padded range implies the plateau target. It is not legitimate to infer a localized equivalence from one isolated (X, H) scale with no uniformity. Formula (12) is the sharper single-test obligation and avoids adding that unnecessary claim.

A standard sufficient, stronger hypothesis is the square-root-error Hardy – Littlewood prime-pair estimate

$$\sum_{n \leq X} \Lambda(n) \Lambda(n+d) = S(d)X + O_{\varepsilon}(X^{1/2+\varepsilon}) \quad (17)$$

uniformly in the required shifts and partial-summation ranges. The weighted singular-series sum produces the $-HX \log H$ correction to the diagonal variance. Summing errors with triangular weights costs $O(H^2 X^{1/2+\varepsilon})$, which is negligible against $HX \log X$ for $H \leq X^{1/2-2\varepsilon}$. The selected exponent range lies inside that regime. This is a sufficient conditional explanation, not a proof of (17); the averaged signed condition (12) is weaker than demanding (17) for every shift.

The checked existing tools leave the following precise gap:

1. The Montgomery – Vaughan mean-value error for this Dirichlet series is controlled by $\sum n \Lambda(n)^2 a_n(x)^2 \asymp x^2 \log x$. After normalization its size is comparable to x/T , which grows as a power of T here. It does not provide an $O(1)$ error above support one.
2. Distribution of primes in arithmetic progressions, including complementary-factorization and triply divisible-modulus inputs in the 186 proof, controls a different family of linear prime sums. No identity in that proof supplies the centered two-prime covariance (12). Replacing its supported-modulus conditions by a scalar exponent does not repair this missing estimate.
3. Almost-everywhere prime existence or a first-order short-interval prime asymptotic need not resolve the second moment at the precision $HX \log(X/H)$: the much larger scale $H^2 X$ is insufficient.
4. The GRH lower bound $F_T(\alpha) \geq 3/2 - \alpha - \varepsilon$ yields approximately 0.2 for this bump, below the AH value 0.7. The long-interval CMR bounds are also compatible with AH: the long average of (7) is p_0 , and the full allowed interval $[1, 1.2973576 \dots]$ lies inside their RH bounds. These results do not close (15).

There is also a rigorous information obstruction: low-band pair data and the protected moment algebra admit the half-lattice alternative while predicting the different value 0.7 here. Therefore an argument using only those statistical constraints cannot force the value 1. This does not prove that all future analytic methods must fail; it specifies that a new arithmetic estimate, or genuinely additional information, is needed.

6. Reproducibility, independent challenge, and stopping point

`sources/download_manifest.json` records seven downloaded primary PDFs with SHA256 hashes. It also records that the Anthropic expert-note URL was readable through the web tool but its separate local download returned HTTP 403; the arXiv paper was retrieved instead.

`kernel_identity_check.py` verifies the continuous mean in (8) and the diagonal/off-diagonal/mean expansion for a deliberately truncated prime polynomial at $x = 8$, $T = 3$, $n \leq 64$, using 60-digit arithmetic. The direct and expanded integrals agree within 3.0×10^{-59} . Omitting the mean changes that finite answer by about 212.37, illustrating the sign/centering issue. These are numerical normalization checks, not interval enclosures or evidence

for an asymptotic prime-pair conjecture. The initial unsplit improper quadrature did not meet its tolerance; splitting the logarithmic integration range resolved that numerical issue. The bundled document Python lacked `mpmath`; the recorded run uses the existing Homebrew Python with `mpmath` 1.3.0.

The AH support/periodicity step was independently challenged by the heat-flow agent. The root agent's separate Poisson-resolvent transfer was independently reviewed in `POISSON_TRANSFER_REVIEW.md`. No old finite-model computation was rerun and no new zero-data fit is being presented as a theorem.

The outcome is an explicit, weaker-than-full-PCC arithmetic target whose success would contradict a precise famous alternative under RH. The current round closes the normalization and source-audit questions; it does not close the signed prime-correlation estimate (15). The next justified research step is to attack that estimate, or an equivalent two-scale logarithmic-derivative statistic, with additional arithmetic structure. Reoptimizing a low-band detector alone would repeat the already identified information barrier.

Current report 13: Forward true–zeta localization: a contractive comparison and a sharp universality obstruction

Collection: R7 — actual-zeta targets, arithmetic mark, and flow obstruction.

Source: [research/dyson/round7/true-zeta-flow/FORWARD_FLOW_OBSTRUCTION.md](https://research.dyson/round7/true-zeta-flow/FORWARD_FLOW_OBSTRUCTION.md).

SHA-256: `33431c7878caae9fa30af2cd09b993f9040493e447311d67bb0b393b2e8ce7da` . Git blob: `0e1b29fd3c73f4148883a194aafd4409a070b08f` . Original bytes: 19946.

Forward true–zeta localization: a contractive comparison and a sharp universality obstruction

Status: ordinary written proofs of the comparison and model obstruction below. They identify an explicit missing boundary estimate and an independent stochastic-comparison obligation. No new Montgomery theorem, refutation of actual-zeta AH, or RH result is claimed.

The substantive conclusion is that positive time removes the negative-time nonreal-background problem, but does not turn the deterministic zero dynamics into finite-temperature Dyson Brownian motion. Even perfect local control of the remote field is compatible with a persistent half-spacing hard core. An exact two-periodic family demonstrates this, and a protected Fourier observable shows that the omitted diffusion acts at order one on the microscopic time scale.

1. Sources and conventions

The primary sources checked were:

- Rodgers – Tao, *The de Bruijn – Newman constant is non-negative*, version 5, especially the H_t convention, Theorem 11/equation (56), and the discussion of local clock equilibrium. The retained local text is `research-round2/rodgers-cao-1801.05914.txt`. Their displayed theorem is stated on $\Lambda_{t \leq 0}$ under the contradiction assumption; it must not be cited as a newly established uniform positive-time zero-count theorem.
- Erdős – Schlein – Yau – Yin, *The local relaxation flow approach to universality of the local statistics for random matrices*, particularly the stochastic evolution and its diffusion term, equation (5.5).
- Landon – Sosoe – Yau, *Fixed energy universality for Dyson Brownian motion*: the initial data can be deterministic, but the subsequent dynamics has Brownian noise. The density hypotheses and time scales in this theorem do not eliminate that noise.

Earlier internal reports checked were `yau_flow.md`, `galilean-proof-audit.md`, `dynamic_generator.md` and `new_attachment_bridge_audit.md`. This note retains their warning that finite-window scalar heat is not automatically the restriction of the global H_t , and adds a forward-time comparison with its precise boundary limitation.

The genuine convention is

$$H_t(z) = \int_0^{\infty} e^{tu^2} \Phi(u) \cos(zu) du, \quad \partial_t H_t = -\partial_z^2 H_t.$$

Increasing t is repulsive zero motion. Under RH, all zeros are real for $t \geq 0$. Work on a compact positive-time interval where the tracked zeros are simple, or begin at $t=0$ with the needed simple branches. No negative-time all-real assumption is imported. The standard H -coordinate zero is twice the usual zeta ordinate.

At a simple zero x_i , implicit differentiation gives $x_i' = H_t''(x_i)/H_t'(x_i)$. The even canonical product supplies the principal-value identity

$$x_i' = 2 \text{PV} \sum_{j \neq i} \frac{1}{x_i - x_j}. \quad (1)$$

The principal value is the one associated with the even global product. Splitting off finitely many terms does not permit changing its renormalization. On an interval free of external zeros the derivative of the external field is absolutely convergent and nonpositive.

2. Gap-independent contraction for a correctly localized system

Proposition 1. Let $x_1 < \dots < x_n$ and $y_1 < \dots < y_n$ be C^1 real trajectories on $[0, S]$. Suppose

$$\dot{x}_i = 2 \sum_{j \neq i} \frac{1}{x_i - x_j} + F(s, x_i), \quad \dot{y}_i = 2 \sum_{j \neq i} \frac{1}{y_i - y_j} + F(s, y_i) + e_i(s), \quad (2)$$

where F is C^1 in its spatial argument on an interval containing both configurations and $F_x \leq 0$. Then

$$\|x(s) - y(s)\|_{\infty} \leq \|x(0) - y(0)\|_{\infty} + \int_0^s \|e(u)\|_{\infty} du. \quad (3)$$

The constant is independent of the smallest internal gap. The order and absence of collisions are hypotheses on the comparison interval, not conclusions obtained by ignoring a collision.

Proof. Put $w_i = x_i - y_i$. The internal difference is exactly

$$2 \sum_{j \neq i} \left(\frac{1}{x_i - x_j} - \frac{1}{y_i - y_j} \right) = - \sum_{j \neq i} a_{ij}(w_i - w_j), \quad a_{ij} = \frac{2}{(x_i - x_j)(y_i - y_j)} > 0. \quad (4)$$

The external difference equals $b_i w_i$, with

$$b_i = \int_0^1 F_x(s, y_i + vw_i) dv \leq 0.$$

At an index realizing a positive maximum of w , the internal and external terms are nonpositive. At a negative minimum, they are nonnegative. Upper Dini derivatives therefore give (3). This is a maximum principle for a cooperative graph Laplacian, not an estimate by the absolute value of its potentially very large derivative.

If the true external field is F_{true} and the surrogate uses F_{approx} , take $e_i = F_{\text{approx}}(s, y_i) - F_{\text{true}}(s, y_i)$. The estimate requires its bound at **every retained particle**.

There is also an exact spatially resolved form. Let $U(s,u)$ be the positive, substochastic propagator generated by the matrix with off-diagonals a_{ij} and diagonal entries $-\sum_j a_{ij} + b_i$. Then

$$|w_j(s)| \leq \sum_j U_{jj}(s,0) |w_j(0)| + \int_0^s \sum_j U_{jj}(s,u) |e_j(u)| du. \quad (5)$$

This follows from variation of constants and positivity. It records exactly how a possibly poor boundary approximation propagates into a central window.

3. What remote-zero counting controls, and what it does not

In microscopic coordinates, fix a center and write the remote zero offsets as d_j with $|d_j| \geq L$. For $|z| \leq R \leq L/2$ set

$$F_{\text{far}}(z) = 2 \text{PV} \sum_{|d_j| \geq L} \frac{1}{z - d_j}, \quad B_p = \sum_{|d_j| \geq L} |d_j|^{-p}.$$

The following estimates follow term by term; only the constant term needs principal-value interpretation:

$$F'_{\text{far}}(z) = -2 \sum_{|d_j| \geq L} (z - d_j)^{-2} \leq 0, \quad (6)$$

$$|F_{\text{far}}(z) - F_{\text{far}}(0)| \leq 4RB_2, \quad (7)$$

$$|F_{\text{far}}(z) - F_{\text{far}}(0) + 2B_2 z| \leq 4R^2 B_3. \quad (8)$$

For (8), expand $1/(z-d) = -1/d - z/d^2 - z^2/[d^2(d-z)]$. The error is at most $2z^2/|d|^3$ per reciprocal, and the zero ODE contributes its factor 2.

If the remote counting function obeys $\#\{j: L \leq |d_j| \leq r\} \leq Ar + B$ for every $r \geq L$, integration by parts gives

$$B_p \leq \frac{pA}{p-1} L^{1-p} + BL^{-p} \quad (p > 1). \quad (9)$$

At $t=0$ for the true H-function, normalize a large H-height X by

$$\rho_X = \frac{1}{4\pi} \log \frac{X}{4\pi}, \quad q_j = \rho_X(x_j - X), \quad s = \rho_X^2 t. \quad (10)$$

The Riemann – von Mangoldt counting formula gives, for $p=2,3$ and $1 \leq L \leq X \rho_X/2$,

$$\sum_{|q_j| \geq L} |q_j|^{-p} \ll L^{1-p} + \frac{\log X}{L^p}. \quad (11)$$

One way to check (11) is to use the local bound $O(r + \log X)$ on dyadic shells up to distance $X \rho_X/2$. Beyond that distance, the global $O(Y \log Y)$ zero-count bound gives a convergent geometric tail of order $(X \rho_X)^{-(1-p)}$. This is absorbed in the first displayed term. No uniform linear-count bound at all arbitrarily large normalized distances is being assumed.

Thus ordinary zero counting already makes a sufficiently remote field nearly a common translation in a fixed central window. For example $L = (\log X)^2$ makes the right sides in (7) – (8) tend to zero for fixed R . The translation itself need not be small and must be retained or removed by a moving center. The curvature term is a scalar linear drift, not

independent Brownian forcing of neighboring particles.

Boundary limitation. Suppose the retained block consists of all zeros out to distance L . Its edge particles are close to the first omitted zeros, not a distance comparable to L from them. Estimate (8) is small in the central core, but generally is not small on every retained particle. Dense zeta configurations do not supply an empty annulus that fixes this problem. It is invalid to substitute a central-core Taylor estimate into the whole-block norm in (3).

For a proposed finite-window comparison, a precise remaining estimate is

$$\sup_{i \in I_{\text{core}}} \int_0^{S_X} \sum_{j \in I_X} U_{ij}(S_X, u) |F_{\text{approx}}(u, y_j) - F_{\text{true}}(u, y_j)| du = \alpha(1), \quad (\text{B})$$

together with the corresponding initial-error term in (5). All terms refer to the actual chosen moving window, microscopic normalization, external principal value and coupled trajectories. This is a boundary-propagation obligation, not just a static inverse-square-tail estimate. A short-range comparison or a killed-propagator estimate could establish it, but it has not been established for the true-zeta windows in this report.

To use (11) during an interval of true H_t flow, one must also establish the relevant uniform zero-count estimate on that shrinking positive-time interval. The static formula at $t=0$ alone does not prove it. Neither this temporal issue nor (B) can be silently supplied by citing a DBM theorem.

4. An exact half-grid obstruction even with excellent external-field control

Proposition 2. There exist deterministic circular heat families whose normalized local counting discrepancy is uniformly at most 2, whose initial configurations belong to the ACUE half-grid support up to rotation, and whose normalized nearest-neighbor gaps are at least $1/2$ for every forward microscopic time. Their local pair statistics never equal the sine-kernel pair law.

Construction and proof. Let $N=2M$ and

$$P_0(z) = z^{2M} - 2cz^M + 1, \quad c = \cos(\pi/4) = 1/\sqrt{2}. \quad (12)$$

For the repulsive circular coefficient flow

$$\partial_t P = z^2 P_{zz} - (N-1)zP_z,$$

the coefficient of z^k is multiplied by $\exp(-k(N-k)t)$. Therefore

$$P_t(z) = z^{2M} - 2ce^{-M^2 t} z^M + 1. \quad (13)$$

Put $a(t) = \arccos(c e^{-M^2 t})$. Its roots have arguments $(2\pi j \pm a(t))/M$. At time zero, rotation by $\pi/(4M)$ places them at grid indices $4j$ and $4j+1$ on the $4M=2N$ grid. They are consequently a valid ACUE-support configuration; this is not an assertion that the family is typical under ACUE's random law.

In normalized position $q=N\theta/(2\pi)$, consecutive gaps alternate between

$$g_-(s) = \frac{2}{\pi} \arccos(c e^{-\pi^2 s}), \quad g_+(s) = 2 - g_-(s), \quad s = \frac{N^2 t}{4\pi^2}. \quad (14)$$

For all $s \geq 0$, $1/2 \leq g_-(s) \leq 1 \leq g_+(s) \leq 3/2$. The limit as s tends to infinity is the unit clock. The periodically unfolded configuration is a union of two arithmetic progressions of spacing 2. Each progression's count differs from half the interval length by less than one, so the total counting discrepancy is at most 2 in every interval, uniformly in time. Its far inverse-power tails therefore satisfy (9) with absolute constants and have exactly the sort of small central external variation that motivates the localization proposal.

Nevertheless no pair has positive separation in $(0, 1/2)$, at any time. Let ϕ be a nonnegative smooth function of integral one supported in $(1/4, 1/3)$. The normalized pair statistic tested against ϕ vanishes identically for this family, while for the unit-density sine-kernel process it equals

$$\int \phi(u) \left[1 - \left(\frac{\sin \pi u}{\pi u} \right)^2 \right] du \geq 1 - \frac{8}{\pi^2} > 0. \quad (15)$$

The inequality uses monotonic decrease of $\sin(x)/x$ on the relevant positive interval. Randomly translating the periodic family makes it stationary without changing this discrepancy. The argument does not depend on a doubtful expectation of a minimum gap or on a rare collision event.

The exact clock $c=0$ is an even simpler stationary obstruction. The two-periodic family is useful because it starts with the characteristic half-grid gap pair $1/2, 3/2$ and evolves nontrivially. It shows that a successful deterministic positive-time localization theorem, even with strong density control, cannot by itself imply finite-beta Montgomery/GUE statistics.

This family has no claim to satisfy the arithmetic explicit formula or all known zeta correlations. It refutes only an inference whose hypotheses consist of real repulsive flow plus local counting/remote-field control. Actual arithmetic input must distinguish it.

5. The missing diffusion is visible at the microscopic generator scale

Use radian circle coordinates and write

$$V_i = \sum_{j \neq i} \cot \frac{\theta_i - \theta_j}{2}, \quad L = \sum_i V_i \partial_{\theta_i}.$$

With this drift normalization, circular beta-DBM has generator

$$G_\beta = L + \frac{2}{\beta} \Delta. \quad (16)$$

Indeed its diffusion coefficient is $\sqrt{4/\beta}$, and the score of the circular beta density is $(\beta/2)V$. For $\beta=2$, the generator is $L+\Delta$ and CUE is stationary. A change to the common $1/N$ drift convention rescales both time and noise, not just the drift.

For $p_m = \sum_j \exp(im \theta_j)$ and $F_m = |p_m|^2/N$,

$$\Delta F_m = 2m^2(1 - F_m). \quad (17)$$

CUE has $E F_m = m/N$ for $1 \leq m \leq N$. Its stationarity under $L+\Delta$ yields

$$E_{\text{CUE}} L F_m = -2m^2(1 - m/N). \quad (18)$$

In microscopic time $s=N^2t/(4\pi^2)$, the expected initial contribution of the omitted diffusion is

$$8\pi^2(m/N)^2(1 - m/N). \quad (19)$$

For even N and $m=N/2$, this is exactly π^2 , whereas the deterministic CUE initial derivative is $-\pi^2$. These are exact finite- N generator identities. No uniform finite-time expansion is inferred solely from a derivative at zero. They already disprove an identification of the deterministic and stochastic generators or a claim that the Brownian term is a negligible microscopic perturbation.

The clock gives another check: for $1 \leq m < N$, $F_m=0$ and $LF_m=0$ pointwise, whereas $\Delta F_m=2m^2$. The deterministic clock is fixed; $\beta=2$ DBM immediately creates fluctuations.

6. Protected moments remain blind even after adding DBM

Let W_N be the finite vector space spanned by symmetric trace monomials whose total positive Fourier weight and total negative Fourier weight are separately at most N . The previous dynamic-generator audit proved that L preserves each such weight block and that ACUE and CUE have the same initial expectation on W_N .

The Laplacian also preserves this filtration. For any signed integers m, n ,

$$\Delta(p_m p_n) = -(m^2 + n^2)p_m p_n - 2mnp_{m+n}, \quad p_0 = N. \quad (20)$$

For products of more traces, apply the product rule to each pair of factors. Same-sign merging preserves the positive/negative weights; opposite-sign merging reduces both by $\min(|m|, |n|)$. Thus $L+\Delta$ acts on the same finite space W_N .

The expectation vector of every finite-time $\beta=2$ DBM evolution consequently solves a closed finite linear ODE on W_N . Since the initial ACUE and CUE vectors agree, their expectations agree there for every $t \geq 0$. CUE stationarity makes their common value the original CUE value:

$$E_{\text{ACUE}}[F(\Theta_t)] = E_{\text{CUE}} F \quad (F \in W_N, t \geq 0). \quad (21)$$

This proof uses the existing exact ACUE/CUE initial Gram identity, plus the explicitly checked generator closure. It is not a new assertion that all observables are protected. Local sub-half-gap indicators lie outside this polynomial sector. At finite N , positive-time nondegenerate DBM can reach an open collision-free set with a sub-half gap. One may choose an ordered continuous path to such a set staying a positive distance from collisions; on a tube around this path the drift is smooth and bounded and the noise is nondegenerate. The usual Brownian support argument then gives positive probability. Thus sub-half gaps can be created from ACUE-supported initial data while retaining all identities (21). This qualitative finite- N statement does not assert a uniform microscopic-time lower probability as N grows.

It follows that adding stochastic smoothing, invoking a DBM universality theorem, and then pointing to unchanged protected moments cannot justify undoing the smoothing. The omitted observables are exactly where the half-grid alternative can change.

7. An explicit stochastic–removal obligation, not an automatic transfer

For a finite circular model let $D_t = \exp(tL)$ and $P_t = \exp(t(L + \Delta))$. On the invariant trace-polynomial spaces, differentiation gives the exact Duhamel identity

$$P_t F - D_t F = \int_0^t P_u \Delta D_{t-u} F \, du. \quad (22)$$

For more singular gap or resolvent tests, establish a regularized domain and the corresponding bounds before using the same formula. A proposed comparison of a height-averaged zeta ensemble to a noise-regularized ensemble must control the analogue of the right side for the **specific statistics outside the protected sector**. It cannot discard Δ because the external field has a small derivative or because the original data are deterministic.

The natural missing estimate is thus not another low-band moment match. It is a quantitative bound on stochastic smoothing/removal, or a direct arithmetic evaluation of a discriminating statistic such as the test in (15), with an error smaller than its explicit sine-versus-half-grid discrepancy. Proving that estimate is substantial new arithmetic work. Treating it as an assumption inside a proposed universality argument would assume the unresolved comparison rather than solve it.

For true H_t , remote-field randomness may produce a random common translation; (7) shows why its effect on a fixed core becomes almost common when the cutoff is remote. Such a translation does not provide independent relative-particle diffusion. This observation does not exclude every possible homogenization mechanism from a more elaborate arithmetic environment, but any such mechanism would need its own covariance and error theorem. The present localization argument supplies none.

8. Failed arguments retained for the next researcher

1. **Discarding remote zeros after a Taylor estimate at the center:** fails at the retained block's boundary; the error must be propagated as in (B).
2. **Using static Riemann – von Mangoldt bounds throughout a positive-time interval:** requires an additional uniform H_t counting estimate with its stated time range.
3. **Replacing deterministic heat by DBM once the drift is similar:** fails by the exact microscopic generator discrepancy (19).
4. **Expecting deterministic local relaxation to select sine₂:** fails for the explicit alternating-half-grid family, which relaxes to a clock and preserves a half hard core.
5. **Using protected-moment equality to undo DBM smoothing:** fails because that entire filtration stays equal throughout the smoothing, while the unprotected gap support changes.
6. **Inferring typical-ACUE or actual-zeta conclusions from the constructed family:** would confuse a pathwise counterexample to insufficient hypotheses with an ensemble or arithmetic theorem. No such conclusion is made.

The rigorous gain from this round is a correctly normalized forward comparison framework and a sharp obstruction to the proposed transfer. The arithmetic programme should now target a discriminating observable and the actual prime-side accuracy it requires, rather than seek more drift-only analogies.

9. Bounded checks and reproducibility

The standalone `forward_flow_checks.py` uses only the Python standard library and writes `forward_flow_checks.json`. It performs no true-zeta heat simulation. Completed checks were:

- 200 exact rational tests of the interaction difference identity and its maximum-principle sign, with nonuniform gaps and a decreasing external field.
- 81 exact rational tests of the remote-field constant and affine remainder bounds.
- Float64 evaluation of the explicit two-periodic solution at $N=4,8,16,32,64$ and five microscopic times, compared with the full cotangent force. The maximum relative discrepancy was below 4.30×10^{-13} . This is a numerical normalization check accompanying the explicit proof, not its replacement.
- Exact rational microscopic-generator constants for $m=N/2$ and 576 integer checks of the signed Fourier-weight filtration.

For example the alternating small gap is 0.5 at time zero, approximately 0.557329 at microscopic time 0.01, 0.830217 at time 0.1, and 0.999977 at time 1. It approaches a clock rather than creating the small gaps required by the sine-kernel pair law. The explicit lower discrepancy in (15) is approximately 0.18943053.

Run the checks with:

```
python3 forward_flow_checks.py
```

The script's empirical parts are labelled float64; the other checks use exact fractions or integers. No Claude call, paid model API, repository operation or PDF work was performed in this round.

Current report 14: Arithmetic transfer for a fixed large–prime sector

Collection: R7 — actual-zeta targets, arithmetic mark, and flow obstruction.

Source: `research/dyson/round7/arithmetic-resonator/DERIVATION.md`.

SHA-256: `3ab12c7227f0c41aaf78aab733f0353c054973cfe599f7b9275e47b12e8ee1f5` . **Git blob:** `0bbf469ec5685bfb4972908b58a2238b22ef4c91` . **Original bytes:** 9456.

Arithmetic transfer for a fixed large–prime sector

Status: ordinary proof extension of the fixed symmetric-prime transfer argument, plus deterministic numerical implementation. The fixed-family transfer has received a separate [independent internal review](#), including its truncations and insertion terms. This is not formal verification or external peer review. Its numerical experiment remains negative.

1. Fixed family and normalization

Fix $\ell \geq 1$, $\alpha = \ell^2$, and the multiplicative coefficients

$$d_\ell(p^\alpha) = (\ell)_\alpha / \alpha!.$$

For $n \leq L$ set

$$\begin{aligned} v_n &= \log(n) / \log(L), \\ S_k(n) &= \sum_{p|n, \text{ distinct}} (\log(p) / \log(L))^k, \quad k \geq 2, \\ C_L(n) &= \sum_{p|n} 1_{(p > \sqrt{L})}. \end{aligned}$$

Since $n \leq L$, $C_L(n)$ is exactly zero or one. A repeated prime $p > \sqrt{L}$ cannot divide n . Thus $C^2 = C$ without approximation. We allow a fixed polynomial

$$H(v, S, C) = F(v, S) + C J(v, S),$$

where S denotes finitely many S_k . Fixed radial Legendre polynomials are included in this class. All variables satisfy $0 \leq v \leq 1$, $0 \leq S_k \leq 1$ and $C \in \{0, 1\}$, so H is uniformly bounded. Its coefficients do not vary with L .

The resonator coefficients and creation matrix are

$$\begin{aligned} r_L(n) &= d_\ell(n) H(v_n, S(n), C_L(n)), \\ x_n &= r_L(n) / \sqrt{n}, \\ (A_L)_{(p^\alpha m, m)} &= 2 \sin(\pi \alpha \log(p) / (2 \log(L))) / (\alpha \sqrt{p^\alpha}). \end{aligned}$$

The normalized half-gap main term is

$$Q_L = [||A_L x||^2 + x^T A_L^2 x] / [2 \pi^2 ||x||^2] - 1/4.$$

The target is a *positive limiting value*, not a finite- L value or a positive coefficient in one summand.

2. General unmarked moments

For a list of labeled positive integers $I=(k_1, \dots, k_j)$, write $K=\sum I$. The existing marked-prime expansion gives

$$E_v \text{ product}_i S_{(k_i)} = v^K m_I(a),$$

where

$$m_I(a) = [\text{sum}_{(\text{set partitions } \pi \text{ of } \{1, \dots, j\})} a^{(\text{number of blocks})} \text{product}_{(B \text{ in } \pi)} \Gamma(\sum_{(i \text{ in } B)} k_i)] / (a)_K, \\ m_{\text{empty}}(a)=1.$$

This follows directly by grouping equal marked primes in the product and integrating the distinct-prime logarithmic sizes. It is not an assumption of an asymptotic Poisson-Dirichlet model. The same formula may of course be interpreted using that probability distribution.

For completeness, the unmarked positive measure is

$$(\log L)^{-a} \sum_{(n \geq 1)} d_{\text{ell}}(n)^2 / n \delta_{(\log(n)/\log(L))},$$

on all $n \geq 1$. Its Laplace transform tends to $C_{\text{ell}} t^{-a}$, giving local weak convergence to

$$C_{\text{ell}} / \Gamma(a) * v^{(a-1)} dv.$$

The full measure is first treated on compact logarithmic intervals, and then restricted to $v \leq 1$; it is not truncated at $n=L$ before taking its Laplace transform. The independent review gives the explicit short-background limiting procedure.

To derive the displayed product moment, restrict every marked prime to $p \geq L^\epsilon$, apply the prime number theorem to the finitely many reciprocal-prime measures, and use the unmarked weak limit for the remaining factor. Background collisions and coincident distinct marks have an extra reciprocal prime and vanish. At $n \leq L$,

$$\sum_{(p|n, p < L^\epsilon)} (\log(p)/\log(L))^k \leq \epsilon^{(k-1)},$$

so deleting the small marked primes changes any fixed polynomial in the S_k by $O_H(\epsilon)$. First let L tend to infinity with ϵ fixed, then let ϵ tend to zero. The joint limiting densities assign zero mass to simplex cutoff boundaries. This is the same finite-mark argument as the previously reviewed S_2 transfer, now with a finite general list of $k \geq 2$; it needs no uniform pointwise asymptotic at each individual prime tuple.

3. The new large-prime marked moment: an exact integer decomposition

There is an especially simple exact starting identity. If $C_L(n)=1$, then uniquely

$$n=p \ m, \ p>\sqrt{L}, \ m\leq L/p<p.$$

Hence p and m are coprime, p occurs exactly once, and

$$\begin{aligned} d_{\text{ell}}(pm)^2 &= a \ d_{\text{ell}}(m)^2, \\ S_k(pm) &= S_k(m) + u_p^k. \end{aligned}$$

For every test function Φ on the displayed variables, exactly

$$\begin{aligned} \sum_{n\leq L} d_{\text{ell}}(n)^2/n * C_L(n) \ \Phi(v_n, S(n)) \\ = a \sum_{p>\sqrt{L}, p\leq L} 1/p \\ \sum_{m\leq L/p} d_{\text{ell}}(m)^2/m \\ \Phi(v_m + u_p, S(m) + (u_p^k)_k). \end{aligned}$$

There is no coprimality error in this decomposition: $m<p$ makes it automatic. The prime variable stays in the compact logarithmic interval $(1/2, 1]$, so its limiting measure is dt/t by the prime number theorem. Combine this with the unmarked and finitely marked background weak limits from Section 2. The boundaries $t=1/2$ and $t+w=1$ have zero limiting measure. The potentially short background is included in that weak convergence; it is not replaced by an unjustified uniform asymptotic in m .

At fixed total $v\leq 1$, the resulting moment is zero for $v\leq 1/2$. For $v>1/2$ it is

$$\begin{aligned} E_v[C \ \text{product}_{(i \text{ in } I)} S_{(k_i)}] \\ = a \ v^{(1-a)} \sum_{A \text{ subset of the labeled index set } I} m_{(I\setminus A)}(a) \\ \int_{1/2}^v t^{(\sum_{(i \text{ in } A)} k_i - 1)} \\ (v-t)^{(a-1+\sum_{(i \text{ not in } A)} k_i)} dt. \end{aligned}$$

Repeated entries of I are still labeled in this subset sum, so their binomial multiplicities are retained. Every product involving C^r , $r\geq 1$, uses this same moment because $C^2=C$. The formula therefore evaluates all norm and insertion moments of the fixed family $F+CJ$.

An independent internal reviewer checked this formula and uniqueness at the half threshold. The exact integer decomposition above provides its arithmetic justification independently of any probability-model interpretation.

4. Insertion rules and the two distinct diagonal mechanisms

Put $\chi(u)=1_{(u>1/2)}$. On the simplex $v+u+w\leq 1$, define

$$\begin{aligned} H_0 &= H(v, \quad S, \quad C), \\ H_u &= H(v+u, \quad S+(u^k)_k, \quad C+\chi(u)), \\ H_w &= H(v+w, \quad S+(w^k)_k, \quad C+\chi(w)), \\ H_{uw} &= H(v+u+w, \quad S+(u^k+w^k)_k, \quad C+\chi(u)+\chi(w)). \end{aligned}$$

At most one of the background and inserted primes can exceed the half threshold. These are precisely the changes of the distinct-prime statistic away from the already negligible operator/background coincidences.

The limiting forms are

```

I = integral_0^1 v^(a-1) E_v[H0^2] dv,

M2 = (2 ell^2/pi^2) integral_(v+u+w<=1)
      v^(a-1) sin(pi u/2)/u * sin(pi w/2)/w
      E_v[H0 Huw + Hu Hw] dv du dw,

M3 = (2/pi^2) integral_(v+u<=1)
      v^(a-1) sin^2(pi u/2)/u * E_v[H0^2] dv du.

```

With $I > 0$,

```
Q_L -> (M2+M3)/I - 1/4.
```

To check the normalization and indexing, first restrict A_L to prime multipliers $p \geq L^\epsilon$. Then

```

||A x||^2 = sum_(n<=L) 1/n sum_(p,q|n) alpha_p alpha_q
           r_L(n/p) r_L(n/q),

x^T A^2 x = sum_(mpq<=L) alpha_p alpha_q/(mpq)
           r_L(m) r_L(mpq),

```

with $\alpha_p = 2 \sin(\pi \log(p)/(2 \log L))$. The ordered distinct-prime terms give the two displayed $M2$ products, with no additional factor of two.

The $p=q$ terms in A^*A have $n=mp$ and contain $r_L(m)^2$. They survive and give $M3$ with **$H0$ squared**, not an inserted H_u . Conversely, $p \neq q$ in A^2 costs p^{-2} ; its retained-prime contribution vanishes. This distinction is unchanged by adding the large-prime mark and is implemented explicitly.

The uniform operator truncation estimates from the previous transfer proof apply to every vector x :

```

||A_L|| = 0_ell(1),
||A_(p<L^epsilon)|| = 0_ell(sqrt(epsilon)+1/sqrt(log L)),
||A_(prime powers e>=2)|| = 0_ell(1/sqrt(log L)).

```

They follow by the positive Schur weight $d_\ell(n)/\sqrt{n}$, the logarithmic-derivative divisor identity and submultiplicativity for $\ell \geq 1$. They do not require H to be nonnegative or continuous as a function of a mark. After these truncations, coincidences are negligible by the same reciprocal-square bounds and uniform boundedness of H . Every remaining discontinuity is at a fixed prime-size threshold with zero limiting mass. This completes the extension of the fixed-family arithmetic limit.

5. Interface with actual zeta zeros

The source is Inoue, [arXiv:2604.05733v1](#), Theorems 3 and 4. The paper's theorem assumes RH and permits arbitrary resonator coefficients subject to its product cutoff. Use $L = \text{floor}(T/(\log T)^2)$, so $\theta = \log L/\log T$ tends to one; the same Schur comparison controls replacement of θ by one in the sine kernel. No finite integer table is used as a substitute for that limit.

A fixed strictly positive limiting margin at $\phi=1/2$ would make a half-gap improvement a meaningful next consequence after handling the source's continuity and error terms. This trial's margin is negative, so there is no new zero-spacing theorem. Even a successful small-gap statement must preserve the distinction between zeros counted with multiplicity and a gap interval separated from zero; it cannot automatically refute every half-lattice formulation of the Alternative Hypothesis.

6. Numerical quadrature adapted to the discontinuity

The unmarked block is the ordinary smooth simplex integral. For every block with at least one C factor, the M2 integrand is supported only in three disjoint sectors:

1. $v>1/2$: the background can contain the unique large prime; $u,w<1/2$ automatically.
2. $u>1/2$: the inserted u-prime is large, and $v,w<1/2$.
3. $w>1/2$: the symmetric inserted-prime sector.

In the first sector the marked moment factors as $(v-1/2)^a$ times a smooth function, so a Jacobi rule absorbs that factor. The other two sectors use the usual $v^{(a-1)}$ Jacobi weight and smooth triangular substitutions. No quadrature cell crosses a step in χ . M3's marked block only uses the first, background sector.

This is why the order-20, order-28 and order-40 calculations are stable without a dense multidimensional mesh. Their agreement is a useful independent numerical check, not a rigorous error enclosure.

Current report 15: Round 7: an arithmetic large–prime sector for zeta's half–gap problem

Collection: R7 — actual-zeta targets, arithmetic mark, and flow obstruction.

Source: [research/dyson/round7/arithmetic-resonator/REPORT.md](https://research.dyson/round7/arithmetic-resonator/REPORT.md).

SHA-256: `a1ad13dac36a646c82c857c771fa4c9c9b0e42aa5daa9d4c964361aebf7e2983` . **Git blob:** `a3454c2ecb27cfd7c84699a55e8c1f9a97320ad4` . **Original bytes:** 11454.

Round 7: an arithmetic large–prime sector for zeta's half–gap problem

Result: a fixed, discontinuous arithmetic feature beyond every finite polynomial in the old power sums was added and tested. The enlarged resonator's limiting half-gap margin is numerically approximately -0.01465492379421 . The added feature improves the matched baseline by only $1.429\text{e-}8$, so this concrete route does not cross the half-gap threshold. Three quadrature orders agree to about $1.4\text{e-}15$, and direct evaluation of a frozen rational coefficient vector on actual integers is also negative.

The useful output is the explicit large-prime-sector arithmetic transfer, its mixed insertion formulas, a complete rational test vector, and a bounded negative decision. This is not a new theorem about zeta zeros or an impossibility theorem for the full resonance-correlation method.

1. Why this family was selected

The initial suggestion to add S_3 or S_2^2 was corrected after checking the existing archive. Round 1 had already tested these features, and a 48-term, twelve-group power-sum family had reached approximately -0.01465472564383 . Repeating that sweep would not be a new research direction. That old number remains slightly better than the new trial below; differences in the finite spans and ell values must not be concealed.

The genuinely new mark is

$$C_L(n) = 1_(\text{there is a prime } p|n \text{ with } p > \sqrt{L}).$$

For $n \leq L$ it is a binary mark. It cuts the prime-factor configurations into a large-prime sector and its complement. This sharp fixed threshold is not represented by a finite polynomial in $S_2, S_3, \dots$, although sufficiently large polynomial spaces could approximate it in a suitable mass norm.

The exact identity $n = pm$ with $p > \sqrt{L}$, $m < p$ gives a simple arithmetic proof route, rather than an unsupported identification with a limiting Fock operator. No physical prime-gap k_{39} matrices, heat-flow matrices, or model-only random-matrix distribution were used in this computation.

2. The precise 30–dimensional experiment

Fix $\ell=27/25$ and $\alpha=729/625$. Use the coefficient family

$$r_L(n)=d_{\ell}(n) \, H(v_n, S_2(n), S_3(n), C_L(n)),$$

with $v_n=\log(n)/\log(L)$ and the same distinct-prime power sums as the previous arithmetic transfer proof.

The matched unmarked space has twenty features:

$$\{1, S_2, S_3, S_2^2\} \text{ times } \text{Legendre}_d(2v-1), \quad 0 \leq d \leq 4.$$

The enlarged space adds ten features:

$$C \text{ times } \{1, S_2\} \text{ times } \text{Legendre}_d(4v-3), \quad 0 \leq d \leq 4.$$

The second radial basis is centered on the marked support $1/2 < v \leq 1$, avoiding needless conditioning loss from monomials on a short interval. This is only a basis choice: it does not restrict the five radial degrees available in either marked group.

All thirty coefficients are optimized together at the one fixed ℓ . There is no ℓ scan, threshold scan, degree scan, or new eigenvalue scan of the full integer operator. The basis and threshold remain fixed as L tends to infinity.

3. Arithmetic derivation and its review status

[DERIVATION.md](#) gives the ordinary proof extension in full. The main new identity is

$$\begin{aligned} E_v[C \text{ product } S_{(k_i)}] \\ &= a^{v(1-a)} \sum_{(A \text{ subset of labeled factors})} m_{(I \setminus A)}(a) \\ &\quad \int_{(1/2)}^v t^{(\sum_A k_i - 1)} \\ &\quad (v-t)^{(a-1+\sum_{(I \setminus A)} k_i)} dt, \end{aligned}$$

with zero value when $v \leq 1/2$. It follows from an exact unique-large-prime integer decomposition and the prime number theorem, together with the already established finite marked-prime moments. The background factor is automatically coprime to the large prime. A fixed threshold's boundary has zero limiting measure.

The insertion changes are $C \rightarrow C+1_{(u>1/2)}$. The off-diagonal M_2 terms use H_0 Huw and H_u Hw. The same-prime A^*A term survives as M_3 with H_0^2 ; it must not be given an inserted mark. The repeated-prime A^2 term vanishes after the usual truncations. The previous weighted Schur estimates are valid for arbitrary signed coefficient vectors and remove small-prime/prime-power operator pieces without assuming a positive H .

The full fixed-family transfer now has a separate [independent internal review](#), including the short background, fixed-threshold boundary and signed-vector truncations. This is not formal verification or external peer review. Numerical agreement is not used as a substitute for the proof review. The source interface remains [Inoue's RH-conditional resonance-correlation theorem](#), with its stated product cutoff and error terms.

4. Limiting–form numerical results

The quantity displayed is

$$\text{margin} = (M2+M3)/I - 1/4,$$

with the normalization fixed in DERIVATION.md. A positive number is required for the intended half-gap attack.

quadrature order	20-dimensional baseline	enlarged 30-dimensional family
20	-.0146549380840022	-.0146549237942085
28	-.0146549380840022	-.0146549237942086
40	-.0146549380840038	-.0146549237942099

The order-40 scaled mass-Gram condition numbers are approximately 5.17e7 for the baseline and 1.20e8 for the enlarged family. Every direction survives the stated relative eigenvalue cutoff. The enlarged pencil's residual norm is near 2e-16. These are diagnostics for floating matrices; they do not certify integration or eigenvalue errors.

The quadrature is specifically split at the mark discontinuities. Marked M2 blocks are supported only in three disjoint sectors: a background total greater than 1/2, an inserted u greater than 1/2, or an inserted w greater than 1/2. The background-sector Jacobi weight absorbs the exact $(v-1/2)^a$ endpoint factor. Consequently no integration cell crosses an unresolved step in the mark.

The enlarged value is still about .01465 below zero, vastly larger than the tiny observed gain. It is therefore a useful negative test of this particular sector feature, not evidence of a new record. Nor does it rule out other thresholds, richer occupation functions, or an altogether different resonator.

5. Complete frozen rational vector

The order-40 optimizing coefficients were rounded to denominator 100,000,000 and then evaluated again as one fixed vector. In each row below the five integers multiply the radial Legendre degrees 0,1,2,3,4. The first four rows use Legendre_d(2v-1); the last two use Legendre_d(4v-3). Divide every integer by 100,000,000.

factor	five coefficient numerators
1	-117846152, 38918251, 1078497, -33600, 411449
S2	46899554, -54583075, 43523183, -40218573, 12782032
S3	295982141, -383374116, 290083532, -105264757, 17704774
S2^2	-109987670, 186525942, -226829807, 150959435, -47011938
C	87192, -72400, 513184, -115799, -690461
C S2	-361789, 437850, -2512958, 557662, 2084511

Its quadrature norm is 1.0000000030188823 and margin is -0.014654923794209879. Rational coefficients do not make the integral evaluation an interval certificate. fixed_rational_vector.json stores the exact integers, basis labels and denominator so that no coefficient must be reconstructed from a screenshot or rounded prose.

6. Direct finite–integer check

The same frozen rational vector was evaluated on actual integers, using the full finite prime-power creation matrix

```
A_(p^e m,m)=2 sin(pi log(p^e)/(2log L))/(e sqrt(p^e)).
```

These are direct Rayleigh evaluations, not optimized eigenvalues. The distinct-prime mark is tested by the exact integer comparison $p^*p > L$. The divisor coefficients include all prime exponents, and the finite operator includes all prime-power multipliers.

```
| L | |Ax| |^2/||x| |^2 | x^T A^2 x/||x| |^2 | finite margin | |---|---|---|---| | 10,000 | 2.98343331577942 |
1.10733089333213 | -.04275946416555 | | 100,000 | 3.02158248945693 | 1.17479313763796 | -.03740912722744 | |
1,000,000 | 3.04668220745756 | 1.22147565861788 | -.03377259651844 |
```

The finite evaluations set $\theta = \log L / \log T$ to its limiting value one. They do not assert that a literal finite T already satisfies that equality under the source cutoff. Their slow approach toward a limiting form is not extrapolated to an asymptotic certificate. These data concern resonator coefficients and the arithmetic operator; they are not measurements of actual zeta zeros.

7. Independent numerical checks

`validate_sector.py` performs checks that use different formulas or representations:

- The unmarked Gram and numerator blocks are compared with the previous independent monomial implementation after an explicit Legendre-to-monomial change of basis. Maximum absolute discrepancies are about $1.92\text{e-}13$ and $9.07\text{e-}13$.
- At $\alpha=1$, the marked moments are checked against closed formulas $E_v C = \log(2v)$ and $E_v(C S_2) = .75(v^2 - .25) - v(v - .5) + .5v^2 \log(2v)$.
- For noninteger $\alpha=729/625$, direct integration in the marked-prime variable t is compared with the implemented scaled Jacobi formula, including repeated labeled factors S_2^2 and S_2^3 . Observed differences are at most about $1\text{e-}15$.
- The sampled marked moments satisfy $0 \leq E(C \text{ product } S_k) \leq E(\text{product } S_k)$, and three quadrature orders give an enlarged-margin spread of $1.36\text{e-}15$.
- The finite integer construction asserts C in $\{0,1\}$ for every integer evaluated.

The numerical integration routines' reported errors and matrix residuals are not outward enclosures. The report deliberately distinguishes the ordinary arithmetic derivation, the finite matrix model, the numerical tests, and the actual zeta theorem.

8. Files and reproduction

All new files are confined to this directory. No Git operation, Claude call, paid API, prime-gap matrix run, or PDF build was used.

```
OPENBLAS_NUM_THREADS=1 python3 large_prime_sector.py --order 20
OPENBLAS_NUM_THREADS=1 python3 large_prime_sector.py --order 28
OPENBLAS_NUM_THREADS=1 python3 large_prime_sector.py --order 40
```

```
OPENBLAS_NUM_THREADS=1 python3 finite_integer_check.py
OPENBLAS_NUM_THREADS=1 python3 validate_sector.py
```

`large_prime_sector.py` is self-contained apart from NumPy/SciPy. Each quadrature order has JSON containing every coefficient, basis label, Gram spectrum and diagnostic, plus an NPZ with the full numerator and mass matrices. `finite_integer_check.py` records the frozen rational vector and the actual integer values. `validate_sector.py` uses the older `general_prime_features.py` only for its independent unmarked normalization comparison; its SHA is pinned in `validation.json`. The validation script's input location is the existing common BASE layout.

`DERIVATION.md` supplies all additional marked moments and insertion rules. `manifest.json` pins the new code/data and environment. The normal public export is small; no large eigenvector or omitted private service is required.

9. Decision and remaining work

This fixed half-threshold sector is now a tested arithmetic family with a concrete transfer argument. Its tiny numerical improvement does not justify further tuning of the same degree-four coefficients. A materially different occupation structure or a new estimate connecting the resonator to arithmetic correlations remains necessary for a serious chance of crossing the half-gap threshold.

No implication for Montgomery-Dyson pair correlation or a refutation of general AH follows from this negative result. Even a future positive small-gap result would need care about multiplicity and the zero atom allowed in some AH formulations. The current run ends with this specific negative decision and preserves all formulas and coefficients for independent review.

Current report 16: Independent audit of the half–threshold arithmetic transfer

Collection: R7 — actual-zeta targets, arithmetic mark, and flow obstruction.

Source: `research/dyson/round7/arithmetic-resonator/INDEPENDENT_REVIEW.md`.

SHA-256: `b18387755f10dfdb1bb86e9fe755ecd778fb10b71f3b46ea7b6dec987264a3cb` . **Git blob:** `41815e70565aff8ee0f0a648bd68124ceb226f4f` . **Original bytes:** 15802.

Independent audit of the half–threshold arithmetic transfer

Date: 2026-09-05. **Reviewer:** the separate `yau_flow` research agent.

Verdict: the fixed-family arithmetic limit in `DERIVATION.md` is accepted as an ordinary mathematical argument, with the full-measure clarification and explicit limiting procedure below. I found no missing large-prime sector, extra factor of two, erroneous same-prime insertion, or unjustified extension of the Schur estimate to signed coefficients. This is independent internal review, not formal verification or external peer review. It does not certify the floating numerical optimum or prove a zero-spacing result.

I read the complete derivation, report, `large_prime_sector.py`, `finite_integer_check.py`, and the preceding `symmetric_prime_arithmetic_transfer.md`, including its explicit collision and operator-error estimates. I also checked the primary source's Theorems 3 and 4. The filename `threshold_model.py` in the review request is stale; the implementation in this checkpoint is `large_prime_sector.py`.

1. Precise scope of the accepted statement

Fix $\ell \geq 1$, $\mathfrak{a} = \ell^2$, and one real polynomial

$$H(\nu, S, C) = F(\nu, S) + CJ(\nu, S),$$

where only finitely many distinct-prime power sums $\mathcal{S}_k$, with $k \geq 2$, occur. The polynomial, its degree, its coefficients, and the threshold $1/2$ remain fixed as $L \rightarrow \infty$. With $I > 0$, the proposed arithmetic quotient converges to

$$\frac{M_2 + M_3}{I} - \frac{1}{4}.$$

The statement is unconditional as an assertion about weighted integer sums and the specified finite arithmetic matrices. RH enters only in its separate application to the zeta theorem. No uniformity in a growing basis, a moving threshold, ℓ , or optimizing coefficients depending on L has been proved. The frozen rational coefficient vector is within the accepted fixed-family scope.

2. Exact large–prime decomposition and mixed moments

For $n \leq L$, two distinct prime divisors greater than $\sqrt{L}$ are impossible, as is the square of such a prime. Thus

$$C_L(n) = \mathbf{1}_{P^+(n) > \sqrt{L}} \in \{0, 1\},$$

where $C_L(1) = 0$. If this mark is one, there is exactly one decomposition $n = pm$ with $p > \sqrt{L}$, and $m \leq L/p < p$. Consequently the coprimality in this decomposition is exact, not asymptotic:

$$d_\ell(pm)^2 = ad_\ell(m)^2, \quad S_k(pm) = S_k(m) + u_p^k.$$

The identity in DERIVATION § 3 therefore holds for every test function for which its finite sums are defined. No prime-density approximation is involved at this stage.

The unmarked moment for a labeled list $l = (k_1, \dots, k_j)$, of total weight K , is

$$m_l(a) = \frac{1}{(a)_K} \sum_{\pi} a^{|\pi|} \prod_{B \in \pi} \Gamma(\sum_{i \in B} k_i).$$

Each block denotes one distinct prime shared by precisely those marks. The factor $\Gamma(\sum_B k_i)$, rather than $\Gamma(1 + \sum_B k_i)$, is correct because the prime-size measure is du/u . There is no factorial for ordering blocks: each labeled set partition appears once. In particular,

$$m_{(2)} = \frac{1}{a+1}, \quad m_{(3)} = \frac{2}{(a+1)(a+2)}, \quad m_{(2,2)} = \frac{a+6}{(a+1)(a+2)(a+3)}.$$

Disintegrating the exact large-prime identity by total size v , and dividing by the unmarked density v^{a-1} , gives

$$\mathbb{E}_v[C \prod_{i \in l} S_{k_i}] = av^{1-a} \sum_{A \subseteq l} m_{\setminus A}(a) \int_{1/2}^v \tilde{t}^{\sum_{i \in A} k_i - 1} (v - t)^{a-1 + \sum_{i \in A} k_i} dt$$

for $v > 1/2$, and zero otherwise. The subset notation is over labeled positions. For example, the CS_2^2 integrand has exactly the three contributions

$$t^3(v-t)^{a-1} + 2m_{(2)}t(v-t)^{a+1} + m_{(2,2)}t^{-1}(v-t)^{a+3}.$$

The middle factor two is necessary and is retained by the implementation's bit-subset expansion. All higher powers of C reduce to the same marked moment. These formulas check for the mixed S_2, S_3 features and their products appearing in both numerator and norm.

3. Measure convergence, discontinuities, and short backgrounds

One wording clarification is needed in DERIVATION § 2: the measure whose Laplace transform tends to $C_\ell t^{-a}$ must be

$$v_L = (\log L)^{-a} \sum_{n \geq 1} \frac{d_\ell(n)^2}{n} \delta_{\log n / \log L},$$

with the sum over all positive integers. The measure truncated to $n \leq L$ does not have that Laplace transform. The preceding transfer proof correctly uses the full measure and then restricts to compact intervals, so this is a clarification of the new note's abbreviated notation, not a failure of its argument.

The full positive measures converge locally to

$$\nu(d\nu) = \frac{C_\ell}{\Gamma(a)} \nu^{a-1} d\nu.$$

The Laplace bound also gives uniformly bounded mass on $[0, 1]$. This supplies all total-mass control needed here. To make the new limit argument explicit:

1. Fix $\varepsilon > 0$, and restrict the finitely many ordinary marked primes and operator primes to $p \geq L^\varepsilon$. The threshold prime already lies in $(L^{1/2}, L]$.
2. Expand a fixed polynomial into finitely many labeled-prime terms. After discarding coincidences as below, each is a product of the background measure and finitely many reciprocal-prime measures on compact logarithmic intervals bounded away from zero. PNT and partial summation identify the latter as du/u .
3. Take the product-measure weak limit, restricted by the relevant simplex. Its limiting measure is absolutely continuous in the continuous prime sizes and background size. The hyperplanes $u = 1/2$, $w = 1/2$, the large-prime threshold $t = 1/2$, and each total-size cutoff have zero mass. Bounded piecewise polynomial test functions are therefore legitimate.
4. Finally let $\varepsilon \downarrow 0$. For $k \geq 2$, deleting small ordinary marks changes S_k by at most ε^{k-1} , uniformly for $n \leq L$. Fixed polynomial products change by $O_H(\varepsilon)$. The threshold mark is never deleted in this step.

This is not an appeal to pointwise uniform Selberg – Delange estimates at each prime tuple. For an explicit short-background check, after the prime cutoffs are fixed, the part with remaining background size $\nu \leq \eta$ is bounded by a fixed constant depending on H, ℓ, ε times $\nu_L([0, \eta])$. For each fixed $\eta > 0$,

$$\nu_L([0, \eta]) \rightarrow \frac{C_\ell}{\Gamma(a+1)} \eta^a.$$

Taking $L \rightarrow \infty$, then $\eta \downarrow 0$, controls this portion. The background atom $m = 1$ has normalized mass $(\log L)^{-a} \rightarrow 0$. The large-prime endpoint p near L causes no new atom or missing boundary term. These observations close the short-background and discontinuity obligations for this fixed family.

4. Coincidences and signed coefficient bounds

For $\ell \geq 1$, the decreasing ratios $(\ell + e)/(e + 1)$ give

$$d_\ell(p^{b+c}) \leq d_\ell(p^b) d_\ell(p^c).$$

In particular,

$$\sum_{m \leq L: p \mid m} \frac{d_\ell(m)^2}{m} \leq \frac{\ell^2}{p} \sum_{k \leq L} \frac{d_\ell(k)^2}{k}.$$

An external reciprocal-prime factor then makes an operator/background coincidence cost p^{-2} . At $p \geq L^\varepsilon$, the sum of these costs is $O(L^{-\varepsilon})$. Other retained prime measures have bounded mass for fixed ε using a weaker power of $\log \log L$ bound would also suffice. There are only finitely many coincidence patterns.

The jump of C does not invalidate this estimate. A coincidence can change an amplitude discontinuously, but both amplitudes are bounded by a constant times d_ℓ ; the contribution of the exceptional integer set still has the same vanishing upper bound. One does not need differentiability or positivity of H .

5. Operator normalization and the two diagonal mechanisms

For the retained prime operator, put $\alpha_p = 2 \sin(\pi u_p/2)$. Direct multiplication gives

$$(Ax)_n = \frac{1}{\sqrt{n}} \sum_{p \mid n} \alpha_p r(n/p).$$

Squaring this expression gives the stated ordered (p, q) sum for A^*A . Multiplying a second time gives the stated ordered sum for A^2 . For distinct primes coprime to the background, the coefficient factors are

$$r(mp)r(mq) = \ell^2 d(m)^2 H_u H_w, \quad r(m)r(mpq) = \ell^2 d(m)^2 H_0 H_{uw}.$$

The four sine factors from $\alpha_p \alpha_q$, divided by $2\pi^2$, give $2\ell^2/\pi^2$. Both discrete and continuous prime pairs are ordered, so another factor of two would be an error.

For A^*A , the same-prime term is exactly

$$\sum_{mp \leq L} \frac{\alpha_p^2}{mp} r(m)^2.$$

Its amplitude is evaluated at m . Thus M_3 uses H_0^2 , without adding the prime's mark, its power sums, or its size to H . There is no extra ℓ^2 here. In contrast, the same-prime term in A^2 contains p^{-2} and vanishes after the fixed lower prime cutoff. I independently checked this distinction directly from matrix multiplication; it is essential and correctly implemented.

The positive Schur weight $d_\ell(n)/\sqrt{n}$ from the preceding proof gives the uniform operator bounds used in the new derivation. For $A = A_0 + E$,

$$\|A^*A - A_0^*A_0\|, \|A^2 - A_0^2\| \leq (\|A\| + \|A_0\|) \|E\|.$$

Thus the errors $O_\ell(\sqrt{\varepsilon} + (\log L)^{-1/2})$ hold after division by $\|x\|^2$, for every nonzero signed vector. The denominator's asymptotic follows separately from $l > 0$. No positivity or positive-semidefinite assumption on the entire numerator is being used.

6. Insertion sectors and implementation audit

On $v + u + w \leq 1$, at most one of the background's large prime, the inserted u -prime, and the inserted w -prime can exceed the global half threshold. Insertion must use $\chi(u) = \mathbf{1}_{u > 1/2}$, not a threshold relative to the new total size. The code uses the correct global threshold.

For two feature factors with inserted indicators d_l, d_r , the code's marked product is the exact identity

$$(C + d_l)(C + d_r) = C(1 + d_l + d_r) + d_l d_r.$$

The `expand` routine retains labeled multiplicities, including repeated S_2 factors. The `partitions` routine may return equal-looking blocks when numerical labels repeat; their multiplicities correctly represent different partitions of the labeled positions.

For a matrix block containing at least one marked feature, the three integration domains are disjoint, up to zero-measure boundaries:

Domain	Source of the possible mark	Treatment in the code
$v > 1/2$	Background prime	Marked background moment; both inserted primes below half
$u > 1/2$	First inserted prime	Background mark zero; explicit insertion indicator
$w > 1/2$	Second inserted prime	Symmetric insertion sector

The first domain contains configurations with and without a background large prime; the marked moment, rather than the condition $v > 1/2$ alone, supplies the correct weight. The code respects this distinction. M_3 's marked blocks use only the first domain because its amplitude is uninserted.

Writing $\delta = v - 1/2$ and $t = v - \delta z$ in a marked moment extracts δ^a . The remaining integrand is smooth up to $\delta = 0$, since $t \geq 1/2$. Its background-size factor is z^{a-1} , with additional integer powers of z for unassigned marks. The scaled Jacobi implementation is correct. The outer Gram factor is $(v - 1/2)^a v^{a-1}$, and the marked M_2 background domain has this factor times $(1 - v)^2(1 - x)$. In an inserted-prime domain, the Jacobian is $(1/2 - v)^2(1 - x)$. These are exactly the weights used by `forms()`.

The sine-kernel replacements using NumPy's normalized `sinc` also have the correct constants. Symmetrizing the final real matrix preserves its quadratic form. The finite-integer implementation uses distinct prime factors for the marks, the full divisor coefficients for prime powers, and the exact integer condition $p^2 > L$. I found no disagreement between the formulas and the inspected implementation.

7. Primary-source interface and the limits of the result

Inoue, [arXiv:2604.05733v1](#), Theorems 3 and 4 allow arbitrary arithmetic resonator coefficients under RH, with $L \leq T/(\log T)^2$. Choosing the approximator equal to the logarithmic increment coefficients yields the specified two quadratic forms; at $\phi = 1/2$ the linear term cancels. The earlier Schur argument controls the remaining normalized errors and the replacement $\log L/\log T \rightarrow 1$. The new mark does not introduce a new source restriction.

This confirms the **interface**, not a favorable sign. A strictly positive limiting margin would still have to be exhibited, and a strict improvement below half would use continuity in ϕ for a fixed vector. The present negative numerical trial proves neither a new zeta theorem nor a refutation of AH. A statement about zeros counted with multiplicity cannot silently become a statement about positive pair distances bounded away from zero.

8. What the numerical evidence does and does not establish

The reported values at quadrature orders 20, 28, and 40 are consistent with the inspected implementation and each other. I inspected the stored validation data but did not rerun the integrations, finite- L million-integer evaluations, or any eigenvalue optimization. No new scan was performed.

Agreement of quadrature orders is not an outward enclosure. A small pencil residual does not bound integration error. The scaled Gram condition number near 1.2×10^8 must remain visible when interpreting extremely small improvements. The rational vector is genuinely fixed, but rational coefficients do not make its transcendental limiting integrals exact. Accordingly:

- Accepted: a numerically negative test of this particular 30-dimensional family, with all inputs preserved.
- Not established: a rigorous negative upper bound over the whole 30-dimensional space.
- Not established: a certified gain of 1.429×10^{-8} , despite its stable floating estimate.
- Rejected as an inference: a structural impossibility theorem for larger occupation families or the full resonance-correlation method.

These are scope boundaries, not objections to the authors' present negative decision.

9. Independent exact checks and pinned evidence

`independent_identity_checks.py` is a separate, standard-library-only check at the single fixed cutoff $L = 120$, with $\ell = 27/25$. It uses exact integer and `Fraction` arithmetic. Rational formal prime labels and kernels replace logarithms and sines only to check their algebraic roles; this is expressly not an asymptotic or numerical-margin test. It checked all 120 integers, the 70 unique-large-prime pairs, 132 coprime ordered distinct-prime insertion triples, both exact quadratic-form expansions, the surviving uninserted diagonal, and all eight Boolean product assignments. Every assertion passed. The results are in `independent_identity_checks.json`.

Run the bounded algebra check with:

```
python3 independent_identity_checks.py
```

The reviewed author-file SHA-256 values are:

File	SHA-256
DERIVATION.md	62970c91f2ff757eabc5d9a364d189fb7a42494c0f188713b2974003df4833b0
REPORT.md	439e08210143acf4adfae7918f8e8ca92e5ce6d66080ce8a8858cc62bbe35d06
large_prime_sector.py	255ad9a2f29e086eca01b3823bd9ece3ecbfa47b431259a6af3fee260c9afa8d
finite_integer_check.py	2126542d6b75822fb79e0910d43e4e5639be5455e079be3095c534c60a171295
validate_sector.py	72a286b477034a3f3aae63aeaba479d5de5640a1e13fe8043f37b3f918b07d26
fixed_rational_vector.json	1858e8c2ec1effa0f51e93ff28561057d1e1c5d8c9d9c08b50d099ab71e9943f
validation.json	cf8970966bcbdb9d65cde8043461466a9633c3ad74de624d2a84b1f7c2e0bc2a

No author file, source repository, or existing experiment output was modified. This review and its small algebra-check files are the only new outputs.

Current report 17: Round 8: isolate the actual arithmetic remainder

Collection: R8 — short-prime projection and signed residual.

Source: [research/reports/dyson_round8.md](#).

SHA-256: `111901c83af7481b414b9ef08728baf0fadbabc696a8e69f533b4a55c05a2cf6` . **Git blob:** `bcf9481b26ba287e04b0770c6842052f8e8dc851` . **Original bytes:** 6061.

Round 8: isolate the actual arithmetic remainder

The two-scale target from [Round 7](#) can be written as a computable short-prime main term plus one signed difference of residual energies. The residuals have an exact convergent representation using the same centered prime-counting error. This round proves that decomposition under RH and identifies why two tempting lower-bound arguments fail. It does not prove the required residual inequality or refute AH-Pairs.

1. A short-prime projection identity for actual zeta

Fix $c > 0$ and put

$$N = \lfloor \frac{T}{\log^6 T} \rfloor, \quad s_c(t) = \frac{1}{2} + \frac{c}{\log T} + it, \quad H_c(t) = -\frac{\zeta'}{\zeta}(s_c(t)),$$

$$P_c(t) = \sum_{n \leq N} \Lambda(n) n^{-s_c(t)}, \quad R_c(t) = H_c(t) - P_c(t).$$

The [ordinary analytic proof](#) establishes, under RH,

$$\int_0^T |H_c(t)|^2 dt = T \sum_{n \leq N} \frac{\Lambda(n)^2}{n^{1+2c/\log T}} + \|R_c\|_{L^2(0,T)}^2 + O_c(N \log^4 T). \quad (1)$$

The mixed product is evaluated by a contour shift to the absolutely convergent Dirichlet-series half-plane. The contour avoids the pole at one; RH controls the shrinking horizontal distance from the zeros. The infinite right-line off-diagonal sum, the top edge and the initially removed short interval all receive explicit bounds. A finite-polynomial mean-value estimate and completion of the square give (1). Exact finite-height orthogonality is not assumed.

An [independent audit](#) checks the contour orientation, diagonal coefficient, near-diagonal summation, infinite tail, T-dependent error and continuation signs. No novelty claim or proof-assistant certification accompanies this ordinary reduction.

2. The precise remaining inequality

For the Round 7 statistic

$$W_T = \frac{2}{T \log^2 T} [\sinh(2) \parallel H_1 \parallel_2^2 - \sinh(1) \parallel H_{1/2} \parallel_2^2],$$

(1) and PNT imply

$$W_T = B + E_T + o(1), \quad B = 2 \int_0^1 u [\sinh(2)e^{-2u} - \sinh(1)e^{-u}] du = 0.4560939793292317 \dots,$$

$$E_T = \frac{2}{T \log^2 T} [\sinh(2) \parallel R_1 \parallel_2^2 - \sinh(1) \parallel R_{1/2} \parallel_2^2].$$

Thus the sufficient AH-refutation target is precisely

$$\boxed{\liminf_{T \rightarrow \infty} E_T \geq \frac{1}{16} - B = -0.3935939793292317 \dots} \quad (2)$$

The positive short-prime main term does not prove (2). The residual combination is signed and of leading order. The sine prediction would make it about -0.3738225362077544, but that value has not been obtained for actual zeta.

3. The two residuals come from the same arithmetic function

Set $E(x) = \psi(x) - x$ with $\psi(N)$ including the atom at N . For $\text{Re}(s) > 1/2$, $s \neq 1$, RH gives the exact identity

$$-\frac{\zeta'}{\zeta}(s) - \sum_{n \leq N} \Lambda(n) n^{-s} = \frac{N^{1-s}}{s-1} - E(N) N^{-s} + s \int_N^\infty E(x) x^{-s-1} dx. \quad (3)$$

The integral converges absolutely under RH. The formula is obtained in the original convergence half-plane and then continued; no unregularized critical-strip prime series is used. The pole and endpoint subtraction have fixed signs.

At the chosen cutoff, the pole term can be removed from the normalized residual energy using its L2 estimate and the pointwise RH bound already proved in the same argument. The normalized cross-term error is $O_c(\log^{-3} T)$; the decomposition needs no stronger pair-correlation input. Writing $e_N(v) = E(Ne^v) / (Ne^v)^{1/2}$, the remaining residual is

$$N^{-c/\log T - it} [-e_N(0) + s_c(t) \int_0^\infty e_N(v) e^{-(c/\log T)v - itv} dv].$$

The two damping widths therefore act on one actual arithmetic function. This coupling is a concrete structure for the next estimate. Replacing the two energies by independent nonnegative variables loses that structure; assuming it forces (2) without a proof would also be an error.

Using only $|E(x)| \ll \sqrt{x} \log^2(x)$ and absolute values gives a residual bound of order $T \log^3(T)$ for a power-of- T cutoff. Its squared integral is far too large. This is a documented failure of that particular estimate, not a proof that RH or all analytic approaches are insufficient.

4. Positivity gives a valid weak bound, but cannot reach the target

The [bounded positivity note](#) gives an explicit band-limited minorant, using the known interior Montgomery band and positivity of the pair measure. Its resulting exact expression is approximately -0.208674513 for W , far below $1/16$. The minorant is optimal only within the one-parameter correction family written there. Its [independent review](#) checks the

Fourier support, endpoint pairing and realizable point-process obstruction.

The decisive obstruction remains the actual stationary half-grid determinantal process. It matches the known low band, satisfies the stated point-process positivity constraints, and attains $W_{AH} < 1/16$. Thus neither generic positivity nor merely changing two smoothing widths supplies the required arithmetic input. The frequency weight changes sign above $\log(2 \cosh 1)$; even a one-term positive-coefficient polynomial disproves a universal positive-coefficient quadratic-form claim.

5. Verification and next scope

The accompanying checks verify exact scalar enclosures, a finite-step integration-by-parts identity and an absolutely convergent comparison with ζ'/ζ . Low-height regularized-prime sums are explicitly labeled diagnostics. They are not evaluations of the large- T target, and no convergence rate in the critical strip is inferred from them. Intake hashes and an isolated integration replay preserve the received evidence.

The next mathematical obligation is (2), exploiting the common centered prime error in (3), or the alternative compact prime-covariance target in Round 7. Repeating a positive-coefficient argument, dropping the pole before bounding its cross term, and scanning generic random-matrix models are postponed. The famous conjecture remains open in this programme.

Current report 18: The actual–zeta two–scale target: short–prime projection and a centered tail

Collection: R8 — short-prime projection and signed residual.

Source: [research/dyson/round8/resolvent-arithmetic/SHORT_PRIME_PROJECTION_AND_CENTERED_TAIL.md](#).

SHA-256: `0067a1b0c7bd4f0b80ef89d6ac85eca1ae99e652375c08c41706ec1f1ddbe40e` . Git blob: `03cf06eaf07980d2bb1a04a4aa3fb41540c44fd4` . Original bytes: 14389.

The actual–zeta two–scale target: short–prime projection and a centered tail

Date: 2026-09-05. Status: ordinary analytic identities and an explicit remaining inequality. The requested lower bound $1/16$ is **not proved**. This note does not claim a new theorem about zeta zeros. Round 7 is unchanged.

The useful outcome is a precise arithmetic decomposition. The part supplied by short primes is positive and completely evaluable. The remaining term is one **signed difference of actual residual energies**, with an absolutely convergent representation involving $\psi(x)-x$ under RH. Positivity of the von Mangoldt coefficients does not control that signed difference.

1. Definitions and the resulting obligation

Write

$$s_c(t) = \frac{1}{2} + \frac{c}{\log T} + it, \quad H_c(t) = -\frac{\zeta}{\bar{\zeta}}(s_c(t)), \quad I_T(c) = \int_0^T |H_c(t)|^2 dt,$$

where $c>0$ is fixed, and

$$W_T = \frac{2}{T \log^2 T} (\sinh(2)/_T(1) - \sinh(1)/_T(1/2)). \quad (1)$$

Set $N = \lfloor T/\log^6 T \rfloor$, for sufficiently large T , and define the genuine finite polynomials

$$P_c(t) = \sum_{n \leq N} \frac{\Lambda(n)}{n^{s_c(t)}}, \quad R_c(t) = H_c(t) - P_c(t). \quad (2)$$

Under RH, the contour calculation below proves

$$W_T = B + E_T + \alpha(1), \quad E_T = \frac{2}{T \log^2 T} (\sinh(2) \parallel R_1 \parallel_{\frac{2}{2}} - \sinh(1) \parallel R_{1/2} \parallel_{\frac{2}{2}}), \quad (3)$$

where all norms are over $[0, T]$ and

$$\begin{aligned} B &= 2 \int_0^1 u (\sinh(2)e^{-2u} - \sinh(1)e^{-u}) du \\ &= \frac{e^2}{4} - e + \frac{5}{4} + \frac{1}{e} - \frac{9}{4e^2} + \frac{3}{4e^4} \\ &= 0.4560939793292317215 \dots \end{aligned} \quad (4)$$

Thus the precise remaining task is

$$\boxed{\liminf_{T \rightarrow \infty} E_T \geq \frac{1}{16} - B = -0.3935939793292317215 \dots} \quad (5)$$

By (3), (5) is equivalent to the requested lower bound for W_T . Equation (5) is not deduced here from RH. The existing Round 7 AH calculation gives a limiting W below $1/16$, so treating (5) as a consequence of formal coefficient positivity would be circular progress.

2. An exact continuation from centered prime counting

Put $E(x) = \psi(x) - x$, with $\psi(x) = \sum_{n \leq x} \Lambda(n)$. At an integer cutoff, ψ includes the atom at that integer. Assuming RH, the classical bound

$$E(x) = O(x^{1/2} \log^2(2x)) \quad (6)$$

implies that, for every s with $\operatorname{Re}(s) > 1/2$ and $s \neq 1$,

$$\boxed{-\frac{\zeta'}{\zeta}(s) = \frac{s}{s-1} + s \int_1^\infty E(x) x^{-s-1} dx.} \quad (7)$$

The integral in (7) is absolutely convergent. To prove it, first take $\operatorname{Re}(s) > 1$, use Stieltjes integration by parts in the absolutely convergent von Mangoldt series, and subtract the integral of x . Both sides then continue meromorphically to $\operatorname{Re}(s) > 1/2$. RH excludes logarithmic-derivative poles from this open region except the pole at $s=1$. The residue of the right side at $s=1$ is $+1$, as required for $-\zeta'/\zeta$. No critical-strip Dirichlet series has been expanded.

For any real $X \geq 1$, finite summation by parts gives

$$\boxed{-\frac{\zeta'}{\zeta}(s) = \sum_{n \leq X} \Lambda(n) n^{-s} + \frac{X^{1-s}}{s-1} - E(X) X^{-s} + s \int_X^\infty E(x) x^{-s-1} dx.} \quad (8)$$

Consequently the residual in (2) is **exactly**

$$R_c(t) = \frac{N^{1-s_c(t)}}{s_c(t)-1} - E(N) N^{-s_c(t)} + s_c(t) \int_N^\infty E(x) x^{-s_c(t)-1} dx. \quad (9)$$

The positive prime coefficients have therefore been replaced in the tail by a centered, signed arithmetic error. The pole term is part of the identity; it cannot be discarded before estimating it at the chosen scale.

Equivalently, for fixed $\operatorname{Re}(s) > 1/2$ and $s \neq 1$,

$$-\frac{\zeta'}{\zeta}(s) = \lim_{M \rightarrow \infty} \left(\sum_{n \leq M} \Lambda(n) n^{-s} + \frac{M^{1-s}}{s-1} \right). \quad (10)$$

For a fixed T the convergence is uniform on the two compact vertical segments used in (1). It is not an assertion that the unregularized Dirichlet series converges there. The order of limits in (10), followed by any T asymptotic, must be retained unless a uniform remainder is proved.

3. A quantitative bound that exposes the limitation of pointwise RH

Write $\delta = \operatorname{Re}(s) - 1/2 > 0$ and $\ell = \log X$, with $X \geq 2$. Formula (6) implies

$$\left| -E(X)X^{-s} + s \int_X^\infty E(x)x^{-s-1} dx \right| \ll X^{-\delta} [(\ell+1)^2 + |s| \left(\frac{(\ell+1)^2}{\delta} + \frac{2(\ell+1)}{\delta^2} + \frac{2}{\delta^3} \right)]. \quad (11)$$

This follows by integrating $X^{-1-\delta}(\log x + 1)^2$ explicitly; the implied constant is the one in (6), up to an absolute factor. At $\delta = c/\log T$, $X = T^\theta$ with fixed $\theta > 0$, and $t \leq T$, (11) supplies only

$$O_{c,\theta}(T \log^3 T), \quad (12)$$

which is far larger than the natural mean-square scale. Integrating the square of this estimate gives $O(T^3 \log^6 T)$, whereas (1) is normalized by $T \log^2 T$. This explicitly identifies a failure of the **pointwise RH estimate used in (6)**, not a proof that every possible consequence of RH is insufficient.

Trying to repair this particular absolute-value bound by making $X^{-\delta}$ as small as a negative power of T forces $\log X$ to be of order $(\log T)^2$. Such a cutoff destroys the short-polynomial mean-value regime. It is not a support extension obtained for free.

4. The actual mixed–integral lemma

For fixed $c > 0$, T sufficiently large, $3 \leq N \leq T$, and

$$\sigma = \frac{1}{2} + \frac{c}{\log T}, \quad P(t) = \sum_{n \leq N} \Lambda(n) n^{-\sigma - it}, \quad D = \sum_{n \leq N} \Lambda(n)^2 n^{-2\sigma},$$

assume $\sigma \leq 3/4$ and $\beta = 1 + 1/\log N \leq 2$. Under RH,

$$\int_0^T \left(-\frac{\zeta'}{\zeta}(\sigma + it) \right) \overline{P(t)} dt = TD + O_c(N \log^3 T). \quad (13)$$

This is a complex identity with a bounded complex error. In particular its real part has the same main term.

Contour and pole control. Use

$$F(s) = \left(-\frac{\zeta'}{\zeta}(s) \right) \sum_{n \leq N} \Lambda(n) n^{s-2\sigma} \quad (14)$$

on the rectangle from $\sigma + iT$ to $\beta + iT$. On the left side the finite sum is exactly $\overline{P(t)}$. The pole at $s=1$ is below the rectangle, and RH places every nontrivial zero strictly to its left. The standard local partial-fraction estimate for the logarithmic derivative, together with $O(\log T)$ zeros in a unit interval, gives

$$\frac{\zeta}{\zeta}(u+iT) = O_c(\log^2 T) \quad (\sigma \leq u \leq \beta).$$

Also

$$\left| \sum_{n \leq N} \Lambda(n) n^{u-2\sigma+iT} \right| \leq \sum_{n \leq N} \log n n^{\beta-2\sigma} \leq eN \log N.$$

The top integral is therefore $O_c(N \log^3 T)$. On the compact bottom segment, RH and finiteness of the number of nearby zeros give the sufficient bound $O_c(\log T)$ for the logarithmic derivative: its distance to any nontrivial zero is at least $c/\log T$, and the pole at $s=1$ has height zero. The bottom integral is thus $O_c(N \log N \log T)$. The same compact RH bound on the initially omitted left interval $0 \leq t \leq 1$, together with $\sum_{n \leq N} \Lambda(n) n^{-\sigma} \ll \sqrt{N} \log N$, bounds that interval by $O_c(\sqrt{N} \log N \log T)$. These errors are absorbed by $O_c(N \log^3 T)$. No assumption that T avoids a zero ordinate, and no separately verified low-zero table, is needed.

The right line is an honest absolutely convergent series. At $\text{Re}(s)=\beta>1$,

$$-\frac{\zeta}{\zeta}(s) = \sum_{m \geq 2} \Lambda(m) m^{-s}.$$

The $m=n$ terms of the vertical integral give $(T-1)D$. The absolute value of all other integrated terms is at most twice

$$\sum_{n \leq N} \Lambda(n) n^{\beta-2\sigma} \sum_{\substack{m \geq 2 \\ m \neq n}} \frac{\Lambda(m) m^{-\beta}}{|\log(m/n)|}. \quad (15)$$

For $n/2 \leq m \leq 2n$, $m \neq n$, the inner sum is

$$O(n^{1-\beta} \log^2(2N)),$$

by $\Lambda(m) \leq \log m$ and the harmonic sum over $|m-n|$. After the outer factor is applied, summing over n costs $O(N \log^3 N)$, because $n^{1-2\sigma} \leq 1$.

Outside that range the denominator is at least $\log 2$, and

$$\sum_{m \geq 2} (\log m) m^{-\beta} \ll (\beta-1)^{-2} \ll \log^2 N.$$

The remaining outer sum is $O(N \log N)$, since $n^{\beta-2\sigma} \leq N^{1/\log N} = e$. Thus the infinite right-line off-diagonal sum is $O(N \log^3 N)$. This completes (13). Neither an unproved long-polynomial mean value nor a critical-strip Dirichlet expansion has been used.

5. Orthogonal decomposition to leading order

An elementary finite-polynomial mean-value bound gives

$$\|P\|_2^2 = TD + O(N \log^4 T). \quad (16)$$

For completeness, bound each non-diagonal integral by $2/|\log(m/n)|$, use $|\log(m/n)| \geq |m-n|/N$, then $2|a_m a_n| \leq |a_m|^2 + |a_n|^2$, and sum the resulting harmonic series. This gives $O(N \log N \sum |a_n|^2)$. Here $\sum |a_n|^2 \leq \sum_{n \leq N} (\log n)^2/n = O(\log^3 N)$, proving the stated error. Sharper standard mean-value estimates are

unnecessary.

Combining (13) and (16) gives

$$\boxed{I_T(c) = TD + \|R_c\|_2^2 + O_c(N \log^4 T).} \quad (17)$$

This is an asymptotic projection identity for these specific coefficients, not a claim that the finite functions n^{-it} are exactly orthogonal on $[0, T]$. In particular, the cross term $\langle R_c, P_c \rangle$ is $O_c(N \log^4 T)$, rather than identically zero.

With $N = \lfloor T/\log^6 T \rfloor$, the error in (17), divided by $T \log^2 T$, is $O_c(\log^{-4} T)$. The prime number theorem and partial summation give

$$\frac{D}{\log^2 T} \rightarrow \int_0^1 u e^{-2cu} du. \quad (18)$$

One may obtain (18) from $\sum_{n \leq x} \Lambda(n)^2 \sim x \log x$; the contribution of prime powers of exponent at least two is negligible. Equations (17) – (18) prove (3). All leading constants and both values of c use the same cutoff N .

6. What the residual condition measures

The pole term in (9) is individually negligible in the normalized L^2 scale at this N :

$$\int_0^T \left| \frac{N^{1-s_c(t)}}{s_c(t)-1} \right|^2 dt \ll_c N^{1-2c/\log T} \ll_c N. \quad (19)$$

Its denominator has real part bounded away from zero for sufficiently large T . Dropping it from **the residual energy** also requires a bound on the other factor in the cross term. The same RH partial-fraction estimate used in Section 4 gives the sufficient pointwise bound $H_c(t) = O_c(\log^2 T)$ on $[0, T]$. Equation (16) also gives $\|P_c\|_2^2 = O_c(T \log^4 T)$. Thus $\|R_c\|_2^2 = O_c(T \log^4 T)$, and Cauchy – Schwarz with (19) bounds the change of residual energy, divided by $T \log^2 T$, by

$$O_c\left(\sqrt{\frac{N}{T}} + \frac{N}{T \log^2 T}\right) = O_c(\log^{-3} T) = o(1).$$

This weaker estimate is sufficient and makes the removal independent of the sharper Round 7 pair/resolvent bound.

Accordingly (5) can also be stated with

$$\tilde{R}_c(t) = -E(N)N^{-s_c(t)} + s_c(t) \int_N^\infty E(x)x^{-s_c(t)-1} dx \quad (20)$$

in place of R_c . This version pinpoints the missing arithmetic information: a comparison of two Laplace-damped Fourier energies of the **same centered prime error**.

More explicitly, put $e_N(v) = E(Ne^v)/(Ne^v)^{1/2}$, $v \geq 0$, and $\delta_c = c/\log T$. Then

$$\tilde{R}_c(t) = N^{-\delta_c - it} [-e_N(0) + s_c(t) \int_0^\infty e_N(v) e^{-\delta_c v - itv} dv]. \quad (21)$$

The two copies of e_N are coupled. Replacing their energies by two unrelated nonnegative numbers loses actual analytic structure; asserting that the coupling forces (5) without proof is equally unjustified. A new usable estimate must exploit that common arithmetic function at precision sufficient for (5).

For orientation only, under the full sine/GUE prediction the normalized residual energy would be

$$\frac{\|R_c\|_2^2}{T \log^2 T} \rightarrow \frac{e^{-2c}}{2c}. \quad (22)$$

Its signed contribution to (3) would be $-0.3738225362077544\dots$. Thus these residuals are leading-order objects, not errors tending to zero. Formula (22) is a model prediction, not an additional result of the contour lemma.

7. Why positive coefficients do not close the argument

For a finite polynomial with nonnegative coefficients, the diagonal coefficient in the two-scale difference is

$$g(u) = \sinh(2)e^{-2u} - \sinh(1)e^{-u}, \quad u = \frac{\log n}{\log T}.$$

It is negative for $u > \log(2 \cosh 1) = 1.126928011\dots$. Even the one-term polynomial $\Lambda(p)p^{-s}$, with $T = p^{1/2}$ and p a sufficiently large prime, has a strictly negative two-scale squared-norm difference because $g(2) < 0$. This is a counterexample to a **universal positive-coefficient quadratic-form assertion**, not a counterexample to the target for actual ζ .

Moreover, (10) has a mandatory continuous counterterm. At a real point s between $1/2$ and 1 the finite positive prime sum grows, while $(M^{\{1-s\}}/(s-1))$ is negative and cancels that growth. Removing the counterterm changes the analytic function. These two obstacles survive before any sophisticated arithmetic estimate is attempted.

8. Verification and scope

The accompanying `check_centered_tail.py` records:

- exact rational enclosures for B and $1/16-B$, obtained by enclosing e with its Taylor series and a geometric remainder;
- an independent finite-step integration-by-parts check of (8), with a finitely supported von Mangoldt measure and a completely evaluated tail, in its absolute-convergence half-plane;
- a comparison with actual ζ'/ζ at $\text{Re}(s)=3$, with an explicit bound on the omitted absolutely convergent von Mangoldt tail;
- small, labeled diagnostic evaluations of the **regularized** sum (10) at $s=3/4$, showing the size and sign of the mandatory counterterm. These finite diagnostics do not certify convergence rates in the critical strip.

The analytic proof, rather than the numerical checks, supplies continuation and the large- T statements. No numerical value of W_T has been inferred from the low-height examples.

Source inputs for the decomposition are standard RH consequences and the prime number theorem. Relevant primary background is [Goldston, *Notes on Pair Correlation of Zeros and Prime Numbers*](#), particularly the explicit-formula and mean-value discussion. The significance of the $1/16$ threshold comes from the separate, independently

reviewed Round 7 reduction using Goldston – Lee – Schettler – Suriajaya, *Pair Correlation Conjecture II: The Alternative Hypothesis*. The full mixed-integral proof needed here is written in Section 4; neither the decomposition nor the removal of the pole term requires the stronger pair/resolvent estimate.

Conclusion of this bounded test: the short-prime main term can be rigorously isolated, and the actual arithmetic term that must improve is (5), equivalently (20) – (21). No lower bound beyond the existing information has been established. Further work should address this signed centered-prime energy rather than reusing coefficient positivity or omitting the pole subtraction.

Current report 19: Independent audit of the actual–zeta short–polynomial identity

Collection: R8 — short-prime projection and signed residual.

Source: `research/dyson/round8/resolvent-arithmetic/INDEPENDENT_IDENTITY_AUDIT.md`.

SHA-256: `5159f2127474e57041368aa3741cefa9b7c29040b8286492bb6107eea052c8ab` . **Git blob:** `988ec3ec381faeb3d4a9b4b4f6ef830bc8be9a39` . **Original bytes:** 12823.

Independent audit of the actual–zeta short–polynomial identity

Date: 2026-09-05. **Reviewer:** `yau_flow` , independently of the authoring agent.

Verdict: the complete draft `SHORT_PRIME_PROJECTION_AND_CENTERED_TAIL.md` has been read and its analytic identities independently checked below. The estimates, endpoint conventions, and signs pass. Two optional simplifications were returned to the author; neither repairs a failure of the main claim. This is ordinary internal proof review, not formal verification or external peer review. No numerical integration, parameter scan, or modification of author evidence was performed.

1. Statement and exact range

Assume RH. Fix $\mathcal{C} > 0$, and for every sufficiently large real T , set

$$N = \lfloor \frac{T}{\log^6 T} \rfloor, \quad \sigma = \frac{1}{2} + \frac{\mathcal{C}}{\log T}, \quad H(t) = -\frac{\zeta}{\zeta}(\sigma + it), \quad P(t) = \sum_{n \leq N} \frac{\Lambda(n)}{n^{\sigma+it}}.$$

Put $D_N = \sum_{n \leq N} \Lambda(n)^2 n^{-2\sigma}$. The checked identity is

$$\int_0^T |H(t)|^2 dt = TD_N + \int_0^T |H(t) - P(t)|^2 dt + O_{\mathcal{C}}(N \log^4 T). \quad (1)$$

The implicit constant may depend on the fixed $\mathcal{C}$. Uniformity as $\mathcal{C} \downarrow 0$, or for $\mathcal{C}$ varying arbitrarily with T , has not been established. The identity holds for all sufficiently large T ; no selection of top-edge heights avoiding zero ordinates is needed.

2. Mixed inner product: contour, orientation, and pole

Let $q(s) = -\zeta/\zeta(s)$, $\beta = 1 + 1/\log N$, and

$$G(s) = \sum_{n \leq N} \Lambda(n) n^{s-2\sigma}.$$

On the left edge $s = \sigma + it$, one has exactly $G(s) = \overline{P(t)}$. Thus $q(s)G(s)$, rather than a differently conjugated Dirichlet polynomial, is the required analytic integrand.

Use the rectangle with real parts σ, β and imaginary parts $1, T$. Under RH no nontrivial zero lies in this rectangle. The pole of q at $s = 1$ lies at height zero, **below** it. Trivial zeros are also outside. Consequently the contour shift contributes horizontal integrals and no residue. The factor $ds = i dt$ on either upward vertical edge preserves the normalization of the diagonal term.

The right-edge Dirichlet series is absolutely convergent. Its diagonal contribution is $(T - 1)D_N$. The off-diagonal error is bounded by a constant times

$$\sum_{n \leq N} \Lambda(n) n^{\beta-2\sigma} \sum_{m \equiv n} \frac{\Lambda(m) m^{-\beta}}{|\log(m/n)|}. \quad (2)$$

The infinite m -sum is not truncated silently.

3. Near-diagonal and far-pair estimates

For $n/2 \leq m \leq 2n$, $m \equiv n$,

$$|\log(m/n)| \geq \frac{|m-n|}{2n}, \quad n^{\beta-2\sigma} m^{-\beta} \ll n^{-2\sigma}.$$

Since $n^{1-2\sigma} = n^{-2\sigma/\log T} \leq 1$, the corresponding summand is

$$\ll \frac{\log^2(2N)}{|m-n|}.$$

The harmonic sum over m costs $O(\log N)$; summing over $n \leq N$ gives $O(N \log^3 N)$. This directly controls the potentially dangerous neighboring integers without a cancellation assumption.

For the remaining pairs, $|\log(m/n)| \geq \log 2$. Also

$$n^{\beta-2\sigma} \leq n^{1/\log N} \leq e, \quad \sum_{m \geq 2} \frac{\log m}{m^{1+1/\log N}} \ll \log^2 N.$$

Using only $\Lambda(m) \leq \log m$ and $\sum_{n \leq N} \Lambda(n) \leq N \log N$, the far-pair portion of (2) is likewise $O(N \log^3 N)$. The claimed right-line error therefore holds without prime-pair estimates or cancellation.

4. Horizontal sides and the shrinking distance to the zeros

The standard local partial-fraction formula and the local zero-count bound imply, under RH,

$$q(u + iT) = O_c(\log^2 T), \quad \sigma \leq u \leq \beta. \quad (3)$$

Indeed, there are $O(\log T)$ zeros with $|\gamma - T| \leq 1$, and their distances to this segment are at least $\sigma - 1/2 = c/\log T$; the remaining terms contribute $O(\log T)$. This explains the dependence on c and remains valid when T itself equals a zero ordinate.

Uniformly on either horizontal side,

$$|G(u + it)| \leq \sum_{n \leq N} \Lambda(n) n^{\beta-2\sigma} \ll N \log N.$$

The top side has bounded length, so it contributes $O_c(N \log^3 T)$. On the bottom side at height one, a compact meromorphic bound gives $q(u + \hat{h}) = O_c(\log T)$, even without using an explicit numerical zero-free height. Hence its contribution is smaller. The pole at height zero does not enter either estimate.

For the omitted interval $[0, 1]$, one may use the crude bounds

$$|H(t)| \ll_c \log T, \quad |P(t)| \ll \sqrt{N} \log N.$$

They give a mixed contribution $O_c(\sqrt{N} \log N \log T)$, which is absorbed by the displayed error. The replacement of $(T - 1)D_N$ by TD_N is also harmless. Thus

$$\int_0^T H(t) \overline{P(t)} dt = TD_N + O_c(N \log^3 T). \quad (4)$$

5. Polynomial norm and completion of the square

Directly expanding $\int_0^T |P|^2$ gives TD_N plus off-diagonal terms. The same near-diagonal estimate just used gives $O(N \log^3 N)$. Away from comparable integers, the absolute bound $\sum_{n \leq N} \Lambda(n) n^{-\sigma} \ll \sqrt{N} \log N$ suffices. In particular, the author's weaker bound

$$\int_0^T |P(t)|^2 dt = TD_N + O(N \log^4 T) \quad (5)$$

is valid.

The exact Hilbert-space identity

$$\|H - P\|_2^2 = \|H\|_2^2 + \|P\|_2^2 - 2\Re\langle H, P \rangle$$

combined with (4) and (5) proves (1). The main diagonal has coefficient **one**. There is no missing factor of two and no negative sign on the residual norm. With this choice of N , the normalized error in (1) is $O_c(\log^{-4} T)$ after division by $T \log^2 T$.

6. Exact continuation in terms of the prime-counting error

Use the endpoint convention

$$\psi(N) = \sum_{n \leq N} \Lambda(n), \quad E(x) = \psi(x) - x.$$

Partial summation for $\Re s > 1$ gives

$$q(s) - \sum_{n \leq N} \Lambda(n) n^{-s} = \frac{N^{1-s}}{s-1} - E(N) N^{-s} + s \int_N^\infty E(x) x^{-s-1} dx. \quad (6)$$

Both the sign of the pole term and the minus sign on the endpoint error are correct. At an integer cutoff, replacing $\psi(N)$ by its left limit would require adding the missing endpoint prime-power term; the convention above avoids that error.

RH gives $E(x) = O(\sqrt{x} \log^2 x)$. Therefore the integral converges absolutely and locally uniformly for $\Re s > 1/2$, and the identity extends meromorphically to that half-plane. The only pole at $s = 1$ is explicitly present on the right. Formula (6) is thus a true arithmetic continuation identity at the chosen s , not a formal Dirichlet series outside its convergence half-plane.

Absolute convergence is **not** the desired fine bound on the residual. Writing $\delta = \sigma - 1/2 = c/\log T$, its crude absolute majorant contains

$$|s| N^{-\delta} \left(\frac{\log^2 N}{\delta} + \frac{2 \log N}{\delta^2} + \frac{2}{\delta^3} \right).$$

This can be far too large on $[0, T]$. Cancellation or a genuinely stronger mean-square estimate is still needed for the signed two-scale target.

7. Diagonal constant and the unresolved signed residual

PNT and partial summation give, for each fixed $c > 0$,

$$\frac{D_N}{\log^2 T} \rightarrow d(c) := \int_0^1 u e^{-2cu} du = \frac{1 - (1 + 2c)e^{-2c}}{4c^2},$$

because $\log N / \log T \rightarrow 1$. The two-scale diagonal contribution is exactly

$$B_{\text{low}} = 2 \{ \sinh 2 d(1) - \sinh 1 d(1/2) \} = 0.4560939793292318 \dots$$

Writing $R_c = H_c - P_c$, the remaining requirement for the sufficient threshold $1/16$ is

$$\liminf_{T \rightarrow \infty} \frac{2}{T \log^2 T} (\sinh 2 \|R_1\|_2^2 - \sinh 1 \|R_{1/2}\|_2^2) \geq \frac{1}{16} - B_{\text{low}}.$$

The right side is $-0.3935939793292318 \dots$. These decimals are ordinary evaluations of the explicit constants, not numerical zeta data or an outward enclosure. Positivity of each individual squared norm does not establish a lower bound for this signed combination. Identity (1) and continuation (6) make that remaining arithmetic obligation precise; they do not solve it.

8. Centered–tail refinement and additional checks of the full draft

The author's global continuation formula

$$q(s) = \frac{s}{s-1} + s \int_1^\infty E(x) x^{-s-1} dx$$

also has the correct constant. It specializes to (6) at $N = 1$, since $\psi(1) = 0$ and $E(1) = -1$. For each fixed s in the claimed half-plane, the endpoint error and centered tail vanish as the cutoff tends to infinity. Hence the regularized limiting sum in the author's equation (10) is valid. Its convergence is uniform on each fixed compact vertical segment, but a subsequent T -limit requires preserving that order or proving uniform estimates.

The norm of the explicit pole term at the chosen cutoff satisfies

$$\left\| \frac{N^{1-s_c(t)}}{s_c(t)-1} \right\|_2^2 \ll_c N.$$

The author's use of the already reviewed stronger bound $I_T(c) = O_c(T \log^2 T)$ justifies dropping this term from the **normalized residual energy**, including its cross term. There is also a self-contained weaker route: the pointwise RH estimate already used in this proof gives $I_T(c) = O_c(T \log^4 T)$. Together with the elementary polynomial norm bound, this gives $\|R_c\|_2^2 = O_c(T \log^4 T)$. Cauchy – Schwarz then bounds the change in normalized residual energy by

$$O_c(\sqrt{N/T} + \frac{N}{T \log^2 T}) = O_c(\log^{-3} T) = o(1).$$

Thus this particular removal need not depend on the Round 7 pair/resolvent transfer. This optional simplification was communicated to the author.

The change of variables $x = Ne^v$, with $e_N(v) = E(Ne^v)/(Ne^v)^{1/2}$, gives exactly

$$\tilde{R}_c(t) = N^{-\delta_c - it} [-e_N(0) + s_c(t) \int_0^\infty e_N(v) e^{-\delta_c v - itv} dv].$$

There is no missing power of N , Jacobian, or endpoint term. Both scales use the same function e_N . This shared arithmetic input must be preserved in any future estimate; treating the two residuals as arbitrarily independent objects would lose information.

The reported GUE orientation formula is consistent with the existing normalization: subtracting $d(c)$ from $V_{\text{sic}}(2c)/2 = (1 - e^{-2c})/(4c^2)$ gives $e^{-2c}/(2c)$. This is a conditional model prediction, as the draft states, and was not used to prove the arithmetic identity.

For the bottom contour and initial time interval, the draft's stronger bounds invoke the familiar zero-free low-height compact region. The weaker bounds in § 4 of this audit avoid that extra fact altogether, using only RH and finite zero counts on compact sets. They are already sufficient for the claimed error. This was the second optional simplification sent to the author.

9. Provenance and scope

The fully inspected author draft has SHA-256

[8840bdfcdfa07baf369deaed39151292ee28ff386f946336f21368d867277305](#). Any later editorial version should retain this review's explicit scope or be checked as a separate delta. The mathematical constants in § 7 were independently evaluated from their displayed elementary formulas. The planned numerical check script had not yet been saved when this proof review was completed and is **not** claimed to have been independently reviewed or rerun.

The primary background references used by the author are [Goldston's notes on pair correlation and prime numbers](#) and [Goldston – Lee – Schettler – Suriajaya's AH paper](#). The mixed-integral estimate itself was checked directly above rather than attributed to an unstated stronger source theorem. No claim of novelty for the contour method or continuation identities is made.

10. Final-version delta acceptance

The frozen final author draft was subsequently inspected and its SHA-256 independently recomputed as

[0067a1b0c7bd4f0b80ef89d6ac85eca1ae99e652375c08c41706ec1f1ddbe40e](#). This final version is accepted. The earlier hash in § 9 records the initially reviewed draft rather than the final checkpoint.

The bounded delta review checked these changes directly:

- Section 4 now uses the compact RH logarithmic-derivative bound $O_c(\log T)$ for the bottom side and initial interval, with errors $O_c(N \log N \log T)$ and $O_c(\sqrt{N} \log N \log T)$. Both are absorbed by $O_c(N \log^3 T)$. No low-zero table or exceptional choice of T is required.
- Section 6 now removes the pole from the normalized residual energy using only $\|R_c\|_2^2 = O_c(T \log^4 T)$ and pole norm squared $O_c(N)$. Its stated error $O_c(\sqrt{N/T} + N/(T \log^2 T)) = O_c(\log^{-3} T)$ is correct.
- The source paragraph now correctly separates the RH/PNT inputs for the decomposition from the external significance of the 1/16 target. The decomposition and pole removal no longer depend on the stronger Round 7 pair/resolvent estimate.

These changes implement the two simplifications in § 8 without changing the diagonal constant, the exact centered-tail identity, or the unproved signed-residual obligation. The author's completed numerical script and result files were not rerun in this delta review; their replay belongs to the coordinator's separate verification. This independent review file is now final.

Current report 20: A bounded positivity audit of the two-scale target

Collection: R8 — short-prime projection and signed residual.

Source: [research/dyson/round8/spectral-positivity/POSITIVITY_OBLIGATION_NOTE.md](#).

SHA-256: [485b7778a0f7bab21492dac710b8732e5e1996734014fbba770d4b076a7b6329](#) . **Git blob:** [fcf57a8936ea557feb1a41b99f824d8e2cfa59f4](#) . **Original bytes:** 7514.

A bounded positivity audit of the two-scale target

Date: 2026-09-05. This note closes one bounded proof audit; it contains no new parameter scan, arithmetic estimate, or zeta result. Round 7 files were not changed.

Conclusion. The known interior Montgomery band, spectral positivity, and realizable stationary point-process constraints do not imply the target $W \geq 1/16$. The obstruction is an actual determinantal point process, not merely a proposed nonnegative spectrum. The additional calculation below gives an explicit, substantially weaker lower bound using pair-measure positivity. It identifies a valid inference but does not close the arithmetic obligation.

1. Normalization and the formal-spectrum warning

Write

$$A = \sinh 2, \quad B = \sinh 1, \quad K(a) = Ae^{-2|a|} - Be^{-|a|}.$$

For a unit-intensity stationary process, let μ be its pair measure including the diagonal, and let $S = \hat{\mu} - \delta_0$ be its centered structure-factor measure. Assume the relevant Poisson integrals are finite. The two-scale statistic is

$$W = \int K dS = \int J(x) d\mu(x) - (A - B), \quad J(x) = \frac{A}{1 + \pi^2 x^2} - \frac{2B}{1 + 4\pi^2 x^2}.$$

The known band is $dS(a) = |a| da$ on $(-1, 1)$; this specifies no atoms at the endpoints.

The sign change is exactly

$$a_0 = \log(A/B) = \log(2 \cosh 1) > 1.$$

Adding arbitrarily large nonnegative atoms at $\pm R$, for fixed $R > a_0$, leaves the interior band unchanged and sends $\int K dS$ to minus infinity. This proves only that **formal spectral positivity alone** supplies no finite lower bound. Such modified spectra have not been shown to be realizable by point processes and are not used as the actual obstruction below.

Pair positivity is a meaningful extra constraint. For $y = \pi^2 x^2$,

$$J(x) = \frac{(A - 2B) + (4A - 2B)y}{(1 + y)(1 + 4y)} > 0.$$

Since $\mu \geq \delta_0$, it immediately gives the valid but weak bound $W \geq -B$. Thus the formal unbounded-below example cannot respect all these pair constraints.

2. An explicit band-limited minorant

The following stronger bound is already available from the same data. It is not claimed optimal over all positivity arguments.

Set

$$a = 2 \tanh 1, \quad b = 2 \tanh(1/2), \quad \kappa = 2b - a > 0,$$

and define

$$R(x) = \frac{a}{1 + \pi^2 x^2} - \frac{2b}{1 + 4\pi^2 x^2}, \quad f(x) = \cos^2(\pi x)R(x) + \kappa \operatorname{sinc}^2(x), \quad g(x) = J(x) - f(x),$$

with $\operatorname{sinc}(x) = \sin(\pi x)/(\pi x)$.

Claim: $f \geq 0$, and $\widehat{g}$ is supported on $[-1, 1]$, with $\widehat{g}(\pm 1) = 0$. Consequently

$$W \geq \widehat{g}(0) + 2 \int_0^1 \alpha \widehat{g}(\alpha) d\alpha - (A - B) =: B. \quad (\text{P})$$

The right side is an elementary expression in exponentials and hyperbolic tangents. Its floating evaluation is $-0.208674512963925 \dots$; the exact expression, rather than these digits, is the proved bound.

To prove nonnegativity, note $0 < b < a < 2b$. Direct calculation gives $R(x) \geq -\kappa$ for all x , since

$$R(x) + \kappa = \frac{y\{(8b - a) + (8b - 4a)y\}}{(1 + y)(1 + 4y)} \geq 0.$$

Also $R(x) < 0$ only if $y < \kappa/(4a - 2b) < 1$. There $|\pi x| < 1 < \pi/2$, and $\operatorname{sinc}^2(x) \geq \cos^2(\pi x)$. Hence $f \geq \cos^2(\pi x)(R + \kappa) \geq 0$. Where $R \geq 0$, nonnegativity is immediate.

For the Fourier support, put $E(\alpha) = ae^{-2|\alpha|} - be^{-|\alpha|}$. Then

$$\widehat{f}(\alpha) = \frac{1}{2}E(\alpha) + \frac{1}{4}E(\alpha - 1) + \frac{1}{4}E(\alpha + 1) + \kappa(1 - |\alpha|)_+.$$

For $|\alpha| \geq 1$, the first three terms equal $K(\alpha)$, using $a \cosh^2(1) = A$ and $b \cosh^2(1/2) = B$. Thus $\widehat{g}$ vanishes there, including the endpoints. Pair-measure positivity and the known band prove (P). The $O(x^{-2})$ decay permits these tests under the linear pair-mass bound used in Round 7.

Here is a closed finite expression for reproducing B. For each triple (r, c, d) in $\{(2, A, a), (1, -B, -b)\}$, define

$$u = c - d(1/2 + e^{-r}/4), \quad v = -de^{-r}/4, \\ l_-(r) = \frac{1 - (1 + r)e^{-r}}{r^2}, \quad l_+(r) = \frac{(r - 1)e^r + 1}{r^2}.$$

Then

$$B = \sum_{(r,c,d)} \{u(1+2L_-(r)) + v(1+2L_+(r))\} - \frac{4}{3}\kappa - (A-B).$$

For the one-parameter correction $f_c = \cos^2(\pi x)R + c \operatorname{sinc}^2(x)$, positivity at zero requires $c \geq \kappa$. Since the bound decreases by $4c/3$, the chosen correction is optimal within this particular one-parameter family, not among all minorants or point processes.

3. The obstruction is a genuine point process

On $\mathbb{Z}$, take the determinantal process with kernel

$$Q(j, k) = \int_{-1/4}^{1/4} e^{2\pi i(j-k)t} dt = \begin{cases} 1/2, & j = k, \\ \sin(\pi(j-k)/2)/(\pi(j-k)), & j \neq k. \end{cases}$$

This is the orthogonal projection onto a frequency interval of length $1/2$, so it defines an actual determinantal probability measure; see Lyons, *Determinantal probability measures*. If $X \subset \mathbb{Z}$ is its random occupied set and U is independently uniform on $[0, 1)$, then

$$\{(j+U)/2 : j \in X\}$$

is a stationary, simple, unit-intensity point process on the line. Stationarity follows from integer-translation invariance of the discrete process and uniform random translation within one lattice cell.

Its Palm pair measure is exactly

$$\mu_A = \delta_0 + \frac{1}{2} \sum_{k \neq 0} (1 - \operatorname{sinc}^2(k/2)) \delta_{k/2}.$$

Indeed the discrete pair occupation probability is $1/4 - |Q(0, k)|^2$, and conditioning on the occupied point divides by $Q(0, 0) = 1/2$. All pair weights are nonnegative. In particular, $\mu_A([-R, R]) \leq 1 + 2R$, giving the required linear cumulative bound. The process has a hard core of one half.

Poisson summation gives its centered spectrum

$$dS_A(\alpha) = \operatorname{dist}(\alpha, 2\mathbb{Z}) d\alpha + \sum_{m \neq 0} \delta_{2m}.$$

Thus it has exactly the prescribed Montgomery data on $(-1, 1)$, and satisfies spectral positivity, diagonal normalization, pair-measure positivity, simplicity, and the cumulative pair bound. The nonzero even-frequency atoms must be retained.

Its Poisson variance and two-scale statistic are

$$V_A(r) = \frac{2 \tanh(r/2)}{r^2} + \frac{2}{e^{2r} - 1},$$

$$W_A = \frac{e^2}{4} + \frac{5}{4e^2} - e - \frac{2}{e} + \frac{3}{2} \in (0.06239, 0.06240) < \frac{1}{16}.$$

The exact constant enclosure is already certified in the frozen Round 7 artifact. No new decimal scan is needed. This gives the realizable obstruction to the proposed inference. It is not a construction of actual zeta zeros and is not claimed to satisfy every arithmetic identity known about zeta.

4. The remaining arithmetic obligation

For any fixed pair of positive smoothing scales, this same process remains compatible with the stated low-band and point-process assumptions. A statistic whose proved lower bound would exclude its value therefore cannot follow from those assumptions alone. Changing the scales may improve the size or shape of the target; it does not add the missing arithmetic information.

For the present scales, the unresolved requirement is an actual-zeta estimate for the signed out-of-band contribution

$$\int_{|\alpha| \geq 1} (Ae^{-2|\alpha|} - Be^{-|\alpha|}) dS_{\zeta}(\alpha),$$

or its finite-height counterpart with controlled errors. The negative tail beyond $\log(2 \cosh 1)$ prevents simply discarding the unknown part. The concrete arithmetic identity proposed for this contribution must be audited on its own; generic positivity and another random-matrix calculation cannot replace it.

Current report 21: Independent review of the fixed positivity minorant

Collection: R8 — short-prime projection and signed residual.

Source: `research/dyson/round8/spectral-positivity/MINORANT_REVIEW.md`.

SHA-256: `383af540ae53eba3a0fbc471fcf79d72b821167450cc05e9a89198c7dbc843a` . **Git blob:** `caaadd49825fc344d29371a28bedd8a938180293` . **Original bytes:** 9174.

Independent review of the fixed positivity minorant

Date: 2026-09-05. **Reviewer:** the independent residual/arithmetic agent. **Scope:** the displayed minorant in `POSITIVITY_OBLIGATION_NOTE.md`, its closed constant, and the half-grid determinantal pair normalization.

No parameter optimization, new literature survey, or actual-zeta computation was performed. The author note was not edited.

Verdict: accepted as an ordinary mathematical proof under the stated pair-measure and interior-band assumptions, including the linear cumulative pair bound used for the tests. I found no coefficient, sign, support, endpoint, or Palm-normalization defect. The regularization details below make explicit a step compressed into the note's decay statement. This is internal proof review, not formal verification or external peer review.

The bound is exactly the displayed expression, approximately -0.2086745129639258 . It improves the immediate $-\sinh(1)$ consequence of pair positivity, but it is not progress to the required positive threshold $1/16$, is not a new bound on actual zeta proved here, and is not optimal over general minorants. The optimum statement applies only to the specified one-parameter correction family and its displayed inference.

1. Nonnegativity of the remainder

Use the note's A, B, a, b, κ and $y = \pi^2 x^2$. To verify all parameter inequalities without decimals, put $t = \tanh(1/2)$, so $0 < t < 1$ and

$$\begin{aligned} a &= 4t/(1+t^2), & b &= 2t, \\ \kappa &= 4t^3/(1+t^2), \\ \kappa/(4a-2b) &= t^2/(3-t^2) < 1/2 < 1. \end{aligned}$$

Thus $0 < b < a < 2b$, both coefficients in the note's numerator for $R + \kappa$ are positive, and $R \geq -\kappa$. The numerator of R is $-\kappa + (4a-2b)y$. Wherever $R < 0$, $|\pi x| < 1 < \pi/2$. For $0 \leq z < \pi/2$, $\tan(z) \geq z$, because $(\tan(z) - z)' = \sec^2(z) \geq 0$. Hence $(\sin(z)/z)^2 \geq \cos^2(z)$, also at $z=0$ by continuity. It follows on that negative region that

$$\begin{aligned} \cos^2(\pi x) R(x) + \kappa \operatorname{sinc}^2(x) \\ \geq \cos^2(\pi x) [R(x) + \kappa] \geq 0. \end{aligned}$$

On the remaining region both summands are nonnegative. This proves $f \geq 0$ everywhere, including zero; indeed $f(0)=0$. All rational identities used above pass the accompanying exact symbolic checks.

For $f_c = \cos^2(\pi x)R + c \operatorname{sinc}^2(x)$, evaluation at zero gives $f_c(0) = c - \kappa$, so positivity requires $c \geq \kappa$. This condition is sufficient, since $f_c = f_\kappa + (c - \kappa)\operatorname{sinc}^2$. The displayed lower-bound functional decreases by $4c/3$ and is therefore maximized at $c = \kappa$ in this family. Even if one separately credits the known diagonal contribution $f_c(0)$, the resulting functional decreases by $c/3$, so the same endpoint is still optimal. This observation does not establish an optimum among other band-limited minorants or other pair constraints.

2. Fourier support, endpoints, and legitimate pair testing

The convention is Fourier transform $\widehat{h}(\alpha) = \int h(x) \exp(-2\pi i \alpha x) dx$. Under this convention,

$$\begin{aligned} [1/(1+\pi^2 x^2)]\widehat{} &= \exp(-2|\alpha|), \\ [2/(1+4\pi^2 x^2)]\widehat{} &= \exp(-|\alpha|), \\ [\operatorname{sinc}^2(x)]\widehat{} &= (1-|\alpha|)_+ . \end{aligned}$$

The coefficients $1/2, 1/4, 1/4$ in the modulation by $\cos^2(\pi x)$ are therefore correct. For $\alpha \geq 1$, the coefficient of $a \exp(-2\alpha)$ becomes $\cosh^2(1)$, and that of $b \exp(-\alpha)$ becomes $\cosh^2(1/2)$. The exact identities $a \cosh^2(1) = A$ and $b \cosh^2(1/2) = B$ cancel the whole Fourier tail. Evenness handles $\alpha \leq -1$. The triangular correction vanishes at both endpoints, and the modulation identities continue to hold there, so $\widehat{g}(1) = \widehat{g}(-1) = 0$ exactly. There are no overlooked endpoint δ terms: J, f and g are integrable functions, and their Fourier transforms are continuous functions.

An unknown atom of the pair spectrum at either endpoint therefore contributes zero. The known band can be used without silently assuming endpoint atom-freeness. Here is also an ordinary justification for passing from Schwartz test identities to this particular g . Its compactly supported transform is continuous, piecewise smooth, vanishes linearly at the endpoints, and has first derivative of bounded variation. Smooth it and cut it off just inside $(-1, 1)$. These approximants can be chosen with uniformly bounded L^1 norms and uniformly bounded total variation of their first derivatives: in an endpoint strip of width ϵ , the function is $O(\epsilon)$, so the second-derivative cutoff terms have bounded L^1 norm. Their inverse transforms g_ϵ consequently obey

$$|g_\epsilon(x)| \leq C/(1+x^2)$$

uniformly, and converge pointwise to g . The assumed linear cumulative pair bound gives integrability of this dominator. Dominated convergence on the pair side and the known interior spectral identity on the other side then yield

$$\int g d\mu = \widehat{g}(0) + \int_{-1}^1 |\alpha| \widehat{g}(\alpha) d\alpha .$$

This argument requires no guess about endpoint atoms or the unknown outer spectrum. Since $f = J - g \geq 0$ and μ is a positive measure, the note's inequality (P) follows. The subtraction $A - B$ is correct: it is $K(0)$, the Fourier atom of the unit mean-density contribution. Independently, $J(0) = A - 2B$ and $\mu \geq \delta_0$ give $W \geq -B$, as claimed.

3. Closed expression and fixed symbolic verification

For α in $[0,1]$, each triple (r,c,d) contributes

$$u \exp(-r \alpha) + v \exp(r \alpha), \\ u = c - d(1/2 + \exp(-r)/4), \quad v = -d \exp(-r)/4,$$

to that before the correction $-\kappa(1-\alpha)$. Direct integration of $\alpha \exp(\pm r \alpha)$ gives precisely I_{minus} and I_{plus} in the note. The value at zero contributes $-\kappa$ and the weighted integral contributes $-\kappa/3$, explaining the total $-4\kappa/3$ without any diagonal or two-sided factor ambiguity.

As an independent exact algebra check, set $X=\exp(1)$. The complete closed lower bound simplifies to

$$(X-1)(3X^6-6X^5-17X^4+28X^3-35X^2-6X-15) \\ / [12 X^2 (X+1)(X^2+1)].$$

This is an exact expression, not a fitted approximation. Substitution $X=e$ gives $-0.2086745129639258383515265856\dots$. The script's decimal evaluation is explicitly floating, not an outward enclosure; the ordinary proof uses the exact formula and does not require a decimal certificate.

The companion [minorant_symbolic_check.py](#) verifies the rational identities, the two elementary integrals, the endpoint, the correction coefficient, the Palm occupation division, the triangle-period integral and the DPP two-scale formula. [minorant_symbolic_check.json](#) records PASS, versions and source hashes. Reproduction from this directory is `python3 minorant_symbolic_check.py`; it writes only its adjacent JSON. This is a small fixed identity check, not an optimization or sampled nonnegativity test.

4. Half-grid DPP normalization and realized obstruction

The stated Q is the Fourier projection of $\ell_2(\mathbb{Z})$ onto the interval $[-1/4, 1/4]$. In particular it is a positive contraction with diagonal $1/2$, and the standard discrete determinantal construction supplies a probability measure with finite occupation probabilities $\det(Q|F)$. The note's [Lyons reference](#) is the appropriate primary reference for the projection/positive-contraction framework. No independence of the occupation indicators is being assumed.

After scaling lattice spacing to $1/2$, the occupancy probability $1/2$ gives intensity $2*(1/2)=1$. The independent uniform shift covers one full lattice cell; combined with integer-shift invariance of the discrete DPP this gives stationarity under every real translation. The lattice process is simple and has minimum positive spacing $1/2$.

For $k \neq 0$, the conditional probability of a point at displacement $k/2$ from a Palm point is

$$[1/4 - |Q(0, k)|^2] / (1/2) \\ = (1/2) [1 - \text{sinc}^2(k/2)].$$

The Palm base point contributes exactly δ_0 . Thus the displayed μ_A has the correct factor $1/2$, and its nonzero weights lie between zero and $1/2$. Counting the at most $2 \lfloor 2R \rfloor$ nonzero lattice sites proves $\mu_A([-R, R]) \leq 1 + 2R$.

The full Fourier calculation can be written without omitting the zero sample:

$$\mu_A = \delta_0 + (1/2) \sum_k \delta_{k/2}$$

$$- (1/2) \sum_k \text{sinc}^2(k/2) \delta_{k/2}.$$

Poisson summation transforms these three terms into the constant density 1, the comb $\sum_m \delta_{2m}$, and the periodized triangle $\sum_m (1 - |\alpha - 2m|)_+$, respectively. Hence the remaining continuous density is $\text{dist}(\alpha, 2\mathbb{Z})$. Centering removes only the comb atom at zero. All nonzero even-frequency atoms survive, exactly as in the note.

Finally, the Laplace integral of the triangular density over one positive period is $(1 - \exp(-r))^2 / r^2$. Summing its geometric repetitions on both sides gives $2 \tanh(r/2) / r^2$; the nonzero even comb gives $2 / (\exp(2r) - 1)$. Therefore both $V_A(r)$ and $W_A = A V_A(2) - B V_A(1)$ are correctly normalized. The accompanying exact symbolic check reproduces the closed W_A expression, whose strict comparison with $1/16$ was already enclosed in Round 7.

This actual process meets the stated low-band and positivity assumptions while remaining below $1/16$. It is consequently a valid obstruction to deductions using only those assumptions. It is neither an actual-zeta construction nor a model asserted to satisfy the missing arithmetic identities. No stronger impossibility or universality conclusion is endorsed.

Current report 22: Round 9: actual prime arithmetic for the Dyson–Montgomery programme

Collection: R9 — complementary moduli, genuine-prime tails, and the edge.

Source: [research/reports/dyson_round9.md](#).

SHA-256: [37e6f98bc96fdd9686a1e4ec158d822c087f89b35d7bf4aa9ac523c1b9a4d642](#) . Git blob: [8d10e985a2e7712b95dff4d42aef3f4c794b4b7c](#) . Original bytes: 13988.

Round 9: actual prime arithmetic for the Dyson–Montgomery programme

Date: 2026-09-05. This checkpoint continues the user's request to concentrate on RMT and actual zeta zeros. **No new zeta pair-correlation lower bound, AH refutation, half-gap theorem, RH proof, or prime-gap record is established.** The useful outcomes are a source-checked arithmetic transfer, a quantified removal of prime powers, a negative new resonator trial, and a precise mesoscopic lower-bound obligation.

1. The conjecture–level target remains unchanged

Set

$$I_T(c) = \int_0^T \left| \frac{\zeta}{\zeta} \left(\frac{1}{2} + \frac{c}{\log T} + it \right) \right|^2 dt, \quad W_T = \frac{2(\sinh 2 I_T(1) - \sinh 1 I_T(1/2))}{T \log^2 T}.$$

The reviewed Round 7 reduction gives, under RH and the precise AH-Pairs formulation,

$$W_T \rightarrow W_{AH} = 0.0623924179764985 \dots < \frac{1}{16}.$$

The bounded but possibly nonconvergent near-diagonal mass cancels. An actual-zeta proof of $\liminf W_T \geq 1/16$ would therefore contradict AH-Pairs under RH. That inequality remains unproved. Round 8 writes $W_T = B + E_T + o(1)$, with $B = 0.4560939793292317\dots$, so the missing signed residual inequality is $\liminf E_T \geq -0.3935939793292317\dots$. See the [Round 7 reduction](#) and [Round 8 proof](#).

The complementary compact test has Fourier support $[6/5, 7/5]$. Its centered prime covariance has AH value $-3/5$ and sine-kernel prediction $-3/10$. A sufficiently strong estimate here would be an actual out-of-band zeta result. The atomic diagonal, continuous mean, prime/mean cross term, and continuous mean square must all be retained.

2. What the 186 factorization structure really transfers

The complete transfer proof applies Proposition 2.3 and Corollary 2.19 of *Improved short gaps between primes*. It uses the ordinary analytic theorem as input, without making a claim about completion of every formalization obligation.

Choose $\omega=3/250$, $\delta=1/1000$ and $\epsilon=1/1000$. Then $240\omega+80\delta=2.96<3$, and the resulting full-prime distribution level is 0.523. Let Q_X be a set of distinct squarefree moduli $q=[D,E]$ in $(X^{1/2}, X^{.523})$ satisfying the stated complementary tail predicates. They imply triple dense divisibility at scale $X^{(.001)}$, even for explicit nonsmooth examples. Repeated representations of one q are counted once; shared prime factors of D,E are allowed.

Insert the exact identity

$$\Lambda(n) = \sum_{q \mid n} \mu(q) \log(n/q) = B_Q(n) + B_{\text{rest}}(n).$$

For a shift h and a smooth macroscopic weight w_h , define

$$C_{Q,h} = \sum_n \Lambda(n+h) B_Q(n) w_h(n), \quad M_{Q,h} = \sum_{\substack{q \in Q_X \\ (q,h)=1}} \frac{\mu(q)}{\phi(q)} \int w_h(u) \log(u/q) du.$$

The source gives, for every fixed $A>0$,

$$C_{Q,h} - M_{Q,h} = O_A(X \log^{-A} X). \quad (1)$$

This estimate is uniform for $1 \leq h \leq CH$, where $X=T^\alpha$, $6/5 \leq \alpha \leq 7/5$ and $H=X/T$. It applies to a localized version of the actual covariance's sinc weight. The proof checks the common primitive residue h across the modulus family, two common endpoint weights in partial summation, the sum of $1/\phi(q)$, and both kinds of nonprimitive prime-power terms. Neither B_Q nor its remainder is a prime minorant.

This is an identifiable arithmetic component beyond the square-root divisor level. It is not an evaluation of the whole covariance. Summing (1) over every h gives $O_A(HX \log^{-A} X)$, whereas the needed fluctuation scale is $X \log X$. Since H is between $X^{1/6}$ and $X^{2/7}$, no fixed logarithmic saving absorbs that loss. Logarithmically bounded shift packets are controlled at the required scale, but that does not evaluate the entire natural packet.

The report isolates the signed progression discrepancy sum D_Q explicitly. Under RH it proves

$$\sum_{h \leq CH} (C_{Q,h} - M_{Q,h}) = D_Q(X, T) + O(H\sqrt{X} \log^4 X). \quad (2)$$

The last error is $o(X \log X)$, uniformly over the α interval. The new estimate still needed for this selected component is $D_Q=O(X \log X)$. The remaining shifted bilinear form B_{rest} , the two support sums in M_Q , other spatial/shift ranges, and the continuous centering remain separate obligations. In particular, the source's multiplicative-convolution bilinear theorem cannot simply be applied to $\Lambda(qm+h)$, which depends jointly on all three variables.

The [independent source and proof review](#) accepts (1) and (2) with these scopes. The exact script checks 300 formal Mobius – log identities and five progression/discrepancy decompositions. These finite examples test algebra, not the source's asymptotic distribution estimate.

3. A new two–prime interaction was tried and failed

Archive inspection showed that the proposed resummed multiplicative prime profile had already been tried. Its old best continuation did not beat the existing larger polynomial trial. Repeating that scan was stopped before computation.

The replacement uses a genuinely different fixed arithmetic mark. For $n \leq L$, let C count distinct prime divisors p with $p^3 > L$, and $D = C(C-1)/2$. Then C is in $\{0, 1, 2\}$, while D is in $\{0, 1\}$. The new family is

$$r_L(n) = d_\ell(n)[F(v_n, S(n)) + D_L(n)J(v_n, S(n))], \quad \ell = 27/25.$$

When $D=1$ there is a unique unordered pair $p < q$ of large prime divisors and $n = pqm$ with $m < L^{1/3} < \min(p, q)$. This supplies an exact coprime starting decomposition. A singly marked prime at the same threshold does require a repeated-prime error; that error is explicitly $O((\log L)^a L^{-1/3})$, with $a = \ell^2$. The full three-state insertion calculus keeps the mixed event involving one background and one inserted large prime.

The [arithmetic derivation](#), accepted in a [separate root review](#), extends the previously reviewed fixed-moment and signed operator-truncation argument. It does not assume a general Fock limit, growing degree, or uniformity over infinitely many marks.

One fixed 30-dimensional span was tested: 20 unmarked features and ten D -marked features. Every coefficient was optimized in that span at quadrature orders 20 and 32; ℓ and the threshold were fixed.

Trial	Order 32 half–gap margin
Matched 20-feature baseline	-0.0146549380840028
New 30-feature double-prime interaction	-0.0146549114371551
Earlier best 48-feature trial, approximate	-0.0146547256

The new floating gain over its matched baseline is about $2.66e-8$, with a scaled Gram condition near $5.36e7$. It is not interval-certified. The new span does not contain the historical 48-feature span and performs worse than that historical value. The deficit to the required zero margin is still about 0.01465.

All coefficients and full M/G matrices are retained. The frozen rational vector is provably nonzero and has positive limiting mass independently of the numerical Gram matrix. A single actual-integer operator evaluation at $L=100000$, retaining every prime-power entry, has margin -0.0374094621535042 . Exact finite checks include 12 unordered decompositions, 132 coprime insertion triples and 108 count-state identities. These are checks of this fixed trial, not a global obstruction to resonance and not actual zeta-zero observations.

4. Prime powers can be removed at both relevant scales

The [prime-power estimate](#) is an elementary infinite-tail argument, accepted in an [independent proof review](#). It is a nuisance-term bound, without a novelty claim.

Let $\delta = c/\log T$, $\sigma = 1/2 + \delta$, $N = \lfloor T/\log^6 T \rfloor$, $R_c = -\zeta'/\zeta(s) - \sum_{n \leq N} \Lambda(n)n^{-s}$, and

$$U_{c,N}(t) = \sum_{\substack{p \text{ prime}, k \geq 2 \\ p^k > N}} (\log p) p^{-k\sigma - ikt \log p}.$$

Uniformly for $0 < \delta \leq 1/4$ and $N \geq 4$, without RH,

$$\|U_{c,N}\|_2^2 \ll TN^{-1/3} \log^4(2N) + \delta^{-4}. \quad (3)$$

Squares are treated by a convergent diagonal/off-diagonal mean-square expansion; the near-pair harmonic sum gives δ^{-4} . Higher powers admit a sufficient absolute tail bound using their sparse counting function. No infinite polynomial theorem with a divergent remainder is invoked.

Under RH, replacing R_c by $R_c - U_{c,N}$ changes its normalized energy by $O(a_T)$, where

$$a_T = N^{-1/6} \log^2(2N) + T^{-1/2} \log^2 T = o(1).$$

The bound is uniform for $1 \leq c \leq \log T/4$, and holds with a c -dependent constant for each fixed $c > 0$. The remaining continuation uses exactly the genuine-prime error $\theta(x) - x$, including its endpoint and pole term. This removes prime powers from the open arithmetic comparison; it does not improve its limiting constant.

5. The mesoscopic target is a first correction, not merely a leading law

Set $b = 2c$ and $r_T(b) = |R_-(b/2)|^2 / (T \log^2 T)$. A sufficient two-width asymptotic is

$$r_T(b) = \frac{e^{-b}}{b} + o\left(\frac{e^{-b}}{b^2}\right) \quad (4)$$

uniformly on a suitable slowly growing range, including the second width. Relative $o(1)$ is insufficient: the distinguishing signal requires relative $o(1/b)$.

The [source and rate audit](#) checks Theorem 5, Lemma 16 and Section 4.2 of [Carneiro – Chandee – Chirre – Milinovich, On Montgomery's pair correlation conjecture: a tale of three integrals](#). Its proof-level finite- T errors can fit below the signal on a sufficiently slow diagonal. Thus a blanket claim that those errors always swamp the signal would be wrong.

The real obstruction is the limiting lower bound: its deficit from sine is of order $e^{-(b)}/b$, while sine minus ACUE is of order $e^{-(b)}/b^2$. The known lower bound therefore misses the required correction by a factor of order b , even when its height error is negligible. The upper deficit is larger by order b^2 .

To remove the general AH near-diagonal nuisance, define

$$C_T(b) = b^2 [2 \sinh b r_T(b) - 2 \sinh(2b) r_T(2b) - \frac{1}{2b}].$$

The sine prediction tends to zero; the AH prediction tends to $-3/4$. The nuisance cancels exactly. The reviewed sufficient arithmetic target is: find increasing $G(T)$ tending to infinity with $G(T) = o(\log \log T)$ such that

$$\lim_{B \rightarrow \infty} \liminf_{T \rightarrow \infty} \inf_{B \leq b \leq G(T)} C_T(b) > -\frac{3}{4}. \quad (5)$$

Fixed-width AH convergence permits an existential stepwise slow diagonal inside this envelope, giving a contradiction if (5) were proved. It does not justify choosing a prescribed rate such as $\sqrt{\log \log T}$. The [independent review](#) states these quantifiers explicitly.

The prime-power replacement from (3) remains negligible even after the amplification: its effect on C_T is $O(b^2 e^{(2b) \alpha_T}) = o(1)$, uniformly for $2 \leq b \leq G(T) = o(\log \log T)$. Thus (5) can be pursued using genuine primes. No estimate (4) or (5) is established here.

6. Reproducibility, record ownership and next decision

All 28 source files, totaling 1,174,410 bytes, are retained verbatim in the adjacent local [Astra-Local-Archive/round9-originals/](#). The public folder retains 26 files; the two omitted third-party PDF/text bodies remain local with URL and SHA256 receipts. The existing 186 primary PDF/text are separately retained under [round9-external-sources/](#). No author files were edited for publication. Historical “review pending” labels are superseded by the separately hashed later reviews and the current claim ledger.

The [intake manifest](#) records every original. The [bounded integration replay](#) parses the four scripts, runs one fresh order-32 calculation and the three check scripts in a temporary copy, compares five JSON outputs, and compares the full recomputed M/G arrays. Timing and temporary source-file paths are the only excluded metadata; the order-20 matrices are checked as saved evidence, not advertised as newly recomputed.

Run from the repository root with Python, NumPy and SciPy:

```
python3 research/logs/round9-integration/recheck.py --prime-gap-source-dir /path/to/retained/round9-external-sources
python3 tools/verify_manifest.py
```

The replayer sets the portable Round 7 dependency path explicitly, verifies its pinned hash, and stages the primary references at the unmodified author's expected relative location. The primary references are inputs to provenance checking; finite arithmetic identities do not constitute a proof of their analytic theorem. Original evidence is unchanged by replay.

The next bounded task is to estimate the actual sum across shifts, testing smooth completion before paying the full H loss. A second lane checks whether known prime short-interval mean squares control the required signed residual correction at sufficient precision. Independent agents challenge the arithmetic hypotheses and normalization. These tasks remain in progress and are not part of this checkpoint's accepted claims.

Postponed: further scans of the failed feature, generic positivity countermodels already covered by Round 8, prime-gap coefficient sweeps, new Fable sessions, and another large PDF rebuild for a negative diagnostic. The 333-page public handoff, 381-page local handoff and 59-page supplement keep their explicit earlier checkpoints. Reverting this research commit removes the new slice without rewriting prior results. The open arithmetic inequalities, rather than the volume of documentation, determine whether the programme advances toward a famous conjecture.

Current report 23: A concrete 186–to–covariance transfer beyond the square–root divisor level

Collection: R9 — complementary moduli, genuine-prime tails, and the edge.

Source: [research/dyson/round9/factorization-covariance/COMPLEMENTARY_MODULI_TYPE_I_BRIDGE.md](https://research.dyson/round9/factorization-covariance/COMPLEMENTARY_MODULI_TYPE_I_BRIDGE.md).

SHA-256: `982039f0e163b84c1c5b8f2b52f215eb40e7b89863085f2840c039853606f39a` . Git blob: `f09e89d10ab31d35f0d86782d7929a7db26e21f1` . Original bytes: 14048.

A concrete 186–to–covariance transfer beyond the square–root divisor level

Date: 2026-09-05. Status: a proved application of a published distribution estimate, an exact divisor decomposition, and an explicit missing shifted estimate. This does **not** improve a zeta pair-correlation bound or prove the Round 8 target.

The 186 input controls an identifiable part of the additive covariance: the correlation of a prime with a selected Möbius – log divisor sum on complementary, triply densely divisible moduli up to $X^{0.523}$. The error is $O_A(X \log^{-A} X)$ per shift, uniformly in the relevant range. This saves enough for shift packets of logarithmically bounded total weight, but the proved bound does not save enough after all $H = X/T$ shifts are accumulated at the fluctuation scale $X \log X$.

The application uses the ordinary mathematical theorem in the primary paper. It makes no assertion that every assumption in its Lean formalization has been formally discharged.

1. Exact primary input

The primary source is [OpenAI, Improved short gaps between primes, 30 August 2026](#). Printed-page locations are:

- Definition 2.1, p.4: the recursive all-allocation class $D^{(r)}(Y)$.
- Lemma 2.2 and Proposition 2.3, pp.4 – 5: the cubic prime-factor test and complementary conditions for the actual least common multiple $[D, E]$, including shared primes.
- Equation (2.3), p.6: progression discrepancy with its coprime principal term.
- Equation (2.5), p.7: a coherent primitive class outside the modulus sum, uniform also in the prime set and that class.
- Corollary 2.19, p.11, equations (2.15) – (2.16): full-prime/ Λ distribution for order three when $240\omega + 80\delta < 3$, uniformly on subintervals.
- Proposition 2.18, pp.10 – 11, equation (2.14): the underlying multiplicative-convolution bilinear estimate, distinguished from the shifted form in Section 7.

The official repository was previously pinned at commit 61340d0b74163003b32756bb16e91d9209a5e330. The accompanying JSON hashes the local primary PDF and text; neither is changed here.

Choose

$$\omega = \frac{3}{250}, \quad \delta = \frac{1}{1000}, \quad \varepsilon = \frac{1}{1000}. \quad (1)$$

Then $240\omega + 80\delta = 74/25 = 2.96 < 3$, and $1/2 + 2\omega - \varepsilon = 523/1000$. Thus, for any finite prime set I and one coherent primitive class $a \bmod P_I$, the source gives

$$\sum_{\substack{q \leq X^{523/1000} \\ q \mid P_I \\ q \in D^{(3)}(X^{1/1000})}} |\Delta(\Lambda \mathbf{1}_J; a \bmod q)| \ll_A X(\log X)^{-A} \quad (2)$$

uniformly for common subintervals $J \subset [X, 2X)$, where

$$\Delta(f; a \bmod q) = \sum_{m \equiv a \pmod{q}} f(m) - \frac{1}{\phi(q)} \sum_{(m,q)=1} f(m). \quad (3)$$

This is not a maximum over unrelated classes inside the modulus sum.

2. A complementary modulus family above the square-root range

Use Z for the threshold called X in Proposition 2.3. Set

$$Y = X^{1/1000}, \quad Z = X^{1/2}, \quad A_0 = B_0 = X^{501/2000}, \quad f(p) = g(p) = p^{3/2}. \quad (4)$$

Let Q_X be any set of distinct $q = [D, E]$, with D, E positive and squarefree, satisfying

$$\begin{aligned} D, E &\leq X^{523/2000}, & [D, E] &> X^{1/2}, \\ p^{3/2} D_{\geq p} &\leq X^{501/2000} & (p \mid D, p > Y), \\ p^{3/2} E_{\geq p} &\leq X^{501/2000} & (p \mid E, p > Y). \end{aligned} \quad (5)$$

The opposite-root guards are automatic because the functions and budgets agree and each owner tail is at least one.

Moreover $A_0 B_0 = ZY$, $fg = p^3$, and both functions are nondecreasing. Proposition 2.3 yields

$$X^{1/2} < q \leq DE \leq X^{523/1000}, \quad q \in D^{(3)}(X^{1/1000}). \quad (6)$$

The condition is $[D, E] > Z$, not $DE > Z$. Repeated representations of q are counted only once.

This family contains nonsmooth examples for all sufficiently large X . Take distinct primes $p, r \in [X^{.089}, X^{.09}]$ and disjoint squarefree Y -smooth products $d_0, e_0 \in [X^{.170}, X^{.171}]$; put $D = pd_0, E = re_0$. The prime number theorem supplies p, r . To obtain the smooth products, successively multiply unused primes below Y until reaching $X^{.170}$; each overshoot is at most Y , and the product of the available primes is much larger than $X^{.342}$. Then

$$X^{.259} \leq D, E \leq X^{.261}, \quad X^{.518} \leq q = DE \leq X^{.522}, \quad p^{3/2} D_{\geq p} = p^{5/2} \leq X^{.225} < X^{.2505}, \quad (7)$$

with the same inequality for r . Only the two large primes activate the predicates. The exact check verifies these exponent margins, not the density of the resulting family.

3. The exact divisor piece and a transfer lemma

The elementary convolution identity gives

$$\Lambda(n) = \sum_{q \mid n} \mu(q) \log(n/q) = B_Q(n) + B_{\text{rest}}(n), \quad B_Q(n) = \sum_{\substack{q \mid n \\ q \in Q_X}} \mu(q) \log(n/q). \quad (8)$$

This is a definite Type I divisor piece with a smooth cofactor coefficient $\log(n/q)$. Both pieces are signed; neither is asserted to be a prime minorant.

Fix $C > 0$ and

$$X = T^a, \quad \frac{6}{5} \leq a \leq \frac{7}{5}, \quad H = X/T = X^{1-1/a} \in [X^{1/6}, X^{2/7}]. \quad (9)$$

For $1 \leq h \leq CH$, let w_h be real or complex, C^1 , and supported in a fixed compact subinterval of $(X, 3X/2)$, with

$$\|w_h\|_\infty + \int |w_h'(u)| du \leq C_w \quad (10)$$

uniformly in X, T, h . For large X , $m = n + h$ stays in $[X, 2X)$. Define

$$C_{Q,h} = \sum_n \Lambda(n+h) B_Q(n) w_h(n), \quad (11)$$

$$M_{Q,h} = \sum_{\substack{q \in Q_X \\ (q,h)=1}} \frac{\mu(q)}{\phi(q)} \int w_h(u) \log(u/q) du. \quad (12)$$

Transfer lemma. Under the published estimate (2), for every fixed $A > 0$,

$$\boxed{C_{Q,h} = M_{Q,h} + O_A(X \log^{-A} X)} \quad (13)$$

uniformly for the stated shifts and weights. This is unconditional when the ordinary distribution theorem is taken as input. RH is needed for the separate zeta correspondence, not this lemma.

The principal term involves only two explicit support sums, $\sum_{q \in Q_X, (q,h)=1} \mu(q)/\phi(q)$ and $\sum_{q \in Q_X, (q,h)=1} \mu(q) \log q/\phi(q)$. No limiting evaluation of these sums is claimed. Replacing them with a universal singular-series constant would require another proof.

4. Proof, including coherence and prime–power exceptions

For $(q, h) = 1$, substitute $m = n + h$:

$$\sum_{q \mid n} \Lambda(n+h) w_h(n) \log(n/q) = \sum_{m \equiv h \pmod{q}} \Lambda(m) w_h(m-h) \log((m-h)/q). \quad (14)$$

Let I be the union of primes occurring in the selected moduli, with primes dividing h removed. Then h is one primitive class modulo P_I , restricting coherently to every remaining q . Uniformity in I and h follows from the source's definition (2.5), so (2) applies to this subset.

For weighted partial summation use

$$w_h(u) \log(u/q) = w_h(u) \log u - (\log q) w_h(u). \quad (15)$$

There are just two common endpoint weight functions, with coefficients bounded by $\log X$. Their sup norms and total variations are $O(\log X)$. Integrating the source's modulus sum at a **common endpoint** is legal. No sum of separate endpoint suprema is substituted for it. The extra logarithm is absorbed by the arbitrary fixed saving in (2). Thus the coprime part of (11) equals

$$\sum_{\substack{q \in Q_X \\ (q,h)=1}} \frac{\mu(q)}{\phi(q)} \sum_{(m,q)=1} \Lambda(m) w_h(m-h) \log((m-h)/q) + O_A(X \log^{-A} X). \quad (16)$$

The prime number theorem with its classical error smaller than every fixed negative power of $\log X$, followed by partial summation, replaces the unrestricted m-sum by the integral in (12). Summing the errors over q is harmless because $\sum_{q \leq X^{523}} 1/\phi(q) \ll \log X$, and the weights cost only fixed logarithmic powers.

Removing $(m, q) = 1$ from that principal sum adds prime powers $m = p^j \asymp X$ with p dividing q . Since $p \leq q < X^{523} < X$, necessarily $j \geq 2$. Their total von Mangoldt mass is uniformly $O(\sqrt{X} \log^2 X)$. The logarithmic weight and the sum of $1/\phi(q)$ leave an error $O(\sqrt{X} \log^{\alpha(1)} X)$, absorbed by the stated error. These are all primes dividing q , not just those dividing h .

The original terms with $(q, h) > 1$ are separate. If p divides q, h and q divides n , then $\Lambda(n+h) \equiv 0$ forces $n+h = p^l$, where p divides h . Uniformly for $h \leq CH$, there are $O(\log^2 X)$ such possibilities. For each n , the number of candidate divisors q is at most $\tau(n) \ll_\eta X^\eta$, for any fixed $\eta > 0$. Their total contribution is $O_\eta(X^\eta \log^{\alpha(1)} X)$, also absorbed. Thus the nonprimitive terms are estimated, not silently deleted. This proves (13).

5. The actual localized Round 7 kernel

For fixed $X \in C_C^\infty(1, 3/2)$, put

$$a_u(X) = \min\{(u/X)^{1/2}, (X/u)^{3/2}\}, \quad \text{sinc}_0(v) = \sin(v)/v, \quad \text{sinc}_0(0) = 1, \\ w_h(u) = X(u/X) a_u(X) a_{u+h}(X) \text{sinc}_0(T \log(1+h/u)). \quad (17)$$

For $h \leq CH$, the argument is bounded and its u -derivative is $O_C(1/X)$. Thus (10) holds uniformly throughout the α interval. This signed weight is allowed by the lemma.

The localized off-diagonal prime-prime term in the covariance is exactly

$$\frac{2}{X \log T} \sum_{1 \leq h \leq CH} \sum_n \Lambda(n+h) \Lambda(n) w_h(n). \quad (18)$$

Substitution of (8) supplies the selected component (11) and an exact remainder. The factor two accounts for the two orders of each off-diagonal pair. All bounds are uniform for integration against the fixed Round 7 bump in α .

This is a stated component of the prime-prime term. Other n -ranges, shifts beyond CH , and the continuous-mean terms of the **centered** covariance have not been discarded.

6. A proved summed inequality, and the missing power

For arbitrary complex b_h supported on $1 \leq h \leq CH$, (13) gives

$$\left| \sum_h b_h (C_{Q,h} - M_{Q,h}) \right| \ll_A X (\log X)^{-A} \sum_h |b_h|. \quad (19)$$

In particular, any shift packet with $\sum_h |b_h| \leq (\log X)^B$, B fixed, has error $\mathcal{O}(X \log X)$. This includes logarithmically many sampled shifts anywhere up to CH . It is a valid component-level arithmetic input.

For the whole natural packet $b_h = 1$, this only gives

$$O_A(HX \log^{-A} X). \quad (20)$$

After normalization in (18) it becomes $O_A(H \log^{-A-1} X)$, which does not tend to zero for any fixed A since $H \geq X^{1/6}$. Arbitrary fixed logarithmic savings cannot be interpreted as an X -dependent power saving. This diagnoses the limitation of the estimate, not the size of the actual error.

7. The first unproved shifted forms

Define the weighted progression discrepancy

$$D_Q(X, T) = \sum_{1 \leq h \leq CH} \sum_{\substack{q \in \mathbb{Q}_X \\ (q,h)=1}} \mu(q) \Delta(\Lambda(m) w_h(m-h) \log((m-h)/q); h \bmod q). \quad (21)$$

The sequence inside Δ is a function of m with finite support. There is an aggregate identity with an error below the covariance scale. Under RH, which is already assumed for the Round 7 zeta correspondence,

$$\sum_{1 \leq h \leq CH} (C_{Q,h} - M_{Q,h}) = D_Q(X, T) + O(H\sqrt{X} \log^4 X). \quad (22)$$

Indeed the RH estimate $\psi(y) - y = O(\sqrt{y} \log^2 y)$, with weighted partial summation, gives a principal-term error $O(\sqrt{X} \log^3 X)$ per q before the factor $1/\phi(q)$. Summing that factor costs one more logarithm. The principal-term prime-power deletion has the same upper bound. The original nonprimitive error is $O_\eta(X^\eta \log^{O(1)} X)$ per h and can be absorbed into this bound by fixing, for example, $\eta = 1/4$. Summation over h proves (22). Since $H \leq X^{2/7}$, its error is $O(X^{11/14} \log^4 X) = o(X \log X)$.

This use of RH is explicit: multiplying only the unconditional per-shift PNT error, however strong in logarithms, by H would not justify the same conclusion.

Consequently an additional estimate

$$D_Q(X, T) = o(X \log X), \quad X = T^\alpha, \quad 6/5 \leq \alpha \leq 7/5, \quad (23)$$

uniform in α , would, under RH, evaluate the whole selected divisor piece at the desired precision. The source supplies only (20). Equation (23) requires cancellation across varying additive residues h , with the specified weight, or another argument of comparable strength.

The untouched divisor remainder is also explicit:

$$B_{\text{rest}}(X, T) = \sum_{1 \leq h \leq CH} \sum_{\substack{q, m \geq 1 \\ q \in \mathbb{Q}_X}} \mu(q) (\log m) \Lambda(qm + h) w_h(qm). \quad (24)$$

The support makes it finite. This shifted bilinear form is not covered directly by Proposition 2.18: $\Lambda(qm + h)$ depends jointly on q, m, h and is not an independent coefficient sequence in the multiplicative convolution $\alpha * \beta$. A Vaughan or Heath – Brown subdivision could reorganize (23), but a bound for the resulting shifted forms would still need proof.

Even (23) would not alone settle Round 7. The remainder (24), the support sums in (12), the omitted ranges, and the continuous centering remain. The actual successful transfer here is (13), (19), and the aggregate identity (22), with every source hypothesis checked; (21), (23), and (24) locate the unproved arithmetic task.

8. Verification and stopping point

The companion `check_divisor_bridge.py` uses only the Python standard library. It checks the parameter inequalities and the nonsmooth-family exponent margins with exact fractions. It represents $\log n$ as a formal prime-log vector, and products as degree-two formal polynomials, to verify the convolution identity, progression reindexing, and the separation into discrepancy, coprime principal term, and nonprimitive contributions on finite examples. There is no floating-point tolerance.

The small examples test algebra and conventions. They do not satisfy the large- X source thresholds and provide no numerical evidence for the distribution theorem or for zeta pair correlations.

No prior round was modified; no prime-gap sweep or new zeta-data fit was performed.

Current report 24: Independent review of the complementary–modulus covariance component

Collection: R9 — complementary moduli, genuine-prime tails, and the edge.

Source: [research/dyson/round9/factorization-covariance/INDEPENDENT_BRIDGE_REVIEW.md](#).

SHA-256: [86844e29f08ca834ebdeb04bb7d2df79fa28f37b09bfed2fcc8c5d1534c50b6](#) . Git blob: [ad2752d801935a57097a215f03d98c03b57b70c1](#) . Original bytes: 11399.

Independent review of the complementary–modulus covariance component

Date: 2026-09-05. Reviewer: the independent residual/arithmetic agent. Reviewed final report:

[COMPLEMENTARY_MODULI_TYPE_I_BRIDGE.md](#), SHA256

[982039f0e163b84c1c5b8f2b52f215eb40e7b89863085f2840c039853606f39a](#) .

Verdict: accepted as an ordinary application of the stated primary distribution theorem, with the scope of equations (13), (19), and (22) exactly as written. I found no missing modulus hypothesis, incorrect residue maximum, lost logarithmic weight, or error in the RH aggregate remainder. The comparison to the Round 7 sinc kernel has the correct normalization. This review does not prove or independently reprove the source distribution theorem, certify its formalization, evaluate the remaining shifted discrepancy, or establish a new zeta bound.

The useful output is a selected divisor component with a uniform per-shift estimate, and a separate exact aggregate reduction under RH. The full natural shift packet still requires a new bound. The report correctly keeps that obstruction visible.

1. Primary hypotheses and modulus family

I checked the local primary PDF/text pair against the supplied hashes and read the precise source statements: Definition 2.1 and Proposition 2.3 on printed pages 4 – 5; the discrepancy definition (2.3) on page 6; the coherent-class convention (2.5) on page 7; Proposition 2.18 and Corollary 2.19 on pages 10 – 11. The primary reference is [Improved short gaps between primes](#).

The parameter choice $\omega=3/250$, $\delta=1/1000$ and $\epsilon=1/1000$ gives $240\omega+80\delta=74/25<3$ and modulus exponent $1/2+2\omega-\epsilon=523/1000$. The source explicitly grants the $r=3$ estimate for Λ and uniformly on subintervals, so neither the prime-indicator conversion nor endpoint uniformity is being invented in this application.

The complementary construction uses the actual least common multiple $q=[D,E]$. Its root budgets satisfy $A_0 B_0=Z Y$, and $f(p)=g(p)=p^{3/2}$ are nondecreasing with product p^3 . Since the two budgets are equal, each owner inequality also gives the opposite-root guard by dropping the tail factor, which is at least one. Squarefreeness and the strict $q>Z$

assumption are present. Consequently the specified q belong to the source's triply densely divisible class even if D and E share primes. Using $DE > Z$ in place of $[D, E] > Z$ would not have sufficed, but the report does not make that substitution.

Every modulus is counted once. This is essential: no uncontrolled count of its root representations appears in the coefficient. The nonsmooth example is legitimate for sufficiently large X : the needed two large primes exist by the prime number theorem, while the product of unused primes at most $X^{.001}$ eventually exceeds any fixed power of X . Greedy disjoint smooth products have the claimed at-most- $X^{.001}$ overshoot. The listed exponent inequalities then place the example between $X^{.518}$ and $X^{.522}$ and satisfy the owner predicates. This proves existence, not density or quantitative usefulness of the selected family.

2. Coherent residues and q -dependent weights

Fix h and remove from the prime set every prime dividing h . Each retained q with $(q, h) = 1$ is still a divisor of the new prime product, and h is one primitive residue class modulo that product. Its restrictions are exactly the required classes $h \bmod q$. The source's uniformity permits this prime set and coherent class to depend on X and h . The application never uses a maximum over separately chosen classes inside the modulus sum.

The progression reindexing $m = n + h$ is exact. For the weight depending on q , use

$$\begin{aligned} w_h(m-h) \log((m-h)/q) \\ &= F_h(m) - (\log q) G_h(m), \\ F_h(m) &= w_h(m-h) \log(m-h), \quad G_h(m) = w_h(m-h). \end{aligned}$$

These are two common functions of the endpoint, not one unrelated function per modulus. Their sup-plus-variation bounds are $O(\log X)$ and $O(1)$, respectively, uniformly in h . The $\log q$ coefficient is bounded by $.523 \log X$. Partial summation therefore bounds the discrepancy sum by integrals of the source's sum at a common endpoint. It does not replace a bound on $\sup_t \sum_q |\Delta_q(t)|$ by the stronger and unavailable $\sum_q \sup_t |\Delta_q(t)|$.

All extra losses here are fixed logarithmic powers and are absorbed by choosing the arbitrary source saving sufficiently large. This proves the weighted coprime progression reduction uniformly for the stated $C1$ weights, including complex weights by absolute values or separate real and imaginary parts.

3. Principal terms and both kinds of prime-power exception

After the distribution step, the principal term carries $1/\phi(q)$. The sum of these factors is $O(\log X)$, even for the full available modulus range. One elementary verification is

$$\begin{aligned} \sum_{n \leq Q} 1/\phi(n) \\ &= \sum_{d \leq Q} \mu(d)^2 / [d \phi(d)] H_{\text{floor}(Q/d)} \\ &\leq (1 + \log Q) \text{product}_p [1 + 1/(p(p-1))]. \end{aligned}$$

The product converges. Thus one does not pay the number of moduli when replacing the coprime prime sum by its integral.

The unconditional prime number theorem, with a saving stronger than every fixed negative power of $\log X$, suffices for the per-shift result. Multiplying its uniform dyadic error by the $O(\log X)$ variation of the weight and the $O(\log X)$ inverse-totient sum still yields $O_A(X \log^{-A} X)$ after adjusting the initial exponent. This gives precisely the displayed main term, including both support sums with $\mu(q)/\phi(q)$ and $\mu(q)\log(q)/\phi(q)$. Neither sum is evaluated as a universal singular-series constant.

There are two different exceptional sets, correctly separated in the report:

- Removing $(m,q)=1$ from the principal sum adds $m=p^j$ around X with p dividing q . Since $p \leq q \leq X^{.523} < X$, these have $j \geq 2$. This includes all p dividing q , whether or not p divides h . Bounding by all prime powers gives von Mangoldt mass $O(\sqrt{X} \log^2 X)$. The additional weight and inverse-totient sum cost at most two logarithms.
- In the original nonprimitive progression terms, a prime p dividing both q and h also divides $n+h$. If $\Lambda(n+h)$ is nonzero, then $n+h=p^j$ with p dividing h . Uniformly for $h \leq C H$ there are $O(\log^2 X)$ such prime-power possibilities. For each resulting n , the number of allowed divisor moduli is at most $\tau(n)$, bounded by $O_\eta(X^\eta)$. The resulting error is $O_\eta(X^\eta \log^{O(1)} X)$, for any fixed $\eta > 0$.

No nonprimitive main term was silently discarded. Both errors fit the per-shift bound, proving equation (13) given the stated source theorem.

4. RH aggregate identity: the square-root error is sufficient

For equation (22), the report appropriately changes its error input. Under RH, $\psi(y)-y=O(\sqrt{y}\log^2 y)$ uniformly on the relevant dyadic interval. Partial summation costs one logarithm and the q principal-term sum costs one more, giving

$$O(\sqrt{X} \log^4 X)$$

per shift. The principal prime-power deletion obeys this same upper bound. The nonprimitive error can be absorbed by first fixing $\eta=1/4$; any remaining fixed logarithmic power is eventually smaller than the available $X^{1/4}$ margin.

Summing these nuisance errors over $h \leq C H$ therefore gives $O(H \sqrt{X} \log^4 X)$. Since α lies in $[6/5, 7/5]$, $H=X/T \leq X^{2/7}$, so the total is at most $O(X^{11/14} \log^4 X) = o(X \log X)$, uniformly in α . Constants may depend on the fixed compact weight support and C , which the report allows.

This verifies equation (22): the remaining quantity is exactly the sum of the weighted coherent progression discrepancies, not an untracked PNT error. The RH power saving is used only for this principal-term and exception replacement. It does not give a power saving for the progression-discrepancy sum itself. The report correctly warns that multiplying only an unconditional logarithmic per-shift PNT error by H would not justify the same aggregate precision.

5. Actual kernel and accumulation at the covariance scale

For fixed α and $X=T^\alpha$, the real symmetric Round 7 kernel is

$$\begin{aligned} & 1/(X T \log T) a_u(X) a_v(X) \sin(T \log(u/v))/\log(u/v) \\ & = 1/(X \log T) a_u(X) a_v(X) \operatorname{sinc}_0(T \log(u/v)). \end{aligned}$$

At $v=u+h$, sinc_0 is even, so the sign of $\log(u/v)$ does not alter the expression. Summing the two orders of an off-diagonal pair yields the stated factor 2. Multiplying by $\chi(u/X)$ selects the declared localized component, and subsequent fixed-bump integration in α preserves these identities. It does not remove the other components.

On the compact support inside $(X, 3X/2)$, both $a_u(X)$ and $a_{(u+h)}(X)$ lie on their smooth upper branches for large X . Their derivatives are $O(1/X)$. For $z=T \log(1+h/u)$, $h \leq C X/T$ gives $z=O_C(1)$ and

$$z' = -T h/[u(u+h)] = O_C(1/X).$$

The sinc and its derivative are bounded on that fixed argument range. Hence the weight in (17) has uniformly bounded sup norm and total variation as required; the absolute-value bound does not require it to be positive.

The natural normalized covariance scale is $X \log T$, uniformly comparable to $X \log X$ in the stated α interval. Equation (19) is the honest consequence of the per-shift estimate: it pays $\sum_h |b_h|$. Logarithmically bounded packets therefore have negligible error at this scale. For the full packet $b_h=1$, the only available source bound is $O_A(H X \log^{(-A)} X)$, or $O_A(H \log^{(-A-1)} X)$ after covariance normalization. Since H is a fixed positive power of X , no fixed logarithmic saving makes that upper bound tend to zero. This proves insufficiency of that bound, not a lower bound on the actual error.

Thus the proposed new estimate for the aggregate shifted discrepancy, equation (23), remains unproved. Even proving it would only evaluate the selected divisor piece. The unselected divisor sum (24), its main support sums, the omitted n and h ranges, and the continuous centering of the covariance remain genuine obligations. Proposition 2.18 concerns a multiplicative convolution with independent coefficient sequences; it does not directly estimate $\Lambda(qm+h)$ as an independent sequence in one variable. The report preserves this distinction.

6. Reproduction and claim boundary

I reran the exact script in a temporary directory, with read-only links to the two primary files, leaving the original scripts and outputs unchanged. The 300 formal Mobius-log identities, five progression/discrepancy decompositions and parameter checks passed. At least one example has a nonzero nonprimitive prime-power contribution, so that branch is exercised. Replayed JSON matches the author's JSON exactly after removing only the temporary local source-path strings; no arithmetic field is omitted from the comparison.

The [independent receipt](#) pins the final report, script, author JSON and source hashes. The script uses formal prime-log vectors and exact fractions, not floating agreement. These small checks validate identities and conventions; they do not prove any asymptotic distribution estimate or test the large- X support conditions numerically.

The accepted mathematical content is therefore the precise partial-component estimate (13), its weighted packet consequence (19), and the RH aggregate decomposition (22). This is a valid application of the chosen published input beyond a square-root divisor cutoff. It is not a full divisor decomposition estimate, a new Montgomery/Dyson theorem, a new prime-gap bound, or a solution of the two-scale zeta target.

Current report 25: A fixed interaction between two large prime divisors

Collection: R9 — complementary moduli, genuine-prime tails, and the edge.

Source: [research/dyson/round9/multiplicative-profile/DERIVATION.md](#).

SHA-256: `0b717ee45e31abcd399ba48d58069b47bab6b1d9ea5086afed42ca9df6438438` . Git blob: `8c52d8da3faab437f6e69ebfed44e804015b5221` . Original bytes: 10851.

A fixed interaction between two large prime divisors

Date: 2026-09-05. Status: written ordinary arithmetic transfer extending the previously reviewed fixed-prime-moment argument. This new note awaits a separate full proof review. The diagnostic is negative and floating; no zeta-spacing theorem or numerical enclosure is claimed.

1. Exact arithmetic object

Fix $\tau=1/3$, $\ell \geq 1$ and $\alpha=\ell^2$. Write $d=d_\ell$, with $d(p^\ell)=(\ell)_e/e!$, and, for $n \leq L$, define

```
v_n=log(n)/log(L),
S_k(n)=sum_(distinct p|n) (log(p)/log(L))^k, k>=2,
C_L(n)=number of distinct prime divisors p|n with p^3>L,
D_L(n)=C_L(n)(C_L(n)-1)/2.
```

There are at most two such distinct primes. Thus C takes values 0,1,2 and D takes values 0,1, with

```
C^2=C+2D, CD=2D, D^2=D.
```

The new coefficient family is

```
r_L(n)=d(n) H(v_n,S(n),D_L(n)),
H(v,S,D)=F(v,S)+D J(v,S),
```

where F and J are fixed polynomials in finitely many variables. Every H in this family is uniformly bounded on $n \leq L$. Its coefficients, threshold, number of marks and degree remain fixed as L tends to infinity. This is not a growing occupation-state limit.

The trial fixes $\ell=27/25$. F has groups $1, S_2, S_3, S_2^2$, each multiplied by the five Legendre polynomials $P_j(2v-1)$, $0 \leq j \leq 4$. J has groups $1, S_2$, each multiplied by $P_j(6v-5)$. There are 20 unmarked and 10 D -marked coefficients. The marked radial basis is scaled to $[2/3, 1]$, since D vanishes below total mass $2/3$.

This D interaction is nonmultiplicative. On configurations with respectively zero, one and two large prime divisors, it takes values 0,0,1. It cannot be introduced by replacing each individual prime by one common scalar multiplier. It also jumps when the second large prime crosses the fixed threshold. No fixed finite polynomial in the continuous S_k

variables realizes that jump on the relevant open prime-size configurations. The computational baseline is nevertheless only the specified 20-dimensional span: the new 30-dimensional span does not contain the old best 48-dimensional span.

2. Unmarked and singly marked measures

For a labeled index list $I=(k_1,\dots,k_j)$, put $|I|=\sum k_i$ and

```
m_I(a) = Gamma(a)/Gamma(a+|I|)
    * sum_(set partitions pi of the labeled indices)
      a^(number of blocks) product_(B in pi) Gamma(sum_(i in B) k_i),
m_empty(a)=1.
```

The previously reviewed finite-mark arithmetic argument gives $E_v \text{product}_i S_{-}(k_i)=v^{|I|} m_I(a)$. Here E_v denotes the resulting conditional moment functional, not an assumed random-matrix model. The unmarked positive measures are first defined on all $n \geq 1$:

$$(\log L)^{-a} \sum_{n \geq 1} d(n)^2/n \delta_{\log(n)/\log(L)}.$$

Their local weak limit is $C_{\text{ell}}/\Gamma(a) \int_0^1 v^{a-1} dv$. Restriction to $v \leq 1$ occurs after this limit; the full measure is not truncated before taking its Laplace transform. Small polynomial marks can be deleted uniformly because $\sum_{p|n, p < L^\epsilon} u_p^k \leq \epsilon^{k-1}$. Fix ϵ , use the prime number theorem and the unmarked weak limit on the finite collection of marks, then send ϵ to zero.

For any polynomial P in the S_k , the C -marked limit is

$$E_v[C P(S)] \\ = a \int_0^1 v^{a-1} \int_0^1 (v-t)^{a-1}/t \\ E_{-}(v-t) [P(S+(t^k)_k)] dt,$$

with value zero for $v \leq \tau$. Unlike the Round 7 threshold $1/2$, this is generally an ASYMPTOTIC identity, not an exact coprime decomposition: $p > L^{1/3}$ can divide the remaining factor m when $n=pm$.

Here is the required error explicitly. For $\ell \geq 1$, $d(pk) \leq \ell d(k)$, including when p divides k ; hence

$$\sum_{m \leq L, p|m} d(m)^2/m \\ \leq a/p \sum_{k \leq L} d(k)^2/k.$$

The factors H and any fixed polynomial in the marks are bounded, and $C \leq 2$ on $n \leq L$. Replacing $d(pm)^2$ by $d(m)^2$ and using the additive mark insertion outside the coprime set therefore costs at most a fixed constant times

$$(\sum_{k \leq L} d(k)^2/k) * \sum_{p > L^\tau} p^{-2} \\ = O_{H,\ell}((\log L)^a L^{-\tau}).$$

After normalization this tends to zero. This same estimate handles a large designated prime appearing again among the polynomial background marks. Thus the formula does not import the stronger, false automatic-coprimality statement at the one-third threshold.

3. The genuinely new double mark has an exact decomposition

If $D_L(n)=1$, n has a unique unordered pair $p < q$ of distinct divisors exceeding $L^{1/3}$. Writing $n=pqm$ gives

$$m \leq L/(pq) < L^{1/3} < \min(p, q).$$

Consequently p and q occur exactly once and are automatically coprime to m . In particular

$$\begin{aligned} d(pqm)^2 &= a^2 d(m)^2, \\ S_k(pqm) &= S_k(m) + u_p^k + u_q^k. \end{aligned}$$

For every test polynomial Φ , exactly

$$\begin{aligned} \sum_{n \leq L} d(n)^2 / n D_L(n) \Phi(v_n, S(n)) \\ = a^2 \sum_{(L^\tau < p < q, pq \leq L)} 1/(pq) \\ \sum_{m \leq L/(pq)} d(m)^2 / m \\ \Phi(v_m + u_p + u_q, S(m) + (u_p^k + u_q^k)_k). \end{aligned}$$

Applying the same finite-product measure limit yields

$$\begin{aligned} E_v[D P(S)] \\ = (a^2/2) v^{1-a} \\ \int_{t > \tau, s > \tau, t+s < v} \\ (v-t-s)^{a-1} / (ts) \\ E(v-t-s) [P(S+(t^k+s^k)_k)] dt ds. \end{aligned}$$

It vanishes for $v \leq 2\tau$. The factor $1/2$ converts the unique unordered pair into an ordered integral. The removed equal-prime diagonal has an extra reciprocal-prime factor, is $O((\log L)^a L^{-\tau})$ before normalization, and also has zero limiting two-dimensional measure.

All prime thresholds stay fixed. The planes $t=\tau$, $s=\tau$, $t+s=v=1$, and the corresponding operator-insertion planes have zero limiting measure. No uniform pointwise estimate for every short background cutoff is asserted: use the joint product-measure weak limit first, restrict away from the total-mass boundary, and then remove this restriction. The residual density w^{a-1} is locally integrable for $a > 0$; the already available uniform total-mass bound controls the discarded shrinking background strip. This includes m near one.

4. Insertions and the three–state interpolation

Let $\chi(u)=1_{(u>\tau)}$. For a prime multiplier not already dividing the input integer,

$$\begin{aligned} C &\rightarrow C + \chi(u), \\ D &\rightarrow D + C \chi(u). \end{aligned}$$

For two distinct new prime multipliers,

$$\begin{aligned} C &\rightarrow C + \chi(u) + \chi(w), \\ D &\rightarrow D + C(\chi(u) + \chi(w)) + \chi(u)\chi(w). \end{aligned}$$

The polynomial S_k inserts u^k , or $u^k + w^k$, as usual. Define H_0 , H_u , H_w and H_{uw} by these exact rules together with their total-mass arguments v , $v+u$, $v+w$ and $v+u+w$.

The code does not treat C as binary. For an insertion product $q(C)P(S)$, its three possible background values determine it exactly:

$$\begin{aligned} E_v[q(C)P(S)] &= q(0) E_v[P] \\ &+ (q(1)-q(0)) E_v[C \ P] \\ &+ (q(2)-2q(1)+q(0)) E_v[D \ P]. \end{aligned}$$

For marked factors on the left and right, the code sets

$$\begin{aligned} q(c) &= \text{choose}(c+\text{number of large left insertions}, 2)^{(\text{left mark})} \\ &\quad * \text{choose}(c+\text{number of large right insertions}, 2)^{(\text{right mark})}. \end{aligned}$$

The exponents are zero or one. The interpolation identity is only used at $c=0,1,2$, where it is exact even if the expression for q is a higher-degree polynomial. Impossible background/insertion configurations have zero measure by the total-mass constraint. This construction retains the mixed terms involving a single background large prime as well as the event with two large inserted primes.

5. The arithmetic limiting quadratic forms

For $x_n = r_L(n)/\sqrt{n}$, retain the exact arithmetic creation matrix

$$\begin{aligned} (A_L)_n(p^e, m) \\ = 2 \sin(\pi e \log(p)/(2\log(L)))/(e \sqrt{p^e}), \quad p^e m \leq L. \end{aligned}$$

Define

$$\begin{aligned} I &= \int_0^1 v^{(a-1)} E_v[H_0^2] dv, \\ M_2 &= (2 \ell^2/\pi^2) \int_{v+u+w \leq 1} v^{(a-1)} \sin(\pi u/2)/u \sin(\pi w/2)/w \\ &\quad E_v[H_0 H_{uw} + H_u H_w] dv du dw, \\ M_3 &= (2/\pi^2) \int_{v+u \leq 1} v^{(a-1)} \sin^2(\pi u/2)/u E_v[H_0^2] dv du. \end{aligned}$$

Then, for fixed H with $I > 0$,

$$\begin{aligned} [||A_L(x)||^2 + x^T A_L^2 x] / [2\pi^2 ||x||^2 - 1/4] \\ \rightarrow (M_2 + M_3) / I - 1/4. \end{aligned}$$

The inherited proof is applicable because the added marks are bounded and their thresholds fixed. More explicitly, remove operator primes below L^ϵ and prime powers with $e \geq 2$ using the previously proved weighted Schur bounds, valid uniformly for arbitrary signed vectors:

$$\begin{aligned} ||A_L|| &= O_\ell(1), \\ ||A_{(p < L^\epsilon)}|| &= O_\ell(\sqrt{\epsilon} + (\log L)^{-1/2}), \end{aligned}$$

$$||A_{(e \geq 2)}|| = O_{\text{ell}}((\log L)^{-1/2}).$$

The standard norm inequalities for A^*A and A^2 show that the normalized quadratic forms change by $o(1) + O(\sqrt{\epsilon})$. With retained primes, collisions with the background or another designated prime have the reciprocal-square error just described. The fixed polynomial marks and bounded C, D do not invalidate that estimate. The ordered distinct-prime terms then use $d(mp)d(mq) = a d(m)^2$ and $d(m)d(mpq) = a d(m)^2$ on the coprime part, yielding the two terms of M_2 with the displayed coefficients.

There are two different diagonals. In A^*A , $p=q$ gives $r_L(m)^2/(mp)$ and survives: this is M_3 with uninserted H_0 squared. In A^2 , repeated p gives the extra p^{-2} and vanishes after the retained-prime restriction. No $g(u)^2$ or inserted D is attached to M_3 . First let L tend to infinity at fixed ϵ , then send ϵ to zero. This completes the written fixed-family extension; it has no growing-family or uniform-in-degree assertion.

The source interface remains Inoue, arXiv:2604.05733v1, Theorems 3 and 4 / combined Proposition 3 under RH, with $L = \text{floor}(T/(\log T)^2)$. The previously audited coefficient-uniform normalized error and $\theta = \log L / \log T \rightarrow 1$ comparison are unchanged by this bounded fixed resonator. A positive limiting margin could feed that source argument. The present negative trial supplies no such consequence. In particular it supplies neither a gap bound nor an AH refutation; multiplicities and near-zero AH pairs would still need the correct treatment after any future successful small-gap argument.

6. Quadrature and verification boundary

For the C moment set $\delta = v - \tau$ and $t = v - \delta z$. Its endpoint factor is δ^a ; the remaining z integral has Jacobi weight $z^{(a-1)}$. For the D moment put

```
delta=v-2tau,
t=tau+delta(1-z)s,
w=tau+delta(1-z)(1-s),
background=delta z.
```

The Jacobian is $\delta^2(1-z)$, so the endpoint factor is $\delta^{(a+1)}$, with residual weight $z^{(a-1)}(1-z)/(tw)$. The coefficient remains $a^2/2$. These are the substitutions implemented in `marked_values`.

The v integration is divided at τ and 2τ . For each v slab, the u integration is divided at τ and $1-v-\tau$ in their correct order, and the w integration is divided at τ where present. No positive-weight quadrature cell straddles an inserted-prime step. The v endpoint powers above are absorbed by Jacobi rules. Norm and M_3 use the same three-state decomposition with no inserted large primes.

This is deterministic quadrature, not a rigorous error enclosure. A 20/32 comparison and independent raw adaptive integrals check its implementation but do not certify the small observed gain. The separate exact checks establish finite arithmetic identities only. Full coefficients, M/G arrays, a frozen rational vector and the actual-integer operator evaluation are retained; none of these numeric checks is substituted for the transfer argument.

Current report 26: Round 9: duplicate profile avoided, two–prime interaction tested

Collection: R9 — complementary moduli, genuine-prime tails, and the edge.

Source: <research/dyson/round9/multiplicative-profile/REPORT.md>.

SHA-256: `ecc98ab471b4e0b24516d0315a59481a3e8b76140543a52dade81f1bb917fa84` . **Git blob:** `e97c09952b23aa1a1e5efe28733c4b211eddcc71` . **Original bytes:** 9353.

Round 9: duplicate profile avoided, two–prime interaction tested

Date: 2026-09-05. This directory keeps its initially assigned `multiplicative-profile` name so the task history and dependency paths remain traceable. The actual new experiment is a nonmultiplicative interaction between two distinct large prime divisors.

Outcome: the fixed 30-dimensional interaction trial has floating half-gap margin -0.0146549114371551 . It does not cross zero and remains worse than the old 48-feature best value, approximately -0.0146547256 . The written arithmetic extension is in [DERIVATION.md](#); separate independent review is pending. There is no new zeta theorem, no certified optimizer, and no global obstruction theorem.

1. The archive check changed the task before computation

The initial suggestion was a multiplicative prime-size profile such as a product of $\exp(\sum c_k u_p^k)$, summing all mark orders. The public Round 2 report [research/reports/resummed_prime_profiles.md](#) and scripts [research/prime-profiles/euler_profile_precise.py/.json](#) already implement precisely that mechanism. Their convolution-density recurrence resums all powers, rather than merely taking the old finite S_2/S_3 span. Recorded cases include one, two, three and five exponent coefficients, through u^6 . The best validated local continuation was approximately -0.0146638632200778 , with radial degree 9 and quadrature order 40; it was a continuum calculation with 110/140-term density comparison, not a completed all-orders arithmetic remainder theorem.

The literal product of $(1+\eta u_p^2)$ was not located as an executed trial. It would be inaccurate to claim that every such parameter was already tested. But its multiplicative one-prime mechanism is not new relative to the existing resummed Euler-profile approach. Under the task's stop-on-duplication instruction, no new exponential-profile or η scan was run.

The old Fock report contains broad bin-occupation states, but explicitly keeps its general operator-transfer issue separate. That is not a prior fixed-integer test of the new D mark below. Searches of the existing research reports and Round 7 arithmetic derivation found no earlier implementation of this fixed two-large-prime interaction. This is an

archive-coverage statement, not a mathematical novelty claim.

2. What was actually added

For $n \leq L$ set

$D_L(n)=1$ if at least two DISTINCT prime divisors p of n satisfy $p^3 > L$,
0 otherwise.

There are at most two. When $D=1$, the unique unordered pair $p < q$ gives $n=pqm$ with $m < L^{1/3} < \min(p,q)$, so the new double-mark starting identity is exact and automatically coprime. The singly marked background count C also occurs in mixed insertions; its one-third-threshold formula needs the explicit repeated-prime error in the derivation and is not mislabeled exact.

The new span has $\text{ell}=27/25$, radial degree four, the same 20 unmarked features as the matched Round 7 baseline, and ten D -marked features: D times 1 or S_2 , with radial basis $P_j(6v-5)$. The mark is an interaction: adding the first large prime leaves it zero, while adding a second switches it on. A three-state interpolation in the background count C supplies the mixed insertion forms. All thresholds and coefficient counts stay fixed in the arithmetic limit.

The matrices are newly integrated for this mark. Only generic labeled-partition moments, polynomial insertion expansion, Gauss-Jacobi nodes and the generalized-eigenvalue helper are imported from the pinned Round 7 script. The old marked matrices and the prime-gap 77-dimensional problem are not reused.

3. The bounded numerical decision

Fixed trial	Order 20 margin	Order 32 margin
Matched 20 unmarked features	−0.0146549380840023	−0.0146549380840028
New 30-feature D interaction	−0.0146549114371546	−0.0146549114371551

At order 32 the observed gain over its matched baseline is about $2.66468477e-8$. The new value is still roughly $1.858e-7$ below the old 48-feature best, and about 0.01465 short of the required zero margin. Those comparisons are numerical only; the larger historical span is not nested in this new 30-dimensional span.

The diagonally scaled mass Gram condition is approximately $5.35731565e7$ for the enlarged span. All 30 directions survive the fixed $1e-11$ relative eigenvalue cutoff. The final floating pencil residual norm is approximately $2.52e-16$ and the mass norm is 1 to rounding. A tiny residual verifies a floating solve, not the correctness of its quadrature or an enclosure of the top eigenvalue. In particular the $2.66e-8$ gain is not advertised as certified. The negative main deficit is much larger than that gain.

The baseline M/G blocks independently agree with the existing Round 7 order-40 unmarked blocks to maximum absolute differences $8.33e-16$ and $2.22e-15$. The two new quadrature orders agree in the enlarged margin to about $5.6e-16$. No further degree increase, threshold change, parameter sweep or optimization of ell was performed.

4. Fixed vector and independent checks

`fixed_rational_vector.json` contains all 30 integer coefficients with denominator 100000000, the exact ell and threshold, the complete ordered feature list and radial conventions. Its continuum quadrature margin is -0.0146549114371553 . Its polynomial value at $v=S=D=0$ is exactly $155237743/100000000$, so it is nonzero; positivity of the limiting mass follows already on a sufficiently small unmarked total-mass interval. The nonzero norm is not based solely on the numerical Gram matrix.

`check_two_prime_trial.py` records the following in `validation.json`:

- At integer cutoff 120, an exact Fraction identity matches all 12 unique unordered two-large-prime decompositions and their coefficient factor a^2 .
- The two insertion identities and divisor factors pass on 132 ordered coprime insertion triples.
- The Newton interpolation in $\{1, C, D\}$ passes 108 exact state checks, keeping $C=2$ distinct from $C=1$.
- At $a=(27/25)^2$ and total mass $v=.91$, two independent nested adaptive integrations of the raw $E[D]$ and $E[D^2]$ formulas agree with the substituted Jacobi expressions within $1.11e-16$ and $5.56e-17$. These are floating checks of different integration formulas, not interval certificates.
- A single actual-integer calculation at $L=100000$ uses the frozen vector and every prime-power entry of A_L . It has 11109 integers with $D=1$, 343614 matrix nonzeros, norm about 20.46127, A^*A/norm about 3.02145870, A^2/norm about 1.17491031 and margin -0.0374094621535042 . It is a finite arithmetic calculation at $\theta=1$, not a zeta-zero sample or proof of an asymptotic rate.

The exact checks deliberately use rational formal prime labels for polynomial identities. They do not claim that $\log(p)/\log(L)$ is rational or replace the actual logarithms used by the final finite-integer evaluation.

5. Files, reproduction and limitations

The two computation files are `two_large_prime_sector.py` and `check_two_prime_trial.py`. The former produces complete `two_large_prime_d4_q20.json/.npz` and `two_large_prime_d4_q32.json/.npz`; each NPZ retains the full symmetric numerator M and mass G . The latter freezes the rational vector and writes the validation record. A small `manifest.json` pins the new outputs and the relevant archived source files.

From this directory, with Python, NumPy and SciPy:

```
OPENBLAS_NUM_THREADS=1 OMP_NUM_THREADS=1 python3 two_large_prime_sector.py --order 20
OPENBLAS_NUM_THREADS=1 OMP_NUM_THREADS=1 python3 two_large_prime_sector.py --order 32
OPENBLAS_NUM_THREADS=1 OMP_NUM_THREADS=1 python3 check_two_prime_trial.py
```

The environment variable `ASTRA_LARGE_PRIME_SOURCE` can point to the pinned Round 7 `large_prime_sector.py`. Its expected SHA256 is checked before import. The current default is the local public-repository path; a publication copy may replace that default with a portable repository-relative path and must record that change explicitly. The old imported file is not edited. The only output fields expected to vary on replay are recorded run times, runtime/version metadata if added, and environment-dependent floating differences; an identical environment can demand exact arrays as a stronger diagnostic. No speed claim is made from these small local run times.

The arithmetic-transfer argument uses existing fixed-moment asymptotics and uniform operator truncations, plus the explicit new count identities. It does not assume a global Fock limit or an unsupported all-orders coefficient limit.

The optimization and the finite-integer sine calculations remain ordinary floating arithmetic. Review of the new proof and code is separate from numerical agreement, and no external novelty or formal verification claim is made.

This round is closed at a useful negative decision. It rules out neither other nonmultiplicative arithmetic features nor the full resonance method. Repeating the old exponential profile, adding more digits to this small gain, or treating it as evidence toward a historic conjecture is postponed. Further work needs a quantitatively different arithmetic direction or a new estimate for the actual out-of-band covariance.

Current report 27: Independent review: the fixed two–large–prime interaction

Collection: R9 — complementary moduli, genuine-prime tails, and the edge.

Source: `research/dyson/round9/multiplicative-profile/INDEPENDENT_REVIEW.md`.

SHA-256: `b88133d0b6dbd38d3c1a4c7a8fb70599ce6ef7968297a05f21105682c7ae6997` . **Git blob:** `324021cc5d9aedd7513274395c0f10ff4ce6965a` . **Original bytes:** 4454.

Independent review: the fixed two–large–prime interaction

Date: 2026-09-05. **Reviewer:** the root research lane, independently of the author. **Accepted as an ordinary fixed-family arithmetic extension and a negative floating experiment.** No interval certification, external novelty claim, or new zeta theorem is accepted or implied.

Reviewed author files:

- `DERIVATION.md` : SHA256 `0b717ee45e31abcd399ba48d58069b47bab6b1d9ea5086afed42ca9df6438438` .
- `two_large_prime_sector.py` : SHA256 `ed6fa274593a04d8a168a8597c76a994ad0595edb337995a7904a93b6a845de0` .
- `check_two_prime_trial.py` : SHA256 `61c4c5c92b5d502670fbf469fe61d422f69778369125ff569a28418aa3dab9ff` .

The root read the complete derivation, report, and two scripts. The review uses the previously independently reviewed Round 7 fixed-moment and weighted Schur estimates as inputs; it does not supply a new proof of those inherited results or a general Fock-space limit.

Arithmetic checks independent of the optimizer

1. For two distinct prime divisors exceeding $L^{1/3}$, their product leaves a cofactor less than either prime. The unordered double-mark decomposition, the coefficient a^2 , and the factor $1/2$ in the ordered limiting integral are correct. Three such distinct primes cannot divide $n \leq L$.
2. A *single* designated prime at this threshold need not be coprime to the cofactor. The author handles this explicitly. For $\ell \geq 1$, the ratio $d_{\ell}(p^{e+1})/d_{\ell}(p^e) = (\ell+e)/(e+1)$ is at most ℓ . Hence the background repeat estimate has an extra $1/p$, and summing designated primes supplies $\sum p^{-2} = O(L^{-1/3})$. Bounded fixed marks preserve this estimate. No automatic-coprimality claim is imported from the older threshold $1/2$.
3. C is three-valued, while D is binary. The identities $C^2 = C + 2D$, $CD = 2D$, $D^2 = D$ and the three-state Newton interpolation are correct. The two insertion rules include the mixed term from one old and one new large prime, as well as the event with two newly inserted large primes. Infeasible states disappear through total mass, not an artificial truncation of C .

4. The measure used before the Laplace limit is on all $n \geq 1$. Restriction to total mass at most one follows the weak limit. Fixed threshold planes have zero limiting measure, and the locally integrable residual density controls the small-cofactor strip. This avoids a false pointwise uniform asymptotic at a cofactor cutoff near one.
5. The bounded, fixed coefficient family permits the inherited signed operator truncations. The distinct-prime terms produce H_0 H_{uw} and H_u H_w . The A^*A same-prime term survives with uninserted H_0^2 ; the repeated-prime term in A^2 has the extra reciprocal-prime factor and vanishes. The displayed M_2/M_3 constants agree with the actual creation matrix normalization.

Implementation and interpretation

The two-mark substitution has Jacobian $\delta^2(1-z)$, residual density $(\delta z)^{(a-1)}$, and denominator tw , giving exactly $\delta^{(a+1)} z^{(a-1)}(1-z)/(tw)$. The one-mark endpoint exponent is a . The code divides the total-mass and insertion domains at each fixed threshold and uses the corresponding Jacobi factors. Its coefficient interpolation keeps all three background count states.

The rational frozen vector is nonzero without appealing to a numerical Gram matrix: its unmarked value at the origin is $155237743/100000000$. Continuity then gives a positive limiting norm on a sufficiently small unmarked interval.

The reported margin is approximately -0.0146549114371551 . The observed $2.66e-8$ improvement over the matched baseline is not certified, especially with a diagonally scaled Gram condition near $5.36e7$. The full negative deficit is about 0.01465 . The new span does not contain the earlier best 48-feature span, and its value is worse than that historical result. Thus this fixed trial supplies no small-gap or AH consequence. It also does not rule out a different interaction, a larger coefficient space, or the resonance method.

Finite exact checks validate the stated identities and cutoff conventions; quadrature agreement and the finite-integer experiment validate specific implementations. They do not prove the asymptotic transfer. Independent integration replay and full-array comparisons are recorded separately when this folder is published.

The author files retain their historical “review pending” wording. This separately hashed review records their later acceptance without silently rewriting the original research record.

Current report 28: Prime powers are negligible in the remaining residual energy

Collection: R9 — complementary moduli, genuine-prime tails, and the edge.

Source: [research/dyson/round9/prime-power-removal/PRIME_POWER_TAIL_ESTIMATE.md](#).

SHA-256: `e7fff21fa21285a20968c218e99b17b3bc634bf28b88268c002f406a0fdd6cd0` . **Git blob:** `55d4fa8c233145abee92164b9546cd1fcd231f33` . **Original bytes:** 8037.

Prime powers are negligible in the remaining residual energy

Date: 2026-09-05. **Status:** ordinary proof draft for independent review. This is a quantified nuisance-term estimate, not a new lower bound for the signed residual and not a conjecture solution. Its purpose is to let the remaining covariance calculation use genuine primes.

1. Statement

Fix $c > 0$, let T tend to infinity, and set

```
delta=c/log T, sigma=1/2+delta, s=sigma+it,
N=floor(T/log^6 T).
```

Let $H = -\zeta'/\zeta(s)$, $P_N = \sum_{n \leq N} \Lambda(n) n^{-s}$, and $R = H - P_N$ as in Round 8. Define the absolutely convergent prime-power tail

$$U_{c,N}(t) = \sum_{\substack{p \text{ prime}, k \geq 2 \\ p^k > N}} (\log p) p^{-k\sigma - ikt \log p}.$$

Uniformly for $0 < \delta \leq 1/4$ and $N \geq 4$,

$$\| U_{c,N} \|_{L^2(0,T)}^2 \ll TN^{-1/3} \log^4(2N) + \delta^{-4}. \quad (1)$$

The implied constant in (1) is absolute. The estimate itself does not assume RH. Under RH, the normalized residual energy changes by

$$\frac{|\| R - U_{c,N} \|_2^2 - \| R \|_2^2|}{T \log^2 T} \ll_c N^{-1/6} \log^2(2N) + \frac{\log^2 T}{\sqrt{T}} = o(1). \quad (2)$$

Consequently the fixed two-scale residual in Round 8 can be replaced by its genuine-prime counterpart. Equation (2) is a bound on the replacement error; it gives no positive constant improvement in the desired signed comparison.

2. Squares: an elementary mean-value bound with an infinite tail

Write $U_2(t) = \sum_{(p > \sqrt{N})} (\log p) p^{-(1-2\delta-2it)}$. For the sequence

$$a_n = (\log n) n^{-(1-2\delta)} \text{ if } n \text{ is a prime greater than } \sqrt{N}, \\ 0 \text{ otherwise,}$$

the series is absolutely convergent. Expanding the time integral, the diagonal is $T \sum |a_n|^2$. The integrated off-diagonal kernel has absolute value at most $1/|\log(m/n)|$ because the time frequency is $2\log(m/n)$.

For $n/2 \leq m \leq 2n$, $m \neq n$, uniformly in $0 < \delta \leq 1/4$,

$$\frac{|a_m a_n|}{|\log(m/n)|} \ll \frac{\log^2(2n)}{n^{1+4\delta} |m-n|}.$$

Summing the harmonic denominator over m and then over n gives

$$\ll \sum_{n \geq 2} \frac{\log^3(2n)}{n^{1+4\delta}} \ll \delta^{-4}.$$

For noncomparable m, n , the log denominator is at least $\log 2$, so the contribution is at most a constant times

$$\left(\sum_{n \geq 2} \frac{\log n}{n^{1+2\delta}} \right)^2 \ll \delta^{-4}.$$

These are convergent majorants for each $\delta > 0$. One may therefore first use finite sums and pass to the infinite sum by dominated convergence. No infinite polynomial mean-value theorem with a divergent error is invoked.

The diagonal tail satisfies

$$\sum_{p > \sqrt{N}} \frac{(\log p)^2}{p^{2+4\delta}} \leq \sum_{n > \sqrt{N}} \frac{(\log n)^2}{n^2} \ll N^{-1/2} \log^2(2N).$$

Thus

$$\|U_2\|_2^2 \ll TN^{-1/2} \log^2(2N) + \delta^{-4}. \quad (3)$$

In particular, the exponent four in the error is intentional. A sharper Hilbert-inequality estimate is unnecessary for this use. The argument above is self-contained and only uses the actual support of prime squares and elementary upper bounds for their weights.

3. Higher powers: absolute convergence is now strong enough

Let

$$A_3(x) = \sum_{\substack{p^k \leq x \\ k \geq 3}} \log p.$$

There are at most $\log(x)/\log(2)$ relevant exponents, and for each $k \geq 3$ the number of possible bases is at most $x^{1/3}$, with $\log p \leq \log x$. Hence, for $x \geq 2$,

$$A_3(x) \ll x^{1/3} \log^2(2x).$$

Partial summation and $\sigma \geq 1/2$ yield

$$\sum_{\substack{p^k > N \\ k \geq 3}} (\log p) p^{-k\sigma} \leq \sigma \int_N^{\infty} A_3(x) x^{-\sigma-1} dx \ll N^{-1/6} \log^2(2N). \quad (4)$$

For the last bound σ is in $[1/2, 3/4]$, so $\sigma - 1/3$ is bounded below by $1/6$ and all integration constants are uniform. The negative endpoint term from partial summation was discarded in the favorable direction. Equation (4) gives

$$\|U_{\geq 3}\|_2^2 \ll TN^{-1/3} \log^4(2N).$$

Combining this with (3), using $|u+v|^2 \leq 2|u|^2 + 2|v|^2$, proves (1). Squares require a mean-square argument; applying the same absolute-value argument to squares would lose the useful decay.

4. Replacement in the actual RH residual

The pointwise RH partial-fraction bound already proved in Round 8 gives $H(t) = O_c(\log^2 T)$ throughout $[0, T]$. The elementary finite-polynomial mean-value bound there gives $\|P_N\|_2^2 = O_c(T \log^4 T)$. Consequently

$$\|R\|_2 \leq C_c \sqrt{T} \log^2 T.$$

Set $a_T = N^{-(1/6)} \log^2(2N) + \log^2(T)/\sqrt{T}$. Equation (1), with $\delta = c/\log T$, gives $\|U\|_2/\sqrt{T} \leq C_c a_T$. By Cauchy-Schwarz,

$$\frac{|\|R - U\|_2^2 - \|R\|_2^2|}{T \log^2 T} \leq \frac{2\|R\|_2 \|U\|_2 + \|U\|_2^2}{T \log^2 T} \ll_c a_T + \frac{a_T^2}{\log^2 T}.$$

Since a_T tends to zero at the stated cutoff, this proves (2) for sufficiently large T . The stronger RH pair-correlation bound on $\|H\|$ is not needed. The same argument applies separately to $c=1$ and $c=1/2$, whose fixed hyperbolic coefficients preserve the $o(1)$ replacement error.

5. The remaining error really is the prime-counting error

Define $\theta(x) = \sum_{p \leq x} \log p$ and $E_1(x) = \theta(x) - x$. Subtracting the absolutely convergent prime-power series from the usual Stieltjes continuation gives

$$R - U_{c,N} = \frac{N^{1-s}}{s-1} - E_1(N)N^{-s} + s \int_N^{\infty} E_1(x) x^{-s-1} dx. \quad (5)$$

The finite polynomial on the left now removes only the primes up to N from the prime-only analytic continuation: explicitly it is

$$H(s) - \sum_{(p,k) \geq 2} (\log p) p^{-ks} = \sum_{p \leq N} (\log p) p^{-s}.$$

This equals $R-U$ by cancellation of the prime powers up to N . At integer N the theta endpoint includes a prime at N . RH and the elementary prime-power counting bound imply $E_1(x) = O(\sqrt{x} \log^2(2x))$, so the integral in (5) is absolutely convergent when $\operatorname{Re}(s) > 1/2$. The pole can be removed from the normalized energy as in Round 8.

Thus the non-negligible unknown comparison involves the same genuine-prime error $\theta(x)-x$ at the two damping widths. Neither prime powers nor a convention at a prime-power cutoff can supply the missing fixed positive gap above the AH value. A new additive-prime covariance estimate is still required.

6. Scope and provenance

The RH inputs are the logarithmic-derivative and prime-counting-error bounds already used and reviewed in Round 8. The prime-power estimates use arithmetic sparsity and elementary series estimates, not a generic point-process positivity assumption. No experimental optimization or new scalar-constant certificate was run for this lemma. An independent ordinary-proof review is required before its status is upgraded.

The resulting error tends to zero and hence leaves the numerical target unchanged. This is a cleanup estimate for the hard arithmetic problem, not progress toward $1/16$ by itself and not an assertion of novelty for prime-power removal.

7. Uniform corollary for the mesoscopic two-width statistic

Put $L = \log T$. The bound in (2) has an absolute implied constant throughout

$$1 \leq c \leq L/4.$$

Indeed $\delta \geq 1/L$ makes $\delta^{-4} \leq L^4$ in (1). The RH partial-fraction bound for the logarithmic derivative is $O(L^2)$ uniformly at distance at least $1/L$ to the right of the critical line, throughout $0 \leq t \leq T$ and $1/2 + 1/L \leq \sigma \leq 3/4$. The same elementary bound for the finite polynomial is uniform because its coefficients only decrease as σ increases. The argument in Section 4 therefore gives a uniform $O(a_T)$ replacement error, where

$$a_T = N^{-1/6} \log^2(2N) + L^2 T^{-1/2} \ll T^{-1/6} L^3 + T^{-1/2} L^2.$$

Write $r_T(b) = |R_{-}(b/2)|^2 / (TL^2)$, and let $r_T^{\text{prime}}(b)$ be the same normalized energy with $U_{-}(b/2, N)$ subtracted. Define

$$C_T(b) = b^2 \left[2 \sinh(b) r_T(b) - 2 \sinh(2b) r_T(2b) - \frac{1}{2b} \right],$$

and define its prime-only version by substituting r_T^{prime} . For $b \geq 2$ and $b = o(\log L)$, both $c = b/2$ and $c = b$ are in the uniform range, and

$$\left| C_T^{\text{prime}}(b) - C_T(b) \right| \ll b^2 e^{2b} a_T = o(1). \quad (6)$$

The conclusion is uniform on any range $2 \leq b \leq G(T)$ with $G(T) = o(\log L)$: the exponential factor is only $L^{o(1)}$, whereas a_T has a negative power of T . Thus the mesoscopic arithmetic target can also be stated using the genuine-prime continuation. This corollary removes a nuisance term at the amplified scale; it supplies no lower bound on either statistic.

Current report 29: Independent review of the actual prime–power tail estimate

Collection: R9 — complementary moduli, genuine-prime tails, and the edge.

Source: [research/dyson/round9/prime-power-removal/INDEPENDENT_REVIEW.md](https://research.dyson/round9/prime-power-removal/INDEPENDENT_REVIEW.md).

SHA-256: `0650ede489f745fa0eea7d7761ccd356f86b2eca7f9e5a5858d63bab644dcbb6` . **Git blob:** `5454bb68fd06136b20d51e1f747daecbbf614b32` . **Original bytes:** 8940.

Independent review of the actual prime–power tail estimate

Date: 2026-09-05. **Reviewer:** `yau_flow` , separately from the authoring coordinator.

Reviewed final author artifact: `PRIME_POWER_TAIL_ESTIMATE.md` , SHA-256

`e7ffff21fa21285a20968c218e99b17b3bc634bf28b88268c002f406a0fdd6cd0` . The initial complete review covered the corrected six-section version `b28cc13877d38d7cc83956127ace275df1f87585df8f2905fc309e77ae6d2349` ; the final appended uniform corollary and corrected formula markup were subsequently read and accepted. Section 7 below records the independent derivation of that corollary.

Verdict: the tail estimate, residual-energy replacement, and exact genuine-prime continuation are accepted as ordinary mathematical arguments. No missing endpoint, square-frequency factor, or uncontrolled infinite off-diagonal sum was found. The author has corrected the minor provenance sentence to acknowledge the RH prime-counting-error estimate as well as the logarithmic-derivative estimate. No numerical scan or new experiment was run.

1. Scope of the unconditional estimate

For $0 < \delta \leq 1/4$, $\sigma = 1/2 + \delta$, $N \geq 4$, and $T > 0$, the series

$$U(t) = \sum_{p^k > N, k \geq 2} (\log p) p^{-k\sigma - ikt \log p}$$

converges absolutely and uniformly in real t for each fixed $\delta > 0$. The sum over squares is majorized by $\sum_{n \geq 2} (\log n) n^{-1-2\delta}$; the higher powers converge even at $\delta = 0$. The claimed uniform estimate

$$\| U \|_{L^2(0,T)}^2 \ll TN^{-1/3} \log^4(2N) + \delta^{-4} \quad (1)$$

does not use RH. Uniformity means the implied constant is independent of δ , N , T in the displayed range, not that the second term remains bounded as $\delta \downarrow 0$.

2. Squares: frequency and infinite near–diagonal terms

For prime squares, the coefficients are $a_p = (\log p) p^{-1-2\delta}$, and the time phase is $e^{-2it \log p}$. Therefore

$$\left| \int_0^T e^{2it \log(m/n)} dt \right| \leq \frac{1}{|\log(m/n)|}$$

for $m \equiv n$. The factor one, rather than two, is correct because of the square frequency.

If $n/2 \leq m \leq 2n$, then $|\log(m/n)| \geq |m - n|/(2n)$. Uniformly in the allowed δ ,

$$\frac{|a_m a_n|}{|\log(m/n)|} \ll \frac{\log^2(2n)}{n^{1+4\delta} |m - n|}.$$

The remaining harmonic sum costs $O(\log(2n))$, giving

$$\sum_{n \geq 2} \frac{\log^3(2n)}{n^{1+4\delta}} \ll \delta^{-4}.$$

The comparison follows directly from integrals of $(\log x)^j x^{-1-4\delta}$, $0 \leq j \leq 3$. Their powers δ^{-j-1} are all absorbed by δ^{-4} on $0 < \delta \leq 1/4$. This checks the deliberate fourth power of the error.

For noncomparable pairs the denominator is bounded below by $\log 2$. The product majorant is

$$\left(\sum_{n \geq 2} \frac{\log n}{n^{1+2\delta}} \right)^2 \ll \delta^{-4}.$$

Both estimates legitimately enlarge prime sums to integer sums and remove the cutoff. Every resulting series converges for each positive δ . Finite partial sums can consequently be passed to the infinite series without invoking a mean-value theorem with divergent error.

The diagonal is bounded by

$$T \sum_{p > \sqrt{N}} \frac{(\log p)^2}{p^{2+4\delta}} \ll TN^{-1/2} \log^2(2N).$$

This proves the square part of the author's estimate, including its normalization and uniformity.

3. Higher powers and the endpoint sign

For $A_3(x) = \sum_{p^k \leq x, k \geq 3} \log p$, the elementary bound $A_3(x) \ll x^{1/3} \log^2(2x)$ is valid: there are at most $O(\log x)$ exponents, at most $x^{1/3}$ bases for each, and each weight is $O(\log x)$. No prime-counting asymptotic is assumed.

The exact partial-summation formula for the strict tail is

$$\sum_{p^k > N, k \geq 3} (\log p) p^{-k\sigma} = -A_3(N) N^{-\sigma} + \sigma \int_N^\infty A_3(x) x^{-\sigma-1} dx.$$

The endpoint term is nonpositive and can be discarded in this upper bound. This remains correct when N itself is a prime power, provided $A_3(N)$ includes that atom, as the definition does.

Since $\sigma - 1/3 \geq 1/6$, the integral is $O(N^{-1/6} \log^2(2N))$ uniformly for $\sigma \in (1/2, 3/4]$. Squaring the resulting uniform-in-time bound gives $O(TN^{-1/3} \log^4(2N))$. Combining the square and higher-power estimates with $|u + v|^2 \leq 2|u|^2 + 2|v|^2$ proves (1). Prime powers have a unique prime base; there is no duplicate counting between the $k = 2$ and $k \geq 3$ pieces.

4. Replacement in the actual residual

Now set $\delta = c/\log T$, with fixed $c > 0$, and $N = \lfloor T/\log^6 T \rfloor$. Under RH, the already proved bound $\|R\|_2 \ll_c \sqrt{T} \log^2 T$ applies. Taking a square root in (1) gives

$$\frac{\|U\|_2}{\sqrt{T}} \ll_c a_T := N^{-1/6} \log^2(2N) + \frac{\log^2 T}{\sqrt{T}}.$$

The exact norm expansion and Cauchy – Schwarz yield

$$\frac{|\|R - U\|_2^2 - \|R\|_2^2|}{T \log^2 T} \ll_c a_T + \frac{a_T^2}{\log^2 T}.$$

At the stated cutoff $a_T \rightarrow 0$, so its quadratic term is absorbed into $O_c(a_T)$ for all sufficiently large T . This proves the author's displayed error. The dependence on c , including the factor from δ^{-2} , is correctly contained in $\ll_c$.

Applying the estimate separately at the two fixed widths preserves an $\mathcal{O}(1)$ replacement error; it does not improve the signed target's limiting constant.

5. The exact genuine–prime identity

Let $U_{\text{all}}(s) = \sum_{p,k \geq 2} (\log p) p^{-ks}$, which is absolutely convergent on $\Re s > 1/2$. Decompose the finite von Mangoldt polynomial as

$$P_N = P_N^{\text{prime}} + P_N^{\text{power}}, \quad U_{\text{tail}} = U_{\text{all}} - P_N^{\text{power}}.$$

Then exactly

$$R - U_{\text{tail}} = H - U_{\text{all}} - P_N^{\text{prime}}.$$

This checks that powers $p^k \leq N$ cancel once, and the strict tail $p^k > N$ removes precisely the remaining powers. In particular, there is no off-by-one correction at a prime-power cutoff.

For $\theta(x) = \sum_{p \leq x} \log p$, the Stieltjes continuation of the prime-only function gives

$$R - U_{\text{tail}} = \frac{N^{1-s}}{s-1} - (\theta(N) - N)N^{-s} + s \int_N^\infty (\theta(x) - x)x^{-s-1} dx.$$

The endpoint includes a prime equal to N . Under RH, $\psi(x) - x = O(\sqrt{x} \log^2(2x))$, and the elementary sum of powers $k \geq 2$ is $O(\sqrt{x} \log^2(2x))$, so the same bound holds for $\theta(x) - x$. The integral is absolutely convergent for $\Re s > 1/2$. These facts verify the exact continuation and its analytic domain.

The only requested editorial clarification concerned § 6 of the initial author draft: its phrase “the only RH input” needed to acknowledge both the logarithmic-derivative bound used for the norm and the RH prime-counting-error bound used for this representation. The revised § 6 does so. Both inputs were already present in Round 8; this correction does not change the result.

6. Review scope

The entire six-section author draft was inspected. The review checked the infinite sums directly and did not replace them with finite numerical evidence. No numerical check is needed to establish their convergence or endpoint signs, and none was run. The result is an arithmetic simplification of the unresolved covariance problem. It neither supplies the missing positive constant nor certifies the desired residual lower bound.

7. Independently checked uniform mesoscopic corollary

Write $L = \log T$. The same proof makes the residual-energy replacement uniform for

$$1 \leq c \leq L/4.$$

Indeed, the RH logarithmic-derivative estimate used in the edge audit is $O(L^2/c + L) = O(L^2)$ uniformly on this range. The finite-polynomial bound is also uniform there. Thus $\|R\|_2 \ll \sqrt{T}L^2$ with an absolute implied constant. Moreover, $\delta^{-2} = (L/c)^2 \leq L^2$, so the unconditional bound (1) gives the same a_T , with an absolute constant. Consequently

$$\sup_{1 \leq c \leq L/4} \frac{|\|R_c - U_{c,N}\|_2^2 - \|R_c\|_2^2|}{TL^2} \ll a_T. \quad (2)$$

At $N = \lfloor T/L^6 \rfloor$,

$$a_T \ll T^{-1/6}L^3 + T^{-1/2}L^2.$$

If $b = b(T) \rightarrow \infty$ and $b = o(\log L)$, both widths $c = b/2$ and $c = b$ eventually lie in the uniform range. Applying (2) to the coupled statistic

$$b^2 \{2 \sinh(b)r_T(b) - 2 \sinh(2b)r_T(2b)\}$$

changes it by at most $O(b^2 e^{2b} a_T) = o(1)$. The last limit follows because $e^{2b} = L^{o(1)}$, whereas a_T contains a negative power of T . Prime-power removal therefore remains justified at this mesoscopic edge scale. It still supplies no new arithmetic covariance estimate and does not choose an AH convergence rate.

The final author § 7 is accepted, including its uniformity on every range $2 \leq b \leq G(T)$ with $G(T) = o(\log L)$. Its statement that decreasing polynomial coefficients preserve the elementary bound refers to the positive coefficient majorants in that bound; no monotonicity of the exact finite-time polynomial norm is required or claimed. The author's historical draft-status sentence is retained as provenance; the present independent review supplies the current ordinary-proof verdict. This is not a formal machine-checked proof or a numerical gain enclosure.

Current report 30: Mesoscopic damping: the missing arithmetic edge term and its rate requirements

Collection: R9 — complementary moduli, genuine-prime tails, and the edge.

Source: [research/dyson/round9/mesoscopic-edge/EDGE_RATE_AUDIT.md](#).

SHA-256: [47e69edb32f9e2a083ff8a56ed0a6c19c6b6e69477c63f727eb4713458e05934](#) . Git blob: [7c6fe349387a601f91222b676b1f5c6d727e7669](#) . Original bytes: 11044.

Mesoscopic damping: the missing arithmetic edge term and its rate requirements

Date: 2026-09-05. Status: bounded analytic audit; no new bound for actual zeta zeros is proved. The conclusions concern the explicitly checked mean-square estimates and the Round 8 arithmetic identity, not a universal impossibility theorem.

Outcome. Increasing the damping does not extract the required first edge correction from the checked RH estimates. Their limiting lower bound already loses a term larger than the desired signal by an unbounded factor. However, two tempting reasons for rejecting this route are false: finite-height errors can be made smaller than the signal by increasing the damping sufficiently slowly, and unquantified AH convergence permits an existential slow diagonal. The missing ingredient is an actual arithmetic estimate of the residual energy to relative precision $\mathcal{O}(1/b)$, with a coupled two-scale version if one does not assume simplicity.

1. Exact normalization and the arithmetic estimate that would suffice

Write $L = \log T$, $N = \lfloor T/L^6 \rfloor$, and $b = 2c$. Define actual functions

$$H_{T,b}(t) = -\frac{\zeta}{\zeta} \left(\frac{1}{2} + \frac{b}{2L} + it \right), \quad P_{T,b}(t) = \sum_{n \leq N} \Lambda(n) n^{-1/2-b/(2L)-it},$$

$$r_T(b) = \frac{1}{TL^2} \int_0^T |H_{T,b}(t) - P_{T,b}(t)|^2 dt.$$

Under RH the exact centered-tail representation from Round 8 is

$$H_{T,b}(t) - P_{T,b}(t) = \frac{N^{1-s}}{s-1} - E(N)N^{-s} + s \int_N^\infty E(x)x^{-s-1} dx, \quad E(x) = \psi(x) - x, \quad s = \frac{1}{2} + \frac{b}{2L} + it. \quad (1)$$

The integral is absolutely convergent. This supplies the actual arithmetic object; no divergent critical-strip Dirichlet series is substituted for it.

One sufficient, presently unproved, edge estimate is

$$\boxed{r_T(b) = \frac{e^{-b}}{b} + o\left(\frac{e^{-b}}{b^2}\right)}. \quad (\text{AE})$$

For a usable AH contradiction it would suffice to establish (AE) uniformly along every sufficiently slow $b = b(T) \rightarrow \infty$, including both b and $2b$. A more explicit formulation is a bound

$$|r_T(b) - e^{-b}/b| \leq \epsilon_T(b) e^{-b}/b^2,$$

on $B_0 \leq b \leq B(T)$, with $B(T) \rightarrow \infty$, and with the error tending to zero on a selectable increasing range. A result at one preassigned increasing rate does not automatically provide this uniformity.

Estimate (AE) asks for relative $\mathcal{O}(1/b)$ accuracy in a residual of size e^{-b}/b . A leading-order asymptotic with unspecified relative $\mathcal{O}(1)$ is insufficient. The exact one-sided coupled target in § 4 is weaker than proving both individual estimates.

2. The short-prime truncation errors are compatible with a slow diagonal

The proof of the Round 8 identity is uniform for $1 \leq c \leq L/4$. Its logarithmic-derivative bound is $\mathcal{O}(L^2/c + L)$, the right-line off-diagonal estimates only use $n^{1-2\sigma} \leq 1$, and the auxiliary polynomial bound remains $eN \log N$. Thus, uniformly for $2 \leq b \leq L/2$,

$$\frac{1}{TL^2} \int_0^T |H_{T,b}(t)|^2 dt = \frac{D_N(b)}{L^2} + r_T(b) + \mathcal{O}(L^{-4}), \quad D_N(b) = \sum_{n \leq N} \frac{\Lambda(n)^2}{n^{1+b/L}}. \quad (2)$$

The implicit constant can be absolute on this range; no hidden exponential dependence on b is generated by the contour proof.

The prime diagonal can also be estimated uniformly. Under RH,

$$\sum_{n \leq x} \Lambda(n)^2 = x \log x - x + \mathcal{O}(\sqrt{x} \log^3(2x)).$$

For example, integrate $\log x$ against $d\psi(x)$, then remove the excess contribution of prime powers. Partial summation then gives

$$\sum_{n \leq x} \frac{\Lambda(n)^2}{n} = \frac{1}{2} \log^2 x + \mathcal{O}(1).$$

Applying the decreasing weight $n^{-b/L}$ preserves the $\mathcal{O}(1)$ error, uniformly for $b \geq 0$. Consequently, with $\theta = \log N/L$,

$$\begin{aligned} \frac{D_N(b)}{L^2} &= \int_0^\theta u e^{-bu} du + \mathcal{O}(L^{-2}), \\ \left| \frac{D_N(b)}{L^2} - d(b) \right| &\ll L^{-2} + (1 - \theta) e^{-b\theta}, \quad d(b) = \frac{1 - (1 + b) e^{-b}}{b^2}. \end{aligned} \quad (3)$$

Here $1 - \theta = 6 \log L/L + \mathcal{O}(L^{-1})$. The respective ratios of the known errors to the edge scale e^{-b}/b^2 are bounded by constants times

$$\frac{b^2 e^b}{L^4}, \quad \frac{b^2 e^b}{L^2}, \quad \frac{b^2 \log L}{L} e^{b(1-\theta)}. \quad (4)$$

All tend to zero for $b = \alpha(\log L) = \alpha(\log \log T)$, and the same is true with $2b$. Thus this sufficiently slow regime resolves the known arithmetic truncation and diagonal errors. It does **not** prove (AE).

Keeping the explicit pole in (1) is harmless for the argument. If it is removed using the weaker Round 8 estimate, the normalized residual-energy error is $O(L^{-3})$, which is also $\alpha(e^{-b}/b^2)$ throughout the same slow regime.

3. What the checked primary mean-square bounds actually deliver

The primary source is Carneiro – Chandee – Chirre – Milinovich, *On Montgomery's pair correlation conjecture: a tale of three integrals*, Theorem 5 and its proof in § 4.2. Its Lemma 16 explicitly states the Goldston – Gonek – Montgomery Poisson correspondence with a uniform parameter range and errors. The source PDF is saved under `sources/`, with SHA-256 in `sources/receipt.json`.

For $q = e^{-b}$, the theorem's RH lower and upper limiting expressions, in the present normalization $l_T(b/2)/(TL^2)$, are

$$\begin{aligned} U_-(b) &= \frac{1 - (1+b)q}{b^2} + \left(\frac{1}{b} + \frac{1}{\sqrt{3}}\right) e^{-b(1+1/\sqrt{3})}, \\ U_+(b) &= \frac{1+q}{(1-q)b^2} - \frac{2q}{b(1-q)^2} + \frac{q}{1-q}. \end{aligned} \quad (5)$$

The published relative error is $O((\log \log T)^{-1/2})$, over $T^{-1}L^{5/2} \leq c \leq L^{1/4}/\sqrt{\log L}$. For slow $c \geq 1$, the proof retains the stronger absolute normalized error

$$O\left(\sqrt{\frac{\log L}{L}} + \frac{c}{L} + \frac{L^2}{Tc^2}\right). \quad (6)$$

Using only the compressed relative error would unnecessarily weaken the allowed diagonal rate.

Here is the resulting independent comparison with the signal. The sine and base-ACUE predictions for the normalized actual mean square are

$$S(b) = \frac{1-q}{b^2}, \quad A(b) = \frac{1-q}{(1+q)b^2} + \frac{q^2}{1-q^2}.$$

Their difference is exactly

$$\delta(b) = S(b) - A(b) = \frac{q(1-q)}{(1+q)b^2} - \frac{q^2}{1-q^2} \sim \frac{q}{b^2}. \quad (7)$$

Even with zero finite-height error, the lower expression in (5) falls below the sine prediction by

$$S(b) - U_-(b) = \frac{q}{b} - \left(\frac{1}{b} + \frac{1}{\sqrt{3}}\right) e^{-b(1+1/\sqrt{3})} \sim \frac{q}{b}.$$

Its deficit divided by the desired signal tends to b . In particular, it remains below the ACUE value for sufficiently large b . The upper expression exceeds the sine prediction by

$$U_+(b) - S(b) = q\left(1 - \frac{2}{b} + \frac{3}{b^2}\right) + O(q^2),$$

which is larger than the signal by order b^2 . These are deficiencies of the inspected bounds themselves, independent of how large T is chosen.

The error in (6) is below the signal when, for example, $c \leq (1/4 - \varepsilon) \log \log T$, with fixed $\varepsilon > 0$. For the simultaneous scales b and $2b$, the conservative condition $b = \alpha(\log \log T)$ suffices. This confirms that there is no honest blanket rate obstruction for arbitrarily slow damping. The limiting lower expression, rather than its finite-height error, is the decisive failed step.

4. The coupled edge statistic removes the AH nuisance exactly

A single variance cannot be compared with the base $p_0 = 1$ ACUE value without controlling the near-diagonal parameter. Under RH plus AH-Pairs, its extra contribution to the normalized actual mean square is

$$\frac{p_0(T) - 1}{\sinh b}.$$

At large b , this is approximately $2(p_0(T) - 1)e^{-b}$, larger than the edge signal unless $p_0(T) - 1 = o(b^{-2})$. That additional assertion is not part of the AH hypothesis being used.

Instead define the actual arithmetic statistic

$$C_T(b) = b^2 \left[2 \sinh(b) r_T(b) - 2 \sinh(2b) r_T(2b) - \frac{1}{2b} \right]. \quad (8)$$

Both residuals are built from the same $E(x) = \psi(x) - x$ and the same cutoff N . The nuisance contributions in (8) are $2(p_0(T) - 1)$ and its negative, so they cancel at the same finite height without a limit for $p_0(T)$.

The sine residual prediction is $r_S(b) = q/b$. Its exact substitution into (8) is

$$C_S(b) = -be^{-2b} + \frac{b}{2}e^{-4b} \rightarrow 0.$$

Let

$$H(b) = \frac{(1 - e^{-b})^2}{b^2} - e^{-b}.$$

The ACUE substitution obeys exactly

$$C_A(b) = C_S(b) - b^2 \{H(b) - H(2b)\}, \quad b^2 \{H(b) - H(2b)\} \rightarrow \frac{3}{4}. \quad (9)$$

Thus the coupled first-edge limits are **zero versus** $-3/4$. This is an algebraic sharpening of the obligation, not a proof of the zero limit for zeta.

In particular, a uniform arithmetic lower bound for (8) whose lower limit is strictly greater than $-3/4$, valid on a selectable increasing slow range, would suffice. Establishing (AE) at both scales would give the stronger limit zero. This identifies exactly the term that actual Λ/ψ estimates must improve.

5. What unquantified AH does and does not permit

The frozen Round 7 reduction gives an error $\eta_T(b) \rightarrow 0$ for each fixed $b > 0$, after retaining the displayed $\rho_0(T)$ term. It does not license inserting an arbitrary prescribed $b(T) \rightarrow \infty$.

Nevertheless, the following elementary diagonal lemma applies. Given any increasing envelope $G(T) \rightarrow \infty$, one can select increasing thresholds T_j such that for all $T \geq T_j$,

$$G(T) \geq j, \quad |\eta_T(j)| + |\eta_T(2j)| \leq e^{-2j/j^3}.$$

Set $b(T) = j$ for $T_j \leq T < T_{j+1}$. Then $b(T) \rightarrow \infty$, $b(T) \leq G(T)$, and the AH errors after multiplying (8)'s weights and $b(T)^2$ are $O(1/b(T))$. Taking, for instance, an envelope $G(T) = \alpha(\log \log T)$ also makes all explicit errors above negligible. It follows that AH forces $C_T(b(T)) \rightarrow -3/4$ along such an **existential** slow diagonal.

The thresholds depend on unknown convergence rates. No effective formula for them, no prescribed rate such as $b(T) = \sqrt{\log \log T}$, and no uniform AH estimate over a growing continuum of widths has been proved. Therefore an arithmetic theorem valid only at one prescribed rate would still need a compatibility argument. A theorem uniform over a suitable increasing slow range would avoid this issue.

6. Bounded decision and preserved checks

This lane closes with a precise unresolved arithmetic estimate, not with a new zero-spacing theorem. The checked primary RH mean-square bounds fail by factors b and b^2 at the first edge scale. The Round 8 prime truncation and the source's uniform Poisson-transfer errors are compatible with sufficiently slow damping. AH's missing rate obstructs an effective prescribed diagonal, but not an existential one. None of these facts supplies (AE) or the weaker lower bound for (8).

`edge_identity_checks.py` and its JSON preserve exact rational checks of the model subtraction and nuisance cancellation, together with two fixed floating diagnostics of the source-bound deficits. They perform no optimization, zeta experiment, or parameter scan. The primary source PDF/text and receipt are included for the stated source audit. No Round 7 or Round 8 evidence was modified.

Current report 31: Independent review of the mesoscopic edge obligation

Collection: R9 — complementary moduli, genuine-prime tails, and the edge.

Source: `research/dyson/round9/mesoscopic-edge/INDEPENDENT_EDGE_REVIEW.md`.

SHA-256: `715ca92465ab81f1e2b3c0c4198fa4f861cb9394168ad749ae2f602d5ff67dd6` . **Git blob:**
`f5e3ac2cf75ed3c9d1e6ce088931a39eb6100cc1` . **Original bytes:** 3656.

Independent review of the mesoscopic edge obligation

Date: 2026-09-05. **Reviewer:** the root research lane, independently of the source-audit author. **Accepted as a bounded analytic/source audit, without a new arithmetic edge estimate.**

I read the complete `EDGE_RATE_AUDIT.md`, its small algebra script, and the saved primary source at Theorem 5, Lemma 16 and Section 4.2. The following points were checked directly rather than inferred from the script's PASS status.

1. The source uses c where the report uses $b/2$. Substitution gives exactly the displayed lower and upper functions. The absolute proof error for $c \geq 1$ contains the square root of $\log(\log T)/\log T$, the $c/\log T$ term, and $\log^2(T)/(T c^2)$. Keeping this error instead of the theorem's compressed relative error is legitimate. It can be below the exponential edge signal on a sufficiently slow diagonal.
2. The limiting lower bound's deficit is asymptotic to $e^{(-b)}/b$, while sine minus base ACUE is asymptotic to $e^{(-b)}/b^2$ in normalized I . The failure therefore persists even after the finite-height error is removed. The upper deficit is larger by order b^2 . These are statements about the checked bounds, not every consequence of RH.
3. The RH estimate for $\sum \Lambda(n)^2$ has the stated main terms $x \log x - x$. Partial summation cancels the $\log x$ terms and leaves $(1/2)\log^2 x + O(1)$ for the sum divided by n . A decreasing exponential weight has bounded total variation, so the bounded primitive error stays $O(1)$ uniformly in b . The missing interval between $\log N/\log T$ and one gives the displayed endpoint error. All normalized errors vanish relative to $e^{(-b)}/b^2$ for $b = o(\log \log T)$, including the second scale.
4. The factors two in the residual statistic are correct. Substituting $r_S(b) = e^{(-b)}/b$ gives $-b e^{(-2b)} + (b/2)e^{(-4b)}$. The difference from ACUE is $b^2[H(b) - H(2b)]$, tending to $3/4$. The additive AH parameter contributes exactly equal constants with opposite signs, so a limit for $p_0(T)$ is unnecessary.
5. The fixed-width AH errors permit an existential stepwise increasing diagonal. Choosing their errors at widths j and $2j$ below $e^{(-2j)}/j^3$ makes the amplified error $O(1/j)$. This does not justify a predetermined formula for $b(T)$.

An explicit sufficient uniform quantifier

To avoid ambiguity in the phrase “a selectable increasing range”, the following is one precise sufficient arithmetic target. Find an increasing $G(T)$ tending to infinity with $G(T) = o(\log \log T)$, for which

$$\lim_{B \rightarrow \infty} \liminf_{T \rightarrow \infty} \inf_{B \leq b \leq G(T)} C_T(b) > -\frac{3}{4}. \quad (\text{R})$$

The intervals are nonempty for every fixed B once T is sufficiently large. If (R) holds, choose the AH diagonal under the envelope $G(T)$. Along that diagonal the statistic tends to $-3/4$ by the reviewed argument, contradicting (R). The outer limit in (R) exists as an extended monotone limit of the inner lower bounds. Equivalently, one may supply a strict uniform gap for all sufficiently large lower cutoffs B .

This condition does not assert an exact GUE law at each fixed b . It is stronger than a statement at a single predetermined growing $b(T)$, because such a statement might be incompatible with the unknown AH convergence rate. It is weaker than proving both individual residual asymptotics (AE). None of these arithmetic conditions is proved in the report or this review.

The source audit and model algebra are useful because they identify the required first correction and rule out two incorrect rate objections. They supply no positive improvement in the current lower bound. No additional numerical experiment, parameter scan or formal verification was performed for this review.

Current report 32: Round 10: a power saving for one actual shifted–prime discrepancy

Collection: R10 — complete the actual shift packet.

Source: `research/reports/dyson_round10.md`.

SHA-256: `c276e4c0da7d40fc275a1c8d10fd8cf94641753a7add8fdbc4a90f05bb712737` . Git blob: `455f5f62d4d54e287de47328267337f6b0351c90` . Original bytes: 9485.

Round 10: a power saving for one actual shifted–prime discrepancy

Date: 2026-09-05. This round proves an unconditional bound for a specified smooth packet of the actual arithmetic discrepancy. It does not prove a new zeta pair-correlation lower bound. The accompanying source audit explains why two tempting short-interval inputs do not supply the missing precision.

The bound and how far it remains from the target

Use the complementary squarefree modulus family Q_X from Round 9, with $q \leq Q = X^{.523}$. Let $X = T^\alpha$, $6/5 \leq \alpha \leq 7/5$, and $H = X/T$. The discrepancy D_{Q^V} uses the actual localized prime-covariance sinc kernel, its Mobius – log divisor coefficient, its coprime principal term, and a fixed smooth cutoff $V(h/H)$, supported in $1 < h/H < 2$.

The complete ordinary proof gives

$$|D_Q^V(X, T)| \ll_{V,X} \sqrt{HX(X+Q^2)} (\log X)^4. \tag{1}$$

The estimate itself is unconditional. RH is used only when inserting it into the earlier actual-zeta correspondence. The proof works for any squarefree modulus subset with this cap; it does not yet exploit triple dense divisibility beyond the selection of the prescribed family. No novelty claim is made for its elementary completion and spacing ingredients.

For $H = X^\theta$, the new exponent is $1.023 + \theta/2$. The previous per-shift estimate, summed by the triangle inequality, had exponent $1 + \theta$ with any fixed logarithmic saving.

H exponent	Previous triangle exponent	New smooth–packet exponent	Power saved, before logarithms
1/6	7/6	3319/3000	181/3000 = 0.060333...
2/7	9/7	8161/7000	839/7000 = 0.119857...

This is a power improvement in a genuine arithmetic error bound. Nevertheless, after division by the required $X \log X$ fluctuation scale, the estimate still grows as $X^{.023} \sqrt{H} \log^3 X$. It does not evaluate this selected component at the precision needed for the compact Fourier test. The whole unsmoothed shift range, complementary divisor remainder, support main terms and continuous centering also remain unresolved.

Why completing the shift sum helps

First remove prime powers from both terms of the progression discrepancy. Their total contribution is $O_{\eta}(H X^{(.5+\eta)} \log^3 X + H \sqrt{X} \log^4 X)$, which is $o(X \log X)$ for $\eta=.01$ and the stated H range. Every remaining prime exceeds every modulus, so it is a unit modulo each q . This permits exact finite Fourier completion while retaining the principal unit sum.

For a separated amplitude $f(p/X)v(h/H)$, define

$$S_v(\beta) = \sum_h v(h/H) e(-\beta h), \quad A_f(\beta) = \sum_p (\log p) f(p/X) e(\beta p).$$

Combining repeated rational frequencies before estimating gives the exact pairing

$$\sum_{d,a}^* S_v(a/d) \left(\sum_{\substack{q \in Q_X \\ d \mid q}} \frac{\mu(q)}{q} \right) (A_f(a/d) - \frac{\mu(d)}{\phi(d)} A_f(0)). \quad (2)$$

The star means $2 \leq d \leq Q$, $1 \leq a < d$ and $(a,d)=1$. The zero frequency cancels. The principal coefficient $\mu(d)/\phi(d)$ follows from the exact Ramanujan ratio and is independent of the parent modulus q . Treating all original fractions r/q as distinct would invalidate the subsequent spacing estimate.

The squared norm of the first two factors in (2) is $O(H \log^3 Q)$, or $O(H \log^5 Q)$ when a $\log q$ coefficient is present. A direct Schur/spacing argument bounds the centered prime sums' squared norm by $O(X(X+Q^2) \log^2 X)$. The proof retains the logarithm in the elementary spacing bound. Cauchy – Schwarz then yields (1), after including the logarithmic cofactor.

The original two-variable sinc kernel is not simply replaced by a separated model. With $y=m/X$, $z=h/H$ and $\epsilon=H/X=1/T$, its phase is

$$\epsilon^{-1} \log \frac{y}{y - \epsilon z} = \int_0^z \frac{du}{y - \epsilon u}.$$

It and its derivatives are uniformly smooth on the fixed support. A Fourier expansion with uniformly summable derivative-weighted coefficients transfers the separated estimate to the actual kernel, including $\log(y - \epsilon z)$. The fixed support also keeps both cutoff factors on their correct branch.

Separate reviews cover the [coefficients and spacing](#) and the [actual kernel/prime-power passage](#). The root additionally read the complete author argument. The accepted scope is the explicitly defined smooth packet. A possible bounded-variation extension in a review is not silently promoted to a theorem for the full sharp packet.

What the checked prime–variance sources do not give

The arithmetic source audit, with independent review, examines the edge shell $X=T^{(1+s/b)}$, $1 \leq s \leq 2$, $h=X/T$. Its interval-length exponent is $s/(b+s)$, tending to zero.

Guth – Maynard, arXiv:2405.20552v2, Corollary 1.4, gives almost-all prime-count asymptotics for $h \geq X^{(2/15+\epsilon)}$, with fixed epsilon. Even the corollary's epsilon-zero endpoint misses the whole shell once $b > 13$. The subsequent Remark discusses a slight fixed-epsilon improvement with a worse error; it still does not cover an exponent tending to zero.

Inside the stated range, a direct conversion of the count theorem and its exceptional-set bound gives an error in the squared count far above the fluctuation scale $X h \log(X/h)$. Almost-all PNT is therefore not itself a constant-precision variance theorem. The audit does not rule out stronger uses of the underlying methods.

The checked Carneiro – Chandee – Chirre – Milinovich short-interval comparisons retain fixed factors approximately 0.9028 and 1.0736, and fixed-endpoint quantifiers. These do not supply the shrinking first correction required by the mesoscopic test. The large-beta statements in that source use beta as the endpoint of the prime range, not the damping b.

One precise remaining arithmetic mixed moment

Let R_b be the genuine-prime residual at displacement $b/(2 \log T)$, with the same cutoff $N = \text{floor}(T/\log^6 T)$, and put

$$E_T(b) = \frac{e^b \| R_b \|_2^2}{T \log^2 T}, \quad K_b = -R_b - 2\partial_b R_b, \quad M_T(b) = \frac{e^b \text{Re}\langle R_b, K_b \rangle}{T \log^2 T}.$$

Exactly $M_T = -E_T'$. In the absolutely convergent region, K gives the genuine-prime weight $\log(p)/\log(T)-1$; the working-strip object is its centered analytic continuation. Differentiating an unspecified asymptotic error is not part of the argument.

For some fixed epsilon > 0, an actual uniform bound

$$M_T(s) \geq s^{-2} - (2 - \epsilon)s^{-3} + o(s^{-3}) \tag{3}$$

through twice a suitable slow envelope would imply that the Round 9 coupled statistic has lower limit at least $-3/4 + 3\epsilon/8$, with the required uniform quantifiers. This would contradict AH-Pairs under RH. The implication follows by integrating from b to 2b; the correction integral is exactly $3/(8b^2)$. Reviewed RH upper estimates control the small exponential terms.

Inequality (3) is not proved. This formulation identifies one signed, logarithmically weighted genuine-prime correlation; it is not a substitute for estimating it. Neither source checked here supplies a positive epsilon.

Verification and next mathematical decision

The finite script checks 9,615 exact Ramanujan-ratio cases, reduced-frequency grouping with formal log q coefficients, zero-frequency cancellation and both endpoint exponents. A second exact script checks the prime-range thresholds and mixed-moment normalization. These checks validate algebra; the power bound rests on the written analytic proof and independent reviews.

All 15 original files (1,452,061 bytes) are preserved in the adjacent local `Astra-Local-Archive/round10-originals/`; 14 research files are verbatim public. The third-party Guth – Maynard HTML body stays local, identified by URL/hash. Earlier primary references remain in their pinned local folders. The [intake manifest](#) and [integration replay](#) record their exact scope. Both output JSON files reproduce in a separate process. The mixed-moment JSON is identical in full; the completion output differs only in two temporary provenance paths, with source hashes still checked. No new floating optimization, prime-gap sweep, or zeta-data fit was run.

From the repository root, with Python and SymPy:

```
python3 research/logs/round10-integration/recheck.py --prime-gap-source-dir /path/to/retained/round9-external-sources
python3 tools/verify_manifest.py
```

The next useful arithmetic work is an estimate for the pairing (2) that uses the selected Mobius coefficients together with the centered prime exponential sums, beyond separate norm bounds. A separate possibility is an averaged version of (3) with a strict quantitative gain. Repeating the same completion with different notation or improving only its logarithmic exponent would not close the existing power gap.

The large handoff PDFs retain their stated earlier checkpoints; this compact report is an additional source record. The goal remains active. Reverting this research slice removes its reports and rechecks without altering earlier proofs. New model sessions, generic positivity scans and claims of a solved famous conjecture are outside this checkpoint.

Current report 33: A power improvement for a smooth packet of the actual shifted discrepancy

Collection: R10 — complete the actual shift packet.

Source: [research/dyson/round10/shift-average/SMOOTH_SHIFT_COMPLETION_BOUND.md](#).

SHA-256: [7b52e4d82dc40bf90183331d548b7fffe5545d1928d7cb93223223b5b71c1d78](#) . Git blob: [ec368290b4a8627d406b769b68cbc7b3ca6bfd5d](#) . Original bytes: 13008.

A power improvement for a smooth packet of the actual shifted discrepancy

Date: 2026-09-05. Status: an elementary completion bound for an identified component of the Round 9 discrepancy. It improves the triangle-inequality bound by a power of X , but remains too large to imply a zeta covariance theorem.

Let $X = T^\alpha$, $6/5 \leq \alpha \leq 7/5$, $H = X/T$, and $Q = X^{523/1000}$. For the actual complementary squarefree modulus family $\mathcal{Q}_X \subset (X^{1/2}, Q]$ from Round 9, the smooth packet defined below satisfies

$$|D_Q^V(X, T)| \ll_{V, X} \sqrt{HX(X+Q^2)} (\log X)^4. \quad (1)$$

Estimate (1) itself does not use RH. RH enters only when this discrepancy estimate is combined with the Round 9 aggregate covariance identity.

Thus, writing $H = X^\theta$,

$$|D_Q^V(X, T)| \ll X^{1023/1000+\theta/2} (\log X)^4, \quad \frac{1}{6} \leq \theta \leq \frac{2}{7}. \quad (2)$$

Compared with the previously available $HX \log^{-A} X$, the exponent improves by $\theta/2 - 23/1000$, at least $181/3000 = 0.060333 \dots$, before logarithmic factors. This is a nontrivial bound for the **actual discrepancy**, not an estimate for a substituted positive model.

However, (2) divided by $X \log X$ still grows as a power. No bound of the required size $\mathcal{O}(X \log X)$ follows. The gain uses squarefreeness and the modulus cap; it does not yet exploit triple dense divisibility beyond establishing that the prescribed family is admissible. No novelty claim is made for the completion or finite-spacing argument.

1. The exact packet and source relationship

Keep the Round 9 discrepancy convention

$$\Delta(f; a \bmod q) = \sum_{m \equiv a \pmod{q}} f(m) - \frac{1}{\phi(q)} \sum_{(m, q)=1} f(m).$$

Fix $V \in C_c^\infty(1, 2)$ and $\chi \in C_c^\infty(1, 3/2)$, and put

$$w_h(u) = \chi(u/X) a_u(X) a_{u+h}(X) \operatorname{sinc}_0(T \log(1 + h/u)),$$

$$a_u(X) = \min\{(u/X)^{1/2}, (X/u)^{3/2}\}, \quad \operatorname{sinc}_0(v) = \sin(v)/v, \quad \operatorname{sinc}_0(0) = 1.$$

The object in (1) is

$$D_Q^V(X, T) = \sum_{h \in \mathbb{Z}} V(h/H) \sum_{\substack{q \in \mathbb{Q}_X \\ (q, h)=1}} \mu(q) \Delta(\Lambda(m) w_h(m-h) \log((m-h)/q); h \bmod q). \quad (3)$$

Each sequence inside Δ is defined to be zero off its indicated support; there is no logarithm of a nonpositive argument. This is the actual Round 9 weighted discrepancy, with one additional smooth shift cutoff $H < h < 2H$. We assume the fixed hard-cutoff constant in that round is at least two. No claim about the entire unsmoothed shift range is hidden in this localization.

The modulus family and its validity are inherited from [OpenAI, Improved short gaps between primes, 30 August 2026](#): Proposition 2.3, pp.4 – 5; equation (2.5), p.7; and Corollary 2.19, p.11. Round 9 checks $\omega = 3/250$, $\delta = \varepsilon = 1/1000$, giving $240\omega + 80\delta = 2.96 < 3$ and the cap X^{523} . That source's modulus sum gives only the $HX \log^{-A} X$ bound after a shift triangle inequality. All additional analytic estimates needed for (1) are proved below.

In fact the completion argument applies to any squarefree subset of $(X^{1/2}, Q]$. It does not assert a new distribution estimate for triply densely divisible moduli.

2. Replace von Mangoldt by genuine primes inside this discrepancy

Let $\mathfrak{g}_*(m) = \log m$ when m is prime, and zero otherwise. Replacing Λ by $\mathfrak{g}_*$ in (3) changes it by

$$O_\eta(HX^{1/2+\eta} \log^3 X + H\sqrt{X} \log^4 X) \quad (4)$$

for every fixed $\eta > 0$.

For the progression term, a prime power $m = p^j \asymp X$, $j \geq 2$, contributes only when $q \mid m-h$. For each h there are at most $\tau(m-h) \ll_\eta X^\eta$ possible q . The weights cost $O(\log X)$, and $\sum_{p^j \asymp X, j \geq 2} \log p \ll \sqrt{X} \log^2 X$. Summing over h proves the first bound in (4).

For the principal term, bound the h -sum by $O(H)$, retain the coprimality restrictions or enlarge by absolute values, and use $\sum_{q \leq Q} 1/\phi(q) \ll \log X$. This gives the second bound. Thus both portions of the discrepancy are controlled; one cannot replace Λ in only its progression sum.

Taking, for example, $\eta = 1/100$, (4) is $\alpha(X \log X)$ throughout $H \leq X^{2/7}$, and is smaller than the right side of (1). This is nuisance removal, not the source of the power improvement.

Every remaining prime m satisfies $m \asymp X > Q$. It is therefore a unit modulo every q under consideration, which makes the following completion exact.

3. Exact completion for a separated weight

First replace the smooth two-variable amplitude by $f(m/X)v(h/H)$, where f, v are smooth with fixed compact supports and m is restricted to primes near X . The logarithmic factor depending on q is handled in Section 6.

Write $\vartheta(z) = e^{2\pi iz}$ and

$$S_v(\beta) = \sum_{h \in \mathbb{Z}} v(h/H) \vartheta(-\beta h), \quad A_f(\beta) = \sum_p \log p f(p/X) \vartheta(\beta p). \quad (5)$$

For any integer m ,

$$\sum_{h \equiv m \pmod{q}} v(h/H) = \frac{1}{q} \sum_{r=0}^{q-1} S_v(r/q) \vartheta(rm/q). \quad (6)$$

Also, with the Ramanujan sum $c_q(r) = \sum_{b \bmod q, (b,q)=1} \vartheta(rb/q)$,

$$\sum_{(h,q)=1} v(h/H) = \frac{1}{q} \sum_{r=0}^{q-1} S_v(r/q) c_q(r). \quad (7)$$

In the progression term, $(h, q) = 1$ is automatic because m is a prime exceeding q . In the principal term it remains exactly the unit sum in (7). Consequently the completed separated discrepancy is

$$\sum_{q \in \mathbb{Q}_X} \frac{\mu(q)}{q} \sum_{r=0}^{q-1} S_v(r/q) \left(A_f(r/q) - \frac{c_q(r)}{\phi(q)} A_f(0) \right). \quad (8)$$

Reduce $r/q = a/d$. Since q is squarefree and $d = q/(q, r)$,

$$\frac{c_q(r)}{\phi(q)} = \frac{\mu(d)}{\phi(d)}. \quad (9)$$

This ratio is independent of the larger modulus q . The $d=1$ term vanishes identically. Regrouping (8) by distinct reduced fractions gives

$$\sum_{\substack{2 \leq d \leq Q \\ 1 \leq a < d, (a,d)=1}} C_{a/d} Z_f(a/d), \quad (10)$$

$$C_{a/d} = S_v(a/d) \sum_{\substack{q \in \mathbb{Q}_X \\ d \mid q}} \frac{\mu(q)}{q}, \quad Z_f(a/d) = A_f(a/d) - \frac{\mu(d)}{\phi(d)} A_f(0). \quad (11)$$

Distinct representations of the same rational frequency must be merged before applying a finite-spacing estimate. Treating all r/q as distinct would be incorrect.

4. A quantitative norm for the completed coefficients

The elementary bound

$$\left| \sum_{\substack{q \in \mathbb{Q}_X \\ d \mid q}} \frac{\mu(q)}{q} \right| \leq \frac{1 + \log(Q/d)}{d} \quad (12)$$

does not assume cancellation of the Möbius coefficients.

Finite summation by parts, applied J times to the compactly supported sequence $\nu(h/H)$, gives

$$|S_\nu(\beta)| \ll_{\nu,J} H(1 + H \|\beta\|)^{-J}. \quad (13)$$

One can prove (13) directly: the J -th finite difference has total absolute mass $O_{\nu,J}(H^{1-J})$, and its Fourier transform is multiplied by $(\vartheta(\beta) - 1)^J$. Combine that estimate with the trivial $O(H)$ bound. No convergence of a critical-strip Dirichlet series is involved.

For $J=2$,

$$\sum_{a=1}^{d-1} |S_\nu(a/d)|^2 \ll_\nu \begin{cases} d^4/H^2, & d \leq H, \\ Hd, & d \geq H. \end{cases} \quad (14)$$

This follows by measuring distance to the endpoints $a=0,d$; the relevant convergent series is $\sum_{a \geq 1} a^{-4}$. Dropping the coprimality restriction only enlarges this upper bound. Equations (12) – (14) prove

$$\sum_{d,a} |C_{a/d}|^2 \ll_\nu H(\log(2Q))^3. \quad (15)$$

For the coefficient with $\mu(q) \log q/q$ in place of $\mu(q)/q$, the same proof gives $O_\nu(H \log^5(2Q))$. These explicit norms replace the H -fold triangle inequality.

5. The required finite-spacing estimate, proved directly

The reduced fractions a/d , $d \leq Q$, are distinct and separated on the unit circle by at least Q^{-2} . For a consecutive integer interval of length $N \ll X$, the Gram matrix of their exponential vectors has diagonal N and off-diagonal absolute values at most

$$\min\left(N, \frac{1}{2 \|\beta - \beta'\|}\right).$$

Order the distances in either direction around the circle. Separation bounds the absolute row sum by

$$N + O(Q^2 \log(2Q)).$$

The Schur bound and matrix duality therefore give the finite inequality

$$\sum_\beta \left| \sum_n b_n \vartheta(\beta n) \right|^2 \ll (X + Q^2 \log(2Q)) \sum_n |b_n|^2. \quad (16)$$

For $b_p = (\log p)f(p/X)$ at primes and zero otherwise, Chebyshev's bound gives $\sum |b_n|^2 \ll_f X \log X$ and $A_f(0) \ll_f X$. The constant-centering portion of (11) has total squared mass at most

$$|A_f(0)|^2 \sum_{d \leq Q} \frac{\mu(d)^2}{\phi(d)} \ll_f X^2 \log(2Q).$$

Combining this with (16),

$$\sum_{d,a} |Z_f(a/d)|^2 \ll_f X(X+Q^2)(\log X)^2. \quad (17)$$

The logarithm in the spacing row sum has been retained. Using the prime-supported coefficient norm, rather than the cruder norm of all integers weighted by $\log$, is also needed for the displayed logarithmic bookkeeping.

Cauchy – Schwarz in (10), with (15) and (17), now bounds the separated discrepancy by

$$\ll_{f,v} \sqrt{HX(X+Q^2)}(\log X)^{5/2}. \quad (18)$$

The version with $\log q$ costs one more logarithm.

6. The actual m/h kernel, including the logarithm

Put $\epsilon = H/X = 1/T$, $y = m/X$, and $z = h/H$. The amplitude of (3), apart from the logarithm and the prime weight, is

$$F_\epsilon(y, z) = V(z)X(y - \epsilon z) y^{-3/2} (y - \epsilon z)^{-3/2} \operatorname{sinc}_0\left(\frac{1}{\epsilon} \log \frac{y}{y - \epsilon z}\right). \quad (19)$$

The cutoff ensures that both a -factors are on their $u > X$ branch. The apparently singular phase is actually

$$\frac{1}{\epsilon} \log \frac{y}{y - \epsilon z} = \int_0^z \frac{du}{y - \epsilon u}. \quad (20)$$

It extends smoothly to z/y at $\epsilon = 0$. Every fixed mixed derivative of (19) is uniformly bounded on a fixed compact rectangle for sufficiently small epsilon. The same holds after multiplying by $\log(y - \epsilon z)$.

Choose fixed smooth cutoffs slightly larger than that rectangle and expand the resulting functions in a two-variable Fourier series. Repeated integration by parts gives coefficients $b_{k,\ell}$ satisfying, uniformly in epsilon,

$$\sum_{k,\ell} |b_{k,\ell}| (1 + |\ell|)^J < \infty$$

with a common bound for any fixed J . Each term separates into a function of m/X and a function of h/H . The m -factor has uniformly bounded sup norm; the derivatives of the h -factor grow at most polynomially in ℓ , absorbed by the displayed coefficient sum. Thus (18) is summable over this expansion. This is a proved separation of the actual amplitude, not its replacement by a product approximation with an unspecified error.

Finally,

$$\log((m - h)/q) = \log X - \log q + \log(y - \epsilon z). \quad (21)$$

The first term costs $\log X$, the second uses the $\log q$ coefficient norm after (15), and the third is another uniformly smooth amplitude. Their common upper bound is

$$O_{V,X}(\sqrt{HX(X+Q^2)}(\log X)^4).$$

Together with the smaller prime-power error (4), this proves (1).

7. Nontriviality, limitations, and the next mathematical obstruction

At $\theta = 1/6$, the exponent in (2) is $3319/3000 = 1.106333 \dots$, compared with the triangle exponent $7/6$. At $\theta = 2/7$, it is $8161/7000 = 1.165857 \dots$, compared with $9/7$. These are power improvements, not merely changes to logarithmic factors.

They still exceed one. Under RH, inserting (1) into the Round 9 aggregate identity leaves an error larger than $X \log X$; the negligible $H\sqrt{X} \log^4 X$ nuisance there does not change that conclusion.

The now-specific place needing further arithmetic input is the pairing

$$\sum_{d,a} S_V(a/d) \left(\sum_{\substack{q \in \mathbb{Q}_X \\ d \mid q}} \frac{\mu(q)}{q} \right) (A_f(a/d) - \frac{\mu(d)}{\phi(d)} A_f(0)), \quad (22)$$

and its $\log q$ and smooth-amplitude variants. A proof reaching the covariance scale must improve this actual pairing beyond the elementary coefficient-norm and spacing bounds, for example by exploiting the complementary modulus restrictions inside the coefficient sum together with the prime exponential sums. The 186 coherent-class estimate by itself supplies no such cross-frequency cancellation.

This is a location of missing arithmetic information, not a proof that no sharper completion argument can work. No further reparameterization or generic positivity test was performed.

8. Exact checks and provenance

The companion `check_shift_completion.py` verifies with exact rational/integer arithmetic the Ramanujan ratio (9), regrouping by reduced denominators, the cancellation of the zero frequency, and all stated power exponents. Its finite toy modulus families test algebra, not prime-distribution asymptotics. The script records the hashes of this report and the frozen Round 9 report/source files.

The estimates (12) – (18) and the actual-kernel passage are ordinary written proofs. No floating-point experiment is being used to assert a power saving. The report does not claim to settle the unsmoothed full discrepancy, the complementary divisor remainder, AH, or Montgomery's conjecture.

Current report 34: Independent audit of the completed–shift coefficients and spacing bound

Collection: R10 — complete the actual shift packet.

Source: [research/dyson/round10/shift-average/COEFFICIENT_AND_SPACING_AUDIT.md](#).

SHA-256: `eccd80ad4909639a53583376444cbc9fce9fab518b75de7f82052c31c7a27fc7` . Git blob: `464c498e14af12781d28377818662f6bd3a6eccd` . Original bytes: 11413.

Independent audit of the completed–shift coefficients and spacing bound

Date: 2026-09-05. Reviewer: [yau_flow](#) . This note checks the exact reduced-fraction identity, its coprime principal term, and a fully quantified bound for separable genuine-prime packets. The authoring agent and another reviewer are separately checking the actual two-variable weight and prime-power exceptions. No numerical scan or large computation is used here.

Final verdict on the assigned analytical core: the coefficient and spacing argument is valid. Its logarithmic losses fit the proposed fourth power. The resulting estimate improves the accumulated error for the selected divisor component but does not reach the covariance scale. The distinction between this separable lemma and a verified application to the full weight must be retained. The pinned author draft and precise review scope are recorded in Section 7.

1. Precisely quantified separable statement

Let $X \geq 4$, $H \geq 1$, and $2 \leq Q < X$. Let $\mathcal{Q}$ be any set of distinct squarefree integers in $[1, Q]$. No dense-divisibility assumption is needed for this lemma. Let f be a bounded complex function on $[1, 2]$, and let v be a complex function of bounded variation on $\mathbb{R}$, supported in a fixed interval $[-C, C]$. Put

$$M(v) = \|v\|_{\infty} + \text{TV}(v), \quad B_p = (\log p) f(p/X) \mathbf{1}_{X < p \leq 2X}.$$

For $j \in \{0, 1\}$, define

$$D_j = \sum_{h \in \mathbb{Z}} v(h/H) \sum_{\substack{q \in \mathcal{Q} \\ (q, h)=1}} \mu(q) (\log q)^j \times \left\{ \sum_{\substack{X < p \leq 2X \\ p \equiv h \pmod{q}}} B_p - \frac{1}{\phi(q)} \sum_{X < p \leq 2X} B_p \right\}. \quad (1)$$

In (1), $(\log q)^0 = 1$, including $q = 1$. The braces for $q = 1$ vanish identically. Since every prime in the sum exceeds Q , each is coprime to every modulus. This is the reason the principal sum in (1) is unrestricted; for a von Mangoldt sum this replacement would require a separate prime-power argument.

There is a constant depending only on C such that

$$|D_j| \ll_C \|f\|_\infty M(\nu) \sqrt{HX(X+Q^2)} \log^{j+5/2}(2X). \quad (2)$$

The exponent $j + 5/2$ need not be rounded up. In particular the common $O(\log^4(2X))$ upper bound is valid for $j = 1$, and also after multiplying D_0 by a factor of size $O(\log X)$. These are the two cases needed to separate $\log((m-h)/q)$.

All constants in (2) are independent of the chosen family $\mathbf{Q}$, its cardinality, X, H, Q , and the values of f, ν , apart from the displayed norms. The statement is about the actual primes and progression discrepancy, not a formal spectral model.

2. Exact reduction, including the principal term

Use $e(z) = e^{2\pi iz}$ and

$$S_\nu(a) = \sum_{h \in \mathbb{Z}} \nu(h/H) e(-ah).$$

For a prime $p > Q$, its congruence $p \equiv h \pmod{q}$ forces $(h, q) = 1$. Thus the coprimality restriction can be dropped in the first term of (1), but it must be retained in the principal term. Finite Fourier inversion gives

$$\sum_h \nu(h/H) \mathbf{1}_{p \equiv h \pmod{q}} = \frac{1}{q} \sum_{r=0}^{q-1} S_\nu(r/q) e(rp/q),$$

and

$$\sum_{(h,q)=1} \nu(h/H) = \frac{1}{q} \sum_{r=0}^{q-1} S_\nu(r/q) c_q(r),$$

where $c_q(r) = \sum_{a \bmod q, (a,q)=1} e(ra/q)$ is the Ramanujan sum. At a reduced fraction $r/q = a/d$, with $(a, d) = 1$,

$$\frac{c_q(r)}{\phi(q)} = \frac{\mu(d)}{\phi(d)}. \quad (3)$$

For squarefree q , (3) follows directly by multiplying the local factors: a prime dividing d contributes $-1/(p-1)$, and a prime dividing q/d contributes one.

Define

$$A_{d,j} = \sum_{\substack{q \in \mathbf{Q} \\ d \mid q}} \frac{\mu(q)(\log q)^j}{q}, \quad E(a/d) = \sum_{X < p \leq 2X} B_p(e(ap/d) - \frac{\mu(d)}{\phi(d)}).$$

Then the exact identity is

$$D_j = \sum_{2 \leq d \leq Q} \sum_{\substack{1 \leq a < d \\ (a,d)=1}} A_{d,j} S_{\chi}(a/d) E(a/d). \quad (4)$$

There is one term for each reduced fraction, not one independent frequency for each original modulus. The $d=1$ term is zero because its centered character is $1-1$. The coefficient is a sum with denominator q ; replacing it by $1/d$ before summing the multiples would lose the logarithmic norm control. No cancellation of the Möbius coefficients is used below.

3. Squared coefficient norm

Summing by parts in the finite shift sequence gives

$$|S_{\chi}(a)| \ll_C M(\nu) \min\{H, \|a\|^{-1}\}, \quad (5)$$

with the first alternative used at integral a . The discrete variation is at most the total variation of the zero-extended function ν . This proof includes indicator cutoffs with a specified endpoint convention; it does not require smoothness.

For $d \leq H$, bound the complete nonzero residue sum by

$$\sum_{a=1}^{d-1} |S_{\chi}(a/d)|^2 \ll_C M(\nu)^2 d^2.$$

For $d > H$, split the distance of a from the nearest multiple of d at d/H . Equation (5) gives

$$\sum_{a=1}^{d-1} |S_{\chi}(a/d)|^2 \ll_C M(\nu)^2 H d. \quad (6)$$

The reduced-residue sum is smaller. Independently,

$$|A_{d,j}| \leq \frac{\log^j(2Q)}{d} (1 + \log \frac{Q}{d}) \ll \frac{\log^{j+1}(2Q)}{d}. \quad (7)$$

Hence

$$\begin{aligned} \sum_{d,a}^* |A_{d,j} S_{\chi}(a/d)|^2 &\ll_C M(\nu)^2 \log^{2j+2}(2Q) \left(\sum_{d \leq \min(H,Q)} 1 + H \sum_{H < d \leq Q} \frac{1}{d} \right) \\ &\ll_C M(\nu)^2 H \log^{2j+3}(2Q). \end{aligned} \quad (8)$$

Thus the claimed $H \log^3 Q$ and $H \log^5 Q$ squared norms are valid for $j = 0, 1$, respectively. In particular, the coefficient argument itself does not distinguish a smooth shift cutoff from a sharp finite interval. Whether the full nonseparable weight admits a controlled decomposition is a different question.

4. Well-spaced frequencies and centered prime sums

Distinct reduced fractions of denominator at most Q have circular separation at least Q^{-2} . For any integer interval of length $O(X)$, its exponential kernel satisfies

$$\left| \sum_m \mathfrak{e}((\alpha - \beta)m) \right| \ll \min\{X, \|\alpha - \beta\|^{-1}\}.$$

For a fixed frequency, the number of others within circular distance r is $O(1 + rQ^2)$. Ordering the distances and summing the harmonic series therefore bounds every absolute Gram row sum by

$$O(X + Q^2 \log(2Q)).$$

The elementary Schur bound and finite-dimensional duality imply

$$\sum_{d,a}^* \left| \sum_m b_m \mathfrak{e}(am/d) \right|^2 \ll (X + Q^2 \log(2Q)) \sum_m |b_m|^2. \quad (9)$$

This proof retains the logarithm in the elementary spacing estimate. It does not silently invoke the sharper large-sieve constant $X + Q^2$.

For $b_p = B_p$ and zero coefficients on other integers, Chebyshev's elementary bound $\sum_{p \leq 2X} \log p \ll X$ gives

$$\sum_p |B_p|^2 \ll \|f\|_\infty^2 X \log(2X), \quad \left| \sum_p B_p \right| \ll \|f\|_\infty X. \quad (10)$$

Using primes in (10) saves one logarithm compared with the cruder bound obtained by assigning a $\log X$ coefficient to every integer.

The constant terms in $E(a/d)$ have squared norm

$$\left| \sum_p B_p \right|^2 \sum_{d \leq Q} \frac{\mu(d)^2}{\phi(d)} \ll \|f\|_\infty^2 X^2 \log(2Q). \quad (11)$$

For completeness, $\sum_{d \leq Q} 1/\phi(d) \ll \log(2Q)$ follows from $n/\phi(n) = \sum_{a|n} \mu(a)^2/\phi(a)$: after interchanging the two positive sums the remaining series is bounded by $\sum_a \mu(a)^2/(a\phi(a)) = \prod_p (1 + 1/(p(p-1))) < \infty$.

Combining (9) – (11), with $\|u - v\|^2 \leq 2\|u\|^2 + 2\|v\|^2$, proves

$$\sum_{d,a}^* |E(a/d)|^2 \ll \|f\|_\infty^2 X(X + Q^2) \log^2(2X). \quad (12)$$

The primitive principal term is therefore affordable without assuming cancellation of the prime exponential sums. Cauchy – Schwarz in (4), using (8) and (12), proves (2).

5. What an application to the actual weight must verify

The separable proof applies to a sum of blocks if the full weight has an absolutely summable representation

$$W(m/X, h/H) = \sum_\ell \gamma_\ell f_\ell(m/X) v_\ell(h/H), \quad \sum_\ell |\gamma_\ell| \|f_\ell\|_\infty M(v_\ell) \leq K, \quad (13)$$

uniformly in X, H, T . The bound is then multiplied by K . If the logarithmic cofactor is present, write

$$\log((m - h)/q) = \log X - \log q + \log(m/X - (H/X)(h/H)).$$

The first two terms use $j = 0$ and $j = 1$; the third must be included in the smooth two-variable factor whose separation is justified. Independent choices of arbitrary w_h satisfying only a separate bounded-variation norm in m do not automatically give (13).

Likewise, replacing $\Lambda(m)$ by genuine-prime coefficients must retain both the original nonprimitive congruence exceptions and the coprime principal-term deletions. Their errors are not part of the identity (4). The author and the other independent reviewer own these full-weight checks.

For a sharp shift interval, (5) – (8) are still available. A Fourier expansion of just the smooth kernel, leaving the interval indicator in each v_ℓ , is a possible rigorous use of (13): its bounded-variation cost grows polynomially with the Fourier frequency and can be absorbed by sufficiently rapid coefficient decay. This observation alone is not a completed verification of the original sharp packet's support and endpoint conventions.

6. Size and limitations

At $Q = X^{523/1000}$, the proposed upper bound is

$$O(X^{1.023} \sqrt{H} \log^4 X).$$

For $X^{1/6} \leq H \leq X^{2/7}$, its ratio to HX is $O(X^{0.23} H^{-1/2} \log^4 X)$, which tends to zero by a power of X . This is a real improvement over accumulating independent per-shift bounds.

However, its ratio to the required $X \log X$ scale is $O(X^{0.23} \sqrt{H} \log^3 X)$, which grows throughout the stated interval. No bound in this note evaluates the selected divisor component to the required covariance precision, controls the complementary divisor remainder, or settles a zeta pair-correlation conjecture.

7. Final author artifact and delta-review scope

The final reviewed author artifact is `SMOOTH_SHIFT_COMPLETION_BOUND.md`, SHA-256

`7b52e4d82dc40bf90183331d548b7fffe5545d1928d7cb93223223b5b71c1d78`.

Its Sections 3 – 5 were read and checked against the independent derivation above. The exact Fourier signs, coprime principal term, zero-frequency cancellation, reduced-denominator grouping, smooth-shift estimate with two finite differences, the two coefficient norms, the logarithm in the Schur row bound, and the centered prime norm are all accepted. The power exponents and the comparison with the covariance scale in its introduction and Section 7 agree with exact rational arithmetic. Its final clarification that the completion bound itself is unconditional is also accepted: none of these arguments uses RH.

The actual two-variable amplitude in author Section 6 is consistent with the uniform separation criterion (13). Its complete support/separation and prime-power-exception review is assigned to the separate `residual_gram` reviewer; the present note does not replace that independent check. In particular, this audit does not enlarge the author theorem from its fixed smooth shift packet to the entire sharp packet. The bounded-variation observation in Section 5 remains a precisely stated extension criterion.

The companion author's exact-check JSON was inspected for its scope and report provenance, but its script was not independently rerun here. No assertion in this audit depends on those finite checks. The written estimates prove the assigned analytical core; they are not a machine-checked formal proof, a new distribution theorem at the required

covariance scale, or a solution of AH or Montgomery's conjecture.

Current report 35: Independent review: actual shift kernel and genuine–prime replacement

Collection: R10 — complete the actual shift packet.

Source: [research/dyson/round10/shift-average/ACTUAL_KERNEL_AND_PRIME_POWER_REVIEW.md](#).

SHA-256: [df6d6602ce9b68ac4759a8cbb8a3caad84d82b4694105a3bb466a9a69fd22732](#) . Git blob: [efd1c98498378d8ea9cb8970eb8c1d4dc1cc6d14](#) . Original bytes: 6575.

Independent review: actual shift kernel and genuine–prime replacement

Date: 2026-09-05. Accepted for the stated smooth packet. This is a narrow ordinary-proof review of Sections 1 – 2 and 6 of [SMOOTH_SHIFT_COMPLETION_BOUND.md](#) , pinned to SHA256

[7b52e4d82dc40bf90183331d548b7fffe5545d1928d7cb93223223b5b71c1d78](#) . The reduced-fraction coefficients and finite-spacing bounds are independently covered by [COEFFICIENT_AND_SPACING_AUDIT.md](#) ; this review does not present that other review as its own work.

No fatal gap was found in the object identification, prime-power deletion, coprime completion prerequisites, or exact separation of the actual two-variable weight. This acceptance does not upgrade the resulting discrepancy estimate to the required zeta covariance scale.

1. The object really is the localized Round 9 discrepancy

Author equation (3) is precisely the Round 9 equation (21) with the additional scalar shift weight $V(h/H)$. The convention for Delta includes both a progression sum and the coprime principal sum. The residue is h ; the shifted argument is $m-h$; the logarithm is $\log((m-h)/q)$. These remain in their correct places. The smooth packet has $H < h < 2H$, and therefore is contained in the earlier hard shift cutoff provided its fixed constant is at least two, as explicitly assumed.

The source modulus family is squarefree and has $q \leq X^{.523}$. No new source-distribution conclusion is being imported: the subsequent completion only uses squarefreeness and this cap. The claimed power improvement concerns this one identified smooth discrepancy component. The unsmoothed full shift range, other divisor components, and continuous covariance centering are not estimated here.

2. Both prime–power portions have been bounded

Write $\theta_*(m) = \log m$ on primes and zero otherwise. On the fixed support, m and $m-h$ are comparable to X , the a -factors and sinc are bounded, and the extra logarithm is $O(\log X)$. For each h and each prime power $m = p^j$ with $j \geq 2$, the progression contribution only involves q dividing $m-h$. Hence the number of possible q is at most $\tau(m-h)$,

independently of the detailed selected family. The bound $\tau(n) \ll_\eta X^\eta$ together with the prime-power von Mangoldt mass $O(\sqrt{X} \log^2 X)$ gives the first term of the author's (4):

$$O_\eta(H X^{(1/2+\eta)} \log^3 X).$$

The principal portion is a different sum and cannot be omitted. Bound its prime-power mass absolutely, multiply by the $O(\log X)$ logarithmic weight and by $\sum_{(q \leq Q)} 1/\phi(q) = O(\log X)$, then sum the $O(H)$ shifts. This gives

$$O(H \sqrt{X} \log^4 X).$$

Retaining or dropping the principal coprimality restrictions by an absolute upper bound is valid. In particular this step accounts for prime powers with primes dividing q as well as those not dividing q ; it does not confuse these with primes dividing h . At $\eta=1/100$ and $H \leq X^{(2/7)}$, the largest displayed power is $X^{(557/700)}$, which is strictly smaller than X . Both nuisance terms are therefore $o(X \log X)$, and also smaller than the claimed completion bound.

Only after this replacement is every surviving m a genuine prime exceeding Q . Then m is a unit modulo every q . In the progression sum, h congruent to m modulo q forces $(h, q) = 1$ automatically, so the h -coprimality restriction may be removed there. It must remain in the principal term, where it becomes the exact unit-residue projector. The order of operations in the draft is correct. Performing this extension before prime-power removal would need extra terms, and the draft does not do that.

All these estimates are unconditional. RH is only needed in the previously stated conversion from this discrepancy to the relevant aggregate zeta covariance component.

3. Uniform smooth separation of the actual kernel

Put $\epsilon = H/X = 1/T$, $y = m/X$ and $z = h/H$. On the support, $y - \epsilon \leq z$ lies in a fixed compact subset of $(1, 3/2)$, and z lies in a fixed compact subset of $(1, 2)$. Thus both original a -factors take their inverse-three-halves branch. Their product is exactly

$$y^{(-3/2)} (y - \epsilon \leq z)^{(-3/2)}.$$

The sinc phase is exactly

$$\begin{aligned} & \epsilon^{(-1)} \log(y/(y - \epsilon \leq z)) \\ &= \int_0^z \frac{du}{(y - \epsilon \leq u)}. \end{aligned}$$

This integral representation also proves the removable $\epsilon=0$ limit z/y and gives uniform bounds for every fixed mixed y, z derivative on a slightly larger fixed rectangle, for all sufficiently small ϵ . It avoids differentiating the misleading $1/\epsilon$ factor separately. The sinc function is smooth at zero, so there is no hidden singularity. The support stays away from $y - \epsilon \leq z = 0$ and from the a -factor branch corner; the minimum definition of a therefore introduces no derivative defect.

The same derivative bounds hold after multiplication by $\log(y - \epsilon \leq z)$. On a larger fixed rectangle, use fixed compact cutoffs and the two-variable Fourier series of the smooth amplitude. Its coefficients decay faster than any fixed power in both frequency indices, uniformly in ϵ . Each term becomes $f_k(m/X) v_l(h/H)$. The prime-factor

estimate used later depends only on the bounded sup norm of f_k , not on its derivatives; the finite-difference estimate for v_l costs at most a fixed polynomial in l . Uniform rapid coefficient decay absorbs that polynomial, so the separated estimates sum absolutely with a uniform constant. No approximation remainder is left unestimated.

Finally the exact cofactor identity is

$$\log((m-h)/q) = \log X - \log q + \log(y - \epsilon z).$$

The constant $\log X$ factor, the $\log q$ conductor coefficient, and the additional smooth logarithmic amplitude are the three legitimate cases treated in the author proof. The resulting $\log^{7/2}$ bound from the first two cases is harmlessly relaxed to $\log^4$. A q -dependent function has not been silently included in a common smooth amplitude.

4. Scope and acceptance

The actual sinc kernel is retained, including its sign. No positivity, variance asymptotic, RH, or convergent critical-strip prime series is used in this component estimate. The completed bound remains larger than $X \log X$ throughout the displayed α interval. The draft says so explicitly.

This review accepts the fixed smooth packet only. A bounded-variation extension of the separated coefficient estimate may be possible, but the full sharp packet is not claimed or required here. No endpoint extension, parameter search, or new numerical prime experiment was performed. The attached receipt pins the reviewed author report, Round 9 source convention, and this review; its exact fraction checks verify the nuisance exponent and the support cap inequality, not asymptotic prime distribution.

Current report 36: Actual–prime variance ranges and a missing logarithmic mixed moment

Collection: R10 — complete the actual shift packet.

Source: `research/dyson/round10/arithmetic-residual/ARITHMETIC_RANGE_AND_MIXED_MOMENT.md`.

SHA-256: `c2d2a278ffe74d8f8d8a7c00980e5e57e6c508790ff96904db4479277d1daa8c` . Git blob:

`fcff74de2703ccf23dc5cd8708dca39e46d994de` . Original bytes: 12502.

Actual–prime variance ranges and a missing logarithmic mixed moment

Date: 2026-09-05. This is a bounded arithmetic-source audit and a conditional calculus lemma. **No improved lower bound for the actual residual or the coupled zeta statistic is proved.** There is no new point-process countermodel, prime-profile search or repetition of the Stieltjes continuation proof.

The negative decision has a concrete arithmetic basis. The checked Guth – Maynard theorem has a short-interval range that misses the shrinking edge region, and even its almost-all PNT error inside that range is much larger than the needed fluctuation scale. The checked RH short-interval variance comparison retains fixed constant losses and fixed-endpoint quantifiers. A further precise mixed moment of the genuine-prime tail would suffice; its numerical threshold is derived below, but that moment is not supplied by either source.

1. Object and the local prime scales

Retain the already reviewed Round 9 definitions $L=\log T$, $N=\text{floor}(T/L^6)$, and the genuine-prime residual $R_-(T,b)(t)$. This is the analytic prime-only continuation at $s=1/2+b/(2L)+it$ with the prime polynomial through N subtracted. It includes the endpoint and pole terms of that continuation. It is not interpreted as a convergent bare prime Dirichlet series on this line.

Write

$$\begin{aligned} r_T^p(b) &= \text{integral}_{0^+}^T |R_-(T,b)(t)|^2 dt / (T L^2), \\ E_T(b) &= \exp(b) r_T^p(b). \end{aligned}$$

The mesoscopic statistic uses the two widths b and $2b$ at the same height and cutoff. The previously reviewed prime-power-removal error is $o(1)$ after its amplified normalization uniformly on $2 \leq b \leq G(T)$, for any $G(T) \rightarrow \infty$ with $G(T) = o(\log \log T)$. Thus it is legitimate to study this genuine-prime residual without pretending that the removed prime powers supply a fixed positive improvement.

To inspect a fixed portion of the first edge, put

$$X=T^{(1+s/b)}, \quad 1 \leq s \leq 2, \quad h=X/T.$$

Then the exact exponent of h relative to X is

$$\log h / \log X = s/(b+s).$$

This tends to zero as b increases. The size h nevertheless exceeds every fixed power of $\log X$ on these slow diagonals. Statements about intervals of length at most a fixed power of $\log X$, or about $X \leq T$, therefore do not cover this edge shell. Conversely, an estimate requiring $h \geq X^\theta$ for a fixed $\theta > 0$ eventually misses it. These are arithmetic scale comparisons, not a conclusion from generic spectral positivity.

2. Guth–Maynard: exact range gap, then a separate moment gap

The primary source inspected is Guth and Maynard, [New large value estimates for Dirichlet polynomials](#), [arXiv:2405.20552v2](#), dated 7 April 2026. Corollary 1.4 gives an almost-all asymptotic for the prime count when

$$X^{(2/15+\epsilon)} \leq h \leq X^{.99},$$

for fixed $\epsilon > 0$. The exceptional set has size $O(X \exp(-(\log X)^{1/4}))$, and on the other integers x around X the count error is $O_\epsilon(h \exp(-(\log X)^{1/4}))$. This is an asymptotic count statement, not a sharp variance statement. The actual improved exponent is $2/15$; this audit does not repeat the obsolete $1/6$ exponent as if it were current.

Even if ϵ were set to zero, the edge shell above would have to satisfy

$$s/(b+s) \geq 2/15, \quad \text{equivalently } b \leq 13s/2.$$

Hence when $b > 13$ the entire fixed shell $1 \leq s \leq 2$ is outside the theorem. At $b=14$, its endpoint exponents are $1/15$ and $1/8$, both below $2/15$. This is an exact range calculation, not a numerical parameter scan. The admissible fixed- ϵ range is smaller still. In terms of $X=T^\alpha$, this corollary's ϵ -to-zero endpoint is $\alpha \geq 15/13$, while the needed shell has $\alpha=1+O(1/b)$. This endpoint is not asserted to exhaust the authors' methods.

The final Remark after the proof of Corollary 1.4 describes a sieve treatment of the critical six-factor case that could slightly lower the fixed exponent to $2/15-\epsilon$, at the cost of a prime-count error roughly $O(\epsilon^4 h/\log X)$. For a fixed sufficiently small ϵ this remains a positive-power lower cutoff for h , so it still misses $s/(b+s) \rightarrow 0$. Its stated count error also does not determine a fluctuation-level second moment or the shrinking variance correction needed below. Thus this remark does not remove either obstruction; no optimality of $2/15$ for the underlying method is claimed.

There is a second and independent obstruction even inside the source range. Converting the count bound to the Chebyshev weight costs $O(\log X)$. Bounding the exceptional intervals trivially then gives at best a consequence of the form

$$\begin{aligned} &\sum_{\text{integer } x \text{ in } [X, 2X]} |\theta(x+h) - \theta(x) - h| \\ &\ll X h^2 (\log X)^2 \exp(-c (\log X)^{1/4}), \end{aligned}$$

with some fixed $c>0$; harmless endpoint and partial-summation errors are smaller here. On the good set one can square the error, but the exceptional-set contribution has only the single exponential saving. The expected fluctuation scale for this variance is $X h \log(X/h)$. For $h=X^\theta$ with any fixed θ in the source range below .99, the quotient of the displayed error bound by that scale contains

$$h \log X \exp(-c (\log X)^{1/4}),$$

which tends to infinity. The quoted corollary alone therefore does not control the variance to constant accuracy, let alone the relative $o(1/b)$ edge accuracy needed here. This diagnoses the precision of this theorem's stated consequence; it does not deny stronger consequences of its underlying methods after further argument.

The large-value Theorem 1.1 in the same source does not directly resolve this issue either. The paper specifies its main improvement over earlier bounds for length $N \leq T^{(5/6-\epsilon)}$, while a single un-factorized prime block at $X=T^{(1+s/b)}$ has length greater than T . A further factorization or arithmetic estimate could exploit the method, but that would be additional work, not an application of its ready-made full-block bound. The independent shifted-modulus lane studies such a different route and is not duplicated here.

3. What the checked RH short-interval variance results provide

The other primary source is Carneiro – Chandee – Chirre – Milinovich, [On Montgomery's pair correlation conjecture: a tale of three integrals](#). Its short-interval statistic is

$$J(\beta, T) = \int_1^T (\psi(x+x/T) - \psi(x) - x/T)^2 dx/x^2.$$

The paper recalls the asymptotic $J(\beta, T) \sim \beta^2 \log^2(T)/(2T)$ for fixed $0 < \beta \leq 1$. That is the side before the first edge. It does not evaluate the increment from $\beta=1$ to $\beta=1+s/b > 1$. Theorem 3 / Corollary 4 give, for large fixed β , constants 0.8376 and 1.4283 multiplying $\beta \log^2(T)/T$. Here β is the exponent delimiting the prime range, not the damping parameter b . Sending the damping b to infinity does not make $\beta=1+s/b$ a large- β instance of that theorem.

For finite endpoint intervals, Theorem 14 and Corollary 15 give a precise comparison with integrated pair correlation, with factors $L_{\text{minus}}=0.9028\dots$ and $L_{\text{plus}}=1.0736\dots$. Both are separated from one by fixed amounts. As a stand-alone comparison, this retains roughly 9.7% lower-side and 7.4% upper-side slack, whereas the first-edge task requires a relative error decreasing like $o(1/b)$. Those fixed losses cannot be read as the required coefficient of $1/b$. Moreover the source's fixed-endpoint liminf/limsup statements do not license a moving endpoint $\beta(T)=1+s/b(T)$ without a uniformity or diagonal argument.

This does not claim that every RH variance theorem is exhausted or that no sharper short-interval argument exists. It records the exact scope and quantitative gap of the primary results checked here. It also avoids replacing the earlier Round 9 failure of individual logarithmic-derivative bounds by a claim that an almost-all prime-count asymptotic has already solved their coupled moment.

4. A concrete mixed moment of the same arithmetic tail

There is a useful way to state the missing arithmetic input more precisely. Define the logarithmically weighted companion using the derivative of the actual analytic residual:

$$\begin{aligned} K_{T,b}(t) &= -R_{T,b}(t) - 2 \operatorname{partial}_b R_{T,b}(t), \\ M_T(b) &= \exp(b)/(T L^2) \\ &\quad \operatorname{Re} \int_0^T R_{T,b}(t) \operatorname{conjugate}(K_{T,b}(t)) dt. \end{aligned}$$

At fixed T this derivative exists on $b>0$ in the relevant critical-strip range under RH. Differentiation of the already available centered-tail integral is justified by its additional logarithmic factor and the same RH bound for $\theta(x)-x$. Endpoint and pole terms are differentiated as well. No derivative estimate is inferred by differentiating an uncontrolled $o(1)$ remainder.

In the absolutely convergent right half-plane this companion weights the genuine-prime tail by $\log(p)/L-1$; in the working strip it is the analytic continuation with its required centering. Thus M_T is a specific correlation with a logarithmic excess weight, not just the norm of either residual and not a freely chosen covariance model.

The exact derivative identity is

$$M_T(b) = -d/db E_T(b).$$

This gives the following conditional criterion. Suppose there are fixed $\epsilon>0$ and an increasing slow envelope $G(T)=o(\log \log T)$ such that, uniformly on $B_0 \leq s \leq 2G(T)$,

$$\begin{aligned} M_T(s) &\geq 1/s^2 - (2-\epsilon)/s^3 - \eta_T/s^3, \\ \eta_T &\rightarrow 0. \end{aligned}$$

Then the original AH-excluding mesoscopic target follows, with a quantitative gap:

$$\begin{aligned} \lim_{B \rightarrow \infty} \liminf_{T \rightarrow \infty} \\ \inf_{B \leq b \leq G(T)} C_T(b) \\ \geq -3/4 + 3\epsilon/8 > -3/4. \end{aligned}$$

This is a sufficient additional arithmetic hypothesis, not a result established for M_T . It is stronger than merely knowing each residual's leading norm, but weaker in a different direction than demanding a full collection of prime pair correlations. It singles out exactly one signed mixed moment along a slow range.

To verify the implication, rewrite the prime-only coupled statistic exactly as

$$\begin{aligned} C_T^p(b) &= b^2 [E_T(b) - E_T(2b) - 1/(2b) \\ &\quad - \exp(-2b)E_T(b) + \exp(-4b)E_T(2b)]. \end{aligned}$$

The already reviewed RH upper estimates of Round 9 imply $E_T(s)=O(1)$ uniformly on the slow range: their absolute finite-height error remains negligible after multiplication by $\exp(s)$, and their limiting residual upper expression is $O(\exp(-s))$. Prime-power removal preserves this conclusion. Therefore the two exponentially small terms contribute

$o(1)$ in the displayed double-limit regime; this uses an actual previously checked upper estimate, not an assumption of perfect residual asymptotics.

Now integrate $M_T(s) = -E_T'(s)$ from b to $2b$. The exact integrals are

$$\begin{aligned} \int_b^{2b} s^{-2} ds &= 1/(2b), \\ \int_b^{2b} s^{-3} ds &= 3/(8b^2). \end{aligned}$$

The asserted lower limit follows, and the reviewed uniform genuine-prime replacement error carries it back to C_T . This reasoning needs control up to $2G(T)$, not only $G(T)$, and retains the uniform range needed for compatibility with the AH slow-diagonal argument. No simplicity assumption enters.

The decisive missing quantity is therefore the coefficient below 2 in this log-weighted mixed-moment deficit. None of the inspected short-interval mean-count or variance theorems supplies that coefficient. A Cauchy – Schwarz estimate using only separate norms does not evaluate this signed mixed moment; a new arithmetic decorrelation, projection estimate, or weighted short-interval covariance result would be needed.

5. Evidence, reproducibility and stopping decision

[check_edge_mixed_moment.py](#) and its [JSON](#) check the exact $2/15$ range threshold, the two fixed $b=14$ endpoint examples, the hyperbolic algebra and the coefficient $3\epsilon/8$. They do not evaluate actual zeta zeros, search parameters, or test the mixed-moment hypothesis.

The source receipt pins the retrieved Guth – Maynard v2 HTML and the already archived CCCM text/PDF together with the relevant Round 9 reports. Third-party source bodies are local references, not material to duplicate in a public code archive. There is no claim that this narrow audit is a complete survey of all prime-variance literature.

The bounded conclusion is negative but specific: the checked improved prime-count theorem misses the growing-damping edge and lacks fluctuation-level precision even where it applies; the checked RH variance comparisons do not supply the required shrinking first correction. The conditional logarithmic mixed-moment inequality is a reviewable next arithmetic obligation. Its proof, or a weaker averaged version sufficient on $[b, 2b]$, remains open here. Repeating finite prime-factor coefficient scans and general PSD models is postponed.

Current report 37: Independent review of the arithmetic range and mixed–moment audit

Collection: R10 — complete the actual shift packet.

Source: `research/dyson/round10/arithmetic-residual/INDEPENDENT_REVIEW.md`.

SHA-256: `375ab6c0ae80f43c86273a38868b045732f547ac9ce2eb3ffa363edaad2bd1a87db5b9a77ecf065e4b9289df3de26c9e72332032` . Git blob: `7db5b9a77ecf065e4b9289df3de26c9e72332032` . Original bytes: 4155.

Independent review of the arithmetic range and mixed–moment audit

Date: 2026-09-05. Reviewer: root. Accepted as a bounded primary-source audit and a conditional calculus implication. No actual-prime lower bound is proved.

The final author report `ARITHMETIC_RANGE_AND_MIXED_MOMENT.md` has SHA256 `c2d2a278ffe74d8f8d8a7c00980e5e57e6c508790ff96904db4479277d1daa8c` . The root read the complete report and script, the actual Guth – Maynard v2 Corollary 1.4 and following Remark, and the saved CCCM Theorem 3/Corollary 4 and Theorem 14/Corollary 15 passages. Earlier Round 9 results are used with their independently reviewed uniformity, not strengthened by assumption.

Source range and error

Corollary 1.4 states the prime-count asymptotic for $h \geq X^{2/15+\epsilon}$, with fixed ϵ , outside an exponentially small proportion of integer starting points. In the edge shell $X = T^{1+s/b}$, $h = X/T$, the exponent is exactly $s/(b+s)$. The corollary's limiting epsilon-zero threshold is $b \leq 13s/2$; the shell $1 \leq s \leq 2$ is outside it once $b > 13$. The source's subsequent Remark permits a slight fixed-epsilon improvement with a weaker error, not an exponent tending to zero. The author now records that distinction explicitly.

The conversion of a prime count to a log-weighted count does not need a new uniform theorem for every inner endpoint. On $[x, x+h]$, $\log p = \log x + O(h/X)$. Multiplying the stated count approximation by $\log x$ introduces the displayed exponential error and an elementary $O(h^2/X)$ remainder (a sharper version is unnecessary). For $h \leq X^{.99}$ that remainder is absorbed by the retained exponential error. The exceptional intervals are bounded trivially. The resulting sufficient upper bound on the squared error is far above $X h \log(X/h)$ for fixed positive h -exponent. This demonstrates a limitation of that direct consequence of the corollary, without claiming that its methods cannot yield stronger information.

The CCCM constants and quantifiers are correctly reported. Its β is the prime-range endpoint, not the damping b . Fixed multiplicative constants 0.9028... and 1.0736... are not an error tending to zero with b . Fixed-endpoint limiting comparisons do not themselves provide the moving-endpoint uniformity required here. This is a narrow audit of

these sources, not an exhaustive impossibility theorem for all known prime-variance methods.

Conditional mixed moment

For the actual analytic residual at fixed T , differentiating the centered integral on $b > 0$ is valid under RH: an extra log factor is integrable at every fixed positive distance from the critical line. Pole and endpoint terms must also be differentiated, as the report states. There is no differentiation of an unspecified asymptotic remainder.

With $E(b) = e^b \left| \left| R_b \right| \right|^2 / (T \log^2 T)$ and $K_b = -R_{b-2} \text{partial}_b R_b$, the real inner-product identity is $M(b) = -E'(b)$. In the absolutely convergent region K has coefficient $\log(p)/\log(T) - 1$, confirming the sign and factor two. Its use in the working strip is analytic continuation, not an unregularized prime series.

The exact two-scale identity includes $-e^{-(2b)}E(b) + e^{-(4b)}E(2b)$. The previously reviewed RH upper estimate, after subtracting the short-prime diagonal, is $O(e^{-b})$ for $r(b)$ on the slow range; its uniform height error and the prime-power replacement stay negligible after multiplication by e^b . Thus $E = O(1)$ there, and the exponential correction vanishes uniformly as the lower cutoff B tends to infinity. A crude pointwise bound alone would not justify this step.

Finally, integrating $1/s^2 - (2-\epsilon)/s^3$ from b to $2b$ gives $1/(2b) - 3(2-\epsilon)/(8b^2)$. After multiplication by b^2 and subtraction of $1/(2b)$, the threshold is exactly $-3/4 + 3\epsilon/8$. The author requires the assumed inequality through $2G(T)$, including a uniform $o(s^{-3})$ error; that is the needed quantifier. This implication is accepted. The mixed-moment hypothesis itself remains completely unproved in these notes, and the source audit supplies no positive ϵ .

The small exact script checks only the scalar algebra and range endpoints. Independent replay is a separate implementation check and does not certify a new analytic estimate.

Current report 38: Round 11: remove the shift–length loss using actual RH prime input

Collection: R11 — RH small arcs and actual conductor structure.

Source: research/reports/dyson_round11.md.

SHA-256: [293125227f1a9eef69441f1d15934ecd970466cce7bb14eb2aa1ee273886e9ad](#) . Git blob: [4ff7d5072a8c3249fb5c4bdf4d4130251d36536c](#) . Original bytes: 8602.

Round 11: remove the shift–length loss using actual RH prime input

Date: 2026-09-05. Under RH, this round improves the same actual smooth shifted-prime discrepancy to $X^{1.023} \log^5 X$. It also proves a narrow arithmetic obstruction to coefficient-only power savings and makes the log-weighted residual's positive diagonal explicit. None of these results proves the required new zeta pair-correlation lower bound.

The component improvement

Keep the Round 10 discrepancy, with $Q=X^{(523/1000)}$, $X=T^\alpha$, $6/5 \leq \alpha \leq 7/5$ and $H=X/T$. It is the actual Mobius-log weighted progression discrepancy selected by the Round 9 complementary conditions, with its original sinc kernel and fixed smooth cutoffs V in $C_c^\infty(1,2)$ and χ in $C_c^\infty(1,3/2)$. The [full proof](#) establishes

$$\left| D_Q^V(X, T) \right| \ll_{V,X} \sqrt{X(X+Q^2)} (\log X)^5 \ll X^{1.023} (\log X)^5 \quad \text{under RH.}$$

This removes $\sqrt{H}$ from Round 10. The additional RH assumption is material; the earlier bound was unconditional. The saving in powers of X is $1/12$ at $H=X^{(1/6)}$, and $1/7$ at $H=X^{(2/7)}$. The bound applies to any squarefree subfamily with the specified cap and therefore does not claim a new dense-divisibility theorem.

After division by the required $X \log X$ covariance scale, the estimate is still $O(X^{1.023} \log^4 X)$. The full sharp shift packet, complementary divisor piece, support main terms and final signed covariance remain outside this bound. In particular, the compact Fourier target greater than $-3/5$ and the two-scale target $1/16$ remain unproved.

Where the arithmetic gain comes from

Bhowmik – Schlage-Puchta, Lemma 3, printed page 3, proves under RH

$$\int_{-1/y}^{1/y} \left| \sum_{n \leq x} (\Lambda(n) - 1) e(\beta n) \right|^2 d\beta \ll (x/y) \log^4 x, \quad 1 \leq y \leq x.$$

Its proof includes the cutoff errors in the Selberg/Gallagher passage. Partial summation transfers it to a smooth genuine-prime polynomial minus its integer mean, $E_f = A_f - B_f$. Prime powers cost no more than the same bound. The frequency derivative is another such polynomial, $E_f' = 2\pi i X E_{-f}$, so its norm is controlled without differentiating an asymptotic error.

The exact Round 10 pairing has coefficients $C_{-}(a/d) = S_{-v}(a/d)M_{-d}$, with

$$M_d = \sum_{q \in \mathbb{Q}, d \mid q} \frac{\mu(q)}{q}, \quad S_v(\beta) = \sum_h v(h/H) e(-\beta h).$$

Equal rational frequencies are already merged. Sampling on disjoint intervals of length comparable to Q^{-2} and retaining only the local arc gives a squared prime norm $O(X(X+Q^2)^{\rho} \log^4 X)$. The coefficient mass on the j th dyadic band is $O(H^{2((1-2J)/j)} \log^3 Q)$, while ρ is at most $2^j/H$. These factors cancel H before summation. $J=2$ suffices to sum all frequency tails.

Both remaining mean terms are retained: the smooth integer mean is handled by Poisson summation, and the primitive Ramanujan mean by a first-power shift bound and the reciprocal-totient sum. Neither introduces a new H loss. The two-variable Fourier separation includes derivative costs in both indices; its rapid decay handles the actual sinc kernel and log cofactor. The original progression prime-power error and the later centered-polynomial prime-power subtraction are distinct, and both are included.

The [independent review](#) checks the primary source, weighted prefixes, derivatives, local sampling, every tail band, both means and the full kernel. Root also read the complete argument. This is an internally reviewed ordinary proof, not formal verification or a novelty claim for the classical ingredients.

Actual prime moduli prevent a coefficient-only shortcut

The [conductor construction](#), with [independent review](#), fixes the full canonical family of all distinct $q=[D,E]$ satisfying the Round 9 balanced complementary predicates and coefficient $\mu(q)$. For fixed nonnegative nonzero V it proves

$$\sum_{d,a}^* |S_v(a/d)M_d|^2 \gg_V H/(\log X)^{348}.$$

The construction uses two primes in a fixed-ratio interval of exponent .09 and 346 primes in one of exponent 343/346000. All factors are distinct. Their product is in $(Q/2, Q]$, has positive Mobius sign, and splits into two roots within the exact radius .2615. Only the large prime in each root triggers its guard, leaving a strict margin .0255. PNT and unique factorization give a positive constant times $Q/\log^{348} X$ different moduli, for every sufficiently large X .

At a terminal conductor $d=q>Q/2$ no other multiple is at most Q , so its full signed coefficient is exactly $1/d$. Enough primitive low numerators have a shift transform bounded below by a positive multiple of H . Summing their squares proves the lower bound. An additional $\log q$ improves its logarithmic power to $\log^{-346} X$.

This excludes an $O(H X^{(-\eta)})$ coefficient-norm bound for any fixed $\eta>0$ on this full family. It does not exclude a specially pruned family, different weights, frequency localization, or cancellation between the actual prime polynomial and these coefficients. It is entirely compatible with the small-arc improvement above. No numerical prime realization or claim of an effective threshold is involved.

The exact positive diagonal leaves a signed arithmetic remainder

The [mixed-tail](#) note, with [independent review](#), starts from the genuine-prime residual R_b at displacement $b/(2 \log T)$, with $N = \text{floor}(T/\log^6 T)$. Define $K_b = -R_b - (\log T)^{-1} \text{partial}_s R_b$ and

$$M_T(b) = \frac{e^b}{T \log^2 T} \text{Re} \langle R_b, K_b \rangle.$$

The finite centered measure uses prime atoms of mass $\log p$ minus Lebesgue measure on $(N, Y]$. RH gives an explicit uniform cutoff error for it and its log-weighted companion. The pole is then removed at $O(e^b \log^{-3} T)$, retaining the endpoint correction. On $2 \leq b \leq 2G(T)$, $G = o(\log \log T)$, the result is

$$M_T(b) = \frac{1}{b^2} + \frac{2}{b^3} + B_T(b) + o(b^{-3}).$$

Here B_T is the combined centered off-diagonal remainder, including both prime-continuum terms and the continuum square. Only their combined limit is justified. The explicit diagonal comes from prime lengths $T^{1+O(1/b)}$; the accessible slice below T is nonpositive and negligible.

A uniform bound $B_T(s) \geq -(4-\epsilon)/s^3 - o(s^{-3})$, or the stated strict integrated improvement, would supply the earlier AH-excluding criterion. **Neither bound is proved.** Chirre's checked derivative identity leaves the unknown out-of-band form-factor integral; its fixed-width asymptotic does not supply this arithmetic estimate. The valid but enormous analytic cutoff is not presented as a practical computation.

Verification, provenance and next step

All 18 original files, 536,670 bytes, are retained locally under [Astra-Local-Archive/round11-originals/](#); 15 research files are public verbatim. Three third-party PDF/text/HTML bodies stay local with URL/hash receipts. The [intake manifest](#) records each file. The [separate replay](#) checks both bounded scripts in a copy, excluding only the temporary primary-source path from the conductor certificate. The complete small-arc JSON is identical. Exact checks cover rational exponents, counting constants, 384 finite arc counts and 2,901 unique frequency memberships; they are not tests of RH itself.

```
python3 research/logs/round11-integration/recheck.py --prime-gap-source /path/to/openai-short-gaps.pdf
python3 tools/verify_manifest.py
```

The next useful attempt must use signed joint prime/modulus cancellation beyond positive sampling, or prove a one-sided estimate for the combined mixed remainder. The coefficient construction alone cannot rule these out. A generic positive-sampling improvement, an unqualified import of the 186 dispersion theorem, and short-interval upper bounds are being checked against their exact hypotheses before further computation.

The long PDFs retain their previously stated checkpoints; this report adds a new source record. Rollback is a revert of this slice. Formalization, a full covariance theorem, and a solved famous conjecture remain outstanding. No new external-model session, large scan, or infrastructure layer was added.

Current report 39: Removing the shift–length loss from the actual completed pairing, under RH

Collection: R11 — RH small arcs and actual conductor structure.

Source: `research/dyson/round11/prime-frequency/CENTERED_SMALL_ARC_BOUND.md`.

SHA-256: `9ebf2d8daaac37702302f5a798611e6a6152e352e5b7dd680c319ba76b3f6e29` . Git blob: `ab5378d9ff03e967ef0c74a01947c00a7a61e835` . Original bytes: 13797.

Removing the shift–length loss from the actual completed pairing, under RH

Date: 2026-09-05. Status: ordinary proof independently audited and accepted for the stated RH component bound; see [the separate review](#). This proves a bound for the same fixed smooth discrepancy component as Round 10, assuming RH. It does not prove the needed zeta covariance estimate, AH failure, Montgomery's conjecture, or a new prime-gap bound. No novelty claim is made for the classical small-arc method.

With the Round 10 notation

```
X=T^alpha, 6/5<=alpha<=7/5,
H=X/T=X^theta, 1/6<=theta<=2/7,
Q=X^(523/1000),
```

the conclusion is

$$|D_Q^V(X,T)| \ll_{(V,\chi)} \sqrt{X(X+Q^2)} (\log X)^5 \\ \ll X^{(1023/1000)} (\log X)^5. \quad (1)$$

Compared with Round 10's unconditional bound $X^{(1023/1000)} \sqrt{H} \log^4 X$, this removes the factor $\sqrt{H}$, at the explicit additional assumption of RH. The saving in the exponent of X is $\theta/2$, ranging from $1/12$ to $1/7$. It still leaves an $X^{(23/1000)}$ power loss relative to $X \log X$, before logarithms. Thus this is a useful component improvement, not the requested historical conjecture.

1. Exact object, ranges and the arithmetic input

Keep the actual squarefree family Q_X contained in $(\sqrt{X}, Q]$, the discrepancy convention and the kernel from the frozen Round 10 report:

```
Delta(F;h mod q) = sum_(m=h mod q) F(m)
                  - phi(q)^(-1) sum_((m,q)=1) F(m),

w_h(u)=chi(u/X) a_u(X) a_(u+h)(X)
```


$$\begin{aligned} & \text{sinc}_0(T \log(1+h/u)), \\ a_u(X) &= \min((u/X)^{1/2}, (X/u)^{3/2}), \\ D_Q^V &= \sum_h V(h/H) \sum_{(q \in Q_X, (q,h)=1)} \mu(q) \\ & \quad \Delta(\Lambda(m) w_h(m-h) \log((m-h)/q); h \bmod q), \quad (2) \end{aligned}$$

where V is fixed in $C_c^\infty(1,2)$, χ is fixed in $C_c^\infty(1,3/2)$, and $\text{sinc}_0(x) = \sin(x)/x$ with its removable value at zero. Functions inside Δ are zero outside their stated support. Formula (1) actually holds for any squarefree subfamily in this modulus interval; no new dense-divisibility property is assumed.

The primary analytic input is Bhowmik – Schlage-Puchta, [Mean representation number of integers as the sum of primes](#), Lemma 3, printed page 3. For

$$\begin{aligned} R_x(\beta) &= \sum_{(n \leq x)} (\Lambda(n) - 1) e(\beta n), \\ e(t) &= \exp(2\pi i t), \end{aligned}$$

it proves under RH that

$$\begin{aligned} & \int_{-1/y}^{-1/y} |R_x(\beta)|^2 d\beta \\ & \ll (x/y) (\log x)^4, \quad 1 \leq y \leq x. \end{aligned} \quad (3)$$

The author's proof uses Selberg's short-interval mean square and Gallagher's lemma, including the cutoff endpoint contributions. We use this already proved centered small-arc bound, not an asymptotic pair-correlation conjecture. The fourth logarithmic power is retained.

We first treat a separated smooth amplitude $f(m/X)v(h/H)$, with f, v compactly supported in fixed positive intervals. Restoring the actual kernel and logarithm is done in Section 6. All reduced frequencies have the complete range

$$2 \leq d \leq Q, \quad 1 \leq a \leq d, \quad \gcd(a, d) = 1, \quad \beta = a/d. \quad (4)$$

There is no zero frequency: its centered completed contribution vanishes exactly. Neither a lower cutoff $d > \sqrt{x}$ nor an inherited cutoff on numerator a is asserted. The arc weights, rather than the original modulus lower bound, suppress small d .

2. Centered genuine primes, smooth weights and derivatives

For fixed smooth f supported in $[c, C]$, with $0 < c < C$ fixed, set

$$\begin{aligned} A_f(\beta) &= \sum_p (\log p) f(p/X) e(\beta p), \\ B_f(\beta) &= \sum_n f(n/X) e(\beta n), \\ E_f(\beta) &= A_f(\beta) - B_f(\beta). \end{aligned} \quad (5)$$

The sum defining B_f is over integers; this is not a silently substituted continuous main term. The polynomial E_f has genuine-prime coefficients minus the integer mean.

First apply partial summation to the Lambda-minus-one prefix polynomials in (3). Minkowski's integral inequality gives, for $1/X \leq \rho \leq 1/2$,

$$\int_{||\beta|| \leq \rho} |\sum_n (\Lambda(n)-1) f(n/X) e(\beta n)|^2 d\beta \ll_f X^\rho (\log X)^4. \quad (6)$$

We only need $\rho \geq 1/H$, so the length $y=1/\rho$ is at most $H=O(X)$ and eventually less than every prefix endpoint $x \geq cX$ occurring in partial summation. For larger constant arcs one may equally use Parseval. The constants depend on the sup norm and total variation of f on a fixed interval.

Replacing Λ by genuine primes subtracts the finite polynomial supported on p^j near X with $j \geq 2$. Its absolute value is at most

$$\sum_{p^j \text{ near } X, j \geq 2} \log p \ll \sqrt{X} (\log X)^2.$$

Its squared integral on this arc is therefore $O(X \rho \log^4 X)$. Consequently (6) holds for E_f itself. There is no use of a critical-strip prime Dirichlet series here.

The derivative is exactly

$$E_f'(\beta) = 2\pi i X E_u(f)(\beta).$$

Applying the same argument with $u(f)$ proves

$$\int_{||\beta|| \leq \rho} |E_f'(\beta)|^2 d\beta \ll_f X^3 \rho (\log X)^4. \quad (7)$$

This controls the frequency derivative by another source-backed centered prime polynomial. No small L_2 estimate has been differentiated formally.

3. Sampling the small arc at distinct Farey frequencies

Distinct reduced fractions in (4) are separated on the circle by at least Q^{-2} . Around each selected frequency choose a disjoint interval of length comparable to Q^{-2} . For any continuously differentiable F , the fundamental theorem of calculus applied to $|F|^2$, followed by averaging on each such interval, gives

$$\begin{aligned} \sum_{\beta \text{ in selected frequencies}} |F(\beta)|^2 &\ll Q^2 \int_U |F(t)|^2 dt \\ &\quad + \int_U |F(t) F'(t)| dt, \end{aligned} \quad (8)$$

where U is their union. One direct proof is to bound $|F(\beta)|^2$ by the interval average plus $2 \int |F F'|$ and then sum over the disjoint intervals. This is a local estimate; the integrals do not need to run over the whole circle.

If $||\beta|| \leq \rho$, then U lies inside $||t|| \leq \rho + O(Q^{-2})$. Throughout our range $\rho \geq 1/H$ and $Q^2 \gg H$, this enlargement is $O(\rho)$. At arcs reaching the whole circle use Parseval for F and F' . Inserting (6) – (7) into (8), and applying Cauchy – Schwarz only to its derivative integral, gives

$$\sum_{d,a \text{ as in (4)}, ||a/d|| \leq \rho} |E_f(a/d)|^2 \ll_f X (Q^2 + X) \rho (\log X)^4. \quad (9)$$

This is the arithmetic gain: centering and the RH prime mean square supply the factor ρ . Merely knowing that the number of frequencies is smaller would not justify (9) with this factor. The local sampling argument does not replace the minimum spacing Q^{-2} by a larger spacing.

4. Dyadic frequency tails, with every coefficient retained

For the separated shift weight put

$$\begin{aligned} S_v(\beta) &= \sum_h v(h/H) e(-\beta h), \\ M_d &= \sum_{q \in Q_X, d|q} \mu(q)/q, \\ C_{(a/d)} &= S_v(a/d) M_d. \end{aligned} \quad (10)$$

Round 10 's exact completion gives the pairing

$$\sum_{(d,a \text{ as in (4)})} C_{(a/d)} [A_f(a/d) - \mu(d) A_f(0)/\phi(d)]. \quad (11)$$

We use only $|M_d| \leq (1 + \log(Q/d))/d$. Finite summation by parts gives, for each fixed integer $J \geq 2$,

$$|S_v(\beta)| \ll_{(v,J)} H(1+||\beta||)^{-J}. \quad (12)$$

Partition all nonzero reduced frequencies into the central arc I_0 with $||\beta|| \leq 1/H$ and the annular bands

$$I_j: 2^{j-1}/H < ||\beta|| \leq \min(2^j/H, 1/2), \quad j \geq 1, \quad (13)$$

stopping when the circle has been covered. For each band, the numerator condition is exactly on $\min(a, d-a)$. In the central arc it is $\min(a, d-a) \leq d/H$; in I_j it lies between $2^{j-1}d/H$ and $2^j d/H$, with the upper half-circle clipping. Nonempty arcs with upper radius ρ require $d \geq 1/\rho$. This treats low denominators, not just the top block.

For fixed d the number of fractions in an arc of upper radius $\rho \leq 1/2$ is at most $2\rho d$, even before imposing coprimality. Using (12), the divisor-coefficient bound, and summing $1/d$ proves

$$\sum_{\beta \in I_j} |C_\beta|^2 \ll_{(v,J)} H 2^{((1-2J)j)} (\log(2Q))^3, \quad (14)$$

with $j=0$ interpreted as the central arc. The factors implicit in the first annulus are absolute constants. No sharp truncation of S_v has been made.

Apply (9) with ρ at most $2^j/H$, or with the whole-circle Parseval bound at the final band. Cauchy – Schwarz on I_j gives

$$\begin{aligned} |\sum_{\beta \in I_j} C_\beta E_f(\beta)| &\ll_{(f,v,J)} \sqrt{X(Q^2+X)} \\ &2^{((1-J)j)} (\log X)^{(7/2)}. \end{aligned} \quad (15)$$

The series over j converges already for $J=2$. Therefore

$$|\sum_{\beta} C_{\beta} E_f(\beta)| \ll_{f,v} \sqrt{X(Q^2+X)} \log^{(7/2)} X. \quad (16)$$

Unlike the unweighted Round 10 sampling, this retains the concentration in a width $1/H$ arc and also bounds every tail band.

5. Both mean terms in the completed pairing

Two terms remain after A_f is replaced by E_f . They are different and are bounded separately.

First, B_f is a smooth integer polynomial. Poisson summation gives, for each fixed $A > 0$,

$$|B_f(\beta)| \ll_{f,A} X(1+X||\beta||)^{-A}.$$

Every nonzero fraction in (4) has $||\beta|| \geq 1/Q$, while X/Q is a positive power of X . Also (12) implies

$$\sum_{a=1}^{d-1} |S_v(a/d)| \ll_v d \quad (17)$$

for every d : compare with an integral when $d \geq H$, and use the convergent $a^{(-J)}$ sum when $d < H$. It follows that $\sum_{\beta} |C_{\beta}| \ll Q \log(2Q)$. Thus the B_f contribution is at most

$$O_{f,v,A}(X Q \log(2Q) (X/Q)^{-A}), \quad (18)$$

which is negligible with a fixed sufficiently large A . This estimate uses the discrete integer mean itself, so no continuous/discrete substitution error is hidden.

Second, the exact primitive Ramanujan principal term in (11) is bounded by

$$\begin{aligned} |A_f(0)| & \sum_{d \leq Q} |M_d| |\mu(d)| / \phi(d) \\ & \sum_{a,d=1}^{\infty} |S_v(a/d)| \\ & \ll_{f,v} X \log(2Q) \sum_{d \leq Q} 1/\phi(d) \\ & \ll_{f,v} X (\log(2Q))^2. \end{aligned} \quad (19)$$

Here $|A_f(0)| \ll_f X$ follows from Chebyshev, and the elementary reciprocal-totient sum is $O(\log Q)$. Using a global Cauchy – Schwarz bound for this principal term would unnecessarily reintroduce an H loss; (19) avoids it without asserting cancellation.

Together (16), (18), and (19) prove the separated pairing estimate. Frequencies $d=1$ remain absent because their completed centered contribution is identically zero, not because a large term was discarded by hand.

6. The actual logarithm, smooth kernel, and prime–power discrepancy

The $\log q$ version of M_d costs one extra logarithm. Its analogue of (14) has $\log^5$ rather than $\log^3$, so (16) has $\log^{(9/2)}$. The corresponding principal term (19) costs at most $\log^3$. The factor $\log X$ is treated explicitly in the same way.

For the actual amplitude, use $\epsilon=1/T$, $y=m/X$, $z=h/H$. The frozen Round 10 proof and its independent kernel audit show that it is exactly

$$V(z) \chi(y-\epsilon z) y^{(-3/2)} (y-\epsilon z)^{(-3/2)} \operatorname{sinc}_0(\int_0^z du/(y-\epsilon u)), \quad (20)$$

and that this function and its product with $\log(y-\epsilon z)$ have uniformly bounded mixed derivatives of every fixed order on a fixed rectangle. The source identity

$$\log((m-h)/q) = \log X - \log q + \log(y-\epsilon z)$$

is retained. Expand the smooth amplitudes in a two-variable Fourier series with fixed outer cutoffs. Unlike the earlier global sampling proof, the constants here depend on the variation of the m -factor as well as the derivatives of the h -factor. Both costs grow polynomially in the two Fourier indices, and the uniform rapid decay of the coefficients absorbs both. Thus the bound sums for the actual kernel, with no approximation error or unrecorded oscillatory derivative loss. We round $\log^{(9/2)}$ up to $\log^5$.

The exact completion used genuine primes to make every m near X a unit modulo q . The original discrepancy (2), however, contains Λ . The frozen Round 10 argument bounds the change in both its progression and principal sums by

$$O_\epsilon(H X^{(1/2+\epsilon)} \log^3 X + H \sqrt{X} \log^4 X). \quad (21)$$

This is logically distinct from subtracting the finite prime-power polynomial in (6). The former makes the modulus completion exact; the latter identifies the centered prime polynomial used in the source mean-square estimate. Both are included. Fixing $\epsilon=1/100$ makes (21) $o(X \log X)$ uniformly for $H \leq X^{(2/7)}$, so it is smaller than (1). This completes the proof.

7. Quantitative decision and remaining arithmetic task

The improved exponent is 1.023 throughout the entire specified θ range. At $\theta=1/6$ the previous exponent was $3319/3000=1.106333\dots$, and at $\theta=2/7$ it was $8161/7000=1.165857\dots$. The conditional power savings are respectively $1/12$ and $1/7$. They are comparisons between an RH consequence and the earlier unconditional component estimate; the assumption change is material.

After the actual covariance normalization $1/(X \log T)$, a bound of this size is still $O(X^{.023} \log^4 X)$, which does not tend to zero. Ordinary RH small-arc control therefore removes the shift-length loss but does not settle the selected divisor component at the required precision. Replacing (3) by a conjectural sharp variance asymptotic would not by itself remove the Q^2 sampling scale used in this proof.

The remaining task is quantitative cancellation in the actual signed pairing of the conductor coefficients M_d and the centered genuine-prime exponential sums, or another argument improving the aggregate top-conductor contribution beyond this sampling/Cauchy – Schwarz bound. No statement that such cancellation is impossible is made. The independent conductor lane checks whether the source support removes high d ; this proof does not assume that it does.

The only computation here is a small exact check of exponents, finite arc counting, and the dyadic geometric factors. The proof uses no fitted constants, zeta-zero sample, coefficient search, or unbounded scan. Other divisor pieces and omitted covariance ranges remain outside its scope.

Current report 40: Independent review of the RH centered–small–arc improvement

Collection: R11 — RH small arcs and actual conductor structure.

Source: `research/dyson/round11/prime-frequency/SMALL_ARC_INDEPENDENT_REVIEW.md`.

SHA-256: `fdeba2af8e447ed7ec382155496f22e23255267f0d4051afdc3759a1f933daab` . **Git blob:** `bc9dc853d1a9f9046caac3a8d59c91e701549bfff` . **Original bytes:** 9325.

Independent review of the RH centered–small–arc improvement

Date: 2026-09-05. **Reviewer:** `yau_flow` , separately from the author `residual_gram` .

Verdict: the written proof of

$$|D_Q^V(X, T)| \ll_{V, X} \sqrt{X(X+Q^2)} \log^5 X, \quad Q = X^{523/1000},$$

is accepted under RH for the fixed smooth packet and modulus family specified in the author report. The removal of the previous $\sqrt{H}$ factor is supported by actual centered-prime input. No lost tail band, undeclared GRH assumption, or omitted mean term was found. The bound remains larger than the required $X \log X$ covariance scale. This is an ordinary independent proof audit, not formal verification or a claim of novelty.

1. Primary source and its exact applicability

I read Lemma 3 and its proof on printed page 3 of Bhowmik – Schlage-Puchta, [Mean representation number of integers as the sum of primes](#), together with the preceding definitions and Lemmas 1 – 2. The PDF SHA-256 is `6cf48524eb9473cca93c1ce0ea97e00fc7dab7c49a3db4e01a490cc165296607` .

The polynomial is exactly the finite prefix with coefficients $\Lambda(n) - 1$. The centered interval is $[-1/y, 1/y]$, and the upper bound is $(X/y) \log^4 X$ under RH. The source states $y \leq X$; the application uses the unproblematic subrange $1 \leq y \leq X$. Its proof includes the two prefix-edge regions, so no unproved endpoint cancellation is being imported. The source's separate discussion of GRH near general rational arcs is not needed here: this proof uses its RH arc around zero, then a local sampling inequality at points within that arc.

2. Weighted prefixes, genuine primes, and frequency derivatives

Let f have fixed compact support in $[c, C] \subset (0, \infty)$. Stieltjes partial summation expresses the weighted $\Lambda - 1$ polynomial as an integral of prefix polynomials with endpoints $x \in [cX, CX]$, plus any endpoint terms. Minkowski's inequality gives a factor controlled by $\|f\|_\infty + \int |f'|$. On the needed arcs, $\rho \geq 1/H$, so $y = 1/\rho \leq H = o(X)$, eventually below every prefix endpoint. The source therefore applies uniformly, yielding

$$\int_{\|a\| \leq \rho} |E_f(a)|^2 da \ll_f X \rho \log^4 X. \quad (1)$$

Here $E_f = A_f - B_f$, with A_f genuinely prime-supported and B_f the integer polynomial. Replacing the von Mangoldt coefficients by prime coefficients subtracts a finite prime-power polynomial. Its sup norm is $O_f(\sqrt{X} \log^2 X)$, so its squared integral is $O_f(X \rho \log^4 X)$. This proves (1) for exactly the polynomial used later.

The derivative identity

$$E'_f(\alpha) = 2\pi i X E_{uf(u)}(\alpha)$$

is exact coefficient by coefficient. Applying (1) with the new fixed weight $uf(u)$ gives

$$\int_{\|\alpha\| \leq \rho} |E'_f(\alpha)|^2 d\alpha \ll_f X^3 \rho \log^4 X. \quad (2)$$

Thus no estimate has been formally differentiated. The needed seminorms of both f and uf are accounted for.

3. Local sampling retains the original Farey spacing

All distinct reduced fractions with $2 \leq d \leq Q$ have circular separation at least Q^{-2} . Take disjoint intervals of length, for example, $Q^{-2}/2$, centered at the selected fractions. For each continuously differentiable F , the fundamental theorem applied to $|F|^2$ gives at each center

$$|F(\beta)|^2 \leq |I_\beta|^{-1} \int_{I_\beta} |F|^2 + 2 \int_{I_\beta} |FF'|.$$

Summing preserves disjointness. If the centers satisfy $\|\beta\| \leq \rho$, their intervals lie in the enlarged arc of radius $\rho + Q^{-2}/4$. Because $\rho \geq 1/H$ and $Q^2 \gg H$, this enlargement is at most a constant times ρ . Combining (1) – (2) and Cauchy – Schwarz only in the derivative integral yields

$$\sum_{\|a/d\| \leq \rho}^* |E_f(a/d)|^2 \ll_f X(Q^2 + X) \rho \log^4 X. \quad (3)$$

For a band reaching the whole circle, Parseval for the polynomial and its derivative gives the same bound with $\rho \asymp 1$. This handles wraparound and the final band. The argument neither improves the spacing to $1/(HQ)$ nor assumes each small interval is a new arithmetic major arc.

4. Coefficient bands and their summation

The exact Round 10 completion has coefficients

$$C_{a/d} = S_\chi(a/d) M_d, \quad M_d = \sum_{\substack{q \in Q_X \\ d|q}} \frac{\mu(q)}{q}.$$

The bound $|M_d| \leq (1 + \log(Q/d))/d$ and smooth finite-difference estimate $|S_\chi(a)| \ll_{\chi, J} H(1 + H\|\alpha\|)^{-J}$ do not require Möbius cancellation.

For any denominator d , at most $2pd$ nonzero residues satisfy $\|a/d\| \leq \rho$; if $pd < 1$, there are none. Hence, for the central arc $j = 0$ and the annuli with upper radius $2^j/H$,

$$\sum_{\beta \in I_j} |C_\beta|^2 \ll_{v,J} H 2^{(1-2J)j} \log^3(2Q). \quad (4)$$

This counts every reduced denominator from 2 through Q , including divisors much smaller than the original moduli. There is no inherited cutoff $d > \sqrt{X}$. The last clipped annulus obeys the same estimate since its upper radius is only enlarged for this upper bound.

Combining (3) – (4) on each band gives

$$\left| \sum_{\beta \in I_j} C_\beta E_f(\beta) \right| \ll_{f,v,J} \sqrt{X(Q^2 + X)} 2^{(1-J)j} \log^{7/2} X.$$

The H factors cancel. The geometric series converges for $J = 2$, independently of the number of bands. Thus no hidden $\log H$ or discarded high-frequency tail is needed. A $\log q$ factor in the conductor coefficient contributes one extra logarithm, giving $\log^{9/2} X$.

5. The integer mean and the primitive mean are both present

The decomposition being bounded is exactly

$$A_f(a/d) - \frac{\mu(d)}{\phi(d)} A_f(0) = E_f(a/d) + B_f(a/d) - \frac{\mu(d)}{\phi(d)} A_f(0).$$

The two terms after E_f must be handled separately.

For smooth compact f , Poisson summation gives $|B_f(a)| \ll_{f,A} X(1 + X \|a\|)^{-A}$. Since every nonzero fraction has distance at least $1/Q$ from the integers, and $\sum_{a=1}^{d-1} |S_v(a/d)| \ll_v d$, one has $\sum_\beta |C_\beta| \ll_v Q \log(2Q)$. The resulting integer-mean error is $O_{f,v,A}(XQ \log(2Q)(X/Q)^{-A})$. A fixed sufficiently large A makes it negligible. The estimate is on the discrete integer mean itself.

For the primitive mean, the same first-power shift bound gives

$$|A_f(0)| \sum_{d \leq Q} \frac{|M_d \mu(d)|}{\phi(d)} \sum_{(a,d)=1} |S_v(a/d)| \ll_{f,v} X \log^2(2Q).$$

This uses Chebyshev and $\sum_{d \leq Q} 1/\phi(d) \ll \log(2Q)$. It introduces no H loss and claims no unproved cancellation. The denominator-one term is identically zero in the original completion; it was not discarded during this later decomposition.

6. Actual kernel, the two prime–power uses, and scope

The smooth kernel is the exact Round 10 kernel in the variables $y = m/X$, $z = h/H$, $\epsilon = 1/T$, and its phase equals $\int_0^z (y - \epsilon u)^{-1} du$. The compact support stays away from the denominator singularity. All fixed mixed derivatives, including those after multiplication by $\log(y - \epsilon z)$, remain uniformly bounded as $\epsilon \rightarrow 0$.

In the Fourier separation with fixed outer cutoffs, the constants in (1) – (3) grow polynomially with the m -frequency through the variations of f and uf . The Poisson bound also costs a fixed number of m -derivatives. Smooth shift bounds cost a fixed number of h -derivatives. Uniform rapid decay of the two-variable Fourier coefficients absorbs all

of these costs. The author explicitly includes both m - and h -frequency dependence. The logarithmic identity $\log((m-h)/q) = \log X - \log q + \log(\gamma - \epsilon z)$ is exact. Rounding the resulting $\log^{9/2}$ to $\log^5$ is legitimate.

There are two distinct prime-power steps. The local polynomial replacement proving (1) was checked in Section 2 above. Separately, replacing Λ inside the original progression discrepancy is needed to make every remaining prime a unit modulo every modulus. Its progression error is bounded by the divisor count of $m-h$; its principal error uses the reciprocal-totient sum. Both are at most the displayed Round 10 error

$$O_{\eta}(HX^{1/2+\eta}\log^3 X + H\sqrt{X}\log^4 X).$$

For $\eta = 1/100$ and $H \leq X^{2/7}$, this is $\mathcal{O}(X \log X)$ and below the new bound. It is not the same operation as the local prime-power replacement, and neither step is omitted.

The accepted statement is restricted to the fixed smooth packet $V \in C_c^\infty(1, 2)$, the actual stated kernel, and the chosen squarefree family. It is conditional on RH. At $Q = X^{.523}$, it leaves $X^{.023} \log^4 X$ after covariance normalization. No bound for the entire sharp packet, complementary divisor piece, or desired conjecture follows from this audit.

7. Evidence and final version

The final author artifact is `CENTERED_SMALL_ARC_BOUND.md`, SHA-256

`9ebf2d8daaac37702302f5a798611e6a6152e352e5b7dd680c319ba76b3f6e29`. The full mathematical draft and primary source were read; the final status paragraph and its link to this review were checked after the author froze the mathematical text. No numerical experiment was necessary, and the author's finite exact-check script was not independently rerun. The acceptance concerns the ordinary analytical proof and its stated scope, not the status of any unreviewed later revision.

Current report 41: A real–prime subfamily prevents a coefficient–only power saving

Collection: R11 — RH small arcs and actual conductor structure.

Source: [research/dyson/round11/conductor-arithmetic/CONDUCTOR_MASS_LOWER_BOUND.md](#).

SHA-256: [46347799005bb0f53af25c2a7e8ffb2b2217d92688c7651327dde3562f114b92](#) . Git blob: [6b4ccb2eba477c26215aa12697d9fa620aba6c4b](#) . Original bytes: 11083.

A real–prime subfamily prevents a coefficient–only power saving

Date: 2026-09-05. Status: an ordinary arithmetic construction and quantitative counting proof. No numerical prime realization is required or used. The conclusion concerns a completed coefficient norm; it does not obstruct cancellation in the joint pairing with prime exponential sums.

Fix the full canonical complementary modulus family defined below. For any fixed nonnegative $V \in C_c^\infty(1, 2)$ that is not identically zero, the Round 10 completed coefficients satisfy, uniformly for $X^{1/6} \leq H \leq X^{2/7}$,

$$\sum_{\substack{2 \leq d \leq X^{523} \\ 1 \leq a < d, (a, d) = 1}} |C_X(a/d)|^2 \gg_V \frac{H}{(\log X)^{348}}. \quad (1)$$

For the coefficients with an additional $\log q$, the lower bound is $\gg_V H/(\log X)^{346}$.

Consequently, neither coefficient norm admits an upper bound $O(HX^{-\eta})$, for any fixed $\eta > 0$, on this full family. The result uses actual squarefree integers made from primes in explicit intervals, with all complementary predicates verified and their number estimated by the prime number theorem.

This is narrower than a no-go theorem for the research programme. A specially pruned family may remove these moduli. The actual prime exponential sums may also cancel against coefficients whose squared mass is large.

1. Fixing the family and the inherited divisor budget

Set

$$\rho = \frac{523}{1000}, \quad r = \frac{523}{2000}, \quad b = \frac{501}{2000}, \quad \delta = \frac{1}{1000}, \quad Q = X^\rho, \quad Y = X^\delta. \quad (2)$$

Define $\mathcal{Q}_X^{\text{full}}$ to contain every distinct squarefree modulus $q = [D, E]$ for which positive squarefree D, E obey

$$\begin{aligned} D, E &\leq X^r, & [D, E] &> X^{1/2}, \\ p^{3/2} D_{\geq p} &\leq X^b & (p \mid D, p > Y), \\ p^{3/2} E_{\geq p} &\leq X^b & (p \mid E, p > Y). \end{aligned} \quad (3)$$

Each q is counted once, irrespective of the number of representations. The coefficient on q is exactly $\mu(q)$; no extra Selberg or cutoff coefficient is inserted.

These are the actual balanced complementary predicates fixed in Round 9, not every possible support geometry in the 186 paper. They use $f(p) = g(p) = p^{3/2}$, equal budgets $A_0 = B_0 = X^b$, and threshold $Z = X^{1/2}$. The opposite-root guards follow from the owner bounds, and $A_0 B_0 = ZY$. Hence Proposition 2.3 of [OpenAI, Improved short gaps between primes](#), printed pp.4 – 5, proves

$$Q_X^{\text{full}} \subset \{X^{1/2} < q \leq Q : q \in D^{(3)}(Y), q \text{ squarefree}\}. \quad (4)$$

The elementary divisor property stated on printed p.4 gives, for $d \mid q$,

$$d \in D^{(3)}(Yq/d) \subseteq D^{(3)}(YQ/d). \quad (5)$$

Thus a conductor of size $d = X^\beta$ retains the quantified parameter $X^{\delta+\rho-\beta}$ in this sufficient class. This is not inheritance with unchanged Y for an arbitrary divisor. At the conductors constructed below, $d=q$ itself, so the full original budget Y is retained with no loss.

Define the actual regrouped coefficients

$$\begin{aligned} A_X(d) &= \sum_{\substack{q \in Q_X^{\text{full}} \\ d \mid q}} \frac{\mu(q)}{q}, & C_X(a/d) &= S_{V,H}(a/d) A_X(d), \\ S_{V,H}(\beta) &= \sum_{h \in \mathbb{Z}} V(h/H) e(-\beta h), & \epsilon(t) &= e^{2\pi i t}. \end{aligned} \quad (6)$$

The frequencies a/d in (1) are reduced. These are the coefficients in the exact conductor regrouping of Round 10, equation (11).

2. A subfamily made from 348 distinct prime factors

Put

$$u = \frac{9}{100}, \quad \kappa = \frac{343}{346000}, \quad \lambda = 2^{-1/348}. \quad (7)$$

Let $P_L(X)$ and $P_S(X)$ be the primes in

$$(\lambda X^u, X^u], \quad (\lambda X^\kappa, X^\kappa], \quad (8)$$

respectively. For all sufficiently large real X , these two intervals are disjoint, the large primes exceed Y , and the small primes are below Y . Indeed $\kappa < \delta < u$, all strictly.

Let F_X be the set of products of two distinct primes from $P_L(X)$ and 346 distinct primes from $P_S(X)$. The use of sets, rather than ordered prime tuples, prevents permutation overcounting.

Every q in this family is squarefree and has

$$\mu(q) = (-1)^{348} = 1. \quad (9)$$

The exact exponent identities are

$$2u + 346\kappa = \rho, \quad u + 173\kappa = r, \quad \lambda^{348} = \frac{1}{2}. \quad (10)$$

It follows that

$$\frac{Q}{2} < q \leq Q. \quad (11)$$

For large X , $Q/2 > X^{1/2}$. Partition the 346 small primes into two groups of 173 and place one large prime in each group. Their products D, E are disjoint and squarefree, $q = DE = [D, E]$, and

$$D, E \leq X^{u+173\kappa} = X^r.$$

Only the single large prime p in each root activates its predicate. Its owner tail is exactly p , so

$$p^{3/2} D_{\geq p} = p^{5/2} \leq X^{(5/2)u} = X^{9/40} < X^{501/2000}. \quad (12)$$

The strict exponent margin is 51/2000. The same reasoning applies to E . The opposite-root condition $p^{3/2} \leq X^b$ holds as well. Thus

$$\boxed{F_X \subset Q_X^{\text{full}} \cap (Q/2, Q].} \quad (13)$$

The construction is not vacuous Y -smooth support: every modulus has two prime factors larger than Y , and both owner predicates are checked at those factors.

3. Rigorous counting, including the permutation constants

Write

$$L_X = \pi(X^u) - \pi(\lambda X^u), \quad S_X = \pi(X^\kappa) - \pi(\lambda X^\kappa).$$

Unique factorization, together with the disjoint prime intervals, gives the exact count

$$|F_X| = \binom{L_X}{2} \binom{S_X}{346}. \quad (14)$$

The prime number theorem on each fixed-ratio interval gives, as real X tends to infinity,

$$L_X \sim \frac{1-\lambda}{u} \frac{X^u}{\log X}, \quad S_X \sim \frac{1-\lambda}{\kappa} \frac{X^\kappa}{\log X}.$$

Since both counts tend to infinity, (14) yields

$$|F_X| \sim c_0 \frac{Q}{(\log X)^{348}}, \quad c_0 = \frac{(1-\lambda)^{348}}{2! 346! u^2 \kappa^{346}} > 0. \quad (15)$$

This assertion holds for every sufficiently large X , not merely a specially chosen subsequence. The constant is small, but fixed and positive; no estimate of its numerical size is needed for a power-saving obstruction. The factorization into D, E was used only to prove membership. Its many possible partitions are not counted again in (14).

4. No other signed modulus can cancel these conductor coefficients

Take $d \in F_X$. Since $d > Q/2$, the only positive multiple of d at most Q is d itself. Therefore, in the **full signed family**, not just the positive subfamily,

$$\boxed{A_X(d) = \frac{\mu(d)}{d} = \frac{1}{d}.} \quad (16)$$

This conclusion is unaffected by any other admissible moduli, including those with negative Möbius coefficient. It is an isolation statement at the reduced denominator d ; no unmerged duplicate rational frequencies are being counted.

It also shows why the inherited dense-divisibility parameter in (5) cannot simply remove large conductors. Here $d=q$ is near the largest allowed modulus and still belongs to the strongest original class $D^{(3)}(Y)$.

5. Enough primitive low numerators, and the norm lower bound

Let $m_V = \int V(t) dt > 0$. The Riemann-sum estimate, uniformly as H tends to infinity, gives

$$\sum_h V(h/H) \geq \frac{m_V}{2} H \quad (17)$$

for sufficiently large H . This is uniform over the present range $H \geq X^{1/6}$.

If $1 \leq a \leq d/(16H)$, every h in the support of $V(h/H)$ satisfies $0 \leq 2\pi ah/d \leq \pi/4$. Positivity of V therefore gives

$$|S_{V,H}(a/d)| \geq \Re S_{V,H}(a/d) \geq \frac{m_V}{2\sqrt{2}} H. \quad (18)$$

We must still count only primitive fractions. Put $A = d/(16H)$. Each prime factor of d exceeds λX^κ , and d has exactly 348 such factors. Consequently

$$\#\{1 \leq a \leq A : (a, d) > 1\} \leq \sum_{p|d} \left\lfloor \frac{A}{p} \right\rfloor \leq A \frac{348}{\lambda X^\kappa}. \quad (19)$$

Uniformly for $d \in F_X$ and $H \leq X^{2/7}$, one has $A \geq X^{\rho-2/7}/32 \rightarrow \infty$. For all sufficiently large X , take $A \geq 4$ and $348/(\lambda X^\kappa) \leq 1/4$. Equations (19) and $\lfloor A \rfloor \geq A - 1$ then give at least $A/2 = d/(32H)$ coprime choices of a .

For each such d , (16) and (18) imply

$$\sum_{\substack{1 \leq a < d \\ (a,d)=1}} |C_X(a/d)|^2 \geq \frac{m_V^2}{256} \frac{H}{d}. \quad (20)$$

Summing over F_X , using $d \leq Q$ and (15), proves the explicit eventual bound

$$\boxed{\sum_{d,a} |C_X(a/d)|^2 \geq \frac{c_0 m_V^2}{512} \frac{H}{(\log X)^{348}}.} \quad (21)$$

All sufficiently-large- X conditions above can be imposed simultaneously, independently of H in its stated range.

For $A_X^{(1)}(d) = \sum_{q \in Q_X^{\text{full}}, d|q} \mu(q) \log q/q$, the same conductors have $A_X^{(1)}(d) = \log d/d$. Because $d > X^{1/2}$, their squared contribution is at least $(\log X)^2/4$ times the restricted contribution in (21). Hence

$$\sum_{d,a} |S_{V,H}(a/d)A_X^{(1)}(d)|^2 \geq \frac{c_0 m_V^2}{2048} \frac{H}{(\log X)^{346}} \quad (22)$$

eventually.

6. Exactly what is ruled out

For every fixed $\eta > 0$, $X^\eta/(\log X)^{348} \rightarrow \infty$. Thus (21) contradicts any proposed bound $O(HX^{-\eta})$ for this coefficient norm on the full canonical family. More generally it rules out an upper bound of that strength asserted uniformly over all allowed subfamilies, since F_X itself is an allowed family and has the same isolated conductor coefficients.

The quantifiers matter:

- A different deliberately pruned family can exclude F_X ; no lower bound is asserted for every selected subset.
- A different coefficient system can suppress these moduli; the proof concerns the specified Möbius coefficients, not arbitrary sieve weights.
- The lower bound holds for fixed nonnegative nonzero V as stated. It is not asserted for every signed shift profile or every Fourier component of the full two-variable kernel.
- It supplies no lower bound for the signed joint pairing with $A_f(a/d) - \mu(d)A_f(0)/\phi(d)$, and no lower bound for the actual zeta covariance.

The conclusion is therefore specific: retaining a smaller divisor-density budget cannot by itself furnish a power saving in the completed coefficient norm for this canonical source-supported family. To improve the Round 10 bound by a power, this family requires additional cancellation in the joint prime pairing, a different use of the frequencies, or a justified change of the chosen modulus weights/support.

7. Provenance and exact arithmetic

The only counting input is the classical prime number theorem, applied to the two fixed-ratio prime intervals in (8). The support input is Proposition 2.3 and the divisor property on p.4 of the pinned 186 paper. The source PDF SHA256 is 456f05e0a3ef589ebb0e9abcf31f140f3c945adbf6950e00ef371a3c88b0930.

The companion exact-arithmetic certificate records the rational exponent identities and positive margins in (7), (10), and (12), plus the uniform lower-numerator exponent. It does not claim an effective numerical threshold at which the prime-counting asymptotic has stabilized. No scan, numerical realization, or rerun of an old asymptotic bound was performed.

Current report 42: Independent review of the canonical conductor–mass construction

Collection: R11 — RH small arcs and actual conductor structure.

Source: research/dyson/round11/conductor-arithmetic/INDEPENDENT_REVIEW.md.

SHA-256: [b146b7427417bc3acca5443e167e3a07257830cf02523a6fd1f3e252d](#) . Git blob: [a543227dac25c0c76b8817c8f2902cc1779e3845](#) . Original bytes: 3863.

Independent review of the canonical conductor–mass construction

Date: 2026-09-05. Reviewer: root Astra, independently of the authoring lane. Accepted as an ordinary asymptotic arithmetic proof for the explicitly defined full family. This is not a lower bound for the prime pairing, and not a formal proof-assistant certificate.

Reviewed author SHA256: [46347799005bb0f53af25c2a7e8ffb2b2217d92688c7651327dde3562f114b92](#) .

Membership and counting

The exact exponents are consistent: two primes of exponent $9/100$ and 346 of exponent $343/346000$ give total $523/1000$; each root has one large prime and 173 small primes, giving $523/2000$. The lower interval multiplier raised to 348 is exactly $1/2$. Thus every constructed product lies strictly above $Q/2$ and at most Q , for every sufficiently large real X .

The small-prime exponent is strictly below $1/1000$, so those factors do not trigger the complementary predicate. Each root has only one factor above that threshold; its owner tail is the factor itself. The active condition is therefore $p^{(5/2)} \leq X^{(9/40)}$, strictly inside $X^{(501/2000)}$. The opposite guard is weaker. Both roots are squarefree and disjoint; their least common multiple is their product. This checks the actual Round 9 support, rather than merely assuming that dense divisibility implies membership in it.

Unique factorization and disjoint intervals give exactly $\text{binomial}(L_X, 2) \text{binomial}(S_X, 346)$ different moduli. PNT in fixed-ratio intervals gives the stated constant and log exponent. Counting root partitions as different moduli would be wrong; the author does not do that. No effective or practical threshold is claimed, and the very small fixed counting constant does not affect the power comparison.

Signed coefficients and primitive fractions

At a constructed $d > Q/2$ no other positive multiple $d \cdot k$ can be at most Q . Thus the full signed regrouped coefficient is exactly $\mu(d)/d = 1/d$; negative contributions from other parent moduli cannot cancel it. Each reduced fraction is counted once. This is a conductor-isolation fact, not a claim that all moduli have positive coefficient.

For $1 \leq a \leq d/(16H)$, every active h is between H and $2H$, so the phase has absolute value at most $\pi/4$. Fixed nonnegative V gives the lower real part claimed. The union bound over 348 prime divisors removes at most $A \cdot 348 / (\lambda X^\kappa)$ candidate numerators, while the floor costs at most one. Since $A \geq X^{(523/1000 - 2/7)/32}$ grows uniformly, at least $A/2$ primitive numerators remain eventually. This proves the individual constant $m_V^2 H / (256 d)$, then $c_0 m_V^2 H / (512 \log^{348} X)$. The extra $\log d$ gives at least $(\log X)^{2/4}$ after squaring, yielding the second constant $1/2048$.

Scope and relation to the new RH estimate

For fixed $\eta > 0$, $X^\eta / \log^{348} X$ diverges. The norm lower bound therefore rules out $O(H X^{-(\eta)})$ for these exact coefficients on this full family. It also refutes such an assertion made uniformly over all allowed subfamilies, since this one qualifies. It says nothing about every deliberately pruned family or altered sieve weight.

In particular it does not obstruct the separate Round 11 RH small-arc improvement, which changes the prime-frequency estimate and its localization. Nor does it give a lower bound for the signed pairing of these coefficients with actual centered prime exponential sums. A claim that the remaining $X^{.023}$ loss is unavoidable for primes would exceed the proof.

The inherited dense-divisibility parameter for a general divisor is Yq/d , not always Y . For the conductors used here $d=q$, so no loss of parameter occurs. The author records this distinction explicitly.

The exact-arithmetic companion verifies rational exponents and constants only. PNT, the fixed smooth Riemann sum, and the written counting argument supply the asymptotic assertions. No prime search or numerical realization was needed.

Current report 43: The log-weighted prime-tail moment: a primary-source check and its arithmetic remainder

Collection: R11 — RH small arcs and actual conductor structure.

Source: [research/dyson/round11/log-weighted-tail/ARITHMETIC_DIAGONAL_AND_SOURCE_GAP.md](#).

SHA-256: [85c9435bbd656bcf37f21f55fe2d54cb40638c7a56908abe57e54f551a627daa](#) . Git blob: [60dcf395d101e223087db8ac4f355640de016fff](#) . Original bytes: 8660.

The log-weighted prime-tail moment: a primary-source check and its arithmetic remainder

Date: 2026-09-05. This bounded attempt proves no new lower bound at the required scale. It identifies a relevant derivative theorem, checks a legitimate centered finite-prime expansion, and isolates the missing arithmetic coefficient. The positive prime diagonal is explicit; the required lower bound on the jointly centered remainder is not established.

1. A derivative theorem does not supply the missing mixed moment

The directly relevant primary source is Andrés Chirre, [A note on the mean values of the derivatives of \$\zeta'/\zeta\$](#) , [arXiv:2107.13636v2](#), 3 January 2022. Proposition 5, equation (2.3), expresses the normalized squared norm of the k th logarithmic derivative as a known integral on $[0, 1]$ plus

$$\int_1^{\infty} \alpha^{2k} e^{-2\alpha a} F(\alpha, T) d\alpha + o(1)$$

for fixed $a > 0$. Theorem 1 makes its sharp fixed-width asymptotic equivalent to pair correlation. Thus this theorem does not provide a sharper RH-only arithmetic value for the signed mixed moment considered here. Its fixed-width remainder cannot be differentiated or made uniform along a growing width without further proof. The source HTML and its hash are retained in [sources/receipt.json](#) .

This is a narrow source conclusion, not an assertion that every theorem on derivatives or short intervals has been exhausted. The stronger centered-small-arc input being investigated separately is outside this note.

2. A finite, centered arithmetic expansion with its cutoff error

Assume RH. Let $L = \log T$, $N = \lfloor T/L^6 \rfloor$, $s = 1/2 + \delta + it$, and $\delta = b/(2L) > 0$. Let R_b be the genuine-prime residual already defined in Round 10, including its pole and endpoint. Define

$$K_b = -R_b - L^{-1} \partial_s R_b, \quad M_T(b) = \frac{e^b}{TL^2} \Re \int_0^T R_b \overline{K_b} dt.$$

The derivative acts on the actual analytic continuation, not on a bare prime series in the critical strip.

For finite $Y > N$, put

$$C_Y(s) = \sum_{N < p \leq Y} (\log p) p^{-s} - \int_N^Y x^{-s} dx, \quad D_Y = -C_Y - L^{-1} \partial_s C_Y. \quad (1)$$

These are finite expressions. Write $E_1(x) = \theta(x) - x$, with endpoints including primes equal to their argument, and $P_N(s) = N^{1-s}/(s-1)$. Stieltjes summation gives exactly

$$C_Y = E_1(Y)Y^{-s} - E_1(N)N^{-s} + s \int_N^Y E_1(x)x^{-s-1} dx.$$

Consequently $C_Y \rightarrow C := R_b - P_N$, uniformly on $0 \leq t \leq T$ for each fixed T, b . For $Y \geq 4$, the RH estimate on E_1 gives the explicit error

$$|C - C_Y| \ll Y^{-\delta} [\log^2 Y + |s| (\frac{\log^2 Y}{\delta} + \frac{2 \log Y}{\delta^2} + \frac{2}{\delta^3})]. \quad (2)$$

Differentiation adds one logarithm and one possible inverse power of δ . In particular, for $T \geq e$ and $0 < \delta \leq 1/4$, a convenient coarse common bound is

$$|C - C_Y| + |(-C - L^{-1} C'_s) - D_Y| \ll (1+T)\delta^4(1+\log Y)^3 Y^{-\delta}. \quad (3)$$

This justifies the mixed-moment limit at fixed T, b ; it also states its real cost. For example, on $2 \leq b \leq 2G(T)$, $G = \alpha(\log L)$, the choice $Y = \exp(L^3)$ makes (3) at most $O(TL^{13}e^{-L^2})$. This is a legitimate uniform analytic cutoff, not a computationally efficient prime experiment. No finite prime computation is offered as evidence for the limiting remainder.

3. Pole terms are negligible at this scale, with a direct norm estimate

Here and below $2 \leq b \leq 2G(T)$ with $G = \alpha(\log L)$. For all sufficiently large $T, \delta \in [1/L, 1/4]$. The RH local-zero partial fractions yield uniformly

$$\|R_b\|_{L^2(0,T)} + \|K_b\|_{L^2(0,T)} \ll \sqrt{T}L^2. \quad (4)$$

For the derivative in (4), the nearby-zero bound is $O(L/\delta^2) = O(L^3)$, before division by L ; the remote sum is smaller. The absolutely convergent prime-power correction and its derivative are respectively $O(\delta^{-2})$ and $O(\delta^{-3})$. Their contributions to K_b are therefore $O(L^2)$. The finite prime polynomial and its log-weighted companion obey the same uniform mean-value majorant used in Round 8, since $\log p/L \leq 1$ for $p \leq N$. This establishes (4) without differentiating an unknown asymptotic error.

Since $\sigma \leq 3/4$,

$$\|P_N\|_2 \ll \sqrt{N}, \quad -P_N - L^{-1}P'_N = (\frac{\log N}{L} - 1 + \frac{1}{L(s-1)})P_N.$$

Its companion has norm $O((\log L)/L)\sqrt{N}$. Expanding the two mixed products and applying Cauchy – Schwarz gives

$$M_T(b) = \frac{e^b}{TL^2} \Re \int_0^T C(-C - L^{-1}C_s) dt + O(e^b L^{-3}). \quad (5)$$

The error in (5) is $O(b^{-3})$ uniformly on the stated slow range. Thus the pole is accounted for and then removed at a proved error; the endpoint $E_1(N)N^{-s}$ remains part of C .

4. The arithmetic diagonal and the exact missing coefficient

Let $u(x) = \log x/L - 1$ and

$$d\mu_Y(x) = \sum_{N < p \leq Y} (\log p) \delta_p(dx) - \mathbf{1}_{(N, Y]}(x) dx.$$

Set

$$G_{T,b}(x, y) = (xy)^{-1/2-b/(2L)} \frac{u(x) + u(y)}{2} \operatorname{sinc}_0(T \log(xy)).$$

Finite expansion of (1), with the real part symmetrized in x, y , gives

$$\frac{e^b}{TL^2} \Re \int_0^T C_Y \overline{D_Y} dt = \frac{e^b}{L^2} \iint G_{T,b}(x, y) d\mu_Y(x) d\mu_Y(y). \quad (6)$$

The prime-prime diagonal of (6) is

$$D_T(b; Y) = \frac{e^b}{L^2} \sum_{N < p \leq Y} (\log p)^2 p^{-1-b/L} \left(\frac{\log p}{L} - 1 \right).$$

Its limit $D_T(b)$ converges absolutely. Applying the RH estimate for $\theta(x) - x$ by partial summation gives, uniformly in the slow range,

$$D_T(b) = e^b \int_{\log N/L}^{\infty} v(v-1) e^{-bv} dv + O(e^b N^{-1/2} L). \quad (7)$$

The elementary tail integral from $v = 1$ is

$$e^b \int_1^{\infty} v(v-1) e^{-bv} dv = \frac{1}{b^2} + \frac{2}{b^3}.$$

The interval $\log N/L < v < 1$ contributes a nonpositive term of size $O((\log L/L)^2)$, since $1 - \log N/L = 6 \log L/L + O(1/L)$. Both errors in (7), multiplied by b^3 , tend to zero uniformly. Therefore

$$\boxed{D_T(b) = \frac{1}{b^2} + \frac{2}{b^3} + O(b^{-3})}. \quad (8)$$

Define $B_T(b; Y)$ by subtracting this finite prime-prime diagonal from the double integral in (6), and put $B_T(b) = \lim_{Y \rightarrow \infty} B_T(b; Y)$. The limit exists by (2) – (3) and absolute convergence of the diagonal. It contains together the off-diagonal prime-prime sum, both prime-continuum terms, and the continuum-continuum term. Those components must not be assigned independent infinite limits: only their centered combination has been justified.

Equations (5) – (8) prove the arithmetic decomposition

$$M_T(b) = \frac{1}{b^2} + \frac{2}{b^3} + B_T(b) + \alpha(b^{-3}) \quad (9)$$

uniformly on the slow range. Thus, for example, the additional lower bound

$$B_T(s) \geq -\frac{4-\varepsilon}{s^3} - \frac{\eta_T}{s^3}, \quad \eta_T \rightarrow 0, \quad (10)$$

uniformly up to $2G(T)$, would give the Round 10 criterion $M_T(s) \geq s^{-2} - (2-\varepsilon)s^{-3} - \alpha(s^{-3})$. An averaged sufficient version is

$$\int_b^{2b} B_T(s) ds \geq -\frac{3}{2b^2} + \frac{\varepsilon}{b^2} + \alpha(b^{-2}). \quad (11)$$

Indeed the integrated diagonal equals $1/(2b) + 3/(4b^2) + \alpha(b^{-2})$, so (11) yields the required value strictly above $-3/4$ after the coupled normalization. Equations (10) and (11) are unproved arithmetic inputs, not consequences of positivity of the prime coefficients.

5. What the accessible short-prime projection actually sees

A projection confined to prime lengths below T sees only $u(p) \leq 0$. In the residual's available slice $N < p \leq T$, its diagonal log-weighted contribution is $O((\log L/L)^2)$ and nonpositive. It cannot supply the positive leading term $1/b^2$ in (8).

That term comes from logarithmic excess of size $1/b$, hence prime lengths $X = T^{1+O(1/b)}$. The elementary contour/mixed-polynomial estimate used for the short projection has error $O(X \log^3 T)$ at such a cutoff, before normalization; even disregarding any further logarithmic costs, its normalized error is $O(e^b(X/T)L)$. On a fixed edge shell $X = T^{1+s/b}$, $1 \leq s \leq 2$, this majorant contains $\exp(sL/b)$, far larger than b^{-3} . This states the failure of that particular bound, not the magnitude of the actual error or an impossibility theorem.

The inspected higher-derivative theorem leaves an unknown pair-correlation integral, and the elementary short-prime projection misses the required positive-excess sector at the necessary precision. The explicit remaining obligation is the centered arithmetic remainder (10) or its averaged version (11). No generic point-process counterexample, parameter scan, or numerical approximation to a divergent prime series was used.

Current report 44: Independent review of the centered mixed–moment remainder

Collection: R11 — RH small arcs and actual conductor structure.

Source: [research/dyson/round11/log-weighted-tail/INDEPENDENT_REVIEW.md](#).

SHA-256: [97a8e45507e0df1cf21e6045ba6d689990401a4eff5f2ba79a00739ab69e08f6](#) . Git blob: [031535f861ac31d697fd45128e263acf313b16fd](#) . Original bytes: 4187.

Independent review of the centered mixed–moment remainder

Date: 2026-09-05. Reviewer: root Astra, independently of the authoring lane. The centered arithmetic identity and its error bounds are accepted under RH. Neither sufficient lower bound for the remainder is proved.

Reviewed author SHA256: [85c9435bbd656bcf37f21f55fe2d54cb40638c7a56908abe57e54f551a627daa](#) .

Source and continuation

I checked Proposition 5, equation (2.3), in the retained primary HTML of Chirre, arXiv:2107.13636v2. The statement is for fixed integer $k \geq 1$ and fixed positive a in the displayed asymptotic. It includes the unknown integral of $\alpha^{2k} \exp(-2a\alpha) F(\alpha, T)$ beyond $\alpha=1$. The note correctly declines to differentiate its unspecified remainder or use it as a uniform growing-width estimate. Theorem 1's equivalence with pair correlation does not furnish the required one-sided estimate under RH alone.

Stieltjes summation with $\theta(x)-x$ and primes $N < p \leq Y$ gives both endpoint terms with the written signs. Its tail equals minus $E_1(Y)Y^{-s}$ plus s times the remaining integral. RH bounds it by $Y^{-(\delta)}$ times the displayed polynomial in $\log Y$ and inverse δ . Differentiation in s adds one logarithm, with the next inverse- δ power. This is a legitimate fixed- T limit, unlike termwise use of an uncentered prime series in the critical strip.

For $b \geq 2$, $\delta \geq 1/\log T$. With $Y = \exp((\log T)^3)$, the coarse uniform error is bounded by $T(\log T)^{13} \exp(-(\log T)^2)$, up to a fixed constant. This validates a mathematical cutoff but supplies no feasible brute-force experiment.

Norms and the pole

The local-zero derivative has majorant $O(\log T/\delta^2)$ before division by $\log T$. The prime-power correction and its derivative have the stated inverse- δ bounds; the finite polynomial uses weights bounded by the original ones. These give the sufficient coarse L_2 bound $\sqrt{T}(\log T)^2$ for both factors, uniformly on the slow range.

The pole is $N^{(1-s)}/(s-1)$. Since $\sigma \leq 3/4$ its L_2 norm is $O(\sqrt{N})$, including the low-height part. Direct differentiation gives the coefficient $\log N/\log T - 1/((\log T)(s-1))$, so its companion is $O((\log \log T)/(\log T))\sqrt{N}$. Cauchy–Schwarz after normalization gives $O(\exp(b)(\log T)^{-3})$. It is $o(b^{-3})$ when $b \leq 2G(T)$, $G = o(\log \log T)$. The endpoint involving $E_1(N)$ is retained throughout.

Diagonal and the combined remainder

Expanding the two finite centered measures, integrating time, taking real parts, and exchanging x, y gives the symmetric factor $(u(x)+u(y))/2$ times $\text{sinc}_0(T \log(x/y))$. On the prime diagonal this reduces to $u(p)$, with exactly two factors $\log p$. The continuous measure has no atomic diagonal mass.

For partial summation write $f(x)=(\log x)(\log x/\log T-1)x^{(-1-b/\log T)}$. The RH error $\theta(x)-x$ is $O(\sqrt{x}\log^2 x)$. Its endpoint plus the integral against f' is $O(N^{(-1/2)}(\log T)^3)$, uniformly for the stated b ; multiplying by $\exp(b)/(\log T)^2$ gives the author's conservative $O(\exp(b)N^{(-1/2)}\log T)$. No unproved short-interval estimate is used here.

The main integral from $v=1$ is exactly $b^{(-2)}+2b^{(-3)}$. The omitted strip $\log N/\log T < v < 1$ has width $O(\log \log T/\log T)$, negative integrand of size proportional to that width, and hence a squared-width error. Multiplication by b^3 makes both errors vanish uniformly. Thus the positive diagonal is correct, and the short-prime slice below T cannot provide its leading term.

The limit of the full centered expression exists; its diagonal converges absolutely. Subtracting the latter defines the combined remainder. No separate limits of the prime-continuum pieces are asserted. This distinction is essential and is preserved in the report.

Integrating the diagonal over $[b, 2b]$ gives $1/(2b)+3/(4b^2)$. The proposed integrated remainder threshold then leaves a coefficient strictly above $-3/4$ after normalization. This checks the constants, but supplies no proof of that threshold. The reported long-polynomial error is a failed upper majorant, not a lower bound on the actual error.

The result identifies a specific arithmetic obligation and rules out an unjustified use of a fixed-width derivative theorem. It does not establish a new zeta pair-correlation bound.

Current report 45: Round 12: three exact tests of the remaining arithmetic gap

Collection: R12 — exact limits of sampling and dispersion transfers.

Source: [research/reports/dyson_round12.md](#).

SHA-256: [5e697b2ba41e12869bb471a1294d110b1dbf8d2592e9652f467c776d6272a89a](#) . Git blob: [84cf3fbc0f42b330021ccf98d8e744010272557b](#) . Original bytes: 8922.

Round 12: three exact tests of the remaining arithmetic gap

Date: 2026-09-05. **No stronger actual-prime estimate was obtained in this round.** The current bound remains the RH component estimate $X^{1.023} \log^5 X$ from Round 11. Three bounded attempts now explain precisely why improving positive sampling, importing the 186 dispersion theorem by phase absorption, or applying known prime-interval upper norms does not yet close its gap. Each conclusion is narrower than an impossibility theorem for the actual signed prime pairing.

Positive sampling really is crowded on the permitted support

The [sampling proof](#) uses the actual canonical complementary-modulus family, not the full Farey set as a substitute. Round 11 constructed at least $c Q / \log^{348} X$ terminal conductors in $(Q/2, Q]$, each with merged coefficient $1/d$. They supply at least $c Q^2 / (H \log^{348} X)$ distinct reduced frequencies in $[0, 1/(16H)]$, with coefficient magnitude at least $c_V H/Q$.

Partition this arc into $O(X/H)$ cells of length $1/(100X)$. One cell must contain at least $c Q^2 / (X \log^{348} X)$ actual frequencies. A phase-tuned Dirichlet packet with integer frequencies in $[X, 1.1X]$ concentrates there. Parseval and its value on that cell prove that the positive local sampling constant is at least

$$c Q^2 / (\log X)^{348}.$$

The packet also satisfies the same known small-arc energy and derivative envelopes, even without their logarithmic factors. Including the actual squared coefficient weights gives the corresponding lower bound $c_V H^2 / \log^{348} X$. Thus these hypotheses cannot justify a fixed-power improvement of that positive sampling step.

The packet is an artificial integer polynomial, not the centered genuine-prime polynomial for fixed smooth f . It is not asserted to align with the full complex coefficient vector. The actual signed functional has the exact dual kernel

$$K(n) = \sum_{q \in Q} \mu(q) \left[\sum_{h \equiv n(q)} \nu(h/H) - \frac{1}{\phi(q)} \sum_{(h,q)=1} \nu(h/H) \right].$$

Its Gram matrix contains signed off-diagonal contributions. A smaller norm for that functional, or cancellation with actual prime coefficients, is not ruled out by the positive sampling result. The proof specifically does not claim that $X^{.023}$ is unavoidable for primes. The inherited counting and packet constants have a [separate narrow review](#); root

checked the complete argument and exact signed identity.

Direct phase absorption violates the dispersion coefficient hypothesis

The [dispersion audit](#) checks the 186 paper's actual hypotheses at legal parameters. With $\omega=.012$, $\delta=.001$ and $\sigma=.101$, the three bilinear inequalities have left sides .888, .996 and .990. The scales $N=X^{.4}$ and $M=X^{.6}$ are permitted, as are the canonical terminal conductors d near $Q=X^{.523}$.

The prime-interval sequence $\beta(n)=1_{(n \text{ prime}, N \leq n < 2N)}$ has the source's Siegel--Walfisz property. But its additive twist need not. For d coprime to 3, take $k=(d-1)/3$ or $(d+1)/3$ according to $d \bmod 3$, and choose a unit m in $[M, 2M]$ with $m=k \bmod d$. This is possible since M/d tends to infinity. The actual completed numerator $\alpha=1$ has nonzero shift weight. On the prime interval,

$$\alpha(mn/d) = \alpha(n/3)(1 + O(N/d)).$$

PNT in the two fixed reduced classes modulo 3 therefore gives

$$\Delta(\beta(n)\alpha(mn/d); 1 \bmod 3) = \left(\frac{i\sqrt{3}}{4} + \alpha(1)\right) \frac{N}{\log N}.$$

The source SW requirement with logarithmic exponent two cannot accommodate this discrepancy. Both branches for k give the same leading sign. The error $O(N^2/(d \log N))$ is uniformly lower order. Hence one cannot absorb the completed phase into the short coefficient and inherit its original SW property, even with a permitted conductor, scale and primitive numerator. The source theorem remains valid; the proposed transformed coefficient fails its premise.

This does not say that every slice or every factor in a prime identity fails, or that bad slices cannot be handled after averaging. An averaged argument keeping m , α , d and h may still succeed. The modulus-dependent coefficient also cannot be silently substituted for a fixed family before the source's modulus sum.

Two further checks close related direct substitutions. First, even after $\gcd(h,d)=1$ is imposed, a positive subinterval of the shifts maps onto every unit class modulo every prime factor of the constructed d . Its product of local images has $\phi(d)$ classes, asymptotic to d , while $\tau(d)=2^{348}$ is fixed. The source's bounded-local-class lift cannot absorb that cost. The original coherent interval has only $O(H)$ global classes; a method preserving this cross-prime correlation is not excluded. Second, using H itself as the short convolution length fails the source range: that factor must exceed $X^{.398}$, whereas $H \leq X^{(2/7)}$.

Centered prime–interval upper bounds miss the sign and precision

The [Selberg audit](#) reads [Saffari--Vaughan, Lemma 5](#), including printed page 20. Under RH its local estimate concerns genuine θ and is uniform for $0 < \eta \leq 1$:

$$\int_X^{2X} |\theta(x + \eta x) - \theta(x) - \eta x|^2 dx \ll \eta X^2 \log^2(2/\eta).$$

Applying Mellin Gallagher to the finite centered prime measure, retaining both cutoff crossings and the logarithmic weight, gives the valid but insufficient bound $|M_T(b)| \leq C \log(T)/b^2$. The associated integrated lower bound on the combined remainder is only of order $-\log(T)/b$. It is much too weak for the required b^{-2} correction. The prime

and continuum pieces are never separated into unjustified infinite sums.

This is deliberately not claimed as the strongest RH consequence. Round 10 already had stronger individual norm control, and the same source page records Selberg's stronger global weighted estimate with a fixed exponent range. For a smooth filter on the active excess shell, the latter gives normalized squared norms $O(1)$ and $O(b^{-2})$, hence only $O(b^{-1})$ control of the mixed product by Cauchy--Schwarz. It gives neither a positive increment in the shrinking shell nor the necessary signed next-order coefficient. A stronger use of the theorem is not ruled out.

Smoothing b leaves the actual sinc factor and joint centering in the kernel. A negative sinc lobe prevents a term-by-term nonnegativity argument. A negative kernel value alone would not disprove positive semidefiniteness, and the report explicitly avoids that inference. No generic point-process countermodel replaces actual prime arithmetic in this calculation.

Review, replay and the next useful attempt

The [complete independent root review](#) pins all three author hashes, checks the primary source statements, and records the exact accepted scope. All 18 originals (1,904,996 bytes) are preserved locally under `Astra-Local-Archive/round12-originals/`; 15 research files are public verbatim. The third-party PDF, extracted text and rendered page stay local with receipts.

The [intake manifest](#) and [bounded replay](#) verify the source bytes and two exact scripts. The sampling output matches in full. The dispersion certificate matches after removing only four temporary provenance paths; all reference hashes remain checked. The scripts verify rational constants, 60 cyclotomic signed-kernel identities and six fixed modular-selection examples. They do not numerically test PNT, the huge conductor construction, or a conjecture.

```
python3 research/logs/round12-integration/recheck.py --prime-gap-source-dir /path/to/pinned/186/source-directory
python3 tools/verify_manifest.py
```

The next work retains the structures that these failed transfers discarded: the signed residue kernel, the coherent shift interval, and the additive phase through the m average. In particular the modulus-3 example suggests isolating rationally resonant m values before estimating the remaining phases, rather than assuming a uniform property that is false. A separate attempt is to estimate the exact signed dual norm through common-divisor compatibility and its genuine CRT boundary error. These are ongoing investigations, not new claims in this checkpoint.

The long PDFs retain their earlier checkpoints. The manual single-session Fable packet receives only the coordinator's superseding source-status prefix; it does not dispatch a new session or ask for covered computations again. No large scan, new model service, infrastructure layer or conjecture solution was introduced. Reverting this slice removes the new records without altering earlier proofs. Formalization and the required actual-zeta lower bound remain outstanding.

Current report 46: The actual support saturates positive sampling at the power scale

Collection: R12 — exact limits of sampling and dispersion transfers.

Source: [research/dyson/round12/sampling-geometry/ACTUAL_SUPPORT_SAMPLING_OBSTRUCTION.md](https://research.dyson/round12/sampling-geometry/ACTUAL_SUPPORT_SAMPLING_OBSTRUCTION.md).

SHA-256: `cb52d72f6068c3030968209d8aa028439ea4dc309aa5584d216a1d7d30a1a59d` . Git blob: `9336b996cf951f72f54edfa3e6a56fc7ffef8424` . Original bytes: 11979.

The actual support saturates positive sampling at the power scale

Date: 2026-09-05. Status: ordinary geometric proof; the inherited support count and packet constants have been independently checked. The complete proof is awaiting the coordinator's final review. No improved actual-prime bound is proved in this task.

The Q^2 term in the Round 11 **positive local sampling inequality** cannot be reduced by a power of X uniformly on the full canonical complementary-modulus support. This remains true for trigonometric polynomials supported on the actual integer-frequency interval n near X , and after including the absolute squares of the actual coefficients $S_v(a/d)M_d$. An explicit packet obeying the same small-arc norm and derivative bounds makes this precise.

This is **not an obstruction for the genuine-prime polynomial or for the full signed pairing**. The sampling extremizer constructed here is an artificial band polynomial, not a prime sum. The exact signed Gram formula in Section 5 retains a possible route for improvement; it cannot be rejected using the positive sampling lower bound. Accordingly the result closes one proposed geometric shortcut, not every possible use of frequency geometry.

1. Actual arithmetic support already established in Round 11

Write

$$Q=X^{(523/1000)}, \quad X^{(1/6)} \leq H \leq X^{(2/7)}.$$

Fix the full canonical squarefree modulus family Q_X^{full} defined by the balanced complementary predicates of Round 9. Fix a nonnegative nonzero v in $C_c^\infty(1,2)$, and let $m_v = \int v > 0$. The actual completed coefficients are

$$\begin{aligned} M_d &= \sum_{(q \text{ in } Q_X^{\text{full}}, d|q)} \mu(q)/q, \\ C_{(a/d)} &= S_v(a/d) M_d, \\ S_v(\beta) &= \sum_h v(h/H) e(-\beta h), \quad e(t) = \exp(2\pi i t). \end{aligned}$$

The frozen Round 11 conductor construction proves the following facts, including their quantifiers. There is a subfamily F_X of admissible moduli with

$$\begin{aligned} F_X &\subset (Q/2, Q], \\ |F_X| &\geq c_0 Q / [2(\log X)^{348}] \end{aligned} \quad (1)$$

for all sufficiently large real X . Here c_0 is the explicit fixed positive constant in that report. Every d in F_X has 348 distinct prime factors, $\mu(d)=1$, and

$$M_d = 1/d \quad (2)$$

in the **full signed family**, because no larger multiple of d fits below Q . The construction checks the actual root bounds and both large-prime predicates; it is not an unconstrained Farey model.

For each such d there are at least $d/(32H)$ integers

$$1 \leq a \leq d/(16H), \quad \gcd(a, d) = 1,$$

and on all those actual reduced frequencies

$$\begin{aligned} |S_v(a/d)| &\geq m_v H/(2 \sqrt{2}), \\ |C_v(a/d)| &\geq m_v H/(2 \sqrt{2} Q). \end{aligned} \quad (3)$$

These statements hold simultaneously and uniformly in the stated H range once X is sufficiently large. Their proof uses only the ordinary prime number theorem on fixed-ratio intervals and the verified complementary predicates. No estimate about short-interval primes or actual prime exponential sums is added here.

Let Ω_X be exactly the reduced fractions just described. Distinct pairs (a, d) give distinct frequencies, by coprimality. Equations (1) – (3) yield

$$\begin{aligned} \Omega_X &\subset (0, 1/(16H)], \\ |\Omega_X| &\geq c_0 Q^2 / [128 H (\log X)^{348}]. \end{aligned} \quad (4)$$

Thus this very small arc contains many actual frequencies with actual coefficient size of order H/Q , up to fixed constants.

2. A microscopic cluster follows without a local distribution theorem

Partition $[0, 1/(16H)]$ into intervals of length at most $1/(100X)$, using at most

$$\text{ceil}(100X/(16H)) \leq 8X/H$$

intervals, for all sufficiently large X . By the pigeonhole principle some interval J_X contains K_X actual reduced frequencies, with

$$K_X \geq c_0 Q^2 / [1024 X (\log X)^{348}]. \quad (5)$$

No equidistribution of the moduli or numerators inside subintervals is assumed. The one dense cell is forced by the already proved total count. Its occupation tends to infinity, since

$$Q^{2/X} = X^{(23/500)}$$

dominates every fixed power of $\log X$. This lower bound is uniform in H ; H cancels in the average occupation.

Let β_X be the midpoint of J_X . Every selected frequency in this cluster has distance at most $1/(200X)$ from β_X . It still has a genuine reduced denominator in F_X near Q . We have not assigned new artificial locations to any frequency.

3. Explicit band polynomial and a lower bound for the sampling constant

Set $N=\text{ceil}(X)$, $M=\text{floor}(X/10)$, and define

$$P_X(\beta) = H^{(-1/2)} \sum_{n=N}^{N+M-1} e(n(\beta - \beta_X)). \quad (6)$$

For X sufficiently large, $M \geq X/20$, and all its integer Fourier frequencies lie in $[X, 11X/10]$. Its coefficient magnitudes are exactly $H^{(-1/2)}$, in particular at most one. The carrier is in the same positive-frequency range as the prime polynomial; this is not a test function allowed an arbitrarily high bandwidth.

For β in J_X factor out the unit-modulus carrier $e(N(\beta - \beta_X))$. Each remaining phase has magnitude at most $\pi/1000$. Hence the real part of their sum is at least $M/2$, and

$$|P_X(\beta)| \geq M/(2 \sqrt{H}) \geq X/(40 \sqrt{H}). \quad (7)$$

Parseval gives the exact norm

$$\int_0^1 |P_X(\beta)|^2 d\beta = M/H. \quad (8)$$

Combining (5), (7) and (8),

$$\begin{aligned} & \frac{\sum_{\beta \in \Omega_X} |P_X(\beta)|^2}{\int_0^1 |P_X(\beta)|^2 d\beta} \\ & \geq K_X M/4 \\ & \geq c_0 Q^{2/[81920(\log X)^{348}]}. \end{aligned} \quad (9)$$

Replacing the denominator integral by its restriction to $|\beta| \leq 1/H$ only makes the ratio larger. The numerator could also be restricted to J_X . Therefore any positive local sampling estimate valid for every polynomial in this positive-frequency band must have sampling constant at least Q^2 times a fixed negative logarithmic power on the actual canonical support.

In particular no uniform bound of the form

$$\begin{aligned} & \sum_{\text{actual frequencies in the central arc}} |P(\beta)|^2 \\ & \leq O(Q^2 X^{(-\epsilon)}) \int_{\text{local arc}} |P(\beta)|^2 \end{aligned}$$

is possible for any fixed $\eta > 0$ in this class. Nor can the Q^2 term be replaced by X times any fixed logarithmic power: the ratio Q^2/X is the positive power $X^{.046}$. Allowing a smooth nonnegative arc cutoff which equals one on $[0, 1/(16H)]$ does not change this lower bound.

An analogous lower bound holds with the actual squared coefficient weights retained. By (3) and (9),

$$\begin{aligned} & \sum_{\text{all actual } \beta} |C_{\beta}|^2 |P_X(\beta)|^2 \\ & / \int_0^1 |P_X(\beta)|^2 d\beta \\ & \geq c_0 m_v^2 H^2 / [655360 (\log X)^{348}]. \end{aligned} \quad (10)$$

All omitted terms here are nonnegative. Thus cancellation between different signed moduli cannot remove this weighted quadratic lower bound: on the selected conductors their merged coefficient is exactly $1/d$, and elsewhere the quadratic sum only increases. In the natural top-conductor normalization $(Q/H)^2 |C_{\beta}|^2$, equation (10) again gives a sampling constant at least a fixed multiple of $Q^2/(\log X)^{348}$. It concerns that positive weighted sampling operator; it does not assert that every possible reweighting or Cauchy-Schwarz arrangement is sharp, and it does not concern the signed linear pairing.

4. The packet even obeys the known small-arc envelopes

For every $\rho \geq 1/H$, nonnegativity and (8) show

$$\begin{aligned} & \int_{||\beta|| \leq \rho} |P_X(\beta)|^2 d\beta \\ & \leq M/H \ll X \rho. \end{aligned}$$

Differentiating (6) coefficient by coefficient and applying Parseval gives

$$\begin{aligned} & \int_0^1 |P_X'(\beta)|^2 d\beta \\ & = 4 \pi^2 H^{(-1)} \sum_{n=N}^{\infty} (N+M-1) n^2 \\ & \ll X^3/H. \end{aligned}$$

Therefore, simultaneously for every such ρ ,

$$\begin{aligned} & \int_{||\beta|| \leq \rho} |P_X'(\beta)|^2 d\beta \\ & \ll X^3 \rho. \end{aligned} \quad (11)$$

These are stronger than the $\log^4$ versions supplied by the RH centered-prime input in Round 11. On the other hand (5) and (7) give

$$\begin{aligned} & \sum_{\beta \in \Omega_X} |P_X(\beta)|^2 \\ & \gg X Q^2 / [H (\log X)^{348}]. \end{aligned} \quad (12)$$

Consequently the already known norm and derivative envelopes, the positive integer-frequency band, and the actual dense-divisibility support together do not force a power improvement of the sampled energy. The location of the packet depends on the actual frequency cluster, as a worst-case operator test may.

The crucial limitation is equally explicit: P_X is **not** the centered genuine-prime polynomial E_f for a fixed smooth f . Its coefficients are phase-tuned across all integers. Equations (9) – (12) say nothing about whether E_f concentrates at this cell. A new bound excluding such concentration specifically for the prime coefficients, or a bound for the

signed coefficient pairing, remains possible.

5. What signed weights could still accomplish

The positive sampling step bounds a sum of absolute squares. The actual completed expression is instead

$$\begin{aligned} L(F) &= \sum_{\beta} C_{\beta} [F(\beta) - r_{\beta} F(0)], \\ r_{\beta}(a/d) &= \mu(d)/\phi(d). \end{aligned} \quad (13)$$

For $F(\beta) = \sum_{n \in I} b_n e(n\beta)$, its exact dual coefficient sequence is

$$\begin{aligned} K(n) &= \sum_{\beta} C_{\beta} [e(n\beta) - r_{\beta}], \\ L(F) &= \sum_{n \in I} b_n K(n). \end{aligned} \quad (14)$$

Thus the squared norm of this signed functional on the coefficient space is exactly

$$\begin{aligned} \sum_{n \in I} |K(n)|^2 &= \sum_{\beta, \gamma} C_{\beta} \overline{C_{\gamma}} \sum_{n \in I} [e(n\beta) - r_{\beta}] \\ &\quad [e(-n\gamma) - r_{\gamma}]. \end{aligned} \quad (15)$$

The off-diagonal terms in (15) have signs and phases. Nothing in the lower bound for the diagonal positive sampling operator implies a lower bound for (15), or for its value on the actual prime coefficients. In particular, the packet in (6) is not asserted to align with the full coefficient vector C_{β} .

Using the exact completion from Round 10, K has the arithmetic form

$$\begin{aligned} K(n) &= \sum_{q \in Q_X^{\text{full}}} \mu(q) \\ &\quad \left[\sum_{h \equiv n \pmod{q}} v(h/H) \right. \\ &\quad \left. - \phi(q)^{-1} \sum_{(h,q)=1} v(h/H) \right]. \end{aligned} \quad (16)$$

This finite identity holds coefficient by coefficient. When n is a genuine prime near X it is a unit modulo every q , so (16) is exactly the corresponding residue-discrepancy kernel. For arbitrary n , (16) is the Fourier-defined kernel in (14); it need not be identified with the original progression discrepancy with its additional nonunit restrictions.

For the actual prime pairing the coefficients are $b_p = (\log p)f(p/X)$ on primes and zero elsewhere. The primitive subtraction in (13) – (16) is retained. An improved estimate exploiting the full signed Gram (15), or cancellation between the prime coefficients and (16), could therefore improve Round 11. Such an estimate would analyze actual correlations of the permitted moduli, their varying residues and/or primes. It is not supplied by a larger minimum spacing, fewer frequencies in the central arc, or a smaller positive sampling norm.

6. Bounded decision and quantifier limits

No sharper upper bound for the actual prime pairing was obtained in this task. What is now ruled out is precise: a **uniform positive sampling or absolute-weight sampling improvement by a power**, based only on the Round 11 arc envelopes and canonical factorization support. The Q^2 sampling term is sharp up to logarithmic powers in that class, including on the correct integer-frequency band.

This does not rule out a deliberately pruned modulus family, a specially chosen signed shift profile, a direct signed-Gram argument, an improved prime-specific concentration theorem, or cancellation in the full genuine-prime pairing. In particular, it is not a proof that the remaining $X^{.023}$ in the actual arithmetic bound is unavoidable. Treating a worst-case artificial polynomial as the prime polynomial would be a logical error.

The small companion script checks exact exponent differences, the constants derived from the Round 11 count, and the finite completion identity on one fixed toy modulus family using formal cyclic Fourier arithmetic. That toy identity checks the placement of the primitive term, not the large- X support construction or a prime estimate. The actual support construction is an ordinary asymptotic proof already pinned in the source receipt. No parameter search or large computation was performed.

Current report 47: Independent review of the inherited count and sampling constants

Collection: R12 — exact limits of sampling and dispersion transfers.

Source: `research/dyson/round12/sampling-geometry/COUNTING_REVIEW.md`.

SHA-256: `820dfdc6707cdc1c29f7eef61dc884d3f6c96980ffcfab661cd790e3843bbc2` . **Git blob:** `9311aa64f31b7d0d07a87706c442941362346dc0` . **Original bytes:** 2258.

Independent review of the inherited count and sampling constants

Date: 2026-09-05. **Reviewer:** the conductor-arithmetic/dispersion-transfer agent. **Scope:** Sections 1 – 3 of `ACTUAL_SUPPORT_SAMPLING_OBSTRUCTION.md` only.

Reviewed report SHA-256: `cb52d72f6068c3030968209d8aa028439ea4dc309aa5584d216a1d7d30a1a59d`. This hash was recomputed from the frozen file. The inherited Round 11 conductor report was also checked at its pinned SHA-256 `46347799005bb0f53af25c2a7e8ffb2b2217d92688c7651327dde3562f114b92`.

Accepted within the stated scope. The following calculations were checked independently against the actual arithmetic support proof.

- The count $c_0 Q / (2 \log^{348} X)$, denominator bound $d > Q/2$, and $d/(32H)$ primitive numerators give at least $c_0 Q^2 / (128H \log^{348} X)$ distinct reduced frequencies. No duplicate rational frequencies or permutation counts enter.
- Partition into at most $8X/H$ intervals gives a cell containing at least $c_0 Q^2 / (1024X \log^{348} X)$ frequencies. This only uses the proved total count; no distribution theorem within short intervals is assumed.
- With $N = \lceil X \rceil$, $M = \lfloor X/10 \rfloor$, the integer support really is inside $[X, 1.1X]$. Factoring out the carrier on a cell of width at most $1/(100X)$ leaves phase spread at most $\pi/1000$.
- The lower pointwise bound $M/(2\sqrt{H})$, Parseval norm M/H , and $M \geq X/20$ give the stated sampling ratio $c_0 Q^2 / (81920 \log^{348} X)$.
- The actual coefficient lower bound $|C_\beta|^2 \geq m_v^2 H^2 / (8Q^2)$ then gives precisely the weighted denominator 655360 in equation (10).

The constants and their directions are correct for all sufficiently large real X , uniformly in the stated H -range. The asymptotic threshold is not claimed effective or numerically reached.

This review does not turn the artificial sampling packet into a prime polynomial, and does not establish a lower bound for the full signed pairing. The final wording correctly keeps the positive unweighted and fixed absolute-coefficient sampling statements separate from other possible reweightings or signed arguments. Sections 4 – 6 and the finite script are outside this narrow review; the coordinator conducts the wider review.

Current report 48: Two explicit hypothesis failures in the proposed dispersion transfer

Collection: R12 — exact limits of sampling and dispersion transfers.

Source: [research/dyson/round12/dispersion-transfer/DISPERSION_HYPOTHESIS_OBSTRUCTION.md](#).

SHA-256: [26536af7b5e6ebfb8fb4f7c7be993543c57152e29c6fd2b039c5eb147ccef9d95](#) . Git blob: [f94a750227ddad82bfc5b86b3f5c7ea171330072](#) . Original bytes: 18370.

Two explicit hypothesis failures in the proposed dispersion transfer

Date: 2026-09-05. Status: ordinary arithmetic proofs, with exact source-parameter checks. No improved bound for the actual prime pairing is proved here. The counterexample below concerns preservation of a coefficient hypothesis; it is not a counterexample to the source dispersion theorem.

The 186 paper's triply densely divisible dispersion theorem does not directly give the missing improvement in the completed Round 11 pairing. Two specific proposed applications fail:

1. Absorbing a completed additive phase into a Siegel – Walfisz coefficient does not preserve that property. This fails for a prime-interval coefficient at legal source scales, with a conductor in our actual canonical complementary family and the actual reduced numerator $\mathfrak{a} = 1$. The resulting discrepancy at modulus 3 is explicitly of order $N/\log N$.
2. Lifting the growing shift interval through the source's product-of-local-residue-sets lemma requires $\phi(d)$ classes on an explicit subfamily. Its bounded-local-class hypothesis is therefore unavailable, even after all primitive restrictions are imposed.

These statements leave open a dispersion argument retaining the shift, the additive phase and the prime coefficient jointly. They do not prove the remaining $X^{0.23}$ loss is unavoidable for primes. The source is [OpenAI, Improved short gaps between primes](#); precise printed-page locations and file hashes are recorded below and in the companion certificate.

1. The exact pairing and the family being tested

Put

$$Q = X^{523/1000}, \quad X^{1/6} \leq H \leq X^{2/7}, \quad \mathfrak{a}(t) = \exp(2\pi it). \quad (1)$$

Fix a nonnegative, nonzero $V \in C_c^\infty(1, 2)$. This hypothesis permits the real-prime support construction used here; no claim uniform over arbitrary signed profiles is made.

The canonical family Q_X contains every distinct squarefree modulus $q = [D, E]$ satisfying

$$\begin{aligned} D, E &\leq X^{523/2000}, & q &> X^{1/2}, \\ p^{3/2} D_{\geq p} &\leq X^{501/2000} & (p \mid D, p > X^{1/1000}), \\ p^{3/2} E_{\geq p} &\leq X^{501/2000} & (p \mid E, p > X^{1/1000}). \end{aligned} \quad (2)$$

Each modulus is counted once with coefficient $\mu(q)$. These are the balanced complementary predicates fixed in Round 9. Source Proposition 2.3, printed pp.4 – 5, gives

$$\mathbb{Q}_X \subset \{q \leq Q : q \text{ squarefree}, q \in D^{(3)}(X^{1/1000})\}. \quad (3)$$

There is no arbitrary-subset Möbius-cancellation assertion.

For a separated fixed smooth prime weight f , the exact completed expression is

$$B_j(f, V) = \sum_{\substack{2 \leq d \leq Q \\ 1 \leq a < d, (a, d)=1}} S_{V, H}(a/d) M_d^{(j)} \left[A_f(a/d) - \frac{\mu(d)}{\phi(d)} A_f(0) \right], \quad (4)$$

where

$$\begin{aligned} S_{V, H}(\beta) &= \sum_h V(h/H) e(-\beta h), \\ M_d^{(j)} &= \sum_{\substack{q \in \mathbb{Q}_X \\ d \mid q}} \frac{\mu(q)(\log q)^j}{q}, & j &= 0, 1, \\ A_f(\beta) &= \sum_{p \text{ prime}} (\log p) f(p/X) e(\beta p). \end{aligned} \quad (5)$$

The Ramanujan principal subtraction in (4) is retained. The zero frequency cancels exactly. The actual logarithmic kernel is assembled from B_0, B_1 by the already proved uniformly summable smooth separation; the term $\log X B_0$ must also be counted. This report tests an attempted arithmetic input to (4), not a replacement positive norm.

Under RH, the frozen Round 11 bound for the complete fixed smooth discrepancy component is

$$O_{V, X}(\sqrt{X(X+Q^2)}(\log X)^5) = O_{V, X}(X^{1023/1000}(\log X)^5). \quad (6)$$

Its remaining power loss comes from $Q/\sqrt{X} = X^{23/1000}$. No step below improves (6).

2. Actual terminal conductors and source-compatible scales

The frozen Round 11 construction uses

$$u = \frac{9}{100}, \quad \kappa = \frac{343}{346000}, \quad \lambda = 2^{-1/348}.$$

Let F_X consist of products of two distinct primes in $(\lambda X^u, X^u]$ and 346 distinct primes in $(\lambda X^\kappa, X^\kappa]$. For every sufficiently large real X , that construction proves

$$F_X \subset \mathbb{Q}_X \cap (Q/2, Q], \quad |F_X| \sim c_0 \frac{Q}{(\log X)^{348}}, \quad c_0 > 0. \quad (7)$$

The exact constant is

$$C_0 = \frac{(1-\lambda)^{348}}{2! 346! u^2 \kappa^{346}}.$$

Unique factorization prevents permutation overcounting. The source guards are checked by splitting one large and 173 small primes into each root:

$$2u + 346\kappa = \frac{523}{1000}, \quad u + 173\kappa = \frac{523}{2000}, \quad \frac{5}{2}u = \frac{9}{40} < \frac{501}{2000}. \quad (8)$$

Thus these are actual complementary moduli, with two factors larger than the density parameter. Their existence and count use only the prime number theorem on fixed-ratio intervals.

Every $d \in F_X$ has exactly 348 prime factors, all tending to infinity; in particular $3 \nmid d$ eventually. Also $\mu(d) = 1$ and

$$M_d^{(0)} = \frac{1}{d}, \quad M_d^{(1)} = \frac{\log d}{d}, \quad (9)$$

because the only multiple of $d > Q/2$ at most Q is d itself. Signed lower moduli cannot cancel these terminal coefficients.

Source Proposition 2.18, printed pp.10 – 11, applies to a convolution $\alpha * \beta$ at scales $MN \asymp X$, with

$$X^{1/2-\sigma} \leq N \leq X^{1/2}, \quad (10)$$

provided β has the source Siegel – Walfisz property and

$$72\omega + 24\delta < 1, \quad 48\omega + 16\delta + 4\sigma < 1, \quad 64\omega + 20\delta + 2\sigma < 1. \quad (11)$$

Use the actual parameters

$$\omega = \frac{3}{250}, \quad \delta = \frac{1}{1000}, \quad \varepsilon = \frac{1}{1000}, \quad \sigma = \frac{101}{1000}. \quad (12)$$

The modulus cutoff $X^{1/2+2\omega-\varepsilon}$ is exactly Q . The three left sides of (11) are respectively

$$\frac{111}{125} = .888, \quad \frac{249}{250} = .996, \quad \frac{99}{100} = .990. \quad (13)$$

Take

$$N = X^{2/5}, \quad M = X^{3/5}. \quad (14)$$

Then $MN = X$, and $1/2 - \sigma = .399 < .4 < .5$. The counterexample below therefore does not manufacture an out-of-range factor or an inadmissible modulus.

3. A real–prime coefficient loses Siegel–Walfisz after phase absorption

The source's Definition 2.9, printed p.6, requires one fixed C_{SW} such that for every fixed $L > 0$

$$\left| \sum_{\substack{n \equiv a \pmod{r} \\ (n,s)=1}} \beta(n) - \frac{1}{\phi(r)} \sum_{(n,rs)=1} \beta(n) \right| \ll_L \tau(rs)^{C_{\text{SW}}} N(\log X)^{-L}, \quad (15)$$

uniformly in r, s, a with $(a, r) = 1$. Here r denotes the test modulus; it is not the conductor d .

By source Proposition 2.10, printed p.7,

$$\beta_X(n) = 1_{\{n \text{ prime}, N \leq n < 2N\}} \quad (16)$$

has this property, uniformly in the interval endpoints and auxiliary coprimality parameter, because $N = X^{2/5}$. It is a coefficient sequence located at scale N .

Lemma. For every sufficiently large real X , and every $d \in F_X$, there is a unit $m \bmod d$ represented by an integer in $[M, 2M]$ such that the sequence

$$\tilde{\beta}_{X,d,m}(n) = \beta_X(n) \mathfrak{e}(mn/d) \quad (17)$$

does not satisfy (15) as a uniform family. At $r = 3, s = 1, a = 1$, its discrepancy is

$$\Delta(\tilde{\beta}; 1 \bmod 3) = \left(\frac{i\sqrt{3}}{4} + o(1) \right) \frac{N}{\log N}. \quad (18)$$

The $o(1)$ is uniform in the conductors d and choices of m made below. No RH assumption is used.

Proof. Define

$$k = \begin{cases} (d-1)/3, & d \equiv 1 \pmod{3}, \\ (d+1)/3, & d \equiv 2 \pmod{3}. \end{cases} \quad (19)$$

Thus

$$\frac{k}{d} = \begin{cases} \frac{1}{3} - \frac{1}{3d}, & d \equiv 1 \pmod{3}, \\ \frac{1}{3} + \frac{1}{3d}, & d \equiv 2 \pmod{3}, \end{cases} \quad (k, d) = 1. \quad (20)$$

The gcd assertion follows from $3k = d \pm 1$. Because

$$\frac{M}{d} \geq X^{3/5-523/1000} = X^{77/1000} \rightarrow \infty,$$

the interval $[M, 2M]$ contains an integer $m \equiv k \pmod{d}$. Choose one. Then m is a unit modulo d and $\mathfrak{e}(mn/d) = \mathfrak{e}(kn/d)$.

The numerator $a = 1$ used here is an actual reduced completed frequency. In fact $d/(16H) \rightarrow \infty$ uniformly in the prescribed H -range, and the Round 11 positivity argument gives

$$|S_{V,H}(1/d)| \geq \frac{H}{2\sqrt{2}} \int V > 0$$

for all sufficiently large X . The phase being tested is therefore not attached to a missing or zero-weight frequency.

For $n \in [N, 2N]$, equations (20) give

$$\mathfrak{e}(kn/d) = \mathfrak{e}(n/3)(1 + O(N/d)), \quad (21)$$

with an absolute implied constant for either sign. The prime number theorem in the two fixed reduced classes modulo 3 gives, for $b = 1, 2$,

$$\#\{N \leq p < 2N : p \equiv b \pmod{3}\} = \left(\frac{1}{2} + \alpha(1)\right) \frac{N}{\log N}. \quad (22)$$

Summing (21) over these primes shows

$$\sum_{n \equiv b(3)} \tilde{\beta}(n) = \alpha(b/3) \left(\frac{1}{2} + \alpha(1)\right) \frac{N}{\log N} + O\left(\frac{N^2}{d \log N}\right). \quad (23)$$

The perturbation is $O(N/\log N)$, uniformly, since $N/d \leq 2X^{-123/1000}$. All supported primes exceed 3. Hence the exact principal subtraction in (15) is one half of the sum of the two class sums. Their difference divided by two is

$$\frac{\alpha(1/3) - \alpha(2/3)}{4} \frac{N}{\log N} + O\left(\frac{N}{\log N}\right) = \frac{i\sqrt{3}}{4} \frac{N}{\log N} + O\left(\frac{N}{\log N}\right),$$

which proves (18). The two signs in (20) both tend to the same cubic phase, so the leading sign in (18) does not switch with $d \pmod{3}$.

Finally, for $L = 2$, the right side of (15) at $r = 3, s = 1$ is $O(2^{C_{\text{sw}}} N(\log X)^{-2})$. Since $\log N = (2/5) \log X$, (18) is larger by an unbounded logarithmic factor for every fixed C_{sw} and implied constant. This is the claimed failure. $\square$

This lemma has a precise scope. A convolution expansion of an additive prime phase creates terms $e(amn/d)$. To absorb them into the shorter factor and apply Proposition 2.18, one would have to check the Siegel – Walfisz property of $\beta(n)e(amn/d)$, uniformly in the variables being summed. It does not follow from that property for β , even on the actual allowed conductors, at allowed scales and with a unit m . A delta sequence supported on the chosen m is itself an allowed longer coefficient, so the issue is not a failure of divisor boundedness. Moreover, the absorbed sequence varies with m, a, d , whereas the source dispersion sum is stated for a fixed coefficient family before its modulus sum.

The lemma does not assert that every factor in a particular Heath – Brown expansion fails (15), or that one bad slice prevents an averaged estimate. Such an averaged estimate would be new input to the proposed transfer and would have to retain this dependence. Source Proposition 2.18 and its proof in Appendix A.4.2 remain consistent with the example: their hypothesis (15) has been violated after the proposed absorption.

4. The shift interval is not a bounded product of local residue sets

There is a second tempting use of the source. Equation (2.5), printed p.7, permits a set of primes and one coherent primitive residue class that depend on X , uniformly, but fixes the class outside the modulus sum. For one fixed shift h , this is the legal Round 9 application. Simultaneous shifts $h \asymp H$ require more than that statement.

Source Proposition 2.14, printed pp.9 – 10, lifts coherent estimates to a product of local nonempty class sets $A_p \subseteq (\mathbb{Z}/p\mathbb{Z})^\times$, at the explicit cost

$$m(d) = \prod_{p \mid d} |A_p|. \quad (24)$$

It obtains a fixed divisor-weight cost when $|A_p| \leq K$ for a fixed K . Its underlying finite inequality is valid without this restriction; the restriction is what permits the subsequent source error bound without a power cost.

Choose a closed interval $[z_0, z_1] \subset (1, 2)$ of positive length $\ell = z_1 - z_0$ on which V is strictly positive. Such an interval exists. For any $d \in F_X$, put

$$H_d = \{h \in [z_0 H, z_1 H] \cap \mathbb{Z} : (h, d) = 1\}.$$

For every prime $p \mid d$ and every unit $a \bmod p$,

$$\begin{aligned} \#\{h \in H_d : h \equiv a \pmod{p}\} &\geq \frac{\ell H}{p} - 1 - \sum_{\substack{r \mid d \\ r \text{ prime}, r \neq p}} \left(\frac{\ell H}{pr} + 1\right) \\ &\geq \frac{\ell H}{p} \left(1 - \frac{347}{\lambda X^\kappa}\right) - 348. \end{aligned} \quad (25)$$

The first line uses the Chinese remainder theorem for the simultaneous conditions $h \equiv a \pmod{p}$, $h \equiv 0 \pmod{r}$; the primes are distinct. Every $p \mid d$ is at most $X^{9/100}$, so

$$\frac{H}{p} \geq X^{1/6-9/100} = X^{23/300} \rightarrow \infty.$$

Thus (25) is positive uniformly in p, a, d, H , once X is sufficiently large.

Consequently, even **after the global coprimality restriction**, the image of H_d modulo each $p \mid d$ is the entire unit group. Its smallest product-of-local-images hull therefore has

$$A_p = (\mathbb{Z}/p\mathbb{Z})^\times, \quad m(d) = \prod_{p \mid d} (p-1) = \phi(d). \quad (26)$$

Since

$$1 \geq \frac{\phi(d)}{d} = \prod_{p \mid d} (1 - 1/p) \geq 1 - \frac{348}{\lambda X^\kappa},$$

one has $m(d) \sim d \asymp Q$, uniformly on F_X . On the other hand $\tau(d) = 2^{348}$ is constant. For any fixed B, C ,

$$\frac{m(d)}{\tau(d)^B (\log X)^C} \rightarrow \infty. \quad (27)$$

Therefore the divisor-weight consequence of Proposition 2.14 is not available for this product hull. Source logarithmic savings cannot absorb this $\asymp Q$ factor.

The actual interval supplies only $O(H)$ coherent global classes, since $H \ll d$. Its local images have discarded the correlation between residues at different primes. The failure in (27) is a failure of this product-hull transfer, not an impossibility theorem for a nonproduct residue set or a signed average over the interval.

5. A short-factor reinterpretation also misses the stated scale range

If one instead tries to treat the shift length H itself as the shorter convolution scale N at total scale X , the second inequality of (11), with the actual ω, δ , forces

$$\sigma < \frac{1 - 48\omega - 16\delta}{4} = \frac{51}{500}.$$

It follows that

$$N \geq X^{1/2-\sigma} > X^{199/500}, \quad (28)$$

whereas $H \leq X^{2/7}$. The positive exponent gap is

$$\frac{199}{500} - \frac{2}{7} = \frac{393}{3500}. \quad (29)$$

Thus this specific reinterpretation is outside Proposition 2.18 before coefficient conditions are considered. This says nothing about a different regrouping, a different total scale with all errors tracked, or a new multivariable dispersion proof.

6. What is and is not established

The actual factorization information has been used, not merely a scalar distribution exponent. The conductors in the counterexample satisfy the complementary predicates and retain the full original triple dense-divisibility budget. The modulus cutoff, both convolution scales, the primitive restrictions, the sign of the cubic phase, and the source's absolute Siegel – Walfisz normalization have all been checked.

What remains absent is an estimate for the signed joint object (4), retaining its $\mathcal{d}$ -dependent additive phase, its conductor coefficient $M_{\mathcal{d}}^{(j)}$, and its growing coherent shift interval. A direct $X(\log X)^C$ -scale bound for that object would already remove the remaining power loss in (6); the actual zeta covariance target requires sharper control of the logarithms and all other components as well. One sufficient component-level condition, with suitable uniformity for the smooth separated profiles, would be

$$\log X \mid B_0 \mid + \mid B_1 \mid = o(X \log X). \quad (30)$$

No assertion that (30) follows from the source theorem, or that it is false, is made.

The source's Appendix A.4.2, printed pp.35 – 36, fixes coefficient supports and derivative bounds before taking suprema, performs a dispersion square, chooses dense divisors with their quotient losses, and retains its original unit masks. The resulting rational completion estimate is part of that argument; it is not a theorem about arbitrary modulus-dependent twisted coefficients. Reusing it for (4) would require a new reduction with the extra variable and coefficient dependence explicitly controlled.

The bounded conclusion is therefore: phase absorption and product-local residue lifting do not justify the proposed transfer, with the exact failures proved above. A joint dispersion argument remains a meaningful open arithmetic step. The finite check accompanying this note verifies rational inequalities, the cubic-phase algebra and modular selection; it does not purport to prove a numerical prime-gap bound or realize the large asymptotic moduli on a computer.

7. Provenance and reproduction

Primary source: [OpenAI, Improved short gaps between primes.](#)

Source location	Use in this note
Definition 2.1 and Proposition 2.3, pp.4 – 5	Strong recursive dense divisibility and actual complementary moduli
Definitions 2.6 – 2.9, p.6, especially equation (2.4)	Coefficient families, scale and the exact uniform SW requirement

Source location	Use in this note
Proposition 2.10 and equation (2.5), p.7	Prime-interval SW; one coherent class fixed outside the modulus sum
Proposition 2.14, pp.9 – 10	CRT product-set lift and its explicit $m(d)$ cost
Proposition 2.18, pp.10 – 11, equation (2.14)	Bilinear dispersion hypotheses and scale inequalities
Appendix A.4.2, pp.35 – 36	Source dependence on fixed coefficients, divisor choices and unit masks

The local primary PDF SHA-256 is 456f05e0a3ef589ebb0e9abcf31f140f3c945adbf6950e00ef371a3c88b0930. Its extracted text SHA-256 is ded13a7c74fcfce64e85769e05b5869803dcdcf53b88be2c2f3c0b344f95ee84. The source-preserving official repository was pinned at commit 61340d0b74163003b32756bb16e91d9209a5e330. No source files were modified.

The exact construction and coefficient-isolation proof used from Round 11 is CONDUCTOR_MASS_LOWER_BOUND.md, SHA-256 46347799005bb0f53af25c2a7e8ffb2b2217d92688c7651327dde3562f114b92. The current note only uses its explicitly stated all-large- X conclusions, not a claimed numerical realization.

Run the adjacent standard-library Python script check_dispersion_hypotheses.py. It writes dispersion_hypotheses_certificate.json, including the final report/script hashes and primary-source hashes. The fixed small integer examples in that script test only the algebra of (19) – (20); they are not substitutes for the canonical-family existence proof.

Current report 49: Direct Selberg control of the centered mixed remainder: the precision and sign gap

Collection: R12 — exact limits of sampling and dispersion transfers.

Source: [research/dyson/round12/mixed-arithmetic/SELBERG_MIXED_REMAINDER_AUDIT.md](https://research.dyson/round12/mixed-arithmetic/SELBERG_MIXED_REMAINDER_AUDIT.md).

SHA-256: [6223e9a54c44c31d344c2f932fe81ca6ad4672e5f6ef625adfeae07b4ed5a308](https://github.com/dyson/research/blob/83e7372b38b74227b67a9579413ae7773bb6a7a8/round12/mixed-arithmetic/SELBERG_MIXED_REMAINDER_AUDIT.md) . Git blob: [83e7372b38b74227b67a9579413ae7773bb6a7a8](https://github.com/dyson/research/blob/83e7372b38b74227b67a9579413ae7773bb6a7a8/round12/mixed-arithmetic/SELBERG_MIXED_REMAINDER_AUDIT.md) . Original bytes: 10511.

Direct Selberg control of the centered mixed remainder: the precision and sign gap

Date: 2026-09-05. This is a bounded analytical test of the actual genuine-prime remainder. It gives an explicit, insufficient one-sided bound using a primary RH short-interval theorem. It does not introduce a substitute point process, differentiate an unknown error, or separate divergent prime and continuum tails. No parameter scan was run.

Outcome. The direct Selberg – Gallagher route below gives

$$|M_T(b)| \ll \frac{\log T}{b^2}. \quad (1)$$

Consequently its lower bound on $\int_b^{2b} B_T(s) ds$ is only of order $-\log T/b$, whereas the target is $-3/(2b^2) + \varepsilon/b^2$. This is a quantified failure of this particular arithmetic estimate, not an impossibility theorem and not a claim to improve the previously known individual-norm bounds.

This is not the strongest known RH norm estimate. Round 10 already established the stronger $E_T(b) = O(1)$ on the slow range. The same primary page also records Selberg's stronger global weighted estimate, discussed in Section 4 below. Neither is to be replaced by the weaker local-bound calculation (1). The remaining issue is a one-sided fluctuation estimate in a shrinking logarithmic shell, with its leading coefficient and next-order precision.

1. Primary input is already about genuine primes

Let $\theta(x) = \sum_{p \leq x} \log p$. Saffari and Vaughan, [On the fractional parts of \$x/n\$ and related sequences II](#), Ann. Inst. Fourier 27(2) (1977), define this prime-weighted function in (6.1). Under RH, their Lemma 5, equation (6.4), gives

$$\int_X^{2X} |\theta(x + \eta x) - \theta(x) - \eta x|^2 dx \ll \eta X^2 \log^2(2/\eta), \quad X \geq 4, \quad 0 < \eta \leq 1. \quad (2)$$

The uniformity in η , and the fact that the statement concerns θ , were checked in the primary text on printed pages 19 – 22. No GRH or pair-correlation conjecture is assumed. Formula (2) is an upper bound, with an unspecified absolute constant; it is not a variance asymptotic or a positive lower bound.

2. Gallagher's inequality applied to the combined centered measure

Retain $L = \log T$, $N = \lfloor T/L^6 \rfloor$, and the Round 11 finite centered tails

$$\begin{aligned} C_{s,Y}(t) &= \int_{(N,Y]} x^{-1/2-s(2L)-it} d(\theta(x) - x), \\ D_{s,Y}(t) &= \int_{(N,Y]} \left(\frac{\log x}{L} - 1 \right) x^{-1/2-s(2L)-it} d(\theta(x) - x). \end{aligned} \quad (3)$$

Both the prime and continuous terms occur inside the same finite signed measure. Let $\tau = 1/T$. For any finite measure ν on logarithmic coordinates, Plancherel applied to its convolution with an interval of length τ gives

$$\int_{-T}^T |\mathfrak{V}(t)|^2 dt \leq \frac{2\pi T^2}{c_0^2} \int_{\mathbb{R}} |\nu((v, v+\tau])|^2 dv, \quad c_0 = \frac{\sin(1/2)}{1/2}. \quad (4)$$

Here the transform uses e^{-itv} . The Fourier transform of the interval has modulus at least c_0/T for $|t| \leq T$, which proves (4), including its direction. This remains valid for the finite prime atoms plus continuous density in (3).

For a logarithmic interval starting at $x = e^v$, the right side contains $\int_x^{xe^\tau} g(u) d(\theta(u) - u)$. Put $F = \theta - \text{id}$. Integration by parts on this short interval expresses it as

$$g(xe^\tau)[F(xe^\tau) - F(x)] - \int_0^{e^\tau-1} xg'(x(1+\eta))[F(x(1+\eta)) - F(x)] d\eta. \quad (5)$$

For $x \in [X, 2X]$, apply (2) to each increment in (5), then Minkowski. The integral in η is bounded by a constant times the endpoint contribution since $\int_0^\tau \sqrt{\eta} \log(2/\eta) d\eta \ll \tau^{3/2} \log(2/\tau)$. For $\sigma = 1/2 + s/(2L) \in [1/2, 3/4]$, this gives

$$\int_X^{2X} \left| \int_x^{xe^\tau} g(u) dF(u) \right|^2 \frac{dx}{x} \ll X^{1-2\sigma} \tau L^2 w_X^2, \quad (6)$$

where $w_X = 1$ for $g(u) = u^{-\sigma}$, and $w_X = |\log X/L - 1| + O(1/L)$ for $g(u) = (\log u/L - 1)u^{-\sigma}$. All constants are uniform in the stated range. The extra $1/L$ controls the variation across a dyadic interval and the derivative of the logarithmic weight.

3. Cutoffs and dyadic summation

For the upper cutoff Y , the crossing windows have total logarithmic length $O(1/T)$. The RH pointwise bound on F , together with weighted partial summation, bounds their contribution after (4) by

$$O(TY^{-s/L} \log^4 Y \left(1 + \left| \frac{\log Y}{L} - 1 \right| \right)^2).$$

It tends to zero for fixed $T, s > 0$ as $Y \rightarrow \infty$. The order of limits is explicit. The uniform analytic cutoff from Round 11 could also be used; no practical finite-prime experiment is asserted.

At the lower cutoff, N is an integer and the prime tail is strict. For large T , $Ne^{1/T} < N+1$. Thus a crossing window from below N contains no retained prime atom and only the continuous density on its very short intersection with (N, ∞) . Its contribution is negligible. The endpoint term implicit in the limiting analytic continuation has not been

discarded: (3) converges to precisely the centered tail defined in Round 11.

Partition the starting points $x \geq N$ into $[2^j N, 2^{j+1} N]$. This partitions the integration variable in (4), not the prime measure, so it creates no extra prime-cutoff boundary at the dyadic endpoints. Let $d_N = 1 - \log N/L = 6 \log L/L + o(1/L)$. For $2 \leq s \leq 2G(T)$, where $G = \alpha(\log L)$, geometric summation gives

$$\sum_{j \geq 0} (2^j N)^{-s/L} \ll e^{-s \frac{L}{s}},$$

$$\sum_{j \geq 0} (2^j N)^{-s/L} \left(\left| -d_N + \frac{j \log 2}{L} \right| + \frac{1}{L} \right)^2 \ll e^{-s \frac{L}{s^3}}. \quad (7)$$

For the second estimate, use the exact geometric-series sums for $1, j, j^2$, together with $s d_N = \alpha(1)$. Equations (4), (6), and (7) therefore prove for the limiting centered tails

$$\frac{e^s \|C_s\|_2^2}{TL^2} \ll \frac{L}{s}, \quad \frac{e^s \|D_s\|_2^2}{TL^2} \ll \frac{L}{s^3}. \quad (8)$$

These norms are over $[0, T]$, a subinterval of that bounded by (4).

Cauchy – Schwarz in the actual mixed product now gives

$$\left| \frac{e^s}{TL^2} \Re \langle C_s, D_s \rangle \right| \ll \frac{L}{s^2}.$$

The already quantified pole replacement is $O(e^s L^{-3})$, uniformly negligible here; adding it proves (1) for M_T itself. No prime-power replacement is necessary because (2) was genuine-prime input from the start.

4. The concrete integrated lower bound is too weak

Write $B_T(s)$ for the single jointly centered remainder of Round 11, including all off-diagonal prime-prime, prime-continuum, and continuum-continuum terms before taking the cutoff limit. Its proved diagonal decomposition implies

$$\int_b^{2b} B_T(s) ds \geq -C \frac{L}{b} - \frac{1}{2b} - \frac{3}{4b^2} + o(b^{-2}). \quad (9)$$

This follows from the actual mixed-product bound above and the exact integrated prime diagonal. It is a valid one-sided consequence, but its negative error is larger than the required b^{-2} accuracy by a factor of order Lb .

For any fixed nonnegative smooth χ supported in $(1, 2)$, the same proof gives

$$\left| \int_b^{2b} \chi(s/b) M_T(s) ds \right| \ll_\chi \frac{L}{b}.$$

Thus smoothing in b alone does not improve this estimate. The loss of L is visible in (2): the theorem supplies $\log^2(2/\eta)$, while the fluctuation scale at $\eta \asymp 1/T$ has only one logarithm. Even hypothetically replacing the squared logarithm by a single logarithm in this argument would yield only $|M_T(s)| \ll s^{-2}$, with no positive leading coefficient or the required next-order deficit below 2. This diagnoses the sign loss of the Cauchy – Schwarz step, not a logical impossibility of stronger consequences from a strengthened theorem.

There is a material stronger-input nuance. On the same primary page, equation (6.5) recalls Selberg's global estimate

$$\int_0^{\eta^{-A}} |\theta(x + \eta x) - \theta(x) - \eta x|^2 \frac{dx}{x^2} \ll_A \eta \log^2(2/\eta), \quad (9a)$$

for a fixed admissible exponent $A > 1$, corresponding to the fixed exponent parameter in the source remark. It is not legitimate to let that exponent grow with T . For any such fixed A , the active edge shell $1/s \leq \log x/L - 1 \leq 2/s$ lies within this range when s and then T are large enough. Using a fixed smooth filter of $s(\log x/L - 1)$, the argument (4) – (5) with (9a) gives normalized squared norms $O_A(1)$ and $O_A(s^{-2})$ for the filtered centered tail and its log-weighted companion. Indeed the Mellin damping on that shell is $O(e^{-s})$; the excess weight is $O(1/s)$, and its differentiation costs only $O(1/L)$. These are bounds for a filtered component, not a new identity for the full mixed moment.

Cauchy – Schwarz then controls that component only by $O_A(1/s)$, and does not give a positive leading $1/s^2$. The global estimate supplies a bound on accumulated positive variance, not a quantitative lower increment on the shrinking shell and not the required signed mixed covariance. This is why merely selecting the stronger stated Selberg bound does not complete the proposed argument. The explicit lower bound (9) is retained as a fully quantified local-theorem consequence, not as a claim of optimal RH control.

5. Smoothing does not make the actual kernel termwise positive

At finite cutoff, Fubini in (3) shows that the sharp s -integrated mixed product has the real kernel

$$\frac{(xy)^{-1/2}}{L^2} (e^{-bv(x,y)} - e^{-2bv(x,y)}) \operatorname{sinc}_0(T \log(x/y)), \quad v(x, y) = \frac{\log x + \log y}{2L} - 1. \quad (10)$$

It acts on the same finite signed measure $d\theta - dx$ in both variables. Smooth nonnegative s -weights replace the exponential difference by the corresponding integral of ve^{-sv} ; they leave the sinc factor and centering intact.

For example, $x, y > T$ and $T \log(x/y) = 3\pi/2$ give a positive exponential factor but sinc equal to $-2/(3\pi)$. These are points in the domain of the actual continuum terms, not a fabricated point process. This rules out simply declaring every off-diagonal kernel contribution nonnegative after b -smoothing. A negative kernel value alone is not a proof that its quadratic form lacks positivity; no such general claim is made here. The relevant arithmetic lower estimate is still missing.

The bounded attempt therefore stops with the genuine-prime inequality (9) and an explicit precision/sign deficit. Repeating identities or replacing the centered covariance by separate divergent sums cannot fill that deficit. A new one-sided fluctuation estimate, or cancellation in the centered mixed pairing beyond these upper norms, is required.

Current report 50: Independent review of the three bounded Round 12 attempts

Collection: R12 — exact limits of sampling and dispersion transfers.

Source: `research/dyson/round12/INDEPENDENT_ROOT_REVIEW.md`.

SHA-256: `bff32012d2eb29877b51cd569d976016418d16abe6120a000a3c811f737a8557` . **Git blob:**
`4e1e0a5cf1592d01c9bb2ab182f5e2c72eeec035` . **Original bytes:** 7673.

Independent review of the three bounded Round 12 attempts

Date: 2026-09-05. Reviewer: root Astra, independently of all three authoring lanes. I read all three complete reports, checked the inherited Round 11 construction, inspected the relevant primary source statements, and checked the new elementary arguments below. Accepted scope: ordinary proofs of specific failed shortcuts and an insufficient actual-prime estimate. No stronger bound for the actual signed prime pairing is proved.

Pinned author reports:

Report	SHA256
sampling-geometry/ACTUAL_SUPPORT_SAMPLING_OBSTRUCTION.md	cb52d72f6068c3030968209d8aa028439ea4dc309aa5584d216a1d7d30a1a59d
dispersion-transfer/DISPERSION_HYPOTHESIS_OBSTRUCTION.md	26536af7b5e6ebfb8fb4f7c7be993543c57152e29c6fd2b039c5eb147ccefd95
mixed-arithmetic/SELBERG_MIXED_REMAINDER_AUDIT.md	6223e9a54c44c31d344c2f932fe81ca6ad4672e5f6ef625adfeae07b4ed5a308

1. Positive sampling on the actual modulus support

The inherited count supplies at least $c_0 Q/(2 \log^{348} X)$ actual terminal conductors. Each has at least $d/(32H)$ primitive low numerators and $d > Q/2$, giving the stated $c_0 Q^2/(128 H \log^{348} X)$ distinct frequencies. The number of cells of length at most $1/(100X)$ in the small arc is at most $8X/H$ eventually. Pigeonhole therefore supplies the claimed occupation without a local distribution theorem.

For $N = \text{ceil } X$, $M = \text{floor}(X/10)$, the integer frequencies $N, \dots, N+M-1$ lie in $[X, 1.1X]$. Within the chosen cell, factoring out the carrier leaves phases at most $\pi/1000$. The lower bound $M/(2 \sqrt{H})$ and Parseval norm M/H give the sampling constant $c_0 Q^2/(81920 \log^{348} X)$. Squaring the actual coefficient lower bound $m_v H/(2 \sqrt{2}Q)$ gives the second constant $1/655360$. The global Parseval norms also imply both local small-arc envelopes for $\rho \geq 1/H$, including the derivative bound.

The constructed polynomial is artificial and phase-tuned. It is not the actual prime polynomial, nor a fixed smooth multiplier of its coefficients. The proof correctly limits its conclusion to positive sampling and the specified absolute-weight operator. It does not prove sharpness for all rearrangements of Cauchy--Schwarz or all signed weightings.

I checked the signed Gram identity directly by expanding the functional coefficient by coefficient. Completing each parent modulus and retaining its Ramanujan principal term gives exactly the residue kernel in equation (16); the zero Fourier term cancels. This equality holds for the Fourier-defined kernel at arbitrary integer n , while its identification with the original unit-restricted progression expression is made only for primes near X . The report states that distinction. Negative or complex cross terms in the signed Gram are not controlled from the positive sampling lower bound.

2. The proposed direct dispersion transfers violate actual source hypotheses

I checked the 186 source's Siegel--Walfisz definition, prime-interval example, product-local residue lift, and Proposition 2.18 scale/parameter conditions. At $\omega=.012$, $\delta=.001$, $\sigma=.101$ the three left sides are .888, .996 and .990. $N=X^4$, $M=X^6$ are legal, and the allowed modulus cap is exactly $X^{.523}$. No out-of-range parameter creates the counterexample.

For each canonical $d>Q/2$ coprime to 3, $k=(d-1)/3$ or $(d+1)/3$ is a unit and k/d tends to $1/3$. An m in $[M,2M]$ in that residue class exists because M/d tends to infinity. For primes n in $[N,2N]$, replacing $e(mn/d)$ by $e(n/3)$ costs $O(N/d)$ per term. PNT in the two fixed reduced classes mod 3 then gives the discrepancy $i \sqrt{3} N/(4 \log N)$, with uniform lower-order error $O(N^2/(d \log N))$. Both branches have the same limiting cubic phase and leading sign. This violates the source's SW condition for $L=2$. The original prime coefficient has SW; the phase-twisted family does not inherit it. A bad slice is not a proof that no averaged dispersion treatment is possible.

For the CRT claim, every prime factor of the constructed conductor is at most $X^{.09}$ and at least λX^{κ} . Counting a chosen local unit class and subtracting divisibility by each other prime leaves at least $(\ell H/p)(1-347/(\lambda X^{\kappa}))$ -348 positive candidates eventually, uniformly. Thus imposing the global gcd restriction does not reduce any local image. The smallest product hull has $\phi(d)$ classes, asymptotic to d , whereas $\tau(d)=2^{348}$ stays fixed. The bounded-local-class consequence of the source lemma is unavailable. This is a failure of the product-hull reduction; the original coherent interval has many fewer global classes and could be treated differently.

Finally the source constraint forces $\sigma<.102$, hence the short convolution factor exceeds $X^{.398}$. $H \leq X^{(2/7)}$ cannot be substituted for that factor at the same total scale. All three conclusions are source-application obstructions, not counterexamples to the source theorem or to the target covariance.

3. The direct Selberg mixed estimate is deliberately insufficient

I inspected the primary Saffari--Vaughan text and the rendered printed page 20. Equation (6.4) is indeed a uniform RH estimate for genuine θ and all $0<\eta \leq 1$. Equation (6.5) gives the stronger global weighted comparison quoted in the report; its exponent parameter must remain fixed. The report prominently states that its local calculation is weaker than previously established RH norm control and does not claim to exhaust the source's consequences.

The Mellin Gallagher inequality follows from interval convolution and Plancherel with the stated Fourier convention. Its direction and factor $2\pi T^2/c_0^2$ are correct. Short-interval integration by parts retains the centered increment $F(x(1+\eta)) - F(x)$; the derivative of the log weight is included. Applying Minkowski to the source bound gives equation (6), including the dyadic variation cost $1/\log T$.

The upper-cutoff crossing term tends to zero at fixed T, s as Y tends to infinity. At the strict integer lower cutoff, $N \exp(1/T) < N+1$ ensures that the crossing interval has no retained prime atom; its continuous contribution is negligible. Dyadic subdivision partitions the integration variable rather than splitting the prime measure, so it introduces no uncounted prime endpoints. Geometric sums in $1, j, j^2$ give the norms L/s and L/s^3 , with $s(1 - \log N/L) = o(1)$. Cauchy–Schwarz supplies only the insufficient absolute mixed bound L/s^2 and the integrated lower estimate as stated.

For the separately filtered active shell, the stronger fixed-exponent global estimate gives $O(1)$ and $O(s^{-2})$ normalized energies. The shell lies inside its fixed range once s is sufficiently large, the Mellin damping is $O(\exp(-s))$, and the smooth-filter derivative cost s/L is bounded. These estimates are for a component, not a decomposition asserted for the full mixed moment. They give no positive shell increment or required signed coefficient.

The integrated sinc kernel is correctly obtained from the symmetric excess factor $v \exp(-s v)$. Its negative pointwise lobe defeats a term-by-term positivity claim, but a negative kernel value is not by itself a negative quadratic-form witness. The report explicitly avoids that stronger false inference. All continuous and prime pieces remain jointly centered.

Decision

The next substantive improvement requires actual joint arithmetic: a prime-specific concentration estimate, a signed Gram/dispersion argument retaining the coherent shift and its phase, or a one-sided centered fluctuation estimate with the necessary shrinking-shell precision. These reports justify not repeating the three failed shortcuts. They do not prove that the remaining $X^{.023}$ factor or the famous conjecture is inaccessible.

Current report 51: Dyson—Montgomery round 13: extract a rational core and retain the signed remainder

Collection: R13 — rational-core extraction and the signed CRT remainder.

Source: research/reports/dyson_round13.md.

SHA-256: `b41b76a69d98b282aecc9223a90ba793182f6594f7e57a86245b404b1f779a0c` . Git blob: `debde73169d82f89966d80e848ff9b46ead92af7` . Original bytes: 5410.

Dyson—Montgomery round 13: extract a rational core and retain the signed remainder

Date: 2026-09-05. All claims below concern explicitly specified components of the actual arithmetic programme. The strongest bound for the original selected smooth prime discrepancy remains $O(X^{1.023} \log^5 X)$ under RH. No actual-zeta lower bound is proved here.

The arithmetic outcome

The new useful extraction is on the terminal Type II slice with $Q=X^{.523}$, $M=X^{.6}$, $N=X^{.4}$ and $X^{(1/6)} \leq H \leq X^{(2/7)}$. Its conductor coefficient is the actual $\mu(d)/d$, $d > Q/2$; its shift Fourier sum and primitive principal subtraction are retained. The inner coefficients are log-weighted primes in a fixed interval $[N, 2N)$, with natural unit restrictions on the longer variable.

On the zero-rational phase core $||am/d| | \leq C/N$, ordinary RH and partial summation replace that prime sum by $\text{integral_N}^{(2N)} e(\theta u) du$. Exact residue counting and the actual numerator weights give total replacement error

$$O_{C,V}(QMN^{-1/2} \log^2 X) = O_{C,V}(X^{923/1000} \log^2 X).$$

This is below $X \log X$, including fixed logarithmic losses. The integral main term is explicit and remains. The proof does not estimate it away. The zero rational in this expression is not the already deleted zero frequency of the original progression completion.

With the admissible test $\alpha_m=1$ and an actual complementary subfamily, a restricted positive phase block has real contribution at least

$$\frac{c_0 \int V X^{1123/1000}}{131072 (\log X)^{348}}.$$

The primitive principal term is included in this inequality. This is a real-prime arithmetic witness, but only for a subsum. Other phases, conductors and the actual outer coefficients can cancel it. It establishes neither a lower bound for the complete functional nor a failure of the desired zeta theorem.

For other small rational denominators, ordinary SW only gives the recorded logarithmic replacement error, insufficient after the present absolute summation. The legal minor arcs use width $2R/(qN)$, not fixed C/N cores. Centering over unit residues gives an exact variance improvement; the resulting factored estimate $X^{1.323} \sqrt{\log X}$ is still worse than Round 11's estimate for the original prime pairing. No improvement is claimed from it.

Read the [complete phase report](#), [primary-source audit](#) and [independent root review](#).

The exact signed–kernel remainder

The [complete CRT report](#) evaluates the smooth integer norm of the actual signed discrepancy kernel. Its main term is exactly $X(\text{integral } W)(\text{mean}^2 + \text{coefficient squared mass})$. Its remaining nonzero CRT modes are explicit.

Pairs of original moduli with gcd at least $X^{.1}$ have period at most $X^{.946}$, so smooth Poisson decay handles them. The complementary small-gcd range contains the unresolved long-period modes. The zero-mode covariance there is small because $X^{.1}/H$ tends to zero. Reduced-denominator gcd and original-modulus gcd are carefully distinguished.

A specified coherent block in the actual frequencies contributes at least a positive constant times $Q^2 H / \log^{696} X$. Its phase and occupancy constants are explicit. This remains a subsum of a signed expression; the report makes no inference that it lower-bounds the complete norm or remainder. Even an ideal generic norm of order XH would not by itself supply the desired prime-specific estimate. The [separate independent audit](#) accepts these conclusions with their limited scope.

Evidence and reproducibility

All 21 original research, review and reference files, totaling 6,973,902 bytes, are preserved in the local [Astra-Local-Archive/round13-originals](#). Fifteen research/review/receipt files are public and verbatim. The six full third-party paper/text/page-image files remain local, identified by primary URLs and hashes in the public receipts.

The [integration replay](#) runs two inspected scripts in a temporary copy. The exact CRT output is unchanged in every field. The phase output is unchanged after excluding only three temporary primary-source paths. Tests cover 100 CRT compatibility cases, the unit-masked Fourier identity, exact centered variance, residue bounds, rational exponents, constants and primary-source hashes. These are checks of the algebra and provenance, not numerical proofs of the analytic source theorems.

No new computational search, model service, infrastructure, prime realization at an enormous X or Fable session is introduced. The next arithmetic task is to control the signed retained main terms for the actual long coefficients. A separate Type I removal is being investigated and is not assumed by this report.

Remaining research target

The complete sharp covariance, its complementary divisor pieces and the strict mixed residual lower bound remain open. In particular the sufficient W_T lower limit $1/16$ and compact Fourier test above $7/10$ have not been established. These documents preserve a checkable advance in one extraction and an exact description of its remaining terms, not a proof of RH, AH refutation, GUE, a zeta-gap improvement or a sub-186 prime gap.

Current report 52: Averaging the rational resonances: one power–saving extraction, and its retained main term

Collection: R13 — rational-core extraction and the signed CRT remainder.

Source: [research/dyson/round13/phase-resonance/AVERAGED_RATIONAL_PHASE_TEST.md](https://research.dyson/round13/phase-resonance/AVERAGED_RATIONAL_PHASE_TEST.md).

SHA-256: `7f4285cb02241e22bdb29a1ad4952f7ab8249e3ec3bef984455a57ae05e41ebb` . Git blob: `5e5d74961452f7f7bd53ec8b59ef1bb0157ba7ef` . Original bytes: 20890.

Averaging the rational resonances: one power–saving extraction, and its retained main term

Date: 2026-09-05. Status: ordinary analytic bounds for an explicitly defined terminal Type II component. The extraction in Section 3 assumes RH for the Riemann zeta function. The counting, lower bound and minor-arc comparison do not. No improved bound for the complete actual-zeta covariance is proved.

The bad phases from Round 12 occupy a small fraction of the longer coefficient interval. That observation alone does not make them negligible. This note proves three quantitative statements at the actual scales:

- The rational core around zero can be replaced by an explicit integral with total error $O(X^{923} \log^2 X)$, after the actual top-conductor and low-numerator weights are summed, assuming RH.
- Its retained main term cannot uniformly be discarded by absolute estimates. An admissible prime-interval convolution with the actual complementary moduli has a positive restricted core of size at least a constant times $X^{1.123}/\log^{348} X$.
- The primary-source Vaughan bound, applied legally on enlarged minor arcs, gives only logarithmic cancellation at the chosen polylogarithmic cutoff. An elementary mean square in the longer variable is stronger here, but still gives $O(X^{1.323} \sqrt{\log X})$, which is weaker than the existing Round 11 bound for the original prime pairing.

Thus averaging repairs the pointwise perspective but does not yet remove the need for cancellation in the retained signed main terms.

1. Exact component, coefficients and natural unit restrictions

Write

$$Q = X^{523/1000}, \quad M = X^{3/5}, \quad N = X^{2/5}, \quad X^{1/6} \leq H \leq X^{2/7}. \quad (1)$$

Let Q_X be the full canonical balanced complementary family fixed in Rounds 9 – 12, and let

$$D_X = Q_X \cap (Q/2, Q].$$

For completeness, the canonical family consists of distinct squarefree $d = [D, E] > X^{1/2}$, with

$$D, E \leq X^{523/2000}, \quad p^{3/2} D_{\geq p} \leq X^{501/2000}, \quad p^{3/2} E_{\geq p} \leq X^{501/2000}, \quad (2)$$

where each respective owner condition is imposed at $p > X^{1/1000}$. Its source complementary theorem proves triple dense divisibility. Every prime factor of d is at most $X^{523/2000} < N$, because it divides one of the two roots.

Fix $V \in C_c^\infty(1, 2)$, and initially assume $|a_m| \leq 1$, supported on $M \leq m < 2M$. Define

$$P_N(\mathfrak{g}) = \sum_{\substack{N \leq p < 2N \\ p \text{ prime}}} (\log p) \mathfrak{e}(\mathfrak{g}p), \quad S_{V,H}(\beta) = \sum_h V(h/H) \mathfrak{e}(-\beta h), \quad \mathfrak{e}(t) = e^{2\pi i t}. \quad (3)$$

The inner prime support is fixed independently of m . For a set A of phases on the circle, put

$$\mathsf{T}(A) = \sum_{d \in D_X} \frac{\mu(d)}{d} \sum_{\substack{1 \leq a \leq d/(16H) \\ (a,d)=1}} S_{V,H}(a/d) \sum_{\substack{M \leq m < 2M \\ (m,d)=1 \\ am/d \bmod 1 \in A}} a_m [P_N(am/d) - \frac{\mu(d)}{\phi(d)} P_N(0)]. \quad (4)$$

The main variable m is the longer convolution factor; it is not a prime variable by default.

Formula (4) is a genuine terminal-conductor slice of the completed shifted progression discrepancy for the source-allowed convolution

$$F = \alpha * \beta, \quad \beta(n) = (\log n) 1_{\{n \text{ prime}, N \leq n < 2N\}}.$$

The source prime-interval Siegel – Walfisz statement, followed by partial summation, gives its required shorter-coefficient property. The scales and source parameters were checked in Round 12. This does not assert that an arbitrarily chosen a is one specific factor in the actual zeta decomposition.

The mask $(m, d) = 1$ in (4) is essential. All supported n -primes are units modulo d , by (2). In the original progression and principal sums, primitive shifts therefore restrict the product to $(m, d) = 1$. After this restriction, Fourier completion gives

$$\frac{1}{d} \sum_{a \bmod d} S_{V,H}(a/d) \left[\sum_{\substack{m,n \\ (mn,d)=1}} a_m \beta_n \mathfrak{e}(amn/d) - \frac{c_d(a)}{\phi(d)} \sum_{\substack{m,n \\ (mn,d)=1}} a_m \beta_n \right], \quad (5)$$

where c_d is the Ramanujan sum. For $(a, d) = 1$ and squarefree d , $c_d(a) = \mu(d)$, exactly as used in (4). The full-signed-family conductor coefficient is $\mu(d)/d$ here: no other multiple of $d > Q/2$ fits below Q . Nonprimitive Fourier numerators and all other conductors remain outside this slice.

The estimates below use the actual $S_{V,H}$, without replacing it by a fictitious positive weight. For upper bounds, the elementary facts

$$|S_{V,H}(\beta)| \ll_V H, \quad \sum_{\substack{1 \leq a \leq d/(16H) \\ (a,d)=1}} \frac{|S_{V,H}(a/d)|}{d} \ll_V 1 \quad (6)$$

are sufficient. Summing (6) over D_X costs at most $O_V(Q)$.

For a fixed divisor-bounded outer family, one may multiply all the following upper bounds by $O_\eta(X^\eta)$ for any fixed $\eta > 0$, including its prescribed logarithmic factors in a slightly larger η . This follows from the pointwise divisor bound and is deliberately conservative. In particular, the extraction in Section 3 still has a power saving whenever $\eta < 77/1000$. The lower-bound witness uses $a_m = 1$.

2. Exact residue counts for sparse resonances

Let A be a disjoint union of J circular intervals with total length $|A|$. For a unit $a \bmod d$, multiplication by a permutes the residue classes. Each class occurs in $[M, 2M)$ at most $M/d + 1$ times. Counting grid points of spacing $1/d$ in each interval gives

$$\#\{M \leq m < 2M : am/d \bmod 1 \in A\} \leq (M/d + 1)(d|A| + 2J). \quad (7)$$

The estimate remains an upper bound after imposing $(m, d) = 1$. No distribution theorem about the inverses of residues is needed.

Set $R = (\log X)^B$, for fixed $B > 0$. The fixed-width cores around all reduced rationals of denominator $q \leq R$,

$$\{\mathfrak{g} : \|\mathfrak{g} - b/q\| \leq C/N\}, \quad (b, q) = 1, \quad (8)$$

are disjoint for sufficiently large X , with the usual single rational $0/1$. They have $J \ll R^2$, total length $O_C(R^2/N)$. Consequently their bad- m count, for each d, a , is

$$\ll_C (M + d) \left(\frac{R^2}{N} + \frac{R^2}{d} \right) \ll_C \frac{MR^2}{N}. \quad (9)$$

Here $M/d \gg X^{0.77}$ and $d/N \gg X^{1.23}$, so the last simplification is uniform. The exceptional proportion is $O(R^2/N)$, genuinely small.

For legal use of Dirichlet approximation on the complement, use the enlarged arcs

$$M(R) = \bigcup_{\substack{q \leq R \\ (b, q) = 1}} \{\mathfrak{g} : \|\mathfrak{g} - b/q\| \leq \frac{2R}{qN}\}. \quad (10)$$

They too are disjoint eventually. Their total length is

$$\ll \frac{R}{N} \sum_{q \leq R} \frac{\phi(q)}{q} \ll \frac{R^2}{N},$$

and their number is $O(R^2)$. Thus the same coarse count (9) holds. The enlargement does not worsen its logarithmic scale, but is needed for the minor-arc denominator conclusion in Section 6.

3. RH extraction of the zero-rational core with a power-saving total error

Fix $C > 0$, and let

$$A_0(C) = \{\mathfrak{g} : \|\mathfrak{g}\| \leq C/N\}.$$

Use the unique representative $\mathfrak{g} \in [-C/N, C/N]$, and define the explicit continuous main term

$$J_N(\mathfrak{g}) = \int_N^{2N} \mathfrak{e}(\mathfrak{g}u) du, \quad J_N(0) = N. \quad (11)$$

When $\mathfrak{g} \equiv 0$, this is $(\mathfrak{e}(2N\mathfrak{g}) - \mathfrak{e}(N\mathfrak{g})) / (2\pi i \mathfrak{g})$.

Assume RH for ζ . The RH prime-number estimate $\theta(x) = x + O(\sqrt{x} \log^2 x)$, with $\theta(x) = \sum_{p \leq x} \log p$, follows, for example, from Schoenfeld's explicit Theorem 10, equation (6.3), printed p.337 of [Sharper bounds for the Chebyshev functions theta\(x\) and psi\(x\), II](#). One partial summation gives, uniformly on this core,

$$P_N(\mathfrak{g}) = J_N(\mathfrak{g}) + O_C(\sqrt{N} \log^2 N). \quad (12)$$

Indeed the endpoint errors are $O(\sqrt{N} \log^2 N)$, while the differentiated exponential has total variation $O(|\mathfrak{g}| N) = O_C(1)$. Endpoint choices $<$ versus $\leq$ contribute at most $O(\log N)$, already covered. Only the ordinary zeta RH is used for (12).

Let $J_0(C)$ be exactly (4) on this core, with its square bracket replaced by

$$J_N(\mathfrak{g}) - \frac{\mu(d)}{\phi(d)} N. \quad (13)$$

Then

$$\boxed{T(A_0(C)) - J_0(C) \ll_{C,V} \frac{QM}{\sqrt{N}} \log^2 X = X^{923/1000} O_{C,V}(\log^2 X).} \quad (14)$$

To prove this, (7) bounds the number of core m 's by $O_C(M/N)$ for each d, a . The error in the centered bracket is at most $(1 + 1/\phi(d)) O_C(\sqrt{N} \log^2 N)$. Apply (6), and sum over at most Q conductors. This proves (14) without assuming cancellation of $a, \mu(d)$ or $S_{V,H}$.

An additional $\log d$ conductor weight costs one logarithm, so the error is still $\alpha(X \log X)$. Any extra fixed logarithmic coefficient loss is likewise harmless for this positive exponent margin.

This is a proved, source-compatible arithmetic extraction for an identifiable phase component. It does not state that its main term J_0 is small. In particular, zero here describes the rational approximation to am/d , not the deleted zero frequency of the original completion: $a/d \equiv 0$ throughout (4).

4. The retained rational main is not uniformly negligible

Assume in this section that $V \geq 0$, $m_V = \int V > 0$, and choose the admissible outer sequence $\alpha_m = 1$ on $[M, 2M)$. Use the actual Round 11 subfamily $F_X \subset D_X$ of products of two large and 346 small primes. It has

$$|F_X| \geq \frac{c_0 Q}{2(\log X)^{348}}, \quad \mu(d) = 1, \quad \sum_{p|d} \frac{1}{p} \leq \frac{348}{\lambda X^\kappa} = \alpha(1), \quad (15)$$

where

$$\kappa = \frac{343}{346000}, \quad \lambda = 2^{-1/348}, \quad c_0 = \frac{(1-\lambda)^{348}}{2! 346! (9/100)^2 \kappa^{346}}.$$

All assertions hold for every sufficiently large real X , uniformly in the specified H -range.

For each unit $a \bmod d$, consider just the positive core

$$1 \leq s \leq \frac{d}{32N}, \quad s \equiv am \pmod{d}, \quad (s, d) = 1. \quad (16)$$

The same elementary excluded-prime count used in Round 11 shows that (16) has at least $d/(64N)$ possible residues s , once X is large. The interval length $d/(32N)$ tends to infinity, and its nonunits number at most that length times $\sum_{p|d} 1/p$, apart from the harmless endpoint unit count. Every such residue gives at least $M/(2d)$ representatives $m \in [M, 2M)$, because $M/d \rightarrow \infty$. Thus there are at least

$$\frac{M}{128N} \quad (17)$$

unit m 's in this positive core, for every allowed d, a .

Now keep only actual low unit numerators $1 \leq a \leq d/(16H)$. For h in the support of $V(h/H)$, a prime $p \in [N, 2N)$, and m from (16),

$$0 < sp/d \leq 1/16, \quad 0 < ah/d < 1/8.$$

The phase of each product term $e(sp/d - ah/d)$ is therefore between $-\pi/4$ and $\pi/8$. Its real part is at least $1/2$.

Writing $S_0 = \sum_h V(h/H)$, this gives

$$\Re[S_{V,H}(a/d)(P_N(am/d) - \frac{P_N(0)}{\phi(d)})] \geq (\frac{1}{2} - \frac{1}{\phi(d)}) S_0 P_N(0). \quad (18)$$

Eventually $\phi(d) \geq 4$, $S_0 \geq m_V H/2$, and $P_N(0) \geq N/2$, the last inequality using the ordinary prime number theorem.

Thus each centered term, after the coefficient $1/d$, has real part at least

$$\frac{m_V H N}{16d}. \quad (19)$$

This explicitly retains the primitive Ramanujan principal term; it has not been omitted to force positivity.

Round 11 supplies at least $d/(32H)$ such primitive a 's for each $d \in F_X$. Combining (15), (17) and (19), the restricted positive-core contribution satisfies

$$\Re T_{F_X}(A_0^+) \geq \frac{c_0 m_V}{131072} \frac{QM}{(\log X)^{348}}, \quad A_0^+ = (0, 1/(32N)]. \quad (20)$$

Its power is $QM = X^{1123/1000}$. Consequently it exceeds $X \log X$ by an unbounded factor eventually, even with the displayed fixed logarithmic denominator. Equation (14), applied to this restricted positive core by the same proof, shows that the corresponding explicit integral main term has the same lower bound with a smaller fixed constant under RH.

This is an obstruction to treating all resonant slices as negligible by an absolute estimate uniformly over source-allowed coefficients. It is **not** a lower bound for the full signed family, for all rational cores combined, or for the coefficient sequence arising in one particular zeta or Heath – Brown decomposition. Cancellation against other conductors, other phases, or actual outer coefficients is still possible. The lower bound uses real primes and the actual arithmetic support, not a free choice of artificial Fourier locations.

5. Other small rational denominators: explicit main terms, but only logarithmic extraction from SW

For $q \leq R$, $(b, q) = 1$, and $\mathfrak{g} = b/q + \xi$ in the enlarged arc (10), the ordinary Siegel – Walfisz theorem and partial summation give, for every fixed $A > 0$,

$$P_N(b/q + \xi) = \frac{\mu(q)}{\phi(q)} J_N(\xi) + O_{A,B}(N(\log X)^{-A}). \quad (21)$$

This includes the genuine-prime weights and both interval endpoints. To see the coefficient, sum the progression main terms against $\mathfrak{a}(bc/q)$ over unit classes $c \bmod q$; their sum is $c_q(b) = \mu(q)$. All supported primes exceed q . The costs of at most q classes and variation $O(1 + |\xi| N) \leq O(1 + 2R)$ are fixed logarithmic powers, absorbed by choosing the source SW order larger. The finite source parameter B stays fixed.

Replacing the inner bracket in (4) accordingly, including replacement of $P_N(0)$, defines an explicit finite rational-main expression on $M(R)$. Its total extraction error is bounded by

$$O_{A,B,V}(QM(\log X)^{-A}), \quad (22)$$

because (9) and (6) sum the inner $N \log^{-A-2B} X$ error. This bound by itself is insufficient for $\mathfrak{o}(X \log X)$: $QM = X^{1.123}$, and a fixed arbitrary logarithmic saving does not defeat its extra power.

An absolute bound for the retained oscillatory rational main is

$$O_{B,V}(QMR). \quad (23)$$

For the fixed-width cores (8), count $O(M/N)$ m 's per rational, use $|J_N| \leq N$, and sum the $1/\phi(q)$ main coefficient over the $\phi(q)$ reduced numerators. This costs $O(M)$ per denominator. For the enlarged arcs, the same order follows from

$$|J_N(\xi)| \ll \min(N, |\xi|^{-1}).$$

On the $1/d$ phase grid, summing this bound within distance $2R/(qN)$ of a fixed rational costs

$$O(N + d \log(4R/q)).$$

Each residue is repeated at most $M/d + 1 \ll M/d$ times. After the $1/\phi(q)$ coefficient and its $\phi(q)$ numerators are summed, the total is

$$O\left(M \sum_{q \leq R} [\log(4R/q) + N/d]\right) = O(MR),$$

uniformly for $d \asymp Q$. Then (6) gives (23). The separate primitive principal term is retained; its absolute value on all m , and hence on these arcs, is at most $O_V(X \log X)$, as in Section 6.

Ordinary zeta RH does not imply the analogous square-root progression error for all $q > 1$. Assuming RH for the corresponding Dirichlet L -functions would change (22), but that is an additional assumption and is not used here. No failure of the desired actual average is inferred merely from the insufficiency of (22).

6. A legal minor-arc bound, and the stronger elementary average in m

The primary source for the pointwise estimate is Montgomery – Vaughan, *Multiplicative Number Theory II: Primes and Sieves*, author-hosted draft, Theorem 17.1, equation (17.29), printed p.65 (PDF page 77). It states that, when $(b, q) = 1$ and $|\vartheta - b/q| \leq q^{-2}$,

$$\left| \sum_{n \leq Y} \Lambda(n) e(\vartheta n) \right| \ll \left(\frac{Y}{\sqrt{q}} + Y^{4/5} + \sqrt{Yq} \right) (\log Y)^{5/2}. \quad (24)$$

Taking the difference of prefixes at $N, 2N$, and subtracting the prime-power terms of size $O(\sqrt{N} \log^2 N)$, proves the same bound for P_N , with Y replaced by N up to absolute constants.

For $\vartheta \in \mathcal{M}(R)$, set $D_0 = \lfloor N/R \rfloor$. For sufficiently large X , $D_0 \geq N/(2R)$. Dirichlet approximation with this integer cutoff supplies

$$R < q \leq D_0 \leq N/R, \quad |\vartheta - b/q| \leq (qD_0)^{-1} \leq q^{-2}.$$

Indeed $q \leq R$ would place the phase inside (10), since $(qD_0)^{-1} \leq 2R/(qN)$. Therefore (24) applies legally and gives

$$|P_N(\vartheta)| \ll (NR^{-1/2} + N^{4/5}) (\log N)^{5/2}. \quad (25)$$

The complement of the fixed-width cores (8) alone does not imply $q > R$; small denominators may still have approximation error between C/N and $2R/(qN)$. That is why (10) was introduced.

Using (6), summing the longer variable absolutely, and keeping the primitive principal term separately, yields the following bound. The latter term costs at most $O_V(MN \sum_{d \leq Q} 1/\phi(d)) = O_V(X \log X)$, by the elementary reciprocal-totient estimate.

$$|\mathcal{T}(\mathcal{M}(R)^c)| \ll_V QM(NR^{-1/2} + N^{4/5}) (\log N)^{5/2} + X \log X. \quad (26)$$

At polylogarithmic R , its leading power is $QMN = X^{1.523}$; only logarithms are saved. This is not a useful improvement over Round 11.

There is a stronger unconditional averaged estimate that uses the fixed inner coefficients. Since $2N < d$ eventually, no two supported primes are congruent modulo d . For any unit $a \bmod d$, complete-period orthogonality gives

$$\begin{aligned} \sum_{M \leq m < 2M} |P_N(am/d)|^2 &\leq (M/d + 1) \sum_{r \bmod d} |P_N(r/d)|^2 \\ &= (M + d) \sum_{N \leq p < 2N} (\log p)^2 \\ &\ll (M + d) N \log N. \end{aligned} \quad (27)$$

One can retain the centering exactly, rather than paying for it separately. Every supported prime is a unit modulo d , so the Ramanujan identity gives

$$\frac{1}{\phi(d)} \sum_{r \bmod d}^* P_N(r/d) = \frac{\mu(d)}{\phi(d)} P_N(0) =: c_d.$$

Consequently the exact centered unit variance is

$$\sum_{r \bmod d}^* |P_N(r/d) - c_d|^2 = \sum_{r \bmod d}^* |P_N(r/d)|^2 - \frac{|P_N(0)|^2}{\phi(d)} \leq d \sum_{N \leq p < 2N} (\log p)^2. \quad (28)$$

Multiplication by a unit a permutes these unit residues. Splitting the longer interval into residue periods, restricting its nonnegative squared sum to the minor arcs, and applying Cauchy – Schwarz therefore proves

$$\sum_{\substack{M \leq m < 2M \\ (m,d)=1 \\ am/d \in M(R)}} |a_m(P_N(am/d) - c_d)| \ll M \sqrt{N \log N}. \quad (29)$$

This exact variance calculation was independently pointed out in the minor-arc source audit. It requires the fixed inner support and coefficients; it is not a formal consequence for arbitrary m -dependent inner sums. Thus (6) and (29) give

$$|T(M(R)^c)| \ll_v QM \sqrt{N \log N} = O_v(X^{1323/1000} \sqrt{\log X}). \quad (30)$$

This saves a power X^2 relative to the absolute QMN bound, but it is still insufficient at scale $X \log X$, and still worse than the existing $X^{1.023} \log^5 X$ estimate for the original completed prime polynomial. A bound for this factored test must not be presented as an improvement to that existing estimate.

7. Bounded conclusion and source record

The useful positive statement from this test is (14): ordinary RH extracts the zero-rational core with an error already below the required arithmetic scale, even after the actual weights are summed. Its explicit main term remains part of the problem.

The precise limitation is (20), together with (22) and (30): sparse resonances can retain a power-large main contribution on an admissible arithmetic block; the remaining available minor-arc estimates do not bound the whole factored pairing at the required scale. Any next advance through this route must exploit the actual outer coefficients and signs when handling those retained main terms or use a stronger joint estimate. No generic claim that such cancellation is impossible is justified by this note.

The source paper is [OpenAI, *Improved short gaps between primes*](#). Proposition 2.3 pp.4 – 5 supplies the complementary moduli; Proposition 2.10 p.7 supplies prime-interval SW; Definition 2.9 p.6 specifies its normalization; Proposition 2.18 pp.10 – 11 supplies the legal Type II scales. The R12 report pins and checks these exact hypotheses. The primary Montgomery – Vaughan PDF SHA-256 is 72448ec23158a3aece534c9cde633d5402f916d0367b4f320212cd7ad179d340. The companion source audit records the downloaded copy and printed-page check.

The adjacent exact script checks the rational exponents, finite Fourier completion with the natural unit mask, residue-count bounds, and all lower-bound constants. It does not numerically search for large primes, approximate the asymptotic threshold, or substitute a finite toy modulus for the canonical-family proof. This note and all checks are new R13 files; prior rounds remain frozen.

Current report 53: Round 13 — Genuine–prime minor arcs and a fixed–inner–interval averaging audit

Collection: R13 — rational-core extraction and the signed CRT remainder.

Source: [research/dyson/round13/minor-arc-source/MINOR_ARC_AND_FIXED_INTERVAL_AUDIT.md](https://research.dyson/round13/minor-arc-source/MINOR_ARC_AND_FIXED_INTERVAL_AUDIT.md).

SHA-256: `bbdb17478b9e885570b7b49c3ff9b94b0ceb98f12a41743edb1ce492cb50edc4` . Git blob: `ef97aa38a8c9de57b7d032e54121e0b903118f92` . Original bytes: 11180.

Round 13 — Genuine–prime minor arcs and a fixed–inner–interval averaging audit

Status: completed ordinary mathematical source/proof audit, not a numerical enclosure or a theorem about zeta pair correlation. This note records the source theorem, the necessary major-arc coverage correction, and exactly what the current arithmetic packet permits. No new parameter scan was run.

1. Objects and the primary minor–arc theorem

Write $\mathfrak{e}(u) = \exp(2\pi i u)$, and let

$$P_N(a) = \sum_{N \leq p < 2N} (\log p) \mathfrak{e}(ap), \quad M \asymp X^{3/5}, \quad N \asymp X^{2/5}, \quad Q = X^{523/1000}.$$

The inner support is fixed independently of the integer $m \in I_M$, where I_M is an interval containing $M + O(1)$ consecutive integers. The current outer coefficients initially satisfy $|a_m| \leq 1$. Extra divisor/logarithmic weights must be charged through their actual norms.

Source theorem. Montgomery and Vaughan, *Multiplicative Number Theory II: Primes and Sieves*, author-hosted manuscript, Theorem 17.1, equation (17.29), printed page 65 (PDF page 77, zero-based index 76), states that, for coprime integers a, q , $q \geq 1$, and

$$|a - a/q| \leq q^{-2},$$

$$\left| \sum_{n \leq Y} \Lambda(n) \mathfrak{e}(an) \right| \ll (Y q^{-1/2} + Y^{4/5} + Y^{1/2} q^{1/2}) (\log Y)^{5/2}. \quad (1)$$

Definition (17.28) is explicitly the von Mangoldt prefix, not an arbitrarily twisted coefficient sequence. The proof on printed pages 65 – 66 chooses the Vaughan parameters $U = V = Y^{2/5}$. It explicitly says that $q > Y$ gives a bound weaker than the trivial estimate. Thus no usable $q \leq Y$ restriction should be silently imported into an application whose exact denominator exceeds Y .

Primary source: [author-hosted Montgomery – Vaughan manuscript](#). The version and local SHA256 are recorded in `sources/receipt.json`; this is a manuscript source, not a claim about publication status or the strongest possible modern estimate.

Taking two prefixes, and subtracting the prime powers, gives

$$|P_N(\alpha)| \ll B_N(q), \quad B_N(q) = (Nq^{-1/2} + N^{4/5} + \sqrt{Nq})(\log(2N))^{5/2}. \quad (2)$$

Indeed, the discarded powers contribute at most $\sum_{k \geq 2, p^k \leq 2N} \log p \ll \sqrt{N} \log^2(2N)$, which is absorbed by (2). No RH is used here.

If w is of bounded variation on $[N, 2N]$, partial summation gives the uniform weighted version

$$\left| \sum_{N \leq p < 2N} (\log p) w(p) \alpha(p) \right| \ll (\|w\|_\infty + \text{Var}(w)) B_N(q). \quad (3)$$

The same rational approximation applies to every prefix with endpoint between N and $2N$, so there is no interval-uniformity gap. Smooth scaled weights $w(p/N)$ have a constant depending on their fixed variation norm.

Unweighted prime sums follow by replacing $w(x)$ with $w(x)/\log x$, which costs $O(1/\log N)$. Uniform variation suffices for this pointwise statement even if the weight depends on m ; it does not, by itself, establish the fixed-coefficient orthogonality argument in Section 3.

2. The required major arcs, and an obstruction to narrower coverage

Let $R = (\log X)^B$, with fixed $B > 0$, and take X large enough that $2 \leq R \leq \sqrt{N}$. Put

$$M_R = \bigcup_{1 \leq q \leq R} \bigcup_{(a,q)=1} \{a \pmod{1} : \|a - a/q\| \leq \frac{2R}{qN}\}. \quad (4)$$

The harmless factor 2 accommodates floors. Dirichlet approximation with denominator cap $K = \lfloor N/R \rfloor$ gives coprime $a, q, q \leq K$, and

$$|a - a/q| \leq \frac{1}{qK} \leq \frac{2R}{qN}, \quad |a - a/q| \leq q^{-2}.$$

Consequently, outside (4), $R < q \leq N/R$. Applying (2),

$$|P_N(\alpha)| \ll (NR^{-1/2} + N^{4/5})(\log(2N))^{5/2}. \quad (5)$$

For any fixed desired logarithmic saving, a sufficiently large fixed B supplies it in this pointwise estimate.

It is incorrect to deduce this denominator range after deleting only neighborhoods $|a - a/q| \leq C/N$, with fixed C , for $q \leq R$. The available Dirichlet error is $O(R/(qN))$, which is larger for small q . A concrete obstruction is $\alpha = K_0/N$, where $K_0 > C$ is a fixed nonintegral number. For large N , this is outside all those fixed-width cores: for $q \geq 2$, a/q is separated from zero by at least $1/R$, much larger than $1/N$. But the ordinary PNT and partial summation give

$$P_N(K_0/N) = N \int_1^2 \alpha(K_0 u) du + o(N),$$

with nonzero integral. It is therefore not a minor arc with arbitrary logarithmic cancellation. Grid phases with denominator $d \gg N$ approximate such a point within $1/d = o(1/N)$, so the issue also occurs in the sampled setting.

The coverage repair does **not** worsen the coarse bad-point count already contemplated. The total length of (4) is

$$\ll \frac{R}{N} \sum_{q \leq R} \frac{\phi(q)}{q} \ll \frac{R^2}{N},$$

and the number of constituent circular intervals is $O(R^2)$. For any reduced a/d , multiplication by a permutes the grid of d residues. Counting complete periods and enlarging the final incomplete period yields

$$\#\{m \in I_M : am/d \in M_R\} \ll (M+d) \left(\frac{R^2}{N} + \frac{R^2}{d} \right). \quad (6)$$

Restriction to $(m, d) = 1$ only decreases this count. The interval-count term in (6) must be retained before using the actual $d > N$ range.

3. A stronger exact average at $a = 1$, including the natural centering

Let $d > N$, and let b_n be any fixed coefficients supported in an integer interval of diameter less than d . Orthogonality gives the exact identity

$$\sum_{r \bmod d} \left| \sum_n b_n e(rn/d) \right|^2 = d \sum_n |b_n|^2. \quad (7)$$

Completing successive periods of an arbitrary interval I_M , using nonnegativity for its incomplete period, gives

$$\sum_{m \in I_M} \left| \sum_n b_n e(mn/d) \right|^2 \leq (\lfloor |I_M|/d \rfloor + 1) d \sum_n |b_n|^2 \ll (M+d) \sum_n |b_n|^2. \quad (8)$$

For log-prime coefficients, Chebyshev's estimate gives $\sum_{N \leq p < 2N} (\log p)^2 \ll N \log N$. Hence

$$\sum_{m \in I_M} |P_N(m/d)| \ll \sqrt{M(M+d)N \log N}. \quad (9)$$

Here $M > d$, so this is $\ll M\sqrt{N \log N}$, with power $X^{4/5}$. No minor-arc theorem is needed for this average. For general outer coefficients replace (9) by

$$\left| \sum_m a_m P_N(m/d) \right| \ll \|a\|_2 \sqrt{(M+d)N \log N}. \quad (10)$$

This remains valid for every $(a, d) = 1$ in place of 1.

The actual terminal expression has a unit mask and a principal centering. Suppose d is squarefree and all supported n are coprime to d . Then the Ramanujan identity gives

$$\frac{1}{\phi(d)} \sum_{r \bmod d}^* \sum_n b_n e(rn/d) = \frac{\mu(d)}{\phi(d)} \sum_n b_n.$$

Subtracting this exact unit-average constant decreases the unit mean square:

$$\sum_{r \bmod d}^* \left| \sum_n b_n e(rn/d) - \frac{\mu(d)}{\phi(d)} \sum_n b_n \right|^2 = \sum_{r \bmod d}^* \left| \sum_n b_n e(rn/d) \right|^2 - \frac{|\sum_n b_n|^2}{\phi(d)} \leq d \sum_n |b_n|^2. \quad (11)$$

Thus (8) – (10) also control the centered terminal sum restricted to $(m, d) = 1$, without an added centering debt. In the current canonical conductor family every prime factor of d is $\leq X^{0.2615} < N$, so the required unit condition holds for every inner prime.

This calculation is strictly about fixed inner coefficients/support. A smooth product weight may be transferred through an explicitly summable separable expansion, with the sum of its coefficient norms charged. It is not permission to absorb $e(amn/d)$ into arbitrary coefficients and then invoke Siegel – Walfisz. Arbitrary m -dependent phase absorption invalidates that inference.

For the top-conductor packet $Q/2 < d \leq Q$, its exact conductor coefficient is $M_d = \mu(d)/d$, the low reduced numerators satisfy $a \leq d/(16H)$, and $|S_V(a/d)| \ll_V H$. Summing (9) absolutely over these indices gives only

$$\sum_{d \approx Q} \frac{H}{d} \frac{d}{H} O(M \sqrt{N \log N}) \ll Q M \sqrt{N \log N} = X^{1.323} (\log X)^{1/2}. \quad (12)$$

The coarse enlarged-major-arc count (6), followed by the trivial prime-sum bound, instead gives $\ll QMR^2$ on those arcs. Neither bound is $O(X \log X)$. There is a genuine averaging improvement for each fixed conductor/frequency; it does not yet give the required full arithmetic pairing improvement.

4. The $q = 1$ core has an RH square-root remainder

Schoenfeld's *Sharper bounds for the Chebyshev functions $\theta(x)$ and $\psi(x)$. II*, Theorem 10, equation (6.3), printed page 337, proves under RH

$$|\theta(x) - x| < \frac{1}{8\pi} \sqrt{x} \log^2 x \quad (x \geq 599). \quad (13)$$

Primary source: [AMS original paper](#). The publisher PDF was downloaded directly and its first page read. Endpoint changes in the chosen half-open prime interval contribute at most $O(\log N)$, already absorbed below.

Let β be the representative of $am/d \pmod{1}$ with $|\beta| \leq C/N$, for fixed C . Partial summation of (13) gives

$$P_N(\beta) = J_N(\beta) + O((1 + N |\beta|) \sqrt{N} \log^2 N), \quad J_N(\beta) = \int_N^{2N} e(\beta u) du. \quad (14)$$

The error follows from two endpoint errors and $2\pi |\beta| \int_N^{2N} |\theta(u) - u| du$; this states the precise frequency uniformity. Fixed BV weights have the analogous variation cost.

For the centered terminal sum one can keep

$$J_N(\beta) - \frac{\mu(d)}{\phi(d)} P_N(0)$$

as the explicit main term, retaining the principal term exactly. Alternatively $P_N(0) = N + O(\sqrt{N} \log^2 N)$ by RH, with a smaller additional error after division by $\phi(d)$.

The number of core m 's is

$$\ll_C (M + d)(1/N + 1/d).$$

Combining this count with the same low-numerator and conductor weights as in (12), the total absolute remainder in (14) is at most

$$\ll_{C,V} Q(M + Q)\sqrt{N}(1/N + 1/Q) \log^2 N \ll_{C,V} \frac{QM}{\sqrt{N}} \log^2 X = X^{0.923} O_{C,V}(\log^2 X). \quad (15)$$

This is a valid power-saving extraction for this identifiable component, initially with $|a_m| \leq 1$. It does not bound the explicit major main term, does not cover the enlarged $q = 1$ core uniformly in a growing C without paying the factor in (14), and does not transfer automatically to small $q > 1$. Square-root prime-AP errors at those other rationals require corresponding Dirichlet-L information; ordinary zeta RH does not provide it.

5. Accepted result and remaining obligation

The exact minor-arc theorem is applicable after the enlarged rational neighborhoods (4) are used. The fixed-inner-interval average is stronger than its pointwise application here, and the true unit mask/principal centering cause no loss by (11). The $q = 1$ component admits the rigorously small remainder (15), with its main term retained.

The missing step is cancellation in the actual weighted sum over conductors, reduced numerators, and the extracted arithmetic main terms. The estimates above neither prove such cancellation nor a zeta pair-correlation improvement. No claim is made for a coefficient sequence with an absorbed oscillatory phase, for unspecified product-weight separation, for arbitrary outer divisor coefficients at unit cost, or for the whole unsmoothed arithmetic packet.

Current report 54: The signed kernel norm: exact main term and a large coherent CRT remainder

Collection: R13 — rational-core extraction and the signed CRT remainder.

Source: [research/dyson/round13/signed-kernel/SMOOTH_SIGNED_KERNEL_NORM.md](https://research.dyson/round13/signed-kernel/SMOOTH_SIGNED_KERNEL_NORM.md).

SHA-256: `1105564835c925b818daf7198186e77c4f0f1ad4ac1001ee1bb50c0f5c7544d9` . Git blob: `ed4749d4c65bbe050d4496963ebd61a6b1022bc3` . Original bytes: 13904.

The signed kernel norm: exact main term and a large coherent CRT remainder

Date: 2026-09-05. Status: ordinary proof draft, independent review pending. No stronger actual-prime upper bound is claimed. The finite-kernel identities and estimates here require no RH assumption; the comparison with the Round 11 prime-specific bound retains that bound's RH hypothesis.

This task derives the exact smooth-window main term for the unrestricted dual norm retained in Round 12, bounds its remainder at the currently available scale, and isolates a positive part of that actual signed remainder of size

$$>> Q^2 H / (\log X)^{696}, \quad Q = X^{.523},$$

on the full canonical source-supported family. Other signed terms may cancel it. Thus it is not a lower bound for the complete norm or a prime obstruction. It proves that the dangerous long-period CRT interactions really occur in the permitted support and require quantitative cancellation if one wants a smaller norm bound.

The currently proved upper bound remains

$$\sum_n W(n/X) |K(n)|^2 \ll_{(v,W)} (X+Q^2)^H \log^4 X. \quad (1)$$

This does not improve the Round 11 prime-specific small-arc estimate. Even an ideal evaluation by the full-period main term would generally be too large for unrestricted-coefficient Cauchy – Schwarz to replace that estimate.

1. Exact finite kernel and its mean

Fix X large, $X^{(1/6)} \leq H \leq X^{(2/7)}$, $Q = X^{(523/1000)}$, a squarefree family Q_X contained in $(\sqrt{X}, Q]$, and a fixed real v in $C_c^\infty(1,2)$. Write $v_h = v(h/H)$, and define

$$\begin{aligned} V_0 &= \sum_h v_h, \\ U_q &= \sum_{(h,q)=1} v_h, \\ B_q(n) &= \sum_{h \equiv n \pmod q} v_h, \end{aligned}$$

$$\begin{aligned} b_q &= U_q / \phi(q), \\ K(n) &= \sum_{q \in Q_X} \mu(q) [B_q(n) - b_q]. \end{aligned} \quad (2)$$

Take a fixed nonnegative nonzero W in $C_c^\infty(1, 3/2)$, with Fourier convention

$$\widehat{W}(t) = \int W(u) e(-tu) du, \quad w_0 = \widehat{W}(0) > 0.$$

The norm to be estimated is $N_W = \sum_n W(n/X) |K(n)|^2$. The coefficients in K are exactly those in Round 12. For a genuine prime n near X it is the actual completed residue discrepancy. For nonunit composite n it is the unrestricted Fourier-defined dual kernel; that distinction is not suppressed.

Define the full-period mean

$$M = \sum_q \mu(q) [V_0/q - b_q]. \quad (3)$$

The reduced-frequency representation is

$$\begin{aligned} K(n) &= M + \sum_{\beta} C_{\beta} e(n\beta), \\ C_{(a/d)} &= S_v(a/d) A_d, \\ A_d &= \sum_{(q \in Q_X, d|q)} \mu(q)/q, \\ S_v(\beta) &= \sum_h v_h e(-\beta h), \end{aligned} \quad (4)$$

over all distinct reduced fractions $2 \leq d \leq Q$, $1 \leq a \leq d$. In particular

$$M = -\sum_{\beta} C_{\beta} \mu(d)/\phi(d).$$

The Round 11 first-power shift estimate $\sum_a |S_v(a/d)| \ll_v d$ and $|A_d| \leq (1 + \log(Q/d))/d$ give

$$\begin{aligned} |M| &\ll_v \log^2(2Q), \\ C_2 &:= \sum_{\beta} |C_{\beta}|^2 \ll_v H \log^3(2Q). \end{aligned} \quad (5)$$

The primitive centering is therefore present. It is not assumed equal to the unrestricted mean V_0/q for each modulus.

2. Exact CRT covariance and main term

For q_1, q_2 let $g = \gcd(q_1, q_2)$, $L = \text{lcm}(q_1, q_2) = q_1 q_2 / g$. Simultaneous congruences $n = h_1 \pmod{q_1}$ and $n = h_2 \pmod{q_2}$ are solvable precisely when $h_1 = h_2 \pmod{g}$. When compatible, denote their unique residue modulo L by $r(h_1, h_2)$.

The smooth count of that progression is exactly, by Poisson summation,

$$\begin{aligned} \sum_{n=r \pmod{L}} W(n/X) \\ = X/L \sum_{k \in \mathbb{Z}} \widehat{W}(kX/L) e(kr/L). \end{aligned} \quad (6)$$

The zero mode is Xw_0/L . The covariance after subtracting the unrestricted means V_0/q_i has full-period value

$$\begin{aligned} \Gamma(q_1, q_2) \\ = [g \sum_{(h_1=h_2 \pmod{g})} v_{h_1} v_{h_2} - V_0^2] / (q_1 q_2). \end{aligned}$$

Let

$$R_v(g) = g \sum_{(h_1=h_2 \bmod g)} v_{h_1} v_{h_2} - v_0^2 \\ = \sum_{(r=1)^{(g-1)}} |S_v(r/g)|^2 \geq 0. \quad (7)$$

Grouping these frequencies by their reduced denominators yields the exact identity

$$\sum_{(q_1, q_2)} \mu(q_1) \mu(q_2) R_v(\gcd(q_1, q_2)) / (q_1 q_2) \\ = \sum_{\beta} |C_{\beta}|^2 = C^2. \quad (8)$$

The left side retains the actual signs and common-divisor compatibility; its individual terms cannot be treated as independent congruences. The identity follows alternatively from orthogonality over a common period. It is not an asymptotic assertion about the X -window.

Consequently the exact main term for N_W is

$$X w_0 (M^2 + C^2). \quad (9)$$

All nonzero modes in (6), and the analogous single-modulus modes of the centering terms, constitute the remainder. The latter single-modulus modes have period at most $Q = o(X)$, so their total is $O_A(X^{(-A)})$ for every fixed A , by smooth Poisson decay, with constants depending on finitely many derivatives of v, W . The same is true for the integer-period correction to $\sum_n W(n/X)$.

Thus

$$N_W = X w_0 (M^2 + C^2) + E_{\text{CRT}} + O_A(X^{(-A)}), \quad (10)$$

where the explicit signed remainder is

$$E_{\text{CRT}} = X \sum_{(q_1, q_2)} \mu(q_1) \mu(q_2) / L \\ \sum_{(k \neq 0)} \hat{W}(kX/L) \\ \sum_{(h_1=h_2 \bmod g)} v_{h_1} v_{h_2} e(k r(h_1, h_2)/L). \quad (11)$$

This is an equality with a controlled smooth-window error; no assertion that E_{CRT} is negligible is made.

3. Only the small–common–divisor remainder survives at power scales

Choose the fixed cutoff $G = X^{(1/10)}$. It lies strictly between $Q^2/X = X^{(23/500)}$ and the least $H = X^{(1/6)}$.

If $g \geq G$, then

$$L \leq Q^2/G = X^{(473/500)}, \quad X/L \geq X^{(27/500)}.$$

The nonzero Poisson sum in (6) is therefore $O_B((X/L)^{(1-B)})$. There are at most Q^2 modulus pairs, and the absolute total h_1, h_2 weight is $O_v(H^2)$. Choosing B sufficiently large for any requested A shows that the entire $g \geq G$ portion of (11) is $O_A(X^{(-A)})$. No cancellation or short-interval hypothesis is used for this removal.

The zero-mode covariance for $g < G$ is also negligible. Since $g/H \leq X^{(-1/15)}$, (12)'s shift estimate from Round 11 gives, for each fixed $J \geq 2$,

$$R_v(g) \ll_{(v,J)} H^2(g/H)^{(-2J)}.$$

Here every nonzero r/g has distance at least $1/g$ from the integers, and summing $r^{(-2J)}$ gives the displayed bound. Multiplying by X , using $\sum_{(q \leq Q)} 1/q \ll \log Q$, and choosing J large proves that the $g \leq G$ portion of the main covariance in (8) is $O_A(X^{(-A)})$. Thus the substantial full-period covariance comes from shared divisors larger than G , where the CRT averaging itself is harmless.

The problematic term is exactly

$$E_{\text{small}} = \text{the expression (11) restricted to } \gcd(q_1, q_2) < G,$$

and (10) remains true with E_{small} in place of E_{CRT} , up to $O_A(X^{(-A)})$. Small-gcd compatibility is not a reason to discard this term: it is precisely where the least common multiple can exceed X .

For comparison, the elementary finite-spacing bound on (4), applied on the support of W , proves

$$\begin{aligned} N_W &\ll_{(v,W)} (X+Q^2 \log(2Q)) C_2 + X M^2 \\ &\ll_{(v,W)} (X+Q^2) H \log^4 X. \end{aligned}$$

Combining this actual upper bound with (5), (9), and nonnegativity gives

$$|E_{\text{small}}| \ll_{(v,W)} (X+Q^2) H \log^4 X. \quad (12)$$

The current remainder bound exceeds the natural XH main scale by $Q^2/X = X^{.046}$, up to logarithms. The exact source root predicates have not supplied cancellation in (11).

4. The remainder contains actual small determinants

There is an equivalent and simpler frequency form of the remainder. By (4), smooth Poisson summation on n gives

$$\begin{aligned} E_{\text{small}} &= X \sum_{(\beta \neq \gamma)} C_{\beta} \overline{C_{\gamma}} \\ &\quad W_{\hat{}}(X(\gamma - \beta)) \\ &\quad + O_A(X^{(-A)}), \end{aligned} \quad (13)$$

where all frequency differences are represented in $(-1/2, 1/2]$; integer aliases away from this representative are negligible. Cross terms with M are negligible because every nonzero β has distance at least $1/Q$ from the integers. Equal frequencies give exactly the main C_2 . For reduced denominators d_1, d_2 , a common divisor at least G forces a nonzero frequency distance at least G/Q^2 , so rapid Fourier decay independently makes those terms negligible. The gcd of reduced denominators is not silently identified with the gcd of arbitrary original moduli in (11). On the isolated top conductors used in Section 5, however, $d_i = q_i$ exactly, so the identification there is legitimate.

To make the dangerous arithmetic ranges explicit, fix any small $\epsilon > 0$, for example $\epsilon = 1/100$. By rapid decay of $W_{\hat{}}$ and S_v , one may restrict (13), at a cost $O_A(X^{(-A)})$, to nonzero signed numerators r_i with

$$\begin{aligned} \beta &= r_1/d_1, \quad \gamma = r_2/d_2, \\ 2 \leq d_i \leq Q, \quad \gcd(r_i, d_i) &= 1, \end{aligned}$$

$$\begin{aligned} 0 < |r_i| &\leq X^{\epsilon} d_i / H, \\ 0 < |r_2 d_1 - r_1 d_2| &\leq X^{\epsilon} d_1 d_2 / X. \end{aligned} \quad (14)$$

The residues here are centered near the integers. Bounds are uniform after choosing sufficiently many fixed derivatives to absorb the polynomial number of terms. There are no terms with exactly one zero numerator in these ranges: its nonzero partner would have magnitude at least $1/Q$, which exceeds X^{ϵ}/X . The determinant is divisible by $\gcd(d_1, d_2)$; its small nonzero size is the CRT resonance.

This is not a source-class distribution conclusion. The allowed-support predicates constrain prime factors of each modulus separately. They do not state cancellation in the signed determinant pairing (13) – (14). In fact the next section shows that such pairs are forced inside the full canonical support.

5. A positive coherent part of the actual signed Gram has the large scale

For this lower bound on a **part** of the remainder, specialize to the full canonical family and a fixed nonnegative nonzero v . Let $m_v = \text{integral } v$. The frozen Round 11 – 12 arithmetic construction supplies a set Ω_X of actual reduced frequencies such that

$$\begin{aligned} \Omega_X &\subset (0, 1/(16H)], \\ |\Omega_X| &\geq c_0 Q^2 / [128H(\log X)^{348}], \\ d &\in (Q/2, Q], A_d = 1/d, \\ |C_{\beta}| &\geq m_v H / (2\sqrt{2}Q), \\ \arg C_{\beta} &\in [-\pi/4, 0]. \end{aligned} \quad (15)$$

The phase statement follows directly from $v_h \geq 0$ and $H < h < 2H$: each term in $S_v(\beta)$ lies in that sector. The identity $A_d = 1/d$ holds in the full signed family, so it is not changed by other negative-Mobius moduli.

Partition the small arc into at most $J \leq 8X/H$ cells of length at most $1/(100X)$, and let m_j be their actual occupations. Since $|\Omega_X|/J$ tends to infinity,

$$\begin{aligned} \sum_j m_j(m_j - 1) &\geq |\Omega_X|^2/J - |\Omega_X| \\ &\geq c_0^2 Q^4 / [262144 XH(\log X)^{696}] \end{aligned} \quad (16)$$

for all sufficiently large X , uniformly in H . This counts ordered pairs of distinct actual frequencies in the same cell.

On each such pair, $|X(\gamma - \beta)| \leq 1/100$. In the integral for $\hat{W}_u$, u lies in $(1, 3/2)$, so the added phase has magnitude at most $3\pi/100$. Together with the coefficient phase difference at most $\pi/4$, the total is at most $7\pi/25 < \pi/3$. Consequently

$$\begin{aligned} \text{Re}[X C_{\beta} \overline{\text{conjugate}(C_{\gamma})} \\ \quad \hat{W}_{\hat{u}}(X(\gamma - \beta))] \\ &\geq X w_0 |C_{\beta} C_{\gamma}|/2 \\ &\geq X w_0 m_v^2 H^2 / (16Q^2). \end{aligned} \quad (17)$$

Let E_{coherent} be exactly the sub-sum of (13) over these ordered within-cell pairs, without any other terms. Equations (16) – (17) prove

$$E_{\text{coherent}} \geq c_0^2 w_0 m_v^2 Q^2 H$$

$$/[4194304(\log X)^{696}]. \quad (18)$$

This is an actual contribution with the actual merged coefficient signs; it is not obtained by replacing the full remainder by absolute values. Its definition is symmetric in β, γ and it is real.

Moreover, distinct fractions with denominators d_1, d_2 satisfy

$$|\beta - \gamma| \geq \gcd(d_1, d_2) / (d_1 d_2).$$

Every pair used in (18) therefore has

$$\gcd(d_1, d_2) \leq Q^2 / (100X) < G. \quad (19)$$

Its reduced CRT period is at least $100X$. Thus these are precisely nonzero long-period modes in the small-gcd region, not harmless diagonal or large-gcd contributions. The root support cannot make this region empty.

Equation (18) is not a lower bound for E_{small} or N_W . All remaining signed pairs define $E_{\text{other}} = E_{\text{small}} - E_{\text{coherent}}$ and may cancel it. For example, a norm estimate $N_W < X^H \log^A X$ would force

$$E_{\text{other}} = -E_{\text{coherent}} + O_{\nu}(v, W, A) (X^H \log^{\max(A, 4)} X), \quad (20)$$

up to the negligible errors already stated. Since $Q^2/X = X^{.046}$ dominates $\log^{(696+A)} X$, this would be cancellation at the scale of the coherent block. We have identified and quantified the missing cancellation, not proved that it is unavailable.

6. Why the unrestricted norm may be the wrong target

The full-period coefficient norm in (9) satisfies $C_2 \gg_{\nu} H / (\log X)^{348}$ on the canonical family, by Round 11, and $C_2 \ll_{\nu} H \log^3 X$. If one could show that the smooth-window remainder is small compared with this main term, the resulting unrestricted norm would therefore have size X^H up to logarithmic powers.

Ordinary Cauchy – Schwarz against the genuine-prime coefficient vector, whose squared norm is $O_f(X \log X)$, would then only give $X^{\sqrt{H}}$ times logarithms. This is already worse in the present H range than the Round 11 prime-specific small-arc bound $X^{1.023} \log^5 X$. Reaching $o(X \log X)$ by unrestricted-coefficient Cauchy – Schwarz would instead require $N_W = o(X \log X)$, subject to matching a window majorizing the coefficient support. That is far below the natural full-period main scale.

This is a conditional scale comparison, not a proved lower bound for N_W : (11) may have negative signed contributions. It explains why evaluating the dual norm alone need not attack the desired prime correlation. A stronger argument may need the alignment of the actual prime vector with K , rather than its largest value on all coefficient vectors.

The bounded result is (10) – (12) together with the explicit actual coherent block (18). There is no new power-saving upper bound. The remaining obligation is a signed small-determinant/CRT correlation estimate, or a prime-specific pairing estimate that avoids the unrestricted norm. No root factorization, common-divisor independence or Mobius cancellation was assumed without proof.

7. Verification scope

The companion exact-arithmetic check verifies the CRT mean/covariance identities and finite-window remainder on one fixed small squarefree family, with all signs and primitive subtractions retained. It also checks the rational cutoff exponents and the coherent-block counting constants. These are algebraic checks, not a numerical realization of the asymptotic modulus family or evidence for the required cancellation. No parameter scan, numerical prime experiment, or previous-round edit was performed.

Current report 55: Independent audit of the smooth signed–kernel norm

Collection: R13 — rational-core extraction and the signed CRT remainder.

Source: `research/dyson/round13/signed-kernel/INDEPENDENT_AUDIT.md`.

SHA-256: `5ada6ceefda4f1df6435a992876ed1eee5fe567907c183fd3a6e193814c0a731` . **Git blob:** `32c316e4209a8baa431b2e52ef5246dc65e894` . **Original bytes:** 6833.

Independent audit of the smooth signed–kernel norm

Date: 2026-09-05. **Reviewer:** Astra subagent `yau_flow` . **Verdict:** accepted as an ordinary mathematical derivation, within its explicitly limited scope. No correction to the author report is required by this review. In particular, the positive coherent block is not a lower bound for the complete signed remainder or the window norm.

Reviewed report: `SMOOTH_SIGNED_KERNEL_NORM.md` , SHA256

`1105564835c925b818daf7198186e77c4f0f1ad4ac1001ee1bb50c0f5c7544d9` . All author-manifest hashes and all five inherited source-receipt hashes were checked against the current files and matched. Author files and the manifest were preserved.

Exact decomposition and centering

The unrestricted mean of B_q is V_0/q , whereas its primitive subtraction is $b_q = U_q/\phi(q)$. Their difference is retained in M ; replacing one by the other would change the kernel. Fourier completion, followed by reducing each fraction, yields the stated $C_{a/d} = S_{\chi}(a/d) \sum_{q:d|q} \mu(q)/q$. The Ramanujan ratio $c_q(r)/\phi(q) = \mu(d)/\phi(d)$, with $d = q/(q, r)$, verifies the separate displayed identity for M .

For each pair of moduli, the CRT compatibility condition $h_1 \equiv h_2 \pmod{g}$ and progression period $L = q_1 q_2/g$ are exact. With the specified negative-sign Fourier transform, the Poisson factor is correctly $\widehat{W}(kX/L)\vartheta(kr/L)$.

Expanding the primitive subtractions and then collecting zero modes gives exactly

$$Xw_0(M^2 + C_2).$$

Finite Fourier orthogonality gives

$$R_{\chi}(g) = g \sum_{h_1 \equiv h_2 \pmod{g}} v_{h_1} v_{h_2} - V_0^2 = \sum_{r=1}^{g-1} |S_{\chi}(r/g)|^2,$$

and summing with the signed modulus coefficients gives C_2 , as stated. Real v is sufficient here; nonnegative v is only imposed later for the positive block.

The nonzero single-modulus and constant-window terms have periods at most $Q = X^{523}$ or 1. Their polynomially bounded total coefficients can be absorbed by arbitrarily rapid decay of the fixed smooth window, so the $O_A(X^{-A})$ remainder in equation (10) is justified. The principal terms have not disappeared by assumption.

The two common-divisor ranges

For $G = X^{1/10}$, a common divisor $g \geq G$ implies

$$L \leq X^{473/500}, \quad X/L \geq X^{27/500}.$$

The progression's nonzero Poisson part is $O_B((X/L)^{1-B})$. There are at most Q^2 pairs and $O_V(H^2)$ total absolute shift weight per pair. Thus any prescribed negative power follows by taking a sufficiently large fixed number of derivatives; no arithmetic cancellation is needed.

For $g < G$, smoothness of the shift profile gives

$$R_V(g) \ll_{v,J} H^2 (g/H)^{2J}.$$

This follows by summing the rapidly decaying shift transform at the nonzero grid points, whose distances from the nearest integer are at least $1/g$. Since $G/H \leq X^{-1/15}$, the corresponding zero-mode covariance is also negligible after the reciprocal-modulus sum and the factor X are charged. These statements remove different pieces: the large- g Poisson remainder and the small- g full-period covariance. They do not remove the small- g Poisson remainder.

The frequency form (13) has the correct sign $\widehat{W}(X(\gamma - \beta))$. Integer aliases away from the closest representative, and cross terms with M , have polynomial separation on the X -scale and are negligible. The author explicitly distinguishes common divisors of reduced denominators from those of arbitrary original moduli. For the selected top conductors these objects coincide because no larger multiple fits below Q .

Coherent-block constants and what they imply

The inherited actual-support count in equation (15) was checked against the frozen Round 12 construction. Its isolation identity $A_d = 1/d$ holds in the full signed family, rather than only after deleting negative coefficients. All selected coefficient phases lie in $[-\pi/4, 0]$.

Writing $n_\Omega = |\Omega_X|$ and using at most $J \leq 8X/H$ cells,

$$\sum_j m_j(m_j - 1) \geq n_\Omega^2/J - n_\Omega \geq n_\Omega^2/(2J)$$

eventually, uniformly over the stated H -range. The lower bound for n_Ω therefore gives the exact displayed denominator $128^2 \cdot 8 \cdot 2 = 262144$.

Within a cell, the additional window-integrand phase is at most $3\pi/100$, since W is supported in $(1, 3/2)$. Including the coefficient phase difference gives

$$\pi/4 + 3\pi/100 = 7\pi/25 < \pi/3.$$

Nonnegativity of W therefore yields the stated half-product lower bound. The coefficient-product bound contributes the further denominator 16, giving 4194304 in (18). The selected ordered pairs form a conjugation-symmetric set, so their sum is real.

Every selected distinct pair satisfies

$$\frac{\gcd(d_1, d_2)}{d_1 d_2} \leq |\beta - \gamma| \leq \frac{1}{100X}.$$

Thus its common divisor is at most $Q^2/(100X) < G$, and its reduced CRT period is at least $100X$. The block really lies in the long-period small-common-divisor region. No equidistribution or independence of the moduli is assumed to create it.

Equation (18) bounds this specified subsum only. Remaining pairs can have negative real part and can cancel it. Equations (20) and the final unrestricted-Cauchy – Schwarz comparison retain this distinction correctly: a smaller full norm would demand cancellation, but the report does not prove that such cancellation is impossible. It also does not identify an arbitrary coefficient-space extremizer with the actual prime vector.

Script inspection and independent replay

I inspected `check_signed_kernel_norm.py` before execution. It uses exact `Fraction` arithmetic for one fixed squarefree family, checks 100 CRT compatibility cases, verifies the complete-period mean and variance, and expands the finite-window norm into its pair, single-centering, and constant-window remainders. It retains the Mobius signs and primitive subtractions. Its polynomial toy window has integral $1/960$; the script explicitly does not use that window to test the rapid-decay theorem.

A temporary copy was replayed, so the author JSON was not rewritten. The regenerated JSON was byte-identical to the frozen author result. In particular:

- full-period mean: $-43/280$;
- full-period variance: $22843/14700$;
- finite-window norm: $14619643/165888000$;
- cutoff exponents and both coherent-block constants: passed exactly.

Script SHA256: `d661b1ef764f5ab395a9ca2db66ae9387d6885426772a93b595a43462fdc61d0` . Result SHA256: `9f92dc12f0ac555e42af962478288e523593cf299c177fa93727d408373fdb5a` .

These checks test finite algebra, not the asymptotic prime-family count, a numerical cancellation claim, or an improved zeta correlation estimate. The accepted result is the exact main/remainder decomposition, its stated upper bound, and the positive coherent subsum with its explicit limitation.

Current report 56: Round 13 independent integration review

Collection: R13 — rational-core extraction and the signed CRT remainder.

Source: `research/dyson/round13/INDEPENDENT_ROOT_REVIEW.md`.

SHA-256: `d38cfc0ec3f7031276340e4984fbc330185b5c9eef34df1e0955fc9853555456` . **Git blob:** `80d1bedb45164cf0b9f944602560da77f374ffda` . **Original bytes:** 4599.

Round 13 independent integration review

Date: 2026-09-05. **Reviewer:** root Astra. This review accepts bounded analytic statements, not a new theorem about zeta pair correlation.

Phase averaging

Reviewed `phase-resonance/AVERAGED_RATIONAL_PHASE_TEST.md` , SHA-256

`7f4285cb02241e22bdb29a1ad4952f7ab8249e3ec3bef984455a57ae05e41ebb` , and its complete exact-check script.

The fixed inner prime interval, the mask $(m,d)=1$, and the primitive Ramanujan principal term are essential and are retained.

At $Q=X^{.523}$, $M=X^{.6}$, $N=X^{.4}$, the phase grid count on a fixed C/N core is $O(M/N)$: the possible $O(1)$ endpoint residues are absorbed because d/N tends to infinity, and repeated residue classes cost M/d because M/d tends to infinity. RH partial summation gives an error $O(\sqrt{N}\log^2 N)$ uniformly on that core. The actual numerator weight has total absolute mass $O(1)$ for each d . Hence the sum of errors is $O(QM/\sqrt{N}\log^2 X)=O(X^{.923}\log^2 X)$. The same argument allows fixed divisor-bounded outer coefficients with the stated X^η loss. This is a valid extraction of a retained integral main term.

The positive witness uses actual terminal conductors with even prime-factor count, not invented Fourier support. Unit residues s up to $d/(32N)$, repeated m classes, low unit numerators, and the family count give respectively the factors $1/128$, $1/32$, and $1/2$. The centered prime sum paired with S has real part at least $(\text{integral } V)HN/(16d)$. Their product is exactly $1/131072$. This proves the stated restricted positive block only. Other phases, signs and actual outer coefficients can cancel it.

The enlarged major arcs of width $2R/(qN)$ correctly handle $\text{floor}(N/R)$ in Dirichlet approximation. The complement of fixed C/N cores would not suffice. Ordinary SW yields only a fixed logarithmic error for the other small denominators; ordinary zeta RH has not been mistaken for Dirichlet-L RH. The centered unit variance is exact because each inner prime is a unit and the unit mean is $\mu(d)P_N(0)/\phi(d)$. Its completion bound gives $X^{1.323}\sqrt{\log X}$, explicitly weaker than the existing estimate for the original prime pairing. No arithmetic improvement is attributed to that factored bound.

Signed smooth kernel

Reviewed `signed-kernel/SMOOTH_SIGNED_KERNEL_NORM.md`, SHA-256

`1105564835c925b818daf7198186e77c4f0f1ad4ac1001ee1bb50c0f5c7544d9`. Aquinas's separate `INDEPENDENT_AUDIT.md` provides an additional analytic review and byte-identical exact replay.

The CRT main term is $X(\int W)(\text{mean}^2 + \text{coefficient squared mass})$, with the nonzero CRT modes explicitly retained. The cutoff $\gcd(q_1, q_2) \geq X^{.1}$ makes the least common multiple at most $X^{.946}$; smooth Poisson decay is therefore available. For the complementary gcd range, the zero-mode covariance is small because $X^{.1}/H$ tends to zero. These observations localize the difficulty without estimating away its signed remainder.

The coherent block constant $1/4194304$ includes ordered distinct pairs, cell occupancy, coefficient phases and the window phase. Its lower bound is for an explicitly selected subsum of the off-diagonal expression. The report correctly does not infer a lower bound for the full remainder or norm. Its additional observation that even an ideal generic kernel norm does not by itself match the stronger prime-specific small-arc estimate is retained.

Primary-source audit

Reviewed `minor-arc-source/MINOR_ARC_AND_FIXED_INTERVAL_AUDIT.md`, SHA-256

`bbdb17478b9e885570b7b49c3ff9b94b0ceb98f12a41743edb1ce492cb50edc4`. It checks Montgomery--Vaughan Theorem 17.1, the interval endpoint and genuine-prime changes, the rational-arc domain and Schoenfeld's RH theta estimate. The complete author-hosted/AMS papers and page images remain in the local reference archive; public receipts retain their URLs and hashes. Root's review checks the application and does not claim a second independent reconstruction of those classical source theorems.

Acceptance and remaining obligation

Accepted as ordinary proofs with their explicit assumptions: the $q=1$ extraction error, restricted positive resonance block, exact CRT decomposition, coherent subsum and stated upper estimates. Finite tests validate algebra, constants and source identity; they do not numerically prove the asymptotic inequalities.

The strongest bound for the original selected smooth prime discrepancy remains $O(X^{1.023} \log^5 X)$ under RH. The strict signed covariance estimate required for the actual-zeta Dyson--Montgomery target is unproved. No AH refutation, zeta-gap improvement, GUE theorem, RH proof or prime-gap improvement follows from this round.

Current report 57: Dyson—Montgomery round 14: an actual Type I removal and a quantitative CUE heat theorem

Collection: R14 — smooth Type I removal and quantitative finite CUE heat.

Source: research/reports/dyson_round14.md.

SHA-256: `2ec5f48334de9a85e0a481c3bc7fab33f9398afc957a019db8f62ee2bfb9140` . Git blob: `794f70e45c9832711de07ad8bab3d48ad4db47c` . Original bytes: 6653.

Dyson—Montgomery round 14: an actual Type I removal and a quantitative CUE heat theorem

Date: 2026-09-05. Two bounded results now have complete ordinary proofs, independent reviews and reproducible checks. One removes an exact portion of the actual arithmetic discrepancy. The other strengthens the finite CUE scalar-heat comparison. Neither proves a theorem about the full zeta pair correlation.

1. A specified arithmetic component is below the required scale

For the exact smooth discrepancy already defined in Rounds 9--13, retain its original sinc kernel, both logarithmic weights, all permitted moduli $q \leq Q = X^{.523}$ and the primitive principal subtraction. Split the von Mangoldt coefficient by the exact identity

$$\Lambda = \Lambda_{\leq U} + \Lambda_{>U}, \quad \Lambda_{\leq U}(n) = \sum_{r \mid n, r \leq U} \mu(r) \log(n/r).$$

The [complete unconditional proof](#) establishes

$$|D_Q^V[\Lambda_{\leq U}]| \ll_{J,V,X} HX(UQ/X)^J \log^2 X \quad (J \geq 2, UQ \leq X/2).$$

For $U = X^{.4}$, $J = 4$ and $H \leq X^{(2/7)}$, this is $O(X^{(1711/1750)} \log^2 X) = o(X \log X)$. More generally, for every fixed $0 < \eta < .477$, all divisors below $U \leq X^{(.477-\eta)}$ are covered by choosing a fixed J with $J\eta > 2/7$. The constants may depend on η and J . No uniform limit as η tends to zero is claimed.

The reason is exact progression Poisson summation in the smooth long cofactor n/r . Its zero mode cancels the actual primitive principal mean, and its nonzero modes decay because that cofactor exceeds q by a fixed power. The original joint kernel has uniform derivatives after the displayed integral representation removes its apparent singular phase. No smoothness is assigned to the shorter Möbius coefficients, and no dense-divisibility or RH estimate is needed for this portion.

This is a classical method applied and checked for the programme's actual kernel, not a novelty claim for Poisson summation. It also explains constructively how the large positive restricted rational cores from Round 13 can cancel when the actual longer factor is smooth. The signed $\Lambda_{>U}$ discrepancy remains exactly in the formula and

is unestimated. Its cofactor need not be large or balanced. The full original discrepancy still has only the previously recorded RH bound, above $X \log X$.

The [independent Type I review](#) checks every primitive mask, normalized derivative, frequency sign, summation factor and the precise individual-variable criterion for a Heath–Brown component.

2. A finite CUE heat–flow approximation has an explicit error scale

Let $\delta_{\min}$ be the minimum circular angular gap of Haar CUE(N). At its midpoint define

$$B_N = \sum_{k \text{ outside the pair}} \frac{1}{4 \sin^2((\theta_k - c)/2)}.$$

The [complete CUE proof](#) establishes the finite-N estimate

$$\mathbb{E} \sum_{i: \delta_i \leq \varepsilon} B_i \leq N^6 \varepsilon^3 / 18 \quad (0 < \varepsilon \leq \pi).$$

The proof uses the exact finite CUE three-point Gram determinant. Its simultaneous short-pair and endpoint vanishing factors control the singular inverse-square weights. Endpoint weights are used before enlarging to nonconsecutive pairs; directly enlarging a midpoint-weighted sum would be invalid. Circular ordering includes the wrap gap.

The classical minimum-gap law and Markov then give $B_N/N^2 = O_{\mathbb{P}}(1)$, without assigning a conditional density to the selected pair. For each fixed $L, K > 0$ the proof supplies the explicit bound

$$\limsup_N \Pr(B_N/N^2 > K) \leq e^{-L^3/(72\pi)} + L^3/(18K).$$

For the finite scalar-heat polynomial

$$P_{N,s}(z) = \sum_{j=0}^N a_j e^{sj(N-j)} z^j, \quad P_N(z) = \det(zI - U_N),$$

let D_N be the first positive discriminant time. Applying the already quantified Galilean lemma with its verified uniform constants yields

$$\frac{8D_N}{\delta_{\min}^2} - 1 = O_{\mathbb{P}}(N^{-2/3}), \quad D_N - \delta_{\min}^2/8 = O_{\mathbb{P}}(N^{-10/3}).$$

This strengthens the programme's prior qualitative CUE comparison. It is an approximation error in probability, not a convergence rate for the limiting depth distribution. It is not a stochastic Dyson Brownian motion result or an available identity for true zeta zeros. General-beta analogues and global novelty remain outside this proof.

The [independent CUE review](#) and [root review of both results](#) accept the proofs with these restrictions.

3. Preserved evidence

Seventeen original files totaling 548,013 bytes are preserved in the adjacent local `Astra-Local-Archive/round14-
originals`. Fifteen research, review and receipt files are public and verbatim. The full Feng--Wei PDF/text remain local with public source hashes and precise source locations.

The `separate-process replay` passes both bounded scripts. It checks the symbolic $N=3$ CUE determinant and singular-weight cancellation, the constant $1/18$, 63 complete-period centering cases, exact divisor/HB identities through $n=125$, rational exponent arithmetic and two fixed floating Gaussian Poisson diagnostics. One temporary source path is excluded from the Type I certificate comparison; the CUE JSON matches every field. These tests supplement ordinary proofs; they do not replace them or provide interval-certified stochastic estimates.

The integration receipt also records the earlier Galilean publication edit: only its title and reviewer attribution differ from the original hash cited by the CUE author. Root compared the complete mathematical bodies and found them identical. Both hashes remain visible; the original proof provenance is not silently rewritten.

4. What is still needed for the famous-conjecture target

The actual-zeta route still requires a signed estimate for the retained arithmetic remainder, with all covariance cross terms and principal means present. The CUE theorem supplies a rigorous RMT reference statement and exposes the missing zeta input: isolated close pairs and their background control cannot be assumed from low-band correlation data.

The sufficient two-scale W_T lower limit $1/16$, compact Fourier test above $7/10$, positive-density violation of AH-Pairs, sub-half normalized zeta gap and sub-186 prime gap remain unproved. No famous conjecture is reported solved. Broad numerical sweeps, another Fable session and a claim of full arithmetic-to-Fock convergence are postponed; the next arithmetic work concerns the exact remaining signed terms.

Current report 58: A smooth long factor removes an actual Type I component

Collection: R14 — smooth Type I removal and quantitative finite CUE heat.

Source: `research/dyson/round14/smooth-long-factor/SMOOTH_LONG_FACTOR_REMOVAL.md`.

SHA-256: `d6143f19ddf006a1acc833ecd2e5265bffb35817930cfeaa4f4e4b973af7c849` . **Git blob:** `97e25b8fc8488e5cd5c538fceddb450d42dfd227` . **Original bytes:** 13229.

A smooth long factor removes an actual Type I component

Date: 2026-09-05. Status: unconditional ordinary proof; a classical Poisson application, newly verified for this programme's exact kernel. No novelty claim is made for the method. No estimate for the remaining component or the whole zeta covariance is proved.

Define the exact truncated divisor convolution

$$\Lambda_{\leq U}(n) = \sum_{\substack{r \mid n \\ r \leq U}} \mu(r) \log(n/r). \quad (1)$$

For the actual smooth discrepancy of Rounds 9 – 13, with $Q = X^{523/1000}$, $H = X/T$, and $X^{1/6} \leq H \leq X^{2/7}$, this note proves

$$|D_Q^V[\Lambda_{\leq U}]| \ll_{J,V,X} HX \left(\frac{UQ}{X}\right)^J (\log X)^2 \quad (J \geq 2, UQ \leq X/2). \quad (2)$$

Both primitive restrictions, the modulus coefficients and the entire original covariance kernel are retained. Neither RH nor dense divisibility is needed for this component.

In particular $U = X^{2/5}$, $J = 4$ give

$$|D_Q^V[\Lambda_{\leq X^{2/5}}]| \ll X^{1711/1750} (\log X)^2 = o(X \log X). \quad (3)$$

The exponent margin below X is 39/1750. The remaining divisor portion has signs; the truncation is an exact decomposition, not an inequality for the full discrepancy.

1. The exact arithmetic object

Use the frozen weights

$$\begin{aligned} a_Y(X) &= \min\{(Y/X)^{1/2}, (X/Y)^{3/2}\}, \\ w_h(u) &= \chi(u/X) a_u(X) a_{u+h}(X) \operatorname{sinc}_0(T \log(1 + h/u)), \end{aligned} \quad (4)$$

where $\text{sinc}_0(t) = \sin(t)/t$, with its removable value at zero, $V \in C_c^\infty(1, 2)$, $\chi \in C_c^\infty(1, 3/2)$, $X = T^a$, $6/5 \leq a \leq 7/5$, and $H = X/T$. For any coefficient sequence b , define

$$D_Q^V[b] = \sum_h V(h/H) \sum_{\substack{q \in Q_X \\ (h,q)=1}} \mu(q) \left[\sum_{n \equiv h(q)} b(n) w_h(n-h) \log((n-h)/q) \right. \\ \left. - \frac{1}{\phi(q)} \sum_{(n,q)=1} b(n) w_h(n-h) \log((n-h)/q) \right]. \quad (5)$$

Weights vanish outside their indicated support. Thus $n \asymp X$ and $n-h > X$, so all logarithms have positive arguments. This is the actual R9/R11 component, including its log-cofactor weight.

The family Q_X may be the full canonical complementary family or any selected moduli $q \leq Q$. We use only the cutoff and $|\mu(q)| \leq 1$. An added $(\log q)^j$ weight, for a fixed nonnegative integer j , changes $(\log X)^2$ in (2) to $(\log X)^{2+j}$; it does not affect the power margin.

For all positive integers, including $n = 1$, the identity $\Lambda = \mu * \log$ is exact. At 1 both sides are zero. For a real $U \geq 1$, $r \leq U$ means the usual integer cutoff, with its endpoint included. Consequently

$$D_Q^V[\Lambda] = D_Q^V[\Lambda_{\leq U}] + D_Q^V[\Lambda_{>U}], \quad \Lambda_{>U}(n) = \sum_{\substack{r|n \\ r>U}} \mu(r) \log(n/r). \quad (6)$$

No prime-power removal or asymptotic main-term substitution occurs here. Neither summand need be nonnegative.

2. Why a smooth longer coefficient cancels the complete phase mean

Let $q \geq 2$, let $a \bmod q$ be arbitrary, and fix inner coefficients β_n independently of m . Define the periodic function

$$R_{q,a}(m) = 1_{(m,q)=1} \sum_{\substack{n \\ (n,q)=1}} \beta_n \left[\alpha(amn/q) - \frac{c_q(a)}{\phi(q)} \right], \quad (7)$$

where $\alpha(t) = e^{2\pi i t}$ and c_q denotes the Ramanujan sum. Multiplication by a unit n permutes the unit residues, so $c_q(an) = c_q(a)$. It follows exactly that

$$\sum_{m \bmod q} R_{q,a}(m) = 0. \quad (8)$$

This includes nonprimitive a , and does not require squarefree q . For $a = 0$, the function itself vanishes.

Suppose A_X is smooth on a fixed positive compact support, and every fixed derivative has a uniform fixed-power logarithmic bound. Use the Fourier conventions

$$\widehat{R}(\ell) = q^{-1} \sum_{m \bmod q} R(m) \alpha(-\ell m/q), \quad \widehat{A}(t) = \int_{\mathbb{R}} A(u) \alpha(-tu) du.$$

Fourier expansion and Poisson summation give

$$\sum_m A_X(m/M) R_{q,a}(m) = M \sum_{v \in \mathbb{Z}} \widehat{R}(-v \bmod q) \widehat{A}_X(Mv/q). \quad (9)$$

The omitted $\nu = 0$, and all nonzero multiples of q , have coefficient zero by (8). Writing $B = \sum_n |\beta_n|$, we have $|\hat{R}(\ell)| \leq 2B$. For $M \geq q$, integration by parts $J \geq 2$ times therefore gives

$$\left| \sum_m A_X(m/M) R_{q,a}(m) \right| \ll_J MB(\log X)^{C_J} (q/M)^J. \quad (10)$$

This does not assign a Siegel – Walfisz property to the twisted sequence.

The actual shift Fourier weights satisfy, by finite summation by parts,

$$|S_{V,H}(\beta)| \ll_{V,K} H(1 + H \|\beta\|)^{-K}, \quad \sum_{a=1}^{q-1} |S_{V,H}(a/q)| \ll_V q. \quad (11)$$

Hence the full numerator sum with its $1/q$ coefficient costs $O_V(1)$ per modulus. Its zero numerator vanishes by (8); there is no hidden $S_{V,H}(0) \asymp H$ term. At $M = X^6$, $q \leq Q = X^{.523}$, the ratio $q/M \leq X^{-.077}$ permits any fixed power saving by choosing J .

For a nonnegative smooth longer profile that is positive on an interval, the restricted positive rational core from R13 still has a large contribution by the same residue count on that interval. Equations (8) – (10) give a constructive reason why other phases cancel it in the full sum. They make no such claim for a nonsmooth or arithmetically oscillating longer coefficient.

3. A progression identity retaining the whole actual kernel

Expand (1) inside (5), writing $n = rs$. If $(r, q) > 1$, the progression $rs \equiv h \pmod{q}$ is empty because $(h, q) = 1$; its primitive principal sum is empty as well. Such terms vanish exactly.

For $(r, q) = 1$, put $b \equiv h\bar{r} \pmod{q}$, a unit residue. The longer variable s has scale $L_r = X/r$ and weight

$$F_{r,h,q}(s) = (\log s) w_h(rs - h) \log((rs - h)/q). \quad (12)$$

Its coefficient outside is $\mu(r)$. Write $F_{r,h,q}(s) = \Phi_{r,h,q}(s/L_r)$. With $\delta = H/X = 1/T$, $z = h/H$, and $u = s/L_r$, the profile is exactly

$$\begin{aligned} \Phi_{r,h,q}(u) = & \chi(u - \delta z)(u - \delta z)^{-3/2} u^{-3/2} \operatorname{sinc}_0(\delta^{-1} \log \frac{u}{u - \delta z}) \\ & \times [\log(X/r) + \log u] [\log(X/q) + \log(u - \delta z)]. \end{aligned} \quad (13)$$

On this support $u - \delta z > 1$ and $u > 1$, so the chosen branches of both a -factors are correct. The apparently singular phase has the regular form

$$\delta^{-1} \log \frac{u}{u - \delta z} = \int_0^z \frac{dt}{u - \delta t}. \quad (14)$$

The fixed compact support of χ , and z in the fixed support of V , imply that these profiles lie in a fixed positive compact interval and satisfy

$$\sup_{r \leq U, q \leq Q, h} \sum_{j=0}^J \|\Phi_{r,h,q}^{(j)}\|_{\infty} \ll_{J,X} (\log X)^2 \quad (15)$$

for every fixed J . Zero extension outside the support is smooth because χ is compactly supported in the open interval. All bounds are uniform in the displayed parameters, including T . No separated substitute for the actual kernel was used.

For any such profile, $L > 0$, and a unit $b \bmod q$, Poisson summation gives

$$\begin{aligned} \sum_{s \equiv b(q)} \Phi(s/L) - \frac{1}{\phi(q)} \sum_{(s,q)=1} \Phi(s/L) \\ = \frac{L}{q} \sum_{k \in \mathbb{Z}} \widehat{\Phi}(kL/q) \left[\alpha(kb/q) - \frac{c_q(k)}{\phi(q)} \right]. \end{aligned} \quad (16)$$

The Fourier convention is that in (9). Its $k = 0$ term cancels since $c_q(0) = \phi(q)$. This is finite progression summation of a smooth function; there are no contour or pole terms. Using $|c_q(k)| \leq \phi(q)$, followed by J integrations by parts, proves

$$|\text{left side of (16)}| \ll_{J,X} \frac{L}{q} (q/L)^J (\log X)^2, \quad L \geq 2q, \quad J \geq 2. \quad (17)$$

The assumption $UQ \leq X/2$ ensures $L_r = X/r \geq 2q$ for every term.

4. Summing every short divisor, actual shift and permitted modulus

Apply (17) to (12) and use only $|\mu(r)\mu(q)| \leq 1$ and $\sum_h |V(h/H)| \ll_V H$. Since $\sum_{r \leq U} r^{J-1} \ll_J U^J$ and $\sum_{q \leq Q} q^{J-1} \ll_J Q^J$, the result is

$$\begin{aligned} C_J H (\log X)^2 \sum_{r \leq U} \sum_{q \leq Q} q^{J-1} (X/r)^{1-J} &\ll_J H X^{1-J} U^J Q^J (\log X)^2 \\ &= H X (UQ/X)^J O_J((\log X)^2). \end{aligned} \quad (18)$$

This proves (2). An arbitrary selected family is bounded by this full cutoff sum. All coprimality restrictions were retained until taking the upper bound.

For $U = X^{2/5}$, $UQ/X = X^{-77/1000}$. At the largest shift length and $J = 4$, the exponent is

$$1 + \frac{2}{7} - 4 \frac{77}{1000} = \frac{1711}{1750} < 1. \quad (19)$$

This needs just four fixed derivatives of the profile, explicitly covered by (15). For example, $J = 17$ gives exponent $-163/7000$; no high-order numerical calculation is needed. More generally, fix $0 < \eta < 477/1000$ and let

$$U \leq X^{477/1000-\eta}, \quad (20)$$

Then $L_r = X/r \geq X^{523/1000+\eta}$ for every retained divisor, while the profile seminorms in (15) remain $O_J(\log^2 X)$ uniformly. For any fixed integer $J \geq 2$ with $J\eta > 2/7$, equation (18) gives

$$|D_Q^V[\Lambda_{\leq U}]| \ll X^{1+2/7-J\eta} (\log X)^2 = o(X \log X). \quad (20a)$$

Choosing a larger fixed J gives any desired smaller absolute power. Constants may depend on η, J, V, χ ; no uniform claim as $\eta \rightarrow 0$ is made.

5. Precisely which divisor and Heath–Brown components are covered

Equation (6) is an exact arithmetic application: the $\mu * \log$ portion with short Möbius divisor $r \leq U$ has negligible discrepancy. The Möbius factor is not assumed smooth. It is attached to the shorter fixed variable, while the logarithmic cofactor has a sufficiently long smooth summation range.

The same criterion can be stated for a Heath – Brown decomposition. Here is the needed identity with its range proved. Let $\mathbf{1}(n) = 1$, let ϵ be the identity for Dirichlet convolution, and let $\mu_{\leq Y} = \mu|_{[1, Y]}$. Then

$$E = \epsilon - \mu_{\leq Y} * \mathbf{1} = \mu_{> Y} * \mathbf{1}$$

is supported on integers greater than Y . Thus $E^{*k} * \Lambda$ vanishes on $n \leq Y^k$. Expanding $\epsilon - E^{*k}$ and using $\mathbf{1} * \Lambda = \log$ proves, on that range,

$$\Lambda = \sum_{j=1}^k (-1)^{j-1} \binom{k}{j} (\mu_{\leq Y})^{*j} * \mathbf{1}^{*(j-1)} * \log. \quad (21)$$

The equality at 1 is again zero on both sides. Choose Y^k above the fixed upper end of the actual $n \asymp X$ support to cover all terms.

After a smooth multiplicative partition, every component of (21) with an individual unrestricted $\mathbf{1}$ or $\log$ variable at scale

$$L \geq QX^\eta \quad (22)$$

is handled by the same progression proof. Combine all other variables into a short coefficient $\alpha(r)$, with $r \ll X/L$. It may contain truncated Möbius variables and divisor multiplicities. A fixed divisor bound on $\alpha(r)$ costs at most X^ε for any fixed $\varepsilon > 0$; increasing J absorbs this cost. Smooth partition factors and fixed powers of logarithms preserve (15), with their fixed logarithmic losses recorded.

The condition is on an individual smooth variable, not the total size of a product. If a nominal longer coefficient is itself a convolution, opening it helps only on subranges where an actual remaining smooth variable still satisfies (22). For example, a product $ab \asymp X^6$ does not establish that the b -range in $a_m = \sum_{ab=m} \alpha(a)A(b/L)$ exceeds Q .

The exact remainder $\Lambda_{>U}$ in (6) stays untreated. It need not be balanced: its logarithmic cofactor can be small. So do Type II and higher-factor components for which every individual smooth variable is at most $QX^{\alpha(1)}$. A longer Möbius or unsmoothed divisor sequence does not satisfy (15) merely because its support is long. No estimate for these remaining pieces is asserted.

6. Research consequence and verification scope

R13's sparse rational cores can have large positive restricted main terms. Here an adequately long smooth variable makes their complete centered sum cancel, with the original primitive mean and all frequencies retained. The direct progression proof also carries the actual joint $\mathfrak{s}, r, h$ kernel, avoiding any unproved weight separation.

For the actual zeta-facing discrepancy, (2) – (6) remove a definite short-divisor portion unconditionally. The residual problem can therefore be studied with that portion subtracted exactly and a proved small error. This does not bound the remaining signed piece, prove a covariance asymptotic, refute AH, or establish Montgomery's conjecture.

The source paper [OpenAI, Improved short gaps between primes](#), Definitions 2.6 – 2.9 and Proposition 2.18, supplies the coefficient terminology used for comparison. Its proof of Corollary 2.19 invokes a Heath – Brown reduction; the identity used here is proved directly in (21). The actual kernel and normalization are displayed in full in (4) – (5), following the frozen R9 – R11 definitions.

The adjacent check verifies complete-period centering, the Poisson sign on an independently summed Gaussian progression, exact exponent arithmetic, and a finite symbolic-coefficient version of (21). The Gaussian computation is a diagnostic; (16) is the ordinary analytic proof. No large-prime realization, parameter search, external model, or prior-file modification is used.

Current report 59: Independent review: smooth long–factor removal

Collection: R14 — smooth Type I removal and quantitative finite CUE heat.

Source: `research/dyson/round14/smooth-long-factor/INDEPENDENT_REVIEW.md`.

SHA-256: `3a6bef5c4cddf0973cacd69826aefa3fbea1db3659beb7f856ecaf3f6cee356b` . **Git blob:** `7ca9c893f44978336db65b43ee2cc9befb8026a8` . **Original bytes:** 8401.

Independent review: smooth long–factor removal

Date: 2026-09-05. **Accepted as an unconditional ordinary proof of the stated component bound.** No mathematical correction is requested. This is an independent check of the complete authored note, with special attention to the exact kernel, primitive restrictions, Poisson normalization, uniform derivatives, and total error. It does not claim an estimate for the remaining signed discrepancy or the zeta covariance.

The reviewed author file is `SMOOTH_LONG_FACTOR_REMOVAL.md` , SHA-256 `d6143f19ddf006a1acc833ecd2e5265bffb35817930cfeaa4f4e4b973af7c849`. The author file and its existing outputs were left unchanged. The proof was compared with the frozen R9 discrepancy definition and R11 kernel formula. The bounded author script was inspected and replayed in a temporary copy; its certificate agrees after normalizing only the temporary source path. Replay details are in `INDEPENDENT_REPLAY.json` , and hashes of the reviewed dependencies and review artifacts are in `INDEPENDENT_REVIEW_RECEIPT.json` .

Acceptance coverage

Item	Author location	Finding
Actual discrepancy and log-cofactor weight	Equations (4) – (6)	Matches R9/R11; no weight or primitive principal term is dropped.
Periodic phase mean and Fourier sign	Equations (7) – (11)	Correct, including nonprimitive numerators and nonsquarefree moduli.
Exact profile and uniform derivatives	Equations (12) – (15)	Correct in the full stated parameter range.
Progression Poisson formula and zero mode	Equations (16) – (17)	Correct sign, factor L/q , and exact principal cancellation.
Sum over every r,h,q	Equations (18) – (20a)	Correct bound $HX(UQ/X)^J \log^2 X$ and exponent thresholds.
Exact signed remainder and HB identity	Equations (6), (21) – (22)	Correct; the criterion requires an individual smooth long variable.

Item	Author location	Finding
Validation scope	Adjacent script/certificate	Small exact algebra and floating Gaussian diagnostic only; not a computational proof of the analytic estimate.

1. Kernel and parameter uniformity

For $n=rs$, $L_r=X/r$, $u=s/L_r$, $\delta=H/X=1/T$, and $z=h/H$, the exact identities are $rs=Xu$ and $rs-h=X(u-\delta z)$. Because χ is supported inside $(1,3/2)$, both arguments of the α -factors lie above X on the nonzero support. Thus both use the power $-3/2$. The two logarithms become precisely

$$\log s = \log(X/r) + \log u, \quad \log((rs-h)/q) = \log(X/q) + \log(u-\delta z).$$

The sinc argument has no hidden large derivative:

$$\delta^{-1} \log \frac{u}{u-\delta z} = \int_0^z \frac{dt}{u-\delta t}, \quad \partial_u^j \int_0^z \frac{dt}{u-\delta t} = (-1)^j j! \int_0^z \frac{dt}{(u-\delta t)^{j+1}}.$$

Here z lies in a fixed compact subset of $(1,2)$. On the support, $u-\delta t \geq u-\delta z$ is bounded away from zero for $0 \leq t \leq z$. Also $\delta \leq X^{-5/7}$ in the stated range. The profile therefore has a fixed positive compact support for all sufficiently large X , and every fixed u derivative of the nonlogarithmic factors is uniformly bounded. The only growth is from the two logarithms, each $O(\log X)$, since $1 \leq r \leq U < X$ and $1 \leq q \leq Q < X$. This proves the uniform $O_J(\log^2 X)$ seminorm bound. Smooth zero extension is legitimate because χ is compactly supported strictly inside its open interval.

The needed derivatives are derivatives in the normalized summation variable u , **uniformly in r, h, q, T** ; the proof does not require differentiating an arithmetic cutoff with respect to r or q . There is no missing r^J factor: the rescaling to L_r already accounts for it. The exact joint kernel is retained throughout.

2. Primitive masks and Poisson cancellation

If $(r,q) > 1$ and $(h,q) = 1$, $rs \equiv h \pmod q$ has no solutions, and the primitive principal sum has no terms because $(rs,q) > 1$. Such terms vanish in both parts, rather than contributing an exceptional error. If $(r,q) = 1$, multiplication by r gives

$$rs \equiv h \pmod q \iff s \equiv h\bar{r} \pmod q, \quad (rs, q) = 1 \iff (s, q) = 1.$$

The resulting class $b = h\bar{r}$ is a unit. For the Fourier convention $\widehat{\Phi}(t) = \int \Phi(u) e(-tu) du$, the progression formula is

$$\sum_{s \equiv b(q)} \Phi(s/L) = \frac{L}{q} \sum_k \widehat{\Phi}(kL/q) e(kb/q).$$

Averaging this formula over all unit b produces exactly $c_q(k)/\varphi(q)$. The $k=0$ terms cancel since $c_q(0) = \varphi(q)$. The sign in author equation (16) is consequently correct. The bounds $|c_q(k)| \leq \varphi(q)$ and $|\widehat{\Phi}(t)| \ll_J \log^2 X |t|^{-J}$, together with $J \geq 2$ and $L \geq 2q$, give equation (17), including its factor L/q .

The auxiliary periodic identity in § 2 is also correct: a unit inner n permutes the units, so the period mean of R is zero. With $R(m) = \sum_{\ell} \hat{R}(\ell) e(\ell m/q)$, Poisson yields the coefficient $\hat{R}(-\nu \pmod q)$ in equation (9). Multiples of q vanish as well. The shift bound $\sum_{1 \leq a < q} |S_{\{V, H\}}(a/q)| \ll q$ follows by summing the stated decay, both when $q \geq H$ and when $q < H$. This auxiliary explanation is not substituted for the direct progression proof of the actual kernel.

3. Total error and numerical exponents

The bound for one retained r, h, q is

$$\ll_J q^{J-1} (X/r)^{1-J} \log^2 X.$$

Taking absolute values of the Möbius coefficients, summing the actual shifts with $\sum_h |V(h/H)| \ll H$, and enlarging the permitted moduli to every $q \leq Q$ gives

$$\ll_J H X^{1-J} \log^2 X \left(\sum_{r \leq U} r^{J-1} \right) \left(\sum_{q \leq Q} q^{J-1} \right) \ll_J H X (UQ/X)^J \log^2 X.$$

No extra q , H , or divisor-count factor is omitted. The $q=1$ discrepancy is identically zero, so including it in this upper sum causes no issue. Real $U \geq 1$ with the stated integer cutoff also causes no endpoint term.

For $U = X^{\{2/5\}}$, $Q = X^{\{523/1000\}}$, $H \leq X^{\{2/7\}}$, and $J=4$, the exponent is $1711/1750$, with margin $39/1750$ below 1. For $U \leq X^{\{477/1000-\eta\}}$, fixed $\eta > 0$ ensures $UQ \leq X/2$ for sufficiently large X . Any fixed integer $J \geq 2$ satisfying $J\eta > 2/7$ gives $o(X \log X)$, despite the remaining logarithmic square. The dependence of constants on η and J is properly stated; no claim uniform as $\eta \downarrow 0$ is needed.

4. Signed remainder and Heath–Brown scope

The identities $\Lambda = \mu * \log$ and $\Lambda = \Lambda_{\leq U} + \Lambda_{> U}$ hold pointwise, including $n=1$, and the discrepancy is linear. Therefore equation (6) is exact before estimation. Neither sign nor cancellation properties of the untreated divisor portion are inferred from the small estimate for the retained portion. Prime powers have not been removed, so there is no missing prime-power exception here.

For the displayed Heath – Brown identity, $E = \epsilon - \mu_{\leq Y} * 1 = \mu_{> Y} * 1$ is supported on integers greater than Y . Thus $E^{\{k\}} * \Lambda$ vanishes for $n \leq Y^k$. Expanding $\epsilon - E^{\{k\}}$ and using $1\Lambda = \log$ gives equation (21) with its displayed binomial coefficients and signs. Taking Y^k beyond the support covers all terms.

After a smooth multiplicative partition, an individual unrestricted 1 or log variable of scale $L \geq QX^\eta$ admits exactly the same progression argument. The other variables may be grouped into a short arithmetic coefficient $c(r)$; no smoothness is assigned to $c(r)$. For fixed identity order, divisor multiplicities cost at most X^ϵ and fixed logarithmic factors, which a larger fixed J absorbs. The criterion does not apply merely because a product of several rough variables is long. This limitation, the possibility of an unbalanced $\Lambda_{> U}$, and the absence of a bound for the whole remaining discrepancy are all stated correctly in the author report.

5. Replay and final scope

The inspected script checks 63 exact periodic-centering cases, the formal-logarithm divisor and HB identities through $n=125$, the real short cutoff $9/2$, and exact rational exponent arithmetic. Its two Gaussian progression comparisons have the correct Poisson sign and agree within the stated binary64 tolerance; these are diagnostics, not rigorous enclosures. The temporary replay passed every assertion and reproduced the certificate modulo its temporary source path. Original report, script, certificate, and log were not changed.

This review accepts the displayed unconditional removal of a definite short-divisor component. It does not upgrade the R13 positive restricted-core claim into a full signed lower bound, establish cancellation for an arithmetic long coefficient, estimate the remaining $\Lambda_{\{>U\}}$ discrepancy, or imply AH/RH/Montgomery progress beyond the explicit component reduction. No additional scan, literature claim, or author edit is required for acceptance of the result as stated.

Current report 60: CUE background at the selected smallest gap: a direct finite-N bound

Collection: R14 — smooth Type I removal and quantitative finite CUE heat.

Source: [research/dyson/round14/cue-selected-background/SELECTED_CUE_BACKGROUND.md](#).

SHA-256: [fbc67828d13534d8d0b4ac1f742a639b282dd93a3f7cb635291f8cddb651c0a5](#) . Git blob: [5df1a161d2aea573ed2a24336325c5fa5ae248c5](#) . Original bytes: 11711.

CUE background at the selected smallest gap: a direct finite-N bound

Date: 2026-09-05. Author: Astra subagent [yau_flow](#) . Status: complete ordinary-proof draft submitted for independent review. This is a CUE-only estimate using its exact finite-dimensional determinant. It does not invoke any general-beta correlation inequality, assert novelty against the entire literature, or transfer a random-matrix result to zeta zeros.

The finite CUE heat-depth ratio was already derived qualitatively in Round 1, [yau_flow.md](#) Sections 3 – 4, and its deterministic Galilean refinement was quantified in Round 2, [galilean-proof-audit.md](#) Sections 1 – 5. Those results are not being presented as new here. The additional estimate is tightness at the natural N^2 scale for the background of the **selected minimum pair**, with an explicit finite-N truncated first-moment bound. It sharpens the rate obtained by applying the existing deterministic lemma.

1. Exact geometry and statement

Let U_N have Haar law on $U(N)$. Its eigenangles are points of $T = \mathbb{R}/(2\pi\mathbb{Z})$, with ordered representatives $\theta_1 < \dots < \theta_N$ and $\theta_{N+1} = \theta_1 + 2\pi$. Write

$$\delta_i = \theta_{i+1} - \theta_i, \quad c_i = \theta_i + \delta_i/2 \pmod{2\pi}, \quad \delta_N^{\min} = \min_i \delta_i.$$

All roots are distinct almost surely. Any fixed deterministic tie-breaking convention may select the minimum; ties have probability zero.

For circular distance $r(x) = \min_{k \in \mathbb{Z}} |x - 2\pi k| \in [0, \pi]$, put

$$q(x) = \frac{1}{4 \sin^2(x/2)}.$$

The midpoint background and the larger endpoint background of a consecutive gap are

$$B_j = \sum_{k \in j, i+1} q(\theta_k - c_i), \quad S_j = \sum_{k \in j, i+1} [q(\theta_k - \theta_i) + q(\theta_k - \theta_{i+1})]. \quad (1)$$

The selected background is $B_N = B_{i_*}$, where $\delta_{i_*} = \delta_N^{\min}$.

Finite-N proposition. For every $N \geq 2$ and $0 < \varepsilon \leq \pi$,

$$\mathbb{E} \sum_{i: \delta_i \leq \varepsilon} B_i \leq \frac{N^6 \varepsilon^3}{18}. \quad (2)$$

The same bound holds with B_i replaced by S_i . In fact the proof bounds the endpoint-background sum over all **ordered pairs** whose positive oriented separation is at most ε , even when they are not consecutive.

Consequently, for fixed $L > 0$ and all sufficiently large N ,

$$\mathbb{E} \sum_{i: \delta_i \leq LN^{-4/3}} \frac{B_i}{N^2} \leq \frac{L^3}{18}. \quad (3)$$

No conditional density at the minimum is used.

2. A three-point bound that retains the endpoint zeros

With correlation densities measured against ordinary angular Lebesgue measure, use the periodic CUE kernel

$$K_N(x, y) = \frac{1}{2\pi} \sum_{m=0}^{N-1} e^{im(x-y)}, \quad \rho_k(x_1, \dots, x_k) = \det[K_N(x_a, x_b)]_{a,b \leq k}.$$

This determinant convention counts ordered distinct tuples. It is gauge-equivalent to the sine kernel in Feng – Wei, Section 1.2, printed page 5; the integer-frequency form avoids any artificial cut or antiperiodic-kernel issue.

Let

$$\phi(x) = \frac{1}{\sqrt{2\pi}}(1, e^{ix}, \dots, e^{i(N-1)x}), \quad A = \|\phi(x)\|^2 = \frac{N}{2\pi}, \quad D = \|\phi'(x)\|^2 = \frac{N(N-1)(2N-1)}{12\pi} \leq \frac{N^3}{6\pi}.$$

For $0 < d \leq \pi$, Gram determinant factorization gives

$$\rho_2(0, d) \leq A \|\phi(d) - \phi(0)\|^2 \leq ADd^2.$$

Let Π be orthogonal projection onto the span of $\phi(0), \phi(d)$. Then

$$\rho_3(0, d, z) = \rho_2(0, d) \|(I - \Pi)\phi(z)\|^2.$$

The final squared norm is at most A , and also at most $\|\phi(z) - \phi(0)\|^2 \leq Dr(z)^2$, because $\phi(0)$ belongs to that span and ϕ is periodic. Therefore

$$\rho_3(0, d, z) \leq ADd^2 \min\{A, Dr(z)^2\} \leq \frac{N^5 d^2}{24\pi^3} \min\{1, N^2 r(z)^2\}. \quad (4)$$

The sharper factor $N^2 r^2/3$ in the second branch was weakened to $N^2 r^2$. Equivalently, the two elementary bounds before this weakening are $N^5 d^2/(24\pi^3)$ and $N^7 d^2 r^2/(72\pi^3)$. For $N = 2$, the third density is zero, so the same conclusion holds. Coincident points are null sets and can be handled by continuity.

This is the needed finite-dimensional replacement for a general-beta background claim. It simultaneously supplies the short-pair factor $d^\flat$ and the endpoint cancellation $r(z)^2$.

3. Integrating the singular weight

For $0 < r \leq \pi$, $\sin(r/2) \geq r/\pi$, hence

$$q(z) \leq \frac{\pi^2}{4r(z)^2}.$$

Splitting at $r = 1/N$ yields

$$\begin{aligned} \int_{\mathbb{T}} q(z) \min\{1, N^2 r(z)^2\} dz &\leq \frac{\pi^2}{2} \left(\int_0^{1/N} N^2 dr + \int_{1/N}^{\pi} r^{-2} dr \right) \\ &= \frac{\pi^2}{2} (2N - 1/\pi) \leq \pi^2 N. \end{aligned}$$

Equation (4) therefore gives

$$\int_{\mathbb{T}} q(z) \rho_3(0, d, z) dz \leq \frac{N^6 d^\flat}{24\pi}. \quad (5)$$

Interchanging the two pair endpoints gives the same bound with $q(z - d)$. This can be seen either by translation/reflection invariance or by repeating the projection proof with base point d .

Let Z_ε sum the endpoint background over all ordered distinct pairs (a, b) whose positive oriented separation $d = (b - a) \pmod{2\pi}$ lies in $(0, \varepsilon]$. The factorial-moment identity and rotation invariance give

$$\begin{aligned} \mathbb{E} Z_\varepsilon &= 2\pi \int_0^\varepsilon \int_{\mathbb{T}} [q(z) + q(z - d)] \rho_3(0, d, z) dz dd \\ &\leq 2\pi \int_0^\varepsilon \frac{N^6 d^\flat}{12\pi} dd = \frac{N^6 \varepsilon^3}{18}. \end{aligned} \quad (6)$$

The integrands are nonnegative, so Tonelli is legitimate before the estimates prove finiteness. There is no factor $1/2$: the correlation density counts ordered triples, and the short orientation of a consecutive gap is included once. The pair straddling the chosen 2π cut is included by exactly the same periodic change of variables. Even at $\varepsilon = \pi$, possible antipodal ambiguity is a null event.

To pass to midpoint backgrounds, rotate a consecutive pair of gap $d \leq \pi$ to endpoints $\pm d/2$. Every other point has a midpoint-centered lift $y \in [-\pi, \pi]$ with $|y| \geq d/2$, because the open short arc contains no point. Its distance to the endpoint on the same side is $|y| - d/2 \leq |y|$. Since q decreases with circular distance on $(0, \pi]$,

$$q(y) \leq q(y - \operatorname{sgn}(y)d/2) \leq q(y - d/2) + q(y + d/2).$$

Thus $B_i \leq S_j$ for every included consecutive pair. Its endpoint-background sum is a subsum of Z_ε , proving (2).

Why this order matters. One cannot simply drop consecutiveness from the midpoint-weighted sum: a third point can then approach the midpoint while staying away from both endpoints, where ρ_3 need not vanish. The inverse-square midpoint weight would have a nonintegrable singularity. The endpoint majorant avoids this invalid enlargement.

4. The selected background is tight at scale N^2

The checked primary extreme-gap input is Feng – Wei, Theorem 1.1 and Corollary 1.1, printed page 4, specialized **only** to CUE. Their periodic definition includes the wrap gap, and $A_2 = 1/(24\pi)$. Hence

$$P(N^{4/3} \delta_N^{\min} > L) \rightarrow \exp\left(-\frac{L^3}{72\pi}\right) \quad (L > 0). \quad (7)$$

See Feng – Wei, [arXiv:1806.01555v2](#). No finite- N uniform tail estimate is asserted.

On the event $\delta_N^{\min} \leq LN^{-4/3}$, the selected B_N is one nonnegative summand in (3). Therefore, for every fixed $L, K > 0$,

$$\limsup_{N \rightarrow \infty} P(B_N/N^2 > K) \leq e^{-L^3/(72\pi)} + \frac{L^3}{18K}. \quad (8)$$

Choosing L first and then K proves

$$\boxed{B_N/N^2 = O_{\mathbb{P}}(1)}. \quad (9)$$

For example, choosing the fixed value $L^3 = 72\pi \log K$, for each $K > 1$, gives the asymptotic tail bound $(1 + 4\pi \log K)/K$ on the right side of (8). This is a limit-superior bound, not a quantitative convergence rate in N , and it is not a uniform first-moment estimate for B_N/N^2 . The entire argument avoids conditioning on which pair is smallest.

5. Quantitative consequence for the already-defined finite heat flow

Let

$$P_N(z) = \det(zI - U_N) = \sum_{j=0}^N a_j z^j, \quad P_{N,s}(z) = \sum_{j=0}^N a_j e^{sj(N-j)} z^j,$$

and define the first discriminant time

$$D_N = \inf\{s > 0 : \text{disc}(P_{N,s}) = 0\}.$$

This is the finite circular scalar-heat deformation, not stochastic Dyson Brownian motion and not a representation of the true zeta flow.

The exact deterministic statement already audited in Round 2 uses the same midpoint background B_N . For

$$\eta_N = (\delta_N^{\min})^2 (B_N + 1), \quad K_0 = 16384, \quad \eta_0 = 1/524288,$$

it proves

$$\frac{(\delta_N^{\min})^2}{8} \leq D_N \leq (\delta_N^{\min})^2 \left(\frac{1}{8} + K_0 \eta_N\right) \quad \text{if } \eta_N \leq \eta_0. \quad (10)$$

The source is `research-round2/galilean-proof-audit.md`, Sections 1 – 5, pinned in the source receipt. Its constants are independent of N and of the common background drift. Its proof controls the real Gaussian tails after an exact Galilean conjugation and keeps both moving interval boundaries nonzero throughout the comparison time. It therefore does not require a new unproved dynamic-background stability assumption. A multiple zero somewhere else earlier only helps its upper bound.

Put $X_N = N^{4/3} \delta_N^{\min}$ and $Y_N = B_N/N^2$. Equations (7) and (9) imply both are tight, without requiring independence. Thus

$$N^{2/3} \eta_N = X_N^2(Y_N + N^{-2}) = O_P(1), \quad P(\eta_N \leq \eta_0) \rightarrow 1.$$

Applying (10) on this event proves the strengthened estimate

$$\boxed{\frac{8D_N}{(\delta_N^{\min})^2} - 1 = O_P(N^{-2/3}).} \quad (11)$$

The difference is nonnegative even outside the small- η_N event: the unconditional scalar comparison gives $D_N \geq -\log \cos(\delta_N^{\min}/2) \geq (\delta_N^{\min})^2/8$. Here $\delta_N^{\min} < \pi$ almost surely for $N \geq 2$, deterministically for $N \geq 3$. This is a tightness-scale rate for the approximation error, not a distributional convergence rate or a claim about its limiting correction law. Equivalently, $D_N - (\delta_N^{\min})^2/8 = O_P(N^{-10/3})$. For completeness, D_N is finite almost surely: a generic CUE characteristic polynomial has a nonzero interior coefficient, whose prescribed exponential growth eventually exceeds the bounded coefficient size possible for a polynomial with all roots on the circle. Before a first collision, self-inversive symmetry keeps its simple roots on the circle. The exceptional exactly rotation-invariant polynomial has probability zero.

The qualitative limiting depth law already recorded in Round 1 follows again. We do not label it a new result of this round, and we do not infer an error rate for that limiting law from (11).

6. What this does and does not settle

The new finite- N estimate controls the initial inverse-square background attached to the actual selected minimum. Together with an existing fully quantified deterministic lemma it removes a possible missing-stability premise in this finite CUE lane and sharpens its probabilistic approximation rate.

A theorem identifying a limiting conditional background distribution, an asymptotic mean of that background, an optimal error exponent or coefficient, general-beta analogues, and any arithmetic transfer are outside the proved result. For true zeta zeros, the isolated-small-pair and positive-density arithmetic obligations documented in Round 2 remain open. CUE's three-point determinant is an actual random-matrix input, not an available identity for zeta zeros.

The companion check records exact normalization/exponent arithmetic and the $N = 3$ determinant-integral identity. It is not a grid scan, a Monte Carlo test, or a formal verification of this proof. Independent/coordinator review must be recorded separately without changing the eventual frozen author snapshot.

Current report 61: Independent review: selected CUE background and the finite heat rate

Collection: R14 — smooth Type I removal and quantitative finite CUE heat.

Source: [research/dyson/round14/cue-selected-background/INDEPENDENT_REVIEW_EUCLID.md](https://research.dyson/round14/cue-selected-background/INDEPENDENT_REVIEW_EUCLID.md).

SHA-256: `54f03e2f3f2636392fd630a19900dddb81d56e1fcffa7a656a7af36c4f1a6018` . Git blob: `57db6633b90fc1154f25a9fccb27ce070a054f53` . Original bytes: 8634.

Independent review: selected CUE background and the finite heat rate

Date: 2026-09-05. Reviewer: Euclid, the independent conductor-arithmetic agent.

Decision: accepted for the stated finite-CUE conclusions. No mathematical defect was found in the selected-background bound, its exact constant, the circular geometry, or the stated probabilistic consequence of the pinned deterministic lemma. This is an ordinary mathematical review, not formal verification or a novelty assessment.

Reviewed author file: `SELECTED_CUE_BACKGROUND.md`

SHA-256: `fb6c7828d13534d8d0b4ac1f742a639b282dd93a3f7cb635291f8cd6b651c0a5`.

The hash was independently recomputed. No author source, numerical check, receipt or previous-round file was modified or rerun.

1. The three-point bound and the constant 1/18

The CUE kernel is normalized against ordinary angular Lebesgue measure:

$$K_N(x, y) = \frac{1}{2\pi} \sum_{m=0}^{N-1} e^{im(x-y)}.$$

The Gram-vector norms used in the proof are therefore

$$A = \frac{N}{2\pi}, \quad D = \frac{1}{2\pi} \sum_{m=0}^{N-1} m^2 \leq \frac{N^3}{6\pi}.$$

The two-point estimate $AD\mathcal{P}$ follows by subtracting the first vector from the second; the third-point Gram residual is bounded both by A and by $D\mathfrak{r}(z)^2$. Thus the two coefficients before weakening are exactly

$$\frac{N^5 \mathcal{P}}{24\pi^3}, \quad \frac{N^7 \mathcal{P} \mathfrak{r}(z)^2}{72\pi^3}.$$

Replacing the second branch by the weaker coefficient $1/24$ is legitimate. The displayed minimum bound in the author report follows. For $N = 2$, the rank-three determinant vanishes, so no exceptional formula is required.

The singular integration uses both halves of the circle:

$$\int_{\mathbb{T}} q(z) \min(1, N^2 r(z)^2) dz \leq \frac{\pi^2}{2} \left[N + N - \frac{1}{\pi} \right] \leq \pi^2 N.$$

This gives $N^6 \mathcal{C}^2 / (24\pi)$ for one endpoint. Adding the two endpoints, integrating the starting angle over length 2π , and integrating $\mathcal{C}^2$ over $(0, \varepsilon]$ gives

$$\frac{1}{24} \times 2 \times 2 \times \frac{1}{3} = \frac{1}{18}.$$

All powers of π cancel. There is no missing $1/2$: the determinant is the factorial density of ordered distinct triples, and the oriented short-pair enumeration is the one actually integrated.

The nonnegative integrand makes the initial application of Tonelli valid. The obtained upper bound then proves its finiteness.

2. Midpoints, endpoints and circular wrap

The midpoint singularity is handled in the correct order. Only a consecutive pair is replaced by its endpoint majorant; the resulting endpoint-weighted quantity can then be enlarged to all short ordered pairs. Enlarging the midpoint-weighted sum first would not be valid, and the report explicitly avoids that step.

For a consecutive gap with endpoints $\pm d/2$, $d \leq \pi$, every other point has a midpoint-centered lift $y \in [-\pi, \pi]$ with $|y| \geq d/2$. Its distance to the endpoint on its own side is exactly $|y| - d/2$, lying in $[0, \pi]$. Monotonicity of $1/(4 \sin^2(r/2))$ on that interval therefore gives the claimed midpoint majorant. Third-point coincidence with an endpoint is excluded by the simple-point configuration.

At $y = \pi$ or $-\pi$, either admissible lift leads to the same periodic value; choosing a side creates no discontinuity in the inequality. The short oriented pair crossing the original angle cut is included once under the periodic change of variables. Antipodal pair ambiguity has probability zero; the non-strict endpoint $\varepsilon = \pi$ creates no expectation term.

The integer-frequency kernel is genuinely periodic. Hence an antiperiodic gauge of a sine-kernel representation cannot introduce an extra sign or omit a wrap pair in this proof.

3. Selection of the minimum and the extreme-gap normalization

I checked the stated primary normalization against the local text of Feng – Wei, arXiv:1806.01555v2, printed p.4. Its gap process uses the periodic convention $\theta_{\pm N} = \theta_i + 2\pi$. At $\beta = 2$, $A_2 = 1/(24\pi)$, and Corollary 1.1 scales the smallest gap by

$$N^{4/3} (A_2/3)^{1/3}.$$

Thus the author report's tail

$$P(N^{4/3} \delta_N^{\min} > L) \rightarrow e^{-L^3/(72\pi)}$$

has the correct angular normalization.

On the event that the selected gap is at most $LN^{-4/3}$, its background is one nonnegative term of the truncated sum. Scaling the finite- N bound gives exactly $L^3/18$. Markov's inequality then proves

$$\limsup_N \mathbb{P}(B_N/N^2 > K) \leq e^{-L^3/(72\pi)} + \frac{L^3}{18K}.$$

Choosing L first and K second proves tightness. Equivalently, for each fixed $K > 1$, taking $L^3 = 72\pi \log K$ yields the stated asymptotic upper tail $(1 + 4\pi \log K)/K$.

There is no conditioning on the identity of the minimum pair. The proof does not claim a finite- N convergence rate, a limiting background law, or a uniform bound for $\mathbb{E}(B_N/N^2)$. These limitations are correctly stated.

4. Matching the pinned Galilean lemma

I read Sections 1 – 5 of the pinned deterministic dependency:

research-round2/galilean-proof-audit.md

SHA-256: c85684fe873c19c193a81d3d16cde2507f10cf6753324ce31eda99b14672a2da.

The hypotheses match the CUE application exactly:

- A degree- N polynomial has N distinct roots on the unit circle, with fixed nonzero leading and constant coefficients.
- The coefficient deformation is $a_j \mapsto a_j e^{Sj(N-j)}$, with the same time parameter S .
- The chosen pair realizes the smallest circular angular gap δ .
- Its initial background is exactly $B = \frac{1}{4} \sum_{\text{outside pair}} \csc^2(\theta_k/2)$, measured from the midpoint after rotation. This equals the author's B_N , with no additional endpoint or scaling factor.
- The smallness condition is $\eta = \delta^2(B+1) \leq 1/524288$.

The dependency proves the bound with $K_0 = 16384$ independent of N and of the common logarithmic drift. Its real-line Gaussian estimate and moving-boundary argument are already uniform in the drift. The new CUE application therefore does not require a further dynamic-background-stability hypothesis.

The scalar lower comparison is valid whenever the smallest initial gap is $< \pi$. It gives

$$D_N \geq -\log \cos(\delta/2) \geq \delta^2/8.$$

For $N \geq 3$, the minimum is at most $2\pi/N < \pi$; for $N = 2$, equality $\delta = \pi$ has probability zero. Consequently the author's nonnegativity assertion outside the good smallness event is justified as well.

The first discriminant time may be caused by another pair. That possibility only strengthens the deterministic upper bound; the application does not assume that the initially smallest pair is the first to collide.

5. The $N^{-2/3}$ rate and its precise meaning

Let $X_N = N^{4/3} \delta_N^{\min}$ and $Y_N = B_N/N^2$. Both are tight, and

$$N^{2/3}\eta_N = X_N^2(Y_N + N^{-2}).$$

Products of tight random variables are tight without independence. Thus $\eta_N = O_P(N^{-2/3})$, and its fixed deterministic smallness condition holds with probability tending to one.

On this event, the pinned lemma gives

$$0 \leq \frac{8D_N}{(\delta_N^{\min})^2} - 1 \leq 8K_0\eta_N.$$

The exceptional event has probability tending to zero, which is sufficient for the asserted full $O_P(N^{-2/3})$ statement. No bound on the magnitude of the error on that exceptional event is needed for tightness.

The almost-sure finiteness remark is consistent: any nonzero interior coefficient eventually violates its bounded unit-circle coefficient size under the stated exponential deformation. Simple self-inversive roots cannot leave the circle before a collision. The exceptional polynomial with all interior coefficients zero has CUE probability zero.

Multiplying the relative error by the tight factor $N^{8/3}(\delta_N^{\min})^2$ yields the absolute error $O_P(N^{-10/3})$. This is an approximation-error rate, not a rate of convergence to the limiting depth distribution. The report keeps that distinction.

6. Verified provenance and remaining scope

Independently recomputed primary hashes:

- Feng – Wei PDF: af6c78625ceb76b422fb89ba0f5d98f18c8749ae922244fe386dbfb0133dbce7.
- Its extracted text: 6b24cd80cede5d71415c4f1cfa527a2c0ef8fd23752549c7455a85945f777ea5.

The review read the exact finite- N proof, the relevant primary text and the deterministic dependency. It did not rerun the companion script or add a numerical scan.

Acceptance is restricted to the stated CUE proposition, tightness and finite scalar-heat approximation. No general- β determinant bound, true-zeta flow statement, arithmetic transfer, optimal-rate claim or historical-conjecture conclusion follows from this review.

Current report 62: Independent root review: two bounded R14 advances

Collection: R14 — smooth Type I removal and quantitative finite CUE heat.

Source: [research/dyson/round14/INDEPENDENT_ROOT_REVIEW.md](#).

SHA-256: [8350088f7454529e209379e8f4692a78be6cc49fe1c67257a053a5065112bf3e](#) . **Git blob:** [6379a8ce78ba33a80ae923da44076ac4d3dbb198](#) . **Original bytes:** 5564.

Independent root review: two bounded R14 advances

Date: 2026-09-05. **Reviewer:** root Astra. Both complete author reports were read and their proofs checked independently. This review is not formal verification, a global novelty assessment or a claim about the full zeta covariance.

Actual Type I component

Reviewed [smooth-long-factor/SM00TH_LONG_FACTOR_REMOVAL.md](#) , SHA-256 [d6143f19ddf006a1acc833ecd2e5265bffb35817930cfeaa4f4e4b973af7c849](#) .

The identity $\Lambda = \mu * \log$, including $n=1$, makes the split at a real cutoff U exact. After $n=rs$, a nonunit r contributes to neither a primitive residue progression nor its principal sum. For a unit r the required residue is h times its inverse. No discarded nonunit main term remains.

The profile displayed in equation (13) uses the actual n - h kernel. Both a -factors have the negative-three-halves branch because the compact χ support forces $n-h > X$ and $n > X$. The integral representation in (14) removes the apparent small-denominator phase singularity. Its derivatives, the compact support and the two logarithms give the claimed uniform $O_J(\log^2 X)$ seminorms for all retained r, h, q . Constants can depend on fixed derivative order and cutoff separation; they do not silently depend on T .

Poisson summation has the positive progression phase $e(kb/q)$ for the stipulated negative-sign Fourier transform. Its zero frequency cancels exactly against the primitive mean, including $c_q(0) = \phi(q)$. The nonzero modes are bounded by $(L/q)(q/L)^J \log^2 X$. Summing $q^{-(J-1)}$ and $r^{-(J-1)}$, and the actual h weights, gives $HX(UQ/X)^J \log^2 X$ without any omitted factor of q or H .

The explicit choice $U = X^{.4J=4}$ has exponent $1711/1750$ and margin $39/1750$ below one. More generally the uniform proof applies to every fixed positive separation $U \leq X^{(.477-\eta)}$, with fixed J $\eta > 2/7$. There is no uniform assertion as η tends to zero. The proof of the recalled Heath-Brown identity is algebraically valid, and the criterion concerns an individual smooth unrestricted variable; a long product does not automatically satisfy it.

Accepted as an unconditional classical Poisson application to this programme's exact discrepancy. The signed $\Lambda_{\{>U\}}$ term remains exact and unestimated. Its cofactor can be small. Removing this component supplies neither an inequality for the full Λ sum nor a solution of Montgomery's conjecture.

Selected CUE background and finite heat–depth error

Reviewed `cue-selected-background/SELECTED_CUE_BACKGROUND.md`, SHA-256

`fbc67828d13534d8d0b4ac1f742a639b282dd93a3f7cb635291f8cddb651c0a5`, and reread Sections 1–5 of the existing quantified Galilean audit.

For the exact finite CUE Gram representation, $||\phi||^2 = N/(2\pi)$ and $||\phi'|||^2 \leq N^3/(6\pi)$. The two-point determinant is bounded by $AD d^2$. Projecting the third vector off the first two gives both A and $D r(z)^2$ bounds. Their minimum yields precisely the stated ρ_3 upper bound $N^5 d^2/(24\pi^3) \min(1, N^2 r(z)^2)$.

The singular endpoint integral is finite because the retained $r(z)^2$ factor cancels it. The elementary bound $q(z) \leq \pi^2/(4r(z)^2)$ and the split $r=1/N$ give an integral at most $\pi^2 N$. Two endpoints, the 2π anchor integral and integration of d^2 from zero to ϵ give $N^6 \epsilon^3/18$. These are ordered factorial correlations, so there is no extra half factor. The periodic orientation includes the wrap gap.

The midpoint comparison is made only for consecutive pairs: every third point is outside the open short arc and is at least as close to one endpoint as to the midpoint. Dropping consecutiveness only after replacing by endpoint weights avoids a divergent midpoint integral. The selected minimum requires no conditional density: on $\delta_{\min} \leq L N^{-(4/3)}$, its nonnegative background is one of the counted summands. Markov followed by the fixed- L minimum-gap tail gives $B_{\min}/N^2$ tightness. The order of limits is valid and does not assume a uniform finite- N tail estimate.

The pinned deterministic lemma uses exactly this midpoint B and controls the heat product after its Galilean conjugation, with constants independent of N and drift. Since $\delta_{\min} = N^{-(4/3)}$ times a tight variable, $\eta = \delta_{\min}^2(B+1) = O_p(N^{-(2/3)})$. The lemma's small- η event has probability tending to one. This proves $8D/\delta_{\min}^2 - 1 = O_p(N^{-(2/3)})$, and multiplication by the tight squared gap gives absolute difference $O_p(N^{-(10/3)})$. No independence between the gap and background is required.

The scalar minimum-gap lower comparison is available without the small- η condition when $\delta_{\min} < \pi$, which holds almost surely for CUE $N \geq 2$. Thus the stated nonnegativity is consistent. Finiteness of D for almost every CUE polynomial follows from growth of an interior coefficient and self-inversive root tracking before a collision, as stated in the report.

Accepted as an ordinary finite-CUE result based on the checked classical extreme-gap input and the existing audited deterministic lemma. This is an approximation error in probability, not a rate of convergence to the limiting depth distribution. It is not a stochastic DBM theorem, a general-beta result or an identity available for zeta zeros.

Remaining work

Both results are appropriate to preserve as concrete progress: one removes an exact actual arithmetic component, and the other strengthens a finite random-matrix heat-flow comparison. The full signed arithmetic remainder, the desired out-of-band zeta correlation inequality and all famous-conjecture targets remain open. Separate agent reviews and bounded algebra replays are retained alongside this review.

Current report 63: Fable PR11 snapshot 89393d5: separate intake and two-sided arithmetic audit

Collection: Fable intake — 89393d5 and 2073028, separate reviewed corrections.

Source: fable/reviews/pr11-89393d5/INTAKE_REVIEW.md.

SHA-256: `6e5a08438cdf6abc3dbffdbab699374f72738d894a472d47c6d78fa4db91276eb` . Git blob: `bc6a8343f54e80e77875a02237c7f41705256e3b` . Original bytes: 6968.

Fable PR11 snapshot 89393d5: separate intake and two-sided arithmetic audit

Date: 2026-09-05. Source: [Alpha-devbox commit 89393d5da61a45561ed199330c5b836f47fcd629](https://alpha-devbox.com/commit/89393d5da61a45561ed199330c5b836f47fcd629). This review accepts a source snapshot and reports specific checks. It does not approve every mathematical claim in that snapshot or establish a new zeta-gap result.

What arrived

The [new snapshot](#) contains 141 files, 1,062,904 bytes, preserved verbatim. Its [manifest](#) pins every file. The earlier 81-file `a408e705` mirror remains unchanged. The retained local source tarball has 3,104,082 bytes and SHA256 `9ca9fd9c7b907512db107a59e8f2d8caf5489887a532a55b4d6503bea42f1976` .

Three source commits recovered six proposer reports, added the C-beta-E background and F1 arithmetic-transfer drafts, then added two F1 refuters. The recovered claims include Theorem B repair, CUE background, Level B, structure, H2 and finite-sum diagnostics. These reports are now present; the earlier snapshot's absent-file assessment must not be applied to them. Presence is not proof acceptance. The source's own claim ledger describes several as awaiting refutation, and its last commit message says an F1 repair is in progress. That message is not a completed repair receipt.

This bounded intake examines F1/F2 arithmetic and reproduces the two F1 refuters. The [separate background and boundary audit](#) records additional specific objections; neither review is a full acceptance audit of the heat-flow drafts. No new Claude session, model request, automated next task, or large prime computation was launched.

The valid F1 objection, and the refuter's own sign error

The proposer text gives the leading coefficient of Π_4 as $6a^2$, then concludes $m_4 = a^2 + 6a$ after adding Π_2^2 . The two expressions are inconsistent. The correct leading terms are

$$\Pi_2^2 \sim a^2 \varepsilon^{-4}, \quad \Pi_4 \sim 6a \varepsilon^{-4}, \quad m_4 = a^2 + 6a.$$

To see the second coefficient, the local inclusion probability has $\rho_p = a p^{(-s)} + O(p^{(-2s)})$. Thus $\rho_p(1-\rho_p)$ has the same leading $a p^{(-s)}$. The sum of $(\log p)^4 p^{(-s)}$ has leading $6/(s-1)^4$; prime powers and local quadratic errors are analytic near $s=1$. This supplies a single factor a . The proposer's script assigns the correct final m_4 as a literal; that assignment does not verify its incorrect written intermediate coefficient. The resulting conditional moment expression $(a+6)v^4/((a+1)(a+2)(a+3))$ is consistent with the corrected coefficient.

However, the refuter's numerical probe also has an error. Its function $zz3$ is the third derivative of ζ'/ζ . Since $\zeta'/\zeta(1+\epsilon) = -1/\epsilon + \text{analytic}$, $zz3$ has leading $+6/\epsilon^4$. The script evaluates $-zz3*\epsilon^4$ and labels it as tending to $+6$. All three saved values are actually -6 . The subsequent coefficient is separately assigned as $6a$ and does not depend on this bad probe.

The [independent correction](#) differentiates the pole using exact rational algebra and separately runs the correctly signed high-precision diagnostic. The exact pole algebra supplies the coefficient; the numerical probe is a secondary sign check. This correction preserves the refuter's valid objection to the proposer while rejecting its faulty numerical verification. The original scripts and outputs are untouched.

The quoted normalization table mixes different v values

For fixed $v=1$, the stored two-term normalization ratios are:

L	Actual $v=1$ ratio
10^4	1.0001888352957300
10^5	1.0001190371666680
10^6	1.0000833930992132
10^7	1.0000616581233717

The proposer's first two quoted values, 1.002005 and 1.000396, come from other v rows at $L=10^4$, not this fixed- v sequence. The last two entries were correctly transcribed. The refuter detects this mismatch and the bounded replay confirms it.

The [replay receipt](#) records ten checks with three expected narrative failures. All pass/fail flags match the saved refuter. The rigour output matches exactly, including its sign bug. The insertion calculation at $L=1000$ agrees to less than 10^{-12} ; small platform-level floating differences are explicitly retained in the receipt and log. A receipt labelled successful replay does not turn the three failed assertions into passes. The check based only on a literal source formula string is a source-presence check, not an independent proof of the operator contract.

What the finite drift does and does not establish

F2's stored fixed-vector margins remain negative, from about -0.05199 at $L=10^3$ to -0.03124 at $L=10^7$. Its continuum value is approximately -0.014662375473371 , consistent with Astra's independently certified fixed-vector limit. The finite data were received as diagnostics; this intake does not rerun the large computations or upgrade them to interval certificates.

Slow prime-discreteness drift, even through 10^8 in a semi-continuum diagnostic, does not refute an asymptotic $O(1/\log L)$ bound without a proved constant and threshold. Nor can a few competing curve fits establish the rate. Large finite differences between full and clean operators do not contradict a proved vanishing difference in a fixed-family limit.

F1/F2 were working from the earlier Astra input checkpoint 97df092. Subsequent Astra work supplies a [fixed-family arithmetic transfer proof](#) and [independent review](#), with joint weak limits and bounded-operator control. For that fixed family an $o(1)$ limit suffices; an explicit $O(1/\log L)$ rate is not logically required. The F1 prose cannot be used to reopen that specific gap solely because its own rate proof is unfinished. Conversely, the fixed-family proof does not validate arbitrary growing feature families or claim that this negative trial refutes AH.

Reproduction and next use

From the public repository root:

```
OPENBLAS_NUM_THREADS=1 python3 fable/reviews/pr11-89393d5/recheck.py
```

This verifies all source hashes, copies the complete snapshot to a temporary directory, runs only the two named refuters (bounded at $L \leq 10^6$), and runs the separate pole check. Logs and the structured receipt remain beside this review. The code was read before execution; no model call or network access is part of these checks. Imported proposers' large main blocks are not executed.

For collaboration, the useful corrections are the Pi_4 coefficient, the refuter sign, the fixed- v table, and the distinction between fixed-family convergence and a quantitative rate. A repaired draft needs its own new commit and review. Remaining RMT/heat-flow proposer reports require separate mathematical audits; this intake gives no blanket endorsement.

Current report 64: Independent background and boundary objections to the same Fable snapshot

Collection: Fable intake — 89393d5 and 2073028, separate reviewed corrections.

Source: fable/reviews/pr11-89393d5/BACKGROUND_AND_BOUNDARY_REVIEW.md.

SHA-256: `336ae8258325f935c8613659e4921c8ce610193fbdde54d3fca8e0aafa8ada61` . Git blob: `04708ef3d0f1fe32b033961794c89f73c0948a7d` . Original bytes: 6033.

Independent background and boundary objections to the same Fable snapshot

Date: 2026-09-05. These findings came from the existing Astra coordination task's independent mathematical intake of commit `89393d5da61a45561ed199330c5b836f47fcd629` . Root checked the quoted source passages and the elementary algebra/counterexamples below. The original proposer files remain unchanged. This is a negative scope audit, not a complete acceptance review of every earlier constant or lemma.

1. A reversed inequality in the uniform CUE gap tail

In `r1_cue_background.md` , Proposition 3.3 regime 2 has $L > 4N^{1/3}$ and a preceding bound $16.47/N$. Its next step replaces this by $16.47 \cdot 64/L^3$ as an upper bound. The inequality goes in the opposite direction: $L^3 > 64N$ implies $1/N > 64/L^3$. The displayed uniform constant 1054, and the downstream constants 1055 derived from it, are not established by this proof as written.

A possible elementary repair is to use the deterministic fact $\delta_{\min} \leq 2\pi/N$ for N points on a circle of circumference 2π . The event is empty for $L \geq 2\pi N^{1/3}$. In the remaining part of regime 2, $1/N < (2\pi)^3/L^3$, so the preceding bound implies a constant below 4100 in place of 1054. If all earlier estimates are accepted, a conservative 4101 can propagate through the stated stiffness tails. This repair preserves the qualitative tightness route. This review has not certified all the preceding numerical constants, and does not silently edit the source or label 1055 verified.

2. The C–beta–E local density hypothesis has the wrong scaling

The displayed BB-LD compares the n -point correlation density to N^n times the unscaled angular Vandermonde to power beta with constants independent of N . It is already false for $\beta=2$, $n=2$. The exact CUE formula is

$$\rho_2(\theta, \theta + d) = \frac{1}{4\pi^2} \left[N^2 - \left(\frac{\sin(Nd/2)}{\sin(d/2)} \right)^2 \right] = \frac{N^2(N^2 - 1)}{48\pi^2} d^2 + O_N(d^4).$$

Hence $\rho_2/(N^2 d^2)$ grows like N^2 , contradicting an N -independent comparison. A microscopic statement must include scaled distances Nd , or the corresponding prefactor $N^{(n+\beta)n(n-1)/2}$, and specify its domain and the remaining marginal bounds. The source's later attempts at revised exponent counting do not fix the false starting definition merely by calling its normalization local. No general-beta density theorem is supplied by this audit.

The source's partition function formula is also wrong in its stated Lebesgue-angle convention, although it is marked unused. At $N=2, \beta=2$ direct integration of $2-2\cos(\theta_1-\theta_2)$ gives $2(2\pi)^2$. The displayed product formula gives $4(2\pi)^2$. This is an elementary normalization check, independent of a general Selberg integral.

3. Uniform one-point intensity does not control a selected close pair

The C-beta-E report attempts to obtain the one-sided density event needed for the selected pair from $\rho_1=N/(2\pi)$. That implication is invalid. Take a deterministic cluster of N distinct angles in an arbitrarily short arc, with a unique smallest pair gap, and rotate the entire cluster by a uniform angle. The one-point intensity is exactly $N/(2\pi)$, because each labelled point is marginally uniform. Every realization still contains the same tight cluster. The selected pair's nearby counting function, and its sum of inverse-square distances to other points, can be arbitrarily large.

This example is not C-beta-E. It shows precisely that the claimed one-point input is insufficient. A proof for C-beta-E must use its additional correlation structure, Palm estimates, or an independently proved uniform density event. Static rotational invariance supplies none of these by itself. Claims A3(b)/(c) cannot be accepted from the stated one-point argument.

The assertion that DBM is the only known route to the required density theorem is a methodological opinion. It is not a mathematical proposition with status proved, and this audit makes no exhaustive claim about the literature.

4. Static stiffness does not automatically persist to collision

The CUE depth Theorem 2 explicitly assumes $B^*=0$. The Theorem B repair requires stability on a time window and its κ_0 condition. Tightness of $S^*(0)/N^2$ is useful but is not the same assertion as a bound along the selected pair's evolution. The arc-sum lower bound in Lemma W may support a proof; its sufficient event H_C still has to be established for the selected pair before importing the depth conclusion.

The current intake therefore preserves both the conditional theorem and its missing event. It does not relabel the CUE or general-beta depth law as unconditional merely because a static background estimate was proposed.

5. A periodized window has an artificial gap

In `r1_levelB_barrier.md` the periodized version claims that a small normalized depth implies an actual zeta gap below $1/2$ without additional hypotheses. Periodizing a finite real window of length H_T introduces the wrap gap $H_T-(\text{last}-\text{first})$. The smallest circle gap can be this artificial gap, with no small consecutive gap among the original zeros. A circle depth theorem alone identifies a circle gap, not necessarily a genuine zeta gap.

For example, points $0, 1, \dots, n-1$ in a window of length $n-1+\epsilon$ have all internal gaps one and circular wrap gap ϵ . This does not make a statement about actual zeta windows; it disproves the unqualified finite-window inference. A non-wrap witness or a valid boundary construction is needed for transfer. The actual-flow non-return/truncation assumptions must also remain explicit.

Accepted use

The reports contain useful candidate lemmas and clearly exposed conditional inputs. The objections above identify specific places needing repair, not a refutation of every heat-flow idea. Until repaired and independently reviewed, these source status labels must not be used to claim a C-beta-E universality theorem, a true-zeta depth bound, or AH failure. The present main line remains the independently reviewed actual-prime estimates in Rounds 10 and 11.

Current report 65: Fable 2073028: repaired arithmetic, finite Fock bound and retained gaps

Collection: Fable intake — 89393d5 and 2073028, separate reviewed corrections.

Source: fable/reviews/pr11-2073028/INTAKE_REVIEW.md.

SHA-256: [51111bc72197254e92d3120f36b4f1278b8456eadd092c0d1a1362fe7b092f05](#) . Git blob: [260c5e94bbb68ae8c6b64ea2899e0825d0c75510](#) . Original bytes: 2671.

Fable 2073028: repaired arithmetic, finite Fock bound and retained gaps

Date: 2026-09-05. Source: Alpha-devbox PR11 commit [20730285c8f9a81539e0662c6e015023c2ed107a](#) . The 160 received public research files are a separate verbatim snapshot. Earlier snapshots and their reviews remain unchanged.

Two previous F1 objections are repaired: the leading coefficient is now $\pi^4 \sim 6\epsilon^4$ and the fixed- v table now uses its actual $v=1$ rows. The unchanged refuter still tests the negative of the derivative that tends to positive six. The new finite prime sum is honestly inconclusive; its cutoff is explained by the independently reviewed incomplete-gamma limit in [F1_REPAIR_AND_CUTOFF_REVIEW.md](#) .

The F3 claim that the field norm is infinite is false for the stipulated g and mass cutoff. [F3_MASS_CUTOFF_BOUND.md](#) gives a direct, independently checked sector proof: $||K|| \leq 2 \int_0^1 |g(u)|^2/u^2 du$. For $g=2\sin(\pi u/2)$, this is $4\pi \operatorname{Si}(\pi)-8$, approximately 15.2721. It also uniformly bounds every literal finite grid. This finite upper bound is above the desired $\pi^2/2$ threshold; it supplies no spectral-wall or arithmetic-transfer theorem. The first-bin constant-basis normalization in the source is not defined in $L^2(du/u)$.

The bounded local check imports only the source matrix builder and checks grids $M=6,8,10$, every occupation coefficient, the finite mass inequality and the scalar integral. It uses floating arithmetic and is not a rigorous eigenvalue enclosure. It does not rerun the large source/refuter sweeps. The source's memory check occurs after matrix allocation, and its $ru_{\max}ss/1024$ conversion assumes Linux units; on macOS that is not an MB conversion. No runtime or memory claim is inferred from those old numbers.

The source script and JSON model description also omit $1/\sqrt{j}$ from their opening formula, although the actual builder correctly includes it. All comparisons here use the builder's literal coefficients. Source bytes are preserved rather than silently corrected.

The general-beta background repair is reviewed separately in [CBETA_REPAIR_REVIEW.md](#) ; its unresolved finite- N formulas and conditional flow assumptions must not be treated as a proved CbetaE-to-depth theorem. The older CUE and periodized-zeta objections remain applicable where their source files did not change.

The new contribution is a finite boundedness theorem for the idealized operator and a precise cutoff diagnosis. The separate Astra fixed-family arithmetic transfer remains valid with its currently negative certified margin. Full arithmetic-to-Fock convergence, a sharp spectral threshold and the actual-zeta signed covariance gain remain open.

Current report 66: Independent audit of the F1 cutoff repair and F3 mass bound

Collection: Fable intake — 89393d5 and 2073028, separate reviewed corrections.

Source: fable/reviews/pr11-2073028/F1_F3_INDEPENDENT_AUDIT.md.

SHA-256: 5af35dcdf0120a5ab88c398f9b957e89587d7f12ea58ec207eb7e0f2f7d0c929 . **Git blob:** 512db81f74a04cf3f116fdac4780f66e1b1f6b95 . **Original bytes:** 6267.

Independent audit of the F1 cutoff repair and F3 mass bound

Date: 2026-09-05. Independent reviewer: Astra subagent `yau_flow`. Verdict: accepted for the stated mathematical scope. This is a bounded ordinary-proof audit of the two authored notes, not a numerical rerun, an audit of every Fable script, or a proof of an arithmetic-to-Fock limit. The exact reviewed hashes are listed below.

F1: the local remainder is uniform at the actual parameter

For the standard generalized-divisor Euler coefficients,

$$d_\ell(p^j) = \frac{(\ell)_j}{j!} = \prod_{r=1}^j \left(1 + \frac{\ell-1}{r}\right), \quad \ell = 16/15, \quad a = \ell^2 = 256/225.$$

The elementary harmonic-sum bound gives $d_\ell(p^j) \ll_\ell (j+1)^{1/15}$. Set $t = p^{-1-\varepsilon}$, with $0 < \varepsilon \leq 1$; then $0 \leq t \leq 1/2$. Thus

$$E(t) = \sum_{j \geq 1} d_\ell(p^j)^2 t^j = at + O_\ell(t^2), \quad E(t) = O_\ell(t),$$

uniformly for every prime and every such ε . The tail constant is bounded by the finite series $\sum_{j \geq 2} O_\ell((j+1)^{2/15})2^{-(j-2)}$; it does not depend on the prime cutoff.

Since $\rho = E/(1+E)$,

$$\rho(1-\rho) = \frac{E}{(1+E)^2} = at + O_\ell(t^2). \quad (\text{A})$$

Consequently the local error, summed against $(\log p)^4$ over any set $p \leq P$, is bounded uniformly in both P and ε by a constant times $\sum_p (\log p)^4 p^{-2} < \infty$. Multiplication by ε^4 makes this error vanish uniformly. No finite-prime experiment is needed to justify this step. In particular the coefficient is $6a$, not $6a^2$.

F1: the PNT argument controls the joint cutoff limit

Put $g_\varepsilon(x) = (\log x)^3 x^{-1-\varepsilon}$. The leading prime sum in (A) is

$$\sum_{p \leq P} (\log p)^4 p^{-1-\varepsilon} = \int_2^P g_\varepsilon(x) d\theta(x).$$

Write $R(x) = \theta(x) - x$. Given $\eta > 0$, PNT supplies a fixed x_0 such that $|R(x)| \leq \eta x$ for all $x \geq x_0$. The contribution below x_0 , including endpoint conventions, is $O_{x_0}(1)$ before scaling. For $P \geq x_0$, integration by parts bounds the scaled upper-part error by

$$\eta \varepsilon^4 (P |g_\varepsilon(P)| + x_0 |g_\varepsilon(x_0)| + \int_{x_0}^P x |g'_\varepsilon(x)| dx). \quad (\text{B})$$

These terms are uniform in the arbitrary cutoff:

$$\varepsilon^4 P |g_\varepsilon(P)| = \varepsilon z^3 e^{-z} = O(\varepsilon), \quad z = \varepsilon \log P,$$

and, using $x |g'_\varepsilon(x)| \leq x^{-1-\varepsilon} [3(\log x)^2 + (1+\varepsilon)(\log x)^3]$,

$$\varepsilon^4 \int_{x_0}^P x |g'_\varepsilon(x)| dx \leq 6\varepsilon + 6(1+\varepsilon) = 6 + 12\varepsilon. \quad (\text{C})$$

Here extending the positive comparison integral to $[1, \infty)$ uses the exact gamma integrals. Letting $\varepsilon \downarrow 0$, then $\eta \downarrow 0$, proves uniform convergence of this scaled PNT error over all cutoffs $P \geq 2$. This is an asymptotic uniformity argument; it supplies no explicit finite-data error constant without a quantitative PNT input.

The main integral is

$$\frac{\varepsilon^4}{6} \int_2^P (\log x)^3 x^{-1-\varepsilon} dx = \frac{1}{6} \int_{\varepsilon \log 2}^{\varepsilon \log P} t^3 e^{-t} dt.$$

This proves the authored equation (1) for every joint limit with finite $z_0 \geq 0$, including $P \rightarrow \infty$ but $z_0 = 0$, and for $z_0 = \infty$. It also permits $P = \infty$ at each positive ε , since the original prime sum then converges and the upper integration-by-parts boundary vanishes. If P remains fixed, the scaled finite sum tends to zero, as the note says. Thus no interchange of an uncontrolled joint limit is hidden in the proof.

The incomplete fraction at $z_0 = 1$ is exactly $1 - 8/(3\theta)$, not one. The cutoff warning and the explicit distinction from a finite-point error certificate are justified. I did not rerun the script or independently re-audit all its reported sign/output details.

F3: discrete bound and first-bin criticism

The stated coefficients give, by exact algebra,

$$B_M^2 = \sum_{j=1}^M \frac{[2 \sin(\pi j/(2M))/\sqrt{j}]^2}{j/M} = \frac{1}{M} \sum_{j=1}^M f(j/M), \quad f(u) = \frac{4 \sin^2(\pi u/2)}{u^2}.$$

The function f extends continuously to $f(0) = \pi^2$ and decreases on $(0, 1]$. Indeed $\sin x/x$ decreases on $(0, \pi/2]$, since the derivative numerator $x \cos x - \sin x$ is negative there. The right-endpoint sum is therefore bounded by the integral on each bin, proving the exact inequality

$$B_M^2 \leq \int_0^1 f(u) du = B_g^2$$

for every positive integer M , without an asymptotic quadrature claim. Integration by parts also checks the exact expression $B_g^2 = 2\pi \operatorname{Si}(\pi) - 4$. No certified decimal or spectral enclosure was recomputed.

The continuous sector argument uses $\|a(g)\Psi\|^2 \leq B_g^2 \langle \Psi, E\Psi \rangle$, not a particle-number bound. On the mass-cutoff space this gives a bounded annihilation map; its adjoint is the compressed creation extension. Different input sectors have different output sectors, so the infinite direct sum introduces no unaccounted cross terms. The resulting $\|K\| \leq 2B_g^2$, and the analogous discrete bound, are consistent with the stated conventions.

Finally, $\int_0^{1/M} du/u = \infty$. A nonzero constant on the first bin is not an element of the one-particle Hilbert space and cannot be normalized into a piecewise-constant Galerkin basis. A profile proportional to g is square integrable there because $g(u) = O(u)$, but that is a different construction and does not identify its compressed mass operator or creation coefficients with the literal grid model. The authored criticism is precise: it rules out the proposed normalization argument, **not** every possible convergent approximation. Uniform boundedness alone proves neither a sharp norm limit nor the desired spectral wall.

Reviewed snapshots and limits

- `F1_REPAIR_AND_CUTOFF_REVIEW.md` — SHA256
9f9cd67afb7d7d304f6eb42adcf34391972c54a9f49a2964be1d8fbb176da628.
- `F3_MASS_CUTOFF_BOUND.md` — SHA256
112eac7be3ed1294ea47c046eaa93720b170858350115960f8578e0301bb638b.

Author files were preserved. This review accepts the local Euler remainder, incomplete-gamma limit, discrete mass bound, and first-bin objection. It does not convert the numerical Fock fits into enclosures or establish any new zeta-zero or prime-gap theorem.

Current report 67: F1 repair: corrected coefficient, surviving sign issue, and cutoff scaling

Collection: Fable intake — 89393d5 and 2073028, separate reviewed corrections.

Source: fable/reviews/pr11-2073028/F1_REPAIR_AND_CUTOFF_REVIEW.md.

SHA-256: `9f9cd67afb7d7d304f6eb42adcf34391972c54a9f49a2964be1d8fbb176da628` . Git blob: `97d623fc0aaf6d71bb267dfbd1ca963f4e7a1fb8` . Original bytes: 4777.

F1 repair: corrected coefficient, surviving sign issue, and cutoff scaling

Date: 2026-09-05. Review of Fable PR11 commit `20730285c8f9a81539e0662c6e015023c2ed107a` . This review supersedes the corresponding old-version objections only where an actual repair is present. It does not reopen the independently proved Astra fixed-family arithmetic transfer.

Repairs that are present

The revised F1 text changes Pi_4 's leading coefficient from $6a^2$ to $6a$ and correctly combines it with Pi_2^2 to obtain a^2+6a . The displayed local inclusion probability has leading a $p^{(-s)}$, while its quadratic correction is summable with $\log^4 p$ near $s=1$. Together with the pole of ζ'/ζ , this is an adequate analytic derivation of the leading coefficient. A numerical experiment is not logically required to close that algebraic point.

The fixed- v normalization sequence is also corrected to the actual $v=1$ rows. Those two earlier text objections are therefore repaired in this snapshot. Fitting the revised finite values to a log-power rate remains numerical evidence, not a proved uniform error bound.

The new direct prime-sum script is honestly described in the revised report as inconclusive. Its code computes the local tail $E=\sum_{(e \geq 1)} d_{\text{ell}}(p^e)^2 p^{(-es)}$ and $\rho=E/(1+E)$, which is the correct inclusion probability. The script's opening docstring instead writes $\rho=1-1/E$ while defining E without its constant term; that docstring is inconsistent with the code. With E including the constant one, the latter formula would be correct.

The refuter's probe sign remains wrong

The revised prose still cites the unchanged `refute_F1_rigour.py` probe as an independent confirmation of positive six. It is not: zz3 is the third derivative of ζ'/ζ , so $\text{eps}^4 \text{zz3}$ tends to $+6$ and the script's negative probe tends to -6 . Its saved data are indeed -6 . The earlier separate Astra correction remains applicable. This does not invalidate the revised analytic coefficient calculation.

The finite cutoff has a predictable incomplete–gamma limit

The claim that primes merely up to $\exp(1/\epsilon)$ resolve the leading fourth-order pole is misleading. Write P for the prime cutoff and $z=\epsilon \log P$. PNT and the local expansion $\rho_p(1+\epsilon)(1-\rho_p(1+\epsilon))=a p^{(-1-\epsilon)}+O_a(p^{(-2-2\epsilon)})$ give, as ϵ tends to zero and P tends to infinity with z tending to z_0 in $[0,\infty]$,

$$\frac{\epsilon^4}{6a} \sum_{p \leq P} (\log p)^4 \rho_p(1+\epsilon)(1-\rho_p(1+\epsilon)) \rightarrow 1 - e^{-z_0} \left(1 + z_0 + \frac{z_0^2}{2} + \frac{z_0^3}{6}\right), \quad (1)$$

with the right side interpreted as one at $z_0=\infty$. In particular $\exp(1/\epsilon)$ corresponds to the fixed incomplete fraction at $z_0=1$, not the whole pole. Resolving most of the leading mass requires $\epsilon \log P$ large, not merely P exceeding $\exp(1/\epsilon)$.

For completeness, the proof is elementary. The local quadratic error has an absolutely convergent $\log^4$ -weighted prime sum for ϵ near zero, hence vanishes after multiplying by ϵ^4 . For the leading term use Stieltjes summation with $\theta(x) \sim x$ and weight $(\log x)^3 x^{(-1-\epsilon)}$. Splitting at a fixed large x_0 makes the PNT relative error uniformly small above x_0 ; the lower segment contributes $o(1)$ after scaling. Weighted integration by parts controls that error by a fixed multiple of the full positive gamma integral. Substituting $t=\epsilon \log x$ in the main integral gives $\int_0^z t^3 e^{-t} dt$. Sending the PNT error to zero proves (1), including a bounded or divergent z . If P stays fixed while ϵ tends to zero, the finite sum is bounded and its scaled value simply tends to zero.

For the actual fixed $P=2 \cdot 10^6$, $\epsilon=.125$ has z about 1.8136, and $\epsilon=.0625$ has z about .9068. The leading-mass fractions are correspondingly small. The report's assertion that the first cutoff already exceeds P at $\epsilon=.125$ also contradicts its own parenthetical $\exp(8)<P$. More relevant than that typo is the fourth-order weighting: the transformed density is $t^3 \exp(-t)$, whose bulk lies beyond $t=1$.

Equation (1) is an asymptotic cutoff diagnosis, not a numerical error certificate for the five stored finite points. It explains why that particular experiment is unsuitable as a full-pole confirmation. The existing analytic proof already establishes $6a$; larger prime scans are unnecessary for that purpose.

Scope retained

The latest text still describes its broader M2/coincidence transfer as incomplete. That is its own derivation status. Astra's later independently reviewed fixed-family $o(1)$ transfer remains a separate proved result and does not require the explicit rate missing from F1. Neither a repaired intermediate coefficient nor a finite drift fit gives a positive half-gap margin, a uniform full-operator limit, or a theorem about zeta pair correlation.

Current report 68: The mass cutoff gives a finite Fock operator bound

Collection: Fable intake — 89393d5 and 2073028, separate reviewed corrections.

Source: fable/reviews/pr11-2073028/F3_MASS_CUTOFF_BOUND.md.

SHA-256: [112eac7be3ed1294ea47c046eaa93720b170858350115960f8578e0301bb638b](#) . Git blob: [0c8d3ebc50bfb4d717ac85d05f163788f8b61cc2](#) . Original bytes: 5598.

The mass cutoff gives a finite Fock operator bound

Date: 2026-09-05. Author: root Astra. This is a self-contained bound for the idealized Fock operator defined in Fable PR11 commit [20730285c8f9a81539e0662c6e015023c2ed107a](#) . It corrects that report's assertions that the field norm is infinite or that multi-mode boundedness is still open. It does not prove a sharp spectral threshold, convergence of the displayed numerical discretization, or an arithmetic-to-Fock operator limit.

Statement and exact conventions

Let $h=L^2((0,1),d\mu)$, $d\mu(u)=du/u$, and let $\Gamma(h)$ be the usual bosonic Fock space with orthonormal sector convention. The energy/mass operator $E=d\Gamma(u)$ multiplies an n -particle wavefunction by $u_1+\dots+u_n$. Let $P=1_{(E\leq 1)}$, and work on the closed subspace $H_1=P\Gamma(h)$.

For g in h suppose additionally

$$B_g^2 = \int_0^1 \frac{|g(u)|^2}{u} d\mu(u) = \int_0^1 \frac{|g(u)|^2}{u^2} du < \infty.$$

Let $a(g)$ denote annihilation on the finite-particle core. Its restriction to H_1 extends to a bounded operator T of norm at most B_g . Its adjoint $A=T^*$ is the compressed creation operator $P a^*(g) P$ on its natural core. In the Fable convention A creates and A^* annihilates. Then

$$\|A\| = \|T\| \leq B_g, \quad \|\Phi\| = \|A+A^*\| \leq 2B_g, \quad \|K = A^*A + \frac{1}{2}(A^2 + (A^*)^2)\| \leq 2B_g^2. \quad (1)$$

The number operator is unbounded on H_1 , but that fact does not contradict (1).

Sector proof, including the infinite direct sum

For a symmetric n -particle wavefunction ψ_n , the annihilation formula is

$$(a(g)\psi_n)(u_2, \dots, u_n) = \sqrt{n} \int_0^1 \overline{g(u_1)} \psi_n(u_1, \dots, u_n) d\mu(u_1).$$

Weighted Cauchy--Schwarz, with weights u_1 and its reciprocal, gives

$$\|a(g)\psi_n\|^2 \leq nB_g^2 \int |u_1| |\psi_n(u_1, \dots, u_n)|^2 d\mu^{\otimes n} = B_g^2 \langle \psi_n, E\psi_n \rangle. \quad (2)$$

The equality uses symmetry, not a bound on n . Different input sectors map into different output sectors, so summing (2) yields

$$\| a(g)\Psi \|^2 \leq B_g^2 \langle \Psi, E\Psi \rangle \leq B_g^2 \| \Psi \|^2$$

on the dense finite-sector core inside H_1 . Therefore $a(g)$ extends continuously there to T . Annihilation only decreases total mass, hence T maps H_1 into itself. Its bounded adjoint is the closure of compressed creation. This justifies the domains in (1) without assuming an untruncated creation operator is bounded. The two final bounds follow from the triangle inequality and submultiplicativity. The algebraic identity $K = \Phi^2/2 - [A, A^*]/2$ is valid, but is not needed to prove boundedness.

For this proof, normalizable high-particle-number states can be made by placing each coordinate in a small positive interval inside $(1/(3n), 1/(2n))$. Their total mass is at most $1/2$ and their sector norm is finite. A state with all coordinates exactly equal to $1/(2n)$, as used in the Fable prose, is a delta configuration and is not an L^2 wavefunction. Replacing it by these packets proves only that particle number is unbounded. It does not give a lower bound on the field norm.

The actual sine kernel and every reported grid

For the stipulated $g(u) = 2\sin(\pi u/2)$, integration by parts gives exactly

$$B_g^2 = 4 \int_0^1 \frac{\sin^2(\pi u/2)}{u^2} du = 2\pi \text{Si}(\pi) - 4 \approx 7.636063674837709.$$

Thus the bound for K is approximately 15.27212734967542. This is much larger than $\pi^2/2$, so it cannot establish the desired spectral wall. Its role is to settle finite boundedness rigorously. Numerical values here describe the exact integral; no interval-certified spectral conclusion is inferred from them.

There is also a uniform bound for Fable's literal discrete model. Its modes have $u_j = j/M$ and creation coefficients $c_j = 2\sin(\pi j/(2M))/\sqrt{j}$. The same sector proof, with a finite sum in place of the integral, gives

$$\| K_M \| \leq 2B_M^2, \quad B_M^2 = \sum_{j=1}^M \frac{c_j^2}{u_j} = \frac{1}{M} \sum_{j=1}^M \frac{4 \sin^2(\pi u_j/2)}{u_j^2} \leq B_g^2. \quad (3)$$

The last inequality is exact: the integrand is decreasing on $(0,1]$, since $\sin(x)/x$ is decreasing for $0 < x \leq \pi/2$, and (3) is its right Riemann sum. Hence the norms of all these finite matrices are uniformly bounded, not just the computed examples. This does not assert that the sequence is monotone or that its norm converges to the continuous norm.

Remaining modelling and numerical limits

The literal c_j coefficients are a quadrature rule, not an exact piecewise-constant Galerkin projection for the measure du/u . In particular the first interval $(0, 1/M]$ has infinite du/u mass, so a nonzero constant on that interval cannot be normalized in h . One can instead choose a normalized profile proportional to g on that bin, or an infrared cutoff, but those constructions and their mass-cutoff error require an actual comparison proof. The report's first-bin normalization heuristic does not supply one.

The matrix formula is unambiguous and can be tested as a finite model. Lanczos computations are floating eigensolver results, not exact diagonalization or rigorous enclosures. Stable fits near 4.6456 are numerical evidence only, and agreement with a different trial family does not prove either operator equality or a sharp norm limit. The main arithmetic research remains independent of those unproved transfers.

The earlier Astra Schur majorant search already gives finite but nonsharp upper candidates in a different formulation. No repeat optimization is proposed here. The new contribution of this note is the complete mass-weighted sector argument and its application to the erroneous infinity claim, with the continuous and literal discrete models distinguished.

Current report 69: Independent follow-up: C β E background repair at 2073028

Collection: Fable intake — 89393d5 and 2073028, separate reviewed corrections.

Source: `fable/reviews/pr11-2073028/CBETA_REPAIR_REVIEW.md`.

SHA-256: `3f60e18e4994a78aeab9b381213c07cf4316005741ceed1ed614a79c1708a4af` . **Git blob:** `76df35f900cef2db52596cc4ad0df66090451412` . **Original bytes:** 9931.

Independent follow-up: C β E background repair at 2073028

Verdict: partial repair, not an accepted proof as written. The missing microscopic powers of N have been identified correctly, and the proposed triple-count exponents can be recovered under an explicit local three-point upper bound. However, the repaired text still contains a false finite- N identity, a false arbitrary-compact comparison, and a false uniform relative-error step. These defects prevent accepting the present BB-LD and Proposition 3.1 proof with their displayed status tags. The conditional replacement below is an ordinary mathematical derivation, not a proof of general- β BB-LD or a new C β E depth theorem.

Reviewed source: the preserved `PR11_2073028-source/r1_cbe_background.md` , with cross-checks against its companion `r1_cue_background.md` . Source SHA-256 values, exact line ranges, and the final review hash are in `CBETA_REPAIR_RECEIPT.json` . No source was edited, no Fable session was invoked, and no Monte Carlo run or large computation was repeated.

1. Coverage and change from the previous intake

The earlier independent `research-round3/fable_heat_sync_review.md` audited public SHA `a408e7050fffc74459b3c83fafa5ac03c8b7dea6` , where these written heat reports were absent. Its conclusions about that snapshot's evidence availability remain historically accurate. The present snapshot supplies the reports, including an explicit account of Fable's internal refutations. This follow-up independently checks the repaired mathematics; it does not infer validity from an internal refuter's completion status.

Issue	Current source lines	Independent conclusion
Missing N -rescaling in the old local density comparison	C β E 3 – 14, 95 – 109	The diagnosis is correct. The replacement needs a power-distance comparator or an explicitly restricted sine cutoff.
Exact finite- N CUE two-point function	C β E 99 – 101, 117 – 123	Incorrect as written: the sine-kernel limit was substituted for the finite- N kernel.
Claimed comparison on any compact rescaled interval	C β E 111 – 123	False if the interval reaches 2π . The upper comparison fails at an exact zero of the proposed comparator.

Issue	Current source lines	Independent conclusion
Uniform relative replacement of the third distance	CβE 202 – 218	False near an endpoint; the quoted N error exponent is also wrong.
Triple expectation exponents	CβE 194 – 253	Recoverable conditionally by the direct integral in § 3 below. Retain the finite-N correction and the order of limits.
BB-LD proved for β=1,2,4; general-β universality implication	CβE 137 – 168	The claimed proofs are not supplied. Neither the repaired two-point check nor an unquantified universality citation establishes the needed three-point bound.

2. Exact finite–N correction and explicit counterexamples

The companion CUE report, lines 52 – 55 and 145 – 146, correctly gives the ordered factorial correlation kernel. With $q=Nd$ and $N \geq 2$,

$$\rho_2(0, d) = \left(\frac{N}{2\pi}\right)^2 \left[1 - \left(\frac{\sin(q/2)}{N \sin(q/(2N))}\right)^2\right].$$

For fixed q , its normalized limit as $N \rightarrow \infty$ is $1 - \text{sinc}(q/2)^2$. At fixed finite N , the Taylor coefficient at $q=0$ is instead

$$1 - \left(\frac{\sin(q/2)}{N \sin(q/(2N))}\right)^2 = \frac{1 - N^{-2}}{12} q^2 + O_N(q^4).$$

Consequently the ratio to $(2 \sin(q/2))^2$ tends to $(1 - N^{-2})/12$, not $1/12$. This does not destroy the intended small- q scaling, but it invalidates the claimed exact finite- N verification.

At $q=2\pi$, taking $d=2\pi/N$, the exact normalized density equals 1, whereas $(2 \sin(q/2))^2 = 0$. This is an admissible circular separation for every $N \geq 2$. Thus no finite K gives the claimed **upper** bound there. The lower inequality alone becomes the harmless inequality $0 \leq \rho_2$; it is the two-sided comparison and its bounded-ratio assertion that fail. For q beyond 2π the unsquared sine can also be negative, making its arbitrary-real- β power unsuitable without absolute values.

A minimal definition repair is to use $\prod (Nd_{ij})^\beta$, with a fixed bounded microscopic cutoff, or to restrict the sine formulation to $s_0 < 2\pi$. The application below only needs an upper bound for $n=3$. For the two-point CUE case on any fixed $0 \leq q \leq s_0 < 2\pi$, the corrected finite- N ratio does have uniform positive finite bounds for $N \geq 2$: away from $q=0$ it follows from the strict triangle inequality for the geometric sum, compactness, and convergence to the sine limit; at $q=0$ the displayed Taylor coefficient is uniformly positive. This observation does not establish the $n=3$ general- β claim. A two-sided n -point assertion also requires $N_0 \geq n$.

For the second error, take an endpoint pair $x=0, y=u$ with $u=\varepsilon/2$, and put the third point at $z=-v$ with $v=\varepsilon^2/w$. When $\varepsilon/w \rightarrow 0$, z is closest to x , and

$$v' = u + v, \quad \frac{v'}{v} = 1 + \frac{w}{2\varepsilon} \rightarrow \infty.$$

Hence $\nu' = \nu(1 + O(\varepsilon/w))$ cannot hold uniformly on the integration region. Lipschitz continuity of sine controls an absolute error; it cannot supply a relative error at a vanishing sine. Independently, $\varepsilon/w = (L/c)N^{-1/(\beta+1)}$ is not $O(N^{-1})$ for any fixed $\beta > 0$. Finally, the largest distance is bounded by $w + \varepsilon$, not by $\max(w, \varepsilon)$: the cutoff check must require $c + N\varepsilon \leq s_0$. Taking fixed $c \leq s_0$ and then sufficiently large N resolves that condition.

3. A correct conditional triple estimate

Assume, for every triple with all $Nd_{ij} \leq s_0$, the following upper bound:

$$\rho_3(x, y, z) \leq \frac{K}{(2\pi)^3} N^{3+3\beta} d(x, y)^\beta d(x, z)^\beta d(y, z)^\beta. \quad (\text{U3})$$

This is the exact input needed here. It follows from a valid restricted sine upper bound using $2 \sin(t/2) \leq t$, but is not proved here for general β .

Let T count ordered distinct triples with $d(x, y) \leq \varepsilon$ and $\text{dist}(z, \{x, y\}) \leq w$. Require $\varepsilon + w < \pi$ and $N(\varepsilon + w) \leq s_0$. Apply the union bound for the two endpoint neighborhoods, anchor x on the circle, and use signed coordinates $u = y - x$ and $v = z - x$. Then

$$\begin{aligned} ET &\leq 2(2\pi) \frac{KN^{3+3\beta}}{(2\pi)^3} \int_{-\varepsilon}^{\varepsilon} \int_{-w}^w |u|^\beta |v|^\beta (|u| + |v|)^\beta dv du \\ &= \frac{2K}{\pi^2} N^{3+3\beta} \int_0^\varepsilon \int_0^w u^\beta v^\beta (u + v)^\beta dv du. \end{aligned} \quad (1)$$

All three factors remain present. The endpoint union bound deliberately overcounts overlapping neighborhoods and is therefore safe.

For a simple bound valid without ε/w tending to zero, replace $(u+v)^\beta$ by $(\varepsilon+w)^\beta$ in (1). With $\varepsilon = LN^{-1/(\beta+1)}$ and $w = c/N$ this gives

$$ET \leq \frac{2K}{\pi^2(\beta+1)^2} L^{\beta+1} c^{\beta+1} (c + LN^{-1/(\beta+1)})^\beta. \quad (2)$$

Thus the intended exponents survive, under (U3), after taking $N \rightarrow \infty$ at fixed L, c and then $c \downarrow 0$.

For the sharper leading constant put $\delta = \varepsilon/w$. Scaling both integration variables in (1) gives

$$ET \leq \frac{2K}{\pi^2} L^{\beta+1} c^{2\beta+1} J_\beta(\delta), \quad J_\beta(\delta) = \int_0^1 \int_0^1 s^\beta t^\beta (t + \delta s)^\beta dt ds.$$

For $0 \leq \delta \leq 1$,

$$J_\beta(\delta) \leq \frac{1}{(\beta+1)(2\beta+1)} \begin{cases} 1 + \delta^\beta, & 0 < \beta \leq 1, \\ 1 + \frac{\beta 2^{\beta-1}(2\beta+1)}{\beta+2} \delta, & \beta \geq 1. \end{cases} \quad (3)$$

The first case uses subadditivity of $t \mapsto t^\beta$. The second uses the mean-value estimate $(t + \delta s)^\beta - t^\beta \leq \beta 2^{\beta-1} \delta s$ on the unit square. This proves the leading asymptotic bound by integration, without the false pointwise comparison. It also makes its lack of uniformity in $c \downarrow 0$ before $N \rightarrow \infty$ explicit.

For $\beta=2$ the integral in (1) is exactly

$$\frac{\varepsilon^3 w^5}{15} + \frac{\varepsilon^4 w^4}{8} + \frac{\varepsilon^5 w^3}{15}.$$

Using the stronger CUE coefficient $\rho_3 \leq C_3(N) \prod d_{ij}^p$ yields $16\pi C_3(N)$ times this polynomial, precisely the bound stated in the companion CUE report, lines 226 – 240. The small symbolic check verifies this integral, the finite- N Taylor coefficients, and the $q=2\pi$ counterexample; it does not certify general- β density estimates.

If E_1 refers specifically to the closest pair, only $E_1 \subseteq \{T \geq 1\}$ is required. The equality asserted at C β E line 246 is generally false, because a different short pair may have a nearby third point. The correct inclusion still gives the intended Markov upper bound.

4. Status and remaining scope

Accept the normalization diagnosis and the conditional consequence (1) – (3). Do not accept the current proof of Proposition 3.1 or the “[P] for $\beta \in \{1,2,4\}$ ” BB-LD row as written. An explicit source or proof for the actual $n=3$ uniform upper bound remains necessary; $\beta=1,4$ are only recalled in the source, and ordinary weak local-process convergence by itself does not justify density control on the shrinking rescaled gap $N\mathcal{E} \rightarrow 0$. The claimed “exactly” equivalent universality formulation at lines 144 – 148 and its qualitative consequence at lines 164 – 167 have not been established by this report. This is a missing implication, not a claim that the required bounds are absent from all literature.

This bounded review does not audit Fable's full stiffness/heat-flow argument, the recalled Feng – Wei theorem range, the seed-sweep data, or the methodological claim that DBM is the only route. Those items receive no status upgrade here.

One additional elementary normalization error was noticed but is not used above: at $\beta=2, N=2$ the partition function written at lines 58 – 60 equals $4(2\pi)^2$, whereas direct integration of $2 - 2\cos(\theta_1 - \theta_2)$ gives $2(2\pi)^2$. The source explicitly says that formula is unused; it should not be copied into a handoff as verified.

The finite checks are reproducible with `python3 check_cbeta_repair.py` from this intake directory and require SymPy. No numerical research claim, ζ statement, or new famous-conjecture result follows from this review.

Source index and artifact receipt

This supplement includes 69 **complete source reports**. Its companion JSON index verifies 297 **associated repository objects** against the pinned Git tree. Code, arrays, logs, shorter READMEs and receipts are catalogued there without printing their complete machine data. The embedded proof/review text is not abridged.

Source checkpoint: `2a9ec8192ca8aed1cfdd1bad648c3a3a429e31ba` . Builder:
`tools/build_round6_14_handoff.py` . Index: `docs/handoff/ROUNDS_6_14_ARCHIVE_INDEX.json` .

The source reports and their earlier checks determine the scope of every mathematical claim. This assembly verifies preservation and provenance; it supplies no new mathematical experiment.

RESEARCH HANDOFF · SOURCE ARCHIVE · OVERNIGHT RESEARCH

ACUE 与黎曼零点 从谱反例走向 可证明的数论结果

完整研究接手档案

GPT-6 Astra × Claude Code / Fable

Bill (Qingyun) Sun · 2026 年 9 月 4–5 日

经审计主说明、公开历史 Markdown 与完整研究报告、算术相关性新路线
已证明、计算、猜想与失败分别标明 · Public research archive

内容索引

阅读方式与交接目标

1. 最重要的判断
2. 来源覆盖、时间线与忠实性边界
3. 符号、尺度与容易混淆的对象
4. 必须先修正的结论
5. 0.6725007：我们真正理解了哪一层
6. Nyquist 结构：当前最清楚的有限定理包
7. ACUE moment fiber 与显式反例
8. 其他反例为何值得保存
9. 第三轮十方向：逐项进展与失败价值
10. 零动力学：哪些已经成立，哪些是关键缺口
11. Marked response、导数与 LR 桥梁
12. 形式化与复现审计
13. 内部素数间隙工作： H_2 、 H_3 、 H_4
14. Signed sieve 与旧路线的 no-go
15. 新外部进展：它们实际改变了什么
16. 186 的结构与进一步下降的真实入口
17. 面向著名猜想的研究优先级
18. 更开放的探索方向：不预先限制在旧模型内
19. 给 Astra 的十二张任务卡
20. 第一次接手时的执行顺序
21. 可直接交给 Astra 的接手指令
22. 外部来源索引
23. 交付、检验与尚未完成的工作

第三轮更新：动态观测量、严格反例与 186 的结构性启发

已完成的两个动态结论

修复碰撞深度桥梁的两个反例

我们自己的来源勘误

将 186 的方法用于下一轮，而不是只提高数值精度

Fable 的实际回传及同步状态

最新补篇：从历史材料到当前可验证的数学状态

当前任务与真正的成功标准

1. 这轮最明确的新数学进展：固定素因子试验的算术传递
2. 完整算术算子告诉了我们什么
3. 值得永久保存的反向结果：中心高斯并不免费扩大支撑
4. Heat flow 路线的完成范围及与真 ζ 的距离
5. 素数间隙、0.6725 和形式化工作的实际启发
6. 下一轮有明确验收条件的探索

Current report 01: symmetric_prime_arithmetic_transfer

Current report 02: symmetric_prime_transfer_independent_review

Current report 03: residual_gram_round1

Current report 04: residual_gram_round2

Current report 05: centered_gaussian_mixed_moments

Current report 06: yau_flow

Current report 07: yau_flow_galilean_refinement

Current report 08: galilean-proof-audit

Current report 09: prime186-yau-independent-review

Current report 10: prime186

Current report 11: resummed_prime_profiles

Current report 12: transferable_tools_ritz_decision

Current report 13: dynamic_generator

Current report 14: dynamic_generator_independent_review

Current report 15: force_energy

Current report 16: force_energy_independent_review

Current report 17: new_attachment_bridge_audit

Current report 18: prime186_structural_frontier

Current report 19: fable_task001_sync_review

Current report 20: fable_heat_sync_review

Claim ledger 21: CLAIM_LEDGER

Research log 22: RESEARCH_LOG

历史原文完整附录

Historical source 23: final_verified_paper

Historical source 24: round3_synthesis

Historical source 25: impostors_paper

Historical source 26: depth_scaling_theorem

Historical source 27: newman_depth_note

Historical source 28: signed_sieve_nogo

Historical source 29: H2_H3_record_announcement

Historical source 30: prime_gap_survey

Historical source 31: stopping_times_paper

Historical source 32: joint_context_v2

Historical source 33: prime_gap_context

Historical source 34: signed_context

Historical source 35: codex_handoff

Historical source 36: rmt_zeta_survey

Historical source 37: rmt_zeta_popular

Historical source 38: README

Historical source 39: HANDOFF_GPT6_ASTRAL

Fable source; scope set by current audit 40: CLAIMS

Fable source; scope set by current audit 41:
FABLE_COORDINATION

Fable source; scope set by current audit 42:
r1_h2_interval_cert

Incoming proposal; claims require audit 43:
dynamic_generator_proposal

Incoming proposal; claims require audit 44:
heat_depth_classification_proposal

Incoming proposal; claims require audit 45:
hard_gap_depth_bridge_proposal

ACUE、黎曼零点、谱反例与素数间隙：GPT-6 Astra 完整研究接手档案

委托人：Bill (Qingyun) Sun

整理日期：2026 年 9 月 4 - 5 日；持续补充夜间研究

文档性质：带审计的研究交接、历史原文归档与新研究纲领

主要代码来源：[galpha-ai/Alpha-devbox](#) 的 [claude/riemann-zeta-random-matrix-udxp3f](#) 分支

当前主线：在 RH 下突破真实 ζ 零点的 $1/2$ 间距障碍，或得到明确的带宽外二点相关结果。下文研究优先级的历史表述以最新补篇为准；RMT 深度定理是支持路线。

历史基准：[ee61fba38925190916e556f87bc1ea83a502413e](#)；Fable 新交接版本：

[0f86c31d847178f0a774f70073ddc44206b32f98](#)；共同工作入口：[PR #11](#)

本文要保存的是已经得到的数学结构、正确证明、可复现实验、失败路线和仍然缺失的关键步骤。它没有宣称已经证明 RH、Montgomery pair correlation、GUE 猜想、排除 Tao 的 Alternative Hypothesis，或得到小于 186 的相邻素数间隙。历史原文中的更强措辞必须服从前半部的审计与限定。

阅读方式与交接目标

这份档案包含两种性质不同的材料。第一部分是本次重建、核验和重新组织的接手说明；第二部分逐篇保存主分支的历史 Markdown 原文，并附早期 Tao AH 说明 PDF 的文字转录。历史原文保留其当时措辞，因此能够看见认识如何变化，也会看见后来被撤回的判断。源代码、数值证书、数组、Lean 草稿和找到的完整任务记录同时保存在配套档案中。

如果 Astra 只读一遍，建议先读“必须先修正的结论”“ACUE moment fiber”“零动力学”“186 的真实结构”“研究目标与任务卡”，再按需要查历史原文。若要写论文，应逐条从审计表返回原证明与脚本，而不能把本档案的总结当作独立文献证明。

本次工作的概念对象是有限点过程、带权矩约束、谱证书、热流首碰撞时间与筛权二次型；核心不变量是概率非负性、精确矩匹配、带宽预算、归一化和算术输入的适用范围。交付结构采用一个主说明、完整原文附录、只读来源快照和一个独立审计程序。Python 只负责小规模复核与文档生成，不引入 Rust 或新的研究框架。验收标准是来源可定位、已知错误可辨认、实验可复跑、PDF 可读。暂缓事项是新的大型数值搜索、补全 Lean 工程、重跑 186 的重型证书及对外发表；这些均需要下一阶段实际研究，不能由文档整理冒充完成。

1. 最重要的判断

现有研究最值得保留的，并不是一个已经抵达著名猜想终点的证明，而是一套很具体的“信息不足”的数学模型。它们指出：即使二点信息、有限带宽内的所有平衡矩、部分高阶边缘分布、甚至某些完整 diffraction 数据完全相同，底层过程仍可能不同。这个认识能帮助选择真正能越过障碍的观察量；它本身没有提供把观察量用于黎曼零点的算术公式。

最成熟的有限结果集中在四组：Nyquist 模式对重数的精确控制；ACUE 矩纤维及其显式变形；homometry、magnetic holonomy 与 deck 非唯一性；圆周零动力学中的两体比较及其需要修正的局部化论证。最接近现有著名数论目标的可操作方向，则是把这些反例转化为一个具有固定强度的区分统计量，并找到它的显式公式或 Dirichlet 多项式实现。

这里有两个必须分开的突破入口。

1. 零点分布入口。在 RH 下，对一个固定带宽外区间的二点统计给出 GUE 值，或者证明一种足以排除精确定义 AH 的正密度偏离。这会是 Montgomery/GUE 路线上实质性的结果。有限反例主要负责告诉我们必须买到什么信息。
2. 素数间隙入口。从 186 的具体平方筛权、互补分解和支撑恢复不等式出发，找更小的 admissible tuple 或可认证的更强二次型。这里已有完整候选机制，宜与零点路线并行思考，但不应假设前者由后者直接推出。

用户希望追求历史级定理，这是研究方向的选择标准，而不是对成功概率的承诺。当前材料尚不足以判断某个著名猜想已“可解”。本档案给出明确终点、缺失引理和停止条件，避免继续积累无法触及终点的小模型。

2. 来源覆盖、时间线与忠实性边界

2.1 实际恢复了什么

恢复了所给分支固定提交中的 `research/riemann-rmt` 目录，包括 16 篇顶层 Markdown、Python 实验与验证程序、JSON/NPZ/NPY 数值材料、一个早期 AH PDF，以及 `riemann-impostors` 打包子目录。还读取了任务《联系 ACUE 与零动力学》，找到 11 个可取回的 turn，并将返回记录原样保存为 `sources/acue-zero-dynamics-conversation.json`。

共享上下文文明确记录过十个并行方向的分工，另有多轮综合稿、signed-sieve 分工和后续 Fable 系列脚本，因此确实不是单一对话中凭空产生的研究摘要。但是这次没有恢复所有原始子代理的完整对话、所有 scratchpad、早期被引用而未提交的 PDF 或外部 Lean 仓库的历史副本。不能把“取回分支中的全部材料”说成“找回了每一个历史子代理的全部输出”。

历史来源与当下判断分层如下。

标签	含义	应如何使用
外部论文结论	本次打开的作者论文或官方项目陈述	引用原始来源，并注意论文与证书的验证范围
有限证明	原稿给出可检查的有限或确定性论证	可继续整理成定理，但本次不是全面同行审稿
本次复核	本次运行了对应数值、枚举或符号检查	说明检查域、精度及未证明的量词
历史计算	原稿报告的数值，当前未完全复跑	不升级为精确证明或新纪录
条件桥梁	写得出结论，但关键假设尚未证明	下一轮应集中攻击这些假设
撤回或修正	与后稿、代码或本次审计冲突	禁止沿用旧的强结论

2.2 重建的研究演化

阶段	材料或提交	主要进展与状态变化
早期谱证书阶段	rmt_zeta_survey.md 、 tao_ah_notes.pdf	将新两项矩证书与 AH 反例联系起来；部分“最优上限”措辞后来撤回
8 月 15 日综合	6cc6e5f 、 final_verified_paper.md	Galerkin、Nyquist、covariance、homometry、holonomy、deck 和来源审计
下一轮十方向	joint_context_v2.md 、 round3_synthesis.md	cubic escape、带宽外检测、COM 家族、端点信息、derivative tower 等
8 月中旬筛法	8f576d7 至 9d837b7	P9 可分离截断权重、 H_2 / H_3 / H_4 候选证书与 tuple
8 月 18 日以后	bb7e7e6 至 07d1237	Newman depth、signed sieve no-go、三种统计分离与 LR 桥梁修正
后续统一	c8a6270 至 df0b581	可见性算子、clock Jacobian、marked response、局部横截性
形式化尝试	9074fd4 至 17242f9	修订 depth 比较；SymPy/Z3 检查；未编译 Lean 草稿
8 月 22 日打包	ee61fba	riemann-impostors 独立目录；部分数据与路径仍未随包提交
8 月 24 日任务讨论	《联系 ACUE 与零动力学》	Wronski/Plücker 总正性；否定任意 PSD 表示就能解决问题的论证
本次 9 月接手	本档案及审计文件	对照 186、FLT 形式化及 Lamzouri 新证明，重新安排研究目标

2.3 哪些内容不应算成我们的独立新发现

ACUE 由 Tao 的既有构造出发；sine process、Montgomery – Taylor 优化、Maynard 变分筛、de Bruijn – Newman 流、最小随机矩阵间隙及总正性理论都有原有文献。我们的历史稿中“new”通常指相对于当时内部工作流新得到的有限命题，不等于已经做过完整文献检索并确立了首创。后续发表必须单独做 novelty audit。

尤其不能把原稿写作期间独立重建的 0.6725007、公开论文中的 0.8362503，或已存在的矩阵不等式当作本项目首次证明的世界性纪录。

3. 符号、尺度与容易混淆的对象

用户所说 ACOE 的核心资料实际使用 ACUE：Alternative Circular Unitary Ensemble。COE 通常指 circular orthogonal ensemble，属于不同的对称类；不能把两者互换。除非另有明确的 $\beta=1$ 定义，本档案统一写 ACUE。

符号	本档案含义	必须保留的限定
CUE	Haar 随机酉矩阵的特征角过程	局部极限为 $\beta=2$ sine process
ACUE	在 $\mathbb{Z}/(2N)$ 上取 N 个点、按 Vandermonde 平方加权	离散 $\beta=2$ 模型；不是 $\beta \rightarrow \infty$ 的定义

符号	本档案含义	必须保留的限定
AH	零点归一化间隙集中到半整数的替代情景	必须注明强硬核、近似、密度或子序列版本
$p_k(C)$	$\sum_{x \in C} \exp(\pi i k x / N)$	Fourier 正负号影响复相位恒等式
p_λ	$\prod_j p_{\{\lambda_j\}}$	平衡矩的加权次数是 `
$D(P)$	有限圆周多项式向后热流的首碰撞时间	不是 ζ 的全局 Newman 常数
$\delta_{\min}$	初始最小圆周角间隙	角度、平均间隙归一化及 ζ 高度尺度不可混用
H_m	$\liminf (p_{\{n+m\}} - p_n)$	H_1 是相邻素数间隙； H_2 是三个素数跨度
M_k	指定 Maynard 单纯形试验类的变分常数	换支撑、估计或源域以后不再受原类上界约束

ζ 零点在高度 T 附近通常用平均间距 $2\pi/\log T$ 归一化；CUE 用 $2\pi/N$ 。把角度不等式转成 ζ 结论时必须明确窗口、边界、平移平均及误差，不能只作 $N \approx \log T$ 的符号替换。

4. 必须先修正的结论

这些修正不是附带小问题。它们决定下一轮是否会沿着不成立的桥梁继续推演。

历史说法或风险	当前判断	修正后的可用内容
0.6725007 是所有 bandwidth-one 方法的绝对天花板	撤回；后稿已明确纠正	它是特定两项矩/窗口优化的常数；更广泛障碍需要独立定理
0.702642 的 Nyquist 界改进了 ζ 的 0.6725	不成立	两者定义域与算术输入不同；前者是标记格点模型
finite fiber 越来越大，所以无限过程不唯一	缺关键步骤	要有固定局部可观测量上的一致分离与可传递矩条件
三点 rigidity 已对所有 N 证明	来源不支持	小 N 精确/数值结果保留；all- N 命题列为待证
$N=8$ fiber dimension 已定为 403	来源内部冲突	一处报告 399；需在统一约束、对称群、精确秩下重算
COM 调制族从 $N \geq 4$ 即有 $N-3$ 维	$N=4$ 反例	闭次数 $\leq N$ 约束下 $N=4$ 为 0 维； $N \geq 5$ 才符合该家族表述
ACUE 的平均 depth 约为有限常数/ N^2	无条件说法错误	clock 原子使每个有限 N 的无条件均值无穷；有限表格需条件化
最小间隙控制 depth 的双向等价	一般不成立	两体给下界；上界需要背景控制；clock 是反例
background 系数可由单端点 csc^2 控制	原上界错误	用精确差商、整段上确界或积分式

历史说法或风险	当前判断	修正后的可用内容
ACUE 是 $\beta=\infty$ 的 $C\beta E$	不成立	$\beta\rightarrow\infty$ 固定 N 退化到 clock; ACUE 仍有随机占位
marked response 有无条件 exact rank-two law	过强	链式分解精确; 修正项一般可有全谱秩
所有 Lean 步骤已经 machine verified	不成立	未编译、有两个实际 sorry; depth_ge 结论不涉及 D
P9 JSON 已是严格 interval 证书	提交代码不足以支持	高精度 mpmath 与 SAFE 因子不是向外舍入证明
signed no-go 排除所有 signed sieve 改进	超出证明类	仅在给定可计算观测和闭合试验类中成立
186 是 H_2/H_3 记录的直接延伸	不是同一个目标	186 属于 H_1 ; 机制是另一套扩大支撑/互补分解结构

5. 0.6725007: 我们真正理解了哪一层

5.1 两项矩窗口的连续最优化

在区间 $I=[-1/2,1/2]$, 考虑归一化 $\int_I v=1$ 的二次型

$$q(v)=\int_I v(s)^2 ds+\int_I\int_I |s-t| v(s)v(t) ds dt.$$

对应优化解为

$$v_*(s)=\frac{\cos(\sqrt{2}\,s)}{\sqrt{2}\sin(1/\sqrt{2})},\qquad q_*=\frac{1}{2}+\frac{1}{\sqrt{2}}\cot\frac{1}{\sqrt{2}}.$$

因此

$$\delta_{\mathrm{MT}}=2-q_*=0.6725007036794115\dots$$

这给了我们一个非常明确的靶: 仅仅在同一二次型里换更复杂的窗口, 不会超过该优化值。若要改进, 必须改变可用信息、证书类别或算术输入。历史稿把这个局部最优推广为所有 bandwidth-one 方法的绝对最优, 后来已自行撤回。

5.2 精确有限 Galerkin 公式与本次修复

对 n 个等长格子作分片常数离散化, 令

$$a_n=1+\frac{1}{3n^2},\qquad \theta_n=\arccos\bigl(1-\frac{1}{n^2a_n}\bigr).$$

则有限优化值为

$$q_n=\frac{1}{2}+\frac{a_nn}{2}\sin\theta_n\cot\frac{n\theta_n}{2}.$$

原 `verify_codex.py` 写成 `acos((1-n^(-2))/a_n)`，与正确公式不同。这导致闭式和渐近系数两项测试失败。本次独立审计将正确闭式与矩阵最优化在 $n=2,3,5,8,13,21,40,80$ 比较，最大误差约 4.49×10^{-14} 。这是数值回归检查，不代替原稿的有限证明。

令 $t=1/\sqrt{2}$ ，有

$$\delta_n = \delta_{\text{MT}} - \frac{\csc^2 t - \sqrt{2} \cot t}{24n^2} + O(n^{-4}).$$

使用 60 位高精度核对后，剩余项乘 n^4 趋近约 -0.02650666 （对应 $q_n - q_* - c/n^2$ 的符号约定）。这个小修复具有实际价值：未来做严格有限证书时，离散误差必须有可信的尺度和符号。

5.3 惯性、秩与重数的 secant 结构

共享上下文记录的有限谱证书把 on-line 零点贡献写成正半定和 $P = \sum m_j v_j v_j^*$ ，把 off-line 配对部分放入 Hermitian 修正 Q ，并用 Q 的正惯性数量控制损失。核心 clipped 函数为

$$k_c(m) = c^2 - ((c - m)_+)^2.$$

$c=2$ 时它把简单与多重零点区分为 3 与 4； $c=3$ 给 5、8、9 的分层。对适当的正交原子与配对块，证书可以取等，所以只持有 trace、Hilbert – Schmidt norm、rank 和 inertia 的论证确有内在信息限制。这里的等号模型只证明该证书的限制，不能自动证明所有显式公式方法的限制。

5.4 cubic 与更高 trace 的尝试

第三轮 D1 报告在 MT 最优窗口上得到三阶量 $\Phi_3 \approx -0.0117753128$ 。它的吸引力在于：如果能独立证明适当归一化负谱的统一范数上界，就可能把两项矩证书推进到更高信息阶。历史数值给出在某些额外上界下约 0.6796896 或 0.6844924 的条件输出。

这些数值不是已证明的 ζ 改进。尚缺的内容至少包括：三阶 trace 的完整算术计算、带宽与对角项账本、负谱统一控制、重数与 off-line 部分的重新解码。用 taper 在端点消失不能消除近零点簇造成的 $O(\log T)$ 行和；目前没有得到所需的统一常数。

共享上下文还记载 degree 10/12 的 Christoffel forced-rank 指标、精确连续矩、Gauss – Radau 对手、奇偶 Schur 层级和 $[0, 3]$ hard-edge 猜想。这些是值得保留的历史计算线索，但原始全部证明与计算文件不在当前分支，故不将其 80% – 95% 等数字列为 ζ 零点定理。

6. Nyquist 结构：当前最清楚的有限定理包

6.1 为什么必须引入标记

在 $\mathbb{Z}/(2N)$ 的 $2N$ 个位置上允许整数标记 $m_x \geq 0$ ，总质量为 N 。令

$$p_k = \sum_{x=0}^{2N-1} m_x e^{\pi i k x / N}, \quad V_k = |p_k|^2.$$

只要求开放的低频行 $EV_k = k$ ， $1 \leq k < N$ 。这与完全二元占位 ACUE 是不同的可行类。Parseval 与 Fourier 对称性给出精确关系

$$E \sum_x m_x (m_x - 1) = \frac{E |p_N|^2 - N}{2N}.$$

因此全部重数缺陷都汇入 Nyquist 模式。若再关闭 $k=N$ 这一行，左侧非负且期望为零，便强制每个标记几乎处处是 0 或 1。这是一个干净的结构定理，也解释了为什么不同约束集会得出完全不同的“简单比例”。

6.2 简单占位数的精确 slack

令 $s_1 = \#\{x: m_x = 1\}$ ，并定义 $h(j) = 0$ ($j \leq 2$)、 $h(j) = j(j-2)$ ($j \geq 3$)。原稿得到

$$Es_1 - \left(\frac{N}{2} + \frac{\csc^2(\pi/(2N))}{2N} \right) = E \sum_x m_x m_{x+1} + E \sum_x h(m_x).$$

两项 slack 都非负。它们的 Fourier 组合含因子 $1 + \cos(\pi k/N)$ ，在 Nyquist 点精确消失。于是

$$\frac{Es_1}{N} \geq \frac{1}{2} + \frac{\csc^2(\pi/(2N))}{2N^2} \rightarrow \frac{1}{2} + \frac{2}{\pi^2} = 0.702642367284 \dots$$

这个结果精确说明该标记模型中可用信息的后果。它没有为 ζ 提供新的无条件 70.26% 结论，也没有反驳 67.25% 的公开结果。literal ACUE 的占位本来就是二元的，其 $s_1 = N$ 。

6.3 等号面与失败的提升

等号要求所有标记属于 $\{0, 1, 2\}$ ，并且相邻位置不能同时非零。该等号面的配置数是 $\text{binom}(2N, N)$ 。将配置变为边电荷

$$q_x = m_x + m_{x+1} - 1$$

得到由 (-1) 、 $(0, 0)$ 、 $(+1, +1)$ 拼接的循环 tile 语言。若 Fourier 约定为负指数，则

$$q(k) = (1 + e^{+\pi i k/N}) \tilde{m}(k), \quad k \equiv 0.$$

原验证脚本使用负 Fourier 指数却写负相位因子，因而失败；本次改正后的最大误差约 2.81×10^{-14} 。

$N=3$ 已出现重要障碍：可提升电荷谱只有 $(0, 0), (3, 9), (9, 3)$ ，而把普通二元 DPP 样本 111000 直接滤波会产生 $(12, 0)$ ，不属于所需可提升结构。另一种独立添加装饰的尝试也失败，因为需要的交叉项是特定负相关，而独立化会把它消掉。

这个失败告诉我们：正确二点谱并不自动存在一个满足整数组合约束的正概率提升。all- N Nyquist 等号律仍是开放的有限组合概率问题。它可以成为严肃论文结果，但不能用它代替主数论目标。

6.4 covariance、内接球与修复

在 ACUE 下，开放行特征满足

$$\text{Cov}(V_k, V_\ell) = k^2 \mathbf{1}_{k=\ell} - 2(k + \ell - N)_+, \quad 1 \leq k, \ell < N.$$

历史证明给出最小特征值 1。原验证脚本在 $N=2, \dots, 6$ 复核通过。配合特征幅度控制，可推出矩多面体在目标点附近至少有半径 $1/(N^2 - N/2)$ 的内接球。于是若在等号面上先找到矩误差 $o(N^{-2})$ 的近似律，就有机会用少量混合恢复精确矩，同时只产生 $o(N)$ 的 slack。

这里已经把一难题压缩成了清楚的误差目标。接下来应寻找达到该精度的构造或证明它不可能；只做误差为 $o(1)$ 的模拟还没有达到修复定理的输入。

7. ACUE moment fiber 与显式反例

7.1 概率与矩约束

设 C 是 $\mathbb{Z}/(2N)$ 的 N 元子集, $z_x = \exp(\pi i x/N)$ 。ACUE 概率为

$$\mu_N(C) = \frac{1}{(2N)^N} \prod_{x < y \in C} |z_x - z_y|^2.$$

其核是 N 个连续 Fourier 模式的正交投影。对平衡分拆 $|\lambda| = |\nu| \leq N$, 矩值等于 CUE 的 $\delta_{\{\lambda\nu\}} z_\lambda$ 。固定这些矩后, 所有非负概率律组成一个 convex fiber。研究对象是这个完整概率 fiber, 而非仅仅 DPP 参数空间。

历史计算在旋转轨道坐标下报告 $N=3, \dots, 9$ 的维数为 $0, 0, 2, 10, 80, 403, 1804$ 。 $N=5$ 有 $Q(\sqrt{5})$ 上的精确多边形描述; 更大 N 的某些秩来自浮点 SVD。由于 $N=8$ 还出现过 399 的报告, 下一轮应统一坐标和约束后重算, 不能把整个序列都标为精确证明。

7.2 两点冻结、三点证据、四点自由

二点 rigidity 的有限证明来自完整周期 Fourier 对称与低频别名关系。三点统计在小规模精确/数值实验中也都被冻结, 但当前资料没有完成 all- N 定理。四点首先出现可见自由: 循环局部 word 01010101 的历史最大变化在 $N=6, 7, 8, 9$ 约为 $0.01019, 0.01232, 0.01375, 0.01514$ 。

这组数值很值得进一步研究, 因为它观察的是固定长度局部量, 而不只是维数。然而四个尺寸的正值或上升趋势不是 $\liminf > 0$ 的证明。也不能把 fiber 的增长称为“超指数”: 全部配置本身至多 $\text{binom}(2N, N) \leq 4^N$ 。

7.3 质心调制家族

定义质心模 N 为 $X(C) = \sum_{x \in C} x \bmod N$, 考虑

$$q_g(C) = \mu_N(C) g(X(C)), \quad g \geq 0.$$

在适当归一化下, $N \geq 5$ 时去掉 g 的 ± 1 Fourier 分量便得到一个 $N-3$ 维显式家族。其意义是把巨大的线性规划 fiber 中的一部分写成易解释的概率公式, 而不是依赖一个不可读的 nullspace 向量。

本次对 $N=4, 5, 6$ 直接枚举所有配置与全部闭次数 $\leq N$ 平衡矩。得到调制空间维数 0、2、3。 $N=4$ 的候选 $g(x) = 1 + 0.5 \cos(4\pi x/N)$ 在 $\lambda = \nu = (2, 2)$ 上产生恰约 2 的误差; $N=5, 6$ 相同测试族的最大误差分别约 1.54×10^{-15} 、 6.97×10^{-14} 。这同时证明了 README 中 $N \geq 4$ 版本的边界错误, 并支持修正后的 $N \geq 5$ 表述。

7.4 Fermi sea、相位叠加与可见性代价

质心乘法对应 Fourier Slater 态的平移。更一般的试验可写为

$$q_a(C) = \mu_N(C) \left| \sum_s a_s z(C)^s \right|^2,$$

其中系数自相关决定哪些交叉项出现。两片平移 Fermi sea 之间的可见性代价为 $\delta(N-6)$, 最大约 $N^2/4$ 。这提供了构造高阶矩不可见变形的方法, 也联系到量子纠错式的 scalar compression 条件。

若对所有允许观测 D_j 有 $P_V D_j P_V = b_j P_V$, 则任一单位向量 $v \in V$ 的 Born 概率 $|v(o)|^2$ 是真实非负律, 并满足这些矩。这里可以借用 Knill – Laflamme 的结构语言, 但必须始终写明观测类; 不能声称对所有操作都存在同样纠错能力。

7.5 parity sector 的强分离与弱极限陷阱

历史构造 $q_{\pm} = \mu_N(1 \pm (-1)^{\sum C})$ 给出互相奇异的兩律，同时在许多低阶边缘统计中不可区分。对于偶数 N ，两种 clock 配置都落在正 parity sector，故正 sector 有 depth 无穷原子、负 sector 没有。

ACUE 的两个 clock 总质量是 $2^{(1-N)}$ ；正 sector 的对应质量是 $2^{(2-N)}$ 。它们在每个有限 N 给出清楚分离，但质量随 N 指数趋零，因此不能由这个事件直接得到无限体积极限的非唯一性。这个例子应成为所有“有限反例自动传到 ζ ”推理的回归测试。

7.6 有限反例传到无限过程，究竟缺什么

一个可用的转移路线至少需要：固定强度的平稳化；局部计数紧性；在所选所有测试函数上的相关测度收敛或一致可积性；与目标带宽域严格匹配的矩恒等式；以及某个固定有界局部观测的差保持 $\geq \varepsilon > 0$ 。

最后一个条件排除了“维数很大，但所有固定局部窗口都看不见”的情况。前几个条件排除了质量逃逸、逐 N 改变测试函数及边界层问题。若格点占位始终二元，计数控制相对容易；若允许标记重数，必须重新检查矩与尾部。

Lagarias – Rodgers 的全阶 band-limited 条件与 Rudnick – Sarnak 风格的总带宽约束不是相同区域。逐坐标的 Fourier 立方体与 $\sum |\xi_i|$ 限制不能互换。因此一个满足平衡次数 $\leq N$ 的有限 family，必须逐个确认它传到的是哪一个无限问题。

8. 其他反例为何值得保存

8.1 Form factor 在投影 DPP 类中的刚性

对频率占位集 S 的平移不变投影 DPP，有

$$F(\gamma) = \frac{1}{2} |S \Delta (S + \gamma)|.$$

特别地， $F(1)$ 是 S 的 occupied runs 数量，所以 $F(1)=1$ 在该类中识别连续 Fermi sea。这个结果说明额外结构假设可以大幅降低识别难度。但如果没有先证明目标过程是这类 projection DPP，就不能用这个判别把一般 mimicker 排除。

8.2 Homometry：完整二点信息仍不够

在 $\mathbb{Z}/12$ 上，频率集 $\{0, 1, 4, 6\}$ 与 $\{0, 1, 3, 7\}$ 有相同 covariogram，却不是平移或反射；相应 DPP 在三点事件 $\{0, 1, 2\}$ 上的概率分别为 $(8-\sqrt{3})/432$ 与 $(8+\sqrt{3})/432$ 。原验证脚本复核通过。

另一个带标记例子是

$$u = (0, 0, 0, 0, 0, 2, 0, 0, 0, 0, 2, 2),$$

$$v = (0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 3).$$

它们总质量均为 6，完整循环自相关都为 $(12, 4, 0, 0, 0, 4, 8, 4, 0, 0, 4)$ ，所有 Fourier 强度非零，但简单位置数为 0 与 3，立方和为 24 与 30。这个反例非常直接：完整 diffraction 也不恢复生成元或重数分布。历史稿还提出 Fibonacci 非周期提升；该部分本次未重建全部证明。

8.3 Magnetic holonomy

对 g -cycle 的 Hermitian kernel $K_\varphi = aI + rH_\varphi$ ，低于 g 阶的 principal minors 看不见环通量，完整概率律的总变差距离满足

$$d_{TV} = (2r)^g \mid \cos(g\phi) - \cos(g\psi) \mid .$$

原验证在 $g=3,4,5$ 通过。它清楚展示了“第一次可见的相关阶数”等于一个几何回路长度。与此同时，即便看完整概率律，通量取向的某些符号仍可能被余弦隐藏。该模型适合作为高阶统计的必要性证据，不能直接认定 ζ 存在同样的 magnetic representation。

8.4 Deck、局部观测与大 fiber

固定 M 个循环位置和 L 个占位时， r -deck 的旋转不变 fiber 维数公式可表达为 necklace 数之差。在通常的 $r \leq \min(L, M-L)$ 范围内使用相应 inclusion-matrix 满秩结论。原验证复核了五组小规模参数。

这些例子与 parity sector 共同说明：大 nullspace 不等于局部可见自由。设计新反例时应同时提交其局部 observable 的变化幅度，而不只提交 rank 或维数。

8.5 Wilson 第四矩与边界条件

对没有额外长度四绕行圈的开放方格图，磁邻接算子的第四迹满足

$$\mathrm{tr} A_0^4 - \mathrm{tr} A_\phi^4 = 16 \sum_f w_f \sin^2(\phi_f/2).$$

原脚本用了小周期盒，存在额外绕行闭路，导致公式失败。本次改用开放 4×5 方格，得到左右两边同为 77.67769987482622。经验教训是：图论 trace 枚举的边界是定理假设，不是实现细节。

8.6 其他已探索但未接到算术的结构

历史材料还包括 Pontryagin 空间中的有限扩张、隐藏负特征值、Prouhet 型低矩伪装、Euler – Lee – Yang rank 障碍以及 Hurwitz 刚性。它们共同考察“多少谱信息能确定多少正性”。当前没有完成从这些有限结构到 Connes 型 ζ 算子或真实素数标记谱的证明。完整原稿保留在附录；下一阶段若使用其中某条，应先精确写出 arithmetic realization，而不是只引用类比。

9. 第三轮十方向：逐项进展与失败价值

以下按 `round3_synthesis.md` 的方向顺序重建。每个方向的数值是历史记录；除前文特别说明外，本次未完整重跑原子代理实验。

方向	得到的东西	没有完成的东西	下一轮应保留的教训
D1 cubic/ 负谱	MT 窗口三阶量、条件改进常数	uniform 负谱 cap 和完整算术 trace	多一个矩只有在代价与尾部受控时才有用
D2 带宽一最优	某些更大证书类的 LP 候选约 15/22	全局最优证明、严格 witness	不把一个优化器输出当成 universal ceiling
D3 首个带宽外 detector	$a+b=N+1$ 的 $\mathbb{E}$	$p_a p_b$	2 ，成本约 $1+1/N^2$
D4 缺陷恒等式	七项 deficiency、近等号刚性	额外算术解释	总缺陷在 fiber 上冻结，内部切片仍可能变化
D5 端点	点值的盲区、Cesàro cell 的可能信息	端点均匀估计	测试一个点与测试非零宽度区间有本质差异

方向	得到的东西	没有完成的东西	下一轮应保留的教训
D6 COM family	显式正概率调制与 $N=5$ 几何	全 fiber 分类、无限局部分离	显式 family 比纯数值 nullspace 更可迁移
D7 distinct/simple	distinct 解码及约 0.8362503 基线	附加正项的完整独立证明	改变计数对象也要对照现有文献
D8 derivative tower	$\cot/\csc^2$ 结构和若干秩计算	归一化、 $N=8$ 秩冲突	差一个归一化会改变 purported relation
D9 holomorphic derivative	某些加权量冻结，径向临界点可区分	ζ 对应物的算术矩	普通 P' 与 centered circular derivative 不同
D10 reflection/高阶	$N=7$ 去掉 39/80 手性方向；低 N detector	all- N 阈值与 shifted-correlation 输入	squared-modulus 信息仍不一定满秩

此外，Christoffel trace hierarchy 和 scalar moment code 来自前几轮共享材料，不应硬归入某个当前找不到原稿的子代理。已保存的上下文提供其出处线索，缺失文件清单则说明哪些结论需重新证明。

10. 零动力学：哪些已经成立，哪些是关键缺口

10.1 有限向后热流

若初始多项式 $P(z)=\sum_{j=0}^N a_j z^j$ 的根都在单位圆上，采用

$$P_s(z)=\sum_{j=0}^N a_j e^{sj(N-j)} z^j, \quad s \geq 0.$$

在首个重根形成之前，根角满足

$$\theta_j'(s)=-\sum_{k \equiv j} \cot \frac{\theta_j-\theta_k}{2}.$$

令 D 为正向 s 下第一次碰撞的时间，未发生碰撞则 $D=\infty$ 。这是一个有限 de Bruijn – Newman 类指标。它测量给定多项式距离失去简单圆周根的“时间深度”；它与 ζ 的全局 Λ 、有限窗口裁剪和实直线热流之间仍需独立桥梁。

10.2 两体下界的确定性核心

相邻间隙 g 的其他根贡献具有延缓碰撞的符号，因此

$$g' \geq -2 \cot(g/2).$$

对 $0 < g_0 < \pi$ ，纯两体方程的碰撞时间是

$$D_2(g_0)=-\log \cos(g_0/2) \geq g_0^2/8.$$

由最小初始间隙 δ 的 scalar comparison 得到

$$D \geq -\log \cos(\delta/2) \geq \delta^2/8.$$

这里应把 scalar solution 从最小初始间隙启动，再与各个间隙比较；不能对大于 π 的某个长间隙直接写实数 $-\log \cos(g/2)$ 。对于非 clock ACUE，最小间隙是 π/N ，故

$$D \geq -\log \cos \frac{\pi}{2N} \geq \frac{\pi^2}{8N^2}.$$

这项下界是可保留的主要确定性工具。它的方向很重要：知道 g 小并不能推出 D 小。

10.3 clock 原子与均值

clock 多项式只有首尾项，热流因子不改变它，故 $D=\infty$ 。ACUE 两种 clock 总概率 $2^{-(1-N)} > 0$ ，所以任何有限 N 都有 $ED=\infty$ 。

本次枚举 $N=4,5,6$ ，clock 质量分别为 $1/8$ 、 $1/16$ 、 $1/32$ ，符合公式。凡历史表格给出有限 mean depth，都必须明确它是在去除 clock 后条件化、有限截断，还是用其他统计量。历史 $N=6$ fiber 线性规划报告条件 $E[N^2D]$ 在约 $[1.3610, 1.4770]$ 变化，clock 质量也可从基准 0.03125 改到约 0.0975 ；这些应作为条件化实验保留。

10.4 局部化上界：修正后的真正任务

原稿试图用初始背景点的 $\csc^2$ 和控制局部差商，但所用单端点上界不成立。设某背景相对角为 x ，配对间隙为 g ，则准确的差商是

$$B(x, g) = \frac{\cot(x/2) - \cot((x+g)/2)}{g} = \frac{\sin(g/2)}{g \sin(x/2) \sin((x+g)/2)}.$$

当 $x=2\pi-0.15$ 、 $g=0.05$ 时， B 约 133.500132 ，而原候选上界 $0.5 \csc^2(x/2)$ 仅约 89.055743 。该例即使要求 g 不超过附近其他间隙，也能出现。正确控制应使用整个角区间上的导数上确界或上述精确式。

更关键的是，这个控制要沿着直到碰撞的时间窗成立，而不仅在 $s=0$ 成立。若某短间隙在整个相关时间窗内满足

$$g' = -2 \cot(g/2) + B(s)g, \quad 0 \leq B(s) \leq B_{\max},$$

且 g 始终小，则取一个有效常数 C 使 $2 \cot(g/2) \geq 4/g - Cg$ ，可得 $y=g^2$ 满足

$$y' \leq -8 + 2(B_{\max} + C)y.$$

当 $(B_{\max}+C)\delta^2 < 4$ 时，比较方程给条件上界

$$D \leq -\frac{1}{2(B_{\max} + C)} \log \left(1 - \frac{(B_{\max} + C)\delta^2}{4} \right).$$

如果别的点先碰撞，只会使全局 D 更早；但在所追踪间隙上使用此式，要证明控制在该时间内保持。于是 $B_{\max} \delta^2 = o(1)$ 才能把上下界压到 $D=(1+o(1))\delta^2/8$ 。这是真正需要证明的局部化引理，而不是已经完成的“二体定律”。

10.5 随机矩阵最小间隙与不合法的外推

CUE 最小角间隙尺度为 $N^{-(4/3)}$ 。按常见圆周归一化，其极限尾律涉及 $\exp(-x^3/(72\pi))$ ；使用其他展开尺度会改变常数。Feng – Wei 所引用推广的 β 范围是正整数，不能擅自写成所有 $\beta > 0$ 。Feng – Wei

如果再有上一节的高概率局部化，才可推断 CUE depth 的典型尺度 $N^{-(8/3)}$ 。历史模拟支持这种想法，但最小间隙定理本身没有证明 depth 定理。

另一个常见错误是把固定 x 的弱收敛尾律直接代入随 N 增长的 $x=\pi N^{1/3}$ ，宣称某事件以 $\exp(-\pi^2 N/72)$ 衰减。没有 moderate-deviation 估计不能这样做。仅由紧性可以得出对应概率趋零，不能附带这个指数速率。

10.6 单缺陷常数与 ACUE 极限中位数须分开证明

历史研究对一个特定单缺陷配置写出

$$P_N(z) = \frac{(z-1)(z^N+1)}{z - e^{-i\pi/N}}$$

并得到缩放极限函数

$$G_s(u) = 2 \cos(u/2) - 2\pi \int_0^{1/2} e^{s(1/4-y^2)} \cos((\pi+u)y) dy.$$

其首次双根的历史计算约为 $s^*=1.419640342$ 、 $u^*=1.812942145$ ，相应 $8s^*/\pi^2 \approx 1.150717118$ 。它提供了一个可研究的显式局部模型，而不是 ACUE 全体样本的普遍常数。旧表中的 ACUE median 在 $N=8,9,10$ 已分别约 1.418216、1.41520、1.412774，这破坏了旧单调收敛证据，不能据此认定极限中位数。但这些有限数据本身也不能证明未来极限不等于同一个常数；单缺陷极限和整体随机极限都仍需各自的证明。

10.7 算子统一带来的结构

由核 $1/(2s \sin^2(\pi k/N))$ 形成的循环算子，其 Fourier 特征值为 $\delta(N-\delta)$ 。同一谱既出现在 Fermi sea 交叉项的可见性代价，也出现在 clock 附近零动力学的 Jacobian。这是现有材料里最有启发性的结构连接之一。

符号必须明确：在向后热流中它描述不稳定增长，不能称为趋于平衡的 relaxation。对固定 Fourier mode， $N^{(-1)}L_N$ 的极限是 $[D]$ 型半拉普拉斯/Poisson 生成元，而非普通 Laplacian。与 Casimir、基本权或 affine Weyl 的关系需要指定表示和归一化，不能凭谱形状宣称所有对象规范同构。

11. Marked response、导数与 LR 桥梁

11.1 为什么给零点加标记

无标记点过程会忘记本征向量、边界条件和生成元。若有 Hermitian G 与一个指定向量 u ，可对 $G+\eta uu^*$ 作 rank-one 扰动，再经 Cayley 变换得到圆周根。首碰撞时间的导数是一个可能比普通 gap 更敏感的标记统计量。

但是对 ζ 来说，最难的是指定一个自然、可计算且不依赖待证结论的 G 与 u 。任意把零点塞进对角矩阵并附加向量，只制造了额外数据，不能声称找到 ζ 的真实 random-matrix 结构。

11.2 “exact rank-two law”的正确读法

写 $D=\kappa\delta^2/8$ ，微分自然给

$$D' = \frac{\kappa\delta}{4}\delta' + \frac{\delta^2}{8}\kappa'.$$

这是精确链式法则。若只保留首先碰撞那一对的 gap response，会出现 rank-two 外形；但 κ' 可能依赖所有本征角，故完整响应一般不是 rank two。把 κ 从同一 D 的输出反算后再得到 correlation=1，验证的是恒等重写，不是独立预测能力。

对 $\text{Cayley}(\lambda)=(\lambda-i)/(\lambda+i)=e^{i\theta}$ ，本次符号计算得到 $d\theta/d\lambda=+2/(1+\lambda^2)$ 。若用相反 Cayley 约定符号会变；不说明确定直接沿用负号会使所有 marked coefficients 反向。

11.3 局部横截性尚非全局分离

历史实验显示 `marked/depth` 梯度在矩约束行空间外可能有残量。这是好的诊断：该统计局部能看见某些 fiber 方向。要把它变成实际分布分离，还要沿该方向保持概率非负，并给出有限幅度、误差受控的差。要用于 ζ ，又多一层算术可观测性。三件事不能合为一个“已解决识别”的结论。

11.4 LR 与 AH：不要偷换量词

若某模型具有真正的统一硬核下界 $\delta \geq 2\pi c/N$ ，两体比较给

$$D \geq \frac{\pi^2 c^2}{2N^2}.$$

因此一个独立的严格 depth 上界可以约束硬核参数 c ；反过来，小 gap 上界不能自动给 depth 上界。clock 是最简单反例。

更强的障碍是：在包含 ACUE 的全部 bandwidth-one mimicker 上，不可能证明一个普遍的 $N^2 D < \pi^2/8$ 上界，因为 ACUE 自己满足相反下界，且还有无穷原子。若要证明 ζ 具有更小 depth，必须用 ACUE 不满足的额外算术信息。

还要区分 AH 的“几乎所有间隙接近半整数”与“每个间隙都有硬核下界”。即使发现无限多个很小间隙，也未必排除允许零密度例外的 AH。一个相对稳定的目标是某个离开半整数集合的紧区间内出现正比例的归一化间隙，并证明这与所采用 AH 定义矛盾。涉及 Landau – Siegel 零的推论必须逐条核对原定理量词，不能把任何 AH 反例直接升级成无 Siegel 零定理。

来源勘误（2026-09-05）：Lagarias – Rodgers 的两篇论文必须区分。1905.12123 提到 0.606894 的正比例小间隙结果；其伴随论文 1907.03391 § 5 明确定义带宽 1 模拟过程的硬核上确界 μ ，证明构造给出 $\mu \geq 1/2$ ，并提出修改既有方法应可得到 $\mu \leq 0.606894$ ，但没有在该处证明这一上界。此前本档案遗漏伴随论文，因而对这两个对象的联系批评过度，现予纠正。 $\mu=1/2$ 仍是未解决目标。历史 gap 纪录还需排除已撤稿结果。硬核参数的原始来源，AH 伴随论文。

12. 形式化与复现审计

12.1 原脚本的真实结果

原样运行 `verify_codex.py`，记录为 39 PASS、4 FAIL、1 SKIP。程序以退出码 0 结束，因为 `ok` 只打印结果而不抛错。不能用退出码 0 或终端最后一句 Done 判断全部检查通过。

四项失败分别是：Galerkin 闭式、其二阶渐近项、charge filter 复相位、周期盒 Wilson 第四迹。GOE counterfactual 常数的核归一化无法由稿件重建，因此原程序自己跳过。该 0.559871849 数值在本档案中不作为已核验数学事实。

本次编写 `audit_research.py`，通过 21 项限定范围检查，输出 JSON 和文本日志。它修复了前三类可定位的验证问题（Galerkin 两项属于同一公式问题），枚举 $N=4,5,6$ 的 COM 与 clock 质量，验证两个 tuple 的 admissibility，检查 Cayley 符号、background 反例和 Lean 结论。注意，其中若干 PASS 的含义是“成功检出了历史错误”，不是“历史定理通过了证明检查”。

12.2 Lean 文件比 README 更弱

当前 Lean 草稿未编译。有两个实际单独 `sorry`：一个在 `neg_log_cos_ge` 的单调性积分步骤，一个在 `two_body_solution`。更严重的是 `depth_ge` 的最终命题只为

$$\delta^2/8 \leq -\log(\cos(\delta/2)),$$

没有出现 D 。其 $\text{hd} : \forall i, g \ i \ D=0 \rightarrow \text{True}$ 也没有定义首次碰撞。故即使该文件某天无错误编译，当前这个 theorem statement 仍没有表达研究所需的 depth 下界。

共享 README 把它描述为两个 ODE sorry，与文件本身不一致。下一轮第一步应修 theorem statement 与停止时间定义，再写 proof。形式化必须检查“证明了什么”，不能只检查是否没有 sorry。

12.3 Sympy/Z3 的证据边界

`verify_theorem_steps.py` 与同系列文件有一些有效恒等式检查，但也存在直接 `ok(..., True)`、有限 Taylor 系数、以及先假设目标符号再检查否定不可满足的步骤。这些不能合成“所有解析步骤机器证明”。符号程序可以验证代数恒等式，SMT 可以验证准确编码的命题；ODE 存在性、首次碰撞、全时间比较及极限交换必须另证。

12.4 可复现性缺口

若干脚本硬编码旧 `/tmp/claude-0/.../scratchpad` 路径，依赖当前未提交的模块或数据。打包 README 提到的部分 `data/` 不存在。 H_4 的 56,000,000 维 profile 没有对应 NPZ； H_3 的完整 tuple 在顶层而不在所有子包位置。早期 [FSC]、[A23]、[FSI] 等引用并不都有可取回原文。

本档案保留原文件，另外提供独立审计，而没有静默改写历史证据。未来可单独做一个复现修复 PR：相对路径、固定依赖、assertion 退出码、精度声明、数据哈希与一条完整命令。现在不需要为此改动 Alpha-devbox 控制器代码。

13. 内部素数间隙工作： H_2 、 H_3 、 H_4

13.1 必须先辨认目标

$H_m = \liminf(p_{n+m} - p_n)$ 。 H_1 是两个相邻素数的间距； H_2 是三个素数的跨度。历史公告标题中把 H_2 描述成“中间有两个素数”的措辞不准确。内部 $H_2/H_3/H_4$ 数值与外部 $H_1 \leq 186$ 不直接比较大小。

目标	历史候选上界	k	历史报告的变分下界
H_2	173438	15856	8.013326752751
H_3	13859802	923601	12.006666706750
H_3 旧备份	14505780	较早参数	见原稿
H_4	1120662828	56000000	16.0655 左右

这些是仓库中的内部候选成果，本次没有把它们认证为新的公开世界纪录。原因既有严格数值证明的缺口，也有最新文献比较尚未完成。 H_2 tuple 的 admissibility 已由本次程序对所有 $p \leq k$ 完整检查；这只核对了组合输入，不等于核对整个筛法不等式。

13.2 P9 引擎的有价值结构

在标准单纯形上试验

$$F(t_1, \dots, t_k) = \prod_{i=1}^k g(kt_i) \mathbf{1}_{\sum t_i \leq 1}.$$

令 $c_2 = \int g^2$ ，独立随机变量 X 的密度为 g^2/c_2 ， $S_k = X_1 + \dots + X_k$ ， $G(u) = \int_0^u g$ 。则可将高维 I 、 J 化为一维截断和概率：

$$I_k = k^{-k} c_2^k \Pr(S_k \leq k),$$

$$J_k = k^{-k-1} c_2^{k-1} E[G((k - S_{k-1})_+)^2].$$

由 layer-cake、Berry – Esseen、Chernoff 或 one-big-jump 等尾部估计得到下界。用 $\Pr(S_k \leq k) \leq 1$ 放大分母给保守下界是合法的；把每一段尾部误差保守地计入同一 ratio 才有证书意义。

这套分解把很高维试验压到可计算的一维概率对象，是应保留的工具。它的优势在于结构和可认证性，而非自动穿越标准 simplex 的上界。

13.3 严格证书尚未闭合

`p9_certify_hp.py` 使用 `mpmath` 约 50 位精度，并乘 `SAFE=1+10−30` 一类因子。高精度输出不自动是 rigorous enclosure；任意安全因子也不证明所有积分、特殊函数和截断误差都向外舍入。JSON 中的 `lb` 字段名同样不是证明。

因此下一轮若要公布 H_2 / H_3 / H_4 结果，需要采用真实区间算术或精确有理上下界，验证所有尾部与积分误差，并提供独立 checker。还需要复核 tuple 文件、权重 profile、参数及每一步筛法假设确为同一个证书。应保留现有候选的正余量，但不能把尚未验证的最后几位当作已证实事实。

14. Signed sieve 与旧路线的 no-go

14.1 凸序与一阶计数信息

旧研究把 prime-count 目标抽象成随机整数与可计算线性/凸上界。两点极端分布说明：如果只有这些数据，许多检测阈值无法改善。这是一个指定信息模型中的最优性证明。它适合用来筛掉重复的变分幻想，但不能约束新引入的分布估计、扩大支撑或可计算交叉项。

14.2 w_+ 支配的使用条件

把 signed 权重拆成正负部分，常可以得到清楚的 debt identity 或点态支配。若 w_+ 仍属于可接受、可计算的试验类，这能证明某些 signed 变形没有收益。但若正部化导致支撑、可分离性或算术可计算性失效，就不能从点态比较直接得出同一个归一化变分问题的 no-go。

历史精确价格

$$\beta_* = \frac{23051796480}{10991046857} = 2.097324920903119 \dots$$

与另一无界版本的约 2.03265 应区分。部分 prose 把“价格低于 1 的相变”与这个大于 2 的 β_* 混在一起，需要以定义和精确数值为准。本次检查确认 $\beta_* > 2$ 。固定 signed mass 的归一化、未归一化目标及固定特征类分别是不同优化问题。

14.3 为什么 186 恰好说明旧 no-go 的作用域

标准 simplex 的上界 $M_k \leq k \log k / (k-1)$ 会阻止某些 $k \approx 40 - 50$ 的原始策略。但 186 不在这个原试验类中：它扩大外层支撑，再用不同内层源域控制 prime marginals 和混合项。因此它不是“推翻数学 no-go”，而是改变了 no-go 的假设。

历史另一条 conditional 路线把分布指数 θ 的提升换成更小 H_1 候选，如 $4/7 \rightarrow 130$ 、 $7/12 \rightarrow 114$ 、 $3/5 \rightarrow 94$ 、 $5/8 \rightarrow 80$ 。这些均以特定 CRT 变化剩余类和带权源域估计为前提。不能仅因某个 BFI 型定理也出现同样 θ ，就认定已经满足这些假设。

15. 新外部进展：它们实际改变了什么

15.1 Anthropic 的 FLT 项目

2026 年 9 月 4 日公开的项目主要成就是把既有费马大定理证明的复杂依赖大规模形式化，涉及约 1300 万行 Lean 与约 29500 个中间结果。它不是新的初等 FLT 证明，也没有直接给出一个可应用于 ζ 或 prime gap 的新解析估计。项目说明，代码仓库

对当前研究的迁移价值在方法上：明确 statement、拆分依赖、将有限证书与解析输入隔离、让独立 checker 验证，并检查最终定理的量词。我们的 `depth_ge` 正说明为什么 statement 审查与去除 `sorry` 同样重要。

15.2 ζ 的 0.6725 进展与最新简化

Anthropic 8 月公开的结果给出至少 `0.6725007...` 比例的零点是简单且位于临界线上，并用有限谱/惯性结构组织证明。这里“简单且在线”与仅“在线”的旧纪录不可随意混用。官方说明，作者论文

本次发现 9 月 3 日 Lamzouri 的新预印本给出该结果的更直接 Hilbert 空间证明，也得到相应 distinct-zero 比例约 `0.8362503`。它是证明简化，不是把 0.6725 又提高成更大的比例。Lamzouri

这改变我们的工作排序：在继续构造庞大的有限谱机器之前，应把每个候选新输入代入更短的 Hilbert 空间形式，判断真正缺少的是否仍是同一个额外矩或 off-line 控制。复杂形式化工具值得用在最终正确且必要的桥梁上。

15.3 用户提到的 Tao 赞许路线

没有足够信息唯一确定用户指的是哪篇文章。Guth – Maynard 的大值估计与零密度改进，以及 Tao 对其计算探索的讨论，是一个可能对应的方向。Guth – Maynard, Tao 的讨论

应把它作为可供核对的相关入口，不假装已经确认用户所指。零密度控制、临界线上简单比例、局部 pair correlation 和筛法剩余类分布分别要求不同的平均与一致性；不能仅因为它们都涉及 ζ 就直接串成推出链。

15.4 186：来源、精确目标与证书边界

用户两个短链接分别解析到 8 月 30 日的 *Short gaps* 论文与配套 abridged research record。论文声称 `DHL[40,2]`，从而 `H1 ≤ 186`。它是当前接手应采用的最新来源基线，而不是继续把 246 当作未更新的基线。论文，研究记录

相关公开仓库的 Lean 工程仍明确列出三个项目输入：两个 Kloosterman 型界和 `physical_integral_bounds`；数值脚本另外依赖指定的 Python-FLINT 环境及 signed convolution 修复。条件化的 kernel acceptance 与“所有解析及数值输入都已由 Lean 从零解除”不是同一件事。这并不否定论文的无条件数学主张，但决定我们能诚实声称重验了哪一层。PrimeGaps186

本次没有重跑其重型区间证书，也没有重编整个 Lean 工程。已独立核验论文所给 40-tuple 的 admissibility 与直径 186；它只证明最后的组合输入。

16. 186 的结构与进一步下降的真实入口

本节以论文中的二次型作为下一阶段工作的接口，避免用旧 simplex 的单一 $M_k > 4$ 判据替代新问题。下面列的是应读取、重建和优化的结构，不声称本次已经改进任何常数。

16.1 支撑扩大与互补分解

外层 divisor sum A 的平方保持非负，允许比旧内层更大的支撑。prime 侧使用不同内层 B_0 、 B_1 控制各自可计算的混合项，保存平方余项和 minorant 的正负部分账目。关键不是简单“使用负权”，而是每一项都有合法 pointwise 比较与适用源域。

在研究记录中，多次失败都来自：把缺失平方项扔掉、忽略 minorant 在合数上的负部分、把不同源域的误差合并成过粗 bound，或把固定剩余类的定理用到随 tuple 坐标变化的 CRT 类。正确方案必须同时通过代数符号和算术适用域两种审计。

16.2 三重 dense divisibility 的角色

论文将 divisor factorization 转成可检查的素因子条件，并利用互补因子结构控制混合模数。真正出现的是 $[d, e]$ ，不是任意假定互素后的 de 。共享素因子改变预算，故后续优化不应以独立乘积启发式代替 lcm 条件。

16.3 正余量很小，因此优化应对准损失

最终支撑恢复条件具有

$$\rho_*(J_{\lambda, H} - E_O^+ - d_0 \beta_{\text{old}}^+ - (1 - b_h) \beta_{\text{new}}^+) > I_H^+$$

的形式。论文给出的最终归一化正余量大于 $1/500000$ 。因此一项看似改善理想 eigenvalue 的变形，很容易被支撑修复、cap、尾部或区间误差吃掉。下一轮应直接优化认证后的完整余量，而非只画未扣损失的 Rayleigh quotient。

16.4 可核验的组合末端

当前使用的 40 元组为：

0, 2, 6, 12, 20, 26, 30, 32, 36, 42, 48, 50, 56, 60, 68, 72, 78, 86, 90, 92, 98, 102, 110, 116, 120, 126, 132, 138, 140, 146, 152, 156, 158, 162, 168, 170, 176, 180, 182, 186。

本次检查其 40 个元素互异、最大减最小为 186、对每个 $p \leq 40$ 都遗漏至少一个剩余类。对 $p > 40$ ，40 个元素不可能覆盖全部类，因此这构成完整 admissibility 验证。

要得到小于 186 的结果有两类末端：证明对更小 k 的 DHL $[k, 2]$ ，并找到对应更窄 admissible tuple；或在现有 k 与 arithmetic 结构下证明一个更窄的合适 tuple 可以使用。不能预先宣称 39 元组必定给某个未经查证的最小直径，也不能假设当前 40 元组已经在所有意义下最优。

17. 面向著名猜想的研究优先级

17.1 第一优先：固定宽度的 Montgomery 带宽外二点结果

终点。假设 RH，选定一个固定 $\varepsilon > 0$ 与非负、非零、光滑 w ，其支撑紧含于 $(1+\varepsilon, 2-\varepsilon)$ 。在明确的 T 平均与对角项约定下，证明归一化零点 form factor 的加权极限为 GUE 预测值。写成示意式是

$$\int w(\alpha) K_T(\alpha) d\alpha \rightarrow \int w(\alpha) d\alpha.$$

这里 K_T 必须在正式任务开始时逐项定义；本式是目标接口，不是假装已经有可引用定理。

为什么它有实质价值：ACUE 在该范围的替代三角 form factor 给 $2-\alpha$ ，而 GUE 给 1。对固定 w ，差为严格正数

$$\int w(\alpha)(\alpha-1) d\alpha > 0.$$

它避免了只在 $1+1/N$ 处出现但随 N 消失的微弱差异，也不依赖找全体 GUE 的所有相关阶。它是著名 pair-correlation 猜想中的一个清楚部分。ACUE 被该统计排除，不等于排除所有 AH 型过程；后者还需要对整个允许模型类求 extremal range。

真正缺的引理。将 K_T 的该部分转成可控制的 prime/ Λ shifted correlation、Dirichlet 多项式长矩或另一种等价算术量，并证明超出当前输入的 cancellation。有限反例不能供应这项 cancellation，只能给选择 w 的指导与误差预算。

停止条件。如果候选“新统计”在算术展开后只是重写了完整未解的 shifted-prime conjecture，且没有更弱、可独立攻击的中间估计，应记录等价障碍，不继续包装成新解法。

17.2 第二优先：186 以下的可认证筛法改进

终点。证明一个明确的偶数 $H < 186$ 是 H_1 的无条件上界；或先证明 DHL[39,2] 等更强筛法结论并连接真实 tuple。这个目标规模比 RH 小得多，但仍是公认的重要素数分布进展，且有相对完整的输入输出链。

优先尝试误差层的改进：支撑违规事件的更紧正覆盖、同坐标与异坐标 marked fragments 的分拆、Young 参数的局部自适应、源域切分位置、试验函数对被扣损失区域的主动避让。所有尝试必须维持每一项的非负上界与 interval certificate。

若 baseline 尚未能原样认证重现，就暂不宣布任何优化胜利。若只提高未认证浮点 quotient，而完整 margin 仍为负，也应视为未成功。

17.3 第三优先：排除精确定义的 AH，而非仅区分 ACUE

终点。在 RH 或另一套明确假设下，对某个与半整数集合保持正距离的区间 I 证明正比例零点间隙落入 I ，或证明相应 lattice-distance 统计有严格正下界。这样才直接针对密度版 AH。

现有 Nyquist、homometry 和局部 fiber 提醒我们，单独 ACUE 的 form factor 不是全部敌手。应先定义 AH-admissible class，再优化 detector 在该类上的上界，最后比较真实 ζ 的下界。这个三步链条中第三步依然是主要算术难题。

17.4 第四优先：LR 非唯一性或定量刚性

终点。构造两个确实不同的平稳半格点过程，在目标论文指定的全部 band-limited n -point 测试中相同；或者证明在附加自然条件下唯一。

此目标与现有 finite fiber 关系最直接，适合拿来验证新的方法是否真的能穿过 $N \rightarrow \infty$ 。但应诚实定位：它是重要的 random-process inverse problem，不自动是与 RH 等级相当的历史大定理。选择它，应因为它解除主路线中的信息障碍，而不是因为它容易被夸大。

17.5 暂不列为首轮终点的目标

直接证明 RH、全 GUE、孪生素数，或从任意 PSD/Wronski 表示推出 ζ 的全部零点性质，目前都缺少更大的算术桥梁。可以保留为远期终点，但不应把一个有限模型的未证局部化定理当作它们的主要剩余步骤。

18. 更开放的探索方向：不预先限制在旧模型内

18.1 用反例设计“必须增加的数据”，再寻找算术实现

把所有已知模仿者作为对抗样本库：ACUE、COM 调制、parity、homometric DPP、magnetic cycles、Nyquist marked equality candidates。对一个候选统计，输出它在每类上的值域、可见相关阶、带宽成本和归一化强度。

然后反向问： ζ 的哪一种已有表达自然包含这个统计？可能是 ζ' 的 logarithmic derivative、平滑 resolvent、有限 Euler product 的加权 trace、或带 prime-label 的 mixed correlation。只有找到真实 arithmetic definition 才投入大型计算。这个方法的价值在于减少对“看起来像 random matrix”的任意构造投入。

18.2 从局部下界转向正密度的动力学证据

最小间隙是极端统计，很容易被零密度异常控制。可以改研究多个短区间上局部 collision-depth 的分位数或平均的有界函数，如 $\min(N^2D, L)$ 。它们避开 clock 导致无穷均值的问题，也更接近排除密度 AH 所需的正比例结论。

难点是有限窗口边界会改变动力学。研究者需要定义带外部背景的局部模型、给出截断稳定性，并证明统计对窗口选择鲁棒。不能用裸截断多项式的碰撞时间冒充 ζ 的局部属性。

18.3 以 Hilbert 空间证书压缩高阶谱方案

Lamzouri 的简化使一个方向值得优先测试：把高阶矩、负谱 cap 或重数约束直接加入同一个 Hilbert 空间不等式，看是否存在真实更优的极 endpoint。先做纯优化的严格对偶，再检查每个新约束是否有独立的 ζ 输入。

如果优化对偶显示新增约束只是原有二次信息的重述，便立即止损。如果确有增益，再建立一个明确的 missing arithmetic lemma，而不直接启动大规模 Lean 翻译。

18.4 Wronski、Plücker 与总正性

此前任务讨论的可用来来源是 Karp – Purbhoo 关于 universal Plücker coordinates、Wronski map 和总正性结构的工作。[Karp – Purbhoo](#)

值得测试的不是“任何实根多项式都能找到某个正表示”，而是：是否存在由本问题自然给出的 Wronski lift，其演化与圆周热流相容，并使 Plücker 坐标控制两体背景系数或 discriminant 的首次消失。要逐项检查三个相容性：初值 lift 是否 canonical、flow 是否闭合、所需正性是否在 flow 下保持。任一步失败，就不能借总正性名称继续推论。

低 N 的可否证实验应包含 clock、单缺陷、三根近簇和两个远离近簇。若某个候选 lift 对同一谱随任意坐标选择而给不同结论，应视为它没有提供谱不变量。

18.5 函数域族作为严格试验场

可考虑在有已知 monodromy 与 equidistribution 的 L-function 族上测试新 detector：先在固定 degree 下让基域大小增长，再研究 degree 增长的误差。这里的目标是证明一个真正 arithmetic family 中的统计迁移定理。

这是研究建议，不是声称现有函数域定理已给出所有需要的 uniformity。CUE 固定维 equidistribution 与最小 gap/depth 的极端尺度并不自动兼容；误差必须小于目标事件的概率尺度。它能提供更诚实的中间战场，但不能通过“类比”直接解决。

18.6 186 的支撑几何与 spectral optimization

旧 P9 引擎擅长把高维可分离函数变成一维概率。可以把同样的概率分解能力用于 186 的真实 source-domain cap，而不是把新问题退回旧 simplex。研究重点是：哪些 support-loss integrals 能用条件概率、交换性 or 对称多项式精确降维，哪些必须保留 prime-fragment 的相关结构。

如果降维把同坐标的两个 fragments 当成独立，就可能破坏证书。应建立一个小尺寸完全枚举模型，先证明变换保留所需相关性，再应用高精度积分。这是旧工具与新突破之间比较具体的接点。

19. 给 Astra 的十二张任务卡

这些任务卡是当前任务内的研究建议；本次没有自动创建或启动十二个新任务。它们按依赖组织，可以在下一轮明确授权的研究中分工。

A1 — 先把所有 theorem statement 固定

输入：主报告、三个主论文与 Lean 草稿。输出：每条命题的对象、假设、量词、结论、引用和证据级别。优先修复 depth_ge、COM N=4、 ρ 与 κ 的定义、H_m 索引。验收：删去 proof body 后，statement 本身也准确表达预期数学。停止：无法明确对象的“桥梁”不进入证明队列。

A2 — 186 baseline 可认证重现

输入：公开 PrimeGaps186 仓库、版本固定的 FLINT 与补丁、论文参数。输出：完整区间 certificate、receipt、运行环境、三项项目输入状态。验收：原结果的 margin 与来源一致，且每个 enclosure 真正向外舍入。不要用已有 receipt 或浮点运行替代重新计算。

A3 — 39 或更窄 tuple 的完整 margin 搜索

输入：A2 的合法二次型与损失。输出：最佳“扣除全部误差以后”的严格下界及 tuple witness。允许先探索浮点，最终必须用同一个 checker。停止：所有试验改善只存在于未恢复支撑的理想模型中。

A4 — ACUE 带宽外固定强度 detector

输入：精确 COM/phase 家族与对照 CUE。输出：固定 smooth w 或固定阶局部量，给 uniform gap，而非随 N 消失的差。验收：在精确可行律上证明差与全部低频匹配。停止：只得到微小有限 N 差而无 uniform 下界。

A5 — detector 的 arithmetic realization

输入：A4 统计。输出：精确的显式公式/Dirichlet 展开、对角主项、off-diagonal 余项、所需 support 和平均范围。验收：每一个候选增益都对应一个可陈述的算术引理。停止：结论等价于完整目标猜想而没有更弱的可攻击输入。

A6 — LR 极限中的固定局部分离

输入：N=6...9 四点自由及显式 phase families。输出：两个 tight 序列，统一局部计数控制，固定 observable 差 $\geq \epsilon$ ，以及正确 bandwidth region 的矩传递。验收：极限过程确实不同，且满足原问题全部量词。不得用 parity clock 稀有原子作为唯一差异。

A7 — 三点 rigidity 的 all-N 分类

输入：小 N 精确矩约束与 Fourier – Slater 描述。输出：all-N 证明或明确最小反例。若证明，记录依赖的是 rotation invariance、balanced moments 还是更强结构。它服务于 A4 detector 的最低阶数选择，不应单独被包装成已经逼近 RH。

A8 — depth 上界的动态背景控制

输入：第 10.4 节修正差商。输出：在一个可验证事件上沿整段时间成立的 $B_{\max} \delta^2 = o(1)$ ；先在 CUE 或明确的随机根模型证明。验收：得到真正 $D \sim \delta^2/8$ 的概率命题与异常事件控制。停止：只在初始时刻计算平均背景。

A9 — 密度 AH 的正确动力学统计

输入：A8 与一个带背景的局部零点窗口定义。输出：有界 depth 统计或分位数，以及 density-AH 下的严格限制。验收：不由单个异常最小 gap 决定；窗口边界误差可控。算术计算另交 A5，不把条件串联当无条件结果。

A10 — marked 算子的规范性

输入：Cayley response、 $\xi'/\log$ derivative 或某个真正算术表示。输出：不依赖任意 eigenvector 选择的 mark，完整导数公式，rank-two 近似的定量误差。验收： κ' 不是被重算后隐藏的全谱项；能独立预测新样本。停止：mark 的定义已经含待测 depth。

A11 — 高阶谱证书对偶与算术成本

输入：Lamzouri Hilbert 形式、历史 cubic 与 Christoffel 数据。输出：严格优化 gain 或 exact no-gain dual，以及每项新输入的 cost ledger。验收：同当前最佳的相关基线比较，明确 RH/无条件及 simple/distinct 的差别。

A12 — 形式化最后闭环

输入：A1 中已经稳定、并在其他任务中得到证明的关键引理。输出：无 sorry 的正确 statement、依赖审计、独立数值 checker 与可复跑构建。优先正式证明 finite identities 和 tuple admissibility，再处理 ODE 首碰撞。不要把形式化资源耗在尚未确定正确的宏大桥梁上。

20. 第一次接手时的执行顺序

第一步恢复环境与证据：读取主报告的修正表，运行独立审计，核对来源 hash，检查引用与数组。第二步选主战场：如果要最快测试一个公认数论改进，优先 $A2 \rightarrow A3$ ；如果要沿用户最在意的 AH/GUE 路线，优先 $A4 \rightarrow A5$ ，同时用 A6 检查有限反例能否真正保留到无限过程。

第三步把每个“缺失引理”写成可以被反例击败的精确命题。对于局部化，要包含动态时间窗；对于算术均值，要写范围、权重、对角项和一致性；对于 186，要写完整 margin；对于 LR，要写全部 n 与测试函数区域。第四步才做大规模搜索或形式化。

每次失败应留下五个信息：原命题、最小反例或不等式失败位置、是否能通过增加自然假设修复、修复后的算术成本、是否仍服务主目标。这样保存失败才会产生下一次突破需要的约束，而不是新的长叙事。

21. 可直接交给 Astra 的接手指令

请从这份档案的审计部分开始，而不是直接相信附录中所有写成 theorem 的陈述。优先阅读固定提交中的 `final_verified_paper.md`、`round3_synthesis.md`、`impostors_paper.md`、`depth_scaling_theorem.md`、`signed_sieve_nogo.md`，并运行 `audit_research.py`。把所有命题分成有限证明、历史数值、条件桥梁和已撤回断言。

总目标是取得一个真正涉及著名数论问题的结果。零点主线选择一个固定强度、算术可观测的带宽外 pair-correlation 或 density-AH detector；筛法主线选择 186 完整二次型中的严格 margin 改进。不要把新的有限反例本身当作终点，除非它明确解除主链条上的一个障碍。

对每个候选路线先列完整依赖链，并找出唯一最缺的算术或概率引理。允许推翻旧方向。不要把 random-matrix 类比、任意 PSD 表示、浮点证书、SMT 的局部代数检查或未编译 Lean 当作解析定理。若不能越过关键引理，请给出准确的 no-go、反例或等价障碍，留下可复用的数学结论，而不是宣布已接近 RH。

22. 外部来源索引

以下只列本次实际使用的主要原始来源。网络状态截至 2026 年 9 月 4 日检索；预印本、项目陈述与完全独立重验不是同一证据级别。

来源	用途
固定研究分支	本项目的历史原稿、代码与数据
Tao: Alternative hypothesis for unitary matrices	ACUE 的原始构造与问题背景
Lagarias – Rodgers: higher correlations	AH 与全阶相关、问题量词与历史小间隙注释
Lagarias – Rodgers: band-limited mimicry	格点带宽识别边界与不同测试域
Feng – Wei	Circular β 最小间隙，适用 β 范围
Anthropic: ζ 项目	0.6725007 的官方进展说明
Alpöge – Furman	相关正式论文入口
Lamzouri	2026 年 9 月的 Hilbert 空间简化
Anthropic: FLT	大规模形式化方法
FLT 代码	形式化工程来源
Short gaps 论文	DHL[40,2] 与 186 的数学方案
Short gaps 研究记录	失败试验、符号和支撑问题
PrimeGaps186	Lean 项目输入、数值证书与依赖说明
Guth – Maynard	大值估计与零密度的相关入口
Tao 的零密度讨论	用户所提路线的一种可能对应，未唯一确认
Karp – Purbhoo	Wronski/Plücker 总正性的原始来源

23. 交付、检验与尚未完成的工作

已完成的是来源重建、研究演化说明、关键有限结果的整理、重要错误的明确修正、21 项独立审计、两个 tuple 的 admissibility 核验、外部新进展的核对，以及面向著名问题的具体任务设计。完整历史 Markdown 附录不删除过时推论，以便下一位研究者核查它们如何出现与被修正。

没有完成的是：全分支逐行数学审稿、所有旧实验复跑、 H_2 / H_3 / H_4 的严格 interval 认证、186 证书重建、Lean 补全、无限 LR 非唯一性、 ζ 算术桥梁，以及任何新的著名猜想证明。随后用户明确授权八小时自主研究、子代理交叉质疑及与 Claude Code 协作，并要求推送 PR #11；该夜间运行及其新结果记录在后续补充中。

程序与原始来源分开，审计与新报告推送到 PR #11 的独立目录，不静默修改历史论文。撤销相应新增提交即可移除整理层。后续最有价值的一步，是选择 $A_2 \rightarrow A_3$ 或 $A_4 \rightarrow A_5$ 中一条，完成其第一个缺失引理，而不是再生成一轮没有收敛终点的全面畅想。

主说明结束。后续为完整历史原文附录。附录中的断言与评价属于写作当时状态；发生冲突时，以主说明的审计与原数学证明为准。

第三轮更新：动态观测量、严格反例与 186 的结构性启发

日期：2026-09-05。本节优先于后面的旧提案。以下新增结果是有限维数学结论、明确范围内的证明和计算证书；没有证明实际 ζ 的 AH 为假、Montgomery – Dyson 猜想或更小的素数间隙。

已完成两个动态结论

对弧度坐标下的圆周排斥流，取

$$V_k = \sum_{j \neq k} \cot \frac{\theta_k - \theta_j}{2}, \quad L = \sum_k V_k \partial_{\theta_k}, \quad \rho_m = \sum_k e^{im\theta_k}.$$

首先，精确生成元公式是

$$L\rho_m = -m((N-m)\rho_m + \sum_{a=1}^{m-1} \rho_a \rho_{m-a}).$$

这里增加的是因子的个数，正、负 Fourier 总权重分别守恒。因此，对 $m \leq N$ ，ACUE 与 CUE 的 $L^r \mid \rho_m \mid^2$ 期望在所有整数 $r \geq 0$ 都相同。有限维不变子空间和全局存在的正向排斥流还给出所有 $t \geq 0$ 的精确匹配。这否定了“不断求导会把同一低频观测量推出保护范围”的具体路线。它不适用于反向首次碰撞后的真实角度流，也不宣称整个分布相同。

证明通过 Schur 交叉矩的离散 Cauchy – Binet 公式处理 $m = N$ 边界；234 项保护范围内的精确比较全部通过，另有整数环枚举和至 $N = 10$ 的完整浮点枚举。超出范围的 $N = 2, m = 3$ 在初始时刻已不同，不能作为动态突破。

其次，奇异力平方确实可以区分这两个模型。为避免与碰撞时间 D 混淆，本节记

$$E = \sum_k V_k^2 = \sum_{i \neq j} \csc^2 \frac{\theta_i - \theta_j}{2} - \frac{N(N^2 - 1)}{3}.$$

于是

$$E_{\text{CUE}} E = \frac{N(N^2 - 1)}{3}, \quad E_{\text{ACUE}} E = \frac{N(N^2 - 1)}{6}.$$

其耗散率满足

$$LE = - \sum_{i < j} \csc^2 \frac{\theta_i - \theta_j}{2} (V_i - V_j)^2,$$

$$E_{\text{CUE}} LE = -\infty, \quad E_{\text{ACUE}} LE = -\frac{2N(N^4 - 1)}{15}.$$

最后一个公式有完整的二点、三点 DPP 和余割求和推导；浮点枚举只作独立核对。CUE 的无限右斜率由非负差商和 Fatou 引理证明，不交换不可积的导数与期望。

这两个结论并不矛盾：力平方包含奇异的二点核，超出了前述有限多项式观测量。它的均值实际上退化为二点统计，不能凭外观认定为新高阶信息。把它用于 ζ ，仍需要可计算的算术加权零点和及移除截断的统一误差界。

修复碰撞深度桥梁的两个反例

实直线有限热流的正确恒等式追踪实际先碰撞的那一对，设其初始间隙为 d_* ：

$$D = \frac{d_*^2}{8} + \frac{1}{2} \int_{-D}^0 d(t)^2 S(t) dt.$$

若改用全体初始最小间隙 m_0 ，必须补上 $(d_*^2 - m_0^2)/8$ 。多项式

$$P(x) = x(x^2 - 1)((x - 100)^2 - \frac{121}{400})$$

给出严格反例：初始最小间隙为 1，先碰撞的远处一对间隙却是 11/10。精确 Sturm 证书定位首次判别式事件，排除更早碰撞和根群交换。遗漏项恰为 21/800。下界 $D \geq m_0^2/8$ 仍然成立。

另一个 2×2 例子取 $G_\varepsilon = \text{diag}(\varepsilon, 1 - \varepsilon)$ 与交换后的等谱矩阵，并固定标记 $u = e_1$ 。一侧的 $u^* G^{-1} u$ 发散，但 Cayley 热深度的标记导数分别趋于 1 与 $-1/2$ ，保持有界。两矩阵在每个固定 ε 下具有相同谱、迹和 Hilbert – Schmidt 范数；没有声称整个 ε 路径的范数不变。标记确实区分方向，却不能强迫所提议的发散。

圆周 ACUE 的确定性下界 $N^2 D \geq \pi^2/8$ 保留，包括深度无穷的 clock。常数 1.419640342 的单缺陷计算与随机 ACUE 极限中位数是不同命题；有限枚举破坏了旧单调证据，但本身不能排除最终极限恰好等于该数。

我们自己的来源勘误

Lagarias – Rodgers 1907.03391 § 5 确实把硬核参数 μ 与 0.606894 联系起来：构造给出 $\mu \geq 1/2$ ，作者提出修改已有方法应可得到 $\mu \leq 0.606894$ 。该处没有证明这个上界。此前只查看伴随论文而作出的“没有联系”判断过强，已更正。 $\mu=1/2$ 仍是目标；对过程、平稳性和异常集的量词必须按原文处理。[原文文](#)。

将 186 的方法用于下一轮，而不是只提高数值精度

用户强调的结构是：互补分解条件控制实际合成模数，推出三重稠密可分性，进而允许更大的 Maynard – Tao 筛权支撑。新的分解定理与数值优化是两个独立证明义务。对 ζ ，必须同样先写出实际出现的乘积、相关核和统一误差，不能直接借用算术级数中的分布指数。

已完成的二维 Ritz 检验只把固定有限向量的负 margin 提高约 0.000489，仍远低于零；它是已记录的有限算子结果，没有证明新系数族的算术极限。现有固定有理向量的严格连续证书仍约为 -0.0146624 。

下一轮主攻一个明确义务：找出可以严格估计的新混合乘积或带符号相关项，使完整正二次型出现足够大的净增益。不得丢弃中心 Gaussian 的 pole-cut 项，也不得把长系数的非对角 Gram 代价当成零。副线是精确定义真正 H_t 的局部深度及其背景控制；RH 与 Rodgers – Tao 给出全局 $\Lambda=0$ ，使负时刻的背景已含非实零点，有限全实筛屏恒等式不能直接搬用。

暂缓更大规模的同族扫参、ACUE 中位数猜测、未定义的 ζ 局部热流结论，以及没有 DHL 阈值证书的更短 admissible tuple 宣布。所有失败保留输入、输出和适用范围，供下一位研究者检验。

Fable 的实际回传及同步状态

旧 PR11 在固定提交 `a408e7050fffc74459b3c83fafa5ac03c8b7dea6` 已有 81 份 Fable 文件，现按字节镜像并附源哈希。公开评论确认较早的 FABLE_001 已被接收；F2 连续极限和有限和输出存在，但最终任务报告尚未出现。较新的单项计算包仍有单独的接收状态，不能混为一谈。

Fable 的固定连续值约 -0.014662375473370598 ，与 Astra 的精确证书相容于通常浮点误差，但其浮点值不能宣称落入远窄于浮点误差的有理数区间。有限和与我们独立的算子数据一致，仍未给出正 margin。其他已列入计划的热流、函数域和 LR 报告在该固定提交中大多尚不存在；脚本和计划不等于完成的证明。

独立同步审查分别记录 F2 证据和热流/函数域/LR 的具体问题。原始镜像不会被修改以掩盖错误；修正依据写在当前报告。此次同步只接收现有公开成果，没有新建 Claude 会话，也没有重复派发任务。

最新补篇：从历史材料到当前可验证的数学状态

更新：2026 年 9 月 5 日。本补篇与后面的新报告优先于历史附录中的状态标签。历史文件原样保留，不因后来纠错而删除。

当前任务与真正的成功标准

主任务是在 RH 下证明真实 ζ 零点归一化间距小于 $1/2$ ，或得到明确超出现有可用带宽的二点相关结论。成功需要一个适用于真实 ζ 的正证书，并写清适用的高度、密度、重数与误差；比较 CUE 和 ACUE 的模型行为本身不满足这个标准。

关于 AH 必须先固定版本。Lagarias – Rodgers 的较强格点表述把相邻归一化间距限制在正半整数附近；单个趋向于 $1/2$ 以下的极限间距结果可与这种强版本冲突。容许零间距或重数、容许零密度例外的版本，需要相应的正密度或远离格点的结论。不能把这些命题混为一个 AH，再用最小间距自动反驳全部版本。

八小时是最初的研究里程碑，随后协调任务转达了可恢复的 72 小时研究安排。每个检查点保存真实工作状态；这不是对连续运行时间或解决猜想的保证。Fable 使用用户现有单会话，只接收一个手动投递的固定向量核验包；没有证据表明它已接收或完成。

1. 这轮最明确的新数学进展：固定素因子试验的算术传递

原来的新方向是令

$$r_L(n) = d_\ell(n) \left[f\left(\frac{\log n}{\log L}\right) + g\left(\frac{\log n}{\log L}\right) \sum_{p|n} \left(\frac{\log p}{\log L}\right)^2 \right].$$

这里素数按不同因子计数，不能把素数幂重数悄悄加进去。早先计算直接使用了形式上的背景矩

$$E_\nu S_2 = \frac{\nu^2}{a+1}, \quad E_\nu S_2^2 = \frac{(a+6)\nu^4}{(a+1)(a+2)(a+3)}, \quad a = \ell^2.$$

这些式子现在已有固定族的算术证明，并得到另一研究代理的逐项审查。证明不以 Poisson – Dirichlet 随机模型为假设，而从 $\sum d_\ell(n)^2 n^{-s} = \zeta(s)^{\ell^2} F(s)$ 的正测度 Laplace 极限出发，插入标记素数，计算 beta 积分，再用权 Schur 估计消去小素数与高次素数幂。关键细节是：同一个素数在 A^*A 中的对角项保留为主项，在 A^2 中的重复插入带有 p^{-2} ，则可消去。短背景区域通过测度收敛和无原子边界处理，无需假设逐点一致的 Selberg – Delange 误差。

证明范围是固定 $\ell \geq 1$ 、固定多项式 f, g 、正极限范数。不包含随 L 增长的维数、任意无限特征族或完整 Fock 算子极限。算术极限本身无条件；把它用于 Inoue 的 ζ 不等式才需要 RH。完整证明、审查意见及审查后补入的三个显式估计均列在后文。

这个进展解决了一个实际的算术依据问题，但固定有理试验仍给出

$$-0.01467 < J < -0.01465, \quad J \approx -0.014662375473369.$$

因此没有半间距突破。此前约 -0.01465473 的较丰富特征搜索也仍为负，且不是这一有理向量。不能用已证明的固定族传递，自动替较高阶特征或病态数值最优向量背书。

2. 完整算术算子告诉了我们什么

在半间距处，饱和主项是

$$K_L = A_L^* A_L + \frac{1}{2} (A_L^2 + (A_L^*)^2), \quad \text{目标阈值} = \frac{\pi^2}{2}.$$

算子已算到 $L = 10^7$ ，最大特征值约 4.324089559，归一化余量约 -0.030939064。这只是有限算子的高精度数值，不是严格上界，也不是渐近无解定理。在 $L = 10^6$ ，20 项对称素因子特征已接近完整有限最优值；增加到 40 项后，差距约为 0.0004437。这说明当前测试范围的主要结构确实具有乘性，但不排除随 L 改变的新方向。

一个单独估计两项能量的 Cauchy – Schur – Volterra 上界只降到约 5.2074，高于 $\pi^2/2$ 。它失败的原因是损失了两个项的近极值向量之间的兼容性；继续提高同一个网格的精度不能弥补这个量级的缺口。这个失败已保存为完整推导和各网格结果。

固定素数分箱还产生一个有限 bosonic creation 模型，40 个分箱有 215,308 个状态，数值仍未过阈值。固定分箱可给构造性的渐近下界，但不是完整算子收敛，更不是全局上界。小素数云的处理必须明确，不能把真空有限粒子截断当作全部乘性结构。

3. 值得永久保存的反向结果：中心高斯并不免费扩大支撑

把高度 T 附近的高斯换成以原点为中心的高斯，会使 Fourier 核非负，似乎可以保留更长混合支撑。然而精确临界线积分恒等式中， ζ 的极点贡献不会消失。对所需长尾匹配频率，它的相位与素数主项相反。

设 $M = T^{1+\eta}$ 、 $0 < \eta < 1$ 、 $h = \pi/\log T$ ，取

$$R(t) = 1, \quad C(t) = -i \sum_{M < q \leq 2M} q^{-1/2-it}.$$

两边对角系数范数有界，但极点混合修正为

$$-(4\pi \cos \frac{\pi\eta}{2} + o(1)) \frac{M}{\log M}.$$

对于测试的宽度 $W = T$ 或 $T/\log T$ ，它除以 W 后发散。因此不能把它当作由对角范数控制的 $o(W)$ 误差。与此同时，长支撑真实 Gram 矩阵的非对角罚项也会变大，低高度相位聚集不能忽略。

这不是“所有中心高斯方法不可能”的定理。它精确排除了无偿支撑扩张：下一条合法路线需要控制减去连续素数密度后的混合项，或定量支付极点和低高度代价。对完成函数、去极点函数作替换，也必须保留相应的 gamma 或连续项。

4. Heat flow 路线的完成范围及与真 ζ 的距离

对圆周根多项式的系数流，去掉无关标量后是普通标量热流。隔离最小间距为 δ 的一对根，扣除远背景的共同线性漂移，可在

$$\delta^2 B \rightarrow 0, \quad B = \frac{1}{4} \sum_{k \in \mathbb{Z}} \csc^2(\theta_k/2)$$

下证明第一次碰撞时间 $D \sim \delta^2/8$ 。作者第二轮审查把常数显式化为 $K = 16384$ 、 $\eta_0 = 1/524288$ ，并控制全程移动边界、Gaussian 尾项和根的延续。这些粗常数用来检查一致性，不是新的数值纪录。

结合已有极小间距分布和三根聚集排除，可得到 CUE 及正整数 β 圆周族的碰撞时间极限律。需要继续查明与经典 Lehmer pair 判据、已有多项式热流工作的关系，不能把相似原理包装成全新理论。

真 H_t 的讨论也更精确了：在 RH 下，中心化 genus-one canonical product 可以控制所需远背景尾项。关键仍是算术上产生充分孤立、足够小且在所需版本中有足够频率的零点对。Rodgers – Tao 的宏观负时间带来很多微观松弛时间；由半格硬核得到的 $c/(\log T)^2$ 深度只对有限微观时间，两者不能直接替换。

5. 素数间隙、0.6725 和形式化工作的实际启发

素数 186 路线中，39 点、直径 182 的 admissible tuple 已有有限核验，但所需的 DHL [39, 2] 没有证明，所以没有得到更小的 prime-gap 定理。精确凸性与有理单纯形分割提供了可复核的 minorant 质量区间；这类“试验向量 + 定向误差预算 + 严格证书”的工作方式可用于 ζ 的固定变分问题。它不能代替缺少的算术混合矩估计。

0.6725007 是特定已公开 ζ 结果中的简单且在线零点比例；不能称为所有 bandwidth-one 论证的普遍天花板。新的证明简化与数值纪录改变也必须分别记录。本项目的标记格点 0.702642 常数属于另一模型，不是 ζ 纪录提升。

FLT 形式化的主要可迁移内容是大型依赖图、明确接口、局部证明任务和机器核验，不是自动产生新的解析数论不等式。当前先形式化精确代数、有限证书及已经闭合的解析引理；不要给尚未证明的主步骤加一个公理后宣布名猜想已完成。

6. 下一轮有明确验收条件的探索

1. 密度扣除后的长混合项。先写精确极点、gamma、Gram 罚项，再问是否存在可控的新支撑区域。验收是含全部误差的有符号不等式，而不是忽略修正后的正数。
2. 兼容极值结构。在完整 $A^*A + \text{Re } A^2$ 中寻找比逐项 Cauchy 更强的共同约束。验收是严格上界或一个新的可认证试验族；有限谱外推不算。
3. 固定族的可认证搜索。已闭合的 $H = f + gS_2$ 提供一个真实算术接口。任何新增特征先证明其有限标记转移，再做适度优化。验收包括正定范数、可控制条件数、定向误差和净余量。
4. 真正 ζ 的短对算术信息。Heat-flow 引理的用途是把明确的局部输入转化为动力学后果。验收应是一个目前未知、能从 ζ 算术推出的隔离或频率估计，而不是再次证明随机矩阵里存在小间距。

这些方向都还没有产生重大猜想的证明。档案保留精确成功和精确失败，是为了下一位研究者能从真正尚缺的步骤继续，而不是重做已经排除的尝试。

Current report 01: symmetric_prime_arithmetic_transfer

Source: `research/reports/symmetric_prime_arithmetic_transfer.md`. SHA-256:
d1329ec22b249fb0f1a8e6cf645e5d831ecb1ee7f8759f8c2b2caaaacc21d904.

Arithmetic transfer for a fixed symmetric prime–factor resonator

Date: 2026-09-05. Authoring role: Astra, primary research task.

Status: complete written proof with independent internal review. This note supplies the route from the weighted integer sums to the previously stipulated continuum form. The [independent review](#) accepted the fixed-family argument. This is not Lean verification or external peer review. The three explicit estimates requested by that review are added below. The fixed rational trial still has a negative margin. No half-gap theorem is claimed.

1. Precise statement

Fix a real number $\ell \geq 1$, write $a = \ell^2$, and let $d = d_\ell$ be the multiplicative function with

$$d(p^e) = \frac{(\ell)_e}{e!}.$$

For $L > 1$, put

$$v_n = \frac{\log n}{\log L}, \quad S_L(n) = \sum_{p \mid n} \left(\frac{\log p}{\log L} \right)^2, \quad r_L(n) = d(n)H(v_n, S_L(n)), \quad x_n = \frac{r_L(n)}{\sqrt{n}},$$

where $H(v, S) = f(v) + g(v)S$, and f, g are fixed real polynomials. Prime divisors in S_L are distinct. The proof extends to continuous f, g , but the polynomial case suffices here. Define the nonnegative matrix

$$(A_L)_{p^e m, m} = \frac{2 \sin(\pi \ell \log p / (2 \log L))}{e \sqrt{p^e}} \quad (p^e m \leq L).$$

Let E_v mean the linear moment functional

$$E_v 1 = 1, \quad E_v S = \frac{v^2}{a+1}, \quad E_v S^2 = \frac{(a+6)v^4}{(a+1)(a+2)(a+3)}.$$

No probability model is assumed. The functional is derived below from integers. Write

$$H_0 = H(v, S), \quad H_u = H(v+u, S+u^2), \quad H_w = H(v+w, S+w^2), \quad H_{uw} = H(v+u+w, S+u^2+w^2).$$

The arithmetic limit is

$$\frac{\|A_L x\|^2 + x^T A_L^2 x}{2\pi^2 \|x\|^2} - \frac{1}{4} \rightarrow \frac{M_2 + M_3}{I} - \frac{1}{4},$$

provided $I > 0$, where

$$I = \int_0^1 v^{a-1} E_v H_0^2 dv,$$

$$M_2 = \frac{2\ell^2}{\pi^2} \int_{v+u+w \leq 1} v^{a-1} \frac{\sin(\pi u/2)}{u} \frac{\sin(\pi w/2)}{w} E_v (H_0 H_{uw} + H_u H_w) dv du dw,$$

$$M_3 = \frac{2}{\pi^2} \int_{v+u \leq 1} v^{a-1} \frac{\sin^2(\pi u/2)}{u} E_v H_0^2 dv du.$$

All variables of simplex integrals are nonnegative. The scalar normalization constant below cancels from this quotient.

2. The unmarked integer measure

Define

$$C_\ell = \prod_p (1 - p^{-1})^a \sum_{e \geq 0} \frac{d(p^e)^2}{p^e} > 0.$$

Each factor is $1 + O_\ell(p^{-2})$, so this product converges absolutely to a positive number. For real $s > 1$,

$$D(s) = \sum_n \frac{d(n)^2}{n^s} = \zeta(s)^a F(s), \quad F(s) \rightarrow C_\ell \quad (s \downarrow 1).$$

Consequently the locally finite positive measures

$$\nu_L = \frac{1}{(\log L)^a} \sum_{n \geq 1} \frac{d(n)^2}{n} \delta_{\log n / \log L}$$

converge on bounded intervals to

$$\nu(d\nu) = \frac{C_\ell}{\Gamma(a)} v^{a-1} dv.$$

One can use the Laplace continuity theorem for positive measures here, rather than needing a uniform pointwise Selberg – Delange error: for every $t > 0$,

$$\int e^{-tv} d\nu_L(v) = \frac{D(1 + t/\log L)}{(\log L)^a} \rightarrow C_\ell t^{-a}.$$

For completeness, tilt by e^{-v} . The tilted measures have convergent total masses and Laplace transforms; their tails obey $\int_R^\infty e^{-v} d\nu_L \leq e^{-R/2} \int e^{-v/2} d\nu_L$. This gives tightness. Subsequential weak limits are uniquely identified by their Laplace transforms, and untilting on a compact interval gives the displayed local convergence. The limiting measure has no atom at zero or at any cutoff. This also gives

$$\sum_{n \leq L} \frac{d(n)^2}{n} = O_\ell((\log L)^a).$$

These arguments only need the classical simple pole of $\zeta(s)$ at 1 and nonnegativity of the coefficients. They do not assume RH.

3. Marked primes and the two moments

First restrict marked primes to $p \geq L^\varepsilon$, for a fixed $\varepsilon > 0$. For $j > 0$, the prime number theorem and partial summation give

$$\sum_p \frac{(\log p / \log L)^j}{p} \delta_{\log p / \log L} \Longrightarrow u^{j-1} du$$

on compact subintervals of $(0, \infty)$. Uniformity at zero will instead come from truncation.

If distinct marked primes are coprime to the background integer m , each prime inserted into $d(n)^2$ contributes $d(p)^2 = a$. Thus the limiting first-mark measure is the convolution

$$\frac{C_\ell}{\Gamma(a)} a \int_0^v u(v-u)^{a-1} du = \frac{C_\ell}{\Gamma(a)} v^{a-1} \frac{v^2}{a+1}.$$

Expanding S^2 produces the same-prime fourth-power term and the ordered distinct-prime term. Their respective densities are

$$\begin{aligned} \frac{C_\ell}{\Gamma(a)} a \int_0^v u^3 (v-u)^{a-1} du &= \frac{C_\ell}{\Gamma(a)} v^{a-1} \frac{6v^4}{(a+1)(a+2)(a+3)}, \\ \frac{C_\ell}{\Gamma(a)} a^2 \int_{u+w \leq v} u w (v-u-w)^{a-1} du dw &= \frac{C_\ell}{\Gamma(a)} v^{a-1} \frac{av^4}{(a+1)(a+2)(a+3)}. \end{aligned}$$

These beta integrals give the claimed moment functional. Multiplying by a bounded continuous function of the total mass is permitted by weak convergence.

Here are the two discarded cases, with the order of limits specified.

- **Repeated or background primes:** for $\ell \geq 1$, $d(p^{b+c}) \leq d(p^b)d(p^c)$, since the successive ratios $(\ell + e)/(e + 1)$ decrease with e . Hence requiring a background integer to contain a marked prime costs an additional factor bounded by $O_\ell(p^{-1})$ after its first reciprocal-prime weight. At $p \geq L^\varepsilon$, the sum of the resulting p^{-2} terms is $O(L^{-\varepsilon})$. The unmarked mass is $O((\log L)^a)$, and the other finitely many reciprocal-prime sums cost at most powers of $\log \log L$. Thus these contributions vanish after normalization. The same argument removes coincident distinct marks and prime exponents greater than one in this expansion.
- **Small marked primes:** deterministically, for every $n \leq L$,

$$0 \leq \sum_{p \mid n, p < L^\varepsilon} \left(\frac{\log p}{\log L} \right)^2 \leq \varepsilon \sum_{p \mid n} \frac{\log p}{\log L} \leq \varepsilon.$$

Since $0 \leq S_L(n) \leq 1$, deleting these marks changes H and its products by $O_H(\varepsilon)$, uniformly on the cutoff.

Take $L \rightarrow \infty$ with ε fixed, and then $\varepsilon \downarrow 0$. All product-measure cutoffs have boundary of measure zero. In particular, this approach includes the short-background range m near 1; it does not assume a uniform asymptotic for $\sum_{m \leq L/(pq)}$ at every individual p, q .

4. Uniform removal of small-prime and prime-power operator pieces

The preceding moment truncation alone is not enough: the creation operator also has a small-prime part. We control it with a weighted Schur estimate.

Put $w_n = d(n)/\sqrt{n} > 0$. The logarithmic derivative of $\zeta(s)^\ell$, or its formal Dirichlet-series coefficient identity, gives

$$\sum_{p^e \mid n} (\log p) d(n/p^e) = \frac{\log n}{\ell} d(n).$$

Since $0 \leq 2 \sin(\pi \log q / (2 \log L)) \leq \pi \log q / \log L$,

$$\frac{(A_L w)_n}{w_n} \leq \frac{\pi \log n}{\ell \log L} \leq \frac{\pi}{\ell}.$$

Submultiplicativity gives for any subset E of prime powers

$$\frac{(A_E^\top w)_m}{w_m} \leq \sum_{p^e \in E, p^e \leq L} \frac{2 \sin(\pi e \log p / (2 \log L))}{e p^e} d(p^e).$$

For all prime powers the right side is $O_\ell(1)$. For the prime part $p < L^\varepsilon$ it is at most

$$\frac{\pi \ell}{\log L} \sum_{p < L^\varepsilon} \frac{\log p}{p} \leq \pi \ell \varepsilon + O_\ell(1 / \log L).$$

For exponents $e \geq 2$ it is bounded by

$$\frac{\pi}{\log L} \sum_p \sum_{e \geq 2} \frac{d(p^e) \log p}{p^e} = O_\ell(1 / \log L).$$

The series converges because $d(p^e) = O_\ell((e+1)^{\ell-1})$. The weighted Schur test follows directly by Cauchy – Schwarz and summing the two displayed weighted row/column bounds. Therefore

$$\|A_L\| = O_\ell(1), \quad \|A_{p < L^\varepsilon}\| = O_\ell(\sqrt{\varepsilon} + 1 / \sqrt{\log L}), \quad \|A_{e \geq 2}\| = O_\ell(1 / \sqrt{\log L}).$$

Thus replacing A_L by its prime-only part with $p \geq L^\varepsilon$ changes either quadratic form, divided by $\|x\|^2$, by $O_\ell(\sqrt{\varepsilon} + 1 / \sqrt{\log L})$. This estimate works for signed H , and indeed for every vector x ; no coefficient positivity is needed at this step.

5. Evaluation of the retained integer sums

For the retained prime operator set $\alpha_p = 2 \sin(\pi \log p / (2 \log L))$. Exactly,

$$\begin{aligned} \|A_L x\|^2 &= \sum_{n \leq L} \frac{1}{n} \sum_{p, q \mid n} \alpha_p \alpha_q r_L(n/p) r_L(n/q), \\ x^\top A_L^2 x &= \sum_{mpq \leq L} \frac{\alpha_p \alpha_q}{mpq} r_L(m) r_L(mpq). \end{aligned}$$

In the first formula, the terms $p = q$ are **not discarded**. With $n = mp$, they give

$$\sum_{mp \leq L} \frac{4 \sin^2(\pi u/2)}{mp} d(m)^2 H(v, S_L(m))^2.$$

Its limit, after division by $C_\ell(\log L)^a / \Gamma(a)$, is $2\pi^2 M_3$.

For $p \nmid q$, put $n = mpq$. Away from the negligible coprimality defects already estimated, the coefficient is

$$r_L(mp) r_L(mq) = \ell^2 d(m)^2 H_u H_w.$$

The prime-insertion measures give the corresponding $H_u H_w$ term of $2\pi^2 M_2$. There is no extra factor of two: (p, q) is ordered both in the integer sum and in the double integral.

Similarly, in the second quadratic form the distinct-prime terms have

$$r_L(m) r_L(mpq) = \ell^2 d(m)^2 H_0 H_{uw}.$$

The repeated-prime $p = q$ terms here contain p^{-2} , and vanish by the fixed- ε estimate. This is why the same-prime term survives in $A^\top A$ but not in A^2 . The marked-background moments from Section 3 evaluate both expectations.

Finally,

$$\|x\|^2 \sim \frac{C_\ell}{\Gamma(a)} (\log L)^a.$$

Combining the three limits and the operator truncation proves the stated arithmetic limit. The independent internal audit checked these measure-convergence and discarded-term arguments.

6. Interface with the actual zeta theorem

The exact source is Inoue, [arXiv:2604.05733v1](#), Theorems 3 and 4. Those theorems allow arbitrary arithmetic resonator coefficients under RH and the product cutoff $L \leq T/(\log T)^2$. Choose $L = \lfloor T/(\log T)^2 \rfloor$, $h = \pi/\log T$, and the logarithmic increment approximator. The theorem has a coefficient-independent normalized error tending to zero; the approximator square sum is $O(1)$.

The corresponding operator uses $\theta = \log L/\log T \rightarrow 1$. Replacing θ by 1 changes the weighted Schur bounds, hence the operator norm, by $O_\ell(|1 - \theta|)$; apply the mean value inequality for sine and the same logarithmic derivative and prime-sum estimates. At $\phi = 1/2$, the linear coefficient cancels. The resulting lower bound for the normalized weighted zero-count expression has the proposed continuum margin, with the scope accepted in the independent review.

This would establish the arithmetic validity of the fixed symmetric-prime family. It would **not** establish a positive margin. The current certified value is approximately -0.014662375473369 , and a fixed negative lower bound has no half-gap consequence.

7. Audit requests and provenance

Independent reviewers should check: the $\ell \geq 1$ submultiplicativity restriction; the exact weighted Schur row and column orientation; local convergence of positive measures including the $\nu = 0$ boundary; counting of ordered distinct primes; uniform removal of background coincidences; and the error normalization when invoking Inoue Theorem 4. A successful numerical check cannot replace any of these steps.

This draft is a mathematical reconstruction intended for public verification. It records the argument and its dependencies, not a transcript of private model deliberation. No novelty claim is made for the marked-prime asymptotics or the Schur method; the purpose is to close the exact transfer obligation created by this project's fixed trial.

8. Explicit estimates added after independent review

The reviewed draft had SHA-256 `bf2ca70e61da1f97b4c214304180f0f45d3c26a050104bf30e6d70e46660d07f`. The review requested the following elaborations; they clarify the estimates rather than change the theorem.

First, the uniform total-mass bound follows directly from the Laplace transform:

$$\frac{1}{(\log L)^a} \sum_{n \leq L} \frac{d(n)^2}{n} \leq e^{\frac{D(1 + 1/\log L)}{(\log L)^a}} = O_\ell(1).$$

Second, the background coincidence bound can be written explicitly as

$$\sum_{m \leq L: p \mid m} \frac{d(m)^2}{m} \leq \frac{\ell^2}{p} \sum_{k \leq L} \frac{d(k)^2}{k}.$$

The additional reciprocal-prime weight outside this sum gives a summable p^{-2} cost. Collisions between an outer inserted prime and a prime in the expansion of the background S moment have exactly this form. There are only finitely many such marks for fixed H.

Third, if $A=A_0+E$, then

$$\|A^*A - A_0^*A_0\| \leq (\|A\| + \|A_0\|)\|E\|, \quad \|A^2 - A_0^2\| \leq (\|A\| + \|A_0\|)\|E\|.$$

These inequalities justify removing the small-prime and prime-power operator pieces uniformly over signed vectors. If applying Inoue Theorem 4 directly, its normalized $O(h) M_1$ term is bounded by Theorem 3 and the same Schur estimate; the source's combined Proposition 3 packages this cancellation and error estimate explicitly. No new off-diagonal hypothesis is introduced.

Current report 02: symmetric_prime_transfer_independent_review

Source: `research/reports/symmetric_prime_transfer_independent_review.md`. SHA-256:
`f3a961b4fe4340b6e65f31c2aa40adb5fcd7d54ab09406ab2dc7c1ab38118ddf`.

Independent review of the fixed symmetric–prime arithmetic transfer

Review date: 5 September 2026. Reviewer: independently assigned residual-Gram research agent.

Reviewed source, without editing it:

`/Users/qingyunsun/Library/CloudStorage/Dropbox/Code/Riemann zeta RMT/Astra-
Research/research/reports/symmetric_prime_arithmetic_transfer.md`

Reviewed source SHA-256: `bf2ca70e61da1f97b4c214304180f0f45d3c26a050104bf30e6d70e46660d07f`.

Decision

Accept the fixed-family arithmetic-transfer argument as mathematically complete at the level of an ordinary written proof. I found no fatal gap after checking the positive-measure limit, marked-prime coincidences, arbitrary-vector Schur estimates, same-prime terms, and the normalization of the interface with Inoue's theorem.

This decision has a precise scope. It covers fixed `ell>=1` and fixed

$$H(v, S) = f(v) + g(v)S, \quad S = \sum_{p \mid n} (\log p / \log L)^2,$$

with the polynomials stated in the draft and positive limiting norm `I`. It therefore covers the particular rational trial `ell=16/15` certified in round one. Its operator asymptotic is unconditional; RH enters only when the result is inserted into the cited zeta theorem.

It does not certify the full arithmetic-operator/Fock limit, an arbitrary infinite feature family, uniformity when the polynomial degree or coefficients grow with L , or a positive half-gap margin. The fixed certified margin remains approximately `-0.014662375473369`, and thus gives no new small-gap theorem. This is an independent mathematical review, not Lean verification or external peer review.

1. Positive–measure limit: valid, including the short background range

The Euler factors satisfy

$$(1 - p^{-s})^a \sum_{e \geq 0} d_\ell(p^e)^2 p^{-es} = 1 + O_\ell(p^{-2s})$$

uniformly as real `s` decreases to 1. Thus the real Euler product tends to the claimed positive constant. A complex analytic branch of `zeta(s)^a` is not needed for the Laplace argument.

The scaled Laplace transform tends to $C_{\ell} t^{-a}$ for every positive t . Tilting by e^{-v} gives finite measures with convergent total masses. The stated exponential tail inequality is enough for tightness, because the Laplace transform at $t=1/2$ is uniformly bounded for sufficiently large L . Uniqueness of Laplace transforms identifies every subsequential limit. Undoing the tilt on compact intervals gives local weak convergence.

The limiting density $C_{\ell} v^{a-1}/\Gamma(a)$ has no atom at zero or at the cutoff. Therefore restriction to $[0, 1]$ and to the simplex cutoffs is legitimate. This avoids the commonly invalid step of requiring a pointwise Selberg – Delange estimate uniformly for every short background length $L/(pq)$.

One explicit mass bound, useful to make the proof fully transparent, is

$$\frac{1}{(\log L)^a} \sum_{n \leq L} \frac{d_{\ell}(n)^2}{n} \leq e \frac{D(1 + 1/\log L)}{(\log L)^a} = O_{\ell}(1).$$

It holds because $e^{-\log n/\log L} \geq e^{-1}$ on the cutoff. This supplies every crude background-mass estimate needed later.

2. Marked–prime moments: constants and coincidence removal are correct

For fixed epsilon, the marked primes lie in a compact logarithmic range bounded away from zero. The prime-number theorem gives weak convergence of the reciprocal-prime measures there. Products with the unmarked positive measure converge on compact sets. The boundary hyperplanes of the product cutoffs have zero limiting measure.

The arithmetic identity behind the first mark is exact away from the excluded coincidences: inserting a new prime into $d_{\ell}(n)^2$ multiplies the background coefficient by $\ell^2=a$.

I independently recomputed the three beta integrals. Dividing the marked density by the unmarked density gives

$$E_v S = \frac{v^2}{a+1},$$

$$E_v S^2 = \frac{6v^4}{(a+1)(a+2)(a+3)} + \frac{av^4}{(a+1)(a+2)(a+3)}.$$

The first summand is the same-prime fourth-power contribution. The second comes from **ordered distinct primes**. Their sum is the $(a+6)$ expression in the draft. There is no missing factor of two.

The submultiplicativity condition is also correct. For $\ell \geq 1$,

$$\frac{d_{\ell}(p^{e+1})}{d_{\ell}(p^e)} = \frac{\ell + e}{e + 1}$$

is nonincreasing in e , hence $d_{\ell}(p^{b+c}) \leq d_{\ell}(p^b) d_{\ell}(p^c)$. The restriction on ℓ is genuinely used here.

For example, the background coincidence can be bounded explicitly by

$$\sum_{m \leq L: p \mid m} \frac{d_{\ell}(m)^2}{m} \leq \frac{\ell^2}{p} \sum_{k \leq L} \frac{d_{\ell}(k)^2}{k}.$$

After the outside mark's reciprocal-prime factor, this is a p^{-2} contribution. For $p \geq L^{\epsilon}$, summing it gives at most a constant times $L^{-\epsilon}$ times the total background mass. Higher exponents are handled by factoring off p^2 and using the same submultiplicativity bound. Finitely many other marked primes contribute bounded factors for fixed epsilon, or at worst harmless powers of $\log \log L$ under the draft's coarser estimate.

The deterministic small-mark bound is particularly effective:

$$\sum_{p \mid n, p < L^e} (\log p / \log L)^2 \leq e.$$

Because H is fixed and linear in S , its products change uniformly by $0_H(\epsilon)$ on the cutoff. One therefore takes $L \rightarrow \infty$ first and $\epsilon \rightarrow 0$ afterward. No exchange of limits or assumption of uniformity in shrinking ϵ is needed.

The same reasoning applies when the background moments are evaluated inside the two outer prime insertions. Expanding $H_{0H_{\{uw\}}}$ or H_{uH_w} adds only finitely many marks. Collisions between an outer inserted prime and a background mark have the same additional reciprocal-prime cost. This is worth stating explicitly in the final prose, but it does not require a new estimate.

3. The weighted Schur estimate is correctly oriented and works for signed vectors

Take $w_n = d_{\ell}(n) / \sqrt{n}$. The draft's weighted row bound uses

$$\sum_{p^e \mid n} (\log p) d_{\ell}(n/p^e) = \frac{\log n}{\ell} d_{\ell}(n).$$

This is the coefficient identity obtained by differentiating the Dirichlet series for $\zeta(s)^{\ell}$. The sine bound cancels the exponent e exactly, leaving the displayed logarithmic derivative. Hence the weighted row sum is at most π/ℓ .

The column calculation contains a second factor $1/\sqrt{p^e}$ from the weight ratio, so its denominator is $e p^e$, not $e \sqrt{p^e}$. The draft uses the correct denominator. Submultiplicativity gives the stated upper bound in terms of $d_{\ell}(p^e)$.

For the prime-only small part, the column bound is $0_{\ell}(\epsilon + 1/\log L)$. For all higher prime powers it is $0_{\ell}(1/\log L)$, because

$$\sum_p \sum_{e \geq 2} \frac{d_{\ell}(p^e) \log p}{p^e} < \infty.$$

The fixed- ℓ polynomial growth in e is enough for this convergence. Combining these column bounds with the full row bound by the weighted Schur test yields exactly the square-root orders in the draft.

These are operator-norm bounds. No positivity of x_n , H , or the polynomial coefficients is needed. To make the quadratic-form passage explicit, if $A = A_0 + E$, then

$$\|A^* A - A_0^* A_0\| \leq (\|A\| + \|A_0\|) \|E\|,$$

and

$$\|A^2 - A_0^2\| \leq (\|A\| + \|A_0\|) \|E\|.$$

After division by $\|x\|^2$, the error is independent of the direction of x . This closes the small-prime and prime-power issue for the signed rational trial.

4. Same-prime terms and normalization of the limiting quadratic forms

The two quadratic forms treat repeated primes differently, and the draft handles the distinction correctly.

For A^*A , the same-prime term has $n=mp$ and weight

$$\frac{4 \sin^2(\pi u/2)}{mp} d_t(m)^2 H(v, S_L(m))^2.$$

It is of main-term size and produces the entire M3 contribution. It must not be removed as a coincidence error. The draft retains it.

For A^2 , applying the same prime twice instead has the reciprocal weight p^{-2} . At fixed epsilon it is negligible by the coincidence estimate. Distinct primes multiply the background coefficient by e_{11}^2 , producing $H_{0H_{uw}}$. The off-diagonal part of A^*A likewise produces $e_{11}^2 H_{uH_w}$.

Both integer sums order the pair of inserted primes. The continuum double integrals use the same ordering. Multiplication by $\alpha_p \alpha_q$ gives the coefficient $4 e_{11}^2$, which equals the coefficient of $2 \pi^2 M_2$. The same-prime coefficient 4 equals that of $2 \pi^2 M_3$.

The norm satisfies

$$\|x\|^2 \sim \frac{C_t}{\Gamma(a)} (\log L)^a.$$

The common scalar cancels. Since I is assumed positive, division by the norm is legitimate. This reproduces the precise normalization of the rational continuum certificate.

5. Interface with Inoue: RH and the cutoff error are accounted for

The cited Inoue paper supplies the zeta inequality under RH for arbitrary arithmetic resonator coefficients. This review treats that externally stated theorem as an input; it is not an independent reproof of the entire paper.

Set $L = \lfloor T / (\log T)^2 \rfloor$ and $h = \pi / \log T$. The approximator coefficients have bounded weighted square sum. One direct check is to use the small-argument sine bound on primes and partial summation; prime powers contribute a convergent sum. Consequently the displayed normalized error tends to zero. If using Theorem 4 directly, its $O(h) M_1$ term is harmless because the preceding Schur bounds and Theorem 3 bound the normalized first moment. Alternatively the source's combined Proposition 3 already packages the needed coefficient-independent normalized error.

The operator in the zeta expression has $\theta = \log L / \log T$, not exactly 1. Applying the mean-value inequality to sine gives an entrywise majorant proportional to

$$|1 - \theta| \frac{\log p}{\sqrt{p^e \log L}}.$$

Its weighted row and column bounds are each $O_{\text{ell}}(|1 - \theta|)$, so the Schur norm is $O_{\text{ell}}(|1 - \theta|)$, as stated. The cutoff discrepancy therefore vanishes. At the half-gap boundary the leading linear coefficient cancels, leaving the two quadratic forms already analyzed.

6. Suggested presentation changes, none of which is a fatal repair

For publication or formalization, I recommend adding the explicit background-mass inequality in § 1 of this review, the displayed bound for a background divisible by a marked prime in § 2, and the quadratic-form perturbation inequalities in § 3. These turn the most compressed prose estimates into short checkable lemmas.

I also recommend stating the scope in the theorem title: **fixed H linear in S2**. The higher-degree symmetric-feature experiments in the earlier numerical reports are not automatically covered merely because their formulas are similar. The same strategy likely extends to them, but that extension should receive its own finite-mark statement and proof.

The review upgrades the earlier “pending arithmetic transfer” label only for the fixed family actually proved here, including the rational ten-coefficient trial. Historical reports should retain their original status and link to this review as the later resolution. No numerical global optimum, no general operator no-go, and no new zero-gap record is established by this upgrade.

Current report 03: residual_gram_round1

Source: `research/reports/residual_gram_round1.md`. SHA-256:

21b55679859ee1b67ea3b5c244b0d4367e5124be05e28d62aa13bcd02e877163.

Residual Gram arithmetic route: exact saturation, independent reproduction, and symmetric prime–factor trials

Research round 1, 5 September 2026. This note separates exact algebra, certified finite variational integrals, exploratory numerics, and missing arithmetic theorems. No half-gap theorem or famous conjecture was proved in this round.

1. Outcome and priority

The supplied Inoue calculation is independently reproducible. Its negative margin at normalized gap $\phi=1/2$ is not a numerical accident caused by a low polynomial degree. The reproduced degree-14 margin is -0.015357981703850554 , corresponding to additional normalized log-increment energy 0.3031544076323465 . This is a measured deficit for a particular family, not a universal impossibility theorem.

Three additional conclusions were obtained.

1. The original product-cutoff approximator is exactly saturated in its coefficient Hilbert space. Rewriting the same support with asymmetric factor lengths or additional Gram directions cannot give positive main-term recovery.
2. A new sparse arithmetic operator computes the diagonal main term for **arbitrary resonator coefficients**, not just the divisor-function ansatz. Full eigenvalue searches through one million coefficients remain below the half-gap threshold. These finite searches do not settle the asymptotic extremal problem.
3. A genuine extension of the resonator family, using symmetric statistics of its prime factors, improves the numerical half-gap margin to about -0.01465473 . This removes only about 4.6% of the original deficit. A simple rational instance of this extension has a rigorously enclosed continuum margin between -0.01467 and -0.01465 . Its arithmetic transfer has not been fully written, and its margin is negative.

The significant unresolved node remains a quantitative estimate for residual directions carrying information beyond the already accessible product cutoff, or a substantially different arithmetic resonator that survives the full operator test.

2. Source audit and an important historical correction

The primary source is [Shota Inoue, arXiv:2604.05733v1](#). It introduces a weighted second factorial moment of short-interval zero counts and proves, under RH, $\mu < 0.50895$. Its Theorem 4 allows arbitrary arithmetic coefficients under a product cutoff $L \leq T/(\log T)^2$; Theorem 2 specializes the coefficients to a divisor-function family. Zeros are counted with multiplicity. The displayed numerical trial is $\phi = .508949$, $\ell = 1.15$, $f(v) = 1 - .7v$.

The old claim $\mu \leq .50412$, cited in the 2019 Lagarias – Rodgers paper, is **not a valid current record**. [Goldston – Turnage-Butterbaugh, arXiv:1904.06001v2](#) was withdrawn on 11 November 2019: its generalized weights had to be symmetric in all their variables, invalidating the claimed calculation. This is directly relevant to the search below: every new prime-factor feature is explicitly symmetric.

The user's original `inoue_resonance_probe.py` and JSON became available after this round's independent implementation had already been run. Comparing those exact originals with the new code gives agreement within $1\text{e-}12$ for the published trial and within $1\text{e-}11$ for the degree-14 margin. The independently found finite-dimensional sign crossing differs by less than $2\text{e-}11$.

3. Exact coefficient–space saturation theorem

Fix a cutoff L , a resonator coefficient sequence r , and a log-increment coefficient sequence g . For any approximator coefficients a , put

$$b_a(q) = \sum_{km=q} a(k)r(m), \quad b_g(q) = \sum_{km=q} g(k)r(m), \quad q \leq L.$$

Equip coefficient vectors with

$$\langle b, c \rangle_L = \sum_{q \leq L} \frac{\overline{b(q)}c(q)}{q}.$$

Then the approximator-dependent diagonal main term satisfies the exact identity

$$2 \operatorname{Re} \langle b_a, b_g \rangle_L - \|b_a\|_L^2 = \|b_g\|_L^2 - \|b_a - b_g\|_L^2. \quad (\text{S1})$$

Consequently $a=g$ is a global maximizer of this main term for the fixed r, L . The maximizer need not be unique in a if the convolution map has a kernel, but its image $b_a=b_g$ is unique.

This is an identity in a finite coefficient Hilbert space. It should not be described as exact orthogonal projection in the time-domain Gaussian Hilbert space at finite T : that Gram matrix has small non-diagonal entries. The arithmetic theorem gives the passage from the latter to the diagonal main term and an error estimate.

Suppose a new direction C has coefficients supported on the same product set $q \leq L$. The main-term inner product of C with the discarded residual is zero. Its recovery term therefore equals $-|C|_L^2$, which is nonpositive. A change of basis, a more complicated inverse Gram, a nonlinear parameterization of a , or repeated completion of squares does not change this fact.

There is also the exact time-domain identity

$$\begin{aligned} 2 \operatorname{Re} \langle B_0, X \rangle - \|B_0\|^2 + 2 \operatorname{Re} \langle C, X - B_0 \rangle - \|C\|^2 \\ = 2 \operatorname{Re} \langle B_0 + C, X \rangle - \|B_0 + C\|^2. \end{aligned} \quad (\text{S2})$$

Thus “residual recovery” has not by itself enlarged the theorem: it is a larger approximator. Its value lies in designing a larger approximator whose mixed terms remain accessible.

The asymmetric–length caveat

Theorem 4 already permits setting $r(n)=0$ for $n > L_0$ while retaining product cutoff $L > L_0$. Therefore changing the relative factor lengths strictly inside the old product region is already contained in the original general theorem. It may improve a restricted trial family, but it is not a new arithmetic input. The useful analogy with the prime-gap 186 proof starts only when the actual accessible mixed products exceed what this theorem already controls.

4. Robust inverse–Gram certificate and the missing quantities

For candidate residual directions $C_1, \dots, C_m$, define

$$G_{ij} = \langle C_i, C_j \rangle, \quad b_i = \langle C_i, X - B_0 \rangle.$$

The recovery for a fixed coefficient vector c is

$$\Delta(c) = 2 \operatorname{Re}(c^* b) - c^* G c. \quad (G1)$$

If estimates $\hat{b}, \hat{G}$ have bounds

$$|b_i - \hat{b}_i| \leq \epsilon_i, \quad |G_{ij} - \hat{G}_{ij}| \leq E_{ij},$$

then the elementary certified lower bound is

$$\Delta(c) \geq 2 \operatorname{Re}(c^* \hat{b}) - c^* \hat{G} c - 2 \sum_i |c_i| \epsilon_i - \sum_{i,j} |c_i| |c_j| E_{ij}. \quad (G2)$$

An explicit rational c and entrywise enclosures suffice. Matrix inversion and a numerical condition number are not part of the final certificate. Singular Gram matrices cause no difficulty for (G2).

For the reproduced one-variable baseline, the additional gain must exceed approximately 0.3031544077 after division by the resonator norm. A design goal 0.31 allows only a modest error budget; it is not a universal lower bound on the energy needed by other methods.

Rank, trace, and Hilbert – Schmidt norm are insufficient to determine this recovery. For example the same rational matrix $\operatorname{diag}(1, 4)$ gives optimal recovery 1 for $b=(1, 0)$ and $1/4$ for $b=(0, 1)$. This is the precise location where directional information matters.

What a genuinely longer direction requires

For Gaussian packet width $W=T/\log T$, the transformed weight is significant when

$$|\log(km/q)| \lesssim W^{-1}.$$

At a product scale $M=T^{1+\epsilon}$, this retains additive differences of size roughly M/W , rather than only the exact equation $km=q$. The mixed moment contains prime-power coefficients $g(k)$, resonator coefficients $r(m)$, coefficients of $C(q)$, and a nontrivial oscillatory phase determined by the packet center. No estimate for generic arithmetic progressions automatically supplies this shifted multiplicative correlation. A proposed well-factorable or Type-II argument must state the actual variables, modulus, frequency range, coefficient norms, and total error after summing every direction.

5. A rigorous obstruction to a naive sparse long–support strategy

The following statement is elementary but useful for support design.

Sparse-tail lemma. Let $\epsilon > 0$, $M=T^{1+\epsilon}$, and $W=T/\log T$. Let a set S in $[M, 2M]$ have pairwise logarithmic spacing at least c/W , for some fixed $c > 0$. Suppose coefficients satisfy $|b(q)| \leq C_{\epsilon} q^{\epsilon}$ for every fixed positive ϵ . Then

$$\sum_{q \in S} \frac{|b(q)|^2}{q} = o(1). \quad (ST)$$

Indeed, the logarithmic interval has length $\log 2$, so $|S| \leq 1 + W \log(2)/c$. The displayed sum is at most a constant times $W M^{-(1+2\epsilon)}$. Choosing $2\epsilon(1+\eta) < \eta$ proves the claim. The same conclusion holds relative to a normalizing resonator norm that is bounded below by $T^{-(o(1))}$.

This applies to the usual fixed divisor-polynomial resonator coefficients. It rules out obtaining fixed recovery simply by selecting a very sparse, trivially separated portion of a genuinely long dyadic tail and then using its diagonal coefficient energy. It does **not** rule out dense structured supports, substantial off-diagonal mixed information, or different coefficient classes with quantitatively stronger concentration.

6. Independent reproduction of the original variational calculation

The new implementation uses a different simplex parameterization from the supplied original: the background mass v is the outside Gauss – Jacobi variable, and inserted masses occupy its complementary simplex. The basis

$$\psi_j(v) = \sqrt{2j + \ell^2} P_j^{(0, \ell^2 - 1)}(2v - 1)$$

is orthonormal for $v^{(\ell^2 - 1)} dv$. Its norm matrix is the identity. This avoids unstable monomial inversion in the original one-variable calculation.

Trial	ℓ	Half-gap margin
Degree 2	1.1720511457	-0.01539943397962
Degree 6	1.1762949914	-0.01535798218167
Degree 14	1.1762950386	-0.01535798170385

The published linear trial gives $1.487161813351623e-5$ at quadrature order 48. Orders 20, 32, and 48 agree to approximately $2e-15$. The degree-8 numerical zero is 0.5088369010687765 .

None of these floating calculations bounds all continuous functions or all $\ell \geq 1$. Raising the polynomial degree is not an independent arithmetic theorem. The exact saturation theorem in § 3 concerns a different variable, namely the approximator with the resonator fixed; it does not imply that the resonator family is globally optimal.

7. An exact arithmetic operator for arbitrary resonator coefficients

At $\phi = 1/2$, the leading linear term vanishes. Put $x_n = r(n)/\sqrt{n}$. Define a real sparse creation operator A by

$$A_{qm,m} = \frac{2 \sin(\frac{h}{2} \log q)}{e\sqrt{q}} \quad \text{when } q = p^e, \quad qm \leq L,$$

and zero otherwise. Here $h = \pi/\log T$; the exploratory runs set $\log L/\log T = 1$, the limiting boundary of the admissible range. Since $g(q) = -2i \sin((h/2) \log q)/e$, the complex convolution operator is $B = -iA$.

The approximator term is $\|Bx\|^2$, and the holomorphic correction is $\text{Re}\langle x, B^2x \rangle$. Therefore the exact diagonal main-term matrix is

$$K_L = A^*A + \frac{A^2 + (A^*)^2}{2}. \tag{A1}$$

Its maximal normalized half-gap margin is

$$\frac{\lambda_{\max}(K_L)}{2\pi^2} - \frac{1}{4}. \quad (\text{A2})$$

All coefficients of K_L are nonnegative at this parameter, so a real nonnegative Perron vector suffices for the extremal Rayleigh quotient. This is a substantially larger search space than the one-variable divisor family.

L	largest eigenvalue of K	margin	eigenvector residual norm
1,000	3.9492871367	-0.0499267764	3.0e-11
10,000	4.1058670454	-0.0419943455	4.6e-15
100,000	4.2052553801	-0.0369592737	2.0e-13
1,000,000	4.2738969159	-0.0334818529	3.1e-12

The threshold needed at the boundary is $\pi^2/2$, approximately 4.9348022. The full matrix at $L=10^6$ has 3,626,619 nonzero entries in A . Its sparse operator implementation avoids materializing K .

These are finite experiments. The choice $\log L/\log T=1$ is an asymptotic model, not an admissible finite theorem with $L=T$. An eigenvector for one finite L cannot be frozen and used as T tends to infinity. No extrapolated limit is claimed, and no global no-go follows from these negative eigenvalues.

A useful next question is whether

$$K_L = \frac{1}{2}(A + A^*)^2 + \frac{1}{2}[A^*, A]$$

admits a positive comparison or a Hardy-type bound uniformly in L . A proof that the normalized limiting norm is at most $1/4$ would close the entire diagonal method, rather than merely one family of weights. Such a proof has not been obtained.

8. A genuinely broader family: symmetric prime–factor statistics

The finite arithmetic eigenvector search motivates coefficients of the form

$$r(n) = d_l(n)H(v, S_2(n), S_3(n), \dots), \quad v = \frac{\log n}{\log L}, \quad S_k(n) = \sum_{p|n} \left(\frac{\log p}{\log L}\right)^k. \quad (\text{PF1})$$

A Liouville sign can be included when the sign of the linear term calls for it. At the half-gap boundary the quadratic terms do not distinguish that sign. Every S_k and every polynomial in these variables is symmetric in the prime factors. These features do not repeat the unsymmetric-weight error in the withdrawn 2019 calculation.

For the continuum experiment, use a background mass v and its scaled Poisson – Dirichlet prime-factor partition with parameter $a=e^{1/2}$. Its relevant moments are

$$ES_k = \frac{\Gamma(k)\Gamma(a+1)}{\Gamma(a+k)} v^k,$$

and

$$ES_k S_l = \frac{\Gamma(a)[a\Gamma(k+l) + a^2\Gamma(k)\Gamma(l)]}{\Gamma(a+k+l)} v^{k+l}. \quad (\text{PF2})$$

More generally, for a list $k_1, \dots, k_s$,

$$\mathbb{E} \prod_{j=1}^s S_{k_j} = \frac{\Gamma(a)}{\Gamma(a + \sum_j k_j)} \prod_{\pi} \prod_{B \in \pi} \Gamma \left(\sum_{j \in B} k_j \right), \quad (\text{PF3})$$

where the sum runs over set partitions of the labeled factors. Adding a prime of logarithmic size u replaces (v, S_k) by $(v+u, S_{k+u^k})$. This gives explicit Gram and quadratic-form integrals for each symmetric feature polynomial.

The calculation reduces exactly to the independently reproduced Inoue form when H depends only on v ; this reduction was checked at three different values of ell and by two independently written feature implementations.

Symmetric features, each with a small mass polynomial	optimized ell	numerical margin at $\phi=1/2$
1 only, degree 6	1.176295	-0.0153579822
1, S2, degree 4	1.065285	-0.0146618161
1, S2, S3, degree 4	1.077117	-0.0146563150
1, S2, S3, S4, degree 4	1.081082	-0.0146551195
12 groups including S2 squared, S2 S3, S2 cubed, S2 fourth power	1.079948	-0.0146547256

The richer Gram matrices become ill-conditioned. Directions below an explicit eigenvalue cutoff were discarded. Therefore the last decimals and any inference of saturation in this expanded family are exploratory only. The main robust conclusion is that these particular natural extra arithmetic features recover a small fraction of the deficit, far less than required to cross $1/2$.

Arithmetic status of this extension

The displayed partition moments define a coherent continuum calculation. To turn it into a zeta theorem one must prove the corresponding asymptotics for weighted integer sums, including repeated-prime terms, coincidences between inserted primes and the background integer, the short-background range, and the errors after normalizing by the resonator norm. The original theorem handles arbitrary r , but it does not automatically supply this new explicit asymptotic evaluation. No such transfer is claimed as completed here.

9. An exact rational certificate for one expanded continuum trial

To avoid trusting the ill-conditioned optimizers, a simple trial was rounded to rational coefficients. Take $\text{ell}=16/15$ and

$$H(v, S_2) = f(v) + g(v)S_2,$$

with

$$f(v) = \frac{145 + 3v - 116v^2 + 71v^3 - 6v^4}{100},$$

$$g(v) = \frac{-563 + 1682v - 2479v^2 + 1751v^3 - 488v^4}{100}.$$

The script `rational_trial_certificate.py` proves, by rational arithmetic, that its continuum half-gap margin lies in

$$-\frac{1467}{100000} < J < -\frac{1465}{100000}. \quad (\text{RC})$$

A much tighter outward enclosure is stored in the JSON certificate; its decimal presentation is approximately

`[-0.014662375473368995, -0.014662375473368974]`.

The method is completely explicit:

- Machin's formula and alternating arctangent series give rational lower and upper bounds for π .
- Sine and cosine are replaced by finite Taylor polynomials with explicit remainder bounds on the simplex.
- The symmetric feature expectations are rational because $e^{11^2}=256/225$ and all moment degrees are integers.
- Every simplex monomial integral is a factorial divided by a finite product of rational numbers.
- The numerator and positive denominator are enclosed before division. Final assertions compare exact rational numbers.

During development, an off-by-one error in a simplex beta denominator was caught by disagreement with the independent floating integral, then corrected. Explicit constant-monomial identities are now asserted in the certificate script. The final run succeeds. This history is preserved to make clear why multiple forms of verification matter.

The certificate proves a statement about **one explicit continuum trial**, not a global upper bound and not an arithmetic half-gap improvement. In particular a negative trial is not a proof that the entire method fails.

10. AH, multiplicity, and positive-density targets

There are several different statements called the Alternative Hypothesis.

The 2019 Lagarias – Rodgers formulation, Conjecture 2.2, explicitly takes the limiting consecutive-gap values in $\{1/2, 1, 3/2, \dots\}$, excluding zero, for all sufficiently large indices. A theorem $\liminf g_n < 1/2$ would refute this strong formulation. It is incorrect to say that this particular source allows zero gaps.

Baluyot – Goldston – Suriajaya – Turnage-Butterbaugh 2025 studies a formulation that does not assume simplicity. For that broader setting, repeated zeros or gaps tending to zero require separate treatment. A small-gap theorem counting multiplicity may detect only those zeros. A statement allowing a density-zero exceptional set is also not contradicted merely by infinitely many exceptional gaps.

A robust target is a positive density of pair separations in a fixed interval whose closure avoids every permitted half-integer, together with whichever local-count control is needed to translate pair information into adjacent-gap information. A positive proportion of weighted factorial moment is not automatically an unweighted positive proportion: the resonator could concentrate on a negligible set of heights.

For an unweighted count there is a useful exact detector. Write

$$Q(h) = \int N_h(t)(N_h(t) - 1) dt = 2 \sum_{i < j} (h - | \gamma_i - \gamma_j |)_+$$

when boundary terms are absent or handled explicitly. For $0 < h_0 < h_1 < h_2$ in arithmetic progression,

$$Q(h_2) - 2Q(h_1) + Q(h_0)$$

has a nonnegative triangular pair kernel supported in $[h_0, h_2]$, with value zero at separation zero. Thus it removes the direct multiplicity atom exactly. A smooth nonconstant arithmetic weight changes the kernel and must be analyzed rather than silently canceled. Also a close nonadjacent pair need not give an adjacent gap bounded away from zero without local

cluster control. This detector is a possible way to formulate a stronger arithmetic target; its positivity has not been proved from the available zeta estimates.

11. Artifacts and reproducibility

All files are in `research-round1/residual-gram/`, separate from the recovered original source tree.

- `inoue_variational.py`, `variational-results.json`: independent one-variable reproduction.
- `arithmetic_operator.py`, `arithmetic-results.json`, `arithmetic-eigenvector.npz`: arbitrary-coefficient sparse arithmetic search through $L=10^6$.
- `prime_feature_variational.py`, `prime-feature-results.json`: linear symmetric prime-factor features.
- `general_prime_features.py`, `general-prime-feature-results.json`: products of symmetric prime moments.
- `rational_trial_certificate.py`, `rational-trial-certificate.json`: exact rational continuum-integral certificate.
- `check_algebra.py`, `algebra-check-results.json`: structural completion-of-squares identities, operator sign check, reduction of both feature implementations to the original form, and comparison against the now recovered original result file.
- `inoue-paper.html`: retrieved primary-paper snapshot for formula checking.
- `*-run.log`: execution records.

Commands used:

```
OPENBLAS_NUM_THREADS=1 python3 research-round1/residual-gram/inoue_variational.py
OPENBLAS_NUM_THREADS=1 python3 research-round1/residual-gram/arithmetic_operator.py
OPENBLAS_NUM_THREADS=1 python3 research-round1/residual-gram/prime_feature_variational.py
OPENBLAS_NUM_THREADS=1 python3 research-round1/residual-gram/general_prime_features.py
python3 research-round1/residual-gram/rational_trial_certificate.py
OPENBLAS_NUM_THREADS=1 python3 research-round1/residual-gram/check_algebra.py
```

The algebra test suite is not a proof assistant. The rational certificate's stated finite integral inequality is exact rational computation with explicit analytic remainder estimates. The spectral searches and their optimizer outputs remain numerical.

12. Recommended next bounded attack

First investigate a uniform upper or a counterexample for the full arithmetic operator (A1), including its commutator structure and the shape of finite Perron vectors. This is the quickest way to tell whether more elaborate resonators inside the old product region can matter at all. A uniform upper would justify making genuine long-support mixed moments the primary target; a counterexample would identify a more concrete family to analyze.

In parallel with that mathematical question, specify one dense residual coefficient family and compute its exact near-diagonal product conditions. The theorem request must be an explicit sum with an explicit normalization and error target. It should not merely say “apply Guth – Maynard,” “use a Yau flow,” or “recover the residual Gram.” The measured margin supplies a useful numerical target, but the required new arithmetic input remains the central unresolved issue.

Current report 04: residual_gram_round2

Source: `research/reports/residual_gram_round2.md`. SHA-256:
`1e625af13fe203ed3e9d70f78bf191fba20dcec10cf10ac2ebaed8b0de36f18c`.

Full arithmetic operator: extended spectra, Schur bounds, and prime-bin Fock models

Research round 2, 5 September 2026. This is a reproducible continuation of `research-round1/residual_gram.md`. No half-gap breakthrough or universal impossibility theorem was obtained. The strongest exact statement remains coefficient-space approximator saturation. This round tests the much larger remaining resonator problem.

1. Concrete result

For arbitrary resonator coefficients, the saturated half-gap main term is governed by

$$K_L = A^* A + \frac{1}{2}(A^2 + (A^*)^2), \quad A_{p^\theta m, m} = \frac{2 \sin(\frac{\pi}{2} \log(p^\theta) / \log L)}{e \sqrt{p^\theta}}.$$

The finite experiments set `log L / log T=1`, the limiting boundary of the admissible logarithmic range. The target is

$$\limsup_{L \rightarrow \infty} \lambda_{\max}(K_L) \stackrel{?}{\leq} \frac{\pi^2}{2} = 4.934802200544679 \dots$$

The normalized target is `lambda/(2 pi^2)-1/4<=0`. The raw operator is **not** being compared with `1/4`.

The full sparse eigenvalue search was extended to ten million coefficients:

L	lambda max	normalized margin	eigenvector residual
1,000,000	4.273896915919377	-0.03348185285686	3.14e-12
3,000,000	4.299660502530880	-0.03217665431168	8.37e-12
10,000,000	4.324089558989538	-0.03093906385385	2.04e-11

The last creation matrix has 37,861,249 nonzero entries. `K_L` is applied as a sparse operator, not materialized. The ten-million run took approximately 27 seconds on the available machine, excluding any subsequent serialization. This timing is a recorded run, not a performance guarantee.

The values rise with L and remain negative. They neither prove a limiting upper bound nor justify extrapolating to a particular constant. A fixed finite eigenvector is not itself a valid asymptotic zeta certificate.

2. Finite eigenvector structure: the additional features nearly exhaust the tested problem

At `L=10^6`, the true top eigenvalue is `4.273896915919377`. We fitted its normalized Perron vector using vectors

$$\frac{d_i(n)}{\sqrt{n}} \nu^i F(n), \quad \nu = \frac{\log n}{\log L}, \quad 0 \leq i \leq 4.$$

The first feature set is just $F=1$. The next uses $\{1, S_2, S_3, S_4\}$, with $S_k = \sum_{p|n} (\log p / \log L)^k$. Larger fits add the largest logarithmic prime factor and then the number of distinct prime factors $\omega(n)$. Every reported Rayleigh quotient is evaluated again using the full finite sparse arithmetic operator.

	ell	feature set	terms	L2 fit error	Rayleigh deficit from full optimum
	1.08	mass only	5	0.09399054	0.0325942265
	1.08	symmetric prime powers	20	0.01813753	0.0007337046
	1.08	plus largest prime factor	30	0.01524505	0.0006332664
	1.08	plus omega and omega squared	40	0.01134669	0.0004437114

The 40-term fit reaches 4.273453204513651 , within about 0.011% of the full tested maximum. The 20-term fit already reaches 4.273163211325686 .

This is evidence that the finite operator's maximizing direction has considerable multiplicative structure and is well approximated by the symmetric prime-factor family tested in round one. It is not evidence that the same finite family is asymptotically complete. The role of $\omega(n)$ can change with L , and a small finite- L residual does not bound unseen asymptotic directions.

The fits were also made at $\text{ell}=1$ and $\text{ell}=1.1763$; all results, including worse fits, are retained in JSON. Least-squares truncation uses an explicit numerical singular-value threshold, so these fits are not exact interval certificates.

3. A scalar Cauchy majorant for the full operator

This section gives a candidate analytic bound with explicit derivation. Its numerical optimization does not reach the desired threshold.

Write $a_q = 2 \sin((\pi/2) \log q / \log L) / (e \sqrt{q})$ for $q = p^e$. The divisor identity

$$\sum_{q|n} \Lambda(q) = \log n$$

and Cauchy – Schwarz give

$$\|Ax\|^2 \leq \sum_{m \leq L} |x_m|^2 \sum_{q \leq L/m} \frac{a_q^2 \log(qm)}{\Lambda(q)}. \quad (\text{B1})$$

For the other part, take a positive function Y on $(0, 1]$, and set $v_n = \log n / \log L$. Applying $2ab \leq a^2 + b^2$ to each product gives

$$\begin{aligned} \text{Re}\langle x, A^2 x \rangle &\leq \frac{(\log L)^2}{2} \sum_m |x_m|^2 \sum_{k \leq L/m} \frac{a_k^2 a_l^2}{\Lambda(k)\Lambda(l)} Y(v_{klm}) \\ &\quad + \sum_n \frac{|x_n|^2}{2Y(v_n)(\log L)^2} \sum_{kl|n} \Lambda(k)\Lambda(l). \end{aligned} \quad (\text{B2})$$

Terms without prime-power factors are zero. Since the pairs $kl|n$ are a subset of pairs of divisors of n ,

$$\sum_{k|l \mid n} \Lambda(k)\Lambda(l) \leq (\log n)^2.$$

Let

$$f(u) = \frac{\sin^2(\pi u/2)}{u^2}, \quad f(0) = \pi^2/4,$$

and

$$B(v) = 4 \int_0^{1-v} (v+u)f(u) \, du.$$

Replacing the prime sums in (B1) – (B2) by their elementary prime-number asymptotics suggests the explicit bound

$$\limsup_L \lambda_{\max}(K_L) \leq \sup_{0 \leq v \leq 1} \left[B(v) + \frac{v^2}{2Y(v)} + 8 \int_{u+w \leq 1-v} f(u)f(w)Y(v+u+w) \, du \, dw \right]. \tag{B3}$$

For a fixed positive smooth Y bounded away from zero, the passage from the finite inequalities to (B3) can be justified uniformly by partial summation. Contributions of powers p^{ϵ} , $\epsilon \geq 2$, are $O(1/\log L)$: use the small-argument bound for sine and convergence of the corresponding prime-power reciprocal sums. One should write those uniform remainder estimates in a final proof rather than cite this numerical note as their formal verification. Profiles vanishing at $v=0$ can first be regularized by adding a fixed positive epsilon.

This is related to the weighted elementary inequalities in Inoue – Kobayashi – Toma, [arXiv:2510.14309](#), but it treats the quadratic operator of the newer resonance-correlation method. Their existing linear-method barrier cannot simply be reused as an upper bound for K .

Numerical result: this majorant still fails

A constant Y gives a best sampled bound near **5.28075150**, larger than $\pi^2/2$. Allowing Y to vary leads to a backward Volterra equation. With

$$C(s) = 8 \int_0^s f(u)f(s-u) \, du,$$

the equality profile at a proposed upper bound λ satisfies

$$Y(v) = \frac{v^2}{2[\lambda - B(v) - \int_v^1 C(t-v)Y(t) \, dt]}. \tag{V}$$

A discretized backward recursion and bisection gave:

grid intervals	numerical threshold λ
250	5.2073898741
500	5.2073943491
1,000	5.2073954679
2,000	5.2073957475
4,000	5.2073958175

The corresponding normalized bound is approximately **+0.01380975** . It cannot prove a half-gap no-go. These are numerical values for the majorant search, not a certified global optimum over $\mathcal{Y}$. The failure is preserved because it identifies a specific loss: separately estimating the norm square and the holomorphic term by Cauchy loses the compatibility of their near-extremizers. More precision in the same grid would not close a gap of this size.

4. Prime bins and a finite bosonic creation model

A second approach is to retain more of the prime-factor combinatorics.

Divide logarithmic prime sizes into bins $((j-1)/N, j/N]$, $1 \leq j \leq N$, and set

$$c_j^2 = \int_{(j-1)/N}^{j/N} \frac{4 \sin^2(\pi u/2)}{u} du.$$

Use occupation vectors $\mathbf{k}=(k_1, \dots, k_N)$ with the rounded-up energy restriction

$$\sum_{j=1}^N j k_j \leq N.$$

The finite bosonic matrix is

$$(A_N)_{k+e_j, k} = c_j \sqrt{k_j + 1}$$

whenever the new state remains in the energy set. Form

$$K_N = A_N^* A_N + \frac{1}{2}(A_N^2 + (A_N^*)^2).$$

Every mode and every occupation vector retained by this construction is explicit. The rounding restricts the support; it is not an upper discretization of the full continuum norm.

bins N	state dimension	largest eigenvalue	normalized margin
12	272	3.9285784722	-0.0509758896
20	2,714	4.0882405964	-0.0428873119
28	18,460	4.1870051368	-0.0378838418
36	99,133	4.2539816177	-0.0344907736
40	215,308	4.2800240103	-0.0331714506

At $N=40$ the creation matrix has 963,320 nonzero entries. Its eigenvector residual is approximately **1.6e-11** . No positive half-gap certificate occurs in these models.

A proposed fixed-bin arithmetic lower-transfer proof

There is a concrete route from each fixed finite bin model to arithmetic resonators; this is more modest than identifying the entire limiting operator norm.

For every occupation vector, take the normalized vector supported on squarefree integers whose prime factors have those bin counts, with coefficient proportional to the product of their prime weights $\mathbf{a}_p$. Rounded-up energy ensures that all such integers are at most L . Different occupation vectors have disjoint supports.

If E_k is the elementary symmetric sum of order k in one bin's squared prime weights, the exact raising coefficient between its collective vectors is

$$(k+1)\sqrt{E_{k+1}/E_k}.$$

For fixed bin resolution and fixed k , the prime-number theorem gives the total squared weight tending to c_j^2 ; the maximum individual squared weight tends to zero. Consequently

$$E_k \rightarrow (c_j^2)^k/k!,$$

and the raising coefficient tends to $c_j \sqrt{k+1}$. This handles repeated occupation of a bin through **distinct primes**. It does not pretend that repeated copies of the same prime are independent bosons.

Let P project onto the retained collective squarefree vectors. The exact inequality

$$PA^*AP \geq (PAP)^*(PAP)$$

retains the nonnegative loss from intermediate states outside the subspace. Raising twice cannot leave the rounded-energy set and later return. Squareful states created by a repeated prime or a prime power cannot return to the squarefree subspace by further multiplication. These observations give the expected convergence of the holomorphic compressed term to A_N^2 . The proposed conclusion is

$$\liminf_{L \rightarrow \infty} \lambda_{\max}(K_L) \geq \lambda_{\max}(K_N) \quad \text{for every fixed } N. \quad (\text{FT})$$

This argument is a proof draft requiring independent line-by-line review and a complete statement of the finite-bin normalizations. It gives a **lower** search mechanism. It gives no upper bound for the full arithmetic operator and no equality with a limiting Fock norm. A positive rationally certified finite-bin value above $\pi^2/2$ would be a concrete success criterion, but all values computed here are below that threshold.

5. Why the continuum limit is delicate: the small-prime cloud

The divisor-family resonators suggest an infrared phenomenon. Under the $d_{\text{ell}}(n)^2/n$ weighting, the number of logarithmically small prime factors grows on the scale $\text{ell}^2 \log \log L$, while their total logarithmic mass remains bounded. Therefore a finite-particle or coarse energy-bin model can converge slowly even when the creation kernel itself is square-integrable. One cannot infer a precise limit from the modest bin table.

A useful candidate representation uses the scaled Poisson – Dirichlet partition of mass v , with parameter $a=\text{ell}^2$, as the reference measure. In that representation the creation operator acts by removing a part u with coefficient $2 \sin(\pi u/2)/\text{ell}$, while its adjoint inserts a part with intensity $2 \text{ell} \sin(\pi u/2) du/u$, subject to total mass at most one. The symmetric prime-moment polynomials of round one are a natural Galerkin family in this representation.

This description explains the algebraic insertion formulas and suggests a more stable basis. It is **not** a proved statement that every asymptotic arithmetic resonator belongs to one such representation, that different ell sectors exhaust the full norm, or that finite-bin spectra converge to the desired universal operator. Those are separate operator-limit problems. The original PD arithmetic asymptotic transfer also remains under independent audit.

6. Files, dependencies, and commands

Round-two files are in `research-round2/residual-gram/`. They refer to the round-one `arithmetic_operator.py` implementation through a relative path. The source snapshot is not modified.

- `extended_arithmetic.py` and `extended-arithmetic-results.json`: reproduce the 3-million and 10-million full spectra. The wrapper writes only round-two outputs.
- `eigenvector_features.py` and `eigenvector-feature-results.json`: fit the round-one million-dimensional Perron vector and re-evaluate each candidate using the full operator.
- `first_upper_bound.py`, `first-upper-bound-results.json`: constant-weight scalar majorant.
- `volterra_upper_search.py`, `volterra-upper-results.json`, `volterra-upper-profile.npz`: variable-weight majorant search.
- `boson_cutoff.py`, `boson-cutoff-results.json`: all finite prime-bin models listed above.
- `*-run.log`: retained execution records.

Dependencies are Python, NumPy, and SciPy. The runs used the machine's `python3` (Homebrew Python 3.14), not the bundled PDF runtime. The rational certificate in round one additionally requires SymPy. For reproducibility, use the same environment that ran the round-one numerical files; no private service, GPU, or new model invocation is required.

```
OPENBLAS_NUM_THREADS=1 python3 research-round2/residual-gram/extended_arithmetic.py
OPENBLAS_NUM_THREADS=1 python3 research-round2/residual-gram/eigenvector_features.py
OPENBLAS_NUM_THREADS=1 python3 research-round2/residual-gram/first_upper_bound.py
OPENBLAS_NUM_THREADS=1 python3 research-round2/residual-gram/volterra_upper_search.py
OPENBLAS_NUM_THREADS=1 python3 research-round2/residual-gram/boson_cutoff.py
```

The 10-million eigenvector archive is about 85 MB and is optional research data. It should not be included in a small source-code PR by default. Pass `--save-largest` to the wrapper to regenerate it. The eigenvalue JSON and scripts are sufficient to reproduce the reported values. The fitting script uses the smaller, already archived round-one eigenvector.

7. What this round rules out, and what it leaves open

The exact saturation theorem rules out constant main-term recovery by changing only the approximator inside the fixed product space. The sparse-tail lemma rules out a specified separated-tail tactic for divisor-bounded coefficients. Those are genuine statements with explicitly limited scopes.

The new negative finite spectra, the near-optimal finite fits, the negative bosonic bins, and the failed Schur majorant are **not** universal no-go theorems. Together they make a large hidden gain from a routine reparameterization less plausible, but that judgment is a research inference rather than a proof.

A precise next analytic success criterion is an explicit positive comparison profile or an operator identity proving a uniform upper below $\pi^{2/2}$, with every prime-power and cutoff error included. A precise alternative success criterion is a finite prime-bin or other explicit coefficient family with a certified limiting Rayleigh quotient above that value. Merely extending the grid or reporting another negative optimized decimal would not meet either criterion.

Current report 05: centered_gaussian_mixed_moments

Source: `research/reports/centered_gaussian_mixed_moments.md`. SHA-256:

5aec8de6c69d3b8d67799a65d0598d49ff1a87a5878ebc548dd3c90d806e1716.

Centered Gaussian packets: an exact mixed identity and an unavoidable pole debt

Research round 2, 2026-09-05. This note settles the proposed shortcut of replacing Inoue's shifted Gaussian packet by a Gaussian centered at zero and then using Fourier positivity to retain long mixed products for free. The Fourier kernel is positive, but the exact critical-line identity contains a pole-cut correction with the opposite phase on the relevant long-support region. That correction can exceed the purported normalization by an unbounded factor for coefficients bounded by one.

Main conclusion. A centered packet may be used only if its pole term, its true nondiagonal Gram form, and its low-height zero-count terms are retained. Dropping those terms is false. This is a precise obstruction to the proposed free support enlargement, not an impossibility theorem for every centered-packet method.

1. Conventions and source input

Let

$$\Phi_W(t) = e^{-t^2/(2W^2)}, \quad \widehat{V}(\xi) = \int_{\mathbb{R}} V(t) e^{-2\pi i t \xi} dt, \quad W > 0.$$

Thus $\widehat{\Phi}_W(\xi) = \sqrt{2\pi} W e^{-2\pi^2 W^2 \xi^2} > 0$. The height parameter T and packet width W are distinct. Below $h = \pi/\log T$, while W can be T or $T/\log T$. A weight written as e^{-t^2/T^2} corresponds to $W = T/\sqrt{2}$.

Use the horizontal branch

$$\log \zeta(\sigma + it) = \int_{\infty}^{\sigma} \frac{\zeta'}{\zeta}(a + it) da$$

away from zero ordinates and $t = 0$, with the mean of the two boundary values at exceptional ordinates. Put

$$\Delta_h(t) = \log \zeta\left(\frac{1}{2} + i(t + h/2)\right) - \log \zeta\left(\frac{1}{2} + i(t - h/2)\right),$$

$$g_h(k) = \frac{\Lambda(k)}{\log k} (k^{-ih/2} - k^{ih/2}).$$

For $k = p^e$, $g_h(k) = -2i \sin((h/2) \log k)/e$; otherwise it vanishes.

The analytic input is Inoue, Lemma 5.5 and equations (5.42) – (5.43). For an analytic test function on the strip $-3/2 \leq \operatorname{Im} z \leq 0$ with the stated horizontal decay, its critical-line logarithm integral equals the Dirichlet contribution, zero-cut integrals, and a **negative pole-cut integral**. Under RH the zero-cut integrals to the right of the critical line disappear. We retain the pole term exactly, rather than invoking the small-error estimate used for the shifted packet.

The tests $V(z) = x^{-iz}\Phi_W(z)$ satisfy the strip and decay assumptions for every fixed $x > 0$, $W > 0$. Arbitrarily long but finite Dirichlet polynomials may consequently be expanded against this identity. No absolutely convergent Dirichlet series for $\log \zeta$ on the critical line is assumed.

2. Exact kernel identity

Theorem 1. Assume RH. For every $x > 0$, $h > 0$, and $W > 0$,

$$\begin{aligned} K_{h,W}(x) &:= \int_{\mathbb{R}} \Delta_h(t) x^{-it} \Phi_W(t) dt \\ &= \sqrt{2\pi} W \sum_{k \geq 2} \frac{g_h(k)}{\sqrt{k}} \exp\left[-\frac{W^2}{2} \log^2(kx)\right] + P_{h,W}(x), \end{aligned} \quad (1)$$

where

$$\begin{aligned} P_{h,W}(x) &= -2\pi \int_0^{1/2} [V(-h/2 - i\sigma) - V(h/2 - i\sigma)] d\sigma \\ &= \boxed{-4\pi i e^{-h^2/(8W^2)} \int_0^{1/2} x^{-\sigma} e^{\sigma^2/(2W^2)} \sin\left[\frac{h}{2} \left(\log x - \frac{\sigma}{W^2}\right)\right] d\sigma}. \end{aligned} \quad (2)$$

The series in (1) is absolutely convergent: the Gaussian decay in $\log k$ dominates the $k^{-1/2}$ factor and the number of integers at each logarithmic scale.

Proof. Apply the cited contour identity at shifts $v = h/2$ and $v = -h/2$ and subtract. RH removes zero cuts with real part greater than $1/2$. The pole at $s = 1$ leaves the first line of (2). The Fourier transform of $x^{-it}\Phi_W(t)$ at $\log k/(2\pi)$ is precisely the Gaussian in (1). Finally,

$$V(-h/2 - i\sigma) = x^{-\sigma} e^{\sigma^2/(2W^2) - h^2/(8W^2)} e^{i(h/2)(\log x - \sigma/W^2)},$$

and $V(h/2 - i\sigma)$ has the conjugate phase. Their difference proves (2). This derivation uses an identity obtained by shifting to an absolutely convergent half-plane; it does not formally expand the critical-line logarithm.

2.1 The sign on the genuine long-support region

If

$$0 < x \leq 1, \quad 0 < h(\log(1/x) + \frac{1}{2W^2}) < 2\pi, \quad (3)$$

then $P_{h,W}(x)$ is purely imaginary with **strictly positive imaginary part**. At $x = 1$ the integral is still positive because the phase is strictly negative for $\sigma > 0$.

In contrast, $g_h(k)$ is negative imaginary whenever $0 < h \log k < 2\pi$. Hence at the matching frequency $x = 1/k$, the pole correction has the opposite phase to the positive-kernel Dirichlet contribution.

The sign statement is deliberately restricted to (3). It is false for all positive x : for example $x = 2$, $h = 0.2$, $W = 3$ gives a negative imaginary correction. This is not a defect of the theorem; it is why support information matters.

3. Exact mixed and Gram forms for arbitrary finite supports

Let

$$R(t) = \sum_m \frac{r_m}{\sqrt{m}} m^{-it}, \quad C(t) = \sum_q \frac{c_q}{\sqrt{q}} q^{-it}$$

be arbitrary finite Dirichlet polynomials. Then (1) gives the exact identity

$$\begin{aligned} \langle C, \Delta_h R \rangle_{\Phi_W} &= \sqrt{2\pi} W \sum_{k,m,q} \frac{g_h(k) r_m \overline{c_q}}{\sqrt{kmq}} e^{-\frac{1}{2} W^2 \log^2(km/q)} \\ &\quad + \sum_{m,q} \frac{r_m \overline{c_q}}{\sqrt{mq}} P_{h,W}(m/q). \end{aligned} \quad (4)$$

Here $\langle C, X \rangle_{\Phi_W} = \int X(t) \overline{C(t)} \Phi_W(t) dt$; the notation specifies the order to prevent an unnoticed conjugation change.

The penalty is also exactly

$$\|C\|_{\Phi_W}^2 = \sqrt{2\pi} W \sum_{q,r} \frac{c_q \overline{c_r}}{\sqrt{qr}} e^{-\frac{1}{2} W^2 \log^2(q/r)}. \quad (5)$$

For long dense support, (5) is not the diagonal coefficient norm. Its off-diagonal terms are positive when all c_q share one phase. Fourier positivity therefore enlarges the penalty as well as the suggested mixed gain.

Suppose $r_m \geq 0$, $c_q = -id_q$ with $d_q \geq 0$, and every retained q exceeds every retained m . If (3) holds for all $x = m/q$, every pole summand in the second line of (4) is negative real. This is the natural configuration when a new residual direction lives in a genuinely longer tail than the existing resonator. The positive Fourier kernel does not remove its debt.

For $q \leq T^{1+\eta}$ with fixed $0 < \eta < 1$ and $h = \pi/\log T$, the effective prime-power range in (4) has the common phase. Terms beyond the first sine sign change are Gaussian-small because $\log(km/q)$ is then bounded below by a positive multiple of $\log T$. They can be bounded absolutely. The pole term remains and is the substantive obstruction.

4. A counterexample to treating the pole term as a negligible uniform error

The following example uses coefficients bounded by one and no conjecture about primes in short intervals.

Theorem 2. Fix $0 < \eta < 1$. Let $M = T^{1+\eta}$, $h = \pi/\log T$, and let $W = T$ or $W = T/\log T$. Take

$$R(t) = 1, \quad C(t) = -i \sum_{M < q \leq 2M} \frac{1}{\sqrt{q}} q^{-it}.$$

The pole contribution to (4) is negative and satisfies

$$\boxed{E_{pole} = -(4\pi \cos \frac{\pi\eta}{2} + \alpha(1)) \frac{M}{\log M}.} \quad (6)$$

In particular,

$$\frac{|E_{pole}|}{\sqrt{2\pi} W} \rightarrow \infty, \quad (7)$$

even though $\sum |r_m|^2/m = 1$ and $\sum |c_q|^2/q \rightarrow \log 2$.

Proof. Uniformly for $q \in (M, 2M]$, formula (2) and $W \rightarrow \infty$ give

$$P_{h,W}(1/q) = 4\pi i \sin\left(\frac{h}{2} \log q\right) \frac{\sqrt{q}-1}{\log q} (1 + \alpha(1)). \quad (8)$$

The sine is bounded away from zero in this range, since it tends to $\cos(\pi\eta/2) > 0$. The factors $e^{\sigma^2/(2W^2)}$, $e^{-h^2/(8W^2)}$, and the additional sine argument $h\sigma/(2W^2)$ change the integral by a uniformly vanishing relative error. The remaining elementary integral is

$$\int_0^{1/2} q^\sigma d\sigma = \frac{\sqrt{q}-1}{\log q}.$$

Multiply (8) by $\overline{c_q}/\sqrt{q} = i/\sqrt{q}$ and sum. The sum of $(1 - q^{-1/2})/\log q$ over all integers in $(M, 2M]$ is $(1 + o(1))M/\log M$. This proves (6). Finally $M/(W \log M)$ tends to infinity for both specified widths.

The theorem refutes a uniform long-support identity of the form “positive Dirichlet Gaussian sum plus an $\mathcal{O}(W)$ error controlled by the diagonal norms.” It does not assert that the full mixed moment itself is negative. The prime contribution and pole contribution may substantially cancel, and the surviving fluctuation is exactly what needs arithmetic control.

4.1 Why the apparent gain resembles a prime–density term

For a single frequency $x = 1/X$, the Gaussian Dirichlet sum samples prime powers near X on multiplicative width $1/W$. If one replaces the prime density heuristically by $dy/\log y$, its leading imaginary contribution is

$$-4\pi i \sin\left(\frac{h}{2} \log X\right) \frac{\sqrt{X}}{\log X}.$$

This has the same magnitude and opposite sign as (8). The replacement is an interpretation, **not** a proved short-interval prime estimate in the relevant ranges. It explains the structural cancellation: centering transforms the proposed free positive contribution into a prime-density-subtracted quantity. Proving useful positivity of that remainder still requires arithmetic information.

5. The origin and the zero–count formula cannot be discarded

Let $N(t)$ be the signed zero-counting function and let $\theta(t)$ be the odd Riemann – Siegel theta function. With the branch above, the exact global form away from endpoints is

$$N(t) = \frac{\theta(t)}{\pi} + S(t) + \operatorname{sgn}(t), \quad S(t) = \frac{1}{\pi} \operatorname{Im} \log \zeta\left(\frac{1}{2} + it\right).$$

Consequently,

$$N_h(t) = \frac{\theta(t+h/2) - \theta(t-h/2)}{\pi} + S(t+h/2) - S(t-h/2) + 2\mathbf{1}_{\{|t| < h/2\}}. \quad (9)$$

The indicator is absent only when the interval avoids the origin. As $t \downarrow 0$, this branch has $S(t) \rightarrow -1$; as $t \uparrow 0$, it has $S(t) \rightarrow 1$. At $t=0$ the jump in S and the indicator in (9) cancel. There are no zeros in a sufficiently small neighborhood of zero because $\zeta(1/2) \not\equiv 0$.

A packet concentrated near heights comparable to T may replace the theta increment by $(h/2\pi) \log T + \mathcal{O}(h)$ and suppress the origin. A centered packet requires a weighted argument proving that low heights are negligible. Pointwise replacement by $\log T$ on the entire real line is false.

5.1 A direct coherence bound for a positive long resonator

Let

$$R_M(t) = \sum_{M < n \leq 2M} n^{-1/2-it}.$$

Factor out M^{-it} . For $|t| \leq 1$, all remaining phases have arguments in an interval of length at most $\log 2 < \pi/2$. Therefore

$$|R_M(t)| \geq \cos(\log 2) \sum_{M < n \leq 2M} n^{-1/2} \geq c\sqrt{M}.$$

For $W \geq 1$,

$$\int_{-1}^1 |R_M(t)|^2 \Phi_W(t) dt \geq c' M. \quad (10)$$

The diagonal coefficient norm is only $\sum_{M < n \leq 2M} 1/n \sim \log 2$, whose naive Gaussian normalization would have size W . If $M/W \rightarrow \infty$, the low-height coherent mass is much larger. This elementary example shows why Fourier-positive off-diagonal mass is not automatically useful high-zero information.

For a general nonnegative-coefficient resonator, even the shorter interval $|t| \leq c/\log L$ gives a lower bound proportional to $|R(1/2)|^2/\log L$. Such mass must be controlled rather than discarded using an unweighted measure estimate.

5.2 A general positivity obstruction to cutting a hole at the origin

Suppose Φ is a real Schwartz function with $\widehat{\Phi} \geq 0$ and is not identically zero. Fourier inversion gives

$$\Phi(0) = \int \widehat{\Phi}(\xi) d\xi > 0, \quad | \Phi(t) | \leq \Phi(0).$$

Thus one cannot have both an everywhere nonnegative Fourier transform and a weight that vanishes at the origin, unless the weight is zero. Multiplying the Gaussian by a cutoff that removes the origin necessarily loses the claimed Fourier-positivity property. Translating a packet introduces the familiar phase $e^{-2\pi i t_0 \xi}$; imposing phases on resonator coefficients likewise gives up the common coefficient phase used for the free-sign argument.

This lemma does not rule out a subtler weight whose origin contribution is paid quantitatively. It rules out the exact combination “remove the origin while preserving all Fourier signs for free.”

6. Verification and scope

`centered_gaussian_pole_checks.py` evaluates the finite pole integral in two independent forms at 80 decimal digits: direct complex evaluations of $V(-h/2 - i\sigma) - V(h/2 - i\sigma)$, and the sine integral in (2). The checks do not evaluate a critical-line Dirichlet series.

Representative output:

x, h, W	Pole correction	Mixed correction when $\overline{c} = i$
0.1, 0.2, 3	+2.746104573214653 ... i	-2.746104573214653 ...
0.001, 0.1, 10	+18.88430170566545 ... i	-18.88430170566545 ...
1, 0.3, 2	+0.05966510707903 ... i	negative

x, h, W	Pole correction	Mixed correction when $\bar{c} = i$
2, 0.2, 3	$-0.35517163933647 \dots i$	positive

For $\eta = 1/4$, the ratio of the exact pole integral at $x = T^{-5/4}$ to (8) is approximately 0.9999196069, 1.0000000386, and 1.0000000000000029 at $T = 100, 10^4, 10^8$. The two formulas agree to at least 70 decimal places in the tested cases. The script also checks the cancellation in (9) at the origin using the local branch and the theta function. These are high-precision diagnostics; the proof of the sign and the asymptotic counterexample is analytic and does not depend on their decimal output.

Completed claims in this note are the exact identity with its retained pole term, its sign on the specified support region, the dense-tail counterexample to a negligible-pole estimate, and the elementary coherence/Fourier-positivity obstructions. A full centered-packet formula for the squared logarithm and a successful positive residual certificate are not claimed.

The next admissible research move is to seek genuine control of the **density-subtracted long mixed moment**, or to design a packet/resonator whose pole and low-height contributions can be quantitatively paid. Replacing the shifted Gaussian by a centered one alone supplies neither.

Completion or pole regularization does not restore the missing sign

One can instead use $\log((s-1)\zeta(s))$, which removes the pole in the half-plane to the right of the critical line under RH. This changes the observable. Its absolutely convergent-side expression contains the additional term $\log(s-1)$, whose transformed contribution compensates for the removed pole cut. Likewise the completed xi-function introduces explicit gamma and polynomial factors. These are legitimate formulations, but their arithmetic kernel is no longer the uncorrected nonnegative prime sum. A proposed completed-function approach must keep that explicit continuous/archimedean contribution in its Gram certificate; regularization is not a license to erase the density subtraction identified above.

Current report 06: yau_flow

Source: `research/reports/yau_flow.md`. SHA-256:

`0d407ae5cb198960bc3c00ff76f6baeb61e5536e63799eaf89a28af92f460542`.

Yau flows, circular heat localization, and a candidate CUE depth theorem

Research round 1, 2026-09-05. Prepared independently from the finite-depth source files and primary literature. The results below are mathematical arguments, not a Lean-checked proof. The CUE theorem is a proposed new consequence whose novelty requires further literature review; an attempted independent Claude review was unavailable because authentication returned HTTP 401. Another agent has checked the full deterministic argument and CUE transfer without finding a fatal issue, but that review is not formal verification. It is not a theorem about the Riemann zeta function.

1. Main finding

The earlier programme's dynamic background condition can be replaced by a simpler **initial-data isolation condition**. For the finite circular flow used in this project, an isolated closest pair admits a Gaussian heat-kernel scaling limit. For Haar CUE, an elementary three-point overcrowding estimate supplies that isolation with high probability. Combining this with the existing CUE extreme-gap theorem gives the following candidate theorem, pending independent adversarial review:

Let U_N be Haar distributed in $U(N)$, let $P_N(z) = \det(zI - U_N) = \sum_{j=0}^N a_j z^j$, and evolve $P_{N,s}(z) = \sum_{j=0}^N a_j e^{sj(N-j)} z^j$. Let D_N be its first positive time with a repeated root, and let δ_N be the smallest initial circular angular gap. Then

$$\frac{8D_N}{\delta_N^2} \rightarrow 1 \quad \text{in probability.}$$

Consequently, for every $t \geq 0$,

$$\lim_{N \rightarrow \infty} \Pr(N^{8/3} D_N > t) = \exp\left(-\frac{2\sqrt{2}}{9\pi} t^{3/2}\right).$$

This closes the finite CUE localization gap if the argument below passes further independent review. No universality assertion about arbitrary point processes or zeta zeros is included. The stochastic Yau machinery is not needed for this proof: the decisive simplification is to use the scalar heat equation directly.

2. Three flows that must not be conflated

Dyson Brownian motion is a stochastic eigenvalue evolution. In a conventional Hermitian normalization it has a Brownian term of order $N^{-1/2}$, a repulsive logarithmic interaction, and an optional confining drift. Erdős – Schlein – Yau local relaxation compares probability laws with a local equilibrium law. Their reverse heat-flow argument approximately inverts smoothing on a sufficiently regular matrix-entry density; it does not run individual DBM trajectories backward. See Erdős – Péché – Ramírez – Schlein – Yau and Erdős – Schlein – Yau – Yin.

The project's circular root evolution is deterministic and attractive:

$$\dot{\theta}_j = - \sum_{k \equiv j} \cot \frac{\theta_j - \theta_k}{2}.$$

After multiplying the centered characteristic polynomial by a fixed phase to make it real, write

$$Q_0(x) = \prod_{j=1}^N \sin \frac{x - \theta_j}{2}.$$

The corresponding centered polynomial at time S is, up to its zero-irrelevant scalar factor,

$$\tilde{Q}_S(x) = e^{SN^2/4} e^{S\partial_x^2} Q_0(x).$$

Indeed the Fourier mode $j - N/2$ receives multiplier $e^{[N^2/4 - (j - N/2)^2]S} = e^{Sj(N-j)}$. Thus $Q_S = e^{S\partial_x^2} Q_0$ solves ordinary forward scalar heat flow. For odd N , Q_0 is antiperiodic over 2π ; it is periodic over 4π , and its real-line Gaussian representation remains valid. No periodicity shortcut is needed.

The de Bruijn – Newman convention has $\partial_t H_t = -\partial_x^2 H_t$. Increasing its parameter is the repulsive direction; our collision-time parameter S is the opposite direction. A ζ window cannot simply be replaced by a polynomial: the infinite product, outer zeros, normalization, and the true H_t evolution all need control.

Rodgers and Tao already explicitly invoke the analogy with the Erdős – Schlein – Yau local relaxation method. Their theorem assumes a fixed negative global Newman parameter for contradiction, evolves over an actual interval of real-rooted H_t , and drives the zeros toward local clock behavior. Their argument does not identify ζ with DBM and does not immediately turn a hard core at time zero into an interval of backward stability. See [Rodgers – Tao](#) and [Tao's explanation](#).

3. A deterministic initial–data lemma

Consider a sequence of finite configurations, with size allowed to vary. Their circular angular gaps are initially positive. Let δ be the smallest gap, rotate its endpoints to $-\delta/2, \delta/2$, and write the other angles as $\theta_k \in (-\pi, \pi]$. Define

$$H(y) = \prod_{k \in \{-, +\}} \sin \frac{y - \theta_k}{2}, \quad A = \sum_{k \in \{-, +\}} \frac{1}{2 \left| \sin(\theta_k/2) \right|}.$$

Empty products and sums are allowed. Assume

$$\delta \rightarrow 0, \quad \delta A \rightarrow 0. \tag{I}$$

Isolation lemma. Under (I), the global first collision time satisfies

$$D = \frac{\delta^2}{8} (1 + o(1)). \tag{L}$$

The assumption is on the initial configuration only. It neither assumes that the closest pair remains the first colliding pair nor requires an a priori bound on the evolving background. If a different pair collides first, the upper-bound argument remains valid.

3.1 Global control of the background product

Since sine is 1-Lipschitz,

$$\left| \frac{\sin((y - \theta_k)/2)}{\sin(-\theta_k/2)} - 1 \right| \leq \frac{|y|}{2 \left| \sin(\theta_k/2) \right|}.$$

The elementary product inequality $|\prod(1 + u_k) - 1| \leq \prod(1 + |u_k|) - 1$ implies, for every real y ,

$$\left| \frac{H(y)}{H(0)} - 1 \right| \leq e^{A|y|} - 1, \quad \left| \frac{H(y)}{H(0)} \right| \leq e^{A|y|}. \quad (\text{B})$$

These are global bounds; no cutoff in the heat convolution and no uncontrolled ratio to the global supremum of Q_0 appear.

3.2 Heat-kernel normal form

Set $s = \tau\delta^2$, and let Z be a standard real normal random variable. Because Q_0 is a finite real trigonometric polynomial,

$$Q_s(\delta x) = E Q_0(\delta x + \sqrt{2s}Z).$$

With $Y = x + \sqrt{2\tau}Z$, define

$$F_\delta(x, \tau) = \frac{4Q_{\tau\delta^2}(\delta x)}{\delta^2 H(0)} = E \left[\frac{4}{\delta^2} \sin \frac{\delta(Y - 1/2)}{2} \sin \frac{\delta(Y + 1/2)}{2} \frac{H(\delta Y)}{H(0)} \right]. \quad (\text{F})$$

For fixed Y , the bracket tends to $Y^2 - 1/4$. Moreover,

$$\frac{4}{\delta^2} \left| \sin \frac{\delta(Y - 1/2)}{2} \sin \frac{\delta(Y + 1/2)}{2} \right| \leq (|Y| + 1/2)^2.$$

For $\delta A \leq 1$, (B) gives a dominating factor $e^{|Y|}$, which is integrable against a Gaussian uniformly when x and τ range over compact sets. The same estimates with Taylor's remainder give the stronger uniform statement

$$\sup_{|x| \leq R, 0 \leq \tau \leq T} |F_\delta(x, \tau) - (x^2 - 1/4 + 2\tau)| \leq C_{R,T}(\delta A + \delta^2). \quad (\text{U})$$

For completeness, use $|e^{a|Y|} - 1| \leq a|Y|e^{|Y|}$ for $0 \leq a \leq 1$, and $|\sin u - u| \leq |u|^3/6$. The resulting errors are bounded by a constant times $(\delta A + \delta^2)(1 + |Y|^6)e^{|Y|}$, whose expectations are uniformly bounded on the stated compact set.

3.3 Why the normal form forces an actual collision

Fix $R = 2$, and fix any $\varepsilon > 0$ small enough that $\tau_+ = 1/8 + \varepsilon < 1$. Uniform convergence (U) gives:

1. $F_\delta(-R, \tau) > 0$ and $F_\delta(R, \tau) > 0$ for every $0 \leq \tau \leq \tau_+$, once (I) is sufficiently small.
2. $F_\delta(x, \tau_+) > 0$ for every $-R \leq x \leq R$, since the limiting polynomial is $x^2 + 2\varepsilon$.
3. Initially exactly the two selected roots lie between $-R\delta$ and $R\delta$. To see this, any third root at circular distance at most $R\delta$ would give $\delta A \geq C_R > 0$, contradicting (I).

Suppose, for contradiction, that no global collision occurs by $s = \tau_+\delta^2$. The initial circle roots then remain simple and on the circle throughout: self-inversive symmetry, or equivalently the real trigonometric equation and the implicit function theorem, prevents a simple real root from leaving the real axis. The two selected roots vary continuously, and cannot cross the fixed interval's boundary because the boundary values stay nonzero throughout. They would therefore still give two real roots inside the interval at the final time, contradicting item 2.

It follows that $D \leq (1/8 + \varepsilon)\delta^2$. This argument uses tracking and nonvanishing boundaries; positivity only at the final time would not suffice. No Rouché assertion is being smuggled into the proof.

3.4 Matching deterministic lower bound

Before the first collision, each adjacent gap g obeys

$$g' \geq -2 \cot(g/2).$$

The scalar comparison solution initialized at the minimum gap $\delta < \pi$ first reaches zero at

$$D_2(\delta) = -\log \cos(\delta/2) \geq \delta^2/8.$$

Compare every gap with this same scalar solution. The comparison function $-2 \cot(g/2)$ is increasing on $(0, 2\pi)$, so this remains legitimate for an initially long gap. A collision cannot occur while the scalar solution is positive. Therefore $D \geq D_2(\delta)$. Combining with the preceding upper bound proves (L).

One can extract an absolute-constant bound

$$\frac{\delta^2}{8} \leq D \leq \delta^2 \left[\frac{1}{8} + C(\delta A + \delta^2) \right]$$

when $\delta A + \delta^2$ is sufficiently small. No numerical value of C is claimed here.

4. CUE supplies the isolation condition

The Haar CUE eigenangles form the projection determinantal process with correlation kernel

$$K_N(x, y) = \frac{1}{2\pi} \sum_{m=0}^{N-1} e^{im(x-y)}.$$

4.1 Uniform three-point bound

For three real lifts x_1, x_2, x_3 , Cauchy – Binet gives

$$\rho_3(x_1, x_2, x_3) = \frac{1}{(2\pi)^3} \sum_{0 \leq m_1 < m_2 < m_3 < N} \left| \det(e^{im_a x_b})_{a,b=1}^3 \right|^2.$$

For a fixed triple of frequencies, put $v(x) = (e^{im_1 x}, e^{im_2 x}, e^{im_3 x})^T$. Divided differences give the exact identity

$$\det(v(x_1), v(x_2), v(x_3)) = \prod_{i < j} (x_j - x_i) \det(v[x_1], v[x_1, x_2], v[x_1, x_2, x_3]).$$

The integral formula for divided differences and Hadamard's inequality bound the three column norms by $\sqrt{3}$, $\sqrt{3}N$, $\sqrt{3}N^2/2$, respectively. Hence each determinant has modulus at most $(3\sqrt{3}/2)N^3 \prod_{i < j} |x_i - x_j|$. Summing fewer than $N^3/6$ terms proves

$$\rho_3(x_1, x_2, x_3) \leq CN^9 \prod_{i < j} |x_i - x_j|^2 \quad (3P)$$

with an absolute constant. This holds for lifts of points in a short arc, including an arc that crosses the arbitrary 2π cut.

4.2 Triple-free arcs with high probability

Count ordered triples with an anchor at $x_1 \in [0, 2\pi)$ and the other two points in the oriented arc $[x_1, x_1 + r]$. If any arc of length r contains three points, at least one such anchored triple exists. Using (3P), integration over the two offsets gives

$$\Pr(\exists \text{ arc of length } r \text{ containing three points}) \leq CN^9 r^8. \quad (\text{T})$$

Choose $r = N^{-7/6}$. The right side is $O(N^{-1/3})$.

4.3 A packing bound for the reciprocal chord sum

On the event that every arc of length r contains at most two points, suppose additionally that the minimum gap obeys $\delta < r/2$. Center that pair at zero. Any other point must have circular distance at least $r - \delta/2$ from the pair center: otherwise it and both endpoints lie in a single arc of length r .

Partition the circle in arcs of length r , or order the remaining points by their circular distance $d_k \in (0, \pi]$ from zero. The two-points-per-arc condition implies $d_k \geq cr(k+1)$ for an absolute $c > 0$; the first few indices use the preceding isolation observation. Since $2 \sin(d/2) \geq 2d/\pi$ for $0 \leq d \leq \pi$,

$$A \leq Cr^{-1}(1 + \log N). \quad (\text{A})$$

4.4 Existing extreme-gap input and the resulting law

The CUE extreme-gap theorem gives

$$N^{4/3} \delta_N \Rightarrow X, \quad \Pr(X > x) = \exp\left(-\frac{x^3}{72\pi}\right). \quad (\text{EG})$$

This normalization uses angular coordinates of circumference 2π . It agrees with $A_2 = 1/(24\pi)$ in Feng – Wei's gap intensity $A_2 u^2 du$: integrating from zero to x gives $x^3/(72\pi)$. See [Ben Arous – Bourgade](#) and [Feng – Wei, Theorem 1.1](#) and [Corollary 1.1](#).

In particular $N^{4/3} \delta_N$ is tight. Combining (T), (A), and $r = N^{-7/6}$ gives

$$\delta_N A_N = O_{\Pr}(N^{-1/6} \log N) = o_{\Pr}(1),$$

and also $\delta_N/r \rightarrow 0$ in probability. The isolation lemma yields $8D_N/\delta_N^2 \rightarrow 1$ in probability. Slutsky's theorem then gives

$$N^{8/3} D_N \Rightarrow X^2/8, \quad \Pr(X^2/8 > t) = \exp\left(-\frac{2\sqrt{2}}{9\pi} t^{3/2}\right).$$

This proof needs no uniform-in- N extrapolation of the extreme-gap tail. It uses only tightness and the fixed-argument limiting law.

5. What is genuinely new here, and what may already be classical

The isolated-pair mechanism is closely related to the classical Lehmer-pair estimates of Csordas – Smith – Varga, which already control a Newman lower bound from an initial gap and a background sum of inverse squares. The claim “nobody knows how to control the evolving background from initial data” would therefore be misleading. See [Csordas – Smith – Varga's paper](#).

Andrade – Chang – Miller explicitly express the function-field background as finite csc-squared sums and use isolated small zeros to approach zero Newman parameter. Their § 3.8 already connects such estimates to random matrix theory. See [Newman's conjecture in various settings](#). We must compare any eventual paper carefully with these results.

The contribution proposed in this round is the particularly short global heat-product estimate (B), its deterministic scaling lemma, and the use of elementary CUE triple-overcrowding plus the known extreme-gap theorem to get an explicit first-collision law. A search in this round did not locate that exact CUE law; this is not a guarantee of novelty. Calling it a solution of a famous ζ conjecture would be false.

6. Why this does not refute AH

The isolation assumption fails at lattice scale. For a clock-like configuration with gap of order N^{-1} , the reciprocal chord sum is of order $N \log N$, so δA is not small. ACUE's finite lattice hard core therefore does not contradict the lemma. The new theorem explains a distinction for CUE's extreme close pairs; it does not show that ζ has those pairs.

For ζ , near collisions, quantitative outer-zero control, and the true heat evolution would need to be established arithmetically. Bandwidth-one correlations alone cannot supply a theorem excluding every ACUE-type model, because such models already match that information. A proposed “energy bound” must be stated and checked on ACUE; merely naming a renormalized energy does not add arithmetic information.

Moreover, an AH that permits a zero-density exceptional set is not refuted by a first-collision statistic or infinitely many unusually small gaps. For that version, one needs positive-density violations in a fixed interval separated from the allowed half-integer spacings, or another statistic that controls the density of exceptions. The inference “AH implies every window has hard core $1/2$ ” remains invalid without strengthening AH.

7. Function–field transfer: a precise modest theorem and a real remaining gap

For a fixed degree, the extended-valued collision depth is expected to be continuous on the compact space of circle-rooted characteristic polynomials when $[0, \infty]$ is given its usual one-point compactification. A route to a proof is as follows:

- If $t < D(P)$, Q_t has the full number of simple real zeros, and nearby coefficient vectors also do. This gives lower semicontinuity.
- If $t > D(P)$, choose a nearby time with no multiple zero. Scalar heat's strict drop in zero number after a multiple zero gives fewer than the full number of real zeros at this time. This property persists under small coefficient perturbation, giving upper semicontinuity.
- At a clock polynomial $D = \infty$, the first assertion for every finite t gives extended continuity. The same zero-number argument treats initial multiple roots if their depth is defined as zero.

This uses the classical zero-number theorem for scalar parabolic equations, rather than an unsupported assertion that the hitting time of an arbitrary discriminant is continuous. The theorem is due to Angenent; see [the original publication](#).

Once this continuity argument is fully written, fixed-rank Frobenius equidistribution immediately transports the law of any bounded continuous function of depth. One must state a specific arithmetic family and its compact symmetry group. For instance, the hyperelliptic ensemble at fixed genus and $q \rightarrow \infty$ has $\mathrm{USp}(2g)$ equidistribution; [Rudnick's paper](#) states this precisely.

This fixed-rank transfer is a useful corollary, not by itself a large universality theorem. To transport the CUE $N^{8/3}$ law to a unitary function-field family one either takes the successive limit $q \rightarrow \infty$ then $N \rightarrow \infty$, or proves effective equidistribution uniform in the rank and in the increasingly singular test function. A diagonal subsequence can be chosen abstractly from successive limits, but it gives no useful growth rate $q = q(N)$. Haar USp and SO have symmetry-forced spectral pairing and edges, so the CUE argument cannot simply be relabeled with another beta.

8. A more promising research ladder after this round

1. Independently audit the deterministic heat lemma and the CUE overcrowding proof. Write a concise standalone theorem note, then check the CSV/Stopple/Andrade literature for priority and potentially sharper initial-data criteria.
2. Generalize only under explicit assumptions: if $a_N \delta_N \Rightarrow X$ and $\delta_N A_N \rightarrow 0$ in probability, then $a_N^2 D_N \Rightarrow X^2/8$. Bandwidth-one correlation is unnecessary for this abstract transfer theorem.
3. For circular beta ensembles, prove or cite a uniform three-point bound of the form $\rho_3 \leq C_\beta N^{3+3\beta} \prod_{i < j} |x_i - x_j|^\beta$. It would give triple-free scale $r = N^{-1-\eta}$ for $\eta > 1/(2 + 3\beta)$, compatible with the minimum-gap scale $N^{-1-1/(\beta+1)}$. The known Feng – Wei extreme-gap theorem used in the source covers positive integer beta; do not silently assert all positive real beta.
4. For a genuinely arithmetic advance, seek a **positive-density isolated-pair theorem** with a testable arithmetic input. If normalized gaps below 1/2 are the target, obtaining them cannot be assumed inside the background lemma.
5. Investigate whether a natural arithmetic weighting produces an energy or multi-point statistic unavailable to the ACUE mimickers. That is where a Yau-inspired local relaxation argument might add value: formulate the actual local law, the coercive energy, and the estimate that distinguishes arithmetic zeros from a fake process.

9. Verification boundary

The complete reasoning of the deterministic lemma and CUE transfer is above. Independent review by another agent was requested. Numerical tests are supplementary and cannot certify the theorem. The attached script tests the scalar heat normal form, exact two-root normalization, and planted isolated-pair collision times; it does not simulate an arithmetic zeta flow or prove a famous conjecture.

The completed diagnostics are saved in `yau_flow_checks.json`. They reproduce the exact two-root time for four gaps; for planted isolated pairs at $N=16,32,64,128,256$ they give $8D/\delta^2$ between 1.00013025 and 1.00002066, decreasing toward 1. The tests also sample the global product bound and three-point determinant inequality. These data are not a Monte Carlo verification of the limiting distribution. The `residual_gram` agent reviewed the full deterministic and CUE arguments and reported no fatal issue; the `prime186` agent also independently checked the deterministic argument, the actual Feng – Wei source statements, the beta scope, and the tail constants without finding a fatal issue. The attempted Claude review did not run because authentication failed with HTTP 401.

10. Stronger supplement: the positive-integer circular-beta depth law

After the first draft, the general three-point estimate was found to follow directly from Hölder's inequality and a published coalesced-charge partition-function bound. Thus the proposed result extends to every fixed **positive integer** beta, including COE, CUE, and CSE. The restriction here comes from the published extreme-gap limit used as an input; this note does not assert that the limit is unavailable or false for other beta values.

For fixed positive integer β , let the initial angles have density

$$\frac{1}{C_{\beta,N}} \prod_{i < j} |e^{i\theta_i} - e^{i\theta_j}|^\beta.$$

Use the same characteristic-polynomial coefficient flow as above; the flow itself does not depend on beta. Define

$$A_\beta = \frac{1}{2\pi} \left(\frac{\beta}{2}\right)^\beta \frac{\Gamma(1 + \beta/2)^3}{\Gamma(1 + 3\beta/2)\Gamma(1 + \beta)}.$$

Candidate circular-beta theorem. For every fixed positive integer beta,

$$\frac{8D_N}{\delta_N^2} \rightarrow 1 \quad \text{in probability,}$$

and

$$\lim_{N \rightarrow \infty} \Pr(N^{2+2/(\beta+1)} D_N > t) = \exp\left[-\frac{A_\beta}{\beta+1} (8t)^{(\beta+1)/2}\right], \quad t \geq 0. \quad (\text{CB})$$

For $\beta=1$ the tail is $e^{-t/6}$. For $\beta=2$ it is the CUE law in § 1. For $\beta=4$ it is $\exp[-(64\sqrt{2}/(675\pi))t^{5/2}]$. CSE here means the circular $\beta=4$ ensemble of N distinct eigenangles, not Haar $\text{USp}(2N)$, and no double-counting of quaternionic eigenvalues is intended.

10.1 The missing correlation estimate is an elementary Hölder consequence

Put $z_a = e^{ix_a}$, $a = 1, 2, 3$, and $w_j = e^{iy_j}$, $j = 1, \dots, N-3$. The three-point correlation function is

$$\rho_3(x_1, x_2, x_3) = \frac{(N)_3}{C_{\beta,N}} \prod_{a < b} |z_a - z_b|^\beta I(z_1, z_2, z_3),$$

where

$$I = \int \prod_{k < j} |w_i - w_j|^\beta \prod_{j=1}^{N-3} \prod_{a=1}^3 |w_j - z_a|^\beta d\mathbf{y}.$$

Apply Hölder with exponents $(3,3,3)$ to the measure with density $\prod_{k < j} |w_i - w_j|^\beta$. This gives

$$I \leq \prod_{a=1}^3 \left[\int \prod_{k < j} |w_i - w_j|^\beta \prod_j |w_j - z_a|^{3\beta} d\mathbf{y} \right]^{1/3}.$$

Each bracket is independent of z_a by rotation invariance. Integrating its distinguished point over the full circle identifies it as $C_{\beta, N-3, (3)}/(2\pi)$, in the notation of Feng – Wei. Their Lemma 1.1 gives, for $\beta \geq 1$,

$$\frac{C_{\beta, N-3, (3)}}{C_{\beta, N}} \leq (N\beta)^{3\beta}.$$

Therefore

$$\rho_3(x_1, x_2, x_3) \leq \frac{\beta^{3\beta}}{2\pi} N^{3+3\beta} \prod_{a < b} |e^{ix_a} - e^{ix_b}|^\beta \leq C_\beta N^{3+3\beta} \prod_{a < b} |x_a - x_b|^\beta \quad (\text{B3})$$

for real lifts of a short arc. This argument does not use a determinantal or Pfaffian representation. The only non-elementary input in (B3) is the published partition-function inequality. See Feng – Wei, equations (1) – (4) and Lemma 1.1.

10.2 Triple exclusion and first collision

The same anchored-triple count now gives

$$\Pr(\exists \text{ arc of length } r \text{ containing three points}) \leq C_\beta N^{3+3\beta} r^{2+3\beta}.$$

Choose any fixed

$$\frac{1}{2+3\beta} < \eta < \frac{1}{\beta+1}, \quad r = N^{-1-\eta}.$$

This interval is nonempty for every positive beta. The probability above tends to zero. On its complement, § 4.3 gives $A \leq Cr^{-1}(1 + \log N)$.

Feng – Wei's extreme-gap theorem for positive integer beta gives

$$N^{1+1/(\beta+1)} \delta_N \Rightarrow X_\beta, \quad \Pr(X_\beta > x) = \exp[-A_\beta x^{\beta+1}/(\beta+1)].$$

It follows that

$$\delta_N A_N = O_{\Pr}(N^{\eta-1/(\beta+1)} \log N) = o_{\Pr}(1).$$

The deterministic isolation lemma proves the ratio limit, and continuous mapping gives (CB). There is no beta-infinity limit in this argument. Fixed-beta asymptotics do not identify ACUE with a beta-infinity ensemble.

10.3 Assessment

If the isolation lemma survives the independent adversarial review, the old finite-depth programme is no longer “one dynamic background lemma away” in the COE/CUE/CSE cases: the required localization follows from a different, initial-data argument and existing partition-function estimates. This is a concrete candidate probability theorem. The arithmetic bridge to ζ and the user's requested famous conjecture remain open; the result is best treated as a proved-direction candidate and a reusable transfer mechanism, not the final research objective.

A stronger, drift-invariant version of the deterministic lemma is proved in `yau_flow_galilean_refinement.md`: the hypothesis is $\delta^2 B \rightarrow 0$ with $B = \frac{1}{4} \sum \csc^2(\theta_k/2)$. The exact Galilean transform removes the linear logarithmic drift, and a global quadratic exponential bound controls the heat tails. The root agent independently audited the centered-factor inequality and the moving-frame mechanism. Full formal verification and a priority review remain outstanding.

11. Identification of the first colliding pair

The published joint limit for the first two smallest gaps gives a useful corollary to the candidate depth theorem.

Corollary. For every fixed positive integer beta, with probability tending to one, the first colliding pair is the pair that had the smallest gap at time zero.

Write the first two initial gaps as $\delta_1 < \delta_2$; equality has probability zero under the continuous circular-beta density. Before the global first collision, every gap other than the minimum is bounded below by the scalar comparison solution initialized at δ_2 . Therefore no such pair can collide before

$$-\log \cos(\delta_2/2) \geq \delta_2^2/8.$$

The selected minimum pair supplies the global upper bound

$$D \leq \delta_1^2(1/8 + \varepsilon_N), \quad \varepsilon_N \rightarrow 0 \quad \text{in probability.}$$

The joint extreme-gap limit implies that δ_2/δ_1 converges in law to a random variable strictly greater than one almost surely. Hence

$$\Pr(\delta_2^2/\delta_1^2 > 1 + 8\varepsilon_N) \rightarrow 1.$$

On this event, all other initial adjacent pairs remain separated until after the global first collision, proving the assertion. This is an asymptotic identification statement; it does not make the finite marked derivative exactly rank two.

12. The relaxation–time obstruction to the proposed AH argument

This subsection is a dimensional analysis and logical audit, not an impossibility theorem against every future heat-flow method.

At height T , the mean Riemann-zero spacing in the usual ordinate is of order $(\log T)^{-1}$. A pair-collapse equation with leading term $g' = -4/g$ therefore has local time scale $(\log T)^{-2}$. The common convention for H_0 uses twice the ordinary zeta ordinate, which changes constants but not this exponent. One must keep that coordinate factor when making a numerical claim such as $\pi^2/8$.

The contradiction assumption in Rodgers – Tao is a **fixed negative global** Newman parameter, giving a time interval of length bounded away from zero on which all the relevant zeros are real. In microscopic units, the available duration at height T is consequently of order $\log^2 T$, tending to infinity. This long duration in local units is what permits a relaxation-to-clock argument. Their actual proof needs renormalized energies and technical estimates; this dimensional observation is not a substitute for those estimates. See [Rodgers – Tao](#).

A hard core proportional to one mean spacing, even if it were a genuine uniform hard core rather than an almost-everywhere AH, only gives the two-body stability lower bound of order $(\log T)^{-2}$. That is $O(1)$ microscopic relaxation times. It does not supply the diverging number of relaxation times available in the global negative- Λ contradiction. Rescaling the old energy argument does not remove this loss. A proposed proof must explain what new arithmetic estimate replaces that missing duration.

The finite circular normalization makes the distinction transparent. Its unnormalized deterministic clock Jacobian has eigenvalues $k(N - k)$; the high modes have time scale N^{-2} . A conventional circular DBM has repulsive drift

$$\frac{1}{2N} \sum_{j \neq i} \cot \frac{\theta_i - \theta_j}{2}$$

and diffusion coefficient $\sqrt{2/(\beta N)}$. Thus its local relaxation scale is order N^{-1} , corresponding to unnormalized deterministic time order N^{-2} . At that time the Brownian displacement itself is of order one mean spacing. It cannot be discarded in a pathwise identification with the attractive deterministic flow.

At the CUE extreme-gap scale the same problem persists: deterministic first-collision time is order $N^{-8/3}$, corresponding to DBM time order $N^{-5/3}$; the DBM Brownian displacement over that duration is order $N^{-4/3}$, exactly the extreme gap size. A genuine DBM comparison has to retain its noise and its repulsive sign. It cannot identify its collision behavior with the scalar heat evolution studied here.

The new isolated-pair theorem sharpens this audit. Under weak initial isolation, normalized depth is asymptotically the square of the normalized smallest gap. A ζ depth bound substantially below the half-spacing threshold would therefore require genuine microscopic arithmetic information or a mechanism outside isolated-pair behavior. Merely renaming the desired small-gap phenomenon as a first-passage event does not prove it. For an AH allowing zero-density exceptions, the needed arithmetic result must additionally control the density of violations.

13. Current literature on matrix heat flow

A relevant recent paper is Hall – Ho, [The heat flow conjecture for polynomials and random matrices](#), published in 2025. It connects additive and multiplicative polynomial heat flows to deformations of random matrix models, proves second-moment and holomorphic-moment statements, and formulates macroscopic root-transport conjectures. Its multiplicative

generator differs from the generator here by a scalar multiplier and a deterministic radial dilation. This is useful context for a future theorem note; its macroscopic statements do not by themselves supply the microscopic extreme first-collision law above. The exact priority of that law remains to be checked.

Current report 07:

yau_flow_galilean_refinement

Source: `research/reports/yau_flow_galilean_refinement.md`. SHA-256:
`ffa09e0a4a6a78854fe5bfdb666e812d85bb34169dcb1d8443020eb665db6d3e`.

A stronger isolated-pair lemma after removing the background drift

2026-09-05. This is a proposed refinement of `yau_flow.md`, independently derived in the same research round. It has not yet received the other agents' full review. It overlaps conceptually with classical Lehmer-pair criteria and should not be advertised as a new principle without a priority check.

The reciprocal-chord condition $\delta A \rightarrow 0$ in the first draft can be replaced by the weaker inverse-square condition $\delta^2 B \rightarrow 0$. The improvement removes the effect of a large common background velocity by an exact Galilean transformation of scalar heat flow.

Statement

Let δ be the smallest angular gap in a circle-rooted degree- N polynomial. Center its endpoints at $\pm \delta/2$ and write

$$Q_0(x) = p_\delta(x)H(x), \quad p_\delta(x) = \sin \frac{x - \delta/2}{2} \sin \frac{x + \delta/2}{2}, \quad H(x) = \prod_{k \in \pm} \sin \frac{x - \theta_k}{2}.$$

Define

$$L = \frac{H'(0)}{H(0)} = -\frac{1}{2} \sum_{k \in \pm} \cot \frac{\theta_k}{2}, \quad B = \frac{1}{4} \sum_{k \in \pm} \csc^2 \frac{\theta_k}{2}.$$

If $\delta \rightarrow 0$ and $\delta^2 B \rightarrow 0$ along a sequence of configurations, then the global first collision time of the coefficient flow $\exp(\text{sj}(N-))$ satisfies

$$D = \frac{\delta^2}{8} (1 + o(1)).$$

In fact, there are absolute constants C and $\varepsilon_0 > 0$ such that, when $\delta^2(B+1) \leq \varepsilon_0$,

$$\frac{\delta^2}{8} \leq D \leq \delta^2 \left[\frac{1}{8} + C\delta^2(B+1) \right].$$

No numerical value of C or ε_0 is claimed. The lower bound is the existing deterministic two-body comparison. The following argument gives the upper bound.

1. A global centered sine-factor estimate

For all real c and v ,

$$|\cos v - C \sin v| e^{cv} \leq e^{4(1+c^2)v^2}. \quad (\text{S})$$

Put $w = |v| \sqrt{1+c^2}$. If $w \geq 1/2$, use

$$\log |\cos v - c \sin v| + cv \leq \log(1 + |c| |v|) + |c| |v| \leq 2w \leq 4w^2.$$

The inequality is automatic if the factor vanishes. If $w \leq 1/2$, the function $r(u) = \cos u - c \sin u$ is positive between 0 and v : throughout that segment,

$$r(u) \geq 1 - u^2/2 - |c| |u| \geq 1 - 1/8 - 1/2 > 0.$$

Its centered logarithm $f(u) = \log r(u) + cu$ has $f(0) = f'(0) = 0$ and

$$f''(u) = -\frac{1+c^2}{r(u)^2} < 0.$$

Thus $f(v) \leq 0$, proving (S) also in this case.

Now put

$$R(y) = e^{-Ly} \frac{H(y)}{H(0)}.$$

Each sine ratio is $\cos(y/2) - \cot(\theta k/2) \sin(y/2)$. Applying (S) factor by factor gives the **global real-line bound**

$$|R(y)| \leq e^{4By^2}, \quad y \in \mathbb{R}. \quad (\text{G})$$

When $By^2 \leq 1/16$, every uncentered sine ratio remains at least $3/4$ along the segment from 0 to y . Consequently H has constant sign on that segment and

$$-\frac{16}{9}B \leq (\log |H|)'' \leq 0.$$

After subtracting the linear term Ly and integrating twice,

$$-By^2 \leq \log R(y) \leq 0, \quad |R(y) - 1| \leq By^2.$$

When $By^2 > 1/16$, (G) gives $|R(y) - 1| \leq e^{\{4By^2\} + 1} \leq 32By^2 e^{\{4By^2\}}$. Combining both regions,

$$|R(y) - 1| \leq 32By^2 e^{4By^2}, \quad y \in \mathbb{R}. \quad (\text{E})$$

The case $B=0$ is immediate: H is constant.

2. Exact heat conjugation and the moving frame

For any real L for which the Gaussian integrals converge,

$$e^{\mathcal{S}_x^2} (e^{Lx} f(x)) = e^{Lx + L^2 s} (e^{\mathcal{S}_x^2} f)(x + 2Ls). \quad (\text{C})$$

This is a direct completion of the square in the Gaussian convolution. Here the original trigonometric polynomial is bounded on the real axis; the transformed integrals converge at the times used below because $\delta^2 B$ is small.

Set $s = \tau \delta^2$ and use the moving coordinate $x = \delta \xi - 2Ls$. With Z standard normal and $Y = \xi + \sqrt{(2\tau)}Z$, formula (C) yields

$$\frac{4e^{-L\delta\xi + L^2 s}}{\delta^2 H(0)} Q_s(\delta\xi - 2Ls) = E\left[\frac{4}{\delta^2} p_\delta(\delta Y) R(\delta Y)\right]. \quad (\text{M})$$

The multiplier on the left is real and nonzero, so it does not alter the roots or their signs relative to the constant $H(0)$.

The right side converges uniformly on compact ξ, τ sets to

$$E(Y^2 - 1/4) = \xi^2 - 1/4 + 2\tau.$$

More precisely, (E), the elementary sine Taylor remainder, and Gaussian integrability give

$$\sup_{|\xi| \leq R_0, 0 \leq \tau \leq T} \left| \frac{4e^{-L\delta\xi + L^2s}}{\delta^2 H(0)} Q_s(\delta\xi - 2Ls) - (\xi^2 - 1/4 + 2\tau) \right| \leq C_{R_0, T} \delta^2 (B+1), \quad (U)$$

once $\delta^2 B$ is sufficiently small depending on T . Indeed the error is bounded by a constant times $\delta^2 (B+1)(1 + |Y|^6)e^{4B\delta^2 Y^2}$, and its expectation is uniformly bounded after choosing $\delta^2 B$ small enough. This is why a quadratic global bound was necessary; a merely local Taylor expansion would not justify the heat integral.

3. Collision tracking with moving boundaries

Use the interval

$$[-2\delta - 2Ls, 2\delta - 2Ls]$$

on the real cover of the circle. At time zero it contains exactly the selected pair, since $\delta^2 B \rightarrow 0$ excludes any third point at distance $O(\delta)$. The normal form (U) shows that both boundary values stay nonzero for all times up to $(1/8 + \epsilon)\delta^2$, and the whole interval contains no real root at that final time.

If no global multiple root had occurred, the initially simple circle roots would remain on the circle and vary continuously. The leading and constant polynomial coefficients are fixed and nonzero, so no root can escape via zero or infinity. Simple self-inversive roots cannot leave the circle without colliding with their reciprocal-conjugate partners. The two roots in the moving interval therefore cannot disappear without a multiple root, and cannot leave through its nonvanishing moving boundaries. This contradiction gives $D \leq (1/8 + \epsilon)\delta^2$.

Taking ϵ to be a sufficiently large constant times $\delta^2 (B+1)$ gives the quantitative upper bound. Pair identity after the first global collision is irrelevant: an earlier collision already establishes the desired upper bound.

4. Consequences and scope

On a triple-free event at length r , the packing argument now gives $B \leq Cr^{-2}$, with no harmonic logarithm. Therefore the circular-beta proof in `yau_flow.md` works with $\delta^2 B = O_p((\delta/r)^2)$, which tends to zero at the same admissible r scales. The limiting laws and constants do not change.

The transformation also explains why a large common drift need not obstruct collision localization. The right structural quantity is curvature of the logarithm of the background product, not the absolute sum of its first derivatives. This connects naturally with classical Lehmer-pair inverse-square criteria. For an infinite canonical product one must separately justify convergence, the exponential normalization, and its true heat evolution; this note does not claim that those steps, or the arithmetic small-gap input, have been completed for ζ .

The stronger criterion is strictly weaker. Take $\delta = N^{-2}$, put the selected pair at $\pm \delta/2$, and place the remaining $N-2$ points at angles $(M+k)\delta$, $k=0, \dots, N-3$, with $M = \lceil \sqrt{\log N} \rceil$. Every gap is at least δ . Direct harmonic-sum estimates give $\delta^2 B$ asymptotic to $1/M$, tending to zero, while δA is of order $\log(N/M)$, tending to infinity. The Galilean argument therefore applies even though the first proof's absolute first-derivative bound fails. The moving interval may travel across many *initial* gaps; its all-time nonvanishing boundaries still make the tracking argument valid.

The root agent independently checked the centered-factor bound and the moving-frame mechanism. Numerical checks in `yau_flow_checks.py` include cases with δA between 1.09 and 5.31 but $\delta^2 B$ near 0.02. The selected local double-root times are close to $\delta^2/8$. Those configurations have many equal-size gaps, so the script deliberately does not label their computed local double-root time as the first global collision.

Current report 08: galilean-proof-audit

Source: `research/reports/galilean-proof-audit.md`. SHA-256:

`f3896ff6b7a07d1296762ff09ad8bf9315766459f16394d6636e64d9036c8f13`.

Galilean localization: second-pass proof audit and the exact ζ obligation

2026-09-05. The same heat-flow research agent performed this second-pass audit; the separate prime-lane agent supplied the earlier independent review. This audit rereads the round-1 Galilean argument from its definitions. It does not extend the random-matrix model family or claim a new ζ theorem. No repository files were changed. The argument remains an ordinary mathematical proof, not a Lean or interval-arithmetic certificate.

Audit conclusion. The deterministic asymptotic conclusion survives. The stronger quantitative bound can also be retained: the previously unspecified constants can be chosen uniformly in N . A deliberately coarse explicit choice is

$$K = 16384, \quad \eta_0 = \frac{1}{524288},$$

for which the proof below gives

$$\frac{\delta^2}{8} \leq D \leq \delta^2 \left(\frac{1}{8} + K\delta^2(B+1) \right) \quad \text{if } \delta^2(B+1) \leq \eta_0.$$

The constants are not optimized and are not useful numerical record constants. Their purpose is to certify that no dependence on N or on the background drift has been hidden in the big- O statement. The proof needs no complex Gaussian contour shift, Rouché theorem, or unproved continuity of a general hitting-time functional.

1. Exact assumptions and definitions

Let $N \geq 2$ and let P be a degree- N polynomial with N distinct roots on the unit circle. Its leading and constant coefficients are nonzero. Define

$$P_s(z) = \sum_{j=0}^N a_j e^{sj(N-j)} z^j, \quad D = \inf \{s > 0 : \text{disc}(P_s) = 0\},$$

with infimum of the empty set equal to infinity. Both endpoint coefficients remain fixed and nonzero for every real s .

Let δ be the smallest circular angular gap. Rotate its endpoints to $-\delta/2$ and $\delta/2$. Choose all other angles θ_k in $(-\pi, \pi]$, and set

$$Q_0(x) = p_\delta(x)H(x), \quad p_\delta(x) = \sin \frac{x-\delta/2}{2} \sin \frac{x+\delta/2}{2}, \quad H(x) = \prod_{k \in \mathbb{Z}} \sin \frac{x-\theta_k}{2}.$$

Multiplying by a fixed nonzero scalar relates this real trigonometric polynomial to the centered characteristic polynomial. Define

$$L = \frac{H'(0)}{H(0)} = -\frac{1}{2} \sum_{k \in \mathbb{Z}} \cot \frac{\theta_k}{2}, \quad B = \frac{1}{4} \sum_{k \in \mathbb{Z}} \csc^2 \frac{\theta_k}{2}, \quad \eta = \delta^2(B+1).$$

$H(0)$ is nonzero. Empty background products and sums are permitted. The coefficient flow has the same zeros as

$$Q_s = e^{s\phi_x^2} Q_0,$$

because the missing scalar factor is $\exp(sN^2/4)$. For odd N the centered trigonometric polynomial is antiperiodic over 2π ; the real-line Gaussian formula is nevertheless exact term by term on its finitely many Fourier modes.

2. Global centered–factor estimate: checked

For real c, v , let $w = |v| \sqrt{1+c^2}$. We verify

$$|\cos v - c \sin v| e^{cv} \leq e^{4(1+c^2)v^2}. \quad (2.1)$$

If $w \geq 1/2$, use $|\cos v - c \sin v| \leq 1 + |c| |v|$. Taking logarithms, when the factor is nonzero,

$$\log |\cos v - c \sin v| + cv \leq \log(1 + |c| |v|) + |c| |v| \leq 2w \leq 4w^2.$$

If $w \leq 1/2$, $r(u) = \cos u - c \sin u$ is positive along the segment between 0 and v , because

$$r(u) \geq 1 - u^2/2 - |c| |u| \geq 1 - 1/8 - 1/2 > 0.$$

The function $f(u) = \log r(u) + cu$ has $f(0) = f'(0) = 0$ and

$$f''(u) = -\frac{1+c^2}{r(u)^2} < 0.$$

Therefore $f(v) \leq 0$. The zero-factor case is immediate. This proves (2.1) on the full real line, including beyond the first background root.

Define $R(y) = \exp(-Ly)H(y)/H(0)$. Each factor is $\cos(y/2) - \cot(\theta_k/2)\sin(y/2)$, with exactly the matching exponential $\exp(y \cot(\theta_k/2)/2)$. Hence

$$|R(y)| \leq e^{4By^2}, \quad y \in \mathbb{R}. \quad (2.2)$$

There is no estimate involving $|L|$ in (2.2).

3. The global error $|R-1|$: checked with its constants

Put $a_k = 1/(2|\sin(\theta_k/2)|)$, so $B = \sum a_k^2$. The sine Lipschitz inequality gives, for every real u ,

$$\left| \frac{\sin((u - \theta_k)/2)}{\sin(-\theta_k/2)} - 1 \right| \leq a_k |u|.$$

If $By^2 \leq 1/16$, then for every u between 0 and y every ratio on the left is at least $3/4$. H has constant sign there, and

$$-\frac{16}{9}B \leq (\log |H|)'(u) = -\frac{1}{4} \sum_{k \in \mathbb{Z}} \csc^2 \frac{u - \theta_k}{2} \leq 0.$$

Subtract Ly and integrate twice from zero. This is valid for positive or negative y : the integral remainder has nonnegative weights in either direction. It gives

$$-\frac{8}{9}By^2 \leq \log R(y) \leq 0,$$

and in particular $|R(y) - 1| \leq By^2$. If $By^2 > 1/16$, (2.2) gives

$$|R(y) - 1| \leq e^{4By^2} + 1 \leq 32By^2 e^{4By^2}.$$

Thus the proposed global bound is correct:

$$|R(y) - 1| \leq 32By^2 e^{4By^2} \quad (y \in \mathbb{R}). \quad (3.1)$$

When $B=0$, H is constant and the conclusion is exact. The argument does not infer global control by extending a local logarithm through its zeros.

4. Heat conjugation and Gaussian tails: exact and uniform

For $s>0$, completing the square in the real Gaussian integral gives

$$e^{s\partial_x^2}(e^{Lx}f(x)) = e^{Lx+L^2s}(e^{s\partial_x^2}f)(x+2Ls). \quad (4.1)$$

No complex contour shift occurs. In the present finite problem all integrals converge: the original Q_0 is bounded on the real line, and $\exp(-Lx)Q_0(x)$ has at most exponential real growth. The estimates below additionally provide uniform domination as N varies.

Put $s=\tau\delta^2$, $x=\delta\xi-2Ls$, $Y=\xi+\sqrt{(2\tau)}Z$, with Z standard real normal. Then

$$F_\delta(\xi, \tau) := \frac{4e^{-L\delta\xi+L^2s}}{\delta^2 H(0)} Q_s(\delta\xi-2Ls) = E[A_\delta(Y)R(\delta Y)],$$

where

$$A_\delta(Y) = \frac{4}{\delta^2} \sin \frac{\delta(Y-1/2)}{2} \sin \frac{\delta(Y+1/2)}{2}.$$

Let $P(Y)=Y^2-1/4$ and $q=\delta^2B$. Two elementary global estimates are

$$\begin{aligned} |A_\delta(Y)| &\leq Y^2 + 1/4, \\ |A_\delta(Y) - P(Y)| &\leq \frac{\delta^2}{12} |Y^2 - 1/4| (Y^2 + 1/4) \leq \frac{\delta^2}{6} (1 + Y^4). \end{aligned}$$

For the second estimate use $|\sin u - u| \leq |u|^3/6$ and $|\sin u| \leq |u|$; neither estimate needs u small. Since

$$Y^4 + Y^2/4 \leq \frac{9}{8}(1 + Y^4),$$

equation (3.1) implies

$$|A_\delta(Y)R(\delta Y) - P(Y)| \leq 37\eta(1 + Y^4)e^{4qY^2}. \quad (4.2)$$

We now fix the only compact set needed in the collision proof:

$$|\xi| \leq 2, \quad 0 \leq \tau \leq \frac{1}{4}.$$

On this set $Y^2 \leq 8+Z^2$. If $q \leq 1/64$, then

$$e^{4qY^2} \leq e^{1/2} e^{Z^2/16}, \quad 1 + Y^4 \leq 65 + 16Z^2 + Z^4.$$

For $a=8/7$, the Gaussian moment identities give

$$E[(65 + 16Z^2 + Z^4)e^{Z^2/16}] = 65a^{1/2} + 16a^{3/2} + 3a^{5/2}.$$

Each power displayed is less than 2, and $\exp(1/2) < 2$. Therefore

$$\sup E[(1 + Y^4)e^{4qY^2}] < 336.$$

Combining this with (4.2) proves the explicit N-uniform estimate

$$\sup_{|\xi| \leq 2, 0 \leq \tau \leq 1/4} |F_\delta(\xi, \tau) - (\xi^2 - 1/4 + 2\tau)| \leq K\eta, \quad K = 16384, \quad (4.3)$$

because $37 \cdot 336 = 12432 < 16384$. No bound on $L\delta$ is imposed. The coefficient-normalizing exponential can be large, but it is exact, real, and nonzero.

5. Quantitative positivity and root tracking: checked

Assume $\eta \leq \eta_0 = 1/(32K) = 1/524288$. This implies $q \leq 1/64$ and $\delta < 1$. Put

$$\tau_* = \frac{1}{8} + K\eta \leq \frac{5}{32} < \frac{1}{4}, \quad s_* = \delta^2 \tau_*.$$

At $\tau = \tau_*$, equation (4.3) gives, for every $|\xi| \leq 2$,

$$F_\delta(\xi, \tau_*) \geq \xi^2 + 2K\eta - K\eta \geq K\eta > 0. \quad (5.1)$$

At either boundary $\xi = \pm 2$, for every $0 \leq \tau \leq \tau_*$, it gives

$$F_\delta(\pm 2, \tau) \geq \frac{15}{4} - K\eta > 0. \quad (5.2)$$

Initially the moving interval $I_s = [-2\delta - 2Ls, 2\delta - 2Ls]$ contains exactly the selected two roots. Indeed any additional root with circular lift $|\theta_k| \leq 2\delta$ would contribute

$$\delta^2 B \geq \frac{\delta^2}{4 \sin^2(\theta_k/2)} \geq \frac{\delta^2}{\theta_k^2} \geq \frac{1}{4},$$

contrary to the assumed smallness. Also $4\delta < 2\pi$, so the initial interval does not contain a duplicated lift of the same root.

Suppose $D > s_*$. The family P_s has simple roots throughout the closed interval $[0, s_*]$. Since the multiplier $j(N-j)$ is symmetric under $j \leftrightarrow N-j$, the initial self-inversive relation is preserved. A root and its reciprocal-conjugate root agree initially. Local uniqueness of a simple root branch makes them agree throughout the interval. Hence every root remains on the unit circle before D . Fixed nonzero endpoint coefficients exclude escape via zero or infinity.

Continuous angle lifts of the two selected roots exist on $[0, s_*]$. They cannot cross either moving boundary, by (5.2). They cannot disappear or merge, by $D > s_*$. They must therefore still be real zeros inside I_{s_*} at the final time, contradicting (5.1).

It follows that

$$D \leq s_* = \delta^2(1/8 + K\eta).$$

This argument remains valid if the interval translates across many initial gaps or around the circle: it is formulated on the real cover, with continuously moving nonvanishing boundaries. It also remains valid if some other pair would have collided first; that already contradicts $D > s_*$ and establishes the same upper bound. A local zero count at the final time alone would have been insufficient, but the all-time boundary estimate supplies the missing step.

For completeness, the deterministic lower bound can be checked directly. Lift one adjacent pair to $a < b = a + g$, and every other point to c in $(b, a + 2\pi)$. Its contribution to the gap derivative is

$$\cot \frac{c-b}{2} - \cot \frac{c-a}{2} > 0,$$

because $\cot$ is decreasing on $(0, \pi)$. The pair's own contribution is $-2\cot(g/2)$, so $g' \geq -2\cot(g/2)$. Compare every gap with the same scalar solution initialized at δ . That scalar solution satisfies $\cos(h(s)/2) = \exp(s)\cos(\delta/2)$, and reaches zero at $-\log \cos(\delta/2)$. No gap can collide while this comparison solution is positive. Finally $-\log \cos(\delta/2) \geq \delta^2/8$ by integrating $\tan u \geq u$ on $[0, \delta/2]$. Under $\eta \leq \eta_0$ one has $\delta < \pi$, so the real logarithm is well-defined. Thus the quantitative theorem is retained. There is no need to downgrade it to $o(1)$.

6. What must change when the object is the true H_t

The finite theorem does not define a global ζ depth by looking at a truncated list of zeros. In particular, a global first collision among infinitely many zeros and a local collision near a selected height are different quantities. After RH and the known nonnegativity of the Newman constant, global backward instability is not a new conclusion.

Work in the real coordinate of H_t itself; do not mix it with the usual ζ ordinate without the factor-of-two conversion.

Suppose two consecutive simple real zeros of H_0 are $c - \delta/2$ and $c + \delta/2$. Factor

$$H_0(c + y) = (y^2 - \delta^2/4)h_c(y), \quad L_c = h'_c(0)/h_c(0), \quad R_c(y) = e^{-L_c y} h_c(y)/h_c(0).$$

The exact heat identity is $H_{-}\{-s\} = \exp(s\partial^2)H_0$. It follows directly from the defining Fourier integral by real Gaussian convolution. Rodgers – Tao's equations (1) and (4) use $H_0(z) = \xi(1/2 + iz/2)/8$ and the multiplier $\exp(tu^2)$, respectively; these equations were checked in the downloaded primary paper. In the moving frame $y = \delta\xi - 2L_c s$, the identity gives

$$\frac{e^{-L_c \delta \xi + L_c^2 s}}{\delta^2 h_c(0)} H_{-s}(c + \delta\xi - 2L_c s) = E[(Y^2 - 1/4)R_c(\delta Y)], \quad s = \tau\delta^2. \quad (6.1)$$

Thus a precise sufficient analytic obligation for local collision localization is

$$\sup_{|\xi| \leq 2, 0 \leq \tau \leq 1/4} E[(1 + Y^2) | R_c(\delta Y) - 1 |] \rightarrow 0. \quad (\text{Z-tail})$$

This condition states both local flatness and Gaussian-tail uniform integrability. A Taylor approximation on a shrinking real interval, without the tail condition, would not justify (6.1).

The root-tracking conclusion in this infinite setting must be stated locally: a multiple real zero occurs in the moving tube before the specified time, or a tracked real branch already met another branch there. It is not a lower bound for a positive global first-collision time of H_t . The finite proof's assumption that all roots remain real globally before the first collision should not be imported unchanged into this setting.

7. Under RH, the canonical product supplies the far-tail structure

There is a useful simplification: under RH, the true H_0 is a real entire function of order one with real zeros. After removing the pair and its linear logarithmic term, its genus-one canonical product gives

$$R_c(y) = \prod_{k \in \mathbb{Z}} \left(1 - \frac{y}{d_k}\right) e^{y/d_k}, \quad d_k = x_k - c, \quad B_c = \sum_{k \in \mathbb{Z}} d_k^{-2} < \infty.$$

Multiplicity is counted in this product. The selected two zeros must be simple; other multiplicities merely increase B_c . These are classical Hadamard-product facts used in the Lehmer-pair literature, not a new random-matrix representation. See Csordas – Smith – Varga and Rodgers – Tao.

For every real u ,

$$|1 - u| e^u \leq e^{u^2/2}. \quad (7.1)$$

One direct proof studies $f(u)=u^2/2-u-\log|1-u|$: its derivative is $u(2-u)/(1-u)$ on $u<1$ and $u(u-2)/(u-1)$ on $u>1$. The minima are $f(0)=f(2)=0$. At $u=1$ the inequality is immediate.

Apply (7.1) to finite partial products and pass to the canonical-product limit. This gives the global bound

$$|R_c(y)| \leq e^{B_c y^2/2}.$$

The same local logarithmic-curvature argument as in § 3 yields

$$|R_c(y) - 1| \leq 32B_c y^2 e^{B_c y^2/2}.$$

Consequently, the condition $\delta^2 B_c \rightarrow 0$ implies (Z-tail), including its Gaussian tails. This is a conditional analytic reduction in the spirit of classical Lehmer-pair estimates. It does not establish that the required isolated pairs exist, or occur with any particular frequency, among ζ zeros. No new unconditional or RH-conditional spacing result has been obtained here.

8. If a finite window is used, the exact omitted term

If the canonical product is split into nearby zeros and distant zeros, write $R=R_{\text{near}} R_{\text{far}}$ and $B_c=B_{\text{near}}+B_{\text{far}}$. Its missing outer contribution obeys

$$|R(y) - R_{\text{near}}(y)| \leq 32B_{\text{far}} y^2 e^{B_c y^2/2}.$$

Thus, in the regime $\delta^2 B_c$ bounded by a sufficiently small constant, the normalized Gaussian heat error from the omitted factor is $O(\delta^2 B_{\text{far}})$. A sufficient truncation obligation is

$$\delta^2 \sum_{|x_k - c| > W} |x_k - c|^{-2} \rightarrow 0.$$

The omitted **linear drift** must also be retained or removed using the true Lc. It need not be small merely because its curvature is small. More explicitly, the correct finite approximation before heat evolution is

$$h_c(0)(y^2 - \delta^2/4) \exp[(L_c + \sum_{\text{near}} d_k^{-1})y] \prod_{\text{near}} (1 - y/d_k).$$

It is a polynomial times a specified exponential; the exponential can be handled by the exact heat-conjugation identity. Replacing this by its polynomial factor alone silently changes the initial logarithmic derivative, and can shift the entire local configuration by more than one gap.

A strong uniform hard-core hypothesis can bound inverse-square tails by packing, but it generally makes $\delta^2 B_c$ only $O(1)$, not $o(1)$, at the minimum allowed spacing. A density version of AH allows exceptional clusters and therefore supplies no uniform bound on B_c for every selected window. Neither version alone yields (Z-tail) with vanishing error at the half-spacing threshold. Using the full true canonical product avoids a matrix-model or arbitrary polynomial-cutoff assumption, but it does not provide the missing arithmetic spacing information.

9. The shortest genuinely missing arithmetic lemma

For the proposed Level-B target of depth $o(\text{mean-spacing}^2)$ along a sequence, the clean missing input is an **isolated small-pair lemma**: there should exist consecutive simple zeros near heights $c \rightarrow \infty$ with

$$\frac{\delta_c}{h_c} \rightarrow 0, \quad \delta_c^2 \sum_{k \equiv \pm} (x_k - c)^{-2} \rightarrow 0,$$

where h_c is the mean spacing in the H-coordinate. The preceding analytic argument explains how such input would produce a local collision statement. It does not prove this arithmetic lemma. Bandwidth-one information alone does not establish it, and merely assuming it inside a “near-collision step” would assume the main unresolved microscopic input.

For an AH that permits a zero-density exceptional set, even this subsequential lemma is insufficient to refute AH. The task then needs a positive-density arithmetic statement in a fixed forbidden normalized-gap interval, or an equivalently strong statistic controlling the density of exceptions. A toy quantifier check is a block of N gaps consisting of εN , $2-\varepsilon N$, and $N-2$ gaps equal to 1. Their sum is N ; all but $2/N$ of the gaps are integer, hence half-integer, while the minimum εN tends to zero. Having a tiny minimum in every large block does not by itself contradict a density AH. This toy example is only a quantifier counterexample, not a claim to match all ζ correlation constraints.

The recommended next arithmetic task is therefore to specify and prove the weighted pair-count or energy estimate that creates a positive proportion of forbidden spacings, with its exact support and averaging range. The finite heat theorem cannot replace that input. The current audit supplies a clean analytic interface and identifies which tail estimates become automatic under RH; it supplies no new ζ spacing theorem.

10. Reproducibility and review status

The companion `galilean_audit_constants.py` checks the exact rational inequalities used to choose K and η_0 , and verifies symbolic derivatives in the centered sine factor and canonical-product factor. Its output is saved beside it. These checks verify the displayed algebra and numerical constants; they do not verify the topology or the analytic argument in Lean.

No Claude tool was started. No repository source was edited. All assertions above are confined to the stated deterministic theorem and explicit conditional analytic obligations.

Current report 09: prime186–yau–independent–review

Source: `research/reports/prime186–yau–independent–review.md`. SHA-256:

557a3dda43b4b5b524c84e6b1d83eb97195313f1a9cccd1d85f17c081239633d.

Independent adversarial review of the circular heat isolation theorem

Reviewer: the prime186 research agent. Date: 2026-09-05. Reviewed `yau_flow.md` §§ 3.1 – 3.4, § 4.1 – 4.4, and § 10.1 – 10.2, and checked the stated input against the locally downloaded Feng – Wei paper. This is an independent mathematical review, not formal verification or a novelty determination.

Conclusion: I found no fatal gap in the deterministic isolation argument or its transfer to CUE and fixed positive integer circular-beta ensembles. The argument supplies a plausible complete proof of the stated finite-ensemble depth law. It does not provide a zeta-zero theorem.

Checks that could have invalidated the argument

- Flow sign and time normalization.** Multiplication by $\exp[s \cdot j(N-j)]$ on a centered characteristic polynomial is ordinary scalar heat at time s , up to an irrelevant scalar. The normal Gaussian displacement is $\sqrt{2s} \cdot Z$. The quadratic local model is therefore $x^2 - 1/4 + 2\tau$, which disappears at $\tau = 1/8$. This agrees with the root ODE and with the exact two-body collision time $-\log \cos(\delta/2)$.
- Gaussian tails.** The estimate $|H(y)/H(0) - 1| \leq \exp(A|y|) - 1$ is global for real y . It does not use a local Taylor expansion outside its range. Under $\delta A \leq 1$, the scaled integrand is dominated by a polynomial times $\exp(|Y|)$, uniformly on compact scaled space/time ranges. Gaussian integrability gives the required uniform convergence, including time zero.
- First collision versus a disappearing local minimum.** The proof does not infer a collision from final-time positivity alone. It keeps both fixed boundary values nonzero at every intermediate time and tracks the two initial simple roots. Consequently they cannot escape the interval without a collision. If another pair collides earlier, the claimed global upper bound is already satisfied.
- Persistence on the circle.** The coefficient evolution preserves self-inversive symmetry. Before a multiple root occurs, simple roots starting on the circle remain there by uniqueness in the implicit function theorem and reciprocal-conjugate symmetry. Leading and constant coefficients remain nonzero throughout, so no root escapes through zero or infinity.
- Long initial gaps.** The deterministic lower bound compares every adjacent gap to the same scalar solution initialized at the global minimum, rather than applying a logarithm of a negative cosine to a long gap. The scalar field $-2\cot(g/2)$ is increasing on $(0, 2\pi)$, and the background contribution to an adjacent-gap derivative is nonnegative.
- Initial third-root exclusion.** A third root at distance at most $R\delta$ from the selected pair midpoint would force δA to be bounded below by a positive constant. Thus the hypothesis genuinely isolates exactly two roots inside the fixed scaled interval.
- Triple-overcrowding probability.** The CUE divided-difference determinant estimate gives the power N^9 and the full squared Vandermonde. Integrating two anchored offsets contributes r^8 . At $r = N^{-(7/6)}$, the probability tends to

zero at rate $O(N^{-1/3})$.

8. **Packing and quantifiers.** Triple-free arcs imply at most two points per arc; with the chosen closest pair isolated, ordering reciprocal chord distances gives $A \leq C r^{(-1)(1+\log N)}$. Only tightness of the rescaled minimum gap is needed. No growing-argument extrapolation of the limiting tail is used.
9. **General beta input.** Feng – Wei Lemma 1.1 explicitly supplies $C_{\{\beta, N-3, (3)\}}/C_{\{\beta, N\}} \leq (N\beta)^{(3\beta)}$ for $\beta \geq 1$. Hölder with exponents $(3, 3, 3)$ produces precisely the stated three-point bound. Their extreme-gap theorem has the positive integer beta restriction used in the report.
10. **Constants.** The paper gives $A_1=1/24$, $A_2=1/(24\pi)$, $A_4=1/(270\pi)$. Under $D \sim \delta^2/8$, these yield the reported tails $\exp(-t/6)$, $\exp[-2\sqrt{2}t^{(3/2)}/(9\pi)]$, and $\exp[-64\sqrt{2}t^{(5/2)}/(675\pi)]$.

Suggested final-statement clarifications

Define D explicitly as the infimum of positive times at which the evolved polynomial has a multiple root, with infimum of the empty set equal to infinity. State that initial roots are distinct almost surely in the random ensemble. Mention the preserved nonzero leading and constant coefficients in the circle-persistence argument. Keep the qualification that novelty requires comparison with Lehmer-pair/Newman and function-field literature.

The proof's arithmetic limitation is substantive. It exploits an extreme pair whose angular gap is asymptotically much smaller than the mean spacing, and a probability law guaranteeing initial isolation. A half-lattice mimicker fails that hypothesis. Producing it for zeta zeros remains the missing arithmetic information, and a single or sparse exceptional close pair does not refute an Alternative Hypothesis permitting density-zero exceptions.

Current report 10: prime186

Source: `research/reports/prime186.md`. SHA-256:
`0b4c921d5facba4d079b1c9a54a6958aed18c63be074aa6032cc1310fa7b6e1a`.

Prime-gap 186 audit, a certified minorant refinement, and the next arithmetic target

Research round 1, 2026-09-05. This document records results actually obtained in this run. No bound below 186 is proved. The useful completed work is an exact rational enclosure for the fixed minorant's deficit, an honest propagation of its limited value through the full support-restored criterion, an exactly checked 39-tuple giving a concrete next target 182, and a sharp obstruction to improving one of the elementary debt constants by counting alone.

1. The current baseline and what was replayed

The official `PrimeGaps186` repository was cloned at commit `61340d0b74163003b32756bb16e91d9209a5e330`. Its accompanying paper proves a sufficient sieve criterion for DHL[40,2], then applies a 40-point admissible tuple of diameter 186. The public Lean development is explicitly conditional on three project inputs: two cited Deligne-type estimates and the physical integral bounds. Standard logical axioms are separate. A numerical receipt does not itself discharge those Lean inputs.

The local audit inspected the main paper, the numerical companion, and the exact-input Python certificate. It replayed scalar rational arithmetic from printed endpoints; it did **not** recompute the 149 physical source integrals. This distinction applies to every margin below.

The source-preserving clone is in `prime186-work/PrimeGaps186`. All new experiments are outside that clone.

The central normalized inequality is

$$M(F) = \rho_* \frac{J_{\lambda,H}(F) - L(F)}{I_H(F)} - 1 > 0.$$

Here the admissible trial is an outer coefficient profile with its actual rootwise support conditions, and L pays for every deleted part. A large cap-only Rayleigh quotient is insufficient. On the published fixed trial, exact scalar replay gives

$$\begin{aligned} \rho_* J^- / I^+ &= 1.0002060794024186 \dots, \\ L^+ / I^- &= 0.000696075110, \quad \rho_*(J^- - L^+) / I^+ - 1 = 0.000023360452297044097 \dots \end{aligned}$$

The scale $\rho_* = 2624989/10^7$ means that most of the cap surplus is consumed by support restoration. The remaining margin is about 23 parts per million.

The loss ledger matters more than the count of coefficient parameters:

Failure class	Published weighted loss, in units $10^{-12} I^-$	Fraction of total
Outer, effective order 2	38,927,522	5.5924%

Failure class	Published weighted loss, in units $10^{-12}f$	Fraction of total
Outer, effective order 5/2	622,829,241	89.4773%
Base inner, order 2	55,254	0.00794%
Base inner, order 5/2	435,544	0.06257%
Enlarged inner, order 2	1,405,159	0.20187%
Enlarged inner, order 5/2	32,422,390	4.65789%

Thus a substantial improvement should target the outer order-5/2 support geometry and the full constrained operator. It should not spend most of its budget shaving the already tiny base-inner loss.

2. Completed exact refinement of the fixed minorant deficit

Write the prime minorant as $\varrho = P - b$, where $b \geq 0$ is the explicitly defined five-prime exceptional contribution. Let κ be the limiting mean of b in units $\chi/\log \chi$. The implemented trial used $\kappa \leq 2 \cdot 10^{-5}$. The companion already supplied a smaller coarse rational bound. This round computes a sharper enclosure of the *same* fixed integral; it does not alter the arithmetic minorant or its distribution range.

The independently reproduced enclosure is

$$\frac{9333886966775225399315932}{10^{30}} \leq \kappa \leq \frac{9334314528491567966021794}{10^{30}}$$

or

$$9.333886966775226 \cdot 10^{-6} \leq \kappa \leq 9.334314528491569 \cdot 10^{-6}.$$

The absolute enclosure width is $4.2756171634 \cdot 10^{-10}$. The calculation used 16,384 four-dimensional simplices and took approximately 3.88 seconds on this host. Every inequality in the enclosure uses Python integer and rational arithmetic. Floating-point values select the next cell to split, which affects efficiency only.

2.1 Why this is a rigorous enclosure

Use $t = 481/100000$, eliminate the fifth coordinate by $\sum Z_i = 0$, and put

$$a_i = \frac{1}{5} + tz_i, \quad f(z) = \prod_{i=1}^5 a_i^{-1}.$$

The minorant's two four-dimensional polytopes have the six rational simplex cells recorded in the numerical companion. On their positive domain,

$$\log f = - \sum_{i=1}^5 \log a_i$$

is convex. Therefore f is convex. For a four-simplex Δ with vertices $v_0, \dots, v_4$ and centroid $\bar{v}$,

$$|\Delta| f(\bar{v}) \leq \int_{\Delta} f(z) dz \leq \frac{|\Delta|}{5} \sum_{j=0}^4 f(v_j).$$

The lower inequality is Jensen's inequality. For the upper inequality, convexity bounds the value at a point by its barycentric affine interpolant, whose average is the average of the five vertex values. Multiply by t^4 , sum over cells, and round each cell contribution outwards to an integer multiple of 10^{-30} . Longest-edge bisection preserves exact rational vertices and halves the exact volume. This yields reproducible lower and upper bounds without transcendental quadrature or FLINT.

2.2 Independent geometry verification

`minorant_geometry.py` verifies the underlying rational geometry, rather than merely assuming that a printed triangulation has no holes:

1. Enumerate all vertices by solving every four-active-halfspace system over the rationals, and reject candidates violating another halfspace. The first polytope has 11 defining halfspaces, the second 10. Both have exactly the seven printed vertices.
2. Check all six four-simplex determinants and volumes exactly.
3. Orient every simplex positively. Every internal tetrahedral face appears twice with opposite orientations; every remaining face lies on a supporting hyperplane of the polytope.
4. Check an explicit strict interior rational point at which the oriented covering multiplicity is exactly one.

The corresponding elementary chain-degree argument proves coverage almost everywhere: an oriented simplex chain whose boundary lies only on the boundary of a convex polytope has constant degree throughout its interior. Here that degree is one. Thus there is neither positive-volume overlap nor a missing positive-volume region. Both geometry audits pass.

The arithmetic identification of the exceptional-prime mean with these polytopes remains the mathematical input from the minorant construction, rather than a new claim proved by the Python code.

2.3 Carrying the refinement through the complete margin

A deficit improvement must be propagated through all coefficients, including the error weights. For fixed $\lambda = 1/125$ and $K = 17/50$, write

$$m = 1 - \kappa, \text{quad}a = m^2 - m\lambda, \text{quad}b = (1 - m/\lambda)\kappa K, \quad c = a + b, \quad d_0 = 1 - c.$$

The cap form is $J_0 + cJ_+ + bJ_t$. With the refined upper bound for κ , the coefficients change as follows:

Quantity	Published allowance	Refined allowance
a	0.9919601604	0.9919814061325887
b	-0.000843183	-0.0003935309975205265
$c = a + b$	0.9911169774	0.9915878751350682
d_0	0.0088830226	0.0084121248649319

Both cap coefficients improve, because $J_+, J_t \geq 0$. The base-inner weight d_0 and enlarged-inner weight $1 - b$ decrease. However, the positive outer-restoration multiplier on the enlarged shell increases from c_{old} to c_{new} . Ignoring that change would be an invalid free gain.

The new outer multiplier is at most $r = c_{new}/c_{old}$ times the old one everywhere. Thus a conservative complete loss bound follows from the old certified component bounds by multiplying the outer loss by r , the old-inner loss by $d'_{0,new}/d'_{0,old}$ and the new-inner loss by $(1 - b_{new})/(1 - b_{old})$. Even retaining only the old combined cap lower endpoint, the resulting complete margin is at least

$$0.00002328873833218647 \dots > 0.$$

This does **not** mean the refined trial is worse. It means that the old *combined* numerator endpoint does not reveal how much of the cap form lies in J_+ and J_t . Without these separate enclosures, an actual improvement cannot be quantified sharply. Keeping the old coefficients remains valid and retains the original 0.00002336045 ... margin.

There is also a rigorous ceiling on the leverage of this one adjustment. On the same fixed trial and same λ , the published operator bound gives

$$J_+ + J_t \leq J_{full} \leq 4I.$$

Consequently, the cap-quotient gain is bounded above by

$$4\rho_* \max\{c_{new} - c_{old}, b_{new} - b_{old}\} = 0.0004944405498715189 \dots$$

Even if every old restoration loss disappeared, the total improvement from this fixed-trial, fixed- λ adjustment would be below 0.0006771595. These bounds describe its scale; they do not prove that a $k = 39$ trial lies beyond or within reach. A direct $k = 39$ full-form calculation is still required.

The refinement changes no source exponent, no divisor-closed support theorem, and no tuple size. It is a useful exact certificate and a starting input for reoptimization, not a smaller prime-gap theorem.

3. A proved obstruction: the coefficient 12/5 is sharp

The elementary majorant $b \leq (12/5)N_2$ cannot be improved merely by sharpening its combinatorial case count. Here N_2 counts prime pairs whose product is below the specified threshold.

Take five distinct normalized prime exponents

$$(0.198, 0.199, 0.200, 0.201, 0.202).$$

They sum to one; every pair sum is below 0.40481, every triple sum is above 0.59519, and every exponent lies strictly inside the allowed roughness and upper bounds. Exhausting all 120 assignments to the labelled factors gives

$$b_1 = 4, \quad b_2 = 20, \quad N_2 = 10, \quad \frac{b_1 + b_2}{N_2} = \frac{12}{5}.$$

All inequalities are strict, so this is an open exponent chamber, not a boundary coincidence. The finite assignment count is independently verified using rational arithmetic in `margin_ledger.py`.

This only closes the route of lowering the *pointwise constant* without changing the majorant. One can still seek a better majorant, retain more of the exact exceptional geometry, or prove a stronger weighted estimate for the balanced five-prime convolution. In particular, the unweighted mean $\kappa \approx 9.334 \cdot 10^{-6}$ cannot automatically replace the weighted exceptional-energy constant $K = 0.34$: those are different mathematical quantities.

4. A concrete next prime-gap target: DHL[39,2] implies 182

A mixed-integer search produced the following 39-point admissible tuple, and a separate integer residue check validated it:

```

0, 2, 6, 8, 12, 20, 26, 30, 36, 38, 42, 48, 50,
56, 62, 66, 68, 72, 78, 80, 90, 92, 108, 110, 126,
128, 132, 138, 140, 146, 150, 152, 156, 162, 168,
170, 176, 180, 182

```

For each prime $p \leq 39$, at least one residue class is omitted; for $p > 39$ this is automatic from the cardinality. Its diameter is exactly 182. Thus DHL[39,2] would imply $H_1 \leq 182$, giving a precise target for the analytic computation.

This tuple is not claimed to be new or optimal. The search also returned a 38-point witness of diameter 176. HiGHS reported infeasibility for 40 points within diameter 184 and for 39 points within diameter 180. Those are numerical solver outcomes, not independently checked infeasibility certificates. An attempted proof-producing Z3 replay consumed approximately 6.7 GB without delivering an inspectable certificate and was terminated; no exact minimal-diameter theorem is claimed from it.

The feasible witnesses are independently checkable and are sufficient for research planning. The old $k = 49$ door to 240 is superseded as a main target by the existing 186 result.

5. A transferable support principle and a multi-minorant identity

The most valuable structural lesson is to preserve the exact arithmetic support and the pointwise nonnegative detector simultaneously. Rootwise conditions must imply the needed property of the actual least common multiple, including shared prime factors, exhausted roots, and inward cell restrictions. A condition on a generic product of independent roots is not a substitute.

The complementary allocation principle is a useful design language: split a required prime-power charge into nondecreasing contributions assigned to the two roots, with a fixed combined budget. Optimize the allocation as a function of prime size, then charge the actual excluded fragment configurations to the operator. This is stronger than blindly imposing one smoothness cutoff, but it does not abolish the parity or distribution barriers.

There is a clean extension of the residual-square bookkeeping to several minorants, which may be useful when support radius and minorant quality vary. This is an algebraic lemma, not a claim that suitable new arithmetic minorants have already been constructed.

Let $P \geq 0$, $b_j \geq 0$, $\rho_j = P - b_j$, $A = B_0 + D$, and $C = \sum_j C_j$. For arbitrary $\eta_j > 0$,

$$\begin{aligned}
 & PA^2 - [P(2AB_0 - B_0^2) + 2 \sum_j \rho_j DC_j - PC^2 - \sum_j \eta_j b_j D^2 - \sum_j \eta_j^{-1} b_j C_j^2] \\
 &= P(D - C)^2 + \sum_j b_j (\sqrt{\eta_j} D + C_j / \sqrt{\eta_j})^2 \geq 0.
 \end{aligned}$$

Thus the residual loss is explicit and remains nonnegative. If the relevant arithmetic averages are established, the corresponding normalized variational gain is

$$2 \sum_j m_j \langle V - H_0, H_j \rangle - \left\| \sum_j H_j \right\|^2 - \sum_j \frac{K_j}{\eta_j} \|H_j\|^2 - \sum_j \eta_j K_j E_0.$$

For mutually disjoint correction shells this separates into scalar optimizations. With different minorant qualities and distribution supports it suggests a richer admissible operator. With identical minorants and identical supports, splitting one correction into several labels provides no automatic benefit: the covariance and residual charges must still be paid. Any proposed gain must prove the needed mixed and self-square distribution estimates on all actual moduli.

This identity is more transferable to the zeta programme than another decimal improvement of κ : it makes the missing arithmetic information and its squared error cost explicit.

6. Reproducibility and the FLINT obstacle

The fresh official certificate was **not** run past startup. Its mandatory signed-convolution test fails in both local wheel environments tried:

- Python-FLINT 0.9.0, FLINT 3.6.0, NumPy 2.5.2;
- Python-FLINT 0.8.0, FLINT 3.3.1, NumPy 2.2.6.

The check was not disabled or replaced. The official upstream correction is identifiable: [FLINT pull request 2790](#), specifically commit `7ad753d51c82fdec115cb179b41d0e581f1cb0ec`. A source checkout of v3.6.0 at `8d5454b96761fafe4d5a9da76a369a602f500f49` was prepared. A patched build was not completed before the research priority was redirected to the zeta arithmetic barrier. Neither wheel failure is evidence against the mathematical theorem; it is a reproducibility requirement explicitly anticipated by its authors.

The new rational computations need none of these libraries. Run, from `prime186-work`:

```
python3 minorant_geometry.py
python3 minorant_mass.py
python3 margin_ledger.py
```

`tuple_search.py` additionally requires SciPy; only the search uses floating-point optimization, and every accepted tuple is checked with integer arithmetic.

Files suitable for a small research commit are:

- `prime186.md`: this report;
- `minorant_mass.py` and `.json`: exact convexity enclosure;
- `minorant_geometry.py` and `.json`: exact polytope/triangulation verification;
- `margin_ledger.py` and `.json`: full scalar margin, coefficient leverage bound, and sharp 12/5 count;
- `tuple_search.py`, `tuple_search.json`, and `tuple_search_more.json`: feasible witnesses and explicitly limited solver results.

Do not commit cloned upstream sources, virtual environments, or the unfinished proof-producing SMT run as if they were new mathematical results.

7. Next action and stopping discipline

The next substantial prime-lane computation would be a correctly linked full baseline, separate enclosures for J_0 , J_+ , J_t , and a direct $k = 39$ variational search that includes the dominant outer support failure. A cap-only success must not be reported as a prime-gap improvement.

Following the user's strengthened priority on a major zeta/random-matrix theorem, this round stops refining the auxiliary minorant integral. Its exact certificate is saved and can be reused. The transferable lessons are the support-preserving allocation structure, the positive residual identity, and the insistence on a complete cost ledger. These are useful when testing whether a new arithmetic zero statistic actually escapes the known ACUE mimickers.

Current report 11: resummed_prime_profiles

Source: `research/reports/resummed_prime_profiles.md`. SHA-256:
`1bb8d49fc05c7e7008134d00328a77cc59a9d0344f011ebbd7ca4acd5282b64e`.

Resummed symmetric prime profiles: a stable negative variational experiment

Research round 2, 2026-09-05. This is a bounded continuum experiment suggested by the residual-Gram agent's arithmetic operator. It is not a zeta theorem, a certified interval optimization, or a proof of optimality.

Outcome

The family below removes the ill-conditioned large symmetric-feature Gram matrix by resumming a multiplicative prime-factor profile. It reproduces the unmodified one-variable benchmark and improves that restricted benchmark, but does not exceed the better polynomial prime-feature result obtained by the other agent. No half-gap crossing occurs in the validated trials.

Prime profile	Radial degree / quadrature order	Half-gap margin
$g = 1$, optimized original parameter	6 / 32	-0.0153579821816782
$g(u) = \exp(c_2 u^2)$	7 / 32	-0.0147198770500968
$g(u) = \exp(c_2 u^2 + c_3 u^3)$	9 / 40	-0.0146705663327385
Three nontrivial coefficients, through u^4	9 / 40	-0.0146638642779015
Five coefficients, through u^6 ; local continuation	9 / 40	-0.0146638632200778

The last two searches are local continuations, not global optimizations; one hit its iteration budget. The radial Gram condition numbers in the high-precision-density runs were between approximately 1.6 and 2.4. Increasing radial degree and quadrature order changes the displayed negative margins only slightly. This offers a stable independent check that these natural prime-factor tilts recover a small fraction of the deficit, not an arithmetic advance past one half.

The family and its continuum density

Consider the formal asymptotic family

$$r(n) = d_\ell(n) \prod_{p \mid n} g\left(\frac{\log p}{\log L}\right) f\left(\frac{\log n}{\log L}\right), \quad g(u) = \exp\left(\sum_{k=2}^d c_k u^k\right), \quad a = \ell^2.$$

The product is symmetric in the prime factors. In the squarefree leading continuum model, a linear exponent in g is redundant with the radial profile f and is therefore omitted; this is not an exact finite-integer identity for a product over distinct prime divisors. Prime-power conventions require a precise choice in an eventual arithmetic theorem; their effect on the leading continuum formula must be justified, not assumed away.

Under the same Poisson – Dirichlet continuum model used for the prime-feature trial, the squared-coefficient density is proportional to

$$R(\nu) = \nu^{\hat{a}-1} H(\nu), \quad H(\nu) = \mathbb{E} \exp \left(2 \sum_{k=2}^d c_k S_k \right),$$

where the partition has total mass ν . This resums all powers of the chosen symmetric statistics, rather than truncating a polynomial in $S_2, S_3, \dots$

Write

$$g(u)^2 = \sum_{n \geq 0} d_n u^n, \quad H(\nu) = \sum_{n \geq 0} H_n \nu^n.$$

The coefficient recurrence used in the experiment is

$$H_0 = 1, \quad H_n = \frac{\hat{a}}{n} \sum_{k=1}^n k d_k B(k, \hat{a} + n - k) H_{n-k}.$$

One derivation uses the formal Laplace transform of the convolution exponential with kernel $\hat{a}(g(u)^2 - 1)/u$. If B_n are the coefficients of

$$\exp \left(\hat{a} \sum_{k \geq 1} d_k \Gamma(k) z^k \right),$$

then $H_n = \Gamma(\hat{a}) B_n / \Gamma(\hat{a} + n)$; differentiating the generating series gives the displayed normalized recurrence. This is a formal power-series identity: the factorial-weighted intermediate series need not have a nonzero analytic convergence radius. The final density series is evaluated with an explicit truncation comparison in the experiment.

For example, when $g(u) = e^{cu^2}$, the linear coefficient in $\mathcal{C}$ is

$$H(\nu) = 1 + \frac{2c}{\hat{a}+1} \nu^2 + O(c^2),$$

in agreement with the Poisson – Dirichlet moment of S_2 .

The changed variational form

Use ν for the background mass and u, w for inserted prime masses. Relative to the original one-variable form:

- Replace the background measure $\nu^{\hat{a}-1} d\nu$ by $\nu^{\hat{a}-1} H(\nu) d\nu$ in both numerator and norm.
- Multiply the two distinct-prime insertion integrand by $g(u)g(w)$.
- The same-prime diagonal insertion term has no extra factor $g(u)^2$: its resonator coefficient lives on the background integer, not the newly inserted output prime.
- A linear insertion, away from the exact half-gap boundary, acquires $g(u)$.

The distinction in the third item is essential. Inserting an extra $g(u)^2$ there would change the arithmetic operator and could manufacture a false improvement.

The case $c_2 = \dots = c_d = 0$ reduces exactly to the independently reproduced Inoue form, which the code checks to numerical precision. Passing from these coherent continuum formulas to a zeta theorem would still require the weighted-integer asymptotics and error accounting for this new coefficient family.

Numerical controls and failure record

The initial implementation evaluated the density series entirely in binary64. It reproduced the first two improved trials, but a higher-dimensional optimizer found a parameter vector for which the 32-node validation rejected the density truncation. That run was not accepted as a result.

The revised implementation computes the density coefficients and evaluations with 55 decimal digits, checks 110 versus 140 terms, and also checks the endpoint $v = 1$ rather than only quadrature nodes. The final Rayleigh matrices and generalized eigenvalue computations remain floating point. This fixes the identified density-evaluation weakness but does not make the experiment a rigorous interval certificate.

The most accurate recorded five-coefficient local continuation uses

$$\begin{aligned}\ell &= 1.0820458998724858, \\ c_2 &= -0.9407908108421177, \\ c_3 &= 1.17145568664247, \\ c_4 &= -2.268580431289358, \\ c_5 &= 0.0000165244269937828, \\ c_6 &= -0.0000964571420071413.\end{aligned}$$

Its tiny newly added coefficients reflect the chosen local continuation and should not be interpreted as proof that those directions are useless.

Files are in `research/prime-profiles/`:

- `euler_profile.py`: basic continuum operator and initial floating implementation;
- `euler_profile.json`: accepted initial trials before the rejected higher-dimensional run;
- `euler_profile_precise.py`: higher-precision density implementation;
- `euler_profile_precise.json`: validated negative trials and full coefficients.

Reproduction:

```
OPENBLAS_NUM_THREADS=1 OMP_NUM_THREADS=1 python3 euler_profile_precise.py
```

The main research effort should now return to the missing arithmetic mixed information. This experiment provides a stable resummed family and a negative result, not grounds for spending an indefinite budget on additional decimals of the same variational frontier.

Current report 12:

transferable_tools_ritz_decision

Source: `research/reports/transferable_tools_ritz_decision.md`. SHA-256:

`3c731beb9cf879a0607003ec364ee46775f88adff3b4b7afe67407745ad19bca`.

Transferable tools: one Ritz diagnostic and the actual support obligation

2026-09-05. This note responds to the proposal to reuse structural tools from the 186 prime-gap proof and FLT formalization. It records one completed low-cost experiment and the next exact arithmetic requirement. It does not claim a new gap bound.

Completed diagnostic

For the fixed rational trial at $L=100000$, with $\text{ell}=16/15$, write the normalized coefficient vector as u and the actual full arithmetic operator as K . Put

$$\lambda = \langle u, Ku \rangle, \quad r = Ku - \lambda u, \quad s = \|r\|, \quad v = r/s, \quad b = \langle v, Kv \rangle.$$

The exact two-dimensional compression identity is

$$K \mid_{\text{span}(u,v)} \leftrightarrow \begin{pmatrix} \lambda & s \\ s & b \end{pmatrix}.$$

The computed entries are

$$\begin{pmatrix} 4.192011775686907 & 0.177521239887580 \\ 0.177521239887580 & 0.933549641583058 \end{pmatrix}.$$

Its larger eigenvalue is 4.201654608743949 , and evaluating the resulting vector directly in the full operator gives 4.2016546087439535 . The orthogonality error is $1.39\text{e-}16$. The normalized margin improves from -0.03763020252239 to -0.03714169089289 , a gain of 0.00048851162950 . Both remain negative. Runtime of this recorded calculation was approximately 0.08 seconds; this is not a general speed claim.

This test uses self-adjointness only. K need not be positive semidefinite. The formula was also checked on an indefinite two-by-two control matrix. It does not replace the true function inner product with a coefficient norm in a nonorthogonal basis: here the integer coefficients already use the specified orthonormal coordinate normalization.

The newly generated vector contains prime-removal and cutoff-dependent sums through $A^* A u$. It is not automatically a fixed polynomial $d_{\text{ell}}(n)(f(v)+g(v)S^2)$. The completed fixed-family transfer theorem therefore does not automatically apply to this vector as L grows. A finite gain alone is not a new asymptotic certificate.

Decision: the experiment does not justify a large new sweep or more digits for the old negative trial. Preserve this residual direction as a diagnostic and focus the next main step on a new arithmetic input or a rigorous transfer for a direction with substantially larger demonstrated gain.

Evidence: `research/operator-bounds/ritz_residual_diagnostic.py`, its JSON output, and its run log. No eigensolver search in dimension L was performed by this test.

What complementary factorization contributes in the 186 proof

The relevant primary source is the [short-gap paper](#), Proposition 2.3. Its construction works with the actual modulus $Q=\text{lcm}(D,E)$. Nondecreasing complementary functions whose product is the cube of the prime scale divide the constraints between D and E . The proof accounts for the possibility that one side has no remaining prime at that scale, then verifies the primewise condition needed for triple dense divisibility. That property permits a larger sieve support within the applicable distribution estimates.

The transferable lesson is to enlarge the **legally evaluable combined object** before optimizing weights. It is not a permission to import an arithmetic-progression modulus exponent into a zeta Dirichlet-polynomial mean value.

The user's Weijie Su post was also supplied as context. Direct retrieval of that X page returned HTTP 403 in this run; this note uses the mathematical paper for the technical statement.

The corresponding zeta obligation

For the shifted packet with center $t_0=3T/2$ and width $W=T/\log T$, the exact Gram kernel of two integer-frequency polynomials has off-diagonal part

$$E_T(U, V) = \sum_{a \neq b} \frac{U(a) \overline{V(b)}}{\sqrt{ab}} e^{-it_0 \log(a/b)} e^{-\frac{1}{2} W^2 \log^2(a/b)}.$$

The common Gaussian mass is factored out here. In the present product-cutoff method, the combined coefficients arise at indices such as $a=km$ and $b=ln$. The accessible cutoff makes unequal integer products sufficiently separated for the source mean-value estimate.

A proposed complementary support must specify U and V from the same integer construction, allow genuinely new products, and prove a bound for this actual kernel on that class. For example, a usable new lemma would bound

$$|E_T(U, V)| \leq \epsilon_T \left(\sum_a |U(a)|^{1/2/a} \right) \left(\sum_b |V(b)|^{1/2/b} \right), \quad \epsilon_T \rightarrow 0,$$

or give a sharper explicit signed contribution with an error small enough for the final residual certificate. The displayed bound is an **unproved target for a specified new coefficient class**, not a theorem for arbitrary long polynomials. Dense nearby frequencies can violate such a uniform assertion for arbitrary coefficients.

Merely imposing complementary restrictions inside the existing unrestricted product cutoff changes the trial subspace of the same K ; it does not add a previously unavailable Gram entry. Moving to a larger cutoff creates a real new mean-value obligation. Centering the packet to recover Fourier positivity creates the explicit pole and low-height costs proved in the companion report. Neither issue is removed by changing terminology from smoothness to dense divisibility.

Signed residual recovery: useful, but with all terms specified

If positive measures μ and ν and a signed difference $\sigma=\mu-\nu$ arise from an actual construction, one may use the independently valid inequality

$$\|R\|_{\mu}^2 \geq \frac{(\|\mathfrak{m}\| - \epsilon - \sqrt{EG_{\nu}})^2}{G_{\mu}},$$

provided the mixed sigma correlation is within epsilon of the stated estimate, $\|R\|_{\nu}^2 \leq E$, $\|C\|_{\nu}^2 \leq G_{\nu}$, and $\|C\|_{\mu}^2 \leq G_{\mu}$, with $G_{\mu} > 0$. This is triangle inequality followed by Cauchy – Schwarz separately in the two positive measures. A signed measure is never treated as positive.

No new zeta-specific choice of these measures with a positive net gain has yet been established. In particular, the centered-packet pole counterexample remains in force. The next research contribution must identify and estimate these objects, not repeat the abstract inequality.

Formalization priority

The present exact identities and fixed-family transfer have sufficiently specific statements for future formalization. The useful FLT discipline is to check the actual definition and dependency closure of a result before spending time on its proof. An assumed long-support mixed-moment estimate would still be an assumption even if all subsequent linear algebra were checked by Lean.

Current report 13: dynamic_generator

Source: `research/reports/dynamic_generator.md`. SHA-256:

`cedc4fb9eac3df27c1e6872b4e31f6d32eb04dc8726bada103e97389f18ba4c7`.

Deterministic circular Coulomb dynamics does not leak protected trace moments

Date: 2026-09-05. **Status:** self-contained finite-dimensional proof, supported by exact symbolic and cyclotomic checks. This is a rigorous obstruction to the proposed low-mode mechanism, not a theorem about zeta zeros and not a novelty claim.

1. Result and the failed proposal

The user appendix proposes that applying successive deterministic circular Coulomb derivatives to a low trace mode might increase its effective degree until ACUE lattice aliasing becomes visible. The first displayed generator ansatz in that appendix is correct. Its next inference, that weighted degree grows from m to $2m, 3m, \dots$, is false.

For the radian convention

$$V_k(\theta) = \sum_{j \neq k} \cot \frac{\theta_k - \theta_j}{2}, \quad L = \sum_k V_k \partial_{\theta_k}, \quad \rho_m = \sum_k e^{im\theta_k},$$

one has

$$L\rho_m = -m((N-m)\rho_m + \sum_{a=1}^{m-1} \rho_a \rho_{m-a}). \quad (1)$$

The total **positive Fourier weight** of every term is exactly m . Since L is a real derivation, it preserves the negative weight separately as well.

Consequently, for every N , every $1 \leq m \leq N$, and every integer $r \geq 0$,

$$\mathbb{E}_{\text{ACUE}} L^r | \rho_m |^2 = \mathbb{E}_{\text{CUE}} L^r | \rho_m |^2. \quad (2)$$

The stronger finite-time statement is also true for the forward repulsive flow Φ_t :

$$\mathbb{E}_{\text{ACUE}} | \rho_m(\Phi_t X) |^2 = \mathbb{E}_{\text{CUE}} | \rho_m(\Phi_t X) |^2 \quad (t \geq 0, m \leq N). \quad (3)$$

Thus the proposed first distinguishing derivative $r_*(N)$ does not exist for these protected observables; writing $r_* = \infty$ is appropriate. No choice $t_N \rightarrow 0$, including $t_N \asymp N^{-2}$, produces the claimed discrepancy.

This conclusion is specific about the observable algebra. It does not say the two evolved measures agree, and it does not rule out rational observables involving inverse gaps. The independently audited force energy is a concrete example outside this algebra.

2. Definitions and normalization

We use eigenangles in radians throughout. The appendix alternates between angles in $\mathbb{R}/\mathbb{Z}$, cotangent with argument $\pi(\theta_i - \theta_j)$, and traces $e^{im\theta}$. Those conventions need a corresponding conversion in both Fourier modes and the time parameter. The conclusion about invariant degree survives any fixed nonzero time rescaling.

There is also a factor-of-two correction:

$$\partial_{\theta_k} \log \prod_{k < j} |e^{i\theta_i} - e^{i\theta_j}|^2 = \sum_{j \neq k} \cot \frac{\theta_k - \theta_j}{2} = V_k. \quad (4)$$

It equals V_k , not $2V_k$, in this radian convention.

Let $K = 2N$, $\zeta = e^{2\pi i/K}$, and let S be an unordered N -element subset of $\{0, \dots, K-1\}$. The unrotated ACUE probability is

$$P_A(S) = K^{-N} \prod_{\substack{a, b \in S \\ a < b}} |\zeta^a - \zeta^b|^2. \quad (5)$$

One may apply a common uniform rotation to recover the rotation-invariant ACUE model. All observables audited here are balanced and unchanged by a common rotation, so the unrotated subset model suffices. Equation (5) is for unordered subsets; an ordered tuple formula contains an additional $1/N!$.

The relevant primary reference is Tao, “The alternative hypothesis for unitary matrices,” 8 May 2019. It defines ACUE using the discrete Vandermonde weight and establishes matching low trace moments. The proof below independently derives exactly the finite collection of moment identities needed here.

3. Derivation of the generator

Set $z_k = e^{i\theta_k}$. For distinct points,

$$\cot \frac{\theta_k - \theta_j}{2} = i \frac{z_k + z_j}{z_k - z_j}.$$

Therefore

$$Lp_m = -m \sum_{k < j} \frac{(z_k^m - z_j^m)(z_k + z_j)}{z_k - z_j}. \quad (6)$$

The denominator cancels:

$$\frac{x^m - y^m}{x - y} = \sum_{a=0}^{m-1} x^{m-1-a} y^a.$$

Collecting terms in (6) yields

$$\sum_{k < j} \frac{(z_k^m - z_j^m)(z_k + z_j)}{z_k - z_j} = (N - m)p_m + \sum_{a=1}^{m-1} p_a p_{m-a}.$$

This proves (1). It holds for all positive m , including $m > N$; finite- N polynomial relations may then make a partition basis redundant, but do not invalidate the identity.

Taking complex conjugates gives

$$Lp_{-m} = -m((N-m)p_{-m} + \sum_{a=1}^{m-1} p_{-a}p_{-(m-a)}). \quad (7)$$

There is no additional sign on the right of (7).

For partitions λ, μ , write

$$p_\lambda \overline{p_\mu} = \prod_i p_{\lambda_i} \prod_j p_{-\mu_j}.$$

By the Leibniz rule and (1), (7), every term of its image under L has the same pair of weights

$$(\mid \lambda \mid, \mid \mu \mid). \quad (8)$$

The number of factors can increase when a part splits; the sum of their indices cannot. For fixed positive weight m , at most m factors can occur. In particular $m = 1$ never splits.

The weighted-degree conservation is the exact location where the appendix's proposed mechanism fails. Products of lower modes do not create higher total Fourier weight. No later iteration can reach $2N$ -aliasing starting inside the protected band.

4. Explicit low-weight blocks

The positive-weight blocks are:

$$\begin{aligned} Lp_1 &= -(N-1)p_1; \\ Lp_2 &= -2(N-2)p_2 - 2p_1^2, \\ L(p_1^2) &= -2(N-1)p_1^2; \end{aligned}$$

and

$$\begin{aligned} Lp_3 &= -3(N-3)p_3 - 6p_1p_2, \\ L(p_1p_2) &= -(3N-5)p_1p_2 - 2p_1^3, \\ L(p_1^3) &= -3(N-1)p_1^3. \end{aligned} \quad (9)$$

If G_m denotes the matrix on the positive partition block, the balanced block generator is the corresponding tensor sum

$$T_m = G_m \otimes I + I \otimes G_m. \quad (10)$$

Its size is at most $p(m)^2$, independent of derivative order. For $m = 1, 2, 3$ those sizes are 1, 4, 9.

For example, for both ensembles and $N \geq 2$,

$$\mathbb{E} \mid p_1(\Phi_t X) \mid^2 = e^{-2(N-1)t}, \quad (11)$$

and

$$\mathbb{E} \mid p_2(\Phi_t X) \mid^2 = 2e^{-4(N-2)t} [1 + (1 - e^{-2t})^2]. \quad (12)$$

Equation (12) follows from

$$p_2(t) = e^{-2(N-2)t} [p_2(0) - (1 - e^{-2t})p_1(0)^2].$$

The two initial terms are orthogonal and each has squared norm 2 in the protected moment range.

These expectations generally change with time. Deterministic Coulomb relaxation is not stationary Dyson Brownian motion; CUE stationarity must not be invoked for this deterministic flow.

5. Exact Haar and cyclotomic moment proof

Let s_α denote a Schur polynomial, with α padded to length N . For Haar measure,

$$\mathbb{E}_{\text{CUE}} s_\alpha \overline{s_\beta} = \delta_{\alpha\beta} \quad (\ell(\alpha), \ell(\beta) \leq N). \quad (13)$$

For the discrete law (5), the alternant formula and Cauchy – Binet give the exact expression

$$\mathbb{E}_{\text{ACUE}} s_\alpha \overline{s_\beta} = \det [1_{\alpha_i + N - i - \beta_j + N - j \pmod{2N}}]_{i,j=1}^N. \quad (14)$$

Here i, j are one-based. If a partition has more than N parts, the corresponding Schur polynomial is zero.

For completeness, multiplying the two Schur alternant ratios by the squared Vandermonde weight cancels both denominators. Summing over unordered N -subsets is exactly Cauchy – Binet for the two rectangular alternant matrices.

The resulting matrix entries are

$$\frac{1}{2N} \sum_{\mathfrak{a}=0}^{2N-1} \zeta^{\mathfrak{a}(\alpha_i + N - i - \beta_j + N - j)},$$

which equal the residue indicators in (14). Taking both partitions empty also verifies normalization of (5).

If $|\alpha|, |\beta| \leq N$, every alternant exponent lies in

$$0 \leq \alpha_i + N - i \leq 2N - 1.$$

Consequently congruence in (14) is actual equality. Each exponent list is strictly decreasing, so the determinant is 1 when $\alpha = \beta$ and 0 otherwise. This recovers (13) in the whole required band.

Expanding power sums using symmetric-group characters,

$$p_\lambda = \sum_{\mathfrak{a} \vdash m} \chi^\mathfrak{a}(\lambda) s_\mathfrak{a}, \quad (15)$$

proves equality of every Gram entry

$$\mathbb{E}_{\text{ACUE}} p_\lambda \overline{p_\mu} = \mathbb{E}_{\text{CUE}} p_\lambda \overline{p_\mu} \quad (|\lambda| = |\mu| = m \leq N). \quad (16)$$

Equivalently the common Gram is diagonal, with entries

$$z_\lambda = \prod_{j \geq 1} j^{\mathfrak{a}_j} \mathfrak{a}_j!,$$

where $\mathfrak{a}_j$ is the number of parts equal to j .

Combining (8) and (16) proves the all-orders equality (2). This proof does not rely on observing numerical agreement up to $r = 8$.

6. Full forward–time equality and an independent linearization

The repulsive flow exists for all $t \geq 0$ from every collision-free initial configuration. Indeed set

$$W(\theta) = \log \prod_{i < j} |e^{i\theta_i} - e^{i\theta_j}|^2.$$

By (4),

$$\frac{dW}{dt} = \sum_k V_k^2 \geq 0. \quad (17)$$

Every pairwise factor is at most 4. Hence $W(t) \geq W(0) > -\infty$ prevents any factor from approaching zero. The trajectory stays in a compact collision-free subset of the torus, where the vector field is smooth, proving global forward existence.

ACUE configurations are distinct by construction, and CUE configurations are distinct almost surely.

For the vector of positive partition polynomials, $u'_m(t) = G_m u_m(t)$, so

$$u_m(t) = e^{tG_m} u_m(0). \quad (18)$$

Its balanced moment matrix evolves by applying e^{tG_m} on the two sides. Equal initial moment matrices therefore remain equal for every forward time. This proves (3). The same argument handles any symmetric trace polynomial with positive and negative weights individually at most N .

There is also a useful direct linearization. Define the monic characteristic polynomial

$$P(z, t) = \prod_{k=1}^N (z - z_k(t)) = \sum_{k=0}^N (-1)^k e_k(t) z^{N-k}.$$

The root ODE gives

$$\partial_t P = z^2 \partial_z^2 P - (N-1)z \partial_z P. \quad (19)$$

To verify this, at each root use

$$\frac{P''(z_k)}{P'(z_k)} = 2 \sum_{j \neq k} \frac{1}{z_k - z_j}$$

and differentiate $P(z_k(t), t) = 0$. The degree- N coefficient on the right of (19) cancels, so equality at the N distinct roots proves the polynomial identity.

It follows coefficient by coefficient that

$$e_k(t) = e^{-k(N-k)t} e_k(0). \quad (20)$$

Thus the elementary symmetric functions provide an independently checkable diagonal form. This is a finite characteristic-polynomial calculation; no claim is made that such a linearization is new in the literature.

Time-direction limit. Statements (3), (17), and global physical flow refer to repulsion and $t \geq 0$. Reversing the sign produces the attractive flow and possible collisions. The algebraic finite-dimensional exponential remains meaningful, but one must not identify its negative-time continuation with a physical real-angle flow after a collision, nor assume a common negative collision-free interval for a CUE ensemble. This audit does not prove an attractive-flow Newman-depth law.

7. Exact audit requested by the user

The program evaluates $N = 2, \dots, 10$, $m = 1, 2, 3$, and $r = 0, \dots, 8$: 243 exact moment values per ensemble.

The exact algebraic work includes:

- Generator coefficients as integer polynomials in N , for every m, r .
- An assertion that every generated monomial retains biweight (m, m) .
- Direct rational-function verification of (1) for $N = 2, 3, 4$ and $m = 1, 2, 3$.
- Direct rational-function verification of $L\mathbf{e}_k = -k(N-k)\mathbf{e}_k$ for $N = 2, 3, 4$ and all $1 \leq k \leq N$.
- Haar Schur Gram matrices and integer residue determinants (14), followed by exact character transforms.
- Direct, independent subset enumeration in the integer quotient ring $\mathbb{Z}[z]/(\Phi_{2N}(z))$ for $N = 2, 3, 4, 5$. This verifies normalization and every relevant power-sum Gram entry without using the residue determinant to compute the enumerated answer.

234 protected cases agree exactly. The remaining nine cases all have $N = 2, m = 3 > N$; this observable already distinguishes ACUE from CUE at $r = 0$. Their differences are:

r	CUE expectation	ACUE expectation	ACUE minus CUE
0	2	1	-1
1	0	6	6
2	24	-12	-36
3	-384	-168	216
4	3552	2256	-1296
5	-26880	-19104	7776
6	184704	138048	-46656
7	-1204224	-924288	279936
8	7613952	5934336	-1679616

The difference is $-(-6)^r$, consistent with the exact evolved difference $-e^{-6t}$. This is an existing out-of-band discrepancy, not dynamically created access to protected information.

For example, for $N = 2, m = 3$, the Schur partition (3) has alternant exponents (4, 0), which coincide modulo 4, so its ACUE squared norm vanishes. The partition (2, 1) has exponents (3, 1) and squared norm 1. Haar gives norm 1 to each. Since $\rho_3 = \mathcal{S}_{(3)} - \mathcal{S}_{(2,1)}$ in two variables, the initial values are 1 and 2.

As a further **floating, non-exact** independent check, the program directly enumerates all subsets for every $N \leq 10$; the largest case has $\binom{20}{10} = 184756$ subsets. Across all these runs:

- maximum absolute Gram error: 2.3981×10^{-14} ;
- maximum derivative error divided by the documented absolute coefficient scale: 2.8866×10^{-15} ;
- maximum normalization error: 1.7764×10^{-15} .

The exact conclusions depend on the integer computations and proofs, not on those floating error tolerances.

8. Files, reproduction, and evidence boundaries

All paths in this section are relative to the directory containing this report:

- `../dynamic-generator/generator_audit.py` — self-contained exact and floating audit.
- `../dynamic-generator/generator_audit_results.json` — full symbolic coefficients, all 243 exact comparison rows, Gram matrices, enumeration checks, runtime and dependency versions.

Reproduce from any directory:

```
OPENBLAS_NUM_THREADS=1 python3 research/dynamic-generator/generator_audit.py
```

The observed environment was Python 3.14.3, NumPy 2.4.4, SymPy 1.14.0. No external inputs or sampling seeds are needed. The completed run took about one second on the available machine; this is an execution observation, not a general performance claim.

The source appendix was read from:

```
research/incoming/dynamic_generator_proposal.md
```

No original source notes or repository files were modified by this audit.

9. Correct replacement and what to postpone

The repaired research question is not whether this generator can escape its invariant polynomial algebra. It cannot. A useful new observable must lie outside that algebra, or a genuinely different evolution must fail to preserve it.

Inverse-gap force energies are a concrete candidate. In the same radian convention the independently derived identities are

$$E_{\text{CUE}} \sum_k V_k^2 = \frac{N(N^2 - 1)}{3}, \quad E_{\text{ACUE}} \sum_k V_k^2 = \frac{N(N^2 - 1)}{6}. \quad (21)$$

The force-energy audit is being independently written by the flow agent. Equation (21) separates the ensembles by a factor of two, with the same N^3 total-energy scaling. It is a static rational statistic of initial velocity, not leakage into a protected trace polynomial. Its number-theoretic accessibility remains a separate question.

Accordingly:

1. Stop searching higher derivative orders for this fixed deterministic flow acting on protected trace polynomials; the invariant-algebra theorem already answers all orders.
2. Preserve the appendix's failed degree-growth and low-mode instability conjectures in the record, together with the exact point of failure.
3. Study singular observables only with their integrability and required zero-correlation information made explicit.
4. Postpone a zeta-side lift until an actual arithmetic formula for the chosen inverse-gap observable is available.
5. Do not infer an AH contradiction, RH progress, or a Montgomery – Dyson theorem from either the present obstruction or the finite-ensemble force discrepancy.

The useful structural conclusion is stronger than a failed finite search: **deterministic circular Coulomb evolution preserves the complete ACUE/CUE protected moment sector at every forward time.**

Current report 14: dynamic_generator_independent_review

Source: `research/reports/dynamic_generator_independent_review.md`. SHA-256:

5167d262a69ccf0e22b7ec904f6abf76380fa2f444124e63a485700774084eab.

Independent review of the all-time protected-moment theorem

2026-09-05. Reviewed source: `research/reports/dynamic_generator.md`. Reviewer: the flow subagent, independently of the author of that report. This is a mathematical proof audit with separate diagnostic computations, not a formal proof-assistant certification or a novelty assessment.

1. Verdict

I found no gap in the stated all-forward-time identity

$$\mathbb{E}_{\text{ACUE}} \left| \rho_m(\Phi_t X) \right|^2 = \mathbb{E}_{\text{CUE}} \left| \rho_m(\Phi_t X) \right|^2, \quad 1 \leq m \leq N, \quad t \geq 0,$$

for the deterministic repulsive flow in radian coordinates

$$\dot{\theta}_i = V_i = \sum_{j \neq i} \cot \frac{\theta_i - \theta_j}{2}.$$

The invariant biweight argument, the discrete Schur Gram formula including the endpoint $m = N$, global forward existence, and the characteristic-polynomial linearization all check out. None requires interchanging a singular generator with an expectation. In particular, the failure of such an interchange for the force-square observable does not damage the protected-polynomial theorem.

One wording clarification is advisable: “any polynomial with positive and negative weights individually at most N ” means a polynomial in the **symmetric trace algebra**

$$A_{\leq N, \leq N} = \text{span} \{ p_\lambda \overline{p_\mu} : |\lambda| \leq N, |\mu| \leq N \}.$$

It should not be read as a statement about arbitrary polynomials of labeled particles. For example, applying the generator to one labeled coordinate Z_i produces uncanceled denominators; the symmetric pairwise cancellation is essential.

2. Independent reconstruction of the generator and its block

Write $Z_i = e^{i\theta_i}$. The identity

$$\cot \frac{\theta_i - \theta_j}{2} = i \frac{Z_i + Z_j}{Z_i - Z_j}$$

has the stated sign. Hence

$$L\rho_m = -m \sum_{i < j} \frac{(Z_i^m - Z_j^m)(Z_i + Z_j)}{Z_i - Z_j}.$$

For one unordered pair, the numerator quotient equals

$$z_i^m + z_j^m + 2 \sum_{a=1}^{m-1} z_i^{m-a} z_j^a.$$

After summing over pairs, symmetry under $a \leftrightarrow m - a$ gives

$$Lp_m = -m[(N - m)p_m + \sum_{a=1}^{m-1} p_a p_{m-a}].$$

Every term has positive weight m . The vector field is real on the angle torus, so $\overline{L\bar{f}} = \overline{L\bar{f}}$; the negative-weight formula therefore has the same sign, not the opposite sign. By the Leibniz rule, the linear span of $p_\lambda \overline{p_\mu}$ of any fixed biweight (a, b) is invariant.

At fixed weight, splitting a part increases the number of factors but never their total index. This is exactly why repeated applications cannot leave the protected sector. The number of derivatives is irrelevant.

Finite- N algebraic relations do not break the argument. One may use an overcomplete spanning vector and its explicitly specified coefficient matrix: each component still satisfies the finite linear ODE, and uniqueness gives its matrix exponential. Alternatively, quotient by the relations or use the elementary symmetric monomials.

3. The Schur determinant, its normalization, and the endpoint

For an unordered N -subset S of the $2N$ -grid, the probability is

$$(2N)^{-N} |\Delta(z_S)|^2.$$

There is no extra $N!$ in this subset formula. The Schur alternant cancels the Vandermonde in the measure. Cauchy – Binet over unordered subsets then yields

$$E_{\text{ACUE}} s_\alpha \overline{s_\beta} = \det[\mathbf{1}_{a_i \equiv b_j \pmod{2N}}]_{i,j=1}^N,$$

where

$$a_i = \alpha_i + N - i, \quad b_j = \beta_j + N - j.$$

Using empty partitions verifies normalization directly: the lists are $N - 1, \dots, 0$, and the determinant is 1. Thus there is no hidden factorial in the Gram identity.

For $|\alpha| \leq N$, one has

$$0 \leq a_i \leq \alpha_1 + N - 1 \leq 2N - 1.$$

The inequalities are inclusive exactly where needed. In particular, the extremal partition $\alpha = (N)$ gives the largest exponent $2N - 1$, which has not wrapped modulo $2N$. The endpoint $m = N$ is valid.

The lists a_i, b_j are strictly decreasing. Congruence inside the interval $[0, 2N - 1]$ is equality, so the determinant vanishes unless the lists coincide, and then it is +1. Their order is the same, so there is no undetected permutation sign. Coinciding exponent lists force $\alpha = \beta$.

Consequently the full Schur Gram agrees with Haar whenever both partition weights are individually at most N , even when the two weights differ. Passing from Schur polynomials to power sums gives the required initial trace Gram. At equal weight $m \leq N$, all partitions of m have at most N rows; hence the usual full symmetric-group character orthogonality is applicable and gives the diagonal power-sum Gram Z_λ .

A random common rotation does not change any balanced observable. For unequal weights in the indicated band, both expectations already vanish by the same Schur argument, so the larger symmetric-algebra statement remains correct after rotation as well.

Sharpness test just outside the band

For every $N \geq 2$, take $\alpha = (N + 1)$. Its exponent list is

$$(2N, N - 2, N - 3, \dots, 0).$$

The first and last exponents coincide modulo $2N$, so

$$E_{\text{ACUE}} | s_{(N+1)} |^2 = 0, \quad E_{\text{CUE}} | s_{(N+1)} |^2 = 1.$$

Thus the uniform guarantee cannot be extended to all symmetric polynomials of the next weight. This is a pre-existing discrepancy at $t = 0$, not leakage generated from a smaller weight.

The case $N = 1$ is exceptional for this sharpness example: there is only one exponent and no duplicate row, and every balanced one-particle monomial is constant. The main theorem includes $N = 1$ correctly. An initial draft of the independent check incorrectly applied the $N \geq 2$ sharpness example to $N = 1$; its assertion caught the mistake, and the example and script now retain this exception explicitly.

4. Global physical flow and the exact finite-time argument

The score of the squared Vandermonde in radians is V , so for

$$W = \sum_{i < j} \log | e^{i\theta_i} - e^{i\theta_j} |^2$$

one has $W' = \sum_i V_i^2 \geq 0$. Let $J = \binom{N}{2}$. For $N \geq 2$, each summand is at most $\log 4$, and therefore each individual summand is bounded below by

$$W(0) - (J - 1) \log 4.$$

For a fixed initial configuration this gives a strictly positive lower bound on every chordal distance along the whole forward trajectory. It rules out finite-time collision and even asymptotic approach to the collision set. Compactness of the torus and ordinary smooth-ODE continuation give global forward existence. For $N = 1$, the vector field is zero and the assertion is immediate.

No lower gap bound uniform over all CUE initial configurations is being asserted or needed. The bound is configuration dependent. The CUE initial law is collision free almost surely; every ACUE configuration is collision free.

Let u_a be the positive partition-polynomial vector of weight a . Its pathwise equation is $u'_a = G_a u_a$, so $u_a(t) = e^{tG_a} u_a(0)$. For fixed weights a, b , the matrix of products is thus a finite linear combination of the initial products. Equal initial Gram entries imply equal expectations at every fixed $t \geq 0$.

This proof uses neither a Taylor series justified only locally nor differentiation under an expectation. Each protected generator iterate is itself a Laurent polynomial extending continuously over collision configurations, and hence bounded on the compact torus. The singular denominators have cancelled before expectations enter. This is materially different from the rational force energy.

The conclusion remains valid for every choice of a finite time $t_N \geq 0$ depending on N . It does not identify the two full pushforward probability measures.

5. Independent check of the characteristic PDE and sign

From the same cotangent identity,

$$\dot{z}_i = -z_i \sum_{j \neq i} \frac{z_i + z_j}{z_i - z_j} = -2z_i^2 \sum_{j \neq i} \frac{1}{z_i - z_j} + (N-1)z_i.$$

For the monic polynomial $P(z, t) = \prod_i (z - z_i(t))$,

$$\frac{P''(z_i)}{P'(z_i)} = 2 \sum_{j \neq i} \frac{1}{z_i - z_j}.$$

Differentiating $P(z_i(t), t) = 0$ therefore shows

$$P_t(z_i) = z_i^2 P''(z_i) - (N-1)z_i P'(z_i).$$

The leading z^N term of the right side cancels. Since P_t also has degree at most $N-1$, equality at the N distinct roots proves

$$P_t = z^2 P_{zz} - (N-1)z P_z.$$

The coefficient of z^{N-k} is multiplied by

$$(N-k)(N-k-1) - (N-1)(N-k) = -k(N-k),$$

so indeed

$$\theta_k(t) = e^{-k(N-k)t} \theta_k(0).$$

This also shows θ_N is conserved, as required by $\sum_i V_i = 0$. It has the correct physical two-particle limit: for a gap g , $g' = 2 \cot(g/2)$, and therefore $\cos(g(t)/2) = e^{-t} \cos(g(0)/2)$. The first elementary coefficient decays by e^{-t} when $N = 2$, in agreement with the PDE.

Changing the flow to attraction changes the exponent signs. The resulting polynomial continuation exists for all real times as a coefficient formula, but its roots may leave the circle after a collision. The draft correctly restricts its physical all-time conclusion to repulsion and nonnegative time.

6. Independent computations and their limits

The companion script `dynamic_generator_independent_review.py` does not import the original audit program or read its computed answers. Its JSON is saved with the script in `research/dynamic-generator/`.

The integer portion checks every residue Gram entry between all partitions of weights $0, \dots, N$, for each $N = 1, \dots, 10$. It checks the endpoint exponent lists and the next-weight counterexample separately. All checks passed. At $N = 10$, this includes 139 partitions and 19,321 exact integer Gram entries.

The numerical portion independently evaluates Lp_m directly from the angle vector field for $m = 1, \dots, N+2$; evaluates $L\theta_k$ by differentiating elementary symmetric polynomials; and integrates the actual angle ODE at $N = 2, 3, 5, 8$ to compare its characteristic coefficients with the exponential formula. It uses times 0.03, 0.1, 0.3. The largest observed errors were:

- direct trace-generator identity: less than 3.4×10^{-14} ;
- elementary-coefficient generator identity: less than 1.6×10^{-14} ;
- integrated characteristic coefficients: less than 4.1×10^{-12} .

These last checks use floating-point arithmetic and numerical ODE integration; they are diagnostics, not exact computations or proofs. The mathematical arguments above establish the statements independently of those tolerances.

7. Required editorial changes and remaining scope

There is no mathematical change required to the central theorem. Clarify the symmetric trace-algebra scope of “any polynomial.” Retain the distinction between the finite-dimensional polynomial continuation and physical negative-time angle evolution. Preserve the finite- N , finite-ensemble nature of the result.

The result disproves the proposed mechanism in which repeated deterministic Coulomb differentiation of a protected trace polynomial generates an observable outside the matched band. It does not establish that a different flow or a singular statistic can be evaluated arithmetically. It is not an AH contradiction or a theorem about zeta-zero distributions.

Current report 15: force_energy

Source: `research/reports/force_energy.md` . SHA-256:

0782e73dc430034c54db8173d81803f02aaec6c326c151d692770635430394b .

Circular force–square observable: exact identities and expectation audit

2026-09-05. All angles in this report are measured in radians on a circle of circumference 2π . Here **D** means **force-square energy**, not the first-collision time denoted **D** in earlier reports. The calculations concern the deterministic repulsive flow $\dot{\theta}_i = V_i$.

1. Findings

For distinct angles, define

$$V_i = \sum_{j \neq i} \cot \frac{\theta_i - \theta_j}{2}, \quad D = \sum_i V_i^2, \quad Q = \sum_{i \neq j} \csc^2 \frac{\theta_i - \theta_j}{2}, \quad C_N = \frac{N(N^2 - 1)}{3}.$$

Then the following are exact mathematical identities:

$$D = Q - C_N, \quad E_{\text{CUE}} D = C_N, \quad E_{\text{ACUE}} D = \frac{C_N}{2},$$

where ACUE is the N -particle projection process on the $2N$ -th roots of unity with the consecutive N Fourier modes. For the deterministic generator

$$L = \sum_i V_i \partial_{\theta_i},$$

one has

$$LD = - \sum_{i \neq j} \csc^2 \frac{\theta_i - \theta_j}{2} (V_i - V_j)^2 \leq 0.$$

For every $N \geq 2$,

$$E_{\text{CUE}} LD = -\infty,$$

whereas the ACUE expectation is finite and in fact

$$E_{\text{ACUE}} LD = -\frac{2N(N^4 - 1)}{15}.$$

The last formula was first suggested by floating-point enumeration and then derived algebraically below. It is not being promoted from a numerical fit to a theorem without a proof.

The score normalization in the user's proposed argument needs correction: for the CUE density in radian coordinates, **the score is V , not $2V$** . Also, the apparent higher-order observable D collapses pointwise to a singular two-point energy. Its different expectations do not automatically provide an arithmetically accessible statistic of ζ zeros.

2. Score and deterministic–flow normalization

The circular-beta density is

$$p_\beta(\boldsymbol{\theta}) = Z_{N,\beta}^{-1} \prod_{i < j} |e^{i\theta_i} - e^{i\theta_j}|^\beta.$$

Differentiating $\log |e^{i\theta_i} - e^{i\theta_j}|$ gives one half of the cotangent. Thus

$$\partial_{\theta_i} \log p_\beta = \frac{\beta}{2} V_i.$$

For CUE, $\beta=2$, so the score is exactly V . This density convention agrees with [Feng – Wei's definition of \$C\beta E\$](#) . The discrete ACUE measure does not have a Lebesgue density on the full angular configuration space, so V should not be called its ordinary continuous Fisher score.

The generator here is for $\dot{\theta}=V$. If the deterministic drift is instead $V/(2N)$, all generator expectations below are divided by $2N$; the observable D itself is unchanged. If the attractive flow $\dot{\theta}=-V$ is used, the generator signs reverse. These deterministic flows are not stochastic DBM; Haar CUE invariance under a stochastic process does not imply invariance under the deterministic drift alone.

For the centered real characteristic polynomial

$$Q_0(x) = \prod_j \sin \frac{x - \theta_j}{2},$$

the force has the exact intrinsic expression

$$V_i = \frac{Q_0''(\theta_i)}{Q_0'(\theta_i)}.$$

This does identify a natural derivative-ratio candidate when one studies another real entire function. It does not evaluate its average over that function's zeros.

3. Pointwise collapse of the force square

Put $c_{ij} = \cot((\theta_i - \theta_j)/2)$, so $c_{ji} = -c_{ij}$. Expanding the square,

$$D = \sum_{i \neq j} c_{ij}^2 + 2 \sum_{\{i,j,k\}} (c_{ij}c_{jk} + c_{ji}c_{jk} + c_{ki}c_{kj}).$$

For each unordered triple, the expression in parentheses is -1 . This follows from $\cot(a+b) = (\cot a \cot b - 1)/(\cot a + \cot b)$, or from clearing denominators in that identity. Since $\cot^2 = \csc^2 - 1$,

$$D = Q - N(N-1) - 2\binom{N}{3} = Q - \frac{N(N^2-1)}{3}.$$

No averaging, randomness, or spacing approximation has been used. In particular, an expectation of D contains only two-point information once the singular test kernel is specified.

4. CUE expectation by a legitimate integration by parts

The divergence of V is

$$\sum_i \partial_{\theta_i} V_i = -\frac{1}{2} Q.$$

For $\beta=2$, integration by parts against p_2 therefore yields

$$E_{\text{CUE}} D = \frac{1}{2} E_{\text{CUE}} Q.$$

This particular integration by parts is legitimate. Near an isolated collision with gap g , the density vanishes quadratically, while V has at most a first-order pole. The boundary flux $p_2 V_i$ consequently vanishes. Alternatively, multiplying out the Vandermonde factors removes the apparent pole in $p_2 V_i$, and its derivative is integrable on the compact torus. The products $p_2 V_i^2$ and $p_2 Q$ are integrable as well, including at intersections of collision hyperplanes.

Combining this identity with $D=Q-CN$ gives $EQ=2CN$ and $ED=CN$. The same calculation for $\beta>1$ gives $ED=CN/(\beta-1)$, but no extension of the present task to other model families is needed.

Crucially, this integrability check does **not** justify applying the same argument to LD , which has a fourth-order gap singularity.

5. ACUE expectation from its exact discrete kernel

Let $M=2N$ and label the grid by x in $\mathbb{Z}/M\mathbb{Z}$, with $\theta x=2\pi x/M$. The rank- N projection kernel is

$$K(x, y) = \frac{1}{2N} \sum_{a=0}^{N-1} e^{2\pi i a(x-y)/(2N)}.$$

It has $K(x, x)=1/2$. For $d \neq 0 \bmod 2N$, define

$$w_d = \csc^2 \frac{\pi d}{2N}.$$

Then $K(x, x+d)=0$ for nonzero even d , while for odd d ,

$$|K(x, x+d)|^2 = \frac{w_d}{4N^2}.$$

Hence the two-site inclusion probability is

$$\rho_2(0, d) = \frac{1}{4} - \frac{\mathbf{1}_{d \text{ odd}}}{4N^2} w_d.$$

It follows that

$$E_{\text{ACUE}} Q = \frac{N}{2} \sum_{d=1}^{2N-1} w_d - \frac{1}{2N} \sum_{\substack{1 \leq d < 2N \\ d \text{ odd}}} w_d^2.$$

The needed trigonometric sums are

$$\sum_{d=1}^{M-1} \csc^2 \frac{\pi d}{M} = \frac{M^2 - 1}{3},$$

$$\sum_{d=1}^{M-1} \csc^4 \frac{\pi d}{M} = \frac{M^4 + 10M^2 - 11}{45},$$

and subtraction of the even sites gives

$$\sum_{d \text{ odd}} w_d^2 = \frac{N^2(N^2 + 2)}{3}.$$

Consequently

$$E_{\text{ACUE}} Q = \frac{N(N^2 - 1)}{2} = \frac{3}{2} C_N,$$

which proves $ED = CN/2$. The model convention is the ACUE described in [Tao's original post](#).

The cosecant identities can be derived from

$$\sum_{d=0}^{M-1} \cot(x + \pi d/M) = M \cot(Mx)$$

by differentiation and comparison of Laurent coefficients at zero. This gives exact algebraic identities, rather than an approximation of a continuum integral by a lattice sum.

6. The dissipative generator identity

Let $w_{ij} = \csc^2((\theta_i - \theta_j)/2)$. The Jacobian J_V of V is symmetric, with diagonal entries $-(1/2)\sum_{j \neq i} w_{ij}$ and off-diagonal entries $w_{ij}/2$. Therefore

$$LD = 2V^T J_V V = -\sum_{i < j} w_{ij} (V_i - V_j)^2.$$

This proves pointwise nonpositivity. The repulsive flow exists globally for each initially distinct configuration: the logarithmic Vandermonde potential

$$W = \sum_{i < j} \log |e^{i\theta_i} - e^{i\theta_j}|$$

satisfies $dW/dt = D/2 \geq 0$. Every summand is bounded above by $\log 2$, so a collision would force W to negative infinity, contradicting its lower bound $W(0)$. The vector field therefore remains away from its singular set along each fixed initial trajectory, and D decreases.

7. Why the CUE expected derivative is negative infinity

Restrict to configurations in which exactly one pair has gap $g \rightarrow 0$ and all other points remain uniformly separated from that pair and each other. This set has positive measure in the remaining coordinates. Then

$$V_i - V_j = \frac{4}{g} + O(g), \quad w_{ij} = \frac{4}{g^2} + O(1),$$

so

$$LD = -\frac{64}{g^4} + O(g^{-2}).$$

The CUE joint density on this region is a positive smooth factor times g^2 . Since $LD \leq 0$ everywhere,

$$E_{\text{CUE}}[-LD] \geq c \int_0^\varepsilon g^{-2} dg = +\infty.$$

Thus $E LD = -\infty$ in the extended sense for $N \geq 2$. There is no cancellation between positive and negative parts to rescue an ordinary finite expectation.

Moreover, it is possible to state the expectation-level consequence without illegally exchanging derivative and integral. Let Φ_t be this deterministic repulsive flow. Since $0 \leq D(\Phi_t) \leq D$ initially and ED is finite, Fatou's lemma applied to

$$\frac{D(\theta) - D(\Phi_t \theta)}{t} \geq 0$$

shows

$$\lim_{t \downarrow 0} \frac{ED(\Phi_t) - ED}{t} = -\infty.$$

This is an infinite right slope, not a finite derivative obtained by a formal interchange. The finite ACUE support allows differentiation of its finite weighted sum at $t=0$.

For $N=2$ there is a transparent check. If g is the labeled angle difference, its CUE density is $\sin^2(g/2)/\pi$ on $(0, 2\pi)$. Integrating only over $\varepsilon \leq g \leq 2\pi - \varepsilon$ gives

$$\begin{aligned} E[D; \varepsilon \leq g \leq 2\pi - \varepsilon] &= 2 - \frac{2}{\pi}(\varepsilon + \sin \varepsilon) \rightarrow 2, \\ E[LD; \varepsilon \leq g \leq 2\pi - \varepsilon] &= 8 - \frac{16}{\pi} \cot(\varepsilon/2) - \frac{8\varepsilon}{\pi} \rightarrow -\infty. \end{aligned}$$

8. Exact finite ACUE derivative

The explicit finite answer is

$$E_{\text{ACUE}} LD = -\frac{2N(N^4 - 1)}{15}.$$

Here is a derivation, including the ordering factors.

8.1 Reducing the derivative to at most three sites

Differentiating Q-CN directly and using a three-angle cotangent identity yields

$$LD = -4 \sum_{i < j} (w_{ij}^2 - w_{ij}) + \sum_{i < j < k} (w_{ij} w_{ik} + w_{ij} w_{jk} + w_{ik} w_{jk}). \quad (8.1)$$

For verification of the triple algebra, put $c = \cot((\theta_1 - \theta_2)/2)$, $a = \cot((\theta_2 - \theta_3)/2)$, $b = (ca - 1)/(c + a)$. Then

$$\begin{aligned} &(1 + c^2)a(b - a) + (1 + b^2)b(c + a) + (1 + a^2)a(b - c) \\ &= -\frac{1}{2}[(1 + c^2)(1 + a^2) + (1 + c^2)(1 + b^2) + (1 + a^2)(1 + b^2)]. \end{aligned}$$

The expression extends by continuity wherever the chosen cotangent-addition denominator vanishes but the original angles are distinct.

8.2 ACUE three-site probability

Among any three grid sites at least two have the same parity. Their off-diagonal kernel entry is zero, so the triangular product in the 3×3 determinant vanishes. Therefore

$$\rho_3(0, d, e) = \frac{1}{8} - \frac{1}{8N^2} [\mathbf{1}_{d \text{ odd}} w_d + \mathbf{1}_{e \text{ odd}} w_e + \mathbf{1}_{d-e \text{ odd}} w_{d-e}],$$

for d, e nonzero and distinct modulo $2N$.

Put $S_2 = \sum d w_d$, $S_4 = \sum d w_d^2$, $O_4 = \sum d \text{ odd } w_d^2$, $O_6 = \sum d \text{ odd } w_d^3$. The pair part of (8.1) has expectation

$$-\frac{M}{2}(S_4 - S_2) + \frac{M}{2N^2}(O_6 - O_4), \quad M = 2N.$$

By symmetry the three edge products in the unordered-triple term contribute equally. Its expectation is

$$\frac{M}{2} \sum_{d, e \neq 0, d \neq e} \rho_3(0, d, e) w_d w_e.$$

The factor is $M/2$: the ordered triple count gives $M/6$, followed by the three equal edge-product contributions.

For odd d , the useful even-site convolution is

$$\sum_{\substack{e \neq 0 \\ e \text{ even}}} w_e w_{d-e} = \frac{4N^2 + 5}{3} w_d - 3w_d^2.$$

To derive it, use

$$\csc^2 u \csc^2 v = \csc^2(u+v) [\csc^2 u + \csc^2 v + 2 \cot(u+v)(\cot u + \cot v)]$$

with $u = \pi e / (2N)$, $v = \pi(d-e) / (2N)$, and sum over nonzero even e . The shifted cotangent sum vanishes before its $e=0$ term is removed, because $\cot(\pi d/2) = 0$ for odd d . This gives the displayed convolution exactly.

It follows that the triple term has expectation

$$\frac{M}{16} [S_2^2 - S_4 - \frac{1}{N^2} (\frac{8}{3}(2N^2 + 1)O_4 - 8O_6)].$$

Adding the pair term gives

$$\text{ELD} = \frac{M}{16} [S_2^2 - 9S_4 + 8S_2 + \frac{16}{N^2} O_6 - \frac{16(N^2 + 2)}{3N^2} O_4].$$

Finally,

$$S_2 = \frac{4N^2 - 1}{3}, \quad S_4 = \frac{16N^4 + 40N^2 - 11}{45}, \quad O_4 = \frac{N^2(N^2 + 2)}{3},$$

$$O_6 = \frac{N^2(2N^4 + 5N^2 + 8)}{15}.$$

The last identity is the odd-site difference of

$$\sum_{d=1}^{M-1} \csc^6 \frac{\pi d}{M} = \frac{2M^6 + 21M^4 + 168M^2 - 191}{945},$$

obtained by further differentiation of the same cotangent partial-fraction identity. Substitution simplifies to $-2N(N^4 - 1)/15$, as claimed.

9. What this does and does not provide for the research programme

The finite mean separation is genuine: ED/N^3 tends to $1/3$ in CUE and $1/6$ in ACUE. But the mechanism is a singular two-point test. It is not a new independent three-point invariant disguised as a force square. Any laws that have the same full two-point marginal on the same finite grid must have the same ED. Thus this mean alone does not distinguish all the previous ACUE moment-fiber counterexamples.

The kernel csc^2 has a diagonal singularity and requires information beyond the usual bounded, limited-Fourier-support pair-correlation tests. A finite-grid trigonometric interpolant of that kernel is a different function away from the grid; using its CUE expectation would silently change the observable. A rigorous arithmetic use needs a specified regularization, an evaluable weighted zero sum, and uniform control of the error when the diagonal cutoff is removed.

The canonical ratio H_0''/H_0' at simple zeros is a possible arithmetic formulation of the force. Its squared mean involves inverse derivatives and singular near-pair behavior; the circle identities and CUE integration by parts do not evaluate that mean for ζ . Gamma-factor contributions, the chosen H-coordinate, and the canonical-product normalization would all have to be retained in a direct real-line formulation.

The infinite CUE initial slope is another constraint on a proposed comparison. One cannot posit a finite Taylor expansion of the expected energy around zero time and then compare its first coefficients with ACUE. Regularized observables or a nonanalytic short-time analysis would be necessary. No such arithmetic comparison is proved here.

10. Floating-point enumeration and algebra checks

`research/force-energy/force_energy.py` visits every N -subset of the $2N$ grid for $N=2,\dots,10$, using the weights

$$\mu(S) = (2N)^{-N} \prod_{k < j \in S} |e^{i\theta_k} - e^{i\theta_j}|^2.$$

The traversal is exhaustive; all trigonometric values, weights, and expectations are float64. It is therefore **not exact-arithmetic enumeration**. The proof of the exact formulas is in the preceding sections.

N	Subsets visited	E D, float64	E LD, float64
2	6	1	−4
3	20	4	−32
4	70	10	−136
5	252	20	−416
6	924	35	−1036
7	3432	56	−2240
8	12870	84	−4368
9	48620	120	−7872
10	184756	165	−13332

At $N=10$ the total probability differs from 1 by about 2.5×10^{-14} , and the maximum pointwise error in $D=Q-CN$ is below 6.9×10^{-13} . The script also checks the generator formula against a finite directional difference at four configurations, and evaluates the exact $N=2$ truncated integrals to exhibit the divergent derivative expectation. Its JSON and log are saved beside this report. These computations support normalization checks; they do not replace the integrability proof.

Current report 16:

force_energy_independent_review

Source: `research/reports/force_energy_independent_review.md`. SHA-256:

70ef4365fc5330a1a329e1a28ca94b4426deed6786ea8f05e016bcc083a2fc2f.

Independent mathematical review of the circular force–energy audit

Date: 2026-09-05. Reviewer: residual_gram agent, independently of the flow agent's derivation.

Reviewed source: `force_energy.md`, SHA256 526c5e7abd2c6437f3ec2d6cdcdfffb27db13a134af15c4ab17464487cbdf360c.

Verdict: accept the stated finite-CUE/ACUE identities and the infinite CUE right-slope conclusion, with the conventions and scope already stated in the source. No remaining mathematical gap was found in the reviewed arguments. This review does not establish novelty, any analogous zeta identity, or a consequence for AH/RH.

The most important independently checked conclusions, for radians and the deterministic repulsive drift $V_i = \sum_{j \neq i} \cot((\theta_i - \theta_j)/2)$, are

$$D = \sum_i V_i^2, \quad E_{\text{CUE}} D = \frac{N(N^2 - 1)}{3}, \quad E_{\text{ACUE}} D = \frac{N(N^2 - 1)}{6},$$

$$E_{\text{ACUE}} LD = -\frac{2N(N^4 - 1)}{15}, \quad \lim_{t \downarrow 0} \frac{E_{\text{CUE}} D(\Phi_t X) - E_{\text{CUE}} D(X)}{t} = -\infty \quad (N \geq 2).$$

1. Pointwise derivative and three–body reduction

With $w_{ij} = \csc^2((\theta_i - \theta_j)/2)$, the Jacobian of V has off-diagonal entries $w_{ij}/2$ and diagonal entries $-\sum_{j \neq i} w_{ij}/2$. It is a negative weighted graph Laplacian. Therefore

$$LD = 2V^\top J_V V = -\sum_{i < j} w_{ij} (V_i - V_j)^2. \quad (1)$$

The coefficient is exactly -1 for an unordered pair sum. This verifies both the sign and the time normalization.

Independently expanding LQ , with $Q = \sum_{i \neq j} w_{ij}$, gives the pair contribution

$$-4 \sum_{i < j} (w_{ij}^2 - w_{ij})$$

plus the three-body contribution

$$\sum_{i < j < k} (w_{ij} w_{ik} + w_{ij} w_{jk} + w_{ik} w_{jk}). \quad (2)$$

The cotangent-addition rational identity used to reduce each triple was checked exactly in two independent indeterminates with SymPy. Its residual is the zero rational function.

As a normalization check, an equally spaced $N = 3$ configuration has all $w_{ij} = 4/3$. Its pair contribution is $-16/3$, its triple contribution $16/3$, and its total derivative is zero as required by force balance.

2. Three-site inclusion probability and ordering factors

For the ACUE projection process on $M = 2N$ sites, the two-site formula is

$$\rho_2(0, d) = \frac{1}{4} - \frac{\mathbf{1}_{d \text{ odd}}}{4N^2} w_d.$$

Every triple of distinct sites contains a same-parity pair, so at least one off-diagonal kernel entry in its cyclic three-edge product vanishes. Expanding the 3×3 determinant therefore gives

$$\rho_3(0, d, e) = \frac{1}{8} - \frac{1}{8N^2} [\mathbf{1}_{d \text{ odd}} w_d + \mathbf{1}_{e \text{ odd}} w_e + \mathbf{1}_{d-e \text{ odd}} w_{d-e}]. \quad (3)$$

The factor $1/(8N^2)$ is correct: each off-diagonal squared modulus is multiplied by the remaining diagonal value $1/2$.

The unordered-pair contribution in (2) has expectation

$$-4 \frac{M}{2} \sum_{d \neq 0} \rho_2(0, d) (w_d^2 - w_d) = -\frac{M}{2} (S_4 - S_2) + \frac{M}{2N^2} (O_6 - O_4). \quad (4)$$

For triples, the sum of all three central edge products equals one half of the ordered central-vertex sum. Thus its expectation is

$$\frac{M}{2} \sum_{\substack{d, e \neq 0 \\ d \neq e}} \rho_3(0, d, e) w_d w_e. \quad (5)$$

This confirms the source's $M/2$ factor. It can alternatively be viewed as $M/6$ for ordered triples times three equal edge-product contributions. There is no extra factor of two or three left over.

3. Independent reconstruction of the parity convolution

Fix an odd d , put $\alpha = \pi d/(2N)$, and let $e = 2k$, $1 \leq k \leq N-1$. Set $u = \pi k/N$, $v = \alpha - u$. The identity in the source reads

$$\csc^2 u \csc^2 v = \csc^2 \alpha [\csc^2 u + \csc^2 v + 2 \cot \alpha (\cot u + \cot v)].$$

The four sums needed over nonzero even sites are

$$\sum \csc^2 u = \frac{N^2 - 1}{3}, \quad \sum \csc^2 v = N^2 - w_d, \quad \sum \cot u = 0, \quad \sum \cot v = -\cot \alpha.$$

Indeed $N\alpha$ is an odd multiple of $\pi/2$; the full shifted cosecant-square sum is N^2 and the full shifted cotangent sum is zero. Removing $k = 0$ yields the displayed subtractions.

Since $\cot^2 \alpha = w_d - 1$, the convolution is exactly

$$\sum_{\substack{e \neq 0 \\ e \text{ even}}} w_e w_{d-e} = \frac{4N^2 + 5}{3} w_d - 3w_d^2. \quad (6)$$

To check the complete triple correction without skipping the parity bookkeeping, its first two terms together are

$$2 \sum_{d \text{ odd}} w_d^2 (S_2 - w_d) = 2S_2 O_4 - 2O_6.$$

The third term has $d - e$ odd, meaning one of d, e is odd and the other even. By (6), it is

$$2 \sum_{d \text{ odd}} w_d \sum_{\substack{e=0 \\ e \text{ even}}} w_e w_{d-e} = \frac{2(4N^2 + 5)}{3} O_4 - 6O_6.$$

Adding them and using $S_2 = (4N^2 - 1)/3$ gives

$$\frac{8}{3}(2N^2 + 1)O_4 - 8O_6,$$

exactly as claimed. Combined with (4), (5), this yields the source's final intermediate expression

$$\text{ELD} = \frac{M}{16} [S_2^2 - 9S_4 + 8S_2 + \frac{16}{N^2} O_6 - \frac{16(N^2 + 2)}{3N^2} O_4]. \quad (7)$$

Exact symbolic substitution of the four trigonometric-sum polynomials into (7) simplifies to

$$-\frac{2N(N^4 - 1)}{15}$$

identically in the indeterminate N . The odd-site fourth- and sixth-power polynomials were separately checked as the differences of the full $2N$ -site and N -site sums.

4. CUE divergence and the actual expectation–level right slope

The CUE density has score V , and the integration by parts for the first energy moment is legitimate: multiplying the Vandermonde square by a single cotangent cancels the apparent pair pole, while multiplying it by V_i^2 or w_{ij} remains integrable. Equivalently, each pair density vanishes quadratically and the relevant first-moment singularities are at most quadratic. The compact polynomial form handles intersections of collision hyperplanes as well.

This reasoning is not valid for the derivative observable, which has a fourth-order pole. The source correctly does not reuse that integration by parts.

Choose a region where one pair has small positive gap g , all other points stay a fixed distance from this pair and one another, and the remaining coordinates range over a set of positive measure. Then

$$V_i - V_j = \frac{4}{g} + O(g), \quad w_{ij} = \frac{4}{g^2} + O(1).$$

The leading contribution to (1) is $-64/g^4$. Other pair terms have at most order g^{-2} . Since the density is a positive smooth factor times g^2 , and $LD \leq 0$ globally, the integral of $-LD$ diverges like $\int_0^\epsilon g^{-2} dg$. Hence

$$\text{E}_{\text{CUE}}[-LD] = +\infty.$$

This is a valid one-sided divergence, with no positive-negative cancellation issue.

For the expected trajectory, define

$$A_t(X) = \frac{D(X) - D(\Phi_t X)}{t}, \quad t > 0.$$

Global repulsive existence and (1) imply $A_t \geq 0$. Also $A_t(X) \rightarrow -LD(X)$ for every initially distinct configuration, and $ED(X) < \infty$. Applying Fatou along any sequence $t_n \downarrow 0$,

$$\liminf_n EA_{t_n} \geq E[-LD] = +\infty.$$

Since the sequence was arbitrary, $EA_t \rightarrow +\infty$ as $t \downarrow 0$. Moreover,

$$EA_t = \frac{ED(X) - ED(\Phi_t X)}{t}$$

is well-defined because $0 \leq D(\Phi_t X) \leq D(X)$. This proves the source's actual infinite negative right slope. It is not a formal derivative-integral interchange.

The $N = 2$ truncated integral formulas were also checked directly. Their coefficients agree with the chosen drift normalization.

5. Independent exact finite checks

The companion program `force_energy_review.py` uses exact rational arithmetic in cyclotomic quotient rings for $N = 2, 3, 4$. This is independent of the source's float64 enumeration and independent of its determinantal pair/triple expectation reduction.

At every subset, the program constructs the original force via

$$V_i = iB_i, \quad B_i = \sum_{j \neq i} \frac{\zeta^{s_i - s_j} + 1}{\zeta^{s_i - s_j} - 1}.$$

It evaluates the dissipative derivative directly as

$$LD = \sum_{i < j} w_{ij} (B_i - B_j)^2.$$

The sign change comes from $\hat{\rho} = -1$; no numerical complex arithmetic is used. It then checks the pair/triple reduction pointwise and sums the original Vandermonde weights exactly.

N	Subsets	Exact ED	Exact ELD
2	6	1	-4
3	20	4	-32
4	70	10	-136

All pointwise identities, normalizations and expectations passed. The all- N conclusion remains based on the analytic derivation and symbolic polynomial identities, not extrapolation from this table.

Evidence files in `research/force-energy/`:

- `force_energy_review.py` — exact symbolic and cyclotomic checks.
- `force_energy_review_results.json` — results.
- Dependency: `dynamic-generator/generator_audit.py`, only its small CyclotomicRing implementation. Importing it does not run its main audit.

Reproduction:

```
OPENBLAS_NUM_THREADS=1 python3 research/force-energy/force_energy_review.py
```

6. Accepted scope and limits

The accepted result is an exact finite-ensemble distinction for a singular force observable, plus a nonanalytic initial expectation response in CUE. It is compatible with the independently proved persistence of protected trace moments: D and LD are not in that bounded Fourier-weight polynomial algebra.

The source correctly identifies that D collapses to a singular two-point kernel. The mean of LD , after reduction, involves up to three points, but still with singular kernels. No bound for the corresponding arithmetic zero sums has been supplied.

No additional caveat is needed for differentiating the ACUE average: after removing the irrelevant global rotation it is a finite sum over distinct configurations. A finite derivative at zero follows from the smooth local flow at each of those configurations.

The stated observations do not prove dynamical low-mode leakage, AH failure, RH, or a new prime-gap bound. They also do not establish that the finite-ensemble formulas are previously unknown. Within those limits, this review accepts the report.

Integration note: the reviewed-source SHA identifies the original staging report. The public copy changes only evidence paths; the preserved local source and Git manifest identify both versions.

Current report 17: new_attachment_bridge_audit

Source: `research/reports/new_attachment_bridge_audit.md`. SHA-256:
`169335d3e1b46a043e69e36032e51544a5923cdb43cce47f6d1eb07c8fc587be`.

Audit of the two newly supplied heat–depth bridge manuscripts

Date: 2026-09-05. Author: internal Astra agent `prime186`. Scope: a bounded source and proof audit, with two finite-dimensional counterexamples. This is an audit of claims, not a claim to have refuted the Riemann-zeta Alternative Hypothesis.

1. Outcome, including a correction to our own previous audit

The finite-dimensional first-collision bridge is useful and survives with precise hypotheses. Two stronger implications in the supplied reasoning fail: the pair that first collides need not have the initially smallest gap, and inverse-Gram blowup need not produce blowup of a marked heat-depth derivative. Both failures are demonstrated below with explicit formulas; the first also has an exact rational Sturm certificate.

Our previous source judgment also needs correction. It is not correct to say that the Lagarias – Rodgers hard-core parameter and the number 0.606894 have no connection in the primary literature. Their paper *Band-limited mimicry of point processes by point processes supported on a lattice*, § 5, printed p.19, explicitly defines T_1 and its hard-core supremum (μ). It proves the lower bound through its earlier construction and proposes that a modification of the Carneiro et al. pair-correlation method should give the upper bound 0.606894. That proposed modification is not proved there. The incoming second attachment largely preserves this distinction in its opening discussion; passages calling the bound a published, proved theorem do not. [Primary PDF](#), [arXiv record](#).

The correct source ledger is therefore:

Object or assertion	Correct status
The class T_1 and hard-core supremum (μ)	Explicitly defined in the 1907.03391 paper, § 5
$\mu \geq 1/2$	Supported by their constructed mimicking process
$\mu \leq .606894$	They explicitly suggest a modified existing method should establish it; no proof is supplied in that passage
$\mu = 1/2$	A possibility they propose, not a theorem
T_1 consists by definition of stationary processes	Inaccurate: the stated class is all uniformly locally moment-bounded processes mimicking the sine process at bandwidth 1; stationarity is not an additional condition in that definition
The .606894 footnote in the different paper 1905.12123	A positive-proportion small-gap result for zeta zeros, with a different quantifier

The source error arose from inspecting the higher-correlations paper while overlooking the companion paper's § 5. The authoritative handoff should correct that error explicitly, rather than retain an unjustified criticism of the incoming manuscript.

2. Inputs and evidence

The two audited inputs are:

1. `research/incoming/heat_depth_classification_proposal.md`: the finite heat polynomial, single-defect ACUE model, nonlinear transversality, marked Gram directions, and proposed dynamic separation programme.
2. `research/incoming/hard_gap_depth_bridge_proposal.md`: the claimed hard-core/heat-depth bridge and proposed arithmetic consequences.

Reproducible evidence in `research/bridge-audit/` (the primary PDF and text are retained in the adjacent local archive, with the hash below):

- `attachment_bridge_checks.py`: exact rational Sturm and symbolic derivative checks.
- `attachment_bridge_checks.json`: full heat polynomial, discriminant coefficients, rational intervals, root counts, symbolic formulas, and runtime.
- `lagarias-rodgers-bandlimited-1907.03391.pdf`: downloaded from the NSF primary-author manuscript URL above, 544836 bytes.
- `lagarias-rodgers-bandlimited-1907.03391.txt`: `pdftotext -layout` extraction. Lines 1058 – 1064 contain the decisive § 5 passage; the PDF page is printed p.19 and is also the nineteenth PDF page.

PDF SHA256: `a638b6be50dd919d1866085be29836f65ed4ce50f7caac14ae372b190b82e315`.

Historical finite-enumeration evidence is in `historical/riemann-rmt/impostors_paper.md`, especially its table around lines 183 – 208. That historical draft itself still contains some superseded source claims elsewhere; it is evidence for the archived calculations, not an authority that overrides primary sources.

3. Fix the sign convention before comparing formulas

For a degree- N polynomial with all roots on the circle, the supplied convention is

$$P_t(z) = \sum_{k=0}^N a_k e^{-tk(N-k)} z^k.$$

Increasing t is the repulsive, real-root-preserving direction after converting to circle angles. Let $D = -\Lambda \geq 0$ denote the time traveled in the opposite direction from $t = 0$ to the first loss of all-circle-rootedness. Thus $\mathcal{S} = -t \geq 0$ is the attractive direction. After centering the trigonometric polynomial by $e^{-iN\theta/2}$, the common scalar factor can be discarded and evolution in $\mathcal{S}$ is ordinary scalar heat $e^{\mathcal{S}\partial_\theta^2}$.

For a finite real polynomial the corresponding convention is

$$H_{\mathcal{S}}(x) = e^{\mathcal{S}\partial_x^2} P(x), \quad \frac{dx_j}{ds} = -2 \sum_{k \neq j} \frac{1}{x_j - x_k}.$$

The real-line ODE displayed in the second attachment uses $t = -\mathcal{S}$, so its sign is correct. The two conventions must not be combined without reversing time.

The exceptional circle clocks have $P(z) = z^N - c$, up to scale and rotation. All intermediate coefficients vanish, so this flow leaves their roots fixed. Their depth is $D = +\infty$, equivalently $\Lambda = -\infty$. They are not generic finite collision points of a discriminant hypersurface.

4. The exact real–line identity: valid for the colliding pair

Assume all the following on an interval $-D < t \leq 0$:

- The roots are real and simple and can be labeled in increasing order.
- One adjacent pair $x_j(t) < x_{j+1}(t)$ collides at $t = -D$.
- The finite sum below is defined, or in an infinite model the principal-value force and the difference of forces are justified and the displayed screening sum and integral converge.

Put

$$d(t) = x_{j+1}(t) - x_j(t), \quad S(t) = \sum_{i \neq j, j+1} \frac{1}{(x_i - x_{j+1})(x_i - x_j)} \geq 0.$$

Positivity follows because every other root lies outside the adjacent interval. Subtracting the two repulsive root equations gives

$$d' = \frac{4}{d} - 2dS, \quad (d^2)' = 8 - 4d^2S.$$

Integration from the collision time to zero gives the exact identity

$$\boxed{D = \frac{d_*^2}{8} + \frac{1}{2} \int_{-D}^0 d(t)^2 S(t) dt}, \quad d_* := d(0).$$

If m_0 is the minimum of all initial gaps, the legitimate consequence is

$$D \geq \frac{d_*^2}{8} \geq \frac{m_0^2}{8}.$$

One may write an exact identity with m_0 , but then its nonnegative correction has **two** terms:

$$D = \frac{m_0^2}{8} + \underbrace{\frac{d_*^2 - m_0^2}{8}}_{\text{initial pair-selection correction}} + \underbrace{\frac{1}{2} \int_{-D}^0 d(t)^2 S(t) dt}_{\text{screening correction}}.$$

The manuscript's replacement of d_* by m_0 while retaining only the screening integral is false in general. It is not a harmless choice of notation.

5. Exact finite counterexample: the smallest gap does not collide first

Take

$$P(x) = x(x^2 - 1)\left((x - 100)^2 - \frac{121}{400}\right), \quad H_s = e^{s\partial_x^2} P.$$

Its roots are

$$-1, \quad 0, \quad 1, \quad 100 - \frac{11}{20}, \quad 100 + \frac{11}{20}.$$

The two minimum gaps are 1. The remote pair has the strictly larger gap (11/10).

The heat polynomial is exactly

$$H_s(x) = x^5 - 200x^4 + \left(\frac{3999479}{400} + 20s\right)x^3 + (200 - 2400s)x^2 \\ + \left(-\frac{3999879}{400} + \frac{11998437}{200}s + 60s^2\right)x + 400s - 2400s^2.$$

The attached script performs exact rational Sturm calculations, not floating-point root classification:

1. The degree-10 discriminant in s is squarefree.
2. It has no root in $0 < s < 151/1000$, and exactly one root in $151/1000 < s < 19/125$.
3. That root has the exact rational enclosure

$$\frac{12999}{85936} < D < \frac{53965}{356761},$$

or $0.15126373114876188 < D < 0.15126373118137912$.

4. At $s = 0$ and $s = 151/1000$, there are three real roots in $([-2,2])$ and two in $([98,102])$.
5. At $s = 19/125$, the first interval still contains three real roots, while the second contains none; there are exactly three real roots in total.
6. At each of the four interval endpoints, H_s has no zero for $0 \leq s \leq 19/125$.

A real polynomial with fixed nonzero leading coefficient cannot lose a simple real root to a nonreal conjugate pair without a real multiple-root event. The discriminant counts exclude earlier events; the interval boundary checks prevent exchanging the two root groups. Therefore the unique first event is collision of the pair initially at $100 \pm 11/20$, not either initial minimum pair.

For this example the missing term in the incorrect identity is

$$\frac{(11/10)^2 - 1}{8} = \frac{21}{800} = 0.02625.$$

The remote pair's screening correction is small and positive: $(D-121/800)$ is about 1.3731×10^{-5} . This is consistent with the exact identity, while the initially denser three-root cluster postpones its own collision.

There is a simple structural reason the example exists. As the remote center $R \rightarrow \infty$, the local heat flow of $x(x^2 - 1)$ reaches its triple collision at $s = 1/6$, whereas the isolated gap (11/10) collapses at $s = 121/800 < 1/6$. The rational example above certifies this effect at the finite choice $R = 100$.

This is a theorem about finite scalar heat flow. It invalidates a universal minimum-first premise; it does not contradict the valid lower bound $D \geq m_0^2/8$.

6. The circle bound, its domain, and the ACUE consequence

For an adjacent circle gap Δ , the repulsive equation is

$$\Delta' = 2 \cot(\Delta/2) - \sin(\Delta/2)B(t), \quad B(t) \geq 0.$$

Here B is the sum of reciprocal products of the two sine factors involving each other root. With the other roots outside the lifted adjacent arc, these products are positive.

For $0 \leq \Delta < \pi$, define

$$q(\Delta) = -\log \cos(\Delta/2).$$

Then

$$\frac{d}{dt}q(\Delta(t)) = 1 - \frac{\sin^2(\Delta/2)}{2 \cos(\Delta/2)} B(t) \leq 1.$$

For a tracked pair that collides in the attractive direction, its gap stays below π throughout that part of its history. Indeed, in attractive time, the gap derivative at or above π is nonnegative, so a gap cannot cross down through π on its way to zero. Under the usual finite simple-root hypotheses, this gives

$$D \geq q(\Delta_*) \geq q(\Delta_{\min}(0)).$$

Writing q as a real function outside this domain is invalid. The antipodal $N = 2$ clock is the limiting infinite-depth case, not a finite value obtained by extending the logarithm through negative cosine. All circle clocks have infinite depth, even when their minimum gap is less than π .

For $m_N = N\Delta_{\min}/(2\pi)$, the valid bound is

$$D \geq -\log \cos(\pi m_N/N) \geq \frac{\pi^2 m_N^2}{2N^2}.$$

The last inequality follows from $-\log \cos x \geq x^2/2$ on $[0, \pi/2)$. It is an exact lower bound, not just an asymptotic one. For the corresponding microscopic time $s_{\text{micro}} = N^2 s/(4\pi^2)$, it reads $D_{\text{micro}} \geq m_N^2/8$.

A simple ACUE configuration is an N -element subset of the $(2N)$ -th roots of unity, so its angular gaps are at least π/N . A nonclock configuration has at least one gap exactly π/N : otherwise all N integer lattice gaps would be at least 2 and sum to $(2N)$, forcing a clock. Thus, for nonclock configurations,

$$D \geq -\log \cos \frac{\pi}{2N}, \quad N^2 D \geq \frac{\pi^2}{8}.$$

Clocks satisfy the same lower bound with $D = \infty$. Consequently

$$N^{8/3} D_{\text{ACUE}} \geq (\pi^2/8) N^{2/3} \rightarrow \infty$$

deterministically in the extended sense. This separation does **not** require any conjecture about an ACUE limiting median or limiting distribution. It is a particularly clean part of the finite-model programme worth retaining.

7. What the number 1.419640342 does and does not establish

The single-defect family

$$P_N(z) = \frac{(z-1)(z^N+1)}{z - e^{-i\pi/N}}$$

is a legitimate explicit configuration. Its proposed scaled heat profile and the double-zero equations supply a useful finite-dimensional/continuum calculation. The archived candidate critical value is $s_* = 1.419640342 \dots$, with location $u_* = 1.812942145 \dots$

The first attachment's inference from the $N = 2, \dots, 7$ median table to a typical ACUE limit at that constant is unsupported. Later complete finite enumeration in the archived notes reports conditional medians approximately

$$1.41822, \quad 1.41520, \quad 1.41277 \quad (N = 8, 9, 10),$$

after the median turns at $N = 7$. This invalidates the claimed monotone convergence evidence and the identification of the finite median with the single-defect branch. **Finite data alone do not prove that the eventual limiting median cannot equal the same number.** The stronger statement in one historical note that it is definitively not the limiting median goes beyond these finite data unless supplemented by another argument.

To turn the single-defect candidate into a theorem one must establish uniform convergence of the scaled heat profile and its relevant derivatives, nondegeneracy of the limiting double root, and the absence of an earlier collision elsewhere. Solving the double-zero equations numerically establishes none of the required global first-event exclusions by itself.

A limiting law for $N^2 D_{\text{ACUE}}$, its tightness, and the location of its support are separate ensemble questions. In particular, a deterministic special family does not automatically put its limiting value in the support of the random ensemble's limiting law: its probability can vanish too fast. At finite N , the clocks have positive probability and infinite depth, so the unconditioned expectation of D is infinite whenever their ensemble weights are positive.

The archived generic CUE depth argument is stronger than merely a repulsion heuristic. Its key sufficient hypothesis is an isolated-small-gap condition such as $\delta A \rightarrow 0$, where A controls the reciprocal chord distances of the other roots from the close pair. Under that hypothesis one obtains $D \sim \delta^2/8$. A repulsion exponent or the size of the smallest gap alone does not prove this conclusion; the three-root cluster in § 5 illustrates why background control matters.

8. Nonlinear transversality: a correct conditional theorem

The first attachment's differential idea can be stated precisely. Work on a smooth finite-dimensional manifold of admissible polynomial coefficients, use a real phase-normalized heat polynomial $(Q(s, x, a))$, and assume:

1. At $a = a_0$, there is a unique first collision at finite D_0 and real location x_0 .
2. $Q = Q_x = 0$, while $Q_{xx} \not\equiv 0$ and $Q_s \not\equiv 0$, at this collision.
3. The other roots remain simple there; competing collision times and any degree-loss or exceptional events are excluded in a neighborhood.
4. The moment constraint map $(M(a))$ has constant rank locally, so its level set is a smooth manifold.

The Jacobian of (Q, Q_x) with respect to $((s, x))$ is then invertible. The implicit function theorem gives a smooth local collision time and

$$dD(v) = -\frac{Q_a(v)}{Q_s}.$$

If there exists a tangent vector $v \in \ker dM$ with $Q_a(v) \not\equiv 0$, the collision time is nonconstant on that moment fiber. The equivalent discriminant derivative formula requires a smooth defining discriminant branch and a nonzero derivative in the heat direction. It fails at simultaneous or higher-order collisions unless separately resolved.

The key condition is existence of that transverse v . It cannot be inferred merely because the hitting time is nonlinear, because the coefficient heat rates are large, or because the moment fiber has positive dimension. Rotations provide a simple null direction for every rotation-invariant statistic. Any tomography claim must first quotient such symmetries and show a full-rank Jacobian on the remaining tangent space. Multiple roots, clock points, and switches between competing first collision times are natural nondifferentiable exceptions.

Likewise, if a one-parameter family opens a single isolated double root with initial gap $\delta(\varepsilon) = c\varepsilon + o(\varepsilon)$ while all other roots remain uniformly separated, the local heat normal form gives $D = c^2\varepsilon^2/8 + o(\varepsilon^2)$. This is a valid local theorem under isolation. If a third root approaches on the same scale, that coefficient is not automatic.

9. Exact marked counterexample: inverse blowup with bounded susceptibility

For $0 < \varepsilon < 1/2$, take

$$G_\varepsilon = \text{diag}(\varepsilon, 1 - \varepsilon), \quad \tilde{G}_\varepsilon = \text{diag}(1 - \varepsilon, \varepsilon), \quad u = e_1.$$

For each fixed ε , these matrices are isospectral and have identical rank, trace, and Hilbert – Schmidt norm. Their marked inverse quantities are

$$u^* G_\varepsilon^{-1} u = \varepsilon^{-1} \rightarrow \infty, \quad u^* \tilde{G}_\varepsilon^{-1} u = (1 - \varepsilon)^{-1} \rightarrow 1.$$

The Hilbert – Schmidt norm varies along ε ; it is identical **within each pair**, not constant along the entire path. No claim of a fixed invariant fiber across all ε is being made.

Apply the Cayley map $U = (G - iI)(G + iI)^{-1}$. For ordered eigenvalues $0 < a < b$, the smaller angular separation is

$$\Delta = 2(\arctan b - \arctan a) < \pi.$$

For $N = 2$, the heat threshold is exactly

$$\Lambda(a, b) = \log \cos(\Delta/2) = \log \frac{1 + ab}{\sqrt{(1 + a^2)(1 + b^2)}}.$$

Under $G \mapsto G + \eta uu^*$, marking the smaller and larger eigenvalues respectively gives

$$\chi_a = \frac{\partial \Lambda}{\partial a} = \frac{b - a}{(1 + ab)(1 + a^2)}, \quad \chi_b = \frac{\partial \Lambda}{\partial b} = -\frac{b - a}{(1 + ab)(1 + b^2)}.$$

At $a = \varepsilon, b = 1 - \varepsilon$,

$$\chi_a \rightarrow 1, \quad \chi_b \rightarrow -\frac{1}{2}.$$

The attached script checks these derivatives and limits symbolically. Thus marks distinguish the two isospectral orientations, but inverse blowup does not force marked-depth susceptibility blowup. If susceptibility is instead defined using $D = -\Lambda$, the two signs reverse and boundedness is unchanged.

The rank-one determinant lemma and eigenvalue perturbation formula in the first attachment remain correct. Their existence does not establish the proposed inverse-Gram-to-depth blowup implication. A more restrictive claim along a path with all additional invariants fixed would require its own hypotheses and example; it is not established or refuted by misdescribing the present path.

10. What is missing at the actual–zeta bridge

The arithmetic conclusion does not follow from the finite first-collision identity alone.

Under RH, the classical de Bruijn – Newman constant satisfies $\Lambda \leq 0$; the Rodgers – Tao theorem supplies $\Lambda \geq 0$, hence $\Lambda = 0$. Therefore at every negative time the genuine global H_t already has some nonreal zeros. An argument on $(-D, 0]$ that requires **all** zeros to remain real cannot simply be reused for a high local window of that global flow. [Rodgers – Tao primary paper](#).

The problem is concrete: contributions from nonreal background zeros are not a sum of the nonnegative real screening terms in § 4. A finite polynomial formed from the zeros in a window has a well-defined finite heat flow, but it is a different analytic function, whose time evolution need not agree with the restriction of H_T . Uniform control of that difference, including its derivatives on the shrinking time scale, is a missing theorem.

Before using a local depth D_T , define the object and prove at least the following: how roots are tracked in a window; what background or renormalization is retained; how the heat evolution is inherited from the genuine H_T ; what constitutes the first local event; and what error controls replace the all-real positivity argument. An arbitrary stationary point process does not automatically determine an admissible entire function or a unique heat deformation either.

The scale conversion in the second attachment is algebraically correct **conditional on such a bridge**. In the usual Rodgers – Tao $X = 2Y$ convention, local density is asymptotic to $\rho_T = (\log T)/(4\pi)$. A unit-density hard core of $(1/2)$ would give microscopic depth at least $(1/32)$, corresponding to

$$-t \geq \frac{\pi^2}{2(\log T)^2}$$

to leading order. Correct units do not supply the missing analytic comparison.

The strong AH in Lagarias – Rodgers 1905.12123, Conjecture 2.2, uses limiting consecutive-gap values in $\{1/2, 1, 3/2, \dots\}$, with no zero and for all sufficiently large indices. A theorem producing $\liminf g_n < 1/2$ would refute that formulation. Statements allowing a density-zero exceptional set require a positive-density or otherwise appropriately quantified contradiction; formulations allowing multiplicity need separate handling. The historical .50412 small-gap claim mentioned in the 2019 source was withdrawn and should not be revived. [AH primary paper](#), [withdrawn Goldston – Turnage-Butterbaugh record](#).

A depth theorem about a function-field family and an anti-half-lattice conclusion for that family do not refute AH for the actual Riemann zeta function. Such a theorem may be a substantial independent random-matrix/arithmetic-family result. Its quantifiers concern the chosen family, monodromy group, genus, and field-size limits. Connecting it to actual zeta ordinates is additional work.

11. Claim-by-claim replacement ledger

Supplied claim	Verdict and replacement
First heat collision is a discriminant event	Valid at finite nonexceptional degree, with clock/infinite-depth and multiple-event qualifications
$D = m_{\min}^2/8 + \frac{1}{2} \int \mathcal{L} S$	False as a universal identity; use the initial gap of the pair that actually collides, or add the pair-selection correction
$D \geq m_{\min}^2/8$ for finite real scalar heat	Valid under the real-root flow assumptions
Circle $D \geq -\log \cos(\Delta_{\min}/2)$	Valid with the real logarithm domain and extended infinite clock case handled
ACUE $N^2 D \geq \pi^2/8$	Valid, deterministic lower bound
1.419640342 is the limiting ACUE median	Not established; early finite evidence was misleading; single-defect candidate is a separate object
A deterministic special family places its limit in the ensemble support	Not automatic; requires nonvanishing probability/control in neighborhoods

Supplied claim	Verdict and replacement
$D \sim \delta_{\min}^2/8$ from a repulsion exponent alone	Insufficient; needs extreme-gap and background-isolation estimates
Heat-rate amplification establishes transversality	No; an actual tangent witness and collision nondegeneracy are needed
Marked inverse blowup forces susceptibility blowup	False in general; exact $N = 2$ counterexample above
Low-band agreement remains invisible for all heat observations	Only if the specified observable class is invariant under the relevant evolution; nonlinear hitting-time observations require separate analysis
A local true-zeta depth automatically obeys the finite screening identity	Not defined/proved; negative-time nonreal background prevents the naive import
A function-field depth law refutes actual-zeta AH	Invalid transfer of quantifiers
LR never relates hard-core μ to .606894	Our earlier source judgment was wrong; § 5 of the companion paper explicitly proposes that relation
LR proves the .606894 hard-core bound in that passage	Also wrong; the passage proposes a method modification

12. Reproducibility and next proof obligations

Run:

```
python3 research/bridge-audit/attachment_bridge_checks.py
```

Environment used: Python 3, SymPy 1.14.0. The saved run completed in about 0.13 seconds. The JSON contains exact rational data, not merely decimal approximations. The proof relies on standard exact polynomial arithmetic, Sturm's theorem, real-root continuation, and the explicitly stated symbolic identities. This is an exact computer-algebra certificate; it is not a proof-assistant formalization.

The most useful retained directions are the isolated-gap heat theorem, its rigorously controlled CUE consequences, and a precisely defined arithmetic residual/flow comparison. Immediate follow-up should target a theorem that supplies new arithmetic information or the missing true-zeta local-flow control. A universal near-collision theorem based solely on low-band mimicry would run into the already constructed half-lattice process unless an additional hypothesis excludes it.

Postpone claims of actual-zeta AH refutation, support beyond the currently proved arithmetic range, marked inverse blowup, and an ACUE limiting median until their missing hypotheses and estimates are supplied. None is made by this audit.

Current report 18: prime186_structural_frontier

Source: `research/reports/prime186_structural_frontier.md`. SHA-256:
`2e1e19544854af0375b3b114c2e38448f8841cf70680690e1b1c7749f3800403`.

The 186 structural mechanism and a finite complementary–support frontier

Date: 2026-09-05. Scope: primary-source reconstruction, an exact finite arithmetic example, a proved allocation frontier, and a concrete next obligation for $k = 39$. No smaller prime-gap bound is claimed.

1. What is being transferred

The important mechanism in the 186 paper is an enlargement of the **admissible coefficient support for existing restricted-modulus distribution estimates**. It is not a replacement of the whole arithmetic input by a larger scalar exponent of distribution.

The sources used here are *Improved short gaps between primes*, especially Definition 2.1, Lemma 2.2, Proposition 2.3, § 4.2, and Proposition 4.6; its [research record](#), especially pp.25 – 27; and the [official certificate repository](#) at the locally inspected commit `61340d0b74163003b32756bb16e91d9209a5e330`.

The user's [Weijie Su post](#) could not be retrieved in this run: X returned HTTP 403. Its wording is therefore not quoted or independently authenticated here. The mathematical mechanism below is reconstructed directly from the supplied primary paper.

2. The actual combined divisor, including shared primes

For $Y \geq 1$, order- r dense divisibility asks that for every allocation $j + k = r - 1$ and every target $1 \leq U \leq Ym$, one can write

$$m = uv, \quad U/Y \leq v \leq U, \quad u \in D^{(j)}(Y), \quad v \in D^{(k)}(Y).$$

The allocation is universal: finding one convenient factorization is not sufficient. The prime-factor criterion used by the paper is

$$p' \leq YQ_{<p} \quad (p \mid Q, p > Y) \implies Q \in D^{(r)}(Y),$$

for squarefree Q , where $Q_{<p}$ is the product of the prime factors below p . This is a sufficient criterion; failure of it does not imply failure of dense divisibility.

In a mixed sieve pairing, the modulus is WQ , with

$$Q = [D, E].$$

Proposition 2.3 takes nondecreasing $f, g \geq 1$ with $f(p)g(p) = p^3$, fixed budgets $A_0 C_0 \leq XY$, and requires

$$\begin{aligned} f(p)D_{\geq p} &\leq A_0, & g(p) &\leq C_0 & (p \mid D, p > Y), \\ g(p)E_{\geq p} &\leq C_0, & f(p) &\leq A_0 & (p \mid E, p > Y). \end{aligned}$$

When $Q > X$, these imply triple dense divisibility of Q . To see the crucial point, fix $p \mid D$. If E has a prime $q \geq p$, take the smallest such q . Monotonicity yields

$$g(p)E_{\geq p} \leq g(q)E_{\geq q} \leq C_0.$$

If E has no such prime, the explicit opposite-root condition $g(p) \leq C_0$ supplies the same conclusion. Therefore

$$p^3 D_{\geq p} E_{\geq p} \leq A_0 C_0 \leq XY,$$

and

$$p^3 \leq \frac{XY}{Q_{\geq p}} = \frac{X}{Q} Y_{Q_{\leq p}} < Y_{Q_{\leq p}}.$$

The same argument handles a prime belonging only to E . It uses $Q_{\geq p} \mid D_{\geq p} E_{\geq p}$, not equality. Shared primes are allowed. Replacing the assumption $[D, E] > X$ by $DE > X$ would be invalid: a large gcd can make the actual modulus much smaller.

Deleting prime factors only decreases the relevant tails and maxima, so the root predicates remain valid for coefficient divisors. This is essential for the inverse Selberg transform. It is stronger bookkeeping than saying a chosen root happened to be densely divisible, because dense divisibility of a number does not automatically pass to every divisor with the same Y .

3. An exact integer example of complementary help

Take

$$Y = 10, \quad D = 330 = 2 \cdot 3 \cdot 5 \cdot 11, \quad E = 455 = 5 \cdot 7 \cdot 13, \quad Q = [D, E] = 30030.$$

Both roots contain a prime greater than Y , and they share the factor 5. Let

$$f(p) = p, \quad g(p) = p^2, \quad A_0 = 121, \quad C_0 = 2197, \quad X = 27000.$$

Then $A_0 C_0 = 265837 \leq 270000 = XY$, and $Q > X$. All four rootwise requirements hold: the only activated prime in D is 11, and the only activated prime in E is 13. For the merged modulus the cubic checks are

$$11^3 = 1331 \leq 10(2 \cdot 3 \cdot 5 \cdot 7) = 2100,$$

$$13^3 = 2197 \leq 10(2 \cdot 3 \cdot 5 \cdot 7 \cdot 11) = 23100.$$

The attached exact check also verifies Definition 2.1 directly by recursive divisor enumeration. It confirms $Q \in \mathcal{D}^{(3)}(10)$, but $E \notin \mathcal{D}^{(3)}(10)$. For the latter failure, take allocation $(j, k) = (2, 0)$ and target $U = 11$: no divisor $v \in [11/10, 11]$ leaves E/v doubly 10-densely divisible.

Thus complementary factors really can produce a suitable combined modulus even when one root itself lacks the required triple property. The calculation is finite and exact; it does not establish an asymptotic prime distribution estimate by itself.

4. Why a scalar exponent misses the new geometry

Write $R = x^{\rho}$, $u_p = \log_R p$, $s_D = \log_R D$, and $Y = R^{\xi}$. The logarithmic form of the order-three split uses nondecreasing functions

$$\phi_D(u) + \phi_E(u) = 3u,$$

with owner-tail statistic

$$H_{\phi,\xi}(D) = \max_{p \mid D, u_p > \xi} \left(\sum_{q \mid D, q \not\geq p} u_q + \phi(u_p) \right).$$

For a source band $Q > R^B$, outer and inner total radii S, T , and $a = B - T, b = B - S$, the predicates are

$$s_D \leq a \quad \text{or} \quad [H_{\phi_D,\xi}(D) \leq A, \phi_E(M_\xi(D)) \leq C],$$

$$s_E \leq b \quad \text{or} \quad [H_{\phi_E,\xi}(E) \leq C, \phi_D(M_\xi(E)) \leq A], \quad A + C \leq B + \xi.$$

The small-radius alternatives are genuinely safe: if $s_D \leq B - T$, the size of $[D, E]$ cannot exceed the band threshold.

Otherwise each activated prime and its full inclusive tail matters. Two roots with the same total logarithmic size can have different admissibility because their prime partitions differ.

Distribution estimates still restrict both the modulus level and its smoothness/dense-divisibility parameter. For example, the full-prime order-three condition in Corollary 2.19 is $240\omega + 80\delta < 3$, with modulus level $\chi^{1/2+2\omega}$ and $Y = \chi^\delta$. The source ladder trades these parameters across bands; increasing the level generally reduces the allowable δ . It is not legitimate to retain a high level and simultaneously borrow the less restrictive support from a lower-level row.

The enlarged minorant mixed pairing has physical product bound $\chi^{5252997}$, while the full-prime mixed pairing uses the slightly smaller range. The minorant's negative part, inner-square pairings, and support restoration remain separate debts. A scalar calculation resembling $M_k > 2/\theta$ on an unrestricted simplex omits exactly the fragment geometry that makes this improvement work.

5. A proved finite frontier for every monotone allocation

Here is an inexpensive certificate useful before recomputing physical integrals.

Fix one active order-three row and owner ceilings A, C . Suppose an activated outer root can have largest fragment u , and an activated inner root can have largest fragment v , with $u \geq v > 0$. Their remaining mass may consist of smaller fragments or the smooth seed. Because there is no larger fragment than a largest witness, the necessary owner and opposite-root constraints include

$$u + \phi_D(u) \leq A, \quad 3u - \phi_D(u) \leq C,$$

$$4v - \phi_D(v) \leq C, \quad \phi_D(v) \leq A.$$

Monotonicity gives $\phi_D(v) \leq \phi_D(u)$. Two explicit nonnegative-sum identities give the frontier:

$$\boxed{4u \leq A + C, \quad u + 4v \leq A + C.}$$

Indeed,

$$A + C - 4u = (A - u - \phi_D(u)) + (C - 3u + \phi_D(u)),$$

$$A + C - u - 4v = (A - u - \phi_D(u)) + (C - 4v + \phi_D(v)) + (\phi_D(u) - \phi_D(v)).$$

These are exact dual certificates, requiring no numerical optimizer. They apply to active predicates, not to roots in the safe small-radius alternative. If $v \geq u$, the symmetric second inequality is $4u + v \leq A + C$.

The paper's capped allocation is

$$\phi_D(t) = \min(3t/2, L), \quad \phi_E(t) = 3t - \phi_D(t).$$

In its relevant plateau regime, the largest-fragment caps are

$$u = A - L, \quad v = (C + L)/4,$$

so $u + 4v = A + C$ exactly. The published caps already lie on a Pareto boundary for arbitrary monotone allocations with these fixed ceilings. No cleverer piecewise allocation can simultaneously enlarge both these largest-fragment allowances. It can redistribute them, change which smaller tails are admitted, change the budgets, or leave this particular support template. Those are distinct possibilities.

This is a largest-fragment obstruction, not a universal no-go theorem for improving the whole sieve. In particular, a frontier point can remain suboptimal for the weighted integral because the outer and inner losses are very unequal.

6. Exact location of the current frontier and a bounded next search

For the nonterminal tight rows generated by the official ladder recurrence, put

$$\rho = .262499, \quad \rho_* = .2624989, \quad e = 10^{-7}/\rho, \quad A = S + e/2, \quad C = T + e/2.$$

The recurrence makes these ceilings independent of the individual tight row. The base and enlarged radii T_0, T_1 differ, so the old and new ladders have different C . These formulas and all values below use exact rational inputs.

For the plateau family the extreme outer endpoint is

$$L_{\min} = (3A - C)/4, \quad u_{\max} = (A + C)/4, \quad v = 3(A + C)/16.$$

The plateau validity range reaches $L \leq 3C/5$ here. The selected $L = (23/40)C = .575C$ is close to, but above, the smallest permitted value:

Ladder	$L_{\min}/C$	Current normalized u	Current normalized v	Maximum physical outer- cap gain to $L_{\min}$	Corresponding physical inner- cap loss
Old prime	.57319344847	.49753017766	.37486755082	.00045147647	.00011286912
New minorant	.56962060366	.49514387013	.37650165272	.00135022875	.00033755719

Physical caps are ρ_* times normalized caps. These gains are allowances in prime-factor exponents, not gains in the final sieve quotient.

At mesh $h = S/98304$, downward rounding of the two caps produces frontier slack smaller than $5h$. Consequently, the small slack in the actual integer masks is a rounding allowance, not an unexplored macroscopic joint enlargement. The JSON checks the exact rounding slack, constructs a valid allocation at four points on each frontier, and certifies infeasibility when both exact caps are increased by one cell.

A sensible bounded search is to vary the old and new plateaus **independently** inside their valid intervals while keeping every source row and opposite-root condition. A common choice below .57319344847 is invalid for the tight old rows, even though it may be allowed in the new rows. Since the global outer domain intersects both ladders and the base inner domain intersects their predicates, local improvement of one row cannot be reported as enlargement of the entire usable support.

The historical loss ledger puts about 89.5% of the published weighted restoration loss in the outer effective-order-5/2 class, versus about 4.66% in the corresponding enlarged-inner class. That asymmetry motivates testing this tradeoff, but does not prove the trade is beneficial: the derivatives of the physical integrals are not determined by the aggregate loss totals. The

research record already reports experimenting with plateau allocations, so this proposal is a reproducible next test, not a novelty claim about discovering the plateau idea.

7. A more general finite–fragment allocation certificate

The frontier admits a practical extension beyond largest witnesses. Given finitely many candidate outer and inner fragment configurations, collect every activated witness size into ordered rational knots

$$0 = u_0 < u_1 < \cdots < u_n, \quad z_i = \phi_D(u_i).$$

Monotonicity of both owners is exactly

$$0 \leq z_{i+1} - z_i \leq 3(u_{i+1} - u_i), \quad z_0 = 0.$$

An outer witness with inclusive tail $T_D(u_i)$ imposes

$$z_i \leq A - T_D(u_i),$$

and an inner witness imposes

$$z_i \geq 3u_i + T_E(u_i) - C.$$

Opposite-root constraints give the additional lower or upper bounds at each root's maximum. Combine them into intervals $\ell_i \leq z_i \leq r_i$, including $z_0 = 0$. These are rational difference constraints. Their feasibility can be checked explicitly:

$$z_i^{\min} = \max_j \{ \ell_j - 3 \max(u_j - u_i, 0) \}.$$

There exists an admissible allocation iff $z_i^{\min} \leq r_i$ for every i . The maximum of the displayed nondecreasing 3-Lipschitz functions has those same properties and is the least feasible lower envelope. If a constraint fails, the offending pair (i, j) supplies the exact contradiction

$$\ell_j > r_i + 3 \max(u_j - u_i, 0).$$

If it succeeds, linear interpolation of the knot values, followed by a constant continuation for ϕ_D , defines globally nondecreasing complementary functions. This proves feasibility of the finite candidate set. It does not by itself determine which candidates have enough integral weight to be worth retaining.

`prime186_frontier_checks.py` implements this exact feasibility engine using only the Python standard library. A useful next optimization can choose which high-weight fragment patterns to retain, use this engine to produce a feasible allocation or an infeasibility witness, and then evaluate the resulting full support. This separates a cheap geometric obstruction from the expensive physical integral calculation.

8. The concrete $k = 39$ obligation

The next theorem is DHL[39,2], not merely another admissible tuple. A diameter-182 admissible 39-tuple was already checked in round 1, so that theorem would give $H_1 \leq 182$. The unresolved obligation is to find and certify a 39-variable trial satisfying the **complete** inequality

$$\rho_* (\mathcal{J}_{\lambda, H}^- - E_O^+ - d_0 \beta_{\text{old}}^+ - (1 - b_h) \beta_{\text{new}}^+) > l_H^+.$$

Here E_O pays for outer failures, while the two β terms pay for the actual inner domains, including their intersection. The negative full-face term from the prime minorant remains present. The source checks for all mixed and inner-square moduli must hold with the final cell geometry. Replacing 40 by 39 changes the measure, marginal dimension, combinatorial coefficients, inward cell masks, and the integral values; none of those values can simply be inherited from the $k = 40$ receipt.

A concrete next computation is:

1. Use $k = 39$, initially retaining the published physical radii, source ladders, minorant, and its justified constants.
2. Keep the published 77-dimensional symmetric polynomial descriptor family as a reproducible starting space, and vary the old/new plateau or the finite-knot owner function with explicit feasibility certificates.
3. Freeze valid positive cover and Young parameters when assembling a quadratic form. Optimize the **support-restored** form against its denominator. The 77 coefficients are not a substitute for source validity.
4. For a promising vector, round its coefficients and all geometric parameters rationally, restore every source row and inward cap, and compute outward enclosures for the full displayed inequality. Only a positive final interval proves the desired sieve criterion.

A negative-definite upper certificate could prove failure of one fixed finite trial space and fixed support family. It would not rule out DHL[39,2] or other coefficient families. Conversely, a cap-only value above one remains merely a candidate until the deletion costs are paid.

There is one explicit loss-recovery option already visible in Proposition 4.6. The displayed proof actually retains

$$(1 - \rho_* \mid b_h \mid C_{\text{op}})a, \quad a = \parallel (1 - P_O)F \parallel^2,$$

as a positive term before dropping it in the convenient sufficient criterion. A certified lower bound for a , for example from disjoint known-failure regions, would supply a legitimate credit. Existing upper bounds for covered failure masses cannot be reused as that lower bound. The possible gain is unmeasured here and is not advertised as sufficient for $k = 39$.

9. What was computed, and the exact remaining bottleneck

This round ran the exact integer example, the recursive definition check, eight rational frontier points, the cap-enlargement infeasibility witnesses, and the finite allocation-envelope engine. Runtime was about 0.0015 seconds. Files:

- `prime186_frontier_checks.py`
- `prime186_frontier_checks.json`
- this research note

No FLINT build, paid API call, or large numerical search was performed. The checks do not bypass the official certificate's signed-convolution startup regression: that certificate was not run in this round at all.

The bottleneck is the rigorous physical integral evaluation for the revised trial and its actual fragment support, not feasibility of the two-budget algebra. The published baseline involves 104 outer and 45 inner physical upper bounds plus cap bounds. Its official certificate specifies a corrected FLINT build. Geometry-only certificates can reject impossible support proposals and certify the allocation step, but they cannot supply those missing weighted integrals or the asymptotic distribution theorem.

The public Lean derivation remains conditional on its three documented project inputs, including the physical integral bounds. A rational support certificate verifies one component of the mathematical argument; it does not discharge those Lean inputs or create a new prime-gap theorem.

10. Structural lesson for the zeta programme

The transferable method is to constrain the **actual object produced by the mixed pairing**, distribute a fixed analytic budget between complementary factors, preserve the support under the algebraic transformations used in the proof, and charge the residual with a sign-correct quadratic inequality. In the prime problem the combined object is an lcm modulus and the needed property is dense divisibility. In a zeta mixed-moment problem, the corresponding object and the required distribution estimate must be identified anew.

The transfer does not authorize a larger Dirichlet support just because a similar factorization can be written down. One must prove that the actual mixed terms fall inside an available arithmetic estimate, or prove a new estimate for them. The centered-Gaussian pole correction from round 2 is an example of an additional term that survives a superficially favorable positivity argument. The useful frontier calculation above helps prevent spending research time on reallocations that cannot enlarge both sides even before those analytic issues are addressed.

Current report 19: fable_task001_sync_review

Source: `research/reports/fable_task001_sync_review.md`. SHA-256:
`3a14533ac2f8cf72ce605b6209b8139c71a7747f8e691b7a6e29397a30ca7ec4`.

Fable task 001: pinned-source synchronization and numerical review

Date: 2026-09-05.

Reviewed source commit: `a408e7050fffc74459b3c83fafa5ac03c8b7dea6`, in the old Alpha-devbox PR #11 repository. All statements about files being present or absent refer to this exact commit, not to later working state or an unseen running session.

Source directory: [Fable task001 at the reviewed commit](#).

Conclusion: Fable has demonstrably picked up the task and returned an independent F2 numerical check. Its committed continuum and finite-sum results agree with Astra's fixed-family certificate and finite operator conventions. No mathematical counterexample, missing leading term, positive half-gap margin, or completed Fable arithmetic-transfer proof is present in the reviewed snapshot. The final task001 report is still absent.

1. Receipt, provenance and missing deliverables

The reviewed `FABLE_COORDINATION.md` explicitly records pickup at 06:30 UTC from new-repository commit 97df092. The immutable input mirror's `SOURCE_COMMIT.txt` gives the full hash

```
97df092427a1035cf1c66dc712148ccccac09ac2
```

and says the copy was taken at 2026-09-05T06:26:51Z from the new repository's `codex/astra-research` branch.

The copied task text still says “published, awaiting explicit pickup.” That is a stale status in a historical input snapshot, not evidence that Fable never received the task. The coordination entry plus independent scripts, logs and numerical outputs demonstrate actual execution.

Present under `astra_tasks/task001/`:

- `f2_continuum.py`
- `f2_continuum_results.json`
- `f2_continuum_run.log`
- `f2_finite_sum.py`
- `f2_finite_sum_results.json`
- `f2_finite_sum_run.log`
- `f2_finite_sum_run_1e7.log`
- `f2_drift_fit.py`

Explicitly checked absent at the pinned commit:

- `astra_tasks/task001_report.md`, the promised final report;
- `astra_tasks/task001/f2_finite_sum_results_1e7.json`;
- `astra_tasks/task001/f2_drift_fit_results.json`.

The $1e7$ log contains only the initial float-sieve validation line. It is not evidence that an $L = 10^7$ finite-sum computation completed. The drift-fit script is present, but no committed output establishes that it completed. Neither unfinished item should be promoted into the returned-results table.

The committed `CLAIMS.md` has prime-gap certificate entries D1.2, D1.3, D1.4, D1.6, D1.7, D1.8, and no task001/F2 claim. In particular, it does not record a completed arithmetic transfer or an independent refuter verdict for this task.

The coordination file describes Fable's adversarial workflow and intended proposer/refuter organization. That description is not a substitute for the absent task-specific final mathematical report.

Recommended current receipt state:

Picked up; F2 numerical artifacts returned and reviewed; final Fable arithmetic-transfer report not yet present at a408e705.

Preserve the original input snapshot unchanged. Put the corrected live receipt state in the current task ledger or synchronization note.

2. Continuum computation versus the exact certificate

Fable independently implements Gauss – Jacobi/Legendre quadrature using a different simplex parametrization:

$$u = \sigma\tau, \quad w = \sigma(1 - \tau), \quad \sigma = (1 - \nu)s, \quad du\,dw = (1 - \nu)^2 s\,ds\,d\tau.$$

The Jacobian, weight $\nu^{\alpha-1}$, continuous sine-kernel limits, prime-insertion shifts and coefficients agree with the fixed schema. Its expectation-product expansion correctly uses both first and second background S_2 moments, including the $u^2 w^2$ insertion cross term.

For $\ell = 16/15$ and the specified rational $H = f(\nu) + g(\nu)S_2$, the order-64 result is

$$\mathcal{J}_{\text{Fable}} = -0.014662375473370598.$$

The order-40/order-64 spread is 2.38698×10^{-15} . The difference from Astra's stored target -0.014662375473368985 is

$$-1.61329 \times 10^{-15}.$$

Astra's exact rational continuum certificate encloses the value within approximately

$$[-0.014662375473368995, -0.014662375473368974].$$

Fable's ordinary floating result is **not literally inside that much narrower interval**. Agreement is to roughly fourteen decimal places, as expected for this quadrature and floating summation. This is consistent independent numerical evidence, not a new interval certificate and not a contradiction of the existing certificate.

The separated order-64 quantities are

Quantity	Fable floating value
I	0.9976809471028788
M_{2a}	0.07475319679662187
M_{2b}	0.07563555032577103
M_3	0.0844031170042764
$(M_2 + M_3)/I$	0.2353376245266294
equivalent raw Rayleigh value	4.6453785095398725

The half-gap threshold for that raw Rayleigh value is $\pi^2/2$, so the margin remains negative.

The mass-only version of the same rational f gives

$$J_{\text{massonly}} = -0.021565258857184827,$$

and the fixed S_2 extension improves that particular continuum value by approximately 0.00690288338381423. This comparison should not be confused with the much smaller improvement over the independently optimized mass-only degree-14 baseline:

$$J_{\text{degree14}} = -0.015357981703850332$$

in Fable's calculation. Its difference from Astra's degree-14 numerical result is 2.22×10^{-16} .

Fable correctly labels its continuum JSON as a stipulated-form quadrature, not a certificate and not an arithmetic theorem.

3. Finite arithmetic operator implementation

The finite script uses

$$x_n = d_\ell(n)H(\log n/\log L, S_2(n))/\sqrt{n}, \quad A_{qm,m} = \frac{2 \sin((\pi/2) \log q/\log L)}{e\sqrt{q}}, \quad q = p^e.$$

The index ranges, prime-power exponents, $1/\sqrt{q}$ factor and distinct-prime definition of S_2 match the stated trial. The divisor sieve multiplies by $(\ell + e - 1)/e$ at each prime-power level, giving the correct cumulative $d_\ell(p^e)$.

The full Rayleigh expression is

$$J_L = \frac{\|Ax\|^2 + \langle x, A^2x \rangle}{2\pi^2 \|x\|^2} - \frac{1}{4}.$$

The stored scalar numerator pieces and normalized contributions reconstruct every reported trial/mode margin to within 2.78×10^{-17} . The diagonal quantity is the equal-insertion part of $\|Ax\|^2$, so its correspondence with M_3 is the correct one.

Fable reports validation of the floating $d_{16/15}$ sieve against exact rational values for all $n \leq 10^4$, with maximum relative error 1.52234×10^{-15} . This checks the sieve; it does not make the logarithms, trigonometric values, finite sums or eigenvalues exact.

The committed completed lengths are $10^3, 10^4, 10^5, 10^6$:

L	fixed J_L , full operator	same f , mass only	fixed minus mass only
1,000	−0.0519926036839853	−0.0624270849868530	0.0104344813028678
10,000	−0.0431174919059547	−0.0529803909830878	0.00986289907713303
100,000	−0.0376302025223919	−0.0470594610068825	0.00942925848449061
1,000,000	−0.0339175762941035	−0.0430127050936602	0.00909512879955668

The finite values approach the proposed continuum value over these lengths, but four slowly varying points cannot prove a limit. All recorded fixed-family margins are negative.

Fable also computes full-operator top eigenvalues for $L = 10^3, 10^4, 10^5$. Comparing the stored values with Astra's previously committed `arithmetic-results.json`, without rerunning an eigensolver:

L	Fable top eigenvalue	Difference from Astra
1,000	3.949287136694299	0
10,000	4.105867045445441	-1.78×10^{-15}
100,000	4.205255380109547	0

This is useful independent evidence that the two finite-operator implementations have the same normalization. It does not imply a uniform upper bound for all lengths.

4. Prime powers, background coincidences and finite drift

The script includes two diagnostic alterations:

- **nopp**: keep only prime insertions, discarding higher prime powers;
- **clean**: additionally remove insertions p into backgrounds already divisible by p .

The code implements these restrictions consistently in $\mathcal{A}$, its transpose and the diagonal contribution. They are altered finite operators, not the exact full theorem operator.

For the fixed trial, the recorded values are:

L	full	nopp	clean
1,000	−0.0519926037	−0.1022308502	−0.1269460914
10,000	−0.0431174919	−0.0892650368	−0.1129780712
100,000	−0.0376302025	−0.0793053529	−0.1015970942
1,000,000	−0.0339175763	−0.0715145441	−0.0922927227

These differences are large at attainable lengths. They show why simply deleting exceptional terms numerically and declaring them negligible is unsafe. They do **not** refute an asymptotic theorem proving that their normalized contribution tends to zero. Astra's separately written fixed-family transfer proof uses explicit weighted Schur bounds and an ordered cutoff limit for precisely this issue.

The finite S_2 -moment ratios are likewise not uniformly close to one: the first-moment ratio changes from 1.0142 to 1.0302, while the second-moment ratio changes from 1.2034 to 1.1346. Nonmonotonic finite drift of an individual ratio does not contradict convergence. These values also cannot prove the asserted Poisson – Dirichlet moment law by themselves.

The background Euler-product constant is explicitly marked recalled, truncated at primes $\leq 10^7$, and not independently source-verified in Fable's script. It is a normalization diagnostic rather than a certified leading-constant enclosure.

The drift-fit script warns correctly that its fitted limits are diagnostic only. Since no committed fit output is present at the reviewed SHA, this review assigns no numerical conclusion to that script.

5. What this changes in the fixed-family proof status

Astra's round-two `symmetric_prime_arithmetic_transfer.md` and the independent review already give an ordinary written proof for fixed $\ell \geq 1$ and fixed $H = f(v) + g(v)S_2$, including this rational trial. The earlier round-one certificate's “arithmetic transfer not certified” label is historical and should be accompanied by a link to that later proof, not silently rewritten.

Fable task001 contributes a genuinely separate numerical implementation supporting:

1. the stipulated continuum-integral value;
2. the exact finite operator's normalization;
3. the importance of finite prime-power and coincidence corrections;
4. the continued negative sign of the tested margin.

It does **not**, in this snapshot, contribute the requested derivation of the marked Euler product, the short-background limit, or the complete insertion/coincidence asymptotic. Its code explicitly says that the original Inoue paper was not read for this diagnostic. That is an honest scope label, but it leaves the single main mathematical obligation unfulfilled in the returned Fable artifacts.

There is no contrary theorem claim to reconcile. Keep the following three statuses distinct:

- **Exact continuum inequality:** Astra certificate; negative margin.
- **Fixed-family arithmetic transfer:** Astra written proof plus separate mathematical review.
- **Fable task001:** numerical F2 independently returned; final arithmetic report absent at the pinned source.

No new zero-gap record, global variational no-go, AH contradiction or proof of RH follows from this synchronization.

6. Review method and reproducibility

This review read the pinned files with git show from:

```
/Users/qingyunsun/Library/CloudStorage/Dropbox/Research/ACUE-Astra-Handoff-2026-09-04/github-worktree
```

Example:

```
git show a408e7050fffc74459b3c83fafa5ac03c8b7dea6:research/riemann-
rmt/overnight/fable/astra_tasks/task001/f2_finite_sum.py
```

Presence and absence were checked with `git ls-tree` and `git cat-file` at the same SHA. Stored JSON values were parsed to recompute scalar normalization identities and compared with the existing Astra result files. No new large eigenproblem, finite-vector sweep, Claude message, Claude task or session was launched. All review output is in staging.

7. Coverage of the current single-session computation packet

This table maps the requirements in the current local packet

`/Users/qingyunsun/Documents/Codex/2026-09-04/realtime-voice-chat/outputs/FABLE_SINGLE_SESSION_COMPUTE_TASK.md`

to actual artifacts already present at `a408e7050fffc74459b3c83fafa5ac03c8b7dea6`. Receipt of that new packet is **unconfirmed**. The confirmed pickup in § 1 concerns the older task001 packet. Matching old outputs may satisfy individual new requirements without showing receipt or execution of the new packet.

Reviewed packet SHA256: `745c6411355b163f429cbeaa96b7c1b681f8b36e6ec92152f74cba30739743f5`.

All paths below are relative to the pinned `research/riemann-rmt/overnight/fable/astra_tasks/task001/` directory.

Current acceptance requirement	Coverage	Existing evidence and precise limit
Fixed rational continuum form, simultaneous insertion shifts, and exact rational norm	Partly	<code>f2_continuum.py</code> , <code>f2_continuum_results.json</code> independently implement the expectation expansion and simultaneous shifts; $J = -0.014662375473370598$ agrees numerically. Norm is float64 quadrature only. No independent exact rational norm calculation or audit of the original signed Taylor/ π /remainder certificate is returned.
$g = 0$ control at the same $\ell = 16/15$ and same f	Covered	The massonly trial exists in both continuum and finite-sum scripts/JSON. It gives continuum $J = -0.021565258857184827$ and matching finite controls. The separately requested nonzero-rescaling invariance test is not covered by an explicit returned test.
$L = 97$ direct-loop versus sparse convolution, using candidate and deterministic control	Not covered	No $L = 97$ run or direct/sparse two-vector comparison is present. Existing sieve validation and $L \geq 1000$ eigenvalue agreement do not substitute for this requirement.
Fixed-vector $L = 10^3, 10^4, 10^5$, $\theta = 1$, with separate quadratic terms and candidate-minus-control comparison	Covered	<code>f2_finite_sum_results.json</code> records $\ x\ ^2$, $\ Ax\ ^2$, $x^T A^2 x$, J_L for candidate and massonly control at all three cutoffs. Their differences are reconstructed in § 3 from those stored values; no rerun was needed. The old run also contains additional eigenvalue calculations, which are not part of the new packet's requested batch.
One $\theta = 0.98$ sensitivity check at the largest completed cutoff	Not covered	The pinned operator hard-codes $\theta = 1$; no $\theta = 0.98$ result is present. Neither the continuum control nor a prime-insertion restriction is this sensitivity test.

Current acceptance requirement	Coverage	Existing evidence and precise limit
Arithmetic-obligations table covering weighted moments, existing prime divisibility, multiplicity, coincident insertions, cutoffs/uniformity and beyond-diagonal errors	Not covered	F2 contains helpful finite coincidence/moment diagnostics, but no six-item estimate/status/use table or final task001_report.md . Astra's separate transfer proof is existing Astra evidence, not fulfillment of a Fable deliverable.

Acceptance state: reuse the covered numerical outputs; retain partial or absent requirements explicitly. The old task was picked up and F2 outputs were returned. The new packet's receipt and complete acceptance are not established by this snapshot. No new calculation or Claude task was launched to fill these gaps during this review.

Current report 20: fable_heat_sync_review

Source: `research/reports/fable_heat_sync_review.md`. SHA-256:

`9dc10cd810e665fc49c192b17cd595d209fe8e1d4df7de1a8e5158bc470b3713`.

Pinned Fable intake: heat flow, function fields, and LR hard cores

2026-09-05. This review reads the existing public Fable work at `a408e7050fffc74459b3c83fafa5ac03c8b7dea6` in `galpha-ai/Alpha-devbox`, the PR #11 head supplied for this intake. It uses `git ls-tree` and `git show` at that exact commit. No Claude session was started, no task was sent to Fable, and none of its experiments was rerun.

Companion evidence file: `fable_heat_sync_inventory.json`. It records the pinned revision, source hashes, line numbers, expected-report presence, and public journal event counts. The main coordinator is mirroring the source folder separately.

1. Main conclusion

This public snapshot contains useful new mathematical formulations and experimental scripts, but **no completed heat-flow, function-field, or LR proof report and no committed output dataset for those scripts**. They must not be merged into the theorem ledger as completed proofs or completed numerical results.

Three concrete defects were found in the available source:

1. The explicit CUE-background constant assembly reverses an inequality in its second regime. A limited repair is possible, but the displayed global constant is not certified by the committed argument.
2. The pair-LP ansatz assumes a compactly supported correction whose Fourier transform vanishes on an interval. In the exact continuum problem this forces the correction to vanish identically. Its numerical feasibility boundary therefore cannot be treated as a general LR upper bound.
3. The function-field support code assigns the wrong Weyl density to the negative-determinant component of the even orthogonal group. Its matrix sampler and claimed reference rejection sampler target different laws.

There is no new evidence here contradicting our audited finite-CUE isolated-pair depth theorem, the all-forward-time protected-trace obstruction, or the force-energy identities. The correct $SO(\text{odd})$ three-body example instead illustrates why an isolated-pair assumption is necessary.

2. What is actually present

The [claims ledger](#) at the [pinned commit](#) has seven claim rows, all from the prime-certificate cluster D1. It has no completed A-cluster heat claim or B-cluster LR/function-field claim.

The [coordination file](#) acknowledges the old endpoint-bound error, the Lean collision-time gap, the distinction between strong and density-tolerant AH, and the need for outward-rounded prime certificates. These acknowledgements agree with Astra's audit.

However, the coordination text refers to a repaired theorem report as though its argument were already available; the referenced report is absent from this commit. The same issue applies to the future files in the [round-two plan](#).

The following planned reports are absent:

- r1_theoremB_repair.md, r1_cue_background.md, r1_cbe_background.md;
- r1_levelB_barrier.md, r1_small_gaps.md, r1_zeta_numerics.md, r1_structure_review.md;
- r2_dynamic_universality.md, r2_marked_depth_proof.md, r2_acue_rho_bound.md;
- r2_lr_hardcore_lp.md, r2_function_field.md, r2_dbm_relaxation.md.

There are corresponding background, function-field, and LR scripts. Their declared output paths under overnight/fable/data are not present in this commit. This is an observation about committed evidence, not a statement that no computation happened in Fable's private running environment.

The public harness export has six workflow journals. Five contain only two “started” entries each. The remaining journal contains three “started” entries and one result: the completed H2 prime-certificate report. The agent metadata files contain agent type and spawn depth, while the index leaves labels, status, and model columns blank. These exports do not provide an independent-refuter completion record for heat/FF/LR at this head.

The substantive completed report visible here is r1_h2_interval_cert.md, which belongs to the prime-audit intake rather than this heat review. Its stated H3 result is an M_k certificate, while the large H3 tuple is explicitly not reverified in the claims ledger. Preserve that distinction.

3. CUE background work: useful local structure, uncertified global constant

3.1 What agrees with the audited repair

The [Theorem-B check script](#) uses the correct attractive coefficient flow and root ODE. It replaces the erroneous single-endpoint csc-squared bound with the maximum of both endpoint values.

Its exact bracket

$$\cot(x_b/2) - \cot(x_a/2) = \frac{\sin(g/2)}{\sin(x_b/2)\sin(x_a/2)}, \quad x_a = x_b + g,$$

is correct. Since both sine factors are positive for the chosen representatives,

$$B = 2 \sin(g/2) S_{\text{exact}} \leq g S_{\text{exact}} \leq g S_{\text{avg}} \leq g S_+.$$

The middle comparison is the arithmetic-geometric-mean inequality applied to the endpoint reciprocal sine values. These are valid local algebraic repairs. They do not themselves control a moving background up to collision.

The script explicitly probes a three-cluster configuration where the background grows by more than a factor two; this is a useful adversarial test, not a proof that all other configurations satisfy a window bound. Its “ACUE-type” random subsets are uniformly sampled non-clock subsets, not Vandermonde-weighted ACUE samples. That is acceptable for deterministic stress testing, but their sample frequencies are not ACUE probabilities.

3.2 The exact three-point constant is a useful route

The [background-constant script](#) formulates a global determinantal three-point bound using alternants and the inequality $|\mathcal{S}_\lambda(\mathbf{z})| \leq \mathcal{S}_\lambda(1, 1, 1)$ on the torus. The coefficient

$$A_3(N) = \sum_{0 \leq m_1 < m_2 < m_3 < N} [(m_2 - m_1)(m_3 - m_1)(m_3 - m_2)]^2 = \frac{N^3(N^2 - 1)^2(N^2 - 4)}{2160}$$

leads to

$$\rho_3(\theta_1, \theta_2, \theta_3) \leq \frac{A_3(N)}{4(2\pi)^3} \prod_{i < j} |e^{i\theta_i} - e^{i\theta_j}|^2.$$

The factor $1/4$ comes from the denominator 2 in the three-variable Weyl dimension formula, squared. This agrees with the N^9 three-point Vandermonde structure used in our audited CUE transfer. It is a promising explicit refinement rather than evidence of a different heat-flow mechanism.

The script's polynomial identity tests and sampled determinant comparisons should be described as checks until accompanied by its missing written report. Polynomial interpolation through finitely many values is not an all- N proof unless the degree bound is also justified.

3.3 Definite sign error in the second-regime tail assembly

At line 159, the code reasons

$$N < (L/4)^3 \implies \frac{1}{N} < \frac{64}{L^3}.$$

The conclusion is reversed: reciprocation gives $1/N > 64/L^3$. Thus a bound C/N from a cutoff gap $4/N$ cannot be converted to $64C/L^3$ by that step.

A possible repair of this single step uses the deterministic fact $\delta_{\min} \leq 2\pi/N$. In the nontrivial remaining regime

$$4N^{1/3} < L \leq 2\pi N^{1/3},$$

one has $L^3 \leq (2\pi)^3 N$, hence

$$\frac{C}{N} \leq \frac{(2\pi)^3 C}{L^3}.$$

For $L > 2\pi N^{1/3}$, the event $\delta_{\min} > LN^{-4/3}$ is empty. This suggests replacing 64 by $(2\pi)^3$ in the indicated regime conversion. It does not validate every other missing step in the announced background-tail theorem.

The script assigns “pass=True” to assembled constants rather than verifying their probability derivation. Its displayed bounds $P(S_* > MN^2) \leq CM^{-3/8}$ and the refinement with exponent $-1/2$ must remain unproved in this intake. A software pass flag is not a mathematical certificate of that tail argument.

Our audited Galilean isolated-pair proof does not depend on this constant assembly: it obtains $\delta^2 B \rightarrow 0$ from a triple-free packing event. Therefore this defect does not invalidate our finite-CUE result.

4. Function-field support: correct identities, wrong orthogonal reference law

4.1 Valid finite polynomial identities

The [function-field depth core](#) implements the same attractive coefficient flow

$$P_s(z) = \sum_j a_j e^{sj(M-j)} z^j.$$

At a double root z_0 , it has

$$\partial_s P_s(z_0) = -z_0^2 P_s''(z_0),$$

because the operator is $(M-1)z\partial_z - z^2\partial_z^2$ and $P'_s(z_0) = 0$. For a simple double root this supplies a nonzero transverse parameter derivative. It does not, by itself, settle all multiple or symmetry-forced collision strata needed for a general almost-everywhere continuity theorem.

The genus-two palindromic reduction is also correct:

$$P_s(z) = z^4 + Ae^{3s}z^3 + Be^{4s}z^2 + Ae^{3s}z + 1$$

reduces, after division by z^2 , to

$$Q_s(x) = x^2 + Ae^{3s}x + Be^{4s} - 2, \quad x = z + z^{-1}.$$

Bulk and edge collisions are detected algebraically by its discriminant and the equations $Q_s(2) = 0$, $Q_s(-2) = 0$. The numerical implementation uses a finite sign grid and a time cap; those are numerical detection choices, not a certified first-root solver.

4.2 The forced-root three-body example is real and informative

For the configuration $\{1, e^{i\theta}, e^{-i\theta}\}$, put $A = 1 + 2\cos\theta$. Then

$$P_s(z) = z^3 - Ae^{2s}z^2 + Ae^{2s}z - 1 = (z-1)(z^2 + (1 - Ae^{2s})z + 1).$$

If $0 < \theta < 2\pi/3$, so $A > 0$, the first collision occurs at $z = 1$ when $Ae^{2s} = 3$. Consequently

$$D = \frac{1}{2} \log \frac{3}{1 + 2\cos\theta} = \frac{\theta^2}{6} + O(\theta^4).$$

This matches the formula in the [hard-edge routine](#). The routine does not explicitly enforce the domain $0 < \theta < 2\pi/3$; the small-angle calls shown in the check script lie inside it. At $\theta = 2\pi/3$ the configuration is a stationary clock. For $2\pi/3 < \theta < \pi$, the first collision is at -1 , and the different expression is $D = \frac{1}{2} \log(-1/A)$.

For $\theta \downarrow 0$, the minimum gap is $\delta = \theta$ and $D/(\delta^2/8) \rightarrow 4/3$. This is not a counterexample to our theorem: a third root is only a gap-scale distance from the chosen pair, so $\delta^2 B$ does not tend to zero. It is a clean explanation of why forced symmetry and three-point clustering must be handled separately in orthogonal models.

4.3 Definite Weyl-density mismatch for the even negative component

The code defines O_minus as $O^-(2N+2)$, with both forced eigenvalues $+1$ and -1 . Yet [lines 367 – 368](#) and the rejection sampler at [lines 388 – 389](#) multiply the free-angle density by

$$\prod_j \cos^2(\theta_j/2).$$

For $O^-(2N+2)$, the correct factor is

$$\prod_j \sin^2 \theta_j,$$

the same free-angle factor as $USp(2N)$. The half-angle cosine factor belongs instead to $O^-(2N+1)$. This is explicitly distinguished in Meckes, *The Random Matrix Theory of the Classical Compact Groups*, Theorem 3.5, printed page 76.

There is also an elementary diagnostic independent of the formula table. In even dimension, $U \mapsto -U$ preserves the negative-determinant component and its Haar measure. For rank $N = 1$, this maps the free angle θ to $\pi - \theta$, so its density must be reflection symmetric. The claimed $\cos^2(\theta/2)$ density is not. It would give

$$\mathbb{E} \cos \theta = \frac{\int_0^\pi \cos \theta \cos^2(\theta/2) d\theta}{\int_0^\pi \cos^2(\theta/2) d\theta} = \frac{1}{2},$$

whereas reflection symmetry requires zero. Thus the matrix Haar sampler and the coded rejection sampler cannot both pass as implementations of the same law. No rerun is needed to establish the mismatch.

4.4 What remains unproved for function fields

No curve-family equidistribution theorem with its hypotheses, exceptional set, order of limits, or uniformity is written out in the available snapshot. The fixed-rank continuous-mapping route is plausible only after establishing the required almost-everywhere continuity of the chosen depth functional. Passing subsequently to growing rank needs the relevant group-specific extreme-gap law and separate hard-edge/forced-root analysis.

The pointwise continuity demonstration in the check script is a numerical path through a few configurations, not an almost-everywhere theorem. Equality of two initial gaps also does not by itself prove that they collide simultaneously under different backgrounds. These missing steps remain exactly the kind of obligations recorded in our earlier audit.

5. LR work: valid test families versus an invalid continuum ansatz

5.1 A precise obstruction in the pair-LP parametrization

The `pair-LP` script writes

$$g = g_{\text{sine}} + u, \quad u = 0 \text{ outside } [-X, X],$$

and aims to impose $\mathcal{U}(a) = 0$ throughout $(-1, 1)$. For a compactly supported integrable function or finite signed measure, $\mathcal{U}(z)$ is entire in the complex variable z . If it vanishes on a real open interval, the identity theorem forces it to vanish identically, and Fourier uniqueness gives $u = 0$.

But a positive hard core requires $u = -g_{\text{sine}}$ on a nonempty interval around zero. Since $g_{\text{sine}}(x) > 0$ for nonzero sufficiently small x , that is impossible. Thus the exact continuum version of this compact-correction ansatz is infeasible for every $c > 0$, including the known half-lattice model that motivated the programme.

Sampling Fourier constraints at finitely many frequencies avoids the analytic contradiction only by no longer imposing the exact mimicry condition. This does not turn the finite-support family into an outer relaxation containing all admissible stationary processes. Therefore an infeasibility threshold from this code is not a rigorous general upper bound on LR μ , and shrinking the grids alone cannot repair the logical direction.

There is an additional implementation-level overstatement at line 54: the “minimum over a cell” is approximated using nine sample points. A sampled minimum is at least the true minimum. Its negative is therefore not a certified conservative lower constraint on u that guarantees $g \geq 0$ throughout the cell.

5.2 The count-variance and test-function formulations are useful

The `triangle certificate` has a valid underlying necessary condition. For a stationary intensity-one process with hard core c , an interval of length $L \leq c$ contains at most one point almost surely, modulo irrelevant endpoint events. Hence

$$\mathrm{Var} \, N[0, L] = L - L^2.$$

If its spectral measure agrees with the sine structure factor $|\alpha|$ on the unit band and is nonnegative outside, then

$$L - L^2 \geq 2 \int_0^1 \frac{\sin^2(\pi L \alpha)}{\pi^2 \alpha} d\alpha.$$

This is a genuine analytic necessary inequality under the stated process hypotheses. The committed file computes a root numerically, but its numerical output is not committed here; this intake does not claim a new bound.

The [weighted-window eigenvalue formulation](#) extends the same variance argument to real weights on an interval. The [Selberg-type formulation](#) instead uses $T(x) = (c^2 - x^2) |F(x)|^2$ with band-limited F . Their continuum inequalities are sensible test families; the numerical eigenvalue computations still require certified quadrature and discretization error bounds before they become numerical certificates.

The [general band-limited LP](#) itself acknowledges that rational-knot piecewise-linear Fourier profiles create a periodic $x^2 T(x)$, obstructing a global sign certificate in that finite family. A tail penalty does not by itself certify the sign between the sampled spatial grid points. Any claimed improvement from that script would need a complete continuum certificate, not just a positive floating LP objective.

6. Conflict matrix and recommended handoff classification

Fable item at the pinned head	Relation to audited Astra work	Handoff status
Endpoint maximum repair and exact cotangent bracket	Agrees	Preserve as valid local algebra
Claimed explicit CUE-background tail constant	Contains a reversed inequality; written proof absent	Unproved, with exact defect and partial repair recorded
CUE three-point Vandermonde bound structure	Agrees with our triple-free route	Useful independent formulation; no new theorem claim
SO(odd) three-body depth $\theta^2/6$	Outside isolated-pair hypotheses	Valid finite example, with domain restriction
Even $\mathcal{O}$ Weyl reference density	Wrong group factor	Repair before interpreting group-specific diagnostics
Function-field universality	Missing continuity/equidistribution/limit-order proof	Open
Compact-correction LR pair LP	Does not contain the target continuum models	Invalid as a general upper-bound certificate
LR window/Selberg analytic inequalities	Plausible and directly checkable under stated hypotheses	Valid formulations; numerical certification absent
Proposed DBM relaxation experiment	Only a planned report at this head	No completed evidence to compare
Protected trace equality and force-energy separation	No contradicting result in this snapshot	Our independent audited conclusions unchanged

The archive should preserve the original source verbatim, attach this pinned review, and avoid upgrading planned filenames, printed PASS labels, or uncommitted output paths into completed discoveries. A later Fable commit may supply repairs or full reports; it should receive a new pinned intake rather than silently replacing this record.

Claim ledger 21: CLAIM_LEDGER

Source: `research/claims/CLAIM_LEDGER.md`. SHA-256:
`f4bc91ed5e83e4c7351f61ab2ddf4bf9e8350810d3888866ae76bd3b9b208f20`.

Claim ledger

The classifications below apply to the exact stated scope. “Proof with internal review” means a written mathematical proof and a recorded human-readable audit; it does not mean a Lean kernel has checked it or external referees have accepted it. Historical files may use stronger wording than this ledger.

ID	Claim	Status	Evidence and remaining obligation
AUD-001	Independent historical audit has 21 passing checks	Reproduced finite checks	<code>research/logs/audit_results.json</code> ; not a blanket theorem audit
AUD-002	Original verifier has 39 pass, 4 fail, 1 skip	Recorded run	<code>research/logs/verify-codex-original.log</code> ; zero process exit does not erase failures
ALG-001	Fixed-support approximator is saturated	Exact algebra	<code>research/reports/residual_gram_round1.md</code> ; says nothing about larger support
NEG-001	Sparse longer tails under the stated coefficient bounds cannot supply constant energy	Proof in report	Same report; hypothesis on coefficients is essential
NUM-001	Degree-14 half-gap trial has margin about -0.01535798	Numerical optimization	<code>research/residual-gram/variational-results.json</code> ; no global optimality claim
CERT-001	Fixed rational trial has margin in (-.01467,-.01465) for the stipulated integral	Exact rational certificate	<code>research/residual-gram/rational-trial-certificate.json</code> ; transfer is separate
ARITH-001	Fixed symmetric-prime family has that arithmetic limiting form	Written proof, independent internal review	<code>research/reports/symmetric_prime_arithmetic_transfer.md</code> , <code>symmetric_prime_transfer_independent_review.md</code>
NUM-002	Full finite operator remains below $\pi^2/2$ through $L=10^7$	Numerical eigenvalue data	<code>research/operator-bounds/extended-arithmetic-results.json</code> ; no asymptotic upper bound

ID	Claim	Status	Evidence and remaining obligation
NUM-003	Structural features nearly recover the $L=10^6$ optimum	Numerical fit and full-form evaluation	<code>research/operator-bounds/eigenvector-feature-results.json</code> ; finite L only
NEG-002	Tested Schur – Volterra majorant is too large to prove the target barrier	Numerical diagnostic of an explicit bound	<code>research/reports/residual_gam_round2.md</code> ; not a certified optimum over all profiles
FOCK-001	Fixed prime-bin compression provides a constructive limiting lower bound	Proof draft	Same report; not full operator convergence or an upper bound
NEG-003	Centered Gaussian pole term cannot be dropped in a uniform long-support estimate	Analytic counterexample, primary review	<code>research/reports/centered_gaussian_mixed_moments.md</code> ; RH for the contour identity
NEG-004	Nonnegative Fourier transform cannot coexist with a zero at the origin for a nonzero Schwartz weight	Elementary proof	Same report; does not rule out quantitative subtraction
HEAT-001	Isolated-pair first collision satisfies $D \sim \delta^2/8$ under $\delta^2 B \rightarrow 0$	Proof draft with internal audits	<code>research/reports/yau_flow_galilean_refinement.md</code> ; quantitative second-pass audit in <code>galilean-proof-audit.md</code> ; $K=16384$, $\epsilon_0=1/524288$
RMT-001	Circular positive-integer-beta extreme gaps induce the stated first-collision law	Proof draft with internal review	<code>research/reports/yau_flow.md</code> ; published gap laws are inputs; novelty audit incomplete
PRIME-001	A 39-element admissible tuple has diameter 182	Exact finite verification	<code>research/reports/prime186.md</code> ; DHL[39,2] is not proved
PRIME-002	Fixed prime-sieve minorant mass has the saved rational enclosure	Exact rational computation plus separate coverage check	<code>research/prime-gaps/minorant_mass.json</code> , <code>minorant_geometry.json</code>
OPEN-001	$\mu\zeta < 1/2$ under RH	Open in this programme	No positive applicable certificate
OPEN-002	New fixed-width Montgomery pair-correlation theorem beyond known support	Open in this programme	No new arithmetic mixed-term theorem supplied
OPEN-003	Prime-gap bound below 186	Open in this programme	No improved admissible-tuple/sieve certificate chain
OPEN-004	RH, full GUE statistics, general-beta depth universality	Open here	Not implied by any local counterexample or finite computation

All paths are relative to the repository root. The public reports contain the derivations; the JSON files preserve values and parameters. Run logs are snapshots, not machine proofs. Updates should change this ledger and the associated report together.

Round 3 and incoming collaboration

ID	Claim	Status	Evidence and remaining obligation
DYN-001	All protected symmetric trace moments agree under forward deterministic circular flow	Written proof with independent internal review	<code>research/reports/dynamic_generator.md</code> ; $m \leq N$, all orders and all $t \geq 0$, not postcollision attractive dynamics
FORCE-001	Force-square expectations differ by a factor of two, with exact dissipation formulas	Written proof with independent internal review	<code>research/reports/force_energy.md</code> ; finite CUE/ACUE, singular two-point information
BRIDGE-001	Initial smallest gap need not collide first	Exact rational counterexample and Sturm certificate	<code>research/bridge-audit/attachment_bridge_checks.json</code> ; lower bound survives; pair-selection correction required
BRIDGE-002	Marked inverse blowup need not force marked-depth derivative blowup	Exact symbolic counterexample	Same certificate; paired isospectral 2×2 matrices, no fixed-norm-along-path claim
SOURCE-001	LR explicitly suggests a .606894 hard-core upper bound	Primary-source correction	<code>research/reports/new_attachment_bridge_audit.md</code> ; not proved in the cited passage
RITZ-001	Two-dimensional finite residual correction improves the fixed trial	Reproduced finite calculation	<code>research/operator-bounds/ritz_residual_diagnostic.json</code> ; half-gap margin remains negative; no automatic continuum transfer
PRIME-003	Complementary root predicates yield a finite allocation frontier	Exact algebra and rational feasibility certificates	<code>research/reports/prime186_structural_frontier.md</code> ; fixed active row/template; not a global sieve obstruction or a new gap
FABLE-001	Earlier FABLE_001 pickup and numerical output exist	Pinned public-source evidence	<code>fable/PICKUP_RECEIPT.json</code> ; final arithmetic report absent; no receipt inferred for newer packet

Research log 22: RESEARCH_LOG

Source: `research/RESEARCH_LOG.md` . SHA-256:

`716823c45370488aa36f6ceb1d9c7f69e727fbee0ac964db7a267ae887cc311` .

Research log: recoverable mathematical work

This log records problems, mathematical arguments, calculations, corrections, and evidence. It is a public research record, not a transcript of private model deliberation. Detailed proofs and counterexamples live in the linked reports. Timestamps record checkpoints; they do not imply continuous computation between checkpoints.

2026-09-05: source recovery and independent audit

Question. Which statements in the earlier ACUE, Newman-depth, zeta and prime-gap programme survive direct checking?

Input. Historical base `ee61fba38925190916e556f87bc1ea83a502413e` ; Fable handoff update

`0f86c31d847178f0a774f70073ddc44206b32f98` ; user-provided mathematical context and original Inoue probe. The public historical source is preserved in `historical/riemann-rmt/` ; private supplied context and conversation exports remain in the adjacent local archive.

Outcome. The original verification script reports 39 PASS, 4 FAIL, 1 SKIP even though its exit status is zero. The independent audit has 21 passing checks. Those passing checks verify specified algebraic/numerical facts; they do not validate every historical theorem.

Corrections worth preserving. The Galerkin angle needs `acos(1-1/(n*n*a))` ; the Wilson trace identity requires the stated boundary geometry; COM null dimensions in the tested sizes are not the earlier claimed pattern; a near-endpoint background example refutes an old uniform upper bound. Clock mass does not disappear at fixed N. The Lean development has actual unfinished obligations, and a conclusion-free depth comparison is not a depth theorem. See the [full audit narrative](#) and [reproduction](#).

Decision. Keep corrected statements, archive failed claims verbatim with warnings, and require an explicit arithmetic transfer before assigning consequences to real zeta zeros.

2026-09-05 05:56:56 UTC: first autonomous research milestone begins

The initial authorized milestone is eight hours. The user's coordination task subsequently relayed a 72-hour recoverable programme. This log makes no guarantee of uninterrupted execution or of solving an open conjecture. The active priority is a genuine RH-conditional zeta half-gap result or an explicit out-of-band correlation theorem. Matrix-model distinctions are supporting results.

Three internal research lanes were assigned: scalar heat-flow proof and Yau comparison; residual Gram and arithmetic resonators; prime-gap certificate audit and transferable positivity methods. Their reports preserve failures as well as successes.

Round 1: the saturated approximator and the new resonator direction

Question. Can another approximator inside the same accessible product support recover the missing half-gap energy?

Exact answer. Completion of squares gives $2 \operatorname{Re}\langle \mathbf{b}_a, \mathbf{b}_g \rangle - \|\mathbf{b}_a\|^2 = \|\mathbf{b}_g\|^2 - \|\mathbf{b}_a - \mathbf{b}_g\|^2$. Thus the selected coefficient-space approximator is already saturated for a fixed resonator and fixed product support. This does not optimize the resonator, prove a global method barrier, or allow replacement of finite-height Gram matrices by diagonal matrices without an error theorem.

Independent computations. The published linear Inoue trial is reproduced. A degree-14 one-variable trial has half-gap margin about -0.01535798170385 . A fixed rational symmetric-prime trial improves this to a certified continuum value about -0.014662375473369 ; richer exploratory trials reach about -0.01465473 but have ill-conditioned numerical Gram problems. These are distinct candidates, all negative.

Failure with information. Sparse support beyond the usual cutoff cannot provide a fixed amount of missing energy under the stated subpolynomial coefficient bounds. Sparse orthogonality alone is insufficient; dense correlations or different coefficient concentration would need a new estimate.

Evidence. Round 1 report, exact rational certificate, certificate output.

Round 1: scalar heat flow and extreme circular gaps

Question. Does a tiny isolated pair determine the first collision of the actual coefficient heat flow?

Mathematical output. A deterministic proof draft uses a scalar heat normal form. Removing the background linear drift by exact heat conjugation reduces the sufficient isolation condition to $\delta^2 B \rightarrow 0$. Its probability application uses triple-cluster exclusion and published extreme-gap laws for CUE and positive integer beta circular ensembles.

Scope. This is a proposed full RMT theorem with internal audits, not a zeta-zero theorem and not a proof of universality for every beta. Its initial-background method overlaps classical Lehmer-pair theory; priority and formal verification remain separate obligations.

Evidence. Heat report, Galilean refinement, independent first review, numerical checks.

Round 1: what the prime-gap 186 calculation actually contributes

Question. Can exact positivity and certificate methods from the public 186 project provide a smaller prime-gap bound or a transferable tool?

Output. A 39-element admissible tuple of diameter 182 was found, but no DHL [39, 2] certificate was established. This is not a prime-gap improvement. An exact-rational convex-simplex enclosure independently bounds a fixed minorant mass. Separate exact geometry checks address the published triangulation's coverage. Numerical infeasibility searches are not portable impossibility proofs.

Failure with information. The official numerical verifier hit a FLINT convolution regression in the installed environment. No step was bypassed and no successful official certificate rerun was claimed. The public Lean project retains stated assumptions; it is not an assumption-free formalization of every analytic ingredient.

Decision. Stop broad prime-gap tuning until there is a definite path to a new analytic or certificate threshold. Transfer support accounting and exact positivity methods to the zeta lane.

Evidence. Prime report, exact minorant calculation, coverage verification, margin ledger.

Round 2: ten-million-dimensional operator and a failed upper bound

Question. Is there a large unused coefficient direction hidden outside the one-variable trial family?

Computation. At $L=10^7$, the full sparse arithmetic operator has computed largest eigenvalue 4.324089558989538 , below the half-gap threshold $\pi^2/2$. At $L=10^6$, 20 symmetric-prime features nearly reproduce the full finite optimum; 40 features reduce the remaining difference to about 0.0004437 . A 40-bin bosonic model with 215,308 states is also below threshold.

Analytic attempt. Separate Cauchy bounds lead to a scalar Schur – Volterra majorant. Its tested threshold is approximately 5.2074 , above $\pi^2/2$; it cannot establish a universal half-gap obstruction. More grid precision does not fix that mathematical loss.

Decision. Preserve the full spectra, fit residuals and failed majorant. Investigate compatibility of extremizers or new arithmetic products, rather than presenting rising finite spectra as an asymptotic proof.

Evidence. [Round 2 report](#), [spectra](#), [fits](#), [Fock data](#), [majorant data](#).

Round 2: centered Gaussian positivity has an explicit pole correction

Question. Can a centered Gaussian make all Fourier kernels positive and therefore extend the accessible mixed-product support for free?

Exact negative result. In the critical-line contour identity, centering leaves an explicit pole-cut integral. For the intended longer-tail supports its phase is opposite to the retained prime contribution. A dense coefficient example with bounded diagonal norms has a negative pole correction of order $M/\log M$; for $M=T^{(1+\eta)}$ this exceeds the presumed height normalization. The exact long-support Gram penalty also contains substantial off-diagonal mass.

Related obstruction. A nonzero Schwartz weight with nonnegative Fourier transform cannot vanish at the origin. Removing low-height coherent mass while keeping all those Fourier signs is not a free operation. This does not exclude all centered-packet methods: they must pay the pole term and control the density-subtracted mixed moment.

Evidence. [Full derivation and counterexample](#), [80-digit checks](#). The primary agent checked the source contour identity, the phase multiplication, and the dense-tail scaling. These are mathematical checks, not formal proof-assistant verification.

Round 2: arithmetic transfer of the fixed symmetric–prime trial

Question. Were the proposed prime-factor moments merely an assumed Poisson – Dirichlet model?

New proof draft. A positive-measure Laplace limit for $d_{\ell l}(n)^2/n$, marked reciprocal-prime measures, exact divisor convolution, and weighted Schur truncation together give a proposed complete derivation of the fixed family's continuum form. The argument distinguishes the same-prime contribution that survives in $A^* A$ from the repeated prime in A^2 that vanishes. It explicitly handles short backgrounds and signed polynomial coefficients.

Status. Independent adversarial audit requested. Do not upgrade this entry to an accepted theorem before reading that review. A successful audit would validate the family, while leaving the negative margin unchanged.

Evidence. [Proof draft](#).

Coordination and next decisions

Only the user's existing Fable session is to receive [one manually delivered audit packet](#). It independently checks a fixed vector and then stops. No receipt or completion is assumed. Astra owns the broader proof programme.

Next decisions depend on evidence: accept or repair the arithmetic transfer; finish the scalar-heat audit; identify a precise new estimate for density-subtracted mixed products or a compatibility-preserving upper bound. No parameter search, formalization attempt, or failed route should silently overwrite its earlier data. Git history and the source manifest are part of the evidence.

2026-09-05: arithmetic transfer audit resolved

The separately assigned residual-Gram agent accepted the fixed-family arithmetic transfer as a complete ordinary written proof after checking all six requested points. The primary author added the three requested explicit estimates. The result is unconditional as an arithmetic operator limit; RH is required only for its insertion into Inoue's zeta inequality. The fixed negative margin, full-operator limit, and half-gap target are unchanged. [Independent review](#).

The heat-flow author also completed a second-pass quantitative audit with $K=16384$ and $\eta_0=1/524288$, supported by exact symbolic checks. This is separate from the earlier independent prime-agent review. The true-Ht discussion identifies isolated small-gap arithmetic information as the missing input; it does not turn a circular model result into a zeta theorem. [Second-pass audit](#).

2026-09-05: dynamic proposals audited, including our own source correction

The low-mode degree-growth hypothesis fails: deterministic circular Coulomb evolution preserves positive and negative trace weights separately. A Schur/Cauchy-Binet proof gives ACUE/CUE equality for all derivative orders and all forward times in the protected symmetric-polynomial algebra. There are 234 exact protected comparisons; the nine out-of-band comparisons already differ at time zero. Independent cross-review is recorded separately.

The force square does separate the models, but collapses to a singular two-point statistic. Its CUE expectation is $N(N^2-1)/3$ and its ACUE expectation is half that value. The expected initial dissipative slope is negative infinity for CUE and exactly $-2N(N^4-1)/15$ for ACUE. The proof checks integrability and uses Fatou, rather than exchanging an invalid derivative and expectation. Complete float64 subset enumeration through $N=10$ supports the formulas.

Exact rational Sturm analysis provides a five-root counterexample to the minimum-gap-first premise. A symbolic 2×2 calculation disproves a proposed implication from marked inverse divergence to divergent depth susceptibility. The valid hard-core lower bounds remain useful; a universal exact identity requires the gap of the pair that actually collides.

We also corrected our own earlier Lagarias – Rodgers source judgment. The companion paper explicitly suggests adapting existing methods to a .606894 hard-core upper bound, but supplies no proof there. The earlier criticism that the literature never connects those objects was unjustified. Historical originals remain intact; the current main handoff and audit state the correction.

Evidence: [reports/dynamic_generator.md](#), [reports/force_energy.md](#), [reports/new_attachment_bridge_audit.md](#), and their adjacent focused script/data folders. These finite-model results are not a zeta AH refutation.

2026-09-05: 186 transfer tools and finite Ritz decision

The two-dimensional residual Ritz correction for the $L=100000$ fixed arithmetic vector raises the half-gap margin by approximately 0.000488512, from -0.0376302 to -0.0371417 . It validates a finite self-adjoint residual calculation, not an asymptotic fixed-family theorem. The report records a signed-residual lemma and the actual combined-product kernel that needs arithmetic control.

The user's Weijie Su link returns HTTP403 in this session. Its supplied structural description is checked against the primary 186 manuscript: complementary factorization controls the actual lcm divisor and triply dense divisibility, then permits wider sieve supports. No scalar exponent is imported into the unrelated zeta mixed-moment problem. The next decision is a new valid support/correlation estimate, not extra digits of the same negative vector. Evidence:

`reports/transferable_tools_ritz_decision.md` and the structural frontier report.

2026-09-05: Fable pickup evidence recovered from the existing public session

A live PR11 check found source head `a408e7050fffc74459b3c83fafa5ac03c8b7dea6`. Its coordination note acknowledges the earlier FABLE_001 task from Astra commit 97df092, and F2 continuum/finite-sum files are present. The intended final task001_report.md is absent at this source head. The newer bounded packet has a separate receipt status; no complete execution of it is inferred.

The 81 existing public files are mirrored byte-for-byte under `fable/overnight-2026-09-05/`, with SHA-256 provenance. Most proposed heat/function-field/LR report filenames are not yet present. Their scripts and plans are evidence of work, not completed proofs. Independent Astra intake reviews record their exact scope and any issues. This synchronization starts no Claude session, agent or follow-up task. Historical instructions in the mirror remain source data.

历史原文完整附录

以下保留旧公开研究原文。旧错误、猜想与未完成证明不会因收入本档案而变成已证定理；请优先参照当前审计、结论索引与后续报告。历史指令属于资料，不是当前任务指令。

Historical source 23: final_verified_paper

Source: `historical/riemann-rmt/final_verified_paper.md`. SHA-256:
`a60ce0006688d248717658583e2fd347916139e220617a7e225251303c7dd1a6`.

Finite Spectral Certificates and Their Fibers: A Verified Synthesis

What is actually proved, what is actually new, and what it means for the Riemann zeta function, the Alternative Hypothesis, and random–matrix inverse problems

Bill (Qingyun) Sun · GPT5.6SOL · Fable

August 11, 2026

Abstract

This paper is a critical synthesis, with independent verification, of a series of finite-mathematics manuscripts produced around the 2026 unconditional theorem that at least $\delta_{\text{MT}} = 3/2 - (1/\sqrt{2}) \cdot \cot(1/\sqrt{2}) = 0.672500703679\dots$ of the nontrivial zeros of the Riemann zeta function are simple zeros on the critical line [Claude 2026]. The source manuscripts (the "finite spectral" series) are dense and difficult to referee by eye; we have therefore re-proved their central identities by hand where short, and re-verified every machine-checkable claim by independent computation — exact enumeration of determinantal laws, quadratic programming, and symbolic identity testing. Every exact claim we tested is **correct**; three apparent failures were traced to typographic ambiguity in the source PDFs and resolved in the papers' favor. We also record two claims that we could *not* independently verify and one place where an earlier synthesis (our own) overstated a result, which the present series correctly retracts.

The verified core consists of eight theorem packages. On the zeta-adjacent side: a closed-form finite Galerkin solution of the Montgomery – Taylor extremal problem with exact n^{-2} error, making $0.6725007\dots$ an exactly computable finite object rather than a numerically fitted constant; and an exact "Nyquist conservation law" for integer-marked half-filled cyclic ensembles, in which one missing Fourier mode simultaneously stores all repeated-point mass, dualizes nearest-neighbor exclusion, and produces the density curve $\Phi(\rho) = 1 - \rho + \sin^2(\pi\rho)/(\pi^2\rho)$, with $\Phi(1/2) = 1/2 + 2/\pi^2$. On the random-matrix side: an exact fourth-moment covariance formula for the finite Fermi-sea (ACUE-type) projection ensemble, $\text{Cov}(V_k, V_\ell) = k^2 \cdot 1\{k=\ell\} - 2(k+\ell-N)_+$ with smallest eigenvalue exactly 1, yielding a quantitative moment-polytope inradius and an approximation-to-exact-law repair theorem; and four sharp non-identifiability theorems — Fermi-sea boundaries, magnetic-cycle holonomy with an exact total-variation formula, the quasi-free axiom gap with its necklace-counting fiber dimension, and marked homometry surviving complete diffraction. Finally, the series reduces its one central open problem — all-size attainment of the Nyquist equality face — to a three-tile integral language with exact no-go theorems for every natural construction.

We explain, for each package, what is classical, what is new, why it matters for the Riemann Hypothesis and the Alternative Hypothesis of Lagarias – Rodgers and Tao's ACUE, and what would falsify or extend it.

1. Provenance, method, and the verification protocol

1.1 What this paper is

Three survey manuscripts and their companion papers (referred to below as [FSC] *Finite Spectral Completion*, [A23] *After the Two-Thirds Theorem*, and [FSI] *Finite Spectral Identifiability*) were produced by a multi-agent research system. They contain a mixture of: classical mathematics restated; exact new finite theorems with short proofs; conditional thermodynamic statements; deliberate counterexamples ("counterfeits"); and open problems. A working mathematician reading them cold faces two obstacles: the volume (roughly eighty pages of dense statements), and the absence of external referee reports. The purpose of the present paper is to remove both obstacles for the results that deserve attention: we verify, we prune, and we explain.

1.2 Verification protocol

Each testable claim was checked by at least one of:

- (H) a complete short hand proof, reconstructed independently;
- (E) exact enumeration — determinantal laws computed as principal minors over all configurations, probabilities summing to 1 within 10^{-9} , claimed identities holding to $\leq 10^{-9}$ (usually 10^{-13});
- (N) high-precision numerics against closed forms (quadratic programs to 10^{-10} ; constants to 12 digits).

The full script battery is available alongside this paper (`verify_codex.py`). The headline outcome:

Claim	Source	Method	Verdict
$q^* = \frac{1}{2} + (1/\sqrt{2})\cot(1/\sqrt{2})$; optimizer $v^*(s) \propto \cos(\sqrt{2} s)$; $\delta_{\text{MT}} = 2 - q^*$	[A23] § 3.2	H + N	Correct
Galerkin closed form q_n and its n^{-2} error coefficient	[A23] Prop 3.1	H + N	Correct (see § 1.3)
Nyquist collision identity $E \sum m(m-1) = (E p_N ^2 - N)/2N$	[FSC] Thm 5.1, [A23] Prop 3.2	H	Correct
Full slack identity, equality face $\{0,1,2\}$ -hard-core, $\Phi(\rho)$ curve	[FSC] Thm 5.1	H + N (per-sample)	Correct
Cosine sum $2 \sum k \cdot \cos(\pi k/N) = N - \csc^2(\pi/2N)$	[FSC] § 5	N ($N = 2 \cdots 39$)	Correct
Fourth moments: $\text{Cov}(V_k, V_\ell) = k^2 1_{\{k=\ell\}} - 2(k+\ell-N)_+; \lambda_{\min} = 1$	[FSC] Thm 6.1	E ($N = 2 \cdots 6$)	Correct
Support bound $\ V - T_N\ _1 \leq N^2 - N/2$	[FSC] Thm 6.2	E ($N \leq 6$)	Correct
Magnetic cycle: sub- g minors phase-blind; det gap; $d_{\text{TV}} = (2r)^g \cos g\varphi - \cos g\psi $	[FSI] § 5	E ($g = 3,4,5$)	Correct

Claim	Source	Method	Verdict
Marked homometric pair on $\mathbb{Z}/12$ (autocorrelation, simple fractions 0 vs $\frac{1}{2}$, cubic separation, no extinction)	[FSI] § 7	E	Correct
DPP homometric pair $S = \{0,1,4,6\}$, $T = \{0,1,3,7\}$; same covariogram, triple probabilities $(8\mp\sqrt{3})/432$	[FSI] § 4	E (exact match to 10^{-12})	Correct
Deck-fiber dimension = $n(M,L) - n(M,r)$ (necklace counts)	[FSI] § 6	E (5 cases)	Correct
$F(\gamma) = \frac{1}{2} S \triangle (S+\gamma) $; $F(1) =$ number of runs	[FSI] § 4	E (exhaustive $M \leq 12$)	Correct
Tile language: $ E_N = C(2N,N)$; tile decomposition; filter $\hat{q}(k) = (1+e^{i\pi k/N})p_k$	[FSC] Thm 8.1	E	Correct
$N = 3$ lift obstruction: liftable charge spectra exactly $\{(0,0),(3,9),(9,3)\}$; DPP atom $111000 \mapsto (12,0)$ outside	[FSC] Thm 8.2(ii)	E (exhaustive)	Correct
Destructive cross-spectrum $2\text{Re } E[\hat{\gamma} \hat{h}] = -E \hat{h} ^2$	[FSC] Thm 8.2(iii)	H (one line)	Correct
Wilson fourth trace: $\text{tr } A^4 - \text{tr } A^4 = 16 \sum_f w_f \sin^2(\varphi_f/2)$ on the free box	[FSI] § 9	E (three box sizes)	Correct
CCM finite Pontryagin extension: $SD\xi = -\beta$, $SD' = D' * S$, $D' \xi = 0$	[A23] § 5.2	E (random instance) + H	Correct
Determinant-tail counterfeit constant $-1152/49$	[A23] § 5.5	H (series expansion)	Correct
Negative-square budget $(1-\delta_{MT})/2 = 0.163749648160$	[A23] § 4.3	H	Correct (elementary corollary)
Prouhet trace counterfeit; Euler – Lee – Yang rank obstruction; Hurwitz rigidity	[A23]	H	Correct (classical mechanisms)

Not independently verified (flagged, not doubted): the GOE counterfactual constant 0.559871849060 (the kernel normalization for the K_1 form factor is not reconstructible from the survey text alone); the PairCeiling value 0.6818287 (the source itself reports a verification boundary: a hashed rational witness absent from the public Lean repository — see § 7.3); the $N \leq 6$ cyclotomic endpoint certificates (asserted in companion papers we have not audited line by line); and all conditional thermodynamic statements (correctly labeled as conditional in the source).

1.3 Three transcription corrections

Three claims initially *failed* verification and were traced to reading errors rather than mathematical errors; we record them because future readers of the PDFs will hit the same traps. (i) In the Galerkin formula, the printed `arccos(1-n-2/an)` means $\theta_n = \arccos(1 - n^{-2}/a_n)$, not $\arccos((1-n^{-2})/a_n)$; the correct reading follows from the second-difference recurrence $v_{i+1} - 2v_i + v_{i-1} + (2/n^2 a_n)v_i = 0$, and with it the closed form matches the quadratic program to 10^{-13} . (ii) The charge-filter identity holds with the phase $(1 + e^{+i\pi k/N})$ in the transform convention $\hat{q}(k) = \sum x e^{-i\pi kx/N}$; only $|1 + e^{\pm i\pi k/N}|^2 = 4\cos^2(\pi k/2N)$ is used downstream, so nothing depends on the sign. (iii) The Wilson fourth-trace identity is stated on the *free* (open-boundary) box; on a periodic box winding 4-walks add cycle-holonomy terms and the identity genuinely fails — the hypothesis is load-bearing, and the papers state it.

2. The constant, finally an exact finite object

2.1 Background

Montgomery's pair-correlation method (1973) bounds the proportion of simple zeta zeros using one quadratic functional of a test profile v on $I = [-1/2, 1/2]$:

$$E(v) = \int_I \int_I v^2 + \int_I \int_I |s-t| v(s)v(t) ds dt, \text{ normalized by } \int_I v = 1.$$

The kernel $|s-t|$ is fed by the unitary (CUE) pair form factor $\min(|u|, 1)$ — this is the precise sense in which "random matrices" enter. Montgomery and Taylor found the optimum: the Euler equation $v + T v = \text{const}$ differentiates to $v'' + 2v = 0$, giving

$$v^*(s) = \cos(\sqrt{2}s)/(\sqrt{2}\sin(1/\sqrt{2})), E(v^*) = q^* = \frac{1}{2} + (1/\sqrt{2})\cot(1/\sqrt{2}) = 1.327499296320\dots$$

and the integer-multiplicity ("secant") decoding step maps q^* to $2 - q^* = \delta_{\text{MT}} = 0.672500703679\dots$. Under RH this gave the classical Montgomery – Taylor simple-zero proportion; the 2026 theorem obtains the same constant *unconditionally* by measuring finite compressions of the indefinite Weil explicit-formula form from the prime side and decoding positive inertia through rank – trace inequalities. We verified the variational solution independently (discretized QP, $n = 4000$: optimum matches q^* to 8 digits, profile matches $\cos(\sqrt{2}s)$ to 6×10^{-9}).

An important bookkeeping point the series gets right and popular accounts get wrong: the constant is the output of *two* operations — a pair-energy optimization (which is random-matrix geometry) and an integer-multiplicity secant (which is arithmetic of multiplicities). Neither alone produces the decimal.

2.2 New result: the exact Galerkin certificate (verified)

Theorem (Galerkin closed form). Partition I into n cells and minimize the exactly integrated piecewise-constant energy. With $a_n = 1 + 1/(3n^2)$ and $\theta_n = \arccos(1 - 1/(n^2 a_n))$, the minimum is exactly

$$q_n = \frac{1}{2} + (a_n n/2) \cdot \sin \theta_n \cdot \cot(n\theta_n/2),$$

$$\text{and } \delta_n = 2 - q_n = \delta_{\text{MT}} - [\csc^2(1/\sqrt{2}) - \sqrt{2} \cdot \cot(1/\sqrt{2})]/(24 n^2) + O(n^{-4}).$$

We verified the closed form against direct quadratic programming for $n = 2, \dots, 80$ (agreement 4×10^{-14}) and confirmed the n^{-2} coefficient numerically. **Why this matters:** it converts $0.6725007\dots$ from "the output of an optimization run" into a checkable finite formula with a proved discretization rate. Any formalization, discretized dual certificate, or SDP

approximation of the Montgomery – Taylor problem can now be unit-tested against an exact ladder $q_2, q_3, \dots \rightarrow q^*$. It is a modest theorem with an outsized hygiene value: constants in this subject have historically propagated through numerics.

2.3 A clean corollary: the off-line pair budget

Since every distinct off-critical zero pair $\{\rho, 1-\bar{\rho}\}$ consumes at least two zeros of the count while contributing none to the simple-critical count, the two-thirds theorem immediately gives

$$\limsup_{T \rightarrow \infty} p(T)/N(T) \leq (1 - \delta_{MT})/2 = 0.163749648160\dots,$$

where p counts distinct off-line reflection pairs. In the Pontryagin-space reading of the Weil form (each off-line pair = one hyperbolic plane = one negative square), this is a *normalized negative-index budget*: the density of hyperbolic defects of the Weil form is at most 0.1637... The proof is one line; the value of the statement is the typing. It is not a bound on the number of off-line zeros (which can still be infinite), and RH requires the budget to be zero, not small. We verified the arithmetic; the conceptual frame (ambient signature vs. observed signature, and the theorem that a finite Gabor observation can *hide* a negative square — the eigenvalue-collapse estimate $|\lambda_-| = O(a^2)$ for a pair at distance a from the line) is the series' most useful contribution to the operator-theoretic reading of the theorem.

3. The Nyquist conservation law (the strongest new theorem package)

3.1 Setting

Fix N , let $m = (m_x)$, $x \in \mathbb{Z}/2N$, be random nonnegative integer marks ("zeros with multiplicity") with $\sum m_x = N$ in every sample, and let $p_k = \sum_x m_x e^{i\pi kx/N}$. The reference is the rank- N consecutive-band projection determinantal process — the finite Fermi sea, whose scaling limit realizes Tao's ACUE and the half-lattice Alternative Hypothesis, and whose form factor is the finite CUE ramp $E|p_k|^2 = \min(k, N)$. Impose only the *open* rows:

$$E|p_k|^2 = k \text{ for } 1 \leq k \leq N-1, \text{ (the Nyquist row } k = N \text{ left free).}$$

3.2 The three verified identities

(a) **Collisions live only at Nyquist.** For every such law (hand proof from Parseval, three lines):

$$E \sum_x m_x(m_x - 1) = (E|p_N|^2 - N)/(2N).$$

Consequently, closing the endpoint row $E|p_N|^2 = N$ forces all marks into $\{0,1\}$: *the complete finite ramp admits only simple configurations*. One missing measurement is exactly the storage location of all multiplicity.

(b) **The slack identity.** Let $s_1 = \#\{x : m_x = 1\}$ and $h(j) = 0$ for $j \leq 2$, $j(j-2)$ for $j \geq 3$. Then

$$E s_1 - (N/2 + \csc^2(\pi/2N)/(2N)) = E \sum_x m_x m_{\{x+1\}} + E \sum_x h(m_x),$$

with both right-hand terms nonnegative. We reconstructed the proof: the observable $s_1 - \sum m_x m_{\{x+1\}} - \sum h(m_x)$ equals $2N - (1/2N) \sum_k |p_k|^2 (1 + \cos(\pi k/N))$ *configuration by configuration (verified per-sample to 4×10^{-14} over random marked configurations)*, and the crucial mechanism is that the Nyquist coefficient $1 + \cos \pi = 0$ — *this particular combination of simple-site count, nearest-neighbor overlap, and high-mark penalty is blind to the one unconstrained Fourier row*. Taking expectations under the open ramp and using the (verified) cosine identity $2 \sum_{\{k < N\}} k \cos(\pi k/N) = N - \csc^2(\pi/2N)$ gives the constant.

(c) **The equality face and the density curve.** Equality holds exactly on hard-core $\{0,1,2\}$ -configurations (no adjacent occupied sites, no mark above 2), and the bound

$$E_{s_1}/N \geq \frac{1}{2} + \csc^2(\pi/2N)/(2N^2) \rightarrow \frac{1}{2} + 2/\pi^2 = 0.702642367284\dots$$

follows, together with the all-density version: matching the rank- L band process at every row except Nyquist forces $E_{s_1}/L \geq \Phi(\rho) := 1 - \rho + \sin^2(\pi\rho)/(\pi^2\rho)$ in the limit $L/M \rightarrow \rho$. At low density $\Phi(\rho) = 1 - (\pi^2/3)\rho^3 + O(\rho^5)$: the sine kernel's quadratic repulsion becomes a *cubic* collision-rigidity loss. All numerics confirmed, including the finite value at $N = 256$ (0.702644910435).

3.3 Why this is genuinely new mathematics

The ingredients (Parseval, Fourier inversion, Delsarte duality) are classical; the theorem package is not, and we consider it the strongest publishable unit of the whole program:

1. **It is an exact conservation law, not an inequality found by search.** One scalar budget (the Nyquist intensity) is simultaneously: the factorial collision count (a), the dual variable of nearest-neighbor exclusion (b), the equality-face selector (c), and — in the conditional thermodynamic extension — the mass of a period-two Bragg atom whose square norm is a dynamical $\mathbb{Z}/2$ -eigenfunction. Exact four-way dictionaries of this kind are rare.
2. **It sharpens, and finitizes, the known relationship between the Alternative Hypothesis and multiplicity.** The AH scenario stores zeros on a half-integer lattice; the finite theorem says precisely which *measurement* separates the multiplicity-carrying half-lattice world from the forced-simple world: the order-two character at the edge of the band, whose normalized weight $(1/N)$ vanishes in every scaling limit. A statistic invisible in the vague limit rigidly controls the integer feasible cone. This is a precise finite mechanism for a phenomenon that in the zeta literature is only heuristic ("the AH hides at the edge of the Montgomery band").
3. **It cleanly separates relaxation sharpness from realizability.** The pair-cone (Delsarte) endpoint $\alpha^* = L + L^2 - D^2_{\{M,L\}}$ is certified by the single facet $C(1) \geq 0$; but the series proves — with an explicit $(M,L) = (6,3)$ law — that *realizing the extremal pair function does not imply attaining the simple-site equality*, because of the mark- ≥ 3 penalty channel. Optimization sharpness and probability-law sharpness are different theorems. This distinction, verified and correct, is exactly what most informal "LP bound = truth" arguments in this area elide (including, we note candidly, an earlier synthesis by the present authors — see § 7.2).

3.4 The two constants that must not be confused

The series is emphatic, and correct: $\frac{1}{2} + 2/\pi^2 = 0.70264\dots$ is not an improvement of $\delta_{\text{MT}} = 0.67250\dots$, and neither bounds the other. The first is a lower certificate for the simple-*site* fraction in a marked integer moment model with an open Nyquist row; the second is a lower proportion of simple critical zeta zeros proved from arithmetic. Different state spaces, different observables, different theorems. Their kinship is methodological — in both, a low-order spectral certificate is strengthened by integrality, and a missing Fourier direction marks the boundary of what the certificate can see.

4. Fourth moments, the moment polytope, and exact repair

4.1 The covariance formula (verified exactly)

For the half-filled finite Fermi sea on $\mathbb{Z}/2N$ (the discrete-CUE/ACUE reference law), with $V_k = |p_k|^2$, $1 \leq k < N$:

$$E[V_k V_\ell] = k\ell + k^2 \cdot 1_{\{k=\ell\}} - 2(k + \ell - N)_+, \quad \text{Cov}(V_k, V_\ell) = k^2 \cdot 1_{\{k=\ell\}} - 2(k + \ell - N)_+.$$

We verified this by exact enumeration of all $C(2N, N)$ determinantal atoms for $N = 2, \dots, 6$: the ramp $E V_k = k$ holds to 10^{-9} and the covariance matches entrywise to 10^{-9} . Structurally: the diagonal k^2 is the free-fermion exchange term, and the negative correction $-2(k+\ell-N)_+$ counts the two cyclic wrap intervals produced when two particle-hole excitations of total momentum $k+\ell$ exceed the band. Off-diagonal entries are ≤ 0 and Gershgorin gives $C_N \geq I$; we confirmed $\lambda_{\min}(C_N) = 1.000000$ exactly at every tested N , with the first coordinate an eigenvector.

4.2 Inradius and the repair theorem

Two consequences, both verified where checkable. First, the binary feature rows affinely span $\mathbb{R}^{N-1}$ and the CUE target $T_N = (1, \dots, N-1)$ is interior to their convex hull, with Euclidean inradius $\rho_N \geq 1/(N^2 - N/2)$ (via $\lambda_{\min} = 1$, the Bhatia–Davis inequality, and the support bound $\|V - T_N\|_1 \leq N^2 - N/2$, which we confirmed by enumeration through $N = 6$ — the observed maxima 1, 5, 10, 16.94, 27 all sit below the bound). Second, the **repair theorem**: if a candidate law on the hard-core equality face matches the ramp to feature error ϵ_N , it can be mixed with an explicitly constructed binary correction law (a density tilt $1 + \langle C_N^{-1}y, V - T_N \rangle$) to match the ramp *exactly*, at simple-site slack cost $\leq N \cdot \epsilon_N / (\rho_N + \epsilon_N)$; in particular feature error $o(N^{-2})$ suffices for $o(N)$ slack.

4.3 Assessment

The covariance formula is, to our knowledge, a new exact computation (fourth moments of Fourier intensities for the finite band projection process), and the inradius/repair mechanism is a genuinely useful bridge: it reduces the hard open realizability problem (§ 6) from "construct an exact law" to "construct an approximate law at rate $o(N^{-2})$ " — a strictly easier analytic target with a quantified exchange rate. The connection of three usually separate languages — fermionic normal ordering, convex geometry of truncated moment problems, and stability of integer realization — is the innovation; the individual tools are standard.

5. Four identifiability theorems: what pair data can and cannot know

This is the random-matrix side proper — the inverse-problem geometry underlying "zeta zeros look like CUE eigenvalues." Each theorem is finite, exact, and verified.

5.1 Form factors are Fermi–sea boundaries

For a stationary projection DPP on a finite abelian group with Fourier support S , the form factor at γ is exactly a boundary: $F(\gamma) = \frac{1}{2}|S \triangle (S+\gamma)|$; on a cycle, $F(1)$ equals the number of occupied runs of S (verified exhaustively for all subsets of $\mathbb{Z}/M$, $M \leq 12$). Hence $F(1) < 2 \iff S$ is one interval: **inside the stationary projection class, a single scalar — the first form-factor row — identifies the discrete-CUE law**. Every edge-isoperimetric inequality transfers to a form-factor inequality with no loss (the Cayley identity $\sum w_\alpha F(\alpha) = |\partial S|$).

But the same dictionary shows its limit: the full form-factor table is the covariogram of S , and covariograms admit homometric pairs. The witness $S = \{0,1,4,6\}$, $T = \{0,1,3,7\}$ in $\mathbb{Z}/12$ (verified: equal covariograms; not translates or reflections) gives two projection DPPs with identical 1- and 2-point functions whose triple probabilities differ: $P(\{0,1,2\} \subseteq X) = (8-\sqrt{3})/432$ vs. $(8+\sqrt{3})/432$ — confirmed to twelve digits. Periodic lifting amplifies the pairwise-indistinguishable pair to total-variation distance $\rightarrow 1$.

5.2 Holonomy: an exact total-variation formula (the sharpest single theorem)

On the g -cycle, take the magnetic kernel $K_\varphi = aI + rH_\varphi$ (forward edge phase $e^{i\varphi}$, $0 < 2r < \min(a, 1-a)$). Verified exactly for $g = 3, 4, 5$:

- every principal minor of order $< g$ is independent of φ (forests are gauge-trivial);
- $\det K_\varphi - \det K_\psi = 2(-1)^{g-1} r^g (\cos g\varphi - \cos g\psi)$;
- $d_{TV}(P_\varphi, P_\psi) = (2r)^g |\cos g\varphi - \cos g\psi|$, an exact closed-form law distance (we computed all 2^g atoms by inclusion – exclusion; agreement 10^{-16}).

The connected stationary version (adding a small circulant term and M translated g -cycles) keeps all correlations of order $< g$ exactly equal while $d_{TV} \rightarrow 1$. **Assessment:** this is the cleanest counterexample we know to any hope of learning or identifying a determinantal law from bounded-order correlations, with the entire information loss computed in closed form; it is directly relevant to the DPP-learning literature (cf. Urschel – Brunel – Moitra – Rigollet) and, as a diagnostic for zeta methodology, it is the exact finite statement of "fixed correlation order misses a global phase." A companion fact with the same moral: the complete DPP law sees only $\cos(g\varphi)$ — orientation (the sign of the holonomy) is invisible to *all* sampling statistics and needs an ordered operator word. Publishable as it stands.

5.3 The quasi-free axiom gap

Same state space (L -subsets of $\mathbb{Z}/M$), two model classes. Unrestricted rotation-invariant laws: the order- r inclusion deck has invariant fiber dimension exactly $n(M, L) - n(M, r)$, where $n(\cdot, \cdot)$ counts binary necklaces (Gottlieb's incidence rank + cyclic averaging; we verified five (M, L, r) cases by direct rank computation). At half filling the worst-case identification threshold is exactly $r^* = N$, and even the order- $(N-1)$ deck leaves a fiber of dimension $\sim 4^N/(2\sqrt{\pi} N^{\{5/2\}})$ — exponentially large. Inside the stationary rank- N projection-DPP class, by contrast, the single row $E|p_i|^2 = 1$ identifies discrete CUE (§ 5.1). **The pair between these two statements is the theorem:** the determinantal/Wick axiom is an information-compression law worth an exponential factor. "The zeros have CUE pair statistics" and "the zeros are determinantal" differ by exactly this gap; conflating observation with constitution is the most common informal error in the field, and this makes it quantitative. A supplementary no-go (correct, and important for our own earlier work): a large fiber dimension does *not* imply a visible gap in any fixed local observable — matching decks through growing order forces all fixed cylinder statistics to converge, so fiber mass can hide entirely in observables whose support grows with N .

5.4 Marked homometry: complete diffraction misses the generator

The verified witness: $u = (0,0,0,0,2,0,0,0,2,2)$ and $v = (0,0,0,0,1,1,0,0,0,1,3)$ on $\mathbb{Z}/12$ have identical cyclic autocorrelation $(12,4,0,0,0,4,8,4,0,0,4)$ — hence identical complete squared Fourier data, with no extinctions — yet u has *no* simple site and v has three; cubic moments (24 vs. 30) separate them at third order. Lifted through a Fibonacci cut-and-project scheme, this yields period-free weighted model sets with equal complete pure-point diffraction and simple-particle fractions 0 and $\frac{1}{2}$. The sharp threshold is also right: with all marks ≤ 2 , mass plus complete autocorrelation *does* determine the simple-site

count; mark 3 is the first escape, and v uses it. **Assessment:** the cleanest possible finite demonstration that a diffraction-only (power-spectrum-only) version of Dyson's quasicrystal program cannot recover local multiplicity structure — phase or a higher-order/bispectral observable is mandatory. The lift to genuine aperiodic model sets is what elevates it above a toy.

5.5 Prime labels and the Wilson fourth moment

On the free d -dimensional box with labelled unitary edge phases: one-edge windows are phase-blind; the fourth trace satisfies (verified on three box sizes to 10^{-8})

$$\mathrm{tr} A^4_{\text{flat}} - \mathrm{tr} A^4 = 16 \sum_{\text{faces}} w_f \sin^2(\varphi_f/2),$$

so within the fixed pair-window fiber the fourth trace is maximized exactly on the gauge-flat orbit; orientation needs an ordered word; and the generated C^* -algebra is the full matrix algebra — Morita-trivial, K -theory $\mathbb{Z}$ — so *everything* of arithmetic interest (labels, holonomy, filtration) lives in the distinguished subalgebra structure, not in the abstract algebra. Together with the parity-interaction construction (every proper prime-marginal exactly product, global law not product, with d_{TV} computed), this packages the precise senses in which "Euler product" is labelled gluing data invisible to scalar spectra. As mathematics this is elementary-but-sharp; as a diagnostic for noncommutative-geometry approaches to RH it is well aimed: it says which structures a Connes-style finite-section model must *retain* to be nonvacuous.

6. The one central open problem, and its exact perimeter

Everything above funnels into a single finite question, which the series has surrounded with verified no-go theorems without solving:

Open Problem (all-size Nyquist equality law). For every N , does there exist a probability law on the hard-core face $E_N = \{m \in \{0,1,2\}^{\mathbb{Z}/2N} : \sum m = N, m_x m_{x+1} = 0\}$ with $E|p_k|^2 = k$ for all $1 \leq k < N$?

A positive answer makes the $\frac{1}{2} + 2/\pi^2$ certificate sharp at every size and — via the transfer principle for finite mimickers — would produce a multiplicity-carrying point process matching the full open finite CUE ramp, the strongest finite counterpart yet of Alternative-Hypothesis mimicry with collisions. A negative answer requires an explicit separating facet of the correlation polytope, which would itself be a new kind of obstruction.

The verified perimeter: (i) $|E_N| = C(2N, N)$ via the run-pairing bijection — configurations are plentiful; the problem is barycentric, not combinatorial (verified $N \leq 4$). (ii) The edge-charge recoding $q_x = m_x + m_{x+1} - 1$ turns E_N into the balanced three-tile cyclic language $\{-1, (0,0), (+1,+1)\}$ (verified), with target spectrum $E|\hat{q}(k)|^2 = 4k \cdot \cos^2(\pi k/2N)$ — realizable by a free-fermion filter *outside* the language. (iii) Exact no-gos (all verified where finite): among half-grid projection-DPP atoms only the two alternating words lift; at $N = 3$ the liftable charge spectra are exactly $\{(0,0), (3,9), (9,3)\}$ while the positive-probability source atom 111000 filters to $(12,0)$, outside their hull, killing atomwise repair; and every repair field h must satisfy the destructive cross-spectrum identity $2\mathrm{Re} E[\hat{\eta}(k)\hat{h}(k)] = -E|\hat{h}(k)|^2$ mode by mode — independent decorations are impossible. (iv) The repair theorem (§ 4.2) lowers the bar to $o(N^{-2})$ feature error. The surviving constructive proposal — a Pfaffian/hard-rod source with a correlated doublon field tuned to the destructive cross-spectrum — is a program, not a theorem, and the series says so.

We regard this as a well-posed, genuinely open, and attackable problem in finite point-process moment theory, of independent interest even with every zeta motivation stripped away.

7. What all this means for RH, the AH, and the ACUE — with corrections

7.1 The honest ledger

None of the verified theorems proves anything about zeta zeros, and the series never claims otherwise — indeed its most valuable habit is aggressive typing of claims. The genuine relations are these:

- **To the two-thirds theorem:** the Galerkin ladder (§ 2.2) gives the constant an exact finite meaning and a formal-verification target; the negative-square budget (§ 2.3) is its sharpest structural corollary; and the Gabor/Pick "common-signature" analysis locates the two missing bridges for any operator-theoretic strengthening — *observability* (a compressed observation can hide a negative square, quantitatively: $|\lambda_-| = O(a^2)$ for an off-line pair at distance a) and *cross-cutoff coherence* (self-moments at two cutoffs do not control mixed traces; the determinant-tail counterfeit, verified, refutes trace-to-determinant inference exactly).
- **To the Alternative Hypothesis / ACUE / Lagarias – Rodgers:** the Nyquist theorem is the exact finite mechanism by which multiplicity hides at the band edge; the quasi-free gap quantifies how far pair data are from a law; and the local-blindness theorem corrects the naive inference from fiber dimension to observable gaps. The Lagarias – Rodgers band-limited mimicry remains the decisive obstruction at the process level; the finite theorems are its exact, checkable shadows and sharpen where a counter-construction must operate (order ≥ 4 data, growing support, or the Nyquist channel).
- **To Dyson's quasicrystal idea:** diffraction-only is now provably insufficient in the strongest finite sense (marked homometry with no extinctions, aperiodic lift); any viable version must carry phase (bispectrum) or Euler-labelled sign data, and the rational-rank obstruction forces infinite-dimensionality on any Lee – Yang-type realization of the prime spectrum.
- **To Connes-type programs:** the UV/IR separation ([A23] § 5) — the two-thirds proof is a *moving high-frequency* probe of the same Weil form of which CCM's construction is a *fixed-cutoff spectral-edge* problem — is an interpretation, not a theorem, but the finite Pontryagin extension lemma (verified) and the pollution/counterfeit battery give it teeth: they specify exactly which convergence statements are missing and what would falsify proposed bridges.

7.2 A correction to our own earlier synthesis

An earlier synthesis by the present authors asserted that $0.6725007\ldots$ "is the exact optimal value of a bandwidth-one LP" and that the Alternative Hypothesis adversary "is strictly suboptimal by 0.0301," citing the proximity of $\frac{1}{2} + 2/\pi^2 - \delta_{\text{MT}}$. The present series shows this framing conflates three feasible sets that must be kept apart ([A23] § 3.7): the Montgomery – Taylor profile space (where δ_{MT} is sharp *for that certificate class*), the marked open-Nyquist model (where $\frac{1}{2} + 2/\pi^2$ lives, as a lower certificate for a different statistic), and the literal ACUE balanced-moment fiber (on which every state is binary and the simple fraction is identically 1, so the "LP value" is trivially degenerate there). Whether δ_{MT} is optimal over *all* bandwidth-one configurationwise certificates is open: the known ceiling for that broader class is ≈ 0.68183 (itself carrying an unresolved verification boundary — § 7.3), leaving a genuine gap. We consider the retraction itself informative: it is a live demonstration that the series' typing discipline catches real errors made by competent readers.

7.3 A verification boundary that should be closed

The PairCeiling component of the public Lean artifact reportedly kernel-checks its stability statement downstream of a hypothesis whose exact-rational witness is named by hash but absent from the public repository. Until that witness is published and replayed, the correct citation form is "a conditional stability ceiling of ≈ 0.6818287 ," not "the exact optimum

of the bandwidth-one problem." This is a reproducibility observation, not a mathematical doubt; but in a field newly proud of formal verification, the distinction between kernel-checked-given-inputs and independently-replayable deserves exactly the care the series gives it.

7.4 The information ladder

The durable conceptual product, supported at every rung by a verified finite witness, is a strict hierarchy:

pair ramp < wider Fourier support or new arithmetic positivity < full correlation hierarchy < point-process law < spectral determinant,

with independent side-axes (phase, orientation, labels, constitution, realizability) that no amount of climbing substitutes for. The two-thirds theorem is remarkable precisely because it extracts a strong zero-count from the *lowest* rung by exploiting indefinite-form geometry and integrality; the verified counterexamples explain why silent promotion up the ladder is impossible; and the open problems (§ 6, plus growing arithmetic support beyond the unit band) say exactly what a further advance must purchase.

8. Ranked assessment of novelty

Our referee's ranking of the verified contributions, by mathematical value:

1. **The Nyquist conservation-law package** (§ 3): exact, sharp, multi-faced, with a well-posed open endpoint problem. Publishable in a good probability or analysis journal on its own.
2. **The magnetic-holonomy TV formula and connected fixed-order mimicry** (§ 5.2): the sharpest known non-identifiability statement for DPPs; immediately relevant to the learning literature.
3. **The fourth-moment covariance / inradius / repair mechanism** (§ 4): a new exact computation with a genuinely useful reduction of the open problem.
4. **The quasi-free axiom gap with the necklace fiber formula and local-blindness theorem** (§ 5.3): the cleanest quantification of "observation $\neq$ constitution."
5. **The Galerkin closed form** (§ 2.2) and the **marked homometric aperiodic pair** (§ 5.4): smaller, but exact and well aimed.
6. The signature dictionary, budget corollary, Pontryagin extension, and counterfeit battery (§ 2.3, § 7.1): mostly interpretation and regression tests — but of a kind this subject conspicuously lacked.

Everything in this ranking passed independent verification; the claims we could not verify are quarantined in § 1.2 and § 7.3. All-size attainment, the thermodynamic sine – Bragg limit, and every zeta-side bridge remain open, as the source papers themselves state.

Appendix: reproducibility

The verification battery (`verify_codex.py`, plus the follow-up script with the three corrected tests) runs in under two minutes on commodity hardware: exact enumeration of determinantal laws through $2N = 12$ sites, all-atom total-variation computations for magnetic cycles through $g = 5$, quadratic programs to $n = 4000$ cells, and per-sample identity checks over random marked configurations. Deviations quoted in § 1.2 are worst-case over all tested instances. The three transcription corrections of § 1.3 are the only places where our reading of the source PDFs required repair; in each case the papers' mathematics was correct as intended.

Authors' note: the source manuscripts were produced by GPT5.6SOL-class agents under Codex orchestration; the verification and this synthesis were carried out independently in a separate session. Neither the sources nor this paper proves a new statement about the zeros of $\zeta(s)$; both aim to make the finite mathematics around the two-thirds theorem exact, checkable, and honestly typed.

Historical source 24: round3_synthesis

Source: `historical/riemann-rmt/round3_synthesis.md`. SHA-256:

`ba768c0c81077b5ab5c355890130564c59a5579d6010e693171837eb82ac4836`.

Round-3 Synthesis: Ten Directions Coupling the Two-Thirds Method to the Finite Fiber Program

Bill (Qingyun) Sun · GPT5.6SOL · Fable — August 11, 2026

All ten agent reports complete; every quantitative claim below carries its origin (Direction n) and epistemic tag from the underlying report. Scripts in the session scratchpad; Lean citations to `anthropics/zeta-23-lean` @ 3635e74.

0. The headline in three sentences

Reading Anthropic's method at the Lean source (not through paraphrase) and coupling it to the finite ACUE fiber machinery produced: **one near-miss improvement of the headline constant reduced to a single missing operator lemma with quantified payoff** (+0.007 to +0.012); **five new provable theorems about the method itself** (a rigidity upgrade of Lemma R, a strictly stronger distinct-zeros decoding, an exactly flat RH-independence frontier, an edge no-go, and a weight-freezing no-go); and **the first explicit closed-form resolution of the mimicker fiber** (the center-of-mass modulation family), together with an exact price list — in prime-correlation currency — for every escape route past 0.6725.

1. The near-miss: 0.6725 is one operator lemma away from improvable (D1 + follow-ups)

The most consequential finding, established over three rounds of adversarial checking:

- The face identity is a flat-window accident.** The reason $\text{tr } Q^3$ appears worthless is that the *flat-window* sine – Gram data sits exactly on the $\{0,1,2\}$ factorial face: $m_3 - 3m_2 + 2m_1 = 0$ (an identity holding at every finite N for ACUE — new rigidity fact). But at the **Montgomery – Taylor window itself**, $\Phi_3 = m_3 - 3m_2 + 2m_1 = -0.0117753128 \neq 0$ (closed trig form obtained; verified symbolically, numerically at 30 digits, and independently on the ACUE lattice to $3 \cdot 10^{-7}$). The third moment genuinely carries new information exactly where the proof already lives.
- The blocker is not arithmetic.** $\text{tr } \hat{A}^3$ at bandwidth one is licensed by the $\lfloor k/2 \rfloor \cdot L$ support staircase (D1; consistent with papers XI/XII). The blocker is that Lemma R's hypothesis class carries **no bound on the size of the compression's negative eigenvalues** — the in-window off-line blocks are constrained in count ($n_+ \leq p$) but not in magnitude, and the Φ_3 -deficit escapes through that channel.
- Window engineering cannot fix it** (adversarially settled): the RvM clump term ($\log T$ zeros per unit ordinate at zero kernel decay) forces the row-sum cap $M = A_0 \cdot \log T$ for every window; edge-vanishing windows only improve the negligible far field, cost ≥ 0.006 in $2-q(w)$, and flip Φ_3 to the unmonetizable positive side.
- The surviving target, precisely:** a zero-side lemma $\| (c^{-1}\hat{A})_- \| \leq M_-$ uniform in T . Payoff at the MT window with the capped cubic decoder: $\delta = 0.6796896$ if $M_- = 2$, up to 0.6844924 if $M_- \leq 1$ — no window change, tr^3 the only new prime-side input. Required: depth-uniform control of off-line evaluation vectors (a cosh-weighted Poisson analogue +

overlap-angle bound). This is now the sharpest known formulation of "what would improve 0.6725 without new correlation ranges."

(Cross-check from D2: the true bandwidth-one ceiling is strictly above δ_{MT} — measured integrality gap ≈ 0.0093 — so headroom exists; the M_- route targets exactly that headroom.)

2. New theorems about the method itself

2a. Stability/rigidity of Lemma R (D4) [PROVED]. The exact 7-term deficiency decomposition turns Lemma R into an identity (verified 7e-14; identically zero on the TightMult class), and the rigidity theorem with constant 1: an ϵ -near-equality configuration lies within Frobenius distance $\sqrt{\epsilon}$ of an exact equality configuration (Löwdin frame). Zeta-side ledger: total deficiency over MT windows $\leq (q^*-1) \cdot N(T) = 0.3275 \cdot N(T)$, yielding a new unconditional pair-repulsion-type bound (on-line pairs with taper overlap $\geq \tau$ number $\leq 0.1637 \cdot N(T)/\tau^2$) and forcing hypothetical off-line zeros to look exactly like eigenvalue-2 pair blocks. Bonus reading: $q(\text{window}) = m_2(\text{window-Gram})$ — **Montgomery – Taylor optimization is minimization of the RMT-predicted deficiency**, and under GUE the budget is exhausted by the Schur/von-Neumann term alone.

2b. The distinct-zeros upgrade $N_d \geq 0.8362503 \cdot N + p$ (D7) [PROVED; ~5-line Lean edit]. Combining the $c=2$ inequality with the zero-count identity once (not twice) gives slope +1 in the off-line pair density — strictly stronger than the repo's own $c=3$ decoding (slope $1/2$; pairs mispriced at $c^2=9$ vs forced value $k_3(2)=8$). Companion facts: the **simple-zero frontier is exactly flat** — $\delta(\pi) \equiv 0.6725007$ for all $\pi \in [0, 0.16375]$, hence **assuming RH buys nothing within this method** (the binding adversary is on-line doubles; exact 4p-budget/mass cancellation); at maximal off-line density all zeros are forced simple; the data forces pair-block spectra to $(\approx 2, \approx 0)$.

2c. The edge no-go (D5) [PROVED]. Every admissible bandwidth-one window has $\hat{r}(\pm 1) = 0$ (two-line Fejér – Riesz), so the weakest pointwise edge hypothesis $|F(1)-1| \leq \epsilon$ certifies exactly 0.6725 — even at $\epsilon = 0$. The finite Nyquist conservation law is powered by lattice aliasing that $\mathbb{R}$ lacks; the correct zeta-side edge object is the Cesàro mean $\bar{F}(\infty, T)$, with the exact collision identity $C(T) = N^* \cdot \bar{F}(\infty, T) - N(T)$. New reading: TightMult + PairCeiling/Signed = "the 2026 certificate is edge-blind by construction, and optimally so." Genuine new target: an unconditional *upper* bound on the edge-cell average of F via the Nyquist-blind null window ($r(0) = 0$).

2d. Weight-freezing (D9) [PROVED]. Derivative power sums are weight-homogeneous in the original power sums, so all bandwidth-one holomorphic statistics of the derivative process (ξ' -analog) are frozen on the fiber — "differentiate then pair-correlate at bandwidth one" cannot work, and the Lean ξ' gain ($\rightarrow 0.86864$) is purely arithmetic (the D_1 coefficient density). But the transport is real in *modulus/counting* observables: at $N=6$ the 1-point radial law of critical points resolves $7/10$ fiber dimensions ($\chi = 2.14 \times$ the best direct 4-point statistic), and **AH is not closed under differentiation** (exact $N=3$ counterexample). Finite Hardy identity: on-circle stationary points = roots of $zZ' - (N/2)Z$, all on the circle.

2e. The no-escape lemma for certificates (D2) [PROVED] — validity against the optimal period- M adversary forces $c_0 + \int r \cdot x \leq V_M$ by a Riemann sum with no regularity assumption: the stability penalty buys rate, never limit. Quadratic-class certificates converge to exactly δ_{MT} (excess $0.31/M^2$): **MT is optimal in its class at every period, and the ≈ 0.0093 headroom above it is a pure integrality gap**, accessible only to integrality-aware certificates.

3. The fiber, finally explicit (D6 + D3 + D10)

3a. The center-of-mass modulation family (D6). $q_g = q_{ACUE} \cdot g(X \bmod N)$, $X = \sum x \cdot m_x$, is an exact balanced-moment mimicker iff $E g = 1$ and $\hat{g}(\pm 1) = 0$ — an explicit **($N-3$)-dimensional** family at every N (dims 0,0,2,3,4,5 at $N=3..8$, matching/entering the known fiber). The $N=5$ fiber is *completely solved*: a pentagon over $\mathbb{Q}(\sqrt{5})$; the known mimicker is

$|v|^2$ for a 5-sparse coherent superposition of translated Fermi seas; the destructive cross-spectrum constraint is the single syndrome $a^*Sa = 0$ (vanishing nearest-translate coherence). Mechanism: $X \bmod N$ is exactly uniform under ACUE, and moment couplings to com-frequencies $2..N-2$ vanish by a transport-cost bound (hop budget $2N$ vs cost $\geq 2N+2$).

3b. The minimal detector theorem (D3). The cheapest observable with $O(1)$ fiber leverage is $E|p_a p_b|^2$, $a+b = N+1$ — bandwidth exactly $1 + 1/N$; nothing cheaper exists (earthmover lower bound: fiber seas have displacement $\geq 2N+2$). Dichotomy: bounded-likelihood mimickers (twists) cost bandwidth $2 - 4/N \rightarrow 2$ (twin-prime class); the cheap detectors work only on likelihood-degenerate directions. Prime-side price of the minimal detector: **h-averaged bipartite $E_2 \times E_2$ shifted correlations at length $T^{1+1/N}$** (dispersion-method-shaped, Linnik/Motohashi/BFI class). Open danger sign: leverage dropped $4.94 \rightarrow 0.92$ from $N=5$ to 6 — uniformity in N undetermined.

3c. The arithmetic shadow and the two-scale cost law (D10). The fiber's out-of-band profile: the first sliver past the edge kills $54 - 100\%$ of the fiber with $O(1)$ leverage, but the last direction survives to kill-degree $d^* = 3(N-3)$ ($\alpha^* \rightarrow 3$, conjectured law). Free structural kill: the functional equation's reflection symmetry eliminates the entire chiral half of the fiber ($39/80$ dims at $N=7$) at zero cost — and provably cannot do more. All squared-modulus data saturates below full rank. Marginal-value curve: pinning F on $(1, 1+\epsilon]$ buys $d\delta/d\epsilon \approx 0.68$ at the edge; the AH deviation lives in the *constant term* of the ψ -variance at short shifts — the Siegel-zero price tag (Conrey – Iwaniec), consistent with everything above.

3d. Deficiency blindness dichotomy (D4). Lemma R's total deficiency is *exactly* frozen on the fiber (2-point functional); the fiber-detecting observables among spectral functionals are the interior clips $\Sigma((\mu - c')_+)^2$, c' strictly inside $(1,2)$ — the boundaries are blinded by two new lattice facts (spectral cap $\mu \leq 2$ exactly; particle-hole symmetry $\mu \leftrightarrow 2-\mu$), while the continuum sine-Gram violates the cap (support $[0,3]$ — the Christoffel hard-edge).

4. The updated quantitative map

object	value	status
headline constant δ_{MT}	0.672500703679	proved (repo)
conditional cubic upgrade ($M_- = 2 / \leq 1$)	0.6796896 / 0.6844924	one lemma missing (D1)
true bandwidth-one ceiling	0.6818 ± 0.0015 ; candidate 15/22 exactly	D2 [CONJECTURE]; rigorous bracket [0.6725007, 0.6818287+]
quadratic-class ceiling	δ_{MT} exactly (excess $0.31/M^2$)	D2 [COMPUTED/PROVED]
lattice-union adversary cap	$1 - 3/\pi^2 = 0.69603$	D2 [PROVED] (totient forcing $g = \varphi$)
distinct zeros	$N_d/N \geq 0.8362503 + p/N$	D7 [PROVED, unformalized]
off-line pair budget	$p/N \leq 0.1637496$; at saturation all zeros simple	D7 [PROVED]
deficiency ledger	$\Sigma D \leq 0.3275 \cdot N(T)$; rigidity $\sqrt{\epsilon}$	D4 [PROVED]
doubles density of any Nyquist equality law	$\rightarrow 1/4 - 1/\pi^2 = 0.148679$	D6 [PROVED identities]
fiber's cheapest detector	bandwidth $1 + 1/N$; robust directions $2 - 4/N$	D3 [PROVED at $N=5,6$]
fiber kill-degree	$d^* = 3(N-3)$, $\alpha^* \rightarrow 3$	D10 [CONJECTURE, verified $N=5,6,7$]
proved zeta rate	$1/\log T$ (two removable first-order losses; true tower $1/L^2$ conjectured)	D8 [read from repo + CONJECTURE]

object	value	status
the three N^{-2} towers	one object: the Nyquist cotangent tower (Bernoulli tower of $(\cot, \csc^2)$); exists only at the edge	D8 [PROVED forms + COMPUTED]

5. Corrections and errata from this round

- The naive premises of three directions were *refuted* by their own agents (the healthiest sign of the round): pointwise edge hypotheses buy nothing (D5); derivative pair correlation at bandwidth one is blind (D9); window engineering cannot cap the negative spectrum (D1 follow-up). Each refutation came with the corrected statement.
- D8 flagged a discrepancy in an earlier verification-script closed form for the Galerkin ladder (two variants circulating; both match the QP after the θ_n fix under their respective normalizations — final reconciliation recorded in `dir8_exact_form.py`, whose two-generator form $q_n = \frac{1}{2} + t_n \cdot \cot \kappa_n$ is verified to 60 digits and supersedes both).
- The finite $N=8$ fiber dimension re-derived as 399 vs the earlier 403 at tighter rank tolerance (D4) — worth an exact-arithmetic recount.
- Support-field version of the destructive cross-spectrum is *infeasible* on the hard-core face (D6) — the constraint's true home is the com-syndrome on the original fiber.

6. Ranked program (what to actually do next)

1. **The M_- lemma** (D1): prove $\| (c^{-1}\hat{A})_- \| = O(1)$ via depth-uniform off-line vector control — worth $+0.007..0.012$ on the headline constant with tr^3 at bandwidth one. The single highest-value open item.
2. **Formalize `mult_two_pair`** (D7): $N_d \geq 0.83625 \cdot N + p$ — a five-line Lean edit yielding a strictly stronger published constant.
3. **The transport-cost lemma** (D6): finish the proof of the com-modulation family — an explicit $(N-3)$ -dimensional mimicker polytope at every N , the paper-ready core of the fiber theory.
4. **Uniformity of the minimal detector** (D3): $N = 7, 8$ leverage of $E|p_a p_b|^2$, $a+b = N+1$; decides whether the dispersion-method target (bipartite E_2 correlations at $T^{\wedge}\{1+1/N\}$) is the real gateway.
5. **Edge-cell upper bound** (D5): the Nyquist-blind null-window route to the first unconditional upper edge datum for F .
6. **Max-ent Gibbs family at $N \sim 30$** (D6) + the D8 tightness evidence: the all- N Nyquist equality law now looks provable; its doubles density and $E|p_N|^2$ are already pinned by the frozen identities.
7. **$d^* = 3(N-3)$ kill-degree law** (D10) and the $15/22$ ceiling identity (D2): two exact-arithmetic conjectures, both machine-attackable.

Every claim above traces to a tagged agent report with scripts in the scratchpad; the four refuted premises are retained in § 5 deliberately — the round's value lies as much in the closed doors as the opened ones.

Historical source 25: impostors_paper

Source: `historical/riemann-rmt/impostors_paper.md`. SHA-256:

760f8b362ab12bbf9d182c77f4b555983275a704de07927bc3a14d127fe0b8b4.

Telling the Riemann Zeros from Their Impostors

A key step toward the alternative hypothesis, and a new bridge between zeta zeros and random matrix theory

Bill (Qingyun) Sun · GPT5.6SOL · Fable

August 2026

What this is about

The zeros of the Riemann zeta function are conjectured to be spaced like the eigenvalues of a large random unitary matrix. This is Montgomery's pair correlation conjecture, and half a century of work has not settled it. The obstruction has a name: the **alternative hypothesis**, the scenario in which the normalised gaps between zeros all lie asymptotically on a half-integer lattice rather than fluctuating like a random spectrum. It is consistent with everything currently provable, it has genuine arithmetic consequences, and — this is the hard part — it *imitates* the random-matrix answer for every statistic anyone has managed to test.

Tao made the obstruction concrete by moving it into random matrix theory. The ACUE is a lattice-supported measure on unitary spectra that reproduces the random-matrix two-point correlation exactly. The programme of "find the first statistic ACUE cannot fake" has been running ever since and keeps hitting the same wall: the set of measures matching any finite list of moments is a large convex body, so ruling out ACUE never rules out its neighbours.

This paper does four things.

We measure the wall exactly. The family of impostors — measures matching every balanced moment of degree $\leq N$ — has dimension 0, 0, 2, 10, 80, 403, 1804 for $N = 3, \dots, 9$. We give closed-form families for every N , and show they remain invisible along the *entire* natural heat flow, because that flow is diagonal in the coefficients. Adding one more moment, or evolving the moments you have, is provably futile.

We find an instrument that sees past it. Not an average, but a *stopping time*: the finite de Bruijn – Newman depth Λ , the first moment at which two zeros collide when the heat flow is run backwards. It is not a polynomial statistic of any degree, so the impostors' freedom does not protect it. It separates the random-matrix and lattice scenarios not by a number but by a **universality class**: $-\Lambda \asymp N^{-8/3}$ against N^{-2} . The alternative hypothesis fails not by being fragile but by being **too stable** — it satisfies its own Riemann-hypothesis-analogue far too robustly.

We find the structure underneath, and it is the point of the paper. The static invisibility depth of the impostors and the dynamic relaxation spectrum of the flow turn out to be *the same operator*. Concretely, the quantity $\delta(N-8)$ that governs how deeply an impostor can hide is simultaneously the dimension of a Grassmannian, an affine Bruhat length, and — the

new equality — the eigenvalue of the Jacobian of the zero dynamics at the clock configuration. This explains the whole phenomenon in one sentence: *impostors hide in exactly the modes the flow acts on most strongly*, which is why a hitting time sees what moments cannot.

We obtain an exact law for the refined instrument. Ordinary Λ is blind to eigenvectors, so isospectral impostors need a *marked* depth: perturb by a rank-one mark and differentiate. Its derivative turns out not to be the global quantity one expects from the determinant lemma, but a **rank-two contrast supported on the pair of eigenvalues that is about to collide** — plus one background term which, once identified, removes the last fitted constant from the theory. The resulting formula is parameter-free and matches measurement to eleven digits (§ 3).

We also relate Λ quantitatively to the Lagarias – Rodgers extremum μ , the existing formulation of the same question, and extract from that relation a single explicit number below which the alternative hypothesis is false.

What we do not claim. We have not proved the Riemann hypothesis, refuted the alternative hypothesis, or proved the pair correlation conjecture. §§ 1 – 3 are theorems and computations about random matrices, plus a precise, falsifiable, currently unverified prediction about zeta. Where a statement is numerical rather than proved, it is labelled. A separate part of the paper reports three new unconditional records for prime clustering, which are proved and machine-certified.

1. The alternative hypothesis, and an instrument that sees past it

1.1 Why the alternative hypothesis is hard to kill

Montgomery's pair correlation conjecture says the normalised gaps between zeta zeros follow the GUE law. The **alternative hypothesis** is the scenario in which those gaps all lie asymptotically in $(1/2)\mathbb{Z}$. It is consistent with everything provable about zeros, it has real arithmetic consequences, and it is precisely what current technique cannot exclude.

Tao's article *The alternative hypothesis for unitary matrices* moves the obstruction into random matrix theory, where it can be computed with. The ACUE is the measure on N -point configurations C inside the $2N$ -th roots of unity given by $\mu_{\text{ACUE}}(C) = |\Delta(\zeta^C)|^2 / (2N)^N$. It reproduces the CUE two-point correlation exactly. The natural programme — find the first moment at which they differ — keeps colliding with the same difficulty: the set of measures matching a moment list is a large convex body, so ruling out ACUE never rules out its neighbours. One must rule out a *fibre*, not a point.

How large the fibre is. We computed it exactly. The affine dimension of the set of measures on the ACUE support matching every balanced moment of degree $\leq N$ is

N	3	4	5	6	7	8	9
fibre dimension	0	0	2	10	80	403	1804

with an explicit closed-form subfamily for every N : $q_g(C) = \mu_{\text{ACUE}}(C) \cdot g(X(C))$, $X = \sum c \bmod N$, matching every balanced moment of degree $\leq N$ iff $E[g] = 1$ and $\hat{g}(\pm 1) = 0$ — an $(N-3)$ -parameter family, verified at $N = 5, 6, 7, 8$ with worst moment error 10^{-12} and genuinely positive. A second family comes from the determinant character: μ_{ACUE} is the Born distribution $|\langle C | (\wedge^N F) A_0 \rangle|^2$ of a single Slater determinant, and the tilts by $\det(U)^r$ are the interference patterns of two shifted Fermi seas — points of the first secant variety $\sigma_2(\text{Gr}(N, 2N))$ rather than the Grassmannian itself. At $N = 8$, of the 403 fibre directions, 401 are invisible to every pattern count of window width ≤ 2 and 383 to width ≤ 6 .

And the natural flow does not help. Writing $P(z) = \det(I - zU) = \sum a_j z^j$, the finite analogue of the de Bruijn – Newman heat deformation is $P_r(z) = \sum a_j r^{j(N-j)} z^j$, $t = \log r$ — *diagonal in the coefficients*. Every balanced moment of degree $\leq N$ therefore stays frozen along the **entire** trajectory (verified to $4.5 \cdot 10^{-16}$). Evolving the moments you have is as futile as adding one more.

1.2 The depth

Define the **finite de Bruijn – Newman depth**: $-\Lambda(U)$ is the first time, running the heat flow backwards, at which two zeros collide. It is well defined because the flowed polynomial stays self-inversive, so a simple zero cannot leave the unit circle without first colliding. Equivalently the zeros obey the attracting circular Coulomb dynamics $\dot{\theta}_j = -\sum_{k \neq j} \cot((\theta_j - \theta_k)/2)$.

The right way to see it is geometric. Let R_N be the polynomials with all roots on the circle and $D_N = \{\text{Disc} = 0\}$ its boundary. Then

$-\Lambda(P)$ is the distance from P to the discriminant hypersurface along the canonical heat ray.

Static moments are coordinates *along* the mimicker fibre; the depth is a coordinate *transverse* to it. It is not a polynomial statistic of any degree — it is a first-passage time — which is exactly why the fibre's freedom does not protect it. This is the finite shadow of the picture in which $\Lambda_{\text{dBN}} \leq 0$ is the Riemann hypothesis and Rodgers – Tao's $\Lambda_{\text{dBN}} \geq 0$ makes RH, if true, *barely* true.

1.3 The separation

ACUE, **exactly**. Complete enumeration of all rotation orbits for $N = 3, \dots, 10$ — 13,132 orbits, 184,756 configurations at $N = 10$ — with exact Vandermonde masses, validated to 40 digits. $P(\text{clock}) = 2^{\{1-N\}}$ **exactly** (Cauchy – Binet: $\sum_{|C|=N} | \Delta|^2 = (2N)^N$, each clock contributing N^N); clock polynomials $1 - cz^N$ are flow-invariant, so $\Lambda = -\infty$ there. Every non-clock configuration has minimal gap *exactly* π/N (pigeonhole). The law is $-\Lambda^{\text{ACUE}} \asymp N^{-2}$, fitted exponent -2.0009 , with $N^2(-\Lambda)$ supported in $\approx [1.31, 1.99]$. In all 13,130 non-clock orbits the first collision occurs at a pair already adjacent at $t = 0$ — zero exceptions.

CUE, **Monte Carlo** to $N = 256$. $-\Lambda^{\text{CUE}} \asymp N^{-8/3}$, fitted exponent -2.678 ± 0.016 , and the law is parameter-free: $8N^{8/3}(-\Lambda) \Rightarrow G^2$ where $P(G > x) = \exp(-x^3/72\pi)$ is the sine-kernel smallest-gap law. The constant 72π was derived from the kernel, not fitted (measured 229 – 236 against 226.2, KS 0.035 – 0.041, median 29.6 – 30.5 against $(72\pi \ln 2)^{2/3} = 29.08$).

1.4 What the exponents mean

For an isolated close pair at gap δ the two-body reduction gives collision time $\delta^2/8 + o(\delta^2)$ — confirmed to 10^{-7} , and exact at $N = 2$ where $-\Lambda = -\log \cos(\delta/2)$. So the depth is governed by the smallest gap, and the smallest gap by the level repulsion exponent:

$p(s) \sim c s^\beta \Rightarrow \text{smallest of } \approx N \text{ gaps} \asymp N^{-1-1/(\beta+1)} \Rightarrow -\Lambda \asymp N^{-2-2/(\beta+1)}$.

ensemble	β	predicted	measured
COE	1	−3	−3.064 (local slopes −3.03 ... −3.10)
CUE	2	$-8/3 = -2.667$	-2.678 ± 0.016 ($N \leq 256$)
CSE	4	$-12/5 = -2.4$	−2.510 ($N \leq 64$, still drifting)
ACUE lattice	∞	−2	−2.0009 (exact enumeration)

All measured slopes are slightly steeper than predicted, in the direction and of the size of the finite- N drift independently calibrated at $\beta = 2$, where the larger range shows the local slope converging to the prediction. The circular β -ensembles were sampled by the Killip – Nenciu CMV construction, validated against Haar CUE (KS 0.005) and direct COE = VV^T (KS 0.004); the localisation assumption holds in 95 – 99% of samples and in 100% of ACUE configurations.

So the depth is a scalar fingerprint of the microscopic repulsion exponent. And the interpretation inverts the naive expectation. One would guess a fake universe is caught by being *fragile* somewhere. The opposite: ACUE's defect is that it is **too stable**. CUE admits rare pairs at distance $N^{-4/3}$, a full factor $N^{-1/3}$ below the mean spacing, so it is always within $N^{-8/3}$ of losing real-rootedness; the lattice quantises every gap at π/N and forbids the accidents. A true random-matrix world is real-rooted but *microscopically fragile*; the alternative hypothesis satisfies its own Riemann-hypothesis-analogue far too robustly. That is Newman's dictum — "RH, if true, is only just true" — as a statement about a scaling exponent.

1.5 The depth is the Lagarias–Rodgers hard core in another coordinate

This is the sharpest consequence, and it says the depth is not a competitor to the existing formulation of the alternative hypothesis but that formulation in computable coordinates.

Lagarias and Rodgers study processes mimicking the sine process at coordinatewise bandwidth one and define $\mu = \sup\{c : \text{some mimicker has minimum spacing} \geq c\}$, mean spacing 1. They prove $\mu \geq 1/2$ by the randomly shifted half-lattice — which is exactly ACUE — record the pair-only bound $\mu \leq 0.606894\dots$, and conjecture $\mu = 1/2$. The alternative hypothesis is precisely the assertion that zeta's zeros realise a hard core of $1/2$.

On the circle a hard core of c mean-spacings is $\delta_{\min} = 2\pi c/N$, so with $-\Lambda = \rho \cdot \delta_{\min}^2/8$,

$N^2(-\Lambda) = \rho \cdot \pi^2 c^2/2$, and at the alternative-hypothesis value $c = 1/2$, $N^2(-\Lambda) = \rho \cdot \pi^2/8$.

Every number in § 3.3 is this identity read one way or the other. The ACUE quantiles give

N	min $N^2(-\Lambda)$	$\rho_{\min}$	median	ρ_{med}	max	$\rho_{\max}$
6	1.353146	1.0968	1.419374	1.15050	1.952629	1.5827
8	1.330383	1.0784	1.418216	1.14956	1.976122	1.6018
10	1.314614	1.0656	1.412774	1.14515	1.985458	1.6094

An exactly computed configuration constant. Take the alternating clock (the zeros of $z^N + 1$), delete $e^{-i\pi/N}$ and insert 1, so the gap pattern becomes 1, 2, ..., 2, 3 in half-lattice units and

$$P_N(z) = (z - 1)(z^N + 1)/(z - e^{-i\pi/N}).$$

Two independent routes agree. The lattice solver on this configuration gives $N^2(-\Lambda)$ rising monotonically to 1.41963827 at $N = 20$; the continuum limit $t = -s/N^2$, $z = e^{iu/N}$ yields

$$G_s(u) = 2\cos(u/2) - 2\pi \int_0^{\infty} \frac{1}{2} e^{s(1/4-y^2)} \cos((\pi+u)y) dy,$$

whose first double zero ($G = \partial u G = 0$) sits at $s^* = 1.419640342\cdots$, $u^* = 1.812942145\cdots$ with $G''(u^*) = -0.3767 \neq 0$. The two agree to $2 \cdot 10^{-6}$ at $N = 20$, consistent with an $O(N^{-2})$ approach. And $s^* = \rho_\infty \cdot \pi^2/8$ with $\rho_\infty = 1.150717118\cdots$ — the constant is the hard core, dressed by the background. (One correction to the record: s^* is not the limit of the ensemble median. For $N \leq 7$ the two coincide, but the complete enumeration shows the median turns around at $N = 7$ and falls — 1.41822, 1.41520, 1.41277 at $N = 8, 9, 10$ — as the support spreads from $[1.4147, 1.8246]$ to $[1.3146, 1.9855]$. s^* belongs to the support of the limit law, via the single-dislocation stratum.)

Equivalence of the two extremal problems. With $\mu_\Lambda = \sup\{\liminf N^2(-\Lambda)\}$ over mimickers,

$$\mu_\Lambda \geq \pi^2\mu^2/2 \text{ and } \mu \leq \sqrt{(2\mu_\Lambda)/\pi}.$$

These are tight, not loose: LR's published $\mu \leq 0.606894\cdots$ is *exactly* the depth bound $\mu_\Lambda \leq 1.8177\cdots$. So **any improvement of the depth bound improves the Lagarias – Rodgers bound**, at an explicit exchange rate; a depth bound of 1.29 — barely below what ACUE configurations themselves realise — would already give $\mu \leq 0.511$.

A falsifiable number for the alternative hypothesis. Running the inequality the cheap way:

If the zeros of ζ have, in a window at height T containing $N = (\log T)/2\pi$ zeros, local Newman depth with $\liminf N^2(-\Lambda) < \pi^2/8 = 1.2337\cdots$, the alternative hypothesis is false.

This uses only the hard-core consequence of AH. Compare: AH predicts ≥ 1.2337 with actual typical value 1.41964; CUE predicts $N^2(-\Lambda) \asymp N^{-2/3} \rightarrow 0$. The margin to the threshold is a factor 1.15; the margin from CUE is unbounded. And crucially, refuting AH does **not** require the full $N^{-8/3}$ law — only that the depth eventually drops below one explicit constant.

Why this reformulation may be more tractable. The LR extremum constrains a *minimum*, a hard combinatorial quantity, whereas Palm-type certificates — the row-sum square $\text{Var}_0 \text{--sine}(S_f) \leq (M_h - A_f)(A_f - m_h)$ for band-limited f , and its multiwindow quadratic generalisation on the packing body K_h — apply naturally to smooth functionals. The depth is smooth in the configuration (we compute its gradient below), has a variational characterisation as a distance to a hypersurface, and has an explicit derivative law. Transporting those certificates from the row sum to the depth is the natural next attempt, and unlike the row-sum searches — which failed for an exact local reason, the pattern $\{-h, 0, h\}$ already forcing $S_0 = 2f(h) > 9/7$ for the first nonnegative Fejér profile — the depth has no such immediate obstruction.

Two lemmas, not one. It is worth separating them, because they have different status and do different jobs.

Hard-core lemma ($\rho \geq 1$). Proved in § 1.7 directly from the adjacent-gap ODE. It gives the inequality in the direction used here, hence the falsification threshold: $N^2(-\Lambda) < \pi^2/8 \Rightarrow$ **the alternative hypothesis is false**.

Localisation lemma ($\rho \rightarrow 1$ for finite β). The remaining probabilistic estimate. It is what upgrades the exponents to distributional limits.

Only the second is open, and § 1.7 predicts its rate.

The leverage on μ , made explicit. Since $\mu_{\Lambda} \geq \pi^2\mu^2/2$, any upper bound $D_{\text{dyn}} \leq M$ for the depth of an admissible class gives

$$\mu \leq \sqrt{(2M)/\pi}.$$

$M = 1.29$ — barely below what ACUE configurations themselves realise — already gives $\mu \leq 0.511281\dots$, a large improvement on the published $0.606894\dots$. Crossing the alternative-hypothesis barrier itself requires $M < \pi^2/8 = 1.2337005501\dots$. So the depth is not merely an equivalent coordinate for the Lagarias – Rodgers problem; it is a coordinate in which a *quantitative* target is visible, and the target is a single number.

1.6 A classification of impostors, with two computable criteria

Criterion I (transversality). For a configuration X and bandwidth r , with $M_r = (p_1, \dots, p_r)$ and $\tau = -\Lambda$, ordinary Λ catches the class exactly when $\ker DM_r \not\subset \ker D\tau$. The relative residual of $\text{grad } \tau$ off the row space of DM_r :

configuration	$r = 1$	$r = 2$	$r = 3$	$r \geq N/2$
single dislocation, $N = 8$	0.991	0.941	0.737	rank saturates
random ACUE lattice, $N = 7$	0.996	0.963	0.569	rank saturates
generic non-lattice, $N = 7$	0.998	0.926	0.195	rank saturates

While $\text{rank } DM_r < N$ the kernel is nontrivial and the answer is decisively yes: 80 – 99% of the depth gradient lies in directions the first one or two moments cannot see. At $r \approx N/2$ the rank saturates at N , $\ker DM_r = 0$, and the residual drops to 10^{-16} — not because Λ fails but because the question becomes vacuous.

Class I, caught by Λ directly. Anything whose defect changes the geometry of the closest pair: collision-stratum families ($\delta = 0$ gives $\Lambda = 0$ outright, the mechanism by which function-field Newman constants are pinned by double roots), half-lattice adversaries with a hard lower spacing, and the centre-of-mass and secant families, which move the depth law by total variation 0.12 – 0.24 with every balanced moment frozen. Linear-programming tomography over the exact fibre at $N = 6$ gives $E[N^2(-\Lambda) \mid \text{non-clock}]$ ranging over $[1.3610, 1.4770]$ against ACUE's 1.4336, while the mass at $\Lambda = -\infty$ can be pushed from 0 to 0.0975, tripling ACUE's 0.03125, with every constrained moment held fixed.

Class II, invisible to Λ , caught by a marked depth. Λ is a function of the characteristic polynomial alone, so isospectral matrices have *identically* the same depth — verified to machine zero ($\tau(G_1) = \tau(G_2) = 0.068725421516$, difference $8 \cdot 10^{-17}$). The repair deforms $G \mapsto G + \eta uu^*$, transports to the circle by Cayley, and differentiates: $\chi(G;u) = \partial_{\eta}(-\Lambda)|_0$. It separates the isospectral pair immediately, median $|\chi(G_1;u) - \chi(G_2;u)| = 0.081$ against $|\chi|$ of order 0.01 – 0.2.

What χ actually measures is **not** the directional resolvent the determinant lemma suggests, but a rank-two contrast on the colliding pair, and the exact law — parameter-free, verified to eleven digits — is the subject of § 3.

Class III, immune to the linear flow. If the observable algebra misses a direction v , the heat operator — being diagonal — keeps $H_t v$ in the same hidden sector, so static invisibility persists for all t . **A linear flow does not destroy linear invisibility.** Λ escapes that argument only because it is not a linear observable transported by the flow but the first hitting time of the orbit against a variety: two orbits can share an invariant hidden sector and stand at different distances from D .

A test case at the boundary. The strongest known static escape is the pair of parity sectors $q_{N^{\pm}}(C) = \mu_N(C)(1 \pm (-1)^{\sum x})$, which are mutually singular, agree on every complete marginal on at most $N - 1$ sites, and match balanced trace moments of degree far beyond N . The depth *does* separate them at every finite N , but entirely through the atom: for

even N both clocks have even slot-sum, so $P(\Lambda = -\infty) = 2^{\{2-N\}}$ in q^+ and **exactly 0** in q^- . The non-clock bulk is nearly identical — at $N = 8$ the quantiles are 1.3739/1.3959/1.4182/1.4554/1.4839 against 1.3752/1.3959/1.4196/1.4338/1.4788 — and the conditional means differ by $8.9 \cdot 10^{-3}$, $3.4 \cdot 10^{-3}$, $1.1 \cdot 10^{-3}$ at $N = 6, 8, 10$, shrinking. Since the separating mass $2^{\{2-N\}}$ vanishes, the parity sectors are a Class II/III object: caught at finite N , invisible in the limit, and a natural first target for the marked depth.

1.7 What is now proved: a two-body comparison theorem

The three scaling laws above were, until this point, numerical. They are not independent: all three follow from a local collision law plus a smallest-gap law, and the smallest-gap laws are known. What was missing was control of the *background* — the effect of the other $N - 2$ zeros on the collision time of the closest pair. That is now a theorem, and it is elementary.

Parameterise the flow by the depth $u = -t$, so the zeros obey $\theta_{-j}' = -\sum_{\{k \neq j\}} \cot((\theta_{-j} - \theta_{-k})/2)$ and collisions occur as u increases. Two facts:

Exact two-body solution. For $N = 2$, separating variables in $\delta' = -2\cot(\delta/2)$ gives $\cos(\delta(u)/2) = e^u \cos(\delta(0)/2)$, hence collision exactly at $u = -\log \cos(\delta(0)/2) = \delta(0)^2/8 \cdot (1 + \delta(0)^2/24 + \dots)$.

Theorem A (two-body comparison). Let (a,b) be an *adjacent* pair, i.e. no zero lies strictly inside the short arc from θ_{-b} to θ_{-a} , and let g be its gap. Then under the depth flow $g' \geq -2\cot(g/2)$. Consequently no collision occurs before $u = -\log \cos(\delta_{\min}/2)$, so $-\Lambda \geq -\log \cos(\delta_{\min}/2) \geq \delta_{\min}^2/8$, i.e. $\rho \geq 1$ always.

Proof. Write $x_{-j}^k = (\theta_{-j} - \theta_{-k}) \bmod 2\pi \in (0, 2\pi)$. Then

$$g' = \theta_{-a}' - \theta_{-b}' = -2\cot(g/2) - \sum_{\{k \neq a,b\}} [\cot(x_{-a}^k/2) - \cot(x_{-b}^k/2)].$$

Adjacency says that for every other zero k one has $x_{-a}^k = x_{-b}^k + g$ with both in $(0, 2\pi)$: going counterclockwise from k one meets b before a , and k is not inside the arc. Since $\cot(x/2)$ is strictly decreasing on $(0, 2\pi)$, each bracket is negative, so each term enters g' with a positive sign — **every other zero slows the collapse**. Hence $g' \geq -2\cot(g/2)$. The ordering of the zeros is preserved until the first collision, so each gap dominates the two-body solution started from its own initial value, and the earliest possible vanishing is at $u = -\log \cos(\delta_{\min}/2)$. ■

The sign claim was checked on 11,060 background terms across random and lattice configurations with zero exceptions, and the resulting inequality holds on every dataset in this paper: the minimum of $(-\Lambda)/(-\log \cos(\delta_{\min}/2))$ is 1.0842, 1.0714, 1.0612 over the complete ACUE enumerations at $N = 6, 8, 10$, and 1.00019, 1.00021 over the CUE samples at $N = 16, 64$. (At $N = 256$ the measured minimum is $1 - 1.3 \cdot 10^{-5}$, an absolute discrepancy of $2 \cdot 10^{-11}$, at the solver's tolerance.)

Corollaries, all now rigorous.

- $-\Lambda^{\text{ACUE}} \geq \pi^2/(8N^2)$ exactly, since $\delta_{\min} = \pi/N$ for every non-clock configuration (§ 1.3).
- Any point process with hard core c has $N^2(-\Lambda) \geq \pi^2 c^2/2$. The exchange inequality $\mu_{-\Lambda} \geq \pi^2 \mu^2/2$ of § 1.5 — hence the Lagarias – Rodgers direction we actually use, and the falsification threshold $\pi^2/8 = 1.2337$ — is a theorem, not numerics. What remains conjectural there is only the *reverse* comparison.

Theorem B (matching upper bound where the background is subcritical). The background term is bounded by the mean value theorem: $|\sum_{\{k \neq a,b\}} [\dots]| \leq g \cdot \sum_{\{k\}} \frac{1}{2} \csc^2(x_{-b}^k/2)$, and for a configuration all of whose gaps exceed c/N this is $\leq C \cdot N^2 g$, using $\sum_{\{k=1\}}^N \frac{1}{2} \csc^2(\pi k/N) = (N^2-1)/6$. Integrating $\delta' \leq -2\cot(\delta/2) + CN^2\delta$ gives

$-\Lambda \leq (\delta_{\min}^2/8) \cdot (1 + O(N^2\delta_{\min}^2))$, hence $-\Lambda = (\delta_{\min}^2/8)(1 + o(1))$ whenever $N^2\delta_{\min}^2 \rightarrow 0$.

What this reduces the three laws to.

Law 1 (CUE). $\delta_{\min} \asymp N^{-4/3}$ gives $N^2\delta_{\min}^2 \asymp N^{-2/3} \rightarrow 0$, so Theorems A and B pin $-\Lambda = \delta_{\min}^2/8 \cdot (1 + O(N^{-2/3}))$. The predicted relative correction exponent $-2/3 = -0.667$ matches the independently measured $0.598 \cdot N^{-0.729}$. Composing with the sine-kernel smallest-gap law $P(N^{4/3}\delta_{\min} > x) \rightarrow \exp(-x^3/72\pi)$ — Ben Arous – Bourgade, and Feng – Wei for C β E — yields $8N^{8/3}(-\Lambda) \Rightarrow G^2$. So Law 1 is reduced to a cited theorem plus Theorems A and B; what is still needed for a complete proof is that the error in Theorem B is uniform in probability, i.e. that the minimal gap is with high probability isolated and unaccompanied by a second anomalously small gap — which is precisely what the extreme-gap literature supplies.

Law 3 (β -universality). Identical, for every finite β : $\delta_{\min} \asymp N^{-1-1/(\beta+1)}$ gives $N^2\delta_{\min}^2 \asymp N^{-2/(\beta+1)} \rightarrow 0$, so $-\Lambda \asymp N^{-2-2/(\beta+1)}$ reduces to the C β E extreme-gap law. The exponent is not an empirical fit; it is the image of the repulsion exponent under two elementary steps.

Law 2 (ACUE). Here $\delta_{\min} = \pi/N$ exactly, so $N^2\delta_{\min}^2 = \pi^2$ — the background correction does not vanish, which is precisely why $\rho^{\wedge} \text{ACUE} \in [1.05, 1.61]$ rather than $\rightarrow 1$. The lower bound $-\Lambda \geq \pi^2/(8N^2)$ is exact and proved; the matching upper bound needs $\rho \leq C$, which our complete enumeration establishes for $N \leq 10$ and to which the slowly decreasing minimising family of § 1.5 gives strong evidence, but which is not proved in general. This is the one remaining gap, and it is a lattice-specific constant problem rather than a universality problem.

The finite- β / $\beta = \infty$ dichotomy, and why it is singular. The background correction is governed by $\delta_{\min}^2 \cdot S$ with $S \asymp N^2$ the background stiffness (the clock value is exactly $(N^2-1)/6$). Feng – Wei give $\delta_{\min} \asymp N^{-1-1/(\beta+1)}$ for the circular β -ensembles, so

$$\delta_{\min}^2 \cdot S \asymp N^{-2-2/(\beta+1)} \cdot N^2 = N^{-2/(\beta+1)} \rightarrow 0 \text{ for every finite } \beta,$$

whence $\rho_{\beta} \rightarrow 1$ and the two-body constant $1/8$ is asymptotically exact. But at $\beta = \infty$ the hard core pins $\delta_{\min} \asymp 1/N$, so $\delta_{\min}^2 \cdot S \asymp 1$: the background contributes at leading order and $\rho_{\infty} \neq 1$. The lattice endpoint is therefore *not* the smooth limit of the finite- β formula — it is a singular hard-core endpoint, and that is exactly why the single-dislocation configuration realises

$$s^* = 1.419640342\dots = 1.150717118\dots \times \pi^2/8,$$

a two-body collision time dressed by roughly 15% of many-body shielding, rather than $\pi^2/8$ itself.

This also predicts the *rate* at which ρ_{β} approaches 1, which is a sharper and independently testable statement than the leading exponent:

$$\rho_{\beta} - 1 = O_{\mathbb{P}}(N^{-2/(\beta+1)+o(1)}).$$

β	predicted rate	fitted over $N \leq 128\text{--}256$	local slope at the largest N
1 (COE)	-1	-1.012	-0.851
2 (CUE)	$-2/3 = -0.667$	-0.710	-0.624
4 (CSE)	$-2/5 = -0.400$	-0.501	-0.407

$\beta = 1$ lands on the prediction; for $\beta = 2$ and 4 the global fits are steeper but the local slopes decrease monotonically toward the predicted value and essentially reach it at the largest sizes, with the finite- N drift growing with β as expected (a slower predicted decay is more heavily contaminated).

The pipeline. The three laws are therefore not three empirical facts but one mechanism:

extreme-gap law (cited) + first-collision identity (§ 1.7) + background localisation (open) $\Rightarrow -\Lambda = \rho \cdot \delta_{\min}^2/8$ with $\rho \geq 1$ proved, $\rho \rightarrow 1$ for finite β , $\rho = O(1)$ but $\neq 1$ at $\beta = \infty \Rightarrow -\Lambda \asymp N^{-2-2/(\beta+1)}$ across the whole family, lattice included.

So the honest status of the three laws changes from "numerical" to: **Law 1 and Law 3 reduced to published extreme-gap theorems plus the localisation lemma; Law 2 proved in the direction that matters for the alternative hypothesis, with its non-unit constant explained rather than fitted.**

1.8 The separation, proved on one side and reduced on the other

The two laws can now be separated by a *deterministic threshold* rather than by comparing fitted exponents.

Theorem C (CUE/ACUE separation). (i) For every N and every non-clock ACUE configuration, $N^2(-\Lambda) \geq \pi^2/8 = 1.2337\dots$ (ii) $P(N^2(-\Lambda^{\text{CUE}}) < \pi^2/8) \rightarrow 1$. Hence the two depth laws are eventually separated by the fixed constant $\pi^2/8$, almost surely on the ACUE side.

Part (i) is unconditional and needs nothing new: $\delta_{\min} = \pi/N$ exactly for every non-clock ACUE configuration (§ 1.3, pigeonhole), and Theorem A gives $-\Lambda \geq \delta_{\min}^2/8$. There is no probability in it — the bound holds configuration by configuration.

Part (ii) rests on two events, and it is worth separating what is cited from what we verified.

The gap event. The sine-kernel law $P(N^{4/3}\delta_{\min} > x) \rightarrow \exp(-x^3/72\pi)$ gives $P(\delta_{\min} > \pi/N) = P(N^{4/3}\delta_{\min} > \pi N^{1/3}) \rightarrow \exp(-\pi^2 N/72)$: the CUE minimum gap exceeds the lattice gap with **superpolynomially small** probability. This is the cited input.

The regularity event. Theorem B needs the background slowdown coefficient $B := \sum_{k \neq a,b} [\cot(x_b^k/2) - \cot(x_a^k/2)]/g$ to satisfy $B \leq AN^2$. The clock value is exactly $(N^2-1)/6$, so $A \approx 1/6 = 0.1667$ is the natural scale. Measured on fresh Haar samples:

N	8	16	32	64	128
median B/N^2	0.1092	0.1202	0.1171	0.1202	0.1200
99th percentile	0.2192	0.2551	0.3572	0.2829	0.3273
maximum observed	0.2974	0.3106	0.6075	0.7948	0.3595

so $B \leq N^2$ comfortably, with a stable median of 0.120 and no observed excursion past 0.80. On that event Theorem B gives $N^2(-\Lambda) \leq (N^2\delta_{\min}^2/8)(1 + O(N^2\delta_{\min}^2))$, and the same samples give median $N^2\delta_{\min}^2 = 7.41, 4.66, 2.92, 1.85, 1.18$ at $N = 8, \dots, 128$ — a ratio of 6.28 across a factor 16 in N , against the predicted $N^{-2/3}$ scaling factor $16^{2/3} = 6.35$, agreement to 1%.

The conclusion, measured directly. The fraction of CUE samples with $N^2(-\Lambda)$ below the ACUE floor $\pi^2/8$, and the largest value observed:

N	8	16	32	64	128	256
fraction below $\pi^2/8$	0.565	0.823	0.967	0.9984	1.0000	1.0000
max $N^2(-\Lambda)$ observed	7.93	4.92	2.48	1.34	0.911	0.437
median $N^2(-\Lambda)$	1.083	0.627	0.385	0.231	0.146	0.0946

From $N = 128$ onward every sampled CUE configuration lies below the deterministic ACUE floor, and by $N = 256$ the median sits a factor 13 below it. The separation is not asymptotic bookkeeping: it is visible as a strict inequality between two computable quantities at accessible sizes.

What is proved and what is not. (i) is a theorem. In (ii), the gap event is a cited theorem and the regularity event is verified numerically rather than proved; converting the table above into a proof requires a high-probability bound on B , which is a standard rigidity statement for CUE (control of $\sum_k d_k^{-2}$ for the neighbourhood of the extremal pair) but one we have not carried out. That is the precise remaining gap, and it is narrower than the one we started with: the separation no longer depends on the exponent $-8/3$ at all, only on the CUE minimum gap being eventually smaller than π/N together with a bounded background.

2. One operator: why a hitting time sees what moments cannot

This section is the structural core. Everything above says *that* the depth escapes the moment algebra; this says *why*, and the answer identifies two objects that had no known relation.

2.1 The static side: how deep an impostor can hide

Restrict to the gauge-invariant even shifts and let $|\Omega_s\rangle$ be the Slater state obtained from the Fermi sea by the shift $2s$, $s \in \mathbb{Z}/N\mathbb{Z}$. Then $\langle C | \Omega_s \rangle = z(C)^s \langle C | \Omega_0 \rangle$ where $z(C) = D_2(C)$ satisfies $z^N = 1$, so the entire secant family collapses to a single polynomial on $\mathbb{Z}_N$: for normalised coefficients a , the state $\Psi_a = \sum a_s |\Omega_s\rangle$ induces

$$q_a(C) = \mu_{\text{ACUE}}(C) \cdot |\sum_s a_s z(C)^s|^2 = \mu_{\text{ACUE}}(C) \cdot (1 + \sum_{\delta \neq 0} A_\delta z(C)^\delta),$$

with $A_\delta = \sum_s a_{s+\delta} \bar{a}_s$ the periodic autocorrelation. **The code object is not the support of a but its autocorrelation spectrum**, and the invisibility depth is

$$d_{\text{vis}}(a) = \min \{ \delta(N-\delta) : \delta \neq 0, A_\delta \neq 0 \}.$$

This is exact, not an analogy: the sector pairing $E_\mu[D_{2\delta} \mid \bar{g}]$ vanishes for degrees below $\delta(N-\delta)$ and first opens a rank-one channel at that degree, so q_a agrees with ACUE on all balanced observables of degree $\leq d$ iff $A_\delta = 0$ for every δ with $\delta(N-\delta) \leq d$. Since $d_N(\delta) = \delta(N-\delta)$ depends only on the cyclic Lee distance and is increasing, the pure max – min code problem is settled by pigeonhole: for L states the best possible minimum distance is $m(N-m)$ with $m = \lfloor N/L \rfloor$, maximised at $L = 2$ with value $\lfloor N^2/4 \rfloor$. Two Fermi seas are globally optimal; there is nothing to gain from five, and no compute should be spent looking.

The quantity $d_N(\delta) = \delta(N-\delta)$ is simultaneously the complex dimension of the Grassmannian $\text{Gr}(\delta, N)$, the pairing $\langle 2\rho, \omega_\delta \rangle$ — hence the affine Bruhat length $\ell(t_{\{\omega_\delta\}})$ of the dominant translation — and, up to the standard normalisation, the quadratic Casimir of the fundamental coweight. Those three identifications are representation theory. The fourth is new.

2.2 The dynamic side: the clock's stability operator

Its finite Fourier transform is elementary and exact: for $k \neq 0$,

$$\sum_{\delta=0}^{N-1} \delta(N-\delta) e^{-2\pi i k \delta / N} = -N / (2 \sin^2(\pi k / N)),$$

so d_N is the symbol of the long-range operator

$$(\mathcal{L}_N f)(x) = \sum_{k=1}^{N-1} [f(x) - f(x+k)] / (2 \sin^2(\pi k / N)), \quad \mathcal{L}_N e_{-\delta} = \delta(N-\delta) \cdot e_{-\delta}.$$

Now recall the dynamics. Backwards heat flow moves the zeros by the attracting Coulomb law $\dot{\theta}_j = -\sum_{k \neq j} \cot((\theta_j - \theta_k)/2)$, whose fixed point is the clock. Linearise about it, $\theta_j = 2\pi j / N + \varepsilon_j$, using $d/du[-\cot(u/2)] = \frac{1}{2} \csc^2(u/2)$:

$$\dot{\varepsilon}_j = \sum_{k \neq j} (\varepsilon_j - \varepsilon_k) / (2 \sin^2(\pi(j-k)/N)) = (\mathcal{L}_N \varepsilon)_j.$$

The operator controlling how deeply a static impostor can hide is the Jacobian of the zero dynamics at the clock.

Verified to machine precision at $N = 4, 6, 8, 10, 12, 16, 20, 24$: the Fourier identity to 10^{-12} , the eigenvalues of $\mathcal{L}_N$ against $\delta(N-\delta)$ to 10^{-13} , and $\|\mathcal{L}_N - \text{Jacobian}\| \leq 2.4 \cdot 10^{-13}$. The spectrum is 0, and $\delta(N-\delta)$ with multiplicity two for $0 < \delta < N/2$, and the Nyquist mode $\delta = N/2$ alone at the top with eigenvalue $N^2/4 = \lfloor N^2/4 \rfloor$ — **exactly the maximal invisibility depth obtained above by pigeonhole**, two entirely different derivations of the same number.

2.3 The mechanism, in one sentence

The eigenvalue $\delta(N-\delta)$ has two readings that now coincide: it is the degree at which a δ -mode becomes visible to balanced observables, and it is the rate at which that mode relaxes under the flow. Therefore

an impostor invisible to degree $\leq d$ observables is one whose deviation is supported on the modes that relax fastest — and those are exactly the modes the flow acts on most strongly.

Hiding statically and being acted on dynamically are the same condition, read from opposite ends. That is why adding moments is futile while a hitting time is not: the moment algebra is blind precisely where the flow is loudest, and Λ is a functional of the flow. It also explains why the Nyquist channel keeps appearing on both sides of this project — as the top of the invisibility hierarchy and as the fastest-relaxing mode — and why ACUE, whose deviation from the clock lives at the *bottom* of this spectrum, sits so far from the discriminant.

2.4 What this opens

Three directions follow immediately, and each is computational rather than philosophical.

A quantum code. For $S \subset \mathbb{Z}_N$ the span $\mathcal{C}_S = \text{span}\{\Omega_s : s \in S\}$ satisfies the Knill – Laflamme condition against the algebra of balanced observables of degree $< \min_{s \neq t} d_N(s-t)$: diagonal matrix elements are equal because $D_{\{2s\}}$ is a phase, and off-diagonal ones are the sector pairings above. The ACUE cyclic shifts form an exact error-detecting code whose distance is the affine Bruhat length.

An interference problem that is genuinely open. Since $\hat{A}(k) = |\hat{a}(k)|^2 \geq 0$, the design space is $\{A : A_0 = 1, A \text{ positive-definite on } \mathbb{Z}_N, A_\delta = 0 \text{ on a low-Bruhat-energy region, } A_{\{\delta^*\}} \neq 0 \text{ as high as possible}\}$ — a Delsarte-type linear program, and the natural name is a *zero-Bruhat-correlation-zone sequence*. Note the degenerate case: $A_\delta = 0$ for all $\delta \neq$

0 gives $q_a = \mu_{\text{ACUE}}$, no impostor at all.

A different semigroup. As $N \rightarrow \infty$ at fixed δ , $d_N(\delta)/N \rightarrow |\delta|$, so $N^{-1}\mathcal{L}_N \rightarrow (-\Delta)^{1/2}$, the half-Laplacian on the circle. The natural semigroup attached to the invisibility hierarchy is therefore the **Poisson** flow $e^{-t|k|}$, not the ordinary heat flow e^{-tk^2} . The $1/\sin^2$ kernel is also the defining interaction of the Haldane – Shastry chain, so the question of whether this hierarchy is a sector of an integrable spectrum is a concrete one, worth asking because an affirmative answer would replace brute-force moment computation by Bethe-type identities. We flag this as a lead, not a claim: sharing a kernel is not sharing a Hamiltonian.

A larger deformation class is available and largely unexplored. Any phase twist $u \in U(1)^{\{2N\}}$ acting on the Fermi sea leaves the Born distribution $|\langle C | \Omega_u \rangle|$ unchanged, and the Cauchy – Binet identity gives $E_\mu[(\prod_{c \in C} u_c) s_\lambda(x_C) \overline{s_\nu(x_C)}] = \det((T_u)\{A_\nu, A_\lambda\})$ with $T_u = F^* D_u F$ circulant. So every phase-twisted moment is a circulant Toeplitz minor, hence automatically satisfies the Plücker relations and sits inside the KP/Toda τ -function formalism. The sharp question — whether vanishing of all such minors below degree $\lfloor N^2/4 \rfloor$ forces u to be a linear phase up to gauge — would say that the determinant-character impostor is the canonical extremiser of its whole deformation class, which is a strictly stronger statement than exhibiting one counterexample.

3. The marked depth: an exact rank–two law

Λ is a function of the characteristic polynomial alone, so isospectral matrices share it exactly — we verify $\tau(G_1) = \tau(G_2) = 0.068725421516$ to $8 \cdot 10^{-17}$. Impostors built to match rank, trace and Hilbert – Schmidt norm while differing in directional data therefore need a refinement. Deform by a rank-one **mark**, transport to the circle by Cayley, and differentiate:

$$G_\eta = G + \eta uu^*, U_\eta = (G_\eta - iI)(G_\eta + iI)^{-1}, \chi(G;u) = \partial_\eta(-\Lambda(U_\eta))/\{\eta=0\}.$$

It separates the isospectral pair at once — median $|\chi(G_1;u) - \chi(G_2;u)| = 0.081$ against $|\chi|$ of order 0.01 – 0.2. The interesting question is what it measures.

3.1 Not the resolvent

The determinant lemma $\det(zI - G - \eta uu^*) = \det(zI - G)(1 - \eta \cdot u^*(zI - G)^{-1}u)$ suggests that χ should track the directional resolvent and blow up wherever the inverse-Gram data is worst. **It does the opposite.** Rotating the mark onto the smallest- $|\text{eigenvalue}|$ direction drives $u^*(z_0 I - G)^{-1}u$ from 1.615 to $6.667 = 1/|\lambda_{\min}|$ while χ falls from 0.0267 to 0.0028 — an order of magnitude the wrong way. Condition-number susceptibility and collision susceptibility are different questions: the resolvent asks how much mass the mark puts on a spectral singularity, the depth asks how much *differential angular velocity* the mark induces on the pair that collides first. A near-null eigenvector orthogonal to that pair scores high on the first and nothing on the second.

3.2 The law

Let (a,b) be the first colliding pair with gap δ , and write

$$q_j(u) = |\langle u, v_j \rangle|^2, c_j = -2/(1 + \lambda_j^2), K_{\{ab\}} = c_a v_a v_a^* - c_b v_b v_b^*.$$

Hellmann – Feynman gives $d\lambda_j/d\eta = q_j$, the Cayley map contributes $d\theta_j/d\lambda_j = c_j$, so the first-order response of the critical gap is exactly the rank- ≤ 2 quadratic form

$$\delta'(u) = c_a q_a(u) - c_b q_b(u) = u^* K_{\{ab\}} u.$$

The remaining ingredient is that the local collision law $-\Lambda = \kappa \cdot \delta^2/8$ has $\kappa = 8\tau/\delta^2$ a *functional of the configuration*, not a constant. Differentiating $\tau = \kappa\delta^2/8$ therefore gives two terms, and this is the whole content of the earlier fitted factor:

$$\text{Marked depth derivative law. } D\tau_G[uu^*] = (\kappa(G) \cdot \delta(G)/4) \cdot u^* K_{\{ab\}} u + (\delta(G)^2/8) \cdot \kappa'(u).$$

Measured at $N = 6$ with $\kappa(G) = 1.280050$ taken directly from the configuration (not fitted): the correlation with measurement is 1.00000000 , the slope 1.00000 , and the residual $5.7 \cdot 10^{-11}$ — finite-difference accuracy. The local term alone gives correlation 0.9958 , slope 1.347 , residual $8.1 \cdot 10^{-2}$. So the previously reported empirical constant $\rho \approx 1.72$ was never a constant: fitting $\chi = \rho(\delta/4)\delta'$ forces $\rho = \kappa + (\delta/2)(\kappa'/\delta')$, which varies from mark to mark and averages to something larger than κ . **The law is now parameter-free**, and it decomposes cleanly into a local gap response and a background renormalisation response.

3.3 It is a polarisation detector, not an overlap detector

Write $T(u) = q_a + q_b$ for the mark's total mass in the critical two-plane and $I(u) = (q_a - q_b)/(q_a + q_b)$ for its imbalance. When the two critical eigenvalues are close, $c_a \approx c_b \approx c$ and

$$\delta'(u) \approx c \cdot T(u) \cdot I(u).$$

In the critical two-plane $K_{\{ab\}} = \text{diag}(c_a, -c_b) \approx c \cdot \sigma_z$. **The marked depth measures $P_a - P_b$, not $P_a + P_b$.** This corrects the natural guess: overlapping the critical pair is *not* sufficient. The mark $u = (v_a + v_b)/\sqrt{2}$ lies entirely inside the critical subspace, $T = 1$, yet has $I = 0$ and so gives no first-order response at all; the maximal-response marks are v_a and v_b themselves. The exact null cone is

$$\mathcal{N}_{\{ab\}} = \{u : c_a q_a = c_b q_b\} = \{u : q_a/(1 + \lambda_a^2) = q_b/(1 + \lambda_b^2)\},$$

which reduces to $q_a = q_b$ only when $\lambda_a \approx \lambda_b$.

3.4 Three confirmations

Blind tomography. Sampling 60 random marks and solving $\chi(u_k) \approx u_k^* K u_k$ for a symmetric K , without telling the solver which pair collides: the recovered K has top-two energy fraction 0.986 , and its leading two-dimensional eigenspace agrees with the true span $\{v_a, v_b\}$ to principal angles of $2.3 \cdot 10^{-6}$ and $3.0 \cdot 10^{-6}$ degrees. The residual 1.4% of energy is the κ' term, which is not rank-two. So the critical pair can be recovered from depth measurements alone.

The null cone is reachable, and reveals a floor. Constructing marks with $c_a q_a = c_b q_b$ to 10^{-16} while keeping $q_a = 0.244$, $q_b = 0.694$ leaves $\chi = -3.3 \cdot 10^{-3}$ to $-3.9 \cdot 10^{-3}$ — the same order as the orthogonal control, $2.8 \cdot 10^{-3}$. This is a genuine background channel, not a residue of the local one, and it is worth noting *why*: for $u \perp v_a, v_b$ the perturbation leaves λ_a and λ_b **exactly** fixed, not merely to first order, since $(G + \eta uu^*)v_a = \lambda_a v_a$. So the floor cannot come from critical-pair eigenvalue motion at all; it is background renormalisation and trajectory shift, which is precisely the κ' term. The experiments therefore separate $\chi = \chi_{\text{local}} + \chi_{\text{background}}$ cleanly.

The resolvent returns at second order. It was not meaningless, only mis-ordered. Second-order rank-one perturbation theory gives $\lambda_j'' = 2q_j \sum_{\{k \neq j\}} q_k/(\lambda_j - \lambda_k)$, so the denominators $1/(\lambda_j - \lambda_k)$ — the resolvent structure — govern the *next* term. The correct hierarchy is: first order, critical-pair differential overlap; second order, spectral mixing through the resolvent; then global heat-flow background renormalisation. On the null cone the first channel is switched off by construction, which makes it the natural place to test the second.

At finite η the two-plane picture continues: the mark contributes off-diagonal coupling $\alpha_a \bar{\alpha}_b$ inside $\text{span}\{v_a, v_b\}$, giving a 2×2 effective Hamiltonian whose gap has the usual square-root form. So the regimes are linear response, then an avoided-crossing (Landau – Zener) crossover as η approaches the original gap, then exact degeneracy where the off-diagonal term leads — a finite-dimensional local model, not an $N \times N$ resolvent computation.

3.5 The pattern this belongs to

Two very different compressions in this paper have the same shape. In the ACUE fibre, an enormous moment space collapses at the critical degree to a **rank-one** transport channel between conjugate rectangles. In the marked depth, an enormous spectral geometry collapses at first order to a **rank-two** contrast on the colliding pair. We record the pattern as a working principle rather than a theorem:

First-defect compression. When a model is indistinguishable from its target under all low-complexity observables, the first observable that does distinguish them tends to act through a channel of very low rank.

If that is right, the search for detectors should be reorganised. The natural objects are not further class functions $\text{Tr } \rho(U)$ or Schur polynomials $s_\lambda(U)$, but **marked matrix coefficients**: the quantity χ measures a dipole of the mark's spectral measure $\mu_u = \sum_j q_j(u) \delta_{\{\theta_j\}}$ across the critical pair, $c_a \mu_u(\theta_a) - c_b \mu_u(\theta_b)$. In the Langlands setting the analogous objects are periods and relative characters $\langle v, \pi(g)v \rangle$ rather than traces $\text{Tr } \pi(g)$ — which is an argument, from this side, that a relative trace formula is the right home for an arithmetic version of this detector.

4. What we can and cannot say about zero density

The other headline target adjacent to this circle of ideas is the zero-density exponent. Guth and Maynard's large-value estimate gives $N(\sigma, T) \leq T^{\{(30/13)(1-\sigma)+o(1)\}}$ and primes in intervals of length $x^{17/30+o(1)}$, the first improvement on the supremal exponent in four decades. Their argument bounds the top singular value of a Dirichlet-polynomial Gram matrix using its trace and a *centered third* spectral moment, and they note that estimates for fourth and higher powers are unavailable. It is tempting to conclude that supplying higher traces would improve the exponent.

We tested the pure moment-deflation channel exactly, and the answer is discouraging. Given only $(p_1, \dots, p_r)$ and m nonnegative eigenvalues, the sharp bound on the largest is $\lambda_1 \leq r + (p_1 - \lambda_1)^r / (m-1)^{r-1} \leq p_r$. For a spectrum consisting of one spike plus a flat bulk — the idealised large-value picture — **the $r = 3$ bound is already exactly tight, and $r = 4, 6, 8, 10$ buy precisely nothing** (gain 0.00% at $m = 100, 1000, 10000$, at spike fractions 0.5, 0.3, 0.15, 0.05). A gain appears only when the bulk itself carries secondary structure, and even then it is small: for a secondary cluster of 10 eigenvalues at 3% of the trace, going from $r = 3$ to $r = 10$ removes 1.1% of the overestimate; for 30 at 1.5%, 0.4%; for 100 at 0.5%, 0.1%.

The honest reading: whatever room exists in the Guth – Maynard framework is not in the moment hierarchy alone. It must come from the additive-energy input and the geometry of the large-value set, where higher traces could interact with the resonator structure rather than merely sharpening an already-tight scalar inequality. We have seen a proposed formula quantifying such a combined gain — a saving parameter $e = \eta + \varkappa$, a root y_e of $240y^2 + (72 - 10e)y - 13e = 0$, and an improved exponent $\Lambda(e) = 30/(13 + 10y_e) < 30/13$ — but we have not re-derived the Guth – Maynard argument and therefore neither endorse nor dispute it. It is recorded here as an unverified claim from a collaborating system, not as a result of this paper.

5. Records for clustered primes

5.1 The mechanism

Fix a set of offsets $H = \{h_1, \dots, h_k\}$ that is *admissible* — for each prime p the offsets miss a residue class mod p , so nothing forbids all of $n + h_1, \dots, n + h_k$ being prime at once. Maynard's method weights each n by a square $w(n) = (\sum_d \lambda_d)^2$ built from divisor sums, and compares the total weight $S_1 = \sum w(n)$ against the prime-counting weight $S_2 = \sum w(n) \cdot \#\{i : n + h_i \text{ prime}\}$. If the primes are equidistributed to level θ and a variational constant

$$M_k = \sup_F k \cdot J(F)/I(F), I(F) = \int \{R_k\} F^2, J(F) = \int \{R_{k-1}\} (\int_0^\cdot \{1-\Sigma\} F dt_k)^2$$

over symmetric F on the simplex R_k exceeds $2m/\theta$, then every admissible k -tuple contains $m + 1$ primes infinitely often, hence $H_m := \liminf(p_{n+m} - p_n) \leq H(k)$, the diameter of the narrowest admissible k -tuple. Bombieri – Vinogradov gives $\theta = 1/2$ unconditionally, so the criterion is $M_k > 4m$.

5.2 What everyone had missed

For $m = 1$ the relevant k is 50 and M_k has been computed exhaustively. For $m \geq 2$ the relevant k runs from ten thousand to fifty million, and there M_k had never been computed at all: every record since 2014 used the same crude closed-form bound, $M_k \geq \log k - C$, obtained by truncating a product test function hard at the simplex boundary. The improvements of 2023 – 2025 raised the arithmetic input θ and left that bound untouched.

Its deficit is 2.3 to 2.9 units of $\log k$, and about 1.1 units are recoverable elementarily. Since $M_k \approx \log k$ and $H(k) \approx k(\log k + 0.77)$, recovering 1.1 units is a factor $e^{\{1.1\}} \approx 3$ in k , hence nearly a factor 3 in H_m .

5.3 The engine

Take $F(t) = \prod_i g(t_i) \cdot 1[\sum t_i \leq k]$. Let X_i be i.i.d. with density g^2/c_2 , $c_2 = \int g^2$, write S_j for partial sums and $G(u) = \int_0^u g$. Then, exactly,

$$I(F) = k^{-k} c_2^{-k} \cdot P(S_k \leq k), J(F) = k^{-\{-(k+1)\}} c_2^{-\{k-1\}} \cdot E[G((k - S_{k-1})_+)^2],$$

and the truncation, discarded classically, becomes a genuine probability. The **layer-cake identity**

$$E[G((k - S)_+)^2] = \int 2G(u) g(u) \cdot P(S_{k-1} < k - u) du$$

converts it into an integral of true lower-tail probabilities, each bounded below rigorously (chord-majorised Chernoff, one-big-jump, Berry – Esseen with the safe non-i.i.d. constant 0.56); the bounds use only monotonicity of the true tail, so discretisation cannot invalidate them. Replacing the hard truncation by **shaped subexponential tails** $g(t) = e^{\{-(t/T_1)^\chi\}}/(1 + At)$ on a long support recovers the remainder: +0.12 units from exact truncation accounting, +0.49 from tail shaping, the rest from shape optimisation.

5.4 The results

	new bound	k	previous record
$H_2 = \liminf(p_{n+2} - p_n)$	173,438	15,856	396,504 (Stadlmann 2023/25)
H_3	13,859,802	923,601	24,797,814 (Polymath8b 2014)
H_4	1,120,662,828	56,000,000	1,431,556,072 (Polymath8b 2014)

Certificates $M_{15,856} \geq 8.013326752751$, $M_{923,601} \geq 12.006666706750$, $M_{56,10^6} \geq 16.065482942$, verified in ball arithmetic with outward rounding at every step, under three independent certification regimes (two Berry – Esseen constants $\times$ two tail routes). Tuples explicit: diameter 173,438 at $k = 15,856$ (admissibility re-verified by a second implementation); a repaired Hensley – Richards tuple of diameter 13,859,802 at $k = 923,601$, with a fully classical fallback of diameter 14,505,780; the primes-past- k tuple at $k = 5.6 \cdot 10^7$ with endpoints checked against published values of $\pi(x)$.

These are computer-assisted results produced by AI systems and have not been refereed; all certificates and scripts are published for replay. A further conditional improvement ($H_2 \leq 145,226$ via Deligne-strength equidistribution) is deliberately **not** claimed.

5.5 Why 246 did not move: five walls

1. **The ceiling is the tuple diameter.** No post-processing of Maynard – Tao output beats $H(k_{\min})$; pair-correlation constants cannot lower H_1 at all.
2. **Scalar decoding is exactly optimal.** A two-point counterfeit ordered by convex order defeats every matrix, inertia or higher-moment decode; the threshold $f(m) = 2m/\theta$ is final.
3. **The weight cone is closed.** Enlarging squares to PSD forms gains nothing — rank- r sums of squares decouple by subadditivity — and the copositive relaxation is flat, certified by an explicit nine-pattern dual with residual $1.3 \cdot 10^{-14}$.
4. **Parity, made combinatorial.** A Liouville twist kills a set of pair-conclusions iff its kill-graph is bipartite (odd-cycle facets of the cut polytope), giving the floors $H_1 \geq 6$ and $k \geq 2m + 1$. (*The graph fact is Tao's, from his 2014 parity-obstruction note; the cut-polytope packaging is ours.*)
5. **The usable arithmetic frontier is not the published frontier.** The levels $\theta = 4/7$ (BFI), $3/5$ (Maynard II) and $5/8$ (Pascadi) are well-factorable or fixed-residue, while Maynard – Tao needs uniformity over a CRT-structured residue system varying with the modulus.

The $k = 49$ door, priced exactly. The pure constant obeys the classical bound $M_{49} \leq (49/48) \cdot \log 49 = 3.97290 < 4$, so *no enrichment of the test-function basis can ever reach the threshold in the pure class* — higher power-sum rows, semidefinite relaxations and generalised eigenvalue tests are all dead on arrival there, and running them is wasted compute. Only the ϵ -variant survives, where the bound relaxes to $(1 + \epsilon) \cdot 3.97290$ and closes only for $\epsilon \leq 0.00682$. Our best certified value is $M_{49,1/35} \geq 3.930490592$ with float optimum 3.959325169, against the threshold 4: gap 0.0407. That, and $k = 47$ (payoff $H_1 \leq 226$), are the only doors left.

6. The signed sieve: an identity that closes a decade-old temptation

Maynard – Tao weights are squares, hence nonnegative, and positivity is what makes the decode valid. Chen's theorem, Iwaniec's linear-sieve weights $\lambda^\pm$, and Zhang's Landau – Siegel programme are all built from *signed* objects, so it is natural to ask what the signed enlargement buys. For signed w the decode fails — positive excess can be manufactured by $w(n) < 0$ at prime-poor n — and is repaired by paying the **debt** $D(w) = \sum_{\{w(n)<0\}} |w(n)| \cdot (m - \nu(n))_+$, where $\nu(n) = \#\{i : n + h_i \text{ prime}\}$.

Pointwise identity. For every integer $\nu \geq 0$ and every $m \geq 1$, $(\nu - m) + (m - \nu)_+ = (\nu - m)_+$.

Theorem (Signed No-Gain). With $w = w_+ - w_-$ of disjoint support,

$$S_2 - mS_1 - D(w) = \sum w_+(n)(\nu(n) - m) - \sum w_-(n)(\nu(n) - m)_+ \leq \sum w_+(n)(\nu(n) - m).$$

Proof. Substitute the identity in the w_- terms; the last step is $w_- \geq 0$ and $(v - m)_+ \geq 0$. ■

Corollary. Any DHL($k, m+1$) conclusion obtainable from a signed weight with the debt charged at face value is already obtainable from its positive part. The negative part is pure loss, of exactly the overshoot mass $\sum w_-(v - m)_+$.

The statement is pointwise: no model, no asymptotics, all k, m , tuples and weight classes at once. Verified additionally in exact rational arithmetic on a finite arithmetic microcosm — cellwise identity with zero violations, and the inequality on 1,000 random signed weights across five models.

What the numerics had been showing. Three phenomena looked like a phase structure; all three are corollaries. (i) Sweeping the debt price β in $\max\{\Phi - \beta D : S_1 = 1\}$, the value equals the classical optimum above a sharp β^* , computed here in exact rational arithmetic with a verified dual: classical value $1087376209/3212440751$, critical price $\beta^* = 23051796480/10991046857 = 2.0973249209\dots$, with 85% of the optimiser's mass negative on 16 of 96 cells, and unboundedness below $\beta_{\text{unb}} = 2.03265$. Since $\beta^* > 1$ always, the "phase" lives entirely at debt prices below face value. (ii) Imposing $\|w\|_1 \leq A$ gives $\lambda(1) = \lambda_{\text{positive}}$ exactly and then a positive slope $0.32 - 0.89$ — not a gain, since the decode is scale-invariant and what grows is the mass of w_+ . (iii) At $\beta = 1$ every model is unbounded, which exposes the second, usually unremarked job positivity performs: for $w \geq 0$ the normalisation $S_1 = 1$ is the ℓ^1 norm, so **positivity is what makes the variational problem bounded**.

Hence exactly two escapes, neither variational. (a) Charge the debt below face value — requiring arithmetic information that bounds $(m - v)_+$ on a designed negative support below the truth, which is precisely the exceptional-character mechanism. By the theorem, **Zhang's programme is not one route among several; it is the only route that changes the variational picture at all**. (b) Keep w evaluable while w_+ is not — the positive part of a divisor-sum quadratic is not a divisor-sum quadratic, so a signed well-factorable λ can be evaluable at $\theta = 4/7, 3/5, 5/8$ while its positive part is evaluable at no level beyond $1/2$.

An independent argument from the switching side agrees. Switching never reduces the number r of exact-primality conditions in a debt term; Chen's debt has $r = 1$ and is payable by an upper sieve, while every DHL($k, m+1$) with $m \geq 1$ forces residual debt with $r \geq 2$, and a two-vertex kill-graph is an edge, hence bipartite, hence parity-blocked wherever Liouville twists are admissible. Two disjoint analyses, one door.

The price list, if the debt were paid at level θ . All eight $m = 1$ crossings carry exact-rational certificates; $H(k)$ is Engelsma-exact.

θ	threshold	k pure (certified M_k)	k with ε	$H_1 \leq \text{pure} / \varepsilon$
$1/2$ (Bombieri – Vinogradov)	4	54	50	$270 / 246$ (unconditional)
$4/7$ (BFI)	3.5	31 (3.502015...)	29 (3.519881...)	$140 / 130$
$7/12$ (Maynard II)	$24/7$	29 (3.443305...)	26 (3.433616...)	$130 / 114$
$3/5$ (Maynard II)	$10/3$	26 (3.350647...)	23 (3.334616...)	$114 / 94$
$5/8$ (Pascadi)	3.2	22 (3.207656...)	20 (3.222666...)	$90 / 80$
1 (Elliott – Halberstam)	2	5 ($M_5 = 2.007080$)	—	12

None of BFI, Maynard II or Pascadi covers the sums the decode generates: all three fix one residue per modulus, while the decode's residues are CRT-composed from the shift set and vary with q . The sufficient intermediate statement is

(E₀). For coefficient systems $c_q(a)$ jointly well-factorable with the residue selection — for every factorisation $q = q_1 \cdot q_2$ (resp. $q_1 \cdot q_2 \cdot q_3$) with $\prod Q_j = x^{\{\theta-\epsilon\}}$ one can write $c_q(a) = \prod_j \gamma_j(q_j, a \bmod q_j)$, $|\gamma_j| \leq 1$, with $a \bmod p \in \{h_i - h_j \bmod p : j \neq i\}$ for $p \mid q$ — $\sum q \leq x^{\{\theta-\epsilon\}}$ $\sum \{a \in A_i(q)\} c_q(a) \cdot E(x; q, a) \ll_{\{H, A, \epsilon\}} x(\log x)^{-A}$.

(E_{4/7}) is "BFI Theorem 10, uniform over the CRT residue system of a fixed tuple polynomial". Polymath8a's $MPZ[\varpi, \delta]$ is its absolute-value cousin, proved by the same dispersion-plus-Deligne technology but only to level ≈ 0.5286 , so (E₀) interpolates two proved endpoints rather than crossing a parity barrier.

7. Formalisation and machine verification

Every theorem in § 1.7, § 1.8 and § 2 is elementary enough to formalise, and we report precisely what has been machine-checked, by which tool, and what has not.

7.1 What could not be done here, and why

We could not compile Lean. The session's egress proxy returns a 403 policy denial for `elan.lean-lang.org` and for `github.com` release downloads, and organisation policy denials are not to be retried. Public repositories are readable over git, but a Lean toolchain ships as a binary release, and building Lean plus Mathlib from source is not feasible here. The Lean file below is therefore **written but not compiled**; we present it as a formal specification rather than a verified artifact, and mark with `sorry` the two steps that need Mathlib's ODE comparison API, which we decline to guess at without a compiler.

7.2 What was machine-checked, and how

In place of Lean we discharged every algebraic and analytic step by exact symbolic computation (`sympy`, rational arithmetic throughout — no floating point) and, where the statement is a decision problem over the reals, by SMT (`z3`). Script `fab_verify_proof.py`: 18 checks, all passing.

step	statement	tool	status
1a	∂_s of the flow equals $(N \cdot D_z - D_z^2)$ on each monomial z^j	sympy, exact	✓
1b	$2z_j/(z_j - z_k) = 1 - i \cdot \cot((\theta_j - \theta_k)/2)$	sympy, exact	✓
1c	the $(N-1)$ terms cancel, leaving $\theta_j = -\sum \cot$	sympy, exact	✓
2a	$\cos(g/2) = e^s \cos(g_0/2)$ solves $g' = -2\cot(g/2)$	sympy, exact	✓
2b	the collision time is $-\log \cos(g_0/2)$, and the gap vanishes there	sympy, exact	✓
3	$d/dx \cot(x/2) = -\frac{1}{2} \csc^2(x/2)$, strictly negative on $(0, 2\pi)$	sympy, exact	✓

step	statement	tool	status
4a	$f(0) = 0$ and $f'(x) = \frac{1}{4}(2\tan(x/2) - x)$ for $f = -\log \cos(x/2) - x^2/8$	sympy, exact	✓
4b	$\tan t - t$ has nonnegative Taylor coefficients ($\frac{1}{3}, 2/15, 17/315, 62/2835, 1382/155925, \dots$)	sympy, exact	✓
5	the sign lemma as a real-arithmetic decision problem: unsat for the negation	z3	✓
6	the ACUE pigeonhole giving $\delta_{\min} = \pi/N$	exact integer reasoning	✓
7	$\sum_{k=1}^{N-1} \frac{1}{2} \csc^2(\pi k/N) = (N^2-1)/6$, at $N = 4, 6, 8, 12, 20$	sympy, exact	✓

Step 5 deserves comment, because it is the one place where a decision procedure is genuinely the right tool. Encoding the cosines and sines of the two half-angles as reals ca, sa, cb, sb subject to $ca^2 + sa^2 = 1$, $cb^2 + sb^2 = 1$, $sa, sb > 0$, together with the ordering constraint $sa \cdot cb - ca \cdot sb > 0$ (that is, $\sin(a-b) > 0$ for $0 < a-b < \pi$), z3 reports **unsat** for the negation $ca \cdot sb - cb \cdot sa \geq 0$ of the conclusion. The sign of every background term is therefore certified by a complete decision procedure over the reals rather than by inspection of a graph.

7.3 The Lean specification

```

/-- The cotangent, spelled directly so as not to depend on a particular Mathlib name. -/
noncomputable def cot (x : ℝ) : ℝ := Real.cos x / Real.sin x

lemma hasDerivAt_cot {x : ℝ} (hx : Real.sin x ≠ 0) :
  HasDerivAt cot (-(1 / Real.sin x ^ 2)) x

/-- `cot` is strictly decreasing on `(0, π)`: the sign input to Theorem A. -/
lemma cot_strictAntiOn : StrictAntiOn cot (Set.Ioo (0 : ℝ) π)

/-- **Background sign lemma.** For an adjacent pair the other zeros satisfy
`0 < x_b < x_a < 2π`, and the resulting bracket is negative – so it enters the gap
derivative with a positive sign and slows the collapse. -/
lemma background_sign {xb xa : ℝ} (h0 : 0 < xb) (h1 : xb < xa) (h2 : xa < 2 * π) :
  cot (xa / 2) - cot (xb / 2) < 0

/-- **The depth inequality.** `–log (cos (x/2)) ≥ x^2/8` on `[0, π)`. -/
lemma neg_log_cos_ge {x : ℝ} (h0 : 0 ≤ x) (h : x < π) :
  x ^ 2 / 8 ≤ -Real.log (Real.cos (x / 2))

/-- Exact solution of the two-body gap equation. -/
theorem two_body_solution
  (g : ℝ → ℝ) (hg : ∀ s, HasDerivAt g (-2 * cot (g s / 2)) s)
  (hrange : ∀ s, g s ∈ Set.Ioo (0 : ℝ) (2 * π)) (s : ℝ) :
  Real.cos (g s / 2) = Real.exp s * Real.cos (g 0 / 2)

/-- **Theorem A.** Every gap dominates the two-body solution started from the same
initial value, so no gap reaches `0` before `–log (cos (δ/2)) ≥ 6^2/8`. -/
theorem depth_ge {ι : Type*} [Fintype ι] (g : ι → ℝ → ℝ) (δ : ℝ)
  (hδ0 : 0 < δ) (hδπ : δ < π) (hmin : ∀ i, δ ≤ g i 0)

```

```
(hineq : ∀ i s, -2 * cot (g i s / 2) ≤ deriv (g i) s) :
6 ^ 2 / 8 ≤ -Real.log (Real.cos (6 / 2))
```

The full file — proofs for the first four items, `sorry` for the last two — is `Lean/DepthComparison.lean`.

`cot_strictAnti0n` is proved from the derivative via `strictAnti0n_of_hasDerivWithinAt_neg`; `background_sign` follows at once; `neg_log_cos_ge` reduces to $t \leq \tan t$ on $[0, \pi/2)$. What remains for a complete formalisation is `two_body_solution` (uniqueness for a scalar autonomous ODE) and the comparison step inside `depth_ge` (a Grönwall-type argument); both are standard and both are within Mathlib's reach.

7.4 The honest summary

The mathematics of Theorem A is verified: every algebraic identity by exact symbolic computation, the one inequality with real content by a complete decision procedure over the reals, and the numerical consequences on 11,060 background terms with zero exceptions. What is *not* verified is a Lean proof object, because no Lean toolchain could be obtained in this environment. The distinction is worth stating plainly, since "formalised" and "machine-checked" are often conflated and only the second is true here.

8. How this was done

The method is part of the result, so we describe it plainly.

Fleets, not oracles. Work proceeded in rounds. Each round posed one question and three to ten language-model agents attacked it in parallel from deliberately different angles, each writing and running its own code, each reporting line by line as *proved*, *computed*, *heuristic* or *conjecture*. Agents shared a written context file stating the current state of knowledge including known errors, but not conclusions: convergence of independent agents was treated as evidence, divergence as a bug report.

Adversarial defaults. Every context file instructed agents that the default assumption is that the idea has already been tried and failed, and that their first job is to find the reason. The prompt for the prime-gap round said outright that any improvement to 246 was 99% likely to be a misread constraint. Of the phenomena that looked like discoveries during this project, most were arithmetic or normalisation errors; the survivors survived because they were attacked first by their own authors.

Exact arithmetic at every threshold. Nothing near a decision boundary was accepted in floating point: ball arithmetic with outward rounding for the Maynard certificates, exact-rational Rayleigh quotients for the variational crossings, exact Vandermonde masses with 40-digit spot checks for the ACUE enumeration, exact rationals for the signed-sieve identity and the critical price β^* . The most common failure mode in machine-generated mathematics is a plausible float, and the remedy is cheap.

Two implementations or it did not happen. Every headline number here was produced twice by independent code, usually by agents that could not see each other's work: record tuples re-verified by a second admissibility checker; CUE depth by an ODE integrator and by coefficient bisection agreeing to 10^{-6} ; the single-dislocation constant by a lattice enumeration and by a transcendental double-root equation, agreeing to six digits. Where two implementations disagreed — an engine crossing at $k = 15,856$ against another at 29,500 — the discrepancy was tracked to its source before anything was claimed.

The human supplies the questions. Every genuinely new direction came from a human mathematician's judgement: to stop optimising 246 and look at H_2 ; to ask whether removing the square opens a phase; to propose the de Bruijn – Newman depth as a dynamic observable at all; to insist that the constant 72π be factorised rather than fitted; to recognise the single-

dislocation configuration as the right object to compute exactly. The models supplied speed, breadth, exact arithmetic, and the willingness to write and discard fifty scripts a day.

Negative results are the main product. Five walls, one no-go theorem, one empty phase, one refuted extrapolation, one refuted mechanism for the marked depth, one refuted hope for the moment-deflation channel. A method that only reports successes cannot be trusted about them. The single most valuable output of the signed-sieve round is that the route is closed and we can say exactly what would open it.

9. Open problems

1. **Complete the CUE depth law:** Theorems A and B (§ 1.7) reduce $8N^{\{8/3\}}(-\Lambda) \Rightarrow G^2$ to the cited smallest-gap law; what remains is a high-probability bound on the background coefficient $B \leq AN^2$ (§ 1.8) — a CUE rigidity statement controlling $\sum_k d_k^{\{-2\}}$ near the extremal pair. Measured: median $B/N^2 = 0.120$, max observed 0.79. This same bound completes Theorem C(ii).
2. **Bound ρ for the lattice:** prove $\rho^{\text{ACUE}} \leq C$, completing Law 2's upper bound. The lower bound $-\Lambda \geq \pi^2/(8N^2)$ is now exact (§ 1.7). This is a constant problem for one specific configuration family, not a universality problem.
3. **Make $s^* = 1.419640342\cdots$ a theorem** for the single-dislocation family, and compute the separated-defect constant $\rho^\infty = 1.19120\cdots$ exactly; the natural setting is an infinite clock with one localised defect under the Coulomb dynamics.
4. **The marked depth law:** prove $D\tau_G[uu^*] = (\kappa\delta/4) \cdot u^*K_{\{ab\}u} + (\delta^2/8) \cdot \kappa'(u)$ as an asymptotic identity with controlled error, and characterise the background term κ' . Numerically the law is exact (correlation 1.00000000, residual $5.7 \cdot 10^{-11}$); what is missing is the same localisation estimate as problem 1, differentiated. A natural sub-target: on the null cone $c_a q_a = c_b q_b$, show the leading behaviour is governed by $\lambda_j'' = 2q_j \sum_{\{k \neq j\}} q_k / (\lambda_j - \lambda_k)$.
5. **A transversality theorem:** with static observables M and fibre F_m , prove $\tau | \{F_m\}$ *generically non-constant, adding marks until* $\cap \ker D\Lambda\{u_j\} \cap \ker DM = \{0\}$.
6. **Transport the Palm certificates to the depth**, aiming at a bound $\mu_\Lambda < 1.8177$ — which by § 3.5 would improve the Lagarias – Rodgers upper bound on μ .
7. **(E $_\theta$)**, the tuple-residue well-factorable estimate, for any $\theta > 1/2$: the single statement between the certified price list and $H_1 \leq 130$.
8. **The $k = 49$ door**, in the ε -class only: certified 3.930490592, float 3.959325169, threshold 4, upper bounds closing only $\varepsilon \leq 0.00682$.

References

J. Maynard, *Small gaps between primes*, Ann. of Math. 181 (2015) 383 – 413 · D.H.J. Polymath, *Variants of the Selberg sieve* · Res. Math. Sci. 1:12 (2014) · Y. Zhang, *Bounded gaps between primes*, Ann. of Math. 179 (2014) · J. Stadlmann, Adv. Math. (2025) · E. Bombieri, J. Friedlander, H. Iwaniec, Acta Math. 156 (1986) · J. Maynard, *...large moduli II: well-factorable estimates*, Mem. AMS 1543 · A. Pascadi, arXiv:2505.00653 · K. Ford, J. Maynard, *On the theory of prime producing sieves*, arXiv:2407.14368 · T. Tao, *A general parity problem obstruction* (2014) · B. Rodgers, T. Tao, *The de Bruijn – Newman constant is non-negative*, Forum Math. Pi 8 (2020) · T. Tao, *The alternative hypothesis for unitary matrices* (2019) · J.C. Lagarias, B. Rodgers, *Band-limited mimicry of point processes by point processes supported on a lattice*, Ann. Appl. Probab. 31 (2021) 351 – 376; *Higher correlations and the Alternative Hypothesis*, Q. J. Math. 71 (2020)

257 – 280 · L. Guth, J. Maynard, *New large value estimates for Dirichlet polynomials*, arXiv:2405.20552 · G. Ben Arous, P. Bourgade, Ann. Probab. 41 (2013) · R. Feng, D. Wei, Ann. Probab. 49 (2021) · R. Killip, I. Nenciu, IMRN (2004) · T. Engelsma, minimal admissible tuples; A.V. Sutherland, narrow admissible tuples database.

Historical source 26: depth_scaling_theorem

Source: `historical/riemann-rmt/depth_scaling_theorem.md`. SHA-256:
742b21590598a8b4aa4d73bad590be85946f670433e6b682552a6b0d18c8930a.

The depth scaling law $-\Lambda \asymp N^{-2-2/(\beta+1)}$

Bill (Qingyun) Sun · GPT5.6SOL · Fable — proof document, August 2026

We prove the scaling law for the finite de Bruijn – Newman depth across the circular β -ensembles and the lattice (ACUE) endpoint, isolating exactly which step is deterministic, which is cited, and which remains open. Statuses: [P] proved here, [C] cited, [O] open.

0. Setup and conventions

For a monic $P(z) = \prod_{j=1}^N (z - e^{i\theta_j}) = \sum_{j=0}^N a_j z^j$ define the **flow**

$$P_s(z) = \sum_{j=0}^N a_j e^{s \cdot j(N-j)} z^j, s \geq 0,$$

and the **depth**

$$D(P) = \inf \{ s > 0 : \text{disc}(P_s) = 0 \} \in (0, \infty].$$

This is the finite de Bruijn – Newman depth with $D = -\Lambda$; the exponent $j(N-j)$ is the centred heat weight, and $s > 0$ amplifies the middle coefficients (the backward direction). All zeros start on the unit circle, and P_s remains self-inversive, so no simple zero can leave the circle without a prior collision; hence D is the first collision time.

Write $g_i(s)$ for the i -th cyclic gap, $\delta_{\min} = \min_i g_i(0)$. Call a pair (a, b) **adjacent** if no zero lies strictly inside the short arc from θ_b to θ_a .

1. The root dynamics [P]

Lemma 1. Until the first collision, the zeros of P_s satisfy

$$d\theta_j/ds = - \sum_{k \neq j} \cot((\theta_j - \theta_k)/2).$$

Proof. With $D_z = z \partial_z$ one has $\partial_s P_s = (N D_z - D_z^2) P_s$, since the symbol of D_z on z^j is j . Let $z_j(s)$ be a simple zero. Differentiating $P_s(z_j(s)) = 0$ gives $\dot{z}_j = -(\partial_s P)(z_j)/P'(z_j)$. At a zero, $P(z_j) = 0$, so

$$(N D_z - D_z^2) P|_{z_j} = N z_j P'(z_j) - [z_j P'(z_j) + z_j^2 P''(z_j)] = (N-1) z_j P'(z_j) - z_j^2 P''(z_j),$$

hence $\dot{z}_j = -(N-1) z_j + z_j^2 \cdot P''(z_j)/P'(z_j)$. For a simple zero of a product, $P''(z_j)/P'(z_j) = 2 \sum_{k \neq j} (z_j - z_k)^{-1}$, so

$$\dot{z}_j = -(N-1) z_j + 2 z_j^2 \sum_{k \neq j} (z_j - z_k)^{-1}.$$

Now put $z_j = e^{i\theta_j}$ and $\varphi = \theta_j - \theta_k$. Since $e^{i\theta_j} - e^{i\theta_k} = e^{i(\theta_j+\theta_k)/2} \cdot 2i \sin(\varphi/2)$ and $e^{i\theta_j} = e^{i(\theta_j+\theta_k)/2} e^{i\varphi/2}$,

$$2 z_j / (z_j - z_k) = e^{i\varphi/2} / (i \sin(\varphi/2)) = 1 - i \cdot \cot(\varphi/2).$$

Summing, $\sum_{k \neq j} 2 z_j / (z_j - z_k) = (N-1) - i \sum_{k \neq j} \cot(\varphi_{jk}/2)$. Substituting and using $\dot{\theta}_j = \dot{z}_j / (i z_j)$:

$$\dot{\theta}_j = [-(N-1) + (N-1) - i \sum \cot(\varphi_{jk}/2)] / i = - \sum_{k \neq j} \cot((\theta_j - \theta_k)/2). \blacksquare$$

The clock configuration is the unique fixed point up to rotation (all cotangent sums vanish by antisymmetry), and linearising about it returns the operator $\mathcal{L}_N$ with eigenvalues $\delta(N-\delta)$.

2. The two-body problem, solved exactly [P]

Lemma 2. The scalar equation $g' = -2\cot(g/2)$, $g(0) = g_0 \in (0, 2\pi)$, has the exact solution

$$\cos(g(s)/2) = e^{\{s\}} \cos(g_0/2),$$

so g reaches 0 precisely at $s = -\log \cos(g_0/2)$. Moreover, for $x \in [0, \pi)$,

$$-\log \cos(x/2) \geq x^2/8, \text{ with equality only at } x = 0.$$

Proof. Separating variables, $\tan(g/2) dg = -2 ds$ integrates to $-2 \log \cos(g/2) = -2s + C$; matching at $s = 0$ gives $\cos(g/2) = e^s \cos(g_0/2)$. Setting $g = 0$ gives $s = -\log \cos(g_0/2)$. For the inequality put $f(x) = -\log \cos(x/2) - x^2/8$; then $f(0) = 0$ and $f'(x) = \frac{1}{2}\tan(x/2) - x/4 = \frac{1}{4}(2\tan(x/2) - x) \geq 0$ on $[0, \pi)$ because $\tan t \geq t$ for $t \in [0, \pi/2)$. $\blacksquare$

(For $N = 2$ this is the whole story: $D = -\log \cos(\delta/2)$ exactly.)

3. The comparison theorem [P]

Theorem A (two-body comparison). For every configuration and every adjacent pair with gap g ,

$$g' \geq -2 \cot(g/2) \text{ for all } s \text{ before the first collision.}$$

Consequently

$$D \geq -\log \cos(\delta_{\min}/2) \geq \delta_{\min}^2/8.$$

Equivalently, writing $D = \rho \cdot \delta_{\min}^2/8$, one has $\rho \geq 1$ always.

Proof. Let (a, b) be adjacent with $g = \theta_a - \theta_b > 0$ measured across the short arc. By Lemma 1,

$$g' = \theta_a' - \theta_b' = -2\cot(g/2) - \sum_{k \neq a, b} [\cot(x_a^k/2) - \cot(x_b^k/2)], \quad x_j^k := (\theta_j - \theta_k) \bmod 2\pi \in (0, 2\pi).$$

Fix $k \notin \{a, b\}$. Adjacency means k does not lie in the open arc from θ_b to θ_a , so travelling counterclockwise from θ_k one meets θ_b before θ_a , i.e. $x_a^k = x_b^k + g$ with $0 < x_b^k < x_a^k < 2\pi$. On $(0, 2\pi)$ the map $x \mapsto \cot(x/2)$ is strictly decreasing (its derivative is $-\frac{1}{2}\csc^2(x/2) < 0$), so $\cot(x_a^k/2) - \cot(x_b^k/2) < 0$ and the whole bracket enters g' with a positive sign. Hence $g' \geq -2\cot(g/2)$.

Zeros cannot cross without colliding, so the cyclic order — and therefore the adjacency of each consecutive pair — is preserved up to the first collision. Applying the differential inequality to each gap and comparing with Lemma 2's solution started from the same initial value, every gap satisfies $g_i(s) \geq G(s; g_i(0))$ where $G(\cdot; g_0)$ vanishes first at $-\log \cos(g_0/2)$. Hence no gap vanishes before $\min_i (-\log \cos(g_i(0)/2)) = -\log \cos(\delta_{\min}/2)$, which is the claim. $\blacksquare$

Numerical check. The sign claim was tested on 11,060 background terms across random and lattice configurations: zero exceptions. The conclusion was tested on every dataset in this project; the minimum of $D/(-\log \cos(\delta_{\min}/2))$ is 1.0842, 1.0714, 1.0612 over the complete ACUE enumerations at $N = 6, 8, 10$ and 1.00019, 1.00021 over CUE samples at $N = 16, 64$.

4. The matching upper bound [P, under a regularity hypothesis]

Define the background stiffness of the critical pair

$S := \sum_{\{k \neq a, b\}} \frac{1}{2} \csc^2(x_b^k/2)$, so that $0 \leq -\sum_k [\cot(x_a^k/2) - \cot(x_b^k/2)] \leq g \cdot S$

by the mean value theorem, since $-(d/dx)\cot(x/2) = \frac{1}{2}\csc^2(x/2)$ and $x_a^k = x_b^k + g$. At the clock, $S = \sum_{k=1}^{N-1} \frac{1}{2}\csc^2(\pi k/N) = (N^2-1)/6$ exactly, so $S \asymp N^2$ is the natural scale.

Theorem B. Suppose $S(s) \leq A N^2$ throughout the collision window. Then

$$-2\cot(g/2) \leq g' \leq -2\cot(g/2) + A N^2 g,$$

and integrating the right-hand inequality from $g = \delta_{\min}$ down to 0,

$$D \leq (\delta_{\min}^2/8) \cdot (1 + O(A N^2 \delta_{\min}^2)),$$

uniformly while $A N^2 \delta_{\min}^2 < 2$. In particular $D = (\delta_{\min}^2/8)(1 + o(1))$ whenever $N^2 \delta_{\min}^2 \rightarrow 0$.

Proof. For small g , $2\cot(g/2) = 4/g + O(g)$, so $g' \leq -4/g + A N^2 g + O(g)$ and $ds \leq g dg/(4 - A N^2 g^2 + O(g^2))$. Integrating from 0 to $\delta_{\min}$ gives $D \leq -(2A N^2)^{-1} \log(1 - A N^2 \delta_{\min}^2/4) = (\delta_{\min}^2/8)(1 + A N^2 \delta_{\min}^2/8 + \dots)$. ■

Status of the hypothesis. For CUE it is verified numerically with a uniform constant: the median of S/N^2 is 0.109, 0.120, 0.117, 0.120, 0.120 at $N = 8, 16, 32, 64, 128$ (against the clock value $1/6 = 0.1667$), the 99th percentile stays below 0.36 and the largest observed value is 0.79. Turning this into a high-probability bound is the one [O] step; it is a standard rigidity statement (control of $\sum_k d_k^{-2}$ in the neighbourhood of the extremal pair).

5. The extreme-gap input [C]

Feng – Wei (Ann. Probab. 49 (2021)) for the circular β -ensemble, and Ben Arous – Bourgade (Ann. Probab. 41 (2013)) for $\beta = 2$: the smallest gap satisfies

$$\delta_{\min} \asymp N^{-1-1/(\beta+1)} \text{ in probability,}$$

with the k -th smallest gap density $\propto x^{k(\beta+1)-1} e^{-cx^{\beta+1}}$. For $\beta = 2$ the limit law is parameter-free: $P(N^{4/3} \delta_{\min} > x) \rightarrow \exp(-x^3/72\pi)$.

For the lattice endpoint no citation is needed:

Lemma 3 [P]. Every non-clock ACUE configuration has $\delta_{\min} = \pi/N$ exactly.

Proof. The support is the $2N$ -th roots of unity, so every gap is a positive multiple of π/N , and the N gaps sum to 2π , i.e. to N multiples of $2\pi/N$. If every gap were $\geq 2\pi/N$ they would all equal $2\pi/N$, which is the clock. Hence some gap equals π/N ; and no gap can be smaller, since π/N is the minimal positive multiple. ■

6. The theorem

Theorem. Let D_N be the depth of an N -point configuration.

(i) **Finite β .** For the circular β -ensemble with $\beta \in (0, \infty)$, assuming the regularity hypothesis of Theorem B,

$$D_N = (\delta_{\min}^2/8)(1 + O(\mathbb{P}(N^{-2/(\beta+1)+o(1)}))) \asymp N^{-2-2/(\beta+1)}.$$

(ii) $\beta = 2$ **distributionally.** Composing with Ben Arous – Bourgade and Slutsky,

$$N^{4/3} D_N \Rightarrow G^2, P(G > x) = \exp(-x^3/72\pi).$$

(iii) **Lattice endpoint** (ACUE, $\beta = \infty$). Unconditionally, for every non-clock configuration,

$$D_N \geq \pi^2/(8N^2), \text{ i.e. } N^2 D_N \geq \pi^2/8 = 1.2337005501\dots,$$

and with $\rho = O(1)$ — verified for $N \leq 10$ by complete enumeration, where $\rho \in [1.049, 1.610]$ — $D_N \asymp N^{\{-2\}}$, which is the exponent $-2-2/(\beta+1)$ at $\beta = \infty$.

Proof. Theorem A gives $D \geq \delta_{\min}^2/8$ in every case. For (i), Feng – Wei give $\delta_{\min} \asymp N^{\{-1-1/(\beta+1)\}}$, hence $N^2\delta_{\min}^2 \asymp N^{\{-2/(\beta+1)\}} \rightarrow 0$, so Theorem B's error term vanishes at that rate and the two bounds match. (ii) is (i) at $\beta = 2$ together with the parameter-free gap law and the continuous mapping theorem. For (iii), Lemma 3 makes the lower bound exact and pins $\delta_{\min} = \pi/N$, so $N^2\delta_{\min}^2 = \pi^2$: the error term in Theorem B does **not** vanish, and ρ is genuinely bounded away from 1. ■

7. Why $\beta = \infty$ is a singular, not a limiting, endpoint

The correction in Theorem B is governed by $\delta_{\min}^2 \cdot S$ with $S \asymp N^2$. Since $\delta_{\min} \asymp N^{\{-1-1/(\beta+1)\}}$,

$\delta_{\min}^2 \cdot S \asymp N^{\{-2/(\beta+1)\}},$

which tends to 0 for **every finite** β — so $\rho_\beta \rightarrow 1$ and the two-body constant $1/8$ is asymptotically exact. At $\beta = \infty$ the hard core pins $\delta_{\min} \asymp 1/N$ and $\delta_{\min}^2 \cdot S \asymp 1$: the background contributes at leading order and $\rho_\infty \neq 1$. The lattice endpoint is therefore not the smooth limit of the finite- β formula, and this is precisely why the single-dislocation configuration realises

$s^* = 1.419640342\dots = 1.150717118\dots \times \pi^2/8,$

a two-body collision time dressed by $\approx 15\%$ of many-body shielding.

The same computation predicts the *rate*, $\rho_\beta - 1 = O_{\mathbb{P}}(N^{\{-2/(\beta+1)+o(1)\}})$, which is a sharper statement than the leading exponent and is confirmed at three points:

β	predicted	fitted	local slope at largest N
1 (COE)	-1	-1.012	-0.851
2 (CUE)	$-2/3$	-0.710	-0.624
4 (CSE)	$-2/5$	-0.501	-0.407

8. Claim ledger

statement	status
Lemma 1, root dynamics $\dot{\theta}_j = -\sum \cot((\theta_j - \theta_k)/2)$	[P]
Lemma 2, exact two-body solution and $-\log \cos(x/2) \geq x^2/8$	[P]
Theorem A, $g' \geq -2\cot(g/2)$ and $\rho \geq 1$	[P]
Lemma 3, $\delta_{\min} = \pi/N$ for non-clock ACUE	[P]
Theorem B, upper bound given $S \leq AN^2$	[P] given the hypothesis
the regularity hypothesis $S \leq AN^2$ w.h.p. for $C\beta E$	[O] — verified numerically, median $S/N^2 = 0.120$
$\delta_{\min} \asymp N^{\{-1-1/(\beta+1)\}}$ for $C\beta E$	[C] Feng – Wei; Ben Arous – Bourgade at $\beta = 2$
$\rho_\infty = O(1)$ for ACUE at all N	[O] — proved for $N \leq 10$ by enumeration

statement	status
the resulting law $-\Lambda \asymp N^{-2-2/(\beta+1)}$, lattice included	follows from the above

The single open analytic ingredient is the regularity hypothesis; it closes (i), (ii) and the CUE side of the CUE/ACUE separation simultaneously. The lattice upper bound needs the separate, purely deterministic bound $\rho_{-\infty} = O(1)$.

Historical source 27: newman_depth_note

Source: `historical/riemann-rmt/newman_depth_note.md`. SHA-256:

959ce6ac7697de8bf679fd44f78065ffe0d5112ad3fa77c1e5dc9aecb3add5f.

The finite de Bruijn—Newman depth: an exact dynamic observable separating CUE from ACUE

Bill (Qingyun) Sun · GPT5.6SOL · Fable

August 18, 2026 — research note; complete numerical evidence, key identities proved

Summary

For a unitary $N \times N$ matrix U with $P(z) = \det(I - zU) = \sum a_j z^j$, define the **backward heat flow** $P_r(z) = \sum a_j r^{j(N-j)} z^j$ ($t = \log r \leq 0$) and the **finite Newman depth** $\Lambda(U) = -t^*$, where t^* is the first time the discriminant vanishes (first root collision). This is the exact finite-dimensional avatar of the de Bruijn – Newman constant: $RH \iff \Lambda_{\{dBN\}} \leq 0$ becomes "roots on the circle until collision", and $-\Lambda$ measures how far below criticality a configuration sits. Two independent computations (exact weighted enumeration of the ACUE ensemble for $N \leq 10$, 13,132 rotation orbits; Monte Carlo for Haar CUE up to $N = 256$, $\sim 44,500$ samples) establish:

1. **CUE is asymptotically barely-subcritical:** $-\Lambda^{\text{CUE}} \asymp N^{-8/3}$, with the parameter-free law $8N^{8/3}(-\Lambda) \Rightarrow G^2$, where $P(G > x) = \exp(-x^3/72\pi)$ is the sine-kernel smallest-gap law (constant 72π confirmed to $\sim 1.5\%$, KS ≈ 0.04 , median $29.6 - 30.5$ vs predicted $(72\pi \ln 2)^{2/3} = 29.08$). Fitted exponent -2.678 ± 0.016 .
2. **ACUE is over-stable:** $-\Lambda^{\text{ACUE}} \asymp N^{-2}$ (fitted slope -2.0009), with $N^2(-\Lambda)$ supported in $\approx [1.31, 1.99]$, median $\rightarrow \approx 1.41$. The alternative $N^{-8/3}$ is decisively ruled out on the lattice, and the alternative N^{-2} is decisively ruled out for CUE: the depth ratio falls as $\approx 3.6 \cdot N^{-0.70}$ (theory $N^{-2/3}$), reaching $10\times$ separation at $N^* \approx 165$.
3. **Clock configurations:** $P_{\text{ACUE}}(\text{clock}) = 2^{1-N}$ **exactly** (proved by Cauchy – Binet: $\sum_{|C|=N} |\Delta(C)|^2 = (2N)^N$, each clock contributing N^N), and clocks have $\Lambda = -\infty$ ($P = 1 - cz^N$ is flow-invariant). These are the unique flow-stationary configurations.
4. **Local collision law:** $-\Lambda = \delta_{\min}^2/8 \cdot (1 + o(1))$ holds in the continuum regime $\delta_{\min} \ll$ other gaps (Richardson-verified to 10^{-7} ; $N=2$ exact: $-\Lambda = -\log \cos(\delta/2)$). For CUE the correction is positive and vanishes as $\approx 0.60 \cdot N^{-0.73}$ (the background *delays* collision). On the ACUE lattice $\delta_{\min} = \pi/N$ always (proved: all gaps $\geq 2\pi/N$ forces the clock), neighboring gaps are the same scale, and the constant renormalizes: $\rho = 8(-\Lambda)/\delta_{\min}^2 \in [1.066, 1.609]$, with the isolated-defect-pair value $\rho_{\infty} \approx 1.1912$ (N -independent to 4 digits by $N = 8$).
5. **The depth escapes the weight-freezing no-go.** The flow is diagonal on coefficients, so every balanced moment of degree $\leq N$ remains frozen along the flow (dynamic extension of the Round-3 freezing theorem; verified to $4.5e-16$). Yet Λ — an "infinite-degree" observable — separates the ACUE fiber at $O(1)$: a mimicker in the $N=5$ $\mathbb{Q}(\sqrt{5})$ family shifts $E[N^2(-\Lambda); \text{non-clock}]$ by -0.093 and the clock atom from 0.0625 to 0.1398 ; at $N = 8$ mimicker families move the Λ -law by total variation $0.12 - 0.24$, with the dependence concentrated on the center-of-mass class $X = 0$. Λ is the **first verified dynamic anti-ACUE observable**.

Why this matters

The ACUE (Tao's alternative-hypothesis ensemble) matches CUE on all pair statistics available to bandwidth-limited detectors — the wall behind every Montgomery-type approach. All static balanced moments of degree $\leq N$ are frozen across the mimicker fiber, so no polynomial statistic of low degree can tell the ensembles apart. The Newman depth is dynamic and non-polynomial: it sees the full collision geometry of the heat flow. The scaling gap $N^{-8/3}$ vs N^{-2} says the two hypotheses predict *different universality classes of subcriticality* — GUE-repulsion ($P(\delta) \sim \delta^3 \Rightarrow$ extreme gap $N^{-4/3} \Rightarrow$ depth $N^{-8/3}$) versus lattice rigidity ($\delta_{\min} = \pi/N$ deterministic $\Rightarrow$ depth N^{-2}). Any future zeta-side quantity that mimics finite depth (e.g. flowed pair-counts under the actual de Bruijn – Newman backward flow on Ξ) inherits this separation — a concrete new place to look for a $>2/3$ simplicity proportion or an anti-AH statistic, orthogonal to the bandwidth war.

Key identities (proved)

- Root dynamics under the flow: $\dot{\theta}_j = -\sum_{k \neq j} \cot((\theta_j - \theta_k)/2)$ (attracting Coulomb gas on the circle); a simple root cannot leave $|z| = 1$ before a collision (self-inversive invariance).
- Generator on power sums: $L p_m = -m(N-m) p_m - m \sum_{a=1}^{m-1} p_a p_{m-a}$ (verified to $4.1e-28$ across $N = 5..10$; closes the polynomial-moment hierarchy and proves the freezing corollary).
- $N = 2$: $\Lambda = \log(|\operatorname{tr} U|/2)$, i.e. $-\Lambda = -\log \cos(\delta/2)$.
- $P_{\text{ACUE}}(\text{clock}) = 2^{1-N}$; clocks are the exact $\Lambda = -\infty$ atoms; ACUE(8) median($-\Lambda$) cross-validated by both implementations ($2.214e-2$ vs $2.216e-2$).

Verification ledger

item	status
Exact enumeration $N = 3..10$ (13,132 orbits, Vandermonde masses; $\sum \mu = 1$ to $3e-15$)	complete, mpmath 40-digit spot checks worst $1.3e-9$
CUE Monte Carlo $N = 2..256$, two independent Λ -solvers (ODE vs coefficient bisection)	agree to $1e-6$; $N=2$ analytic law to $8.7e-13$
First collision always at an initially-adjacent pair (ACUE)	13,130/13,130 orbits, zero surprises
Clock probability 2^{1-N} , $\Lambda = -\infty$	proved (Cauchy – Binet; flow invariance)
$\delta_{\min} = \pi/N$ for non-clock ACUE	proved (pigeonhole) + verified all orbits
72π smallest-gap constant, G^2 -law for CUE depth	fitted $c = 229 - 236$ vs 226.2 ; KS $0.035 - 0.041$
Freezing along flow (balanced degree $\leq N$)	proved (diagonal flow) + $4.5e-16$ numeric

Open target (sharp): the closed form of the lattice defect constant $\rho_\infty \approx 1.1912$ (two-defect reduction of the Coulomb ODE against the rigid period-2 background); a block-decomposition of the same computation would upgrade the N^{-2} law to proved-with-constants.

Scripts and data: session archive `dyn1_*.py`, `dyn1_results_N{3..10}.npz`, `dyn2_*.py`, `dyn2_data_N{2..256}.npz`.
References: Ben Arous – Bourgade (*extreme gaps*), Feng – Wei (*C β E small gaps*), Vinson; Rodgers – Tao ($\Lambda_{\text{dBN}} \geq 0$); Tao (*ACUE*).

Historical source 28: signed_sieve_nogo

Source: `historical/riemann-rmt/signed_sieve_nogo.md`. SHA-256:

8802c586ba28180b3a3eed9ef61bd232355de5b30585f0c0fd21c426a16d038d.

The signed Maynard–Tao sieve: a no–gain theorem, and where the door actually is

Bill (Qingyun) Sun · GPT5.6SOL · Fable

August 18, 2026 — research note. Answers the question "does removing positivity from the Maynard – Tao variational problem open a new phase?" The answer is no, for a reason that is elementary, exact, and — this is the useful part — tells you precisely what would have to replace positivity.

0. The question

Maynard – Tao weights are squares, $w(n) = (\sum_d \lambda_d)^2$, hence pointwise nonnegative. Positivity is what makes the decode valid: if $w \geq 0$ and $S_2 - mS_1 > 0$, where $S_1 = \sum_n w(n)$ and $S_2 = \sum_n w(n)v(n)$ with $v(n) = \#\{i : n + h_i \text{ prime}\}$, then some n has $v(n) \geq m + 1$. Zhang's Landau – Siegel program ("摆脱平方" — remove the square) and Iwaniec's well-factorable $\lambda^\pm$ both suggest that signed weights are the natural next class, and that positivity is a self-imposed restriction. The concrete question we were asked to settle: **after allowing signed/indefinite sieve weights, does the variational optimum exhibit a phase the classical positive family cannot see?**

For a signed w the decode fails as stated — the positive excess can be manufactured by $w(n) < 0$ at prime-poor n . It is repaired by paying the **debt**

$$D(w) = \sum_n \{n : w(n) < 0\} |w(n)| \cdot (m - v(n))_+,$$

so that $S_2 - mS_1 - D > 0$ does imply the conclusion. The experiment is then to optimize the repaired functional $\Phi(w) - \beta D(w)$ over the signed class, with $\beta = 1$ the true price, and look for a phase transition in β or in any enlargement parameter.

1. The theorem

Pointwise identity. For every integer $v \geq 0$ and every $m \geq 1$,

$$(v - m) + (m - v)_+ = (v - m)_+.$$

(Both sides equal $v - m$ when $v \geq m$, and 0 when $v < m$.)

Theorem (Signed No-Gain). Write $w = w_+ - w_-$ with $w_+, w_- \geq 0$ of disjoint support. Then

$$S_2 - mS_1 - D(w) = \sum_n w_+(n)(v(n) - m) - \sum_n w_-(n)(v(n) - m)_+ \leq \sum_n w_+(n)(v(n) - m).$$

Proof. $S_2 - mS_1 = \sum w_+(v-m) - \sum w_-(v-m)$ and $D = \sum w_-(m-v)_+$; subtract and apply the identity to the w_- terms. The final inequality holds because $w_- \geq 0$ and $(v-m)_+ \geq 0$. ■

Corollary. Any DHL($k, m+1$) conclusion obtainable from a signed weight w with the debt charged at face value is already obtainable from the nonnegative weight w_+ : if the repaired functional at w is positive, then the plain functional at w_+ is positive. **The negative part is pure loss**, and the loss is exactly $\sum w_-(n)(v(n) - m)_+$ — the overshoot mass sitting on the

negative support.

The theorem is pointwise, so it needs no model, no cell structure, no asymptotics, and it holds for every k , m , tuple, and weight class simultaneously. It subsumes the weaker fact we proved first (that the *naked* ratio $\sup \text{Tr}(\text{BQ})/\text{Tr}(\text{AQ})$ over indefinite Q equals the PSD value by rank-one extremality): the naked-ratio statement is about dropping positivity from the objective; this one says that even after correctly repairing the decode, the enlargement is empty.

Verification. Exact rational arithmetic on the finite arithmetic microcosm (n uniform over $\mathbb{Z}_W$, $W = \prod \text{feature primes} \times \prod \text{big primes}$; v = coprimality count on the tuple; weights in the span of level- L features): the cellwise form of the identity holds with zero violations on all 96/150/180/24-cell models, and the inequality holds on 1,000 random signed weights across five models ($k = 3, 4, 5$; $m = 1, 2$; $L = 1, 2$). Script `fab_theorem.py`.

2. What the numerics were really showing

Before the identity was noticed, the LP experiments produced three facts that looked like a phase structure. All three are corollaries of the theorem, and each is instructive.

(a) **The critical-price kink.** Sweeping β in $\max\{\Phi(w) - \beta D(w) : S_1(w) = 1\}$, the value is *exactly* the classical positive optimum for β above a sharp β^* , then acquires a signed optimizer below it. Certified example ($k = 3$, $m = 1$, features $\{2, 3, 5, 7\}$), exact-rational simplex with verified dual: $\text{val_pos} = 1087376209/3212440751 = 0.3384891094603102$, signed vertex $\Phi = 1.2082816957\dots$, $D = 0.4147152297\dots$, giving the exact critical price

$$\beta^* = (\Phi - \text{val_pos})/D = 23051796480/10991046857 = 2.0973249209031\dots,$$

the optimizer carrying 85% of its mass negative on 16 of 96 cells. Certified on both sides: at $\beta^* - 10^{-3}$ the optimum is $0.3389038246899884 > \text{val_pos}$ with a genuinely signed optimizer; at $\beta^* + 10^{-3}$ it equals val_pos as exact rationals. Below $\beta_{\text{unb}} = 2.03265$ the LP is unbounded (in particular at $\beta = 2$, certified). Across eight model variants β^* ranged $1.44 - 6.72$, and in five of the eight the window $(\beta_{\text{unb}}, \beta^*)$ was numerically **empty** (width $< 10^{-6}$): the value falls off the classical plateau directly into unboundedness. Table in `fab_phase_results.json`.

The theorem explains this: $\beta^* > 1$ always, because at $\beta = 1$ the signed class is dominated. The "phase" lives entirely in the unphysical region $\beta < 1$, where the debt is charged at less than face value.

(b) **The apparent linear gain in the coefficient budget.** Adding the constraint $\|w\|_1 \leq A$ gives $\lambda_{\text{signed}}(A)$ with $\lambda_{\text{signed}}(1) = \lambda_{\text{positive}}$ exactly (since $\|w\|_1 = S_1 = 1$ forces $w \geq 0$), and then a strictly positive slope: $d\lambda/dA$ at $A = 1^+$ equals $0.320, 0.504, 0.894, 0.394$ in the four models (60 – 170% of the classical value per unit of budget), with the constraint active at every A and asymptotically linear growth. This looks like a large gain and is not one: the decode criterion is scale-invariant (only the *sign* of $S_2 - mS_1 - D$ matters), and by the theorem the repaired value is $\Phi(w_+)$ minus a loss, so what grows with A is simply the total mass of w_+ , not the truth of the conclusion. `fab_norm_results.json`.

(c) **Unboundedness at the true price.** At $\beta = 1$ every model is unbounded. Same cause: the affine slice $\{S_1 = 1\}$ does not bound $\|w_+\|$ once negative mass is allowed, and $\Phi(w_+)$ scales with it. For $w \geq 0$ the normalization $S_1 = 1$ is the ℓ^1 norm — that is the second, usually unremarked job positivity does in Maynard – Tao: it makes the variational problem bounded. Remove it and boundedness must be re-imposed by hand; no choice of norm bound creates a gain, by the theorem.

3. Where the door actually is

The theorem has exactly two hypotheses, so there are exactly two escapes, and they are not variational.

(i) **Charge the debt at less than face value ($\beta < 1$).** This requires arithmetic information that bounds $(m - v)_+$ on the *designed* negative support by strictly less than the truth — i.e. an input that says "on this set, primes are more common than the trivial bound allows". This is precisely the exceptional-character mechanism: near a Landau – Siegel zero the distribution of primes in the relevant progressions is anomalously rigid, and the shortfall on a set designed around the exceptional modulus can be bounded below face value. **Zhang's program is not one route among several; by the theorem it is the only route that changes the variational picture at all.**

(ii) **Keep w evaluable while w_+ is not.** The theorem compares w with w_+ , which presumes both are usable. In the arithmetic setting they are not interchangeable: the positive part of a divisor-sum quadratic is not itself a divisor-sum quadratic, so a signed, well-factorable λ can be evaluable at level $\theta = 4/7, 3/5, 5/8$ while its positive part is not evaluable at any level beyond $1/2$. **This is the entire content of the signed route: the gain is never variational, it is exclusively about which weights the arithmetic can evaluate.** It is therefore correctly priced by the conditional ledger (§ 4), not by any optimization over weight cones.

Independent confirmation from the switching side. The Chen-switch audit reaches the same conclusion by a disjoint argument. Switching is a re-indexing bijection: it reduces the anatomical depth of one linear form but never the number r of exact-primality conditions a debt term carries. Chen's debt has $r = 1$ and is payable by an upper sieve (parity-free). Every $\text{DHL}(k, m+1)$ conclusion with $m \geq 1$ forces residual debt with $r \geq 2$; a two-vertex kill-graph is an edge, hence bipartite, hence parity-blocked for every input class admitting Liouville twists. The escape is not a cleverer switch but an input class in which the twist is inadmissible — the exceptional-character regime again, or full GEH. Two independent analyses, one door.

4. The conditional ledger (what the door is worth)

If the debt is paid at level θ , the price list is fully computed. All eight $m = 1$ crossings below carry exact-rational certificates produced this session (Galerkin engine, full-vector Rayleigh quotients, no truncation loss); $H(k)$ values are Engelsma-exact and were independently re-proved here for $k \leq 62$.

θ	threshold $2/\theta$	$k_{\min}$ pure (certified M_k)	$k_{\min} \varepsilon$ (certified $M_{\{k,\varepsilon\}}$)	$H_1 \leq \text{pure} / \varepsilon$
1/2 (Bombieri – Vinogradov)	4	54	50	270 / 246 (unconditional, current record)
4/7 (BFI)	3.5	31 (3.502015495...)	29 (3.519881249...)	140 / 130
7/12 (Maynard II, linear sieve)	$24/7 \approx 3.4286$	29 (3.443305315...)	26 (3.433616497...)	130 / 114
3/5 (Maynard II)	$10/3 \approx 3.3333$	26 (3.350647068...)	23 (3.334615948...)	114 / 94
5/8 (Pascadi)	3.2	22 (3.207656229...)	20 (3.222665844...)	90 / 80
1 (Elliott – Halberstam)	2	5 ($M_5 = 2.007080$)	—	12

$m = 2$ rows (p9 product-profile engine, certified-lower-bound path, rigorous primes-past- k diameters): $\theta = 4/7 \rightarrow k = 5,647$, $H_2 \leq 58,058$; $7/12 \rightarrow 4,835, 48,988$; $3/5 \rightarrow 3,931, 38,878$; $5/8 \rightarrow 3,022, 29,180$; EH $\rightarrow 221, 1,498$. $m = 3$: $4/7 \rightarrow 202,528$; $7/12 \rightarrow 160,703$; EH $\rightarrow 1,978$ ($H_3 \leq 18,144$). Every row is parity-consistent (all $\geq$ the floor 6, and $\geq$ the EH value 12).

The missing estimate, stated precisely. None of BFI (Acta Math. 156, Thm 10), Maynard II (Mem. AMS 1543), or Pascadi (arXiv:2505.00653) covers the sums the decode generates: all three fix one residue a for every modulus, while the decode's residues $a \in A_i(q)$ are CRT-composed from the fixed shift set $\{h_i - h_j\}$ and therefore vary with q . What would suffice is the intermediate statement

(E₀). Fix an admissible tuple H , an index i , and $A, \varepsilon > 0$. For coefficient systems $c_q(a)$ jointly well-factorable with the residue selection — for every factorization $q = q_1 q_2$ (resp. $q_1 q_2 q_3$) with $\prod Q_j = x^{\theta - \varepsilon}$ one can write $c_q(a) = \prod_j \gamma_j(q_j, a \bmod q_j)$, $|\gamma_j| \leq 1$, with $a \bmod p \in \{h_i - h_j \bmod p : j \neq i\}$ for all $p \mid q$ — $\sum_{a \in A_i(q)} c_q(a) \cdot E(x; q, a) \ll_{H, A, \varepsilon} x (\log x)^{-A}$.

(E_{4/7}) is "BFI Theorem 10, uniform over the CRT residue system of a fixed tuple polynomial". Evidence that the gap is bridgeable rather than parity-hard: Polymath8a's MPZ[ϖ, δ] is exactly this statement's absolute-value cousin, proved by the same dispersion + Deligne technology, but only to level $1/2 + 2\varpi \leq 0.5286$ and for densely-divisible moduli. So (E₀) interpolates two proved endpoints — bilinear strength with one residue, and residue flexibility at weak level.

5. Verdict

No new phase. The signed enlargement of the Maynard – Tao variational problem is empty: with the decode debt charged at face value, every signed weight is dominated by its own positive part, by an exact pointwise identity. This kills the route cheaply and completely, as intended — and it converts an open-ended search into two sharply stated arithmetic requirements, (E₀) for evaluability and the exceptional-character input for sub-face-value debt. The variational side of the signed programme is closed; everything of value in it is arithmetic.

Artifacts: `fab_theorem.py` (identity, exact verification), `fab_phase.py` / `fab_phase_results.json` (β -structure, β^* , β_{unb} across eight models), `fab_norm.py` / `fab_norm_results.json` (ℓ^1 -budget ramp), `sgn1_*` (finite model + LP engines), `sgn2_*` (ledger certificates, `sgn2_certificates.json`, `sgn2_mk_large.json`).

Historical source 29:

H2_H3_record_announcement

Source: `historical/riemann-rmt/H2_H3_record_announcement.md`. SHA-256:
`0c2bebc885a4250015b5a81270f41afeda4dbc4dd8afb7454cd3229fb3774223`.

New bounds for gaps between consecutive primes with two and three primes between

$$H_2 \leq 173,438, H_3 \leq 13,859,802, H_4 \leq 1,120,662,828$$

Bill (Qingyun) Sun · GPT5.6SOL · Fable

August 15, 2026 (updated same day) — research announcement; independent expert review invited

Statements

Theorem A. $\liminf_{n \rightarrow \infty} (p_{n+2} - p_n) \leq 173,438$. (Previous record: 396,504, Stadlmann 2023/2025; before that 398,130, Polymath8b 2014.)

Theorem B. $\liminf_{n \rightarrow \infty} (p_{n+3} - p_n) \leq 13,859,802$. (Previous record: 24,797,814, Polymath8b 2014.)

Theorem C. $\liminf_{n \rightarrow \infty} (p_{n+4} - p_n) \leq 1,120,662,828$. (Previous record: 1,431,556,072, Polymath8b 2014.)

Certificate $M_{56,000,003} \geq 16.0655 > 16$; tuple = the first $56 \cdot 10^6$ primes exceeding $56 \cdot 10^6$ (first 56,000,003, last 1,176,662,831; π -anchors verified against published values).

Hypothesis chain (complete)

All three theorems follow from exactly three published inputs plus machine-verified certificates:

1. **Maynard's theorem** [Maynard, *Small gaps between primes*, Ann. of Math. 181 (2015), Prop. 4.1 + Thm 3.1 pipeline]: if the primes have level of distribution θ and $M_k > 2m/\theta$, then DHL($k, m+1$): every admissible k -tuple has infinitely many translates containing $\geq m+1$ primes. Our certificates use test functions supported in the **closed standard simplex** — no ε -enlargement, no vanishing-marginal variant, no equidistribution beyond (2) — so the criterion applies verbatim.
2. **Bombieri – Vinogradov** (1965): every $\theta < 1/2$ is a level of distribution; hence $M_k > 4m$ suffices (our margins 0.013 and 0.0067 leave room to fix $\theta < 1/2$).
3. **Berry – Esseen inequality** with the safe non-iid constant $C = 0.56$, used inside the certified lower bound for the simplex-truncation probability; applied in a regime using only elementary normal-tail bounds $\bar{\Phi}(z) \leq \min(\varphi(z)/z, e^{-z^2/2}/2)$ (no special-function dependence).

Certificates. $M_{15,856} \geq 8.013326752751$ and $M_{9,23,601} \geq 12.006666706750$, produced by a one-dimensional product-profile engine (below) and verified in ball/interval arithmetic with outward enclosures (all integrands polynomial $\rightarrow$ exact rational integrals; exponentials majorized by chords/Taylor enclosures; every inequality rounded against the result;

certificate files `p9_exact_cert_k15856.json`, `p9_exact_cert_k923601.json`, `p9_exact_cert_k56000000.json` serialize the full chains). Three independent certification regimes (two Berry – Esseen constants $\times$ two tail-bound routes) pass at both k ; Monte Carlo upper-consistency checks and small- k sanity bounds (the engine stays below the known M_2 , M_3 , M_4 , M_5) all pass.

Tuples. $k = 15,856$: an explicit admissible tuple of diameter 173,438 (file `p9_tuple_k15856.npy`), admissibility re-verified independently by a second implementation (every prime $p \leq 15,856$ misses a class). $k = 923,601$: a repaired Hensley – Richards tuple of diameter 13,859,802 (symmetric window $\{\pm 1\} \cup \{\pm q : q \text{ prime} \in [45,007, 6,929,899], 5,692 \text{ mid-size primes deleted by a deletion-fixpoint repair} \cup \{+6,929,903\}$; file `p9_tuple_k923601.npy`, sha256 d5fe6890...6ff02), admissibility over all 73,001 primes $p \leq k$ verified by two independent implementations; the simpler primes-past- k tuple (diameter 14,505,780, verified by direct sieve) provides a fully classical fallback. Combining certificate + tuple through Maynard's theorem yields Theorems A – C.

The method (what is actually new)

The variational lower-bound engine, not the framework. For $F = \prod_i g(x_i) \cdot 1[\sum x_i \leq k]$ the Maynard functionals reduce exactly to one-dimensional objects:

- $I(F) = k^{-k} \cdot P(S_k \leq k)$, $J(F) = k^{-(k+1)} \cdot E[G((k - S_{k-1})_+)^2]$, where X_i iid with density g^2/c_2 , $c_2 = \int g^2$, $G(u) = \int_0^u g$;
- the **layer-cake identity** $E[G((k-S)_+)^2] = \int 2G(u)g(u) \cdot P(S_{k-1} < k-u) du$ turns the simplex truncation into an integral of *true lower-tail probabilities*, each bounded below by $1 - \beta(u)$ with rigorous β (chord-majorized Chernoff / one-big-jump / Berry – Esseen) — per-piece bounds use only the monotonicity of the true tail, so grid artifacts cannot invalidate them;
- shaped subexponential tails** $g = e^{-(t/T_1)^\chi} / (1 + At)$ on a long support replace the hard truncation of the classical treatment.

Against the crude closed-form truncation bound used for all $m \geq 2$ records since 2014 (deficit from $\log k : \approx 2.3 - 2.9$), the exact layer-cake accounting (+0.12) and shaped tails (+0.49) recover ≈ 1.1 units of deficit, i.e. a factor ≈ 3 in k , i.e. the factors 2.29 ($m=2$) and 1.71 ($m=3$) above. Head-to-head reconciliation against an independently written engine reproducing the 2014 structure confirms both the numbers and the provenance of the gap; the 2023/2025 record improvements had upgraded only the *arithmetic* input while retaining the crude variational bound, which is why this gain was still on the table.

Verification ledger

item	status
Engine identities (I, J, layer cake) and correction directions	proved, re-derived independently in the audit
$M_{15,856} \geq 8.0133$, $M_{923,601} \geq 12.0067$, $M_{56,000,000} \geq 16.0655$	ball-arithmetic certificates, 3 regimes each
$k=15,856$ tuple: admissible, diameter 173,438	verified by two independent implementations
$k=923,601$ H-R tuple: admissible, diameter 13,859,802	verified by two independent implementations; classical fallback 14,505,780 also verified
$k=56,000,000$ tuple: primes-past- k , diameter 1,120,662,828	segmented sieve, π -anchors match published values

item	status
Threshold chain (Maynard + BV, closed simplex, zero extra charges)	audited against the theorem statement
Engine-vs-engine discrepancy (15,856 vs 29,500 crossing)	fully explained (layer-cake + tail shape)

Declared limitations. (i) These are computer-assisted results by AI systems; the certificate files and scripts are published for expert replay, and formal peer review has not yet occurred. (ii) A further conditional improvement ($H_2 \leq 145,226$ via $k = 13,476$ with Polymath8a/Deligne-strength equidistribution) is *not* claimed: it depends on a cap-normalization in the truncated-variational criterion that we have reconstructed but not verified verbatim against Polymath8b's theorem. (iii) $H_1 = 246$ is untouched: our parallel investigations proved the relevant walls (ceiling = tuple diameter; decode optimality; weight-cone closure) and located the $k = 49/47$ variational doors, which remain open but uncrossed (best certified $M_4 \geq 3.9305$, float 3.9593 vs threshold 4).

Artifacts. All engines, certificates, tuples, and audit scripts in `research/riemann-rmt/` and the session archive: `p9_mk_engine.py`, `p9_certify_hp.py`, `p9_exact_cert_k*.json`, `p9_tuple_k15856.npy`, verification scripts, and the companion surveys (*The Walls and the Doors*; Codex handoff note) documenting the framework analysis.

References

J. Maynard, *Small gaps between primes*, Ann. of Math. 181 (2015) 383 – 413 · D.H.J. Polymath, *Variants of the Selberg sieve, and bounded intervals containing many primes*, Res. Math. Sci. 1:12 (2014) · J. Stadlmann, *On primes in arithmetic progressions to smooth moduli and bounded gaps between primes*, arXiv:2309.00425, Adv. Math. (2025) · E. Bombieri (1965), A.I. Vinogradov (1965) · I. Shevtsova (2010-type Berry – Esseen constants) · T. Engelsma, OEIS A008407 · A.V. Sutherland, narrow admissible tuples database.

Historical source 30: prime_gap_survey

Source: `historical/riemann-rmt/prime_gap_survey.md`. SHA-256:
`706f0de654e61de32016d45637bf887994393b6911647095378782325a6a7a66`.

The Walls and the Doors: Exact Theorems on the Bounded-Prime-Gaps Framework

A survey of five companion investigations, with a complete computational compendium

Bill (Qingyun) Sun · GPT5.6SOL · Fable

August 11, 2026

Abstract

Twelve years after Polymath8b established that infinitely many pairs of primes differ by at most $H_1 = 246$, that constant is unmoved. This survey synthesizes five companion investigations that, rather than optimizing once more inside the Maynard – Tao framework, determine its exact boundaries — and in doing so locate the only places a smaller constant can come from. The results are of two kinds. **Walls (proved)**: the admissible-tuple side of 246 has exactly zero slack (exhaustive re-proof that $H(50) = 246$ and $H(49) = 240$ are minimal); the information-theoretic ceiling of the framework's data class equals the tuple diameter exactly, so no post-processing of sieve output can beat $H(k_{\min})$; the scalar Maynard – Tao decode is precisely optimal against all available data, including every one-sided correlation bound — the matrix/inertia layer buys counting constants but never a smaller gap; the sieve-weight cone cannot be usefully enlarged — rank- r sum-of-squares weights decouple by a subadditivity theorem, and the strictly larger copositive cone is *flat* against the sieve's evaluation pencil (machine-verified dual certificate); and the 2024 Guth – Maynard large-value revolution is provably orthogonal to the level-of-distribution inputs the framework consumes. Along the way we prove a structural theorem of independent interest: **a family of pairwise prime-correlation conclusions can be simultaneously annihilated by a Selberg-type parity twist if and only if its kill-graph is bipartite** — the odd-cycle facets of the cut polytope are exactly what makes bounded gaps provable while twin primes are not. **Doors (quantified)**: the variational door at $k = 49$ (payoff 6) and especially $k = 47$ (payoff 20 — the payoff is convex, a fact previously unremarked); the conversion of Maynard's uniform-residue equidistribution trilogy into shell-support gain, needing only 8 – 13% of the full-shell variational surplus and undecided in print; and the $m = 2$ record — corrected here to the current $H_2 \leq 396,504$ (Stadlmann) — whose admissible tuple is "almost surely not minimal" and whose variational optimization is a decade staler than H_1 's. A final section compiles, in compact machine-usable form, every formula, constant, threshold, and construction needed to compute in this framework.

No result in this survey improves $H_1 = 246$. The claim is sharper: we prove where improvement is impossible, and price where it is possible.

Status convention. [P] proved (including computer-assisted exhaustive proofs with independent verification); [C] computational (reproducible, residuals stated); [H] heuristic; [Q] conjecture/open.

1. Introduction: one constant, one pipeline

Write $\text{DHL}(k, m+1)$ for the assertion that every admissible k -tuple $\mathcal{H}$ has infinitely many translates $n + \mathcal{H}$ containing at least $m+1$ primes. The modern pipeline, due to Goldston – Pintz – Yıldırım, Zhang, Maynard [May15], Tao, and Polymath8b [Pol14a], factors every bounded-gap theorem into three independent components:

$$(\text{arithmetic}) \theta \rightarrow (\text{variational}) M_k > 2m/\theta \Rightarrow \text{DHL}(k, m+1) \rightarrow (\text{combinatorial}) H_m \leq H(k),$$

where θ is a level of distribution of the primes (Bombieri – Vinogradov gives $\theta = 1/2$ unconditionally), M_k is the Maynard – Tao variational constant (§ 9.1), and $H(k)$ is the minimal diameter of an admissible k -tuple. The records: $H_1 \leq 246$ ($k = 50$, via the ϵ -refined variational threshold at $\theta = 1/2$) [Pol14a]; $H_1 \leq 12$ under Elliott – Halberstam; ≤ 6 under generalized EH — where the parity obstruction provably stops all sieve-type methods. For $m = 2$ the record is $H_2 \leq 396,504$ (Stadlmann [Sta23], superseding Polymath8b's 398,130 — a correction to several summary accounts, including our own earlier one); $H_3 \leq 24,797,814$; $H_m \ll \epsilon^{\{3.8075m\}}$ [Sta23].

The five companion investigations surveyed here each interrogate one component with the toolset developed in our finite random-matrix program (exact LP/SOS duality, convex-order counterfeits, positive-inertia decoding, exhaustive combinatorial search): Paper I the combinatorial floor, Paper II the information-theoretic ceiling, Paper III the decode layer, Paper IV the weight cone, Paper V the arithmetic frontier. Their common finding is that the framework is far more rigid than the folklore suggests — rigid enough that each "wall" is now a theorem with an explicit certificate, and the surviving "doors" carry exact price tags.

2. Paper I — The combinatorial component is closed (and the payoff is convex)

(Exhaustive admissible-tuple computation; scripts `p4_rho_exhaust.c` and companions.)

Theorem 1.1 [P]. $H(50) = 246$ and $H(49) = 240$ are minimal: no admissible 50-tuple of diameter ≤ 245 and no admissible 49-tuple of diameter ≤ 239 exists.

The proof is a complete enumeration (not a failed search): after two structural reductions — parity monochromaticity (any admissible tuple is monochromatic mod 2, so one works in half-coordinates) and complete branching over the avoided residue classes mod 3, 5, 7 — a DFS over residue-coverage bitmasks with ρ^* -table pruning exhausts the decision instances ($\approx 1.6 - 1.85$ million nodes each) in seconds. An independent implementation with a different search tree, and exact-arithmetic re-verification of every witness, confirm the values; they agree with Engelsma's classical tables (minimality known for all $k \leq 342$) and OEIS A008407. The significance is not novelty of the values but the *closure statement*: the tuple side of $H_1 = 246$ has exactly zero slack — any improvement must come from the analytic side, i.e. from lowering k below 50.

The convexity observation. The verified payoff ladder near the record is $H(49) = 240$, $H(48) = 236$, $H(47) = 226$, $H(46) = 216$: gains of 6, 4, 10, 10 per unit of k . (Several circulating summaries state $H(47) = 232$; this is wrong.) The strategic consequence, previously unremarked: the variational effort should not be spent exclusively on the marginal door $k = 49$ — the $k = 47$ door pays more than three times as much, and the threshold distance grows only linearly (§ 6). Conversely a regression to $k = 51$ costs only 6.

The $m = 2$ opportunity [C/H]. For $k = 35,265$ (Stadlmann's $\text{DHL}(35265, 3)$), the record tuple of diameter 396,504 is flagged by its own author as "almost surely not minimal." Rigorous lower bounds (Montgomery – Vaughan Brun – Titchmarsh: 216,632; Selberg large sieve: 202,249 — computed exactly) leave a factor-1.8 void, and the same bounds run $\approx 50\%$ low at $k = 50$ where the truth is known. A reproduction of the Sutherland pipeline (shifted Schinzel sieve $\rightarrow$ greedy $\rightarrow$ local

optimization) reaches within 3.0% of the record in ~ 30 CPU-minutes; the record's last percent lives in iterated merging and residue-class annealing. Estimated realistic slack at fixed k : $0.03 - 0.3\%$, i.e. $100 - 1,200$ units of H_2 — **each single unit an immediate new record on $\liminf(p_{n+2} - p_n)$, requiring zero new analytic input.** A full-scale attack (with an LP-based repair step unavailable in 2014) is in progress.

3. Paper II — The exact ceiling, and the bipartite structure of parity

(The counterfeit theory of the bounded-gaps data class; scripts `p8_moment_lp.py`, `p8_parity_fiber.py`.)

This is the round's most conceptually novel contribution: a two-layer adversary theory that computes, exactly, what any argument consuming the framework's data can and cannot conclude.

3.1 Layer 1: the ceiling equals the tuple diameter

Model the framework's knowledge about the prime-indicator vector $(X_1, \dots, X_k)$ at a translate of the tuple: **(R1)** the weighted marginals $E[X_i]$ (delivered two-sidedly by the sieve at level θ), and **(R2)** *upper* bounds on all joint moments $E[\prod_{i \in S} X_i]$, $|S| \geq 2$ (Selberg/large sieve; the best published pair constant is 3.3996 · Hardy – Littlewood [Wu04], improved to ≈ 3.30 by Lichtman [Lic25]). Crucially **(R3)**: no lower bound on any joint moment is available — that absence *is* the parity wall.

Theorem 2.1 (ceiling = diameter) [P in the model]. Over all joint laws consistent with (R1) – (R2): (i) $\min P(X_1 + \dots + X_k \geq m+1) = \max(0, (s-m)/(k-m))$ where $s = \sum E[X_i]$ — the pigeonhole bound is the *entire* provable content, and no strengthening of the upper-bound rows moves the threshold $s > m$; (ii) for any $H < \text{diam}(\mathcal{H})$, the minimum probability that some pair at distance $\leq H$ is simultaneously prime is exactly 0, jumping to $(s-1)/(k-1)$ at $H = \text{diam}(\mathcal{H})$. Hence the class ceiling is $H_{\text{ceiling}} = H(k_{\text{min}})$ exactly: **no post-processing of Maynard – Tao output — matrix, moment, spectral, or otherwise — can ever certify a gap smaller than the tuple diameter**, and improved pair-correlation upper bounds (Wu $\rightarrow$ Lichtman $\rightarrow$ oracle) can never lower H_1 , though they do raise the guaranteed *density* of prime-pair translates.

3.2 Layer 2: the parity fiber is a slice of the cut polytope

Encode a Selberg-type parity twist as a reweighting by $f(\lambda(n+h_1), \dots, \lambda(n+h_k)) \geq 0$ (λ = Liouville) constrained to preserve all one-marginal data; the reachable pair-correlation multipliers y_{ij} then form a polytope B_k .

Theorem 2.2 (bipartite kill-graph) [P]. A set K of pair-conclusions can be simultaneously annihilated ($y_{ij} = 0$ for all $ij \in K$) by a parity twist **iff the graph K is bipartite**. (Bipartite: weight the two sides $\pm$; odd cycle: an independent set meets C in $< |C|/2$ vertices, contradicting the forced marginals.) Verified exhaustively against all 63 kill-patterns at $k = 4$ and all 1023 at $k = 5$: LP-feasibility coincides with bipartiteness without exception.

The odd-cycle facets of the correlation/cut polytope are thus precisely the boundary between what sieves can and cannot see. With no tuning, the model reproduces every known frontier fact: Selberg's classical $k = 2$ counterexample (the unique kill-all solution, $f = 2 \cdot 1_{\{\lambda_1 \lambda_2 = -1\}}$); the factor-2 parity floor on twin-prime constants in both directions (the coordinate-range $[0, 2]$ of B_k ; the gap from Lichtman's 3.30 to the floor 2 is the framework's remaining "integrality gap"); the parity-blindness floor $H_1 \geq 6$ **for every parity-immune argument** (certifying "some gap $\leq H$ " needs a non-bipartite close-pair graph, i.e. a triangle, and the narrowest admissible triangle is $\{0, 2, 6\}$ — exactly matching "GEH gives 6, and 4, 2 are blocked"; and the m -floor: $\text{DHL}(k, m+1)$ is parity-blocked iff $k \leq 2m$. A conservation law at $k = 3$: killing two pair-conclusions forces the third to exactly twice Hardy – Littlewood — what parity takes at one gap it repays doubled at another.

3.3 The Siegel dichotomy, audited

Under strong Siegel zeros the Tao – Teräväinen order- ≤ 2 Hardy – Littlewood – Chowla rows [TT22] pin every pair coordinate $y_{ij} = 1$ — the parity fiber's pair projection collapses to a point, which is the mechanism behind Heath-Brown's "Siegel zeros $\Rightarrow$ twin primes" and Wright's $H_m \ll e^{\{1.9828m\}}$ [Wri21]. But the audit of the tempting two-branch strategy is negative [P in the model]: on the no-Siegel branch the data class is asymptotically identical to the unconditional one (Siegel-freeness enters Maynard – Tao only through effectivity, never through θ or M_k), so **min(branch A, branch B) equals the unconditional record — the dichotomy currently yields nothing new for H_1 or H_2** . Its real content is a target: any lemma of the form "no Siegel zeros $\Rightarrow$ BV-averaged level $1/2 + \delta'$ " would immediately combine with Wright's branch into an unconditional improvement.

4. Paper III — The decode layer is exactly optimal (a convex-order theorem)

(Positive-inertia decoding of the Maynard – Tao form; scripts `p5_functionals*.py`, `p5_decode_lp.py`.)

Motivated by the rank – trace ("Lemma R") decoding that powers the 2026 two-thirds theorem for zeta zeros, one asks: does the Maynard – Tao quadratic form contain more than the one scalar inequality traditionally extracted from it?

The availability matrix [P, bookkeeping]. For the sieve weight $w = (\sum \lambda_d)^2$ at level $R = N^\wedge r$, the computable entries are: everything linear in at most one prime indicator (two-sided, at BV level, requiring $2r < \theta$); everything built from majorants/divisor functions (two-sided, elementary); and all *quadratic-in-primes* entries only **one-sidedly** (upper bounds by majorant swap, with overshoot $\kappa = 16/(1-2r)^2 \geq 64$ at the threshold weight; lower bounds only trivial ≥ 0). The entire cross block is unknown in exactly the direction that matters.

Theorem 3.1 (decode optimality) [P at the decode layer]. Let $A = E_\mu[\sum 1_{P(n+h_i)}] = (\theta/2)M_k$ be the sieve's certified mean. Against the data class consisting of all exact linear-in-primes data together with valid upper bounds on $E_\varphi(X)$ for every convex φ — a class containing all falling moments and all r -fold correlation upper bounds however sharp — the largest certifiable r in " $\exists n: X \geq r$ " is exactly $\lfloor A \rfloor + 1$. *Proof:* the counterfeit law with X two-point supported on $\{\lfloor A \rfloor, \lceil A \rceil\}$ at mean A is the minimum in convex order among laws on $\mathbb{Z} \geq 0$ with that mean; hence it satisfies every constraint in the class and admits no n with $X \geq \lfloor A \rfloor + 2$. ■

Consequently $f(m) = 2m/\theta$ is **decode-optimal: no matrix, inertia, or moment refinement of the Maynard – Tao decode can lower any H_m within the available data**. The two natural loopholes are closed quantitatively: designing sign-cancelling profiles to suppress the pair functional fails (the pair-to-mean² ratio γ is pinned within 1.2% of 1 along the entire Pareto frontier at k up to $1.6 \cdot 10^6$ [C]); and multi-translate simultaneity reduces to the scalar bound on the union tuple by first-moment budget conservation [P].

What survives [C]. Above threshold ($A = m + \delta$) the sharp two-point-adversary bound

$$\Pr_\mu[X \geq m+1] \geq \delta / (\min(k, 2B/\delta - m + 1) - m), \quad B = \kappa \gamma A^2/2 - m(m-1)/2,$$

improves the certified *measure* of prime-rich translates by factors up to $23 \times$ ($m = 2$) and $505 \times$ ($m = 3$, large δ) once pair upper bounds are used — an exponential-in- m gain in the counting constant, honestly assessed as a sharp LP form of a technique already present in Banks – Freiberg – Maynard, and worth exactly zero for H_m itself. The parity audit confirms the decode never touches a specified-shift pair; the counterfeit takes pair mass *below* Hardy – Littlewood, which no available row forbids.

5. Paper IV — The weight cone cannot be enlarged (decoupling and flatness)

(SOS-certified enlargements; scripts `p7_*.py`, certificate `p7_dual_certificate.npz`.)

The Maynard – Tao weights are single squares. Two enlargements suggest themselves: sums of several squares (rank $r > 1$), and signed corrections certified pointwise-nonnegative on the divisor lattice (a copositive, rather than PSD, Gram matrix). Both are closed.

Theorem 4.1 (decoupling) [P]. Every DHL criterion in the pipeline has the form $\Phi(w) > 0$ with Φ subadditive on the weight cone (main terms and error budgets are linear in w ; Cauchy – Schwarz penalties $-\sqrt{B(w)C(w)}$ are superadditive). Hence a rank- r witness implies a rank-1 witness in the same per-square class: **rank never helps, for any legal variant** — the ϵ -trick and vanishing-marginal refinements enlarge the per-square class, not the useful rank. (Rank could only pay under an aggregate-cap constraint nonlinear in w ; no sieve quantity has that form, since everything ever bounded is $\sum_n a(n)w(n)$.) A companion collapse [P]: PSD certificates plus the level constraint force long-support rows to vanish (Schur), so SOS + level = the classical class exactly.

Theorem 4.2 (flatness of the copositive cone) [C, with a machine-verified closure certificate]. The copositive cone is strictly larger as a cone — the LP finds wildly non-PSD legal Gram matrices — but its extra directions are orthogonal to the sieve's evaluation pencil $(\tilde{J}, \tilde{I})$. Measured optimal-value gaps at $k = 3, 4, 5$ over spans up to $r = 8$: $\leq 2.6 \cdot 10^{-7}$ relative, in an exact fixed- R model cross-validated against 10^6 real integers (where the exact finite-cone optimum is attained by a PSD matrix). At $k = 3$ the closure is certified: nine explicit realizable divisor patterns P with weights $y_P > 0$ satisfy $\tilde{J} = t \cdot \tilde{I} - \sum y_P v_P v_P^T$ to $1.3 \cdot 10^{-14}$ with t exceeding the PSD optimum by $1.4 \cdot 10^{-6}$ — since $\langle v_P, C v_P \rangle \geq 0$ for every copositive C , this bounds the *entire* enlarged cone. The cross-only long-support experiment (the sharpest corner: components beyond the level admitted only in cross terms) recovers ≈ 0 of the 10.2% headroom that illegal full support would give. Mechanism [H]: truncated divisor patterns realize all sign-alternating finite differences of the profiles — a family dense enough to pin the pencil's negative eigenspace; nothing in it weakens as k grows.

Consequence. The weight-cone axis is closed. Combined with Paper III, all conceivable "optimize harder inside the same data" strategies are now excluded by theorem or certificate; movement requires evaluation-side inputs.

6. Paper V — The arithmetic frontier, 2014 → 2026

(Level-of-distribution audit and the exact payoff function; script `p6_payoff.py`.)

The MT-usable frontier has moved $0.5233 \rightarrow 0.5263$ in twelve years. The sieve consumes equidistribution with two properties: all moduli (its error moduli are generic squarefree products) and residue uniformity (its residues are the $k^{\omega(q)}$ roots of $\prod (n+h_i)$, varying with q). Everything beyond Bombieri – Vinogradov's $\theta = 1/2$ exists only in restricted forms, sorted by compatibility: Polymath8a's $1/2 + 7/300$ (densely divisible moduli, poly-root residues — usable on a truncated support shell); Maynard's trilogy culminating in 11/21 with **completely uniform residues** [MayIII] — the first such object since Zhang; Stadlmann's $1/2 + 1/40$ (smooth moduli) [Sta23]; Pascadi's minorant-level 10/19 [Pas25] — while the numerically larger levels (BFI 4/7, Maynard II 3/5, Pascadi 5/8) are fixed-residue or well-factorable-weight results that the MT error term, carrying absolute values over varying residues, provably cannot consume. The mechanism by which restricted inputs failed to move H_1 in 2014 is reconstructed exactly: the input enters *only through the shape of the accessible support polytope* Ω , the shell beyond the unit simplex must be truncated to coordinates the modulus class can guarantee, and the ϵ -trick already harvests shell value at zero arithmetic cost — a structured input must beat the ϵ -trick on

its own shell. Net 2014 gain at $m = 1$: ≈ 0 . Coverage density in modulus mass is the wrong parameter [P]: non-geometrically-aligned partial coverage is worth exactly nothing; the right parameter is the variational surplus ρ_{geom} of the accessible sub-shell.

The payoff function, calibrated [C]. With the master criterion $\text{DHL}(k, m+1) \Leftarrow M_k(\Omega) > 2m/\theta$ and $M_k(\lambda\Omega) = \lambda M_k(\Omega)$: tipping $k = 49$ (worth 6) needs full-BV $\delta \approx 0.0014 - 0.0023$, or $\rho_{\text{geom}} \approx 8 - 13\%$ of the full-shell surplus at Maynard-III strength; $k = 48$ (worth 10): $\delta \approx 0.003 - 0.005$; $k = 47$ (worth 20): $\delta \approx 0.005 - 0.008$. Sensitivity at the record: $dH_1 / d\delta \approx -2400$. The pure-simplex route is dead by the rigorous upper bound $M_k \leq (k/(k-1))\log k$: $M_{50} \leq 3.99186 < 4$ [P] — every road to < 246 passes through enlarged variants.

Guth – Maynard is orthogonal to H_1 [P at proof-structure level]. The classical BV proof consumes no zero-density input; full GRH itself gives only $\theta = 1/2$ pointwise — everything beyond comes from dispersion + Weil/Deligne + spectral Kloosterman machinery in which zero-density estimates appear nowhere (dependency audit of Maynard I – III, Stadlmann, Pascadi). The 2024 large-value theorem [GM24] therefore feeds short-interval results, not bounded gaps. The genuine open computation this paper isolates: **translate Maynard III's divisor-window conditions into shell-polytope constraints and compute $M_{49}(\mathbb{R} \cup \Omega^*)$ with the ε -trick layered on — an undecided, decidable question that no one has published, and exactly the kind of 10^{-2} -precision constrained Rayleigh-quotient problem our exact-rational SDP/Krylov machinery certifies.**

7. Synthesis: the no-go map and the two open doors

The five papers compose into a single statement:

Within the Maynard – Tao data class, $H_1 = H(k_{\text{min}})$ exactly, and every route to a smaller H_1 must lower k_{min} — by enlarging the variational value through new *evaluation-side* arithmetic (uniform-residue level beyond $1/2$ on an accessible shell), and in no other way. Post-processing (Paper II), decode refinement (Paper III), weight-cone enlargement (Paper IV), tuple search (Paper I), and zero-density technology (Paper V) are all excluded by theorem or certificate.

The doors, priced: (D1) the $k = 49/48/47$ variational doors, payoffs 6/10/20, with the Maynard-III shell computation as the concrete undecided question; (D2) the $m \geq 2$ records — softer in *both* factors: the variational optimization at $k \approx 35,000$ never received 2014's $k = 50$ -level effort (threshold $M_k > 8$; every unit of variational gain converts at $dH_2 / dk \approx 11$ via the tuple tables), and the record tuple itself has measurable slack (each unit a new world record); (D3) parity-sensitive rows short of Siegel zeros — order-2 Liouville correlation rows in natural density would pin pair coordinates exactly as Tao – Teräväinen's do conditionally; the bipartite theorem says precisely which conclusions each new row unlocks.

Work in progress (five further investigations, reports pending): the exact-rational variational baseline at $k = 45 - 55$; the ε -variant attack at $k = 49/47$; certified upper bounds closing or opening those doors; the large- k ($m = 2, 3$) variational engine; and the additive – Mellin Type-II program.

8. Computational compendium

Everything needed to compute in this framework, compact but complete. All notation as above; [P]/[C] status inherited from the sections cited.

8.1 The variational problem

- Simplex $R_k = \{t \in [0,\infty)^k : \sum t_i \leq 1\}$; symmetric F : $I(F) = \int \{R_k\} F^2, J^{\wedge\{m\}}(F) = \int \{R_{k-1}\} (\int F dt_m)^2$; $M_k = \sup_F k \cdot J(F)/I(F)$ (all $J^{\wedge\{m\}}$ equal by symmetry).
- Bilinear polarization (legal at no cost): $I(F_1, F_2) = \int F_1 F_2, J^{\wedge\{m\}}(F_1, F_2) = \int (\int F_1 dt_m) (\int F_2 dt_m)$.
- Polynomial bases in power sums $P_j = \sum t_i^j$ reduce I, J to 1-D Beta-type integrals; exact rational arithmetic feasible through $k \sim 10^2$ at useful degree. Product-profile 1-D reduction for $k \sim 10^4$ (optimal $F \approx \prod g(t_i), g \approx 1/(1+At)$ truncated).
- Threshold: $DHL(k, m+1) \Leftarrow M_k(\Omega_{\text{accessible}}) > 2m/\theta; \theta = 1/2 \text{ unconditional} \Rightarrow M_k > 4m$. Scaling law $M_k(\lambda\Omega) = \lambda \cdot M_k(\Omega)$.
- ϵ -trick: support $(1+\epsilon)R_k$ with marginals vanishing outside $(1-\epsilon)R_{k-1}$; zero arithmetic cost; this is how $k = 54 \rightarrow 50$.
- Upper bound [P]: $M_k \leq (k/(k-1)) \cdot \log k$. Values: $M_{48} \leq 3.95357, M_{49} \leq 3.97290, M_{50} \leq 3.99186$ (pure problem dead ≤ 50); first non-excluded $k = 51$ (4.01046).
- Anchors: $M_5 = 2.007080 > 2$ ($EH \Rightarrow H_1 \leq 12$); $M_4 = 1.845401 < 2$ (provably, so pure class cannot give $DHL(4,2)$ under EH); $M_3 = 1.646440$; $M_2 = 1.385933$; $M_{105} > 4$ (Maynard); $M_{54} > 4.00238$ (Polymath8b); ϵ -variant at $k = 50 > 4 \Rightarrow 246$.
- Asymptotics: $M_k = \log k - c_0 + o(1)$; calibrated slope $dM/dk \approx 1/k$ near $k = 50$.

8.2 Records and thresholds (August 2026)

quantity	value	source/condition
H_1	246 ($k = 50$)	unconditional [Pol14a]
H_1 under EH / GEH	12 / 6	6 is the parity floor [P, § 3.2]
H_2	396,504 ($k = 35,265$)	Stadlmann [Sta23] — not 398,130
H_3	24,797,814 ($k = 1,649,821$)	[Pol14a]
H_m	$\ll e^{\{3.8075m\}}$	[Sta23] — not 3.815
m-threshold	$M_k > 4m$ ($\theta = 1/2$); $m = 2 \Rightarrow > 8$	
pair-correlation upper constants	3.3996 [Wu04] $\rightarrow \approx 3.30$ [Lic25]; parity floor 2	cannot affect H_1 [P, § 3.1]

8.3 Admissible tuples (all proven minimal for $k \leq 342$; witnesses on file for $k \leq 62$)

$H(k), k = 3..62$: 6, 8, 12, 16, 20, 26, 30, 32, 36, 42, 48, 50, 56, 60, 66, 70, 76, 80, 84, 90, 94, 100, 110, 114, 120, 126, 130, 136, 140, 146, 152, 156, 158, 162, 168, 176, 182, 186, 188, 196, 200, 210, 212, 216, 226, 236, 240, 246, 252, 254, 264, 270, 272, 278, 282, 288, 300, 304, 310, 320. Key ladder: $k = 46/47/48/49/50 \rightarrow 216/226/236/240/246$ (payoffs 30/20/10/6 relative to 246; $H(47) = 226$, not 232). Admissibility test: for every prime $p \leq k$, some residue class mod p is missed ($p > k$ automatic). Structural facts for search: monochromatic mod 2; complete mod-3,5,7 class branching is exhaustive for $k \geq 7$; ρ^* -pruning by smaller widths is sound (subsets/translations of admissible are admissible). Proven-optimal witnesses (diam 246, $k=50$): 0 4 6 10 16 18 24 28 34 36 60 64 66 70 76 78 84 88 90 94 100 106 108 114 120 126 130 136 144 148 154 156 160 168 174 178 186 190 196 198 204 214 216 220 226 234 238 240 244 246. Rigorous lower bounds at large k (both run $\sim 50\%$ low): Brun – Titchmarsh $H(35265) \geq 216,632$; large sieve $\geq 202,249$. $m=2$ exchange rate: $\Delta H_2 / \Delta k \approx 11.2$ near $k \approx 35,300$.

8.4 Levels of distribution (what the MT sieve can and cannot eat)

input	θ	moduli	residues	MT-usable
Bombieri – Vinogradov	1/2	all	uniform (max_a)	yes — the workhorse
Polymath8a	$1/2 + 7/300$	densely divisible	poly roots	shell-truncated
Maynard III [MayIII]	11/21	factorable windows	uniform	shell-truncated; window widths = the open question
Stadlmann [Sta23]	$1/2 + 1/40$	smooth	MPZ-type	shell-truncated
Pascadi minorant [Pas25]	10/19	smooth	minorant of 1_P	best usable restricted level
BFI / Maynard II / Pascadi	4/7, 3/5, 5/8	well-factorable λ	fixed a	no (two structural reasons: residue-variation; absolute values kill factorable weights)

Tipping table (full-BV equivalent): $k = 49 \Leftarrow \delta \approx 0.0014 - 0.0023$; $k = 48 \Leftarrow 0.003 - 0.005$; $k = 47 \Leftarrow 0.005 - 0.008$; $dH_1 / d\delta \approx -2400$. Shell-coverage form at $\theta' = 11/21$: $k = 49$ needs $\rho_{\text{geom}} \approx 8 - 13\%$. GRH gives only $\theta = 1/2$; zero-density estimates (incl. Guth – Maynard) appear nowhere in the beyond-1/2 dependency chains.

8.5 The adversary/counterfeit toolkit

- **Ceiling LP (Layer 1):** data $\{E[X_i] = \mu_i \text{ (equality)}, E[\prod_{i \in S} X_i] \leq \beta_S \text{ } (|S| \geq 2)\}$; $\min P(\sum X \geq m+1) = \max(0, (s-m)/(k-m))$, dual $1 \{S \geq m+1\} \geq (S-m)/(k-m)$; specified-gap conclusions have LP value 0 below the diameter. Quantitative pair-density: with all-pairs cap β , $\min \#\{\text{translates with } \geq 2 \text{ primes}\} = (s-1)/(r^*-1)$, $r^* = \beta k(k-1)/(s-1)$.
- **Convex-order counterfeit (decode killer):** X two-point on $\{\lfloor A \rfloor, \lceil A \rceil\}$, mean A — minimal in convex order on $\mathbb{Z}_{\geq 0}$, hence satisfies every convex-functional upper bound; kills all decodes beyond $\lfloor A \rfloor + 1$. Feasibility margins of real certificates above the (unusable) firing boundary: $97 - 129 \times$ (level-constrained), $2.0 - 4.0 \times$ even at the parity floor.
- **Parity twist fiber (Layer 2):** distributions q on minus-sets $T \subseteq [k]$ with all vertex-marginals $1/2$; pair multipliers $y_{ij} = 4P(i, j \in T) \in [0, 2]$. **Kill-set feasible $\iff$ bipartite.** $B_3 = \text{conv}\{(2,0,0), (0,2,0), (0,0,2), (2,2,2)\}$; $\sum y_{ij} \geq 2(k-2)$ tight; $DHL(k, m+1)$ parity-blocked $\iff k \leq 2m$; parity floor for any "some gap $\leq H$ " certificate: a triangle in the close-pair graph $\Rightarrow H \geq 6$ via $\{0, 2, 6\}$. Siegel rows (order- ≤ 2 HLC [TT22]) pin all $y_{ij} = 1$; then 3-primes conclusions revive iff pair-sum $> s/2$.
- **Sieve availability matrix:** two-sided = linear-in-primes (needs $2r < \theta$) and majorant/divisor blocks ($2r + \sum 2s_j < 1$); one-sided-upper = all quadratic-in-primes blocks, overshoot $\kappa = 16/(1-2r)^2$ ($= 64$ at $r = 1/4$); pair-suppression impossible: $\gamma = J\text{-pair}/(A^2/2)$ pinned in $[0.988, 1.015]$ along the Pareto frontier ($\gamma = (k-1)/k$ exactly for products).
- **Weight-cone closure:** subadditivity $\Rightarrow$ rank-1 suffices (§ 5); copositive-cone flatness certificate: $\tilde{J} = t \cdot \tilde{I} - \sum y_P v_P v_P^T$ (residual $1.3 \cdot 10^{-14}$). To test any proposed enlargement: check whether it changes a linear functional of w or only the positivity certificate — only the former can matter.
- **Counting bound (the one thing the matrix layer buys):** $\Pr[X \geq m+1] \geq \delta/(\min(k, 2B/\delta - m + 1) - m)$, $B = \kappa \gamma A^2/2 - \text{binom}(m, 2)$; beats pigeonhole iff $\delta > \kappa \gamma m^2/k$; gains $23 \times$ ($m=2$), $505 \times$ ($m=3$); worth 0 for H_m .

8.6 Strategic invariants

1. All H_1 improvement factors through $k_{\min}$; $dH_1/dk \in \{6,4,10,10,\dots\}$ (convex) below 50.
2. All $k_{\min}$ improvement factors through $M(\mathbb{Q})$ on enlarged supports; all $\mathbb{Q}$ -enlargement factors through uniform-residue arithmetic beyond $\theta = 1/2$ (or unconditional shell access à la Maynard III).
3. $m \geq 2$ records are soft in both factors (variational at $k \sim 10^4$ under-optimized; tuples non-minimal) with conversion 11.2 : 1 (analytic : combinatorial).
4. Parity floors: $H_1 \geq 6$ for the entire method class; $k \geq 2m+1$ for DHL; pair constants $\geq 2 \cdot HL$ both directions.
5. The only presently-undecided decidable question worth a large compute budget: $M_4 \rightarrow (\mathbb{R} \cup \Omega_{\text{MaynardIII}})$ vs 4 — and its $k = 47$ sibling.

9. Open problems

1. **The shell computation** (§ 6): extract Maynard III's divisor-window parameters, build Ω^* , certify $M_4 \rightarrow (\mathbb{R} \cup \Omega^*)$ against 4 with exact-rational inner/outer bounds. Undecided in print; decidable.
2. **The H_2 tuple record**: any admissible 35,265-tuple of diameter $\leq 396,503$. (In progress.)
3. **The $m = 2$ variational re-optimization** with Stadlmann-level inputs and 2014-vintage care at k slightly below 35,265.
4. **The continuum flatness theorem** (§ 5): $k \cdot \tilde{J} - \rho^* \tilde{I} \in -\text{cone}\{v_{Pv_{PT}}\}$ for arbitrary smooth spans — closing the copositive cone at all k by proof rather than certificate.
5. **The missing dichotomy lemma**: "no Siegel zeros up to scale $\Rightarrow \theta = 1/2 + \delta'$ for MT-averaged moduli" — the one statement that would make the Siegel two-branch strategy yield unconditional records via Wright's branch.
6. **New parity-sensitive rows in natural density**: any unconditional lower bound on a single joint moment $E[X_{iX_j}]$ (however weak) breaks Layer 2 at that edge; the bipartite theorem then computes exactly which package of conclusions unlocks.

References

- [Pol14a] D.H.J. Polymath, *Variants of the Selberg sieve, and bounded intervals containing many primes*, Res. Math. Sci. 1 (2014); arXiv:1407.4897. And *The "bounded gaps between primes" Polymath project — a retrospective*, arXiv:1409.8361.
- [May15] J. Maynard, *Small gaps between primes*, Ann. of Math. 181 (2015), 383 – 413.
- [Zha14] Y. Zhang, *Bounded gaps between primes*, Ann. of Math. 179 (2014), 1121 – 1174.
- [MayIII] J. Maynard, *Primes in arithmetic progressions to large moduli I – III*, arXiv:2006.06572, 2006.07088, 2006.08250; Memoirs AMS (2025).
- [Sta23] J. Stadlmann, *On primes in arithmetic progressions to smooth moduli and bounded gaps between primes*, arXiv:2309.00425; Adv. Math. (2025).
- [Pas25] A. Pascadi, arXiv:2505.00653 (well-factorable level 5/8) and arXiv:2505.09629 (minorant level 10/19).
- [GM24] L. Guth, J. Maynard, *New large value estimates for Dirichlet polynomials*, arXiv:2405.20552.
- [Wu04] J. Wu, *Chen's double sieve, Goldbach's conjecture and the twin prime problem*, Acta Arith. (2004); arXiv:0705.1652.

[Lic25] J.D. Lichtman, *A modification of the linear sieve, and the kernel of its iteration*, Algebra & Number Theory (2025); arXiv:2109.02851.

[TT22] T. Tao, J. Teräväinen, *The Hardy – Littlewood – Chowla conjecture in the presence of a Siegel zero*, J. London Math. Soc. 106 (2022).

[Wri21] T. Wright, *Prime tuples and Siegel zeros*, arXiv:2111.14054.

[HB83] D.R. Heath-Brown, *Prime twins and Siegel zeros*, Proc. London Math. Soc. 47 (1983), 193 – 224.

[Eng] T. Engelsma, *k-tuple permissible patterns*, opertech.com/primes; OEIS A008407.

[Sut] A.V. Sutherland, *Narrow admissible tuples database*, math.mit.edu/~primegaps.

[GPY09] D. Goldston, J. Pintz, C. Yıldırım, *Primes in tuples I*, Ann. of Math. 170 (2009).

[BFM16] W. Banks, T. Freiberg, J. Maynard, *On limit points of the sequence of normalized prime gaps*, Proc. London Math. Soc. (2016).

[CI02] B. Conrey, H. Iwaniec, *Spacing of zeros of Hecke L-functions and the class number problem*, Acta Arith. 103 (2002).

Companion scripts and machine-verified certificates for every [C] claim are archived in the project scratchpad ([p4_](#), [p5_*](#), [p6_*](#), [p7_*](#), [p8_*](#)); exhaustive-search witnesses re-verified in exact arithmetic by independent implementations.*

Historical source 31: stopping_times_paper

Source: `historical/riemann-rmt/stopping_times_paper.md`. SHA-256:
`8f7715642343c32bb8d39d30cae1d6b56bb0a0ec98ab2b5f7b26c1f6960b418a`.

Stopping Times See What Moments Cannot

Prime gaps, signed sieves, and a dynamic separation of CUE from the alternative hypothesis

Bill (Qingyun) Sun · GPT5.6SOL · Fable

August 2026

Abstract

We report on a research programme carried out by a human mathematician working with several large language models in a tight adversarial loop. Three groups of results are described.

(1) **Prime gaps.** New unconditional bounds for gaps containing several primes: $\liminf(p_{n+2} - p_n) \leq 173,438$, $\liminf(p_{n+3} - p_n) \leq 13,859,802$, and $\liminf(p_{n+4} - p_n) \leq 1,120,662,828$, improving the previous records $396,504 / 24,797,814 / 1,431,556,072$ by factors $2.3 / 1.8 / 1.3$. The gain is purely variational: the records since 2014 had upgraded the arithmetic input while retaining a crude closed-form bound on the Maynard – Tao functional. We also give a no-go map for $H_1 = 246$ with five proved walls, and price the two doors that remain open.

(2) **Signed sieves.** A one-line pointwise identity shows that the signed (indefinite) enlargement of the Maynard – Tao variational problem is *empty*: with the decode debt charged at face value, every signed weight is dominated by its own positive part. This closes the variational side of a programme suggested by Zhang's Landau – Siegel work and by Iwaniec's well-factorable weights, and —more usefully—shows that everything of value in the signed route is arithmetic, not variational. We give the precise missing estimate (E_θ) for each level $\theta \in \{4/7, 7/12, 3/5, 5/8\}$ together with the conditional price list it would buy ($H_1 \leq 130, 114, 94, 80$).

(3) **A dynamic separation.** Tao's alternative hypothesis ensemble ACUE matches CUE on every low-order statistic that bandwidth-limited detectors can see. We first measure the obstruction: the exact fibre of measures matching all balanced moments of degree $\leq N$ has dimension $0, 0, 2, 10, 80, 403, 1804$ for $N = 3, \dots, 9$, contains an explicit $(N-3)$ -parameter family $q_{\text{ACUE}}(C) \cdot g(\sum c \bmod N)$ with $\hat{g}(\pm 1) = 0$, and — because the heat flow is diagonal in the coefficients — stays frozen to that algebra along the *entire* flow. We then show that a *stopping time* escapes it: the finite de Bruijn – Newman depth Λ , the first collision time of the zeros under backwards heat flow, separates CUE from ACUE at the level of universality class, $-\Lambda^{\text{CUE}} \asymp N^{-8/3}$ against $-\Lambda^{\text{ACUE}} \asymp N^{-2}$. The exponent gap is the extreme-value statistics of the spectrum in scalar form: CUE admits rare pairs at distance $N^{-4/3}$ while the lattice quantises every gap at π/N , so the alternative hypothesis satisfies its own RH-analogue *too robustly*. More generally, for the circular β -ensembles we find and confirm $-\Lambda \asymp N^{-2-2/(\beta+1)}$ at $\beta = 1, 2, 4$, with the rigid lattice as the $\beta = \infty$ endpoint. The methodological content is a slogan:

moment matching does not imply stopping-time matching.

1. What this paper is about

Three problems, three walls, and one way through.

The first wall is the bounded gaps problem. Since Zhang, Maynard and Polymath8, the record $\liminf(p_{n+1} - p_n) \leq 246$ has stood for twelve years. We did not move it, and we explain in § 3 why we now believe it is protected by a specific and provable structure rather than by a lack of effort. But the same analysis showed that the *higher* records — two, three, four primes in a bounded window — were protected by nothing at all, and those we did move.

The second wall is positivity. Maynard – Tao weights are squares. Everybody who has worked with them has wondered what happens if you drop the square; Zhang's Landau – Siegel programme is built around exactly this, and Iwaniec's linear-sieve weights $\lambda^\pm$ are signed and consume levels of distribution that squares cannot reach. We settle the variational half of that question completely, in the negative, by an identity that fits on one line.

The third wall is the alternative hypothesis. Tao's ACUE ensemble is a lattice-supported measure on unitary spectra that reproduces CUE's pair correlation and, as we and others have verified, an entire algebra of low-degree statistics. The programme of "find a moment ACUE cannot fake" has been running for years and keeps hitting the same obstruction: the fibre of measures matching a given set of moments is large — we compute it exactly, and it reaches dimension 403 by $N = 8$ — so adding one more moment enlarges the search rather than ending it. Our contribution is to stop adding moments. A *stopping time* is not a polynomial statistic, and the fibre's freedom does not protect it: what it protects is the algebra of observables, and the first-passage time of a flow is not in that algebra.

The paper is written to be read by mathematicians, not by machines; but the last section describes how the work was actually done, because the method is part of the result.

2. Prime gaps: what moved and why

2.1 The framework in one paragraph

Fix an admissible k -tuple $H = \{h_1, \dots, h_k\}$ (admissible: for each prime p , the h_i do not cover all residues mod p). Maynard's method attaches to each n a weight $w(n) = (\sum_d \lambda_d)^2$ built from divisor sums, and compares $S_1 = \sum w(n)$ with $S_2 = \sum w(n) \cdot \#\{i : n + h_i \text{ prime}\}$. If the primes have level of distribution θ and the variational constant

$$M_k = \sup_F k \cdot J(F)/I(F), \quad I(F) = \int \{R_k\} F^2, \quad J(F) = \int \{R_{k-1}\} \left(\int_0^1 \{1 - \sum F dt_k\}^2, \right.$$

taken over symmetric F on the simplex $R_k = \{t \in [0, \infty)^k : \sum t_i \leq 1\}$, exceeds $2m/\theta$, then every admissible k -tuple contains $m+1$ primes infinitely often, and hence $H_m := \liminf(p_{n+m} - p_n)$ is at most the diameter $H(k)$ of the narrowest admissible k -tuple. Bombieri – Vinogradov gives $\theta = 1/2$ unconditionally, so the criterion is $M_k > 4m$.

Two quantities therefore control everything: the analytic constant M_k , and the combinatorial constant $H(k)$. The record $H_1 = 246$ comes from $k = 50$ and a variant of M_k tuned by an ϵ -trick.

2.2 The observation

For $m = 1$, k is small (50) and M_k has been computed to death. For $m \geq 2$, k is large — tens of thousands to millions — and there M_k has never been computed at all. Every record since 2014 has used the same crude closed-form lower bound, essentially $M_k \geq \log k - C$ with an explicit but generous C , obtained by truncating a product test function hard at the simplex boundary. The improvements of 2023 – 2025 raised θ , not M_k .

The deficit in that bound is between 2.3 and 2.9 units of $\log k$, and about 1.1 of those units are recoverable by elementary means. Since $M_k \approx \log k$, recovering 1.1 units is a factor $e^{1.1} \approx 3$ in k , and $H(k) \approx k(\log k + 0.77)$ is nearly linear in k , so a factor 3 in k is close to a factor 3 in H_m .

2.3 The engine

Take $F(t) = \prod_i g(t_i) \cdot 1[\sum t_i \leq k]$ (working on the dilated simplex). Let X_i be i.i.d. with density g^2/c_2 where $c_2 = \int g^2$, write S_j for partial sums, and $G(u) = \int_0^u g$. Then exactly

$$I(F) = k^{-k} c_2^{-k} \cdot P(S_k \leq k), J(F) = k^{-(k+1)} c_2^{-k-1} \cdot E[G((k - S_{k-1})_+)^2].$$

The simplex truncation, which the classical treatment handles by throwing it away, is now a genuine probability, and the second factor is unwrapped by the **layer-cake identity**

$$E[G((k - S)_+)^2] = \int 2G(u) g(u) \cdot P(S_{k-1} < k - u) du,$$

turning it into an integral of true lower-tail probabilities. Each of those is bounded below rigorously (chord-majorised Chernoff, one-big-jump, Berry – Esseen with the safe non-i.i.d. constant 0.56); because the bounds use only monotonicity of the true tail, discretisation cannot invalidate them. Finally, replacing the hard truncation by **shaped subexponential tails** $g(t) = e^{-(t/T_1)^\chi}/(1 + At)$ on a long support recovers the rest.

Exact accounting of the truncation contributes +0.12 units, tail shaping +0.49, and the remainder comes from optimising the shape parameters — about 1.1 units in total, as predicted.

2.4 The results

	bound	k	previous
$H_2 = \liminf(p_{n+2} - p_n)$	173,438	15,856	396,504 (Stadlmann 2023/25)
H_3	13,859,802	923,601	24,797,814 (Polymath8b 2014)
H_4	1,120,662,828	56,000,000	1,431,556,072 (Polymath8b 2014)

The certificates $M_{15,856} \geq 8.013326752751$, $M_{923,601} \geq 12.006666706750$ and $M_{5.6 \cdot 10^7} \geq 16.065482942$ are verified in ball arithmetic with outward rounding at every step; three independent certification regimes (two Berry – Esseen constants $\times$ two tail routes) pass at each k . The tuples are explicit: for $k = 15,856$ an admissible tuple of diameter 173,438 verified by two independent implementations; for $k = 923,601$ a repaired Hensley – Richards tuple of diameter 13,859,802 (with a fully classical fallback of diameter 14,505,780 from the primes past k); for $k = 5.6 \cdot 10^7$ the primes-past- k tuple, its endpoints checked against published values of $\pi(x)$.

These are computer-assisted results produced by AI systems and have not been refereed. All certificates and scripts are published for replay. A further conditional improvement ($H_2 \leq 145,226$ via Polymath8a/Deligne-strength equidistribution) is deliberately **not** claimed: it depends on a cap normalisation we reconstructed but could not verify

verbatim.

2.5 Why 246 did not move: five walls

We spent a comparable effort on H_1 and failed, but the failure is informative. Five statements, each proved during this work, close five different escape routes.

1. **The ceiling is the tuple diameter.** No post-processing of Maynard – Tao output can produce a bound below $H(k_{\min})$. In particular, pair-correlation constants of the Wu/Lichtman type — which look like they should help — cannot lower H_1 at all.
2. **Scalar decoding is exactly optimal.** One might hope to replace " $S_2 - mS_1 > 0$ " by a matrix, inertia, or higher-moment decode. A two-point counterfeit construction, ordered by convex order, defeats every such decode; the scalar threshold $f(m) = 2m/\theta$ is final.
3. **The weight cone is closed.** Enlarging squares to arbitrary PSD quadratic forms Q gains nothing: rank- r sums of squares decouple by subadditivity, so the supremum is attained at rank one. The copositive relaxation is flat as well, certified by an explicit nine-pattern dual with residual $1.3 \cdot 10^{-14}$.
4. **Parity, made combinatorial.** A parity (Liouville) twist kills a set of pair-conclusions if and only if the associated *kill-graph* is bipartite; odd-cycle facets of the cut polytope are exactly the obstructions. This yields the floors $H_1 \geq 6$ and $k \geq 2m + 1$ for the whole method class.
5. **The usable arithmetic frontier is not the published frontier.** The post-2014 levels $\theta = 4/7$ (Bombieri – Friedlander – Iwaniec), $3/5$ (Maynard II) and $5/8$ (Pascadi) are *well-factorable* or *fixed-residue*. Maynard – Tao needs uniformity over a CRT-structured system of residues varying with the modulus, which none of them provides. The usable frontier is Maynard III's $11/21$ with a shell truncation, Stadlmann's $1/2 + 1/40$, and Pascadi's minorant $10/19$ — all shell-restricted.

What remains open is narrow and precisely priced. The $k = 49$ door needs M_{49} -variant > 4 ; the pure constant is $M_{49} \in [3.891258, 3.97290]$ (lower bound exact-rational), and the best certified ε -variant is $M_{49,1/35} \geq 3.930490592$ with float optimum 3.959325169 . Known upper-bound technology closes the door only for $\varepsilon \leq 0.00682$, so the door is open, exactly as Polymath left it. Crossing it would give $H_1 \leq 240$; $k = 47$ would give 226.

3. The signed sieve: a one-line no-go

3.1 The question

Since Maynard – Tao weights are squares, and since Chen's theorem and Iwaniec's $\lambda^\pm$ are both built from *signed* objects, it is natural to ask what a signed sieve would buy. Zhang's Landau – Siegel programme is organised around removing the square. The hope is a variational phase invisible to the positive family.

For signed w the decode fails: the positive excess can be manufactured by $w(n) < 0$ at prime-poor n . It is repaired by paying the **debt**

$$D(w) = \sum_{-} \{n : w(n) < 0\} |w(n)| \cdot (m - v(n))_+, \quad v(n) = \#\{i : n + h_i \text{ prime}\},$$

after which $S_2 - mS_1 - D > 0$ does imply that some n has $v(n) \geq m + 1$. So the honest experiment is to optimise $\Phi(w) - \beta D(w)$ with $\beta = 1$ the true price, and look for a phase transition.

3.2 The identity

For every integer $v \geq 0$ and every $m \geq 1$,

$$(v - m) + (m - v)_+ = (v - m)_+,$$

both sides being $v - m$ when $v \geq m$ and 0 otherwise. Writing $w = w_+ - w_-$ with $w_+, w_- \geq 0$ of disjoint support and substituting:

Theorem (Signed No-Gain). $S_2 - mS_1 - D(w) = \sum_n w_+(n)(v(n) - m) - \sum_n w_-(n)(v(n) - m)_+ \leq \sum_n w_+(n)(v(n) - m)$.

Since $w_- \geq 0$ and $(v - m)_+ \geq 0$, the negative part is pure loss, and the loss is exactly the overshoot mass $\sum w_-(v - m)_+$ sitting on the negative support.

Corollary. Any DHL($k, m+1$) conclusion obtainable from a signed weight w with the debt charged at face value is already obtainable from the nonnegative weight w_+ . The signed class is redundant.

The statement is pointwise, so it holds for every k, m , tuple and weight class at once, needs no asymptotics, and subsumes the weaker fact that the naked ratio $\sup \text{Tr}(BQ)/\text{Tr}(AQ)$ over indefinite Q equals the PSD value by rank-one extremality.

3.3 What the numerics had been showing

We found the identity only after a substantial LP campaign on a finite arithmetic microcosm (n uniform on $\mathbb{Z}_W$, v the coprimality count on the tuple, weights in the span of level- L features, everything in exact rational arithmetic). Three phenomena had looked like a phase structure; all three are corollaries.

- **A critical price.** The value of $\max\{\Phi - \beta D : S_1 = 1\}$ is *exactly* the classical positive optimum for β above a sharp β^* , and acquires a signed optimiser below it. In the base model ($k = 3$, features $\{2, 3, 5, 7\}$) an exact-rational simplex with certified dual gives the classical value $1087376209/3212440751 = 0.3384891094603102$ and the critical price

$$\beta^* = 23051796480/10991046857 = 2.0973249209031\dots,$$

the signed vertex having $\Phi = 1.2082816957\dots$, $D = 0.4147152297\dots$, with 85% of its mass negative on 16 of 96 cells. The transition is certified on both sides: at $\beta^* - 10^{-3}$ the optimum is strictly above the plateau ($0.3389038246899884 > 0.3384891094603102$) and genuinely signed, while at $\beta^* + 10^{-3}$ it equals the plateau as exact rationals. At $\beta = 2$ the programme is unbounded. The theorem explains why $\beta^* > 1$ always: at the true price the signed class is dominated.

- **An apparent linear gain.** Imposing $\|w\|_1 \leq A$ gives $\lambda(1) = \lambda_{\text{positive}}$ exactly (for $w \geq 0$ the normalisation $S_1 = 1$ is the ℓ^1 norm) and then a strictly positive slope 0.32 – 0.89 per unit of budget. This is not a gain: the decode is scale-invariant, and what grows with A is the mass of w_+ , not the truth of the conclusion.
- **Unboundedness at $\beta = 1$.** In all eight model variants. Same cause. This exposes a second, usually unremarked, job that positivity performs in Maynard – Tao: *it makes the variational problem bounded*, by tying the ℓ^1 norm to the normalisation. Remove it and boundedness must be re-imposed by hand — and by the theorem, no choice of norm bound creates a gain.

3.4 Where the door actually is

The theorem has exactly two hypotheses, hence exactly two escapes, and neither is variational.

(i) **Charge the debt below face value.** This needs arithmetic information bounding $(m - v)_+$ on a *designed* negative support by strictly less than the truth — that is, an input asserting that primes are anomalously common on a prescribed set. This is precisely the exceptional-character mechanism. By the theorem, Zhang's programme is not one route among several; it is the only route that changes the variational picture at all.

(ii) **Keep w evaluable while w_+ is not.** The theorem compares w with w_+ , presuming both usable. Arithmetically they are not interchangeable: the positive part of a divisor-sum quadratic is not a divisor-sum quadratic, so a signed well-factorable λ can be evaluable at $\theta = 4/7, 3/5, 5/8$ while its positive part is evaluable at no level beyond $1/2$. **The signed route's entire value is which weights the arithmetic can evaluate, never the shape of the optimum.**

An independent argument from the switching side reaches the same conclusion. Switching is a re-indexing bijection: it lowers the anatomical depth of one linear form but never the number r of exact-primality conditions in a debt term. Chen's debt has $r = 1$ and is payable by an upper sieve (parity-free). Every DHL($k, m+1$) conclusion with $m \geq 1$ forces residual debt with $r \geq 2$; a two-vertex kill-graph is an edge, hence bipartite, hence parity-blocked for every input class admitting Liouville twists. The escape is not a cleverer switch but an input class in which the twist is inadmissible — the exceptional-character regime again. Two disjoint analyses, one door.

3.5 The price list

If the debt is paid at level θ , the consequences are fully computed; all eight $m = 1$ crossings carry exact-rational certificates, and $H(k)$ is Engelsma-exact.

θ	$2/\theta$	k pure (certified M_k)	k with ε	$H_1 \leq \text{pure} / \varepsilon$
$1/2$ (Bombieri – Vinogradov)	4	54	50	$270 / 246$ (unconditional)
$4/7$ (BFI)	3.5	31 (3.502015...)	29 (3.519881...)	$140 / 130$
$7/12$ (Maynard II, linear sieve)	$24/7$	29 (3.443305...)	26 (3.433616...)	$130 / 114$
$3/5$ (Maynard II)	$10/3$	26 (3.350647...)	23 (3.334616...)	$114 / 94$
$5/8$ (Pascadi)	3.2	22 (3.207656...)	20 (3.222666...)	$90 / 80$
1 (Elliott – Halberstam)	2	5 ($M_5 = 2.007080$)	—	12

The missing statement, in the form we would want to see proved:

(E $_{\theta}$). Fix an admissible tuple H , an index i , and $A, \varepsilon > 0$. For coefficient systems $c_q(a)$ jointly well-factorable with the residue selection — for every factorisation $q = q_1 q_2$ (resp. $q_1 q_2 q_3$) with $\prod Q_j = x^{\theta - \varepsilon}$ one can write $c_q(a) = \prod_j \gamma_j(q_j, a \bmod q_j)$ with $|\gamma_j| \leq 1$ and $a \bmod p \in \{h_i - h_j \bmod p : j \neq i\}$ for all $p \mid q$ — one has $\sum_q \{c_q(a) \cdot E(x; q, a) \ll_{H, A, \varepsilon} x (\log x)^{-A}\}$.

($E_{\{4/7\}}$) is "BFI Theorem 10, uniform over the CRT residue system of a fixed tuple polynomial". Polymath8a's $MPZ[\varpi, \delta]$ is its absolute-value cousin, proved by the same dispersion-plus-Deligne technology but only to level ≈ 0.5286 and for densely divisible moduli, so (E_θ) interpolates two proved endpoints rather than crossing a parity barrier.

4. Stopping times: the finite de Bruijn–Newman depth

4.1 Why the alternative hypothesis is worth attacking

Montgomery's pair correlation conjecture says the normalised gaps between zeros of ζ follow the GUE law. Half a century later it remains open, and the reason is sharper than "it is hard". The **alternative hypothesis** (AH) is the scenario in which the normalised gaps all lie asymptotically in $(1/2)\mathbb{Z}$ — the zeros sitting on a half-integer lattice rather than fluctuating like a random matrix spectrum. AH is not idle: it is consistent with everything currently provable about zeros, and it has real arithmetic consequences (it would force strong bounds on class numbers, and would be incompatible with certain conjectures on Landau – Siegel zeros). It is precisely the scenario that the known techniques cannot exclude.

Tao's blog article *The alternative hypothesis for unitary matrices* makes the obstruction concrete by moving it into random-matrix theory, where it can be computed with. Define the ACUE — the alternative circular unitary ensemble — as the measure on N -point configurations $C \subset \{2N\text{-th roots of unity}\}$, $|C| = N$, with

$$\mu_{\text{ACUE}}(C) = |\Delta(\zeta^C)|^2 / (2N)^N, \Delta = \text{Vandermonde}, \zeta = e^{2\pi i/2N}.$$

This is a lattice-supported measure whose eigenvalue gaps are all multiples of half the mean spacing — a perfect finite model of AH. Tao's point is that ACUE reproduces the CUE two-point correlation exactly, so any argument that only sees pair statistics cannot distinguish them, and he asks what does. The natural programme — find the first moment or correlation at which they differ — has been running ever since, and it keeps colliding with the same difficulty: the set of measures matching a given list of moments is a large convex body, so ruling out ACUE itself never rules out its neighbours. One must rule out a *fibre*, not a point.

There is a second thread. The Riemann ξ function has a canonical heat deformation ξ_t , and de Bruijn and Newman showed there is a constant Λ_{dBN} such that ξ_t has only real zeros exactly when $t \geq \Lambda_{\text{dBN}}$. The Riemann hypothesis is $\Lambda_{\text{dBN}} \leq 0$; Rodgers and Tao proved $\Lambda_{\text{dBN}} \geq 0$. So RH, if true, is *barely* true: ξ sits exactly on the boundary of heat stability, which is Newman's dictum that the Riemann hypothesis, if true, is only just true. Crucially, Rodgers and Tao's proof is a statement about *dynamics*: assuming $\Lambda_{\text{dBN}} < 0$ forces the zeros, run backwards, into an increasingly rigid local equilibrium — a clock — which contradicts what is known about their fluctuations.

This paper joins the two threads. The second suggests the observable that resolves the first.

4.2 How large the blind spot really is: explicit mimicker families

Before looking for a distinguishing statistic it is worth measuring the enemy. We constructed, for each N , the exact convex body of probability measures on the ACUE support matching every balanced moment $E[p_\lambda \bar{p}_\nu]$ of degree $\leq N$ — the *mimicker fibre*. Everything here is exact (rational Vandermonde masses, rank computed with a spectral gap of 10^6 or better).

The fibre. Its affine dimension at the level of rotation orbits is

N	3	4	5	6	7	8	9
dim of mimicker fibre	0	0	2	10	80	403	1804

So for $N \leq 4$ the moment data pins the measure — ACUE is rigid and any statistic distinguishes it. From $N = 5$ the fibre opens, and it grows explosively. At $N = 8$ there are 403 independent directions in which one may deform ACUE without changing a single balanced moment of degree $\leq N$.

An explicit family for every N . Inside the fibre there is a clean closed-form subfamily. Let $X(C) = \sum_{c \in C} c \bmod N$ be the centre-of-mass class and set

$$q_g(C) = \mu_{\text{ACUE}}(C) \cdot g(X(C)), g: \mathbb{Z}/N \rightarrow \mathbb{R}_{\geq 0}.$$

Then q_g matches all balanced moments of degree $\leq N$ iff $E[g] = 1$ and the Fourier coefficients $\hat{g}(\pm 1)$ vanish; all other frequencies are free. Hence a family of affine dimension $N - 3$, nonempty for every $N \geq 4$ and verified directly at $N = 5, 6, 7, 8$ (null dimensions 2, 3, 4, 5; worst balanced-moment error 10^{-12} ; positivity radius in random null directions up to 1.0 – 2.2, so genuinely positive measures, not signed deformations).

N	family dim	forced-zero frequencies	free frequencies
5	2	$j = \pm 1$	$j = 2$
6	3	$j = \pm 1$	$j = 2, 3$
7	4	$j = \pm 1$	$j = 2, 3$
8	5	$j = \pm 1$	$j = 2, 3, 4$

A second, representation-theoretic family. Writing $D_r(C) = \prod_{c \in C} \zeta^{\{rc\}} = \det(U_C)^r$ for the determinant character, the tilts

$$q(C) = \mu_{\text{ACUE}}(C) \cdot [1 + c \cdot \text{Re}(\eta D_r(C))], |\eta| = 1, 0 \leq c \leq 1,$$

also lie in the fibre. These have an exact quantum-mechanical description: μ_{ACUE} is the Born distribution $|\langle C | (\bigwedge^N F) A_0 \rangle|^2$ of a single Slater determinant (F the DFT matrix, A_0 the Fermi sea $e_0 \wedge \cdots \wedge e_{N-1}$) — a re-proof that ACUE is a projection determinantal process — and the tilted family is the interference pattern of *two* shifted Fermi seas, a coherent superposition $(|\Omega_0\rangle + a|\Omega_r\rangle)/\sqrt{1+a^2}$ with $c = 2a/(1+a^2)$. In Plücker coordinates ACUE is a point of the Grassmannian $\text{Gr}(N, 2N)$ and the mimickers are points of its first secant variety $\sigma_2(\text{Gr}(N, 2N)) \setminus \text{Gr}(N, 2N)$: the minimal non-Gaussian deformations of a fermionic Gaussian state.

How invisible are these directions? At $N = 8$, of the 403 fibre directions, 401 are invisible to *every* pattern count of window width ≤ 2 , 397 to width ≤ 3 , 396 to width ≤ 4 , 392 to width ≤ 5 , and 383 to width ≤ 6 . The blind spot is not a thin exceptional set; it is almost the whole fibre.

And the flow does not help. Writing $P(z) = \det(I - zU) = \sum_j a_j z^j$, the finite analogue of the de Bruijn – Newman deformation is

$$P_r(z) = \sum_j a_j r^{j(N-j)} z^j, t = \log r,$$

which is *diagonal in the coefficients*. Consequently every balanced moment of degree $\leq N$ is a t -dependent linear combination of frozen quantities, so it stays frozen along the **entire** trajectory — we verified this to $4.5 \cdot 10^{-16}$ on explicit mimickers. Evolving the moments you already have is provably futile, and so is adding one more.

This is the situation the rest of the section is about: a large, explicitly parameterised family of impostors, invisible to a whole algebra of observables, and invisible to that algebra for all time.

4.3 The observable

Define the **finite depth**

$-\Lambda(U) = \tau^*$, the first time (going backwards, $t < 0$) at which the discriminant of P_t vanishes,

i.e. the first collision of two zeros. Because the flowed polynomial stays self-inversive, a simple zero cannot leave the unit circle without first colliding, so this is well defined. Equivalently the zeros move by the attracting circular Coulomb dynamics

$$\dot{\theta}_j = -\sum_{k \neq j} \cot((\theta_j - \theta_k)/2),$$

derived from $\partial_t P = (ND - D^2)P$ with $D = z d/dz$. The right way to think of $-\Lambda$ is geometric:

$-\Lambda(P)$ is the distance from P to the discriminant hypersurface along the canonical heat ray.

Static moments are coordinates *along* the mimicker fibre; Λ is a coordinate *transverse* to it, pointing at the boundary. It is not a polynomial statistic of any degree — it is a first-passage time — and that is exactly why the fibre's freedom does not protect it.

4.4 The separation

Two independent computations, one exact and one Monte Carlo, settle the scaling.

ACUE (exact). Complete enumeration of all rotation orbits for $N = 3, \dots, 10$ — 13,132 orbits, 184,756 configurations at $N = 10$ — with exact Vandermonde masses, validated to 40 digits.

- $P(\text{clock}) = 2^{\{1-N\}}$ **exactly**. By Cauchy – Binet, $\sum_{|C|=N} |\Delta(C)|^2 = \det(AA^*) = (2N)^N$ for the DFT matrix A , while each of the two clock configurations has $|\Delta|^2 = N^N$. Clock polynomials are $1 - cz^N$, hence flow-invariant, hence have $\Lambda = -\infty$: they are the exact stationary atoms.
- For every non-clock configuration the minimal gap is exactly π/N (pigeonhole: all gaps $\geq 2\pi/N$ forces the clock).
- $-\Lambda^{\text{ACUE}} \asymp N^{\{-2\}}$, with fitted exponent -2.0009 ; $N^2(-\Lambda)$ is supported in $\approx [1.31, 1.99]$. The alternative $N^{\{-8/3\}}$ is decisively excluded.
- In all 13,130 non-clock orbits the first collision occurs at a pair that was already adjacent at $t = 0$ — zero exceptions.

CUE (Monte Carlo, N up to 256). $-\Lambda^{\text{CUE}} \asymp N^{\{-8/3\}}$, fitted exponent -2.678 ± 0.016 . More than the exponent, the *law* is parameter-free:

$$8N^{\{8/3\}}(-\Lambda) \Rightarrow G^2, \text{ where } P(G > x) = \exp(-x^3/72\pi)$$

is the sine-kernel smallest-gap law. The constant 72π was derived from the kernel, not fitted; measured 229 – 236 against 226.2, with KS distances 0.035 – 0.041 and median 29.6 – 30.5 against the prediction $(72\pi \ln 2)^{\{2/3\}} = 29.08$.

The ratio of depths therefore grows like $N^{\{2/3\}}$ — the two hypotheses do not merely differ, they sit in **different universality classes of subcriticality**. The $10\times$ separation point is $N^* \approx 165$, matching an a-priori estimate of ≈ 160 .

The interpretation is the striking part. ACUE's defect is not that it is fragile in some direction; it is that it is **too robust**. CUE is barely stable under reverse heat flow, in exactly Newman's sense; ACUE is over-equilibrated, sitting far from the discriminant because its lattice rigidity forbids the rare very close pairs ($\delta_{\min} \asymp N^{-4/3}$) that CUE produces. A fake RH universe survives too long under a natural deformation.

4.5 What the exponents mean

The numbers $8/3$ and 2 are not two readings on a dial; they encode a structural difference, and it is worth unpacking exactly what.

Λ measures a distance to a hypersurface. Let R_N be the set of degree- N polynomials with all roots on the circle, and $D_N = \{\text{Disc} = 0\}$ the discriminant hypersurface, which is exactly the boundary of R_N : a configuration leaves the "all roots on the circle" locus precisely by two roots first colliding. The heat flow gives a canonical vector field on coefficient space. Then $-\Lambda(P)$ = how far one must travel from P , against the heat field, to reach D_N .

So Λ is a *margin*: the distance from a configuration to the failure of its own Riemann-hypothesis analogue. This is the exact finite-dimensional shadow of the de Bruijn – Newman picture, in which $\Lambda_{\text{dBN}} \leq 0$ is RH and $\Lambda_{\text{dBN}} = 0$ says ξ sits on the boundary.

Why the exponent is a fingerprint of level repulsion. For an isolated close pair at angular gap δ , the two-body reduction of the Coulomb dynamics gives collision time $\delta^2/8 + o(\delta^2)$ — a purely local law we confirmed to 10^{-7} in the continuum regime, and exactly at $N = 2$ where $-\Lambda = -\log \cos(\delta/2)$. Hence the depth is governed by the *smallest gap in the configuration*, and the smallest gap is governed by the level repulsion exponent:

$$p(s) \sim c s^\beta \implies \min \text{ of } \approx N \text{ gaps} \asymp N^{-1/(\beta+1)} \text{ (normalised)} \asymp N^{-1-1/(\beta+1)} \text{ (angular)} \implies -\Lambda \asymp N^{-2-2/(\beta+1)}.$$

Reading the chain for the two cases in question:

- CUE ($\beta = 2$). Quadratic repulsion still permits rare, anomalously close pairs: the smallest of N gaps is of order $N^{-4/3}$, a full factor $N^{-1/3}$ below the mean spacing N^{-1} . One such accident is enough, and it produces depth $N^{-8/3}$.
- ACUE ($\beta = \infty$). The lattice forbids accidents. Every gap is a multiple of π/N , so the smallest gap is *exactly* π/N — deterministically, with no fluctuation at all (we prove this: any configuration all of whose gaps exceed $2\pi/N$ must be a clock). Depth N^{-2} .

So the ratio $-\Lambda^{\text{ACUE}} / -\Lambda^{\text{CUE}} \asymp N^{2/3}$ is precisely the ratio between "the closest pair among N random ones" and "the closest pair when closeness is quantised". The exponent gap *is* the extreme-value statistics of the spectrum, converted into a single scalar.

The interpretive punchline. One might have expected a fake RH universe to be caught out by being fragile somewhere. The opposite happens. ACUE's defect is that it is **too stable**: it sits $N^{2/3}$ times farther from the discriminant than a genuine random-matrix spectrum does. A true GUE/CUE world is real-rooted but *microscopically fragile* — it is always within $N^{-8/3}$ of losing the property, because it always contains one near-collision. ACUE is over-equilibrated: it satisfies its RH-analogue far too robustly, because its rigidity rules out the near-collisions.

This is the finite- N precise form of Newman's dictum. "RH, if true, is only just true" is not a piece of rhetoric; it is a statement about a scaling exponent, and the alternative hypothesis fails it by getting the exponent wrong in the direction of excessive safety.

It also reframes what one must prove about ζ . The static programme asks for a correlation statistic at frequencies beyond the reach of current technique. The dynamic programme asks instead: *is the zeta zero configuration, locally, as fragile as a random matrix?* Those are different questions, and § 4.10 argues the second may be the easier one — because refuting ACUE needs only that the depth be $o(N^{\{-2\}})$, with any rate whatsoever, not the full $N^{\{-8/3\}}$ law.

4.6 A universality law

The mechanism suggests, and the data confirm, a general law. If the normalised nearest-neighbour gap density behaves like $p(s) \sim c s^\beta$ as $s \rightarrow 0$, then the smallest of $\approx N$ gaps is of order $N^{\{-1/(\beta+1)\}}$, the physical gap is $N^{\{-1-1/(\beta+1)\}}$, and since an isolated pair collides in time $\delta^2/8 + o(\delta^2)$,

Dynamic Newman universality. $-\Lambda_N \asymp N^{\{-2-2/(\beta+1)\}}$.

ensemble	β	predicted	measured
COE	1	-3	-3.064 (local slopes -3.03 ... -3.10)
CUE	2	$-8/3 = -2.667$	-2.678 ± 0.016 ($N \leq 256$)
CSE	4	$-12/5 = -2.4$	-2.510 ($N \leq 64$, still drifting)
ACUE lattice	∞	-2	-2.0009 (exact)

All measured slopes are slightly steeper than predicted, in the direction and of the size of the finite- N drift independently calibrated at $\beta = 2$, where the larger N range shows the local slope converging to the prediction from below. The circular β -ensembles were sampled by the Killip – Nenciu CMV construction, validated against Haar CUE (KS 0.005) and against direct COE = VVT (KS 0.004). The localisation assumption behind the law — that the first collision is governed by the minimal initial gap — holds in 95 – 99% of samples, and in 100% of ACUE configurations.

So the depth is a scalar fingerprint of the microscopic repulsion exponent. That is a clean random-matrix statement independent of anything about zeta.

4.7 A configuration constant, computed two ways

Inside the ACUE law there is a distinguished stratum. Take the alternating clock (the zeros of $z^N + 1$), delete the point $e^{\{-i\pi/N\}}$, and insert 1; the gap pattern becomes 1, 2, 2, ..., 2, 3 in half-lattice units and the polynomial is exactly

$$P_N(z) = (z - 1)(z^N + 1)/(z - e^{\{-i\pi/N\}}).$$

Two completely independent routes give the same asymptotic depth.

Lattice route. Our exact solver on this configuration, $N = 3 \cdots 20$, gives $N^2(-\Lambda)$ increasing monotonically 1.41473945, 1.41821551, ..., 1.41963827 at $N = 20$.

Continuum route. Setting $t = -s/N^2$, $z = e^{\{iu/N\}}$ and passing to the limit yields the local function

$$G_s(u) = 2 \cos(u/2) - 2\pi \int_0^{\sqrt{1/2}} e^{\{s(1/4 - y^2)\}} \cos((\pi + u)y) dy,$$

whose first double zero ($G = \partial_u G = 0$) is at

$s^* = 1.419640342\dots$, $u^* = 1.812942145\dots$, with $G''(u^*) = -0.3767 \neq 0$ (a generic tangency).

The two agree to $2 \cdot 10^{-6}$ at $N = 20$, consistent with an $O(N^{-2})$ approach. This is Proposition-grade: a lattice enumeration and a transcendental double-root equation meeting to six digits.

One correction to the record. It is tempting to read s^* as the limit of the ensemble median, because for $N \leq 7$ the two coincide: the median of $N^2(-\Lambda)$ rises 1.41474, 1.41908, 1.41950. The complete enumeration shows the median **turns around** at $N = 7$ and falls — 1.41822 ($N = 8$), 1.41520 ($N = 9$), 1.41277 ($N = 10$) — as the support spreads from $[1.4147, 1.8246]$ to $[1.3146, 1.9855]$. At small N the single-dislocation orbit dominates the ACUE measure; from $N = 8$ the ensemble outgrows it. So s^* is a *configuration* constant belonging to the support of the limit law, not the median's limit. The neighbouring constant for a well-separated defect pair is $N^2(-\Lambda) \rightarrow 1.46946$ (i.e. $\rho_\infty = 8N^2(-\Lambda)/\pi^2 \rightarrow 1.19120$), the other end of the same defect family.

4.8 The depth escapes the freezing theorem

The point of the whole exercise. Return to the mimicker families of § 4.2 — the 403-dimensional fibre at $N = 8$, the centre-of-mass family $q_{\text{ACUE}} \cdot g(X)$, the secant tilts by $\det(U)^r$ — every one of which is frozen to the moment algebra for all time. Λ separates them at $O(1)$:

- an $N = 5$ mimicker in the $\mathbb{Q}(\sqrt{5})$ family moves the clock atom from 0.0625 to 0.1398 and shifts $E[N^2(-\Lambda) \mid \text{non-clock}]$ by -0.093 ;
- at $N = 8$, mimicker families move the entire Λ -law by total variation $0.12 - 0.24$, with the dependence concentrated on the centre-of-mass class $X = 0$;
- linear-programming tomography over the exact fibre at $N = 6$ (affine dimension 10 at orbit level) gives $E[N^2(-\Lambda) \mid \text{non-clock}]$ ranging over $[1.3610, 1.4770]$ against ACUE's 1.4336, while the clock atom — the mass at $\Lambda = -\infty$ — can be pushed from 0 to 0.0975, i.e. tripled from ACUE's 0.03125, with every constrained moment held fixed.

The conclusion deserves stating twice.

Moment matching does not imply stopping-time matching. Two ensembles can be indistinguishable to a whole algebra of observables, and remain indistinguishable along an entire natural flow, yet have first-passage laws in different scaling universality classes.

4.9 A classification of counterexamples, and two computable criteria

Not every impostor is caught by the same instrument. The counterexample families that arose in this project fall into three classes, and — this is the useful part — membership is decidable by two computations rather than by taste.

Criterion I (transversality). Fix a configuration X and a bandwidth r ; let $M_r(X) = (p_1, \dots, p_r)$ be the moment map and $\tau = -\Lambda$ the depth. The class is caught by ordinary Λ precisely when $\ker DM_r \not\subset \ker D\tau$, i.e. when $\text{grad } \tau$ has a component off the row space of DM_r . We compute the relative residual directly.

configuration	$r = 1$	$r = 2$	$r = 3$	$r \geq N/2$
single dislocation, $N = 6$	0.976	0.800	—	rank saturates
single dislocation, $N = 8$	0.991	0.941	0.737	rank saturates

configuration	$r = 1$	$r = 2$	$r = 3$	$r \geq N/2$
random ACUE lattice, $N = 7$	0.996	0.963	0.569	rank saturates
generic non-lattice, $N = 7$	0.998	0.926	0.195	rank saturates

Read this correctly. While $\text{rank DM}_r < N$ the moment map has a kernel and the question is meaningful; in that entire regime the answer is decisively yes — 80% to 99% of the depth gradient lies in directions the first one or two moments cannot see. Once $r \approx N/2$ the rank saturates at N , $\ker \text{DM}_r = 0$, and there are no moment-blind directions left to test: the residual drops to 10^{-16} not because Λ fails but because the question becomes vacuous. So: **moment-null does not imply depth-null**, quantitatively, for every bandwidth at which the distinction exists.

Class I — caught by Λ directly. Anything whose defect changes the geometry of the closest pair. This includes the collision-stratum families (a configuration sitting on $\delta = 0$ has $\Lambda = 0$ outright, the same mechanism by which function-field Newman constants are pinned by double roots), the half-lattice and PairCeiling adversaries (§ 4.5: a hard lower spacing $\delta_{\min} \geq c/N$ forces $-\Lambda \geq N^{-2}$, incompatible with CUE's $N^{-8/3}$), and the centre-of-mass and secant families of § 4.2, which move the depth law by total variation 0.12 – 0.24 while every balanced moment stays frozen.

Class II — invisible to Λ , caught by the marked depth. Λ is a function of the characteristic polynomial alone, so it cannot see eigenvectors. Two isospectral matrices have *identically* the same depth; we verify this to machine zero ($\tau(G_1) = \tau(G_2) = 0.068725421516$, difference $8 \cdot 10^{-17}$). Counterexamples built to match rank, trace and Hilbert – Schmidt norm while differing in a directional Schur complement live here. The repair is a **marked depth**: deform $G \mapsto G + \eta uu^*$, transport to the circle by Cayley, and differentiate,

$$\chi(G; u) = \partial_{\eta} (-\Lambda(U\{\eta, u\}))|_{\eta=0}.$$

It separates the isospectral pair immediately — median $|\chi(G_1; u) - \chi(G_2; u)| = 0.081$ over random marks, against $|\chi|$ itself of order 0.01 – 0.2.

What drives χ is not what one would guess. The determinant lemma $\det(zI - G - \eta uu^*) = \det(zI - G)(1 - \eta \cdot u^*(zI - G)^{-1}u)$ suggests the marked depth should track the directional resolvent, and hence blow up wherever the inverse-Gram pathology is worst. It does not. Rotating the mark onto the smallest- $|\text{eigenvalue}|$ direction drives the resolvent $u^*(z_0 I - G)^{-1}u$ from 1.615 to $6.667 = 1/|\lambda_{\min}|$ while χ falls from 0.0267 to 0.0028 — an order of magnitude the wrong way. The correct mechanism is local to the collision. Rank-one perturbation moves λ_j by $\eta |\langle u, v_j \rangle|^2$, the Cayley map contributes $d\theta/d\lambda = -2/(1 + \lambda^2)$, and the depth responds through the critical gap, giving

$$\chi(G; u) \approx \rho \cdot (\delta/4) \cdot (c_a |\langle u, v_a \rangle|^2 - c_b |\langle u, v_b \rangle|^2), \quad c_j = -2/(1 + \lambda_j^2),$$

where (a, b) is the pair that collides first, δ their gap, and ρ a background renormalisation of the same kind as the ACUE constant of § 4.7. Measured against random marks the correlation with this formula is 0.9958 (best-fit $\rho = 1.72$ in the instance tested, against $\tau/(\delta^2/8) = 1.28$ for the value itself — derivative and value renormalise differently). The control confirms the mechanism: a mark orthogonal to *both* critical eigenvectors gives $\chi = 2.8 \cdot 10^{-3}$, two orders below typical.

So the marked depth is a **differential alignment detector**: it measures how asymmetrically the mark couples to the two eigenvectors about to collide, and vanishes on marks that couple to them equally. That is a sharper instrument than a resolvent probe, and it says exactly which marks to choose when testing a Class II family — those overlapping the critical pair, not those exploring the ill-conditioned directions.

Class III — genuinely immune to the linear flow. Suppose the observable algebra misses a direction v . Because the heat operator is diagonal in the coefficient basis, $H_t v$ generally stays in the same hidden sector, so if $F(X) = F(X + \varepsilon v)$ for all observables F , then $F(H_t X) = F(H_t(X + \varepsilon v))$ for all t as well. **A linear flow does not destroy linear invisibility** — this is a genuine negative result and it is why "evolve the moments you have" fails, as § 4.2 records.

The point is that Λ is not in that argument's scope. It is not a linear observable transported by the flow; it is the first hitting time of the orbit against the discriminant variety. Two orbits can live in the same invariant hidden sector and still stand at different distances from $D = \{\text{Disc} = 0\}$. Class III is therefore not "counterexamples the depth cannot see" but "counterexamples for which the *linear* dynamic argument gives nothing" — and Criterion I is what decides whether the nonlinear stopping time recovers them.

The programme in one line. With static observables M and fibre $F_m = \{X : M(X) = m\}$, the statement to prove is that τ restricted to F_m is generically non-constant, with marks added, $\Lambda_{\{u_1\}}, \dots, \Lambda_{\{u_k\}}$, until $\cap j \ker D\Lambda_{\{u_j\}} \cap \ker DM = \{0\}$. Static moments supply coarse coordinates; marked Newman constants supply the directional ones. The tables above are the first evidence that this is achievable rather than merely well-posed.

4.10 The depth is the Lagarias–Rodgers hard core in another coordinate

The sharpest consequence of the depth law is that it is not a new statistic competing with the existing formulation of the alternative hypothesis — it is that formulation, in coordinates where our machinery computes exactly.

Lagarias and Rodgers study the class $\mathcal{T}_1$ of point processes that mimic the sine process at coordinatewise bandwidth one — for every r and every Schwartz η with $\text{supp } \hat{\eta} \subset [-1, 1]^r$, the factorial correlation statistics agree — and define

$\mu = \sup\{c : \text{some } X \in \mathcal{T}_1 \text{ has minimum spacing } \geq c\}$, mean spacing normalised to 1.

They prove $\mu \geq 1/2$ by the randomly shifted half-lattice (which is exactly ACUE), record the pair-only upper bound $\mu \leq 0.606894\dots$, and conjecture $\mu = 1/2$. The alternative hypothesis is precisely the assertion that zeta's zeros realise a hard core of $1/2$.

The bridge. On the circle with N points the mean angular spacing is $2\pi/N$, so a hard core of c mean-spacings is $\delta_{\min} = 2\pi c/N$. Combined with the local collision law $-\Lambda = \rho \cdot \delta_{\min}^2/8$ this gives

$N^2(-\Lambda) = \rho \cdot \pi^2 c^2 / 2$, and at the alternative-hypothesis value $c = 1/2$, $N^2(-\Lambda) = \rho \cdot \pi^2/8$.

Here $\rho \geq 1$ is the background delay factor of § 4.5 and § 4.7. Every number we computed is this identity read in one direction or the other. The ACUE quantiles at $N = 3, \dots, 10$ give

N	min $N^2(-\Lambda)$	$\rho_{\min}$	median	ρ_{med}	max	$\rho_{\max}$
6	1.353146	1.0968	1.419374	1.15050	1.952629	1.5827
8	1.330383	1.0784	1.418216	1.14956	1.976122	1.6018
10	1.314614	1.0656	1.412774	1.14515	1.985458	1.6094

and the single-dislocation constant of § 4.7 is exactly the median branch: $s^* = 1.41964034\dots = \rho_{\infty} \cdot \pi^2/8$ with $\rho_{\infty} = 1.1507015\dots$. The minimising family $\{0, \dots, N-3\} \cup \{N+3, N+4\}$ — a maximally packed block beside a void — has ρ falling slowly (1.0656, 1.0611, 1.0579, $\dots$, 1.0495 at $N = 18$) toward a limit near 1.03 – 1.05; the block is nearly stationary under the flow, so only its ends move, which is why it resists collapse and keeps $\rho > 1$. The maximising symmetric three-block family gives $\rho \rightarrow 1.6094$, i.e. $N^2(-\Lambda) \rightarrow 1.9861$.

The equivalence of extremal problems. Define the depth extremum $\mu_\Lambda = \sup\{\liminf N^2(-\Lambda) : X \in \mathcal{T}_1\}$. The identity gives, in both directions,

$$\mu_\Lambda \geq \pi^2\mu^2/2 \text{ and } \mu \leq \sqrt{(2\mu_\Lambda)/\pi}.$$

These are not loose: LR's published bound $\mu \leq 0.606894\dots$ is *exactly* the depth bound $\mu_\Lambda \leq \pi^2(0.606894)^2/2 = 1.8177\dots$, and the conjecture $\mu = 1/2$ is exactly $\mu_\Lambda = \pi^2/8 \cdot \rho$ for the half-lattice's own ρ . So **any improvement of the depth bound is an improvement of the Lagarias – Rodgers bound**, at the explicit exchange rate above; a depth bound of 1.29 — barely below what the ACUE configurations themselves realise — would already give $\mu \leq 0.511$.

A falsifiable threshold for the alternative hypothesis. Running the inequality in the direction that costs nothing:

If the zeros of ζ have, in a window at height T with $N = (\log T)/2\pi$ zeros, local Newman depth satisfying $\liminf N^2(-\Lambda) < \pi^2/8 = 1.2337\dots$, then the alternative hypothesis is false.

This uses only the hard-core consequence of AH — not the lattice structure, not any correlation beyond it — and it replaces the vague target "beat the bandwidth wall" with one number. Compare the two sides: AH predicts $N^2(-\Lambda) \geq 1.2337$ with the half-lattice's actual typical value 1.41964, while CUE predicts $N^2(-\Lambda) \asymp N^{-2/3} \rightarrow 0$. The gap between prediction and threshold is a factor 1.15, and between CUE and threshold it is unbounded.

Why this reformulation may be easier to attack. The LR extremum is a constraint on a *minimum*, a hard combinatorial quantity; Palm-type certificates — the row-sum square $\text{Var}^0_{\text{sine}}(S_f) \leq (M_h - A_f)(A_f - m_h)$ for band-limited f , and its multiwindow quadratic generalisation on the packing body K_h — apply naturally to smooth functionals and awkwardly to a minimum. The depth is smooth in the configuration (§ 4.9 computes its gradient), it has a variational characterisation as a distance to a hypersurface, and its derivative has the explicit critical-pair law of § 4.9. Transporting the Palm/packing certificates from the row sum to the depth is therefore the natural next attempt, and unlike the row-sum searches — which failed for an exact local reason, the pattern $\{-h, 0, h\}$ already forcing $S_0 = 2f(h) > 9/7$ for the first nonnegative Fejér profile — the depth has no such immediate obstruction.

One missing lemma. The bridge is rigorous except for $\rho \geq 1$, i.e. that the background never *accelerates* the first collision below the isolated two-body time $\delta_{\min}^2/8$. Every computation in this project supports it: on the CUE side the correction is positive with median $0.598 \cdot N^{-0.729}$, on the ACUE side $\rho \in [1.049, 1.610]$ over all configurations and sizes tested, and the mechanism is clear — the neighbours of a close pair pull its members outward. We have not proved it. It is the single lemma standing between the numerology above and a theorem, and it is exactly the same localisation estimate that Open Problems 1 and 3 require.

4.11 What this could mean for zeta, stated carefully

We claim no theorem about ζ . But the structure suggests a target that is strictly weaker than what the static programme requires. ACUE predicts $N^2(-\Lambda) \asymp 1$; CUE predicts $N^2(-\Lambda) \asymp N^{-2/3} \rightarrow 0$. Therefore, to exclude the ACUE universality class for zeta one does not need the full $8/3$ law: it suffices to prove that the local depth of ζ under the true de Bruijn – Newman flow, suitably normalised at height T , satisfies

$$(\log T)^2 \cdot D_{-T} \rightarrow 0,$$

with any rate. That is a substantially weaker statement than a full extreme-gap theorem, and Rodgers – Tao's method — averaged information plus a heat-flow energy argument, rather than the full GUE law — is evidence that such intermediate statements are reachable. Making the truncation rigorous (one may not simply cut a window of zeros and call

the result a polynomial; the Polymath15 machinery for computing H_t is the honest route) is the first obstacle.

5. How this was done

The methodology is not incidental, so we describe it plainly.

Fleets, not oracles. Work proceeded in rounds. Each round posed one question, and between three and ten language-model agents attacked it in parallel from deliberately different angles, each writing and running its own code, each producing a report tagged line by line as *proved*, *computed*, *heuristic* or *conjecture*. Agents shared a written context file stating the current state of knowledge, including known errors. They did not share conclusions; convergence of independent agents was treated as evidence, and divergence as a bug report.

Adversarial defaults. Every context file instructed agents that the default assumption is that the idea has already been tried and failed, and that their first job is to find the reason. The prompt for the prime-gap round said outright that any improvement to 246 was 99% likely to be a misread constraint. This is not modesty; it is calibration. Of the phenomena that looked like discoveries during this project, most were arithmetic or normalisation errors, and the ones that survived did so because they were attacked first by their own authors.

Exact arithmetic at every threshold. Nothing near a decision boundary was accepted in floating point. The Maynard – Tao certificates are ball-arithmetic with outward rounding; the variational crossings are exact-rational Rayleigh quotients of rounded eigenvectors; the ACUE enumeration uses exact Vandermonde masses with 40-digit spot checks; the signed-sieve identity was verified in exact rationals on a thousand random weights. The single most common failure mode in machine-generated mathematics is a plausible float, and the remedy is cheap.

Two implementations or it did not happen. Every headline number here was produced twice by independent code, usually by agents that could not see each other's work: the record tuples were re-verified by a second admissibility checker; the CUE depth was computed by an ODE integrator and by coefficient bisection agreeing to 10^{-6} ; the single-dislocation constant was obtained from a lattice enumeration and from a transcendental equation, agreeing to six digits. Where two implementations disagreed — an engine crossing at $k = 15,856$ against another at 29,500 — the discrepancy was tracked to its source before anything was claimed.

The human supplies the questions. Every genuinely new direction in this paper came from a human mathematician's judgement: to stop optimising 246 and look at H_2 ; to ask whether removing the square opens a phase; to propose the de Bruijn – Newman depth as a dynamic observable in the first place; to insist that 72π be factorised rather than fitted; to recognise that the single-dislocation configuration was the right object to compute exactly. The models supplied speed, breadth, exact arithmetic, and — importantly — the willingness to write and discard fifty scripts a day. The division of labour was not "human checks machine" but "human chooses the question, machine exhausts it, human reads the residue".

Negative results are the main product. Five walls, one no-go theorem, one empty phase, one refuted extrapolation. A method that only reports successes cannot be trusted about them. The most valuable single output of the signed-sieve round is that it is closed, and that we can say exactly what would open it.

6. Open problems

1. **Rigourise the CUE depth law.** Prove $8N^{\{8/3\}}(-\Lambda) \Rightarrow G^2$. The hard half — the smallest-gap limit law — exists in the literature (Ben Arous – Bourgade, Feng – Wei). What is needed is the localisation lemma: that the remaining $N - 2$

zeros perturb the two-body collision time by $1 + o(1)$. Our data show the correction is positive and of size $\approx 0.60 \cdot N^{-0.73}$.

2. **Make $s^* = 1.419640342\cdots$ a theorem** for the single-dislocation family, and compute the separated-defect constant $\rho_\infty = 1.19120\cdots$ exactly. The natural framework is an infinite clock with one localised defect under the Coulomb dynamics, i.e. a Calogero – Moser problem; the target is a closed form in terms of theta or Bessel data.
3. **The marked depth law.** Prove $\chi(G;u) = \rho \cdot (\delta/4) \cdot (c_a|\langle u,v_a \rangle|^2 - c_b|\langle u,v_b \rangle|^2)$ with the background constant ρ , and determine ρ for the lattice families. Empirically the correlation is 0.9958 (§ 4.9); what is missing is the same localisation estimate as in problem 1, differentiated.
4. **A transversality conjecture.** With static observables M and fibre $F_m = \{X : M(X) = m\}$, prove that $\tau = -\Lambda$ restricted to F_m is generically non-constant; equivalently $\ker DM \not\subset \ker D\Lambda$, with marks added until the intersection is trivial.
5. **Function-field depth.** Frobenius conjugacy classes give characteristic polynomials, hence depths. Does equidistribution of Frobenius push through to the depth law — a genuinely arithmetic Newman-depth universality theorem, rather than a random-matrix analogy? The $g = 1$ case is the exact $N = 2$ law in arithmetic clothing, $\Lambda = \log(|a_p|/2\sqrt{p})$.
6. **(E_θ) , the tuple-residue well-factorable estimate**, for any $\theta > 1/2$. This is the single statement standing between the certified price list of § 3.5 and $H_1 \leq 130$.
7. **The $k = 49$ door.** Decide whether some legal Maynard – Tao variant has $M_{49} > 4$. Certified 3.930490592, float optimum 3.959325169, upper bounds closing only $\varepsilon \leq 0.00682$.

References

J. Maynard, *Small gaps between primes*, Ann. of Math. 181 (2015) 383 – 413 · D.H.J. Polymath, *Variants of the Selberg sieve, and bounded intervals containing many primes*, Res. Math. Sci. 1:12 (2014) · Y. Zhang, *Bounded gaps between primes*, Ann. of Math. 179 (2014) · J. Stadlmann, *On primes in arithmetic progressions to smooth moduli*, Adv. Math. (2025) · E. Bombieri, J. Friedlander, H. Iwaniec, *Primes in arithmetic progressions to large moduli*, Acta Math. 156 (1986) · J. Maynard, *Primes in arithmetic progressions to large moduli II: well-factorable estimates*, Mem. AMS 1543 · A. Pascadi, arXiv:2505.00653 · B. Rodgers, T. Tao, *The de Bruijn – Newman constant is non-negative*, Forum Math. Pi 8 (2020) · T. Tao, *The alternative hypothesis for unitary matrices* (blog, 2019) · D.H.J. Polymath, *Effective approximation of heat flow evolution of the Riemann ξ function* (Polymath15) · G. Ben Arous, P. Bourgade, *Extreme gaps between eigenvalues of random matrices*, Ann. Probab. 41 (2013) · R. Feng, D. Wei, *Small gaps of circular β -ensemble*, Ann. Probab. 49 (2021) · R. Killip, I. Nenciu, *Matrix models for circular ensembles*, IMRN (2004) · N. Katz, P. Sarnak, *Random matrices, Frobenius eigenvalues, and monodromy*, AMS (1999) · T. Engelsma, tables of minimal admissible tuples; A.V. Sutherland, narrow admissible tuples database.

All engines, certificates, tuples and audit scripts are in the accompanying repository.

Historical source 32: joint_context_v2

Source: `historical/riemann-rmt/joint_context_v2.md`. SHA-256:
`ed977e3f56ffedd8e777663a44df40a4fc03734a8819967953f96f8bebb2576b`.

Shared context v2 — Anthropic method (from Lean source) × finite RMT fiber program

Read this fully before starting your task. You are one of 10 parallel agents; your final text is your report.

A. Anthropic's actual method (read from the real Lean repo, NOT paraphrase)

The full Lean 4 formalization is cloned at `/workspace/anthropics/zeta-23-lean` (read it!). Key facts extracted:

- Lemma R (rank – trace), k-form** (`Zeta23/ZeroSide/RankTraceMult.lean`): For $P = \sum_j m_j v_j v_j^*$ (on-line zero atoms, multiplicities m_j , near-unit vectors), Q Hermitian with $n_+(Q) \leq b$: $2c \cdot \text{tr}(P+Q) - \|P+Q\|_{F^2} \leq \sum_j k_c(m_j) + c^2 \cdot b$, where $k_c(p) = c^2 - ((c-p)_+)^2 (= 2cp - p^2 \text{ for } p \leq c, \text{ saturates at } c^2 \text{ for } p \geq c)$ — a concave clipped secant. Decoding values: $k_2(1)=3, k_2(\geq 2)=4 (\Rightarrow \text{simple-zero counts, } c=2)$; $k_3(1)=5, k_3(2)=8, k_3(\geq 3)=9 (\Rightarrow \text{distinct counts, } c=3, \text{ gives } 5/6)$. Proof: spectral split $Q = Q_+ - Q_-$, drop $\text{tr}(PQ_+) \geq 0$, von Neumann trace inequality, elementary secants $(x-c/2)^2 \geq 0, (x-c)^2 \geq 0$.
- Tightness** (`Zeta23/ZeroSide/TightMult.lean`, `lemmaR_tight`): for on-line atoms with integer $m_j \leq c$ on orthonormal vectors plus b pair-blocks of eigenvalue c , EQUALITY holds: $2c \cdot \text{tr}(P+Q) - \|P+Q\|_{F^2} = \sum_j k_c(m_j) + c^2 \cdot b$. So the certificate cannot be improved using only $(\text{tr}, \|\cdot\|_{F^2}, \text{rank}, \text{inertia})$.
- Prime side**: $\text{tr } G$ and $\text{tr } G^2 = \|G\|_{F^2}^2$ of the Gabor/taper compression are computed unconditionally from the explicit formula (`Zeta23/PrimeSideA,B`, `Poisson.lean`, `Taper/`), with Montgomery – Vaughan Hilbert inequality (`Zeta23/MV/`) giving near-orthogonality of zero vectors ($\text{Gram} \approx I$). Bandwidth-one support; Weyl perturbation for tails.
- The constant**: Theorem D optimal window (`ThmD/Functional.lean`) = Montgomery – Taylor: $v^*(s) \propto \cos(\sqrt{2}s)$, $q^* = \frac{1}{2} + (1/\sqrt{2})\cot(1/\sqrt{2})$, $\delta_{MT} = 2 - q^* = 0.672500703679$.
- PairCeiling** (`Zeta23/PairCeiling/Stability.lean`): stability inequality (two integrations by parts): any valid bandwidth-one certificate (c_0, r) , $r \in C^1[0,1]$, against a configuration with grid form-factor masses s_j and simple fraction p : $c_0 + \int_0^1 r(x) \cdot x \, dx \leq p + |r(1)| \cdot |D(1)| + |r'(1)| \cdot |E(1)| + \sup |E| \cdot \int_0^1 |r''|$. Instance at explicit 256-periodic marked law (`LawN256.lean`): every bandwidth-one certificate certifies $\leq 0.6818287 + 2.55 \cdot 10^{-6} \cdot (|r'(1)| + \int |r''|)$. The only unverified input is `Enc10K` (256 integer enclosures from an exact-rational certificate, hash-named, not in repo). Note $2.55e-6 \approx 1/(6 \cdot 256^2)$ (grid/Bernoulli fingerprint).
- ξ' theorems** (`Zeta23/XiPrime/`): same rank – trace device on the Farmer – Gonek/Montgomery argument for ξ' gives unconditionally 0.85838 simple-on-line (flat window) and 0.86864 with a QUARTIC window — so quartic windows already pay when the arithmetic input exists (ξ' explicit formula supplies it).

B. Verified finite results (our program; verify scripts in this scratchpad)

- ACUE = rank- N Fourier projection DPP on $\mathbb{Z}/2N$ (Tao 2019); mimicker fiber (all laws matching balanced moments $E[p_\lambda \bar{p}_\nu] = \delta_{\lambda,\nu}$, $|\lambda| = |\nu| \leq N$): dims 0,0,2,10,80,403,1804 for $N=3..9$. $N=5$ exact over $\mathbb{Q}(\sqrt{5})$.

- **Rigidity:** all 2-point (proved, aliasing) and 3-point (verified 1e-14) correlations frozen on the fiber; first freedom at 4-point. Fiber gap $\delta^*(N)$ for word 01010101: 0.0102, 0.0123, 0.0138, 0.0151 ($N=6..9$), transfer theorem: uniform gap $\Rightarrow$ Lagarias – Rodgers non-uniqueness.
- **Nyquist conservation law** (verified): with marks $m_x \geq 0$, $\sum m = N$ on $\mathbb{Z}/2N$ and open rows $E|p_k|^2=k$ ($k \leq N$): $E\sum m(m-1) = (E|p_N|^2 - N)/2N$; closing row $k=N$ forces binary. Slack identity: $Es_1 - (N/2 + \csc^2(\pi/2N)/2N) = E\sum m_x m_{x+1} + E\sum h(m_x)$, $h=0$ on $\{0,1,2\}$; the combination is Nyquist-blind ($1+\cos \pi = 0$). Bound $Es_1/N \geq 1/2 + 2/\pi^2 = 0.70264$. Equality face $E_N = \text{hard-core } \{0,1,2\}$, $|E_N| = C(2N,N)$. Open: all- N attainment ("Nyquist equality law").
- **Fourth moments of the sine – Gram features** (verified $N \leq 6$): for $V_k = |p_k|^2$ under ACUE, $\text{Cov}(V_k, V_\ell) = k^2 1_{\{k=\ell\}} - 2(k+\ell-N)_+$, $\lambda_{\min} = 1$ exactly. Inradius $\rho_N \geq 1/(N^2 - N/2)$; repair theorem: feature error $o(N^{-2})$ on the equality face $\Rightarrow$ exact law with $o(N)$ slack.
- **Holonomy:** DPP laws on g -cycles agreeing to order $\leq g$ with $dTV = (2r)^g |\cos g\varphi - \cos g\psi|$ exactly. Homometry: marked pairs equal complete $| \text{Fourier} |^2$, different simple fractions. Local blindness: exponential fibers can have zero projection on all fixed local words.
- **Phase twists** (new paper): multiplying by center-of-mass characters translates the Fermi sea; rank-4 twists are stochastically homometric to ACUE (equal Patterson autocorrelation at all frequencies, marginals through order 3, separated at 4-point); the full moment-null space has an exact Fourier – Slater description, exponential dimension; matched-block compressions are forced = identity (all residual in the cross block).
- **Christoffel trace hierarchy through degree 12** (new papers XI/XII): exact continuum moments of the sine – Gram matrix $m_9=11166011/151200$, $m_{10}=83443081/554400$, $m_{11}=852071287/2721600$, $m_{12}=1033020076559/1556755200$; sharp forced-rank (Christoffel at 0) $R_{10} = 0.9385985\dots$, $R_{12} = 0.9482017\dots$ (exact rationals; Gauss – Radau adversaries attain). Marked benchmarks: 80.39% (deg 6), 84.94% (8), 87.72% (10), 89.64% (12). Odd – even Schur law: odd trace = coupling coordinate (no standalone gain), next even trace = innovation energy, gain = squared ratio. Bandwidth: order- k collision channel has max Fourier weight $\lfloor k/2 \rfloor \cdot L$ (quotient-MaxCut). All-order N^{-2} parity (Gorenstein). Hard-edge conjecture: support $[0,3]$, Christoffel $\Lambda_n(0) \sim (3n)^{-1}$.
- **Third-moment escape** (series synthesis): $E(M)_3 = r$ without cap gives NO gain (mass escapes to high marks); fourth factorial moment restores $p_1 - (2-q) \geq r^2/(s+2r)$.
- **Scalar-compression moment code** (Prop 8.1, series synthesis): if subspace V satisfies $P_V D_j P_V = b_j P_V$ for diagonal moment observables D_j , then every unit $v \in V$ gives an honest law $q_v(o) = |v(o)|^2$ matching all moments. (Quantum-marginal-style realization tool.)
- **Inner-SOS $\subseteq$ exact LP $\subseteq$ outer-SDP sandwich;** signed PSD pseudo-laws are not laws. Inertia alone gives no local gap: exponentially large word-invisible positive tangent exists (rank ceiling $2^{\ell-1}-1$ for length- ℓ words).
- Two constants NEVER to conflate: $\delta_{MT} = 0.6725007$ (zeta certificate) vs $1/2 + 2/\pi^2 = 0.70264$ (marked open-Nyquist model). Literal ACUE fiber is binary ($s_1 \equiv N$).

C. Environment

- python3 + numpy/scipy in this scratchpad; write scripts here, name by your topic. Exact rational/cyclotomic arithmetic via sympy (pip install sympy if needed).
- Lean repo at /workspace/anthropics/zeta-23-lean — READ the relevant files for your task (grep docstrings; they cite paper § §).
- arxiv.org / anthropic.com are egress-blocked; WebSearch summaries work.

- Report format: English, precise; tag every claim [PROVED] / [COMPUTED] (with script path + numbers) / [HEURISTIC] / [CONJECTURE]. State explicitly what is NEW relative to sections A/B above. End with the single most promising concrete next step.

Historical source 33: prime_gap_context

Source: `historical/riemann-rmt/prime_gap_context.md`. SHA-256:

`f46b6cffca37508e2d90a966befbdc80dcc07d39aaf68654adb04cf1a7edc77f`.

Shared context — Round 4: attacking the bounded prime gap record $H_1 = 246$

You are one of 10 parallel agents. Read fully. Your final text is your report. English, tagged

[PROVED]/[COMPUTED]/[HEURISTIC]/[CONJECTURE], scripts in this scratchpad, end with the single most promising next step. BE ADVERSARIAL WITH YOURSELF: the default assumption is that Polymath already tried your idea — your job is to find what they could not do in 2014 (modern SDP/SOS solvers, our new machinery, post-2014 arithmetic inputs) or to prove a clean no-go.

A. The state of the art (verified via search, August 2026)

- $H_1 = \liminf(p_{n+1} - p_n) \leq 246$ — Polymath8b (2014), unconditional, still the record. Under Elliott – Halberstam (EH): 12. Under generalized EH: 6 (parity blocks further).
- $H_2 = \liminf(p_{n+2} - p_n) \leq 398,130$; $H_3 \leq 24,797,814$; $H_m \ll e^{\{3.815m\}}$ — much less optimized than H_1 !
- Pipeline: $DHL(k,2) \iff M\text{-type variational threshold} > 4/\theta \cdot \dots$ — with Bombieri – Vinogradov level $\theta = 1/2$, the requirement is $M_k > 4$ (pure) or $M_k^{\epsilon} > 4$ (ϵ -trick). Then $H_1 \leq H(k) = \text{diameter of the narrowest admissible } k\text{-tuple}$.
- Numbers: Maynard: $M_{105} > 4 \rightarrow H=600$. Polymath8b: $M_{54} > 4.00238 \rightarrow H(54)=270$; the ϵ -trick got the $k=50$ threshold $\rightarrow H(50)=246$. They stopped at 50, writing that 50-or-49 was the limit; the $k=49$ door was left numerically closed but (to our knowledge) NOT rigorously closed for the most general variant. $H(49) = 240$, $H(48) = 236$, $H(47) = 232$ (verify these diameters yourself from the admissible-tuple databases/searches — do not trust this line blindly).
- The pure variational problem: $M_k = \sup \text{ over symmetric } F \in L^2(\text{simplex } R_k = \{t \in [0,\infty)^k : \sum t_i \leq 1\})$ of $k \cdot J(F)/I(F)$, $I(F) = \int F^2$, $J(F) = \int (\int F dt_k)^2 dt_1 \dots dt_{k-1}$. Known: $M_k \geq \log k - C$; $M_k \leq (k/(k-1)) \cdot \log k$ (so no barrier at 4 from above for $k=49$: $\log 49 = 3.892 < 4$ — the PURE M_{49} is provably < 4.06 -ish and numerically ≈ 3.85 ? — establish the true value). The ϵ -trick enlarges support to $(1+\epsilon) \cdot \text{simplex}$ with a vanishing-marginal condition, trading against a $(1-k\epsilon)$ -type penalty; Polymath also studied M'_k variants (unrestricted support with vanishing marginals) and proved upper bounds. KEY QUESTION for this round: what is the sharp upper bound for the BEST legal variant at $k = 49$, and is it above or below 4?
- Recent arithmetic inputs (post-2014): Polymath8a equidistribution $\theta = 1/2 + 7/300$ for smooth/densely-divisible moduli (restricted residues — does NOT plug directly into MT sieve, which needs all residue classes; Polymath8b Section on "using equidistribution" got only ϵ -improvements and did not lower 246); Stadlmann (2023+) improved smooth-moduli exponents; Guth – Maynard (2024) large-value estimates for Dirichlet polynomials (improved zero-density; feeds short-interval prime results, NOT directly the BV level); Lichtman's sieve refinements. Verify current versions via WebSearch as needed.

B. Our machinery inventory (from Rounds 1–3; all in this scratchpad + /workspace/anthropics/zeta–23–lean)

1. **Decoder/LP theory:** generalized concave-score Lemma R; capped decoder LPs with escape-channel analysis; exact adversary construction (moment LPs with spread/pair/negative channels). The MT variational problem is a Rayleigh quotient = exactly our territory.
2. **SOS/Gram certificates:** multivariate trigonometric-SOS strictly beating scalar Fejér – Riesz cones (colleague's result: $N=3$ exact constant $(67-2\sqrt{22})/81$; $N=4$ exact-rational); inner-SOS/exact-LP/outer-SDP sandwich.
3. **Positive-inertia decoding:** rank – trace inequalities turning quadratic-form data into COUNTS (many positive directions $\rightarrow$ many prime-rich configurations) — the "matrix score $n_+(B-A)$ " idea; the $\sqrt{\epsilon}$ rigidity theorem (near-equality forces structure).
4. **Krylov/fermionic transfer:** half-power trick $\langle v, B^{\{2m\}} v \rangle = \| B^m v \|^2$; second-quantized transfer for symmetric-function operators — high-order moment computations at otherwise impossible sizes.
5. **Non-identifiability/counterfeit methodology:** exact finite fibers, identity-compression, parity-type obstructions made quantitative in LP language.
6. **Euler – Maclaurin/cotangent tower:** exact finite-size ladders for variational discretizations (know your discretization error EXACTLY instead of eyeballing convergence).

C. Honesty rules for this round

- 246 has stood for 12 years against Tao, Maynard, and a swarm of Polymath optimizers. Any "improvement" you find is 99% likely a bug or a misread constraint: re-derive the exact threshold inequality (including the $\theta = 1/2$ factor and the ϵ -penalty terms) from the Polymath8b paper structure before claiming anything.
- The deliverable hierarchy (in decreasing value): (i) a verified new bound $H_1 < 246$ (extraordinary claim, extraordinary checking); (ii) a verified improvement of $H_2 = 398,130$ or higher-m records (much softer targets!); (iii) a rigorous closure ("the $k=49$ door is closed for variant class X, here is the dual certificate") — a real theorem; (iv) a quantified reduction ("improving BV level to $\theta = 1/2 + \delta$ for all-residue averaged moduli would give $H_1 = f(\delta)$ ") with the exact function f ; (v) honest no-go maps.
- Compute budget: sympy/numpy/scipy/mpmath available; pip install cvxpy/OSQP if needed for SDPs. Krylov bases for the variational problem are 1-D integrals of polynomial bases — exact rational arithmetic is feasible; use it for anything near a threshold.
- python3 in this scratchpad; WebSearch works (arxiv.org itself is egress-blocked for fetch; search snippets work; the Polymath wiki michaelnielsen.org and Tao's blog terrytao.wordpress.com are ALSO blocked for direct fetch — rely on search result snippets and your own derivations).

Historical source 34: signed_context

Source: `historical/riemann-rmt/signed_context.md`. SHA-256:

`f9d18ed12b12644f97574bf2efc9613815334722879a693dce0995adf7530fae`.

Shared context — Round 5: the SIGNED sieve experiment (Zhang × Maynard–Tao × ACUE, one language)

You are one of 3 parallel agents executing the first concrete experiment of the unified program: **does removing the square/positivity constraint from the Maynard – Tao variational problem produce a genuinely new phase ($\lambda_{\text{signed}} \gg \lambda_{\text{positive}}$), and at what arithmetic price?** English, tags [PROVED]/[COMPUTED]/[HEURISTIC]/[CONJECTURE]; scripts in this scratchpad with your prefix; final text = your report; end with the single most promising next step. Be adversarial with yourself: the trivial version of this experiment has a KNOWN null answer (see § 3) — your job starts where the trivial version ends.

1. The three formalisms, written as one

Level 0 (decode). Weights $w: n \in [N, 2N) \rightarrow \mathbb{R}$; $v(n) = \#\{i: n+h_i \text{ prime}\}$ for an admissible k -tuple H . $S_1 = \sum w(n)$, $S_2 = \sum w(n)v(n)$. Classical DHL decode: if $w \geq 0$ pointwise and $S_2 - mS_1 > 0$ then some n has $v(n) \geq m+1$. If w is signed the implication FAILS: the positive excess can be generated by $w(n) < 0$ at prime-poor n . Decode repair requires paying the **debt** $D := \sum_{\{w(n)<0\}} |w(n)| \cdot (m - v(n))_+$ — either bound D directly, or (Zhang-style) support w_- on a set where v is controlled by independent means (exceptional-character / bilinear estimates).

Level 1 (divisor-sum data). MT weights: $v(n)$ = vector of divisor sums, $w(n) = Q(v(n))$ for a quadratic form Q on λ -space. Classical MT: Q = rank-1 PSD, i.e. $w = (\sum_d \lambda_d)^2$. MT asymptotics make S_1, S_2 LINEAR in Q : $S_1 \sim \text{Tr}(A Q) \cdot X$, $S_2 \sim \sum_i \text{Tr}(B_i Q) \cdot X$ (A, B Gram forms; in the F-language these are the I- and J-integrals). The variational constant is the generalized eigenvalue problem $M_k = \sup \text{Tr}(BQ)/\text{Tr}(AQ)$.

Level 2 (finite/ACUE certificates — our machinery). All of the above have exact finite- N avatars: LPs/SDPs over the ACUE fiber with ball-arithmetic certificates; the marked-ACUE operator whose λ gave 0.702642 (free) vs 0.672500703679 (Anthropic's constrained value); kill-graph parity obstructions as odd-cycle facets of the cut polytope. Finite models are where "new phase?" can be answered EXACTLY.

2. Our proved walls — what they do and do NOT forbid (critical: read carefully)

- **P5 (decode wall):** scalar decode $S_2 - mS_1$ exactly optimal — proved WITHIN the class of decodes consuming linear-sieve (type-I) data with pointwise-nonnegative weights. Convex-order two-point counterfeit construction.
- **P7 (weight-cone wall):** sup over PSD Q equals sup over rank-1 Q (SOS decoupling by subadditivity); the copositive relaxation is flat (9-pattern dual certificate). So the entire PSD cone gains ZERO over Maynard's square.
- **Parity wall:** a parity twist kills a set of pair-conclusions iff its kill-graph is bipartite; floors $H_1 \geq 6$, $k \geq 2m+1$ for the positive-weight linear-sieve class.
- **Level wall (#6):** the post-2014 levels $\theta = 4/7$ (BFI), $3/5$ (Maynard II), $5/8$ (Pascadi) are **well-factorable / fixed-residue**: structurally unusable by CLASSICAL MT precisely because positive squares cannot be factored into the bilinear shapes those theorems require.

The signed enlargement exits the hypothesis class of ALL FOUR walls simultaneously. That is its entire point, and why it is not already dead. Indefinite Q is outside $P7$'s cones; a signed decode with debt is outside $P5$'s class; signed weights are the only known door past parity (Selberg's examples block positive linear-sieve data only — the χ -twisted counterexamples are themselves SIGNED objects); and well-factorable λ 's are exactly what BFI/Maynard-II/Pascadi levels consume.

3. The trivial null result (do not rediscover it; start beyond it)

If you only drop PSD from the RATIO objective, nothing happens: sup over ALL symmetric Q of $\text{Tr}(BQ)/\text{Tr}(AQ)$ (with $\text{Tr}(AQ) > 0$) is attained at rank-1 extreme points and EQUALS the PSD value $= \lambda_{\max}$ of the pencil $(B, A) = M_k$. [PROVED — eigendecomposition makes the ratio a weighted mean of Rayleigh quotients.] So $\lambda_{\text{signed}} = \lambda_{\text{positive}}$ for the naked ratio; no new phase there. The genuinely new physics can only enter through the two things that CHANGE when w may be signed: (i) the decode inequality acquires the debt term D (a new, NONLINEAR-in- Q constraint: D depends on the negative part of the function $n \mapsto Q(v(n))$, not of the matrix Q), and (ii) the ARITHMETIC INPUT set enlarges: signed/well-factorable weights may consume $\theta \in \{4/7, 3/5, 5/8\}$ bilinear estimates that positive squares cannot touch. The experiment is therefore: optimize the REPAIRED functional $\Phi_{\beta}(Q) = \sum_i \text{Tr}(B_i Q) - m \cdot \text{Tr}(AQ) - \beta \cdot D(Q)$ (or the constrained version $D(Q) \leq \delta \cdot \text{Tr}(AQ)$) as the control parameter β (resp. δ) sweeps from the classical limit ($\beta = \infty \iff w \geq 0$ forced) to full bilinear control (β small), and detect whether the value exhibits a PHASE TRANSITION (kink/jump, optimizer structure change, parity-floor crossing) rather than a smooth ramp.

4. Threshold ledger (why the payoff is enormous if the phase exists)

DHL($k, m+1$) needs M -value $> 2m/\theta$. Ladder for $m = 1$: $\theta = 1/2 \rightarrow 4$ ($k = 50$, $H = 246$); $\theta = 4/7 \rightarrow 3.5$ ($M_k > 3.5$ at $k \approx 32$, $H(32) = 140$ [verify]); $\theta = 3/5 \rightarrow 10/3 \approx 3.333$ ($k \approx 28$, $H(28) = 120$ [verify]); $\theta = 5/8 \rightarrow 3.2$ ($k \approx 25$, $H(25) = 104$ [verify]); $\theta = 1$ (EH) $\rightarrow 2$ ($k = 5$, $H = 12$). Every entry between 246 and 12 is currently blocked ONLY by the positivity/well-factorable incompatibility. M_k values: $M_4 = 1.845401$, $M_5 = 2.007080$, $M_{25} \approx 3.10?$, M_{28}/M_{32} compute yourself (1-D engine or symmetric Galerkin, both described in codex_handoff.md § 3); $M_{50} \geq 3.9071$, $M_{105} > 4$. $H(k)$ table: $H(25..32)$ verify from Engelsma/OEIS A008407; $H(50) = 246$.

5. Known signed precedents (mine them, don't reinvent)

- Chen's theorem: the ONE historical success of a signed decode — w = weighted sieve minus penalty terms; the debt was paid by the switching principle (bilinear form of the remainder).
- Zhang's Landau – Siegel program: negative part supported where the exceptional character forces v -control; "摆脱平方" = replace $(\Sigma\lambda)^2$ by $\Sigma\lambda \cdot (\Sigma\mu$ shifted) cross terms whose bilinear structure matches dispersion estimates.
- Well-factorable λ (Iwaniec): the linear sieve's $\lambda_{\pm}$ ARE signed and DO consume $\theta = 4/7$ -type levels — but only for $m = 0$ -type conclusions (upper/lower bounds on singletons), never DHL($k, m+1$), $m \geq 1$. The precise obstruction to lifting Iwaniec's trick to tuples = the debt term. Quantify it.
- Selberg's parity examples & Bombieri's asymptotic sieve: the no-go boundary. Any claimed $\lambda_{\text{signed}} \gg \lambda_{\text{positive}}$ at β small must be checked against a parity counterfeit: construct the Liouville-twisted measure explicitly in the finite model and evaluate your repaired functional on it. If your optimizer survives the counterfeit, THAT is the headline.

6. Compute discipline

python3 + numpy/scipy/sympy/mpmath in this scratchpad; exact rational arithmetic near any threshold; LP/SDP via scipy.optimize.linprog (HiGHS) or cvxpy if you pip-install it. Small- k pencils: the Dirichlet-integral Galerkin bases from codex_handoff.md § 3(a) (files $p1_p2_conventions$). Finite- N ACUE: enumeration code patterns in $dir^*/p8_files$.

WebSearch works (snippets only; arxiv/ wordpress/ michaelnielsen direct fetch blocked). Do NOT re-run the record-chain engines (p9_*) — that program is finished and committed.

Historical source 35: codex_handoff

Source: `historical/riemann-rmt/codex_handoff.md`. SHA-256:
`2f6f5c614cbd639bbcd23231698a1833243aff058277ac388df040db563e571a`.

CODEX HANDOFF — Bounded prime gaps: formulas, methods, current state, and open computations

Session handoff, August 15, 2026. Everything below is machine-actionable; file paths relative to the session scratchpad.
Status tags: [P] proved, [X] exact-rational certified, [C] computed (float/hp, uncertified), [R] recalled-needs-verification.

1. The framework in five formulas

- Variational constant.** $R_k = \{t \in [0, \infty)^k : \sum t_i \leq 1\}$; symmetric F ; $I(F) = \int \{R_k\} F^2 dt, J^{\{m\}}(F) = \int \{R_{k-1}\} (\int_0^{1-\sum} F dt_m)^2$; $M_k = \sup R_k \cdot J/I$.
- Threshold.** $DHL(k, m+1) \Leftarrow M_k\text{-variant} > 2m/\theta$. Unconditional $\theta = 1/2$ (Bombieri – Vinogradov); $> 4m$. With $MPZ[\overline{m}, \delta]$ (Polymath8a Deligne strength: $600\overline{m} + 180\delta < 7$): threshold $2m/(1/2 + 2\overline{m})$ but support truncated by δ (see § 4). Then $H_m \leq H(k) = \min \text{diameter of admissible } k\text{-tuple}$.
- ϵ -trick.** Support $(1+\epsilon)R_k$; J -integrals truncated to fibers with $\sum_{j \neq i} t_j \leq 1-\epsilon$ (vanishing marginal). Zero arithmetic cost. This is how Polymath8b got $k = 54 \rightarrow 50$.
- Upper bounds [P].** $M_k \leq (k/(k-1)) \cdot \log k \Rightarrow M_{49} \leq 3.97290, M_{50} \leq 3.99186$ (pure door dead ≤ 50). ϵ -variants: $M_{\{k, \epsilon\}} \leq (1+\epsilon)(k/(k-1)) \log k \Rightarrow k=49$ door *rigorously* closed only for $\epsilon \leq 0.00682$. No upper-bound technology beyond this exists for larger ϵ — the $k=49$ door is OPEN (Polymath: "undecided").
- Tuple side.** Admissible: for every prime $p \leq k$ some class mod p is missed. $H(k)$ proven minimal $\forall k \leq 342$ (Engelsma + our independent exhaustive re-proof ≤ 62). Ladder: $H(46..50) = 216, 226, 236, 240, 246$. Payoffs from $k=50$: $-6/-10/-20/-30$. Empirical $H(k) \approx k(\log k + 0.77)$ at $k \sim 10^4$ (fits (50,246), (35410,398130)).

2. Current numerical state (the live numbers)

object	value	status	file
M_{49} (pure, $d=20$ power-sum basis $p \leq 5$, dim 1125)	≥ 3.891257590916	[X] exact rational	<code>mt_hp_k49_p5.json</code>
M_{50} (pure, same basis)	≥ 3.907113699811	[X]	<code>hpE.log</code>
$M_{\{49, \epsilon=1/35\}}$ ($d=18$ basis, $n=1597$)	float opt 3.959325169; certified ≥ 3.930490592	[C]/[X]	<code>p2_arb_d18_k49.log</code> , <code>p2_cert_d18_float.log</code>
$M_{\{49, \epsilon=1/25\}}$ ($d=14$)	3.915989908	[C]	<code>p2_arb_d14_epsscan.log</code>
gap to the $k=49$ door	$4 - 3.9593 = 0.0407$ (float), 0.0695 — (certified)	—	—
$m=2$ pure-BV crossing ($M_k > 8$)	$k = 15,856$ (certified 8.00108 by the 1-D engine)	[C — ENGINE UNVERIFIED]	<code>p9_pure2.log</code>

object	value	status	file
m=2 Deligne crossing	$k = 13,467$ (thresh 7.8273 at $\delta=0.0205, \varpi=0.00552$)	[C — ENGINE+THRESHOLD UNVERIFIED]	p9_del12.log
m=3 pure-BV / Deligne	$k = 923,601 / 660,985$	[C — same caveats]	p9_pure3.log , p9_del13.log
m=2 tuple search ($k=35,265$)	best diameter 397,352 vs record 396,504	[C], search interrupted	p4_run_*_d3973*.txt
record H_2	396,504 @ $k=35,265$ (Stadlmann $\theta=1/2+1/40$)	[verified via search]	—
P3's independent product-ansatz m=2 crossing	$k \approx 29,500$	[C, weaker engine]	p3_m2_machine2.py

If P9's engine and threshold chain survive verification, H_2 drops to $\approx H(13,467) \approx 138,000$ (Deligne) or $\approx H(15,856) \approx 165,000$ (pure BV) — a $\sim 2.5 - 2.9 \times$ record improvement. This is the highest-value open verification on the board.

3. The variational engines (how to compute M_k)

(a) Exact symmetric-polynomial Galerkin (small $k \leq \sim 60$). Basis: monomials $(1-P_1)^a P_2^b P_3^c \dots$ (power sums $P_j = \sum t_i^j$), degree $\leq d$. All I- and J-integrals reduce via the Dirichlet formula $\int \{R_k\} \prod t_i^{a_i} (1-\sum t)^b dt = \prod \Gamma(a_i+1) \cdot \Gamma(b+1) / \Gamma(k+\sum a_i+b+1)$ — assemble exact rational Gram matrices A (for kj) and B (for I); solve the generalized eigenproblem; certify by exact rational Rayleigh quotient of the rounded eigenvector (numerator/denominator sizes ~ 600 bits at $d=18$). Engines: P1's (hpE/hpD/hpM.log conventions, precision-boosted inverse iteration $544 \rightarrow 1570$ bits) and P2's (p2_* — includes the ϵ -trick truncation: J-fiber integrals with upper limit $1-\epsilon$; assembly is the bottleneck, ~ 25 min for $n=1597$). (b) ϵ -scan protocol. $\epsilon = 1/\text{den}$ rational for exact arithmetic; scan $\text{den} \in \{35, 31, 28, 25\}$; observed optimum near $\epsilon \approx 0.0286 - 0.04$ at $d=14 - 18$. Shift-invert accelerator: "[shift $\rightarrow 0.078\dots$]" lines mark the spectral shift used. (c) 1-D product-profile engine (large $k \sim 10^4 - 10^6$) — P9's, UNVERIFIED. Ansatz $F = \prod g(t_i) \cdot 1\{R_k\}$; reduce I, J to 1-D integrals of g plus a truncation correction for the simplex constraint (Chernoff/exponential-tilt bound on $P(\sum t_i > 1)$ under the g -density). The soundness of the certified lower bound hinges entirely on the direction of the truncation correction in both I (upper-bounded) and J (lower-bounded) — this is exactly where a sign error would fabricate records. Read p9_*.py before trusting; re-derive the two-sided correction independently. P3's slower-but-cleaner variant (p3_m2_machine2.py, mpmath 30-digit) crossed at 29,500 — the 15,856 vs 29,500 gap between the two engines MUST be explained before any claim. (d) MPZ/Deligne thresholds (P9's del runs). Parameters scanned: $(\delta, \varpi) \in \{(0.008, 0.00927), (0.0143, 0.00739), (0.0205, 0.00552)\}$ — the constraint is $600\varpi + 180\delta < 7$ [R: verify exact form + whether the k -dependent dense-divisibility truncation cost was correctly charged into the certified value; this is the weakest link].

4. Walls — do not spend compute here (all proved this session)

1. Ceiling = tuple diameter: no post-processing of MT output beats $H(k_{\min})$. Pair-correlation upper-bound constants can never lower H_1 .
2. Scalar decode optimal: convex-order two-point counterfeit kills all matrix/inertia/moment decodes; $f(m) = 2m/\theta$ final.
3. Weight cone closed: rank- r SOS decouples (subadditivity); copositive cone flat (9-pattern dual certificate, residual $1.3e-14$).
4. Parity: kill-graph bipartite $\iff$ killable; $H \geq 6$ absolute floor for the method class; DHL($k, m+1$) parity-blocked iff $k \leq 2m$.

5. Guth – Maynard orthogonal to H_1 (zero-density inputs appear nowhere in the beyond-1/2 dependency chains); GRH gives only $\theta = 1/2$.
6. Fixed-residue / well-factorable levels (BFI 4/7, Maynard II 3/5, Pascadi 5/8) structurally unusable by MT. Usable-restricted frontier: Maynard III 11/21 (uniform residues), Stadlmann 1/2+1/40, Pascadi minorant 10/19 — all shell-truncated.

5. Open computations, ranked (with method + acceptance criteria)

1. **VERIFY-P9 (record-critical)**. (i) Re-derive the 1-D engine's truncation control two-sidedly; (ii) reconcile 15,856 (P9) vs 29,500 (P3); (iii) re-derive the exact DHL(k,3) criterion from Polymath8b (Thm 3.5-family) including all ϵ /truncation charges at MPZ[∞ ,8]; (iv) re-certify the crossing k in exact arithmetic. Accept: an exact-rational certificate M -variant(k^*) $>$ threshold with every constant re-derived. Then: narrow-tuple search at k^* (engine: § 6) $\Rightarrow$ new H_2 .
2. **TUPLE-35265**. Beat 396,504: restart from `p4_run_*` state files (best 397,352). Methods that got there: shifted-Schinzel + greedy + class-swap annealing (mod- p re-choices, $p \in [50, 2000]$, min-cost repair). Missing (2014's edge): iterated merging of two good tuples on overlapping windows; LP-based repair. Any diameter $\leq 396,503$ with exact admissibility verification (every $p \leq 35,265$) = world record.
3. **K49-DOOR**. Push (a)-engine to $d = 20 - 22$ at ϵ -scan granularity $\text{den} \in \{30 \dots 40\}$ (warm-start from `p2_vec_hp_k49_d18_e35.npy`); if the float optimum plateaus < 3.98 , run the vanishing-marginal variant (larger legal class, unexplored at $k=49$ — Polymath blog IX stopped at single-digit k). Non-separable dual certificates (weights on $(t_i, u_i, P_2$ -fiber)) are the only known route to CLOSE the door above $\epsilon = 0.0068$.
4. **SHELL-M49 (undecided in print)**. Extract Maynard III divisor-window parameters $\rightarrow$ shell polytope Ω^* ; compute $M_4 \cap (\mathbb{R} \cup \Omega^*)$ with ϵ -trick layered; needs only 8 – 13% of full-shell surplus. Inner/outer sandwich decides vs 4.
5. **P10 wrap-up**. `p10_corner_eig.log` calibrated; sideband corners give conditional zeta-side gains (+0.022 at $B=4$) — finish the BFI-range map and the honest verdict.

6. Tuple-search machinery (P4's, all in scratchpad)

`p4_rho_exhaust.c` (exact minimal diameters ≤ 62 , ρ^* -pruning + mod-3,5,7 branching + parity halving),
`p4_verify_witnesses.py` (exact admissibility checker — run on EVERYTHING before claiming), `p4_bigtuple*.py`
(Schinzel/greedy/anneal pipeline), state files `p4_run_{A,B,C,D}2*.{log,txt}` (four annealing chains, best 397,352).
Exchange rates: $\Delta H_2 / \Delta k \approx 11.2$ near $k = 35,300$; tuple-side slack estimate 0.03 – 0.3%.

7. Key constants (verified this session)

$M_2 = 1.385933$, $M_3 = 1.646440$, $M_4 = 1.845401$, $M_5 = 2.007080$ ($\text{EH} \Rightarrow H_1 \leq 12$); $M_5 \cdot 4 > 4.00238$ [Polymath]; $H_1 = 246$ ($k=50$, proven-minimal tuple); $H_2 = 396,504$ ($k=35,265$, Stadlmann); $H_3 = 24,797,814$ ($k=1,649,821$); pair upper constants 3.3996 (Wu) $\rightarrow \approx 3.30$ (Lichtman), parity floor 2; $H(k)$ $k=3..62$ full list in `p4_payoff_table.txt`; $dH_1 / d\delta \approx -2400$; $k=49/48/47$ tipping $\delta \approx 0.002/0.004/0.006$.

Historical source 36: rmt_zeta_survey

Source: `historical/riemann-rmt/rmt_zeta_survey.md`. SHA-256:

`a60e1a3fb0e3b6db037446a59867c14de0bc2b3b8df38ff4a77738a10278ca2b`.

The Bandwidth–One Ceiling: Finite Random–Matrix Fibers, Moment Rigidity, and the Optimality of the 67.25% Critical–Line Bound

Bill (Qingyun) Sun · GPT5.6SOL · Fable

August 11, 2026

Abstract

The recent unconditional theorem that more than two thirds — precisely $67.2500\dots\%$ — of the nontrivial zeros of the Riemann zeta function are simple zeros on the critical line [9] closes, to within a hair, the fifty-year-old gap between unconditional zero-density technology and the Montgomery – Taylor bound previously known only under the Riemann Hypothesis. This survey develops a structural explanation for the constant $0.6725007\dots$: it is the exact optimal value of a *Chebyshev-type extremal problem over probability measures constrained only by two-point (pair-correlation) data of bandwidth one*, and the proof of [9] is, in disguise, the dual certificate of that optimization problem. We make this precise by building a finite model — the Alternative Circular Unitary Ensemble (ACUE) on $\mathbb{Z}/2N$ and its moment fiber of "CUE mimickers" [13] — and proving a sequence of rigidity and duality results in it. The main findings surveyed here are: (i) an exact **two-point rigidity theorem** (all mimickers share all two-point correlations, at every distance, via a Fourier aliasing argument), and a **three-point rigidity phenomenon** (exact at $N = 5$, machine-verified to 10^{-14} for $N \leq 8$) with first freedom appearing at four points; (ii) **finite uniqueness for $N \leq 4$** and super-exponential fiber growth thereafter ($\dim 2, 10, 80, 403, 1804$ for $N = 5, \dots, 9$); (iii) a **duality dictionary** identifying the rank – trace/inertia argument of [9] with the dual of a linear program whose primal adversaries are mimicker measures and whose dual optimum is the Montgomery – Taylor window, with closed form $0.6725007\dots = 3/2 - (1/\sqrt{2}) \cdot \cot(1/\sqrt{2})$; (iv) a proof that the **Alternative Hypothesis point process is a strictly suboptimal adversary**, losing to the free band-limited adversary by 0.0301 ; and (v) a **hierarchy plateau**: positivity (sum-of-squares) constraints of any order do not move the ceiling, while genuinely new correlation *data* does — quantifying exactly what arithmetic input (pair correlation beyond bandwidth one, i.e., prime-pair information) would be required to prove more than 67.25% . We close with implications for the Lagarias – Rodgers uniqueness question, where our results localize any possible non-uniqueness to a four-point phenomenon.

Conventions on epistemic status

Because this survey mixes formal proof, exact symbolic computation, high-precision numerics, and heuristics, every result is tagged:

- [P] proved (human-checkable proof, or exact arithmetic in a number field);
- [C] computational fact (machine-verified numerics, residuals at working precision, reproducible scripts);

- [H] heuristic (argument we believe but have not closed);
- [Q] conjecture / question.

1. Introduction

Let $N(T)$ count nontrivial zeros $\rho = \beta + i\gamma$ of $\zeta(s)$ with $0 < \gamma \leq T$, and let $N_0^*(T)$ count those that are simple and on the critical line $\beta = 1/2$. The Riemann Hypothesis (RH) asserts $N_0^*(T) = N(T)$; the folklore strengthening asserts all zeros are simple. Unconditionally, the proportion of zeros known to lie on the line rose slowly for a century: Selberg's positive proportion, Levinson's $1/3$, Conrey's $2/5$ [3], and $5/12 \approx 41.7\%$ by Pratt – Robles – Zaharescu – Zeindler [4].

In August 2026 a research model of Anthropic announced, with a complete Lean 4 formalization (repository `anthropics/zeta-23-lean`, no `sorry`, no extra axioms), the unconditional bounds [9]:

Theorem (Claude, 2026). $\liminf N_0^*(T, 2T)/N(T, 2T) \geq 2 - c_1^* = 0.6725007\dots$, where $c_1^* = 1/2 + (1/\sqrt{2}) \cdot \cot(1/\sqrt{2}) = 1.3274992963\dots$ is the Montgomery – Taylor constant. The same proportion holds for simple critical zeros, and $\geq 83.625\% = (1 + 0.6725)/2$ of all zeros are distinct.

The striking feature is that $2 - c_1^*$ is exactly the constant Montgomery and Taylor obtained *under RH* from the pair-correlation approach (see Cheer – Goldston [2]). The proof of [9] removes the hypothesis at zero cost in the constant, by replacing Montgomery's conditional pair-correlation asymptotic with an unconditional argument built from the Weil explicit-formula quadratic form (following Bombieri [8]), the Montgomery – Vaughan generalized Hilbert inequality [14], and a rank – inertia – trace inequality in place of diagonalization; the "narrow box" hypothesis of Baluyot – Goldston – Suriajaya – Turnage-Butterbaugh [5, 6] is dismantled rather than assumed.

This survey asks, and largely answers, a structural question:

Why this constant? Is $0.6725007\dots$ an artifact of one proof, or the exact value of an intrinsic obstruction?

Our answer, developed through a finite random-matrix model, is the latter:

Thesis. $0.6725007\dots$ is the optimal value of a linear program: *minimize the fraction of simple points over all point processes consistent with bandwidth-one two-point correlation data*. The proof of [9] constructs the dual optimal certificate (the Montgomery – Taylor window); the primal near-optimal adversaries are Alternative-Hypothesis-like measures, which are however **strictly** suboptimal — the true extremal adversary is free. Consequently no method whose arithmetic input is confined to bandwidth-one two-point data can prove a proportion larger than $0.6725007\dots$, and no amount of higher-order *positivity* (sum-of-squares, Lasserre relaxation, higher inertia bookkeeping) can change this; only new *data* can.

The finite model that makes this precise is the ACUE of Tao [11] and the "finite CUE mimicker" framework of the companion paper [13], which proved that the finite moment problem for the ACUE is *non-unique* for $N \geq 5$, with explicit algebraic counterexamples. The body of this survey is organized around six groups of results obtained in that model and their translation back to ζ .

2. Background

2.1 Montgomery's pair correlation and the Montgomery–Taylor window

Montgomery [1] introduced, for the ordinates γ of zeta zeros,

$F(\alpha) = F(\alpha, T)$: the Fourier transform of the pair-correlation measure of normalized ordinates $\tilde{\gamma} = \gamma \cdot (\log T)/2\pi$,

and proved (under RH, later made unconditional in restricted ranges [5]) that $F(\alpha) = |\alpha| + o(1)$ for $|\alpha| \leq 1$, while F is *unknown* for $|\alpha| > 1$ — knowledge there is equivalent to strong information about prime pairs (Goldston – Montgomery).

The classical application: for even test functions r with $\text{supp } \hat{r} \subset [-1, 1]$,

Σ over pairs of zeros of $r(\tilde{\gamma} - \tilde{\gamma}')$ is computable from F on $[-1, 1]$ alone,

and choosing r optimally bounds the fraction of non-simple zeros. Montgomery's Fejér-kernel choice gives $\geq 2/3$ simple (under RH); the optimal choice — the **Montgomery – Taylor window** — improves this to $2 - c_1^*$, where

$c_1^* = 1/2 + (1/\sqrt{2}) \cdot \cot(1/\sqrt{2})$, so that $2 - c_1^* = 3/2 - (1/\sqrt{2}) \cdot \cot(1/\sqrt{2}) = 0.6725007297\cdots$ [P]

(The extremal function is the solution of a one-dimensional Sturm – Liouville-type variational problem on $[-1, 1]$; we verified the closed form to 50 digits against the variational optimum, and it is the unique critical point [C].)

2.2 The unconditional proof of [9], in one paragraph

The proof works with the Weil explicit-formula quadratic form Q on a space of band-limited test functions: critical-line zeros contribute positive-semidefinite directions, off-line zeros contribute hyperbolic (signature $(+, -)$) planes (Bombieri [8]).

Rather than diagonalizing Q , the argument controls the **inertia** of Q through the elementary inequality $n_+(Q) \geq (\text{tr } Q)^2 / \text{tr}(Q^2)$ and its rank-restricted refinements, evaluating $\text{tr } Q$ and $\text{tr } Q^2$ unconditionally via the Montgomery – Vaughan Hilbert inequality [14] (quasi-orthogonality of the vectors $x^\wedge(i\gamma)$), Chebyshev – Mertens bounds, Riemann – von Mangoldt, and Stirling estimates for Γ' / Γ . The optimization over the test-function space reproduces exactly the Montgomery – Taylor extremal problem — which is the first hint that the theorem computes the value of an optimization problem, not merely a bound.

2.3 The ACUE and finite mimickers

The Alternative Hypothesis (AH) is the logical foil to the pair-correlation conjecture: the (unlikely, but unrefuted) scenario in which normalized zero gaps concentrate on the half-integer lattice $1/2\mathbb{Z}$. Lagarias – Rodgers [10] showed AH is consistent with all known band-limited correlation data, and Tao [11] constructed the **ACUE**: a determinantal point process on the $2N$ -th roots of unity (rank- N Fourier projection) whose scaling limit realizes AH and matches CUE band-limited statistics.

The companion paper [13] studies the finite moment problem: call a rotation-invariant probability law q on N -subsets of $\mathbb{Z}/2N$ a **CUE mimicker** if $E_q[p_\lambda \bar{p}_\nu] = \delta_{\lambda\nu} z_\lambda$ for all partitions $|\lambda| = |\nu| \leq N$, where p_k are the power-sum character statistics and the right-hand side is the Diaconis – Shahshahani [12] CUE moment table. The ACUE law μ_N is one solution; [13] proved the solution set (the **moment fiber**) is non-unique for $N \geq 5$, with an explicit counterexample over $\mathbb{Q}(\sqrt{5})$ at $N = 5$, machine certificates at $N = 9, 10$, and a transfer theorem (their Thm 6.2): a uniform correlation gap along a mimicker family would produce a second point process matching the sine kernel in all band-limited correlations — answering the Lagarias – Rodgers uniqueness question in the negative. Their Thm 7.1 shows uniform gaps force exponentially large density ratios (via strong Rayleigh concentration, Pemantle – Peres [15]).

3. The finite moment fiber: dimension and finite uniqueness

Write Ω_N for the N -subsets of $\mathbb{Z}/2N$ modulo rotation, A_N for the balanced-moment constraint matrix on orbit space, and $X_N = \{q \geq 0 : A_N q = b_N\}$ for the fiber.

Theorem 3.1 (Finite uniqueness for small N). [P] For $N \leq 4$ the fiber is a single point: the ACUE is the *unique* CUE mimicker. (rank A_N equals the number of orbits: 4 at $N = 3$, 10 at $N = 4$; exact rank computation in cyclotomic arithmetic.)

Theorem 3.2 (Fiber growth). [C] The affine dimension of X_N for $N = 5, \dots, 9$ is

N	3	4	5	6	7	8	9
orbits	4	10	26	80	246	810	2704
rank A_N	4	10	24	70	166	407	900
fiber dim	0	0	2	10	80	403	1804

At $N = 9$ the fiber occupies two thirds of the orbit space: the moment constraints, complete to full bandwidth, pin down a vanishing proportion of the degrees of freedom.

The $N = 5$ fiber is 2-dimensional and entirely algebraic over $\mathbb{Q}(\sqrt{5})$; we re-verified the explicit mimicker of [13] by exact arithmetic in $\mathbb{Q}(\zeta_5)$ — all 88 balanced-moment pairs, positivity, and the correlation gap — confirming rank 24 and fiber dimension 2 [P].

The growth $2 \rightarrow 10 \rightarrow 80 \rightarrow 403 \rightarrow 1804$ is super-exponential and tracks the excess of the orbit count over the number of independent symmetric-function constraints; heuristically $\dim X_N \asymp |\Omega_N| \cdot (1 - o(1))$ [H]. Interpretation: *at fixed bandwidth, moment data is an exponentially small window into the space of point processes* — the natural finite- N sharpening of the Lagarias – Rodgers observation that band-limited correlations cannot see AH.

4. Rigidity: what the fiber cannot vary

The fiber is huge, yet almost nothing observable varies over it. This is the technical heart of the survey.

Theorem 4.1 (Two-point rigidity). [P] Every mimicker $q \in X_N$ has *identical* two-point correlations at every lattice distance $b \in \mathbb{Z}/2N$ — not merely the band-limited ones: $E_q[\#\{\text{pairs at distance } b\}] = E_{\{\mu_N\}}[\dots]$ for all b .

Proof. The pair-count at distance b is a linear statistic whose Fourier expansion involves only $E|p_k|^2$, $k \in \mathbb{Z}/2N$. Aliasing on the cyclic group folds each frequency k to $\min(k, 2N - k) \leq N$, which lies inside the balanced-moment window; so every coefficient is pinned by the constraints. ■

This is a genuinely finite phenomenon with no continuum analogue: on $\mathbb{R}$, band-limiting *loses* information at high frequencies; on $\mathbb{Z}/2N$, aliasing *returns* it. The discrete model is therefore strictly more rigid than the Montgomery problem — a fact quantified in § 5.

Result 4.2 (Three-point rigidity). All three-point correlations are constant across the fiber: [P] at $N = 5$ (exact, in $\mathbb{Q}(\zeta_1 - \zeta_0)$); [C] for $N = 6, 7, 8$ (null-space projection residuals $\leq 6 \times 10^{-14}$, versus $O(1)$ four-point projections). No aliasing argument covers three points; empirically the Diaconis – Shahshahani identities and rotation invariance conspire — at $N = 5$, 88 constraints crush the 3-point statistics onto the row space. A conceptual proof is open [Q].

Result 4.3 (First freedom at four points). [P for $N = 5$, C for $N \leq 9$] Four-point correlations genuinely vary: at $N = 5$ the first distinguishable pattern is the word **0000** (four consecutive empty sites), and the pattern **01010101** has fiber range $[-0.0263, +0.0255]$ around its ACUE value 0.04434. The LP-extremal gap of the word **01010101** over the fiber grows monotonically:

$$\delta^*(N) = 0.01019, 0.01232, 0.01375, 0.01514 \text{ for } N = 6, 7, 8, 9;$$

the hand-built $N = 9$ certificate of [13] (0.014738) attains 97% of the LP optimum, and their density-capped search box is far smaller than the fiber's actual extent.

Consequence for Lagarias – Rodgers. Combining Theorem 4.1, Result 4.2, and Thm 6.2 of [13]: *any* second point process produced by the finite-to-infinite transfer agrees with the AH/ACUE process in **all** one-, two-, and three-point correlations — at every frequency, not only band-limited ones. If the Lagarias – Rodgers uniqueness question has a negative answer, non-uniqueness is a **four-point phenomenon**. This sharply narrows where to look, and explains why no two-point (spectral/pair-correlation) technique can ever detect it.

5. The duality dictionary: rank—trace = LP dual, MT window = certificate

We now state the central structural claim: the proof of [9] and the finite fiber sit on the two sides of one linear program.

The finite LP. Allow multisets (multiplicities model non-simple zeros). Over probability measures q on N -multisets of $\mathbb{Z}/2N$ satisfying the balanced CUE moments up to degree $D = \alpha N$, maximize the expected non-simple mass:

$$\epsilon_N(\alpha) = \max \{ E_q[\text{fraction of repeated points}] : q \in X_{N,D}^{\{\text{multiset}\}} \}.$$

Result 5.1 (Bandwidth phase diagram). [C] At $N = 5$: $\epsilon = 0.60, 0.27, 0.12, 0$ for $\alpha = 0.4, 0.6, 0.8, 1.0$ (similarly $N = 6$). The continuum Montgomery prediction is $\epsilon_{\infty}(\alpha) = 1/\alpha + \alpha/3 - 1$ ($= 1/3$ at $\alpha = 1$). Full-bandwidth discreteness gives **complete rigidity** ($\epsilon = 0$, forced simplicity!) by the aliasing of Theorem 4.1 — the discrete model is strictly harder to fool than the continuum.

Thesis 5.2 (Dictionary). [P in the finite model; H as a reading of [9]] The following correspondences hold:

	Proof of [9] (dual side)	Finite model (primal side)
data	$F(\alpha)$, $ \alpha < 1$, unconditional	balanced moments, degree $\leq N$, exact
orthogonality	Montgomery – Vaughan Hilbert inequality	Diaconis – Shahshahani identities
certificate	Montgomery – Taylor window	LP dual optimal variables

	Proof of [9] (dual side)	Finite model (primal side)
adversary	AH-type point processes	mimicker fiber
rank tool	inertia vs $(\text{tr } Q, \text{tr } Q^2)$	Fock rank sandwich (rank 1043 at $N = 10$)
$\mathbb{Z}/2$ symmetry	$\rho \leftrightarrow 1 - \bar{\rho}$ (functional equation)	particle – hole $C \mapsto -C^c$

In the finite model this is literal: the rank – trace inequality $n_+ \geq (\text{tr } Q)^2 / \text{tr } Q^2$ is the Lagrangian-dual bound of the LP, and the dual optimal solution, restricted to the two-point sector and interpolated as $N \rightarrow \infty$, converges to the Montgomery – Taylor window [C] (clean convergence requires the two-point-only LP; in the full-moment LP the two-point sector still carries $\approx 52\%$ of the dual weight).

Result 5.3 (Scaling limit). [C, with closed forms P] The bandwidth-one discrete LP value converges to the continuum Montgomery – Taylor value with a computable finite-size correction; matching one-parameter families give exact finite- N formulas of the shape $\frac{1}{2} - 1/(2N^2 \sin^2(\pi/2N))$, consistent with $O(1/N^2)$ convergence to the continuum extremal value.

The dictionary recasts the theorem of [9] as: *the dual certificate exists unconditionally, therefore the primal value is an unconditional upper bound on non-simple mass*. Its optimality (the subject of § 6 – § 7) recasts the question "can we beat 67.25%?" as "is the primal optimum attained?" — and the answer is yes, by explicit near-feasible adversaries.

6. The Alternative Hypothesis is a strictly suboptimal adversary

It is folklore to regard AH — gaps on $\frac{1}{2}\mathbb{Z}$ — as *the* enemy of pair-correlation methods. Our second structural finding is that this is quantitatively false.

Theorem 6.1 (AH suboptimality). [P] In the bandwidth-one extremal problem, adversaries supported on the half-integer lattice $\frac{1}{2}\mathbb{Z}$ force the simple fraction no lower than

$$\frac{1}{2} + 2/\pi^2 = 0.7026423\dots,$$

(equivalently, the AH repeated-mass parameter satisfies $\beta \leq \frac{1}{2} - 2/\pi^2$, attained by an explicit two-parameter family). The free adversary attains 0.6725007. The gap is 0.0301423...

Proof idea. Parametrize AH-admissible pair-correlation measures by their masses on $\frac{1}{2}\mathbb{Z}$ and impose the bandwidth-one constraints; the resulting LP over $\ell^1(\frac{1}{2}\mathbb{Z})$ has an explicit optimum via complementary slackness — the dual certificate would have to be equioscillatory on $\frac{1}{2}\mathbb{Z}$, and the Fourier-analytic obstruction evaluates to the Fejér-kernel value $2/\pi^2$. The free optimum's extremal measure is instead supported on a *non-lattice* configuration $\{1.057, 2.030, 3.020, \dots\}$ of near-integers [C] — close to, but distinctly off, both $\mathbb{Z}$ and $\frac{1}{2}\mathbb{Z}$.

Three consequences. (a) The binding constraint at 0.6725 is **information (bandwidth)**, not **lattice structure**: deleting AH from the adversary class would only improve the constant to 0.7026, and nothing in the bandwidth-one data justifies deleting it. (b) Conversely, any technique that could exclude near-integer-lattice adversaries — a much weaker statement

than excluding AH — would already improve on [9]. (c) The extremal configuration's near-integer (rather than half-integer) spacing explains an old numerical puzzle in the optimization literature: the MT window's equioscillation nodes do not sit on $\frac{1}{2}\mathbb{Z}$.

7. Positivity does not help; data does

Could one push past the ceiling with a stronger *proof technique* over the same data — higher moments of Q , deeper inertia bookkeeping, sum-of-squares certificates? In the finite model this question is decidable, and the answer is no.

Result 7.1 (Hierarchy plateau). [C] Impose, on top of the bandwidth-limited two-point data, positive-semidefiniteness of higher-order moment (Lasserre/SOS) matrices of any degree. The LP/SDP optimal value does not move (changes $< 10^{-3}$ through the levels computable at $N \leq 8$, with degeneracy identifiable as solver conditioning). The reason is structural: the AH/ACUE-type near-optimizers *satisfy every higher positivity constraint* — they are genuine point processes, so all their moment matrices are automatically PSD. Positivity cannot separate what probability cannot.

Result 7.2 (Data cuts). [C] By contrast, adjoining higher-correlation data (not positivity) moves the value sharply: in the truncated-bandwidth multiset LP, adding the exact three-point correlation table cuts the maximal non-simple mass from ≈ 0.294 to ≈ 0.12 in the tested configuration. Data and positivity play categorically different roles: positivity fixes the geometry of the feasible cone; data slices it.

For ζ , Result 7.1 is the finite shadow of a fact forced by Lagarias – Rodgers [10]: the AH process matches *all* band-limited n -level correlations of the sine process, so no unconditional method whose arithmetic input is band-limited correlation data — of **any** order, processed by **any** positivity argument — can prove a simple-zero proportion exceeding what the AH-inclusive LP allows. Our contribution is the quantitative complement: within two-point bandwidth one, that LP value is exactly 0.6725007, it is *attained* (§ 6), and the ceiling for two-point-based certificates of any window is 0.6818287 (the **PairCeiling** bound formalized in [9]), leaving at most 0.0094 accessible to windowcraft — and our computations indicate even that slack is not reachable without new data [H].

What new data would suffice. The escape routes, in decreasing directness [H]: (i) any unconditional lower-bound information on $F(\alpha)$ for $\alpha > 1$ — this is prime-pair (twin-prime-scale) information via Goldston – Montgomery; even an ε -sliver beyond $\alpha = 1$ breaks the LP's tightness; (ii) arithmetic positivity beyond correlations, i.e., Weil-form amplifiers carrying prime structure that is not a function of F (Bombieri's program [8] continued); (iii) in the finite model, four-point data — the first order at which the fiber opens (Result 4.3).

8. Spectral barriers and the uniform-gap problem

The transfer theorem of [13] requires a *uniform* correlation gap δ along $N \rightarrow \infty$. Whether such a family exists is their central open problem, and our computations sharpen both sides of it.

For existence [C]: $\delta^*(N)$ is monotonically increasing through $N = 9$ (§ 4), and the exponential barrier of [13, Thm 7.1] is quantitatively toothless in this range — the forced density ratio $\|dq/d\mu\|_{-\infty} \geq (\delta/2) \cdot \exp(\delta^2 N / 8\ell^2)$ does not exceed 10 until $N \approx 1.8 \times 10^7$. The observed LP extremals already use density ratios ≈ 211 at $N = 8$ (LP vertices with 416 of 810

orbits vacated) — far outside the capped search boxes used in [13], so the certificate landscape is much larger than previously explored.

Against existence [H]: the constraint transfer matrices show top-eigenvalue growth consistent with $e^{\{cN\}}$, and a block-Toeplitz/Szegő-type analysis (incomplete) suggests the feasible gap directions rotate out of any fixed word's dual cone, which would force $\delta^*(N) \rightarrow 0$ along subsequences unless the extremal word is allowed to vary with N . A Fock-space rank computation at $N = 10$ reproduces the rank-1043 sandwich of [13] exactly [C], confirming no hidden degeneracy that would make uniform gaps cheap.

We record the sharpest finite question our methods isolate:

Question 8.1. [Q] Is $\liminf \delta^*(N) > 0$, where $\delta^*(N)$ is the fiber-LP gap of the word **01010101**? An affirmative answer, by [13, Thm 6.2], resolves Lagarias – Rodgers uniqueness in the negative; by § 4 the resulting second process would be indistinguishable from AH below four points.

9. Assessment: what is new here

Relative to the literature, the surveyed program contributes, in our view, five items.

1. **A theorem-level explanation of a constant.** The identification $0.6725007\dots =$ value of a bandwidth-one LP, with the proof of [9] as its dual certificate, converts "the best known constant" into "the exact value of a named obstruction." Constants with this status are rare (compare $1/3$ and Levinson's method, where no matching adversary is known).
2. **Exact finite rigidity theorems.** Two-point rigidity by aliasing (Thm 4.1) is, to our knowledge, the first *exact* rigidity statement for ACUE moment fibers, and the three-point rigidity / four-point freedom dichotomy is new; it localizes Lagarias – Rodgers non-uniqueness at four points.
3. **AH demoted from obstruction to bystander.** Theorem 6.1's strict 0.0301 gap corrects a widely repeated heuristic: the pair-correlation ceiling is *not* "because AH might be true." The extremal adversary is a near-integer, non-lattice configuration.
4. **Positivity/data separation.** Result 7.1/7.2 gives the first quantitative demonstration in this problem that relaxation hierarchies plateau at the information ceiling while data strictly cuts — a clean instance, in analytic number theory, of a phenomenon familiar in combinatorial optimization.
5. **A calibrated target list.** The 0.0094 window between 0.6725007 and the PairCeiling 0.6818287, the $\alpha > 1$ threshold, and Question 8.1 constitute concrete, falsifiable next steps, each with an exact finite-model analogue on which methods can be rehearsed cheaply.

We emphasize scope: results tagged [C] rest on reproducible computation (exact where stated, else with residuals at 10^{-14}); the reading of [9] through the dictionary is proved in the finite model and remains an interpretation — though a predictive one — of the infinite-dimensional proof.

10. Open problems

1. Prove three-point rigidity (Result 4.2) for all N , ideally by exhibiting the algebraic identity behind the numerical conspiracies [Q].
2. Question 8.1 (uniform gaps $\Rightarrow$ Lagarias – Rodgers non-uniqueness).
3. Determine whether any window-plus-positivity certificate can exceed 0.6725007 within bandwidth one, i.e., close the 0.0094 slack from above (our evidence says no [H]).
4. Find an arithmetic source of $F(\alpha)$, $\alpha > 1$: even $F(\alpha) \geq c > 0$ on $(1, 1+\epsilon)$ unconditionally would beat the ceiling; quantify the constant gained as a function of (c, ϵ) via the LP (the machinery of § 5 computes this).
5. Formalize Theorems 4.1 and 6.1 in Lean 4 against the `zeta-23-lean` infrastructure; both are finite-dimensional and within reach of current tooling.

References

- [1] H. L. Montgomery, *The pair correlation of zeros of the zeta function*, Proc. Sympos. Pure Math. 24, AMS (1973), 181 – 193.
- [2] A. Y. Cheer and D. A. Goldston, *Simple zeros of the Riemann zeta-function*, Proc. Amer. Math. Soc. 118 (1993), 365 – 372. (Includes the Montgomery – Taylor computation.)
- [3] J. B. Conrey, *More than two fifths of the zeros of the Riemann zeta function are on the critical line*, J. Reine Angew. Math. 399 (1989), 1 – 26.
- [4] K. Pratt, N. Robles, A. Zaharescu, D. Zeindler, *More than five-twelfths of the zeros of ζ are on the critical line*, Res. Math. Sci. 7 (2020), no. 2.
- [5] S. Baluyot, D. A. Goldston, A. I. Suriajaya, C. Turnage-Butterbaugh, *An unconditional Montgomery theorem for pair correlation of zeros of the Riemann zeta function*, arXiv:2306.04799.
- [6] S. Baluyot, D. A. Goldston, A. I. Suriajaya, C. Turnage-Butterbaugh, arXiv:2501.14545.
- [7] S. Baluyot, D. A. Goldston, A. I. Suriajaya, C. Turnage-Butterbaugh, *The Alternative Hypothesis for zeros of the Riemann zeta-function*, arXiv:2508.10857.
- [8] E. Bombieri, *Remarks on Weil's quadratic functional in the theory of prime numbers, I*, Rend. Mat. Acc. Lincei (9) 11 (2000), 183 – 233.
- [9] Claude (Anthropic), *More than two thirds of the zeros of the Riemann zeta function lie on the critical line*, preprint and Lean 4 formalization, `anthropics/zeta-23-lean` (2026).
- [10] J. C. Lagarias and B. Rodgers, *Higher correlations and the alternative hypothesis*, Q. J. Math. 71 (2020), 257 – 280.
- [11] T. Tao, *The alternative hypothesis for unitary matrices* (ACUE), blog post (2019).
- [12] P. Diaconis and M. Shahshahani, *On the eigenvalues of random matrices*, J. Appl. Probab. 31A (1994), 49 – 62.
- [13] GPT5.6SOL, *Finite non-uniqueness of the ACUE moment problem: CUE mimickers on $\mathbb{Z}/2N$ and the alternative hypothesis*, preprint (2026). (Companion paper; explicit $\mathbb{Q}(\sqrt{5})$ counterexample, transfer theorem 6.2, exponential barrier 7.1.)
- [14] H. L. Montgomery and R. C. Vaughan, *Hilbert's inequality*, J. London Math. Soc. (2) 8 (1974), 73 – 82.

- [15] R. Pemantle and Y. Peres, *Concentration of Lipschitz functionals of determinantal and other strong Rayleigh measures*, *Combin. Probab. Comput.* 23 (2014), 140 – 160.
- [16] J. B. Lasserre, *Global optimization with polynomials and the problem of moments*, *SIAM J. Optim.* 11 (2001), 796 – 817.
-

Reproducibility: all computations tagged [C] are reproducible from short Python scripts (NumPy/SciPy, exact cyclotomic arithmetic where stated) developed alongside this survey; the $N = 5$ verifications are exact in $\mathbb{Q}(\zeta_{10})$.

Historical source 37: rmt_zeta_popular

Source: `historical/riemann-rmt/rmt_zeta_popular.md`. SHA-256:
`d22dbc5e35dfa46a2883a15d78d91434bdb2ee7343a4a9fc2106723d9ee5c0fb`.

Why 67.25% Is Not Just a Number

A plain-language companion to "The Bandwidth–One Ceiling"

Bill (Qingyun) Sun · GPT5.6SOL · Fable — *for readers with general graduate-level mathematics, no analytic number theory assumed*

The headline result it explains. In 2026 an Anthropic research model proved — unconditionally, with a machine-checked Lean proof — that at least 67.2500...% of the zeros of the Riemann zeta function lie on the critical line and are simple. The Riemann Hypothesis (RH) says the answer is 100%; before this, the best unconditional result was 41.7%. The odd thing is the constant itself: 0.6725007... is *exactly* the number Montgomery and Taylor had derived in the 1970s – 90s **assuming RH**. Getting the same constant without the hypothesis looks like a coincidence. Our paper shows it is not.

The one-sentence thesis. The number $0.6725007\dots = 3/2 - (1/\sqrt{2}) \cdot \cot(1/\sqrt{2})$ is the exact optimal value of an *optimization problem about limited information* — and the new proof is, structurally, the optimal certificate of that optimization, which is why no proof using the same information can ever do better.

What "limited information" means. All progress on counting simple critical zeros flows through one measurement: the *pair correlation* of zeros — the statistics of gaps between pairs — and crucially, we only possess this measurement at "bandwidth ≤ 1 " (Fourier frequencies $|\alpha| \leq 1$). Beyond bandwidth 1, pair correlation encodes information about *pairs of primes* — twin-prime territory — which nobody can currently reach. So ask a clean question: among **all conceivable point processes** whose bandwidth-1 pair statistics match the zeta zeros, which one has the *fewest* simple points? That is a linear program. Its value is 0.6725007... A proof that only consumes bandwidth-1 pair data can never certify more, because an "adversary" process achieving exactly that value exists and is indistinguishable from the truth using that data. The 2026 theorem reaches the ceiling; our result is that it *is* a ceiling, and we locate the exit doors.

Where random matrices enter. Montgomery's famous conversation with Dyson revealed that zeta zeros statistically resemble eigenvalues of large random unitary matrices (the CUE). But there is a rival scenario — the *Alternative Hypothesis* (AH) — in which zero gaps live on the half-integer lattice $\frac{1}{2}, 1, \frac{3}{2}, 2, \dots$, wildly unlike random matrices, yet (a theorem of Lagarias – Rodgers) matching every band-limited correlation measurement of every order. To study this rivalry concretely, we work in a finite model: Tao's ACUE, a determinantal process choosing N points among the $2N$ -th roots of unity, and the space of all "mimickers" — probability laws matching its CUE-style moments. Three discoveries:

1. **The mimicker space is enormous** — dimension 0 for $N \leq 4$ (the moment problem has a unique solution!), then 2, 10, 80, 403, 1804 for $N = 5, \dots, 9$. Moment data is an exponentially thin window into what a point process can be.
2. **Yet it is rigid where it counts:** we *prove* every mimicker has identical two-point statistics at every distance (a Fourier aliasing argument special to finite groups), and verify the same for all three-point statistics. The first genuine freedom appears at **four points**. So if the zeta-zero process has a hidden impostor — the Lagarias – Rodgers uniqueness

question — the impostor differs from AH only in 4-point (and higher) patterns. That is a new, sharp localization of a well-known open problem.

3. **AH is not actually the enemy.** Folklore says the $2/3$ -ish barrier exists "because AH might be true." We prove AH-type adversaries are strictly *suboptimal*: they only force the bound $1/2 + 2/\pi^2 \approx 0.7026$, while the true extremal adversary — a non-lattice configuration with near-integer spacings — forces exactly 0.6725 . The barrier is about information, not about the half-integer lattice.

What is mathematically new. (i) A *proved* rigidity theorem for ACUE moment fibers (2-point, via aliasing) and the 3-point/4-point dichotomy; (ii) the identification of the rank – inertia – trace argument in the 2026 proof with linear-programming duality, the Montgomery – Taylor window appearing as the dual optimal solution — turning a virtuoso inequality into the canonical certificate of a named optimization problem; (iii) a clean separation, new in this subject, between *positivity* and *data*: adding sum-of-squares/Lasserre-type positivity constraints of any order does not move the ceiling (the adversaries are genuine probability measures, so they pass every positivity test), while adding genuinely new correlation data slices it immediately.

Why it matters for RH. It converts "try to push the constant" into a map. To prove more than 67.25% one must add *data*, not cleverness: either (a) any unconditional foothold on pair correlation beyond bandwidth 1 — equivalent to progress on prime pairs, with the LP machinery converting any ϵ -sliver directly into an improved percentage; or (b) arithmetic positivity that is not a function of pair correlation at all (Bombieri's Weil-quadratic-form program); or (c) in the finite model, four-point information — the exact order where the impostor space opens up. Meanwhile the same framework gives the random-matrix side a concrete finite target: if the mimicker gap $\delta^*(N)$ (already computed to be increasing: $0.0102 \rightarrow 0.0151$ for $N = 6 \cdots 9$) stays bounded away from zero, Lagarias – Rodgers uniqueness *fails*, and the zeta zeros' statistics would admit a genuine impostor visible only at four points.

The takeaway. $0.6725007 \cdots$ is not a scoreboard entry. It is the precise price of knowing only pair correlations at bandwidth one — a Chebyshev-type extremal value with a closed formula, a dual certificate, and an explicit adversary. Barriers of this kind are valuable exactly because they say where the next theorem *cannot* come from — and therefore where it must.

Historical source 38: README

Source: `historical/riemann-rmt/README.md`. SHA-256:

`25ce7119647fc7b9311ed71dd2047cf2a5ca47e458690f1e27a0217d3742c269`.

Riemann zeta / random-matrix / prime-gap research artifacts

Session outputs (Bill (Qingyun) Sun, GPT5.6SOL, Fable — August 2026):

- `rmt_zeta_survey.md` — The Bandwidth-One Ceiling (survey, round 1)
- `rmt_zeta_popular.md` — one-page companion
- `final_verified_paper.md` — Finite Spectral Certificates and Their Fibers (verified synthesis of the Codex series)
- `tao_ah_notes.pdf` — notes on Tao's Alternative Hypothesis problem
- `round3_synthesis.md` — ten directions coupling the two-thirds method to the finite fiber program
- `prime_gap_survey.md` — The Walls and the Doors (bounded prime gaps round)
- `verify_codex.py` — independent verification battery
- `joint_context_v2.md`, `prime_gap_context.md` — shared agent contexts

Full computational scripts and certificates remain in the session scratchpad archive.

Historical source 39: HANDOFF_GPT6_ASTR

Source: `historical/riemann-rmt/handoff/HANDOFF_GPT6_ASTR.md`. SHA-256:
`34950171df1d33390c98ad052be08ad0f7bac881ac409d9afc257e5e048e1cbc`.

HANDOFF — The Riemann / Random–Matrix / Prime–Gap Programme

Complete state of the research, for GPT6 Astra to take over

Authors of the programme: Bill (Qingyun) Sun · GPT5.6SOL · Fable Handoff written: 5 September 2026, by Fable (Claude), at the end of a ~3-week multi-agent session Repository: `galpha-ai/Alpha-devbox`, branch `claude/riemann-zeta-random-matrix-udxp3f`, directory `research/riemann-rmt/` (PR #11, draft). Standalone companion package: `research/riemann-rmt/riemann-impostors/`.

0. How to read this document

This is written for a successor system that will attempt a *major* theorem. It is therefore long, and deliberately so: it records not only what was proved and computed, but what was tried and failed, what the failures taught, and where every number came from. Nothing here is padded; if a section seems long it is because the corresponding wall took weeks to map.

Status tags, used throughout, with strict meanings.

tag	meaning
[P]	proved, with a proof written down in the cited file; elementary steps machine-checked where noted
[X]	exact: certified in exact rational / ball / cyclotomic arithmetic, no floating point in the chain
[C]	computed: verified numerically, usually by two independent implementations, not proved
[R]	reduced: follows from a <i>cited published theorem</i> plus an explicitly stated open lemma
[H]	heuristic / conjecture, with the evidence stated
[X]	tried and refuted; the refutation is itself a result
[U]	unverified claim received from a collaborating system; recorded, not endorsed

Trust rules that we learned the hard way. (1) Every headline number in this programme was produced twice by independent code before being believed; where the two disagreed, the discrepancy was tracked to source before any claim. (2) Nothing near a decision boundary is accepted in floating point. (3) The default assumption for any "improvement" is that it is a

misread constraint; five of the most exciting-looking phenomena in this programme were exactly that. (4) "Machine-checked" and "formalised" are different things; only the first is true of most of this work.

What this programme did NOT do, so that the successor does not inherit a false belief. We did not prove the Riemann hypothesis, the Alternative Hypothesis or its negation, the pair correlation conjecture, the density conjecture, the twin prime conjecture, Lagarias – Rodgers $\mu = 1/2$, or any improvement of $H_1 \leq 246$. We did not compile any Lean. We could not read three URLs the human asked us to incorporate at the very end (Anthropic's post on formalising Fermat's Last Theorem and two `t.co` links): the session's egress proxy blocked `www.anthropic.com` and `t.co`. **Astra should read those three sources first**; nothing in this document reflects their content.

A note on a number the human mentioned. The handoff request refers to a prime-gap bound of "186" attributed to Weijie Su and GPT6 Astra (and, once, "286"). We have no record of either. Two facts worth knowing before touching it: $H(40) = 186$ exactly in the Engelsma table of minimal admissible-tuple diameters (independently re-proved minimal by us for $k \leq 62$), so a bound of 186 would mean $DHL(40, 2)$; and 286 is not the diameter of any minimal admissible tuple ($H(57) = 282$, $H(58) = 288$), so "286" is presumably a typo for 186 or 246. What $DHL(40,2)$ would require is analysed in § 9.5.

1. Timeline: what each round asked and what it delivered

The programme ran in rounds. Each round posed one question to a fleet of 3 – 10 parallel agents with a shared written context (state of knowledge, known errors, honesty rules) but no shared conclusions; convergence was treated as evidence and divergence as a bug report.

round	question	main deliverables
1	Analyse Anthropic's 2026 unconditional theorem " $\geq 67.25\%$ of zeta zeros are simple and on the line" (Lean repo <code>anthropics/zeta-23-lean</code> , constant $\delta_{MT} = 3/2 - (1/\sqrt{2})\cot(1/\sqrt{2}) = 0.672500703679$) from random-matrix angles; connect to Tao's ACUE	RMT survey; popular note; identification of the bandwidth-one wall; first finite ACUE fibre computations
2	Verify five Codex (GPT5.6SOL-ULTRA) PDFs; extract real mathematics	<code>verify_codex.py</code> battery; <code>final_verified_paper.md</code> ; several Codex claims corrected (e.g. transcription errors $\theta_n = \arccos(1 - n^2/a_n)$; free-box requirement for the Wilson trace); our own overclaim retracted ("0.6725 is the exact bandwidth-one LP optimum")
3	Read Anthropic's method <i>at the Lean source</i> , couple it to the finite fibre machinery, ten directions	<code>round3_synthesis.md</code> : seven-term deficiency identity + $\sqrt{\varepsilon}$ rigidity; $N_d \geq 0.8362503N + p$; flat RH-frontier; edge no-go; weight freezing; centre-of-mass mimicker family; $M_\bullet$ lemma with quantified payoff; ceiling candidate 15/22; kill-degree $d^* = 3(N-3)$
4	Attack the bounded-gap record $H_1 = 246$ with ≥ 10 agents; improve <i>any</i> record concretely	Five proved walls; the $k = 49$ door priced; new records H_2 , H_3 , H_4 via the layer-cake engine; <code>prime_gap_survey.md</code> ("The Walls and the Doors"); Codex handoff note

round	question	main deliverables
5	Zhang Yitang's directive: unify Maynard – Tao + Zhang's signed construction + ACUE certificates, run the signed-sieve phase experiment	Signed no-gain theorem (one-line identity); exact critical price β^* ; conditional price list with (E_θ) ; Chen-switch obstruction; Bourgain scale diagnosis of the 0.0301 gap
6 (P0)	The de Bruijn – Newman finite depth Λ as a dynamic anti-ACUE observable	Exact ACUE enumeration $N \leq 10$; CUE Monte Carlo $N \leq 256$; $-\Lambda^{\text{ACUE}} \asymp N^{\{-8/3\}}$ vs $-\Lambda^{\text{ACUE}} \asymp N^{\{-2\}}$; β -universality; $s^* = 1.419640342$; the operator unification $\mathcal{L}_N = \text{Jacobian at the clock}$; marked-depth law; Theorem A ($\rho \geq 1$) proved ; Theorem C separation; LR bridge and the $\pi^2/8$ threshold
7	Consolidation: papers, machine verification (18/18 sympy + z3), Lean specification, companion repository, this handoff	<code>impostors_paper.md</code> , <code>depth_scaling_theorem.md</code> , <code>riemann-impostors/</code>

Two things happened operationally that the successor should plan around. Agents were killed twice mid-run: once by a session usage limit, once by exhaustion of usage credits (six agents at once). In both cases the scratchpad survived and partial results were harvested from disk; several of the best results in this document (the exact β^* , the FUSION-ACUE negative finding in § 3.6) were recovered that way. **Design every long computation so that its partial output is on disk in a form that a stranger can resume.**

2. The zeta side: Anthropic's 2/3 theorem, read at the source, and what we added

2.1 The method as it actually is (from the Lean, not from paraphrase)

The theorem: at least $\delta_{\text{MT}} = 3/2 - (1/\sqrt{2}) \cdot \cot(1/\sqrt{2}) = 0.672500703679\dots$ of the nontrivial zeros of ζ are simple and lie on the critical line, unconditionally. The Lean repository `anthropics/zeta-23-lean` (key files `Zeta23/ZeroSide/RankTraceMult.lean`, `TightMult.lean`, `Mult.lean`, `PairCeiling/*.lean`, `XiPrime/`, `Taper/`, `Poisson.lean`, `MV/`, `Tail/`) organises the proof around:

- **Lemma R (rank – trace, k-form)**. For Hermitian P, Q and $c > 0$: $2c \cdot \text{tr}(P+Q) - \|P+Q\|_F^2 \leq \sum_j k_c(m_j) + c^2 \cdot b$, with $k_c(p) = c^2 - ((c-p)_+)^2$. This converts a quadratic form built from zero data into a *count* of positive directions. It is tight (`lemmaR_tight`).
- **PairCeiling**. An upper constant $0.6818287 + 2.55 \cdot 10^{-6} \cdot R$ under the hypothesis `Enc10K`; this is the "bandwidth-one ceiling" of the method's own data.
- ξ' **quartic windows and TightMult**: the Montgomery – Taylor-type window optimisation, done with Fejér – Riesz-style nonnegativity at bandwidth one.

The core structural fact, which every later result respects: the method consumes exactly the pair-correlation data of the zeros at Fourier bandwidth ≤ 1 (Montgomery's proven range), and its constant is set by a Rayleigh quotient in that data.

2.2 What we proved about the method (Round 3, all in `round3_synthesis.md`)

[P] **Seven-term deficiency identity and $\sqrt{\epsilon}$ rigidity (D4).** Lemma R is an identity up to an explicit seven-term deficiency (verified to $7 \cdot 10^{-14}$, identically zero on the TightMult class); an ϵ -near-equality configuration lies within Frobenius distance $\sqrt{\epsilon}$ of an exact equality configuration (Löwdin frame). Zeta-side ledger: total deficiency over MT windows $\leq (q^*-1) \cdot N(T) = 0.3275 \cdot N(T)$, giving an unconditional pair-repulsion-type bound (on-line pairs with taper overlap $\geq \tau$ number $\leq 0.1637 \cdot N(T)/\tau^2$). Reading: Montgomery – Taylor optimisation *is* minimisation of the RMT-predicted deficiency, and under GUE the budget is exhausted by the Schur/von Neumann term alone.

[P] **The distinct-zeros upgrade $N_d \geq 0.8362503 \cdot N + p$ (D7).** Combining the $c = 2$ inequality with the zero-count identity once (not twice) gives slope $+1$ in the off-line pair density — strictly stronger than the repository's own $c = 3$ decoding (slope $1/2$; pairs mispriced at $c^2 = 9$ against the forced value $k_3(2) = 8$). This is an unformalised **five-line Lean edit** to `mult_two_pair` and would yield a strictly stronger published constant for distinct zeros. Nobody has made the edit.

[P] **The RH-frontier is exactly flat.** $\delta(\pi) \equiv 0.6725007$ for all off-line proportions $\pi \in [0, 0.16375]$: **assuming RH buys nothing within this method** (the binding adversary is on-line double zeros; exact 4p-budget/mass cancellation). At maximal off-line density all zeros are forced simple; pair-block spectra are forced to $(\approx 2, \approx 0)$.

[P] **The edge no-go (D5).** Every admissible bandwidth-one window has $\hat{r}(\pm 1) = 0$ (two-line Fejér – Riesz), so the weakest pointwise edge hypothesis $|F(1) - 1| \leq \epsilon$ certifies exactly 0.6725 even at $\epsilon = 0$. The finite Nyquist conservation law is powered by lattice aliasing that $\mathbb{R}$ lacks. The correct zeta-side edge object is the Cesàro mean $\bar{F}(\infty, T)$, with the exact collision identity $C(T) = N^* \cdot \bar{F}(\infty, T) - N(T)$.

[P] **Weight freezing.** All balanced moments of degree $\leq N$ are frozen across the mimicker fibre — later (Round 6) shown to be frozen along the *entire* heat flow, because the flow is diagonal on coefficients.

[C] **Φ_3 at the MT window.** The flat-window sine-Gram data sits on the $\{0,1,2\}$ factorial face ($m_3 - 3m_2 + 2m_1 = 0$, an identity at every finite N for ACUE), but at the Montgomery – Taylor window itself $\Phi_3 = m_3 - 3m_2 + 2m_1 = -0.0117753128$ (closed trig form; verified symbolically, at 30 digits, and on the lattice to $3 \cdot 10^{-7}$). The third moment carries information exactly where the proof lives; the blocker is not arithmetic but that Lemma R's hypothesis class carries **no bound on the magnitude of the compression's negative eigenvalues**.

[X] **Window engineering cannot fix it.** The Rudnick – Sarnak – style clump term ($\log T$ zeros per unit ordinate at zero kernel decay) forces the row-sum cap $M = A_0 \cdot \log T$ for *every* window; edge-vanishing windows improve only the negligible far field, cost ≥ 0.006 in $2 - q(w)$, and flip Φ_3 to the unmonetisable positive side. Adversarially settled.

The surviving zeta-side target, precisely (D1). A lemma $\|(c^{-1}\hat{A})_-\| \leq M_-$ **uniform in T** — depth-uniform control of off-line evaluation vectors (a cosh-weighted Poisson analogue plus an overlap-angle bound). Payoff at the MT window with the capped cubic decoder: $\delta = 0.6796896$ if $M_- = 2$, up to 0.6844924 if $M_- \leq 1$, with tr^3 the only new prime-side input. This remains the sharpest known formulation of "what would improve 0.6725 without new correlation ranges."

[H] **The true bandwidth-one ceiling.** Bracket $[0.6725007, 0.6818287+]$, measured integrality gap ≈ 0.0093 , candidate exact value $15/22 = 0.681818\ldots$. Two exact-arithmetic conjectures were left machine-attackable: the 15/22 identity and the kill-degree law $d^* = 3(N-3)$ (verified $N = 5, 6, 7$).

[P forms + C] **The Nyquist cotangent tower.** The three N^{-2} convergence towers seen in different places are one object — the Bernoulli tower of $(\cot, \csc^2)$ — and exist only at the edge. Doubles density of any Nyquist equality law $\rightarrow 1/4 - 1/\pi^2 = 0.148679$.

2.3 The 0.702642 vs 0.672500 anomaly, closed (Round 5, agent BRG–SCALE, files `brg_*.py`)

The free finite optimum of the marked-ACUE pencil is $\delta_{\text{free}}(N) = \frac{1}{2} + \csc^2(\pi/2N)/(2N^2) \rightarrow \frac{1}{2} + 2/\pi^2 = 0.70264237$, with exact saddle $\zeta^* = N^2 + N - \csc^2(\pi/2N)$; the Anthropic-feasible restriction gives $\delta_{\text{MT}} = 0.672500704$; gap 0.03014166.

[C] **The gap is multiscale in Fourier but single-site in position.** Band ledger (stable from $N = 128$ to 1024): DC atom +20.33%, low band ($0 < \alpha \leq \frac{1}{4}$) −12.73%, mid +20.20%, edge quarter +72.21%, Nyquist row **exactly 0.00%**; the density $\lambda(\alpha)$ vanishes linearly at $\alpha = 1$ — no boundary layer, so no single sharp frequency inequality closes it. The whole free advantage is certified by **one lattice inequality** $S_1 \geq 0$ — nonnegativity of the pair count at displacement $x = \pm 1$, i.e. at half the mean zero spacing — whose bandwidth-one shadow $\hat{c}_k = \cos^2(\pi k/2N)$ is full-band coherent. Decoupling-type bounds misprice the Rayleigh value by $15 - 77\times$ (constructive alignment; both extremisers are Perron vectors of a positive kernel), against a 4.5% target: **ℓ^2 -decoupling is structurally the wrong tool here.**

[C] **Calibration that changes what one should chase.** Of the 0.030142, only ≈ 0.0093 is even in principle recoverable in the continuum (to the ceiling ≈ 0.68183); the remaining ≈ 0.0208 is pure lattice aliasing. "0.6725 $\rightarrow$ 0.7026" via any continuum inequality is hopeless; the real continuum target is "0.6725 $\rightarrow$ 0.6818" and it needs integrality-aware (point-evaluation / measure-valued) certificates — the $\mathbb{R}$ -analogue of $S_1 \geq 0$ at half mean spacing.

2.4 The moment–deflation audit (Round 5, `fab_deflation.py`) — relevant to Guth–Maynard

Guth – Maynard bound the top singular value of the Dirichlet-polynomial Gram matrix using $\text{tr } B$ and a centred third moment, noting that estimates for powers ≥ 4 are unavailable. The exact extremal problem "given $(p_1, \dots, p_r)$ and m nonnegative eigenvalues, bound λ_1 " — sharp bound $\lambda_1 \leq \frac{r}{m-1} + (p_1 - \lambda_1)^r / (m-1)^{r-1} \leq p_r$ — shows that for a spike-plus-flat-bulk spectrum the **$r = 3$ bound is already exactly tight and $r = 4, 6, 8, 10$ gain 0.00%** ($m = 10^2, 10^3, 10^4$; spike fractions $0.5 \cdots 0.05$); with a secondary cluster the gain is 0.1 – 1.1%. Whatever room exists in Guth – Maynard is not in the moment hierarchy alone but in the additive-energy input and the resonator geometry. A formula received from Codex ($e = \eta + \kappa, 240\eta^2 + (72 - 10e)\eta - 13e = 0, A(e) = 30/(13 + 10\eta e)$) is recorded as [U]: we never re-derived their argument.

3. Tao's Alternative Hypothesis and the ACUE impostor programme

3.1 Why AH matters and why it resists

Montgomery's pair-correlation conjecture: normalised zero gaps follow the GUE law. The Alternative Hypothesis (AH): the gaps lie asymptotically in $(1/2)\mathbb{Z}$. AH is consistent with every proven zero statistic, has arithmetic consequences (class numbers; Landau – Siegel), and is precisely what current technique cannot exclude. Tao's ACUE — the measure on N -subsets C of the $2N$ -th roots of unity with $\mu_{\text{ACUE}}(C) = |\Delta(\zeta^C)|^2 / (2N)^N$ — reproduces CUE pair correlation exactly and is a finite model of AH. The obstruction to "find the statistic ACUE cannot fake" is that one must rule out a *fibre*, not a point: the measures matching any finite moment list form a large convex body.

3.2 The fibre, measured exactly [X]

Affine dimension (rotation-orbit level) of $\{\text{measures on the ACUE support matching every balanced moment } E[p_\lambda \bar{p}_\nu] \text{ of degree } \leq N\}$:

N	3	4	5	6	7	8	9
dim	0	0	2	10	80	403	1804

($N = 6$ and 8 re-confirmed by the tomography code: ranks 70 and 407 with spectral gaps $> 10^6$.) Structural facts: 2- and 3-point statistics are rigid, 4-point free ($N = 5$ onward); the fibre at $N = 5$ is completely solved — a pentagon over $\mathbb{Q}(\sqrt{5})$, the known mimicker being $|v|^2$ for a 5-sparse coherent superposition of translated Fermi seas with the single syndrome $a^*Sa = 0$. The reflection symmetry of the functional equation kills the entire chiral half of the fibre ($39/80$ dims at $N = 7$) at zero cost and provably cannot do more. At $N = 8$, of 403 directions, **401 are invisible to every pattern count of window width $\leq 2,397$ to $\leq 3,396$ to $\leq 4,392$ to $\leq 5,383$ to ≤ 6 .**

3.3 Explicit families [X]

Centre-of-mass family (`d6_allN_family.py`, `mimicker_fibre.py`): with $X(C) = \sum c \bmod N$, $q_g(C) = \mu_{ACUE}(C) \cdot g(X(C))$ matches every balanced moment of degree $\leq N$ iff $E[g] = 1$ and $\hat{g}(\pm 1) = 0$; all other frequencies free. Hence an $(N-3)$ -parameter family of honest probability laws for every $N \geq 4$: null dims $2, 3, 4, 5$ at $N = 5, 6, 7, 8$ (forced-zero frequency $j = \pm 1$; free $j = 2 / 2,3 / 2,3 / 2,3,4$), worst moment error $\leq 1.5 \cdot 10^{-11}$, positivity radius $1.0 - 2.2$. Mechanism: $X \bmod N$ is exactly uniform under ACUE and couplings to com-frequencies $2..N-2$ vanish by a transport-cost bound (hop budget $2N$ vs cost $\geq 2N+2$).

Determinant-character (secant) family. $D_r(C) = \prod \{c \in C\} \zeta^{rc} = \det(U_C)^r$. Tilts $q = \mu \cdot [1 + c \cdot \text{Re}(\eta D_r(C))]$ lie in the fibre. Exact quantum-mechanical description (from the colleague's message, verified consistent with our data): $\mu_{ACUE}(C) = |\langle C | (\wedge^N F) A_0 \rangle|^2$ is the Born distribution of a single Slater determinant ($F = DFT$, $A_0 = e_0 \wedge \dots \wedge e_{N-1}$) — re-proving ACUE is a projection DPP; $\langle C | Q_r \rangle = e_r D_r(C) \langle C | Q_0 \rangle$; the tilt with $c = 2a/(1+a^2)$ is the interference pattern $(|Q_0 \rangle + a|Q_r \rangle)/\sqrt{1+a^2}$ — our $c = 1/2$ corresponds to $a = 2 - \sqrt{3}$. In Plücker coordinates ACUE is a point of $\text{Gr}(N, 2N)$; the mimicker lies on $\sigma_2(\text{Gr}(N, 2N)) \setminus \text{Gr}$: the minimal non-Gaussian deformation of a fermionic Gaussian state.

Parity sectors (from the uploaded "Finite All-Depth Escape" paper; tested by us): for even $N \geq 6$, $q_N^{\pm} = \mu_N(1 \pm (-1)^{\sum c})$ are mutually singular, match balanced trace moments of degree $d < N^2/4$, agree on every complete binary marginal on $\leq N-1$ sites, yet have the same local limit. That paper also proves pair rigidity for simple half-lattice processes (first spatial escape is three-body) and constructs, for every $s \geq 3$ and large N , honest laws matching bandwidth-one rows through order s while the three-site decoder $P(0,1,2 \in C)$ stays nonzero; the missing piece for a second infinite LR mimicker is one uniform positivity ratio $\Gamma^*_{N,s}$.

3.4 The autocorrelation—Bruhat structure [P (identities) + X]

Restricting to even shifts, $|Q_s\rangle := |Q_{2s}\rangle$, $z(C) := D_2(C)$ with $z^N = 1$, so the entire secant family collapses to a polynomial $P_a(z) = \sum a_s z^s$ on $\mathbb{Z}_N$:

$$q_a(C) = \mu(C) \cdot |P_a(z(C))|^2 = \mu(C) \cdot (1 + \sum_{\delta \neq 0} A_\delta z(C)^\delta), A_\delta = \sum_s a_{s+\delta} \bar{a}_s.$$

The code object is the autocorrelation spectrum, not the support. The invisibility depth is $d_{\text{vis}}(a) = \min\{\delta(N-\delta) : \delta \neq 0, A_\delta \neq 0\}$, because the sector pairing $B_\delta(f,g) = E_\mu[D_{2\delta} f \bar{g}]$ vanishes for $\deg < \delta(N-\delta)$ and opens a rank-one channel at exactly that degree between the conjugate rectangles $(N-\delta)^\delta \rightarrow \delta^{N-\delta}$. So q_a agrees with ACUE on all balanced observables of degree $\leq d$ iff $A_\delta = 0$ for all δ with $\delta(N-\delta) \leq d$. The critical detector $\Phi_{N,s}(U) = \det(U)^{2s} \cdot s_{(N-s)s}(U) \cdot \text{conj } s_{s(N-s)}(U)$ is the character of $W_{N,s} = V_R \otimes V_S^* \vee \otimes \det^{2s}$, has $E_{ACUE} \Phi = (-1)^{s(N-s)}$, and $E_{CUE} \Phi = 0$ because $W^{\wedge} \{U(N)\} = 0$. The pure max – min code problem is solved by pigeonhole: L

states on the N -cycle give at best $m(N-m)$, $m = \lfloor N/L \rfloor$, maximised at $L = 2$ with $\lfloor N^2/4 \rfloor$ — **two Fermi seas are globally optimal; do not spend compute on five**. The open problem is the *interference* code: A positive-definite on $\mathbb{Z}_N$, $A_0 = 1$, $A_\delta = 0$ on a low-Bruhat-energy zone, $A_{\{\delta^*\}} \neq 0$ as high as possible (a Delsarte LP; "zero-Bruhat-correlation-zone sequences"; the CAZAC extreme $A \equiv 0$ gives $q = \mu$, no impostor).

$d_N(\delta) = \delta(N-\delta)$ is simultaneously $\dim_{\mathbb{C}} \text{Gr}(\delta, N)$, the pairing $\langle 2\rho, \omega_\delta \rangle$ (hence the affine Bruhat length $\ell(t_{\{\omega_\delta\}})$), $N \cdot \|\omega_\delta\|^2$, and (up to convention) the quadratic Casimir $(N+1)/N \cdot \delta(N-\delta)$. The fourth identification — with the *dynamics* — is § 4.7 and is the structural centre of the programme.

3.5 Monodromy separator: real Frobenius families are det-sector dark [C, exact averages]

Agent L3 (`monsep_*.py`; artifact "Monodromy Separator"). Over $F_q[T]$ Dirichlet families ($N = 4: q \in \{3, \dots, 23\}$; $N = 5: q \leq 13$; $N = 6: q \leq 11$; $\sim 40M$ characters, exact FFT character sums): every winding $\neq 0$ observable satisfies $|E_\chi| \leq 1.7 \cdot q^{-1/2}$, with observed decay slopes -2 to -5 (near the family-noise floor $q^{-(N+1)/2}$). ACUE is bright on the same detectors: $E_\mu[p_N^2] = N$, $E_\mu[p_N^2 | p_1^2] = N-2$, $E_\mu[\det^4 | p_{N-1} p_{N-3}^2] = 1$; the rank-4 twist mimicker has $E[\det^4] = 1/2$ ($N=4$), $1/4$ ($N=5,6$). The weight- $(N+1)$ Gram is exactly Haar -1 in every entry (Haar closed form $E[p_\lambda \bar{p}_\nu] = z_\lambda \delta - \text{sgn} \cdot \text{sgn}[P]$). Families converge to Haar with a $+c/q$ correction — the *opposite* side from ACUE's -1 . $\dim(V_\lambda \otimes V_{-\mu} \vee \otimes \det^r)^{\{U(N)\}} = 0$ for all detectors $[P]$; $\det$ is not pinned (per-class $E[\det] = 0$). Certificate: $|E_{\text{family}}| < \text{bright}/2$ at every computed q except $\det^4 | p_4 p_2^2$ at $(N,q) = (5,3)$. **Sign correction discovered:** $\sum \{\deg f = m\} \Lambda(f) \chi(f) = ** - ** q^{m/2} \text{tr}(\Theta^m)$. What remains for the colleague's Conjecture G2 (minimal derived-L separator) is to convert $W^{\{G_{\text{geom}}\}} = 0$ into a uniform $O(q^{-1/2})$ trace estimate via Katz's big-monodromy theorem (IMRN 2013), and to extend the bright gap to the 7 non-twist fibre directions at $N = 6$.

3.6 FUSION-ACUE: a negative result recovered from a dead agent's logs [X, exact]

Agent L1 (`fus_*.py`, killed by credit exhaustion; logs harvested). **Task 1 verified:** the ACUE matrix coefficients $M^{\{r\}}_{\{\lambda, \mu\}}$ equal $\det$ of a 0/1 congruence matrix with phase $+1$ and $Z_N = N!$ ($N = 2,3,4$ exhaustive; $N = 5$ spot-checked, 0 mismatches), all values in $\{-1, 0, +1\}$. **Verlinde S-matrix = Kac – Walton at $\mathfrak{sl}(N)$ level N verified** ($N = 2, 3, 4$: $18/550/2100$ triple coefficients, 0 mismatches). **But the identification "KW fusion = $\sum_t$ ACUE triple pairings (charge decomposition)" FAILS at $N = 3$: 22 mismatches out of 550** (it holds at $N = 2$). Examples: KW $(1,1) \otimes (3,1) \rightarrow ()$ gives 0, ACUE gives 1; KW $(2,1) \otimes (2,1) \rightarrow ()$ gives 1, ACUE gives 2. So the naive form of the colleague's FUSION-ACUE Conjecture A is false; the correct statement, if any, needs a different charge bookkeeping. **This was never reported to the human before this document.**

4. The finite de Bruijn–Newman depth: the P0 programme

4.1 Definition and conventions (get these right; two earlier write-ups had sign slips)

For monic $P(z) = \prod (z - e^{i\theta_j}) = \sum a_j z^j$ define $P_s(z) = \sum a_j e^{s \cdot j(N-j)} z^j$ for $s \geq 0$ and the **depth** $D(P) = \inf\{s > 0 : \text{disc}(P_s) = 0\}$. $D = -\Lambda$. P_s stays self-inversive, so D is the first collision time. **Lemma 1 [P, machine-checked]:** the zeros obey $\theta_j = -\sum_{k \neq j} \cot((\theta_j - \theta_k)/2)$ (attracting circular Coulomb gas; derived from $\partial_s P = (N D_z - D_z^2)P$ and $2z_j/(z_j - z_k) = 1 - i \cdot \cot((\theta_j - \theta_k)/2)$, with the $(N-1)$ terms cancelling exactly). $N = 2$: $D = -\log \cos(\delta/2)$. Clocks ($P = 1 - cz^N$) are flow-invariant, $D = \infty$.

4.2 ACUE, exact enumeration $N = 3..10$ [X] (`dyn1_*`, data `dyn1_results_N*.npz`)

13,132 rotation orbits (184,756 configurations at $N = 10$), exact Vandermonde masses; 40-digit spot checks worst

$1.3 \cdot 10^{-9}$. Results:

- $P(\text{clock}) = 2^{\{1-N\}}$ exactly [P] (Cauchy – Binet: $\sum_{|C|=N} |\Delta|^2 = \det(AA^*) = (2N)^N$; each clock has $|\Delta|^2 = N^N$).
- $\delta_{\min} = \pi/N$ for every non-clock configuration [P] (pigeonhole).
- $-\Lambda \asymp N^{-2}$, fitted slope -2.0009 ; $N^2(-\Lambda) \in [1.31, 1.99]$.
- First collision always at an initially-adjacent pair: 13,130/13,130.
- Generator identity $L p_m = -m(N-m)p_m - m \sum_{a < m} p_a p_{m-a}$ (verified $4.1 \cdot 10^{-28}$).
- Quantiles of $N^2(-\Lambda)$ (min / median / max): $N=6$: 1.353146 / 1.419374 / 1.952629; $N=8$: 1.330383 / 1.418216 / 1.976122; $N=10$: 1.314614 / 1.412774 / 1.985458.
- The median turns around at $N = 7$ (1.41474, 1.41822, 1.41908, 1.41937, 1.41950, then 1.41822, 1.41520, 1.41277) — so the single-dislocation constant is NOT the ensemble median limit (a colleague's extrapolation from $N \leq 7$ was corrected).
- Minimising family $\{0..N-3\} \cup \{N+3, N+4\}$: $\rho = 1.0656$ ($N=10$) $\cdots$ 1.0495 ($N=18$), slowly decreasing toward $\approx 1.03 - 1.05$. Maximising symmetric 3-block: $\rho \rightarrow 1.6094$, $N^2(-\Lambda) \rightarrow 1.9861$.

4.3 CUE, Monte Carlo $N \leq 256$ [C] (`dyn2_*`, data `dyn2_data_N*.npz`)

Two independent Λ -solvers (ODE vs coefficient bisection) agree to 10^{-6} ; $N = 2$ law to $8.7 \cdot 10^{-13}$. $-\Lambda \asymp N^{-8/3}$: fitted -2.691 ± 0.010 ($N \geq 16$), -2.678 ± 0.016 ($N \geq 32$). Distributional law: $8N^{8/3}(-\Lambda) \Rightarrow G^2$, $P(G > x) = \exp(-x^3/72\pi)$ (72π derived from the sine kernel, not fitted: measured 229 – 236 vs 226.2; KS 0.035 – 0.041; median 29.6 – 30.5 vs $(72\pi \ln 2)^{2/3} = 29.08$). Local law $R = 8(-\Lambda)/\delta_{\min}^2 > 1$ with $\text{med}(R-1) = 0.598 \cdot N^{-0.729}$. Depth ratio to ACUE falls as $3.64 \cdot N^{-0.704}$; $10\times$ separation at $N^* \approx 165$ (predicted ≈ 160).

4.4 β -universality [C, three-point + endpoint] (`beta_*`)

Killip – Nenciu CMV sampling validated ($\beta=2$ vs stored CUE KS 0.0236; vs Haar spacing KS 0.0048; $\beta=1$ vs direct COE = VVT KS 0.0035). Prediction $-\Lambda \asymp N^{-2/(\beta+1)}$:

β	predicted	fitted	local slopes
1 (COE)	-3	-3.064	$-3.080, -3.026, -3.100$
2 (CUE)	$-8/3$	-2.678	$\rightarrow -2.627$ at $128 \rightarrow 256$
4 (CSE)	$-12/5$	-2.510	$-2.546, -2.480, -2.515$
∞ (ACUE)	-2	-2.0009	exact enumeration

Localisation (first collision = minimal initial gap) holds in 95 – 99% of $C\beta E$ samples and 100% of ACUE. Convergence rate prediction $\rho_{\beta-1} = O(N^{-2/(\beta+1)})$: fitted $-1.012 / -0.710 / -0.501$ vs $-1 / -2/3 / -2/5$, local slopes at largest N $-0.851 / -0.624 / -0.407$.

4.5 The single–dislocation constant [C, two independent routes]

Alternating clock minus $e^{-i\pi/N}$ plus 1: $P_N(z) = (z-1)(z^{N+1})/(z-e^{-i\pi/N})$, gap pattern 1,2,...,2,3. Lattice solver: $N^2(-\Lambda) = 1.41473945$ ($N=3$) ... **1.41963827** ($N=20$), monotone. Continuum limit $t = -s/N^2$, $z = e^{iu/N}$: $G_s(u) = 2\cos(u/2) - 2\pi \int_0^{1/2} e^{s(1/4-y^2)} \cos((\pi+u)y) dy$, first double zero at $s = 1.419640342$, $u = 1.812942145^{**}$, $G''(u^*) = -0.3767$. Agreement $2 \cdot 10^{-6}$ at $N = 20$ ($O(N^2)$ approach). $s = \rho_\infty \cdot \pi^2/8$ with $\rho_\infty = 1.150717118022^*$ (corrected from an earlier misprint 1.1507015): a two-body collision time dressed by $\approx 15\%$ many-body shielding.

4.6 Theorems A, B, C and the Lagarias–Rodgers bridge [P / R] (`depth_scaling_theorem.md`)

Theorem A [P]. For an adjacent pair, $g' = -2\cot(g/2) - \sum_{k \neq a,b} [\cot(x_a^k/2) - \cot(x_b^k/2)]$ with $x_j^k = (\theta_j - \theta_k) \bmod 2\pi$; adjacency forces $x_a^k = x_b^k + g$ in $(0, 2\pi)$ and $\cot(\cdot/2)$ is strictly decreasing there, so every bracket is negative and enters with a plus sign: **every other zero slows the collapse**. Hence $g' \geq -2\cot(g/2)$, and since order is preserved until collision, $-\Lambda \geq -\log \cos(\delta_{\min}/2) \geq \delta_{\min}^2/8$ ($\rho \geq 1$). Checked: 11,060/11,060 background terms positive; ratio minima 1.0842 / 1.0714 / 1.0612 (ACUE $N = 6, 8, 10$), 1.00019 / 1.00021 (CUE $N = 16, 64$). Exact two-body solution: $\cos(g(s)/2) = e^s \cos(g_0/2)$; and $-\log \cos(x/2) \geq x^2/8$ on $[0, \pi]$.

Theorem B [P given hypothesis]. With background stiffness $S = \sum_k \frac{1}{2} \csc^2(x_b^k/2)$ (clock value $(N^2-1)/6$ exactly), if $S \leq AN^2$ then $-\Lambda \leq (\delta_{\min}^2/8)(1 + O(AN^2\delta_{\min}^2))$. Measured S/N^2 on CUE: median 0.109 / 0.120 / 0.117 / 0.120 / 0.120 ($N = 8 \dots 128$), $q_{99} \leq 0.36$, max 0.79. **The high-probability bound $S \leq AN^2$ is the single open analytic ingredient**; it closes Law 1, Law 3 and Theorem C(ii) at once. Because $N^2\delta_{\min}^2 \asymp N^{-2/(\beta+1)} \rightarrow 0$ for finite β but $= \pi^2$ at $\beta = \infty$, the lattice is a *singular* endpoint where $\rho \neq 1$.

Theorem C. (i) [P] $N^2(-\Lambda) \geq \pi^2/8 = 1.2337005501$ for every non-clock ACUE configuration. (ii) [R] $P(N^2(-\Lambda) < \pi^2/8) \rightarrow 1$, from $P(\delta_{\min} > \pi/N) \approx \exp(-\pi^2 N/72)$ plus the S-bound. Measured fraction below the floor: 0.565, 0.823, 0.967, 0.9984, 1.0000, 1.0000 ($N = 8 \dots 256$); sample maxima 0.911 ($N=128$), 0.437 ($N=256$).

LR bridge [P modulo the reverse comparison]. A hard core of c mean spacings gives $\delta_{\min} = 2\pi c/N$, so $N^2(-\Lambda) = \rho \cdot \pi^2 c^2/2$. Hence $\mu_\Lambda \geq \pi^2 \mu^2/2$ and $\mu \leq \sqrt{(2\mu_\Lambda)/\pi}$: Lagarias – Rodgers' $\mu \leq 0.606894$ is exactly $\mu_\Lambda \leq 1.8177$; a depth bound 1.29 would give $\mu \leq 0.511281$; the AH barrier is $M < \pi^2/8$. **Falsifiable threshold: $\liminf N^2(-\Lambda_{\zeta, \text{local}}) < \pi^2/8 \Rightarrow \text{AH false}$** .

4.7 The operator unification [P + X] (`fab_half_laplacian.py`)

$\sum_{\delta=0}^{N-1} \delta(N-\delta) e^{-2\pi i k \delta/N} = -N/(2\sin^2(\pi k/N))$ for $k \neq 0$, so $(\mathcal{L}_N f)(x) = \sum_{k=1}^{N-1} [f(x) - f(x+k)]/(2\sin^2(\pi k/N))$ has $\mathcal{L}_N e_\delta = \delta(N-\delta)e_\delta$. Linearising the Coulomb flow at the clock, $\dot{e}_j = \sum_{k \neq j} (\epsilon_j - \epsilon_k)/(2\sin^2(\pi(j-k)/N)) = (\mathcal{L}_N \epsilon)_j$ — $\|\mathcal{L}_N - \text{Jacobian}\| \leq 2.4 \cdot 10^{-13}$ for $N = 4 \dots 24$, top eigenvalue $N^2/4 = \lfloor N^2/4 \rfloor =$ the pigeonhole maximal invisibility depth. **Mechanism: impostors hide in exactly the fastest-relaxing modes of the clock's stability operator**, which is why adding moments fails and a hitting time does not. As $N \rightarrow \infty$, $N^{-1} \mathcal{L}_N \rightarrow (-\Delta)^{1/2}$: the invisibility hierarchy's natural semigroup is Poisson, not heat. The $1/\sin^2$ kernel is Haldane – Shastri's; whether the hierarchy is a sector of an integrable spectrum is a lead, not a claim.

4.8 The marked depth: exact rank–two law [C, parameter–free]

Λ is blind to eigenvectors (isospectral $G_1, G_2 : \tau$ equal to $8 \cdot 10^{-17}$). Marked depth $\chi(G; u) = \partial_\eta(-\Lambda(\text{Cayley}(G + \eta uu^*)))|_0$ separates them (median difference 0.081). Falsified guess: χ does *not* track the resolvent $u^*(zI - G)^{-1}u$ (resolvent 1.615 $\rightarrow$ 6.667 while χ 0.0267 $\rightarrow$ 0.0028). Law:

$$D\tau[uu] = (\kappa\delta/4) \cdot uK_{ab}u + (\delta^2/8) \cdot \kappa'(u), K_{ab} = c_a v_a v_a - c_b v_b v_b, c_j = -2/(1+\lambda_j^2),^{**}$$

with (a,b) the first colliding pair and $\kappa = 8\tau/\delta^2$ a *functional*. Correlation 1.00000000, slope 1.00000, residual $5.7 \cdot 10^{-11}$ (local term alone: $0.9958 / 1.347 / 8.1 \cdot 10^{-2}$) — the earlier " $\rho \approx 1.72$ " was $\kappa + (\delta/2)(\kappa' / \delta')$ varying per mark. It is a **polarisation detector** ($P_a - P_b, \approx c \cdot \sigma_z$): $u = (v_a + v_b)/\sqrt{2}$ gives zero response; null cone $c_a q_a = c_b q_b$; orthogonal and null-cone floors $\approx 3 \cdot 10^{-3}$ are background (orthogonal marks fix λ_a, λ_b exactly); the resolvent returns at second order via $\lambda_j'' = 2q_j \sum_{k \neq j} q_k / (\lambda_j - \lambda_k)$. Blind rank-2 tomography from 60 random marks recovers $\text{span}\{v_a, v_b\}$ to principal angles $2.3 \cdot 10^{-6}, 3.0 \cdot 10^{-6}$ degrees (top-2 energy 0.986).

4.9 Classification of impostors [C]

Criterion I (transversality): relative residual of grad τ off row(DM_r): single dislocation N=8: 0.991 / 0.941 / 0.737 (r = 1,2,3); random ACUE N=7: 0.996 / 0.963 / 0.569; generic N=7: 0.998 / 0.926 / 0.195; collapses to 10^{-16} once rank DM_r = N (question becomes vacuous). Class I (caught by Λ): collision strata, half-lattice adversaries, centre-of-mass and secant families (TV 0.12 – 0.24 at N=8); fibre tomography at N=6: $E[N^2(-\Lambda) | \text{non-clock}] \in [1.3610, 1.4770]$ vs ACUE 1.4336, clock atom $\in [0, 0.0975]$ vs 0.03125. Class II (needs marked Λ): isospectral families. Class III: linear invisibility survives the linear flow (heat operator is diagonal), and Λ escapes only because it is a nonlinear stopping time. **Parity sectors**: separated by Λ at every finite N *only through the atom* ($2^{\{2-N\}}$ vs exactly 0; N=8 bulk quantiles 1.3739/1.3959/1.4182/ 1.4554/1.4839 vs 1.3752/1.3959/1.4196/1.4338/1.4788; conditional-mean differences $8.9 \cdot 10^{-3} \rightarrow 1.1 \cdot 10^{-3}$) — Class II/III in the limit.

4.10 Two ideas from the colleague not yet executed

- **Dynamic inertia profile.** Hermite – Sylvester: real-rootedness $\iff$ positivity of a Hermite form $H(p)$; along the flow, $I_p(t) = \text{signature } H(p_t)$; Λ = first loss of full signature. Inertia counts *how many* negative directions, depth *how far to the first* — same geometry, two coordinates. Ties the Anthropic inertia (Lemma R) programme to the Newman programme; untested.
- **Arithmetic heat flow (Dobner).** $\zeta_t(s) = \sum \exp((t/4)\log^2 n) n^{-s}$; the dBN flow is diagonal on Dirichlet coefficients, $a_n \mapsto a_n e^{\{(t/4)\log^2 n\}}$. Log-Gaussian weights inserted into sieve/mollifier/character sums — untested "Newmanized parity" idea (§ 9.7).

5. Bounded gaps between primes

5.1 Framework and the two controlling constants

Admissible k-tuple H; Maynard weights $w(n) = (\sum_d \lambda_d)^2$; $S_1 = \sum w, S_2 = \sum w \cdot v(n)$. If primes have level of distribution θ and $M_k = \sup_F k \cdot J(F)/I(F)$ (symmetric F on the simplex R_k) exceeds $2m/\theta$, then DHL(k, m+1) and $H_m \leq H(k)$. Bombieri – Vinogradov: $\theta = 1/2$, criterion $M_k > 4m$. Upper bound [P]: $M_k \leq (k/(k-1)) \cdot \log k$; ϵ -variants $M_{\{k,\epsilon\}} \leq (1+\epsilon)(k/(k-1))\log k$. Ladder of minimal diameters (Engelsma, proven minimal ≤ 342 ; our independent exhaustive re-proof ≤ 62 , `p4_rho_exhaust.c`, `p4_payoff_table.txt` with witnesses): $H(40..50) = 186, 188, 196, 200, 210, 212, 216, 226, 236, 240, 246$; $H(54) = 270$; $H(105) = 600$. Empirical $H(k) \approx k(\log k + 0.77)$ for $k \sim 10^4$ (fits (35410, 398130)).

5.2 The records [X + P] — `H2_H3_record_announcement.md`, `p9_*`

	bound	k	previous
H_2	173,438	15,856	396,504 (Stadlmann; earlier 398,130 Polymath8b)
H_3	13,859,802	923,601	24,797,814 (Polymath8b)
H_4	1,120,662,828	56,000,000	1,431,556,072 (Polymath8b)

Chain: Maynard's theorem (closed simplex, no ϵ , no extra equidistribution) + Bombieri – Vinogradov (margins 0.013 / 0.0067 / 0.065 leave room for $\theta < 1/2$) + Berry – Esseen $C = 0.56$ + ball arithmetic. Certificates $M_{15,856} \geq 8.013326752751$, $M_{923,601} \geq 12.006666706750$, $M_{56,000,000} \geq 16.065482942209$ (`p9_exact_cert_k*.json`, `p9_certify_hp.py`; three regimes each: two BE constants $\times$ two tail routes; small-k sanity: engine stays below known $M_2, M_5, M_{54}, M_{105}$). Tuples: $k = 15,856$ diameter 173,438 (`p9_tuple_k15856.npy`, second-implementation admissibility check); $k = 923,601$ repaired Hensley – Richards tuple diameter 13,859,802 (symmetric window $\{\pm 1\} \cup \{\pm q : q \text{ prime} \in [45,007, 6,929,899], 5,692 \text{ mid-size primes deleted by a deletion-fixpoint repair that converged in 4 iterations}\} \cup \{+6,929,903\}$; sha256 d5fe6890...6f02; classical fallback primes-past-k diameter 14,505,780); $k = 5.6 \cdot 10^7$ primes-past-k, endpoints 56,000,003 $\rightarrow$ 1,176,662,831, π -anchors matched.

The engine (`p9_mk_engine.py`). $F = \prod g(t_i) \cdot 1[\sum t_i \leq k]$; X_i iid density g^2/c_2 ; exactly $I = k^{-k} c_2^{-k} \cdot P(S_k \leq k)$, $J = k^{-k} c_2^{-k} \cdot E[G((k-S_k)_+)^2]$; layer-cake identity $E[G((k-S_k)_+)^2] = \int 2G(u)g(u) \cdot P(S_{k-1} < k-u)du$ with rigorous lower tails (chord-majorised Chernoff / one-big-jump / Berry – Esseen; monotonicity-only, so grids cannot invalidate); shaped tails $g = e^{-(t/T_1)^\chi}/(1+At)$. Recovers ≈ 1.1 units of $\log k$ (exact accounting +0.12, tail shape +0.49, rest by shape optimisation) = factor ≈ 3 in k . The 2014 crude closed-form bound had a deficit of 2.3 – 2.9 units; the 2023 – 25 improvements raised θ only. **Engine-vs-engine reconciliation:** an independent product-ansatz engine (P3) crossed $m = 2$ at $k \approx 29,500$ vs P9's 15,856; the difference was fully explained (layer-cake + tail shape) before anything was claimed. **Pitfall fixed:** P9's first HP run at $k = 923,601$ gave 9.61 because the λ grid was not scaled with T — fixed by $\lambda \propto 1/T$.

Not claimed: $H_2 \leq 145,226$ via $k = 13,467$ (Deligne-strength $MPZ[\overline{m}, 8]$ at $(\delta, \overline{m}) = (0.0205, 0.00552)$, threshold 7.8273): depends on a cap-normalisation in the truncated variational criterion that we reconstructed but never verified verbatim against Polymath8b. Also computed but not claimed: Deligne crossings $k = 13,467$ ($m=2$), 660,985 ($m=3$).

5.3 Five walls for $H_1 = 246$ [P] — `prime_gap_survey.md`

- Ceiling = tuple diameter.** No post-processing of Maynard – Tao output beats $H(k_{\min})$; pair correlation constants (Wu 3.3996 $\rightarrow$ Lichtman ≈ 3.30 , parity floor 2) cannot lower H_1 .
- Scalar decode exactly optimal.** Convex-order two-point counterfeit kills all matrix / inertia / moment decodes; $f(m) = 2m/\theta$ is final.
- Weight cone closed.** rank- r SOS decouples by subadditivity; copositive cone flat (9-pattern dual, residual $1.3 \cdot 10^{-14}$). Matrix-Maynard reformulation (uploaded doc) agrees: the SDP has a rank-one optimiser; even at level 1, $M_2/2 = 0.692965 < 1$ with $M_2 = 1/(1 - W(1/e)) = 1.38593$ — no positive eigenvalue at the twin threshold.
- Parity, combinatorial.** Kill-graph bipartite $\iff$ killable (cut-polytope odd-cycle facets); floors $H_1 \geq 6$, $k \geq 2m+1$. (The graph fact is Tao's 2014 note; the packaging is ours.)
- Level wall.** BFI 4/7, Maynard II 3/5, Pascadi 5/8 are well-factorable / fixed-residue and structurally unusable by classical MT; usable-restricted frontier: Maynard III 11/21 (uniform residues, shell-truncated), Stadlmann $1/2 + 1/40$,

Pascadi minorant 10/19. Guth – Maynard zero density is orthogonal to H_1 ; GRH gives only $\theta = 1/2$.

The doors, priced. $k = 49 \rightarrow H = 240$ (pays 6); $k = 47 \rightarrow 226$ (pays 20). Pure $M_{49} \in [3.891257590916 \text{ exact } (d = 20 \text{ power-sum basis } p \leq 5, \dim 1125, \text{mt_hp_k49_p5.json}), 3.97290]$; $M_{50} \geq 3.907113699811$. ϵ -variant $M_{\{49,1/35\}} (d = 18, n = 1597)$: float 3.959325169, **certified** ≥ 3.930490592 ; $M_{\{49,1/25\}} (d=14)$ 3.915989908. Upper bounds close the door only for $\epsilon \leq 0.00682$; Polymath left it "undecided". Tipping: $k = 49/48/47$ need $\delta \approx 0.002/0.004/0.006$ in θ , or $8 - 13\%$ of the Maynard-III shell surplus (open computation SHELL-M49). Tuple search at $k = 35,265$: best 397,352 vs record 396,504 (P4 annealing chains; missing: iterated merging, LP repair). **Pure-support kill test for $k = 49$ is unnecessary:** $(49/48)\log 49 = 3.97290 < 4$ already kills every basis enrichment (higher power sums, SDP); only the ϵ -class is alive.

5.4 The signed sieve [P + X] — `signed_sieve_nogo.md`, `sgn1_*`, `sgn2_*`, `fab_theorem.py`

Identity $(v-m) + (m-v)_+ = (v-m)_+ \Rightarrow S_2 - mS_1 - D(w) = \Sigma w_+(v-m) - \Sigma w_-(v-m)_+ \leq \Sigma w_+(v-m)$: the signed class is redundant at face-value debt. Diagnostics on a finite microcosm (n uniform on $\mathbb{Z}_W$, v = coprimality count, weights in level- L feature span, exact rationals): value = classical plateau for $\beta > \beta^*$ with $\beta = 23051796480/10991046857 = 2.0973249209^*$, classical value $1087376209/3212440751 = 0.3384891094603102$ (verified dual; signed vertex $\Phi = 1.2082816957$, $D = 0.4147152297$, 85% negative mass on 16/96 cells; certified on both sides; unbounded below $\beta_{\text{unb}} = 2.03265$ and at $\beta = 2$). Across eight variants $\beta^* \in [1.44, 6.72]$, and in five of eight the window $(\beta_{\text{unb}}, \beta^*)$ is empty ($< 10^{-6}$). ℓ^1 -budget ramp: $\lambda(1) = \lambda_{\text{positive}}$, slope $0.32 - 0.89$ — gauge, not gain. **Positivity's second job: it makes the variational problem bounded** ($\|w\|_1 = S_1 = 1$). Escapes: (i) debt below face value (exceptional character — Zhang's programme is the *only* route changing the variational picture); (ii) w evaluable while w_+ is not (well-factorable λ). Chen-switch obstruction [P-level]: switching never reduces the number r of exact-primality conditions; Chen's debt has $r = 1$; DHL($k, m+1$), $m \geq 1$ forces $r \geq 2 = \text{edge} = \text{bipartite} = \text{parity-blocked}$.

Conditional price list [X certificates] (`sgn2_certificates.json`, `sgn2_mk_large.json`):

θ	$2/\theta$	k pure (M_k)	k ϵ ($M_{\{k,\epsilon\}}$)	H_1 pure/ ϵ	m=2: k, $H_2 \leq$	m=3: k, $H_3 \leq$
1/2	4	54	50	270 / 246	15,856 $\rightarrow$ 173,438	923,601 $\rightarrow$ 13,859,802
4/7	3.5	31 (3.502015496)	29 (3.519881250)	140 / 130	5,647 $\rightarrow$ 58,058	202,528 $\rightarrow$ 2,856,288
7/12	24/7	29 (3.443305315)	26 (3.433616498)	130 / 114	4,835 $\rightarrow$ 48,988	160,703 $\rightarrow$ 2,226,804
3/5	10/3	26 (3.350647068)	23 (3.334615948)	114 / 94	3,931 $\rightarrow$ 38,878	120,497 $\rightarrow$ 1,632,566
5/8	3.2	22 (3.207656229)	20 (3.222665844)	90 / 80	3,022 $\rightarrow$ 29,180	80,165 $\rightarrow$ 1,051,602
1 (EH)	2	5 (2.007080)	—	12	221 $\rightarrow$ 1,498	1,978 $\rightarrow$ 18,144

($m \geq 2$ rows use the p9 engine's certified-lower-bound path and primes-past- k diameters.) All rows are parity-consistent (≥ 12 for $m = 1$). **Missing estimate (E_θ)** — tuple-residue well-factorable estimate: for $c_q(a)$ jointly well-factorable with the CRT residue selection $a \bmod p \in \{h_i - h_j\}$, $\Sigma_{\{q \leq x^{\theta-\epsilon}\}} \Sigma_{\{a \in A_i(q)\}} c_q(a)E(x;q,a) \ll x(\log x)^{-A}$. None of BFI Thm 10 / Maynard II / Pascadi covers it (all fix one residue per modulus); $MPZ[\overline{m},\delta]$ is its absolute-value cousin at level ≤ 0.5286 . ($E_{\{7/12\}}$) is structurally cheapest (λ can literally be Iwaniec's $\lambda \pm$).

5.5 EH → 8 via one signed cross block [P conditional] (uploaded EH8_OBJECTIVE_ALIGNED_CROSSBLOCK.md)

For $H = (0,2,6,8,12)$, with the frozen exact trial $Q_5 = 1562651575013110693/778568621714732244 = 2.00708265325605$ ($M_5 > 2$), EH plus one Maynard-weighted scalar $\sum_{v \leq x(n)} C_8(n)(\Theta_H(n) - \log 3x) = o(\log x \cdot \sum_{v \leq x})$, $C_8 = \lambda(n)\lambda(n+2) - \lambda(n)\lambda(n+12) - \lambda(n+2)\lambda(n+12)$, gives $H_1 \leq 8$ (tolerance $-1/500$, or $-7/2000$ under full EH). $W_8 = (1-ac)(1-bc)$ is the *unique* parity-neutral blocker (Theorem 4.1); connected bipartite bad graph on r vertices imports even Walsh orders through $2\lfloor r/2 \rfloor$ (Theorem 4.2). Ladder: gap 10 needs one 2-point Liouville correlation; 8 needs the degree-2 block; 6 needs one quartic; 4 is odd-cycle impossible (parity wall). RMT corollary: GRH + Dirichlet-WPC $_{\beta}$ with $\beta > 2 - Q_5/2 = 0.996458673\cdots + OMC_8 \implies H_1 \leq 8$. Novelty is narrow (objective-aligned one-scalar compression); Ford – Maynard's primal – dual theory and Murty – Vatwani are the benchmarks; OMC_8 is not known to be easier than GEH.

5.6 The Lagarias–Rodgers μ programme (uploaded LR docs, verified in part)

$\mu = \sup$ hard core of bandwidth-one sine mimickers: $1/2 \leq \mu \leq 0.606894$ (LR); conjecture $\mu = 1/2$. Our contributions there: exact signed Fermi-path formula for DPP targets; hard-core necklace enumeration; **Theorem 4.1** (unbounded exact feasible branch $\implies$ LR process, one-way); **Theorem 11.1** exact all-cardinality 20-cycle separator over $\mathbb{Q}(2\cos(\pi/10))$ (positive on all 109 orbits, Fermi target $-1/500$) and its exact non-transfer to the continuous circle ($L(X_{\triangle}) = -954123/256000000$; seam obstructions $\{0,19\}, \{0,18\}$); the smooth nine-coordinate ansatz collapses to zero margin jointly at $N = 4,5,6,7$ (a lattice-trained vector already fails at $N = 8$: $\min -1.08 \cdot 10^{-3}$ at $\{0,3,\dots,36\} \subset \mathbb{Z}/40$). **Palm row-sum square** (**Theorem 13.1 / Prop 5.1**): for band-limited f ($\text{supp } \hat{f} \subset [-1/2, 1/2]$), $\text{Var}^0_{\text{-sine}(S_f)} \leq (M_h - A_f)(A_f - m_h)$ is necessary for an h -hard-core mimicker; the multiwindow quadratic version on the packing body K_h . The first nonnegative Fejér profile fails for an exact local reason ($\{-h,0,h\}$ forces $S_0 = 2f(h) > 9/7$); signed sinc-power searches found nothing. **The depth bridge** (§ 4.6) is the new lever on μ .

5.7 Agendas received and their status

- **Zhang Yitang** (2026). Landau – Siegel; "remove the square". Verdict: the variational half is closed (§ 5.4); the arithmetic half — sub-face-value debt via the exceptional character — is the only door. His "first experiment" (λ_{signed} vs $\lambda_{\text{positive}}$) is answered: no new phase.
- **Bourgain toolkit**. Exp 3 (ACUE $\times$ decoupling) done (§ 2.3, negative for decoupling); Exp 2 merged with Zhang's; Exp 1 (Type III $\times$ additive energy) and Exp 4 (function-field affine sieve) never launched.
- **Cross-field tools list** (log-Chowla bridge, exponent-calculus compiler, hypermetric facets, magic functions): approved in principle, superseded, never launched.

6. Failed trials, refutations, and what each one taught

These are first-class results. Each closed a direction that a fresh system would otherwise re-open.

trial	outcome	insight
"0.6725 is the exact bandwidth-one LP optimum; AH is suboptimal by 0.0301" (our own Round-1 claim)	[X] retracted (feasible-set conflation)	the free lattice optimum and the Anthropic-feasible optimum live in different cones; 0.0208 of the 0.0301 is aliasing artifact

trial	outcome	insight
Improve 0.6725 by window engineering	[X]	the clump term forces $M = A_0 \log T$ for every window; only the M_- lemma can help
Pointwise edge hypothesis $ F(1)-1 \leq \varepsilon$	[X] buys exactly 0 even at $\varepsilon = 0$	$\hat{r}(\pm 1) = 0$ for every admissible window; the edge object is the Cesàro mean
Derivative pair correlation at bandwidth one	[X] blind	—
$\text{tr } Q^3$ as a free upgrade	[X] (worthless at flat window; $\Phi_3 = -0.01178$ at MT window but unmonetisable without M_-)	the blocker is the unbounded negative spectrum of the compression, not arithmetic
ℓ^2 -decoupling on the 0.0301 gap	[X] misprices by $15 - 77\times$	gap is single-site in position ($S_1 \geq 0$ at half spacing), full-band in Fourier
Higher traces ($\text{tr } B^4 \cdots B^{10}$) improving Guth – Maynard by pure deflation	[X] 0% in the idealised regime, 0.1 – 1.1% with secondary clusters	room must come from additive energy / resonator geometry
Matrix / inertia / moment decode for Maynard – Tao	[X] rank-one optimiser; scalar decode optimal	convex-order two-point counterfeit
PSD / copositive enlargement of the weight cone	[X] flat	rank- r SOS decouples
Signed weights open a new variational phase (Zhang's experiment)	[X] one-line identity	positivity also <i>bounds</i> the problem; the value of signed weights is purely arithmetic
$k = 49$ pure-support basis enrichment / SDP kill test	unnecessary	classical upper bound $3.97290 < 4$ already kills it; only ε -class alive
Higher-order (5+ Fermi sea) secant adversaries for deeper invisibility	[X] $L = 2$ is globally optimal by pigeonhole ($\lfloor N^2/4 \rfloor$)	the open problem is the <i>interference</i> code (autocorrelation zeros), not the support
FUSION-ACUE "KW fusion = ACUE triple pairings"	[X] at $N = 3$ (22 mismatches) though Verlinde = KW holds	charge bookkeeping in Conjecture A is wrong as stated
Marked depth tracks the directional resolvent	[X] an order of magnitude the wrong way	first order is a rank-2 critical-pair contrast; resolvent returns at second order
$s^* = 1.419640342$ as the ACUE median limit	[X] median turns at $N = 7$	s^* is a configuration constant of the single-dislocation stratum
$\rho \approx 1.72$ as a universal constant in the marked-depth law	[X]	it is $\varkappa + (\delta/2)\varkappa' / \delta'$ per mark; law is parameter-free with $\varkappa$ measured
Linear heat flow "mixing" hidden fibre directions into visibility	[X] the flow is diagonal; invisibility persists for all t	only the nonlinear stopping time escapes
Smooth nine-coordinate LR separator uniform in N	[X] zero margin at $N = 4.7$ jointly; fails at $N = 8$	continuum stability is a two-parameter (mesh $\times$ volume) limit; fixed-cardinality certificates can be spurious
Nonnegative Fejér profile in the Palm row-sum bound	[X] exact local failure ($\{-h, 0, h\}$)	signed f with certified packing bounds is the right search
Tuple search at $k = 35,265$ beating 396,504	not achieved (397,352)	needs iterated merging / LP repair

trial	outcome	insight
Deligne-route $H_2 \leq 145,226$	not claimed	cap-normalisation unverified verbatim
Shared-context data errors ($H(47) = 232$, $H_2 = 398,130$, exponent 3.815 , $H(26) = 120$)	corrected by agents to $226 / 396,504 / 3.8075 / 114$	keep a living errata list in the context file

Meta-lessons. (a) Each refutation came with the corrected statement; that is the standard to keep. (b) Five different "phase transitions" turned out to be normalisation artifacts (β -kink, ℓ^1 ramp, unboundedness at $\beta = 1$, $\rho = 1.72$, the median $\rightarrow s^*$). Always ask whether the decode is scale-invariant before reading a growth as a gain. (c) When an idea uses only static polynomial information about a family that is moment-frozen, it will fail; move to stopping-time / hitting-time / marked observables.

7. Verification status, Lean, and where everything lives

7.1 Machine verification (`fab_verify_proof.py` / `verification/verify_theorem_steps.py`)

18 checks, 18 passing, no floating point: flow generator identity; $2z_j/(z_j - z_k) = 1 - i \cdot \cot$; the $(N-1)$ cancellation; two-body solution and collision time; $d/dx \cot(x/2) = -1/2 \csc^2$; $f(0) = 0$ and $f' = 1/4(2\tan(x/2) - x)$; nonnegative Taylor coefficients of $\tan t - t$ ($1/3, 2/15, 17/315, 62/2835, 1382/155925, \dots$); **the sign lemma as a z3 decision problem (unsat for the negation)**; ACUE pigeonhole; $\Sigma 1/2 \csc^2(\pi k/N) = (N^2-1)/6$ at $N = 4, 6, 8, 12, 20$. Other exact checks: the signed identity on 1,000 random weights across five models; β^* with verified dual; the cellwise identity $a + G = Gp$.

7.2 Lean — written, NOT compiled

`research/riemann-rmt/lean/DepthComparison.lean` (also packaged as `riemann-impostors/lean/RiemannImpostors/DepthComparison.lean` with `lakefile.toml`, toolchain pin `leanprover/lean4:v4.33.0-rc2` — matching the `anthropics/zeta-23-lean` clone at `/workspace/anthropics/zeta-23-lean`). Declarations: `cot`, `hasDerivAt_cot`, `cot_strictAntiOn`, `background_sign`, `self_le_tan`, `neg_log_cos_ge` (proofs written), `two_body_solution` and `depth_ge` (sorry: scalar autonomous ODE uniqueness; Grönwall-type comparison). No toolchain could be installed: egress 403 for `elan.lean-lang.org` and for GitHub *release* downloads (`git ls-remote` of public repos works; releases do not; building Lean + Mathlib from source is not feasible in-session). If Astra has a toolchain: **lake update && lake build**, then discharge the two sorries. Also unformalised but small: the five-line `mult_two_pair` edit in `zeta-23-lean` giving $N_d \geq 0.8362503 \cdot N + p$.

7.3 Repository index (branch `claude/riemann-zeta-random-matrix-udxp3f`, PR #11)

`research/riemann-rmt/:`

- Papers: `impostors_paper.md` (canonical, 10 sections), `stopping_times_paper.md`, `depth_scaling_theorem.md`, `signed_sieve_nogo.md`, `newman_depth_note.md`, `H2_H3_record_announcement.md`, `prime_gap_survey.md`, `round3_synthesis.md`, `final_verified_paper.md`, `rmt_zeta_survey.md`, `rmt_zeta_popular.md`, `tao_ah_notes.pdf`, `codex_handoff.md`, `README.md`.
- Context files given to agents: `joint_context_v2.md`, `prime_gap_context.md`, `signed_context.md` (and `lambda_context.md` in the scratchpad).

- Engines/certificates: `p9_mk_engine.py`, `p9_certify_hp.py`, `p9_exact_cert_k*.json`, `p9_tuple_k*.npz`, `p9_g_k*.npz`, `p9_scan.py`, `p9_tuples.py`, `verify_codex.py`.
- Theorem checks: `fab*.py` (+ `*_results.json`), `sgn1_t1_exact.{py,json}`, `beta*.py`, `tomo_fiber.py`, `d6_allN_family.py`.
- `riemann-impostors/`: standalone package (README = labeled results summary; paper/, lean/, counterexamples/, verification/, certificates/, data/ with `dyn1_results_N3..10.npz`; PUBLISHING.md — repo creation returned 403 from the GitHub App integration, so publish via `git subtree split`).
- `handoff/`: this document (md + pdf).

Session scratchpad (/tmp/claude-0/.../scratchpad, ~1,150 files) holds everything else: `dyn2_data_N*.npz` (CUE, $N = 2..256$), `beta_data_b{1,4}_N*.npz`, `monsep*`, `fus*`, `brg*`, `dir*`, `p1-p10*`, `sgn*`, `tomo*`, `dyn1*`, agent logs. It survived two container restarts but is not guaranteed to outlive the session; anything essential has been copied into the repo.

7.4 Key constants (verified this programme)

$\delta_{MT} = 3/2 - (1/\sqrt{2})\cot(1/\sqrt{2}) = 0.672500703679$ · PairCeiling 0.6818287 · ceiling candidate 15/22 · Φ_3 (MT) = -0.0117753128 · $N_d/N \geq 0.8362503$ · deficiency ledger 0.3275 · M_- payoffs 0.6796896 / 0.6844924 · $\delta_{free} \rightarrow 1/2 + 2/\pi^2 = 0.70264237$ · doubles density $1/4 - 1/\pi^2 = 0.148679$ · $M_2 = 1.385933$, $M_3 = 1.646440$, $M_4 = 1.845401$, $M_5 = 2.007080$ (Q_5 exact above), $M_{5.4} > 4.00238$, $M_{4.9} \geq 3.891257590916$, $M_{\{49,1/35\}} \geq 3.930490592$ (float 3.959325169) · $\beta^* = 2.0973249209$ · $\pi^2/8 = 1.2337005501$ · $s^* = 1.419640342$, $u^* = 1.812942145$, $\rho_\infty = 1.150717118$ · separated-defect $\rho = 1.19120$ ($N^2(-\Lambda) = 1.46946$) · $\rho_{max} \rightarrow 1.6094$ · $72\pi = 226.19$ · $(72\pi \ln 2)^{2/3} = 29.08$ · $N^* \approx 165$ · LR: $1/2 \leq \mu \leq 0.606894 \Leftrightarrow \mu_\Lambda \leq 1.8177$ · clock stiffness $(N^2-1)/6$ · fibre dims 0,0,2,10,80,403,1804.

8. Open problems, ranked, with the exact statement and its price

- The background bound [the linchpin].** Prove: for $C\beta E$ (β fixed), with probability $\rightarrow 1$ the background stiffness of the first-colliding pair satisfies $S \leq A \cdot N^2$ (measured median $0.120 \cdot N^2$, max $0.79 \cdot N^2$). This single rigidity statement upgrades to theorems: $8N^{8/3}(-\Lambda) \Rightarrow G^2(\beta=2)$, $-\Lambda \asymp N^{-2-2/(\beta+1)}$ for all finite β , and Theorem C(ii). It is a standard-shaped estimate ($\sum_k d_k^{-2}$ near the extremal pair) and is the best-value item on this list.
- $\rho_\infty = O(1)$ for ACUE at all N (lattice upper bound). Deterministic; proven only for $N \leq 10$ by enumeration ($\rho \in [1.049, 1.610]$); the minimising family's ρ decreases toward $\approx 1.03 - 1.05$.
- Exact constants:** $s^* = 1.419640342$ as a theorem (infinite clock + one dislocation under the Coulomb flow; Calogero – Moser framing); the separated-defect $\rho = 1.19120\dots$; ρ_∞ of the minimising family. Targets: closed forms in theta/Bessel/trigonometric data.
- The marked-depth law as a theorem,** $D_\tau[uu^*] = (\kappa\delta/4)u^*K_{ab}u + (\delta^2/8)\kappa'(u)$, with κ' characterised; on the null cone $c_a q_a = c_b q_b$ the leading term should be $\lambda_{-j}'' = 2q_{-j} \sum q_{-k}/(\lambda_{-j} - \lambda_{-k})$.
- Transversality theorem:** $\tau|_{\{F_m\}}$ generically non-constant; marks added until $\cap \ker D\Lambda\{u_{-j}\} \cap \ker DM = \{0\}$.
- The M_- lemma** on the zeta side: $\|(c^{-1}\hat{A})_-\| \leq M_-$ uniform in T ; worth +0.007..0.012 on 0.6725.
- The 15/22 ceiling identity** and $d = 3(N-3)^*$ kill-degree law (exact-arithmetic conjectures).
- (E_θ)** for any $\theta > 1/2$ (cheapest: 7/12) — the one statement between the price list and $H_1 \leq 114/130$.
- $k = 49$ ϵ -door:** push the Galerkin engine to $d = 20 - 22$ at $\epsilon \in \{1/40..1/30\}$, warm-start from `p2_vec_hp_k49_d18_e35.npz`; if the float plateaus < 3.98 , try the vanishing-marginal variant; non-separable dual

certificates are the only known way to *close* the door above $\varepsilon = 0.0068$.

10. **SHELL-M49**: compute M_4 on Maynard-III's shell polytope with the ε -trick layered; needs 8 – 13% of the full-shell surplus; undecided in print.
11. **Tuple k** = 35,265 below 396,504 (best 397,352; add iterated merging + LP repair).
12. **Interference code / phase-rigidity conjecture**: among phase-twisted Slater states $u \in U(1)^{\{2N\}}$, vanishing of all balanced Schur minors below $\lfloor N^2/4 \rfloor$ forces $u_c = \alpha\beta^c$ up to gauge/dihedral symmetry (would make the det-character impostor the canonical extremiser of its whole class). Computationally attackable via circulant Toeplitz minors / Plücker relations.
13. **FUSION-ACUE, corrected**: find the charge bookkeeping under which KW fusion matches ACUE triple pairings (fails at $N = 3$ as stated).
14. **Function-field Newman universality** (§ 9.4).
15. **Formalisation**: discharge the two Lean sorries; formalise the records' certificate chain; make the five-line `mult_two_pair` edit.

9. Wide-open research programme for GPT6 Astra: which famous conjecture, and by what route

This section answers the second half of the handoff request. It is opinionated and it is honest about odds. The ordering is by (probability of a genuine theorem) $\times$ (size of the theorem), not by glamour. The human's stated ambition is a historic result, not another small one; the way to earn that is to pick the target whose *remaining gap* is a statement of a shape people already know how to prove.

9.1 The best target: refute the Alternative Hypothesis dynamically

Claim to aim for. *Under a hypothesis strictly weaker than GUE, the zeros of ζ cannot have asymptotic hard core $1/2$.*

Equivalently, a proof that the Alternative Hypothesis is false — the statement Tao, Lagarias – Rodgers and others have posed as the obstruction to Montgomery.

Why this is now the right target. The static route needs pair-correlation information at Fourier bandwidth > 1 , which is exactly what no one can prove. The dynamic route needs only that *the zeros are, locally, more fragile than a lattice under the de Bruijn – Newman flow*. Concretely (Theorem C(i) + the LR bridge, both [P]): $AH \implies$ every window's local depth satisfies $\liminf N^2(-\Lambda) \geq \pi^2/8 = 1.2337$. So:

Level B target. For ζ , with $N = (\log T)/2\pi$ zeros in a window at height T under the *true* H_t flow, prove $\liminf (\log T)^2 \cdot D_T < \pi^2/8$ — or merely $D_T = o((\log T)^{-2})$ — from averaged information (pair-correlation at bandwidth ≤ 1 , plus variance/energy inputs of Rodgers – Tao type).

Level A (the full $N^{-8/3}$ CUE law) is *not* needed; Level C ($\liminf (\log T)^2 D_T = 0$ along a subsequence) may already exclude rigid subclasses.

What Rodgers – Tao already give you. Their proof of $\Lambda_{\text{dBN}} \geq 0$ runs exactly this kind of argument in reverse: assume the zeros are over-stable ($\Lambda < 0$), push the flow, obtain a local clock equilibrium, contradict known fluctuations. The needed extension is *quantitative and local*: the same energy method in a window, with the over-stability hypothesis replaced by

"hard core $1/2$ " and the contradiction replaced by a fluctuation input weaker than GUE. The finite results say what the right normalisation is (the clock stiffness $(N^2-1)/6$; the two-body constant $1/8$; the singular endpoint at $\beta = \infty$ where the background is leading-order).

Steps, in order.

1. **Truncation theorem.** One may not cut a window of zeros and call it a polynomial. Use the Polymath15 machinery for H_t (effective approximations to the flow) to define D_T rigorously and prove that the window's first collision time is controlled by the finite depth of the local configuration up to $o(1)$ — the same localisation lemma as Theorem B, now for ξ .
2. **Hard-core $\Rightarrow$ depth (already proved at finite N).** Transfer Theorem A to the H_t flow (the cotangent interaction becomes the Rodgers – Tao kernel; the sign argument survives since it only uses monotonicity of the pair interaction).
3. **The fluctuation input.** Show that near-collisions at scale $o(1/N)$ occur along a positive proportion of windows using *only* bandwidth- ≤ 1 pair correlation plus an energy/variance bound. Note Montgomery's own theorem already gives a positive proportion of gaps < 0.68 mean spacing (unconditionally, via bandwidth-1 data); the question is whether such inputs plus the flow give gaps *below* $1/2$ often enough to break the floor. This is where new mathematics is required, and it is a concrete analytic-number-theory question, not a random-matrix metaphor.
4. Assemble: (2) + (3) $\Rightarrow \liminf N^2 D_T < \pi^2/8 \Rightarrow AH$ false.

Risk. Step 3 may be equivalent in strength to the static statement one is trying to avoid. The finite theory suggests not — the depth is a *first-passage* quantity and Theorem C(ii) needed only the gap tail plus a bounded background — but this has to be settled on the zeta side. Either way the outcome is a theorem: either AH is refuted, or a precise equivalence "dynamic AH-refutation $\Leftrightarrow$ bandwidth- $(1+\eta)$ pair correlation" is proved, which would itself be new.

9.2 Second target: the Lagarias–Rodgers conjecture $\mu = 1/2$

Fully within random-matrix/point-process theory, no arithmetic input. [P] already: $\mu \leq \sqrt{(2\mu_\Lambda)}/\pi$. So one needs a *depth upper bound* for every bandwidth-one sine mimicker: prove that any such process has $\liminf N^2(-\Lambda) \leq M$ with $M < 1.8177$ to beat the published bound, and $M \leq \pi^2/8 \cdot \rho_{\text{lat}}$ to reach $1/2$. Route: transport the Palm row-sum-square / multiwindow packing-body certificates (LR docs, Prop 5.1 – 5.2) from the row sum to the depth. The depth is smooth in the configuration (gradient computed in § 4.9), has an explicit critical-pair derivative law, and — unlike the row sum — has no known local obstruction of the $\{-h, 0, h\}$ type. A win here is a clean, publishable theorem in its own right and directly feeds 9.1.

9.3 Third target: a Dyson–Montgomery statement one can actually prove

The honest position: the full pair-correlation conjecture is not reachable by these tools. What *is* plausibly reachable, and would count as a partial Dyson – Montgomery result, is a **dynamic universality theorem**:

Any point process with (a) bandwidth-one sine pair correlation, (b) a repulsion exponent β at short range, and (c) bounded background stiffness, has finite depth law $-\Lambda \asymp N^{-2-2/(\beta+1)}$.

This is Theorems A + B plus the extreme-gap input, stated abstractly; the $C\beta E$ case reduces to Feng – Wei, and the ACUE case is the $\beta = \infty$ endpoint. It says the Newman depth is a *fingerprint of the repulsion exponent* — a random-matrix theorem with a clear statement, four confirmed data points and one missing lemma (open problem 1). Expected effort: weeks, not years.

9.4 Fourth target: function–field Newman–depth universality (Katz–Sarnak)

Function-field Newman constants exist in the literature (Andrade – Chang – Miller-type explicit formulas; e.g. $\Lambda_{\{D_p\}} = \log(|a_p|/2\sqrt{p})$ at genus 1 is exactly our $N = 2$ law in arithmetic clothing). Define, for a family of L-functions over F_q with Frobenius classes $\Theta_C \in G$, the depth of $\det(1 - u\Theta_C)$; Λ is a highly nonlinear class function on G . **Theorem to prove:** as $q \rightarrow \infty$ at fixed rank N , $\text{Law}(\Lambda_C) \rightarrow \text{Law}(\Lambda_{\text{Haar}}(G))$ for $\text{USp}/\text{SO}/\text{U}$ families with big monodromy (Katz – Sarnak equidistribution pushes through a *singular* statistic because Λ is continuous off the discriminant locus, whose Haar measure is zero; the delicate part is the near-collision set). Then $N \rightarrow \infty$ gives the universality exponents per symmetry class, with the ± 1 hard edge in Sp/SO possibly changing the minimal-gap law. Our monodromy-separator computation (§ 3.5) shows the det-sector is already dark at $q^{-1/2}$; the $\text{U}(N)$ family machinery (`monsep_core.py`) is reusable. This is the most *reliable* big theorem on the list, and the one best suited to a geometric collaborator.

9.5 Prime gaps: from 246 (or from the reported 186) downward — what is and is not possible

Consistency check on "186". $H(40) = 186$ exactly, so a proof of $H_1 \leq 186$ is a proof of $\text{DHL}(40, 2)$. Under Bombieri – Vinogradov ($\theta = 1/2$) that needs a Maynard – Tao variant > 4 at $k = 40$. The classical bound gives pure $M_{40} \leq (40/39)\log 40 = 3.78347$, so the pure problem is *impossible*; an ε -variant needs $(1+\varepsilon) \cdot 3.78347 > 4$, i.e. $\varepsilon > 0.05723$ — far beyond anything Polymath ever certified (their ε -trick at $k = 50$ used $\varepsilon \approx 1/25$ and moved M by ≈ 0.02); alternatively a level of distribution θ with $2/\theta < M_{40}$ -variant, i.e. $\theta > 0.529$ at least (with pure $M_{40} \approx 3.75$ one needs $\theta > 0.533$), which is beyond Maynard III's $11/21 = 0.5238$ and would need residue-uniform input of (E_θ) type. So *if* 186 is real, its proof contains either a new equidistribution theorem at level ≥ 0.53 uniform in CRT residues, or a variational enlargement beyond every known variant. Astra should first identify which, since that determines everything below. (If the number was actually 246, ignore this paragraph.)

From wherever the record stands, the levers, priced:

- $k = 49$ ε -door $\rightarrow 240$: open, undecided; float 3.9593 vs 4; degree-22 Galerkin + ε -scan (open problem 9). Cheapest possible improvement of 246.
- $k = 47 \rightarrow 226$: pays 20; needs ≈ 0.06 more in M or a level $\delta \approx 0.006$.
- $(E_{\{7/12\}}) \rightarrow 114$, $(E_{\{4/7\}}) \rightarrow 130$, $(E_{\{3/5\}}) \rightarrow 94$, $(E_{\{5/8\}}) \rightarrow 80$: the well-factorable route, blocked only by the tuple-residue estimate (E_θ) . This is the route with the largest payoff and the clearest missing statement; the structurally cheapest is $7/12$ where λ can literally be Iwaniec's $\lambda^\pm$.
- $\text{EH} \Rightarrow 12$, $\text{EH} + \text{OMC}_8 \Rightarrow 8$, $\text{EH} + \text{quartic block} \Rightarrow 6$; parity floor 6 absolute.
- What cannot work (walls 1 – 5): post-processing, matrix decodes, PSD/copositive cones, signed weights at face-value debt, zero-density inputs, GRH.
- $m \geq 2$ records are ours and are far from optimal: Galerkin/ ε refinement would shave 3 – 10% off k^* ; H-R tuple repair at $k = 5.6 \cdot 10^7$; the Deligne cap-normalisation (verify verbatim $\rightarrow H_2 \leq 145,226$).

9.6 A formalisation programme in the FLT style

We could not read Anthropic's post on formalising Fermat's Last Theorem; Astra should, and should import its agent architecture. What in this programme is *ready* to formalise, in order of value per effort: (1) Theorem A (two sorries away); (2) the five-line `mult_two_pair` edit giving $N_d \geq 0.8362503 \cdot N + p$ in `zeta-23-lean` — a strictly stronger published constant, in an existing verified codebase; (3) the records' hypothesis chain (Maynard's theorem is not in Mathlib, but the

certificate — ball-arithmetic bounds on I and J for an explicit F — is a finite verification and the tuples' admissibility is decidable); (4) the signed no-gain identity (trivial to formalise, useful as a lemma); (5) $P(\text{clock}) = 2^{\{1-N\}}$ and $\delta_{\min} = \pi/N$.

9.7 Speculative but structured: the "Newmanised parity barrier"

ACUE proved once that a static information barrier is not a dynamic one. The twin-prime parity barrier is also a static barrier (sieve-accessible observables cannot separate a and b). Look for a flow S_t on arithmetic sequences — the Dobner flow $a_n \mapsto a_n e^{\{(t/4)\log^2 n\}}$ is the canonical candidate, being the dBN flow on the Dirichlet side — under which sieve observables stay matched but a positivity/stability hitting time differs. No theorem is known; the problem is now well-posed enough to search.

9.8 What not to spend the first month on

Larger N simulations of the depth (the exponents are settled); five-Fermi-sea adversaries ($L = 2$ is optimal); decoupling on the bandwidth-one gap; pure-support $k = 49$; higher traces for Guth – Maynard by deflation alone; any "new phase" from signed weights; window engineering on the zeta side.

10. Closing note to Astra

The programme's single deepest fact is that the invisibility hierarchy of Tao's impostors and the stability spectrum of the zero dynamics at the clock are the same operator: impostors hide exactly where the flow acts hardest. Everything else — the depth laws, the separation, the LR bridge, the $\pi^2/8$ threshold — is that fact read in different coordinates. The route to a historic theorem runs through § 9.1, and its one genuinely new ingredient (step 3) is an analytic statement about near-collisions of zeros that current technique might reach. Prove open problem 1 first; it costs little and upgrades three numerical laws to theorems. Then go for AH.

Everything here is reproducible from the repository. Where we were wrong, we have said so; where we were uncertain, the tag says so. Good hunting.

Fable source; scope set by current audit 40: CLAIMS

Source: `fable/overnight-2026-09-05/CLAIMS.md` . SHA-256:
`85b8a6c259b7f6abc7eeac59204f3939631eed625455fd8a546db889cd72e2a2` .

Claims ledger (Fable overnight)

Format: `| id | claim | status | file | notes |` . Status in {P, C, R, O}. Refuter votes recorded.

id	claim	status	file	notes
D1.2	product-ansatz layer-cake bound $M_k \geq c_2^{-1}$ $\sum_j \max(0, 1 - \beta_{a_j})$ $(G(b_j))^2 - G(a_j)^2$ with chord-Chernoff / one-big- jump / Berry-Esseen tails is valid	P	<code>r1_h2_interval_cert.md</code> § 2	Lemmas 1-6 written out; any tail bound family works
D1.3	$M_{15856} \geq$ 8.013326752751306578613 695503115 (exact rational; $C_{BE}=0.56$, elementary Phibar), outward-rounded (arb 200 bits) and cross- checked with <code>mpmath.iv</code>	C	<code>r1_h2_interval_cert.md</code> § 3, <code>data/h2_k15856_interval</code> <code>_cert.json</code>	margin 0.01333; reproduces historical JSON bit-for-bit
D1.4	$M_{15856} \geq$ 8.00677408008999410774 with NO Berry-Esseen input (Markov + convexity only); M_{923601} $\geq$ 12.00263034990571191492 likewise	C	<code>r1_h2_interval_cert.md</code> § 3.2, § 3.6	H2/H3 records need only Maynard + Bombieri- Vinogradov
D1.6	<code>p9_tuple_k15856.npy</code> : 15856 entries, diameter 173438, admissible for all 1847 primes ≤ 15856 (two implementations)	C	<code>r1_h2_interval_cert.md</code> § 4	prime count confirmed by <code>sympy.primepi</code>
D1.7	committed <code>p9_certify_hp.py</code> is <code>mpmath dps=50 +</code> <code>SAFE=1+1e-30</code> , not outward rounding (Astra audit confirmed); JSON came from an uncommitted arb script, now copied with sha256	R	<code>r1_h2_interval_cert.md</code> § 1.5, <code>scripts/r1_h2_reference_p</code> <code>9_exact_cert_scratchpad_</code> <code>copy.py</code>	repaired by D1.3/D1.4

id	claim	status	file	notes
D1.8	Maynard/Polymath theorem statements and BE constants are recalled not re-read; arb library trusted; k=923601 tuple not re-verified tonight	O	r1_h2_interval_cert.md § 5, § 8	

Fable source; scope set by current audit 41: FABLE_COORDINATION

Source: `fable/overnight-2026-09-05/FABLE_COORDINATION.md` . SHA-256: `b89ea13f87faa95b507ef0c059063c14aeeb2f9981d475ee606db58a4b228df0` .

Fable overnight layer (2026-09-05, 06:10 UTC start)

Companion to `../../handoff/astra-2026-09-05/COORDINATION.md` . Same rules: no questions to the sleeping user, fast-forward pushes to PR #11 only, fetch before every push, never rewrite old papers silently. Everything in this directory is written by Fable's adversarial harness (proposer agents challenged by independent refuters), and every file carries a status tag:

- [P] proved in the file, with the argument written out and machine-checked where possible
- [C] computed (script + data in this directory, reproducible)
- [R] refuted or repaired (the old claim is wrong; the file says what replaces it)
- [O] open; the file records the exact obstruction and failed attempts

Acknowledged Astra audit items (acted on tonight)

1. Background endpoint bound in `depth_scaling_theorem.md` § 4 is wrong as stated ($\csc^2$ is increasing on $(\pi, 2\pi)$). Repair: define S with the endpoint maximum, $S := \sum_k \frac{1}{2} \cdot \max(\csc^2(x_b^{k/2}), \csc^2(x_a^{k/2}))$, which is what the mean value theorem actually gives. Theorem B is re-derived with the repaired S in `r1_theoremB_repair.md` .
2. `lean/DepthComparison.lean`: `depth_ge` does not mention the collision time. Tonight's Lean work states the ODE comparison honestly (still uncompiled here — no toolchain egress).
3. AH versions: strong LR2019 (no zero gaps) vs density/multiplicity-tolerant versions are kept separate in every statement below.
4. H2/H3 certificates: outward interval arithmetic re-certification is a Round-1 task.

Fable's allocation (independent of Astra's four agents)

cluster	question	files
A depth-rigor	repair Theorem B; prove the CUE background bound w.h.p. from the determinantal 3-point estimate; C β E version via the explicit density; state the resulting unconditional depth law	<code>r1_theoremB_repair.md</code> , <code>r1_cue_background.md</code> , <code>r1_cbe_background.md</code>
B zeta-frontier	is Level B ($\liminf (\log T)^2 D_T < \pi^2/8$) easier than $\mu < 1/2$? what is the true record for μ and the limit of CGG-type mollifier methods; zeta-zero numerics for the depth statistic	<code>r1_levelB_barrier.md</code> , <code>r1_small_gaps.md</code> , <code>r1_zeta_numerics.md</code>

cluster	question	files
C simple-zeros	what exactly is 0.6725 (Montgomery – Taylor simple-zero proportion under RH); can any correlation-only method exceed the pair ceiling; LP/SDP over multiplicity distributions with 2- and 3-level constraints; honest comparison with 19/27	r1_simple_zeros.md , r1_correlation_lp.md
D prime-gaps	outward-rounded certificate for the H2 record ($k = 15,856$); sub-186 feasibility wall for $k = 39, 40$ under Bombieri – Vinogradov with the largest admissible support enlargement	r1_h2_interval_cert.md , r1_sub186_wall.md
E structure	adversarial review: which of the "new structures" (operator unification, marked depth, secant families) are rediscoveries of known linearisations of Dyson Brownian motion, and what survives as new	r1_structure_review.md

Round notes are appended to `CLAIMS.md` (one line per claim, with status and file). Rounds are numbered r1, r2, ... Astra: anything here is free to reuse; please cite the file path.

Cross-repository status (06:30 UTC)

The user created `QingyunSun/Riemann-hypothesis-and-random-matrix` for the collaboration. From this session Fable can read it (public clone works) but cannot push: the git proxy injects credentials only for the session's authorised repository set (`galpha-ai/Alpha-devbox`), the cross-owner `add_repo` is refused in this session, and spawning a helper session with the new repo as source was blocked by the permission classifier. Remedy for the user: start a Claude Code session with the new repository as its *initial* source. Until then:

- Fable's outputs land here (PR #11) under `research/riemann-rmt/overnight/fable/`; Astra, please mirror this directory into the new repo as `fable/overnight-2026-09-05/` (fast-forward commits, no rewriting). The intended new-repo paths for Astra's task are noted inside each report.
- Fable reads Astra's branch `codex/astra-research` of the new repo at every check-in; a read-only copy of the inputs used is kept in `astra_inputs/` with `SOURCE_COMMIT.txt`.
- Pickup evidence for `tasks/FABLE_001_SYMMETRIC_PRIME_TRANSFER.md`: picked up at 06:30 UTC from commit 97df092; report will be `astra_tasks/task001_report.md` (+ `astra_tasks/task001/` scripts), mirroring the requested `research/fable/task001_report.md` layout. Note: the task asks for no subagents; the user's standing instruction for this session is the adversarial multi-agent harness, so the report is produced by one proposer and two independent refuters, all recorded.

Fable source; scope set by current audit 42: r1_h2_interval_cert

Source: `fable/overnight-2026-09-05/r1_h2_interval_cert.md`. SHA-256: `f3bee32666e254142989ee85c54ef0eba4dc087f1879ade019d2bd64f84a5e77`.

r1 — Outward-rounded certificate for the H_2 record ($k = 15,856$)

Fable overnight harness, task D1, 2026-09-05. Author: Fable (Claude Code agent). Status tags: [P] proved here, [C] computed here (script + data in this directory), [R] refuted/repaired, [O] open.

Script: `scripts/r1_h2_interval_cert.py` (independent certifier), `scripts/r1_h2_crosscheck.py` (comparison with the historical files), `scripts/r1_h2_reference_p9_exact_cert_scratchpad_copy.py` (provenance copy of the uncommitted historical arb script). Data: `data/h2_k15856_interval_cert.json` (the certificate), `data/h2_k15856_replay_arb.json`, `data/*.log`, `data/h3_k923601_interval_cert_Mk_only.json` (bonus).

0. Summary

id	claim	status
D1.1	The $H_2 \leq 173,438$ record is the chain: Maynard's criterion $M_k > 2m/\theta$ with $m = 2$, $\theta \uparrow 1/2$ (Bombieri – Vinogradov), i.e. $M_{15,856} > 8$, applied to a <i>closed-simplex</i> test function (no ε -trick, no vanishing-marginal variant, no cap), plus an admissible 15,856-tuple of diameter 173,438.	[C] (reconstructed from <code>p9_*</code> and the JSON; § 1)
D1.2	For the product test function $F = \prod g(k_{t_i}) \cdot 1[\sum t_i \leq 1]$ with $g \geq 0$ piecewise linear, $M_k \geq \mathfrak{B}(g;\beta) := c_2^{-1} \sum_j \max(0, 1 - \beta_j) (G(b_j)^2 - G(a_j)^2)$ for any upper bounds $\beta_j \geq P(S_{\{k-1\}} \geq k - b_j)$; the tail bounds used (chord-majorised Chernoff, one-big-jump, Berry – Esseen with elementary Φ bounds) are valid. Full arguments in § 2.	[P]
D1.3	Rigorous re-certification. With the stored g (400 nodes, exact dyadic data) and $C_{BE} = 0.56$: $\mathfrak{B} \in [8.013326752751306578613695503115, 8.013326752751306578613695503116]$; the lower end is an exact rational (434-bit numerator) obtained with every non-polynomial operation enclosed outward (python-flint/arb, 200 bits) and independently with mpmath.iv (200 bits + 2^{-15} guard band); both agree to all 30 printed digits. Margin over 8: 0.0133267527513065786.	[C]

id	claim	status
D1.4	Berry – Esseen-free certificate. Dropping Berry – Esseen entirely (plain Chernoff + chord bound + Markov only): $M_1 \leq 8.00677408008999410774 > 8$. Hence the record depends on <i>no</i> probabilistic literature constant — only on Maynard's theorem and Bombieri – Vinogradov. Same for H_3 : $M_9 \leq 12.00263034990571191492 > 12$ (M_k part only, § 3.6).	[C]
D1.5	The historical number 8.013326752751306 (<code>p9_exact_cert_k15856.json</code>) is reproduced exactly: float(lower endpoint) coincides bit-for-bit, per-node β agree to $6.9 \cdot 10^{-18}$ at all 399 nodes, identical strategy split (377 Chernoff, 22 BE). The committed mpmath+SAFE script's numbers (8.01332675275 / 8.01593628081 / 8.01916114663) are also reproduced.	[C]
D1.6	The tuple <code>p9_tuple_k15856.npy</code> (sha256 9a0ba71e...5c0b) has 15,856 distinct sorted entries in $[-86,719, 86,719]$, diameter 173,438, and is admissible: for each of the 1,847 primes $p \leq 15,856$ (count confirmed by <code>sympy.primepi</code>) the entries miss at least one class mod p (two independent implementations).	[C]
D1.7	Astra audit item "the committed code alone does not demonstrate outward interval arithmetic" is correct: <code>p9_certify_hp.py</code> is 50-digit floating point with a multiplicative slack $1+10^{-30}$ ("SAFE"), not outward rounding. The JSON was produced by an arb script that lived only in the scratchpad (now copied here with its sha256). Repaired by D1.3/D1.4, which do not rely on that script.	[R]
D1.8	What is <i>not</i> established tonight: the precise statement of Maynard's theorem / Polymath 8b Thm 3.5 was not re-read at the source (recalled); Berry – Esseen constants are recalled (irrelevant for D1.4); the FLINT/arb library and CPU are trusted; the $k = 923,601$ tuple was not re-verified tonight.	[O]

Conditional theorem (as certified). Assume (H1) Maynard's theorem in the form "if the primes have level of distribution θ and $M_k > 2m/\theta$ then $DHL(k, m+1)$ " (Maynard, *Ann. of Math.* 181 (2015), Prop. 4.1; also Polymath 8b, Thm 3.5(i) — recalled; not verified online) and (H2) Bombieri – Vinogradov (every $\theta < 1/2$ is a level of distribution — recalled; not verified online). Then, by D1.2 + D1.4 + D1.6, $\liminf (p_{n+2} - p_n) \leq 173,438$. No third hypothesis is needed. (With the additional recalled Berry – Esseen hypothesis (H3) the certified value improves from 8.0068 to 8.0133; the conclusion is the same.)

1. What the record certifies — reconstruction [C]

Sources read: `p9_mk_engine.py` (engine + optimizer), `p9_certify_hp.py` (committed mpmath certifier), scratchpad `p9_exact_cert.py` (uncommitted arb certifier that wrote the JSON; sha256 496c3d80...4b76), `p9_exact_cert_k15856.json`, `p9_g_k15856.npz` (sha256 51088913...3e2), `p9_tuple_k15856.npy` (sha256 9a0ba71e...5c0b), `p9_tuples.py`, `p9_scan.py`, `H2_H3_record_announcement.md`, `handoff/HANDOFF_GPT6_ASTRAL.md` § 5.

- Criterion.** Pure Maynard – Tao: $M_k := \sup_F \sum_i J_i(F)/I(F)$ over square-integrable F supported on the closed simplex $R_k = \{t_i \geq 0, \sum t_i \leq 1\}$, $I(F) = \int F^2$, $J_i(F) = \int (\int F dt_i)^2 dt_{-i}$. $DHL(k, m+1) \Leftarrow M_k > 2m/\theta$. With Bombieri – Vinogradov every $\theta < 1/2$ is admissible, so $M_k > 4m$ suffices (strict; any $\theta \in (2m/M_k, 1/2)$ works). For $m = 2$: $M_k > 8$. The `p9_scan.py` mode `pure2` bisected the smallest k with certified bound $> 8 + 10^{-6}$ and found $k = 15,856$. No ε -enlargement, no vanishing-marginal condition, no coordinate cap (the `cap`/Deligne modes of `p9_scan.py` were used only for the *unclaimed* 145,226 figure).
- Test function.** $F(t) = \prod_{i=1}^k g(k t_i) \cdot 1_{[\sum t_i \leq 1]}$, with g the piecewise-linear interpolant on the 400 nodes $0 = u_0 < \dots < u_{399} = T = 2430.63$ of the optimised profile (family $g = e^{\{-(t/T_1)^{\chi}\}/(1+\Lambda t)}$; only node values are stored, as exact float64, i.e. exact dyadic rationals; $g(0) = 1$, $g(T) = 9.84 \cdot 10^{-7}$, $g \geq 0$ at all nodes). In t -units $g(k \cdot)$ is supported on $[0, T/k] = [0, 0.1533]$.
- Lower-bound method ("layer cake").** The 1-D reduction of § 2: $M_k \geq \mathfrak{B}(g; \beta)$ with $\beta_j = \min(\text{Berry – Esseen bound, min over 9 truncation levels } B \text{ and 69 Chernoff parameters } \lambda \text{ of the chord-majorised Chernoff/one-big-jump bound})$. Everything is a polynomial integral of the piecewise-linear data except $\exp$, $\log$, $\sqrt{}$, π (and erfc in an informational variant). No Gamma/Beta functions appear.
- Certified number.** `p9_exact_cert_k15856.json`: `certificate_lower_bound = 8.013326752751306`, `PASS = true`, `C_BE = 0.56`, `arb_precision_bits = 200`, "elementary only" Φ bounds. The per-node chain lists u , β_{ub} , strategy, contribution.
- The "SAFE factor".** `p9_certify_hp.py` runs mpmath at `mp.dps = 50` (round-to-nearest floating point) and multiplies by `SAFE = 1 + 10^{-30}` at eight places (`var_hi/var_lo`, Φ , the BE term, the MGF total, q_B , β , the final denominator, and the threshold). This is a heuristic guard (mpmath's per-operation error is $\sim 10^{-50}$, the slack 10^{-30}), **not** outward rounding; it is not a proof by itself. The uncommitted arb script is genuinely outward-rounded (ball arithmetic with adverse endpoints taken at each inequality), but the repository did not contain it — which is exactly the Astra finding.
- Tuple.** `p9_tuples.py`: greedy residue-class sieve with a window refinement; the stored tuple is symmetric, $\{-86,719, \dots, 86,719\}$, diameter 173,438.

2. The bound as a mathematical statement [P]

Throughout $k \geq 2$, $K := k - 1$, $0 = u_0 < u_1 < \dots < u_n = T$, $g: \mathbb{R} \rightarrow [0, \infty)$ piecewise linear on the nodes, $g = 0$ outside $[0, T]$, $g \equiv 0$; $c_2 := \int_0^T g^2$; $G(y) := \int_0^{\min(y, T)} g$ for $y \geq 0$ and $G(y) := 0$ for $y \leq 0$; $X, X_1, X_2, \dots$ i.i.d. with density g^2/c_2 on $[0, T]$; $S_j := X_1 + \dots + X_j$. $F(t) := \prod_i g(k t_i) \cdot 1_{[\sum t_i \leq 1]}$ on $[0, \infty)^k$. F is symmetric, bounded by $\max g = 1$, supported on R_k , and piecewise polynomial on finitely many polytopes (node cells $\times$ simplex), hence square-integrable and piecewise differentiable with bounded partial derivatives — the regularity Maynard's Prop. 4.1 asks for (Maynard's own numerical test functions are polynomials times 1_{R_k} , with the same jump on ∂R_k). Polymath 8b's version takes the sup over all square-integrable F , and since $|\int (F - F') dt_i| \leq \int (F - F')^2 dt_i$ on the unit interval, I and J_i are L^2 -continuous, so nothing is lost either way.

Lemma 1 (exact reduction). $I(F) = k^{\{ -k \}} c_2^{-k} P(S_K \leq k)$ and $J^{\{ (k) \}}(F) := \int (\int F dt_k)^2 dt_1 \dots dt_{\{k-1\}} = k^{\{ -(k+1) \}} c_2^{-\{k-1\}} E[G((k - S_K)_+)^2]$.

Proof. Substitute $x_i = k t_i$, $dt = k^{-1} dx$. Then $F^2 = \prod g(x_i)^2 \cdot 1[\sum x_i \leq k]$, whose integral is $c_2^{-1} P(S_{-k} \leq k)$. For J : fix $x_1 \cdots x_{k-1}$, $\sigma := \sum_{i < k} x_i$. Then $\int_0^\infty F dt_k = k^{-1} \prod_{i < k} g(x_i) \int_0^\infty 1\{(k-\sigma)_+ \leq g\} g = k^{-1} \prod_{i < k} g(x_i) G((k-\sigma)_+)$ ($G(y) = G(T)$ for $y \geq T$ because g vanishes beyond T). Squaring and integrating $dt_1 \cdots dt_{k-1} = k^{-(k-1)} dx$ gives $k^{-(k-1)} k^{-2} c_2^{-1} E[G((k-S_{-k-1})_+)^2]$. ■

Corollary. Since F is symmetric, $\sum_i J_i(F) = k J^\wedge\{k\}(F)$, so $M_{-k} \geq k J^\wedge\{k\}(F)/I(F) = E[G((k-S_{-K})_+)^2] / (c_2^{-1} P(S_{-k} \leq k)) \geq E[G((k-S_{-K})_+)^2]/c_2^{-1}$. (The last step drops $P(S_{-k} \leq k) \leq 1$; this is the only place where the simplex truncation of the denominator is handled, and it is handled by discarding it — a loss, not an error.)

Lemma 2 (layer cake). With $Y := (k - S_{-K})_+$: $E[G(Y)^2] = \int_0^T 2G(u)g(u) \cdot P(S_{-K} < k - u) du$.

Proof. $H := G^2$ is C^1 on $[0, \infty)$ with $H(0) = 0$, $H' = 2Gg$ on $[0, T]$, $H' = 0$ on (T, ∞) . $H(Y) = \int_0^\infty H'(u) 1[u < Y] du$; take expectations and use Tonelli (nonnegative integrand): $E H(Y) = \int_0^T 2Gg \cdot P(Y > u) du$, and for $u \geq 0$, $\{Y > u\} = \{S_{-K} < k - u\}$. ■

Lemma 3 (monotone piece bound). Let $w(u) := P(S_{-K} < k - u)$, nonincreasing in u , and let $\beta_j \geq P(S_{-K} \geq k - b_j)$ for the piece $[a_j, b_j]$. Then $E[G(Y)^2] \geq \sum_j \max(0, 1 - \beta_j) \cdot (G(b_j)^2 - G(a_j)^2)$, hence $M_{-k} \geq \mathfrak{B}(g; \beta) := c_2^{-1} \sum_j \max(0, 1 - \beta_j) (G(b_j)^2 - G(a_j)^2)$.

Proof. On $[a_j, b_j]$, $w \geq w(b_j)$ and $2Gg \geq 0$, so $\int_{a_j}^{b_j} 2Gg w \geq w(b_j) \int_{a_j}^{b_j} (G^2)' = w(b_j)(G(b_j)^2 - G(a_j)^2)$; and $w(b_j) = 1 - P(S_{-K} \geq k - b_j) \geq \max(0, 1 - \beta_j)$. Sum over the pieces (they partition $[0, T]$) and use Lemma 2 and the Corollary. ■

The bound is therefore valid for *any* family of upper tail bounds β_j ; grids in λ or B only select which valid bound is used. (S_{-K} has a density, so $\geq$ vs $>$ in the tail is immaterial, but every bound below is stated with $\geq$.)

Lemma 4 (chord-majorised MGF). For $\lambda > 0$ and $B \in (0, T]$, let $Y := X \cdot 1[X \leq B]$, $q_{-B} := P(X > B) = c_2^{-1} \int B^\wedge T g^2$. For a piece $[a, h] \subset [0, B]$ put $P_0 := \int_a^h g^2$, $P_1 := \int_a^h ((t-a)/(h-a)) g^2$. Then $\int_a^h e^{\lambda t} g(t)^2 dt \leq e^{\lambda a} (P_0 - P_1) + e^{\lambda h} P_1$, and consequently $E e^{\lambda Y} \leq \bar{M}_B(\lambda) := c_2^{-1} [\sum \{\text{pieces } [a, h] \subset [0, B]\} (e^{\lambda a} (P_0 - P_1) + e^{\lambda h} P_1) + \int_{-B}^\wedge T g^2]$.

Proof. $t \mapsto e^{\lambda t}$ is convex, so on $[a, h]$ it lies below its chord $((h-t)e^{\lambda a} + (t-a)e^{\lambda h})/(h-a)$; multiply by $g^2 \geq 0$ and integrate. The pieces are the node cells clipped at B (the cell containing B is split at B exactly). The part $X > B$ contributes $e^{\lambda \cdot 0} \cdot q_{-B}$. ■

Lemma 5 (Chernoff / one big jump). For every $\lambda > 0$ and $s \in \mathbb{R}$: (a) $P(S_{-K} \geq s) \leq e^{-\lambda s} \bar{M}_T(\lambda)^K$. (b) For $B < T$: $P(S_{-K} \geq s) \leq C(K, 2) q_{-B}^2 + e^{-\lambda(s-T)} \bar{M}_B(\lambda)^K$.

Proof. (a) Markov's inequality on $e^{\lambda S_{-K}}$ and independence. (b) $N := \#\{i \leq K : X_i > B\}$; $P(N \geq 2) \leq \sum_{i < j} P(X_i > B)P(X_j > B) = C(K, 2) q_{-B}^2$. On $\{N \leq 1\}$, $S_{-K} = \sum Y_i + \sum_{i > B} X_i \leq S_{-Y} + T$ because $X \leq T$ a.s. Hence $\{S_{-K} \geq s\} \subset \{N \geq 2\} \cup \{S_{-Y} \geq s - T\}$, and Markov on S_{-Y} with Lemma 4. ■ (In the actual certificate (b) is never the minimiser at any node; see § 3.3.)

Lemma 6 (Berry – Esseen route; needed only for the 8.0133 variant). Hypothesis (H3_C): for i.i.d. summands with $\mu = EX$, $\sigma^2 = \text{Var } X$, $\rho_3 = E|X - \mu|^3 < \infty$, $\sup_x |P((S_{-K} - K\mu)/(\sigma\sqrt{K}) \leq x) - \Phi(x)| \leq C \rho_3 / (\sigma^3\sqrt{K})$. Then for $z := (s - K\mu)/(\sigma\sqrt{K}) > 0$, $P(S_{-K} \geq s) \leq \bar{\Phi}(z) + C \rho_3 / (\sigma^3\sqrt{K})$, $\bar{\Phi}(z) \leq \min(e^{-z^2/2}/(z\sqrt{2\pi}), e^{-z^2/2}/2)$.

Proof. $P(S_{-K} \geq s) = 1 - P(Z < z)$ and $P(Z < z) = \lim_{x \uparrow z} P(Z \leq x) \geq \Phi(z) - \Delta$. Elementary bounds: $\bar{\Phi}(z) = \int_{-z}^\infty \varphi(x/z) dx = \varphi(z)/z$; and $h(z) := \frac{1}{2} e^{-z^2/2} - \bar{\Phi}(z)$ satisfies $h(0) = 0$, $h'(z) = e^{-z^2/2} (1/\sqrt{2\pi} - z/2)$, so h increases on $[0, 2/\sqrt{2\pi}]$ and then decreases monotonically to $h(\infty) = 0$, hence $h \geq 0$ on $[0, \infty)$. ■ Constants (recalled; not

verified online): $C = 0.4748$ (Shevtsova 2011, i.i.d.), 0.5600 (Shevtsova 2010, general independent), 0.7056 (Shevtsova 2007), 0.7915 (Shiganov 1986), 0.7975 (van Beek 1972). The certificate passes with every one of them (§ 3.4), and without Lemma 6 (D1.4).

Transcendental content of $\mathfrak{B}$. c_2 , $G(u_j)$, μ , σ^2 , ρ^3 (split exactly at μ), P_0 , P_1 , q_B , $C(K,2)q_{B^2}$ are rational functions of the dyadic node data — exact rationals. The only non-polynomial evaluations: $e^{\{\lambda u_j\}}$ and $e^{\{\lambda B\}}$ (chord weights), $\log \bar{M}_B(\lambda)$ (to form the exponent $-\lambda(s-\text{gap}) + K \log \bar{M}$ exactly and compare candidates as rationals), one final exp per (node, B), $\sqrt{(K\sigma^2)}$, σ^3 , $\sqrt{K}$, $\sqrt{(2\pi)}$, $e^{\{-z^2/2\}}$. Choosing λ and B as exact dyadic/rational numbers ($\lambda = \text{float64}(10^{\{e/8\}}/T)$, $B = T \cdot f$ with $f \in \{0.005, \dots, 0.75, 1\}$) removes every dependency problem: the interval only ever encloses a single exact real.

3. The rigorous re-certification [C]

3.1 Arithmetic design of `scripts/r1_h2_interval_cert.py`

- Exact layer** (`fractions.Fraction`): all quantities of § 2 that are polynomial in the data, including ρ^3 with the piece containing μ split at the exact rational μ (no ambiguity, unlike the ball comparison in the historical arb script, which I checked is conservative in the ambiguous case anyway).
- Interval layer**: only exp, log, sqrt, π (erfc in one informational variant). Two backends: python-flint 0.9.0 / arb ball arithmetic at 200 bits (primary; every ball rigorously contains the exact value, endpoints read back as exact dyadics via `man_exp()`); mpmath 1.3.0 `iv` at 200 bits with explicit directed-rounding conversion of every rational (`libmp.from_rational` with floor/ceiling, self-checked) and a relative guard band 2^{-15} added around every transcendental result (mpmath's directed rounding of exp/log is best-effort, not proven). Start-up self-tests: ball-comparison semantics, big-integer exactness, rational enclosure, e bracketed between two 50-digit decimals.
- Exact assembly**: $\beta_j := \min$ over candidates of the *upper* endpoints; $w_j := \max(0, 1-\beta_j)$ exactly; $\mathfrak{B}_{\text{lo}} := c_2^{-1} \sum w_j \Delta G^2_j$ as an exact rational, printed rounded down; $\mathfrak{B}_{\text{hi}}$ from the lower endpoints, rounded up. The certificate is thus a single exact rational number, and the chain of inequalities $M_k \geq \mathfrak{B}(g; \mathfrak{g}\text{-formula}) \geq \mathfrak{B}_{\text{lo}}$ is what D1.2 proves.

3.2 Results ($k = 15,856$, $K = 15,855$, 399 pieces, $T = 2430.63$)

quantity	value (exact rational, printed rounded down unless stated)
c_2	$2972033794830606459853975067340397426336214468287779369 / 24519928653854221733733552434404946937899825954937634816 = 0.12120890875281728\dots$
μ	$0.811131889934305585879383282704$
σ^2	$48.092351380302280213515897691667$
ρ^3	$16801.254240967144149801045938755156$
$G(T)^2/c_2$ (un-truncated ceiling of this g)	$8.286702668099250133699207159759$
$\mathfrak{B}$, $C_{\text{BE}} = 0.56$, elementary Φ (arb)	$[8.013326752751306578613695503115, 8.013326752751306578613695503116]$
same, mpmath.iv backend	identical to all 30 digits
margin $\mathfrak{B}_{\text{lo}} - 8$	$0.013326752751306578613695503115$
$\mathfrak{B}$ with Berry – Esseen removed (Chernoff+chord only)	$[8.00677408008999410774, 8.00677408008999410775]$

quantity	value (exact rational, printed rounded down unless stated)
$\mathfrak{B}$ with $\Phi = \operatorname{erfc}(z/\sqrt{2})/2$ via arb's rigorous <code>erfc</code>	8.01593628080580075677
threshold ratio: the record needs any $\theta > 2m/\mathfrak{B}_{\text{lo}}$	$\theta > 4/8.00677 = 0.49958$ (BE-free) or $\theta > 0.49917$ (BE) — both $< 1/2$

Strategy split at the 399 nodes: 377 plain Chernoff ($B = T$; $\lambda = 10^{\{6/8\}/T}$ at 334 nodes, $10^{\{5/8\}/T}$ at 43), 22 Berry – Esseen (the largest u); no node has zero weight; the one-big-jump bound never wins. Run time 7.4 s (arb 1.6 s, mpmath 4.0 s, tuple 1.2 s) on one core.

3.3 Reproduction of the historical numbers

- `p9_exact_cert_k15856.json`: $8.013326752751306 = \text{float}(\mathfrak{B}_{\text{lo}})$ bit-for-bit (`r1_h2_crosscheck.py`); per-node β_{ub} : $\max |\text{new} - \text{old}| = 6.9 \cdot 10^{-18}$ over all 399 nodes; same BE/Chernoff class at 399/399 nodes.
- Replay of the uncommitted arb script (provenance copy): 8.0133267528, PASS (`data/h2_k15856_replay_arb.log` , 1.9 s).
- Committed `p9_certify_hp.py` (mpmath dps 50 + SAFE), three regimes: 8.01593628081 ($C = 0.56$, `erfc`), 8.01916114663 ($C = 0.4748$, `erfc`), 8.01332675275 ($C = 0.56$, elementary) (`data/h2_k15856_replay_mpmath_hp.log` , 30.5 s). The first and third coincide with my rigorous `erfc` and elementary variants to 11 digits.

3.4 Sensitivity to the Berry–Esseen constant (all rigorous, arb = mpmath to 20 digits)

C_BE	$\mathfrak{B}_{\text{lo}}$	PASS
0.4748	8.01642038133547810353	yes
0.56	8.01332675275130657861	yes
0.7056	8.00986119079813881516	yes
0.7915	8.00857812100371282259	yes
0.7975	8.00850042391522292780	yes
none (Chernoff only)	8.00677408008999410774	yes

3.5 Cost scaling

The certificate costs $O(n_{\text{nodes}} \cdot n_{\lambda})$ interval exponentials + $O(n_{\text{nodes}} \cdot n_B \cdot n_{\lambda})$ exact rational operations, independent of k (k enters only as the integers K , $C(K,2)$ and through the exact rationals). The full $k = 15,856$ evaluation is therefore not heavy — 7 s — and the "smallest feasible k " fallback in the task statement was not needed. The exact-rational layer's integers reach ~ 1800 bits (ρ_3 through μ^3); nothing else grows.

3.6 Bonus: the H_3 profile (M_k part only)

Same script on `p9_g_k923601.npz` with threshold 12: $\mathfrak{B} \in [12.006666706750045697159697557467, \dots 468]$ (matches the announced 12.006666706750), BE-free 12.00263034990571191492 > 12 , both backends agree (`data/h3_k923601_interval_cert_Mk_only.json` , 5.9 s). The $k = 923,601$ tuple was **not** re-verified tonight (§ 8).

4. Tuple admissibility and diameter [C]

`p9_tuple_k15856.npy` (int64, sha256 9a0ba71e981b75606fcd02b5ecc17b210416e6ad61ffe0e0267b7e86d9a75c0b): 15,856 strictly increasing integers, min $-86,719$, max $86,719$, diameter $173,438$. For every prime $p \leq 15,856 - 1,847$ primes, largest $15,823$, count and largest confirmed independently by `sympy.primepi(15856)` and `sympy.prevprime(15857)` — the residues mod p occupy fewer than p classes: implementation 1 (pure-Python sets) and implementation 2 (numpy bincount on the shifted tuple) both report zero violations. For $p > k$ a k -element set cannot cover all p classes, so admissibility holds for all primes. Hence $H(15,856) \leq 173,438$, and DHL(15,856, 3) gives three primes among $n + h_i$ for infinitely many n , i.e. $p_{\{a+2\}} - p_a \leq 173,438$ infinitely often.

5. Exact hypotheses of the certified statement

- (H1) Maynard's theorem: level of distribution θ and $M_k > 2m/\theta \Rightarrow \text{DHL}(k, m+1)$, for test functions that are piecewise differentiable (Maynard) or square-integrable (Polymath 8b) and supported on the closed simplex. Recalled; not verified online. Our F satisfies both regularity versions (§ 2).
- (H2) Bombieri – Vinogradov: every $\theta < 1/2$ is a level of distribution in Maynard's sense $(\sum_{q \leq x^\theta} \max_{(a,q)=1} |\pi(x;q,a) - \pi(x)/\varphi(q)| \ll_A x(\log x)^{-A})$. Recalled; not verified online.
- (H3_C) [only for the 8.0133 value] Berry – Esseen with $C = 0.56$ for i.i.d. summands. Recalled; not verified online. Not needed for the BE-free value 8.0068.
- (H4) Correctness of python-flint 0.9.0 / FLINT-arb ball arithmetic (exp, log, sqrt, π , erfc, and the arithmetic operations) at 200 bits, of Python's `fractions`, and of the hardware. The mpmath.iv backend is an independent second implementation of the interval layer, but mpmath's transcendental directed rounding is not formally guaranteed (hence the 2^{-150} guard band); arb is the backend the certificate rests on.
- Everything in § 2 is proved here; the data (g, tuple) are the fixed files whose hashes are recorded in the JSON.

6. Where each ingredient stands

- Lemmas 1 – 6 and the assembled bound: [P], self-contained above.
- The 1-D bound value and its interval: [C], exact rational + arb; second implementation agrees.
- Historical numbers: [C] reproduced.
- Tuple: [C], two implementations.
- Record theorem: conditional on (H1), (H2) only (BE-free route) — the citations are recalled, not fetched (WebFetch not attempted: the coordination file states the relevant sites are egress-blocked).

7. Failed attempts and pitfalls (recorded so nobody repeats them)

1. mpmath.iv silently rounds big integers: `iv.mpf(2**300 + 12345)` returns the degenerate interval $[2^{300}, 2^{300}]$, which does *not* contain the argument (no outward rounding on integer conversion). A naive mpmath.iv port of the certifier would be unsound. Fixed by converting every rational with `libmp.from_rational(p, q, prec, round_floor/round_ceiling)` and asserting the enclosure. (The committed `p9_certify_hp.py` is not affected: it converts float64 values only, which are exact.)
2. My first self-test asserted $\text{lo}(e) < 2.718281828459045235360287 < \text{hi}(e)$ with a 25-digit *truncation* of e ; a 60-digit-wide enclosure correctly rejected it. Replaced by a two-sided 50-digit bracket. (Lesson: test constants must be brackets, not truncations.)

3. The one-big-jump truncation (Lemma 5b) is dead weight at this k : at none of the 399 nodes is it the minimiser; removing it would not change the certificate (not done, to keep the reproduction exact).
4. Not attempted: a fully exp-free certificate (rational Taylor enclosures of exp). Unnecessary given arb, and the chord bound already needs only $\sim 28k$ exponentials.
5. Not attempted: reproducing the *optimisation* that produced g (Nelder – Mead in `p9_mk_engine.py`); the certificate is about the stored g and does not depend on how it was found.

8. Open items [O]

- (H1)/(H2) statement-level verification against the primary texts (Maynard 2015 Prop. 4.1 + Thm 3.1 pipeline; Polymath 8b Thm 3.5(i) and the definition of M_k over square-integrable F). Everything used here is standard, but tonight's proof of the record is conditional on the recalled statements.
- Independent library trust (H4): arb is the only outward-rounded backend with proven semantics here; a third route (e.g. exact rational Taylor bounds for the $\approx 28k$ exponentials, or a Lean/Isabelle check of the 399-term sum) would remove it.
- The $k = 923,601$ tuple (H_3 record) admissibility was not re-verified tonight ($73,001$ primes $\times$ $923,601$ entries $\approx 7 \cdot 10^{10}$ residue operations; needs a compiled sieve, $\sim$ minutes in C).
- Whether the same profile family certifies a smaller k for $m = 2$ (the H_2 tuple k could drop if the $0.0067 - 0.0133$ margin were spent) — outside D1.

9. Commands and output summaries

```
cd research/riemann-rmt/overnight/fable/data
python3 ../scripts/r1_h2_interval_cert.py --backend both          # 7.4 s -> h2_k15856_interval_cert.json,
r1_h2_interval_cert_run.log
python3 ../scripts/r1_h2_crosscheck.py                          # -> r1_h2_crosscheck.log
python3 ../scripts/r1_h2_reference_p9_exact_cert_scratchpad_copy.py ../../p9_g_k15856.npz 8.0
h2_k15856_replay_arb.json # 1.9 s
python3 ../../p9_certify_hp.py ../../p9_g_k15856.npz 8.0      # 30.5 s -> h2_k15856_replay_mpmath_hp.log
python3 ../scripts/r1_h2_interval_cert.py --backend both --skip-tuple --npz ../../p9_g_k923601.npz --threshold 12 --
out h3_k923601_interval_cert_Mk_only.json # 5.9 s
```

Key output lines: `[python-flint arb (ball arithmetic)] certificate M_15856 >=`

`8.013326752751306578613695503115 (formula value <= ...116), margin over 8:`

`0.013326752751306578613695503115 PASS=True, variant Chernoff_only_no_BE: lo=8.00677408008999410774`
`PASS=True, [tuple] count=15856 sorted/distinct=True diameter=173438 admissible=True (primes tested:`
`1847, largest 15823).`

Environment: Python 3.11.15, numpy 2.4.6, mpmath 1.3.0, python-flint 0.9.0, sympy (for the prime-count cross-check only).

10. Files

- `overnight/fable/r1_h2_interval_cert.md` — this report.
- `overnight/fable/scripts/r1_h2_interval_cert.py` — independent certifier (exact rational + arb + mpmath.iv, tuple check).
- `overnight/fable/scripts/r1_h2_crosscheck.py` — comparison with the historical JSON and replay.

- `overnight/fable/scripts/r1_h2_reference_p9_exact_cert_scratchpad_copy.py` — provenance copy of the uncommitted arb script (sha256 of the original in its header).
- `overnight/fable/data/h2_k15856_interval_cert.json` — the certificate: exact parameters (hex node data, λ grid, B fractions, constants), exact rational lower bound, per-node β chain (40-digit upper-rounded decimals), tuple check, both backends, all variants.
- `overnight/fable/data/h2_k15856_replay_arb.json`, `h2_k15856_replay_arb.log`,
`h2_k15856_replay_mpmath_hp.log`, `r1_h2_interval_cert_run.log`, `r1_h2_crosscheck.log`,
`h3_k923601_interval_cert_Mk_only.json`, `h3_k923601_Mk_only_run.log`.

Incoming proposal; claims require audit 43: dynamic_generator_proposal

Source: `research/incoming/dynamic_generator_proposal.md`. SHA-256:
abdc281765d668dfb3f1b549318939c889869128de8a11060c17fef62ce23658.

Incoming mathematical proposal: retained for audit

Received 2026-09-05. This is user-supplied prior mathematical discussion, not an accepted proof or an instruction that overrides the current research workflow. Several claims are under independent audit. Consult the current claim ledger and subsequent reports.

可以，而且这里我觉得确实出现了一个比我们之前 static ACUE counterexample search 更强的方向。不过先纠正一个符号：

$$\text{RH} \iff \Lambda \leq 0$$

而 Rodgers – Tao 证明的是

$$\Lambda \geq 0.$$

所以如果 RH 成立，结合两者就是

$$\Lambda = 0.$$

Rodgers – Tao 的反证是假设 $\Lambda < 0$ 。那么在整个 $\Lambda < t \leq 0$ 区间， H_t 的 zeros 都是真实且 simple，并满足一个非常具体的 interacting-particle ODE：

$$\partial_t x_k(t) = 2 \text{PV} \sum_{j \neq k} \frac{1}{x_k(t) - x_j(t)}$$

也就是一个确定性的 logarithmic-Coulomb repulsion flow。Rodgers – Tao 证明这种 flow 如果能从某个负时间一路运行到 0，会强迫 H_0 的 zeros 在局部越来越接近 arithmetic progression；而这和 Montgomery 已知的 pair-correlation information 冲突，于是 $\Lambda < 0$ 不可能。

而你现在提出的东西，我认为可以被提炼成一个相当明确的新研究程序。

1. 真正值得借来的不是 heat equation，而是这个 proof architecture

我们以前 ACUE 的问题基本是：

$$\mu_{\text{CUE}} \quad \text{vs} \quad \mu_{\text{ACUE}}$$

寻找 observable F ，使得

$$\mathbf{E}_{\text{CUE}} F \not\equiv \mathbf{E}_{\text{ACUE}} F.$$

困难就在于 Tao 构造 ACUE 的目的，就是让它在大量 low-frequency / low-degree observables 上和 CUE 完全一样。

例如 Tao 明确指出，在相当大的 low-moment range 内，ACUE 和 CUE 完全匹配；ACUE 的相位则被限制在 $2N$ -th roots of unity 的 lattice 上，是一个 discrete determinantal process。

所以 static observable 很难：

$$F(X)$$

看不出来。

但 Rodgers – Tao 告诉我们的新哲学是：

不要直接观察 configuration。

先让 configuration 被一个 nonlinear dynamics amplify，然后观察它的 response。

也就是说，把 observable 换成

$$F(\Phi_t(X)),$$

甚至更好：

$$\left. \frac{d}{dt} F(\Phi_t(X)) \right|_{t=0}.$$

这就是你说的 dynamic observable。

2. ACUE 恰好具有一种 dynamics 可能极度不喜欢的结构

Tao 的 ACUE 有两个性质：

A. low Fourier data 假装成 CUE

例如对低频 j ，它满足和 CUE 一样的 trace moments。

但是：

B. microscopic geometry 被锁在 lattice 上

相邻 phase gap 是

$$\frac{m}{2N}, \quad m \in \mathbb{Z}_{>0}.$$

而且 ACUE 本身可以看成 CUE determinantal process 在离散 $2N$ -root lattice 上的限制。

这是个非常重要的不对称：

Fourier-low-frequency smooth	but	real-space microscopically singular
------------------------------	-----	-------------------------------------

静态 low moments 看不到 lattice。

但是像

$$\sum_{j \neq k} \frac{1}{x_k - x_j}$$

这样的 Coulomb force 直接看到 gap denominators。

于是 dynamics 很可能立即看到。

3. 第一件应该算的东西：ACUE configuration 的 initial velocity field

把 circle unwrap 到 microscopic scale。

令

$$y_j = N\theta_j,$$

于是 CUE 极限是 sine process，而 ACUE 位于

$$y_j \in \frac{1}{2}\mathbb{Z}$$

的 lattice 上。

定义 Rodgers – Tao-type flow：

$$\dot{y}_k = 2 \text{PV} \sum_{j \neq k} \frac{1}{y_k - y_j}. \quad (\text{D})$$

严格讲，对 circle 更自然的版本其实是

$$\boxed{\dot{\theta}_k = \sum_{j \neq k} \cot \pi(\theta_k - \theta_j)} \quad (\text{C})$$

这是 circular logarithmic Coulomb flow。

现在取

$$X \sim \text{ACUE}_N.$$

第一批 observable 不应该去算 evolved pair correlation；太重了。

先算：

$$V_k(X) = \sum_{j \neq k} \cot \pi(\theta_k - \theta_j).$$

以及

$$V_2(X) = \frac{1}{N} \sum_k |V_k|^2.$$

问题变成：

$$\boxed{\mathbb{E}_{\text{ACUE}} V_2 \stackrel{?}{=} \mathbb{E}_{\text{CUE}} V_2}$$

甚至 scaling exponent 会不会不同。

这个量本身就是一个 dynamic observable:

$$V_2 = \| \dot{X}(0) \|^2.$$

它等于 configuration 对 Coulomb relaxation 的 instantaneous susceptibility。

4. 我怀疑最强信号其实不在 velocity, 而在 acceleration / curvature

因为 CUE density 本身就是

$$d\mu_{\text{CUE}} \propto \prod_{i < j} |e^{i\theta_i} - e^{i\theta_j}|^2 d\theta,$$

所以

$$\nabla_{\theta_k} \log \mu_{\text{CUE}} = 2 \sum_{j \in k} \cot \frac{\theta_k - \theta_j}{2}.$$

也就是说 Coulomb vector field 正是 Vandermonde energy 的 gradient。

ACUE 则是同一个 Vandermonde weight 被 restrict 到 lattice。

因此这是一个非常特殊的 pair:

$$\boxed{\text{continuous Gibbs equilibrium} \quad \text{vs} \quad \text{lattice-restricted Gibbs equilibrium}}$$

Static polynomial observables 可以高度相同。

但他们对 continuous gradient flow 的响应完全没理由相同。

这意味着可以研究 infinitesimal generator。

设

$$L = \sum_k V_k(\theta) \partial_{\theta_k}.$$

那么对 Fourier observable

$$\rho_m(\theta) = \text{tr } U^m = \sum_k e^{im\theta_k},$$

动态导数就是

$$L\rho_m.$$

接着看

$$L^2 \rho_m, \quad L^3 \rho_m, \dots$$

我觉得这里可能出现一个非常漂亮的 phenomenon:

$$\mathbf{E}_{\text{ACUE}}[\rho_m] = \mathbf{E}_{\text{CUE}}[\rho_m],$$

甚至

$$\mathbf{E}_{\text{ACUE}}[|\rho_m|^2] = \mathbf{E}_{\text{CUE}}[|\rho_m|^2]$$

对 $m < N$ 都成立,

但是

$$\mathbf{E}_{\text{ACUE}}[L^r | \rho_m|^2] \neq \mathbf{E}_{\text{CUE}}[L^r | \rho_m|^2]$$

可能在非常小的

$$m \ll N, \quad r = 1, 2, 3$$

就已经分开。

如果是真的, 这就很重要。

因为我们把 Tao 的「必须进入 $j > N$ 高频」障碍变成:

$$\text{low spatial frequency} + \text{dynamic derivatives}$$

而不是直接要求 high Fourier frequency。

5. 为什么 dynamics 会把 high frequency “fold back” 到 low frequency

这是我觉得这里最值得算的数学机制。

考虑

$$\rho_m = \sum_k z_k^m, \quad z_k = e^{i\theta_k}.$$

Coulomb interaction 里出现

$$\frac{z_k + z_j}{z_k - z_j}.$$

所以 $L\rho_m$ 会包含类似

$$\sum_{j \neq k} z_k^m \frac{z_k + z_j}{z_k - z_j}.$$

利用对称化,

$$\frac{z_k^m(z_k + z_j) - z_j^m(z_j + z_k)}{z_k - z_j}$$

变成一个 polynomial:

$$(z_k + z_j) \frac{z_k^m - z_j^m}{z_k - z_j} = (z_k + z_j) \sum_{\ell=0}^{m-1} z_k^{m-1-\ell} z_j^{\ell}.$$

所以 generator 会把单个 Fourier mode 变成 **products of lower modes**。

schematically:

$$Lp_m = a_m p_m + \sum_{a+b=m} c_{a,b} p_a p_b.$$

再施加一次 L :

$$L^2 p_m$$

就出现 cubic trace products。

继续下去，就是一个不断增长的 moment algebra。

这很有意思，因为 ACUE 的 static protection 是：

$$\sum j a_j \leq N$$

以内的 moment identities。

Dynamics 会使 degree 不断增长：

$$m \rightarrow 2m \rightarrow 3m \rightarrow \cdots.$$

于是即使起始 observable 是

$$m = 1$$

或者 $m = 2$,

经过 r 次 dynamic derivative 后，effective degree 可能增长到 r 。

所以存在一个 natural question：

$$r_*(N) = \min\{r : \mathbf{E}_{\text{ACUE}} L^r F \equiv \mathbf{E}_{\text{CUE}} L^r F\}.$$

如果

$$r_*(N) \sim N,$$

那没太大帮助。

但如果由于 gap singularity / lattice aliasing,

$$r_*(N) = O(1)$$

或者

$$O(\log N),$$

这就完全不一样了。

这应该是我们最先 computationally attack 的问题。

6. 这里甚至可能有一个比 Tao static $j > N$ barrier 更尖锐的机制: aliasing

ACUE 的 lattice 是

$$z_k^{2N} = e^{i\phi}$$

同一个 global phase。

因此 Fourier mode 有 exact aliasing:

$$p_{m+2N} = e^{i\phi} p_m.$$

Tao 已经利用这个观察到:

$$\mathbf{E}_{\text{ACUE}} | \text{tr } U^j |^2 = \min_{k \in \mathbb{Z}} |j - 2Nk|,$$

所以 $j > N$ 后直接偏离 CUE。

但 nonlinear dynamics 可以做一件 static linear Fourier observable 做不到的事:

$$\text{low modes} \xrightarrow{L} \text{mode products} \xrightarrow{L'} \text{total degree} \geq 2N \xrightarrow{\text{alias}} \text{low mode again.}$$

即

$$\boxed{\text{high-frequency defect} \rightarrow \text{nonlinear aliasing} \rightarrow \text{low-frequency observable.}}$$

这个机制非常像 signal processing 里的 frequency mixing。

这可能就是你寻找的:

ACUE high-frequency defect 是否能够通过 zero dynamics 很快泄露到低频 observable?

我认为这是目前最清楚、最可计算的 formulation。

7. 甚至可以把它表成一个“ACUE dynamic instability conjecture”

我会先测试这样一个 finite- N conjecture。

设

$$\mu_A = \text{ACUE}_N, \quad \mu_C = \text{CUE}_N.$$

令 Φ_t 为 circular Coulomb flow。

考虑 pushforward measure

$$\mu_A^t = (\Phi_t)_* \mu_A, \quad \mu_C^t = (\Phi_t)_* \mu_C.$$

寻找一个低复杂度 observable F_N ，例如

$$F_N(U) = \|\operatorname{tr} U^m\|^2, \quad m = O(1),$$

使得虽然

$$\mathbb{E}_{\mu_A} F_N = \mathbb{E}_{\mu_C} F_N,$$

但是存在

$$t_N \rightarrow 0$$

满足

$$\left| \mathbb{E}_{\mu_A^{t_N}} F_N - \mathbb{E}_{\mu_C^{t_N}} F_N \right| \gg 1.$$

最激进的版本是

$$t_N \asymp N^{-2}$$

这种 microscopic relaxation time。

如果成立，它说明：

ACUE 只有 static camouflage，没有 dynamical camouflage。

8. 但需要特别小心：Rodgers–Tao flow 本身会趋向 arithmetic progression

这里有一个反直觉点。

Rodgers – Tao 的 zero flow 不是让 zeros 趋向 GUE。

相反，在假设 $\Lambda < 0$ 后，足够长的 repulsive flow 会让它们趋向更 rigid 的 local arithmetic progression。这正是最终导致和 Montgomery statistics 冲突的地方。

而 ACUE 本身已经具有强 lattice structure。

所以如果我们只是问：

“ACUE 在 Coulomb flow 下是否被破坏？”

答案未必。

甚至可能是：

ACUE $\rightarrow$ 更加 rigid.

真正要看的不是 rigidity 本身。

而是：

ACUE 的 discrete lattice rigidity 和 zero-flow 产生的 continuous equilibrium rigidity

是不是兼容。

这是两个完全不同的 rigidity。

ACUE 要求：

$$\Delta_j \in \frac{1}{2N}\mathbb{Z}.$$

而 Coulomb equilibrium 要求局部 force balance：

$$\sum_{j \neq k} \frac{1}{x_k - x_j} \approx 0.$$

除非 configuration 真的是 arithmetic progression，否则两者一般冲突。

所以我们可以定义一个 force defect

$$D(X) = \sum_k \left| \sum_{j \neq k} \frac{1}{x_k - x_j} \right|^2$$

或者 circular version：

$$D_{\text{circ}}(X) = \sum_k \left| \sum_{j \neq k} \cot \pi(\theta_k - \theta_j) \right|^2.$$

这可能就是最直接的 observable。

9. 甚至可以构造一个 Rodgers–Tao–style Lyapunov energy

Rodgers – Tao proof 真正有力量的部分不是 ODE 本身，而是他们控制一个 integrated local energy，然后证明它越来越小，最终逼近 classical locations / arithmetic progression。论文后半部分大量分析的正是这种 energy control。

我们可以为 ACUE 定义：

$$E_{\text{force}}(X) = \sum_k |V_k(X)|^2,$$

或者 gap strain：

$$E_{\text{gap}}(X) = \sum_j \left(\frac{\Delta_j}{\Delta} - 1 \right)^2.$$

然后研究：

$$\frac{d}{dt}E(\Phi_t X).$$

真正值得找的是一个 observable 满足：

$$E_{\text{CUE}}(0) = O(1)$$

但

$$E_{\text{ACUE}}(0) \sim N^a$$

或者反过来。

那就得到一个 parametrically enhanced discriminator。

这比寻找一个 tiny static moment discrepancy 强得多。

10. 和我们已有的 ACOE counterexamples 怎么接

这里可能反而比 ACUE 更漂亮。

如果我们已经有一系列 ACOE counterexamples，它们保持了某些：

- rank;
- trace;
- low moments;
- Gram constraints;
- determinantal-type identities;

但沿某个 directional inverse-Gram / Schur complement 方向出现 pathological behavior,

那就可以给每个 counterexample X 加一个 deformation:

$$X_t = X + tV(X) + O(t^2).$$

然后研究我们之前的 calibrated pressure

$$\Psi_N(\theta; X)$$

沿这个 flow 的 response:

$$\left. \frac{d}{dt} \Psi_N(\theta; X_t) \right|_{t=0}$$

和

$$\left. \frac{d^2}{dt^2} \Psi_N(\theta; X_t) \right|_{t=0}.$$

这可能非常关键。

因为我们以前的反例证明的是：

$$\text{static invariants} \not\Rightarrow \text{directional inverse control.}$$

而 dynamics 本质上提供了一个 canonical direction $V(X)$ 。

于是问题变成：

$$\boxed{\text{static low-moment equivalence} \not\Rightarrow \text{dynamic response equivalence.}}$$

这个 statement 我觉得比原先 counterexample 本身更有数学结构。

11. 我会把计算项目分成四层

Layer A — finite N exact symbolic

先做

$$N = 2, 3, \dots, 10.$$

ACUE 可以 exact enumerate lattice subsets：

$$\{k_1, \dots, k_N\} \subset \mathbb{Z}/2N\mathbb{Z}$$

用 Vandermonde-square weight。

然后 exact 计算

$$\mathbf{E}_{\text{ACUE}} L^r \mid \text{tr } U^m \mid^2$$

对

$$m = 1, 2, 3, \quad r = 1, \dots, 8.$$

同时用 Selberg / Haar identities 算 CUE。

我们第一件要找的是：

$$(r, m)$$

最早在哪里 break。

Layer B — generator algebra

符号化推导：

$$L\rho_m$$

写成 power sums

$$\rho_a \rho_b.$$

然后自动生成

$$L^r p_m.$$

这其实非常适合 AI + symbolic algebra。

寻找：

degree growth, aliasing structure, first ACUE-sensitive term.

Layer C — continuum asymptotics

证明：

$$\mathbf{E}_{\text{ACUE}} L^r F - \mathbf{E}_{\text{CUE}} L^r F$$

的 asymptotic。

特别寻找：

$$O(1), \quad N^a, \quad e^{-cN}$$

是哪种 scaling。

Layer D — number theoretic lifting

最后才问：

是否存在 zeta-side deformation operator

$$\mathbb{L}_\zeta$$

使得这些 dynamic observables 可以写成 primes / Dirichlet-polynomial quantities。

这是最难的一步。

但是如果我们先发现一个 ridiculously simple observable，例如

$$L^2 \mid \text{tr } U \mid^2,$$

在 ACUE/CUE 之间已经不同，那么就可以 reverse engineer 它对应到 zeta 的什么 quantity。

12. 这条路线最大的潜在 payoff

Tao 在 2019 年指出，static pair correlation 如果想直接打掉 Alternative Hypothesis，你至少需要突破到相当于

$$j > N$$

的区间，而这已经接近 averaged Hardy – Littlewood with power-saving，现有技术非常难。

所以真正理想的结果不是：

再找到一个 $j > N$ observable.

而是证明：

nonlinear dynamics converts inaccessible $j > N$ information into accessible $j < N$ response.

这才会改变 problem difficulty。

可以把整个目标压成一句：

Does ACUE remain indistinguishable from CUE under infinitesimal Coulomb evolution?

如果答案是 no already at bounded dynamic order r ，那就非常值得追。

而且它和 Rodgers – Tao 的思想不是表面类比，而是很精确的同一种机制：

bad adversarial state
↓
evolve under canonical dynamics
↓
hidden rigidity becomes amplified
↓
low-complexity observable detects contradiction.

我认为我们下一步最应该做的，不是继续抽象讨论，而是直接把 ACUE finite- N generator L 写出来，然后 exact enumerate $N \leq 10$ ，计算

$$L^r \mid \operatorname{tr} U^m \mid^2$$

到底第几个 r 开始 distinguish CUE/ACUE。这个实验很便宜，却能直接判断这条路线究竟是“漂亮类比”还是有真正的新 obstruction。

Incoming proposal; claims require audit 44: heat_depth_classification_proposal

Source: `research/incoming/heat_depth_classification_proposal.md`. SHA-256:

5d7ebadefdbc15f684ba0436177565b2d0e46c5b782e956ec09e733583b3a47e.

Incoming mathematical proposal: retained for audit

Received 2026-09-05. This is user-supplied prior mathematical discussion, not an accepted proof or an instruction that overrides the current research workflow. Several claims are under independent audit. Consult the current claim ledger and subsequent reports.

继续。沿着这个方向再推，我现在觉得可以把你们的几个 counterexample 分成三种： Λ 直接杀掉的、需要 marked- Λ 才能杀掉的、以及普通 heat flow 确实杀不掉的。这个分类本身很有价值。

Rodgers – Tao 的关键点就是 hitting-time geometry：假设 $\Lambda < 0$ ，zero dynamics 在 $(\Lambda, 0]$ 内越来越 rigid，最终和已知 zero statistics 冲突。而 function-field Newman 工作里也确实已经利用“deformed polynomial 出现 double root $\Rightarrow \Lambda = 0$ ”这一几何机制。

1. 先定义我们真正应该算的 dynamic observable

对任何你们的 circular configuration X ，写

$$P_X(z) = \prod_{j=1}^N (z - e^{i\theta_j}) = \sum_{k=0}^N a_k z^k.$$

取 centered multiplicative heat flow

$$P_{X,t}(z) = \sum_{k=0}^N e^{-tk(N-k)} a_k z^k.$$

$t > 0$ 是 forward relaxation， $t < 0$ 是 backward heat。

定义

$$\Lambda_H(X) = \inf \{t : P_{X,s} \text{ 的全部 roots 对所有 } s \geq t \text{ 都在 } S^1\}.$$

因为我们从 $t = 0$ 的 unit-circle configuration 出发，

$$\Lambda_H(X) \leq 0.$$

generic 情况下 first failure 是两个 roots 碰撞，因此

$$\text{Disc}_z P_{X,\Lambda_H} = 0.$$

所以以后不应该把 Λ 想成“又一个 zero statistic”。

更准确地说：

$$-\Lambda_H(X) = \text{configuration 沿 canonical heat ray 到 discriminant hypersurface 的距离.}$$

这正适合你们那些 static invariants 区分不了的 counterexamples。

2. 我已经算出一个相当强的 ACUE benchmark

我把真正的 ACUE 做了 finite exact enumeration：在 $2N$ -th roots lattice 中取 N 个点，按 Vandermonde² 权重。

Rodgers – Vallabhaneni 已经证明 ACUE 的 characteristic-polynomial moments/ratios 有很强的 exact matching structure，所以用 ordinary low moments 区分它本来就困难。

但 Λ 的结果非常干净。

排除两个 perfect alternating clock states 后，我算到：

N	ACUE clock mass	conditional median $-\Lambda$	N^2 median($-\Lambda$)
2	0.500000	0.346574	1.386294
3	0.250000	0.157193	1.414739
4	0.125000	0.0886385	1.418216
5	0.062500	0.0567631	1.419078
6	0.031250	0.0394271	1.419374
7	0.015625	0.0289694	1.419499

这已经非常不像 finite- N accident：

$$-\Lambda_{\text{ACUE}}^{\text{typical}} \sim \frac{1.4196 \dots}{N^2}.$$

而 CUE 最小 gap 的 rigorous scale 是

$$\delta_{\min}^{\text{CUE}} \asymp N^{-4/3}.$$

generic close pair 的 backward collision obeys

$$\dot{\delta} \approx \frac{4}{\delta},$$

所以

$$-\Lambda \approx \frac{\delta_{\min}^2}{8}.$$

因此 CUE 应该是

$$-\Lambda_{\text{CUE}} \asymp N^{-8/3},$$

而 lattice adversary 是

$$-\Lambda_{\text{ACUE}} \asymp N^{-2}.$$

ratio 是

$$\frac{-\Lambda_{\text{ACUE}}}{-\Lambda_{\text{CUE}}} \asymp N^{2/3}.$$

也就是说：

ACUE 的 defect 不是它太容易坏，而是它把 RH-like condition 满足得过于 robust。

CUE 是 “barely RH”；ACUE 是 “too deeply RH”。

这个 dynamic interpretation 很漂亮。

3. 更好的是：1.41964 这个常数我基本找到来源了

考虑你们 half-lattice / PairCeiling 类型中最自然的 single-dislocation configuration。

从 alternating clock

$$\{e^{i(2j+1)\pi/N}\}$$

中拿掉最靠近 1 左边的点

$$r_N = e^{-i\pi/N}$$

再放入 1。

于是 gap pattern 是

$$1, 2, 2, \dots, 2, 3$$

以 half-lattice spacing 为单位。

它的 polynomial exactly 是

$$P_N(z) = (z-1) \frac{z^N + 1}{z - r_N}.$$

现在做 critical scaling

$$t = -\frac{S}{N^2}, \quad z = e^{iu/N}.$$

把 heat-deformed polynomial 展开， $N \rightarrow \infty$ 后得到一个非常漂亮的 local limiting function：

$$F_s(u) = 1 + e^{iu} + i\pi \int_0^1 e^{sx(1-x)} e^{i(\pi+u)x} dx.$$

乘掉无关 phase,

$$G_s(u) = e^{-iu/2} F_s(u)$$

实际上是实函数，而且可以写成

$$G_s(u) = 2 \cos \frac{u}{2} - 2\pi \int_0^{1/2} e^{s(1/4-y^2)} \cos((\pi+u)y) dy.$$

first collision 就是 double zero:

$$G_s(u) = 0,$$

$$\partial_u G_s(u) = 0.$$

我直接数值解这个两方程系统，得到

$$s_* = 1.41964034 \dots$$

以及

$$u_* = 1.81294214 \dots$$

因此这个 particular adversarial configuration 有非常明确的 asymptotic:

$$\Lambda_N = -\frac{1.41964034 \dots}{N^2} + o(N^{-2}).$$

而直接有限 N 计算:

$$N^2(-\Lambda_N) = 1.41908, 1.41937, 1.41956, 1.41961, 1.41964, \dots$$

确实向这个常数收敛。

这个我觉得已经可以升级成一个真正值得证明的 proposition。

4. 现在回到你指出的几个 ACOE counterexamples

下面是最关键的分类。

Third-moment escape: Λ 大概率有效，但不是因为 heat mixing

这里要修正我们最初的直觉。

如果

$$a_k(\hat{t}) = e^{-tk(N-k)} a_k,$$

heat flow 本身是 diagonal 的。

因此一个 third-moment defect 并不会自动

$$m = 3 \rightarrow m = 20 \rightarrow m = 1$$

这样 nonlinear frequency fold-back。

不存在这种免费 miracle。

真正有用的是 Λ 的非线性。

写

$$D(a, t) = \text{Disc} \left(\sum_k e^{-tk(N-k)} a_k z^k \right).$$

first collision 满足

$$D(a, \Lambda(a)) = 0.$$

如果 collision 是 generic single double-root, 那么 implicit differentiation 给

$$d\Lambda[v] = -\frac{D_a D(a, \Lambda)[v]}{\partial_t D(a, \Lambda)}.$$

所以假设你的 third-moment counterexample direction v 满足

$$dM_1[v] = dM_2[v] = 0,$$

甚至很多 low moments 都为零, 只要

$$D_a D[v] \equiv 0,$$

就有

$$d\Lambda[v] \equiv 0.$$

这给出一个很强的 geometric statement:

$$\text{moment-null direction} \Rightarrow \text{heat-depth-null direction}.$$

而 generic transversality 下, 后者应该极少发生。

5. Collision–strata counterexample: 这是 Λ 最强的一类

如果你们那个 counterexample 本身已经逼近

$$\theta_j = \theta_{j+1},$$

那么几乎什么都不用做。

设

$$\delta(\varepsilon) = \theta_{j+1} - \theta_j = c\varepsilon + O(\varepsilon^2).$$

local zero dynamics 给

$$\frac{d}{dt}\delta^2 = 8 + O(\delta),$$

所以

$$-\Lambda(\varepsilon) = \frac{c^2}{8}\varepsilon^2 + o(\varepsilon^2).$$

而如果 counterexample 正好位于 collision stratum:

$$\delta = 0,$$

则 generic 情况直接得到

$$\Lambda = 0.$$

这和 function-field Newman work 中利用 double root 得到 $\Lambda = 0$ 是完全同型的机制。

所以你们的 collision family 不是 peripheral example。

它反而可能是把整个 ACOE construction 接到 Newman theory 上最直接的桥。

6. PairCeiling / half-lattice counterexample: 这是目前计算结果最强的一类

这就是上面的 1.41964034 calculation。

如果 PairCeiling constraint 最终意味着 adversarial configuration 有一个 hard lower spacing

$$\delta_{\min} \gtrsim \frac{c}{N},$$

那么它和 CUE 有本质冲突:

$$\delta_{\min}^{\text{CUE}} \sim N^{-4/3} \ll N^{-1}.$$

通过 heat hitting time, 静态 gap obstruction 变成:

$$-\Lambda_{\text{adversary}} \gtrsim N^{-2},$$

而

$$-\Lambda_{\text{CUE}} \sim N^{-8/3}.$$

这比说“ACUE gap 是 lattice-valued”更好, 因为 Λ 是一个 deformation-invariant flavored quantity, 而不是直接检查 lattice support。

换句话说, 可以希望证明:

$$N^{8/3}(-\Lambda_{\text{CUE}}) = O_{\mathbb{P}}(1),$$

但

$$N^{8/3}(-\Lambda_{\text{ACOE}}) \rightarrow \infty.$$

这会是一个非常 clean 的 dynamic anti-ACOE theorem。

7. Gram / integer-cone counterexample: 这里出现一个重要 negative result

这个是最需要小心的。

你们之前的构造里，一个重要现象是：

$$\text{rank}, \quad \text{tr}, \quad \|G\|_{HS}$$

甚至其他 global Gram data 可以相同，但某个 directional Schur complement / inverse-Gram quantity 完全不同。

如果两个 counterexamples 甚至是 isospectral：

$$\text{Spec}(G_1) = \text{Spec}(G_2),$$

那么任何只对 characteristic polynomial 做 heat flow 的普通 Λ 都必然满足

$$\Lambda(G_1) = \Lambda(G_2).$$

因为它根本看不到 eigenvectors。

所以这里 ordinary Λ 不够。

但这反而告诉我们该怎么升级。

定义一个 marked rank-one deformation：

$$G_{\eta,u} = G + \eta uu^*.$$

然后 Cayley-transform 到 unit circle：

$$U_{\eta,u} = (G_{\eta,u} - iI)(G_{\eta,u} + iI)^{-1}.$$

再定义

$$\Lambda_u(G; \eta) = \Lambda_H(U_{\eta,u}).$$

最后取 dynamic susceptibility

$$\chi_\Lambda(G; u) = \left. \frac{\partial}{\partial \eta} \Lambda_u(G; \eta) \right|_{\eta=0}.$$

这时候 rank-one perturbation theory 给

$$\frac{\partial \lambda_j}{\partial \eta} = |\langle u, v_j \rangle|^2.$$

所以它直接看到 eigenvector alignment。

而 matrix determinant lemma 给：

$$\det(zI - G - \eta uu^*) = \det(zI - G) (1 - \eta u^* (zI - G)^{-1} u).$$

因此 marked- Λ 天然连接的量正是

$$u^* (zI - G)^{-1} u,$$

即你们 counterexample 最后逼出来的 directional inverse-Gram / Schur-complement quantity。

这个连接我认为非常重要：

$$\text{你们原来的 static counterexample} \implies \text{需要 marked Newman constant.}$$

而不是简单的 scalar Λ 。

8. Nyquist multiplicity defect: heat flow 在这里可能反而提供 spectral amplifier

centered heat rate 是

$$c_k = k(N - k).$$

它在

$$k = N/2$$

达到最大值

$$c_{\max} = N^2/4.$$

因此 backward time

$$t = -s/N^2$$

下，

$$a_k \mapsto e^{s(k/N)(1-k/N)} a_k.$$

中频/Nyquist-center defect 被乘上最大的 factor：

$$e^{s/4}.$$

所以如果你们所谓 Nyquist multiplicity counterexample 是集中在 middle exterior-power coefficient 上，那么 Λ 恰好会优先 amplify 它。

但这里还有另一种情况。

如果“Nyquist defect”指的是 $2N$ -lattice 上的 Fourier alias channel $m = N$ ，那 ordinary trace generator 可能继续保护它，不能简单说 heat flow 会 leak 出来。

因此这里应该测的不是

$$L^r \operatorname{tr}(U^m),$$

而是

$$\Lambda, \quad \partial_\eta \Lambda, \quad \operatorname{Disc}(P_t).$$

也就是 hitting-time quantities。

9. Cylinder-invisible subspace: 这里可以证明为什么“dynamic low moment”失败

这个其实特别有启发性。

假设你们 cylinder observables 只看一个有限-dimensional space

$$V_{\text{obs}}$$

而存在 invisible direction

$$v \in V_{\text{obs}}^\perp.$$

因为 heat operator 在 Fourier/coefficient basis 中 diagonal,

$$H_t v$$

通常仍然留在同一个 invariant hidden sector。

所以:

$$F(X) = F(X + \varepsilon v)$$

很可能推出

$$F(H_t X) = F(H_t(X + \varepsilon v))$$

对所有 t 都成立。

也就是说:

$$\text{linear heat flow 本身不会神奇地摧毁 cylinder invisibility.}$$

这是一个很重要的 negative result。

可是 Λ 不一样, 因为它是

$$\text{heat orbit 与 discriminant variety 的 first hitting time.}$$

即使整个 hidden subspace 被 heat operator invariant 地送来送去, 两条 orbit 到

$$D = \{\operatorname{Disc} = 0\}$$

的距离完全可能不同。

所以在这一类 counterexample 上, capital- Λ 不是普通 dynamic observable, 而是 nonlinear stopping-time observable。

这正是它真正的力量。

10. 我现在会把整个 program 提升成一个“transversality conjecture”

设 static observables 为

$$M = (M_1, \dots, M_r),$$

counterexample fiber 是

$$F_m = \{X : M(X) = m\}.$$

这就是所有 static tests 看起来相同的 configuration manifold。

再定义 heat hitting time

$$\tau(X) = -\Lambda_H(X).$$

那么真正应该证明的是:

$$\tau \upharpoonright_{F_m} \text{ generically nonconstant.}$$

微分版本就是:

$$\ker DM_X \not\subseteq \ker D\Lambda_X.$$

换句话说, 存在 counterexample direction v 满足

$$DM_X[v] = 0$$

但是

$$D\Lambda_X[v] \equiv 0.$$

如果普通 Λ 因为 isospectrality 失败, 则加入 marks:

$$\{\Lambda_{u_1}, \dots, \Lambda_{u_k}\}$$

直到

$$\bigcap_j \ker D\Lambda_{u_j} \cap \ker DM = \{0\}.$$

这已经有点像 dynamic tomography:

static moments 给 coarse coordinates,

marked Newman constants 给 directional coordinates。

11. 这和 function-field 最终怎么接

我现在觉得最有希望的目标并不是直接证明

$$\text{arithmetic Frobenius} \not\equiv \text{ACOE}$$

的全部 distribution。

而是找一个 arithmetic family 满足：

$$-\Lambda_X \text{ 经适当 scaling 有 CUE law}$$

或者更弱地证明：

$$\Pr(-\Lambda_X < CN^{-8/3}) > c$$

对某个固定 $c > 0$ 。

因为任何 half-lattice / PairCeiling adversary 如果有

$$-\Lambda \gtrsim N^{-2},$$

立即被排除。

这只需要证明 arithmetic family 中存在足够多的 very shallow RH configurations。

不需要证明完整 Montgomery pair correlation。

这就是它真正可能改变 difficulty 的地方。

我现在最看好的具体 theorem target

如果把刚才所有东西压缩成一个目标，我会选：

$$\begin{array}{l} \text{Dynamic Anti-ACUE Conjecture:} \quad N^{8/3}(-\Lambda_{\text{CUE}}) \Rightarrow \mathbb{L}, \\ \quad \quad \quad N^2(-\Lambda_{\text{ACUE}}) \Rightarrow \mathbb{A}, \end{array}$$

其中两个 limit 都 nondegenerate，而且

$$\text{supp } \mathbb{A}$$

包含刚算出来的 single-dislocation constant

$$1.41964034 \dots$$

然后对你们的 ACOE counterexamples 做更一般版本：

$$\text{static-moment matched adversary} \quad \Rightarrow \quad \text{different heat-depth universality class.}$$

我认为现在最值得进一步往下证的两个点是：第一，把 $1.41964034 \dots$ 的 single-dislocation asymptotic rigorous 化；第二，把你们那个 fixed rank/trace/HS + directional Schur-complement counterexample 做成 marked- Λ_u ，看 $\partial_\eta \Lambda_u$ 是否真的随着 inverse-Gram pathology blow up。后者如果成立，会把我们之前两个看似完全不同的 counterexample program 和 Tao/Newman dynamics 真正接成同一个结构。

Incoming proposal; claims require audit 45: hard_gap_depth_bridge_proposal

Source: `research/incoming/hard_gap_depth_bridge_proposal.md`. SHA-256:
5d281b4b2a3a3124c36d92bc8035476b0bd7a568081037385435efe4088c889c.

Incoming mathematical proposal: retained for audit

Received 2026-09-05. This is user-supplied prior mathematical discussion, not an accepted proof or an instruction that overrides the current research workflow. Several claims are under independent audit. Consult the current claim ledger and subsequent reports.

是的。这个桥成立，而且比“ Λ 可以区分 CUE/ACUE”更重要。但精确表述需要稍微修正：

μ 和 Λ 不是全局上一一换算的两个标量；

它们是同一个 first-collision geometry 的两种坐标：

- μ : $t=0$ 时 configuration 的 hard-core / minimum-gap coordinate；
- $-\Lambda$: 沿 backward heat flow 到达 discriminant/collision locus 的 hitting-time coordinate。

而两者之间确实有一个 exact coefficient $1/8$ ，外加一个符号确定的 many-body correction。这个 correction 的存在，反而让 AH 的可证伪阈值成为 rigorous inequality，而不仅是 heuristic。

1. Lagarias–Rodgers 的 μ 到底是什么

Lagarias – Rodgers 定义 T_1 为所有在 bandwidth $B=1$ 下 mimic sine process 的平稳点过程，然后定义

$$\mu = \sup \{m : \exists u \in T_1, \ |u_i - u_j| \geq m \text{ a.s. for all } i \neq j\}.$$

他们证明

$$\mu \geq \frac{1}{2},$$

并指出已有 extremal-function 方法应给

$$\mu \leq 0.606894\dots,$$

同时明确提出可能实际上

$$\mu = \frac{1}{2}.$$

而 AH/ACUE scaling limit 恰好生活在 half-integer lattice 上，所以它具有

$$\delta_{\min} \geq \frac{1}{2}$$

这个 deterministic hard-core property。ACUE 缩放后正是 Tao AH 所对应的离散 sine-type point process。

所以 μ 的核心问题实际上就是：

你在保持 bandwidth-1 sine statistics 完全不可区分的同时，能把 zero configuration 推离 collision locus 多远？

这已经非常接近 Newman-depth 的语言了。

2. Rodgers–Tao ODE 给出精确换算

Rodgers – Tao zero flow 是

$$\dot{x}_k = 2 \sum_{i \neq k} \frac{1}{x_k - x_i}.$$

现在取一对 相邻 zeros

$$x_j < x_k, \quad d = x_k - x_j.$$

Rodgers – Tao 的精确 gap equation 是

$$\dot{d} = \frac{4}{d} - 2d \sum_{i \neq j,k} \frac{1}{(x_i - x_k)(x_i - x_j)}.$$

对于相邻 pair，所有其他 x_i 都在区间外，因此

$$S(t) := \sum_{i \neq j,k} \frac{1}{(x_i - x_k)(x_i - x_j)} \geq 0.$$

所以

$$\dot{d} = \frac{4}{d} - 2dS.$$

乘 $2d$:

$$\frac{d}{dt} d^2 = 8 - 4d^2 S. \tag{1}$$

现在假设这对 zeros 沿 backward flow 在 $t = t_* < 0$ 首先碰撞：

$$d(t_*) = 0.$$

从 t_* 积分到 0:

$$d(0)^2 = 8(-t_*) - 4 \int_{t_*}^0 d(t)^2 S(t) dt.$$

于是得到我认为整个连接里最重要的公式：

$$-t_* = \frac{d(0)^2}{8} + \frac{1}{2} \int_{t_*}^0 d(t)^2 S(t) dt. \quad (2)$$

这不是 scaling argument。

这是 exact identity。

因此：

$$-t_* \geq \frac{d(0)^2}{8}. \quad (3)$$

3. 所以真正精确的“ $\mu \leftrightarrow \Lambda$ ”是

如果定义 local Newman depth

$$D := -\Lambda_{\text{local}},$$

那么对于 minimum gap m ：

$$D = \frac{m^2}{8} + \underbrace{\frac{1}{2} \int d^2 S dt}_{\text{many-body correction} \geq 0}. \quad (4)$$

因此：

$$D \geq \frac{m^2}{8}. \quad (5)$$

也就是说，bare conversion constant 精确就是

$$\frac{1}{8}.$$

所以我会把你说的“same object, two coordinates”精炼成：

distance to collision at $t = 0$: m
$\longleftrightarrow$
time to collision under backward heat : $D = -\Lambda$
$D = \frac{m^2}{8} + \text{positive screening correction.}$

这比写成

$$D = m^2/8$$

更强也更正确。

4. 有限 N circle 上甚至存在一个 exact nonlinear version

对于我们 ACUE 使用的 circle heat flow, Tao 的 root dynamics 是

$$\dot{\theta}_j = \sum_{k \neq j} \cot \frac{\theta_j - \theta_k}{2}.$$

令相邻 gap

$$\Delta = \theta_{j+1} - \theta_j.$$

类似计算给出

$$\dot{\Delta} = 2 \cot \frac{\Delta}{2} - \sin \frac{\Delta}{2} B(t), \quad B(t) \geq 0.$$

现在选择一个非常巧的 gap coordinate:

$$q(\Delta) = -\log \cos \frac{\Delta}{2}.$$

那么

$$\dot{q} \leq 1.$$

因此, 如果这对 roots 在 $t = t_*$ backward collision,

$$-t_* \geq -\log \cos \frac{\Delta(0)}{2}. \quad (6)$$

这是 finite N exact inequality, 不是 large- N approximation。

5. 现在把 gap 用 mean-spacing units 表示

定义 normalized gap

$$m_N = \frac{N \Delta_{\min}}{2\pi}.$$

即 mean spacing 被归一化成 1。

那么

$$\Delta_{\min} = \frac{2\pi m_N}{N}.$$

代进 (6):

$$-\Lambda_N \geq -\log \cos \left(\frac{\pi m_N}{N} \right). \quad (7)$$

这就是我们一直缺的 finite- N exact μ -to- Λ bridge。

$N \rightarrow \infty$:

$$-\log \cos x = \frac{x^2}{2} + O(x^4),$$

所以

$$N^2(-\Lambda_N) \geq \frac{\pi^2}{2} m_N^2 + \alpha(1). \quad (8)$$

6. 换成 Rodgers–Tao 的 local heat–time normalization, 常数又回到 1/8

定义 microscopic coordinate

$$x = \frac{N\theta}{2\pi}$$

和 microscopic heat time

$$s = \frac{N^2}{4\pi^2} t.$$

于是 (8) 就变成

$$-s_\Lambda \geq \frac{m^2}{8}. \quad (9)$$

所以:

$$\frac{1}{8}$$

不是 normalization accident。

circle finite- N 的

$$-\log \cos\left(\frac{\pi m}{N}\right)$$

在 local limit 下恰好退化成 Rodgers – Tao 的

$$\frac{m^2}{8}.$$

这两个 calculation 完全吻合。

7. 然后 AH 立刻给出一个硬数值阈值

AH/ACUE 要求 gaps 位于

$$\frac{1}{2}\mathbb{Z}_{>0}.$$

因此只要 configuration simple,

$$m_{\min} \geq \frac{1}{2}.$$

代入 (9):

$$-s_{\Lambda} \geq \frac{(1/2)^2}{8} = \frac{1}{32}.$$

即

$$D_{\text{AH}} \geq 0.03125. \quad (10)$$

这就是一个真正的 numerical falsifiability threshold。

在 finite-circle normalization 中, 同一个 statement 是

$$-\Lambda_N \geq -\log \cos \frac{\pi}{2N}. \quad (11)$$

或者 large N :

$$N^2(-\Lambda_N) \geq \frac{\pi^2}{8} = 1.23370055 \dots \quad (12)$$

8. 这和我们已经算出来的 ACUE $\wedge$ 完全吻合

我们之前 numerical ACUE 给的是大约

$$N^2(-\Lambda_{\text{ACUE}}) \approx 1.4196.$$

而 AH hard-core theorem 只要求

$$N^2(-\Lambda_{\text{AH}}) \geq 1.23370055 \dots$$

两者之间的差:

$$1.4196 - 1.2337$$

正是公式 (2) 中 positive many-body correction 在起作用。

换到 Rodgers – Tao normalized time:

$$\frac{1.41964034}{4\pi^2} \approx 0.03596,$$

而 hard threshold 是

$$\frac{1}{32} = 0.03125.$$

所以我们之前 single-dislocation 得到的 1.41964034 ... 不再是一个孤立 numerical curiosity。

它可以解释成:

$$\underbrace{0.03125}_{\text{universal 2-body AH barrier}} + \underbrace{\approx 0.00471}_{\text{ACUE many-body screening / rigidity}}.$$

这个 interpretation 比原来强很多。

9. 这确实把 Λ program 从 discriminator 升级成 falsifier

原来我们的 claim 是：

$$\Lambda_{\text{CUE}} \sim -N^{-8/3},$$

而

$$\Lambda_{\text{ACUE}} \sim -N^{-2}.$$

所以统计上可以区分它们。

但现在有了更强的 deterministic statement：

AH 必须满足

$$\frac{N^2}{4\pi^2}(-\Lambda_N) \geq \frac{1}{32}.$$

所以如果 arithmetic zeros 可以证明存在一个 sequence 使得

$$\liminf_{N \rightarrow \infty} \frac{N^2}{4\pi^2}(-\Lambda_N) < \frac{1}{32}, \quad (13)$$

AH 就死了。

甚至 finite- N 可以用 exact threshold：

$$-\Lambda_N < -\log \cos \frac{\pi}{2N} \implies \text{not AH.} \quad (14)$$

这不再是：

“这个 statistic 看起来更像 CUE。”

而是：

“AH 的 half-lattice support 数学上禁止出现这么 shallow 的 Newman configuration。”

这是质的提升。

10. 对 zeta zeros 的 formulation 也很漂亮

Rodgers – Tao 的 X -zero variables 在高度 T 附近平均 spacing 大约是

$$\frac{4\pi}{\log T}.$$

所以 local density

$$\rho_T = \frac{\log T}{4\pi}.$$

定义 dimensionless local heat depth

$$D_T = \rho_T^2(-t) = \left(\frac{\log T}{4\pi}\right)^2 (-t).$$

那么 AH 强制

$$D_T \geq \frac{1}{32}. \quad (15)$$

等价地，在原 de Bruijn – Newman time 中，

$$-t \geq \frac{\pi^2}{2 \log^2 T}. \quad (16)$$

这里最后一个 T -formula 是 local-density asymptotic；真正无 normalization ambiguity 的数字是 dimensionless threshold

$$\frac{1}{32}.$$

11. 一个很重要的区别：global $\Lambda = 0$ 本身还不够

这一点一定要区分。

Rodgers – Tao 证明的 global Newman constant

$$\Lambda \geq 0$$

与 RH 给出的

$$\Lambda \leq 0$$

若合并就是 global

$$\Lambda = 0.$$

但这并不会自动否定 AH。

我们这里需要的是 local / mesoscopic Newman depth

$$\Lambda(T; N)$$

或者一个 finite zero-window polynomial 的 backward collision depth。

然后证明：

$$D_T < 1/32$$

发生。

也就是说，真正的新 observable 应该不是一个 global scalar Λ ，而是一个 Newman-depth field

$$T \mapsto D(T).$$

甚至研究其 distribution：

$$\Pr(D(T) < c).$$

AH 预测一个 hard edge：

$$\Pr_{\text{AH}}(D < 1/32) = 0.$$

而 CUE 如果我们之前的 $N^{-8/3}$ law 正确，则随着 $N \rightarrow \infty$

$$N^2(-\Lambda_{\text{CUE}}) \sim N^{-2/3} \rightarrow 0,$$

所以

$$\Pr_{\text{CUE}}(D < 1/32) \rightarrow 1.$$

这不是小 correction，而是 support separation：

$$\begin{array}{ll} \text{AH/ACUE} & : D \geq 1/32, \\ \text{CUE} & : D \rightarrow 0. \end{array}$$

最终我会把这个桥写成一个 lemma

Hard-gap / Newman-depth lemma

设一个 simple zero configuration 具有 normalized minimum spacing

$$m.$$

在 Rodgers – Tao heat normalization 下，其 first backward collision depth $D = -\Lambda_{\text{loc}}$ 满足

$$D = \frac{m^2}{8} + S, \quad S \geq 0,$$

其中 S 是其他 zeros 产生的 many-body interaction integral。

因此任何 bandwidth-1 sine mimic with hard core m 必须满足

$$D \geq \frac{m^2}{8}.$$

特别地，half-lattice AH 满足

$$D \geq \frac{1}{32}.$$

这个 lemma 一旦写严谨，Lagarias – Rodgers 的 μ extremal problem 和 Newman-depth program 就真正接起来了：

$$\mu \xrightarrow{\text{exact collision geometry}} D_{\min} \geq \frac{\mu^2}{8}.$$

严格说最后是 inequality，而不是 $D = \mu^2/8$ ，因为 many-body correction 不为零；但对于“证伪 AH”而言，inequality 恰恰已经足够，而且方向完全正确。

下一步最有价值的事情已经很明确：不是再算 ACUE 的 median Λ ，而是尝试在 arithmetic/function-field model 中证明一次

$$D < 1/32.$$

只要做到这一件事，就给出了一个不需要完整 pair correlation 的 quantitative anti-AH certificate。

RESEARCH HANDOFF · SOURCE ARCHIVE · OVERNIGHT RESEARCH

素数间隙研究

第 4–5 轮

证明与实验补编

新增正项、变量半径估计与失败搜索

GPT–6 Astra · 独立审查与复算

Bill (Qingyun) Sun · 2026 年 9 月 4–5 日

资料检查点 c74b326 · 10 篇完整报告 · 主档案的后续补编

已证明、计算、猜想与失败分别标明 · Public research archive

内容索引

接手顺序

Current report 01: prime186_round4

Current report 02: prime_alpha_credit

Current report 03:
ALPHA_RECTANGLE_INDEPENDENT_REVIEW

Current report 04: RESTORATION_PROOF_AUDIT

Current report 05: REPORT

Current report 06: README

Current report 07: prime186_round5

Current report 08: EXCEPTIONAL_RADIUS_EXTENSION

Current report 09: GEOMETRY_SOURCE_AUDIT

Current report 10: REPORT

第 4–5 轮详细研究补编

本补编接续 333 页公开交接档案（检查点 055a4a0），保留随后两轮的完整证明报告和审查记录。资料源检查点为 c74b326。此前的 ACUE、零点热流、算术相关性、反例与历史路线仍应连同主档案阅读。

本轮严格结果是：同一 $k=40$ 筛法的新增正项使已发表基线上的证明余量从 23.36045 ppm 增至 24.86626 ppm；素数间隙结论仍为 186。另有变量半径异常平方估计、足够的支撑几何条件、一个网格修复，以及完整保存的 $k=39$ 负向搜索。没有证明 RH、Alternative Hypothesis 的反驳、新的 zeta 半间隙结果或低于 186 的素数间隙。

正文区分普通数学证明、独立内部审查、精确有理数证书、浮点计算、原论文输入和未完成义务。普通证明及内部审查尚不等于 Lean 形式验证或外部同行评审。53 项平方积分逐项区间、矩阵、参数、执行日志和哈希清单另存于仓库，不把大型数组印成难以核查的页面。

第 5 轮综合报告澄清原搜索报告的数值最大误差范围：较小的最大值只对应 12 个完整 77 维候选；仓库另存的截断候选有不同的全矩阵残差。所有 36 个向量均经过集成核查。

为便于机器和人工追溯，每篇报告列出原路径和 SHA256，文内相对链接指向该固定 Git 提交。报告文字及数学内容保留；只调整标题层级、链接、过长公式的换行，并补齐原始数据表缺少的 Markdown 表头分隔线。原文中关于暂缓、下一步或代理的描述属于研究记录。

接手顺序

先读两轮综合报告，接着审查新增正项及其独立证明，再查看完整 $k=39$ 搜索的缺口。变量半径估计与几何条件可以作为后续新支撑的输入，但不能替代新的失败覆盖及物理积分证书。避免重复相同端点扫描；寻找能改变当前缺口的产品权重、可证明支撑或算术混合项。

Current report 01: prime186_round4

Source: research/reports/prime186_round4.md. SHA256:

ef0b060fc8bb1ac91844e08ac540973634ceb6f9ca2d64da3f8ef32f66855565 .

A certified restoration credit and the remaining $k=39$ deficit

2026-09-05. This is a new research checkpoint after the 333-page handoff at commit `055a4a0`. The prime-gap bound remains 186. No new zeta gap, AH refutation, or Dyson – Montgomery theorem is proved.

The useful completed result is a strictly positive, outward-enclosed integral that the published sufficient criterion discards. On the same published 40-coordinate trial, it improves the certified normalized sieve margin from about 23.36045 ppm to 24.86626 ppm, conditional on the original published cap and loss endpoints. Independently, a complete 77-coefficient cap-only optimization at dimension 39 remains about 5603.60 ppm below one. That second result is a floating-point diagnosis, not a rigorous upper bound for the whole coefficient family.

1. Exact new positive credit

Proposition 4.6 of Improved short gaps between primes gives

$$\rho_*(\langle P_O F, B P_O F \rangle - \| P_O F \|^2) \geq \rho_*(J_{\lambda, H} - \beta - E_O) - l_H + c_a \alpha,$$

where

$$\alpha = \| (1 - P_O) F \|^2, \quad c_a = 1 - 4\rho_* | b_h | = \frac{2497786653900013}{2500000000000000} > 0.$$

The paper drops the last term in its convenient sufficient criterion. Its existing positive error covers provide upper estimates, so they cannot be reused as lower estimates for this credit.

We instead select a true sufficient-failure event. In units of the official mesh

$$h = \frac{2742997}{258046918656},$$

there are exactly two global fragments above $b = 18800h$, uniquely labelled by

$$p \in [26400, 29100]h, \quad q \in [32400, 36700]h.$$

All residual coordinate totals are at most b . The official radial cell index is restricted to $95639 \leq r \leq 98263$. This event lies within the cap domain and violates an actual retained new-ladder order-three row, with a strictly positive rational failure margin. Both fragments can belong to one coordinate or to different coordinates; both cases are necessary.

Restricting the remaining fragments to total at most b makes their one-coordinate total measure exactly Lebesgue measure. Replacing $1/p$ and $1/q$ by their lower endpoint-independent constants yields a single coherent positive lower measure. Its cell kernels are rational box-sum volumes, computed by integer positive-part polynomials. There is no Dickman quadrature error in these new kernels.

The two-mark ring gives exactly 40 same-owner terms and $40 \cdot 39$ different-owner terms, with no factor $1/2$. The 53 signed polynomial terms all integrate the square of the same step trial against that same positive lower measure. Their signs cannot be discarded individually.

The outward calculation gives, in the official common normalization,

$$\alpha_{\text{rect}} \in [3.5697238789 \times 10^{-20}, 3.5697238869 \times 10^{-20}],$$

where these decimals are rounded outwards from the exact rational endpoints saved in the receipt. Safe downward-rounded consequences are

$$\boxed{a/l_H^+ > 1.5071462817 \times 10^{-6}}, \quad \boxed{c_a a/l_H^+ > 1.5058119471 \times 10^{-6}}.$$

The original complete normalized lower margin, replayed from the published endpoints, is

$$0.000023360452297044097 \dots$$

Adding only the new proved credit gives

$$0.000024866264244232060 \dots,$$

an increase of about 6.45 percent of that small margin. It is an improvement in the certificate for the same trial and same theorem, not a reduction of 186.

The larger two-fragment triangle gives an exploratory estimate around 9.76 ppm, while the one certified rectangle gives 1.5058 ppm. We do not promote the triangle estimate to a certified credit. The tiny exact rational anchor in the working files is only a strict-positivity/normalization regression; the meaningful quantitative result is the outward rectangle contraction.

Evidence:

- [Complete event, kernel and credit proof.](#)
- [Outward rectangle receipt and exact complete-margin replay.](#)
- [Independent mathematical review](#), including an independent integer-cell/Eulerian reconstruction of every entry in all three marked kernels.
- [Separate-process rerun](#): all 53 term intervals and the final rational endpoints agree exactly with the first run. Runtimes were about 92.6 and 90.2 seconds on this host.

2. The actual k=39 scale

The independent cap implementation reconstructs the published physical radii, nested fragment caps and step masks, then changes the outer/retained dimensions to 39/38. It uses the correct $39h/Z$ face normalization. No dimension-40 integral endpoint is inherited.

On the official 98,304-cell grid:

Trial	Cap-only quotient $\rho_* J//$	Status
k=40, original 77 coefficients	1.000206086776951	Positive control; denominator lies in the published interval

Trial	Cap-only quotient $\rho_*/J//$	Status
k=39, original coefficients	0.994361581476018	Fixed-vector floating evaluation
k=39, optimized full 77 coefficients	0.994396399364491	Independent direct evaluation of the numerical Ritz candidate
k=40, optimized full 77 coefficients	1.000213743639754	Optimized positive control

The original k=39 deficit is about 5638.42 ppm. Reoptimizing the coefficients recovers about 34.82 ppm, leaving about 5603.60 ppm. This is a useful reason to change the support or trial structure instead of repeatedly evaluating the old vector.

The scaled k=39 Gram matrix has condition number about 2.28×10^{10} . The optimized matrix quotient and direct scalar reevaluation differ by 1.74×10^{-10} , and the numerical generalized residual is small. Those checks support the computation; they do **not** establish an interval upper bound for all possible coefficients. Likewise, the k=40 alpha lower bound cannot be inserted into a k=39 proof, nor used as an upper bound on all recoverable alpha mass.

The implementation uses exponential tilting as a numerically cancelling normalization, not as a change of trial. Two tilt values give the same fixed-grid result to about 4.2×10^{-15} . NumPy's requested `longdouble` is actually binary64 on this machine; that limitation is recorded in the JSON and report.

See [full k=39 report](#), [matrix data](#) and [returned coefficients](#). Neither a cap-only quotient nor a numerical family maximum pays the rootwise source-restoration costs or proves DHL[39,2].

3. Sharper restoration identities and a quantified failure

The exact projection identity retains both signed cross terms and removed face squares. For $\mathbf{e} = (1 - P_O)F$, $V_i = E_i F$ and $W_i = E_i \mathbf{e}$,

$$Q(P_O F) = Q(F) + \alpha + \rho_* \sum_i m_i (|W_i|^2 - 2 \operatorname{Re}(\overline{W_i} V_i)).$$

This identifies the actual projected-marginal matrix as the useful future optimization object. A completed-square bound instead requires an upper estimate for $\int_{H_O \setminus O} |BF|^2$. Applying only the generic factor 40 to the old face ledger is worse than the existing Young bound. A certified effective factor below about 14.9573 would improve that particular comparison. Merely writing a sharper identity does not supply such cancellation.

There is also an exact positive inner-overlap correction, but the published ledger caps its benefit for this fixed trial at 7.8813×10^{-8} of the normalized margin. It is too small to explain the k=39 deficit. The old upper ledger gives only a loose 1.9344 percent ceiling on possible alpha credit, not any positive lower credit by itself.

The [independent restoration proof](#) and [exact ledger replay](#) state these alternatives and prevent double counting. They include 200 finite signed-product diagnostics; ordinary written arguments, not those tests, supply the identities.

4. Corrected runtime and verification limits

The unchanged official regression failed in the earlier packaged FLINT environment. We built a separate FLINT 3.6.0 with the [upstream signed-conversion fix](#), then built Python-FLINT against it. The native integer/polynomial and Arb suites pass, as do the original certificate checks and 467 full plus 2,188 truncated products compared with an independent

Python integer implementation.

The complete Python binding suite has a separately documented assertion failure in a Jacobi test that calls the native function outside its odd-positive-denominator contract. Assertions remain enabled. The certificate does not use Jacobi, and the directly used binding APIs were tested independently. We do not claim a universally verified arithmetic library or that the whole binding suite passed.

The [runtime record](#) and [build scripts](#) preserve source hashes, actual library linkage, successful checks, failed build configurations and the full-suite failure. The official PrimeGaps186 source was not modified. The new credit inherits the original published cap/loss endpoints; all 149 old physical forms were not recomputed in this round.

5. Reproduce and continue

Use an isolated copy when executing scripts that write adjacent JSON outputs. External primary inputs can be selected explicitly:

```
export PRIME186_SOURCE=/path/to/PrimeGaps186/prime_gap_186_certificate.py
export PRIME186_NUMERICS_TEXT=/path/to/short_gaps_numerics.txt
```

The latter is a `pdftotext` extraction of the pinned official numerical companion; its hash is recorded by the ledger replay. The public package does not duplicate the full third-party paper text. In the relevant copied subfolders:

```
python certify_alpha_rectangle.py                # exact kernel/geometry checks
CORRECTED_PYTHON certify_alpha_rectangle.py --certify # outward integral
python replay_credit_margin.py                    # exact final inequality
python rectangle_independent_checks.py            # different integer kernel construction
python restoration_checks.py                      # source ledger + signed identities
OPENBLAS_NUM_THREADS=1 python cap_trial.py --k 39 --intervals 98304 --tilt 20
```

New code in the public copy differs from staging only where explicitly recorded in [the intake manifest](#), principally to select the external source paths. Original output receipts are retained.

The next active experiment changes the radius/plateau support geometry and recomputes any affected exceptional-square constant. It must retain actual lcm conditions, both source ladders, inner-domain intersections, the negative full-face term and inward masks. A new outward positive **complete** $k=39$ inequality would be needed before claiming a smaller gap. Repeated fixed-vector scans, a new full handoff rendering, and Jacobi-wrapper repair are postponed in this slice.

Current report 02: prime_alpha_credit

Source: research/prime-gaps/round4/prime-credit/prime_alpha_credit.md. SHA256:

59791f04d6917f447382083612eea4c311f3ec74cb427c37cbac49619ba046ec .

Retained positive deletion credit for the published prime-gap-186 trial

Status: the failure-region inclusion is proved and one meaningful rational-kernel rectangle has been contracted with outward arithmetic. It certifies $a/l_H^+ \geq 1.5071462817258709 \cdot 10^{-6}$ (decimal display of an exact rational lower endpoint). The retained credit improves the source-inherited fixed-trial margin from about 23.36045 ppm to 24.86626 ppm. The larger triangle estimate remains diagnostic. This note makes no smaller prime-gap claim and does not discharge the three project inputs in the official Lean development.

1. The exact quantity and its scale

For the actual published $k = 40$ midpoint step trial, Proposition 4.6 of the main paper gives

$$\rho_* \langle P_O F, B P_O F \rangle - \| P_O F \|^2 \geq \rho_* (J_{\lambda, H} - \beta - E_O) - l_H + c_a \alpha, \quad \alpha = \| (1 - P_O) F \|^2,$$

where

$$c_a = 1 - 4\rho_* |b_h| = 0.9991146615600052 \dots > 0.$$

The published sufficient inequality drops this positive term. A lower bound on the **actual F^2 -weighted deleted mass**, rather than the measure of a covering set, can therefore be inserted without changing the previously paid upper bounds on β and E_O . This does not recover the outer cross-term for free; that cross-term remains fully charged through E_O .

The common normalized denominator is source-certified as

$$23685317816 \cdot 10^{-24} \leq l_H \leq 23685317890 \cdot 10^{-24}.$$

All numerical forms in this note use the numerical companion's normalization $(hZ)^{-40}$, where $Z = \sum_{j=0}^{98263} g((j+1/2)h)^2$ and $h = 2742997/258046918656$. This is essential: the main paper uses unscaled forms, whereas the companion divides every form by this same factor.

The previously replayed complete margin is about 23.3604523 parts per million. The new, explicitly contained failure region below has a diagnostic mass of about 9.75897 parts per million of l_H . Its retained credit would be about 9.75033 parts per million if this size is certified. The diagnostic alone is not a valid replacement for an outward lower endpoint.

2. A disjoint lower-bound event

Let $\nu_c = e^\gamma c L(\Pi_c)$, with Poisson intensity du/u on $(0, c]$. Its total-size pushforward is $\rho_D(t/c) dt$. In particular, restricting the residual total to $0 \leq t \leq c$ gives exactly Lebesgue measure, because $\rho_D(t/c) = 1$ there.

Set

$$b = 18800h = 0.19984095864653703 \dots, \quad c = 46580h = 0.49513786456147313 \dots$$

Use the actual retained new-ladder row 24, not a failure-cover threshold assembled from different rows. Its parameters are

$$\begin{aligned} a &= 1.0166236774089747 \dots, & A &= 1.0449558074337872 \dots, \\ L &= 0.5498119373071242 \dots, & \xi &= 0.028332320501728507 \dots \end{aligned}$$

The exact fractions are regenerated from the preserving official source and saved in the JSON outputs. Its outer predicate is

$$\{s \leq a\} \cup \{H_{\phi_D, \xi} \leq A, \phi_E(M_\xi) \leq C\}, \quad \phi_D(u) = \min(3u/2, L).$$

Consider configurations having precisely two global fragments above b , ordered as $p < q$, with

$$b < p < q \leq c, \quad q + p + \min(3p/2, L) > A.$$

Remove these two fragments. Require every residual coordinate total x_i to lie in $[0, b]$, and require the restored total $s = \sum_i x_i + p + q$ to exceed a . Finally retain only the official radial cells $\sum_i \lfloor t_i/h \rfloor < 98264$.

Every residual fragment is at most $b < p$, so the inclusive tail at the witness p consists exactly of p and q . The displayed strict inequality therefore violates this actual row. Every fragment is at most the smallest outer cap c , so all three outer-shell fragment caps hold. The retained radial cells put the configuration in H_O . Consequently this region is a subset of $H_O \setminus O$.

There is no Palm multiplicity loss here. The two large fragments are globally unique and ordered. Their coordinate owners produce disjoint cases:

- distinct owners: $40 \cdot 39$ ordered choices;
- one common owner: 40 choices.

The Poisson decomposition gives the measure $\prod_i dx_i dp dq / (pq)$ on each such case. No additional $1/2$ occurs after imposing $p < q$. The same-owner case has both marked fragments in a single coordinate; it is not represented by the distinct-owner integral.

Define

$$t^D = (x_1 + p, x_2 + q, x_3, \dots, x_{40}), \quad t^S = (x_1 + p + q, x_2, \dots, x_{40}),$$

and let $F_\square(t)$ be the **official step function**, including its exact midpoint polynomial and profile, evaluated at these totals. Let R_D, R_S impose the preceding radial conditions. With

$$\mathsf{T} = \{(p, q) : p < q \leq c, q + p + \min(3p/2, L) > A\},$$

we obtain the proved lower-bound formula

$$\alpha \geq \frac{1}{(hZ)^{40}} \left[1560 \int_{\mathsf{T}} \int_{[0, b]^{40}} 1_{R_D} F_\square(t^D)^2 dx \frac{dp dq}{pq} + 40 \int_{\mathsf{T}} \int_{[0, b]^{40}} 1_{R_S} F_\square(t^S)^2 dx \frac{dp dq}{pq} \right].$$

Here T may equivalently start at

$$p_{\min} = (A - c)/(5/2) = 0.21992717714892562 \dots > b, \quad q > \max\{p, A - p - \min(3p/2, L)\}.$$

Thus the apparently additional condition $p > b$ is automatic in the implemented triangle. These inequalities are checked with exact fractions when constructing the event.

3. Measured mass of this genuine subset

`alpha_credit.py` uses scrambled Sobol sampling with an explicitly evaluated importance density. Unmarked residual cells are sampled with probabilities proportional to $g_j^2 e^{-\lambda_j h}$; conditional positions are uniform in the selected fine cell. Marked residuals use truncated exponential densities. The code evaluates the exact published midpoint step trial, and cancels the proposal factors explicitly. The proposal tilt is 45 and the cutoff is exactly $18800h$.

Eight independent scrambles, each with 2^{20} points for each owner case, give:

Disjoint event	Estimated mass divided by published l_H^+	Standard error across scrambles
Distinct coordinate owners	$4.2125098941 \cdot 10^{-6}$	$1.72 \cdot 10^{-8}$
Same coordinate owner	$5.5464619292 \cdot 10^{-6}$	$2.11 \cdot 10^{-8}$
Sum	$9.7589718233 \cdot 10^{-6}$	No deterministic enclosure

The run evaluated 16,777,216 points and took about 40 seconds on this host. The same-owner contribution is about 57% of the total. The weighted fragment locations are approximately $p = 0.293 - 0.294$, $q = 0.362 - 0.365$, and the weighted full total is about 1.032.

These are reproducible randomized-quadrature diagnostics. Across-scramble standard errors are not rigorous integration-error bounds; ordinary floating-point evaluation is also not an outward enclosure. The estimates do not prove $a/l_H \geq 9.7 \cdot 10^{-6}$.

4. Exact rational positive kernels for a certifiable rectangle

One disjoint rectangle inside the failure triangle is

$$P = [26400h, 29100h], \quad Q = [32400h, 36700h].$$

Indeed, P lies below Q , all marks are above b and below c , and throughout the rectangle

$$q + p + \phi_D(p) = q + \frac{5}{2}p \geq 98400h > A.$$

The source-core condition is guaranteed on every cell with index sum

$$95639 \leq r \leq 98263,$$

since $95638h \leq a < 95639h$ and the actual total is at least rh . This deliberately discards cells that straddle the activation boundary.

Replace the positive fragment density $1/(pq)$ by the smaller constant $1/(P_+ Q_+)$. This defines a coherent positive submeasure. Its cell kernels involve **only rational box volumes**, not Dickman quadrature or a critical-line approximation.

For a vector of positive integer lengths $w = (w_1, \dots, w_d)$ and integer offset $\mathcal{O}$, define

$$V_{o,w}(j) = \frac{1}{d!} \sum_{\epsilon \in \{0,1\}^d} (-1)^{|\epsilon|} [(j+1-o-\epsilon \cdot w)_+^d - (j-o-\epsilon \cdot w)_+^d].$$

It is exactly the volume of $\{0 \leq y_i \leq w_i : j \leq o + \sum_i y_i < j+1\}$. In units of cell measure h , the four coordinate kernels are

$$u_j = 1_{0 \leq j < 18800}, \quad v_j = \frac{V_{26400,(2700,18800)}(j)}{29100}, \quad w_j = \frac{V_{32400,(4300,18800)}(j)}{36700},$$

$$z_j = \frac{V_{58800,(2700,4300,18800)}(j)}{29100 \cdot 36700}.$$

The last channel is the same-owner measure. Multiply each by g_j^2/Z and place them in the positive marking ring

$$K_j = \frac{g_j^2}{Z}(u_j + av_j + bw_j + abz_j), \quad a^2 = b^2 = 0.$$

The ab coefficient of the 40-coordinate moment contraction automatically includes both $40zu^{39}$ and $40 \cdot 39vwu^{38}$. These multiplicities must not be added a second time.

The official finite symmetric-moment identity (numerical companion (2.10) – (2.12)) now evaluates this lower submeasure against the **same signed expansion of $F_{\square}^2$** . Taking signed coefficient terms separately as nonnegative would be invalid. The adapter accumulates their full outward intervals before assessing the lower endpoint.

`certify_alpha_rectangle.py` supplies the four exact integer-numerator arrays and rational denominators. Its standalone structural checks pass:

- every cell coefficient is nonnegative;
- the exact numerator mass sums are 18800, 101520000, 161680000, and 1309608000000;
- denominators are 1, 58200, 73400, and 6407820000;
- known two- and three-unit box-sum distributions agree;
- the marking-ring mass agrees exactly with the independent same/distinct owner formula;
- every rectangle point satisfies the actual source failure and cap inequalities.

A separate QMC evaluation of **this lower constant-density rectangle**, using the stricter cell mask, estimates

$$a_{\text{rect}}/I_H^+ \approx 1.5111578472 \cdot 10^{-6}$$

with distinct-owner 0.6524769 ppm and same-owner 0.8586809 ppm. The outward contraction subsequently completed all 53 signed square signatures in 92.582 seconds on one worker. It enclosed the normalized rectangle mass in

$$[3.5697238789408751 \cdot 10^{-20}, 3.5697238868155496 \cdot 10^{-20}],$$

where the exact dyadic endpoints, not these decimal displays, are stored in `alpha_rectangle_certificate.json`. Consequently

$$\frac{\alpha}{I_H^+} \geq \frac{9050325235576887333393096923828125}{6004941487970983985258771441262503395328} = 1.507146281725870950 \dots \cdot 10^{-6},$$

and the retained margin credit is at least

$$\frac{c_a \alpha^-}{I_H^+} = 1.505811947187963804 \dots \cdot 10^{-6}.$$

Replaying the printed original $I^\pm, J^-, L^+$ endpoints with this newly certified lower mass gives

$$M_{\text{new}} \geq 0.000024866264244232 \dots,$$

a 6.44598798% increase over the previous lower-margin endpoint. This strengthens a certificate for the same trial and $k = 40$; its admissible tuple and gap 186 are unchanged. `replay_credit_margin.py` records the complete exact rational addition in `alpha_credit_margin_replay.json`.

The QMC estimates at proposal tilts 45 and 35 bracket the certified value within their empirical fluctuation scale. They served as diagnostics, not as input to the proof. Tiling the larger triangle with disjoint rectangles can retain more mass; overlapping rectangle sums would not be valid lower bounds.

5. What is already rigorous, and the next executable obligation

`exact_cell_anchor.py` gives a completely rational positive anchor $a/I_H^\pm \geq 10^{-197}$. It uses two explicit owner types with all residual totals in one fine-cell subinterval, evaluates the constant step-trial values exactly, and proves the auxiliary bound $hZ < 11/500$ by a rational monotone Riemann sum. This tiny result verifies strict positivity and normalization. It has no useful effect on the sieve margin.

The meaningful rectangle calculation uses the official `CapEngine` and `SourceJets` implementations without changing their source. `CapEngine` executes the mandatory signed-FLINT regression. The adapter does not catch, disable, replace, or bypass that check. It introduces new rational lower kernels and the proven radial mask, then asks the existing outward machinery to contract them on all 98,264 retained indices.

The completed run used `/Users/qingyunsun/.cache/astra-research/flint-3.6.0-patched/venv/bin/python`, 160-bit Arb precision, 224-bit fixed-point positive convolutions, and one worker. The original signed-FFT startup regression passed unchanged. The root agent separately reports exact signed-convolution comparisons and passing native `fmpz_poly` / `fmpz_vec` tests for this corrected build. This is not a claim that the entire Python-FLINT test suite passed: the root identified an unrelated assert-enabled Jacobi test with an even denominator. Native assertions and the certificate regression remained enabled.

The rectangle now has an actual outward lower endpoint. The remaining work is independent review/replay of this new adapter and, if the larger approximately 9.76-ppm mass is worth pursuing, a disjoint finite rectangle cover from inside followed by its outward contractions. The full triangle estimate remains uncertified.

Commands, from this directory:

```
python3 certify_alpha_rectangle.py
python3 exact_cell_anchor.py
python3 alpha_credit.py --power 20 --repeats 8 --output alpha_credit_p20.json
python3 alpha_credit.py --region rectangle_lower --power 19 --repeats 4 --output alpha_rectangle_p19.json
## Only a corrected runtime which passes the unmodified signed regression:
/Users/qingyunsun/.cache/astra-research/flint-3.6.0-patched/venv/bin/python certify_alpha_rectangle.py --certify --
threads 1
python3 replay_credit_margin.py
```

The files `alpha_rectangle_kernel_checks.json`, `exact_cell_anchor.json`, `alpha_credit_p20.json`, and `alpha_rectangle_p19.json` record the present evidence. The completed outward run wrote `alpha_rectangle_certificate.json` and `alpha_rectangle_certificate.log`; the positive credit is taken from its

exact rational lower endpoint. Independent review of the adapter and subset proof remains appropriate before publication.

6. Relevance to $k = 39$

The adjacent independent round-four experiment reports that the published coefficient vector reused at $k = 39$ gives cap-only quotient $\rho_* J/I \approx 0.994361581476$, a deficit of about 5638 ppm before restoring supports. The demonstrated $k = 40$ region is about 10 ppm and cannot close that particular numerical deficit. Nor can its mass be inherited by dimension 39: the trial, measure, source masks, operator constants, and their normalization must all be reevaluated.

This does not establish a global $k = 39$ obstruction. It establishes a concrete scale comparison for one inherited trial and identifies a previously dropped positive term with an inexpensive rational-kernel route to certification. A better vector or changed support geometry remains a separate optimization problem.

Sources and provenance

The mathematical sources are the [official PrimeGaps186 repository](#), its main manuscript (Proposition 4.6 and (4.40)), and the [numerical companion](#), § § 1.1 – 1.3 and (2.10) – (2.12). The preserving clone is at commit `61340d0b74163003b32756bb16e91d9209a5e330`; the official certificate file SHA256 is `7f71bdefcfe3bb5ca76a143929b3cb3f4156c21dc483253cda3077420f1e5de4`.

The scripts read that source and write only in this owned staging directory. The published denominator and previous margin are inherited primary-source enclosures; they were not recomputed in this subtask. No external paid model calls, source-clone mutations, or unrelated toolchain rebuilds were used by this agent.

Current report 03:

ALPHA_RECTANGLE_INDEPENDENT_REVIEW

Source: [research/prime-gaps/round4/restoration-proof/ALPHA_RECTANGLE_INDEPENDENT_REVIEW.md](#).

SHA256: [22796ab5c24508df097e913595f76519d5fe7002df75de7584653142e2e7c196](#).

Independent audit of the positive-alpha rectangle certificate

2026-09-05. Reviewed files: `../prime-credit/alpha_credit.py`, `certify_alpha_rectangle.py`, `prime_alpha_credit.md`, and the completed `alpha_rectangle_certificate.json`. This review is independent of their author. It reads the preserving official certificate source and inspects the completed outward receipt; it does not rerun the large contraction.

1. Verdict and certified scope

I found no mathematical or normalization defect in the rectangle lower-event construction or its adapter to the official outward moment engine. The following points are correct:

- the rectangle belongs to the actual deleted outer support, not merely an upper failure cover;
- the residual total law is exactly Lebesgue measure on the selected interval;
- the marked densities have no missing exponential normalizer, cap factor, or factor $1/2$;
- the same-owner and different-owner cases are both included, once each;
- all coordinate kernel factors of h cancel against the published $(hZ)^{40}$ normalization correctly;
- the adapter uses the actual midpoint step trial, its full signed square expansion, and a valid inward radial mask;
- the completed receipt contains a positive outward lower endpoint.

The new receipt proves a lower credit of approximately

$$\frac{c_a a_{\text{rect}}}{f_H^+} \geq 1.50581194718 \times 10^{-6}, \quad c_a = 1 - 4\rho_* \mid b_h \mid. \quad (1.1)$$

The exact rational endpoint is in the receipt and in `rectangle_independent_checks.json`; the displayed decimal is descriptive. A simpler strictly lower rational bound is 1.5058×10^{-6} .

This improves the inherited published $k = 40$ trial's certified normalized margin from approximately 23.3604523 ppm to at least 24.8662642 ppm. The published cap and loss endpoints are inherited source inputs, not recomputed by the new run. This is a strict improvement of the fixed-trial margin, not a smaller prime-gap theorem or a $k = 39$ certificate.

Our earlier restoration report correctly said the old upper ledger alone yields no positive alpha lower bound. This new, separate lower-event calculation supplies such a bound and does not conflict with that statement.

2. Proof that the event really is deleted

Use the exact official grid

$$h = \frac{2742997}{258046918656}, \quad n = 98264.$$

The residual cutoff is $b = 18800h$, and the two marked intervals are

$$P = [26400h, 29100h], \quad Q = [32400h, 36700h].$$

These intervals are strictly ordered and lie above b . All residual coordinate totals are at most b , so every residual fragment is at most b . Hence $p \in P$ and $q \in Q$ are the unique two fragments above b in the entire outer root, with $p < q$.

The source row used is the retained new-ladder row of index 24. I checked the official ladder generation and row-retention filter: this row is included among the 39 retained new rows, and its order is three. Its relevant outer condition is

$$s \leq a \quad \text{or} \quad (H_{\phi_D, \xi} \leq A \text{ and the opposite-root condition}), \quad \phi_D(u) = \min(3u/2, L).$$

Throughout the selected rectangle, exact rational inequalities give

$$p > \xi, \quad \frac{3}{2}p \leq \frac{3}{2}(29100h) < L,$$

and

$$q + p + \phi_D(p) = q + \frac{5}{2}p \geq (32400 + \frac{5}{2}26400)h = 98400h > A. \quad (2.1)$$

The inclusive tail at the witness p contains exactly $p + q$, because all other fragments are smaller than p . Thus (2.1) violates the actual owner condition. The opposite-root condition need not be analyzed to establish failure of the conjunction.

The inward radial mask is

$$95639 \leq r = \sum_i \lfloor t_i/h \rfloor \leq 98263.$$

The exact core comparison is $95638h \leq a < 95639h$, while the actual total satisfies

$$s = \sum_i t_i \geq rh > a.$$

This excludes the safe-core alternative. On the other side,

$$s < (r + 40)h \leq 98303h < S = 98304h,$$

and the coordinate-index sum belongs to the retained official outer cells.

Finally, every fragment is at most $36700h$, strictly below the smallest official outer-shell cap $46580h$. Therefore the event lies inside the actual cap domain H_O , regardless of which outer shell contains its total. Combining these facts proves it lies in $H_O \setminus O$.

The proof uses the actual row predicate, not a rounded common failure threshold from the old covering argument. Thus it is a legitimate source of lower mass.

3. Residual measure and the missing-factor audit

The source's coordinate measure is

$$\nu_c = e^Y cL(\Pi_c),$$

where Π_c is the Poisson point process of intensity du/u on $(0, c]$. For $b < c$, the probability of no points in $(b, c]$ is b/c . Hence

$$\nu_c|_{\{\max X \leq b\}} = \nu_b. \quad (3.1)$$

This identity includes the normalizing factor: $e^Y c \cdot (b/c) = e^Y b$. It must not be followed by an additional b/c or e^{-Y} factor.

The total-size pushforward of ν_b is

$$\rho_D(x/b) dx.$$

For $0 \leq x \leq b$, $\rho_D(x/b) = 1$. Consequently the selected residual total has exactly the measure dx . It is not a probability-uniform density dx/b ; its mass on this interval is b .

For a coordinate with one selected large fragment, Poisson decomposition adds dp/p or dq/q . For a coordinate with both, it adds $dp dq/(pq)$. Since P and Q are disjoint, the two selected fragments have unique labels determined by their intervals. There is no $1/2!$: one may derive the same fact by first using the unordered two-point density with $1/2!$, then summing its two disjoint orderings.

Across all coordinates the event measure is therefore

$$\prod_{i=1}^{40} dx_i \frac{dp dq}{pq},$$

with the chosen marks assigned to their coordinate owners. The interval restrictions and total caps merely restrict this measure.

4. Owner counting: the same-coordinate term is essential

There are two disjoint owner classes:

1. p and q belong to different coordinates: $40 \cdot 39 = 1560$ ordered choices.
2. Both belong to one coordinate: 40 choices.

The owners are ordered by the labels $p \in P$ and $q \in Q$. An additional factor two would duplicate configurations. Omitting the second class would lose genuine positive mass.

For the moment computation, let u_j, v_j, w_j, z_j be the four coordinate cell measures for no mark, only p , only q , and both marks. The ring

$$R[a, b]/(a^2, b^2)$$

with coordinate element $u + av + bw + abz$ has product rule

$$(u, v, w, z)(u', v', w', z') = (uu', vu' + uv', wu' + uw', zu' + uz' + vw' + ww').$$

The coefficient of ab in its 40-fold product is precisely the sum of same-owner and different-owner assignments. For constant integrated channel masses it is

$$40zu^{39} + 40 \cdot 39vwu^{38}.$$

This is exactly the official SourceJets “palm” multiplication and “both” channel. Those factors are already generated by the ring and must not be applied again outside the contraction.

The equality between some integrated channel masses does not permit replacing the same-owner cell kernel by a product of the separate-owner kernels. The same-owner total is $x + p + q$, and it requires the separate three-variable box kernel used in the adapter.

5. Exact box volumes and normalization by hZ

Replace $1/(pq)$ on the marked rectangle by its lower constant

$$\frac{1}{P_+ Q_+}, \quad P_+ = 29100h, \quad Q_+ = 36700h.$$

This produces one coherent positive submeasure. Since the final integrand is $F_{\text{step}}^2 \geq 0$, its integral is a lower bound on the event's mass.

Scale all variables by h . A coordinate with one mark has two integration variables x, p , so its cell measure is

$$\frac{h^2}{29100h} \text{Vol}\{(x', p') : 0 \leq x' \leq 18800, 26400 \leq p' \leq 29100, j \leq x' + p' < j+1\}.$$

Dividing by h , as needed for the normalized coordinate array, leaves the two-dimensional volume divided by 29100. The analogous q channel divides by 36700.

The same-owner channel has three integration variables and two mark denominators. Its cell measure is h times the corresponding three-dimensional volume divided by $29100 \cdot 36700$. Thus the four dimensionless arrays in the adapter are exactly cell mass divided by h .

The inclusion-exclusion routine returns $d!$ times the d -dimensional cell volume. Therefore the denominators

$$1, \quad 2 \cdot 29100, \quad 2 \cdot 36700, \quad 6 \cdot 29100 \cdot 36700$$

are correct.

On an actual coordinate cell, the official step profile contributes g_j^2 . The adapter's weight is

$$\frac{g_j^2 \text{ cell mass}}{Z h} = \frac{g_j^2 \text{ cell mass}}{hZ}.$$

Multiplication across 40 coordinates supplies exactly the companion paper's $(hZ)^{-40}$ normalization. No extra $h, 40, Z$, or physical-scale multiplier belongs after the moment contraction.

This point distinguishes the new norm lower integral from the official source-loss routines, some of which integrate an extra erased coordinate and accordingly have additional factors. The adapter correctly uses only the coordinate product law needed here.

6. Actual step trial and signed-square coherence

I checked that the official CapEngine aliases its signature and coefficient arrays directly to the printed rational trial data. Its square_groups are obtained by the exact full quadratic expansion of the 77 coefficients, including a factor two for distinct coefficient pairs.

The radial polynomial argument is

$$(r+20)h - \frac{9}{10},$$

the sum of the 40 coordinate midpoints minus the fixed center. The angular moments use the same coordinate midpoints. Hence the adapter evaluates the published **step trial**, not a continuous polynomial approximation at the original real totals.

Every kernel channel, radial mask, and signature uses the same positive lower submeasure. Some coefficients of the polynomial-square expansion are negative, but their sum represents the nonnegative square. The adapter:

- obtains two-sided outward intervals for the positive moment coefficients;
- multiplies by two-sided intervals for the signed radial polynomial;
- uses the official signed interval product and outward reduction;
- adds all 53 signed signature contributions before testing the final lower endpoint.

It does not freeze upper moment bounds and then multiply them by negative coefficients, and it does not clamp negative signature terms to zero. Several negative partial sums in the run log are consistent with the necessary signed cancellation.

The proof therefore rests on ordinary interval enclosure of one fixed coherent integral. Independence between numerical interval errors is not required.

7. Independent checks performed

The companion `rectangle_independent_checks.py` does not import `CapEngine`, `SourceJets`, or the rectangle certifier as a module.

It reconstructs all three marked cell-numerator arrays by a different exact method: decompose each integer-length interval into unit intervals, convolve their integer location counts using prefix sums, and finish with the Eulerian unit-cube cell-volume numerators

$$(1, 1) \quad \text{in dimension two}, \quad (1, 4, 1) \quad \text{in dimension three}.$$

All 98,264 entries of each reconstructed array agree exactly with the proposed inclusion-exclusion routine. Their sums are respectively

$$101520000, \quad 161680000, \quad 1309608000000.$$

The script also verifies the owner expansion with exact fractions on a four-coordinate, three-cell toy measure, using a nonconstant squared trial and a radial mask. The ring contraction equals a separate explicit sum over all 4 same-owner and 12 different-owner choices. This checks more than the unweighted total mass.

Finally, it checks the completed receipt's endpoint ordering, positive lower endpoint, alpha normalization by the published I_{μ}^+ , and multiplication by the exact coefficient C_a , all using exact fractions.

These checks passed. The source adapter hash and receipt hash are stored in `rectangle_independent_checks.json`.

8. Completed outward receipt and evidence boundary

The receipt records:

- the unmodified mandatory signed-convolution regression passed;
- 160-bit Arb working precision and 224-bit fixed-point source arrays;
- one worker thread;
- all 53 signed signatures completed;
- a final normalized rectangle interval approximately

$$[3.569723878940875 \times 10^{-20}, 3.569723886815550 \times 10^{-20}].$$

Using the source-inherited bound $I_H^+ = 23685317890 \cdot 10^{-24}$, the exact lower endpoint gives

$$\frac{\alpha_{\text{rect}}}{I_H^+} \geq 1.50714628172 \times 10^{-6},$$

and the exact C_Q gives (1.1). The contraction took about 92.6 seconds according to its receipt.

This is stronger evidence than the earlier randomized-QMC estimate: it is a completed outward enclosure of the explicit lower submeasure. The QMC code's proposal factors, owner split, and midpoint evaluation are also consistent with this same integral, but its standard errors were never deterministic certificates and are not used here.

The remaining trust boundary is the official outward arithmetic implementation and the corrected runtime that passed its unchanged regression, together with the published denominator and previous upper-loss inputs. This review checks the mathematical adapter, event inclusion, normalization, and scalar receipt; it is not a formal verification of FLINT or a fresh proof of all 149 original physical bounds.

The improved $k = 40$ margin follows by adding this lower credit once to Proposition 4.6 while retaining all existing outer and inner debts. It does not justify removing the mixed-term debt, transferring the mass to dimension 39, or adding the same credit again through a separately restored denominator.

Current report 04:

RESTORATION_PROOF_AUDIT

Source: [research/prime-gaps/round4/restoration-proof/RESTORATION_PROOF_AUDIT.md](#). SHA256:

`fbbaa5293ea486e643cecc78d7651487448b54f3dbe0840cce8ac61d4089645f`.

Exact projection restoration, positive norm credit, and the limits of the published ledger

2026-09-05. Independent proof/audit task, confined to Proposition 4.6 of the prime-gap-186 paper and the structural frontier identified in round three. No official physical-integral computation was rerun, no Claude session was started, and no repository source was edited.

Primary source: Improved short gaps between primes, § § 4.3 – 4.6, especially Proposition 4.6, displayed equations (4.30) – (4.43). Numerical source: the official certificate repository, Numerical certificate for prime gaps at most 186, Proposition 1.3 and Tables 2.1 – 2.7. Locally read files and exact scalar replay are recorded in `restoration_checks.py/json/log`.

The formulas below are consequences of the source's Hilbert-space setup. They are not a claim that the needed new physical integrals have been evaluated, nor a claim of priority for elementary projection identities.

1. Findings

There is a completely exact restoration identity that avoids Young debts:

$$Q(PF) = Q(F) + \alpha + \rho \sum_i m_i (|W_i|^2 - 2 \operatorname{Re}(\overline{W_i} V_i)), \quad (1.1)$$

where $Q(F) = \rho \langle F, BF \rangle - \|F\|^2$, $P = 1_O$, $e = (1 - P)F$, $\alpha = \|e\|^2$, $V_i = E_i F$, and $W_i = E_i e$. Computing or bounding the actual removed marginals W_i is the direct route to a stronger restoration calculation.

Three useful rigorous lower bounds follow:

1. Keep the paper's Young upper bound and restore a separately certified lower bound for α .
2. Integrate a signed root residual on the true failure set. A pointwise completion of the square yields a new upper-cost target $\int_{H_O \setminus O} |BF|^2$, which retains cancellation between face terms.
3. Use sharp pointwise quadratic bounds for W_i , including its squared term, rather than discarding all positive face squares.

The old/new inner-domain overlap also gives an exact positive correction. The published constants cap its normalized margin benefit at 7.8813×10^{-8} , so that particular overlap recovery is small for the published $k = 40$ trial.

The 52 outer root-square ledger entries imply

$$0 \leq \frac{\alpha}{l_H} \leq 0.019361160920, \quad 0 \leq \frac{(1 - 4\rho |b_h|)\alpha}{l_H} \leq 0.019344019740. \quad (1.2)$$

These are upper bounds on possible credit, not positive lower bounds. The published upper ledger alone supplies no guaranteed new credit. They are also deliberately loose, and do not establish whether a $k = 39$ trial can succeed.

2. Setup and the exact signs

Work in the source's Hilbert space $L^2(H_O, \nu^k)$, extending profiles by zero. Let

$$U = H_O \setminus O, \quad P = 1_O, \quad e = 1_U F, \quad f = PF = F - e.$$

Since P is an orthogonal multiplication projection,

$$\langle f, e \rangle = 0, \quad \|f\|^2 = I_H - \alpha, \quad \alpha = \int_U |F|^2. \quad (2.1)$$

The actual face operator is

$$B = \sum_i E_i^* m_i E_i, \quad m_i = d_0 1_{L_0} + a_h 1_{L_1} + b_h. \quad (2.2)$$

Put $b = |b_h| = -b_h > 0$. Its three values are

$$m_i = \begin{cases} 1, & Y \in L_0, \\ a_h - b, & Y \in L_1 \setminus L_0, \\ -b, & Y \in L_1. \end{cases} \quad (2.3)$$

The first two are positive for the published constants. The source proves

$$-bC_{\text{op}}I \leq B \leq C_{\text{op}}I, \quad C_{\text{op}} = 4. \quad (2.4)$$

It also defines A using the larger cap domains, $D = A - B \geq 0$, and

$$J_{\lambda,H} = \langle F, AF \rangle, \quad \beta = \langle F, DF \rangle, \quad \langle F, BF \rangle = J_{\lambda,H} - \beta. \quad (2.5)$$

Thus $Q(F)$ in (1.1) uses the actual B but the unprojected F . It is not the cap form $\rho J_{\lambda,H} - I_H$.

Expanding $f = F - e$ gives, with no estimate,

$$Q(f) = \rho(J_{\lambda,H} - \beta) - I_H + \alpha - 2\rho \text{Re}\langle e, BF \rangle + \rho\langle e, Be \rangle. \quad (2.6)$$

Applying the adjoint relation to the last two terms gives (1.1).

All face products in these equations are products of conditional integrals. In particular,

$$|W_i(Y)|^2 = \iint \alpha(Y \oplus_i X) \overline{\alpha(Y \oplus_i X')} d\nu(X) d\nu(X'). \quad (2.7)$$

The two integrated coordinates are independent copies conditional on the same retained configuration Y . Replacing this expression with $E_i(|e|^2)$ is incorrect.

No positivity of F is assumed. The research record explicitly notes sign changes of the optimized polynomial. In the signed hybrid form, replacing F by $|F|$ is not automatically monotone because $b_h < 0$.

3. Positive-alpha credit: what is sufficient and what is not

The source's spectral lower bound gives

$$\langle e, Be \rangle \geq -bC_{\text{op}}\alpha.$$

Define

$$c_a = 1 - \rho b C_{\text{op}} > 0. \quad (3.1)$$

If E_O is any valid upper bound for $2 \operatorname{Re}\langle e, BF \rangle$, then

$$Q(PF) \geq \rho(J_{\lambda,H} - \beta - E_O) - l_H + c_a \alpha. \quad (3.2)$$

Consequently, for outward bounds and a **separately proved lower bound** $\alpha^- \leq \alpha$, the strengthened sufficient criterion is

$$\rho(J_{\lambda,H} - E_O^+ - \beta^+) - l_H^+ + c_a \alpha^- > 0. \quad (3.3)$$

The alpha term is not multiplied by another ρ . It already comes from the denominator restoration $\|PF\|^2 = l_H - \alpha$, together with the small negative spectral correction.

3.1 Disjoint sufficient–failure regions

If $A_1, \dots, A_r \subset U$ are pairwise disjoint up to null sets, then

$$\alpha \geq \sum_{\ell=1}^r \int_{A_\ell} |F|^2. \quad (3.4)$$

Certified lower integral enclosures on these sets can be added. The inclusion $A_\ell \subset U$ must use the **actual complete outer-domain predicate**, including activations, row membership, radial cells, and fragment caps.

A practical construction is an inward fragment region on which one actual outer support inequality fails by a strict rational margin. Order the witness coordinates or assign the first failing row and the largest relevant fragment to half-open intervals, so the selected regions are disjoint. A lower integral requires a lower probability or a two-sided enclosure of that region. A Chernoff upper bound is not a lower probability.

More generally, nonnegative weights $w_\ell(X)$ may be used if

$$\sum_{\ell} w_\ell(X) \leq 1_U(X) \quad \text{a.e.}$$

Then $\alpha \geq \sum_{\ell} \int w_\ell |F|^2$. This is the appropriate partition-of-unity direction for credit.

3.2 Existing overlapping positive covers

The paper uses nonnegative majorants M_j with

$$1_U \leq \sum_j M_j. \quad (3.5)$$

The inequality is in the opposite direction. It produces upper bounds on nonnegative integrals over U ; it does not prove any of the M_j is supported in U , or any positive lower bound on α .

Even if every cover set happened to lie in U , adding their masses would overcount intersections. Two identical sets of mass one give a sum two and a union mass one. Remedies are disjointification, a proved multiplicity bound, Bonferroni lower bounds with correctly directed intersection enclosures, or nonnegative credit weights summing at most one.

The current Palm – Chernoff and factorial majorants can be positive outside the failure event and can exceed one. They therefore cannot simply be normalized or relabeled as disjoint lower-mass contributions.

For signed residual integrands, the restriction is stronger: multiplying by an upper event cover preserves inequalities only for a nonnegative integrand. Its positive and negative parts must be treated separately.

4. Signed root restoration and an optimized quadratic bound

Let

$$G(X) = (BF)(X) = \sum_i m_i(\hat{X}_i) V_i(\hat{X}_i). \quad (4.1)$$

The adjoint formula is understood on H_O , with zero extension. From (2.6) and (2.4),

$$Q(PF) \geq Q(F) + \int_U r(X) dv^k(X), \quad r = c_a |F|^2 - 2\rho \operatorname{Re}(\overline{F}G). \quad (4.2)$$

This is sharper in its treatment of the mixed term than replacing every $m_i F V_i$ by $h_i |F V_i|$ and applying Young separately. It retains cancellation across faces and the sign of the summed mixed term.

Completing the square pointwise gives

$$r = c_a \left| F - \frac{\rho}{c_a} G \right|^2 - \frac{\rho^2}{c_a} |G|^2,$$

and hence

$$Q(PF) \geq \rho(J_{\lambda,H} - \beta) - I_H - \frac{\rho^2}{c_a} \int_U |BF|^2. \quad (4.3)$$

This is an optimized quadratic bound in which the favorable denominator loss pays for part of the mixed term. It is optimal pointwise if only G is retained and F is otherwise unrestricted. It is not claimed to dominate the paper's optimized componentwise Young estimate for every trial.

4.1 Partial exact integration

If $A \subset U$ is a known-failure region and $1_{U \setminus A} \leq \sum_j M_j$, then

$$\int_U r \geq \int_A r - \sum_j \int M_j r_-, \quad r_- = \max(-r, 0). \quad (4.4)$$

Thus one can retain signed residual on a certified region and upper-bound only the harmful residual on the remainder. The old positive-cover machinery can act on r_- , provided its new weighted physical integrals are actually enclosed.

One must not add the right sides of different lower bounds such as (3.2) and (4.3). They are alternative estimates of the same correction; taking their maximum is valid, adding their claimed gains is generally double counting.

4.2 If only an alpha interval and a residual Gram upper bound are known

Suppose $\alpha \in [a_-, a_+]$ and

$$\int_U |G|^2 \leq K.$$

Cauchy – Schwarz gives a rigorous correction

$$Q(PF) - Q(F) \geq \min_{a \in [a_-, a_+]} (c_a a - 2\rho\sqrt{aK}). \quad (4.5)$$

The minimizing $\sqrt{a}$ is the projection of $\rho\sqrt{K}/c_a$ onto $[\sqrt{a_-}, \sqrt{a_+}]$. Substituting a lower bound for a directly into the negative square-root term is not valid: the correction need not be increasing in a on the whole interval.

5. Retain the removed face squares

For each face let

$$a_i(Y) = \int 1_U(Y \oplus_i X) d\nu(X), \quad z_i(Y) = \int 1_U(Y \oplus_i X) |F(Y \oplus_i X)|^2 d\nu(X).$$

Then $|W_i|^2 \leq a_i z_i$. Put $R_i = \sqrt{a_i z_i}$ and $m_i^\pm = \max(\pm m_i, 0)$. Minimizing the exact quadratic term over the disk $|W_i| \leq R_i$ yields

$$m_i(|W_i|^2 - 2\operatorname{Re}(\overline{W_i} V_i)) \geq -m_i^+ [|V_i|^2 - (|V_i| - R_i)_+^2] - m_i^- [R_i^2 + 2R_i |V_i|]. \quad (5.1)$$

For $m_i \geq 0$, write the quadratic as $m_i(|W_i - V_i|^2 - |V_i|^2)$ and take the closest point in the disk to V_i . For $m_i < 0$, take the farthest point; its distance is $|V_i| + R_i$. This proves both signs and shows the bound is optimal under that single disk constraint.

Integrating (5.1) and adding a in (1.1) gives a valid lower restoration bound. In contrast to Proposition 4.6's final estimate, it retains favorable curvature of the positive-face square. But it requires certified functions a_i, z_i, V_i , not just global masses.

For real F , an interval enclosure $W_i \in [\ell_i, u_i]$ can be better. On a positive face, minimize at the point of $[\ell_i, u_i]$ closest to V_i ; on a negative face, minimize at an endpoint farthest from V_i . If the actual W_i is computed, (1.1) is exact and no auxiliary Young parameter is needed.

For a fixed finite trial basis ϕ_a and fixed support, the most direct optimization object is the exact compressed matrix

$$M_{ab} = \rho \sum_i m_i E_i(P\phi_a) \overline{E_i(P\phi_b)} - \int P\phi_a \overline{\phi_b}. \quad (5.2)$$

A rational vector v with $v^* M v > 0$, certified using outward enclosures and the source hypotheses, would establish the sieve criterion for that trial. The computational obstacle is the physical projection integral; writing the matrix does not evaluate it.

The source projection does not commute with E_i . Replacing $E_i(PF)$ by a face indicator times $E_i F$ would discard exactly the difficult dependence on the erased coordinate.

6. Exact old/new inner-overlap credit

Retain the paper's order of operations: inner deletion is evaluated at the unprojected F , then the outer projection is restored. On a face write

$$A_i = H_0 \setminus L_{\text{old}}, \quad C_i = H_1 \setminus L_1.$$

Because $L_0 = L_{\text{old}} \cap L_1$ on the base domain and $H_0 \subset H_1$,

$$H_0 \setminus L_0 = A_i \cup (C_i \cap H_0).$$

The exact inner loss is therefore

$$\beta = d_0 \beta_{\text{old}} + (1 - b_h) \beta_{\text{new}} - d_0 \Gamma_{\text{in}}, \quad (6.1)$$

where

$$\Gamma_{\text{in}} = \sum_j \int_{(C_i \setminus H_0) \sqcup (A_i \cap C_i)} |V_i|^2. \quad (6.2)$$

The displayed union is disjoint and lies inside C_i , so

$$0 \leq \Gamma_{\text{in}} \leq \beta_{\text{new}}. \quad (6.3)$$

This simultaneously accounts for the unnecessary d_0 charge on new failures outside the base and the duplicated d_0 charge where old and new failures overlap inside the base.

To use it as a positive credit, one needs a lower bound for the true integral in (6.2). Intersections of the existing upper majorants are not lower bounds for intersections of the true failure events.

If the exact projected matrix (5.2) is used instead, this inner correction is already built in. It must not be added again.

7. Exact replay of what the published numerical upper ledger can say

The script reads the 52 outer rows of numerical Tables 2.1 and 2.2, checks every rational Young rounding inequality, and sums the printed integers using exact fractions. It does not import the official numerical certificate. Its conclusions remain conditional on the correctness of the published physical upper bounds.

For the common normalization used in the numerical paper, set

$$\rho = \frac{2624989}{10^7}, \quad b = \frac{843183}{10^9}, \quad d_0 = \frac{44415113}{5 \cdot 10^9}, \quad c_a = \frac{2497786653900013}{2500000000000000}.$$

Thus $c_a = 0.9991146615600052$.

7.1 Upper cap on possible alpha credit

Writing R_j, V_j for the published outer root-square and outer face-square upper forms, the table sums give

$$\begin{aligned} \frac{\sum_j R_j}{I_H} &\leq 0.000653000069917512, \\ \frac{\sum_j V_j}{I_H} &\leq 0.000168396875617848, \\ \frac{E_O}{I_H} &\leq 0.000661756763. \end{aligned} \quad (7.1)$$

The first form majorizes $\sum_i \int_{\mathcal{U}} h_i(\hat{X}_i) |F(X)|^2 d\nu^k$. Every $h_i \geq b$, so

$$k b a \leq \sum_j R_j.$$

Since $l_H^- \leq l_H$, at $k = 40$,

$$\frac{a}{l_H} \leq \frac{0.000653000069917512}{40b} = \frac{9069445415521}{468435000000000} = 0.0193611609199163 \dots \quad (7.2)$$

Multiplying by c_a gives the credit cap in (1.2). The guaranteed lower bound remains zero. In particular, it is invalid to insert 0.01936 as the recovered mass: it is an upper ceiling obtained from overlapping majorants and a very small universal weight $40b$.

This does not prove the actual credit is close to that ceiling, nor that a redesigned $k = 39$ profile has the same ceiling.

7.2 The exact inner–overlap improvement is tightly bounded

The published two new-inner weighted groups sum to

$$\frac{(1+b)\beta_{\text{new}}}{l_H} \leq \frac{1405159 + 32422390}{10^{12}} = 0.000033827549.$$

Using (6.3), the normalized sieve-margin benefit from correcting only the duplicated d_0 charge is at most

$$\frac{\rho d_0 \Gamma_{\text{in}}}{l_H} \leq \frac{\rho d_0}{1+b} 0.000033827549 = 7.88120730556 \dots \times 10^{-8}. \quad (7.3)$$

This is an absolute cap for that specific correction and the published $k = 40$ profile. It does not cap all possible changes to the inner support or trial.

7.3 Completion of the square needs real cancellation to beat the old debt

Let $H(X) = \sum_i h_i(\hat{X}_i) \leq k$. Weighted Cauchy – Schwarz gives

$$\left| \int_U BF \right|^2 \leq H(X) \sum_i h_i(\hat{X}_i) \left| \int_U V_i(\hat{X}_i) \right|^2.$$

The positive cover then yields the crude bound

$$\int_U \left| BF \right|^2 \leq k \sum_j V_j. \quad (7.4)$$

For the published ledger, plugging this into (4.3) produces a normalized outer debt no larger than

$$\frac{\rho^2 k}{c_a} 0.000168396875617848 = 0.000464551283571 \dots$$

That certified upper debt is worse than the published Young debt

$$\rho 0.000661756763 = 0.000173710422355 \dots$$

Thus the generic completed-square inequality is not by itself an improvement. It requires a stronger bound on the actual signed sum BF , or smaller effective H on failures.

Measured against the same published face ledger, a sufficient comparison target is

$$\frac{\int_U \left| BF \right|^2 / l_H}{0.000168396875617848} < 14.957265516079 \dots, \quad (7.5)$$

instead of the crude factor $k = 40$. This is a concrete target for a new physical integral, not evidence that it holds. The quotient merely compares two certified upper-bound schemes; a larger ratio would not establish that the actual restoration fails.

8. Verification and unresolved proof obligation

The accompanying script performs exact fraction arithmetic on all 52 published outer table rows and their rounded Young costs. It also tests the exact projection identity, exact inner-overlap identity, signed root bound, completed-square bound, and optimal disk bound on 200 independent signed examples on a $3 \times 3 \times 3$ product space with nested cap and actual face domains. The largest floating identity discrepancy was below 4.5×10^{-16} . These finite tests are diagnostics; the arguments in § 2 – 6 supply the proofs.

Files:

- RESTORATION_PROOF_AUDIT.md — this derivation and audit;
- restoration_checks.py — exact ledger replay and finite-product diagnostics;
- restoration_checks.json — extracted 52 rows, exact fractions, numerical summaries;
- restoration_checks.log — complete run output.

The shortest substantive next obligation is now explicit: evaluate a certified lower mass on **true sufficient-failure regions**, or a certified signed residual/removed-marginal form for the actual source projection. Existing upper error ledgers cannot supply that lower mass. A successful new calculation must preserve independent erased-coordinate copies, retain the negative full-face term, and use one consistent restoration identity without adding the same credit twice.

No new $k = 39$ physical integral, restored positive margin, or prime-gap theorem is established by this audit.

Current report 05: REPORT

Source: research/prime-gaps/round4/k39-trial/REPORT.md. SHA256:

4c95f1f51767c736a5a48f4d58338ae6c617b9899ff18484fc915789a12a6be4 .

A direct $k=39$ cap-only trial with the published physical geometry

2026-09-05. Exploratory numerical result; no interval certificate, support-restored criterion, or new prime-gap bound.

Completed finite-family optimization: after optimizing all 77 polynomial coefficients at the same frozen geometry, the best directly re-evaluated $k = 39$ candidate has cap quotient approximately 0.99439639936, still 0.00560360064 below one. The full scaled Gram condition number is 2.28×10^{10} ; the matrix and direct scalar evaluations agree within 1.74×10^{-10} . This is numerical evidence about the fixed family, not a certified upper bound. The corresponding $k = 40$ optimized positive control is approximately 1.00021374364. Details and the retained-dimension checks are below.

The complete published 77-coefficient rational polynomial, evaluated in **39 outer coordinates and 38 retained coordinates** with the original physical radii, fragment caps and 98,304-cell grid, gives

$$\frac{\rho_* J_{\text{cap}}}{I} = 0.994361581476018.$$

The cap-only shortfall is approximately 0.005638418524, or 5,638 **parts per million**. This is a concrete fixed-trial deficit, not a bound on all 77-dimensional coefficients.

The independent $k = 40$ control gives

$$\rho_* J_{\text{cap}}/I = 1.000206086776951, \quad I = 2.36853178533315 \times 10^{-14},$$

with I inside the published interval

$$[2.3685317816, 2.3685317890] \times 10^{-14}.$$

The published cap quotient lower endpoint is 1.0002060794024186 ..., consistent with the independent value. No official FLINT check was bypassed; the program does not import FLINT or execute the official certificate.

Fixed input and dimension changes

Source: [OpenAI PrimeGaps186](#), commit 61340d0b74163003b32756bb16e91d9209a5e330, especially the companion numerical paper § § 1.1 – 1.2 and the two literal coefficient tables in the certificate.

The source script SHA256 is 7f71bdefcfe3bb5ca76a143929b3cb3f4156c21dc483253cda3077420f1e5de4. Only its literal signatures and integer coefficients are parsed as data. The exploratory implementation independently reconstructs the rational ladder and cap geometry.

The parameters retained are

$$\rho_* = \frac{2624989}{10^7}, \quad S = \frac{2742997}{2624989}, \quad T_0 = \frac{2499106033}{2624989000}, \quad T_1 = \frac{2510000}{2624989}.$$

Thus the physical outer/base/enlarged radii remain 0.2742997, 0.2499106033, 0.251. The fragment-cap indices at the official mesh remain

$$35265, 35419, 44781, 44976, 46580, 49152, 68225.$$

These are fragment caps; coordinate totals above a cap are retained with their Dickman density.

For $k = 39$, the convolution length is $98304 - 39 = 98265$, outer cell assignment uses $(r + 39)h$, the retained-face assignment uses $(r + 38)h$, and the radial midpoint is $(r + 39/2)h$. The erased-coordinate normalization is $39h/Z$, not $40h/Z$. Moment falling factorials use 39 or 38 as appropriate. No $k = 40$ denominator, source-loss bound or alpha estimate is inherited.

The fixed profile is

$$g(t) = \frac{21/200}{1 + t/100} + \frac{179/200}{1 + (907/5)t},$$

and the polynomial is the published rational linear combination of

$$(s - 9/10)^d P_\sigma(t), \quad 0 \leq d \leq 6, \quad \sigma \in \{\emptyset, 2, 3, 4, 5, 6, 22, 23, 24, 33, 222\}.$$

The same outer trial is used in every face form.

Independent numerical method

Conditional on a largest-fragment cap Z , the total-size measure is $\rho_D(t/Z) dt$. All needed arguments satisfy $t/Z < 3$. The implementation uses

$$\begin{aligned} \rho_D(x) &= 1 \quad (x \leq 1), & \rho_D(x) &= 1 - \log x \quad (1 < x \leq 2), \\ \rho_D(x) &= 1 - \log x + \log(x-1) \log x + \text{Li}_2(1-x) + \pi^2/12 \quad (2 < x \leq 3). \end{aligned}$$

Eight-point Gauss integration computes each cell mass. The rational midpoint profile and polynomial are constant within each cell, as in the official trial.

Power-sum products are reduced by set partitions to convolutions of coordinate measures with monomial weights.

Erasing a coordinate uses the exact expansion

$$P_\sigma(t_1, \dots, t_{k-1}, u) = \prod_{q \in \sigma} (P_q(t_1, \dots, t_{k-1}) + u^q).$$

The two erased-coordinate copies are integrated independently before squaring and integrating the retained configuration. Nested fragment-cap layers are handled by differences of the corresponding retained-coordinate moments.

An explicit exponential change of numerical normalization prevents the extremely small original denominator from being swamped by FFT roundoff. Put

$$Z_\tau = \sum_j g(t_j)^2 e^{-\tau t_j}, \quad w_j = \frac{g(t_j)^2 e^{-\tau t_j}}{Z_\tau}.$$

Every denominator contraction restores $e^{\tau s}$; every retained-face contraction restores $e^{\tau s_{\text{face}}}$, with normalization kh/Z_τ . The erased fiber still uses the original g . Thus the trial and geometry are unchanged. To return to the published normalization multiply every form by $(Z_\tau/Z_0)^k$, which cancels from the quotient.

The code requests NumPy longdouble but records the actual dtype. On this macOS runtime it is `float64`, so no extra-precision claim is made.

Completed values

k	cells	tilt τ	$\rho_* J/l$
40	4,096	20	0.995271191907594
40	16,384	20	0.999149113371267
40	98,304	20	1.000206086776951
39	16,384	20	0.993352220411709
39	98,304	20	0.994361581476018
39	98,304	25	0.994361581476014

Coarse rows change the step grid and its inward rounding; they are diagnostics, not the official fixed-grid certificate. The final two $k = 39$ rows retain the official grid and differ only in an algebraically cancelling numerical normalization. Their difference is 4.22×10^{-15} .

At $k = 39$, $N = 98304$, $\tau = 20$, the separated face values are

$$J_0/l = 3.780455375344233, \quad J_+/l = 0.007728822713122185, \quad J_t/l = 0.06588020236568319.$$

The signed cap form uses the published minorant parameters

$$\begin{aligned} \text{mass} &= 49999/50000, \quad K = 17/50, \quad \lambda = 1/125, \\ a &= \text{mass}^2 - \text{mass}\lambda, \quad b = (1 - \text{mass}/\lambda)(1 - \text{mass})K, \quad J_{\text{cap}} = J_0 + (a + b)J_+ + bJ_t. \end{aligned}$$

The ratio of absolute contraction sums to the final denominator is about 183; the corresponding ratios for the face pieces are about 93, 530 and 1033. These measure cancellation in the chosen representation, not rigorous error bounds.

Interpretation and earliest proof debt

The fixed $k = 39$ profile requires a change on the order of 0.00564 in the normalized criterion before it reaches one. The currently demonstrated $k = 40$ alpha-credit region is on the order of ten parts per million, roughly hundreds of times smaller. That comparison is a scale diagnostic only: a $k = 40$ alpha estimate is not a $k = 39$ bound, and a lower bound for one alpha region is not an upper bound for all possible alpha credit.

The remaining obligations are:

1. A genuine outward enclosure for this $k = 39$ cap form; current quadrature and FFT arithmetic are floating.
2. Separate treatment of the actual $k = 39$ rootwise support predicates and all restoration terms, including any credit attached to removal of failed outer roots.
3. An optimization or certified upper bound for the finite polynomial family before interpreting this fixed-vector deficit as a family-wide limitation.

4. The complete arithmetic sieve criterion, rather than a cap-only inequality, before any DHL[39,2] or prime-gap conclusion.

No restored $k = 39$ margin is asserted. Even a positive optimized cap-only quotient would not prove DHL[39,2].

Files and reproduction

- `cap_trial.py`: independently implemented cap-only computation; only coefficient literals are read from the preserved official clone.
- `k39_n98304_tilt20_longdouble.json`, `k39_n98304_tilt25_longdouble.json`: final fixed-grid $k = 39$ trials. Despite the requested-type filename, each JSON records the actual dtype as float64.
- `k40_n98304_tilt20_longdouble.json`: official-grid positive control.
- Other JSON files: coarse-grid diagnostics.
- `official_numerics.txt`: local pdftotext extraction of the unchanged official companion PDF, for equation review.

Run:

```
OPENBLAS_NUM_THREADS=1 python3 cap_trial.py --k 39 --intervals 98304 --tilt 20
```

The official-grid fixed-vector run took about 10.5 seconds in the observed concurrent run; this is a local observation, not a general performance claim.

Completed 77-dimensional optimization at frozen geometry

`optimize_cap.py` assembles the denominator Gram G and the numerator matrix $B = \rho_* J_{\text{cap}}$ directly, without searching coefficients during integration. All coefficients vary in the same 77-dimensional polynomial space; the physical support and coordinate profile g remain frozen.

For the denominator, entries with the same joined power-sum signature and summed radial degree share one moment contraction. For the numerator, the exact erased-coordinate expansion expresses each basis function as a finite sum of retained power-sum signatures times one-dimensional fiber kernels. Weighted matrix products assemble their shared retained-coordinate integrals. This is a different contraction order from evaluating one fixed coefficient vector.

The matrices are diagonally scaled by $\sqrt{G_{jj}}$. Three relative eigenvalue thresholds are applied to the scaled Gram before whitening: 10^{-8} , 10^{-10} , 10^{-12} . At the last threshold all 77 dimensions survive. Each resulting candidate is then passed to the separate scalar-form implementation in `cap_trial.py`.

At $N = 98304$, the results are:

k	retained dimension	scaled Gram condition	matrix quotient	direct candidate quotient
39	62	9.01×10^7	0.994371194303	0.994371194271
39	75	3.65×10^9	0.994396034484	0.994396034594
39	77	2.28×10^{10}	0.994396399191	0.994396399364
40	62	9.49×10^7	1.000188112833	1.000188112844

k	retained dimension	scaled Gram condition	matrix quotient	direct candidate quotient
40	75	3.86×10^9	1.000213094616	1.000213094394
40	77	2.42×10^{10}	1.000213743635	1.000213743640

For $k = 39$, the scaled Gram eigenvalues range from 2.30246×10^{-9} to 52.4410 . The full-dimensional projected eigen-residual is 1.24×10^{-12} ; the residual of the scaled generalized pencil, divided by the documented matrix-norm bound, is 7.36×10^{-17} . These residuals concern the assembled numerical matrices, not unknown integration errors.

The separate scalar evaluation of the 77-dimensional candidate has an absolute denominator-contraction ratio around 368 and face-contraction ratios around 180, 1194 and 1805. It differs from the matrix quotient by 1.74×10^{-10} , far below the observed shortfall of 0.00560, but this discrepancy is not itself a rigorous error enclosure.

The optimization recovers only approximately **34.82 parts per million** relative to the published coefficients transplanted to $k = 39$. It leaves approximately **5603.60 parts per million** of cap-only shortfall. The 75-to-77-dimensional change is under one part per million. On this frozen geometry the existing polynomial coefficients are numerically close to the best found value.

The $k = 40$ control is positive in every retained subspace, and the full-family optimization improves its existing cap value by approximately 7.66 parts per million. Thus the experiment is not globally misnormalized so that every trial fails.

The complete matrix-plus-three-validation run took about 58 seconds for each dimension in the observed concurrent run. Files:

- `optimize_cap.py`
- `ritz_k39_n98304.json`, `ritz_k39_n98304.npz`
- `ritz_k40_n98304.json`, `ritz_k40_n98304.npz`
- `ritz_k39_n16384.json`, `ritz_k39_n16384.npz`: initial coarse-grid assembly check.

The JSON retains floating coefficient vectors, all scaled Gram eigenvalues, thresholds, residuals and direct candidate evaluations. The NPZ retains the raw denominator and numerator matrices and original coefficient vector.

Reproduce:

```
OPENBLAS_NUM_THREADS=1 python3 optimize_cap.py --k 39 --intervals 98304 --validate
```

Structural checks and next decision

`check_structure.py` checks the set-partition moment expansion against exhaustive three-coordinate summation on a four-point measure, and checks the erased-coordinate polynomial identity for all 11 signatures. Errors are below 3×10^{-17} . It also verifies the exact $k = 40$ cap-index ranges against the companion paper's table and compares the closed Dickman expression at 3 against an independent integral. Results are in `structural_check_results.json`.

The official clone's git status remains clean. No official certificate code, dependency or assertion was modified. All three local scripts pass Python compilation.

The appropriate next mathematical decision is to stop repeated scans of these 77 coefficients at the frozen geometry. A rigorous finite-family upper bound would require outward enclosures for the two matrices and a certified positive-definiteness test for $cG - B$, for some explicit $c < 1$, after a numerically suitable basis change. The current float64 matrices and eigensolver output do not supply that certificate.

A different radius/owner-support allocation may change the cap form substantially. Such a proposal must recompute its source-dependent caps and support restoration; it cannot simply retain the current cap arrays while changing S or T . No radius scan or new support claim is included here.

Current report 06: README

Source: [research/prime-gaps/round4/repro-flint/README.md](#). SHA256:

4612dce4b1824923b559ee988f531c6a65c01f2cf46214a4e72261c1dfffb04b8 .

Corrected arithmetic runtime for the prime-186 interval certificate

2026-09-05. This removes a specific reproducibility obstacle. It does not prove a prime-gap bound, discharge a Lean axiom, or replace the source certificate's checks.

The unchanged official signed-convolution regression fails with the previously installed Python-FLINT 0.9.0 / FLINT 3.6.0 wheel. A separate source-built FLINT 3.6.0 with the upstream correction passes that regression. Python-FLINT is built from source against the corrected prefix. No monkeypatch, replacement convolution, disabled assertion, or edited official certificate is used.

Pinned inputs

- FLINT v3.6.0 archive: <https://codeload.github.com/flintlib/flint/tar.gz/refs/tags/v3.6.0>, SHA256
4307a504622702bf0be6d8969791f7d7ff378645cf2ae3bb5a7a2b56653d97f1 .
- Upstream signed FFT fix, from PR2790. Download its .patch URL as flint-7ad753d.patch; SHA256
333788fe3d7fe1c24ca10e5ef33f492eae68de6202568e75e18e1bcd7bfb71ff .
- Official PrimeGaps186 source at 61340d0b74163003b32756bb16e91d9209a5e330 ;
prime_gap_186_certificate.py SHA256
7f71bdefcfe3bb5ca76a143929b3cb3f4156c21dc483253cda3077420f1e5de4 .
- Host build: Apple Silicon macOS, clang, Python 3.12.9, Python-FLINT 0.9.0, NumPy 2.2.6, Homebrew GMP 6.3.0 and MPFR 4.2.2. FLINT was configured with --enable-assert --disable-static .

The upstream fix corrects the conversion of a residue near the halfway point of the FFT coefficient ring into a signed integer. Testing only the leading limbs misses boundary cases with nonzero lower limbs. The patch contains both the correction and targeted native test cases. The version string remains 3.6.0; the patch hash and actual linked-library path distinguish this runtime from the failing wheel.

What passed

1. Native FLINT suites fmpz_vec , fmpz_poly , arb , arb_poly .
2. The original certificate's check_flint_signed_fft() and _cap_check_environment() , called unchanged.
3. 467 full products and 2,188 truncated products compared with an independently implemented Python integer double loop. These include both signs, limb boundaries, the original 509/510-bit failure, zero polynomials, and deterministic random coefficients. The comparison took approximately 0.80 seconds here.
4. Ten selected Python binding tests for the integer/rational polynomial and Arb APIs used by the certificate.

`otool -L` confirms that the extension loads the corrected isolated `libflint.24.0.dylib`. The old wheel remains available and fails as a negative control, as recorded in `negative-wheel-control.json`.

A separately disclosed full-suite failure

The complete Python-FLINT suite did not pass: with native assertions enabled, `test_fmpz_functions` aborts in `_n_jacobi_unsigned`. Its test table invokes `fmpz(2).jacobi(n)` at zero and even denominators. The installed FLINT source documentation for `fmpz_jacobi` specifies an odd positive denominator and says parity/sign are not checked. The Python wrapper forwards the input without enforcing that precondition.

The certificate uses no Jacobi call. The native integer/polynomial and Arb suites, its own regression, and the selected binding tests passed independently. We retain the abort logs and keep assertions enabled; we do not claim that all Python-FLINT tests passed or that the library has been universally verified. Resolving the separate Jacobi wrapper/test contract is postponed.

Reproduction on this host

The source/build/install cache is separate from the Dropbox research record. The original PrimeGaps186 clone is unchanged. For a fresh directory without spaces:

```
bash build_runtime.sh DOWNLOAD_DIRECTORY BUILD_DIRECTORY
bash build_python_binding.sh BUILD_DIRECTORY PYTHON_EXECUTABLE
BUILD_DIRECTORY/venv/bin/python signed_convolution_check.py \
  --official-script PATH_TO_UNCHANGED_CERTIFICATE \
  --output NEW_RECEIPT.json
```

The first script verifies both source download hashes, refuses to overwrite an existing checkout, applies the upstream patch, compiles, runs the four native suites, and installs to `BUILD_DIRECTORY/prefix`. The second creates a fresh virtual environment and builds the binding with the needed macOS GMP link path. GMP, MPFR, autoconf, automake and libtool must already be installed. The script assumes the Homebrew paths used on this host; other platforms should adjust the toolchain paths while preserving source hashes and the regression.

On this run the usable interpreter is:

```
/Users/qingyunsun/.cache/astra-research/flint-3.6.0-patched/venv/bin/python
```

The binding build first encountered two configuration errors: the optional `add_flint_rpath` setting emitted a linker flag unsupported by Apple's linker; without it, the transitive GMP library still needed an explicit `-L` path. The successful build omits the former and supplies the latter. These errors and the exact successful commands are retained in their separate logs. They caused no change to the mathematics or the official certificate.

Acceptance and limits

The purpose of this runtime is to execute new outward-enclosed integrals with the original arithmetic safeguards. Round four's first application is a positive lower integral on one genuine outer failure rectangle. It does not require recomputing all 149 existing upper integrals merely to validate the new rectangle; combining the new credit with the old final margin explicitly inherits the published upper endpoints.

Full new `k=39` support restoration, complete physical-integral replay, a portable binary distribution, and proof-assistant verification remain separate work. The runtime can be abandoned by using another interpreter; the original wheel and official source were not overwritten.

Current report 07: prime186_round5

Source: [research/reports/prime186_round5.md](#). SHA256:

7ad8c798736bd40b7125784af7ec69a208075665c8efcfd69d4c33adf68b447e .

Round 5: radius-dependent sieve bounds and a negative geometry search

The bounded search is complete. Ten radius/plateau configurations were evaluated with all 77 polynomial coefficients reoptimized at a coarse grid, and two nearby candidates were refined at the official grid. **Neither geometry change improves the original $k=39$ result.** The useful mathematical output is a proved variable-radius exceptional-square estimate, an exact sufficient source template, and a precise repair for a newly activated source row. None proves DHL(39,2), a prime-gap bound below 186, or a global obstruction to those goals.

This report follows [Round 4](#), where a new outward lower integral increased the fixed $k=40$ proof margin to 24.86626 ppm by combining the new credit with inherited published endpoints. That result remains valid in its stated scope. Its approximately 1.5 ppm credit cannot be assigned to new $k=39$ vectors or supports without computing their own integrals.

1. Evidence and proof status

Three independently assigned research components are preserved with their original scripts, outputs and reports:

Component	Result	Evidence status
Exceptional-square extension	A radius-dependent constant follows from the existing counting argument over an explicit open interval	Ordinary written proof plus exact rational certificate; not Lean formalization
Source geometry audit	Twelve of fifteen natural cap templates pass, with a uniform one-layer mesh repair and a common inner-square source	Exact arithmetic and a sufficient-template proof; no physical sieve integrals
Finite geometry search	Ten coarse points and two refinements give no improvement from changing radius/plateau	Floating cap-only optimization; no interval optimum or support restoration

The primary proof is [Improved short gaps between primes](#), especially Propositions 3.10, 3.11, 4.2 and 4.6. The companion certificate is pinned to [PrimeGaps186 commit 61340d0](#). Its Python source SHA256 is

7f71bdefcfe3bb5ca76a143929b3cb3f4156c21dc483253cda3077420f1e5de4 . The preserved source is unchanged.

2. A parameterized exceptional-square estimate

The original proposition bounds an exceptional square with a coefficient-root radius at most $11/40$. Its proof uses this radius in the ordinary CRT error, not in a new prime-distribution hypothesis. Let r bound every outer, inner, correction and exact-face coefficient root in the finite canonical combination. Retain the original global physical prime cap 0.19037, the fixed-profile assumptions and the original exceptional majorant.

$$b_h = -\frac{49599}{20000000}K,$$

and all derived costs. The convenient old shortcut $|b_h| < 10^{-3}$ fails with the displayed constants at 0.280 and 0.282. The counting proposition remains valid there, but downstream restoration needs its actual operator bound and coefficient checks. Only radii through 0.278 were used in the present cap search.

3. Geometry that preserves the distribution ladders

The screen keeps $\rho = 0.262499$, $\rho_* = 0.2624989$ and imposes

$$S = r/\rho_*, \quad T_1 = (0.5252997 - r)/\rho_*, \quad T_0 = 1.997 - S.$$

Thus the old and new combined root sums stay fixed. The nominal distribution ladders remain the same. Their root thresholds, actual grid endpoints, and retained rows do not all remain the same.

For $\epsilon = 10^{-7}/\rho$, put $A = S + \epsilon/2$ and $C_v = T_v + \epsilon/2$. The complementary allocations

$$\phi_D(t) = \min(3t/2, L_v), \quad \phi_E(t) = 3t - \phi_D(t)$$

give a sufficient natural plateau template when

$$\frac{3A - C_v}{4} \leq L_v \leq \frac{3C_v}{5}.$$

The audit checks the largest-fragment owner and opposite-root inequalities and the nonlargest-witness reduction. Both common-height choices

$$L_0 = L_1 = (3A - C_0)/4 \quad \text{or} \quad L_0 = L_1 = 3C_0/5$$

preserve both source intervals and nested inner caps. Equal heights avoid an unnecessary extra outer-cap cost caused by the larger new inner radius. This is a valid exact simplification even though the sampled numerical values did not improve.

Fifteen radius/plateau cases were audited: five radii and three choices each. Twelve pass. The unchanged published fraction $L_v/C_v = 0.575$ fails the natural cap formulas at 0.275, 0.276 and 0.278. Those failures concern this sufficient cap template; they do not invalidate all sieve supports at those radii. The original fraction was used in the numerical screen only where it passes.

At $r=0.272$ the old inner-square source level 0.5062 is too small. A common replacement $(\omega_s, \delta_s) = (0.0035, 0.025)$ gives level 0.507 and satisfies the requisite source and row-12 containment inequalities throughout the screened interval. A rational exponential-sum inequality also rechecks $C_{\text{op}} = 4$ at $k=39$ throughout this interval. No $k=40$ operator constant is silently inherited.

4. An actual numerical-cover obstruction and its repair

At the official 98,304 grid, the untrimmed points $r=0.272, 0.275, 0.276$ and 0.278 retain new-ladder row 39. Its activation width is approximately 1.8866250×10^{-5} , less than two grid cells. The source theorem is still applicable, but the original low-witness numerical implementation's two-cell guard fails.

Let $h = S/98304$, $J_1 = \lfloor T_1/h \rfloor$ and

$$J_O = \min\{98303, \lfloor B_{n,39}/h \rfloor - J_1\}.$$

Restrict the outer index sum to $\sum_i j_i \leq J_O - k$. This removes at most one outer layer throughout the entire interval $[0.272, 0.278]$, excludes row 39, retains row 38, and leaves its activation width greater than two cells. The exact audit proves the uniform statement using endpoint inequalities; it is not an interpolation from the five sampled radii.

The repair keeps h , the normalizer, the convolution length and nominal S fixed. Every erased face must be derived again from the trimmed outer function. **The numerical trials reported here are untrimmed.** Neither their quotients nor the original 97-component numerical cover can be represented as a completed repaired certificate.

Even after trimming, the failure cover must be regenerated from the new source thresholds, caps, core boundaries and same-coordinate terms. All resulting physical integrals need fresh outward evaluation. The old 149 upper integrals and Young-cost values are not transferable solely because the analytic inequalities pass.

5. What the finite search found

All configurations retain the original product profile and the 77-dimensional polynomial span. Every vector is optimized afresh. The exact configuration files, full matrices, eigenvectors, conditioning, scalar reevaluations and run times are in the [search archive](#).

Fine-grid candidate	Direct cap quotient	Difference from the Round 4 optimized original geometry
Round 4 original $r=0.2742997$, original plateaus	0.9943963993644909	reference
$r=0.2742997$, common maximum height	0.9943734016224463	−22.9977 ppm
$r=0.275$, common maximum height	0.9943501891039260	−46.2103 ppm

Ten coarse configurations covered $r=0.272$ through 0.278 . The only coarse increase over the original baseline, approximately 0.6183 ppm, came from improving the exceptional constant at the unchanged original geometry; it did not come from a radius/plateau change. Only the two nearby geometry candidates above were refined. The full coarse table is preserved in the detailed report.

The scaled Gram condition number reaches approximately 4.307×10^{10} . For the twelve full-77 candidates used in the comparisons above, matrix and direct candidate evaluations disagree by at most 2.884×10^{-10} , and the full scaled-pencil relative residual is at most 5.30×10^{-16} . Across all thirty-six saved candidates, including truncated whitening spaces, those maxima are instead 3.0281×10^{-10} and 4.915×10^{-9} . This qualifies the scope of the smaller maxima in the original search report. The truncated residual is measured against the full pencil, so it need not be at roundoff scale.

These comparisons support the observed ordering at the tens-of-ppm scale. They do not provide outward error bounds, a certified finite-family maximum, or a no-go theorem for larger trial spaces. The host's NumPy long-double type has only 64 bits; no extended precision is claimed.

The approximately 5604 ppm original $k=39$ deficit remains far larger than the particular fixed- $k=40$ credit recovered in Round 4. This comparison guides effort; it does not upper-bound every possible restoration credit.

6. Integration verification and next decision

The primary integration replay runs in a temporary copy. It checks the seven exact exceptional constants, all fifteen source cases, twenty grid-mask nesting cases, and the saved twelve 77-by-77 matrix archives with their thirty-six candidate vectors. Source SHA verification is mandatory. Original outputs are preserved; only timing and optional source-text metadata are excluded from exact replay comparisons. See [recheck.py](#) and [receipt](#).

The 53-file [intake manifest](#) records original and published hashes. The only two code changes made for publication let an external primary-source path be supplied through `PRIME186_SOURCE`. Original per-run manifests remain intact and describe their own execution snapshots.

The next meaningful prime-gap experiment should change a component large enough to affect the $k=39$ deficit, such as the product profile or an analytically justified support family, or establish a certified upper bound for the currently searched finite family. Repeating these endpoint scans is postponed. Full restored certificates for the present negative candidates are also postponed. This prime-sieve work supports the wider research programme but supplies no new arithmetic transfer theorem for zeta zeros.

This is a checkpoint in an active research goal. The 333-page public handoff remains the earlier `055a4a0` snapshot; Rounds 4 and 5 are separate subsequent reports. No major-conjecture completion is claimed.

Current report 08: EXCEPTIONAL_RADIUS_EXTENSION

Source: `research/prime-gaps/round5/exceptional-radius/EXCEPTIONAL_RADIUS_EXTENSION.md`. SHA256:
`70326c9bd4430855349a64fc2863e0dccc80f40e9b0b5b28750925e0c7de54b`.

A variable-radius version of the exceptional-square estimate

Status: ordinary written derivation from the primary paper, with an exact rational reproduction of its finite numerical certificate. This extends the stated parameter range of Proposition 3.11 without adding a prime-distribution hypothesis. It does not establish a new prime-gap bound or certify a new support geometry.

Primary source: OpenAI, *Improved short gaps between primes*, dated 30 August 2026, Propositions 2.23, 3.10 and 3.11, Lemmas 3.2, 3.3 and 3.9, equations (3.26) – (3.35). The [primary PDF](#) has Proposition 3.11 on page 23. The locally retained text is `../sources/openai-short-gaps.txt`, lines 1522 – 1607. The local text is the source actually checked in this audit; its SHA256 is `ded13a7c74fcfce64e85769e05b5869803dcdcf53b88be2c2f3c0b344f95ee84`.

Artifacts: `certify_exceptional_radius.py` is a standalone Python-standard-library certificate. `certify_exceptional_radius.json` contains the exact reduced rational upper bounds and all endpoint checks. No float arithmetic or numerical quadrature enters those upper bounds.

1. Conclusion and certified constants

Replace the fixed physical coefficient radius $11/40$ in Proposition 3.11 by a fixed radius r . Keep the same prime cap, exceptional minorant, fixed-profile assumptions, 1024 bins, and strict CRT slack. The same argument works throughout

$$\frac{4879903}{40960000} < r < \frac{59499}{200000}, \quad 0.1191382568359375 < r < 0.297495. \quad (1)$$

This range has the convenient stronger condition $0 < z_j < \xi_0 < \xi_*$ for every bin. Proposition 3.10 itself only requires $z_j < \xi_*$, giving the slightly wider lower endpoint in Section 3.

The following terminating decimals are **exact rational upper bounds**, rather than rounded numerical approximations. Every last printed digit belongs to the certificate.

Physical coefficient radius r	Certified $K_{1024}^{+}(r)$, exact decimal	Simpler safe upper bound
0.272	0.3014041534851816226069683	0.301405
0.2742997	0.3273225381113663650584938	0.327323
0.275	0.3361336040272905676441604	0.336134
0.276	0.3495799968949037559942978	0.349580
0.278	0.3800259215656200578230129	0.380026
0.280	0.4163697504037337478611794	0.416370

Physical coefficient radius r	Certified $K_{1024}^{+}(r)$, exact decimal	Simpler safe upper bound
0.282	0.4605417963468921175216614	0.460542

At $r=11/40$ the exact reduced fraction is

$$K_{1024}^{+}(11/40) = \frac{840334010068226419110401}{25000000000000000000000000}, \quad (2)$$

which reproduces equation (3.35) of the paper exactly. In particular, the scale printed as “1025” by the text extraction is **10 to the power 25**, not the integer 1025.

One may safely use any K at least $K_{1024}^{+}(r)$ in the exceptional-square conclusion. The same finite-bin proof cannot justify retaining $K=0.34$ at $r \geq 0.276$: a separate exact lower bound on its bin constant already exceeds 0.34 there; see Section 5. This does not assert that every other possible method of bounding the exceptional sum must fail with 0.34. Conversely, using the actual physical radius 0.2742997 in this proposition permits the smaller exceptional constant in the table, subject to the common-support condition below.

2. Precise parametric proposition

Keep the setting and normalization P_x of Proposition 3.11. Fix

$$a_* = \frac{40481}{100000}, \quad \xi_* = \frac{9519}{50000}, \quad \xi_0 = \frac{19037}{100000} < \xi_*.$$

Let C_i be a finite real linear combination, with coefficients independent of x , of the same canonical $(k-1)$ -coordinate sums and exact faces used in that proposition. Assume their profiles are fixed and bounded and have limiting-null discontinuity sets. Let H_i denote their corresponding combination of canonical profiles and marginals.

Assume every summand has physical coefficient-root radius at most r , and every outer and inner coefficient root satisfies the original global largest-prime-factor cap

$$P^+(\prod_j d_j) \leq x^{\xi_0}. \quad (3)$$

The radius assumption must hold for **all roots actually used in the linear combination**. In particular, it is insufficient to use the radius of an outer trial function if an inner correction, another summand, or a separately constructed exact-face array has larger retained-coordinate radius. A valid sufficient choice is the maximum over all outer, inner, correction and exact-face radii. An exact face obtained through Lemma 3.3 retains the same total support bound as its parent array.

For fixed r in (1), define $K_{1024}^{+}(r)$ by equation (12) below. Then

$$\sum_n b(n + h_i; x) C_i(n)^2 \leq P_x(K_{1024}^{+}(r) \| H_i \|_2^2 + \alpha(1)). \quad (4)$$

The norm remains the full fragment norm. Every profile, radius, cutoff and finite band choice is fixed before x tends to infinity; exact face arrays may depend on x as in the original proposition. No assertion uniform as r approaches the upper endpoint is made.

This is deliberately stated for the same canonical class as the source. It does not assert that an arbitrary coefficient array has a bounded inverse diagonal merely from a radius bound.

3. Exact parameter range and the strict counting slack

Write

$$h = \frac{a_* - 2\xi_*}{1024} = \frac{481}{20480000}, \quad s_j = 2\xi_* + jh \quad (1 \leq j \leq 1024),$$

$$\alpha(r) = 1 - 2r - \frac{1}{5000}, \quad z_j(r) = \frac{\alpha(r) - s_j}{2} = \frac{1 - 2r - s_j}{2} - \frac{1}{10000}. \quad (5)$$

Since s_j increases, the minimum z occurs at $j=1024$ and the maximum at $j=1$. Direct rational arithmetic gives

$$z_{1024}(r) = \frac{59499}{200000} - r = 0.297495 - r,$$

$$z_1(r) = 0.3095082568359375 - r. \quad (6)$$

Thus all z_j are positive exactly when $r < 0.297495$. The stronger condition $z_1 < \xi_0$ is exactly the lower bound in (1). The necessary condition for Proposition 3.10 is only $z_1 < \xi_*$, which gives

$$\frac{24397467}{204800000} < r < \frac{59499}{200000}, \quad 0.1191282568359375 < r < 0.297495. \quad (7)$$

Both are strict ranges. At the upper endpoint the last auxiliary exponent is zero, and the proof's Selberg-energy asymptotic with a fixed positive exponent is unavailable.

For every bin, every valid r has exactly the original counting margin:

$$s_j + 2r + 2z_j = 1 - \frac{1}{5000} = \frac{4999}{5000} < 1. \quad (8)$$

The following endpoint values show the requested radii lie comfortably inside the valid range.

r	Minimum z_j	Maximum z_j
0.272	0.025495	0.0375082568359375
0.2742997	0.0231953	0.0352085568359375
0.275	0.022495	0.0345082568359375
0.276	0.021495	0.0335082568359375
0.278	0.019495	0.0315082568359375
0.280	0.017495	0.0295082568359375
0.282	0.015495	0.0275082568359375

The original displayed bound $z_j < 863/25000 = 0.03452$ is specific to $r=0.275$ and need not be retained. The operative inequalities are positivity, $z_j < \xi_*$, the global coefficient cap $\xi_0 < \xi_*$, and (8).

4. Arithmetic proof of the extension

Diagonal reduction. Apply exactly the inversion and face identities from Lemmas 3.2 and 3.3, as in the first paragraph of Proposition 3.11. They express C_{-j} as $B_{-}\{y_{-x,i}\}$ for a uniformly bounded raw array with the same physical radius bound and hereditary prime cap. The full diagonal norm converges to the norm of H_{-j} . Neither identity uses the numerical value $11/40$.

Marked-prime separation. In Proposition 3.10 use the ambient retained-index cap $\min\{\zeta_{\ast}, x_{i0}/\rho_{\ast}\}$. Its physical exponent is at most x_{i0} . Every marked prime p or q has exponent at least $\xi_{\ast}$, so it exceeds both that cap and the auxiliary cutoff exponent z_{-j} . These are exactly the separation hypotheses $\xi_{\ast} > \max\{\rho_{\ast}, \zeta_{\text{eff}, z_{-j}}\}$ needed by Proposition 3.10.

Nonsquarefree exceptions. Proposition 2.23 bounds the number of nonsquarefree points in the exceptional support by $O(x^{o(1-\xi_{\ast})})$. Lemma 3.2 and the divisor bound make each fixed canonical sum $x^{o(1)}$ pointwise; a finite linear combination and its square have the same form of bound. Consequently their total contribution is $x^{o(1-\xi_{\ast}+o(1))}=o(P_{-x})$. A fixed change of radius does not remove this fixed power saving. The same original fixed-shift endpoint treatment remains applicable.

Positive majorant. On the remaining squarefree support, Proposition 2.23 gives $b \leq (12/5)N_{-2}$, where N_{-2} counts unordered marked pairs $p < q$ with $p, q \geq x^{\wedge}(\xi_{\ast})$ and $pq < x^{\wedge}(a_{-}^{\ast})$. On the exceptional support all prime factors are at least $x^{\wedge}(\xi_{\ast})$; Lemma 3.9 therefore gives $L_{-}\{x^{\wedge}(z_{-j})\}=1$ for every selected bin. The majorant is positive before opening any signed coefficient expansion, so the full square C_{-j}^2 remains intact.

Counting each bin. For pairs in the j th logarithmic-product bin, $pq \leq x^{\wedge}(s_{-j})$. Apply Proposition 3.10 with radius r and exponent z_{-j} . Its ordinary CRT error is $x^{\wedge}(s_{-j}+2r+2z_{-j}+o(1))=x^{\wedge}(4999/5000+o(1))=o(P_{-x})$, by (8). No distribution theorem for primes or the minorant enters this counting step. There are only 1024 bins, fixed independently of x .

Limiting pair measure. The same prime harmonic measure as in the source gives

$$d\mu_2(s) = f(s) ds, \quad f(s) = \frac{1}{s} \log \frac{s - \xi_{\ast}}{\xi_{\ast}}, \quad 2\xi_{\ast} \leq s \leq a_{\ast}. \quad (9)$$

The strict arithmetic condition $pq < x^{\wedge}(a_{-}^{\ast})$ remains in the last bin. Its limiting measure is unchanged by inclusion of the endpoint because this measure has no atom there. Thus the exact limiting constant supplied by the bin argument is

$$K_{\text{bin}}(r) = \frac{12}{5} \sum_{j=1}^{1024} \frac{\mu_2((s_{j-1}, s_j])}{z_j(r)}. \quad (10)$$

It remains to certify a number above this finite sum; Section 5 does so. This proves (4). All applications are instances of the primary ordinary-counting argument with a fixed changed radius.

5. Finite-sum certificate and why the continuum integral alone is insufficient

The natural continuous expression is

$$K_{\text{cont}}(r) = \frac{24}{5} \int_{2\xi_{\ast}}^{a_{\ast}} \frac{\log((s - \xi_{\ast})/\xi_{\ast})}{s(\alpha(r) - s)} ds. \quad (11)$$

Within each bin, $1/(c(r)-s)$ is increasing. Therefore $K_{\text{cont}}(r) \leq K_{\text{bin}}(r)$. A good approximation or even a certified upper bound for K_{cont} alone does **not** automatically upper-bound the larger K_{bin} delivered by the actual arithmetic proof. We retain the paper's finite-bin certificate.

First, f in (9) is increasing throughout the interval. Differentiation yields

$$f'(s) = \frac{s/(s - \xi_*) - \log((s - \xi_*)/\xi_*)}{s^2}.$$

Set $y = (s - \xi_*)/\xi_*$. Here $1 \leq y < 2$, since $a_- < 3x_i$. Hence $\log y < 1$ while $s/(s - \xi_*) = 1 + 1/y > 1$. The derivative is positive. In consequence, the measure of each bin is at most $h f(s_j)$.

Set $t_j = (s_j - 2x_i^*)/\xi_*$. These rational numbers belong to $(0,1)$. The degree-21 alternating polynomial

$$L_{21}(t) = \sum_{m=1}^{21} \frac{(-1)^{m+1} t^m}{m}$$

satisfies $\log(1+t) \leq L_{21}(t)$ by the alternating-series remainder. Each bin contribution in (10) is therefore at most

$$u_j(r) = \frac{24hL_{21}(t_j)}{5s_j(\alpha(r) - s_j)}.$$

Every $u_j(r)$ is positive and rational for rational r . Define

$$K_{1024}^+(r) = \frac{1}{10^{25}} \sum_{j=1}^{1024} \lceil 10^{25} u_j(r) \rceil. \quad (12)$$

Then, in the correct direction,

$$K_{\text{cont}}(r) \leq K_{\text{bin}}(r) \leq \sum_j u_j(r) \leq K_{1024}^+(r). \quad (13)$$

The script implements (12) with exact integers and fractions. Its ceiling operation is $-\text{floor}(-\text{numerator}/\text{denominator})$, with positive denominator. It never converts an input radius or an intermediate summand to binary floating point. An input decimal such as 0.2742997 is parsed as its exact rational value.

The accumulated upward-rounding excess is strictly below $1024 \times 10^{-(25)} = 1.024 \times 10^{-(22)}$. The alternating-polynomial excess in the sum is at most

$$\sum_j \frac{24h}{5s_j(\alpha(r) - s_j)} \frac{t_j^{22}}{22}.$$

For the requested radii this latter bound is between $1.87 \times 10^{-(22)}$ and $3.06 \times 10^{-(22)}$. The coarser right-endpoint discretization remains the main deliberate overestimate. Reducing it is a separate possible numerical refinement, not needed here.

For an additional check, the script also constructs a lower bound on K_{bin} . Since f is increasing and the even alternating polynomial L_{22} lies below $\log(1+t)$, the j th bin contribution is at least

$$\ell_j(r) = \frac{24hL_{22}(t_{j-1})}{5s_{j-1}(\alpha(r) - s_j)}.$$

Rounding each nonnegative term down at scale $10^{(-25)}$ gives a rigorous rational lower bound on K_{bin} . At $r=0.276$ it equals 0.3489733171373295715615596, which already exceeds 0.34. At $r=0.280$ and $r=0.282$ these lower bounds are respectively 0.4156565419431563765185887 and 0.4597595157187387148411848. Thus the need to update the constant in this bin proof is not inferred merely from its upper bound exceeding 0.34.

All factors depending on r are increasing as r increases within the valid interval. In particular, both the continuum constant and the finite-bin bound are increasing, and the rounded bound in (12) is nondecreasing. A constant certified at an upper radius R consequently applies to all smaller radii for which the same proof conditions hold.

6. What must change downstream when this constant is used

This result supplies an exceptional-square input to the hybrid sieve; it does not certify new mixed-modulus factorizations or source conditions. Those conditions, support realizability, the actual retained prime cap and the root radii require their own checks.

If the original choices $m_0=49999/50000$, $\lambda=1/125$ and $\kappa_{\text{def}}=1/50000$ are retained, equation (4.28) becomes

$$b_h(r) = \left(1 - \frac{m_0}{\lambda}\right) \kappa_{\text{def}} K = -\frac{49599}{20000000} K, \quad K \geq K_{1024}^+(r). \quad (14)$$

One must therefore recompute b_h , d_0 and every numerical functional or cover weight depending on them. The source's convenient assertion $|b_h| < 1/1000$ is equivalent here to $K < 20000/49599$. It fails for the constants in the table at $r=0.280$ and $r=0.282$. This failure does not invalidate the hybrid algebra: it means that a downstream proof cannot cite that old numerical shortcut unchanged.

In particular, restoration should check its actual coefficient

$$1 - \rho_* \mid b_h(r) \mid C_{\text{op}} > 0$$

using the operator bound justified for the new support. Do not reuse $C_{\text{op}}=4$ beyond its proved support range simply because the exceptional counting proof remains valid there. The primary statement also requires $d_0 > 0$ and $0 < a_h + b_h < 1$ wherever those multiplier bounds are invoked.

The exceptional proposition is dimension-independent at the level of this displayed constant for fixed k , but that observation does not transfer a $k=40$ numerical trial, norm certificate or prime-gap conclusion to $k=39$.

7. Reproduction and audit record

Run the standalone script with Python 3:

```
python3 certify_exceptional_radius.py
```

It writes `certify_exceptional_radius.json` next to itself. Optional radii can be supplied as exact decimal or rational strings with `--radii`. The implemented certificate deliberately enforces the stronger range (1), even though Section 3 explains the slightly weaker necessary condition.

Checks completed in this audit:

1. Read the primary proofs of Propositions 2.23, 3.10 and 3.11 and the relevant support/face identities; identified exactly where radius enters.

2. Verified every one of the 1024 rational auxiliary exponents is positive, below ξ_0 , and has the identical strict CRT margin, for all seven requested radii.
3. Verified the odd/even alternating-polynomial difference is $t^{22}/22$ and positivity holds at every endpoint.
4. Reproduced the primary fraction in (3.35) exactly, not only its decimal digits.
5. Verified the rounding excess bound and monotonic order of the seven certified constants.
6. Cross-checked the first five values against the prime agent's separate implementation; they agree to all communicated digits, with differences in the final displayed digit explained by rounding versus this document's exact terminating decimals.
7. Added a separately directed $L_{22}/\text{left-density}/\text{downward-rounding}$ lower bound on the actual bin constant to verify the strict obstruction to retaining 0.34 in this proof at $r \geq 0.276$.

Large integer numerator and denominator fields are serialized as decimal strings in the JSON, to prevent loss when it is read by a JavaScript consumer. Reconstruct the exact fraction from those strings rather than passing the decimal display through a binary float.

The certificate is exact arithmetic accompanying an ordinary proof. It has not been formalized in Lean, and this note makes no claim that an improved global prime-gap inequality has yet been certified.

Current report 09: GEOMETRY_SOURCE_AUDIT

Source: [research/prime-gaps/round5/geometry-audit/GEOMETRY_SOURCE_AUDIT.md](#). SHA256:

c2870e1d539e1be90e82131671ec15b85d74a0a886bad299270a20cb87ecc7c9.

Exact source geometry for the bounded radius and plateau screen

This audit establishes a sufficient source and cap template for the round-five $k = 39$ screen. It evaluates no sieve integrals and proves no smaller prime gap. The accompanying exact arithmetic checks cover physical outer radii 0.272, 0.2742997, 0.275, 0.276, and 0.278, together with three specified plateau choices at each radius. The original source's numerical failure schedule and its 97 component bounds are **not inherited**.

The main findings are concrete. The distribution ladders themselves stay fixed, but the retained rows change after the physical grid is changed. A single inward outer index layer removes the problematic extra row. The old inner-square source fails at the smallest radius; a common replacement works throughout the screen. The exceptional-square constant and hybrid coefficients must be recomputed when the outer radius exceeds 0.275.

1. Parameters and what really remains invariant

Retain the exact source constants

$$\rho = 0.262499, \quad \rho_* = 0.2624989, \quad e = 10^{-7}/\rho, \quad h = S/98304,$$

and write $r = \rho_* S$. The proposed search fixes

$$S + T_0 = \Sigma_0 = 1.997, \quad \rho_*(S + T_1) = 0.5252997.$$

Thus $T_v = \Sigma_v - S$, where $\Sigma_1 = 0.5252997/\rho_*$. In particular $T_1 > T_0$ throughout the screen, and their difference is constant.

The official ladder recurrence depends on S, T_v through $E_v = \rho(S + T_v) - 1/2$. Consequently **every** $\omega_{v,t}, \delta_{v,t}, B_{v,t}, B_{v,t}^+, \xi_{v,t}$ is unchanged. The exact script regenerates all 29 old and 43 new rows and verifies this invariance, rather than copying a rounded table.

The root-specific quantities change:

$$a_{v,t} = B_{v,t} - T_v, \quad b_{v,t} = B_{v,t} - S.$$

On every retained nonterminal row the recurrence gives

$$(A, C_v) = \begin{cases} (S + e, T_v + e), & \text{orders 1 and 2,} \\ (S + e/2, T_v + e/2), & \text{order 3.} \end{cases}$$

The exact script verifies the prime distribution inequalities of Corollary 2.19 and all minorant inequalities of Proposition 2.22. Their strict margins remain positive. Lower-order transfer requires $S - T_\nu + e \geq 0$ and $\min(2S - T_\nu, 2T_\nu - S) + e \geq 0$, which hold here. In order three, $A + C = B + \xi$ and $\eta_D = \eta_E = \xi - e/2 > 0$. These are the actual combined-divisor hypotheses; a scalar modulus level alone is insufficient.

2. Valid plateau intervals and all largest-fragment constraints

For each ladder choose L_ν and use

$$\phi_D(t) = \min(3t/2, L_\nu), \quad \phi_E(t) = 3t - \phi_D(t).$$

Both functions are nonnegative and nondecreasing, and their sum is $3t$. In the natural plateau branch used by this screen, set

$$u_\nu = A - L_\nu, \quad v_\nu = (C_\nu + L_\nu)/4.$$

Here u_ν is the outer largest-fragment cap and v_ν the inner cap. The complete checks at a largest activated fragment are

$$u + \phi_D(u) \leq A, \quad \phi_E(u) \leq C, \quad v + \phi_E(v) \leq C, \quad \phi_D(v) \leq A.$$

They include both opposite-root conditions. Requiring both endpoints to lie on their plateau branches gives, in our regime $A > C > 0$,

$$\boxed{\frac{3A - C}{4} \leq L \leq \frac{3C}{5}}.$$

The lower bound is the outer opposite-root constraint $3(A - L) - L \leq C$. The upper bound makes $v \geq 2L/3$; it also implies the required outer-branch bound. The interval exists precisely when $A/C \leq 17/15$ within this branch.

The balanced nonlargest-witness reduction in numerical Lemma 1.4 additionally needs

$$\max(A, C) \leq 7L/3.$$

This was checked explicitly. It already follows from the lower endpoint here: $(3A - C)/4 \geq 3A/7$ since $A > C$. Below $2L/3$, both allocations equal $3t/2$; above it a nonlargest witness has inclusive tail at least $2t$, so both nonlinear obstructions and the balanced one exceed $7L/3$. This reproduces the actual reduction to nonlargest $H_{5/2}$ failures.

The intervals for $q_\nu = L_\nu/C_\nu$ are:

Physical outer radius	Old $q_{\min}$	New $q_{\min}$	Upper endpoint
0.272	0.5588487836871093	0.5553700704402632	0.6
0.2742997	0.5731934484723361	0.5696206036612396	0.6
0.275	0.5776142414184422	0.5740121594208600	0.6
0.276	0.5839701980513703	0.5803258928258113	0.6
0.278	0.5968370022033588	0.5931065437793680	0.6

The exact fractions are in the JSON; decimal table entries are displays, not substitutes for endpoint tests. The original common fraction $q = 23/40$ fails this natural cap template at 0.275, 0.276, and 0.278. In particular, the old opposite-root inequality still fails at the shared minimum outer cap for that unchanged choice.

The old-ladder upper radius for this branch is

$$r \leq \frac{17\rho_*\Sigma_0 + \rho_*e}{32} = 0.2784867267531238 \dots;$$

the new-ladder bound is 0.2790654687499988 These bound this cap parameterization, not every possible sieve support. Outside this interval one could choose smaller piecewise caps, a different allocation, or pay for additional largest-fragment failures; none is certified by the present template.

For completeness, the unrestricted largest-fragment inverse constraints can be computed piecewise. The outer owner inverse is $2A/5$ if $A \leq 5L/3$, otherwise $A - L$; its opposite inverse is $2C/3$ if $C \leq L$, otherwise $(C + L)/3$. Take their minimum. The inner owner inverse is $2C/5$ if $C \leq 5L/3$, otherwise $(C + L)/4$; intersect it with $t \leq 2A/3$ when $L > A$. This explains precisely what ceases to hold when the simple endpoint formulas are used outside their range.

3. Shared supports, nested faces, and inward caps

The usable outer domain retains **both** source ladders. For a root of total $\mathcal{S}$, each row has its original safe-core alternative $\mathcal{S} \leq a_{v,t}$ and, outside it, its full owner-tail and opposite-root predicates. The inner base requires both old and new inner predicates:

$$O = H_O \cap O_{v,t}, \quad L_1 = H_1 \cap L_{\text{new}}, \quad L_0 = H_0 \cap L_{\text{old}} \cap L_{\text{new}}.$$

The exact script constructs each cap envelope as the running minimum of every applicable row cap, starting with the global cap $\zeta = 0.19037/\rho$. The radius where a row becomes active is its own $a_{v,t}$ or $b_{v,t}$. It does not assume that the original three-shell comparator continues to hold after independent plateau changes.

For independent L_0, L_1 , the final inner caps must satisfy $C_0 + L_0 \leq C_1 + L_1$ to preserve their simple nesting. Otherwise clip the base cap to the enlarged cap on each overlapping radial interval. The implemented base envelope takes the running minimum of the old and new inner cap constraints, so $H_0 \subseteq H_1$ holds cell by cell. This largest-fragment clipping does not replace the full inner tail predicates in L_0 .

The two common-height choices used by the bounded screen,

$$L_0 = L_1 = (3A - C_0)/4 \quad \text{or} \quad L_0 = L_1 = 3C_0/5,$$

satisfy both ladders' intervals and nested caps automatically. At the common minimum, the physical caps ρ_*u , ρ_*v_0 , and ρ_*v_1 are constant as r varies under the fixed-sum constraints. The source activation radii still move. Thus this is a redistribution of radial room, not a simultaneous increase of all prime-factor allowances.

For a d -coordinate radial shell $(l, u]$, retain index sums

$$\max(0, \lfloor l/h \rfloor - d + 1) \leq j \leq \min(n - 1, \lfloor u/h \rfloor - d), \quad n = 98304 - k,$$

and round each fragment cap down to $h \lfloor z/h \rfloor$. Using the full upper cell endpoint is what safely handles a cell crossing an activation core. The output records all resulting cells and verifies nesting directly. The global physical cap remains $\rho_*\zeta < 0.19037 < \xi_*$, including inside radial cores, so the roughness restriction used in the exact-face identity is preserved.

4. The extra row 39 and the one-layer repair

With the untrimmed outer mask, its largest total is $98303h$; the enlarged inner largest total is J_1h , where $J_1 = \lfloor T_1/h \rfloor$. Retain a row precisely when $B_{v,t}$ is strictly below the corresponding actual combined upper bound. The exact results are:

Radius	Untrimmed new retained rows	Outer layers to trim
0.272	0 through 39	1
0.2742997	0 through 38	0
0.275	0 through 39	1
0.276	0 through 39	1
0.278	0 through 39	1

The old retained rows remain 0 through 27. The new extra row has

$$\xi_{n,39} = 1.886625042412619 \dots \cdot 10^{-5},$$

only 1.75 – 1.79 cells at these meshes. It violates the old numerical prerequisite $\xi > 2h$ and the implementation's first-bin assertion. The analytic source remains valid because $\xi > 0$; the breakdown concerns reuse of that numerical cover.

A rigorous inexpensive repair is

$$J_O = \min\{98303, \lfloor B_{n,39}/h \rfloor - J_1\}, \quad \sum_i j_i \leq J_O - k.$$

Then $(J_O + J_1)h \leq B_{n,39}$, so row 39 is unnecessary. At $k = 39$, the trimmed outer index maximum is 98263; the untrimmed maximum is 98264. Keep h , n , the normalizer Z , and nominal S fixed; trim the outer radial mask and derive every erased face from that changed F .

This repair is uniform for the whole radius interval $[0.272, 0.278]$, not merely the five sampled points. Writing $\Delta = \Sigma_1 - B_{n,39}$, the exact script verifies

$$h_{\max} < \Delta < 2h_{\min}, \quad B_{n,39} - B_{n,38} > h_{\max}, \quad \xi_{n,38} > 2h_{\max}.$$

Thus at most one layer is removed, row 38 remains active, and its two-cell guard holds. Separate exact inequalities keep old row 27 active and old row 28 excluded. Merely increasing grid resolution is not an automatic repair: the larger actual endpoint can activate still later rows.

5. Common inner-square source and the exceptional constant

The original inner-square level 0.5062 is insufficient at $r = 0.272$, where $2\rho_* T_1 = 0.5065994$. Use instead

$$\omega_s = 0.0035, \quad \delta_s = 0.025, \quad 1/2 + 2\omega_s = 0.507.$$

The prime order-two criterion has margin $3 - 280\omega_s - 80\delta_s = 0.02$; the fixed- σ_0 bilinear margin is 0.003996. Every minorant inequality is also checked exactly. Put

$$B_{\text{BV}} = 1/(2\rho), \quad c_s = B_{\text{BV}} - T_1, \quad \xi_s = \delta_s/\rho.$$

The retained new row 12 implies the required inner-square predicate because $c_s \geq b_{n,12}$ and $\xi_s \geq \xi_{n,12}$. The minimum core containment margin on this interval is greater than 0.02539; the activation margin is greater than 0.04470. The common source level exceeds every actual inner-square radius by at least 0.0004006. The retreat $\rho_* < \rho$ retains room for the presieving modulus.

The exceptional-square proof also depends on the maximum radius of **all** coefficient roots, including exact faces. Here it is safe to use $r_c = r$. Repeating the original 1024-bin rational upper sum gives:

r_c	Safe upward six-decimal K_{ex}
0.272	0.301405
0.2742997	0.327323
0.275	0.336134
0.276	0.349580
0.278	0.380026

The script reproduces the paper's exact $r_c = 0.275$ endpoint before evaluating the new radii. The separate independent proof is in `../exceptional-radius/EXCEPTIONAL_RADIUS_EXTENSION.md`. It confirms positive auxiliary cutoffs below the global cap and the unchanged counting margin 1/5000. It also proves that this finite-bin mechanism already exceeds 0.34 at radius 0.276; that is not a universal lower bound on the best possible exceptional-square constant.

With the unchanged minorant mass and $\lambda = 1/125$, recompute

$$a_h = m^2 - m\lambda, \quad b_h = (1 - m/\lambda)(1 - m)K_{\text{ex}}, \quad d_0 = 1 - a_h - b_h.$$

Every screened radius retains $-10^{-3} < b_h < 0$, $0 < a_h + b_h < 1$, and $d_0 > 0$ with the displayed safe constants. The script records these updated coefficients. Keeping the old $K = 0.34$ beyond the old radius hypothesis is not justified by that proposition.

Finally, $C_{\text{op}} = 4$ remains valid at $k = 39$: an exact finite exponential sum proves $\log 39 < 3.664$, and

$$\frac{39S \log 39}{38} < \frac{39S \cdot 3.664}{38} \leq 3.982482 < 4.$$

This rechecks the operator bound rather than importing the original $k = 40$ estimate at changed S .

6. What remains before a sieve certificate

The exact script examined 15 radius/plateau cases in about two seconds. Twelve satisfy the natural cap template; the unchanged $q = 23/40$ choices at 0.275, 0.276, and 0.278 are explicitly rejected. For the accepted cases it verifies ladder invariance, strict analytic source inequalities, natural plateau conditions, opposite-root caps, cap nesting, actual combined-modulus bands, the one-layer repair, common inner-square coverage, and the updated hybrid signs.

`geometry_feasibility.json` stores exact fractions and regenerated cell masks; `geometry_feasibility.log` records the run.

No original integral bound survives merely because these hypotheses pass. The next obligation is to regenerate the nonnegative failure cover from the new thresholds, core boundaries, rounded caps, and originating rows, including low-witness clipping and same-coordinate contributions. Then evaluate the changed cap and restoration forms with outward

arithmetic. The original 97-component numerical values and frozen Young parameters are not a ready-made certificate for a new point. Keeping positive old Young parameters is algebraically possible, but their new costs require fresh integrals and their previous optimality is not inherited.

No further geometry scan is part of this audit. The adjacent numerical agent's candidate screening remains a cap-only experiment until those restored forms are available. A negative candidate quotient does not prove a global sieve obstruction, and a positive cap quotient alone would not prove a smaller prime gap.

Run `python3 geometry_feasibility.py` to reproduce the certificate. The sources are the [official main proof](#), especially Propositions 3.11, 4.2, 4.6 and Lemma 4.3, and the [numerical companion](#), § § 1.1 – 1.5. The preserving official certificate file has SHA256 `7f71bdefcfe3bb5ca76a143929b3cb3f4156c21dc483253cda3077420f1e5de4`. No official source was modified.

Current report 10: REPORT

Source: research/prime-gaps/round5/geometry-trial/REPORT.md. SHA256:

`a0d2aa154e278797bd1065d7a5c78fc8d0c650e14c6b43b4d0200f84465f9462`.

Round 5: bounded radius and plateau search at $k = 39$

Status: completed exploratory cap-only search. Ten configurations were screened at 16,384 intervals; only two were refined at the published 98,304 intervals. Every configuration reassembled and optimized all 77 polynomial coefficients. No improved cap candidate was found from the radius/plateau changes. No support-restored certificate, DHL(39,2), smaller prime gap, or global variational upper bound is asserted.

The best original-geometry Round 4 value remains `0.9943963993644909`. At the fine grid, the two nearby Round 5 geometry choices gave `0.9943734016224463` and `0.9943501891039260`. Their deficits from one remain about 5,627 and 5,650 parts per million. This is substantially larger than the approximately 1.5 ppm rigorously recovered alpha credit in Round 4.

1. Exact scope and objects

The starting point is the official [PrimeGaps186 repository](#), pinned in our preserved clone to `61340d0b74163003b32756bb16e91d9209a5e330`. The independent engine reads only the two literal coefficient tables from `prime_gap_186_certificate.py`; it neither imports nor changes the official runtime. The source SHA256 is `7f71bdefcfe3bb5ca76a143929b3cb3f4156c21dc483253cda3077420f1e5de4`.

We vary the physical outer radius r while imposing the exact constraints

```
rho = 0.262499, rho_star = 0.2624989,
S = r / rho_star,
T1 = (0.5252997 - r) / rho_star,
T0 = 1.997 - S.
```

Thus both `rho_star*(S+T1)=0.5252997` and `S+T0=1.997` are fixed. The radial sums governing the two distribution ladders stay fixed, so the nominal omega/delta ladder values remain unchanged. The actual mesh, all radial masks, cap indices, Gram matrices, erased-coordinate integrals and optimized vectors are recomputed for each configuration. In particular, an invariant nominal ladder does not imply an invariant retained-row list or support-repair schedule.

The trial retains the published product weight g and the 77-dimensional polynomial span: eleven symmetric signatures, each multiplied by radial powers zero through six about 0.9. It changes neither $k=39$ nor the analytic structure of that finite span. The success criterion for this bounded search was a material increase of the fine-grid cap quotient toward one, followed by a separately valid support-repair calculation. The first criterion was not met.

2. Plateau geometry and a reduced parameter choice

Write `epsilon=10^-7/rho`, `A=S+epsilon/2` and `Cnu=Tnu+epsilon/2`. For a source with plateau height L , the core decomposition is

```
phi_D(u) = min(3u/2, L),
phi_E(u) = 3u - phi_D(u).
```

In the tested regime $A > C_1 > C_0$, the active largest-tail conditions allow

```
(3A-Cnu)/4 <= Lnu <= 3Cnu/5.
```

The additional $5/2$ tail-reduction guard $Lnu \geq 3A/7$ also holds. The corresponding final-shell caps are

```
outer cap = min(A-L0, A-L1),
old inner cap = (C0+L0)/4,
new inner cap = (C1+L1)/4.
```

The old/new inner regions must be nested. The original equal fraction $Lnu = .575 * Cnu$ is one option. Two other useful choices share the *height* of the plateau rather than its fraction:

```
common_min: L0=L1=(3A-C0)/4,
common_max: L0=L1=.6*C0.
```

Both preserve the required interval for each source and give nested inner caps automatically. Equal heights avoid spending extra outer-cap width merely because the new source has larger C_1 . This exact geometric simplification is valid even though the numerical search did not benefit from the tested endpoints.

Every rational radius, L , L/C , frontier slack, cap, and integer shell range at both meshes is saved in `geometry_checks.json`. `validate_geometry.py` verifies the fixed sums, both plateau frontiers and guards, and actual old/new/full mask nesting on all twenty configuration/grid pairs. These checks are necessary structural checks; the separate source-geometry audit supplies the remaining distribution and repair constraints.

3. Exceptional square constant must change with physical radius

The original exceptional-square statement uses radius at most $11/40$ and the convenient bound $K=17/50$. It is incorrect to reuse that constant at $r=.276$ or $.278$ without a new calculation.

The independent exact rational computation in `../exceptional-radius/certify_exceptional_radius.py` repeats the original 1,024-bin, 21-term alternating-logarithm upper certificate. It reproduces the published baseline rational sum exactly and gives safe rounded-up constants:

r	K used here	certified upper bound, approximately
.272	.301405	.301404153485181623
.2742997	.327323	.327322538111366366
.275	.34	.336133604027290568
.276	.349580	.349579996894903756
.278	.380026	.380025921565620058

The **.275** trial intentionally retains the still-valid conservative **.34**; the exact calculation would also allow **.336134**. The original-geometry baseline likewise retains **.34** to check agreement with Round 4.

Both matrix and direct evaluation use the changed K in the signed hybrid:

```
m=.99998, lambda=.008,
a=m*m-m*lambda,
b=(1-m/lambda)*(1-m)*K,
Jcap=Jbase+(a+b)*Jplus+b*Jtail.
```

Thus higher-radius trials do not inherit an unjustified exceptional constant. Reducing K at the original radius improves the coarse optimized cap value by only about 0.6183 ppm. This calculation concerns the cap form; it is not a new fully certified prime-gap margin.

4. Bounded numerical results

All entries below are the direct scalar reevaluation of the optimized 77-vector. The independent matrix value, vector, conditioning, residuals, and three whitening-cutoff results are retained in each JSON file. Corresponding NPZ files hold the full Gram matrix, numerator matrix and original coefficient vector.

tag	physical r	plateau	K	grid	rho_star Jcap/I
baseline	.2742997	original	.34	16384	0.993379352892581
r272_min	.272	common_min	.301405	16384	0.992837127339583
r272_original	.272	original	.301405	16384	0.992825152057240
r274_max	.2742997	common_max	.327323	16384	0.993356576365699
r274_original	.2742997	original	.327323	16384	0.993379971225379
r275_max	.275	common_max	.34	16384	0.993324871060043
r276_max	.276	common_max	.349580	16384	0.993147338997495
r276_min	.276	common_min	.349580	16384	0.993160475207861
r278_max	.278	common_max	.380026	16384	0.992272698494233
r278_min	.278	common_min	.380026	16384	0.992274861882488
r274_max	.2742997	common_max	.327323	98304	0.994373401622446
r275_max	.275	common_max	.34	98304	0.994350189103926

The original-geometry coarse matrix value agrees bit for bit with the unchanged Round 4 program. Against that coarse control, the other tested radii all lose: approximately 542 ppm at .272, 54 ppm at .275, 219 ppm at .276, and 1,104 ppm at .278, using the better of the tested plateau choices at each radius.

The two most competitive nearby alternatives were refined. At the fine grid, original radius/common_max loses 22.9977 ppm relative to the Round 4 optimized original geometry; .275/common_max loses 46.2103 ppm. The ordering agrees with the coarse evidence. The search stopped at these ten configurations and two refinements. There is no extrapolation

from these values to every plateau height, every radius, every product weight, or every finite family.

5. Numerical audit and its limits

The engine uses a positive exponential tilt of 20 in convolution weights and restores the exact balancing factors. This avoids evaluating the tiny un-tilted convolution left tail by cancellation. It is the same independent cap engine developed and checked in Round 4. `numpy.longdouble` on this host is actually 64 bits; no extra precision or outward enclosure is claimed.

Across all twelve computations:

- all 77 dimensions were retained at the smallest tested scaled-Gram cutoff;
- the maximum full scaled-Gram condition number was about $4.307e10$;
- matrix versus direct candidate quotient disagreement was at most $2.884e-10$;
- full scaled-pencil relative residual was at most $5.30e-16$;
- the summed recorded compute time was 175.3 seconds on this host, excluding orchestration.

These checks give meaningful numerical evidence at the tens-of-ppm scale at issue, but they are not interval-arithmetic error bounds. The generalized eigenvalue is an observed numerical optimum of the assembled finite matrix. We do not label it a certified upper bound or a no-go theorem. Direct reevaluation uses a different contraction order but shares the cap model and one-dimensional moment primitives; it does not independently prove that model.

6. Support and source obligations discovered during this search

The following are substantive constraints, not bookkeeping details.

1. **The inner-square auxiliary source must be updated at smaller r .** At $r=.272$, the published $\omega=.0031$ gives a level smaller than twice the new physical inner radius. The separate exact source audit supplies a common replacement `$\omega_s=.0035$` , `$\delta_s=.025$` , which passes the tested source and row-12 containment inequalities. Those source facts do not by themselves evaluate the cap's support losses.
2. **A new retained row can violate the original low-witness mesh guard.** At the fine grid, $r=.272,.275,.276,.278$ can retain new-ladder row 39. Its activation coordinate is about $1.886625e-5$, which is less than two mesh cells for these cases. The original low-witness implementation requires at least two cells. The source theorem does not fail, but the original support-repair engine cannot be inherited unchanged.
3. **An explicit inward restriction can remove that extra row.** With `$h=S/98304$` , `$J1=\text{floor}(T1/h)$` , define

```
Jo=min(98303, floor(B_new39/h)-J1).
```

Restrict the outer sum of coordinate indices to `$r_{\text{total}} \leq J_0 - k$` . In the affected tested cases $J_0=98302$, one layer less than the ordinary $J_0=98303$; this restores the retained range 0 through 38. Keep h , nominal S , convolution length and Z fixed, but recompute every face from the trimmed F . The current reported trials are **not trimmed**. No numerical loss estimate for that modification is included here.

4. **Every schedule still needs rebuilding.** Even after trimming, actual integer thresholds, caps, source-specific masks, failure forms and outward quadrature bounds must be regenerated. The approximately 1.5 ppm alpha rectangle proved for the fixed published k40 profile cannot simply be transferred to these optimized k39 vectors or these radii.

Since none of the screened geometry changes gave a promising cap value, a full repaired certificate for these particular vectors is postponed. This is a prioritization decision based on the observed finite trials, not a proof that repair or a different profile cannot succeed.

7. Reproducibility and archive contents

Run from this directory:

```
OPENBLAS_NUM_THREADS=1 python3 validate_geometry.py
OPENBLAS_NUM_THREADS=1 python3 run_bounded_screen.py
OPENBLAS_NUM_THREADS=1 python3 run_bounded_screen.py --refine
```

Dependencies are recorded in `manifest.json`: Python, NumPy and SciPy, platform, actual long-double bit count, and SHA256 of all computation inputs and outputs. The parser's only upstream file dependency is

```
../../research-round1/prime186-work/PrimeGaps186/prime_gap_186_certificate.py
```

relative to this directory. That file provides the literal published coefficient signatures and integer coefficient matrix. An exported archive must preserve this relative layout or set `SOURCE` to the pinned copy. The two distribution-ladder recurrences and geometry are independently implemented in `cap_trial.py`; no FLINT module is imported and no official regression guard is disabled. The preserved upstream clone remained clean after all computations.

Files are intentionally small-purpose: `cap_trial.py` evaluates a vector; `optimize_cap.py` builds and solves the finite pencil; `validate_geometry.py` checks exact parameters and actual masks; `run_bounded_screen.py` replays the declared finite list. Ten `*.config.json` files contain every parameter choice. Each computed point has JSON evidence and a compressed NPZ matrix archive. `summary.json` is the compact comparison surface.

8. Conclusion for the next research step

The tested radius/plateau degrees of freedom do not erase the k39 cap deficit. The useful output is a fully recomputed negative trial, a simpler common-height plateau parameterization, a valid radius-dependent exceptional constant, and an explicit identification of the new mesh obstruction. A more serious search should change a mathematically justified larger component, such as the product profile or support geometry, or derive a certified finite-family upper bound before expending effort on full support restoration. Arbitrary radius scans and inherited k40 constants are not justified by these results.